Mathematics

Arthur Li

June 13, 2026

Contents

Notation List 1

I Mathematical Language, Sets, and Logic 10

1 Sets 10
1.1 Introduction to Sets . . . 10
1.2 Cartesian Products . . . 10
1.3 Subsets and Power Sets . . . 10
1.4 Union, Intersection, Difference, and Complement . . . 10
1.5 Indexed Sets and Number Systems . . . 11
1.6 Russell's Paradox and Axioms . . . 11

2 Logic and Proofs 11
2.1 Statements and Well-Formed Language . . . 12
2.2 Connectives, Conditionals, and Truth Tables . . . 12
2.3 Quantifiers, Bound Variables, and Negation . . . 12
2.4 Proof Grammar and Inference . . . 12
2.5 Direct, Contrapositive, and Contradiction Proofs . . . 12
2.6 Induction, Recursion, and Counterexamples . . . 12

3 Counting and Finite Sums 13
3.1 Lists and Counting Principles . . . 13
3.2 Permutations, Subsets, and Binomial Coefficients . . . 13
3.3 Inclusion-Exclusion, Pigeonhole, and Combinatorial Proof . . . 13
3.4 Finite Sums . . . 14

4 Relations, Orders, and Lattices 15
4.1 Relations and Equivalence . . . 15
4.2 Partitions, Congruence, and Quotients . . . 15
4.3 Ordered Sets, Well-Orders, and Lattices . . . 15
4.4 Relations Between Sets and Closures . . . 16

5 Functions 16
5.1 Functions, Images, and Preimages . . . 16
5.2 Injective, Surjective, and Bijective Maps . . . 17
5.3 Composition, Restriction, and Inverse Functions . . . 17

6 Cardinality, Choice, and Infinite Sets 18
6.1 Same Cardinality, Finite, Countable, and Uncountable . . . 18
6.2 Diagonalization and Power Sets . . . 20
6.3 Schroeder–Bernstein and Cardinal Comparison . . . 20
6.4 Cardinal Arithmetic . . . 21
6.5 Choice, Zorn's Lemma, and Well-Ordering . . . 21

7 Axiomatic and Advanced Set Theory 22
7.1 Extensionality, Specification, and Replacement . . . 22
7.2 Ordinals, Transfinite Recursion, and Well-Foundedness . . . 22
7.3 Models, Constructibility, and Independence . . . 22
7.4 Filters, Ultrafilters, and Sets of Reals . . . 22

II Linear and Operator Structures 23

1 Structural Linear Algebra 23
1.1 Vector Spaces and Linear Structure . . . 23
1.2 Direct Sums, Quotients, and Decomposition . . . 29
1.3 Linear Maps and Representations . . . 35
1.4 Block Matrices and Matrix Operations . . . 52
1.5 Duality and Multilinear Structure . . . 54

1.6 Finite-Dimensional Topological Vector Spaces . . . 61

2 Geometric Linear Algebra 62
2.1 Inner Product Geometry . . . 62
2.2 Orthogonal Complements and Projections . . . 73
2.3 Least Squares and Metric Optimization . . . 77
2.4 Adjoints and Orthogonal Operators . . . 79
2.5 Metric Vector Spaces and Bilinear Forms . . . 86
2.6 Affine and Projective Geometry . . . 99
2.7 Positive Systems, Convexity, and Separation . . . 104

3 Spectral and Operator Structure 108
3.1 Eigenstructure and Invariant Subspaces . . . 108
3.2 Polynomial Structure of Operators . . . 117
3.3 Jordan and Canonical Forms . . . 128
3.4 Spectral Theory of Structured Operators . . . 134
3.5 Operator Functions . . . 143
3.6 The Umbral Calculus . . . 146

4 Matrix Analysis and Numerical Structure 153
4.1 Matrix Classes and Invariants . . . 154
4.2 Matrix Decompositions . . . 172
4.3 Matrix Norms and Conditioning . . . 179
4.4 Perturbation and Stability . . . 185
4.5 Numerical Linear Algebra . . . 190

III Algebraic Structures 195

1 Algebraic Language and Categories 195
1.1 Sets, Maps, Relations, and Universal Properties . . . 195
1.2 Categories and Functors . . . 195
1.3 Algebraic Structures as Objects . . . 196
1.4 Exactness and Diagrams . . . 196

2 Group Theory 197
2.1 Groups, Subgroups, and Homomorphisms . . . 197
2.2 Group Actions . . . 198
2.3 Normality and Structure Theorems . . . 198
2.4 Sylow Theory and Finite Groups . . . 198
2.5 Representation-Theoretic Preview . . . 199

3 Rings, Ideals, and Modules 199
3.1 Rings and Ideals . . . 199
3.2 Polynomial Rings and Factorization . . . 202
3.3 Modules . . . 205
3.4 Structure Theorems . . . 212
3.5 Linear Operators as Modules . . . 220
3.6 Associative Algebras . . . 227

4 Fields and Galois Theory 231
4.1 Field Extensions . . . 231
4.2 Splitting Fields and Normality . . . 231
4.3 Galois Groups . . . 231
4.4 Classical Applications . . . 231
4.5 Scalar Extension in Linear Algebra . . . 232

5 Tensor, Exterior, and Homological Algebra **236**
5.1 Tensor Products and Universal Properties . . . 236
5.2 Symmetric and Exterior Algebras . . . 239
5.3 Multilinear Algebra Revisited . . . 243
5.4 Chain Complexes and Homology . . . 243
5.5 Derived Functors Preview . . . 243

IV Algebraic Geometry and Commutative Algebra **244**

1 Commutative Algebra **244**
1.1 Prime Ideals and Spectra . . . 244
1.2 Localization and Local Rings . . . 244
1.3 Noetherian Rings and Modules . . . 244
1.4 Integral Dependence and Dimension . . . 244
1.5 Regularity and Singularities Preview . . . 244

2 Algebraic Geometry **244**
2.1 Affine Algebraic Sets . . . 244
2.2 Morphisms and Coordinate Rings . . . 244
2.3 Projective Geometry . . . 244
2.4 Schemes . . . 244
2.5 Sheaves and Cohomology Preview . . . 244
2.6 Smoothness, Tangent Spaces, and Dimension . . . 244
2.7 Divisors, Line Bundles, and Curves . . . 244

V Analysis and Continuous Structure **245**

1 Foundations of Analysis **245**
1.1 Number Systems and Completeness . . . 245
1.2 Cardinality and Infinite Sets . . . 251
1.3 Real Sequences and Completeness Methods . . . 251

2 Metric and Topological Structure **255**
2.1 Metric Spaces and Euclidean Topology . . . 255
2.2 Compactness and Completeness . . . 263
2.3 Connectedness and Continuity . . . 274
2.4 Selected General Topology . . . 283

3 Convergence and Function Spaces **284**
3.1 Infinite Series and Summability . . . 284
3.2 Uniform Convergence . . . 291
3.3 Function-Space Compactness . . . 294
3.4 Power Series and Approximation . . . 297
3.5 Fourier Structure . . . 303

4 Differential and Geometric Analysis **305**
4.1 One-Variable Differentiation and Integration . . . 305
4.2 Several-Variable Differentiation . . . 318
4.3 Inverse, Implicit, and Rank Theorems . . . 322
4.4 Differential Forms and Manifolds . . . 324
4.5 Stokes Theory . . . 329

5 Measure and Integration Structure **331**
5.1 Sigma-Algebras and Measures . . . 331
5.2 Regular Approximation, Additivity, and Nonmeasurable Sets . . . 345
5.3 Measurable Functions . . . 351
5.4 Lebesgue Integration . . . 354
5.5 Convergence Theorems . . . 359
5.6 Product Measures and Fubini Theory . . . 363
5.7 Radon-Nikodym and Differentiation . . . 368

5.8 Lp Spaces . . . 375
5.9 Modes of Convergence . . . 380

6 Functional and Operator Analysis 387
6.1 Normed and Banach Spaces . . . 387
6.2 Hilbert Spaces . . . 388
6.3 Core Functional Analytic Theorems . . . 393
6.4 Weak Topologies and Duality . . . 394
6.5 Compact Operators . . . 394
6.6 Spectral Theory of Operators . . . 395
6.7 Distributions and Sobolev Preview . . . 395

VI Probability and Stochastic Structure 396

1 Probability Spaces and Convergence 396
1.1 Probability Spaces and Random Variables . . . 396
1.2 Expectation and Integration in Probability . . . 402
1.3 Conditional Probability and Conditional Expectation . . . 407
1.4 Modes of Probabilistic Convergence . . . 411
1.5 Laws of Large Numbers . . . 413
1.6 Central Limit Theory . . . 414
1.7 Probability on Metric Spaces . . . 415

2 Stochastic Processes 415
2.1 Process-Level Probability . . . 415
2.2 Filtrations and Adapted Processes . . . 416
2.3 Martingales in Discrete and Continuous Time . . . 416
2.4 Markov Chains and Markov Processes . . . 416
2.5 Stationary Processes and Ergodicity . . . 419
2.6 Gaussian Processes . . . 419
2.7 Point Processes and Renewal Structure . . . 419

3 Stochastic Calculus and Dynamics 421
3.1 Brownian Motion . . . 421
3.2 Stochastic Integration . . . 421
3.3 Itô Calculus . . . 421
3.4 Stochastic Differential Equations . . . 421
3.5 Change of Measure and Girsanov Theory . . . 421
3.6 Diffusions and PDE Connections . . . 421
3.7 Stochastic Flows and Advanced Dynamics . . . 421

VII Dynamical and PDE Systems 422

1 Dynamical Systems and ODEs 422
1.1 ODE Foundations and Flows . . . 422
1.2 Linear Systems and Stability . . . 422
1.3 Nonlinear Stability Theory . . . 422
1.4 Bifurcation Theory . . . 422
1.5 Dynamical Systems Structure . . . 422
1.6 Semigroups and Evolution Viewpoint . . . 422

2 PDE and Evolution Equations 422
2.1 First-Order PDE and Characteristics . . . 422
2.2 Elliptic Equations . . . 422
2.3 Parabolic Equations . . . 422
2.4 Hyperbolic Equations . . . 422
2.5 Sobolev Spaces for PDE . . . 422
2.6 Weak Solutions and Variational Methods . . . 423
2.7 Evolution Equations . . . 423

VIII Numerical and Computational Mathematics 424

1 Numerical Analysis 424
1.1 Floating Point and Error Analysis . . . 424
1.2 Nonlinear Equations and Optimization Basics . . . 424
1.3 Interpolation and Approximation . . . 424
1.4 Numerical Differentiation and Quadrature . . . 424
1.5 Numerical Stability and Convergence . . . 424

2 Numerical Differential Systems 424
2.1 Numerical ODEs . . . 424
2.2 Geometric Numerical Integration . . . 424
2.3 Finite Difference PDE Methods . . . 424
2.4 Finite Element Methods . . . 424
2.5 Numerical SDEs . . . 424
2.6 Computational Stability Across ODE/PDE/SDE . . . 424

3 Monte Carlo and Stochastic Computation 424
3.1 Monte Carlo Foundations . . . 424
3.2 Variance Reduction . . . 425
3.3 Markov Chain Monte Carlo . . . 426
3.4 Sequential Monte Carlo . . . 427
3.5 Rare-Event Simulation . . . 427
3.6 Stochastic Optimization . . . 427

IX Statistical and Learning Systems 428

1 Statistical Learning 428
1.1 Statistical Inference Foundations . . . 428
1.1.1 Models, Estimators, and Confidence . . . 429
1.1.2 Empirical CDFs and Functionals . . . 430
1.1.3 Bootstrap . . . 431
1.1.4 Parametric Inference . . . 432
1.1.5 Hypothesis Testing and p-values . . . 433
1.1.6 Bayesian Inference . . . 435
1.1.7 Statistical Decision Theory . . . 436
1.2 Risk Minimization . . . 437
1.3 Linear and Generalized Linear Models . . . 439
1.4 Linear Regression . . . 441
1.5 Regularization and Model Selection . . . 443
1.6 Regression Model Selection . . . 444
1.7 Kernels and Nonlinear Methods . . . 445
1.8 Trees, Ensembles, and Boosting . . . 447
1.9 Generalization Theory . . . 448
1.10 Association and Independence Inference . . . 449
1.11 Undirected Graphs and Log-Linear Models . . . 453
1.12 Counterfactual Causal Inference . . . 458
1.13 Directed Graphs and Interventions . . . 462
1.14 Nonparametric Curve Estimation . . . 466
1.14.1 Histograms . . . 467
1.14.2 Kernel Density Estimation . . . 468
1.14.3 Nonparametric Regression . . . 470
1.15 Orthogonal Series and Wavelet Smoothing . . . 472
1.15.1 Orthogonal Coordinates in L^2 . . . 472
1.15.2 Orthogonal Series Density Estimation . . . 473
1.15.3 Orthogonal Series Regression . . . 474
1.15.4 Haar Wavelets . . . 475

2 Deep Learning and Representation Systems **478**
2.1 Neural Network Foundations . . . 478
2.1.1 Deep Feedforward Networks . . . 478
2.2 Optimization for Deep Learning . . . 483
2.3 Regularization and Generalization . . . 491
2.4 Representation Learning . . . 494
2.5 Sequence Models . . . 494
2.6 Generative Models . . . 494
2.7 Modern Architecture and Systems View . . . 494
2.7.1 Convolutional Neural Networks . . . 494

3 Time Series and Sequential Systems **496**
3.1 Time Series Foundations . . . 496
3.1.1 Univariate Time Series Models . . . 496
3.2 ARMA and ARIMA Structure . . . 498
3.2.1 Stationary Time Series . . . 498
3.3 Spectral Methods . . . 498
3.4 State-Space Models . . . 498
3.5 Nonlinear and Volatility Models . . . 498
3.6 Long Memory and Multiscale Structure . . . 498
3.7 Sequential Inference and Filtering . . . 498

Notation List

Collision Index and Scope Rules

bare $A, B, \ldots$; $\mathcal{A}, \mathcal{B}, \ldots$	unadorned Latin and calligraphic letters denote local objects unless a decorated form is listed below; local declarations in a theorem, proof, or subsection take precedence
$B,\ B_r(x),\ \overline{B}_r(x),\ \mathrm{B}(x,y),\ \mathcal{B}(\mathbb{R}^n)$	matrix or set variable / open ball / closed ball / beta function / Borel σ-algebra; font, arguments, and ambient subject determine the referent
$I_{\mathcal{V}},\ I_n,\ I^k,\ I_R$	identity operator, identity matrix, k-cell, and power-series interval
$P,\ \mathcal{P}(X),\ \mathscr{P},\ P_{\mathcal{U}},\ P_j$	locally declared projection or map / power set / umbral polynomial algebra / orthogonal projection / spectral projection
$\mathcal{F},\ \mathscr{F},\ \mathcal{F}(K),\ \mathcal{F}(\mathcal{V})$	local family of functions or operators / umbral formal-series algebra / finite subsets of K / finite-rank operators
$\mathcal{L}(\mathcal{V},\mathcal{W}),\ \mathcal{L}^n,\ \mathscr{L}^1,\ \mathrm{LS}(A,\mathbf{y}),\ \mathcal{L}(\theta,\alpha)$	linear maps / Lebesgue measurable sets / integrable functions / least-squares solution set / local Lagrangian or loss notation
$R(f,\Pi,\{c_i\}),\ R,\ R_k(A),\ A=QR$	Riemann sum, locally declared radius or scalar, row-sum quantity, and QR factorization; arity and argument type disambiguate the notation
$\Delta x_i,\ \Delta A,\ \Delta(A),\ \Delta(\mathcal{V}),\ \Delta_{k-1}$	increment, matrix perturbation, defect from normality, metric discriminant, and coefficient simplex
$\Omega,\ \Omega_T,\ \Omega_n,\ \Omega(\theta),\ \mathrm{vol}(\Omega)$	local domain or set / finite spectrum of T / primitive roots of unity or DFT diagonal, as specified by context / learning penalty / volume domain
$Q^k,\ Q_n,\ \Phi_n$	standard k-simplex / normalized Fourier matrix / cyclotomic polynomial; cyclotomic polynomials use Φ_n
$\overline{z},\ \overline{E},\ \overline{\mathcal{A}},\ A^*,\ T^*\omega,\ f',\ \mathcal{V}'$	conjugation, closure, uniform closure, adjoint or conjugate transpose, pullback, derivative, and dual space; the decorated object determines the operation

Global Conventions and Basic Notation

Convention. Notation is scoped by subject, argument type, font, decoration, and any local definition in force. Unadorned letters denote local objects by default; persistent meanings are carried by subscripts, superscripts, operators, fonts, or declared argument patterns. When a glyph is reused, the distinguishing feature is recorded below.

Set, Function, and Logical Notation

$x \in E,\ x \notin E$	membership / non-membership
$E \subseteq F,\ \emptyset$	subset inclusion, not necessarily proper / empty set
$X \times Y,\ f : X \to Y$	Cartesian product / function with domain X and codomain Y
$\Longrightarrow,\ \Longleftrightarrow$	logical implication / logical equivalence
$:=$	definition by equality

Number, Size, and Scalar Notation

$\mathbb{N}, \mathbb{Z}, \mathbb{Q}, \mathbb{R}, \mathbb{C}$	natural numbers, integers, rationals, reals, and complexes
$\mathbb{F},\ \mathbb{F}_0$	ambient scalar field unless fixed as $\mathbb{R}$ or $\mathbb{C}$ / prime subfield of $\mathbb{F}$; probabilities use explicit Pr or locally declared probability spaces
$\lvert x\rvert,\ \lvert E\rvert,\ \lVert x\rVert$	absolute value or complex modulus for scalars; interval length or cardinality according to argument type; norm of a vector, function, or operator
a, b, c	locally declared scalars or interval endpoints

Linear and Operator Structures Notation

Vector Spaces and Lists

$\mathbb{F},\ \mathbb{F}_0$	ambient scalar field, typically $\mathbb{R}$ or $\mathbb{C}$ / prime subfield of $\mathbb{F}$
$V(n,q),\ \mathbb{F}_q^n$	n-dimensional coordinate space over the finite field with q elements
$\mathbb{F}^{(\mathbb{N})}$	vector space of finitely supported sequences over $\mathbb{F}$
$\mathbf{u},\mathbf{v},\mathbf{w},\mathbf{0}$	vectors / zero vector
$\mathcal{U},\mathcal{V},\mathcal{W}$	vector spaces or subspaces; other calligraphic letters such as $\mathcal{A},\mathcal{B},\mathcal{F},\mathcal{M}$ are local unless listed below; umbral algebras use script $\mathscr{P},\mathscr{F}$
$I_{\mathcal{V}},\ I_n,\ 0$	identity operator on $\mathcal{V}$ / identity matrix / zero map or zero matrix; distinct from the cell I^k and interval I_R
$\beta=(\mathbf{v}_1,\dots,\mathbf{v}_m)$	ordered vector list; a basis iff it is linearly independent and spans
$(),\ \beta_{\widehat{j}}$	empty list / list β with the jth entry omitted
$(\beta,\mathbf{v}),\ \beta_{j\leftarrow\mathbf{v}}$	list with $\mathbf{v}$ appended / list with the jth entry replaced by $\mathbf{v}$
$\mathcal{B}$	locally declared basis set, such as a Hamel, Hilbert, module, or affine basis; $\mathcal{B}(\mathbb{R}^n)$ denotes the Borel σ-algebra in analysis
$\operatorname{span}(E),\ \dim(\mathcal{V})$	linear span of E / dimension of $\mathcal{V}$
$\operatorname{codim}_{\mathcal{V}}(\mathcal{U})$	codimension of $\mathcal{U}$ in $\mathcal{V}$; equal to $\dim(\mathcal{V}/\mathcal{U})$ in finite dimension
$\bigwedge_{i\in I}\mathcal{U}_i,\ \bigvee_{i\in I}\mathcal{U}_i$	meet / join of a family of subspaces
$\sum_{i\in I}\mathcal{U}_i$	sum of an indexed family of subspaces; equivalently the join
$\mathcal{U}+\mathcal{W},\ \mathcal{U}\oplus\mathcal{W},\ \mathcal{V}/\mathcal{U}$	sum, direct sum, quotient space
$\mathcal{V}_1\times\cdots\times\mathcal{V}_m$	product vector space with componentwise operations
$\prod_{i\in I}\mathcal{V}_i,\ \bigoplus_{i\in I}^{\text{ext}}\mathcal{V}_i$	external direct product / external direct sum of finite-support families
$\mathcal{V}\cong\mathcal{W},\ \mathbf{e}_{j,k}$	isomorphic vector spaces / basis vector supported in one product component
$\mathbf{v}+\mathcal{U},\ \pi(\mathbf{v})$	coset modulo $\mathcal{U}$ / locally declared quotient map $\mathcal{V}\to\mathcal{V}/\mathcal{U}$

Linear Maps and Duality

$T:\mathcal{V}\to\mathcal{W},\ \mathcal{L}(\mathcal{V},\mathcal{W})$	linear map / vector space of linear maps; distinct from measure-theoretic $\mathcal{L}^n$, least-squares $\operatorname{LS}(A,\mathbf{y})$, and local loss notation $\mathcal{L}(\theta)$
$ST,\ T\vert_{\mathcal{U}},\ T/\mathcal{U}$	composition $S\circ T$ / restriction to $\mathcal{U}$ / induced quotient operator
$T^{-1},\ T^m,\ p(T)$	inverse of an invertible map or operator / operator power / polynomial functional calculus
$\ker(T),\ \operatorname{range}(T)$	kernel and range of T
$\mathcal{V}_{\mathbb{C}},\ \operatorname{cpx},\ T_{\mathbb{C}}$	complexification of a real vector space / canonical embedding $\mathbf{v}\mapsto\mathbf{v}+i\mathbf{0}$ / complexification of a real operator
$\operatorname{graph}(T),\ [S,T]$	graph of a map / commutator $ST-TS$
$\operatorname{Comm}(\mathcal{F}),\ \operatorname{Comm}(A)$	commutant of a family / commutant of a single matrix or operator
$A\sim_{\text{row}}B,\ S\sim T$	row-equivalent matrices / operator similarity when declared; distinct from cardinal, a.e., and coefficient equivalence
$\mathcal{V}',\ \beta',\ T'$	dual space, dual basis, and dual map; prime denotes duality in this context
$\varphi,\psi,\ \mathcal{U}^0$	linear functionals / annihilator
$\mathcal{V}'',\ \Lambda_{\mathcal{V}}$	double dual / natural evaluation map $\mathcal{V}\to\mathcal{V}''$
$[\mathbf{v}]_\beta,\ [T]_{\gamma\leftarrow\beta}$	coordinate vector in basis β / matrix of T from β-coordinates to γ-coordinates
$C_{\gamma\leftarrow\beta}$	change-of-basis matrix sending β-coordinate columns to γ-coordinate columns

Matrices and Entrywise Structure

$A, B, C,\ \mathrm{Mat}(m, n, \mathbb{F})$	matrix variables / space of $m \times n$ matrices over $\mathbb{F}$; decorated forms such as C_f and C_{ij} have separate meanings
$A_{ij},\ A_{j,\cdot},\ A_{\cdot,k}$	entry in row i and column j / jth row / kth column
$[a_{ij}],\ E_{j,k},\ A\mathbf{x}$	matrix with entries a_{ij} / matrix unit / matrix-vector product
$A^{-1},\ A^t,\ A^T,\ A^*$	inverse matrix when it exists / transpose / alternate transpose notation / conjugate transpose or adjoint
$\mathrm{diag}(\lambda_1, \ldots, \lambda_n)$	diagonal matrix with diagonal entries $\lambda_1, \ldots, \lambda_n$
$\overline{z},\ \mathrm{Re}\, z,\ \mathrm{Im}\, z$	complex conjugate / real part / imaginary part; for sets, overline denotes closure
$\mathrm{tr}(A), \det(A)$	trace and determinant
$\mathfrak{S}_n,\ \mathrm{sign}(\sigma),\ \mathrm{adj}(A)$	permutation group / permutation sign / adjugate matrix
$M_{ij},\ C_{ij},\ A_j(\mathbf{b})$	minor matrix / cofactor / matrix obtained from A by replacing column j with $\mathbf{b}$
$\mathrm{col}(A),\ \mathrm{row}(A)$	column space and row space of A
$\mathrm{rank}(A),\ \mathrm{nullity}(A)$	rank and nullity of A
$A \oplus B,\ A \otimes B$	matrix direct sum / Kronecker product; vector-space direct sums use subspace notation
$A \otimes I_q + I_n \otimes B$	Kronecker sum of square matrices A and B in compatible dimensions
$A \circ B,\ \mathrm{vec}(A)$	Hadamard product / column-stacking vectorisation
$M/A,\ M/D$	Schur complement of the block A or D in a compatible block matrix M; distinct from quotient-space notation
$B,\ L = B^t B$	directed incidence matrix / associated incidence Laplacian in a graph context

Inner Product Geometry

$\|\mathbf{v}\|,\ \|T\|_{\mathrm{op}}$	vector norm / operator norm
$\|\cdot\|_1,\ \|\cdot\|_2,\ \|\cdot\|_\infty$	coordinate one-norm / Euclidean or spectral two-norm according to argument type / sup norm
$\|\cdot\|_F$	Frobenius norm
$\langle \mathbf{u}, \mathbf{v} \rangle,\ \langle A, B \rangle_F$	inner product / Frobenius inner product
$\mathbf{u} \perp \mathbf{v},\ \mathcal{U}^\perp,\ \delta_{ij}$	orthogonal vectors, orthogonal complement, Kronecker delta
$\mathrm{proj}_{\mathbf{v}}(\mathbf{u}),\ P_{\mathcal{U}}$	orthogonal projection of $\mathbf{u}$ onto $\mathrm{span}(\mathbf{v})$ / orthogonal projection onto $\mathcal{U}$
$\mathrm{LS}(A, \mathbf{y})$	least-squares solution set $\arg\min_{\mathbf{x} \in \mathbb{F}^n} \|A\mathbf{x} - \mathbf{y}\|_2$ for $A \in \mathrm{Mat}(m, n, \mathbb{F})$; roman notation avoids collision with $\mathcal{L}$ and $\mathcal{L}^n$
$d(\mathbf{v}, \mathcal{U})$	distance $\inf_{\mathbf{u} \in \mathcal{U}} \|\mathbf{v} - \mathbf{u}\|$ from $\mathbf{v}$ to $\mathcal{U}$
$\mathrm{Gram}(\mathbf{u}_1, \ldots, \mathbf{u}_n)$	Gram matrix of $\mathbf{u}_1, \ldots, \mathbf{u}_n$
$\mathrm{GramDet}(\mathbf{u}_1, \ldots, \mathbf{u}_n)$	Gram determinant
$P_{\mathbf{u}},\ U_{\mathbf{w}}$	rank-one projection matrix / Householder matrix, with defining vector indicated by the subscript
$\gamma = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$	ordered basis; orthonormal when declared in an inner-product context
$\mathbf{e} = (1, \ldots, 1)^t$	all-ones vector
$\mathrm{Ray}_A(\mathbf{x}),\ \mathbf{x}^* A \mathbf{x}$	Rayleigh quotient notation; for unit $\mathbf{x}$, $\mathrm{Ray}_A(\mathbf{x}) = \mathbf{x}^* A \mathbf{x}$ / Hermitian quadratic form
$\mathrm{Sym}(B),\ \mathrm{Skew}(B),\ \mathrm{Alt}(B)$	predicates that the bilinear form B is symmetric, skew-symmetric, or alternating
$M_\beta,\ \Delta(\mathcal{V}),\ \mathrm{rad}(\mathcal{V})$	matrix of a metric bilinear form in basis β / discriminant set or square class / radical of a metric vector space
$\mathcal{V} = \mathcal{S} \odot \mathcal{T},\ \mathcal{H},\ w(\mathcal{V})$	orthogonal direct sum / hyperbolic space or plane in metric-vector-space contexts / Witt index; affine hyperplanes use $\mathcal{H}(\mathbf{N}, b)$
$\mathcal{O}(\mathcal{V}),\ \mathrm{Sp}(\mathcal{V}),\ \tau_{\mathbf{v},a},\ \sigma_{\mathbf{u}}$	orthogonal group / symplectic group / symplectic transvection / orthogonal symmetry

Affine, Projective, and Convex Geometry

$\mathcal{A}(\mathcal{V}),\ \mathcal{A}_0(\mathcal{V}),\ \mathrm{Sub}(\mathcal{V})$	flats in $\mathcal{V}$ / $\mathcal{A}(\mathcal{V}) \cup \{\emptyset\}$ / subspaces of $\mathcal{V}$; $\mathcal{A}$ without arguments remains local
$\mathrm{affspan}(E),\ \mathrm{affhull}(E)$	affine span / least flat containing E
$X \parallel Y,\ X \vee Y$	parallel flats / join of flats or projective subspaces; meet is intersection in the relevant lattice
$\mathbb{R}^n_+,\ \mathbb{R}^n_{++},\ \mathbf{x} \geq \mathbf{0},\ \mathbf{x} > \mathbf{0},\ \mathbf{x} \gg \mathbf{0}$	nonzero nonnegative orthant / coordinatewise positive orthant / coordinatewise nonnegative, nonzero nonnegative, and coordinatewise positive inequalities
$\mathrm{conv}(S),\ \Delta_{k-1},\ C_A$	convex hull / coefficient simplex $\{(t_i) : t_i \geq 0,\ \sum_i t_i = 1\}$ / nonnegative cone generated by the columns of A; distinct from cofactor and companion notation
$\mathcal{H}(\mathbf{N}, b),\ \mathcal{H}_{\pm}(\mathbf{N}, b),\ \langle \mathbf{N}, X \rangle < b$	affine hyperplane / closed half-spaces / setwise inequality, meaning $\langle \mathbf{N}, \mathbf{x} \rangle < b$ for all $\mathbf{x} \in X$

Spectral and Polynomial Structure

$E_\lambda(T),\ \mathrm{spec}(T)$	eigenspace $\ker(T - \lambda I_{\mathcal{V}})$ / spectrum
$\mathrm{Inv}(T)$	lattice or set of T-invariant subspaces
$\lambda_i(A),\ m_A(\lambda)$	eigenvalues in the locally declared order / algebraic multiplicity of λ as a root of p_A; distinct from measure notation $m(E)$
$\mathcal{G}_\lambda(T),\ J_k(\lambda)$	generalized eigenspace for λ / Jordan block of size k and eigenvalue λ
$p_A,\ q_T,\ \chi_T$	matrix characteristic polynomial / operator characteristic polynomial / characteristic polynomial determined by elementary divisors
$\mu_A,\ \mu_T,\ w_p(\lambda, A)$	minimal polynomial of a matrix or operator / Weyr count of λ-Jordan blocks of size at least p; distinct from measure notation $\mu(E)$
$P_j,\ T = \sum_j \mu_j P_j$	spectral projections / spectral decomposition with distinct eigenvalues μ_j
$\sqrt{T},\ A^{1/2},\ \rho(A)$	positive square root of a positive operator or matrix / spectral radius of A; distinct from the interval parameter in K_ρ
$\mathbb{F}[z]$	polynomials in z over $\mathbb{F}$
$\mathscr{P} = \mathbb{C}[x],\ \mathscr{P}',\ \mathscr{F} = \mathbb{C}[[t]]$	umbral polynomial algebra / algebraic dual / formal power-series algebra; script notation distinguishes these from $\mathcal{P}(X)$ and local families $\mathcal{F}$
$\mathcal{V}_T,\ \mathcal{C}_T(\mathbf{v}),\ \mathrm{ann}_T(\mathbf{v})$	$\mathbb{F}[z]$-module induced by T / T-cyclic subspace generated by $\mathbf{v}$ / annihilator ideal of $\mathbf{v}$
$m_{\mathbf{v}},\ C_f,\ \mathrm{InvFact}(T),\ \mathrm{ElemDiv}(T)$	order polynomial of $\mathbf{v}$ under T / companion matrix of f / invariant factors / elementary divisors
e^{tA}	matrix exponential $\sum_{k \geq 0} (tA)^k / k!$

Matrix Analysis and Numerical Structure

$A \succeq 0,\ A \succ 0$	Hermitian positive semidefinite / Hermitian positive definite
$\mathrm{in}(A),\ n_+(A), n_-(A), n_0(A)$	inertia of a Hermitian matrix / positive, negative, and zero eigenvalue counts
$A = QR,\ A = LL^*,\ A = LDL^t$	QR factorization / Cholesky factorization / symmetric or Hermitian LDL^t factorization
$A = V\Sigma W^*,\ A = X\Sigma_r Y^*$	singular value decomposition / compact SVD with rank r
$\sigma_j(A),\ \sigma_{\max}(A),\ \sigma_{\min}(A)$	singular values and extreme singular values
$\lvert T\rvert,\ \lvert A\rvert,\ \mathrm{vol}(\Omega)$	operator or matrix modulus / volume of a set; scalar absolute value remains $\lvert x\rvert$
$A^\dagger,\ A_k$	Moore–Penrose pseudoinverse / rank-k truncated SVD approximation in an SVD context
$\kappa_2(A),\ \Delta A,\ \Delta\mathbf{x}$	spectral condition number of an invertible matrix / matrix and solution perturbations
$\Delta(A)$	defect from normality; distinct from perturbation notation ΔA
$R_k(A),\ R_k^{\mathrm{off}}(A),\ R_{\max}(A)$	absolute row sum $\sum_j \lvert a_{kj}\rvert$ / deleted row sum $\sum_{j\neq k} \lvert a_{kj}\rvert$ / maximum absolute row sum
$\mathrm{Ger}_k(A),\ \mathrm{Ger}(A)$	kth Gersgorin disk / union of all Gersgorin disks
$\mathrm{tridiag}(a, b, c)$	tridiagonal matrix with subdiagonal a, diagonal b, and superdiagonal c
$w,\ F_n,\ Q_n,\ \omega_n$	half-bandwidth / unnormalised Fourier matrix / normalised Fourier matrix / primitive nth root used for the DFT
$\mathcal{K}_m(A, \mathbf{b})$	Krylov subspace $\mathrm{span}(\mathbf{b}, A\mathbf{b}, \ldots, A^{m-1}\mathbf{b})$
$Q_m,\ H_m$	Krylov basis matrix / projected Hessenberg or tridiagonal matrix
$A = S - T,\ B = S^{-1}T,\ \mathbf{e}_k$	matrix splitting / stationary iteration matrix / kth error vector
$A = D + L + U,\ B_J,\ B_{\mathrm{GS}}$	diagonal, lower, and upper splitting / Jacobi and Gauss–Seidel iteration matrices
$\omega,\ S_\omega,\ T_\omega,\ B_\omega$	SOR parameter / SOR splitting matrices / SOR iteration matrix; distinct from the differential form ω
$M,\ M^{-1}A,\ AM^{-1}$	preconditioner / left and right preconditioned matrices

Algebraic Structures Notation

Algebraic Language and Categories

$\mathcal{C},\ A \xrightarrow{f} B,\ \mathrm{id}_A$	ambient category / morphism $f : A \to B$ / identity morphism of A; $\mathcal{C}$ remains local unless declared
$\mathrm{Hom}_{\mathcal{C}}(A, B),\ g \circ f$	set of morphisms $A \to B$ in $\mathcal{C}$ / composition for compatible morphisms
$A \simeq B,\ F : \mathcal{C} \to \mathcal{D}$	isomorphic objects / functor with domain $\mathcal{C}$ and codomain $\mathcal{D}$

Groups, Actions, Rings, and Ideals

$G, H, K,\ e_G,\ \langle S\rangle$	groups or subgroups / identity element of G / substructure generated by S in the ambient algebraic category
$H \leq G,\ H \trianglelefteq G,\ G/H,\ [G : H]$	subgroup / normal subgroup / quotient group / subgroup index
$\mathrm{ord}(g),\ \mathrm{Aut}(G)$	order of an element / automorphism group
$R, S,\ 0_R,\ 1_R,\ R^\times$	rings / zero and identity elements of R / unit group of R
$Z(R),\ \mathrm{char}(R)$	center of R / characteristic of R
$I, J,\ \mathfrak{a}, \mathfrak{p},\ R/I$	ideals; fraktur letters denote distinguished ideals, such as prime ideals, when declared / quotient ring
$R[x],\ \mathbb{F}[z],\ \deg f$	polynomial rings in the displayed indeterminate / degree of a polynomial
$U_n,\ \Omega_n,\ \Phi_n(x),\ \varphi(n)$	complex nth roots of unity / primitive nth roots / nth cyclotomic polynomial / Euler totient; Ω_n is algebraic root notation, distinct from domain notation

Modules and Associative Algebras

$M, N,\ S \le M,\ M/N$	modules / submodule S of M / quotient module
$\mathrm{Hom}_R(M,N),\ \mathrm{End}_R(M)$	set of R-module homomorphisms $M \to N$ / endomorphism ring of M
$\ker\tau,\ \mathrm{im}\,\tau,\ \overline{\tau}$	kernel / image / quotient-induced map $M/\ker\tau \to \mathrm{im}\,\tau$ for a homomorphism $\tau : M \to N$
$\bigoplus_i M_i,\ M = N \oplus H$	external direct sum of modules / internal direct-sum decomposition
$\mathcal{C}_S(N),\ \mathcal{L}_\sigma,\ \mathcal{R}_\sigma$	set $\{H \le N : N = S \oplus H\}$ of complements / left inverses of σ / right inverses of σ; subscript distinguishes these from $\mathcal{L}$ and $\mathcal{L}^n$
$\mathcal{F}(\mathcal{V})$	finite-rank operators $\{T \in \mathcal{L}(\mathcal{V}) : \dim \mathrm{im}\, T < \infty\}$ in associative-algebra contexts; distinct from local function-family notation

Algebraic Geometry and Commutative Algebra Notation

Status. This part of the notes is still mostly scaffolded. The entries below reserve the notation promised by the section headings and keep it distinct from already-active algebra and analysis notation.

Commutative Algebra

$\mathrm{Spec}\, R,\ \mathfrak{p} \in \mathrm{Spec}\, R$	set of prime ideals of a commutative ring R / prime ideal $\mathfrak{p} \subseteq R$
$D(f),\ V(I)$	standard open set $\{\mathfrak{p} \in \mathrm{Spec}\, R : f \notin \mathfrak{p}\}$ / Zariski closed set $\{\mathfrak{p} : I \subseteq \mathfrak{p}\}$
$S^{-1}R,\ R_{\mathfrak{p}},\ k(\mathfrak{p})$	localization at a multiplicative set S / localization $(R \setminus \mathfrak{p})^{-1}R$ / residue field $R_{\mathfrak{p}}/\mathfrak{p}R_{\mathfrak{p}}$
$\mathrm{rad}(I),\ \sqrt{I}$	radical ideal $\{r \in R : r^n \in I \text{ for some } n \ge 1\}$; use $\mathrm{rad}(\mathcal{V})$ for metric-vector-space radicals

Affine and Projective Geometry

$\mathbb{A}^n_k,\ \mathbb{P}^n_k$	affine n-space over a field k / projective n-space over k; geometric notation distinct from probability notation
$I(X),\ Z(I),\ k[X]$	vanishing ideal $\{f \in k[x_1,\dots,x_n] : f|_X = 0\}$ / zero locus of an ideal / coordinate ring, typically $k[x_1,\dots,x_n]/I(X)$
$f^{\#},\ \varphi^*$	coordinate-ring pullback, e.g. $f^{\#} : k[Y] \to k[X]$ for $f : X \to Y$; distinct from adjoints and differential-form pullbacks by context

Analysis and Continuous Structure Notation

Sets and Functions in Analysis

$\mathcal{P}(X)$	power set $\{A : A \subseteq X\}$; projection symbols use $P_{\mathcal{U}}$ or P_j
$f(A),\ f^{-1}(B)$	image $\{f(a) : a \in A\}$ / preimage $\{x : f(x) \in B\}$ under f
$\mathrm{range}(f),\ f\|_A$	range $f(\mathrm{dom}\, f)$ / restriction of f to A
$\{A_i\}_{i\in I},\ \bigcup_{i\in I} A_i,\ \bigcap_{i\in I} A_i$	indexed family of sets / set of points lying in some A_i / set of points lying in every A_i
$g \circ f,\ \mathrm{id}_X$	composition of functions with compatible domain and codomain / identity function on X
$J_n,\ \#E,\ A \sim_{\#} B,\ [x]$	finite initial segment $J_0 = \emptyset$, $J_n = \{1,\dots,n\}$ / cardinality of a finite set / existence of a bijection $A \to B$ / equivalence class of x for the specified relation

Order and Completeness

$\sup E,\ \inf E$	least upper bound / greatest lower bound in the ambient ordered set, with values in $\overline{\mathbb{R}}$ when extended values are allowed
$\|x\|,\ \|a-b\|$	real absolute value or complex modulus / usual real-line distance $d(a,b) = \|a-b\|$
$\overline{\mathbb{R}}$	extended real line $\mathbb{R} \cup \{-\infty, +\infty\}$ ordered by $-\infty < x < +\infty$ for $x \in \mathbb{R}$; the bar is conventional, not topological closure
$x < y,\ x \le y$	strict order / non-strict order in the ambient ordered set

Sequences and Series

$(a_n),\ a_n \to a$	sequence indexed by $n \in \mathbb{N}$ / convergence of the sequence to a
$(a_{n_k}),\ n_1 < n_2 < \cdots$	subsequence indexed by a strictly increasing sequence (n_k) of positive integers
$s_n = \sum_{k=1}^{n} a_k$	partial-sum sequence associated to $\sum a_k$
$\limsup a_n = s^*,\ \liminf a_n = s_*$	upper and lower limits in $\overline{\mathbb{R}}$, equivalently $\inf_N \sup_{n\geq N} a_n$ and $\sup_N \inf_{n\geq N} a_n$ for real sequences
$\sum_{n=1}^{\infty} a_n,\ \prod_{n=1}^{\infty} a_n$	limit of partial sums / limit of finite partial products when defined

Metric and Topological Notation

$d(x,y),\ B_r(x)$	metric distance / open ball $\{y : d(x,y) < r\}$
$B_\varepsilon(a) = V_\varepsilon(a)$	real-line open ball $(a-\varepsilon, a+\varepsilon)$; $V_\varepsilon(a)$ is Abbott's notation
$\overline{B}_r(x),\ \{G_\alpha\}_{\alpha\in A}$	closed ball $\{y : d(x,y) \leq r\}$ / indexed open cover
$d_E,\ d_{\max},\ d_{\text{sum}}$	Euclidean, maximum, and sum product metrics on a product space
$E^\circ,\ \overline{E},\ E^{\text{acc}}$	interior / topological closure / set of limit points of E
$\text{bd}(E),\ E^c$	boundary $\overline{E} \setminus E^\circ$ / complement in the ambient space
$\operatorname{diam} E$	diameter $\sup\{d(p,q) : p,q \in E\}$ of a nonempty metric set
$X \simeq Y$	homeomorphic metric spaces, i.e. related by a bijection with continuous inverse
$A \cap \overline{B} = \emptyset,\ \overline{A} \cap B = \emptyset$	definition of separated subsets A and B
$F_\sigma,\ G_\delta,$ meager/thin/first category	countable union of closed sets / countable intersection of open sets / subset of a countable union of nowhere-dense sets
residual/thick, generic	set containing a countable intersection of dense open sets / property whose truth set is residual

Functions, Differentiation, and Integration

$\lim_{x\to p} f(x),\ f(x+),\ f(x-)$	limit of f at p / right and left limits at x
$C(X),\ C([a,b]),\ C_b(X),\ \|f\|_\infty$	continuous functions on X and $[a,b]$ / bounded continuous functions / supremum norm
$f_n \to f,\ f_n \rightrightarrows f$	pointwise convergence / uniform convergence on the ambient domain
$f'(x),\ Df(a),\ D_j f_i,\ \partial_i f,\ J_f(x)$	one-variable derivative, total derivative at a, component partial $D_j f_i$, scalar partial, and Jacobian determinant; J_f is distinct from J_n and $J_k(\lambda)$
$\nabla f,\ D_{\mathbf{u}} f,\ C^r,\ C^\infty,\ C^\omega$	gradient, directional derivative, and classes C^r, C^∞, C^ω
$\Pi = \{x_0, \ldots, x_n\},\ c_i,\ \Delta x_i,\ \Delta\alpha_i$	partition of $[a,b]$, tags $c_i \in [x_{i-1}, x_i]$, and increments $\Delta x_i = x_i - x_{i-1}$, $\Delta\alpha_i = \alpha(x_i) - \alpha(x_{i-1})$
$R(f,\Pi,\{c_i\}),\ \delta,\ \delta(x)$	$R(f,\Pi,\{c_i\}) = \sum_{i=1}^{n} f(c_i)\Delta x_i$ / constant mesh bound / gauge; uppercase $R(\cdots)$ is local to Riemann sums
$U(\Pi,f,\alpha),\ L(\Pi,f,\alpha)$	Riemann–Stieltjes upper and lower sums with integrator α
$f \in \mathcal{R},\ f \in \mathcal{R}(\alpha)$	ordinary Riemann integrability / Riemann–Stieltjes integrability with respect to α
$\int_a^b f,\ \int_a^b f\,d\alpha,\ \int_E f\,d\mu$	ordinary Riemann integral / Riemann–Stieltjes integral / measure integral
$\int_a^\infty f,\ \int_{-\infty}^\infty f$	one-sided and two-sided improper integrals as limits over finite truncations
$\ell(\Pi,\gamma),\ \ell(\gamma)$	polygonal length of γ along Π / curve length
$I^k,\ Q^k,\ \operatorname{supp}(f)$	k-cell, standard k-simplex, and closed support $\overline{\{x : f(x) \neq 0\}}$; I^k is not an identity operator
$dx_I,\ dx_I \wedge dx_J$	basic form indexed by I / wedge product of basic forms
$\omega,\ d\omega,\ T^*\omega,\ \partial_{\text{ch}}\Psi$	differential form, exterior derivative, pullback by T, and chain boundary
$\int_\Psi \omega,\ \partial_{\text{ch}}[\mathbf{p}_0, \ldots, \mathbf{p}_k]$	integral of ω over the chain Ψ / oriented boundary $\sum_{j=0}^{k}(-1)^j[\mathbf{p}_0, \ldots, \widehat{\mathbf{p}_j}, \ldots, \mathbf{p}_k]$
$dV,\ dA,\ ds$	volume, area, and arc-length elements

Power Series, Fourier, and Special Functions

$\sum_{n\geq 0} c_n(x-a)^n,\ R$	power series centered at a / radius of convergence in this context, distinct from $R(\cdots)$ and the R in QR
$I_R,\ K_\rho$	centered interval $I_R = (-R, R)$ / compact subinterval $K_\rho = [-\rho, \rho] \subset I_R$ for $0 < \rho < R$
$\widehat{f}(m),\ s_N(f;x),\ s_N(f)$	$\widehat{f}(m) = (2\pi)^{-1}\int_{-\pi}^{\pi} f(x)e^{-imx}\,dx$ / $s_N(f;x) = \sum_{\lvert m\rvert\leq N}\widehat{f}(m)e^{imx}$, viewed pointwise or as a function
$f \sim_{\text{coef}} \sum \widehat{f}(n)e^{inx},\ f \sim_{\text{coef}} \sum c_n\phi_n$	formal Fourier or orthonormal coefficient expansion; no convergence assertion is implied
$D_N(t)$	Dirichlet kernel $D_N(t) = \sum_{m=-N}^{N} e^{imt}$; uppercase D_N is not derivative notation
$\mathsf{E},\ \mathsf{L},\ \mathsf{C},\ \mathsf{S}$	power-series exponential, inverse logarithm, cosine, and sine
$\Gamma(x),\ \mathrm{B}(x,y)$	gamma function / beta function in roman type, distinct from balls $B_r(x)$ and variables B

Measure and Function Spaces

$E \in \mathcal{M},\ E \in \mathcal{M}(\mu)$	E measurable in the ambient σ-algebra / Rudin measurable set generated by μ
$\mathcal{B}(\mathbb{R}^n),\ \mathcal{L}^n$	Borel σ-algebra on $\mathbb{R}^n$ / Lebesgue measurable subsets of $\mathbb{R}^n$; not a basis $\mathcal{B}$, a linear-map space $\mathcal{L}$, or a least-squares set LS
$\mathcal{E},\ \mathcal{M}_F(\mu),\ A\triangle B,\ d_\mu(A,B)$	elementary sets, finitely μ-measurable sets, symmetric difference, and $d_\mu(A,B) = \mu^*(A\triangle B)$; $\mathcal{M}$ is measure-theoretic here
$\mu^*(E),\ m^*(E),\ \mathbf{1}_E$	μ-outer measure / Lebesgue outer measure / indicator $\mathbf{1}_E(x) = 1$ for $x \in E$ and 0 otherwise
$\mathrm{vol}_n(B),\ m_n^*(\Omega)$	n-dimensional box volume / box-cover outer measure $\inf\{\sum_{j\in J}\mathrm{vol}_n(B_j) : (B_j)_{j\in J} \text{ covers } \Omega,\ J \text{ at most countable}\}$
$m(E),\ \mu(E)$	Lebesgue or elementary volume of E / measure of E; not algebraic multiplicity $m_A(\lambda)$ or minimal polynomial μ_A
$\lambda^+,\ \lambda^-,\ \lvert\lambda\rvert,\ \lambda \ll \mu,\ \lambda \perp \mu,\ \frac{d\lambda}{d\mu}$	positive variation, negative variation, total variation, absolute continuity, mutual singularity, and Radon–Nikodym derivative
$\bowtie,\ \{f \bowtie a\}$	relation in $\{>, \geq, <, \leq\}$ / level set $\{x : f(x) \bowtie a\}$
$f^+,\ f^-,\ \mathcal{S}^+(f),\ \mathcal{I}_E(s)$	$f^+ = \max(f,0),\ f^- = -\min(f,0)$ / simple minorants of f / simple-function integral on E
$f \sim_{\text{a.e.}} g \text{ on } E,\ \text{a.e.}$	$\mu(\{x \in E : f(x) \neq g(x)\}) = 0$ / outside a null set
$E_x,\ E^y,\ f_x,\ f^y$	product slices $E_x = \{y : (x,y) \in E\}$, $E^y = \{x : (x,y) \in E\}$, and $f_x(y) = f^y(x) = f(x,y)$
$\mathcal{U}f,\ F_a$	undergraph $\{(x,t) : 0 < t < f(x)\}$ for $f \geq 0$ / superlevel set $F_a = \{x : f(x) > a\}$; $\mathcal{U}f$ is not a vector subspace
$[f:A],\ [B:Q],\ \mathrm{dp}(E)$	$[f:A] = m(A)^{-1}\int_A f$ / $[B:Q] = m(B\cap Q)/m(Q)$ / density-point set of E
$\mathscr{L}^1(\mu),\ \mathscr{L}^1_{\text{loc}}(\mathbb{R}^n)$	measurable f with $\int \lvert f\rvert\,d\mu < +\infty$ / locally integrable functions; script $\mathscr{L}$ is distinct from $\mathcal{L}^n$ and $\mathcal{L}$
$L^p(\mu),\ L^p(E),\ \lVert f\rVert_p$	Lebesgue space modulo a.e. equality with respect to μ or the ambient measure on E / $\lVert f\rVert_p = (\int \lvert f\rvert^p\,d\mu)^{1/p}$

Probability and Stochastic Structure Notation

Probability Spaces and Distributions

$(\Omega, \mathcal{F}, \mathbb{P})$, $\Pr(A)$	probability space / event probability $\Pr(A) = \mathbb{P}(A)$ for $A \in \mathcal{F}$; explicit probability notation keeps $\mathbb{P}_k^n$ and $\mathbb{F}_0$ unambiguous
X, Y, $X_*\mathbb{P}$, $\mathcal{L}(X)$	random variables / pushforward law $X_*\mathbb{P}(B) = \mathbb{P}(X^{-1}(B))$ / law of X when locally declared; not Lebesgue $\mathcal{L}^n$
$X \sim F$, $X \overset{\text{dist}}{=} Y$, $X_1, \dots, X_n$ iid	X has distribution F / X and Y have the same distribution / independent and identically distributed sample
$F(t)$, $g(t)$, $\Phi(t)$	distribution function $F(t) = \Pr(X \le t)$ / density g / standard normal distribution function in probabilistic contexts
$\mathbb{E}[X]$, $\operatorname{Var}(X)$, $\operatorname{Cov}(X, Y)$	expectation / $\operatorname{Var}(X) = \mathbb{E}[(X - \mathbb{E}X)^2]$ / $\operatorname{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}X)(Y - \mathbb{E}Y)]$, when finite
$X \perp\!\!\!\perp Y$	independence of random variables, equivalently independence of $\sigma(X)$ and $\sigma(Y)$ when defined
$X \perp\!\!\!\perp Y \mid Z$	conditional independence of X and Y given Z
$X_n \xrightarrow{\mathbb{P}} X$, $X_n \rightsquigarrow X$, $X_n \xrightarrow{L^p} X$, $X_n \xrightarrow{\text{a.s.}} X$, $X_n \xrightarrow{qm} X$	convergence in probability / convergence in distribution / convergence in L^p / almost-sure convergence / convergence in quadratic mean

Stochastic Processes and Markov Structure

$(X_t)_{t \in T}$, $\mathcal{F}_t$	stochastic process indexed by T / filtration $(\mathcal{F}_t)$ when locally declared; this $\mathcal{F}_t$ is not the umbral algebra $\mathscr{F}$
$\mathbf{x}$, $\mathbf{e}$, $A\mathbf{e} = \mathbf{e}$	probability vector $\mathbf{x} \ge 0$, $\sum_i x_i = 1$ / all-ones vector $\mathbf{e} = (1, \dots, 1)^t$ / row-stochastic Markov convention $A = [a_{ij}] \ge 0$, $A\mathbf{e} = \mathbf{e}$
π, $\pi A = \pi$, A^p	stationary distribution written as a row probability vector / stationarity equation / p-step Markov matrix
P_t, $p_t(x, y)$	transition semigroup / transition probability, kernel, or density from x to y, when introduced

Numerical and Computational Mathematics Notation

Status. Apart from numerical linear algebra, already covered under Linear and Operator Structures, the numerical and computational chapters remain placeholders. The entries below reserve notation only; they do not introduce PDE or machine-learning conventions.

Floating Point, Approximation, and Error

$\operatorname{fl}(x)$, u, $\varepsilon_{\text{mach}}$	floating-point rounded value, unit roundoff, and machine epsilon
x_h, e_h, $\mathcal{O}(h^p)$	approximation at discretization scale h, corresponding error, and Landau order as $h \to 0$
$\widehat{x}$, Δx, $r = b - A\widehat{x}$	computed approximation, perturbation or error in x, and residual for $Ax = b$

Monte Carlo and Stochastic Computation

$X_1, \dots, X_N$, $\widehat{\mu}_N$	samples / sample-mean estimator $\widehat{\mu}_N = N^{-1} \sum_{i=1}^N X_i$ when estimating $\mu = \mathbb{E}X$
$\operatorname{Var}(\widehat{\mu}_N)$, $\operatorname{SE}(\widehat{\mu}_N)$	estimator variance / standard error $\operatorname{SE}(\widehat{\mu}_N) = \sqrt{\operatorname{Var}(\widehat{\mu}_N)}$
w_i, $\widehat{I}_N$	importance weights / Monte Carlo estimate of an integral or expectation, when locally introduced

Part I
Mathematical Language, Sets, and Logic

Foundational prerequisites collect the logical, set-theoretic, functional, order-theoretic, and cardinal language used before the dedicated algebraic and analytic foundations sections.

1 Sets

Sets provide the membership language used throughout the notes: elements, subsets, products, power sets, operations, indexed families, and the first warnings about unrestricted comprehension.

1.1 Introduction to Sets

Definition 1.1.1. *Sets and Operations*
A *set* is a collection. Write $x \in E$, $x \notin E$, $\emptyset$, and $A \subseteq B$; $A = B$ iff $A \subseteq B$ and $B \subseteq A$. For $A, B \subseteq X$,

$$\begin{aligned} A \cup B &:= \{x \in X : x \in A \text{ or } x \in B\}, \\ A \cap B &:= \{x \in X : x \in A \text{ and } x \in B\}, \\ A \setminus B &:= \{x \in X : x \in A \text{ and } x \notin B\}, \\ A^c &:= X \setminus A. \end{aligned}$$

1.2 Cartesian Products

Definition 1.2.1. *Cartesian Product*
For sets X and Y, the *Cartesian product* is $X \times Y := \{(x, y) : x \in X,\ y \in Y\}$.
The order matters: the first coordinate comes from X, and the second comes from Y.

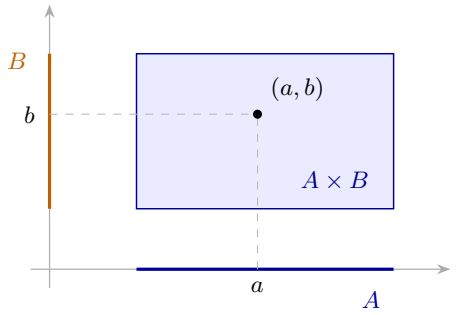


Figure 1: $A \times B$: first coordinate in A, second in B.

1.3 Subsets and Power Sets

Definition 1.3.1. *Power Set*
Let X be a set. The *power set* is $\mathcal{P}(X) := \{A : A \subseteq X\}$.

Method 1.3.2. *Proof by Two Containments*
To prove $A = B$, prove both inclusions.

(i) To prove $A \subseteq B$, let $x \in A$ be arbitrary and use the definitions and hypotheses to show $x \in B$.
(ii) To prove $B \subseteq A$, let $x \in B$ be arbitrary and use the definitions and hypotheses to show $x \in A$.
(iii) Once both containments are proved, conclude $A = B$ by extensionality of sets.

If both sides are described by membership conditions, this is a chain, $x \in A \iff \cdots \iff x \in B$.

Placeholder: Add extensionality examples, proper subset notation, and power-set examples.

1.4 Union, Intersection, Difference, and Complement

Remark 1.4.1. *Set-Builder Notation*
$\{x \in X : P(x)\}$ means all $x \in X$ satisfying $P(x)$.
For sets (A_n), $\bigcup_n A_n$ means membership in some A_n; $\bigcap_n A_n$, membership in every A_n.
Key Insight: Expand membership conditions.

Proposition 1.4.2. *Basic Set Algebra*
Let $A, B, C \subseteq X$. Then

(i) $A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$;
(ii) $A \cup (B \cap C) = (A \cup B) \cap (A \cup C)$;
(iii) $(A \cup B)^c = A^c \cap B^c$;
(iv) $(A \cap B)^c = A^c \cup B^c$.

Proof. It suffices to prove (i) and (iii); the others are similar. Let $x \in X$. Then

(i) $x \in A \cap (B \cup C) \iff x \in A$ and ($x \in B$ or $x \in C$) $\iff$ ($x \in A \cap B$) or ($x \in A \cap C$) $\iff x \in (A \cap B) \cup (A \cap C)$
(iii) $x \in (A \cup B)^c \iff x \notin A \cup B \iff x \notin A$ and $x \notin B \iff x \in A^c \cap B^c$

Structural Commentary: Set algebra is membership logic.
Key Insight: Translate sets into conditions on x. □

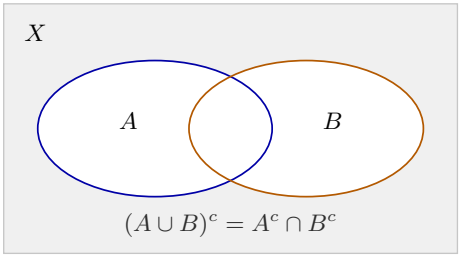


Figure 2: De Morgan's law says that being outside a union is the same as being outside each set.

Theorem 1.4.3. *Arbitrary De Morgan Laws*
Let X be a set. For any collection $\{E_\alpha\}$ of subsets of X, $\left(\bigcup_\alpha E_\alpha\right)^c = \bigcap_\alpha E_\alpha^c$ *and* $\left(\bigcap_\alpha E_\alpha\right)^c = \bigcup_\alpha E_\alpha^c$.

Proof. The first identity follows from membership: $x \in \left(\bigcup_\alpha E_\alpha\right)^c \iff x \notin E_\alpha \ \forall \alpha \iff x \in \bigcap_\alpha E_\alpha^c$.
The second follows by applying the first identity to the complements E_α^c, or by the same membership argument.
Structural Commentary: Complements exchange unions and intersections.
Key Insight: Negating "for some" gives "for every", and conversely. □

1.5 Indexed Sets and Number Systems

Definition 1.5.1. *Indexed Families of Sets*
An *indexed family* $\{A_i\}_{i \in I}$ is one set $A_i \subseteq X$ for each index $i \in I$. Define

$$\bigcup_{i \in I} A_i := \{x \in X : x \in A_i \text{ for some } i \in I\},$$
$$\bigcap_{i \in I} A_i := \{x \in X : x \in A_i \text{ for every } i \in I\}.$$

Key Insight: Union means "some index"; intersection means "every index".

Placeholder: Add the standard number-system inclusions and the convention that indexed notation binds its index variable.

1.6 Russell's Paradox and Axioms

Placeholder: Add Russell's paradox as a motivation for restricted comprehension, then state the informal role of extensionality, separation, pairing, union, power set, infinity, replacement, foundation, and choice.

2 Logic and Proofs

Logic fixes the grammar of mathematical statements before the notes use that grammar for definitions, theorems, and proofs.

2.1 Statements and Well-Formed Language

Principle 2.1.1. *Well-Formed Statements and Bound Variables*
A mathematical statement is well formed only after every symbol has a declared role and every expression has a defined meaning. A variable introduced by a quantifier, summation sign, integral, set-builder condition, or similar construction is *bound* inside that construction; changing its name does not change the statement as long as its scope is unchanged and no free variable is captured.

2.2 Connectives, Conditionals, and Truth Tables

Placeholder: Add the truth-functional meanings of $\wedge, \vee, \neg, \Rightarrow, \iff$, inclusive "or", truth tables, logical equivalence, converse, inverse, and contrapositive.

2.3 Quantifiers, Bound Variables, and Negation

Principle 2.3.1. *Language Carries Structure*
Objects, hypotheses, conclusions, and quantifier order matter. For instance,

$$\forall x \in X\ \exists y \in Y\ P(x,y), \qquad \exists y \in Y\ \forall x \in X\ P(x,y)$$

are different: in the first y may depend on x; in the second one y works for all x.

Method 2.3.2. *Proof Grammar and Negation*
In theorem statements, put the hypotheses before the conclusion and keep quantifiers in the order meant.
In particular, switching $\forall$ and $\exists$ usually changes the statement; see Principle 2.3.1.
Before applying a rule, also check that the displayed expressions are well formed; an undefined expression makes the sentence ill formed, not merely false.
The word "or" is inclusive in mathematics, and the basic negation rules are

$$\neg(a \wedge b) \iff \neg a \vee \neg b, \qquad \neg(a \vee b) \iff \neg a \wedge \neg b, \qquad \neg(a \Rightarrow b) \iff a \wedge \neg b$$

Thus a proof by contradiction of $a \Rightarrow b$ assumes both a and $\neg b$ and derives an impossibility.
Key Insight: Before proving or negating a statement, parse its quantifiers and logical connectives.

Placeholder: Add negation rules for quantified statements and examples translating English into symbolic logic.

2.4 Proof Grammar and Inference

Placeholder: Add inference rules such as modus ponens, universal instantiation, existential introduction, and proof templates for definitions.

2.5 Direct, Contrapositive, and Contradiction Proofs

Placeholder: Add direct proof, proof by cases, contrapositive proof, contradiction, biconditional proof, existence proof, uniqueness proof, and disproof by counterexample.

2.6 Induction, Recursion, and Counterexamples

Principle 2.6.1. *Mathematical Induction*
Let $P_1, P_2, \ldots$ be statements indexed by positive integers. To prove P_n for every $n \in \mathbb{N}$, it is enough to prove:

(i) *basis*: P_1 is true;
(ii) *induction step*: for every $n \in \mathbb{N}$, P_n implies P_{n+1}.

The induction step is an implication; it does not prove P_{n+1} in isolation.
The basis starts the chain, and the implication carries truth from one statement to the next.

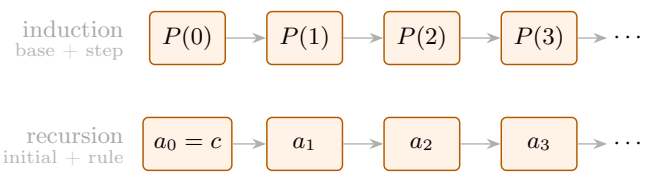


Figure 3: Induction propagates truth, while recursion propagates values along the same index chain.

Principle 2.6.2. *Well-Ordering of $\mathbb{N}$*
Every nonempty subset of $\mathbb{N}$ has a least element, justifying first unused indices, first repeated values, and increasing listings of infinite subsets of $\mathbb{N}$.

Placeholder: Add strong induction, smallest-counterexample arguments, recursion, and the relation between induction and well-ordering of $\mathbb{N}$.

3 Counting and Finite Sums

Counting records finite constructions before cardinality and infinite sets are treated abstractly.

3.1 Lists and Counting Principles

Placeholder: Add lists, ordered choices, multiplication principle, addition principle, subtraction principle, and division principle.

3.2 Permutations, Subsets, and Binomial Coefficients

Definition 3.2.1. *Permutations and Sign*
Let $n \geq 1$. Let $\mathfrak{S}_n$ denote the set of permutations of $\{1, \dots, n\}$.
The *sign* of $\sigma \in \mathfrak{S}_n$ is $\operatorname{sign}(\sigma) = 1$ if number of pairs $j < k$ with $\sigma(j) > \sigma(k)$ is even, and $\operatorname{sign}(\sigma) = -1$ otherwise.

Definition 3.2.2. *Factorials and Binomial Coefficients*

(i) The *factorial* is defined by $0! := 1$ and, for $n \in \mathbb{N}$ with $n \geq 1$, $n! := n(n-1)\cdots 3 \cdot 2 \cdot 1$.
(ii) For integers $0 \leq k \leq n$, the *binomial coefficient* is $\binom{n}{k} := \frac{n!}{k!(n-k)!}$.
(iii) The symbol $\binom{n}{k}$ is read *n choose k* and counts the k-element subsets of an n-element set.

Example 3.2.3. *Committees*
The number of ways to choose a committee of 3 people from a class of 20 people is $\binom{20}{3} = \frac{20!}{3!17!} = \frac{20 \cdot 19 \cdot 18}{3 \cdot 2 \cdot 1} = 1140$.

Proposition 3.2.4. *Edge Values and Symmetry*
For $n \in \mathbb{N}$ and $k \in \mathbb{Z}$ with $0 \leq k \leq n$:

(i) Edge values: $\binom{n}{0} = \binom{n}{n} = 1$.
(ii) Symmetry: $\binom{n}{k} = \binom{n}{n-k}$.

Proof. Let $n \in \mathbb{N}$ and $k \in \mathbb{Z}$ with $0 \leq k \leq n$.

(i) Since $0! := 1$, $\binom{n}{0} = \frac{n!}{0!n!} = 1$ and $\binom{n}{n} = \frac{n!}{n!0!} = 1$.
(ii) Since $0 \leq n-k \leq n$, $\binom{n}{n-k} = \frac{n!}{(n-k)!(n-(n-k))!} = \frac{n!}{(n-k)!k!} = \frac{n!}{k!(n-k)!} = \binom{n}{k}$.

Structural Commentary: Choosing the selected k objects is equivalent to choosing the omitted $n-k$ objects.
Key Insight: Complements turn k-subsets into $(n-k)$-subsets. □

3.3 Inclusion-Exclusion, Pigeonhole, and Combinatorial Proof

Definition 3.3.1. *Multisets*
Let S be a set. A *multiset on S* is a pair (S_0, m) such that

(i) $S_0 \subseteq S$ is the *underlying set*;
(ii) $m : S_0 \to \mathbb{Z}_{>0}$ is the *multiplicity function*;
(iii) $m(s)$ is the *multiplicity* of $s \in S_0$;
(iv) if S_0 is finite, the *size* is $|m| := \sum_{s \in S_0} m(s)$.

Equivalently, record (S_0, m) as $\{(s, m(s)) : s \in S_0\}$.

Example 3.3.2. *Multiset Bookkeeping*
The list $\{a, a, b, b, b, c\}$ is the multiset on $\{a, b, c\}$ with multiplicities $2, 3, 1$. This language is useful whenever order is irrelevant but repetition matters, such as eigenvalues counted with algebraic multiplicity.

Placeholder: Add inclusion-exclusion, pigeonhole arguments, and combinatorial proof as equality by double counting.

3.4 Finite Sums

Definition 3.4.1. *Finite Sums Indexed by Integers*

Let $a_i \in \mathbb{C}$ be assigned for integers $m \le i \le n$. Define $\sum_{i=m}^{n} a_i := a_m + a_{m+1} + \cdots + a_n$.

If $m > n$, the sum over the empty index set is 0.

Proposition 3.4.2. *Basic Laws for Finite Sums*

Let $a_i, b_i \in \mathbb{C}$ and $c \in \mathbb{C}$ be assigned for $m \le i \le n$.

(i) If $m \le r < n$, then $\sum_{i=m}^{n} a_i = \sum_{i=m}^{r} a_i + \sum_{i=r+1}^{n} a_i$.

(ii) $\sum_{i=m}^{n} (a_i + b_i) = \sum_{i=m}^{n} a_i + \sum_{i=m}^{n} b_i$ and $\sum_{i=m}^{n} ca_i = c \sum_{i=m}^{n} a_i$.

(iii) $\left| \sum_{i=m}^{n} a_i \right| \le \sum_{i=m}^{n} |a_i|$.

(iv) If $a_i, b_i \in \mathbb{R}$ and $a_i \le b_i$ for every i, then $\sum_{i=m}^{n} a_i \le \sum_{i=m}^{n} b_i$.

Proof.

(i) By Definition 3.4.1 and associativity of addition in $\mathbb{C}$, $\sum_{i=m}^{n} a_i = \sum_{i=m}^{r} a_i + \sum_{i=r+1}^{n} a_i$.

(ii) This follows from associativity, commutativity, and distributivity in the field $\mathbb{C}$.

(iii) Induct on $n - m$. The case $n = m$ is immediate.
If the claim holds up to n, then the triangle inequality gives

$$\left| \sum_{i=m}^{n+1} a_i \right| \le \left| \sum_{i=m}^{n} a_i \right| + |a_{n+1}| \le \sum_{i=m}^{n} |a_i| + |a_{n+1}| = \sum_{i=m}^{n+1} |a_i| .$$

(iv) Induct on $n - m$. The case $n = m$ is the hypothesis $a_m \le b_m$.
If the claim holds up to n, then compatibility of the order on $\mathbb{R}$ with addition gives

$$\sum_{i=m}^{n+1} a_i = \sum_{i=m}^{n} a_i + a_{n+1} \le \sum_{i=m}^{n} b_i + b_{n+1} = \sum_{i=m}^{n+1} b_i.$$

Structural Commentary: Finite sums are iterated addition, so their laws are inherited from the scalar field.
Key Insight: No convergence hypothesis is needed for a finite sum. □

Definition 3.4.3. *Summation over a Finite Set*

Let X be a finite set with $\#X = n$, and let $f : X \to \mathbb{C}$.
Choose any bijection $g : J_n \to X$, where $J_n = \{1, \ldots, n\}$, and set $\sum_{x \in X} f(x) := \sum_{i=1}^{n} f(g(i))$.

When $X = \emptyset$, this is the empty sum 0.
Key Insight: Proposition 3.4.4 shows that the value is independent of the chosen bijection.

Proposition 3.4.4. *Finite Set Summation Laws*

Let X, Y be finite sets.

(i) Well-definedness: different bijections $J_{\#X} \to X$ give the same value in Definition 3.4.3.

(ii) Substitution: if $h : Y \to X$ is a bijection and $f : X \to \mathbb{C}$, then $\sum_{x \in X} f(x) = \sum_{y \in Y} f(h(y))$.

(iii) Linearity: for $f, g : X \to \mathbb{C}$ and $c \in \mathbb{C}$, $\sum_{x \in X} (f(x)+g(x)) = \sum_{x \in X} f(x) + \sum_{x \in X} g(x)$ and $\sum_{x \in X} cf(x) = c \sum_{x \in X} f(x)$.

(iv) Disjoint unions: if $X \cap Y = \emptyset$ and $f : X \cup Y \to \mathbb{C}$, then $\sum_{z \in X \cup Y} f(z) = \sum_{x \in X} f(x) + \sum_{y \in Y} f(y)$.

(v) Finite Fubini: if $f : X \times Y \to \mathbb{C}$, then $\sum_{x \in X} \sum_{y \in Y} f(x, y) = \sum_{(x,y) \in X \times Y} f(x, y) = \sum_{y \in Y} \sum_{x \in X} f(x, y)$.

Proof.

(i) For $f : X \to \mathbb{C}$, let $n = \#X$ and let $g, k : J_n \to X$ be bijections. Induct on n. The case $n = 0$ is empty. If $n \ge 1$, let $j \in J_n$ satisfy $k(j) = g(n)$. Removing n from g and j from k gives bijections onto $X \setminus \{g(n)\}$; after relabeling by J_{n-1}, the induction hypothesis makes the remaining sums equal. Adding the common last term $f(g(n)) = f(k(j))$ gives equality of the two sums.

(ii) Let $g : J_{\#Y} \to Y$ be a bijection. Then $h \circ g : J_{\#Y} \to X$ is a bijection, so $\#Y = \#X$ by Proposition 6.1.3.
Hence, by (i) and Definition 3.4.3, $\sum_{x \in X} f(x) = \sum_{i=1}^{\#Y} f(h(g(i))) = \sum_{y \in Y} f(h(y))$.

(iii) Choose a bijection $\gamma : J_{\#X} \to X$.
Proposition 3.4.2 gives $\sum_{x\in X}(f(x)+g(x)) = \sum_{x\in X} f(x) + \sum_{x\in X} g(x)$ and $\sum_{x\in X} cf(x) = c\sum_{x\in X} f(x)$

(iv) Let $\alpha : J_{\#X} \to X$ and $\beta : J_{\#Y} \to Y$ be bijections, and enumerate $X \cup Y$ by listing α first and then β.
Since $X \cap Y = \emptyset$, this is a bijection, and Proposition 3.4.2 gives $\sum_{z\in X\cup Y} f(z) = \sum_{x\in X} f(x) + \sum_{y\in Y} f(y)$.

(v) Let $\alpha : J_{\#X} \to X$ and $\beta : J_{\#Y} \to Y$ be bijections.
The maps $\ell((i-1)\#Y + j) = (\alpha(i), \beta(j))$ and $r((j-1)\#X + i) = (\alpha(i), \beta(j))$ for $i \in J_{\#X}$ and $j \in J_{\#Y}$ give the two lexicographic bijections $J_{(\#X)(\#Y)} \to X \times Y$.
Applying Definition 3.4.3, Proposition 3.4.2, and (i) yields $\sum_{x\in X}\sum_{y\in Y} f(x,y) = \sum_{(x,y)\in X\times Y} f(x,y) = \sum_{y\in Y}\sum_{x\in X} f(x,y)$.

Structural Commentary: A finite set sum is an order-free finite sum.
Key Insight: Bijections are bookkeeping devices; the value does not depend on the bookkeeping order. □

4 Relations, Orders, and Lattices

Relations package comparisons, identifications, order structures, and quotient constructions.

4.1 Relations and Equivalence

Definition 4.1.1. *Relations and Equivalence*
A *relation* from X to Y is a subset $R \subseteq X \times Y$. If $(x, y) \in R$, one often writes xRy.
A relation $\sim$ on X is an *equivalence relation* if, for all $x, y, z \in X$,

i. $x \sim x$;
ii. $x \sim y$ implies $y \sim x$;
iii. $x \sim y$ and $y \sim z$ imply $x \sim z$.

The equivalence class of x is $[x] := \{y \in X : y \sim x\}$.

4.2 Partitions, Congruence, and Quotients

Proposition 4.2.1. *Equivalence Classes Partition the Set*
Let $\sim$ be an equivalence relation on X. Then the equivalence classes form a partition of X.

Proof. Every $x \in X$ lies in $[x]$ since $x \sim x$. If $[x] \cap [y] \neq \emptyset$, choose z in the intersection.
Then $x \sim z \sim y$, so $u \in [x]$ implies $u \sim y$ and $u \in [y]$. Hence $[x] \subseteq [y]$, and symmetry gives equality.
Structural Commentary: Equivalence means partition.
Key Insight: Touching classes coincide. □

Placeholder: Add congruence modulo n, quotient sets, and well-defined operations on equivalence classes.

4.3 Ordered Sets, Well-Orders, and Lattices

Definition 4.3.1. *Ordered Set Vocabulary*
Let S be a set, let $E, C \subseteq S$, and let $\leq_S$ be a relation on S. Write $x <_S y$ iff $x \leq_S y$ and $x \neq y$.

(i) $\leq_S$ is a *partial order* if it is reflexive, antisymmetric, and transitive; then $(S, \leq_S)$ is a *poset*.
(ii) $\leq_S$ is an *order* if it is a partial order and any $x, y \in S$ are comparable; then $(S, \leq_S)$ is an *ordered set*.
(iii) C is a *chain* if $\leq_S |_C$ is an order.
(iv) S is *well-ordered* if every nonempty $Y \subseteq S$ has $\ell \in Y$ with $\ell \leq_S y$ for all $y \in Y$.
(v) $\beta \in S$ is an *upper bound* of E if $x \leq_S \beta$ for all $x \in E$; lower bounds reverse the inequality.
(vi) E is *bounded above* if it has an upper bound; bounded below is dual.
(vii) $\sup E$ is the least upper bound of E, and $\inf E$ is the greatest lower bound of E, when these exist.
(viii) $m \in E$ is *maximal* if no $x \in E$ satisfies $m <_S x$; minimal is dual.
(ix) S has the *least-upper-bound property* if every nonempty bounded-above $E \subseteq S$ has $\sup E \in S$.

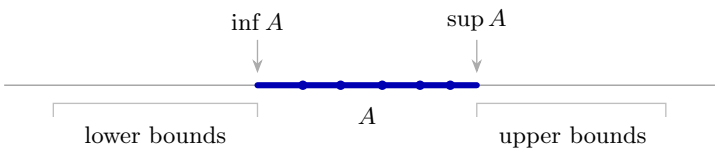


Figure 4: The supremum is the least upper bound, while the infimum is the greatest lower bound.

Example 4.3.2. *Suprema Need Not Be Elements*
In $\mathbb{Q}$, the set $E_1 = \{r \in \mathbb{Q} : r < 0\}$ and the set $E_2 = \{r \in \mathbb{Q} : r \leq 0\}$ both have supremum 0, but $0 \notin E_1$ and $0 \in E_2$. If $E = \{1/n : n = 1, 2, 3, \ldots\}$, then $\sup E = 1$ while $\inf E = 0$, and the infimum is not a member of E.

Theorem 4.3.3. *Least Upper Bounds Give Greatest Lower Bounds*
If S has the least-upper-bound property, then every nonempty bounded-below set $B \subseteq S$ has an infimum.

Proof. Let L be the set of lower bounds of B. Then L is nonempty and every $x \in B$ is an upper bound of L, so $\alpha = \sup L$ exists. Since every $x \in B$ bounds L above, $\alpha \leq x$ for all $x \in B$; hence α is a lower bound of B. If $\beta > \alpha$, then $\beta \notin L$, so β is not a lower bound. Thus $\alpha = \inf B$.
Structural Commentary: Infima are suprema of lower bounds.
Key Insight: Completeness above gives completeness below. □

Definition 4.3.4. *Lattice Vocabulary*
Let $(L, \leq)$ be a partially ordered set.

(i) A *meet* of $x, y \in L$ is $\inf\{x, y\}$ when it exists, often written $x \wedge y$.
(ii) A *join* of $x, y \in L$ is $\sup\{x, y\}$ when it exists, often written $x \vee y$.
(iii) A *lattice* is a partially ordered set in which every two elements have a meet and a join.
(iv) A lattice is *complete* if every subset $E \subseteq L$ has an infimum and a supremum.
(v) A *bottom* element is an element $\bot \in L$ with $\bot \leq x$ for every $x \in L$; a *top* element is an element $\top \in L$ with $x \leq \top$ for every $x \in L$.
(vi) A lattice is *modular* if $x \leq z$ implies $z \wedge (y \vee x) = (z \wedge y) \vee x$ for all $x, y, z \in L$.

4.4 Relations Between Sets and Closures

Placeholder: Add relations between different sets, inverse relations, relation composition, reflexive/transitive closure, and order-isomorphism examples.

5 Functions

Functions are relations with unique outputs; their images, preimages, restrictions, compositions, and inverses organize nearly every later construction.

5.1 Functions, Images, and Preimages

Definition 5.1.1. *Functions, Images, and Preimages*
Let $f : X \to Y$ be a *function*, assigning each $x \in X$ a unique $f(x) \in Y$.

(i) $X = \mathrm{dom}(f)$ is the *domain* and Y is the *codomain*;
(ii) $\mathrm{range}(f) := f(X) := \{f(x) : x \in X\}$ is the *range* of f;
(iii) for $A \subseteq X$, $f(A) := \{f(x) : x \in A\}$ is the *image* of A;
(iv) for $B \subseteq Y$, $f^{-1}(B) := \{x \in X : f(x) \in B\}$ is the *preimage* of B; no inverse function is required;
(v) if Y has distinguished 0, $\mathrm{supp}(f) := \{x \in X : f(x) \neq 0\}$;

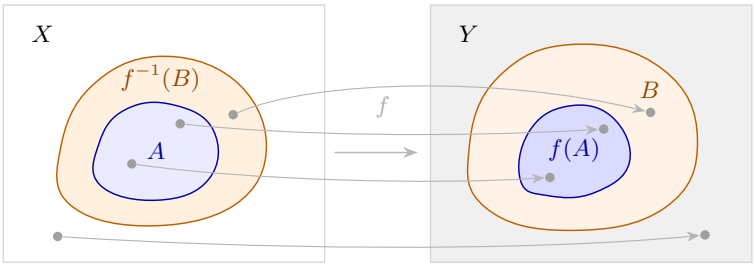


Figure 5: Image–preimage containment: $f(A) \subseteq B \iff A \subseteq f^{-1}(B)$.

Proposition 5.1.2. *Image and Preimage Calculus*
Let $f : X \to Y$, let $\{A_i\}_{i \in I}$ be a nonempty family of subsets of X, and let $\{B_i\}_{i \in I}$ be a nonempty family of subsets of Y. Then

(i) $f\left(\bigcup_{i \in I} A_i\right) = \bigcup_{i \in I} f(A_i)$;
(ii) $f\left(\bigcap_{i \in I} A_i\right) \subseteq \bigcap_{i \in I} f(A_i)$;

(iii) $f^{-1}\left(\bigcup_{i\in I} B_i\right) = \bigcup_{i\in I} f^{-1}(B_i)$;
(iv) $f^{-1}\left(\bigcap_{i\in I} B_i\right) = \bigcap_{i\in I} f^{-1}(B_i)$;
(v) $f^{-1}(Y \setminus B) = X \setminus f^{-1}(B)$ *for every* $B \subseteq Y$;
(vi) $f(f^{-1}(B)) = B \cap \operatorname{range}(f)$ *for every* $B \subseteq Y$.

Proof.

(i) $y \in f(\bigcup_i A_i)$ iff $y = f(x)$ for some $x \in A_i$ and some i, iff $y \in \bigcup_i f(A_i)$.
(ii) If $y = f(x)$ with $x \in \bigcap_i A_i$, then $y \in f(A_i)$ for every i.
(iii) $x \in f^{-1}(\bigcup_i B_i)$ iff $f(x) \in B_i$ for some i, iff $x \in \bigcup_i f^{-1}(B_i)$.
(iv) $x \in f^{-1}(\bigcap_i B_i)$ iff $f(x) \in B_i$ for every i, iff $x \in \bigcap_i f^{-1}(B_i)$.
(v) For $x \in X$, $x \in f^{-1}(Y \setminus B) \iff f(x) \notin B \iff x \notin f^{-1}(B) \iff x \in X \setminus f^{-1}(B)$
(vi) $y \in f(f^{-1}(B))$ iff $y = f(x)$ for some x with $f(x) \in B$, iff $y \in B \cap \operatorname{range}(f)$.

Structural Commentary: Preimages are Boolean-exact; images need not be.
Key Insight: Pull back conditions through f. □

5.2 Injective, Surjective, and Bijective Maps

Definition 5.2.1. *Injective, Surjective, and Bijective Maps*
Let $f : X \to Y$.

i. f is *injective* if $f(x_1) = f(x_2)$ implies $x_1 = x_2$.
ii. f is *surjective* if $\operatorname{range}(f) = Y$.
iii. f is *bijective* if it is both injective and surjective.

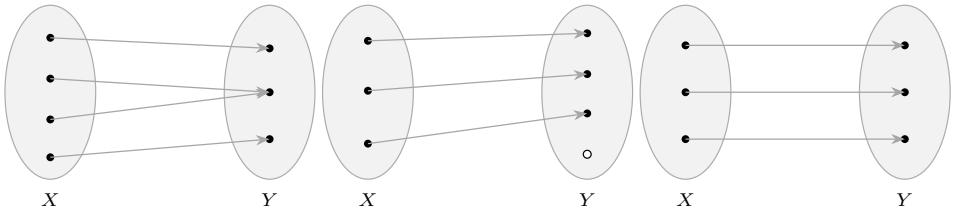


Figure 6: Surjective: covers codomain; injective: no collisions; bijective: both.

Example 5.2.2. *Image and Preimage for the Square Map*
Let $f : \mathbb{R} \to \mathbb{R}$ be given by $f(x) = x^2$. Then $f([-2,1]) = [0,4]$, $f^{-1}([1,4]) = [-2,-1] \cup [1,2]$. The map is not injective because $f(1) = f(-1)$, and it is not surjective as a map $\mathbb{R} \to \mathbb{R}$ because no real number maps to -1.

5.3 Composition, Restriction, and Inverse Functions

Definition 5.3.1. *Composition, Restriction, Extension, and Identity*
Let $f : X \to Y$.

(i) for $A \subseteq X$, $f|_A : A \to Y$ is the *restriction* of f to A;
(ii) if $X \subseteq U$, an *extension* of f to U is $\widetilde{f} : U \to Y$ with $\widetilde{f}|_X = f$;
(iii) if $g : Y \to Z$, then $g \circ f : X \to Z$ is given by $(g \circ f)(x) = g(f(x))$;
(iv) $\mathrm{id}_X : X \to X$ denotes the *identity function*.

Theorem 5.3.2. *Bijective Means Invertible*
A function $f : X \to Y$ is bijective iff there exists a function $g : Y \to X$ s.t. $g \circ f = \mathrm{id}_X$, $f \circ g = \mathrm{id}_Y$. When such a function exists, it is unique and is denoted $f^{-1} : Y \to X$.

Proof.

($\Rightarrow$) If f is bijective, define $g(y)$ as the unique x with $f(x) = y$. Then $g \circ f = \mathrm{id}_X$ and $f \circ g = \mathrm{id}_Y$.
($\Leftarrow$) If such a g exists, applying g to $f(x_1) = f(x_2)$ gives $x_1 = x_2$, and $y = f(g(y))$ gives surjectivity.

If g and h are two-sided inverses, then $g(y) = g(f(h(y))) = h(y)$, so the inverse is unique.
Structural Commentary: A true inverse undoes f on both sides.
Key Insight: Bijective means exactly one preimage per output. □

Placeholder: Add algebra of composition and image/preimage laws under composition.

6 Cardinality, Choice, and Infinite Sets

Cardinality compares sets by functions rather than by informal size.

6.1 Same Cardinality, Finite, Countable, and Uncountable

Definition 6.1.1. *Same Cardinality*
Two sets A and B have *same cardinality*, or are *equivalent*, if $\exists$ a bijection $f : A \to B$. Then $A \sim_{\#} B$.
The relation $\sim_{\#}$ is an equivalence relation: $A \sim_{\#} A$ by identity map, symmetry follows from Theorem 5.3.2, and transitivity follows by composing bijections.

Definition 6.1.2. *Finite, Countable, and Uncountable*
Set $J_0 := \emptyset$, and for $n \geq 1$ set $J_n := \{1, 2, \ldots, n\}$. A set A is

(i) *finite* if $A \sim_{\#} J_n$ for some $n \in \mathbb{N}$;
(ii) *infinite* if it is not finite;
(iii) *countable* if $A \sim_{\#} \mathbb{N}$;
(iv) *uncountable* if it is neither finite nor countable;
(v) *at most countable* if it is finite or countable.

Countable sets are also called enumerable or denumerable.

Proposition 6.1.3. *Finite Cardinality is Well Defined*
If $A \sim_{\#} J_m$ and $A \sim_{\#} J_n$, then $m = n$. Thus $\#A = n$ is well defined for finite sets.

Proof. By transitivity of $\sim_{\#}$, it suffices to prove $J_m \sim_{\#} J_n \Rightarrow m = n$. Induct on m.
For $m = 0$, only $J_0 = \emptyset$ is bijective to J_0.
If $m, n \geq 1$ and $f : J_m \to J_n$ is bijective, remove m and its partner $f(m)$. The remaining sets are bijective and relabel as J_{m-1} and J_{n-1}, so induction gives $m - 1 = n - 1$, hence $m = n$.
Structural Commentary: Finite counting is invariant under bijection.
Key Insight: Remove one matched element and count down. □

Definition 6.1.4. *Sequences as Enumerations*
A *sequence* in set A is $x : \mathbb{N} \to A$, written (x_n), where $x_n := x(n)$. Terms of sequence need not be distinct.
If A is countable, then its elements can be arranged as a sequence with distinct terms. Conversely, a sequence whose terms are distinct and whose range is A gives a bijection $\mathbb{N} \to A$.

Example 6.1.5. *The Integers Are Countable*
The set $\mathbb{Z}$ is countable. One explicit enumeration is given by

$$f(n) = \begin{cases} \frac{n}{2}, & n \text{ even}, \\ -\frac{n-1}{2}, & n \text{ odd}, \end{cases} \qquad n \in \mathbb{N}.$$

This lists $\mathbb{Z}$ as $0, 1, -1, 2, -2, 3, -3, \ldots$.

Theorem 6.1.6. *Infinite Subsets of Countable Sets*
Every infinite subset of a countable set is countable.

Proof. Let $E \subseteq A$ be infinite, and enumerate A by distinct terms (x_n). Let $n_1 < n_2 < \cdots$ be exactly the indices for which $x_{n_k} \in E$, listed in increasing order. Then $k \mapsto x_{n_k}$ is a bijection from $\mathbb{N}$ onto E.
Structural Commentary: The proof keeps only the entries of the original enumeration that land in E.
Key Insight: An infinite subset of an enumeration inherits an enumeration by first appearances. □

Proposition 6.1.7. *Subsets and Images of At Most Countable Sets*
Let A be at most countable.

(i) If $B \subseteq A$, then B is at most countable.
(ii) If $f : A \to Y$ is a function, then $f(A)$ is at most countable.

Proof. If A is finite, both claims are immediate. Otherwise write $A = \{a_n : n \in \mathbb{N}\}$ with the a_n distinct.

(i) If B is infinite, let $I_B := \{n : a_n \in B\} = \{n_1 < n_2 < \cdots\}$.
Then $k \mapsto a_{n_k}$ bijects $\mathbb{N}$ onto B; the finite case is trivial.
(ii) Let $I := \{n : f(a_n) \notin \{f(a_m) : m < n\}\}$. By (i), I is at most countable, $n \mapsto f(a_n)$ bijects I onto $f(A)$.

Structural Commentary: Subsets select indices; images select first appearances.
Key Insight: Subsets and images preserve at-most-countability. □

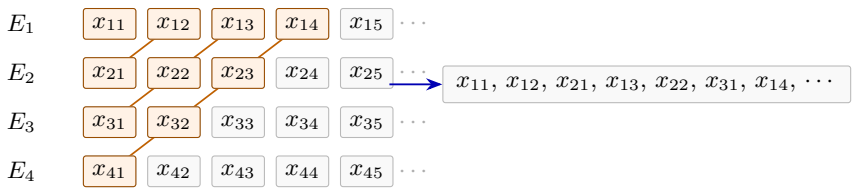


Figure 7: Diagonal scan of countable rows into one sequence, skipping repeats.

Theorem 6.1.8. *Countable Union of Countable Sets*
Let (E_n) be a sequence of countable sets and put $S := \bigcup_{n=1}^{\infty} E_n$. Then S is countable.

Proof. Enumerate each E_n as $(x_{nk})_{k\in\mathbb{N}}$.
Reading the array (x_{nk}) by diagonals, as in Figure 7, gives a sequence whose image is S.
Thus S is at most countable by Proposition 6.1.7.
Since $E_1 \subseteq S$ and E_1 is countable, S is infinite; hence S is countable.
Structural Commentary: The diagonal order turns a two-indexed list into a one-indexed list.
Key Insight: Countably many countable rows can still be scanned by one sequence. □

Corollary 6.1.9. *At Most Countable Unions*
If A is at most countable and B_α is at most countable for each $\alpha \in A$, then $\bigcup_{\alpha\in A} B_\alpha$ is at most countable.

Proof. Let $T := \bigcup_{\alpha\in A} B_\alpha$. If $T = \emptyset$, we are done. Otherwise enumerate the nonempty sets among the B_α as $B_{\alpha_1}, B_{\alpha_2}, \ldots$, repeating indices if needed, and list each B_{α_j} as $(b_{jk})_{k\in\mathbb{N}}$.
A diagonal scan of (b_{jk}) has image T, so Proposition 6.1.7 makes T at most countable.
Structural Commentary: Finite pieces cause no difficulty; they can be embedded into countable rows.
Key Insight: At most countable indexing + at most countable indexed sets gives at most countable total set. □

Theorem 6.1.10. *Finite Tuples over a Countable Set*
Let A be countable. For each fixed $n \in \mathbb{N}$, the set of all n-tuples $(a_1, \ldots, a_n)$ with $a_k \in A$ is countable.

Proof. Let $B_n := \{(a_1, \ldots, a_n) : a_k \in A\}$, $B_1 = A$. If B_{n-1} is countable, $B_n = \bigcup_{b\in B_{n-1}} \{(b, a) : a \in A\}$, each set $\{(b, a) : a \in A\}$ is equivalent to A. Hence B_n is countable by Theorem 6.1.8.
Structural Commentary: Build an n-tuple by appending one coordinate to an $(n-1)$-tuple.
Key Insight: Induction plus countable unions controls finite products. □

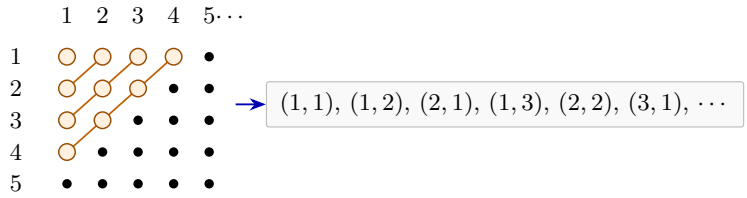


Figure 8: Diagonal enumeration of $\mathbb{N} \times \mathbb{N}$ by finite $i + j$ diagonals.

Corollary 6.1.11. *Countable Cartesian Products*
If A, B are at most countable, then $A \times B$ is at most countable.
If A, B are countable and nonempty, then $A \times B$ is countable. Hence $\mathbb{N} \times \mathbb{N}$ is countable.

Proof. For $a \in A$, $\{a\} \times B \sim B$, and $A \times B = \bigcup_{a\in A} \big(\{a\} \times B\big)$, thus $A \times B$ at most countable (Corollary 6.1.9). If A, B are countable and $b_0 \in B$, then $a \mapsto (a, b_0)$ injects A into $A \times B$, so $A \times B$ is infinite; hence it is countable.
Structural Commentary: A product is a countable union of countable slices.
Key Insight: Two countable coordinates can still be scanned by one sequence. □

Corollary 6.1.12. *The Rationals Are Countable*
The set $\mathbb{Q}$ is countable.

Proof. Let $D := \{(a,b) \in \mathbb{Z} \times \mathbb{Z} : a \neq 0\}$ and $\pi : D \to \mathbb{Q}$ by $\pi(a,b) = b/a$.
Since $\mathbb{Z}$ is countable, Corollary 6.1.11 and Proposition 6.1.7 make D at most countable. Since $\pi(D) = \mathbb{Q}$, Proposition 6.1.7 makes $\mathbb{Q}$ at most countable. As $\mathbb{Z} \subseteq \mathbb{Q}$, it is countable.
Structural Commentary: Repeated fractions only add repetitions to the list.
Key Insight: Countable pairs produce countably many rational values. □

6.2 Diagonalization and Power Sets

Theorem 6.2.1. *Cantor's Theorem*
For every set X, no function $f : X \to \mathcal{P}(X)$ is surjective. Hence X and $\mathcal{P}(X)$ cannot have the same cardinality.

Proof. Suppose $f : X \to \mathcal{P}(X)$ were surjective, and set $D := \{x \in X : x \notin f(x)\}$.
Since $D \subseteq X$, we have $D \in \mathcal{P}(X)$, so surjectivity gives $D = f(a)$ for some $a \in X$.
But then $a \in D \iff a \notin f(a) \iff a \notin D$, a contradiction.
Structural Commentary: The diagonal set is defined to disagree with each proposed row at its own index.
Key Insight: Power sets are too large to be exhausted by their base set. □

$$\begin{array}{lcl} s_1 & = & \boxed{a_{11}}\ a_{12}\ a_{13}\ a_{14}\ a_{15} \\ s_2 & = & a_{21}\ \boxed{a_{22}}\ a_{23}\ a_{24}\ a_{25} \\ s_3 & = & a_{31}\ a_{32}\ \boxed{a_{33}}\ a_{34}\ a_{35}\ \cdots \\ s_4 & = & a_{41}\ a_{42}\ a_{43}\ \boxed{a_{44}}\ a_{45} \\ s_5 & = & a_{51}\ a_{52}\ a_{53}\ a_{54}\ \boxed{a_{55}} \end{array}$$

Figure 9: Cantor diagonalization: $s(n) = 1 - a_{nn}$ differs from row s_n at position n.

Theorem 6.2.2. *Binary Sequences Are Uncountable*
Let A be the set of all sequences whose terms are the digits 0 and 1. Then A is uncountable.

Proof. Suppose A were countable, $A = \{s_1, s_2, \ldots\}$ with $s_n = (s_n(k))_{k \in \mathbb{N}}$. Define $s \in A$ by $s(k) := 1 - s_k(k)$.
Then $s(n) \neq s_n(n)$ for every n, so $s \neq s_n$ for all n. Hence $s \notin A$, a contradiction. Thus A is uncountable.
Structural Commentary: The construction attacks the diagonal of a proposed enumeration.
Key Insight: A sequence can be built to escape every row of a countable list. □

Definition 6.2.3. *Continuum Cardinal*
The *continuum cardinal* is $\mathfrak{c} := 2^{\aleph_0} := \left|\{0,1\}^{\mathbb{N}}\right|$. Equivalently, $\mathfrak{c} = |\mathcal{P}(\mathbb{N})|$ via $s \mapsto \{n \in \mathbb{N} : s(n) = 1\}$.

Perspective 6.2.4. *Continuum Cardinality*
Theorem 6.2.2 and Theorem 6.2.1 give $\aleph_0 < \mathfrak{c}$.
Thus the continuum is the first cardinal jump produced by binary choice over $\mathbb{N}$.

6.3 Schroeder–Bernstein and Cardinal Comparison

Theorem 6.3.1. *Schroeder–Bernstein Theorem*
If there are injections $f : A \to B$ and $g : B \to A$, then there is a bijection $h : A \to B$.

Proof. Let $A_0 := A \setminus g(B)$, and recursively define $B_n := f(A_n)$ and $A_{n+1} := g(B_n)$.
The sets A_n are pairwise disjoint: if $A_i \cap A_j \neq \emptyset$ with $i < j$, then $i = 0$ already contradicts $A_j \subseteq g(B)$; if $i > 0$, injectivity of g and then f steps backward to $A_{i-1} \cap A_{j-1} \neq \emptyset$, and iteration gives the same contradiction.
Since f is injective, the sets $B_n = f(A_n)$ are pairwise disjoint as well.
Let $A_* := A \setminus \bigcup_{n \geq 0} A_n$ and $B_* := B \setminus \bigcup_{n \geq 0} B_n$. The map g restricts to a bijection $B_* \to A_*$.
Indeed, $g(B_*)$ avoids every A_n, while if $a \in A_*$ then $a \notin A_0$, so $a = g(b)$ for a unique $b \in B$; this b cannot lie in any B_n, otherwise $a \in A_{n+1}$.
Define $h(x) := \begin{cases} f(x), & x \in \bigcup_{n \geq 0} A_n, \\ g^{-1}(x), & x \in A_*, \end{cases}$ where g^{-1} denotes the inverse of $g : B_* \to A_*$.
This maps each A_n bijectively onto B_n and maps A_* bijectively onto B_*. Hence h is a bijection $A \to B$.
Structural Commentary: Two injections carve both sets into matched layers plus a common remainder.
Key Insight: Use f on the escaping layers and g^{-1} on the stable remainder. □

Definition 6.3.2. *Cardinal Comparison and Arithmetic*
Let A, B be sets, and let $\kappa = |A|$, $\lambda = |B|$.

(i) *Cardinal comparison*: $|A| \leq |B|$ iff $\exists$ an injection $A \to B$.
(ii) *Strict cardinal comparison*: $|A| < |B|$ iff $|A| \leq |B|$ and $A \not\sim_{\#} B$.
(iii) If $A \cap B = \emptyset$, the *cardinal sum* is $\kappa + \lambda := |A \cup B|$.
(iv) The *cardinal product* is $\kappa\lambda := |A \times B|$.
(v) The *function set* is $A^B := \{f : B \to A\}$, and the *cardinal power* is $\kappa^\lambda := \left|A^B\right|$.

6.4 Cardinal Arithmetic

Proposition 6.4.1. *Cardinal Arithmetic Reference Laws*
Let κ, λ, μ be cardinal numbers.

(i) $\kappa + \lambda = \lambda + \kappa$ *and* $\kappa\lambda = \lambda\kappa$;
(ii) $\kappa + (\lambda + \mu) = (\kappa + \lambda) + \mu$ *and* $\kappa(\lambda\mu) = (\kappa\lambda)\mu$;
(iii) $\kappa(\lambda + \mu) = \kappa\lambda + \kappa\mu$;
(iv) $\kappa^{\lambda+\mu} = \kappa^\lambda \kappa^\mu$, $(\kappa^\lambda)^\mu = \kappa^{\lambda\mu}$, *and* $(\kappa\lambda)^\mu = \kappa^\mu \lambda^\mu$.

Proof. Choose pairwise disjoint sets A, B, C with $|A| = \kappa$, $|B| = \lambda$, and $|C| = \mu$.

(i) Use $A \cup B \cong B \cup A$ and $A \times B \cong B \times A$.
(ii) Use $A \cup (B \cup C) \cong (A \cup B) \cup C$ and $A \times (B \times C) \cong (A \times B) \times C$.
(iii) Use $A \times (B \cup C) \cong (A \times B) \cup (A \times C)$.
(iv) Use $A^{B \cup C} \cong A^B \times A^C$, $(A^B)^C \cong A^{B \times C}$, and $(A \times B)^C \cong A^C \times B^C$.

Taking cardinalities gives (i)–(iv).
Structural Commentary: The laws are canonical bijections.
Key Insight: Cardinal arithmetic is set construction up to bijection. □

Theorem 6.4.2. *Infinite Cardinal Absorption*
Let κ, λ be cardinals and set $\mu := \max\{\kappa, \lambda\}$. If μ is infinite, then $\kappa + \lambda = \mu$.
If also $\kappa, \lambda > 0$, then $\kappa\lambda = \mu$. In particular, $\aleph_0 + \aleph_0 = \aleph_0 \cdot \aleph_0 = \aleph_0$.

Proof. For every infinite cardinal ν, the standard square law gives $\nu + \nu = \nu^2 = \nu$.
Since $\kappa, \lambda \leq \mu$, $\mu \leq \kappa + \lambda \leq \mu + \mu = \mu$, so $\kappa + \lambda = \mu$.
If $\kappa, \lambda > 0$, choose nonempty representatives $|A| = \kappa$ and $|B| = \lambda$; fixing one point in the smaller factor injects the larger factor into $A \times B$, so $\mu \leq \kappa\lambda$. Also $\kappa\lambda \leq \mu^2 = \mu$, so $\kappa\lambda = \mu$.
Structural Commentary: Infinite cardinals absorb smaller summands and nonzero factors.
Key Insight: Squeeze between μ and $\mu^2 = \mu$. □

6.5 Choice, Zorn's Lemma, and Well-Ordering

Principle 6.5.1. *Choice Principle*
For nonempty sets $\{X_i\}_{i \in I}$, the axiom of choice asserts $\exists c : I \to \bigcup_{i \in I} X_i$ such that $c(i) \in X_i$ $\forall i \in I$.
Use direct construction for finite or explicit choices; use choice only for arbitrary simultaneous choices.

Principle 6.5.2. *Zorn's Lemma*
If a nonempty partially ordered set X has the property that every chain in X has an upper bound in X, then X contains a maximal element. This is equivalent in strength to the axiom of choice, Principle 6.5.1.

Method 6.5.3. *Cardinality Proof Patterns*

(i) Same cardinality: construct a bijection.
(ii) Countable: construct a sequence whose range contains every element.
(iii) Uncountable: show every proposed sequence misses some element.
(iv) Choice/maximality: cite Principle 6.5.1 for arbitrary simultaneous choices and Principle 6.5.2 for maximal elements in partially ordered sets.

Key Idea: Use maps: bijections prove equal cardinality, enumerations $\mathbb{N} \to A$ prove countability, and diagonal arguments show every proposed enumeration misses an element.

Placeholder: Add the well-ordering theorem, maximality methods, and examples of where choice is avoidable versus structurally necessary.

7 Axiomatic and Advanced Set Theory

This section reserves space for the deeper set-theoretic material suggested by the set-theory and metatheory contents: axioms, ordinals, models, forcing, large cardinals, and sets of reals.

7.1 Extensionality, Specification, and Replacement

Placeholder: Add ZF-style axiom statements, informal motivations, and the separation/replacement distinction.

7.2 Ordinals, Transfinite Recursion, and Well-Foundedness

Placeholder: Add well-founded relations, transitive sets, ordinal numbers, ordinal arithmetic, transfinite induction, and transfinite recursion.

7.3 Models, Constructibility, and Independence

Placeholder: Add models of set theory, the constructible universe, forcing previews, consistency strength, and independence statements.

7.4 Filters, Ultrafilters, and Sets of Reals

Placeholder: Add filters, ideals, ultrafilters, stationary/closed-unbounded sets, Borel and projective sets, descriptive set theory, and large-cardinal gateways.

Part II
Linear and Operator Structures

Linear algebra moves from structure to geometry, spectra, and computation: vector spaces, maps, and duality; inner products, projections, least squares, and adjoints; invariant subspaces, polynomial structure, and spectral theorems; decompositions, norms, perturbation, and numerics.

1 Structural Linear Algebra

Structural linear algebra develops the coordinate-free language of vector spaces, decompositions, linear maps, and duality.

1.1 Vector Spaces and Linear Structure

Vector spaces encode addition and scalar multiplication. Subspaces, spans, bases, dimension, and coordinates describe how linear structure is generated and measured.

Definition 1.1.1. *Vector Space*
Let $\mathbb{F}$ be a field. A vector space over $\mathbb{F}$ is a set $\mathcal{V}$ with vector addition and scalar multiplication such that, $\forall\, \mathbf{u}, \mathbf{v}, \mathbf{w} \in \mathcal{V}$ and $\forall\, a, b \in \mathbb{F}$, the following axioms hold:

i. vector addition is commutative: $\mathbf{u} + \mathbf{v} = \mathbf{v} + \mathbf{u}$;
ii. vector addition is associative: $(\mathbf{u} + \mathbf{v}) + \mathbf{w} = \mathbf{u} + (\mathbf{v} + \mathbf{w})$;
iii. there exists an additive identity $\mathbf{0} \in \mathcal{V}$ such that $\mathbf{0} + \mathbf{v} = \mathbf{v}$;
iv. for each $\mathbf{v} \in \mathcal{V}$ there exists an additive inverse $-\mathbf{v} \in \mathcal{V}$ such that $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$;
v. scalar multiplication has identity: $1\mathbf{v} = \mathbf{v}$;
vi. scalar multiplication is compatible with multiplication in $\mathbb{F}$: $(ab)\mathbf{v} = a(b\mathbf{v})$;
vii. scalar multiplication distributes over vector addition: $a(\mathbf{u} + \mathbf{v}) = a\mathbf{u} + a\mathbf{v}$;
viii. scalar multiplication distributes over addition in $\mathbb{F}$: $(a + b)\mathbf{v} = a\mathbf{v} + b\mathbf{v}$.

If $\mathcal{V} = \{\mathbf{0}\}$, then $\mathcal{V}$ is a *zero vector space*; otherwise $\mathcal{V}$ is a *nonzero vector space*.

Definition 1.1.2. *Finitely Nonzero Sequence Space*
Let $\mathbb{F}^{(\mathbb{N})}$ denote the set of sequences $(x_1, x_2, \ldots)$ in $\mathbb{F}$ with finite support, i.e. $x_j = 0$ for all but finitely many $j \in \mathbb{N}$. Addition and scalar multiplication are componentwise.

Proposition 1.1.3. *Finitely Nonzero Sequences Form a Vector Space*
The set $\mathbb{F}^{(\mathbb{N})}$ is a vector space over $\mathbb{F}$ under componentwise addition and scalar multiplication.

Proof. The zero sequence has finite support.
If $\mathbf{x}, \mathbf{y} \in \mathbb{F}^{(\mathbb{N})}$ and $a \in \mathbb{F}$, then $a\mathbf{x} + \mathbf{y}$ can be nonzero only where $\mathbf{x}$ or $\mathbf{y}$ is nonzero, a finite set of indices.
Hence $a\mathbf{x} + \mathbf{y} \in \mathbb{F}^{(\mathbb{N})}$, and the vector space axioms are inherited componentwise from $\mathbb{F}$.
Structural Commentary: Finite support is preserved by finite linear combinations.
Key Insight: Componentwise operations cannot create infinitely many new nonzero entries. □

Example 1.1.4. *Standard Vector Space Models*
Let $\mathbb{F}$ be a field and let X be a set. The following are $\mathbb{F}$-vector spaces.

(i) $\mathbb{F}^X = \{f : X \to \mathbb{F}\}$, with pointwise operations $(f + g)(x) = f(x) + g(x)$ and $(af)(x) = af(x)$.
(ii) $\mathrm{Mat}(m, n, \mathbb{F})$ and $\mathbb{F}^n$, with entrywise operations. For $\mathbb{F} = \mathbb{F}_q$, set $V(n, q) := \mathbb{F}_q^n$.
(iii) $\mathrm{Seq}(\mathbb{F}) = \mathbb{F}^{\mathbb{N}}$, with componentwise operations; $\mathbb{F}^{(\mathbb{N})}$ is its finite-support subspace.
(iv) Over $\mathbb{C}$, the sequence spaces $c_0 = \{(z_n) : z_n \to 0\}$, $\ell^\infty = \{(z_n) : \sup_n |z_n| < \infty\}$, and, for $1 \le p < \infty$, $\ell^p = \left\{(z_n) : \sum_{n=1}^{\infty} |z_n|^p < \infty\right\}$ are vector spaces under componentwise operations; for ℓ^p, closure under addition follows from Minkowski's inequality.

Structural Commentary: The vector-space axioms are independent of finite-dimensional coordinates; function and sequence spaces give canonical infinite-dimensional models.
Key Insight: Pointwise or componentwise operations inherit linear structure from the scalar field.

Theorem 1.1.5. *Basic Consequences of Vector Space Axioms*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$. Then:

(i) the additive identity is unique;
(ii) for each $\mathbf{v} \in \mathcal{V}$, the additive inverse of $\mathbf{v}$ is unique;

(iii) $\forall \mathbf{v} \in \mathcal{V},\ 0\mathbf{v} = \mathbf{0}$;
(iv) $\forall a \in \mathbb{F},\ a\mathbf{0} = \mathbf{0}$;
(v) $\forall a \in \mathbb{F}$ *and* $\mathbf{v} \in \mathcal{V},\ a\mathbf{v} = \mathbf{0} \iff a = 0$ *or* $\mathbf{v} = \mathbf{0}$.

Proof.

(i) If $\mathbf{0}$ and $\mathbf{0}'$ are additive identities, then $\mathbf{0} = \mathbf{0} + \mathbf{0}' = \mathbf{0}'$.
(ii) If $\mathbf{w}, \mathbf{w}'$ are additive inverses of $\mathbf{v}$, then $\mathbf{w} = \mathbf{w} + \mathbf{0} = \mathbf{w} + (\mathbf{v} + \mathbf{w}') = (\mathbf{w} + \mathbf{v}) + \mathbf{w}' = \mathbf{0} + \mathbf{w}' = \mathbf{w}'$.
(iii) $0\mathbf{v} = (0 + 0)\mathbf{v} = 0\mathbf{v} + 0\mathbf{v}$; adding $-(0\mathbf{v})$ gives $0\mathbf{v} = \mathbf{0}$.
(iv) $a\mathbf{0} = a(\mathbf{0} + \mathbf{0}) = a\mathbf{0} + a\mathbf{0}$; adding $-(a\mathbf{0})$ gives $a\mathbf{0} = \mathbf{0}$.
(v) ($\Rightarrow$) Suppose $a\mathbf{v} = \mathbf{0}$. If $a = 0$, the first alternative holds.
If $a \neq 0$, then $\mathbf{v} = 1\mathbf{v} = (a^{-1}a)\mathbf{v} = a^{-1}(a\mathbf{v}) = a^{-1}\mathbf{0} = \mathbf{0}$.
(v) ($\Leftarrow$) If $a = 0$ or $\mathbf{v} = \mathbf{0}$, then (iii) or (iv) gives $a\mathbf{v} = \mathbf{0}$.

Structural Commentary: These identities make cancellation, zeros, and inverses available.
Key Insight: The vector space axioms force the zero and inverse behaviour to be rigid, not optional structure. □

Corollary 1.1.6. *Scalar Multiplication by -1 Is Negation*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$. Then $\forall \mathbf{v} \in \mathcal{V},\ (-1)\mathbf{v} = -\mathbf{v}$.

Proof. For $\mathbf{v} \in \mathcal{V}$, $(-1)\mathbf{v} + \mathbf{v} = (-1)\mathbf{v} + 1\mathbf{v} = (-1 + 1)\mathbf{v} = 0\mathbf{v} = \mathbf{0}$.
Hence $(-1)\mathbf{v}$ is an additive inverse of $\mathbf{v}$; by uniqueness, $(-1)\mathbf{v} = -\mathbf{v}$.
Structural Commentary: Scalar multiplication by -1 is exactly the same operation as taking an additive inverse.
Key Insight: Negation is not an extra operation; it is already encoded by the field action. □

Definition 1.1.7. *Subspace*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$. A subset $\mathcal{U} \subseteq \mathcal{V}$ is a *subspace*, written $\mathcal{U} \leq \mathcal{V}$, if $\mathcal{U}$ is a vector space over $\mathbb{F}$ under the operations inherited from $\mathcal{V}$.

Theorem 1.1.8. *Subspace Test*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$ and let $\mathcal{U}$ be a nonempty subset of $\mathcal{V}$.
Then $\mathcal{U}$ is a subspace of $\mathcal{V}$ iff $\forall \mathbf{u}, \mathbf{v} \in \mathcal{U}$ and $\forall c \in \mathbb{F}$, one has $c\mathbf{u} + \mathbf{v} \in \mathcal{U}$.

Proof.

($\Rightarrow$) If $\mathcal{U} \leq \mathcal{V}$, then closure under scalar multiplication and addition gives $c\mathbf{u} + \mathbf{v} \in \mathcal{U}$ for all $\mathbf{u}, \mathbf{v} \in \mathcal{U}$ and $c \in \mathbb{F}$.
($\Leftarrow$) Assume the stated closure condition. Choose $\mathbf{u}_0 \in \mathcal{U}$. Then $\mathbf{0} = (-1)\mathbf{u}_0 + \mathbf{u}_0 \in \mathcal{U}$. For $\mathbf{u}, \mathbf{v} \in \mathcal{U}$ and $c \in \mathbb{F}$, the identities $-\mathbf{u} = (-1)\mathbf{u} + \mathbf{0}$, $c\mathbf{u} = c\mathbf{u} + \mathbf{0}$, and $\mathbf{u} + \mathbf{v} = 1\mathbf{u} + \mathbf{v}$ give closure under additive inverses, scalar multiplication, and addition. The remaining axioms are inherited from $\mathcal{V}$, so $\mathcal{U} \leq \mathcal{V}$.

Structural Commentary: One closure condition captures the subspace axioms.
Key Insight: To prove a subset is a subspace, prove it survives one affine-looking linear combination. □

Definition 1.1.9. *Linear Code and Hamming Weight*

(i) The *ambient code space* is $V(n, q) := \mathbb{F}_q^n$.
(ii) A *linear code* of length n over $\mathbb{F}_q$ is a *subspace* $C \leq V(n, q)$.
(iii) For $\mathbf{x} = (x_1, \ldots, x_n) \in \mathbb{F}_q^n$, its *Hamming weight* is $\mathrm{wt}(\mathbf{x}) := \#\{j : 1 \leq j \leq n,\ x_j \neq 0\}$.

Example 1.1.10. *Even-Weight Code*
In $V(n, 2) = \mathbb{F}_2^n$, let $E_n = \left\{ \mathbf{x} \in \mathbb{F}_2^n : \sum_{j=1}^{n} x_j = 0 \right\}$. The condition $\sum_{j=1}^{n} x_j = 0$ is equivalent to $\mathrm{wt}(\mathbf{x})$ being even.
If $\mathbf{x}, \mathbf{y} \in E_n$ and $c \in \mathbb{F}_2$, then $\sum_{j=1}^{n} (cx_j + y_j) = c\sum_{j=1}^{n} x_j + \sum_{j=1}^{n} y_j = 0$, so $c\mathbf{x} + \mathbf{y} \in E_n$. Hence $E_n \leq V(n, 2)$ by Theorem 1.1.8, and E_n is a linear code.
Structural Commentary: The even-weight code is the parity-zero subspace in $V(n, 2)$.
Key Insight: Parity is linear over $\mathbb{F}_2$.

Proposition 1.1.11. *Intersection of Subspaces*
Let $\mathcal{U}$ and $\mathcal{W}$ be subspaces of $\mathcal{V}$. Then $\mathcal{U} \cap \mathcal{W} = \{\mathbf{v} : \mathbf{v} \in \mathcal{U}$ and $\mathbf{v} \in \mathcal{W}\}$ is a subspace of $\mathcal{V}$.

Proof. Since $\mathbf{0} \in \mathcal{U}$ and $\mathbf{0} \in \mathcal{W}$, we have $\mathbf{0} \in \mathcal{U} \cap \mathcal{W}$, so the intersection is nonempty.
If $\mathbf{u}, \mathbf{v} \in \mathcal{U} \cap \mathcal{W}$ and $c \in \mathbb{F}$, then $c\mathbf{u} + \mathbf{v}$ lies in both $\mathcal{U}$ and $\mathcal{W}$ because each is a subspace.
Thus $c\mathbf{u} + \mathbf{v} \in \mathcal{U} \cap \mathcal{W}$, and Theorem 1.1.8 gives $\mathcal{U} \cap \mathcal{W} \leq \mathcal{V}$.
Structural Commentary: Intersections inherit shared closure.
Key Insight: Common vectors remain closed under linear combinations. □

Definition 1.1.12. *Lists, Linear Combinations, and Span*

(i) A *list* in $\mathcal{V}$ is $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_m)$; the empty list is $()$.

(ii) For $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_m)$ and $a_1, \ldots, a_m \in \mathbb{F}$, a *linear combination* of β is $\sum_{j=1}^{m} a_j \mathbf{v}_j$.

(iii) For $E \subseteq \mathcal{V}$, the *span* of E is $\operatorname{span}(E) := \left\{ \sum_{j=1}^{m} a_j \mathbf{v}_j : m \in \mathbb{N},\ a_j \in \mathbb{F},\ \mathbf{v}_j \in E \right\}$, with $\operatorname{span}(\emptyset) = \{\mathbf{0}\}$.

(iv) For $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_m)$, the *span* of β is $\operatorname{span}(\beta) := \operatorname{span}(\{\mathbf{v}_1, \ldots, \mathbf{v}_m\})$.

(v) $E \subseteq \mathcal{V}$ *spans* $\mathcal{V}$ iff $\operatorname{span}(E) = \mathcal{V}$; a list β *spans* $\mathcal{V}$ iff $\operatorname{span}(\beta) = \mathcal{V}$.

Theorem 1.1.13. *Span is the Smallest Containing Subspace*
Let $E \subseteq \mathcal{V}$. Then $\operatorname{span}(E)$ is the smallest subspace of $\mathcal{V}$ containing E.

Proof. If $E = \emptyset$, then $\operatorname{span}(E) = \{\mathbf{0}\}$, which is a subspace. Suppose $E \neq \emptyset$. For $\mathbf{x}, \mathbf{y} \in \operatorname{span}(E)$ and $c \in \mathbb{F}$, both $\mathbf{x}$ and $\mathbf{y}$ are finite linear combinations of vectors in E, so $c\mathbf{x} + \mathbf{y}$ is also such a combination. Hence $c\mathbf{x} + \mathbf{y} \in \operatorname{span}(E)$, and Theorem 1.1.8 makes $\operatorname{span}(E)$ a subspace.
Each $\mathbf{v} \in E$ satisfies $\mathbf{v} = 1\mathbf{v} \in \operatorname{span}(E)$, so $E \subseteq \operatorname{span}(E)$. If $\mathcal{U} \leq \mathcal{V}$ and $E \subseteq \mathcal{U}$, then closure under finite linear combinations gives $\operatorname{span}(E) \subseteq \mathcal{U}$. Thus $\operatorname{span}(E)$ is the smallest subspace containing E.
Structural Commentary: Span closes a set of vectors under finite linear combinations, giving exactly the subspace forced by those vectors.
Key Insight: Nothing beyond linear closure is added. Every containing subspace must already contain span. □

Method 1.1.14. *Subspace Diagnostic*
Let $S \subseteq \mathcal{V}$.

(i) To prove $S \leq \mathcal{V}$, show $S \neq \varnothing$ and $\forall\, \mathbf{u}, \mathbf{v} \in S,\ \forall\, c \in \mathbb{F},\ c\mathbf{u} + \mathbf{v} \in S$.

(ii) To disprove $S \leq \mathcal{V}$, exhibit $\mathbf{u}, \mathbf{v} \in S$ and $c \in \mathbb{F}$ with $c\mathbf{u} + \mathbf{v} \notin S$; $\mathbf{0} \in S$ alone is insufficient.

(iii) The generated subspace is $\langle S \rangle = \operatorname{span}(S) = \left\{ \sum_{i=1}^{n} a_i \mathbf{s}_i : n \in \mathbb{N},\ a_i \in \mathbb{F},\ \mathbf{s}_i \in S \right\}$.

Proposition 1.1.15. *Monotonicity of Span*
Let $\mathcal{U}$ and $\mathcal{W}$ be subsets of $\mathcal{V}$. If $\mathcal{U} \subseteq \mathcal{W}$, then $\operatorname{span}(\mathcal{U}) \subseteq \operatorname{span}(\mathcal{W})$.

Proof. If $\mathcal{U} = \emptyset$, then $\operatorname{span}(\mathcal{U}) = \{\mathbf{0}\} \subseteq \operatorname{span}(\mathcal{W})$.
Suppose $\mathcal{U} \neq \emptyset$. Every vector in $\operatorname{span}(\mathcal{U})$ is a finite linear combination of vectors in $\mathcal{U}$; since $\mathcal{U} \subseteq \mathcal{W}$, the same expression is a finite linear combination of vectors in $\mathcal{W}$. Hence $\operatorname{span}(\mathcal{U}) \subseteq \operatorname{span}(\mathcal{W})$.
Structural Commentary: Enlarging the generating set can only enlarge the collection of linear combinations.
Key Insight: Span is monotone with respect to inclusion of generating sets. □

Definition 1.1.16. *Linear Independence, Bases, and Dimension*

(i) A list $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_m)$ is *linearly independent* if $\sum_{j=1}^{m} a_j \mathbf{v}_j = \mathbf{0} \Rightarrow a_1 = \cdots = a_m = 0$.
Otherwise it is *linearly dependent*. The empty list is linearly independent.

(ii) A list β is a *basis* of $\mathcal{V}$ if β is linearly independent and $\operatorname{span}(\beta) = \mathcal{V}$.

(iii) $\mathcal{V}$ is *finite-dimensional* if some finite list β satisfies $\operatorname{span}(\beta) = \mathcal{V}$; otherwise $\mathcal{V}$ is *infinite-dimensional*.

Proposition 1.1.17. *Unique Coefficients for Independent Lists*
Let $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ be a linearly independent list in $\mathcal{V}$.
Then $a_1\mathbf{v}_1 + \cdots + a_m\mathbf{v}_m = b_1\mathbf{v}_1 + \cdots + b_m\mathbf{v}_m$ iff $a_j = b_j$ for each $j = 1, \ldots, m$.

Proof.
($\Rightarrow$) Subtracting the two sides gives $(a_1 - b_1)\mathbf{v}_1 + \cdots + (a_m - b_m)\mathbf{v}_m = \mathbf{0}$.
Linear independence forces $a_j - b_j = 0$ for each j, hence $a_j = b_j$.
($\Leftarrow$) If $a_j = b_j$ for each j, then the two linear combinations have identical coefficients and are equal.
Structural Commentary: Linear independence makes coefficients in a linear combination unique.
Key Insight: Independent vectors allow no hidden coefficient cancellation. □

Proposition 1.1.18. *Appending to a Dependent List*
Let $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ be a linearly dependent list in $\mathcal{V}$.
Then for each $\mathbf{v} \in \mathcal{V}$, the list $(\mathbf{v}_1, \ldots, \mathbf{v}_m, \mathbf{v})$ is linearly dependent.

Proof. Since $(\mathbf{v}_1, \dots, \mathbf{v}_m)$ is linearly dependent, $\exists$ scalars $a_1, \dots, a_m$, not all zero, with $a_1\mathbf{v}_1 + \cdots + a_m\mathbf{v}_m = \mathbf{0}$. Appending coefficient 0 gives $a_1\mathbf{v}_1 + \cdots + a_m\mathbf{v}_m + 0\mathbf{v} = \mathbf{0}$, and the scalars $a_1, \dots, a_m, 0$ are not all zero.
Structural Commentary: Longer lists inherit dependence; set new coefficients to 0.
Key Insight: Once a list has a redundancy, adding more vectors cannot remove it. □

Definition 1.1.19. *Omitting a Vector from a List*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_m)$ be a list in $\mathcal{V}$ with $m \geq 2$.
For $j \in \{1, \dots, m\}$, the list obtained from β by omitting $\mathbf{v}_j$ is denoted $(\mathbf{v}_1, \dots, \widehat{\mathbf{v}_j}, \dots, \mathbf{v}_m)$.

Proposition 1.1.20. *Omitting from an Independent List*
If $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_m)$ is linearly independent and $m \geq 2$, then $\beta_{\widehat{j}} := (\mathbf{v}_1, \dots, \widehat{\mathbf{v}_j}, \dots, \mathbf{v}_m)$ is linearly independent for each $j \in \{1, \dots, m\}$.

Proof. Fix j. Any relation $\sum_{i \neq j} a_i\mathbf{v}_i = \mathbf{0}$ on $\beta_{\widehat{j}}$ becomes a relation on β after setting $a_j = 0$. Independence of β gives $a_i = 0$ for all i, so the sublist is independent.
Structural Commentary: Deletion cannot create cancellation.
Key Insight: Every sublist of an independent list remains independent. □

Proposition 1.1.21. *Adjoining and Removing Vectors from Lists*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_m)$ be a list in $\mathcal{V}$. Write $(\beta, \mathbf{v}) := (\mathbf{v}_1, \dots, \mathbf{v}_m, \mathbf{v})$ and $\beta_{\widehat{j}} := (\mathbf{v}_1, \dots, \widehat{\mathbf{v}_j}, \dots, \mathbf{v}_m)$.

(i) If β is linearly independent and $\mathbf{v} \notin \operatorname{span}(\beta)$, then $(\beta, \mathbf{v})$ is linearly independent.

(ii) If $\sum_{i=1}^{m} a_i\mathbf{v}_i = \mathbf{0}$ is nontrivial and $a_j \neq 0$, then $\operatorname{span}(\beta_{\widehat{j}}) = \operatorname{span}(\beta)$.
Hence $\operatorname{span}(\beta) = \mathcal{V}$ implies $\operatorname{span}(\beta_{\widehat{j}}) = \mathcal{V}$.

Proof.

(i) Suppose $\sum_{i=1}^{m} a_i\mathbf{v}_i + a\mathbf{v} = \mathbf{0}$. If $a \neq 0$, then $\mathbf{v} = -a^{-1}\sum_{i=1}^{m} a_i\mathbf{v}_i \in \operatorname{span}(\beta)$, a contradiction.
Hence $a = 0$, and independence of β gives $a_1 = \cdots = a_m = 0$.

(ii) Solving the relation for $\mathbf{v}_j$ gives $\mathbf{v}_j = -a_j^{-1}\sum_{i \neq j} a_i\mathbf{v}_i \in \operatorname{span}(\beta_{\widehat{j}})$. Thus every vector of β lies in $\operatorname{span}(\beta_{\widehat{j}})$, so $\operatorname{span}(\beta) \subseteq \operatorname{span}(\beta_{\widehat{j}})$. The reverse inclusion is immediate because $\beta_{\widehat{j}}$ is a sublist of β.

Structural Commentary: New directions preserve independence; redundant vectors preserve span.
Key Insight: Add new directions; remove redundant ones. □

Lemma 1.1.22. *Replacement Lemma*
Let $\mathcal{V} \neq \{\mathbf{0}\}$, let $\beta = (\mathbf{u}_1, \dots, \mathbf{u}_r)$ span $\mathcal{V}$, and let $\mathbf{0} \neq \mathbf{v} = \sum_{i=1}^{r} c_i\mathbf{u}_i$.
For $j \in \{1, \dots, r\}$, set $\beta_{j \leftarrow \mathbf{v}} := (\mathbf{u}_1, \dots, \mathbf{u}_{j-1}, \mathbf{v}, \mathbf{u}_{j+1}, \dots, \mathbf{u}_r)$.

(i) $c_j \neq 0$ for some j.
(ii) If $c_j \neq 0$, then $\operatorname{span}(\beta_{j \leftarrow \mathbf{v}}) = \mathcal{V}$; if β is a basis, then $\beta_{j \leftarrow \mathbf{v}}$ is a basis.
(iii) If $r \geq 2$, β is a basis, and $\mathbf{v} \notin \operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_k)$ for some $1 \leq k < r$, then $\beta_{j \leftarrow \mathbf{v}}$ is a basis for some $j > k$.

Proof.

(i) If all $c_j = 0$, then $\mathbf{v} = \mathbf{0}$.

(ii) Since $\mathbf{u}_j = c_j^{-1}\left(\mathbf{v} - \sum_{i \neq j} c_i\mathbf{u}_i\right)$, every vector of β lies in $\operatorname{span}(\beta_{j \leftarrow \mathbf{v}})$.
Conversely, $\mathbf{v} \in \operatorname{span}(\beta)$, so every vector of $\beta_{j \leftarrow \mathbf{v}}$ lies in $\operatorname{span}(\beta)$. Hence $\operatorname{span}(\beta_{j \leftarrow \mathbf{v}}) = \mathcal{V}$.
If β is a basis and $\sum_{i<j} b_i\mathbf{u}_i + a\mathbf{v} + \sum_{i>j} b_i\mathbf{u}_i = \mathbf{0}$, substitution gives $\sum_{i \neq j}(b_i + ac_i)\mathbf{u}_i + ac_j\mathbf{u}_j = \mathbf{0}$. Independence of β gives $ac_j = 0$ and $b_i + ac_i = 0$ for $i \neq j$.
Since $c_j \neq 0$, $a = 0$, and then every $b_i = 0$. Thus $\beta_{j \leftarrow \mathbf{v}}$ is linearly independent.

(iii) Since $\mathbf{v} \notin \operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_k)$, some coefficient among $c_{k+1}, \dots, c_r$ is nonzero.
Choose $j > k$ with $c_j \neq 0$; by (ii), $\beta_{j \leftarrow \mathbf{v}}$ is a basis.

Structural Commentary: Replacement preserves span.
Key Insight: Replace where the coefficient is nonzero. □

Lemma 1.1.23. *Linear Dependence Lemma*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_m)$ be linearly dependent with $\mathbf{v}_1 \neq \mathbf{0}$, and set $\beta_{\widehat{j}} := (\mathbf{v}_1, \dots, \widehat{\mathbf{v}_j}, \dots, \mathbf{v}_m)$.
Then some $j \in \{2, \dots, m\}$ satisfies $\mathbf{v}_j \in \operatorname{span}(\mathbf{v}_1, \dots, \mathbf{v}_{j-1})$ and $\operatorname{span}(\beta_{\widehat{j}}) = \operatorname{span}(\beta)$.

Proof. Choose a nontrivial relation $\sum_{i=1}^{m} a_i \mathbf{v}_i = \mathbf{0}$ and let $j = \max\{i : a_i \neq 0\}$.
If $j = 1$, then $a_1 \mathbf{v}_1 = \mathbf{0}$ with $a_1 \neq 0$, so $\mathbf{v}_1 = \mathbf{0}$, contrary to hypothesis. Hence $j \geq 2$.
Since $a_i = 0$ for $i > j$, we have $\mathbf{v}_j = -a_j^{-1} \sum_{i=1}^{j-1} a_i \mathbf{v}_i$, so $\mathbf{v}_j \in \operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_{j-1})$.
Substituting this expression for $\mathbf{v}_j$ rewrites every linear combination of β without $\mathbf{v}_j$, so $\operatorname{span}(\beta) \subseteq \operatorname{span}(\beta_{\widehat{j}})$. The reverse inclusion follows because $\beta_{\widehat{j}}$ is a sublist of β. Hence the spans are equal.
Structural Commentary: The last nonzero coefficient gives the removable vector.
Key Insight: Dependence gives a later vector removable without losing span. □

Theorem 1.1.24. *Independent Lists Are No Longer Than Spanning Lists*
If $\mathcal{V}$ is finite-dimensional, $\alpha = (\mathbf{u}_1, \ldots, \mathbf{u}_m)$ is linearly independent, and $\gamma = (\mathbf{w}_1, \ldots, \mathbf{w}_n)$ spans $\mathcal{V}$, then $m \leq n$.

Proof. Suppose $m > n$.
If $n = 0$, then $\mathcal{V} = \operatorname{span}(\gamma) = \{\mathbf{0}\}$, contradicting independence of the nonempty list α; hence $n \geq 1$.
Construct spanning lists γ_k of length n that contain $\mathbf{u}_1, \ldots, \mathbf{u}_k$ for $0 \leq k \leq n$. Set $\gamma_0 = \gamma$. Given γ_k with $k < n$, write $\mathbf{u}_{k+1}$ in $\operatorname{span}(\gamma_k)$. Since $\mathbf{u}_{k+1} \notin \operatorname{span}(\mathbf{u}_1, \ldots, \mathbf{u}_k)$, some coefficient on an entry of γ_k outside $\mathbf{u}_1, \ldots, \mathbf{u}_k$ is nonzero. Replacing that entry by $\mathbf{u}_{k+1}$ preserves span by Lemma 1.1.22(ii); call the result γ_{k+1}.
Then γ_n spans $\mathcal{V}$ and contains $\mathbf{u}_1, \ldots, \mathbf{u}_n$, so $\operatorname{span}(\mathbf{u}_1, \ldots, \mathbf{u}_n) = \mathcal{V}$.
Since $m > n$, this puts $\mathbf{u}_{n+1}$ in the span of earlier vectors, contradicting independence of α.
Structural Commentary: Independent length is bounded by spanning length.
Key Insight: Independent lists cannot be longer than spanning lists in finite dimensions. □

Proposition 1.1.25. *Unique Coordinate Representation*
$\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ is a basis of $\mathcal{V}$ iff each $\mathbf{v} \in \mathcal{V}$ has a unique coefficient vector $(a_1, \ldots, a_n) \in \mathbb{F}^n$ with $\mathbf{v} = \sum_{j=1}^{n} a_j \mathbf{v}_j$.

Proof.

($\Rightarrow$) Spanning gives a representation of each $\mathbf{v} \in \mathcal{V}$. If $\mathbf{v} = \sum_{j=1}^{n} a_j \mathbf{v}_j = \sum_{j=1}^{n} b_j \mathbf{v}_j$, then Proposition 1.1.17 gives $a_j = b_j$ for each j.

($\Leftarrow$) Existence gives $\operatorname{span}(\beta) = \mathcal{V}$. If $\sum_{j=1}^{n} a_j \mathbf{v}_j = \mathbf{0}$, compare with $\mathbf{0} = \sum_{j=1}^{n} 0 \mathbf{v}_j$; uniqueness gives $a_j = 0$ for each j, so β is independent.

Structural Commentary: A basis is exactly existence and uniqueness of coordinates.
Key Insight: Coordinates are basis coefficients. □

Theorem 1.1.26. *Reducing Spanning Lists*
Every spanning list in a finite-dimensional vector space contains a sublist that is a basis.

Proof. Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a spanning list of $\mathcal{V}$. If β is independent, it is a basis. Otherwise remove one vector without changing span: omit $\mathbf{v}_1$ if $\mathbf{v}_1 = \mathbf{0}$, and otherwise use Lemma 1.1.23 to remove a suitable $\mathbf{v}_j$. The resulting sublist still spans $\mathcal{V}$ and has length one less. Repeat this step. The length strictly decreases, so the process stops with a spanning independent sublist, hence a basis.
Structural Commentary: Delete redundant vectors while preserving span.
Key Insight: Spanning gives enough vectors; reduction deletes exactly the unnecessary ones. □

Theorem 1.1.27. *Extending Independent Lists*
Every linearly independent list in a finite-dimensional vector space can be extended to a basis.

Proof. Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_m)$ be linearly independent in $\mathcal{V}$, and choose a spanning list $(\mathbf{w}_1, \ldots, \mathbf{w}_n)$ of $\mathcal{V}$.
Initialize $\mathcal{B} = \beta$. Inspect $\mathbf{w}_1, \ldots, \mathbf{w}_n$ in order, appending $\mathbf{w}_j$ exactly when $\mathbf{w}_j \notin \operatorname{span}(\mathcal{B})$. Each append preserves independence by Proposition 1.1.21.
After the inspection, each $\mathbf{w}_j$ lies in $\operatorname{span}(\mathcal{B})$: it was either already there or was appended. Since the $\mathbf{w}_j$ span $\mathcal{V}$ and $\operatorname{span}(\mathcal{B})$ is a subspace, every vector of $\mathcal{V}$ lies in $\operatorname{span}(\mathcal{B})$.
Thus $\mathcal{B}$ is a spanning independent list, hence a basis.
Structural Commentary: An independent list has no redundancy, but may miss directions. Appending only vectors outside its span fills those directions while preserving independence.
Key Insight: A finite-dimensional independent list becomes a basis once enough new directions are added. □

Corollary 1.1.28. *Bases and Subspaces in Finite Dimension*
If $\mathcal{V}$ is finite-dimensional, then $\mathcal{V}$ has a basis. Moreover, $\mathcal{U} \leq \mathcal{V}$ implies $\mathcal{U}$ is finite-dimensional.

Proof. Choose a finite spanning list γ of $\mathcal{V}$. By Theorem 1.1.26, some sublist of γ is a basis of $\mathcal{V}$.
Let $\mathcal{U} \leq \mathcal{V}$. If $\mathcal{U} = \{\mathbf{0}\}$, then $()$ is a basis. Otherwise start with $\beta = ()$ and append vectors from $\mathcal{U} \setminus \text{span}(\beta)$ while possible. Proposition 1.1.21 preserves independence, while Theorem 1.1.24 bounds the length by the length of γ; hence the process terminates. At termination, $\mathcal{U} \subseteq \text{span}(\beta)$, so β is a finite basis of $\mathcal{U}$.
Structural Commentary: Finite-dimensional spaces have finite coordinate systems, and subspaces cannot contain longer independent lists than the ambient space permits.
Key Insight: Finite-dimensionality passes downward to subspaces. □

Theorem 1.1.29. *Basis Length Is Well-Defined*
Any two bases of a finite-dimensional vector space have the same length.

Proof. Let $\beta_1 = (\mathbf{v}_1, \ldots, \mathbf{v}_m)$ and $\beta_2 = (\mathbf{w}_1, \ldots, \mathbf{w}_n)$ be bases of $\mathcal{V}$. Theorem 1.1.24 applied to β_1 against β_2 gives $m \leq n$, and reversing the roles gives $n \leq m$. Hence $m = n$.
Structural Commentary: Bases are simultaneously independent and spanning, so the comparison theorem traps their lengths in both directions. This makes dimension well-defined.
Key Insight: Dimension is an intrinsic number, not an artefact of a particular basis. □

Definition 1.1.30. *Dimension*

(i) If $\mathcal{V} = \{\mathbf{0}\}$, then $\dim(\mathcal{V}) := 0$.
(ii) If $\mathcal{V}$ has a basis of length $n \geq 1$, then $\mathcal{V}$ is *n-dimensional* and $\dim(\mathcal{V}) := n$.

Theorem 1.1.29 makes $\dim(\mathcal{V})$ well-defined.

Proposition 1.1.31. *Subspace Dimension Criterion*
Let $\mathcal{V}$ be finite-dimensional and let $\mathcal{U} \leq \mathcal{V}$. Then $\dim(\mathcal{U}) \leq \dim(\mathcal{V})$, and equality holds iff $\mathcal{U} = \mathcal{V}$.

Proof. Let $\delta = (\mathbf{u}_1, \ldots, \mathbf{u}_m)$ be a basis of $\mathcal{U}$ and let β be a basis of $\mathcal{V}$, where $n = \dim(\mathcal{V})$.
Since δ is independent in $\mathcal{V}$ and β spans $\mathcal{V}$, Theorem 1.1.24 gives $m \leq n$.

($\Rightarrow$) If $\mathcal{U} = \mathcal{V}$, then any basis of $\mathcal{U}$ is a basis of $\mathcal{V}$, so $\dim(\mathcal{U}) = \dim(\mathcal{V})$.
($\Leftarrow$) If $\dim(\mathcal{U}) = \dim(\mathcal{V})$, then δ is independent in $\mathcal{V}$ of length n. By Theorem 1.1.27, extending δ to a basis of $\mathcal{V}$ adds no vector, since every basis has length n. Hence $\mathcal{U} = \text{span}(\delta) = \mathcal{V}$.

Structural Commentary: Dimension bounds subspaces; equality forces fullness.
Key Insight: Equal dimension subspaces fill the ambient space. □

Proposition 1.1.32. *Full-Dimension Basis Criteria*
Let $\mathcal{V}$ be finite-dimensional and set $n := \dim(\mathcal{V})$. Then:

(i) if γ spans $\mathcal{V}$ and has length n, then γ is a basis;
(ii) if α is linearly independent in $\mathcal{V}$ and has length n, then α is a basis.

Proof.

(*i*) By Theorem 1.1.26, γ contains a basis sublist. Since every basis has length n, no proper sublist can have length n; hence γ itself is a basis.
(*ii*) Extend α to a basis of $\mathcal{V}$ by Theorem 1.1.27. Since every basis has length n and α already has length n, no vector is added; hence α is a basis.

Structural Commentary: Full-length spanning or independent lists cannot be shortened or extended.
Key Insight: In dimension n, n independent or n spanning vectors already form a basis. □

Theorem 1.1.33. *Basis Characterizations*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a finite list in $\mathcal{V}$, and write $\beta^\frown \mathbf{x} = (\mathbf{v}_1, \ldots, \mathbf{v}_n, \mathbf{x})$. Then the following are equivalent:

(i) β is a basis of $\mathcal{V}$;
(ii) each $\mathbf{x} \in \mathcal{V}$ has a unique coefficient vector $(a_1, \ldots, a_n) \in \mathbb{F}^n$ with $\mathbf{x} = \sum_{j=1}^{n} a_j \mathbf{v}_j$;
(iii) $\text{span}(\beta) = \mathcal{V}$, and no proper sublist of β spans $\mathcal{V}$;
(iv) β is linearly independent, and $\beta^\frown \mathbf{x}$ is linearly dependent for each $\mathbf{x} \in \mathcal{V}$.

Proof.

(*i*) $\Rightarrow$ (*ii*) This is Proposition 1.1.25.
(*ii*) $\Rightarrow$ (*iii*) Existence gives $\text{span}(\beta) = \mathcal{V}$.
If a proper sublist $\gamma \subsetneq \beta$ spanned $\mathcal{V}$ and omitted the occurrence $\mathbf{v}_j$, then $\mathbf{v}_j$ would have both its standard β-coordinate representation and one using only γ, contradicting uniqueness.

(*iii*) $\Rightarrow$ (*iv*) If β were dependent, Proposition 1.1.21 would remove some $\mathbf{v}_j$ without changing $\operatorname{span}(\beta)$, contradicting minimality. Thus β is independent.
Since $\operatorname{span}(\beta) = \mathcal{V}$, writing $\mathbf{x}$ as a linear combination of β gives a nontrivial relation on $\beta^\frown\mathbf{x}$.
(*iv*) $\Rightarrow$ (*i*) If $\operatorname{span}(\beta) \neq \mathcal{V}$, choose $\mathbf{x} \notin \operatorname{span}(\beta)$. Then Proposition 1.1.21 makes $\beta^\frown\mathbf{x}$ independent, contradicting (iv). Hence β spans $\mathcal{V}$, and with independence from (iv), it is a basis.

Structural Commentary: A basis is simultaneously just enough to span and too large to extend independently.
Key Insight: Spanning and independence are dual minimality and maximality conditions. □

Definition 1.1.34. *Hamel Basis*
A subset $\mathcal{B} \subseteq \mathcal{V}$ is a *Hamel basis* of $\mathcal{V}$ if it is linearly independent and $\operatorname{span}(\mathcal{B}) = \mathcal{V}$. Equivalently, each $\mathbf{v} \in \mathcal{V}$ has a unique finitely supported coefficient function $a : \mathcal{B} \to \mathbb{F}$ with $\mathbf{v} = \sum_{\mathbf{b}\in\mathcal{B}} a(\mathbf{b})\mathbf{b}$.

Finite ordered bases are orderings of finite Hamel bases.

Theorem 1.1.35. *Hamel Basis Extension*
Let $I \subseteq S \subseteq \mathcal{V}$, with I linearly independent and $\operatorname{span}(S) = \mathcal{V}$.
Assuming Zorn's lemma, there exists a Hamel basis $\mathcal{B}$ of $\mathcal{V}$ such that $I \subseteq \mathcal{B} \subseteq S$. Consequently, every nonzero vector space has a basis; every independent set extends to a basis; every spanning set contains a basis.

Proof. Let $\mathscr{P} = \{E : I \subseteq E \subseteq S,\ E \text{ linearly independent}\}$, ordered by inclusion. Then $I \in \mathscr{P}$.
If $\mathcal{C} \subseteq \mathscr{P}$ is a nonempty chain, put $U = \bigcup_{E\in\mathcal{C}} E$. In any finite relation $\sum_{k=1}^{m} a_k\mathbf{u}_k = \mathbf{0}$ in U, the finitely many vectors $\mathbf{u}_k$ lie in one member of the chain; independence there gives $a_k = 0$ for all k.
Thus $U \in \mathscr{P}$ and is an upper bound for $\mathcal{C}$. The empty chain is bounded by I.
By Zorn's lemma, Principle 6.5.2, choose maximal $\mathcal{B} \in \mathscr{P}$.
If $s \in S \setminus \operatorname{span}(\mathcal{B})$, then $\mathcal{B} \cup \{s\}$ is linearly independent, contradicting maximality. Hence $S \subseteq \operatorname{span}(\mathcal{B})$, so $\operatorname{span}(\mathcal{B}) = \mathcal{V}$ and $\mathcal{B}$ is a Hamel basis. The consequences follow by taking $S = \mathcal{V}$ or $I = \emptyset$.
Structural Commentary: Finite extension uses exchange; arbitrary extension uses maximality.
Key Insight: Maximal independence inside a spanning set forces spanning. □

Theorem 1.1.36. *Dimension as Basis Cardinality*
If $\mathcal{B}$ and $\mathcal{C}$ are Hamel bases of $\mathcal{V}$, then $|\mathcal{B}| = |\mathcal{C}|$. Hence $\dim(\mathcal{V}) := |\mathcal{B}|$ is well-defined for any Hamel basis $\mathcal{B}$.

Proof. If one basis is finite, say $|\mathcal{B}| = n$, then $\mathcal{C}$ cannot contain $n+1$ distinct vectors, since such a finite sublist would be independent while $\mathcal{B}$ spans $\mathcal{V}$, contradicting Theorem 1.1.24. Hence $\mathcal{C}$ is finite, and Theorem 1.1.29 gives $|\mathcal{C}| = n$. Thus assume both bases are infinite. For $\mathbf{c} \in \mathcal{C}$, let $\operatorname{supp}_{\mathcal{B}}(\mathbf{c}) \subseteq \mathcal{B}$ be its finite $\mathcal{B}$-support.
Every $\mathbf{b} \in \mathcal{B}$ appears in some $\operatorname{supp}_{\mathcal{B}}(\mathbf{c})$: otherwise all vectors in $\operatorname{span}(\mathcal{C})$ would have zero $\mathbf{b}$-coordinate, contradicting $\operatorname{span}(\mathcal{C}) = \mathcal{V}$. Hence $\mathcal{B} \subseteq \bigcup_{\mathbf{c}\in\mathcal{C}} \operatorname{supp}_{\mathcal{B}}(\mathbf{c})$. Thus $|\mathcal{B}| \leq \sum_{\mathbf{c}\in\mathcal{C}} |\operatorname{supp}_{\mathcal{B}}(\mathbf{c})| \leq |\mathcal{C}|\,\aleph_0 = |\mathcal{C}|$ by Theorem 6.4.2.
Interchanging $\mathcal{B}$ and $\mathcal{C}$ gives $|\mathcal{C}| \leq |\mathcal{B}|$; Theorem 6.3.1 gives equality.
Structural Commentary: Even in infinite dimensions, a basis size is intrinsic.
Key Insight: Each basis vector is detected inside finite supports relative to the other basis. □

1.2 Direct Sums, Quotients, and Decomposition

Direct sums split spaces into complementary pieces; quotients collapse selected directions. Dimension formulas keep both constructions controlled.

Definition 1.2.1. *Sums of Subspaces*
Let $\mathcal{U}_1, \ldots, \mathcal{U}_m \leq \mathcal{V}$.

(i) The *sum* of $\mathcal{U}_1, \ldots, \mathcal{U}_m$ is $\mathcal{U}_1 + \cdots + \mathcal{U}_m := \{\mathbf{u}_1 + \cdots + \mathbf{u}_m : \mathbf{u}_j \in \mathcal{U}_j\}$.
(ii) For subspaces $\mathcal{U}, \mathcal{W} \leq \mathcal{V}$, $\mathcal{U} + \mathcal{W} = \operatorname{span}(\mathcal{U} \cup \mathcal{W}) = \{\mathbf{u} + \mathbf{w} : \mathbf{u} \in \mathcal{U},\ \mathbf{w} \in \mathcal{W}\}$.
(iii) The sum is *direct* if every vector in it has a unique decomposition as $\mathbf{u}_1 + \cdots + \mathbf{u}_m$ with $\mathbf{u}_j \in \mathcal{U}_j$; in this case it is denoted $\mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_m$.

Proposition 1.2.2. *Sum of Subspaces Is the Smallest Containing Subspace*
Suppose $\mathcal{U}_1, \ldots, \mathcal{U}_m$ are subspaces of $\mathcal{V}$. Then $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ is the smallest subspace of $\mathcal{V}$ containing $\mathcal{U}_1, \ldots, \mathcal{U}_m$.

Proof. Since $\mathbf{0} \in \mathcal{U}_j$ for each j, we have $\mathbf{0} \in \mathcal{U}_1 + \cdots + \mathcal{U}_m$. If $\mathbf{x} = \mathbf{u}_1 + \cdots + \mathbf{u}_m$ and $\mathbf{y} = \mathbf{v}_1 + \cdots + \mathbf{v}_m$ lie in $\mathcal{U}_1 + \cdots + \mathcal{U}_m$, then for any $c \in \mathbb{F}$, $c\mathbf{x} + \mathbf{y} = (c\mathbf{u}_1 + \mathbf{v}_1) + \cdots + (c\mathbf{u}_m + \mathbf{v}_m) \in \mathcal{U}_1 + \cdots + \mathcal{U}_m$, because $c\mathbf{u}_j + \mathbf{v}_j \in \mathcal{U}_j$ for each j. By Theorem 1.1.8, $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ is a subspace of $\mathcal{V}$.
Each $\mathcal{U}_j$ is contained in $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ by taking all other summands to be $\mathbf{0}$. If $\mathcal{W}$ is any subspace of $\mathcal{V}$ containing $\mathcal{U}_1, \ldots, \mathcal{U}_m$, then closure under addition gives $\mathbf{u}_1 + \cdots + \mathbf{u}_m \in \mathcal{W}$ whenever $\mathbf{u}_j \in \mathcal{U}_j$.
Hence $\mathcal{U}_1 + \cdots + \mathcal{U}_m \subseteq \mathcal{W}$, so the sum is the smallest such subspace.

Structural Commentary: Sums of subspaces are obtained by closing several subspaces under finite addition. Minimality follows because every containing subspace is already closed under those sums.
Key Insight: The sum adds no extra structure beyond the linear combinations forced by the given subspaces. □

Perspective 1.2.3. *Sum Is Not Union*
The union $\mathcal{U} \cup \mathcal{W}$ keeps vectors from either subspace; the sum $\mathcal{U} + \mathcal{W}$ adds vectors from both. Unless one subspace contains the other, $\mathcal{U} \cup \mathcal{W}$ is not closed under addition.

Theorem 1.2.4. *Finite Unions of Proper Subspaces*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be over an infinite field $\mathbb{F}$. If $\mathcal{U}_1, \dots, \mathcal{U}_n \subsetneq \mathcal{V}$ are subspaces, then $\mathcal{V} \neq \bigcup_{j=1}^{n} \mathcal{U}_j$.

Proof. Assume for contradiction that $\mathcal{V} = \bigcup_{j=1}^{n} \mathcal{U}_j$. After deleting redundant subspaces and reindexing, assume the cover is minimal, so $\mathcal{U}_i \not\subseteq \bigcup_{j \neq i} \mathcal{U}_j$. If $n = 1$, then $\mathcal{V} = \mathcal{U}_1$, contradicting $\mathcal{U}_1 \subsetneq \mathcal{V}$; hence $n \geq 2$.

Choose $\mathbf{w} \in \mathcal{U}_1 \setminus \bigcup_{j=2}^{n} \mathcal{U}_j$ and $\mathbf{v} \in \mathcal{V} \setminus \mathcal{U}_1$. Since $\mathbf{0} \in \mathcal{U}_j$ for every j, we have $\mathbf{w} \neq \mathbf{0}$.
Let $A = \{\mathbf{v} + r\mathbf{w} : r \in \mathbb{F}\}$. Then $|A| = |\mathbb{F}| = \infty$, $A \cap \mathcal{U}_1 = \emptyset$, and for $j \geq 2$, $|A \cap \mathcal{U}_j| \leq 1$, since two points $\mathbf{v} + r_1\mathbf{w}, \mathbf{v} + r_2\mathbf{w} \in \mathcal{U}_j$ with $r_1 \neq r_2$ imply $\mathbf{w} \in \mathcal{U}_j$.
Hence $\left| A \cap \bigcup_{j=1}^{n} \mathcal{U}_j \right| \leq n - 1 < \infty$, contradicting $A \subseteq \mathcal{V} = \bigcup_{j=1}^{n} \mathcal{U}_j$.
Structural Commentary: Over an infinite field, finite proper subspaces cannot cover a well-chosen affine line.
Key Insight: A single one-parameter translate has too many points for finitely many proper intersections. □

Perspective 1.2.5. *Sums Always Exist; Directness Is Extra*
For subspaces $\mathcal{U}, \mathcal{W} \subseteq \mathcal{V}$, the sum $\mathcal{U} + \mathcal{W}$ is always a subspace by Proposition 1.2.2. Writing $\mathcal{U} \oplus \mathcal{W}$ asserts the extra condition $\mathcal{U} \cap \mathcal{W} = \{\mathbf{0}\}$, equivalently uniqueness of the decomposition $\mathbf{u} + \mathbf{w}$.

Proposition 1.2.6. *Two-Subspace Directness Criterion*
Let $\mathcal{U}$ and $\mathcal{W}$ be subspaces of $\mathcal{V}$. Then $\mathcal{U} + \mathcal{W}$ is direct if and only if $\mathcal{U} \cap \mathcal{W} = \{\mathbf{0}\}$.

Proof.
(⇒) Suppose $\mathcal{U} + \mathcal{W}$ is direct. If $\mathbf{x} \in \mathcal{U} \cap \mathcal{W}$, then $\mathbf{x} = \mathbf{x} + \mathbf{0} = \mathbf{0} + \mathbf{x}$ gives two decompositions of $\mathbf{x}$ with one term in each subspace. Uniqueness gives $\mathbf{x} = \mathbf{0}$, so $\mathcal{U} \cap \mathcal{W} = \{\mathbf{0}\}$.
(⇐) Suppose $\mathcal{U} \cap \mathcal{W} = \{\mathbf{0}\}$. Every vector in $\mathcal{U} + \mathcal{W}$ has at least one representation $\mathbf{u} + \mathbf{w}$. If $\mathbf{u}_1 + \mathbf{w}_1 = \mathbf{u}_2 + \mathbf{w}_2$, with $\mathbf{u}_1, \mathbf{u}_2 \in \mathcal{U}$ and $\mathbf{w}_1, \mathbf{w}_2 \in \mathcal{W}$, then $\mathbf{u}_1 - \mathbf{u}_2 = \mathbf{w}_2 - \mathbf{w}_1 \in \mathcal{U} \cap \mathcal{W}$, hence equals $\mathbf{0}$.
Thus $\mathbf{u}_1 = \mathbf{u}_2$ and $\mathbf{w}_1 = \mathbf{w}_2$, so the sum is direct.
Structural Commentary: Trivial intersection prevents cancellation between summands.
Key Insight: Directness means unique decomposition. □

Definition 1.2.7. *Subspace Poset, Meet, and Join*
(i) The *subspace poset* of $\mathcal{V}$ is $\mathrm{Sub}(\mathcal{V}) := \{\mathcal{U} \subseteq \mathcal{V} : \mathcal{U} \leq \mathcal{V}\}$, ordered by inclusion.
(ii) For $\{\mathcal{U}_i : i \in I\} \subseteq \mathrm{Sub}(\mathcal{V})$, the *meet* is $\bigwedge_{i \in I} \mathcal{U}_i := \bigcap_{i \in I} \mathcal{U}_i$.
(iii) For $\{\mathcal{U}_i : i \in I\} \subseteq \mathrm{Sub}(\mathcal{V})$, the *join* is $\bigvee_{i \in I} \mathcal{U}_i := \sum_{i \in I} \mathcal{U}_i := \mathrm{span}\left(\bigcup_{i \in I} \mathcal{U}_i\right)$.
(iv) The empty meet is $\mathcal{V}$, and the empty join is $\{\mathbf{0}\}$.

Proposition 1.2.8. *Subspace Meet and Join*
For subspaces $\{\mathcal{U}_i : i \in I\}$ of $\mathcal{V}$, ordered by inclusion, $\bigwedge_{i \in I} \mathcal{U}_i = \inf\{\mathcal{U}_i : i \in I\}$; and $\bigvee_{i \in I} \mathcal{U}_i = \sup\{\mathcal{U}_i : i \in I\}$.

Proof. If $I = \emptyset$, conventions give vacuous bounds $\mathcal{V}$ and $\{\mathbf{0}\}$. Assume $I \neq \emptyset$.
The meet $M = \bigcap_{i \in I} \mathcal{U}_i$ is a subspace by Proposition 1.1.11; moreover $M \subseteq \mathcal{U}_i$ for all i, and any subspace $\mathcal{W}$ with $\mathcal{W} \subseteq \mathcal{U}_i$ for all i satisfies $\mathcal{W} \subseteq M$. The join $J = \mathrm{span}\left(\bigcup_{i \in I} \mathcal{U}_i\right)$ is a subspace by Theorem 1.1.13; moreover $\mathcal{U}_i \subseteq J$ for all i, and any subspace $\mathcal{W}$ with $\mathcal{U}_i \subseteq \mathcal{W}$ for all i contains $\bigcup_{i \in I} \mathcal{U}_i$, hence contains J.
Structural Commentary: Meet enforces common constraints; join closes the combined directions under span.
Key Insight: Meet is intersection, while join is span of the union. □

Theorem 1.2.9. *Subspace Lattice*
The subspace poset $\mathrm{Sub}(\mathcal{V})$ *is a complete modular lattice.*

Explicitly, for every $\{\mathcal{U}_i : i \in I\} \subseteq \mathrm{Sub}(\mathcal{V})$, $\bot = \{\mathbf{0}\}$, $\top = \mathcal{V}$, $\bigwedge_{i\in I} \mathcal{U}_i = \bigcap_{i\in I} \mathcal{U}_i$, $\bigvee_{i\in I} \mathcal{U}_i = \sum_{i\in I} \mathcal{U}_i = \mathrm{span}\left(\bigcup_{i\in I} \mathcal{U}_i\right)$.
Moreover, if $\mathcal{U} \subseteq \mathcal{S}$, *then* $\mathcal{S} \cap (\mathcal{T} + \mathcal{U}) = (\mathcal{S} \cap \mathcal{T}) + \mathcal{U}$.

Proof. Proposition 1.2.8 gives arbitrary meets and joins with the displayed formulas; the empty meet and join give $\top = \mathcal{V}$ and $\bot = \{\mathbf{0}\}$. It remains to prove modularity. Assume $\mathcal{U} \subseteq \mathcal{S}$. If $\mathbf{x} \in (\mathcal{S} \cap \mathcal{T}) + \mathcal{U}$, then $\mathbf{x} = \mathbf{y} + \mathbf{u}$ with $\mathbf{y} \in \mathcal{S} \cap \mathcal{T}$ and $\mathbf{u} \in \mathcal{U}$, hence $\mathbf{x} \in \mathcal{S} \cap (\mathcal{T} + \mathcal{U})$. Conversely, if $\mathbf{x} \in \mathcal{S} \cap (\mathcal{T} + \mathcal{U})$, write $\mathbf{x} = \mathbf{t} + \mathbf{u}$ with $\mathbf{t} \in \mathcal{T}$ and $\mathbf{u} \in \mathcal{U}$. Since $\mathbf{x} \in \mathcal{S}$ and $\mathbf{u} \in \mathcal{U} \subseteq \mathcal{S}$, $\mathbf{t} = \mathbf{x} - \mathbf{u} \in \mathcal{S} \cap \mathcal{T}$, so $\mathbf{x} \in (\mathcal{S} \cap \mathcal{T}) + \mathcal{U}$.
Structural Commentary: The subspace lattice is complete and modular, not generally distributive.
Key Insight: If one summand already lies inside $\mathcal{S}$, intersecting with $\mathcal{S}$ can be pushed through the sum. □

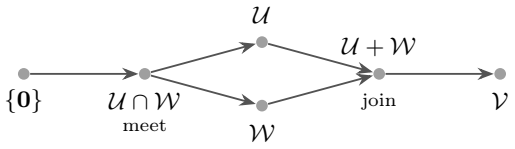


Figure 10: For two subspaces, intersection is the meet and sum is the join under inclusion.

Proposition 1.2.10. *Equivalent Conditions for a Direct Sum*
For subspaces $\mathcal{U}_1, \ldots, \mathcal{U}_m \subseteq \mathcal{V}$, *the following are equivalent:*

(i) $\mathcal{V} = \mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_m$;
(ii) every $\mathbf{v} \in \mathcal{V}$ *has a unique representation* $\mathbf{v} = \mathbf{u}_1 + \cdots + \mathbf{u}_m$ *with* $\mathbf{u}_j \in \mathcal{U}_j$;
(iii) $\mathcal{V} = \mathcal{U}_1 + \cdots + \mathcal{U}_m$, *and if* $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$ *with* $\mathbf{u}_j \in \mathcal{U}_j$, *then* $\mathbf{u}_1 = \cdots = \mathbf{u}_m = \mathbf{0}$.

Proof.

(*i*) ⇒ (*ii*) This follows from the definition of direct sum.
(*ii*) ⇒ (*iii*) Every vector lies in $\mathcal{U}_1 + \cdots + \mathcal{U}_m$, so $\mathcal{V} = \mathcal{U}_1 + \cdots + \mathcal{U}_m$. If $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$, then this representation of $\mathbf{0}$ must equal $\mathbf{0} + \cdots + \mathbf{0}$ by uniqueness, so each $\mathbf{u}_j = \mathbf{0}$.
(*iii*) ⇒ (*i*) Since $\mathcal{V} = \mathcal{U}_1 + \cdots + \mathcal{U}_m$, every $\mathbf{v} \in \mathcal{V}$ has at least one representation. If $\mathbf{v} = \mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{w}_1 + \cdots + \mathbf{w}_m$, then $(\mathbf{u}_1 - \mathbf{w}_1) + \cdots + (\mathbf{u}_m - \mathbf{w}_m) = \mathbf{0}$. By (iii), $\mathbf{u}_j - \mathbf{w}_j = \mathbf{0}$ for every j, so the representation is unique. Hence $\mathcal{V} = \mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_m$.

Structural Commentary: Direct sums require both spanning by the summands and no cancellation among them. The zero-representation condition isolates the no-cancellation part.
Key Insight: Uniqueness is equivalent to saying the only way the summands add to zero is trivially. □

Definition 1.2.11. *Support and Internal Direct Sum of a Family*
Let I be an index set. For each $i \in I$, let $\mathcal{X}_i$ be a vector space with zero vector $\mathbf{0}_i$.

(i) *Support.* For $\mathbf{x} = (\mathbf{x}_i)_{i\in I} \in \prod_{i\in I} \mathcal{X}_i$, define $\mathrm{supp}(\mathbf{x}) := \{i \in I : \mathbf{x}_i \neq \mathbf{0}_i\}$.
(ii) The indexed family $\mathbf{x}$ is *finitely supported* iff $|\mathrm{supp}(\mathbf{x})| < \infty$.
(iii) *Finite sum.* If all $\mathcal{X}_i$ lie in a common vector space and $\mathbf{x}$ is finitely supported, define $\sum_{i\in I} \mathbf{x}_i := \sum_{i\in \mathrm{supp}(\mathbf{x})} \mathbf{x}_i$.
(iv) *Internal direct sum.* For $\{\mathcal{U}_i : i \in I\} \subseteq \mathrm{Sub}(\mathcal{V})$, $\mathcal{V} = \bigoplus_{i\in I} \mathcal{U}_i$ means every $\mathbf{v} \in \mathcal{V}$ has a unique finitely supported $\mathbf{u} = (\mathbf{u}_i)_{i\in I}$ with $\mathbf{u}_i \in \mathcal{U}_i$ and $\mathbf{v} = \sum_{i\in I} \mathbf{u}_i$.

Equivalently, $\mathcal{V} = \sum_{i\in I} \mathcal{U}_i$ and every finite relation $\sum_{i\in F} \mathbf{u}_i = \mathbf{0}$, with $F \subseteq I$ finite and $\mathbf{u}_i \in \mathcal{U}_i$, has $\mathbf{u}_i = \mathbf{0}$ $\forall i \in F$.

Proposition 1.2.12. *Family Direct-Sum Criterion*
For $\{\mathcal{U}_i\}_{i\in I} \subseteq \mathrm{Sub}(\mathcal{V})$, $\sum_{i\in I} \mathcal{U}_i$ *is direct* $\iff \forall i \in I$, $\mathcal{U}_i \cap \sum_{j\in I\setminus\{i\}} \mathcal{U}_j = \{\mathbf{0}\}$.

Proof.

(⇒) If $\mathbf{x} \in \mathcal{U}_i \cap \sum_{j\in I\setminus\{i\}} \mathcal{U}_j$, then $\mathbf{x} = \sum_{j\in F} \mathbf{u}_j$ for some finite $F \subseteq I \setminus \{i\}$ and $\mathbf{u}_j \in \mathcal{U}_j$. Directness applied to $\mathbf{x} - \sum_{j\in F} \mathbf{u}_j = \mathbf{0}$ gives $\mathbf{x} = \mathbf{0}$.

($\Leftarrow$) If $\sum_{k=1}^{m} \mathbf{u}_{i_k} = \mathbf{0}$ with distinct i_k and $\mathbf{u}_{i_k} \in \mathcal{U}_{i_k}$, then for each k, $\mathbf{u}_{i_k} = -\sum_{\ell \neq k} \mathbf{u}_{i_\ell} \in \mathcal{U}_{i_k} \cap \sum_{j \in I \setminus \{i_k\}} \mathcal{U}_j = \{\mathbf{0}\}$.
Thus all finite zero-relations are trivial, so the sum is direct.

Structural Commentary: Family directness requires each summand to miss all others together.
Key Insight: Directness is subspace independence. □

Definition 1.2.13. *Product of Vector Spaces*
Let $\mathcal{V}_1, \dots, \mathcal{V}_m$ be vector spaces over $\mathbb{F}$. Their *product* is $\mathcal{V}_1 \times \cdots \times \mathcal{V}_m := \{(\mathbf{v}_1, \dots, \mathbf{v}_m) : \mathbf{v}_j \in \mathcal{V}_j\}$.
Addition and scalar multiplication are defined componentwise:
$(\mathbf{u}_1, \dots, \mathbf{u}_m) + (\mathbf{v}_1, \dots, \mathbf{v}_m) := (\mathbf{u}_1 + \mathbf{v}_1, \dots, \mathbf{u}_m + \mathbf{v}_m)$ and $a(\mathbf{v}_1, \dots, \mathbf{v}_m) := (a\mathbf{v}_1, \dots, a\mathbf{v}_m)$.

Proposition 1.2.14. *Products Are Vector Spaces*
If $\mathcal{V}_1, \dots, \mathcal{V}_m$ are vector spaces over $\mathbb{F}$, then $\mathcal{V}_1 \times \cdots \times \mathcal{V}_m$ is a vector space under componentwise operations.

Proof. The additive identity is $(\mathbf{0}_1, \dots, \mathbf{0}_m)$, where $\mathbf{0}_j$ is the zero vector of $\mathcal{V}_j$, and additive inverse of $(\mathbf{v}_1, \dots, \mathbf{v}_m)$ is $(-\mathbf{v}_1, \dots, -\mathbf{v}_m)$. All remaining vector space axioms hold componentwise because they hold in each $\mathcal{V}_j$.
Structural Commentary: Product spaces collect several vector spaces into one coordinatewise structure.
Key Insight: Every vector space axiom reduces to the same axiom in each component. □

Theorem 1.2.15. *Dimension of a Product*
If $\mathcal{V}_1, \dots, \mathcal{V}_m$ are finite-dimensional vector spaces over $\mathbb{F}$, then $\mathcal{V}_1 \times \cdots \times \mathcal{V}_m$ is finite-dimensional, and $\dim(\mathcal{V}_1 \times \cdots \times \mathcal{V}_m) = \dim(\mathcal{V}_1) + \cdots + \dim(\mathcal{V}_m)$.

Proof. For each j, let $(\mathbf{v}_{j,1}, \dots, \mathbf{v}_{j,n_j})$ be a basis of $\mathcal{V}_j$. For every j and k, let $\mathbf{e}_{j,k}$ denote the element of $\mathcal{V}_1 \times \cdots \times \mathcal{V}_m$ whose jth component is $\mathbf{v}_{j,k}$ and whose other components are zero.
The list of all $\mathbf{e}_{j,k}$ spans the product because each component can be expanded in its own basis. If a linear combination of the $\mathbf{e}_{j,k}$ is zero, then looking at the jth component gives a linear relation among $(\mathbf{v}_{j,1}, \dots, \mathbf{v}_{j,n_j})$, so all coefficients in that component are zero. Hence the list is linearly independent and therefore is a basis of the product. Its length is $n_1 + \cdots + n_m$.
Structural Commentary: A basis for a product is formed by placing each component basis into its own slot.
Key Insight: Product dimension counts the independent directions contributed by each factor. □

Definition 1.2.16. *External Product and Direct Sum*
Let $\{\mathcal{V}_i : i \in I\}$ be $\mathbb{F}$-vector spaces.

(i) The *external direct product* is $\prod_{i \in I} \mathcal{V}_i = \{\mathbf{v} : I \to \bigcup_{i \in I} \mathcal{V}_i : \mathbf{v}(i) \in \mathcal{V}_i\}$.

(ii) Operations are componentwise: $(\mathbf{u} + \mathbf{v})(i) = \mathbf{u}(i) + \mathbf{v}(i)$ and $(a\mathbf{v})(i) = a\mathbf{v}(i)$.

(iii) The *external direct sum* is $\bigoplus_{i \in I}^{\text{ext}} \mathcal{V}_i = \{\mathbf{v} \in \prod_{i \in I} \mathcal{V}_i : |\operatorname{supp}(\mathbf{v})| < \infty\}$. If $|I| < \infty$, then $\bigoplus_{i \in I}^{\text{ext}} \mathcal{V}_i = \prod_{i \in I} \mathcal{V}_i$.

Proposition 1.2.17. *External Sums Are Finite-Support Products*
With componentwise operations, $\prod_{i \in I} \mathcal{V}_i$ is an $\mathbb{F}$-vector space, $\bigoplus_{i \in I}^{\text{ext}} \mathcal{V}_i = \{\mathbf{v} \in \prod_{i \in I} \mathcal{V}_i : |\operatorname{supp}(\mathbf{v})| < \infty\} \leq \prod_{i \in I} \mathcal{V}_i$.

Proof. Let $P = \prod_{i \in I} \mathcal{V}_i$. For $\mathbf{u}, \mathbf{v} \in P$ and $a, b \in \mathbb{F}$, $(a\mathbf{u} + b\mathbf{v})(i) = a\mathbf{u}(i) + b\mathbf{v}(i) \in \mathcal{V}_i$.
The zero $0_P(i) = 0_i$, inverses $(-\mathbf{v})(i) = -\mathbf{v}(i)$, and all axioms are coordinatewise.
Let $S = \bigoplus_{i \in I}^{\text{ext}} \mathcal{V}_i$. If $\mathbf{u}, \mathbf{v} \in S$ and $a, b \in \mathbb{F}$, then $\operatorname{supp}(a\mathbf{u} + b\mathbf{v}) \subseteq \operatorname{supp}(\mathbf{u}) \cup \operatorname{supp}(\mathbf{v})$, so $a\mathbf{u} + b\mathbf{v} \in S$.
Also $0_P \in S$. Hence $S \leq P$.
Structural Commentary: $\prod_i \mathcal{V}_i$: arbitrary support; $\bigoplus_i^{\text{ext}} \mathcal{V}_i$: finite support.
Key Insight: Finite support is exactly what finite linear combinations preserve. □

Proposition 1.2.18. *Direct Sum Map*
Suppose $\mathcal{U}_1, \dots, \mathcal{U}_m$ are subspaces of $\mathcal{V}$. Define $\Gamma : \mathcal{U}_1 \times \cdots \times \mathcal{U}_m \to \mathcal{U}_1 + \cdots + \mathcal{U}_m$ by $\Gamma(\mathbf{u}_1, \dots, \mathbf{u}_m) = \mathbf{u}_1 + \cdots + \mathbf{u}_m$. Then $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ is a direct sum if and only if Γ is injective.

Proof.

($\Rightarrow$) Suppose $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ is direct. If $\Gamma(\mathbf{u}_1, \dots, \mathbf{u}_m) = \mathbf{0}$, then $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$, so directness gives $\mathbf{u}_j = \mathbf{0}$ for all j. Hence $\ker \Gamma = \{(\mathbf{0}, \dots, \mathbf{0})\}$, and Γ is injective.

($\Leftarrow$) Suppose Γ is injective. If $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$ with $\mathbf{u}_j \in \mathcal{U}_j$, then $\Gamma(\mathbf{u}_1, \dots, \mathbf{u}_m) = \Gamma(\mathbf{0}, \dots, \mathbf{0})$. Injectivity gives $\mathbf{u}_j = \mathbf{0}$ for all j, so every zero decomposition is trivial. Hence the sum is direct.

Structural Commentary: The addition map identifies decompositions precisely through its kernel.
Key Insight: Directness is injectivity of Γ. □

Theorem 1.2.19. *Dimension Formula for a Sum*
If $\mathcal{U}, \mathcal{W}$ are finite-dimensional subspaces of $\mathcal{V}$, then $\dim(\mathcal{U}+\mathcal{W}) = \dim(\mathcal{U}) + \dim(\mathcal{W}) - \dim(\mathcal{U}\cap\mathcal{W})$.

Proof. Choose a basis $(\mathbf{e}_1, \dots, \mathbf{e}_r)$ of $\mathcal{U}\cap\mathcal{W}$, extend it to a basis $(\mathbf{e}_1, \dots, \mathbf{e}_r, \mathbf{u}_1, \dots, \mathbf{u}_p)$ of $\mathcal{U}$, and to a basis $(\mathbf{e}_1, \dots, \mathbf{e}_r, \mathbf{w}_1, \dots, \mathbf{w}_q)$ of $\mathcal{W}$. Then $(\mathbf{e}_1, \dots, \mathbf{e}_r, \mathbf{u}_1, \dots, \mathbf{u}_p, \mathbf{w}_1, \dots, \mathbf{w}_q)$ spans $\mathcal{U}+\mathcal{W}$ because every vector in the sum is the sum of one vector from $\mathcal{U}$ and one vector from $\mathcal{W}$.
To prove linear independence, suppose $a_1\mathbf{e}_1 + \cdots + a_r\mathbf{e}_r + b_1\mathbf{u}_1 + \cdots + b_p\mathbf{u}_p + c_1\mathbf{w}_1 + \cdots + c_q\mathbf{w}_q = \mathbf{0}$. Then $a_1\mathbf{e}_1 + \cdots + a_r\mathbf{e}_r + b_1\mathbf{u}_1 + \cdots + b_p\mathbf{u}_p$ lies in $\mathcal{U}\cap\mathcal{W}$. Since $(\mathbf{e}_1, \dots, \mathbf{e}_r)$ is a basis of $\mathcal{U}\cap\mathcal{W}$ and $(\mathbf{e}_1, \dots, \mathbf{e}_r, \mathbf{u}_1, \dots, \mathbf{u}_p)$ is a basis of $\mathcal{U}$, we get $b_1 = \cdots = b_p = 0$. The remaining relation lies in the basis $(\mathbf{e}_1, \dots, \mathbf{e}_r, \mathbf{w}_1, \dots, \mathbf{w}_q)$ of $\mathcal{W}$, so $a_1 = \cdots = a_r = c_1 = \cdots = c_q = 0$. Thus the displayed list is a basis of $\mathcal{U}+\mathcal{W}$.
Therefore $\dim(\mathcal{U}+\mathcal{W}) = r + p + q = (r+p) + (r+q) - r$.
Structural Commentary: The formula corrects double-counting: vectors in $\mathcal{U}\cap\mathcal{W}$ are counted once in each subspace. Extending an intersection basis separates shared directions from new directions.
Key Insight: The dimension of a sum is the sum of dimensions minus the overlap. □

Corollary 1.2.20. *Dimension of a Direct Sum*
If $\mathcal{U}$ and $\mathcal{W}$ are finite-dimensional subspaces of $\mathcal{V}$ with $\mathcal{U}\cap\mathcal{W} = \{\mathbf{0}\}$, then $\dim(\mathcal{U}\oplus\mathcal{W}) = \dim(\mathcal{U}) + \dim(\mathcal{W})$.

Proof. Since $\mathcal{U}\cap\mathcal{W} = \{\mathbf{0}\}$, the sum is direct and $\dim(\mathcal{U}\cap\mathcal{W}) = 0$. The result follows from Theorem 1.2.19.
Structural Commentary: When the overlap is trivial, there is no double-counting in the dimension formula. Direct sums are exactly the case where the overlap term vanishes.
Key Insight: Dimensions add exactly for direct sums. □

Theorem 1.2.21. *Direct Sums and Dimension*
Suppose $\mathcal{V}$ is finite-dimensional and $\mathcal{U}_1, \dots, \mathcal{U}_m$ are subspaces of $\mathcal{V}$. Then $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ is a direct sum if and only if $\dim(\mathcal{U}_1 + \cdots + \mathcal{U}_m) = \dim(\mathcal{U}_1) + \cdots + \dim(\mathcal{U}_m)$.

Proof. For each j, choose a basis β_j of $\mathcal{U}_j$, and concatenate these bases into a list β.
This list spans $\mathcal{U}_1 + \cdots + \mathcal{U}_m$ and has length $\dim(\mathcal{U}_1) + \cdots + \dim(\mathcal{U}_m)$.

($\Rightarrow$) If the sum is direct, any linear relation among the vectors in β gives a zero decomposition $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$ with $\mathbf{u}_j \in \mathcal{U}_j$. Directness gives $\mathbf{u}_j = \mathbf{0}$ for every j, and independence of each β_j gives all coefficients zero. Thus β is a basis of the sum, and the dimensions add.

($\Leftarrow$) If the dimensions add, then β is a spanning list whose length equals the dimension of the sum, so Proposition 1.1.32 makes β a basis. A zero decomposition $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$ expands to a zero linear relation in this basis, so every $\mathbf{u}_j = \mathbf{0}$. Hence the sum is direct.

Structural Commentary: Direct sums add dimensions without hidden overlap.
Key Insight: Equality of dimensions means the concatenated bases have no linear redundancy. □

Definition 1.2.22. *Basis Adapted to a Direct Sum*
Suppose $\mathcal{V} = \mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_m$ and each β_j is a basis of $\mathcal{U}_j$. The concatenated basis $\beta = (\beta_1, \dots, \beta_m)$ of $\mathcal{V}$ is called a *basis adapted to the direct-sum decomposition*.
Equivalently, the vectors in an adapted basis are grouped by the summands whose direct sum is $\mathcal{V}$.

Corollary 1.2.23. *Large Subspaces Must Intersect*
Let $\mathcal{U}$ and $\mathcal{W}$ be subspaces of a finite-dimensional vector space $\mathcal{V}$, and let k be a positive integer.

(i) If $\dim(\mathcal{U}) + \dim(\mathcal{W}) > \dim(\mathcal{V})$, then $\mathcal{U}\cap\mathcal{W}$ contains a nonzero vector.
(ii) If $\dim(\mathcal{U}) + \dim(\mathcal{W}) \geq \dim(\mathcal{V}) + k$, then $\mathcal{U}\cap\mathcal{W}$ contains k linearly independent vectors.

Proof. Since $\mathcal{U}+\mathcal{W} \subseteq \mathcal{V}$, Theorem 1.2.19 gives
$\dim(\mathcal{U}\cap\mathcal{W}) = \dim(\mathcal{U}) + \dim(\mathcal{W}) - \dim(\mathcal{U}+\mathcal{W}) \geq \dim(\mathcal{U}) + \dim(\mathcal{W}) - \dim(\mathcal{V})$.
If $\dim(\mathcal{U}) + \dim(\mathcal{W}) > \dim(\mathcal{V})$, the right-hand side is at least 1, so the intersection contains a nonzero vector.
If $\dim(\mathcal{U}) + \dim(\mathcal{W}) \geq \dim(\mathcal{V}) + k$, then $\dim(\mathcal{U}\cap\mathcal{W}) \geq k$, so intersection contains k linearly independent vectors.
Structural Commentary: Excess dimension appears as overlap.
Key Insight: Two large subspaces cannot fit inside $\mathcal{V}$ without sharing directions. □

Proposition 1.2.24. *Complements Exist in Finite Dimensions*
If $\mathcal{V}$ is finite-dimensional and $\mathcal{U} \subseteq \mathcal{V}$ is a subspace, then there exists a subspace $\mathcal{W}$ such that $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$.

Proof. Choose a basis $(\mathbf{u}_1, \dots, \mathbf{u}_m)$ of $\mathcal{U}$ and extend it to a basis $(\mathbf{u}_1, \dots, \mathbf{u}_m, \mathbf{w}_1, \dots, \mathbf{w}_n)$ of $\mathcal{V}$.
Let $\mathcal{W} = \operatorname{span}(\mathbf{w}_1, \dots, \mathbf{w}_n)$. Every vector in $\mathcal{V}$ is a linear combination of this extended basis, so $\mathcal{V} = \mathcal{U} + \mathcal{W}$. If $\mathbf{v} \in \mathcal{U}\cap\mathcal{W}$, then $\mathbf{v}$ has one representation using only $\mathbf{u}_1, \dots, \mathbf{u}_m$ and another using only $\mathbf{w}_1, \dots, \mathbf{w}_n$.
Uniqueness of coefficients in the basis of $\mathcal{V}$ forces $\mathbf{v} = \mathbf{0}$. Hence $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$.
Structural Commentary: Complements split a space into a subspace and the remaining basis directions.
Key Insight: Extend a basis of $\mathcal{U}$ to split $\mathcal{V}$. □

Theorem 1.2.25. *Complements Exist*
Assuming Zorn's lemma, for every $\mathcal{U} \leq \mathcal{V}$, $\exists\, \mathcal{W} \leq \mathcal{V}$ such that $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$.

Proof. By Theorem 1.1.35, choose a Hamel basis $\mathcal{B}_\mathcal{U}$ of $\mathcal{U}$ and a Hamel basis $\mathcal{B} \supseteq \mathcal{B}_\mathcal{U}$ of $\mathcal{V}$.
Put $\mathcal{C} = \mathcal{B} \setminus \mathcal{B}_\mathcal{U}$ and $\mathcal{W} = \operatorname{span}(\mathcal{C})$. For $\mathbf{v} \in \mathcal{V}$, its finite $\mathcal{B}$-expansion splits as $\mathbf{v} = \sum_{b \in F_\mathcal{U}} a_b b + \sum_{c \in F_\mathcal{C}} a_c c \in \mathcal{U} + \mathcal{W}$, with finite $F_\mathcal{U} \subseteq \mathcal{B}_\mathcal{U}$ and $F_\mathcal{C} \subseteq \mathcal{C}$. If $\mathbf{x} \in \mathcal{U} \cap \mathcal{W}$, then its $\mathcal{B}$-coordinates are supported in both $\mathcal{B}_\mathcal{U}$ and $\mathcal{C}$; uniqueness gives $\mathbf{x} = \mathbf{0}$. Hence $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$.
Structural Commentary: Extend a basis of $\mathcal{U}$; span the new directions.
Key Insight: The uniqueness of basis expansion becomes the uniqueness of direct-sum decomposition. □

Definition 1.2.26. *Affine Subset and Parallel Subspaces*
Let $\mathcal{U}$ be a subspace of $\mathcal{V}$. For $\mathbf{v} \in \mathcal{V}$, the subset $\mathbf{v} + \mathcal{U} := \{\mathbf{v} + \mathbf{u} : \mathbf{u} \in \mathcal{U}\}$ is an *affine subset* of $\mathcal{V}$ parallel to $\mathcal{U}$.

Proposition 1.2.27. *Parallel Affine Subsets Are Equal or Disjoint*
Let $\mathcal{U}$ be a subspace of $\mathcal{V}$ and let $\mathbf{v}, \mathbf{w} \in \mathcal{V}$. The following are equivalent:

(i) $\mathbf{v} - \mathbf{w} \in \mathcal{U}$;
(ii) $\mathbf{v} + \mathcal{U} = \mathbf{w} + \mathcal{U}$;
(iii) $(\mathbf{v} + \mathcal{U}) \cap (\mathbf{w} + \mathcal{U}) \neq \emptyset$.

Proof.

(i) $\Rightarrow$ *(ii)* If $\mathbf{v} - \mathbf{w} \in \mathcal{U}$ and $\mathbf{u} \in \mathcal{U}$, then $\mathbf{v} + \mathbf{u} = \mathbf{w} + ((\mathbf{v} - \mathbf{w}) + \mathbf{u}) \in \mathbf{w} + \mathcal{U}$. Thus $\mathbf{v} + \mathcal{U} \subseteq \mathbf{w} + \mathcal{U}$, and the reverse inclusion is similar.
(ii) $\Rightarrow$ *(iii)* Equal affine subsets have the same elements.
(iii) $\Rightarrow$ *(i)* If $\mathbf{x} \in (\mathbf{v} + \mathcal{U}) \cap (\mathbf{w} + \mathcal{U})$, then $\mathbf{x} = \mathbf{v} + \mathbf{u}_1 = \mathbf{w} + \mathbf{u}_2$ for some $\mathbf{u}_1, \mathbf{u}_2 \in \mathcal{U}$. Hence $\mathbf{v} - \mathbf{w} = \mathbf{u}_2 - \mathbf{u}_1 \in \mathcal{U}$.

Structural Commentary: Parallel affine subsets either coincide or do not meet.
Key Insight: Intersection is controlled by the difference of base points. □

Definition 1.2.28. *Quotient Space*
Let $\mathcal{U}$ be a subspace of $\mathcal{V}$.
The *quotient space* $\mathcal{V}/\mathcal{U}$ is the set of all affine subsets of $\mathcal{V}$ parallel to $\mathcal{U}$: $\mathcal{V}/\mathcal{U} := \{\mathbf{v} + \mathcal{U} : \mathbf{v} \in \mathcal{V}\}$.

Perspective 1.2.29. *Cosets as Equivalence Classes*
The quotient relation is $\mathbf{v} \sim \mathbf{w}$ when $\mathbf{v} - \mathbf{w} \in \mathcal{U}$. Proposition 1.2.27 shows that the equivalence class of $\mathbf{v}$ is exactly the coset $\mathbf{v} + \mathcal{U}$. Thus the quotient space remembers vectors only up to directions in $\mathcal{U}$.

Proposition 1.2.30. *Quotient Relation Is a Linear Congruence*
Let $\mathcal{U} \leq \mathcal{V}$ and define $\mathbf{v} \equiv_\mathcal{U} \mathbf{w} \iff \mathbf{v} - \mathbf{w} \in \mathcal{U}$.
Then $\equiv_\mathcal{U}$ is an equivalence relation and, for $a, b \in \mathbb{F}$, $\mathbf{v}_1 \equiv_\mathcal{U} \mathbf{v}_2$, $\mathbf{w}_1 \equiv_\mathcal{U} \mathbf{w}_2 \implies a\mathbf{v}_1 + b\mathbf{w}_1 \equiv_\mathcal{U} a\mathbf{v}_2 + b\mathbf{w}_2$.

Proof. Since $\mathcal{U} \leq \mathcal{V}$, $\mathbf{0} \in \mathcal{U}$, $-\mathcal{U} = \mathcal{U}$, and $\mathcal{U} + \mathcal{U} \subseteq \mathcal{U}$; hence $\equiv_\mathcal{U}$ is reflexive, symmetric, and transitive.
If $\mathbf{v}_1 \equiv_\mathcal{U} \mathbf{v}_2$ and $\mathbf{w}_1 \equiv_\mathcal{U} \mathbf{w}_2$, then $(a\mathbf{v}_1 + b\mathbf{w}_1) - (a\mathbf{v}_2 + b\mathbf{w}_2) = a(\mathbf{v}_1 - \mathbf{v}_2) + b(\mathbf{w}_1 - \mathbf{w}_2) \in \mathcal{U}$.
Structural Commentary: Subspace congruence makes quotient operations well-defined.
Key Insight: Subspace closure is exactly the well-definedness mechanism for quotient operations. □

Proposition 1.2.31. *Structure of the Quotient*
Let $\mathcal{U}$ be a subspace of $\mathcal{V}$. Addition and scalar multiplication on $\mathcal{V}/\mathcal{U}$ defined by $(\mathbf{v} + \mathcal{U}) + (\mathbf{w} + \mathcal{U}) := (\mathbf{v} + \mathbf{w}) + \mathcal{U}$ and $a(\mathbf{v} + \mathcal{U}) := a\mathbf{v} + \mathcal{U}$ are well-defined, and with these operations $\mathcal{V}/\mathcal{U}$ is a vector space.

Proof. If $\mathbf{v} + \mathcal{U} = \mathbf{v}' + \mathcal{U}$ and $\mathbf{w} + \mathcal{U} = \mathbf{w}' + \mathcal{U}$, then $\mathbf{v} - \mathbf{v}' \in \mathcal{U}$ and $\mathbf{w} - \mathbf{w}' \in \mathcal{U}$.
Since $\mathcal{U}$ is a subspace, $(\mathbf{v} + \mathbf{w}) - (\mathbf{v}' + \mathbf{w}') \in \mathcal{U}$, and $a\mathbf{v} - a\mathbf{v}' \in \mathcal{U}$ so the operations do not depend on the chosen representatives. The vector space axioms are inherited from $\mathcal{V}$ under these operations.
Structural Commentary: Quotients preserve operations modulo $\mathcal{U}$.
Key Insight: Cosets identify vectors differing by elements of $\mathcal{U}$. □

Theorem 1.2.32. *Dimension of a Quotient*
If $\mathcal{V}$ is finite-dimensional and $\mathcal{U}$ is a subspace, then $\dim(\mathcal{V}/\mathcal{U}) = \dim(\mathcal{V}) - \dim(\mathcal{U})$.

Proof. Choose a basis $(\mathbf{u}_1, \ldots, \mathbf{u}_m)$ of $\mathcal{U}$ and extend it to a basis $(\mathbf{u}_1, \ldots, \mathbf{u}_m, \mathbf{w}_1, \ldots, \mathbf{w}_k)$ of $\mathcal{V}$.
The cosets $\mathbf{w}_1 + \mathcal{U}, \ldots, \mathbf{w}_k + \mathcal{U}$ span $\mathcal{V}/\mathcal{U}$ because the $\mathbf{u}_j$-terms vanish in the quotient.
If $a_1(\mathbf{w}_1 + \mathcal{U}) + \cdots + a_k(\mathbf{w}_k + \mathcal{U}) = \mathcal{U}$, then $a_1\mathbf{w}_1 + \cdots + a_k\mathbf{w}_k \in \mathcal{U}$, so a linear relation in the extended basis forces $a_1 = \cdots = a_k = 0$. Hence these cosets form a basis of $\mathcal{V}/\mathcal{U}$, and its length is $k = \dim(\mathcal{V}) - \dim(\mathcal{U})$.
Structural Commentary: Quotiening removes the directions in $\mathcal{U}$ and keeps the added basis directions.
Key Insight: Passing to $\mathcal{V}/\mathcal{U}$ subtracts the collapsed dimension. □

This block develops two complementary ways to organise a vector space: build from smaller pieces by spans and direct sums, or collapse a chosen subspace by passing to a quotient. Bases and dimension measure exactly what is added, split off, or made invisible.

1.3 Linear Maps and Representations

Linear maps preserve vector-space structure. Kernels, ranges, rank-nullity, matrices, change of basis, and similarity describe them without making coordinates primary.

Definition 1.3.1. *Linear Map*
Let $\mathcal{V}$ and $\mathcal{W}$ be vector spaces over the same field $\mathbb{F}$.
A function $T : \mathcal{V} \to \mathcal{W}$ is *linear* if $T(a\mathbf{u} + b\mathbf{v}) = aT\mathbf{u} + bT\mathbf{v}$ for all $\mathbf{u}, \mathbf{v} \in \mathcal{V}$ and $a, b \in \mathbb{F}$.
The set of all linear maps from $\mathcal{V}$ to $\mathcal{W}$ is denoted $\mathcal{L}(\mathcal{V}, \mathcal{W})$. If $\mathcal{V} = \mathcal{W}$, elements of $\mathcal{L}(\mathcal{V})$ are *operators* on $\mathcal{V}$.

Definition 1.3.2. *Zero and Identity Maps*
For vector spaces $\mathcal{V}$ and $\mathcal{W}$, the *zero map* from $\mathcal{V}$ to $\mathcal{W}$ is $0_{\mathcal{V},\mathcal{W}} : \mathcal{V} \to \mathcal{W}$, where $0_{\mathcal{V},\mathcal{W}}\mathbf{v} = \mathbf{0}_{\mathcal{W}}$ for all $\mathbf{v} \in \mathcal{V}$.
The *identity map* on $\mathcal{V}$ is $I_{\mathcal{V}} : \mathcal{V} \to \mathcal{V}$, where $I_{\mathcal{V}}\mathbf{v} = \mathbf{v}$ for all $\mathbf{v} \in \mathcal{V}$.

Proposition 1.3.3. *Linear Maps Send Zero to Zero*
If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then $T\mathbf{0}_{\mathcal{V}} = \mathbf{0}_{\mathcal{W}}$.

Proof. Linearity gives $T\mathbf{0}_{\mathcal{V}} = T(\mathbf{0}_{\mathcal{V}} + \mathbf{0}_{\mathcal{V}}) = T\mathbf{0}_{\mathcal{V}} + T\mathbf{0}_{\mathcal{V}}$. Adding $-T\mathbf{0}_{\mathcal{V}}$ to both sides gives $T\mathbf{0}_{\mathcal{V}} = \mathbf{0}_{\mathcal{W}}$.
Structural Commentary: Linearity preserves zero.
Key Insight: A linear map cannot move zero away from zero. □

Definition 1.3.4. *Operations on Linear Maps*
For $S, T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ and $a \in \mathbb{F}$, define pointwise $(S + T)\mathbf{v} := S\mathbf{v} + T\mathbf{v}$ and $(aT)\mathbf{v} := a(T\mathbf{v})$ for every $\mathbf{v} \in \mathcal{V}$.

Proposition 1.3.5. *Linear Maps Form a Vector Space*
With the operations above, $\mathcal{L}(\mathcal{V}, \mathcal{W})$ is a vector space over $\mathbb{F}$.

Proof. If $S, T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ and $a, c \in \mathbb{F}$, then for $\mathbf{u}, \mathbf{v} \in \mathcal{V}$,

$$(aS + T)(c\mathbf{u} + \mathbf{v}) = a(cS\mathbf{u} + S\mathbf{v}) + (cT\mathbf{u} + T\mathbf{v}) = c(aS + T)\mathbf{u} + (aS + T)\mathbf{v}.$$

Thus $\mathcal{L}(\mathcal{V}, \mathcal{W})$ is closed under linear combinations. The zero map is the additive identity, $-T$ is the additive inverse of T, and the remaining vector-space axioms hold pointwise because they hold in $\mathcal{W}$.
Structural Commentary: Linear map operations are pointwise.
Key Insight: Pointwise sums and scalar multiples preserve linearity. □

Definition 1.3.6. *Product of Linear Maps*
For $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ and $S \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, the *product* $ST \in \mathcal{L}(\mathcal{U}, \mathcal{W})$ is defined by $(ST)\mathbf{u} := S(T\mathbf{u})$ for all $\mathbf{u} \in \mathcal{U}$.

Proposition 1.3.7. *Algebra of Linear Map Products*
If $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ and $S \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then $ST \in \mathcal{L}(\mathcal{U}, \mathcal{W})$. Whenever defined, $(RS)T = R(ST)$, $TI_{\mathcal{U}} = T = I_{\mathcal{V}}T$, $(S_1 + S_2)T = S_1T + S_2T$, and $S(T_1 + T_2) = ST_1 + ST_2$.

Proof. For $\mathbf{u}, \mathbf{x} \in \mathcal{U}$ and $a \in \mathbb{F}$, $(ST)(a\mathbf{u} + \mathbf{x}) = S(aT\mathbf{u} + T\mathbf{x}) = aS(T\mathbf{u}) + S(T\mathbf{x}) = a(ST)\mathbf{u} + (ST)\mathbf{x}$, hence $ST \in \mathcal{L}(\mathcal{U}, \mathcal{W})$. For each remaining identity $A = B$ and admissible $\mathbf{z}$, $A\mathbf{z} = B\mathbf{z}$ by associativity of evaluation, identity maps, or distributivity in the target space; hence $A = B$.
Structural Commentary: Composition chains linear actions, while addition distributes through evaluation.
Key Insight: Linear map algebra is function composition plus pointwise vector-space algebra. □

Definition 1.3.8. *Graph of a Function*
For a function $T : \mathcal{V} \to \mathcal{W}$, the *graph* of T is $\text{graph}(T) := \{(\mathbf{v}, T\mathbf{v}) : \mathbf{v} \in \mathcal{V}\} \subseteq \mathcal{V} \times \mathcal{W}$.

Proposition 1.3.9. *Linearity and Graphs*
A function $T : \mathcal{V} \to \mathcal{W}$ is linear if and only if $\text{graph}(T)$ *is a subspace of $\mathcal{V} \times \mathcal{W}$.*

Proof.

($\Rightarrow$) If T is linear, then $(\mathbf{0}, \mathbf{0}) \in \text{graph}(T)$.
If $(\mathbf{u}, T\mathbf{u}), (\mathbf{v}, T\mathbf{v}) \in \text{graph}(T)$ and $a \in \mathbb{F}$, then $a(\mathbf{u}, T\mathbf{u}) + (\mathbf{v}, T\mathbf{v}) = (a\mathbf{u} + \mathbf{v}, T(a\mathbf{u} + \mathbf{v})) \in \text{graph}(T)$.
Thus $\text{graph}(T)$ is a subspace.

($\Leftarrow$) If $\text{graph}(T)$ is a subspace, then for $\mathbf{u}, \mathbf{v} \in \mathcal{V}$ and $a \in \mathbb{F}$, the vector $a(\mathbf{u}, T\mathbf{u}) + (\mathbf{v}, T\mathbf{v}) = (a\mathbf{u} + \mathbf{v}, aT\mathbf{u} + T\mathbf{v})$ lies in $\text{graph}(T)$. Hence $T(a\mathbf{u} + \mathbf{v}) = aT\mathbf{u} + T\mathbf{v}$, so T is linear.

Structural Commentary: A graph records a function as a subset of a product space.
Key Insight: Closure of the graph is exactly linearity of the function. □

Principle 1.3.10. *Coordinate-Free Principle*
A linear map is completely determined by what it does to a basis. Coordinates do not create the map; they only encode it once bases have been chosen.

Proposition 1.3.11. *Linear Maps Are Determined by a Basis*
If $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ *is a basis of* $\mathcal{V}$ *and* $\mathbf{w}_1, \dots, \mathbf{w}_n \in \mathcal{W}$ *are arbitrary, then* $\exists$ *a unique linear map* $T : \mathcal{V} \to \mathcal{W}$ *such that* $T(\mathbf{v}_j) = \mathbf{w}_j$ *for* $j = 1, \dots, n$.

Proof. Every $\mathbf{v} \in \mathcal{V}$ has a unique expression $\mathbf{v} = a_1\mathbf{v}_1 + \cdots + a_n\mathbf{v}_n$. Define $T(\mathbf{v}) := a_1\mathbf{w}_1 + \cdots + a_n\mathbf{w}_n$. This is well-defined because basis coordinates are unique, and linearity follows from coordinate expansions. Any linear map with the prescribed basis values must agree with this formula, so it is unique.
Structural Commentary: A basis reduces a linear map to finitely many choices.
Key Insight: Once basis images are fixed, linearity determines every other value. □

Definition 1.3.12. *Kernel and Range*
For $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, define $\ker(T) := \{\mathbf{v} \in \mathcal{V} : T\mathbf{v} = \mathbf{0}_{\mathcal{W}}\}$ and $\operatorname{range}(T) := \{T\mathbf{v} : \mathbf{v} \in \mathcal{V}\}$.
The kernel is also called the *nullspace* of T.

Definition 1.3.13. *Injective and Surjective Linear Maps*
A function $T : \mathcal{V} \to \mathcal{W}$ is *injective* if $T\mathbf{u} = T\mathbf{v}$ implies $\mathbf{u} = \mathbf{v}$. It is *surjective* if $\operatorname{range}(T) = \mathcal{W}$.

Proposition 1.3.14. *Kernel and Range Are Subspaces*
If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, *then* $\ker(T)$ *is a subspace of* $\mathcal{V}$ *and* $\operatorname{range}(T)$ *is a subspace of* $\mathcal{W}$.

Proof. Since $T\mathbf{0}_{\mathcal{V}} = \mathbf{0}_{\mathcal{W}}$, the kernel is nonempty.
If $\mathbf{u}, \mathbf{v} \in \ker(T)$ and $a \in \mathbb{F}$, then $T(a\mathbf{u} + \mathbf{v}) = aT\mathbf{u} + T\mathbf{v} = \mathbf{0}_{\mathcal{W}}$, so $a\mathbf{u} + \mathbf{v} \in \ker(T)$.
For the range, if $T\mathbf{u}, T\mathbf{v} \in \operatorname{range}(T)$ and $a \in \mathbb{F}$, then $aT\mathbf{u} + T\mathbf{v} = T(a\mathbf{u} + \mathbf{v}) \in \operatorname{range}(T)$.
Structural Commentary: The kernel and range are the two canonical subspaces attached to a linear map.
Key Insight: Linearity makes both preimages of zero and images of vectors closed under linear combinations. □

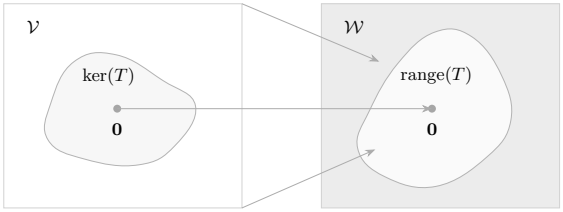


Figure 11: For a linear map $T : \mathcal{V} \to \mathcal{W}$, the range of T may be a proper subspace of $\mathcal{W}$, and the kernel may contain nonzero vectors. The conditions $\operatorname{range}(T) = \mathcal{W}$ and $\ker(T) = \{\mathbf{0}\}$ characterize surjectivity and injectivity, respectively.

Definition 1.3.15. *Projection Along a Complement*
An operator $P \in \mathcal{L}(\mathcal{V})$ is an *algebraic projection*, or *idempotent operator*, if $P^2 = P$.
If $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$, the *projection onto* $\mathcal{U}$ *along* $\mathcal{W}$ is $\rho_{\mathcal{U},\mathcal{W}} \in \mathcal{L}(\mathcal{V})$ where $\rho_{\mathcal{U},\mathcal{W}}(\mathbf{u} + \mathbf{w}) := \mathbf{u}$ for $\mathbf{u} \in \mathcal{U}$, $\mathbf{w} \in \mathcal{W}$.

Proposition 1.3.16. *Algebraic Projections and Direct-Sum Splittings*
For $P \in \mathcal{L}(\mathcal{V})$, P *is an algebraic projection if and only if* $\mathcal{V} = \operatorname{range}(P) \oplus \ker(P)$ *and* $P = \rho_{\operatorname{range}(P),\ker(P)}$.
In particular, if $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$ *and* $\rho = \rho_{\mathcal{U},\mathcal{W}}$, *then* ρ *is an algebraic projection,* $\operatorname{range}(\rho) = \mathcal{U}$, *and* $\ker(\rho) = \mathcal{W}$.

Proof.

(⇒) Suppose P is an algebraic projection, so $P^2 = P$. For every $\mathbf{v} \in \mathcal{V}$, $\mathbf{v} = P\mathbf{v} + (\mathbf{v} - P\mathbf{v})$, with $P\mathbf{v} \in \operatorname{range}(P)$ and $\mathbf{v} - P\mathbf{v} \in \ker(P)$. If $\mathbf{x} \in \operatorname{range}(P) \cap \ker(P)$, then $\mathbf{x} = P\mathbf{y}$ and $\mathbf{0} = P\mathbf{x} = P^2\mathbf{y} = \mathbf{x}$. Thus the sum is direct. Also P fixes $\operatorname{range}(P)$ and kills $\ker(P)$, so $P = \rho_{\operatorname{range}(P),\ker(P)}$.

(⇐) Assume the right-hand side. For $\mathbf{v} = \mathbf{x} + \mathbf{y}$ with $\mathbf{x} \in \operatorname{range}(P)$ and $\mathbf{y} \in \ker(P)$, $P^2\mathbf{v} = P\mathbf{x} = \mathbf{x} = P\mathbf{v}$. Thus P is an algebraic projection. For $\rho_{\mathcal{U},\mathcal{W}}$, the formula $\rho_{\mathcal{U},\mathcal{W}}(\mathbf{u} + \mathbf{w}) = \mathbf{u}$ gives range $\mathcal{U}$ and kernel $\mathcal{W}$, and the equivalence above shows that ρ is an algebraic projection.

Structural Commentary: Idempotents keep one summand and kill its complement.
Key Insight: Direct-sum coordinates turn a projection into $I \oplus 0$. □

Proposition 1.3.17. *Complementary Algebraic Projections*
If $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$, *then* $\rho_{\mathcal{U},\mathcal{W}} + \rho_{\mathcal{W},\mathcal{U}} = I_{\mathcal{V}}$, *and* $\rho_{\mathcal{U},\mathcal{W}}\rho_{\mathcal{W},\mathcal{U}} = \rho_{\mathcal{W},\mathcal{U}}\rho_{\mathcal{U},\mathcal{W}} = 0$.

Proof. For $\mathbf{v} = \mathbf{u} + \mathbf{w}$, $\mathbf{u} \in \mathcal{U}$, $\mathbf{w} \in \mathcal{W}$, $(\rho_{\mathcal{U},\mathcal{W}} + \rho_{\mathcal{W},\mathcal{U}})\mathbf{v} = \mathbf{u} + \mathbf{w} = \mathbf{v}$,
and $\rho_{\mathcal{U},\mathcal{W}}\rho_{\mathcal{W},\mathcal{U}}\mathbf{v} = \rho_{\mathcal{U},\mathcal{W}}\mathbf{w} = \mathbf{0}$; $\rho_{\mathcal{W},\mathcal{U}}\rho_{\mathcal{U},\mathcal{W}}\mathbf{v} = \rho_{\mathcal{W},\mathcal{U}}\mathbf{u} = \mathbf{0}$. Thus the stated operator identities hold.
Structural Commentary: A direct-sum decomposition gives two complementary idempotents.
Key Insight: Projection onto one summand kills the other summand. □

Definition 1.3.18. *Orthogonal Projections in the Algebraic Sense*
Two algebraic projections $\rho, \sigma \in \mathcal{L}(\mathcal{V})$ are *algebraically orthogonal*, written $\rho \perp \sigma$, if $\rho\sigma = \sigma\rho = 0$.

Resolution of identity on $\mathcal{V}$ is a finite family of algebraic projections $(\rho_i)_{i=1}^k$ with $\rho_i \perp \rho_j$ for $i \neq j$ and $\sum_{i=1}^k \rho_i = I_{\mathcal{V}}$.

Theorem 1.3.19. *Resolution Ranges Form a Direct Sum*
If $(\rho_i)_{i=1}^k$ is a resolution of $I_{\mathcal{V}}$ and $R_i = \operatorname{range}(\rho_i)$, then $\mathcal{V} = \bigoplus_{i=1}^k R_i$, and $\ker(\rho_i) = \bigoplus_{j\neq i} R_j$.

Proof. Since $\sum_{i=1}^k \rho_i = I_{\mathcal{V}}$, $\mathbf{v} = \sum_{i=1}^k \rho_i\mathbf{v} \in \sum_{i=1}^k R_i$.
If $\sum_{i=1}^k \mathbf{x}_i = \mathbf{0}$ with $\mathbf{x}_i \in R_i$, then for each $1 \leq m \leq k$, $\mathbf{0} = \rho_m \sum_{i=1}^k \mathbf{x}_i = \mathbf{x}_m$, as ρ_m fixes R_m and kills R_i for $i \neq m$.
Thus $\mathcal{V} = \bigoplus_{i=1}^k R_i$. Also, for $\mathbf{v} = \sum_{j=1}^k \mathbf{x}_j$ with $\mathbf{x}_j \in R_j$, $\rho_i\mathbf{v} = \mathbf{x}_i$, so $\ker(\rho_i) = \bigoplus_{j\neq i} R_j$.
Structural Commentary: A resolution decomposes the identity into independent idempotent pieces.
Key Insight: Pairwise zero products are the operator form of directness. □

Theorem 1.3.20. *Direct Sums Give Resolutions*
If $\mathcal{V} = \bigoplus_{i=1}^k \mathcal{U}_i$ and $\rho_i = \rho_{\mathcal{U}_i,\bigoplus_{j\neq i}\mathcal{U}_j}$, then $(\rho_i)_{i=1}^k$ is a resolution of $I_{\mathcal{V}}$.

Proof. For $\mathbf{v} = \sum_{i=1}^k \mathbf{u}_i$, $\mathbf{u}_i \in \mathcal{U}_i$, $\sum_{i=1}^k \rho_i\mathbf{v} = \sum_{i=1}^k \mathbf{u}_i = \mathbf{v}$. Thus $\sum_{i=1}^k \rho_i = I_{\mathcal{V}}$.
If $i \neq j$, then $\rho_i\rho_j\mathbf{v} = \rho_i\mathbf{u}_j = \mathbf{0}$ and $\rho_j\rho_i\mathbf{v} = \rho_j\mathbf{u}_i = \mathbf{0}$, so $\rho_i \perp \rho_j$.
Structural Commentary: Direct-sum coordinates produce mutually annihilating coordinate projections.
Key Insight: Coordinate projections add to the identity and kill all other summands. □

Theorem 1.3.21. *Algebra of Projections*
Let $\rho, \sigma \in \mathcal{L}(\mathcal{V})$ be algebraic projections and $\operatorname{char}(\mathbb{F}) \neq 2$.

(i) $\rho + \sigma$ is an algebraic projection $\iff \rho \perp \sigma$.
Then $\operatorname{range}(\rho + \sigma) = \operatorname{range}(\rho) \oplus \operatorname{range}(\sigma)$; $\ker(\rho + \sigma) = \ker(\rho) \cap \ker(\sigma)$
(ii) $\rho - \sigma$ is an algebraic projection $\iff \rho\sigma = \sigma\rho = \sigma$.
Then $\operatorname{range}(\rho - \sigma) = \operatorname{range}(\rho) \cap \ker(\sigma)$; $\ker(\rho - \sigma) = \ker(\rho) \oplus \operatorname{range}(\sigma)$.
(iii) If $\rho\sigma = \sigma\rho$, then $\rho\sigma$ is algebraic projection and $\operatorname{range}(\rho\sigma) = \operatorname{range}(\rho) \cap \operatorname{range}(\sigma)$; $\ker(\rho\sigma) = \ker(\rho) + \ker(\sigma)$.

Proof.

(i) Put $A = \rho\sigma$, $B = \sigma\rho$, and $C = \rho\sigma\rho$. Since $(\rho + \sigma)^2 = \rho + \sigma \iff A + B = 0$,
$A + B = 0$ gives $A + C = \rho(A + B) = 0$ and $C + B = (A + B)\rho = 0$, hence $B = A$ and $2A = 0$.
Thus $A = B = 0$. The converse is immediate. Assume $\rho \perp \sigma$. Then $\operatorname{range}(\rho + \sigma) \subseteq \operatorname{range}(\rho) + \operatorname{range}(\sigma)$,
while $(\rho + \sigma)(\mathbf{x} + \mathbf{y}) = \mathbf{x} + \mathbf{y}$ for $\mathbf{x} \in \operatorname{range}(\rho)$, $\mathbf{y} \in \operatorname{range}(\sigma)$. Also $\operatorname{range}(\rho) \cap \operatorname{range}(\sigma) = \{\mathbf{0}\}$, since
$\mathbf{z} \in \operatorname{range}(\rho) \cap \operatorname{range}(\sigma)$ and $\mathbf{z} = \sigma\mathbf{w}$ imply $\mathbf{z} = \rho\mathbf{z} = \rho\sigma\mathbf{w} = \mathbf{0}$. Finally, $(\rho + \sigma)\mathbf{v} = 0 \iff \rho\mathbf{v} = 0 = \sigma\mathbf{v}$.

(ii) Since $P^2 = P \iff (I_{\mathcal{V}} - P)^2 = I_{\mathcal{V}} - P$, $\rho - \sigma$ is an algebraic projection $\iff (I_{\mathcal{V}} - \rho) + \sigma$ is an algebraic projection. By (i), this is equivalent to $(I_{\mathcal{V}} - \rho)\sigma = \sigma(I_{\mathcal{V}} - \rho) = 0$, i.e. $\rho\sigma = \sigma\rho = \sigma$.
Assume $\rho\sigma = \sigma\rho = \sigma$ and set $Q = \rho - \sigma$. Then $Q\mathbf{v} \in \operatorname{range}(\rho) \cap \ker(\sigma)$; $Q\mathbf{x} = \mathbf{x}$, $\mathbf{x} \in \operatorname{range}(\rho) \cap \ker(\sigma)$,
so $\operatorname{range}(Q) = \operatorname{range}(\rho) \cap \ker(\sigma)$. Also Q kills both $\ker(\rho)$ and $\operatorname{range}(\sigma)$, and $\ker(\rho) \cap \operatorname{range}(\sigma) = \{\mathbf{0}\}$.
If $\mathbf{v} \in \ker(Q)$, then $\rho\mathbf{v} = \sigma\mathbf{v}$ and $\mathbf{v} = (\mathbf{v} - \rho\mathbf{v}) + \sigma\mathbf{v} \in \ker(\rho) \oplus \operatorname{range}(\sigma)$. Hence $\ker(Q) = \ker(\rho) \oplus \operatorname{range}(\sigma)$.

(iii) If $\rho\sigma = \sigma\rho$, then $(\rho\sigma)^2 = \rho^2\sigma^2 = \rho\sigma$. Moreover $\mathbf{x} \in \operatorname{range}(\rho\sigma) \iff \rho\mathbf{x} = \mathbf{x} = \sigma\mathbf{x}$,
so $\operatorname{range}(\rho\sigma) = \operatorname{range}(\rho) \cap \operatorname{range}(\sigma)$. Finally, $\ker(\rho) + \ker(\sigma) \subseteq \ker(\rho\sigma)$, and if $\rho\sigma\mathbf{v} = 0$,
then $\mathbf{v} = (\mathbf{v} - \sigma\mathbf{v}) + \sigma\mathbf{v}$; $\mathbf{v} - \sigma\mathbf{v} \in \ker(\sigma)$, $\sigma\mathbf{v} \in \ker(\rho)$. Thus $\ker(\rho\sigma) = \ker(\rho) + \ker(\sigma)$.

Structural Commentary: Operator identities encode range-kernel geometry.
Key Insight: Idempotence is a rigid equation; expanding it forces the cross terms to encode the geometry. □

Proposition 1.3.22. *Shift Operators on Finitely Nonzero Sequences*
Define $R, S : \mathbb{F}^{(\mathbb{N})} \to \mathbb{F}^{(\mathbb{N})}$ by $R(x_1, x_2, \ldots) = (0, x_1, x_2, \ldots)$, $S(x_1, x_2, \ldots) = (x_2, x_3, \ldots)$. Then $R, S \in \mathcal{L}(\mathbb{F}^{(\mathbb{N})})$.
Also, $\ker(R) = \{\mathbf{0}\}$, $\operatorname{range}(R) = \{(x_1, x_2, \ldots) \in \mathbb{F}^{(\mathbb{N})} : x_1 = 0\}$, $\ker(S) = \operatorname{span}((1, 0, 0, \ldots))$, $\operatorname{range}(S) = \mathbb{F}^{(\mathbb{N})}$.

Proof. Both R and S preserve finite support, so they map $\mathbb{F}^{(\mathbb{N})}$ into itself.
Since both maps act componentwise by shifting entries, $R(a\mathbf{x}+\mathbf{y}) = aR\mathbf{x}+R\mathbf{y}$ and $S(a\mathbf{x}+\mathbf{y}) = aS\mathbf{x}+S\mathbf{y}$ for all $\mathbf{x},\mathbf{y} \in \mathbb{F}^{(\mathbb{N})}$ and $a \in \mathbb{F}$. Thus R and S are linear.
The map R has zero kernel as $R\mathbf{x} = \mathbf{0}$ forces every entry of $\mathbf{x}$ to be zero, and its range is exactly the set of finitely nonzero sequences with first entry zero. For S, $S\mathbf{x} = \mathbf{0}$ forces $x_2 = x_3 = \cdots = 0$, so $\ker(S) = \operatorname{span}((1,0,0,\ldots))$.
Finally, for any $(y_1, y_2, \ldots) \in \mathbb{F}^{(\mathbb{N})}$, we have $S(0, y_1, y_2, \ldots) = (y_1, y_2, \ldots)$, so S is surjective.
Structural Commentary: Shift operators expose behavior that finite-dimensional maps cannot have.
Key Insight: The right shift is injective but not surjective, while the left shift is surjective but not injective. □

Example 1.3.23. *Shift Operators Show Noncommutativity*
For the shifts $R, S \in \mathcal{L}(\mathbb{F}^{(\mathbb{N})})$ from Proposition 1.3.22, $SR = I_{\mathbb{F}^{(\mathbb{N})}}$ and $RS(x_1, x_2, \ldots) = (0, x_2, x_3, \ldots)$.
Hence $RS \neq SR$, for instance on $(1,0,0,\ldots)$. Therefore $\mathcal{L}(\mathbb{F}^{(\mathbb{N})})$ is not commutative under composition.

Proposition 1.3.24. *Injectivity and the Kernel*
A linear map $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$ is injective if and only if $\ker(T) = \{\mathbf{0}\}$.

Proof.
($\Rightarrow$) If T is injective and $\mathbf{v} \in \ker(T)$, then $T\mathbf{v} = \mathbf{0} = T\mathbf{0}$, so $\mathbf{v} = \mathbf{0}$.
($\Leftarrow$) If $\ker(T) = \{\mathbf{0}\}$ and $T\mathbf{v} = T\mathbf{w}$, then $T(\mathbf{v}-\mathbf{w}) = \mathbf{0}$. Hence $\mathbf{v}-\mathbf{w} \in \ker(T)$, so $\mathbf{v} = \mathbf{w}$.
Structural Commentary: The kernel measures exactly where a linear map collapses vectors to zero.
Key Insight: Injectivity is equivalent to collapsing only the zero vector. □

Theorem 1.3.25. *Range from Basis Images*
Let $\mathcal{V}$ be finite-dimensional with basis $(\mathbf{v}_1, \ldots, \mathbf{v}_n)$, and let $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$. Then $\operatorname{range}(T) = \operatorname{span}(T\mathbf{v}_1, \ldots, T\mathbf{v}_n)$.
In particular, $\operatorname{range}(T)$ *is finite-dimensional and* $\dim(\operatorname{range}(T)) \leq n$.

Proof. If $\mathbf{v} = c_1\mathbf{v}_1 + \cdots + c_n\mathbf{v}_n$, then $T\mathbf{v} = c_1T\mathbf{v}_1 + \cdots + c_nT\mathbf{v}_n$.
Thus every vector in $\operatorname{range}(T)$ lies in $\operatorname{span}(T\mathbf{v}_1, \ldots, T\mathbf{v}_n)$, and the reverse inclusion is immediate from the definition of range. The dimension bound follows because this spanning list has length n.
Structural Commentary: The range is generated by the images of any basis of the domain.
Key Insight: Linear maps carry a basis-spanning formula directly into a range-spanning formula. □

Definition 1.3.26. *Rank and Nullity*
For a finite-dimensional domain, the rank of T is $\dim(\operatorname{range}(T))$ and the nullity of T is $\dim(\ker(T))$.

Theorem 1.3.27. *Rank-Nullity Theorem*
If $\mathcal{V}$ is finite-dim and $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$, then $\operatorname{range}(T)$ *is finite-dim and* $\dim(\mathcal{V}) = \dim(\ker(T)) + \dim(\operatorname{range}(T))$.

Proof. By Proposition 1.3.14 and Corollary 1.1.28, $\ker(T)$ has a basis $(\mathbf{u}_1, \ldots, \mathbf{u}_m)$.
Extend it to a basis $(\mathbf{u}_1, \ldots, \mathbf{u}_m, \mathbf{w}_1, \ldots, \mathbf{w}_r)$ of $\mathcal{V}$ by Theorem 1.1.27.
If $\mathbf{v} = \sum_{i=1}^{m} a_i\mathbf{u}_i + \sum_{j=1}^{r} b_j\mathbf{w}_j$, then $T\mathbf{v} = \sum_{j=1}^{r} b_jT\mathbf{w}_j$, because $T\mathbf{u}_i = \mathbf{0}$. Hence $(T\mathbf{w}_1, \ldots, T\mathbf{w}_r)$ spans $\operatorname{range}(T)$.
If $\sum_{j=1}^{r} c_jT\mathbf{w}_j = \mathbf{0}$, then $\sum_{j=1}^{r} c_j\mathbf{w}_j \in \ker(T)$, so for some scalars d_i, $\sum_{j=1}^{r} c_j\mathbf{w}_j = \sum_{i=1}^{m} d_i\mathbf{u}_i$.
Thus $\sum_{j=1}^{r} c_j\mathbf{w}_j - \sum_{i=1}^{m} d_i\mathbf{u}_i = \mathbf{0}$, and linear independence of extended basis forces every c_j and d_i to be zero.
Hence $(T\mathbf{w}_1, \ldots, T\mathbf{w}_r)$ is a basis of $\operatorname{range}(T)$. Thus $\dim(\ker(T)) = m$, $\dim(\operatorname{range}(T)) = r$, and $\dim(\mathcal{V}) = m+r$.
Structural Commentary: An adapted basis splits $\mathcal{V}$ into kernel directions and directions detected by T.
Key Insight: Rank-nullity says every domain dimension is either killed by T or appears in its range. □

Corollary 1.3.28. *Maps to Smaller Spaces Are Not Injective*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional vector spaces with $\dim(\mathcal{V}) > \dim(\mathcal{W})$.
Then no linear map from $\mathcal{V}$ to $\mathcal{W}$ is injective.

Proof. Let $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$. By Theorem 1.3.27, $\dim(\ker(T)) = \dim(\mathcal{V}) - \dim(\operatorname{range}(T)) \geq \dim(\mathcal{V}) - \dim(\mathcal{W}) > 0$.
Thus $\ker(T)$ contains a nonzero vector, so T is not injective by Proposition 1.3.24.
Structural Commentary: Excess domain dimension forces collapse.
Key Insight: The dimension deficit must appear inside the kernel. □

Corollary 1.3.29. *Maps to Larger Spaces Are Not Surjective*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional vector spaces with $\dim(\mathcal{V}) < \dim(\mathcal{W})$.
Then no linear map from $\mathcal{V}$ to $\mathcal{W}$ is surjective.

Proof. Let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. By Theorem 1.3.27, $\dim(\text{range}(T)) = \dim(\mathcal{V}) - \dim(\ker(T)) \leq \dim(\mathcal{V}) < \dim(\mathcal{W})$. Hence $\text{range}(T)$ is a proper subspace of $\mathcal{W}$, so T is not surjective.
Structural Commentary: A smaller domain cannot cover a larger codomain.
Key Insight: The range can never have dimension larger than the domain. □

Corollary 1.3.30. *Equal-Dimension Kernel-Range Criterion*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional vector spaces with $\dim(\mathcal{V}) = \dim(\mathcal{W})$, and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. Then $\ker(T) = \{\mathbf{0}\}$ if and only if $\text{range}(T) = \mathcal{W}$. Equivalently, T is injective if and only if T is surjective.

Proof.
($\Rightarrow$) If $\ker(T) = \{\mathbf{0}\}$, then rank-nullity gives $\dim(\text{range}(T)) = \dim(\mathcal{V}) = \dim(\mathcal{W})$.
Since $\text{range}(T)$ is a subspace of $\mathcal{W}$ with the same dimension, $\text{range}(T) = \mathcal{W}$.
($\Leftarrow$) If $\text{range}(T) = \mathcal{W}$, then rank-nullity gives $\dim(\ker(T)) = \dim(\mathcal{V}) - \dim(\mathcal{W}) = 0$, hence $\ker(T) = \{\mathbf{0}\}$.
Structural Commentary: Equal finite dimensions make injectivity and surjectivity equivalent.
Key Insight: In equal dimensions, having no kernel loss is exactly having full range. □

Definition 1.3.31. *Invertible Linear Map and Isomorphism*
A linear map $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ is *invertible* if there exists $S \in \mathcal{L}(\mathcal{W}, \mathcal{V})$ such that $ST = I_{\mathcal{V}}$ and $TS = I_{\mathcal{W}}$.
An invertible linear map is called an *isomorphism*, and if one exists we write $\mathcal{V} \cong \mathcal{W}$.

Proposition 1.3.32. *Inverse Is Unique*
An invertible linear map has a unique inverse. Thus if T is invertible, its inverse is denoted T^{-1}.

Proof. Suppose $S_1, S_2 \in \mathcal{L}(\mathcal{W}, \mathcal{V})$ are inverses of $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.
Then $S_1 = S_1 I_{\mathcal{W}} = S_1(TS_2) = (S_1 T)S_2 = I_{\mathcal{V}} S_2 = S_2$. Hence the inverse is unique.
Structural Commentary: A two-sided inverse is forced by the composition identities.
Key Insight: Associativity lets the two candidate inverses meet around T and collapse to the same map. □

Proposition 1.3.33. *Inverse of a Product Reverses Order*
Let $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ and $S \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ be invertible. Then ST is invertible and $(ST)^{-1} = T^{-1}S^{-1}$.

Proof. Since $T^{-1} \in \mathcal{L}(\mathcal{V},\mathcal{U})$ and $S^{-1} \in \mathcal{L}(\mathcal{W},\mathcal{V})$, the product $T^{-1}S^{-1} \in \mathcal{L}(\mathcal{W},\mathcal{U})$ is defined.
Moreover, $(T^{-1}S^{-1})(ST) = T^{-1}(S^{-1}S)T = T^{-1}T = I_{\mathcal{U}}$ and $(ST)(T^{-1}S^{-1}) = S(TT^{-1})S^{-1} = SS^{-1} = I_{\mathcal{W}}$.
Thus $T^{-1}S^{-1}$ is a two-sided inverse of ST, so $(ST)^{-1} = T^{-1}S^{-1}$.
Structural Commentary: To undo a composition, undo the last map first.
Key Insight: Composition order reverses under inversion. □

Proposition 1.3.34. *Invertibility and Bijectivity*
A linear map is invertible if and only if it is both injective and surjective.

Proof.
($\Rightarrow$) Let S be an inverse of T. If $T\mathbf{v} = T\mathbf{w}$, then $\mathbf{v} = ST\mathbf{v} = ST\mathbf{w} = \mathbf{w}$, so T is injective. Also every $\mathbf{w}$ in the codomain satisfies $\mathbf{w} = T(S\mathbf{w})$, so T is surjective.
($\Leftarrow$) Suppose T is injective and surjective. Define $S\mathbf{w}$ to be the unique vector $\mathbf{v}$ with $T\mathbf{v} = \mathbf{w}$.
Since both sides below map under T to $a\mathbf{w}_1 + \mathbf{w}_2$, $S(a\mathbf{w}_1 + \mathbf{w}_2) = aS\mathbf{w}_1 + S\mathbf{w}_2$. Thus S is linear and satisfies $ST = I$ and $TS = I$, so T is invertible.
Structural Commentary: Bijectivity gives a unique reverse map, and linearity makes it a linear inverse.
Key Insight: A linear map is invertible exactly when it is one-to-one and onto. □

Definition 1.3.35. *Morphism Terminology for Linear Maps*
For $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$:

(i) T is a *homomorphism*.
(ii) If $\mathcal{V} = \mathcal{W}$, then T is an *endomorphism*.
(iii) If $\ker(T) = \{\mathbf{0}\}$, then T is a *monomorphism*.
(iv) If $\text{range}(T) = \mathcal{W}$, then T is an *epimorphism*.
(v) If $\ker(T) = \{\mathbf{0}\}$ and $\text{range}(T) = \mathcal{W}$, then T is an *isomorphism*.
(vi) If $\mathcal{V} = \mathcal{W}$ and T is invertible, then T is an *automorphism*.

Proposition 1.3.36. *Rank-Nullity as a Complement Isomorphism*
Let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, and suppose $\mathcal{V} = \ker(T) \oplus \mathcal{M}$. Then $T|_{\mathcal{M}} : \mathcal{M} \to \text{range}(T)$ is an isomorphism.

Proof. Indeed, $\ker(T|_{\mathcal{M}}) = \mathcal{M} \cap \ker(T) = \{\mathbf{0}\}$, and every $T\mathbf{v}$ equals $T\mathbf{m}$ for the decomposition $\mathbf{v} = \mathbf{k} + \mathbf{m}$.
Thus rank-nullity counts $\mathcal{V} = \ker(T) \oplus \mathcal{M}$ with $\mathcal{M} \cong \text{range}(T)$. □

Definition 1.3.37. *Quotient Map*
Let $\mathcal{U}$ be a subspace of $\mathcal{V}$. The *quotient map* is the map $\pi : \mathcal{V} \to \mathcal{V}/\mathcal{U}$ defined by $\pi(\mathbf{v}) = \mathbf{v} + \mathcal{U}$.

Proposition 1.3.38. *Properties of the Quotient Map*
The quotient map π is linear, surjective, and $\ker(\pi) = \mathcal{U}$.

Proof. For $\mathbf{v}, \mathbf{w} \in \mathcal{V}$ and $a \in \mathbb{F}$, $\pi(a\mathbf{v}+\mathbf{w}) = (a\mathbf{v}+\mathbf{w})+\mathcal{U} = a(\mathbf{v}+\mathcal{U})+(\mathbf{w}+\mathcal{U}) = a\pi(\mathbf{v})+\pi(\mathbf{w})$ so π is linear. It is surjective because every coset in $\mathcal{V}/\mathcal{U}$ has the form $\mathbf{v}+\mathcal{U} = \pi(\mathbf{v})$. Finally, $\ker(\pi) = \{\mathbf{v} \in \mathcal{V} : \mathbf{v}+\mathcal{U} = \mathcal{U}\} = \mathcal{U}$.
Structural Commentary: The quotient map collapses $\mathcal{U}$ to zero and records cosets.
Key Insight: Quotient maps realise subspaces as kernels. □

Theorem 1.3.39. *Universal Property of the Quotient Map*
Let $\mathcal{U} \leq \mathcal{V}$, $\pi : \mathcal{V} \to \mathcal{V}/\mathcal{U}$, and $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ with $\mathcal{U} \leq \ker(T)$.
Then $\exists!\, \overline{T} \in \mathcal{L}(\mathcal{V}/\mathcal{U}, \mathcal{W})$ such that $\overline{T}\pi = T$, $\overline{T}(\mathbf{v}+\mathcal{U}) = T\mathbf{v}$, and $\ker(\overline{T}) = \ker(T)/\mathcal{U}$, $\operatorname{range}(\overline{T}) = \operatorname{range}(T)$.

Proof. Define $\overline{T}(\mathbf{v}+\mathcal{U}) = T\mathbf{v}$. If $\mathbf{v}+\mathcal{U} = \mathbf{w}+\mathcal{U}$, then $\mathbf{v}-\mathbf{w} \in \mathcal{U} \leq \ker(T)$, so $T\mathbf{v} = T\mathbf{w}$; hence $\overline{T}$ is well-defined.
Also $\overline{T}\big(a(\mathbf{v}+\mathcal{U}) + b(\mathbf{w}+\mathcal{U})\big) = T(a\mathbf{v}+b\mathbf{w}) = aT\mathbf{v} + bT\mathbf{w}$.
By definition, $\overline{T}\pi = T$. If $S\pi = T$, then $S(\mathbf{v}+\mathcal{U}) = S\pi(\mathbf{v}) = T\mathbf{v} = \overline{T}(\mathbf{v}+\mathcal{U})$, so $S = \overline{T}$.
Also $\mathbf{v}+\mathcal{U} \in \ker(\overline{T}) \iff T\mathbf{v} = \mathbf{0} \iff \mathbf{v} \in \ker(T)$.
So $\ker(\overline{T}) = \ker(T)/\mathcal{U}$, and $\operatorname{range}(\overline{T}) = \{\overline{T}(\mathbf{v}+\mathcal{U}) : \mathbf{v} \in \mathcal{V}\} = \operatorname{range}(T)$.
Structural Commentary: Maps killing $\mathcal{U}$ factor through $\mathcal{V}/\mathcal{U}$.
Key Insight: Quotients are universal domains for maps with prescribed kernel ambiguity. □

Theorem 1.3.40. *Subspace Correspondence for Quotients*
Let $\pi : \mathcal{V} \to \mathcal{V}/\mathcal{U}$ be the quotient map. The assignments $\mathcal{T} \longmapsto \mathcal{T}/\mathcal{U}$ and $\mathcal{Y} \longmapsto \pi^{-1}(\mathcal{Y})$ are inverse inclusion-preserving bijections between $\{\mathcal{T}$ subspace of $\mathcal{V} : \mathcal{U} \subseteq \mathcal{T}\}$ and $\{\mathcal{Y}$ subspace of $\mathcal{V}/\mathcal{U}\}$.

Proof. If $\mathcal{U} \subseteq \mathcal{T} \subseteq \mathcal{V}$, then $\mathcal{T}/\mathcal{U}$ is a subspace of $\mathcal{V}/\mathcal{U}$.
If $\mathcal{Y} \subseteq \mathcal{V}/\mathcal{U}$ is a subspace, then $\pi^{-1}(\mathcal{Y})$ is a subspace of $\mathcal{V}$ and contains $\ker(\pi) = \mathcal{U}$.
For $\mathcal{U} \subseteq \mathcal{T} \subseteq \mathcal{V}$, $\pi^{-1}(\mathcal{T}/\mathcal{U}) = \mathcal{T}$, since $\mathbf{v}+\mathcal{U} = \mathbf{t}+\mathcal{U}$ with $\mathbf{t} \in \mathcal{T}$ implies $\mathbf{v}-\mathbf{t} \in \mathcal{U} \subseteq \mathcal{T}$.
For $\mathcal{Y} \subseteq \mathcal{V}/\mathcal{U}$, surjectivity gives $\pi\left(\pi^{-1}(\mathcal{Y})\right) = \mathcal{Y}$, that is, $\pi^{-1}(\mathcal{Y})/\mathcal{U} = \mathcal{Y}$
Both assignments preserve inclusion by monotonicity of images and preimages.
Structural Commentary: Subspaces above $\mathcal{U}$ are exactly the subspaces visible after collapsing $\mathcal{U}$.
Key Insight: The quotient map translates containment above $\mathcal{U}$ into containment inside $\mathcal{V}/\mathcal{U}$. □

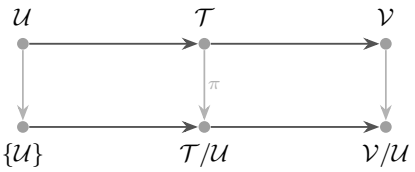


Figure 12: Intermediate subspaces above $\mathcal{U}$ correspond to subspaces of the quotient.

Theorem 1.3.41. *Quotient by the Kernel*
Let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. Then $\widetilde{T}(\mathbf{v} + \ker(T)) := T\mathbf{v}$ determines a well-defined linear map $\widetilde{T} : \mathcal{V}/\ker(T) \to \mathcal{W}$. Also,

(i) $\widetilde{T}$ is injective;
(ii) $\operatorname{range}(\widetilde{T}) = \operatorname{range}(T)$;
(iii) $\mathcal{V}/\ker(T) \cong \operatorname{range}(T)$.

Proof. If $\mathbf{v} + \ker(T) = \mathbf{v}' + \ker(T)$, then $\mathbf{v} - \mathbf{v}' \in \ker(T)$, so $T\mathbf{v} = T\mathbf{v}'$. Thus $\widetilde{T}$ is well-defined.
For cosets and $a \in \mathbb{F}$, $\widetilde{T}\big(a(\mathbf{v} + \ker(T)) + (\mathbf{v}' + \ker(T))\big) = T(a\mathbf{v} + \mathbf{v}') = aT\mathbf{v} + T\mathbf{v}'$, so $\widetilde{T}$ is linear.

(i) If $\widetilde{T}(\mathbf{v} + \ker(T)) = \mathbf{0}$, then $T\mathbf{v} = \mathbf{0}$, so $\mathbf{v} \in \ker(T)$ and $\mathbf{v} + \ker(T) = \ker(T)$, the zero coset.
Hence $\ker(\widetilde{T}) = \{\ker(T)\}$ and $\widetilde{T}$ is injective.
(ii) For every coset $\mathbf{v} + \ker(T)$, the value $\widetilde{T}(\mathbf{v} + \ker(T)) = T\mathbf{v}$ lies in $\operatorname{range}(T)$.
Conversely, every vector in $\operatorname{range}(T)$ has the form $T\mathbf{v} = \widetilde{T}(\mathbf{v} + \ker(T))$. Hence the ranges are equal.
(iii) Let $\widehat{T} : \mathcal{V}/\ker(T) \to \operatorname{range}(T)$ be the same formula as $\widetilde{T}$ with codomain restricted to $\operatorname{range}(T)$.
The linearity proved above and parts (i)–(ii) show that $\widehat{T}$ is linear, injective, and surjective, so Proposition 1.3.34 gives $\mathcal{V}/\ker(T) \cong \operatorname{range}(T)$.

Structural Commentary: Quotienting by the kernel removes exactly the ambiguity of T.
Key Insight: Every linear map factors as a quotient followed by an isomorphism onto its range. □

Proposition 1.3.42. *Quotient by a Complemented Subspace*
Let $\mathcal{U}$ and $\mathcal{W}$ be subspaces of $\mathcal{V}$ with $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$. Then $J : \mathcal{W} \to \mathcal{V}/\mathcal{U}$, $J(\mathbf{w}) = \mathbf{w} + \mathcal{U}$ is an isomorphism. Hence every complement of $\mathcal{U}$ is isomorphic to $\mathcal{V}/\mathcal{U}$.

Proof. The map is linear by the quotient operations from Proposition 1.2.31.
If $J(\mathbf{w}) = \mathcal{U}$, then $\mathbf{w} \in \mathcal{U} \cap \mathcal{W} = \{\mathbf{0}\}$, so J is injective.
If $\mathbf{v} + \mathcal{U} \in \mathcal{V}/\mathcal{U}$, write $\mathbf{v} = \mathbf{u} + \mathbf{w}$ with $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{W}$. Then $\mathbf{v} + \mathcal{U} = \mathbf{w} + \mathcal{U} = J(\mathbf{w})$, so J is surjective.
Structural Commentary: A complement gives canonical representatives for quotient cosets.
Key Insight: Collapsing $\mathcal{U}$ leaves exactly the complementary directions visible. □

Proposition 1.3.43. *Quotient of a Direct Sum*
If $\mathcal{V} = \mathcal{V}_1 \oplus \mathcal{V}_2$ and $\mathcal{U}_j \leq \mathcal{V}_j$ for $j = 1, 2$, then $\frac{\mathcal{V}}{\mathcal{U}_1 \oplus \mathcal{U}_2} \cong \frac{\mathcal{V}_1}{\mathcal{U}_1} \oplus \frac{\mathcal{V}_2}{\mathcal{U}_2}$.

Proof. Define, using the unique decomposition
$\mathbf{v} = \mathbf{v}_1 + \mathbf{v}_2$, $\Phi : \mathcal{V} \to (\mathcal{V}_1/\mathcal{U}_1) \oplus (\mathcal{V}_2/\mathcal{U}_2)$ and $\Phi(\mathbf{v}_1 + \mathbf{v}_2) = (\mathbf{v}_1 + \mathcal{U}_1, \mathbf{v}_2 + \mathcal{U}_2)$.
For $\mathbf{x} = \mathbf{x}_1 + \mathbf{x}_2$, $\mathbf{y} = \mathbf{y}_1 + \mathbf{y}_2$, and $a, b \in \mathbb{F}$, $\Phi(a\mathbf{x} + b\mathbf{y}) = a\Phi(\mathbf{x}) + b\Phi(\mathbf{y})$, so Φ is linear.
It is surjective since $(\mathbf{v}_1 + \mathcal{U}_1, \mathbf{v}_2 + \mathcal{U}_2) = \Phi(\mathbf{v}_1 + \mathbf{v}_2)$. Finally, $\ker(\Phi) = \{\mathbf{v}_1 + \mathbf{v}_2 : \mathbf{v}_1 \in \mathcal{U}_1,\ \mathbf{v}_2 \in \mathcal{U}_2\} = \mathcal{U}_1 \oplus \mathcal{U}_2$.
The result follows from Theorem 1.3.41.
Structural Commentary: Quotienting respects direct-sum coordinates.
Key Insight: Direct-sum coordinates make quotienting componentwise. □

Perspective 1.3.44. *Quotient and Complement Warnings*
If $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$, then $\mathcal{W} \cong \mathcal{V}/\mathcal{U}$, but complements are not functorial in $\mathcal{U}$.
For $\mathcal{V}$ countably infinite-dimensional with basis $(\mathbf{e}_n)_{n \geq 1}$, put $\mathcal{S} = \operatorname{span}(\mathbf{e}_n : n \geq 2)$.
Then $\mathcal{S} \cong \mathcal{V}$ and $\mathcal{V} = \mathcal{V} \oplus \{\mathbf{0}\} = \mathcal{S} \oplus \mathbb{F}\mathbf{e}_1$; $\mathcal{V}/\mathcal{V} = \{\mathbf{0}\}$, $\mathcal{V}/\mathcal{S} \cong \mathbb{F}\mathbf{e}_1$.

Definition 1.3.45. *Codimension*
If $\mathcal{U}$ is a subspace of $\mathcal{V}$ and $\dim(\mathcal{V}/\mathcal{U}) < \infty$, define $\operatorname{codim}_{\mathcal{V}}(\mathcal{U}) := \dim(\mathcal{V}/\mathcal{U})$.
Thus, if $\mathcal{V}$ is finite-dimensional, Theorem 1.2.32 gives $\operatorname{codim}_{\mathcal{V}}(\mathcal{U}) = \dim(\mathcal{V}) - \dim(\mathcal{U})$.
If $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$, Proposition 1.3.42 gives $\operatorname{codim}_{\mathcal{V}}(\mathcal{U}) = \dim(\mathcal{W})$.

Perspective 1.3.46. *Quotients Split Over Fields*
In finite-dimensional vector spaces, choosing a complement $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$ identifies $\mathcal{V}/\mathcal{U}$ with $\mathcal{W}$. The quotient map collapses $\mathcal{U}$ and keeps the complementary directions.

Proposition 1.3.47. *One-Sided Inverses Split in Finite Dimensions*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional vector spaces and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.

(i) If T is injective, then there exists $L \in \mathcal{L}(\mathcal{W}, \mathcal{V})$ such that $LT = I_{\mathcal{V}}$.
(ii) If T is surjective, then there exists $R \in \mathcal{L}(\mathcal{W}, \mathcal{V})$ such that $TR = I_{\mathcal{W}}$.

Proof.

(i) Choose $\mathcal{H} \subseteq \mathcal{W}$ with $\mathcal{W} = \operatorname{range}(T) \oplus \mathcal{H}$. For $\mathbf{w} = T\mathbf{v} + \mathbf{h}$, define $L\mathbf{w} = \mathbf{v}$. This is well-defined since the direct-sum decomposition is unique and T is injective. It is linear, and $LT\mathbf{v} = \mathbf{v}$ for every $\mathbf{v} \in \mathcal{V}$.
(ii) Choose $\mathcal{H} \subseteq \mathcal{V}$ with $\mathcal{V} = \ker(T) \oplus \mathcal{H}$. Then $T|_{\mathcal{H}} : \mathcal{H} \to \mathcal{W}$ is injective since $\mathcal{H} \cap \ker(T) = \{\mathbf{0}\}$, and surjective since each $\mathbf{v} = \mathbf{k} + \mathbf{h}$ satisfies $T\mathbf{v} = T\mathbf{h}$. Let $R = (T|_{\mathcal{H}})^{-1}$, viewed as a map $\mathcal{W} \to \mathcal{V}$. Then R is linear and $TR = I_{\mathcal{W}}$.

Structural Commentary: Complements supply the missing directions on which a one-sided inverse is chosen.
Key Insight: Finite-dimensional vector spaces split kernels and ranges. □

Proposition 1.3.48. *Isomorphisms Preserve Bases*
Let $T : \mathcal{V} \to \mathcal{W}$ be an isomorphism and let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a finite list in $\mathcal{V}$. Set $T\beta = (T\mathbf{v}_1, \ldots, T\mathbf{v}_n)$. The following are equivalent:

(i) β is a basis of $\mathcal{V}$.
(ii) $T\beta$ is a basis of $\mathcal{W}$.

Proof.

(i) $\Rightarrow$ (ii) If $\mathbf{w} \in \mathcal{W}$, choose $\mathbf{v} \in \mathcal{V}$ with $\mathbf{w} = T\mathbf{v}$ and write $\mathbf{v} = \sum_j a_j \mathbf{v}_j$. Then $\mathbf{w} = \sum_j a_j T\mathbf{v}_j$, so $T\beta$ spans $\mathcal{W}$. If $\sum_j a_j T\mathbf{v}_j = \mathbf{0}$, applying T^{-1} gives $\sum_j a_j \mathbf{v}_j = \mathbf{0}$; hence all $a_j = 0$. Thus $T\beta$ is a basis.
(ii) $\Rightarrow$ (i) Assume $T\beta$ is a basis of $\mathcal{W}$. If $\mathbf{v} \in \mathcal{V}$, write $T\mathbf{v} = \sum_j a_j T\mathbf{v}_j$; applying T^{-1} gives $\mathbf{v} = \sum_j a_j \mathbf{v}_j$, so β spans $\mathcal{V}$. If $\sum_j a_j \mathbf{v}_j = \mathbf{0}$, applying T gives $\sum_j a_j T\mathbf{v}_j = \mathbf{0}$; hence all $a_j = 0$. Thus β is linearly independent, and therefore a basis of $\mathcal{V}$.

Structural Commentary: Isomorphisms preserve exactly the spanning and independence data that define a basis.
Key Insight: An isomorphism changes the ambient space, not the linear relations among vectors. □

Theorem 1.3.49. *Finite-Dimensional Invertibility Criteria*
If $\mathcal{V}$ is finite-dimensional and $T \in \mathcal{L}(\mathcal{V})$, then the following are equivalent:

i. *T is invertible;*
ii. *T is injective;*
iii. *T is surjective.*

Proof.

$(i) \Rightarrow (ii)$ follows from Proposition 1.3.34.
$(ii) \Rightarrow (iii)$: If T is injective, then $\ker(T) = \{\mathbf{0}\}$.
Rank-nullity gives $\dim(\text{range}(T)) = \dim(\mathcal{V})$, hence $\text{range}(T) = \mathcal{V}$, so T is surjective.
$(iii) \Rightarrow (i)$: If T is surjective, rank-nullity gives $\dim(\ker(T)) = 0$, so T is injective.
Thus T is bijective, hence invertible by Proposition 1.3.34.

Structural Commentary: For operators on finite-dimensional spaces, kernel and range dimensions trade off exactly. Injectivity and surjectivity therefore force each other.
Key Insight: In equal finite dimensions, losing no information is the same as reaching the whole space. □

Perspective 1.3.50. *Kernel Obstruction to Invertibility*
By Theorem 1.3.49 and Proposition 1.3.24, a finite-dimensional operator is invertible exactly when its kernel is zero. Any nonzero $\mathbf{v} \in \ker T$ therefore blocks invertibility.

Theorem 1.3.51. *Isomorphism and Dimension*
Two finite-dimensional vector spaces are isomorphic if and only if they have the same dimension.

Proof.

$(\Rightarrow)$ If $T : \mathcal{V} \to \mathcal{W}$ is an isomorphism and $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ is a basis of $\mathcal{V}$, then Proposition 1.3.48 shows that $(T\mathbf{v}_1, \dots, T\mathbf{v}_n)$ is a basis of $\mathcal{W}$.
$(\Leftarrow)$ If both spaces have dimension n, choose bases $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ and $(\mathbf{w}_1, \dots, \mathbf{w}_n)$.
By Proposition 1.3.11, there is a linear map T with $T\mathbf{v}_j = \mathbf{w}_j$ for every j. Then $T(\sum a_j \mathbf{v}_j) = \sum a_j \mathbf{w}_j$, so $\ker(T) = \{\mathbf{0}\}$ and $\text{range}(T) = \mathcal{W}$. Hence T is an isomorphism.

Structural Commentary: Dimension classifies finite-dimensional spaces up to isomorphism.
Key Insight: Matching bases give an isomorphism; any isomorphism carries bases to bases. □

Definition 1.3.52. *Matrices and Matrix Operations*

(i) *Matrix space*: for $m, n \in \mathbb{Z}_{>0}$, $\text{Mat}(m, n, \mathbb{F}) = \{[A_{j,k}]_{1 \le j \le m, 1 \le k \le n} : A_{j,k} \in \mathbb{F}\}$.
(ii) *Entries, rows, columns*: for $A = [A_{j,k}] \in \text{Mat}(m, n, \mathbb{F})$,
$A_{j,k}$ is the (j,k)-entry, $A_{j,\cdot}$ is row j, and $A_{\cdot,k}$ is column k.
(iii) *Transpose*: for $A \in \text{Mat}(m, n, \mathbb{F})$, $A^t \in \text{Mat}(n, m, \mathbb{F})$ is defined by $(A^t)_{k,j} = A_{j,k}$.
(iv) *Main diagonal*: for $A = [A_{j,k}] \in \text{Mat}(m, n, \mathbb{F})$ and $r = \min(m, n)$, $\text{diag}(A) = (A_{1,1}, \dots, A_{r,r})$.
(v) *Symmetry*: for $A \in \text{Mat}(n, n, \mathbb{F})$, A is symmetric if $A^t = A$ and skew-symmetric if $A^t = -A$.
(vi) *Triangularity*: for $A \in \text{Mat}(n, n, \mathbb{F})$, A is upper triangular if $A_{j,k} = 0$ for $j > k$,
lower triangular if $A_{j,k} = 0$ for $j < k$, and diagonal if $A_{j,k} = 0$ for $j \neq k$.
(vii) *Linear operations*: for $A, C \in \text{Mat}(m, n, \mathbb{F})$ and $a \in \mathbb{F}$, $(A + C)_{j,k} = A_{j,k} + C_{j,k}$ and $(aA)_{j,k} = aA_{j,k}$.
(viii) *Matrix product*: for $A \in \text{Mat}(m, n, \mathbb{F})$, $B \in \text{Mat}(n, p, \mathbb{F})$, $AB \in \text{Mat}(m, p, \mathbb{F})$ is $(AB)_{j,k} := \sum_{r=1}^{n} A_{j,r} B_{r,k}$.
(ix) *Identity and inverse*: $I_n = [\delta_{j,k}] \in \text{Mat}(n, n, \mathbb{F})$; for $A, B \in \text{Mat}(n, n, \mathbb{F})$, B is a two-sided inverse of A if $AB = I_n = BA$, and A is invertible if such B exists.

Lemma 1.3.53. *Uniqueness of Two-Sided Inverses*
Let $A, B, C \in \text{Mat}(n, n, \mathbb{F})$. If $AB = I_n$ and $BC = I_n$, then $A = C$.

Proof. Using $BC = I_n$ and $AB = I_n$, $A = AI_n = A(BC) = (AB)C = I_nC = C$.
Structural Commentary: A common middle inverse forces the left and right inverses to coincide.
Key Insight: Associativity turns $AB = I_n = BC$ into $A = C$. □

Definition 1.3.54. *Inverse Matrix Notation*
If $A \in \text{Mat}(n, n, \mathbb{F})$ is invertible, Lemma 1.3.53 shows that its two-sided inverse is unique.
This inverse is denoted A^{-1}, and the set of invertible matrices in $\text{Mat}(n, n, \mathbb{F})$ is denoted $\text{Mat}(n, n, \mathbb{F})^\times$.

Proposition 1.3.55. *Matrix Spaces*
For positive integers m, n, $\text{Mat}(m, n, \mathbb{F})$ is a vector space over $\mathbb{F}$ with dimension mn.

Proof. Entrywise addition and scalar multiplication satisfy the vector space axioms because $\mathbb{F}$ does. For $1 \le j \le m$ and $1 \le k \le n$, let $E_{j,k}$ be the matrix with 1 in entry (j,k) and 0 elsewhere.

Every $A = [A_{j,k}]$ has the unique expansion $A = \sum_{j=1}^{m} \sum_{k=1}^{n} A_{j,k} E_{j,k}$.

Thus the mn matrices $E_{j,k}$ form a basis, so $\dim(\mathrm{Mat}(m,n,\mathbb{F})) = mn$.
Structural Commentary: A matrix is a finite array of independent scalar slots.
Key Insight: The standard entry matrices isolate those slots one at a time. □

Proposition 1.3.56. *Matrix Product Entries Are Row Times Column*
If $A \in \mathrm{Mat}(m,n,\mathbb{F})$ and $B \in \mathrm{Mat}(n,p,\mathbb{F})$, then $(AB)_{j,k} = A_{j,\cdot} B_{\cdot,k}$ for $1 \le j \le m$ and $1 \le k \le p$.

Proof. The row-column product $A_{j,\cdot} B_{\cdot,k}$ is exactly $\sum_{r=1}^{n} A_{j,r} B_{r,k}$, which is the definition of $(AB)_{j,k}$.
Structural Commentary: Matrix multiplication is assembled entry-by-entry from row-column pairings.
Key Insight: Each entry of AB sees one row of A and one column of B. □

Proposition 1.3.57. *Matrix-Vector Products as Column Combinations*
If $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(m,n,\mathbb{F})$ and $\mathbf{x} = (x_1,\dots,x_n)^t \in \mathbb{F}^n$, then $A\mathbf{x} = x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n$.

Proof. The jth entry of $A\mathbf{x}$ is $\sum_{k=1}^{n} A_{j,k} x_k$, which is also the jth entry of $x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n$.
Structural Commentary: Multiplying by a column selects a linear combination of the columns of A.
Key Insight: The coordinates of $\mathbf{x}$ are precisely the column coefficients. □

Definition 1.3.58. *Column Space, Matrix Rank, and Nullity*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(m,n,\mathbb{F})$.

(i) $\mathrm{col}(A) := \mathrm{span}(\mathbf{a}_1,\dots,\mathbf{a}_n) \subseteq \mathbb{F}^m$, and $\mathrm{row}(A) := \mathrm{col}(A^t) \subseteq \mathbb{F}^n$.
(ii) $\mathrm{rank}(A) := \dim(\mathrm{col}(A))$; column rank is $\dim(\mathrm{col}(A))$, and row rank is $\dim(\mathrm{row}(A))$.
(iii) A has *full rank* if $\mathrm{rank}(A) = \min(m,n)$, *full column rank* if $\mathrm{rank}(A) = n$, and *full row rank* if $\mathrm{rank}(A) = m$.
(iv) $\ker(A) := \{\mathbf{x} \in \mathbb{F}^n : A\mathbf{x} = \mathbf{0}\}$ is the *nullspace*, and $\mathrm{nullity}(A) := \dim(\ker(A))$.

Proposition 1.3.59. *Solvability and Complete Solution Sets*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(m,n,\mathbb{F})$ and let $\mathbf{b} \in \mathbb{F}^m$. Then $A\mathbf{x} = \mathbf{b}$ is solvable if and only if $\mathbf{b} \in \mathrm{col}(A)$.
If $\mathbf{x}_p$ is one solution, then every solution is of the form $\mathbf{x} = \mathbf{x}_p + \mathbf{x}_0$, where $\mathbf{x}_0 \in \ker(A)$.

Proof.

($\Rightarrow$) If $A\mathbf{x} = \mathbf{b}$ for some $\mathbf{x} = (x_1,\dots,x_n)^T$, then $\mathbf{b} = A\mathbf{x} = \sum_{j=1}^{n} x_j \mathbf{a}_j \in \mathrm{col}(A)$.

($\Leftarrow$) If $\mathbf{b} \in \mathrm{col}(A)$, then $\mathbf{b} = \sum_{j=1}^{n} x_j \mathbf{a}_j$ for some scalars x_j. Hence, for $\mathbf{x} = (x_1,\dots,x_n)^T$, we have $A\mathbf{x} = \mathbf{b}$.

Now suppose $A\mathbf{x}_p = \mathbf{b}$. For any $\mathbf{x}$, $A\mathbf{x} = \mathbf{b} \iff A(\mathbf{x} - \mathbf{x}_p) = \mathbf{0} \iff \mathbf{x} - \mathbf{x}_p \in \ker(A)$.
Thus every solution is exactly of the form $\mathbf{x} = \mathbf{x}_p + \mathbf{x}_0$, where $\mathbf{x}_0 \in \ker(A)$.
Structural Commentary: A consistent system is an affine shift of $\ker(A)$ by any particular solution.
Key Insight: Column space decides existence; kernel decides nonuniqueness. □

Proposition 1.3.60. *Matrix Products by Columns*
If $A \in \mathrm{Mat}(m,r,\mathbb{F})$ and $B = [\mathbf{b}_1 \ \cdots \ \mathbf{b}_n] \in \mathrm{Mat}(r,n,\mathbb{F})$, then $AB = [A\mathbf{b}_1 \ \cdots \ A\mathbf{b}_n]$.

Proof. For each k, the kth column of AB has entries $\sum_{\ell=1}^{r} A_{j,\ell} B_{\ell,k}$, which are exactly the entries of $A\mathbf{b}_k$.
Structural Commentary: Matrix products apply A to columns of B.
Key Insight: The columns of AB are obtained by applying A to the columns of B. □

Proposition 1.3.61. *Matrix Products as Sums of Outer Products*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_r] \in \mathrm{Mat}(m,r,\mathbb{F})$, and write $B \in \mathrm{Mat}(r,n,\mathbb{F})$ by rows as

$$B = \begin{bmatrix} \mathbf{b}_1^t \\ \vdots \\ \mathbf{b}_r^t \end{bmatrix}.$$

Then $AB = \mathbf{a}_1\mathbf{b}_1^t + \cdots + \mathbf{a}_r\mathbf{b}_r^t$.

Proof. The (j,k) entry of $\mathbf{a}_\ell \mathbf{b}_\ell^t$ is $A_{j,\ell} B_{\ell,k}$. Summing over $\ell = 1,\dots,r$ gives $\sum_{\ell=1}^{r} A_{j,\ell} B_{\ell,k} = (AB)_{j,k}$.
Structural Commentary: Column-row products decompose a matrix product into single column-row pieces.
Key Insight: Each shared index ℓ contributes one outer product. □

Proposition 1.3.62. *Elementary Transpose Rules*
Let $A, C \in \mathrm{Mat}(m,n,\mathbb{F})$ and let $a \in \mathbb{F}$. Then $(A^t)^t = A$; $(A+C)^t = A^t + C^t$; and $(aA)^t = aA^t$.

Proof. For all admissible indices, $((A^t)^t)_{j,k} = (A^t)_{k,j} = A_{j,k}$, so $(A^t)^t = A$.
Also $((A+C)^t)_{j,k} = (A+C)_{k,j} = A_{k,j} + C_{k,j} = (A^t + C^t)_{j,k}$ and $((aA)^t)_{j,k} = (aA)_{k,j} = aA_{k,j} = (aA^t)_{j,k}$.
Structural Commentary: Transpose swaps indices; entrywise algebra is unchanged.
Key Insight: Transpose is linear and involutive. □

Example 1.3.63. *Symmetric and Skew-Symmetric Decomposition*
Assume $\mathrm{char}(\mathbb{F}) \neq 2$, $M = \mathrm{Mat}(n,n,\mathbb{F})$, let $\mathrm{Sym}_n(\mathbb{F}) = \{A \in M : A^t = A\}$, $\mathrm{Skew}_n(\mathbb{F}) = \{A \in M : A^t = -A\}$.
For the map $\tau : M \to M$ defined by $\tau(A) = A^t$, linearity follows from Proposition 1.3.62. Hence $\mathrm{Sym}_n(\mathbb{F}) = \ker(\tau - I)$ and $\mathrm{Skew}_n(\mathbb{F}) = \ker(\tau + I)$, so both are subspaces of M.
For every $A \in M$, $A = \frac{1}{2}(A + A^t) + \frac{1}{2}(A - A^t) \in \mathrm{Sym}_n(\mathbb{F}) + \mathrm{Skew}_n(\mathbb{F})$.
Also $A \in \mathrm{Sym}_n(\mathbb{F}) \cap \mathrm{Skew}_n(\mathbb{F})$ implies $A = A^t = -A$, hence $A = 0$. Thus $M = \mathrm{Sym}_n(\mathbb{F}) \oplus \mathrm{Skew}_n(\mathbb{F})$.
Structural Commentary: Transpose splits a matrix into its even and odd parts.
Key Insight: The directness of the split uses the ability to divide by 2.

Proposition 1.3.64. *Transpose Reverses Products*
If $A \in \mathrm{Mat}(m,n,\mathbb{F})$ and $B \in \mathrm{Mat}(n,p,\mathbb{F})$, then $(AB)^t = B^t A^t$.

Proof. For $1 \le j \le p$ and $1 \le k \le m$, $((AB)^t)_{j,k} = (AB)_{k,j} = \sum_{\ell=1}^{n} A_{k,\ell} B_{\ell,j} = \sum_{\ell=1}^{n} (B^t)_{j,\ell} (A^t)_{\ell,k} = (B^t A^t)_{j,k}$.
Structural Commentary: Transposition swaps the order of composition encoded by matrix multiplication.
Key Insight: Rows become columns, so the product order reverses. □

Corollary 1.3.65. *Matrix Inverses and Transposes Reverse Products*
Let $A, B \in \mathrm{Mat}(n,n,\mathbb{F})$ be invertible. Then AB, A^t are invertible, and $(AB)^{-1} = B^{-1}A^{-1}$; $(A^t)^{-1} = (A^{-1})^t$.

Proof. First, $(B^{-1}A^{-1})(AB) = B^{-1}(A^{-1}A)B = B^{-1}B = I_n$; $(AB)(B^{-1}A^{-1}) = A(BB^{-1})A^{-1} = AA^{-1} = I_n$.
Hence $(AB)^{-1} = B^{-1}A^{-1}$. By Proposition 1.3.64, $(A^{-1})^t A^t = (AA^{-1})^t = I_n$; $A^t(A^{-1})^t = (A^{-1}A)^t = I_n$.
Structural Commentary: Inversion and transposition both reverse multiplication order.
Key Insight: Apply the defining inverse equations and reverse products. □

Definition 1.3.66. *Coordinate Vector*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a basis of a finite-dimensional vector space $\mathcal{V}$.
If $\mathbf{v} = c_1\mathbf{v}_1 + \cdots + c_n\mathbf{v}_n$, then the *$\beta$-coordinate vector* of $\mathbf{v}$ is

$$[\mathbf{v}]_\beta := \begin{bmatrix} c_1 \\ \vdots \\ c_n \end{bmatrix} \in \mathbb{F}^n.$$

The function $[\cdot]_\beta : \mathcal{V} \to \mathbb{F}^n$ is the *β-basis representation function*.

Proposition 1.3.67. *Coordinate Maps Are Linear Bijections*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a basis of a finite-dimensional vector space $\mathcal{V}$.
Then the coordinate map $[\cdot]_\beta : \mathcal{V} \to \mathbb{F}^n$ is linear and bijective.

Proof. If $\mathbf{u} = a_1\mathbf{v}_1 + \cdots + a_n\mathbf{v}_n$ and $\mathbf{v} = b_1\mathbf{v}_1 + \cdots + b_n\mathbf{v}_n$, then for $c \in \mathbb{F}$,

$$[c\mathbf{u} + \mathbf{v}]_\beta = \begin{bmatrix} ca_1 + b_1 \\ \vdots \\ ca_n + b_n \end{bmatrix} = c[\mathbf{u}]_\beta + [\mathbf{v}]_\beta$$

Thus $[\cdot]_\beta$ is linear. It is injective because $[\mathbf{v}]_\beta = \mathbf{0}$ forces all basis coefficients of $\mathbf{v}$ to be zero, and it is surjective because every column $(c_1, \ldots, c_n)^t \in \mathbb{F}^n$ equals $[c_1\mathbf{v}_1 + \cdots + c_n\mathbf{v}_n]_\beta$.
Structural Commentary: A basis converts abstract vectors into coordinate columns without losing structure.
Key Insight: Coordinates give a structure-preserving copy of the vector space, not extra information. □

Definition 1.3.68. *Matrix Representation*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a basis of $\mathcal{V}$ and $\gamma = (\mathbf{w}_1, \ldots, \mathbf{w}_m)$ a basis of $\mathcal{W}$. For $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, the matrix of T relative to (β, γ) is $[T]_{\gamma \leftarrow \beta} := \begin{bmatrix} [T(\mathbf{v}_1)]_\gamma & \cdots & [T(\mathbf{v}_n)]_\gamma \end{bmatrix} \in \mathrm{Mat}(m,n,\mathbb{F})$.

Proposition 1.3.69. *Stacks Intersect Kernels, Concatenations Sum Columns*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$, $B \in \mathrm{Mat}(p,n,\mathbb{F})$, $C \in \mathrm{Mat}(m,r,\mathbb{F})$, and $D \in \mathrm{Mat}(m,s,\mathbb{F})$.
Then $\ker \begin{bmatrix} A \\ B \end{bmatrix} = \ker(A) \cap \ker(B)$; and $\mathrm{col} \begin{bmatrix} C & D \end{bmatrix} = \mathrm{col}(C) + \mathrm{col}(D)$.

Proof. For $\mathbf{x} \in \mathbb{F}^n$, $\mathbf{x} \in \ker \begin{bmatrix} A \\ B \end{bmatrix} \iff \begin{bmatrix} A \\ B \end{bmatrix} \mathbf{x} = \mathbf{0} \iff A\mathbf{x} = \mathbf{0}$ and $B\mathbf{x} = \mathbf{0} \iff \mathbf{x} \in \ker(A) \cap \ker(B)$.

For $\mathbf{y} \in \mathbb{F}^m$, $\mathbf{y} \in \operatorname{col} \begin{bmatrix} C & D \end{bmatrix} \iff \exists\, \mathbf{u} \in \mathbb{F}^r,\ \mathbf{v} \in \mathbb{F}^s$ such that $\mathbf{y} = \begin{bmatrix} C & D \end{bmatrix} \begin{bmatrix} \mathbf{u} \\ \mathbf{v} \end{bmatrix}$
$\iff \exists\, \mathbf{u}, \mathbf{v}$ such that $\mathbf{y} = C\mathbf{u} + D\mathbf{v} \iff \mathbf{y} \in \operatorname{col}(C) + \operatorname{col}(D)$.
Structural Commentary: Stacking intersects constraints; concatenation adds column directions.
Key Insight: Stacks encode intersections of nullspaces, while concatenations encode sums of column spaces. □

Proposition 1.3.70. *Matrices Respect Linear Combinations of Maps*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a basis of $\mathcal{V}$ and γ be a basis of $\mathcal{W}$.
If $S, T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ and $a \in \mathbb{F}$, then $[aS + T]_{\gamma \leftarrow \beta} = a[S]_{\gamma \leftarrow \beta} + [T]_{\gamma \leftarrow \beta}$.

Proof. The jth column of $[aS + T]_{\gamma \leftarrow \beta}$ is $[(aS + T)\mathbf{v}_j]_\gamma = [aS\mathbf{v}_j + T\mathbf{v}_j]_\gamma = a[S\mathbf{v}_j]_\gamma + [T\mathbf{v}_j]_\gamma$, which is the jth column of $a[S]_{\gamma \leftarrow \beta} + [T]_{\gamma \leftarrow \beta}$.
Structural Commentary: Matrix representation records basis images column-by-column.
Key Insight: Linear combinations of maps become entrywise linear combinations of their matrices. □

Definition 1.3.71. *Linear Map Induced by a Matrix*
For $A \in \operatorname{Mat}(m, n, \mathbb{F})$, the *linear map induced by A* is the map $T_A : \mathbb{F}^n \to \mathbb{F}^m$ defined by $T_A(\mathbf{x}) = A\mathbf{x}$.

Proposition 1.3.72. *Matrix-Induced Maps Are Linear*
If $A \in \operatorname{Mat}(m, n, \mathbb{F})$, then $T_A : \mathbb{F}^n \to \mathbb{F}^m$ is linear.

Proof. For $\mathbf{x}, \mathbf{y} \in \mathbb{F}^n$ and $a \in \mathbb{F}$, $T_A(a\mathbf{x} + \mathbf{y}) = A(a\mathbf{x} + \mathbf{y}) = aA\mathbf{x} + A\mathbf{y} = aT_A(\mathbf{x}) + T_A(\mathbf{y})$.
Structural Commentary: Matrix multiplication preserves addition and scalar multiplication.
Key Insight: Every matrix defines a canonical linear map between coordinate spaces. □

Corollary 1.3.73. *Matrix Rank-Nullity*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$. Then $\dim(\ker(A)) + \dim(\operatorname{col}(A)) = n$ and $\operatorname{nullity}(A) + \operatorname{rank}(A) = n$.
If $m = n$, then $\ker(A) = \{\mathbf{0}\}$ if and only if $\operatorname{col}(A) = \mathbb{F}^n$.

Proof. Apply Theorem 1.3.27 to the induced linear map $T_A : \mathbb{F}^n \to \mathbb{F}^m$. By definition, $\ker(T_A) = \ker(A)$ and $\operatorname{range}(T_A) = \operatorname{col}(A)$, giving the displayed identities. When $m = n$, the remaining equivalence follows as:
($\Rightarrow$) If $\ker(A) = \{\mathbf{0}\}$, then $\dim(\operatorname{col}(A)) = n$. Since $\operatorname{col}(A) \subseteq \mathbb{F}^n$, Proposition 1.1.31 gives $\operatorname{col}(A) = \mathbb{F}^n$.
($\Leftarrow$) If $\operatorname{col}(A) = \mathbb{F}^n$, then $\dim(\operatorname{col}(A)) = n$, so the identity gives $\dim(\ker(A)) = 0$. Hence $\ker(A) = \{\mathbf{0}\}$.
Structural Commentary: Matrix rank-nullity is ordinary rank-nullity written in coordinates.
Key Insight: Kernel loss plus column-space dimension equals the number of columns. □

Theorem 1.3.74. *Row Rank Equals Column Rank*
For $A \in \operatorname{Mat}(m, n, \mathbb{F})$, $\dim(\operatorname{row}(A)) = \dim(\operatorname{col}(A))$. Equivalently, $\operatorname{rank}(A^t) = \operatorname{rank}(A)$.

Proof. Let $r = \operatorname{rank}(A)$. If $r = 0$, then $\operatorname{col}(A) = \{\mathbf{0}\}$, so every column of A is zero and $A = 0$; the result follows. Suppose $r > 0$. Choose columns $\mathbf{a}_{j_1}, \ldots, \mathbf{a}_{j_r}$ forming a basis of $\operatorname{col}(A)$, and set $B = [\mathbf{a}_{j_1} \ \cdots \ \mathbf{a}_{j_r}] \in \operatorname{Mat}(m, r, \mathbb{F})$. For each column $\mathbf{a}_k$ of A, there is a unique $\mathbf{x}_k \in \mathbb{F}^r$ such that $\mathbf{a}_k = B\mathbf{x}_k$. With $X = [\mathbf{x}_1 \ \cdots \ \mathbf{x}_n] \in \operatorname{Mat}(r, n, \mathbb{F})$, we have $A = BX$. Therefore $A^t = X^t B^t$, so every column of A^t is a linear combination of the r columns of X^t. Hence $\operatorname{rank}(A^t) \leq r = \operatorname{rank}(A)$. Applying the same argument to A^t gives $\operatorname{rank}(A) \leq \operatorname{rank}(A^t)$.
Structural Commentary: Column rank controls row rank by factoring A through a basis of its column space.
Key Insight: Rows and columns measure the same intrinsic rank, viewed from transposed coordinates. □

Perspective 1.3.75. *Row and Column Spaces Have Different Hosts*
For $A \in \operatorname{Mat}(m, n, \mathbb{F})$, $\operatorname{row}(A) \subseteq \mathbb{F}^n$ and $\operatorname{col}(A) \subseteq \mathbb{F}^m$.
Hence, when $m \neq n$, Theorem 1.3.74 asserts only $\dim \operatorname{row}(A) = \dim \operatorname{col}(A)$, not equality of subspaces.

Theorem 1.3.76. *Four Fundamental Matrix Subspaces*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ and let $r = \operatorname{rank}(A)$. Then $\operatorname{row}(A), \ker(A) \subseteq \mathbb{F}^n$; $\operatorname{col}(A), \ker(A^t) \subseteq \mathbb{F}^m$, and $\dim(\operatorname{col}(A)) = \dim(\operatorname{row}(A)) = r$; $\dim(\ker(A)) = n - r$; $\dim(\ker(A^t)) = m - r$.

Proof. The inclusions follow from the ambient dimensions of the rows and columns, and from the domains of A and A^t. By definition, $\dim(\operatorname{col}(A)) = r$; by Theorem 1.3.74, $\dim(\operatorname{row}(A)) = r$.
Rank-nullity gives $\dim(\ker(A)) = n - \operatorname{rank}(A) = n - r$. Since $\operatorname{rank}(A^t) = \dim(\operatorname{col}(A^t)) = \dim(\operatorname{row}(A)) = r$, rank-nullity for $A^t : \mathbb{F}^m \to \mathbb{F}^n$ gives $\dim(\ker(A^t)) = m - \operatorname{rank}(A^t) = m - r$.
Structural Commentary: One rank determines all four fundamental dimensions.
Key Insight: Rank counts what survives; nullity counts what is killed. □

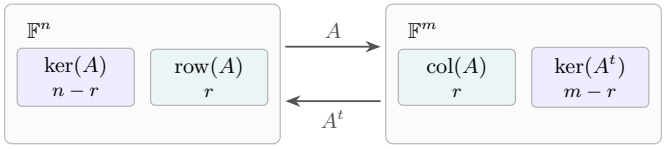


Figure 13: Four fundamental subspaces by host coordinate space.

Corollary 1.3.77. *Homogeneous Systems with More Variables Than Equations*
If $A \in \mathrm{Mat}(m,n,\mathbb{F})$ with $n > m$, then the homogeneous system $A\mathbf{x} = \mathbf{0}$ has a nonzero solution.

Proof. The induced map $T_A : \mathbb{F}^n \to \mathbb{F}^m$ has domain dimension larger than codomain dimension, so Corollary 1.3.28 shows that T_A is not injective.
Hence $\ker(A) = \ker(T_A)$ contains a nonzero vector $\mathbf{x}$, and this vector satisfies $A\mathbf{x} = \mathbf{0}$.
Structural Commentary: Extra variables force a nontrivial kernel.
Key Insight: A nontrivial kernel is exactly a nonzero homogeneous solution. □

Corollary 1.3.78. *Inhomogeneous Systems with More Equations Than Variables*
If $A \in \mathrm{Mat}(m,n,\mathbb{F})$ with $m > n$, then there exists $\mathbf{b} \in \mathbb{F}^m$ such that $A\mathbf{x} = \mathbf{b}$ has no solution.

Proof. The induced map $T_A : \mathbb{F}^n \to \mathbb{F}^m$ has domain dimension smaller than codomain dimension, so Corollary 1.3.29 shows that T_A is not surjective. Choose $\mathbf{b} \in \mathbb{F}^m \setminus \mathrm{range}(T_A)$. Then no $\mathbf{x} \in \mathbb{F}^n$ satisfies $A\mathbf{x} = \mathbf{b}$.
Structural Commentary: Too few variables cannot reach every possible right-hand side.
Key Insight: A missing range vector becomes an inconsistent inhomogeneous system. □

Definition 1.3.79. *Elementary Row Operations and Elementary Matrices*
For row symbols R_i, the *elementary row operations* are:

(i) *Row scaling*: $R_i \leftarrow cR_i$, $c \in \mathbb{F}^\times$;
(ii) *Row swap*: $R_i \leftrightarrow R_j$, $i \neq j$;
(iii) *Row replacement*: $R_i \leftarrow R_i + cR_j$, $i \neq j$, $c \in \mathbb{F}$.

An *elementary matrix* is $E = \epsilon(I_m)$, where ϵ is one elementary row operation.

Proposition 1.3.80. *Elementary Matrices Perform Row Operations*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$.

(i) If ϵ is an elementary row operation and $E_\epsilon = \epsilon(I_m)$, then $E_\epsilon A = \epsilon(A)$.
(ii) $E_\epsilon^{-1} = E_{\epsilon^{-1}}$; hence every elementary matrix is invertible with elementary inverse.
(iii) If $F = E_\epsilon^t \in \mathrm{Mat}(n,n,\mathbb{F})$, then $AF = (\epsilon(A^t))^t$.
Hence right multiplication performs the corresponding column operation.

Proof.

(i) For each row index r, $\mathrm{row}_r(E_\epsilon A) = \mathrm{row}_r(E_\epsilon)A$.
So the rows of $E_\epsilon A$ are exactly the ϵ-linear combinations of the rows of A.
(ii) Since ϵ^{-1} is elementary, $E_{\epsilon^{-1}}E_\epsilon = E_{\epsilon^{-1}}\epsilon(I_m) = \epsilon^{-1}(\epsilon(I_m)) = I_m$, and similarly $E_\epsilon E_{\epsilon^{-1}} = I_m$.
(iii) By (i) and Proposition 1.3.64, $(AF)^t = E_\epsilon A^t = \epsilon(A^t)$, so $AF = (\epsilon(A^t))^t$.

Structural Commentary: Elementary matrices encode row and column operations by multiplication.
Key Insight: Row operations are left multiplication; column operations are right multiplication. □

Definition 1.3.81. *Row Equivalence and Reduced Row Echelon Form*
For $A, B \in \mathrm{Mat}(m,n,\mathbb{F})$, write $A \sim_{\mathrm{row}} B$ if $B = E_k \cdots E_1 A$ for elementary $E_i \in \mathrm{Mat}(m,m,\mathbb{F})$ and some $k \geq 0$.
A matrix $R = [r_{ij}]$ is in *reduced row echelon form* if its nonzero rows have pivot columns $p_1, \ldots, p_s$ such that:

(i) rows $1, \ldots, s$ are nonzero and rows $s+1, \ldots, m$ are zero;
(ii) $r_{i,p_i} = 1$ and $r_{i,j} = 0$ for $j < p_i$;
(iii) $p_1 < \cdots < p_s$;
(iv) $r_{\ell,p_i} = 0$ for all $\ell \neq i$.

Theorem 1.3.82. *Reduced Row Echelon Form*
For every $A \in \mathrm{Mat}(m,n,\mathbb{F})$, there is a unique RREF matrix $\mathrm{rref}(A)$ with $A \sim_{\mathrm{row}} \mathrm{rref}(A)$.
If $R = \mathrm{rref}(A)$, then $R = E_k \cdots E_1 A$ for elementary $E_i \in \mathrm{Mat}(m,m,\mathbb{F})$.

Proof. Gaussian elimination successively chooses a pivot column, swaps/scales to a pivot 1, clears the pivot column, and continues on the remaining submatrix. Proposition 1.3.80 gives $R = E_k \cdots E_1 A$ with R in RREF. For uniqueness, row operations preserve $\ker(A) = \{\mathbf{x} : A\mathbf{x} = \mathbf{0}\}$. For an RREF matrix, the pivot set and pivot-variable formulas are determined by this kernel.
Hence two row-equivalent RREF matrices have the same pivot columns and rows.
Structural Commentary: RREF is the canonical row-operation representative of a matrix.
Key Insight: Elimination gives existence; the solution space pins down the reduced representative. □

Lemma 1.3.83. *Invertible Matrices Row-Reduce to the Identity*
If $P \in \mathrm{Mat}(m, m, \mathbb{F})$ is invertible, then $\mathrm{rref}(P) = I_m$. Hence there are elementary matrices $E_1, \ldots, E_k \in \mathrm{Mat}(m, m, \mathbb{F})$ such that $E_k \cdots E_1 P = I_m$.

Proof. Let $R = \mathrm{rref}(P)$. By Theorem 1.3.82, $R = E_k \cdots E_1 P$ for elementary matrices E_i. Since each E_i is invertible by Proposition 1.3.80, left multiplication by $E_k \cdots E_1$ preserves the kernel, so $\ker(R) = \ker(P)$. Because P is invertible, $P\mathbf{x} = \mathbf{0}$ implies $\mathbf{x} = P^{-1}P\mathbf{x} = \mathbf{0}$; hence $\ker(R) = \{\mathbf{0}\}$.
Suppose the nonzero rows of R have pivot columns $p_1 < \cdots < p_s$. If $s < m$, choose a non-pivot column j. Set the non-pivot variable $x_j = 1$, set all other non-pivot variables equal to 0, and define the pivot variables recursively from $i = s$ down to 1 so that the ith row equation is satisfied. This gives a nonzero $\mathbf{x}$ with $R\mathbf{x} = \mathbf{0}$, contradicting $\ker(R) = \{\mathbf{0}\}$. Thus $s = m$.
Since $p_1 < \cdots < p_m$ are m columns in an $m \times m$ matrix, $p_i = i$ for every i. The RREF conditions then force the ith pivot column to be the ith standard coordinate column, so $R = I_m$. □

Corollary 1.3.84. *Row Equivalence Characterization*
For $A, B \in \mathrm{Mat}(m, n, \mathbb{F})$, the following are equivalent:

(i) $A \sim_{\text{row}} B$.
(ii) $\mathrm{rref}(A) = \mathrm{rref}(B)$.
(iii) $\exists\, P \in \mathrm{Mat}(m, m, \mathbb{F})^{\times}$ *such that* $B = PA$.

Proof.

(i) $\Rightarrow$ *(ii)* Row-equivalent matrices have the same unique RREF by Theorem 1.3.82.
(ii) $\Rightarrow$ *(iii)* If $R = \mathrm{rref}(A) = \mathrm{rref}(B)$, then $R = E_k \cdots E_1 A = F_\ell \cdots F_1 B$.
Since each F_j is invertible, $B = (F_\ell \cdots F_1)^{-1} E_k \cdots E_1 A$.
(iii) $\Rightarrow$ *(i)* If $B = PA$ and $P \in \mathrm{Mat}(m, m, \mathbb{F})^{\times}$, Lemma 1.3.83 gives elementary matrices with $E_k \cdots E_1 P = I_m$.
Hence $P = E_1^{-1} \cdots E_k^{-1}$, so B is obtained from A by elementary row operations.

Structural Commentary: Row equivalence is invertible left multiplication.
Key Insight: Products of elementary matrices are the coordinate record of row reduction. □

Proposition 1.3.85. *Matrix Representation Acts on Coordinates*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a basis of $\mathcal{V}$, let γ be a basis of $\mathcal{W}$, and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.
Then for every $\mathbf{v} \in \mathcal{V}$, $[T(\mathbf{v})]_\gamma = [T]_{\gamma \leftarrow \beta} [\mathbf{v}]_\beta$.

Proof. Write $\mathbf{v} = a_1\mathbf{v}_1 + \cdots + a_n\mathbf{v}_n$. Then linearity gives $T(\mathbf{v}) = a_1 T(\mathbf{v}_1) + \cdots + a_n T(\mathbf{v}_n)$ and the right-hand side is exactly the linear combination of the columns of $[T]_{\gamma \leftarrow \beta}$ with coefficients $a_1, \ldots, a_n$.
Structural Commentary: Matrix representation turns applying T into multiplying coordinate columns.
Key Insight: The columns of $[T]_{\gamma \leftarrow \beta}$ are the coordinate images of the basis vectors. □

Theorem 1.3.86. *Columns of an Invertible Matrix Form a Basis*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(n, n, \mathbb{F})$ be invertible. Then $(\mathbf{a}_1, \ldots, \mathbf{a}_n)$ is a basis of $\mathbb{F}^n$.

Proof. The columns span $\mathbb{F}^n$ because for any $\mathbf{y} \in \mathbb{F}^n$, writing $\mathbf{x} = A^{-1}\mathbf{y}$ gives $\mathbf{y} = A\mathbf{x} = x_1\mathbf{a}_1 + \cdots + x_n\mathbf{a}_n$.
They are linearly independent because if $c_1\mathbf{a}_1 + \cdots + c_n\mathbf{a}_n = \mathbf{0}$, then $A\mathbf{c} = \mathbf{0}$, so $\mathbf{c} = A^{-1}\mathbf{0} = \mathbf{0}$.
Structural Commentary: An invertible matrix is a coordinate isomorphism of $\mathbb{F}^n$. Its columns are the images of the standard coordinate directions.
Key Insight: Invertibility forces the columns to span and prevents nontrivial column relations. □

Theorem 1.3.87. *One-Sided Inverses of Square Matrices*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$. If $AB = I_n$, then $BA = I_n$.

Proof. If $B\mathbf{x} = \mathbf{0}$, then $\mathbf{x} = AB\mathbf{x} = \mathbf{0}$, so the induced map $T_B : \mathbb{F}^n \to \mathbb{F}^n$ is injective. By Theorem 1.3.49, T_B is surjective. For any $\mathbf{y} \in \mathbb{F}^n$, choose $\mathbf{x}$ with $\mathbf{y} = B\mathbf{x}$. Then $BA\mathbf{y} = BAB\mathbf{x} = B\mathbf{x} = \mathbf{y}$.
Thus BA acts as the identity on every vector of $\mathbb{F}^n$, so $BA = I_n$.
Structural Commentary: In finite dimensions, injectivity and surjectivity coincide for square linear maps. A left inverse makes B injective, hence onto, which forces the missing inverse relation.
Key Insight: A one-sided inverse of a square matrix is automatically two-sided. □

Proposition 1.3.88. *Composition and Matrix Multiplication*
If $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ and $S \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then $[ST]_{\delta \leftarrow \alpha} = [S]_{\delta \leftarrow \beta}[T]_{\beta \leftarrow \alpha}$ for any bases α of $\mathcal{U}$, β of $\mathcal{V}$, and δ of $\mathcal{W}$.

Proof. For every $\mathbf{u} \in \mathcal{U}$, $[(ST)\mathbf{u}]_\delta = [S]_{\delta \leftarrow \beta}[T\mathbf{u}]_\beta = [S]_{\delta \leftarrow \beta}[T]_{\beta \leftarrow \alpha}[\mathbf{u}]_\alpha$.
Since coordinate columns $[\mathbf{u}]_\alpha$ range over all of $\mathbb{F}^{\dim(\mathcal{U})}$, the representing matrices are equal.
Structural Commentary: Matrix multiplication is the coordinate shadow of composition of linear maps.
Key Insight: Applying T then S becomes multiplying first by the matrix of T, then by the matrix of S. □

Theorem 1.3.89. *Linear Maps and Matrices Are the Same Data*
If $\dim(\mathcal{V}) = n$ and $\dim(\mathcal{W}) = m$, then choosing bases of $\mathcal{V}$ and $\mathcal{W}$ gives isomorphism $\mathcal{L}(\mathcal{V}, \mathcal{W}) \cong \mathrm{Mat}(m, n, \mathbb{F})$. Consequently, $\dim(\mathcal{L}(\mathcal{V}, \mathcal{W})) = mn$.

Proof. Fix bases β of $\mathcal{V}$ and γ of $\mathcal{W}$. The map $T \mapsto [T]_{\gamma \leftarrow \beta}$ is linear. It is injective because a linear map whose matrix is zero sends every basis vector of $\mathcal{V}$ to $\mathbf{0}$, hence is the zero map. It is surjective because, for any $A = [a_{ij}] \in \mathrm{Mat}(m, n, \mathbb{F})$, there is a unique linear map satisfying $T(\mathbf{v}_j) = a_{1j}\mathbf{w}_1 + \cdots + a_{mj}\mathbf{w}_m$, $j = 1, \ldots, n$.
Thus $\mathcal{L}(\mathcal{V}, \mathcal{W}) \cong \mathrm{Mat}(m, n, \mathbb{F})$, whose standard entry matrices form a basis of length mn.
Structural Commentary: Choosing bases converts every linear map into one and only one array of scalars.
Key Insight: A matrix is exactly the list of coordinate images of the domain basis vectors. □

Motivation 1.3.90. *Coordinate Encoding*
This is the key bridge between abstraction and computation. The linear map is basis-free; the matrix is what the map looks like after coordinates have been chosen.

Definition 1.3.91. *Change of Basis Matrix*
Let β and γ be two bases of the same vector space $\mathcal{V}$. The change-of-basis matrix from β-coordinates to γ-coordinates is $C_{\gamma \leftarrow \beta} := [I_\mathcal{V}]_{\gamma \leftarrow \beta}$. Thus $[\mathbf{v}]_\gamma = C_{\gamma \leftarrow \beta}[\mathbf{v}]_\beta$ for every $\mathbf{v} \in \mathcal{V}$.

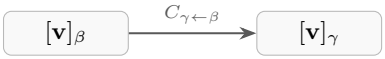


Figure 14: Change of basis: $[\mathbf{v}]_\gamma = C_{\gamma \leftarrow \beta}[\mathbf{v}]_\beta$.

Theorem 1.3.92. *Change-of-Basis Matrices*
Let $\mathcal{V}$ be n-dimensional and let β be a basis of $\mathcal{V}$.

(i) If γ is a basis of $\mathcal{V}$, then $C_{\gamma \leftarrow \beta}$ is invertible and $(C_{\gamma \leftarrow \beta})^{-1} = C_{\beta \leftarrow \gamma}$.
(ii) If $S \in \mathrm{Mat}(n, n, \mathbb{F})$ is invertible, then there is a basis γ of $\mathcal{V}$ such that $S = C_{\gamma \leftarrow \beta}$.

Proof.

(*i*) For any $\mathbf{v} \in \mathcal{V}$, $C_{\beta \leftarrow \gamma}C_{\gamma \leftarrow \beta}[\mathbf{v}]_\beta = [\mathbf{v}]_\beta$, $C_{\gamma \leftarrow \beta}C_{\beta \leftarrow \gamma}[\mathbf{v}]_\gamma = [\mathbf{v}]_\gamma$.
Since $[\cdot]_\beta$ and $[\cdot]_\gamma$ are onto $\mathbb{F}^n$, both products equal I_n. Hence $C_{\gamma \leftarrow \beta}$ is invertible with inverse $C_{\beta \leftarrow \gamma}$.
(*ii*) Write $S^{-1} = [\mathbf{s}_1 \ \cdots \ \mathbf{s}_n]$. For each j, choose the unique $\mathbf{w}_j \in \mathcal{V}$ with $[\mathbf{w}_j]_\beta = \mathbf{s}_j$, and set $\gamma = (\mathbf{w}_1, \ldots, \mathbf{w}_n)$.
Since columns of S^{-1} form a basis of $\mathbb{F}^n$ and $[\cdot]_\beta$ is an isomorphism, γ is a basis of $\mathcal{V}$.
Moreover $C_{\beta \leftarrow \gamma} = S^{-1}$, so $C_{\gamma \leftarrow \beta} = S$.

Structural Commentary: Change of basis is an invertible coordinate rewrite.
Key Insight: Every basis change is reversible, and every invertible matrix is one. □

Corollary 1.3.93. *Basis Columns Give an Invertible Matrix*
If $(\mathbf{a}_1, \ldots, \mathbf{a}_n)$ is a basis of $\mathbb{F}^n$, then $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(n, n, \mathbb{F})$ is invertible.

Proof. Let $\epsilon = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be the standard basis and $\alpha = (\mathbf{a}_1, \ldots, \mathbf{a}_n)$. Then $C_{\epsilon \leftarrow \alpha} = \begin{bmatrix}[\mathbf{a}_1]_\epsilon & \cdots & [\mathbf{a}_n]_\epsilon\end{bmatrix} = A$.
Since α is a basis, Theorem 1.3.92 gives that A is invertible.
Structural Commentary: Basis columns form a change-of-basis matrix.
Key Insight: Changing coordinates is reversible. □

Theorem 1.3.94. *Completing a Full-Column-Rank Matrix*
If $X \in \mathrm{Mat}(m, n, \mathbb{F})$ has rank $n < m$, then $\exists\ X' \in \mathrm{Mat}(m, m-n, \mathbb{F})$ s.t. $\begin{bmatrix}X & X'\end{bmatrix} \in \mathrm{Mat}(m, m, \mathbb{F})$ is invertible.

Proof. Write $X = [\mathbf{x}_1 \ \cdots \ \mathbf{x}_n]$. Since $\mathrm{rank}(X) = n$, the columns $\mathbf{x}_1, \ldots, \mathbf{x}_n$ are linearly independent in $\mathbb{F}^m$. By Theorem 1.1.27, extend this list to a basis $(\mathbf{x}_1, \ldots, \mathbf{x}_n, \mathbf{x}_{n+1}, \ldots, \mathbf{x}_m)$ of $\mathbb{F}^m$.
Set $X' = [\mathbf{x}_{n+1} \ \cdots \ \mathbf{x}_m]$. Corollary 1.3.93 shows that $\begin{bmatrix}X & X'\end{bmatrix}$ is invertible.
Structural Commentary: Full column rank means the existing columns are independent directions. Extending them to a basis turns the rectangular matrix into an invertible square matrix.
Key Insight: Independent columns can always be completed to a basis in finite dimensions. □

Corollary 1.3.95. *Invertibility and Full Rank*
For $A \in \mathrm{Mat}(n,n,\mathbb{F})$, A is invertible $\iff \mathrm{rank}(A) = n$.

Proof.
($\Rightarrow$) If A is invertible, its columns form a basis of $\mathbb{F}^n$ by Theorem 1.3.86; hence $\mathrm{rank}(A) = n$.
($\Leftarrow$) If $\mathrm{rank}(A) = n$, then $\mathrm{col}(A) = \mathbb{F}^n$.
Thus the columns of A form a basis by Proposition 1.1.32, so Corollary 1.3.93 gives that A is invertible.

Structural Commentary: Full rank is square-matrix invertibility.
Key Insight: Invertibility means the columns fill $\mathbb{F}^n$. □

Corollary 1.3.96. *Transpose Preserves Invertibility*
For $A \in \mathrm{Mat}(n,n,\mathbb{F})$, the matrix A is invertible if and only if A^t is invertible.

Proof. By Theorem 1.3.74, $\mathrm{rank}(A) = \mathrm{rank}(A^t)$. Thus $\mathrm{rank}(A) = n$ if and only if $\mathrm{rank}(A^t) = n$.
Corollary 1.3.95 gives the result.
Structural Commentary: Transposition exchanges rows and columns without changing rank.
Key Insight: A square matrix is invertible exactly when either its rows or columns have full rank. □

Theorem 1.3.97. *Full-Rank Products Preserve Rank*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$. If $X \in \mathrm{Mat}(p,m,\mathbb{F})$ has full column rank and $Y \in \mathrm{Mat}(n,q,\mathbb{F})$ has full row rank, then $\mathrm{rank}(A) = \mathrm{rank}(XAY)$. This holds when X and Y are square invertible matrices of appropriate sizes.

Proof. Because X has full column rank, the induced map $L_X : \mathbb{F}^m \to \mathbb{F}^p$, $L_X(\mathbf{u}) = X\mathbf{u}$, has kernel $\{\mathbf{0}\}$.
Thus $\mathbf{z} \in \ker(XA) \iff X(A\mathbf{z}) = \mathbf{0} \iff A\mathbf{z} = \mathbf{0} \iff \mathbf{z} \in \ker(A)$.
Since XA and A have the same number of columns, Corollary 1.3.73 gives $\mathrm{rank}(XA) = \mathrm{rank}(A)$.
Because Y has full row rank, Y^t has full column rank. Applying the left-multiplication result to A^t and Y^t, then using Theorem 1.3.74, gives $\mathrm{rank}(AY) = \mathrm{rank}((AY)^t) = \mathrm{rank}(Y^tA^t) = \mathrm{rank}(A^t) = \mathrm{rank}(A)$.
Applying the left-multiplication result once more to AY gives $\mathrm{rank}(XAY) = \mathrm{rank}(AY) = \mathrm{rank}(A)$.
Structural Commentary: Multiplication by a full-column-rank matrix on the left loses no output information, while full row rank on the right loses no input information.
Key Insight: Rank changes only when multiplication introduces a new kernel or collapses the row side. □

Theorem 1.3.98. *Change of Bases for a Linear Map*
Let $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$. Let β, β' be bases of $\mathcal{V}$ and let γ, γ' be bases of $\mathcal{W}$. Then $[T]_{\gamma' \leftarrow \beta'} = C_{\gamma' \leftarrow \gamma}[T]_{\gamma \leftarrow \beta} C_{\beta \leftarrow \beta'}$.

Proof. For every $\mathbf{v} \in \mathcal{V}$, $[T\mathbf{v}]_{\gamma'} = C_{\gamma' \leftarrow \gamma}[T\mathbf{v}]_{\gamma} = C_{\gamma' \leftarrow \gamma}[T]_{\gamma \leftarrow \beta}[\mathbf{v}]_{\beta} = C_{\gamma' \leftarrow \gamma}[T]_{\gamma \leftarrow \beta} C_{\beta \leftarrow \beta'}[\mathbf{v}]_{\beta'}$.
Since $[\cdot]_{\beta'}$ is onto, the representing matrices are equal.
Structural Commentary: Basis changes act by $C_{\gamma' \leftarrow \gamma}$ on the left and $C_{\beta \leftarrow \beta'}$ on the right.
Key Insight: Domain changes act on input coordinates; codomain changes act on output coordinates. □

Definition 1.3.99. *Matrix Equivalence*
For $A, B \in \mathrm{Mat}(m,n,\mathbb{F})$, write $A \sim_{\mathrm{eq}} B$ if $\exists\ P \in \mathrm{Mat}(m,m,\mathbb{F})^{\times}$ and $Q \in \mathrm{Mat}(n,n,\mathbb{F})^{\times}$ s.t. $B = PAQ$.

Proposition 1.3.100. *Matrix Equivalence and Elementary Operations*
The relation $\sim_{\mathrm{eq}}$ is an equivalence relation on $\mathrm{Mat}(m,n,\mathbb{F})$. Moreover, $A \sim_{\mathrm{eq}} B$ if and only if B can be obtained from A by finitely many elementary row and column operations.

Proof. Reflexivity follows from $A = I_m A I_n$. If $B = PAQ$ with P, Q invertible, then $A = P^{-1}BQ^{-1}$, so symmetry holds. If $B = P_1AQ_1$ and $C = P_2BQ_2$, then $C = (P_2P_1)A(Q_1Q_2)$, and products of invertible matrices are invertible by Corollary 1.3.65; hence transitivity holds.
Each elementary row operation is left multiplication by an invertible elementary matrix by Proposition 1.3.80. Each elementary column operation is right multiplication by the transpose of an invertible elementary matrix, again by Proposition 1.3.80 and Corollary 1.3.65. Thus any finite sequence of elementary row and column operations has the form $B = PAQ$ with P, Q invertible, so $A \sim_{\mathrm{eq}} B$.
Conversely, suppose $A \sim_{\mathrm{eq}} B$, so $B = PAQ$ with P, Q invertible. Lemma 1.3.83 gives products of elementary matrices expressing P and Q as products of elementary factors. Left multiplication by the factors of P is a sequence of elementary row operations, and right multiplication by the factors of Q is a sequence of elementary column operations. Hence B can be obtained from A by finitely many elementary row and column operations. □

Theorem 1.3.101. *Equivalent Matrices Represent the Same Linear Map*
Let $\dim \mathcal{V} = n$, $\dim \mathcal{W} = m$, and $A, B \in \mathrm{Mat}(m,n,\mathbb{F})$.
Then $A \sim_{\mathrm{eq}} B \iff \exists\, T \in \mathcal{L}(\mathcal{V},\mathcal{W})$ with $A = [T]_{\gamma \leftarrow \beta}$, $B = [T]_{\gamma' \leftarrow \beta'}$ for some bases β, β' of $\mathcal{V}$ and γ, γ' of $\mathcal{W}$.

Proof.

($\Rightarrow$) Let $B = PAQ$ with P, Q invertible. Choose bases β, γ and define $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ by $[T]_{\gamma \leftarrow \beta} = A$.
By Theorem 1.3.92, choose basis γ' with $P = C_{\gamma' \leftarrow \gamma}$, and basis β' with $Q^{-1} = C_{\beta' \leftarrow \beta}$. Then $Q = C_{\beta \leftarrow \beta'}$.
Then Theorem 1.3.98 gives $[T]_{\gamma' \leftarrow \beta'} = PAQ = B$.

($\Leftarrow$) If $A = [T]_{\gamma \leftarrow \beta}$ and $B = [T]_{\gamma' \leftarrow \beta'}$, then $B = C_{\gamma' \leftarrow \gamma} A C_{\beta \leftarrow \beta'}$, so $A \sim_{\text{eq}} B$.

Structural Commentary: Equivalence is two-sided basis change.
Key Insight: Same map; different coordinates. □

Theorem 1.3.102. *Rank Classifies Matrix Equivalence*
For $0 \le r \le \min(m, n)$, *set* $J_r^{m,n} := \begin{bmatrix} I_r & 0 \\ 0 & 0 \end{bmatrix} \in \text{Mat}(m, n, \mathbb{F})$.
For $A, B \in \text{Mat}(m, n, \mathbb{F})$, $A \sim_{\text{eq}} J_{\text{rank}(A)}^{m,n}$ *and* $A \sim_{\text{eq}} B \iff \text{rank}(A) = \text{rank}(B)$.

Proof.

($\Rightarrow$) If $B = PAQ$ with P, Q invertible, then $\text{rank}(B) = \text{rank}(A)$ by Theorem 1.3.97.

($\Leftarrow$) For any X with $\text{rank}(X) = r$, let $T_X : \mathbb{F}^n \to \mathbb{F}^m$ be $T_X(\mathbf{x}) = X\mathbf{x}$.
Choose $\mathcal{M}_X$ with $\mathbb{F}^n = \ker(T_X) \oplus \mathcal{M}_X$; then $T_X|_{\mathcal{M}_X} : \mathcal{M}_X \to \text{range}(T_X)$ is an isomorphism.
Choose a basis $(\mathbf{u}_1, \ldots, \mathbf{u}_r)$ of $\mathcal{M}_X$, extend it by a basis of $\ker(T_X)$ to β_X, and choose γ_X beginning with $T_X\mathbf{u}_1, \ldots, T_X\mathbf{u}_r$. Then $[T_X]_{\gamma_X \leftarrow \beta_X} = J_r^{m,n}$, so $X \sim_{\text{eq}} J_r^{m,n}$.
If $\text{rank}(A) = \text{rank}(B) = r$, then $A \sim_{\text{eq}} J_r^{m,n}$ and $B \sim_{\text{eq}} J_r^{m,n}$. By Proposition 1.3.100, symmetry gives $J_r^{m,n} \sim_{\text{eq}} B$, and transitivity gives $A \sim_{\text{eq}} B$.

Structural Commentary: Rank is the invariant of two-sided equivalence.
Key Insight: Adapt bases to kernel and range. □

Corollary 1.3.103. *Rank Normal Form for Linear Maps*
Let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, $\dim \mathcal{V} = n$, $\dim \mathcal{W} = m$, *and* $\text{rank}(T) = r$. *Then* $\exists$ *bases* β, γ *with* $[T]_{\gamma \leftarrow \beta} = J_r^{m,n}$.

Proof. For initial bases β_0, γ_0, let $A = [T]_{\gamma_0 \leftarrow \beta_0}$.
By Theorem 1.3.102, $A \sim_{\text{eq}} J_r^{m,n}$; Theorem 1.3.101 gives bases β, γ with $[T]_{\gamma \leftarrow \beta} = J_r^{m,n}$.
Structural Commentary: Rank gives the two-basis normal form.
Key Insight: Rank normalizes the map. □

Corollary 1.3.104. *Linear Maps Classified by Rank*
For $S, T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, *write* $S \sim_{\text{bc}} T$ *if* $\exists$ *bases* β_S, β_T *of* $\mathcal{V}$ *and* γ_S, γ_T *of* $\mathcal{W}$ *such that* $[S]_{\gamma_S \leftarrow \beta_S} = [T]_{\gamma_T \leftarrow \beta_T}$.
Then $S \sim_{\text{bc}} T \iff \text{rank}(S) = \text{rank}(T)$

Proof.

($\Rightarrow$) If $[S]_{\gamma_S \leftarrow \beta_S} = [T]_{\gamma_T \leftarrow \beta_T}$, then the two maps have the same matrix rank, hence the same rank.

($\Leftarrow$) If $\text{rank}(S) = \text{rank}(T) = r$, Corollary 1.3.103 gives bases with $[S]_{\gamma_S \leftarrow \beta_S} = J_r^{m,n} = [T]_{\gamma_T \leftarrow \beta_T}$.

Structural Commentary: Rank is the invariant of basis-change equivalence.
Key Insight: Same rank means same normal form. □

Theorem 1.3.105. *Change of Basis Formula*
Let $\mathcal{V}$ *be finite-dimensional, let* $T \in \mathcal{L}(\mathcal{V})$, *and let* β *be a basis of* $\mathcal{V}$.

(i) If γ *is a basis of* $\mathcal{V}$ *and* $S = C_{\gamma \leftarrow \beta}$, *then* $[T]_{\gamma \leftarrow \gamma} = S[T]_{\beta \leftarrow \beta} S^{-1}$.
(ii) If $S \in \text{Mat}(n, n, \mathbb{F})$ *is invertible, then there is a basis* γ *of* $\mathcal{V}$ *such that* $[T]_{\gamma \leftarrow \gamma} = S[T]_{\beta \leftarrow \beta} S^{-1}$.

Proof.

(i) Suppose γ is a basis and $S = C_{\gamma \leftarrow \beta}$. By Theorem 1.3.92, $S^{-1} = C_{\beta \leftarrow \gamma}$.
For every $\mathbf{v} \in \mathcal{V}$, $[T\mathbf{v}]_\gamma = S[T\mathbf{v}]_\beta = S[T]_{\beta \leftarrow \beta}[\mathbf{v}]_\beta = S[T]_{\beta \leftarrow \beta} S^{-1}[\mathbf{v}]_\gamma$.
Since $[\cdot]_\gamma$ is onto $\mathbb{F}^n$, the representing matrices are equal.

(ii) If S is invertible, Theorem 1.3.92 supplies a basis γ with $S = C_{\gamma \leftarrow \beta}$, and part (i) applies.

Structural Commentary: Changing basis conjugates operator matrices.
Key Insight: Similarity is a coordinate rewrite. □

Definition 1.3.106. *Similarity*
Matrices $A, B \in \text{Mat}(n, n, \mathbb{F})$ are similar if there exists an invertible matrix P such that $B = PAP^{-1}$.

Proposition 1.3.107. *Similarity Is an Equivalence Relation*
Similarity is an equivalence relation on $\text{Mat}(n, n, \mathbb{F})$.

Proof. For any A, $A = I_n A I_n^{-1}$, so similarity is reflexive. If $B = PAP^{-1}$, then $A = P^{-1}BP$, so it is symmetric. If $B = PAP^{-1}$ and $C = QBQ^{-1}$, then $C = Q(PAP^{-1})Q^{-1} = (QP)A(QP)^{-1}$ so it is transitive.
Structural Commentary: Similarity groups matrices by change of basis.
Key Insight: Identity, inverse, and composition give the three equivalence properties. □

Perspective 1.3.108. *Similarity as Same Operator, New Coordinates*
Similarity does not mean two operators are merely related; it means they are the same operator written in two different bases. This is the structural content behind change of basis.

Remark 1.3.109. *Similarity Convention*
The forms PAP^{-1} and $M^{-1}AM$ are the same convention with $P = M^{-1}$. In Theorem 1.3.105, this means choosing opposite coordinate-change directions.

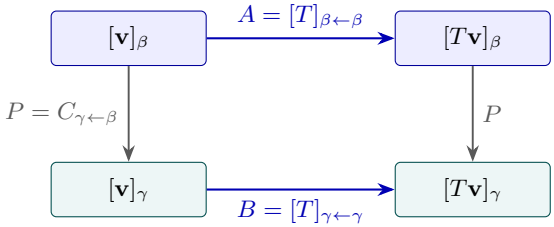


Figure 15: One operator in two bases: $PA = BP$, so $B = PAP^{-1}$.

Corollary 1.3.110. *Similarity as One Operator in Two Bases*
For $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$, the following are equivalent:

(i) A and B are similar.
(ii) There are an n-dim vector space $\mathcal{V}$, bases β, γ of $\mathcal{V}$, and $T \in \mathcal{L}(\mathcal{V})$ such that $A = [T]_{\beta\leftarrow\beta}$ and $B = [T]_{\gamma\leftarrow\gamma}$.

Proof.

(*i*) $\Rightarrow$ (*ii*): Suppose $B = PAP^{-1}$ with P invertible.
Let $\mathcal{V} = \mathbb{F}^n$, let β be the standard basis, and let $T_A(\mathbf{x}) = A\mathbf{x}$. By Theorem 1.3.92, choose γ with $P = C_{\gamma\leftarrow\beta}$. Then $[T_A]_{\beta\leftarrow\beta} = A$, and Theorem 1.3.105 gives $[T_A]_{\gamma\leftarrow\gamma} = B$.
(*ii*) $\Rightarrow$ (*i*): If $A = [T]_{\beta\leftarrow\beta}$ and $B = [T]_{\gamma\leftarrow\gamma}$, then $B = C_{\gamma\leftarrow\beta} A C_{\beta\leftarrow\gamma}$ with $C_{\gamma\leftarrow\beta}$ invertible.

Structural Commentary: Similar matrices are one operator in different bases.
Key Insight: Conjugation records a change of basis. □

Definition 1.3.111. *Similarity of Operators*
For $T \in \mathcal{L}(\mathcal{V})$ and $S \in \mathcal{L}(\mathcal{W})$, $T \sim S \iff \exists\, \Phi : \mathcal{V} \xrightarrow{\sim} \mathcal{W}$ such that $S = \Phi T \Phi^{-1}$.
If $\mathcal{V} = \mathcal{W}$, then $\Phi \in \mathrm{Aut}(\mathcal{V})$ and similarity is conjugacy in $\mathcal{L}(\mathcal{V})$.

Proposition 1.3.112. *Operator Similarity and Common Matrix Representatives*
Let $\dim \mathcal{V} = \dim \mathcal{W} = n$, $T \in \mathcal{L}(\mathcal{V})$, and $S \in \mathcal{L}(\mathcal{W})$.
Then $T \sim S \iff \exists\, A \in \mathrm{Mat}(n, n, \mathbb{F})$ with $A = [T]_{\beta\leftarrow\beta} = [S]_{\gamma\leftarrow\gamma}$ for some bases β of $\mathcal{V}$ and γ of $\mathcal{W}$.

Proof.

($\Rightarrow$) Let $S = \Phi T \Phi^{-1}$ and $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$. Put $\gamma = (\Phi\mathbf{v}_1, \ldots, \Phi\mathbf{v}_n)$.
Then γ is a basis and, for all j, $S(\Phi\mathbf{v}_j) = \Phi T \mathbf{v}_j$, so $[S]_{\gamma\leftarrow\gamma} = [T]_{\beta\leftarrow\beta}$.
($\Leftarrow$) Suppose $[T]_{\beta\leftarrow\beta} = [S]_{\gamma\leftarrow\gamma}$. Define $\Phi : \mathcal{V} \to \mathcal{W}$ by $\Phi(\beta_j) = \gamma_j$.
Then Φ is an isomorphism and $\Phi T \beta_j = S \Phi \beta_j$ for every j, hence $\Phi T = S\Phi$. Thus $S = \Phi T \Phi^{-1}$.

Structural Commentary: Similarity is equality of matrices after identifying bases.
Key Insight: A common matrix is an operator-level conjugacy written in bases. □

Theorem 1.3.113. *Similarity and Scalar Shifts*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$.

(i) If A and B are similar, then $A - \lambda I_n$ and $B - \lambda I_n$ are similar for every $\lambda \in \mathbb{F}$.
(ii) If $A - \lambda I_n$ and $B - \lambda I_n$ are similar for some $\lambda \in \mathbb{F}$, then A and B are similar.

Proof.

(i) If $B = PAP^{-1}$ with P invertible, then $B - \lambda I_n = PAP^{-1} - \lambda PP^{-1} = P(A - \lambda I_n)P^{-1}$.
(ii) If $B - \lambda I_n = P(A - \lambda I_n)P^{-1}$ with P invertible, then adding $\lambda I_n = \lambda PP^{-1}$ gives $B = PAP^{-1}$.

Structural Commentary: Scalar shifts commute with change of basis.
Key Insight: The identity matrix is basis-independent. □

1.4 Block Matrices and Matrix Operations

Definition 1.4.1. *Block Matrices*
Suppose $\mathcal{V} = \mathcal{U}_1 \oplus \mathcal{U}_2$ and $\mathcal{W} = \mathcal{Z}_1 \oplus \mathcal{Z}_2$. If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, relative to bases adapted to these decompositions in the sense of Definition 1.2.22, the matrix of T has block form $[T] = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}$, where each block records the part of T mapping one summand of the domain into one summand of the codomain.

Definition 1.4.2. *Conformal Block Partitions*
For a block matrix $A = [A_{ij}]$, the submatrices A_{ij} are called *blocks*. Two block partitions are *conformal for addition* when corresponding blocks have the same sizes. They are *conformal for multiplication* when every block product appearing in the usual row-column block multiplication is defined.

Proposition 1.4.3. *Block Matrix Arithmetic*
For conformal 2×2 block matrices,

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} + \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} A_{11} + B_{11} & A_{12} + B_{12} \\ A_{21} + B_{21} & A_{22} + B_{22} \end{bmatrix}$$

and

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{bmatrix}.$$

Proof. Each displayed block consists of the same entries obtained by ordinary entrywise addition or ordinary row-column multiplication, with the common indices grouped according to the partition.
Structural Commentary: This is ordinary matrix arithmetic with indices grouped into rectangular packages.
Key Insight: Conformality is exactly the condition that the grouped row-column products make sense. □

Proposition 1.4.4. *Block Centrosymmetry and Commutation*
Let $M = \begin{bmatrix} A & B \\ C & D \end{bmatrix} \in \mathrm{Mat}(2n, 2n, \mathbb{F})$, $L_{2n} = \begin{bmatrix} 0 & I_n \\ I_n & 0 \end{bmatrix}$.
Then M has block centrosymmetric form $A = D$ and $B = C$ if and only if $L_{2n}M = ML_{2n}$.

Proof. Block multiplication gives $L_{2n}M = \begin{bmatrix} C & D \\ A & B \end{bmatrix}$ and $ML_{2n} = \begin{bmatrix} B & A \\ D & C \end{bmatrix}$.
The two matrices are equal exactly when $A = D$ and $B = C$.
Structural Commentary: The matrix L_{2n} swaps the two block coordinates. Commuting with that swap means M treats the two block positions symmetrically.
Key Insight: Compute both block products; equality forces matching opposite blocks. □

Definition 1.4.5. *Direct Sum of Matrices*
For $A \in \mathrm{Mat}(m, n, \mathbb{F})$ and $B \in \mathrm{Mat}(p, q, \mathbb{F})$, the *direct sum* of A and B is $A \oplus B := \begin{bmatrix} A & 0 \\ 0 & B \end{bmatrix} \in \mathrm{Mat}(m+p, n+q, \mathbb{F})$.

Definition 1.4.6. *Polynomial Evaluation on a Square Matrix*
If $A \in \mathrm{Mat}(n, n, \mathbb{F})$ and $p(z) = \sum_{j=0}^{d} c_j z^j \in \mathbb{F}[z]$, then

$$p(A) := \sum_{j=0}^{d} c_j A^j, \qquad A^0 := I_n.$$

Proposition 1.4.7. *Algebra of Matrix Direct Sums*
Whenever the expressions are defined,

(i) $a(A \oplus B) + b(C \oplus D) = (aA + bC) \oplus (aB + bD)$;
(ii) $(A \oplus B)(C \oplus D) = AC \oplus BD$;
(iii) if A and B are square, then $(A \oplus B)^k = A^k \oplus B^k$ for every positive integer k;
(iv) if A and B are square and p is a polynomial, then $p(A \oplus B) = p(A) \oplus p(B)$.

Proof. Block addition gives $a(A \oplus B) + b(C \oplus D) = \begin{bmatrix} aA + bC & 0 \\ 0 & aB + bD \end{bmatrix} = (aA + bC) \oplus (aB + bD)$.
Block multiplication gives $(A \oplus B)(C \oplus D) = \begin{bmatrix} AC & 0 \\ 0 & BD \end{bmatrix} = AC \oplus BD$.

For powers, the case $k = 1$ is the definition. If $(A \oplus B)^k = A^k \oplus B^k$, then part (ii) gives $(A \oplus B)^{k+1} = A^{k+1} \oplus B^{k+1}$. Hence (iii) follows by induction. For the polynomial statement, suppose $A \in \mathrm{Mat}(m, m, \mathbb{F})$ and $B \in \mathrm{Mat}(r, r, \mathbb{F})$. For $j = 0$, Definition 1.4.6 gives $(A \oplus B)^0 = I_{m+r} = I_m \oplus I_r = A^0 \oplus B^0$. If $p(z) = \sum_{j=0}^{d} c_j z^j$,

then $p(A \oplus B) = \sum_{j=0}^{d} c_j (A^j \oplus B^j) = \left(\sum_{j=0}^{d} c_j A^j \right) \oplus \left(\sum_{j=0}^{d} c_j B^j \right) = p(A) \oplus p(B)$.

Structural Commentary: A direct sum matrix keeps the two blocks independent under linear combinations, products, powers, and polynomials.
Key Insight: Block diagonal matrices behave componentwise. □

Proposition 1.4.8. *Diagonal Multiplication Scales Rows and Columns*
Let $A = [a_{ij}] \in \mathrm{Mat}(m, n, \mathbb{F})$, let $\Lambda = \mathrm{diag}(\lambda_1, \ldots, \lambda_n)$, and let $M = \mathrm{diag}(\mu_1, \ldots, \mu_m)$. Then $A\Lambda = [\lambda_j a_{ij}]$ and $MA = [\mu_i a_{ij}]$. Thus right multiplication by Λ scales columns of A, left multiplication by M scales rows of A.

Proof. The (i, j) entry of $A\Lambda$ is $\sum_{r=1}^{n} a_{ir}\Lambda_{rj} = a_{ij}\lambda_j$, and the (i, j) entry of MA is $\sum_{r=1}^{m} M_{ir}a_{rj} = \mu_i a_{ij}$.
Structural Commentary: Diagonal matrices act by independent scalar weights on coordinate directions.
Key Insight: Multiplication on the right scales columns, while multiplication on the left scales rows. □

Proposition 1.4.9. *Inverse of an Upper Block-Triangular Matrix*

Let $Y \in \mathrm{Mat}(m, m, \mathbb{F})$ and $Z \in \mathrm{Mat}(n, n, \mathbb{F})$ be invertible, and let $X \in \mathrm{Mat}(m, n, \mathbb{F})$. Then $\begin{bmatrix} Y & X \\ 0 & Z \end{bmatrix}^{-1} = \begin{bmatrix} Y^{-1} & -Y^{-1}XZ^{-1} \\ 0 & Z^{-1} \end{bmatrix}$.

Proof. Let $M = \begin{bmatrix} Y & X \\ 0 & Z \end{bmatrix}$ and $N = \begin{bmatrix} Y^{-1} & -Y^{-1}XZ^{-1} \\ 0 & Z^{-1} \end{bmatrix}$. Proposition 1.4.3 gives

$$MN = \begin{bmatrix} I_m & -XZ^{-1} + XZ^{-1} \\ 0 & I_n \end{bmatrix} = I_{m+n}, \qquad NM = \begin{bmatrix} I_m & Y^{-1}X - Y^{-1}X \\ 0 & I_n \end{bmatrix} = I_{m+n}$$

Thus $N = M^{-1}$.
Structural Commentary: Upper block-triangular inversion preserves the zero lower-left block. The off-diagonal term is forced by cancelling the original upper-right block.
Key Insight: Inverting a block triangular matrix inverts diagonal blocks and corrects the coupling block. □

Theorem 1.4.10. *Inverse of an Upper Triangular Matrix*
Let $A = [a_{ij}] \in \mathrm{Mat}(n, n, \mathbb{F})$ be upper triangular with nonzero diagonal entries. Then A is invertible, A^{-1} is upper triangular, and the diagonal entries of A^{-1} are $a_{11}^{-1}, \ldots, a_{nn}^{-1}$.

Proof. We argue by induction on n. The case $n = 1$ is immediate. For $n > 1$, write $A = \begin{bmatrix} B & \mathbf{x} \\ 0 & a_{nn} \end{bmatrix}$, where B is $(n-1) \times (n-1)$ upper triangular with nonzero diagonal entries.
By induction, B^{-1} is upper triangular with diagonal entries $a_{11}^{-1}, \ldots, a_{n-1,n-1}^{-1}$.
Proposition 1.4.9 gives $A^{-1} = \begin{bmatrix} B^{-1} & -B^{-1}\mathbf{x}a_{nn}^{-1} \\ 0 & a_{nn}^{-1} \end{bmatrix}$, which is upper triangular and has the asserted diagonal.
Structural Commentary: An upper triangular matrix splits into a smaller upper triangular block and one final scalar block. The block inverse formula preserves this triangular structure.
Key Insight: Nonzero diagonals make blockwise triangular inversion possible. □

Proposition 1.4.11. *Upper Block-Triangular Similarity Shift*
Let $B \in \mathrm{Mat}(m, m, \mathbb{F})$, let $D \in \mathrm{Mat}(n, n, \mathbb{F})$, and let $C, X \in \mathrm{Mat}(m, n, \mathbb{F})$.
Then $\begin{bmatrix} B & C \\ 0 & D \end{bmatrix}$ *is similar to* $\begin{bmatrix} B & C + XD - BX \\ 0 & D \end{bmatrix}$.

Proof. The matrices $\begin{bmatrix} I_m & X \\ 0 & I_n \end{bmatrix}$ and $\begin{bmatrix} I_m & -X \\ 0 & I_n \end{bmatrix}$ are inverses in $\mathrm{Mat}(m + n, m + n, \mathbb{F})$, and block multiplication gives

$$\begin{bmatrix} I_m & X \\ 0 & I_n \end{bmatrix} \begin{bmatrix} B & C \\ 0 & D \end{bmatrix} \begin{bmatrix} I_m & -X \\ 0 & I_n \end{bmatrix} = \begin{bmatrix} B & C + XD - BX \\ 0 & D \end{bmatrix}.$$

Structural Commentary: Unit upper block-triangular conjugation changes only the coupling block.
Key Insight: Similarity can adjust the connecting block while preserving the diagonal blocks. □

1.5 Duality and Multilinear Structure

Duality studies vectors through scalar-valued tests. Functionals, dual maps, annihilators, transposes, bilinear forms, and tensor products extend linear structure into multilinear language.

Definition 1.5.1. *Linear Functional and Dual Space*
A *linear functional* on $\mathcal{V}$ is a linear map from $\mathcal{V}$ to $\mathbb{F}$. The *dual space* of $\mathcal{V}$ is $\mathcal{V}' := \mathcal{L}(\mathcal{V}, \mathbb{F})$.

Example 1.5.2. *Evaluation and Integral Functionals*
The evaluation map $\mathrm{ev}_0 : \mathbb{F}[x] \to \mathbb{F}$, $p \mapsto p(0)$ is a linear functional, since $\mathrm{ev}_0(\alpha p + \beta q) = \alpha p(0) + \beta q(0)$ for all $\alpha, \beta \in \mathbb{F}$ and $p, q \in \mathbb{F}[x]$. If $C[a,b]$ is the real vector space of continuous functions on $[a,b]$, then $I : C[a,b] \to \mathbb{R}$, $I(f) := \int_a^b f(x)\,dx$ is a linear functional, since $I(\alpha f + \beta g) = \alpha I(f) + \beta I(g)$ $\forall \alpha, \beta \in \mathbb{R}$ and $f, g \in C[a,b]$.

Definition 1.5.3. *Dual Basis*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$ be a basis of $\mathcal{V}$. The *dual basis* of β is the list $\beta' = (\varphi_1, \dots, \varphi_n)$ in $\mathcal{V}'$ defined by

$$\varphi_j(\mathbf{v}_k) = \begin{cases} 1, & k = j, \\ 0, & k \neq j. \end{cases}$$

Existence and uniqueness of each φ_j follow from Proposition 1.3.11.

Theorem 1.5.4. *Dual Basis and Dimension of the Dual Space*
Let $\mathcal{V}$ be finite-dimensional with basis $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$, and let $\beta' = (\varphi_1, \dots, \varphi_n)$ be the dual basis. Then β' is a basis of $\mathcal{V}'$. In particular, $\dim(\mathcal{V}') = \dim(\mathcal{V})$.

Proof. First suppose $a_1\varphi_1 + \cdots + a_n\varphi_n = 0$. Evaluating at $\mathbf{v}_j$ gives $a_j = 0$ for each j, so β' is linearly independent. Now let $\varphi \in \mathcal{V}'$. For $\mathbf{v} = c_1\mathbf{v}_1 + \cdots + c_n\mathbf{v}_n$, $\left(\sum_{j=1}^{n} \varphi(\mathbf{v}_j)\varphi_j\right)(\mathbf{v}) = \sum_{j=1}^{n} \varphi(\mathbf{v}_j)c_j = \varphi(\mathbf{v})$.
Thus $\varphi = \sum_{j=1}^{n} \varphi(\mathbf{v}_j)\varphi_j$, so β' spans $\mathcal{V}'$.
Structural Commentary: A basis of $\mathcal{V}$ induces the dual coordinate system.
Key Insight: A linear functional is determined exactly by its values on a basis. □

Proposition 1.5.5. *Coordinate Functionals for a Hamel Basis*
Let $\mathcal{B} = \{\mathbf{b}_i\}_{i \in I}$ be a Hamel basis of $\mathcal{V}$, define $\varepsilon_i \in \mathcal{V}'$ by $\varepsilon_i\left(\sum_{j \in F} a_j \mathbf{b}_j\right) := a_i$, $F \subseteq I$ finite, $a_i = 0$ if $i \notin F$.
Then $\{\varepsilon_i\}_{i \in I}$ is linearly independent in $\mathcal{V}'$. Moreover, the following are equivalent:

(i) $\mathrm{span}\{\varepsilon_i : i \in I\} = \mathcal{V}'$;
(ii) $|I| < \infty$;
(iii) $\dim(\mathcal{V}) < \infty$.

Proof. If $\sum_{r=1}^{n} c_r \varepsilon_{i_r} = 0$ with distinct i_r, then $0 = \left(\sum_{r=1}^{n} c_r \varepsilon_{i_r}\right)(\mathbf{b}_{i_s}) = c_s$ $\forall s$. Hence $\{\varepsilon_i\}_{i \in I}$ linearly independent.

(*i*) $\Rightarrow$ (*ii*). Suppose $\mathrm{span}\{\varepsilon_i : i \in I\} = \mathcal{V}'$. Put $\mathrm{supp}_{\mathcal{B}}(\psi) := \{i \in I : \psi(\mathbf{b}_i) \neq 0\}$.
If I were infinite, define $\varphi \in \mathcal{V}'$ by $\varphi(\mathbf{b}_i) = 1$ for all $i \in I$.
Then $\mathrm{supp}_{\mathcal{B}}(\varphi) = I$, while $\mathrm{supp}_{\mathcal{B}}(\psi)$ is finite for every $\psi \in \mathrm{span}\{\varepsilon_i : i \in I\}$, a contradiction.
(*ii*) $\Rightarrow$ (*iii*). Since $\mathcal{B}$ is a basis, $\dim(\mathcal{V}) = |I| < \infty$.
(*iii*) $\Rightarrow$ (*i*). If $\dim(\mathcal{V}) < \infty$, then Theorem 1.5.4 gives $\mathrm{span}\{\varepsilon_i : i \in I\} = \mathcal{V}'$.

Structural Commentary: Coordinate spans are finite-support.
Key Insight: Infinite-dimensional coordinate functionals are independent, not spanning. □

Theorem 1.5.6. *Dimension of the Algebraic Dual*
For every $\mathbb{F}$-vector space $\mathcal{V}$, $\dim(\mathcal{V}) \leq \dim(\mathcal{V}')$. *Moreover, the following are equivalent:*

(i) $\dim(\mathcal{V}) = \dim(\mathcal{V}')$;
(ii) $\dim(\mathcal{V}) < \infty$.

Proof. Let $\mathcal{B} = \{\mathbf{b}_i\}_{i \in I}$ be a Hamel basis. By Proposition 1.5.5, $\{\varepsilon_i\}_{i \in I} \subseteq \mathcal{V}'$ is linearly independent; hence $\dim(\mathcal{V}) = |I| \leq \dim(\mathcal{V}')$. For the equality criterion:

(*i*) $\Rightarrow$ (*ii*). Suppose $\dim(\mathcal{V}) = \dim(\mathcal{V}')$. If I were infinite, then $\mathcal{V} \cong (\mathbb{F}^I)_0$ and $\mathcal{V}' \cong \mathbb{F}^I$.
Let $\mathbb{F}_0 \subseteq \mathbb{F}$ be the prime subfield. Since $\mathbb{F}_0$ is countable, $\dim_{\mathbb{F}_0}\left((\mathbb{F}_0^I)_0\right) = |I|$ and $\left|\mathbb{F}_0^I\right| > |I|$ by Cantor, so $\dim_{\mathbb{F}_0}(\mathbb{F}_0^I) > |I|$; hence $\dim_{\mathbb{F}_0}\left((\mathbb{F}_0^I)_0\right) = |I| < \dim_{\mathbb{F}_0}(\mathbb{F}_0^I)$.
Every $\mathbb{F}_0$-linearly independent family in $\mathbb{F}_0^I$ is $\mathbb{F}$-linearly independent in $\mathbb{F}^I$, by expanding coefficients in an $\mathbb{F}_0$-basis of $\mathbb{F}$. Thus $\dim_{\mathbb{F}}(\mathcal{V}) = \dim_{\mathbb{F}}\left((\mathbb{F}^I)_0\right) < \dim_{\mathbb{F}}(\mathbb{F}^I) = \dim_{\mathbb{F}}(\mathcal{V}')$, a contradiction. Hence $\dim(\mathcal{V}) < \infty$.

(*ii*) $\Rightarrow$ (*i*). If $\dim(\mathcal{V}) < \infty$, then Theorem 1.5.4 gives $\dim(\mathcal{V}) = \dim(\mathcal{V}')$.

Structural Commentary: Algebraic duals have arbitrary support.
Key Insight: Infinite algebraic duals are bigger than their original spaces. □

Proposition 1.5.7. *Functionals Separate Points*
Let $\mathcal{V}$ be a vector space.

(i) If $\mathbf{v} \neq \mathbf{0}$, then there exists $\varphi \in \mathcal{V}'$ such that $\varphi(\mathbf{v}) \neq 0$.
(ii) A vector $\mathbf{v} \in \mathcal{V}$ is zero if and only if $\varphi(\mathbf{v}) = 0$ for every $\varphi \in \mathcal{V}'$.
(iii) If $\mathcal{U} \leq \mathcal{V}$ and $\mathbf{x} \notin \mathcal{U}$, then there exists $\varphi \in \mathcal{V}'$ such that $\varphi|_{\mathcal{U}} = 0$ and $\varphi(\mathbf{x}) = 1$.
(iv) If $\varphi \in \mathcal{V}'$ and $\varphi(\mathbf{x}) \neq 0$, then $\mathcal{V} = \operatorname{span}(\mathbf{x}) \oplus \ker(\varphi)$.
(v) If nonzero $\varphi, \psi \in \mathcal{V}'$ have the same kernel, then $\varphi = a\psi$ for some nonzero $a \in \mathbb{F}$.

Proof.

(i) Extend $\{\mathbf{v}\}$ to a Hamel basis $\mathcal{B}$ of $\mathcal{V}$ by Theorem 1.1.35. Define φ on $\mathcal{B}$ by $\varphi(\mathbf{v}) = 1$ and $\varphi(\mathbf{b}) = 0$ for $\mathbf{b} \in \mathcal{B} \setminus \{\mathbf{v}\}$, then extend linearly to finite basis expansions.
(ii) If $\mathbf{v} = \mathbf{0}$, then $\varphi(\mathbf{v}) = 0$ for every $\varphi \in \mathcal{V}'$. Conversely, if $\varphi(\mathbf{v}) = 0$ for every $\varphi \in \mathcal{V}'$ and $\mathbf{v} \neq \mathbf{0}$, then (i) gives some $\psi \in \mathcal{V}'$ with $\psi(\mathbf{v}) \neq 0$, a contradiction.
(iii) Choose a Hamel basis $\mathcal{B}_{\mathcal{U}}$ of $\mathcal{U}$. Since $\mathbf{x} \notin \mathcal{U} = \operatorname{span}(\mathcal{B}_{\mathcal{U}})$, the set $\mathcal{B}_{\mathcal{U}} \cup \{\mathbf{x}\}$ is linearly independent. Extend it to a Hamel basis $\mathcal{B}$ of $\mathcal{V}$ by Theorem 1.1.35. Define φ on $\mathcal{B}$ by $\varphi(\mathbf{x}) = 1$ and $\varphi(\mathbf{b}) = 0$ for every $\mathbf{b} \in \mathcal{B} \setminus \{\mathbf{x}\}$, then extend linearly to finite basis expansions.
(iv) Since $\varphi(\mathbf{x}) \neq 0$, $\operatorname{span}(\mathbf{x}) \cap \ker(\varphi) = \{\mathbf{0}\}$.
For any $\mathbf{v} \in \mathcal{V}$, $\mathbf{v} = \frac{\varphi(\mathbf{v})}{\varphi(\mathbf{x})}\mathbf{x} + \left(\mathbf{v} - \frac{\varphi(\mathbf{v})}{\varphi(\mathbf{x})}\mathbf{x}\right)$, and the second summand lies in $\ker(\varphi)$.
(v) Let $\mathcal{K} = \ker(\varphi) = \ker(\psi)$ and choose $\mathbf{x} \notin \mathcal{K}$. By (iv), every $\mathbf{v} \in \mathcal{V}$ can be written $\mathbf{v} = c\mathbf{x} + \mathbf{k}$ with $\mathbf{k} \in \mathcal{K}$. Setting $a = \varphi(\mathbf{x})/\psi(\mathbf{x})$ gives $\varphi(\mathbf{v}) = c\varphi(\mathbf{x}) = a\,c\psi(\mathbf{x}) = a\psi(\mathbf{v})$.

Structural Commentary: Dual vectors separate points and define hyperplanes.
Key Insight: A nonzero functional splits the space into one visible direction plus its kernel. □

Corollary 1.5.8. *Kernels of Nonzero Functionals*
If $0 \neq \varphi \in \mathcal{V}'$, then φ is surjective onto $\mathbb{F}$, $\mathcal{V}/\ker(\varphi) \cong \mathbb{F}$, and $\operatorname{codim}_{\mathcal{V}}(\ker\varphi) = 1$.

Proof. Choose $\mathbf{x}$ with $\varphi(\mathbf{x}) \neq 0$. For each $a \in \mathbb{F}$, $\varphi\big(a\varphi(\mathbf{x})^{-1}\mathbf{x}\big) = a$, so $\operatorname{im}\varphi = \mathbb{F}$.
Hence, by Theorem 1.3.41, $\mathcal{V}/\ker(\varphi) \cong \operatorname{im}\varphi = \mathbb{F}$ and $\operatorname{codim}_{\mathcal{V}}(\ker\varphi) = \dim(\mathcal{V}/\ker\varphi) = 1$.
Structural Commentary: A nonzero functional sees exactly one scalar direction and kills a hyperplane.
Key Insight: Kernels of nonzero functionals have codimension one. □

Definition 1.5.9. *Algebraic Double Dual and Natural Map*

(i) The *algebraic double dual* of $\mathcal{V}$ is $\mathcal{V}'' := (\mathcal{V}')'$.
(ii) The *natural evaluation map* is $\Lambda_{\mathcal{V}} : \mathcal{V} \to \mathcal{V}''$; $\Lambda_{\mathcal{V}}(\mathbf{v})(\varphi) = \varphi(\mathbf{v})$.
(iii) $\mathcal{V}$ is *algebraically reflexive* if $\Lambda_{\mathcal{V}}$ is surjective.

Theorem 1.5.10. *Algebraic Reflexivity*
Natural map $\Lambda_{\mathcal{V}} \in \mathcal{L}(\mathcal{V}, \mathcal{V}'')$ is linear and injective. It is an isomorphism if and only if $\mathcal{V}$ is finite-dimensional.

Proof. For $\mathbf{v} \in \mathcal{V}$, $\Lambda_{\mathcal{V}}(\mathbf{v}) : \varphi \mapsto \varphi(\mathbf{v})$ is linear on $\mathcal{V}'$.
For $a, b \in \mathbb{F}$ and $\mathbf{u}, \mathbf{v} \in \mathcal{V}$, $\Lambda_{\mathcal{V}}(a\mathbf{u} + b\mathbf{v})(\varphi) = a\Lambda_{\mathcal{V}}(\mathbf{u})(\varphi) + b\Lambda_{\mathcal{V}}(\mathbf{v})(\varphi)$, $\varphi \in \mathcal{V}'$, so $\Lambda_{\mathcal{V}} \in \mathcal{L}(\mathcal{V}, \mathcal{V}'')$.
If $\mathbf{v} \in \ker(\Lambda_{\mathcal{V}})$, then $\varphi(\mathbf{v}) = 0$ for all $\varphi \in \mathcal{V}'$, hence $\mathbf{v} = \mathbf{0}$ by Proposition 1.5.7. Thus $\ker(\Lambda_{\mathcal{V}}) = \{\mathbf{0}\}$.
For the isomorphism criterion:

($\Rightarrow$) If $\Lambda_{\mathcal{V}}$ is an isomorphism and $\dim(\mathcal{V}) = \infty$, then Theorem 1.5.6 gives $\dim(\mathcal{V}'') > \dim(\mathcal{V}') > \dim(\mathcal{V})$, contradicting $\mathcal{V} \cong \mathcal{V}''$. Hence $\dim(\mathcal{V}) < \infty$.
($\Leftarrow$) If $\dim(\mathcal{V}) < \infty$, then $\dim(\mathcal{V}'') = \dim(\mathcal{V}') = \dim(\mathcal{V})$ by Theorem 1.5.4.
Since $\Lambda_{\mathcal{V}}$ is injective, it is an isomorphism.

Structural Commentary: $\Lambda_{\mathcal{V}}$ embeds vectors as evaluation functionals.
Key Insight: Finite-dimensional spaces are exactly the algebraically reflexive vector spaces. □

Definition 1.5.11. *Dual Map*
If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, the *dual map* of T is the map $T' : \mathcal{W}' \to \mathcal{V}'$ defined by $T'\varphi := \varphi T$ for every $\varphi \in \mathcal{W}'$.
Equivalently, $(T'\varphi)(\mathbf{v}) = \varphi(T\mathbf{v})$.

Proposition 1.5.12. *Dual Maps Are Linear*
If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then $T' \in \mathcal{L}(\mathcal{W}', \mathcal{V}')$.

Proof. For $\varphi, \psi \in \mathcal{W}'$, $a \in \mathbb{F}$, and $\mathbf{v} \in \mathcal{V}$, $(T'(a\varphi+\psi))(\mathbf{v}) = (a\varphi+\psi)(T\mathbf{v}) = a\varphi(T\mathbf{v})+\psi(T\mathbf{v}) = (aT'\varphi+T'\psi)(\mathbf{v})$.
Structural Commentary: The dual map pulls functionals back along T and preserves linear structure.
Key Insight: Dual maps reverse direction because a functional can test $T\mathbf{v}$ only after T is applied. □

Theorem 1.5.13. *Algebra of Dual Maps*
The dual operation has the following properties:

(i) If $S, T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ and $a \in \mathbb{F}$, then $(aS+T)' = aS'+T'$.
(ii) If $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ and $S \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then $(ST)' = T'S'$.
(iii) $(I_{\mathcal{V}})' = I_{\mathcal{V}'}$.

Proof.

(i) For $\varphi \in \mathcal{W}'$ and $\mathbf{v} \in \mathcal{V}$, $((aS+T)'\varphi)(\mathbf{v}) = \varphi(aS\mathbf{v}+T\mathbf{v}) = a(S'\varphi)(\mathbf{v})+(T'\varphi)(\mathbf{v})$, so $(aS+T)' = aS'+T'$.
(ii) For $\varphi \in \mathcal{W}'$ and $\mathbf{u} \in \mathcal{U}$, $((ST)'\varphi)(\mathbf{u}) = \varphi(S(T\mathbf{u})) = (T'(S'\varphi))(\mathbf{u})$, so $(ST)' = T'S'$.
(iii) For $\varphi \in \mathcal{V}'$ and $\mathbf{v} \in \mathcal{V}$, $(I'_{\mathcal{V}}\varphi)(\mathbf{v}) = \varphi(I_{\mathcal{V}}\mathbf{v}) = \varphi(\mathbf{v})$, so $(I_{\mathcal{V}})' = I_{\mathcal{V}'}$.

Structural Commentary: Duality is linear on maps and contravariant under composition.
Key Insight: Pullbacks compose in the opposite order. □

Corollary 1.5.14. *Double Dual Maps Identify with Original Maps*
Let $\dim \mathcal{V}, \dim \mathcal{W} < \infty$ and $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.
Under $\Lambda_{\mathcal{V}} : \mathcal{V} \cong \mathcal{V}''$ and $\Lambda_{\mathcal{W}} : \mathcal{W} \cong \mathcal{W}''$, $T''\Lambda_{\mathcal{V}} = \Lambda_{\mathcal{W}}T$, so T'' identifies with T.

Proof. For $\mathbf{v} \in \mathcal{V}$ and $\psi \in \mathcal{W}'$, $\big(T''\Lambda_{\mathcal{V}}(\mathbf{v})\big)(\psi) = \Lambda_{\mathcal{V}}(\mathbf{v})(T'\psi) = (T'\psi)(\mathbf{v}) = \psi(T\mathbf{v}) = \Lambda_{\mathcal{W}}(T\mathbf{v})(\psi)$.
Hence $T''\Lambda_{\mathcal{V}} = \Lambda_{\mathcal{W}}T$.
Structural Commentary: Double dualisation commutes with natural evaluation.
Key Insight: The double dual square commutes because both sides evaluate every functional the same way. □

Corollary 1.5.15. *Dual Map of an Inverse*
If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ is an isomorphism, then T' is an isomorphism and $(T^{-1})' = (T')^{-1}$

Proof. Since $T^{-1} \in \mathcal{L}(\mathcal{W}, \mathcal{V})$, Theorem 1.5.13 gives $T'(T^{-1})' = (T^{-1}T)' = (I_{\mathcal{V}})' = I_{\mathcal{V}'}$
and $(T^{-1})'T' = (TT^{-1})' = (I_{\mathcal{W}})' = I_{\mathcal{W}'}$. Thus $(T^{-1})'$ is inverse of T', T' is isomorphism and $(T^{-1})' = (T')^{-1}$.
Structural Commentary: Duality preserves inverses contravariantly.
Key Insight: Invertibility survives duality. □

Definition 1.5.16. *Annihilator*
If $\mathcal{U} \subseteq \mathcal{V}$, the *annihilator* of $\mathcal{U}$ is $\mathcal{U}^0 := \{\varphi \in \mathcal{V}' : \varphi(\mathbf{u}) = 0 \text{ for every } \mathbf{u} \in \mathcal{U}\}$.

Proposition 1.5.17. *Annihilators Are Subspaces*
If $\mathcal{U} \subseteq \mathcal{V}$, then $\mathcal{U}^0$ is a subspace of $\mathcal{V}'$.

Proof. The zero functional lies in $\mathcal{U}^0$. If $\varphi, \psi \in \mathcal{U}^0$, $a \in \mathbb{F}$, and $\mathbf{u} \in \mathcal{U}$, then $(a\varphi+\psi)(\mathbf{u}) = a\varphi(\mathbf{u})+\psi(\mathbf{u}) = 0$.
Thus $a\varphi+\psi \in \mathcal{U}^0$, and Theorem 1.1.8 shows that $\mathcal{U}^0$ is a subspace of $\mathcal{V}'$.
Structural Commentary: The annihilator collects all linear tests that vanish on a chosen set. Closure follows because vanishing is preserved by linear combinations of tests.
Key Insight: Annihilators are kernels of evaluation constraints. □

Theorem 1.5.18. *Dimension of the Annihilator*
Let $\mathcal{V}$ be finite-dimensional and let $\mathcal{U}$ be a subspace of $\mathcal{V}$. Then $\dim(\mathcal{U}) + \dim(\mathcal{U}^0) = \dim(\mathcal{V})$.

Proof. Choose a basis $(\mathbf{u}_1, \dots, \mathbf{u}_m)$ of $\mathcal{U}$ and extend it to a basis $(\mathbf{u}_1, \dots, \mathbf{u}_m, \mathbf{w}_1, \dots, \mathbf{w}_k)$ of $\mathcal{V}$.
Let $(\varphi_1, \dots, \varphi_m, \psi_1, \dots, \psi_k)$ be the corresponding dual basis.
By Theorem 1.5.4, every $\varphi \in \mathcal{V}'$ has the expansion $\varphi = \sum_{j=1}^{m} \varphi(\mathbf{u}_j)\varphi_j + \sum_{\ell=1}^{k} \varphi(\mathbf{w}_\ell)\psi_\ell$.
If $\varphi \in \mathcal{U}^0$, then $\varphi(\mathbf{u}_j) = 0$ for every j, so $\varphi \in \operatorname{span}(\psi_1, \dots, \psi_k)$.
Conversely, each ψ_ℓ vanishes on every $\mathbf{u}_j$, hence on $\mathcal{U} = \operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_m)$. Therefore $\mathcal{U}^0 = \operatorname{span}(\psi_1, \dots, \psi_k)$.
The list $(\psi_1, \dots, \psi_k)$ is linearly independent as a sublist of the dual basis, so $\dim(\mathcal{U}^0) = k = \dim(\mathcal{V}) - \dim(\mathcal{U})$.
Structural Commentary: Extending a basis of $\mathcal{U}$ separates directions inside $\mathcal{U}$ from complementary directions. The annihilator is dual to the complementary part.
Key Insight: To vanish on $\mathcal{U}$, a functional can only use dual coordinates attached to directions outside $\mathcal{U}$. □

Theorem 1.5.19. *Annihilator Lattice Identities*
Let $\mathcal{V}$ be finite-dimensional, and let $\mathcal{S}, \mathcal{T}$ be subspaces of $\mathcal{V}$.

(i) If $\mathcal{S} \subseteq \mathcal{T}$, then $\mathcal{T}^0 \subseteq \mathcal{S}^0$.

(ii) $(\mathcal{S}+\mathcal{T})^0 = \mathcal{S}^0 \cap \mathcal{T}^0$.
(iii) $(\mathcal{S}\cap\mathcal{T})^0 = \mathcal{S}^0 + \mathcal{T}^0$.

Proof.

(i) If $\varphi \in \mathcal{T}^0$ and $\mathbf{s} \in \mathcal{S}$, then $\mathbf{s} \in \mathcal{T}$, so $\varphi(\mathbf{s}) = 0$. Hence $\varphi \in \mathcal{S}^0$.
(ii) For $\varphi \in \mathcal{V}'$, $\varphi \in (\mathcal{S}+\mathcal{T})^0$ if and only if $\varphi(\mathbf{s}+\mathbf{t}) = 0$ for all $\mathbf{s} \in \mathcal{S}$ and $\mathbf{t} \in \mathcal{T}$, if and only if $\varphi \in \mathcal{S}^0 \cap \mathcal{T}^0$.
(iii) Since $\mathcal{S}\cap\mathcal{T} \subseteq \mathcal{S}, \mathcal{T}$, part (i) gives $\mathcal{S}^0 + \mathcal{T}^0 \subseteq (\mathcal{S}\cap\mathcal{T})^0$. By part (ii), Theorem 1.2.19, and Theorem 1.5.18, $\dim(\mathcal{S}^0 + \mathcal{T}^0) = \dim(\mathcal{S}^0) + \dim(\mathcal{T}^0) - \dim(\mathcal{S}^0 \cap \mathcal{T}^0) = \dim(\mathcal{V}) - \dim(\mathcal{S}\cap\mathcal{T}) = \dim\big((\mathcal{S}\cap\mathcal{T})^0\big)$. The inclusion is therefore equality.

Structural Commentary: Annihilation is an order-reversing lattice duality.
Key Insight: Vanishing on more vectors gives fewer functionals. □

Theorem 1.5.20. *Double Annihilator*
Let $\dim \mathcal{V} < \infty$ *and* $\emptyset \neq M \subseteq \mathcal{V}$. *Under* $\Lambda_{\mathcal{V}} : \mathcal{V} \cong \mathcal{V}''$, $M^{00} = \Lambda_{\mathcal{V}}(\operatorname{span} M)$.
In particular, if $\mathcal{S} \leq \mathcal{V}$, *then* $\mathcal{S}^{00} = \Lambda_{\mathcal{V}}(\mathcal{S}) \cong \mathcal{S}$.

Proof. If $\mathbf{v} \in \operatorname{span} M$ and $\varphi \in M^0$, then $\varphi(\mathbf{v}) = 0$, so $\Lambda_{\mathcal{V}}(\mathbf{v})(\varphi) = 0$. Hence $\Lambda_{\mathcal{V}}(\operatorname{span} M) \subseteq M^{00}$.
Conversely, let $\eta \in M^{00}$. By Theorem 1.5.10, $\eta = \Lambda_{\mathcal{V}}(\mathbf{v})$ for a unique $\mathbf{v} \in \mathcal{V}$.
If $\mathbf{v} \notin \operatorname{span} M$, Proposition 1.5.7(iii) gives $\varphi \in \mathcal{V}'$ with $\varphi|_{\operatorname{span} M} = 0$ and $\varphi(\mathbf{v}) = 1$.
Then $\varphi \in M^0$ but $\eta(\varphi) = \Lambda_{\mathcal{V}}(\mathbf{v})(\varphi) = 1$, contradicting $\eta \in M^{00}$. Thus $\mathbf{v} \in \operatorname{span} M$, so $M^{00} \subseteq \Lambda_{\mathcal{V}}(\operatorname{span} M)$.
Structural Commentary: Double annihilation returns the span under $\Lambda_{\mathcal{V}}$.
Key Insight: Linear functionals separate any vector outside a finite-dimensional span. □

Theorem 1.5.21. *Kernel and Range of the Dual Map*
Let $\mathcal{V}$ *and* $\mathcal{W}$ *be finite-dimensional, and let* $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. *Then*

(i) $\ker(T') = (\operatorname{range}(T))^0$;
(ii) $\dim(\ker(T')) = \dim(\ker(T)) + \dim(\mathcal{W}) - \dim(\mathcal{V})$;
(iii) $\operatorname{range}(T') = (\ker(T))^0$;
(iv) $\dim(\operatorname{range}(T')) = \dim(\operatorname{range}(T))$.

Proof.

(i) For $\varphi \in \mathcal{W}'$, $T'\varphi = 0$ if and only if $\varphi(T\mathbf{v}) = 0$ for every $\mathbf{v} \in \mathcal{V}$, if and only if φ vanishes on $\operatorname{range}(T)$.
(ii) By (i), Theorem 1.5.18, and rank-nullity,
$\dim(\ker(T')) = \dim(\mathcal{W}) - \dim(\operatorname{range}(T)) = \dim(\ker(T)) + \dim(\mathcal{W}) - \dim(\mathcal{V})$.
(iii) If $\varphi = T'\psi$ and $\mathbf{v} \in \ker(T)$, then $\varphi(\mathbf{v}) = \psi(T\mathbf{v}) = 0$, so $\operatorname{range}(T') \subseteq (\ker(T))^0$. Also, by rank-nullity for T', (ii), Theorem 1.5.4, and Theorem 1.5.18, $\dim(\operatorname{range}(T')) = \dim(\mathcal{V}) - \dim(\ker(T)) = \dim\big((\ker(T))^0\big)$. Thus $\operatorname{range}(T') = (\ker(T))^0$.
(iv) By (iii), Theorem 1.5.18, and rank-nullity, $\dim(\operatorname{range}(T')) = \dim(\mathcal{V}) - \dim(\ker(T)) = \dim(\operatorname{range}(T))$

Structural Commentary: Duality exchanges kernels and ranges through annihilators.
Key Insight: A functional dies after pulling back by T exactly when it vanishes on everything T can reach. □

Corollary 1.5.22. *Surjectivity and Dual Injectivity*
Let $\mathcal{V}$ *and* $\mathcal{W}$ *be finite-dimensional, and let* $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. *Then* T *is surjective if and only if* T' *is injective.*

Proof.

($\Rightarrow$) Suppose T is surjective. Then $\operatorname{range}(T) = \mathcal{W}$, so Theorem 1.5.21 gives $\ker(T') = \mathcal{W}^0 = \{0\}$.
($\Leftarrow$) Suppose T' is injective. Then $\ker(T') = \{0\}$, so Theorem 1.5.21 gives $(\operatorname{range}(T))^0 = \{0\}$. By Theorem 1.5.18, $\dim(\operatorname{range}(T)) = \dim(\mathcal{W})$, hence $\operatorname{range}(T) = \mathcal{W}$.

Structural Commentary: A map is onto exactly when no nonzero functional can vanish on its range. Dual injectivity detects this absence of invisible codomain directions.
Key Insight: Surjectivity of T becomes injectivity of the pullback map T'. □

Corollary 1.5.23. *Injectivity and Dual Surjectivity*
Let $\mathcal{V}$ *and* $\mathcal{W}$ *be finite-dimensional, and let* $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. *Then* T *is injective if and only if* T' *is surjective.*

Proof.

($\Rightarrow$) Suppose T is injective. Then $\ker(T) = \{\mathbf{0}\}$, so Theorem 1.5.21 gives $\operatorname{range}(T') = (\ker(T))^0 = \mathcal{V}'$.
($\Leftarrow$) Suppose T' is surjective. Then $\operatorname{range}(T') = \mathcal{V}'$, so Theorem 1.5.21 gives $(\ker(T))^0 = \mathcal{V}'$. By Theorem 1.5.18, this forces $\dim(\ker(T)) = 0$, hence $\ker(T) = \{\mathbf{0}\}$.

Structural Commentary: Injectivity says there are no nonzero domain directions killed by T. The dual statement says every test on the domain comes by pulling back a test from the codomain.
Key Insight: No kernel directions means no obstruction to extending functionals backward through T. □

Corollary 1.5.24. *Dual of a Quotient as an Annihilator*
Let $\mathcal{V}$ be finite-dimensional, let $\mathcal{U} \leq \mathcal{V}$, and let $\pi : \mathcal{V} \to \mathcal{V}/\mathcal{U}$ be the quotient map. Then $\ker(\pi') = \{0\}$ and $\operatorname{range}(\pi') = \mathcal{U}^0$. Hence π' identifies $(\mathcal{V}/\mathcal{U})'$ with $\mathcal{U}^0$, so $(\mathcal{V}/\mathcal{U})' \cong \mathcal{U}^0$.

Proof. By Theorem 1.2.32, $\mathcal{V}/\mathcal{U}$ is finite-dimensional. Since π is surjective and $\ker(\pi) = \mathcal{U}$, Corollary 1.5.22 gives $\ker(\pi') = \{0\}$, and Theorem 1.5.21 gives $\operatorname{range}(\pi') = (\ker(\pi))^0 = \mathcal{U}^0$.
Structural Commentary: A functional on the quotient is exactly a functional on $\mathcal{V}$ that cannot see $\mathcal{U}$.
Key Insight: Pulling back along the quotient map identifies $(\mathcal{V}/\mathcal{U})'$ with $\mathcal{U}^0$. □

Theorem 1.5.25. *Direct Sums and Annihilators*
Let $\mathcal{V}$ be finite-dimensional, and suppose $\mathcal{V} = \mathcal{S} \oplus \mathcal{T}$.

(i) Extension by zero on $\mathcal{S}$ gives an isomorphism $\mathcal{T}' \cong \mathcal{S}^0$.
(ii) $\mathcal{V}' = \mathcal{S}^0 \oplus \mathcal{T}^0$.

Proof.

(i) Define $E : \mathcal{T}' \to \mathcal{S}^0$ by $(E\varphi)(\mathbf{s} + \mathbf{t}) = \varphi(\mathbf{t})$. Since $\mathcal{V} = \mathcal{S} \oplus \mathcal{T}$, this is well-defined and linear. Restriction $R : \mathcal{S}^0 \to \mathcal{T}'$, $R\psi = \psi|_{\mathcal{T}}$, satisfies $RE = I_{\mathcal{T}'}$ and $ER = I_{\mathcal{S}^0}$, so E is an isomorphism.
(ii) By Theorem 1.5.19, $\mathcal{S}^0 + \mathcal{T}^0 = (\mathcal{S} \cap \mathcal{T})^0 = \{\mathbf{0}\}^0 = \mathcal{V}'$, and $\mathcal{S}^0 \cap \mathcal{T}^0 = (\mathcal{S} + \mathcal{T})^0 = \mathcal{V}^0 = \{0\}$. Thus $\mathcal{V}' = \mathcal{S}^0 \oplus \mathcal{T}^0$.

Structural Commentary: A direct sum decomposes the dual by complementary annihilators.
Key Insight: Direct sums split the dual space by annihilators. □

Theorem 1.5.26. *Range Membership Dual Alternative*
Let $\mathcal{V}, \mathcal{W}$ be finite-dimensional, $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, and $\mathbf{w} \in \mathcal{W}$. Exactly one holds:

(i) $\mathbf{w} \in \operatorname{range}(T)$;
(ii) there exists $\psi \in \mathcal{W}'$ such that $\psi T = 0$ and $\psi(\mathbf{w}) \neq 0$.

Proof.

(i) The alternatives are incompatible: if $\mathbf{w} = T\mathbf{v}$ and $\psi T = 0$, then $\psi(\mathbf{w}) = \psi(T\mathbf{v}) = 0$.
(ii) Suppose $\mathbf{w} \notin \operatorname{range}(T)$. By Proposition 1.5.7(iii), there exists $\psi \in \mathcal{W}'$ such that $\psi|_{\operatorname{range}(T)} = 0$ and $\psi(\mathbf{w}) = 1$.
Then ψ vanishes on $\operatorname{range}(T)$, so $\psi T = 0$, while $\psi(\mathbf{w}) = 1$.

Structural Commentary: A missing range vector is detected by a dual test that vanishes on the whole range.
Key Insight: Either a vector is reached by the map, or a functional separates it from the range. □

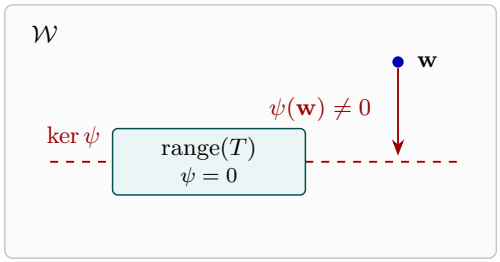


Figure 16: A dual certificate detects a vector outside the range.

Proposition 1.5.27. *Standard Coordinate Dual Identification*
For each positive integer m, the map

$$\Theta : \mathbb{F}^m \to (\mathbb{F}^m)', \qquad \Theta(\mathbf{y})(\mathbf{z}) = \mathbf{y}^t\mathbf{z},$$

is an isomorphism. Hence every $\psi \in (\mathbb{F}^m)'$ has a unique form $\psi(\mathbf{z}) = \mathbf{y}^t\mathbf{z}$ with $\mathbf{y} \in \mathbb{F}^m$.

Proof. Linearity of Θ follows from bilinearity of the coordinate dot product. If $\Theta(\mathbf{y}) = 0$, then evaluating at the standard basis vector $\mathbf{e}_j$ gives $y_j = 0$ for every j, so $\mathbf{y} = \mathbf{0}$. Thus Θ is injective. Since $\dim(\mathbb{F}^m) = m = \dim((\mathbb{F}^m)')$ by Theorem 1.5.4, Corollary 1.3.30 gives that Θ is an isomorphism.

Structural Commentary: Coordinate column vectors represent all coordinate linear functionals by row multiplication.
Key Insight: The standard dual of $\mathbb{F}^m$ is represented by transposed coordinate vectors. □

Corollary 1.5.28. *Matrix Equation Alternative*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$ and $\mathbf{b} \in \mathbb{F}^m$. Exactly one holds:

(i) $A\mathbf{x} = \mathbf{b}$ *for some* $\mathbf{x} \in \mathbb{F}^n$;
(ii) $\mathbf{y}^t A = \mathbf{0}^t$ *and* $\mathbf{y}^t\mathbf{b} \neq 0$ *for some* $\mathbf{y} \in \mathbb{F}^m$.

Proof.

(i) By Theorem 1.5.26, for $T_A : \mathbb{F}^n \to \mathbb{F}^m$ and $\mathbf{w} = \mathbf{b}$, $\mathbf{b} \in \mathrm{range}(T_A)$ is exactly $A\mathbf{x} = \mathbf{b}$ for some $\mathbf{x} \in \mathbb{F}^n$.
(ii) By Proposition 1.5.27, $\psi \in (\mathbb{F}^m)'$ has the form $\psi(\mathbf{z}) = \mathbf{y}^t\mathbf{z}$ for a unique $\mathbf{y} \in \mathbb{F}^m$.
Thus $\psi T_A = 0$ iff $\mathbf{y}^t A = \mathbf{0}^t$, and $\psi(\mathbf{b}) \neq 0$ iff $\mathbf{y}^t\mathbf{b} \neq 0$.

Structural Commentary: Inconsistency is certified by a left-null vector.
Key Insight: Left null vectors are the coordinate form of dual separating functionals. □

Theorem 1.5.29. *Matrix of the Dual Map Is the Transpose*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$, and let $\beta = (\mathbf{v}_1,\dots,\mathbf{v}_n)$ and $\gamma = (\mathbf{w}_1,\dots,\mathbf{w}_m)$ be bases of $\mathcal{V}$ and $\mathcal{W}$. If β' and γ' are the corresponding dual bases, then $[T']_{\beta'\leftarrow\gamma'} = ([T]_{\gamma\leftarrow\beta})^t$.

Proof. Let $A = [T]_{\gamma\leftarrow\beta}$ and $C = [T']_{\beta'\leftarrow\gamma'}$. Write $\gamma' = (\psi_1,\dots,\psi_m)$ and $\beta' = (\varphi_1,\dots,\varphi_n)$.
For $1 \le j \le m$, $T'\psi_j = \sum_{r=1}^{n} C_{r,j}\varphi_r$. Evaluating at $\mathbf{v}_k$ gives $C_{k,j} = (T'\psi_j)(\mathbf{v}_k) = \psi_j(T\mathbf{v}_k)$.
Since $T\mathbf{v}_k = \sum_{r=1}^{m} A_{r,k}\mathbf{w}_r$, we get $\psi_j(T\mathbf{v}_k) = A_{j,k}$. Hence $C_{k,j} = A_{j,k}$ for all j,k, so $C = A^t$.
Structural Commentary: Duality turns the matrix of a map into the matrix of its pullback on coordinate functionals. The row-column indexing is exactly transposition.
Key Insight: Coefficients of $T\mathbf{v}_k$ become values of the dual basis functionals on $T\mathbf{v}_k$. □

Definition 1.5.30. *Bilinear Forms and Matrices*
A map $B : \mathcal{V}\times\mathcal{W} \to \mathbb{F}$ is *bilinear* if it is linear in each variable separately.
For bases $\beta = (\mathbf{v}_1,\dots,\mathbf{v}_m)$ of $\mathcal{V}$ and $\gamma = (\mathbf{w}_1,\dots,\mathbf{w}_n)$ of $\mathcal{W}$, the *matrix of B in the bases β,γ* is $[B]_{\beta,\gamma} = [B(\mathbf{v}_i,\mathbf{w}_j)]$. If $\mathbf{x} = [\mathbf{v}]_\beta$ and $\mathbf{y} = [\mathbf{w}]_\gamma$, then $B(\mathbf{v},\mathbf{w}) = \mathbf{x}^t[B]_{\beta,\gamma}\mathbf{y}$.

Proposition 1.5.31. *Bilinear Forms Linearize through Duals*
Let $B : \mathcal{V}\times\mathcal{W} \to \mathbb{F}$ be bilinear, and define $\Phi_B : \mathcal{V} \to \mathcal{W}'$ where $(\Phi_B\mathbf{v})(\mathbf{w}) = B(\mathbf{v},\mathbf{w})$, and $\Psi_B : \mathcal{W} \to \mathcal{V}'$ where $(\Psi_B\mathbf{w})(\mathbf{v}) = B(\mathbf{v},\mathbf{w})$. Then Φ_B,Ψ_B are linear. For bases β,γ and $M = [B]_{\beta,\gamma}$, $[\Phi_B]_{\gamma'\leftarrow\beta} = M^t$; and $[\Psi_B]_{\beta'\leftarrow\gamma} = M$.

Proof. For fixed $\mathbf{v}$ and $\mathbf{w}$, bilinearity gives $\Phi_B\mathbf{v} \in \mathcal{W}'$ and $\Psi_B\mathbf{w} \in \mathcal{V}'$. For $\mathbf{v}_1,\mathbf{v}_2 \in \mathcal{V}$, $a \in \mathbb{F}$, and every $\mathbf{w} \in \mathcal{W}$, $(\Phi_B(a\mathbf{v}_1+\mathbf{v}_2))(\mathbf{w}) = B(a\mathbf{v}_1+\mathbf{v}_2,\mathbf{w}) = a(\Phi_B\mathbf{v}_1)(\mathbf{w}) + (\Phi_B\mathbf{v}_2)(\mathbf{w})$, so Φ_B is linear.
For $\mathbf{w}_1,\mathbf{w}_2 \in \mathcal{W}$, $a \in \mathbb{F}$, and every $\mathbf{v} \in \mathcal{V}$, $(\Psi_B(a\mathbf{w}_1+\mathbf{w}_2))(\mathbf{v}) = B(\mathbf{v},a\mathbf{w}_1+\mathbf{w}_2) = a(\Psi_B\mathbf{w}_1)(\mathbf{v}) + (\Psi_B\mathbf{w}_2)(\mathbf{v})$, so Ψ_B is linear. Let $\beta = (\mathbf{v}_1,\dots,\mathbf{v}_m)$ and $\gamma = (\mathbf{w}_1,\dots,\mathbf{w}_n)$.
The jth γ'-coordinate of $\Phi_B(\mathbf{v}_i)$ is $(\Phi_B\mathbf{v}_i)(\mathbf{w}_j) = B(\mathbf{v}_i,\mathbf{w}_j) = M_{i,j}$, so column i of $[\Phi_B]_{\gamma'\leftarrow\beta}$ is row i of M transposed, hence $[\Phi_B]_{\gamma'\leftarrow\beta} = M^t$. Likewise, the ith β'-coordinate of $\Psi_B(\mathbf{w}_j)$ is $M_{i,j}$, so $[\Psi_B]_{\beta'\leftarrow\gamma} = M$.
Structural Commentary: A bilinear form becomes linear once one input is frozen. Tracking both freezes gives two dual-valued linear maps whose matrices are transposes of one another.
Key Insight: Bilinearity is linearity into a dual space. □

Definition 1.5.32. *Radicals and Nonsingularity of a Bilinear Pairing*
Let $B : \mathcal{V}\times\mathcal{W} \to \mathbb{F}$ be bilinear.

(i) The *left* and *right radicals* are

$$\mathrm{rad}_L(B) := \{\mathbf{v} \in \mathcal{V} : B(\mathbf{v},\mathbf{w}) = 0 \text{ for every } \mathbf{w} \in \mathcal{W}\},$$
$$\mathrm{rad}_R(B) := \{\mathbf{w} \in \mathcal{W} : B(\mathbf{v},\mathbf{w}) = 0 \text{ for every } \mathbf{v} \in \mathcal{V}\}.$$

(ii) B is *left nondegenerate* if $\mathrm{rad}_L(B) = \{\mathbf{0}\}$ and *right nondegenerate* if $\mathrm{rad}_R(B) = \{\mathbf{0}\}$.
(iii) B is *nonsingular* or *perfect* if $\Phi_B : \mathcal{V} \to \mathcal{W}'$ and $\Psi_B : \mathcal{W} \to \mathcal{V}'$ from Proposition 1.5.31 are isomorphisms.

Proposition 1.5.33. *Radicals, Matrices, and Dual Maps*
Let $B : \mathcal{V}\times\mathcal{W} \to \mathbb{F}$ be bilinear, $\dim\mathcal{V} = m$, $\dim\mathcal{W} = n$, and $M = [B]_{\beta,\gamma}$ in ordered bases β,γ. Then $\{[\mathbf{v}]_\beta : \mathbf{v} \in \mathrm{rad}_L(B)\} = \ker(M^t)$ and $\{[\mathbf{w}]_\gamma : \mathbf{w} \in \mathrm{rad}_R(B)\} = \ker(M)$, therefore $\dim\mathrm{rad}_L(B) = m - \mathrm{rank}(M)$ and $\dim\mathrm{rad}_R(B) = n - \mathrm{rank}(M)$. Moreover, $\ker(\Phi_B) = \mathrm{rad}_L(B)$ and $\ker(\Psi_B) = \mathrm{rad}_R(B)$.

Proof. For $\mathbf{x} = [\mathbf{v}]_\beta$ and $\mathbf{y} = [\mathbf{w}]_\gamma$, $B(\mathbf{v}, \mathbf{w}) = \mathbf{x}^t M \mathbf{y}$.
Thus $\mathbf{v} \in \mathrm{rad}_L(B)$ if and only if $\mathbf{x}^t M \mathbf{y} = 0$ for all $\mathbf{y}$, if and only if $M^t \mathbf{x} = 0$. Similarly, $\mathbf{w} \in \mathrm{rad}_R(B)$ if and only if $M\mathbf{y} = 0$. This proves the coordinate-set identities. Rank-nullity gives the dimension formulas. Finally, $\mathbf{v} \in \ker(\Phi_B) \iff B(\mathbf{v}, \mathbf{w}) = 0$ for every $\mathbf{w} \in \mathcal{W} \iff \mathbf{v} \in \mathrm{rad}_L(B)$; similarly $\ker(\Psi_B) = \mathrm{rad}_R(B)$.
Structural Commentary: Radicals are the kernel obstructions in each slot.
Key Insight: In coordinates, radicals are matrix nullspaces. □

Proposition 1.5.34. *Square Pairing Nonsingularity Criteria*
Let $B : \mathcal{V} \times \mathcal{W} \to \mathbb{F}$ be bilinear with $\dim \mathcal{V} = \dim \mathcal{W} = n$, and let $M = [B]_{\beta,\gamma}$ in ordered bases β, γ. Then the following are equivalent:

(i) B is left nondegenerate;
(ii) B is right nondegenerate;
(iii) M is invertible;
(iv) B is perfect.

Proof.

(*i*) $\Rightarrow$ (*ii*). If $\mathrm{rad}_L(B) = 0$, then by Proposition 1.5.33, $\ker(M^t) = 0 \Rightarrow \mathrm{rank}(M) = n \Rightarrow \ker(M) = 0$.
(*ii*) $\Rightarrow$ (*iii*). If $\mathrm{rad}_R(B) = 0$, then $\ker(M) = 0$ by Proposition 1.5.33; since $M \in \mathrm{Mat}(n, n, \mathbb{F})$, M is invertible.
(*iii*) $\Rightarrow$ (*iv*). If M is invertible, then $\ker(M) = \ker(M^t) = 0$, so $\mathrm{rad}_L(B) = \mathrm{rad}_R(B) = 0$.
Hence $\ker(\Phi_B) = \ker(\Psi_B) = 0$ by Proposition 1.5.33.
Since $\dim \mathcal{V} = \dim \mathcal{W}' = \dim \mathcal{W} = \dim \mathcal{V}' = n$, Φ_B and Ψ_B are isomorphisms; hence B is perfect.
(*iv*) $\Rightarrow$ (*i*). If B is perfect, then Φ_B is an isomorphism. Thus $\mathrm{rad}_L(B) = \ker(\Phi_B) = 0$.

Structural Commentary: In square pairings, one-sided nondegeneracy is invertibility.
Key Insight: Perfect pairings are exactly invertible matrices viewed without coordinates. □

Corollary 1.5.35. *Perfect Pairings Represent Functionals*
Let $B : \mathcal{V} \times \mathcal{W} \to \mathbb{F}$ be perfect.
Then $\forall \psi \in \mathcal{W}'\ \exists! \mathbf{v} \in \mathcal{V} : \psi(\mathbf{w}) = B(\mathbf{v}, \mathbf{w})$, $\mathbf{w} \in \mathcal{W}$; and $\forall \varphi \in \mathcal{V}'\ \exists! \mathbf{w} \in \mathcal{W} : \varphi(\mathbf{v}) = B(\mathbf{v}, \mathbf{w})$, $\mathbf{v} \in \mathcal{V}$.

Proof. Since B is perfect, $\Phi_B : \mathcal{V} \to \mathcal{W}'$ and $\Psi_B : \mathcal{W} \to \mathcal{V}'$ are isomorphisms.
Thus $\psi = \Phi_B(\mathbf{v})$ and $\varphi = \Psi_B(\mathbf{w})$ for unique $\mathbf{v}, \mathbf{w}$, which gives the two formulas.
Structural Commentary: Perfect pairings identify each space with the other's dual.
Key Insight: Functionals are represented by vectors exactly when the appropriate pairing map is onto. □

Definition 1.5.36. *Matrix Congruence*
Matrices $M, N \in \mathrm{Mat}(n, n, \mathbb{F})$ are *congruent* if $\exists\, P \in \mathrm{Mat}(n, n, \mathbb{F})$ invertible such that $N = P^t M P$.

Proposition 1.5.37. *Basis Change for a Square Bilinear Form*
Let $\mathcal{V}$ be finite-dimensional, let $B : \mathcal{V} \times \mathcal{V} \to \mathbb{F}$ be bilinear, and let β, γ be ordered bases of $\mathcal{V}$.
Set $M_\beta := [B]_{\beta,\beta}$; $M_\gamma := [B]_{\gamma,\gamma}$, and $C := C_{\beta \leftarrow \gamma}$.
Then $M_\gamma = C^t M_\beta C$. Hence M_β and M_γ are congruent.

Proof. For $\mathbf{x}, \mathbf{y} \in \mathcal{V}$, $[\mathbf{x}]_\beta = C[\mathbf{x}]_\gamma$ and $[\mathbf{y}]_\beta = C[\mathbf{y}]_\gamma$. Therefore $B(\mathbf{x}, \mathbf{y}) = [\mathbf{x}]_\beta^t M_\beta [\mathbf{y}]_\beta = [\mathbf{x}]_\gamma^t C^t M_\beta C [\mathbf{y}]_\gamma$.
By uniqueness of the representing matrix in the basis γ, $M_\gamma = C^t M_\beta C$.
Since C is invertible by Theorem 1.3.92, Definition 1.5.36 gives congruence.
Structural Commentary: Operators change by similarity; square bilinear forms change by congruence because both input slots are rewritten.
Key Insight: A basis change acts on a square bilinear form from both sides. □

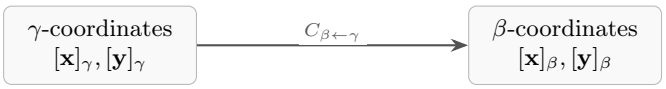


Figure 17: Changing both slots gives $M_\gamma = C_{\beta \leftarrow \gamma}^t M_\beta C_{\beta \leftarrow \gamma}$.

Perspective 1.5.38. *Bilinear Forms, Inner Products, and Tensor Scope*

(i) Bilinear forms: $B : \mathcal{V} \times \mathcal{W} \to \mathbb{F}$; inner products add $\mathbf{v} \neq \mathbf{0} \Rightarrow B(\mathbf{v}, \mathbf{v}) > 0$ and, over $\mathbb{C}$, $B(\mathbf{v}, \mathbf{w}) = \overline{B(\mathbf{w}, \mathbf{v})}$.
(ii) $\Phi_B : \mathcal{V} \to \mathcal{W}'$ and $\Psi_B : \mathcal{W} \to \mathcal{V}'$ encode B in dual spaces, with $[\Phi_B]_{\gamma' \leftarrow \beta} = [B]_{\beta,\gamma}^t$ and $[\Psi_B]_{\beta' \leftarrow \gamma} = [B]_{\beta,\gamma}$.
(iii) $\mathcal{V} \oplus \mathcal{W}$ and $\mathcal{V} \times \mathcal{W}$ use single-coordinate slots; the finite tensor scaffold uses paired slots (i, j).

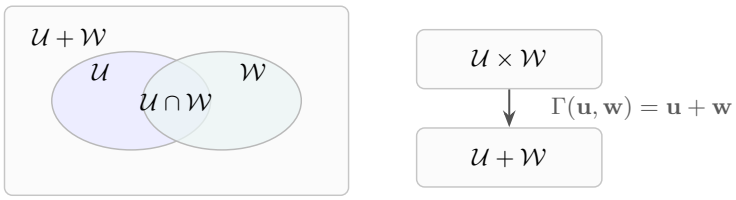


Figure 18: Intersection, sum, and product of subspaces.

1.6 Finite-Dimensional Topological Vector Spaces

Finite-dimensional real vector spaces carry a canonical topology: choose coordinates, use the standard Euclidean topology on $\mathbb{R}^n$, and then forget that coordinates were chosen. All coordinate changes are homeomorphisms.

Definition 1.6.1. *Topological Vector Space*
Let $\mathcal{V}$ be a real vector space and let $\mathcal{T}$ be a topology on its underlying set.
The pair $(\mathcal{V},\mathcal{T})$ is a *topological vector space* if addition $A:\mathcal{V}\times\mathcal{V}\to\mathcal{V}$, $A(\mathbf{u},\mathbf{v})=\mathbf{u}+\mathbf{v}$, and scalar multiplication $M:\mathbb{R}\times\mathcal{V}\to\mathcal{V}$, $M(a,\mathbf{v})=a\mathbf{v}$ are continuous, where products carry the product topology.

Example 1.6.2. *The Standard Topology on $\mathbb{R}^n$*
With the Euclidean topology, $\mathbb{R}^n$ is a topological vector space: the maps $(\mathbf{x},\mathbf{y})\longmapsto\mathbf{x}+\mathbf{y}$, $(a,\mathbf{x})\longmapsto a\mathbf{x}$ have polynomial coordinate functions x_i+y_i and ax_i, hence are continuous.

Proposition 1.6.3. *Real Matrix Maps Are Continuous*
Let $A=[a_{j,k}]\in\mathrm{Mat}(m,n,\mathbb{R})$, and let $T_A:\mathbb{R}^n\to\mathbb{R}^m$ be the induced map $T_A(\mathbf{x})=A\mathbf{x}$. Then T_A is continuous for the Euclidean topologies. If $m=n$ and A is invertible, then T_A is a homeomorphism. In particular, every linear functional on $\mathbb{R}^n$ is continuous.

Proof. For $\mathbf{x}=(x_1,\dots,x_n)^t$, the jth coordinate of $T_A(\mathbf{x})$ is $\sum_{k=1}^{n}a_{j,k}x_k$, hence is continuous. Thus T_A is continuous. If A is invertible, then $(T_A)^{-1}=T_{A^{-1}}$, which is continuous by the first part. Finally, every $f\in(\mathbb{R}^n)'$ equals T_A for a unique $A\in\mathrm{Mat}(1,n,\mathbb{R})$, so f is continuous. *Structural Commentary*: Matrix maps on Euclidean coordinate spaces are continuous coordinate by coordinate.
Key Insight: Invertible real matrices give homeomorphisms of Euclidean space. □

Definition 1.6.4. *Natural Topology*
Let $\mathcal{V}$ be an n-dimensional real vector space, let β be a basis, and let $[\cdot]_\beta:\mathcal{V}\to\mathbb{R}^n$ be the coordinate isomorphism.
The *natural topology* on $\mathcal{V}$ is $\mathcal{T}_\beta:=\{[\cdot]_\beta^{-1}(O):O\subseteq\mathbb{R}^n\text{ is open}\}$.
Equivalently, $U\subseteq\mathcal{V}$ is open if and only if $[\cdot]_\beta(U)$ is open in $\mathbb{R}^n$.

Proposition 1.6.5. *Natural Topology Is Basis-Independent*
If β and γ are bases of $\mathcal{V}$, then $\mathcal{T}_\beta=\mathcal{T}_\gamma$.

Proof. Set $H_{\gamma\leftarrow\beta}:=[\cdot]_\gamma\circ[\cdot]_\beta^{-1}:\mathbb{R}^n\to\mathbb{R}^n$.
Then $H_{\gamma\leftarrow\beta}\mathbf{x}=C_{\gamma\leftarrow\beta}\mathbf{x}$. Since $C_{\gamma\leftarrow\beta}$ is invertible with inverse $C_{\beta\leftarrow\gamma}$, Proposition 1.6.3 shows that $H_{\gamma\leftarrow\beta}$ is a homeomorphism with inverse $H_{\beta\leftarrow\gamma}$.
Hence, for $U\subseteq\mathcal{V}$, $U\in\mathcal{T}_\beta\iff[\cdot]_\beta(U)$ is open $\iff H_{\gamma\leftarrow\beta}([\cdot]_\beta(U))=[\cdot]_\gamma(U)$ is open $\iff U\in\mathcal{T}_\gamma$. *Structural Commentary*: Basis changes are homeomorphisms.
Key Insight: The natural topology is carried by the vector-space structure, not by a preferred basis. □

Theorem 1.6.6. *Characterization of the Natural Topology*
Let $\dim_{\mathbb{R}}\mathcal{V}<\infty$. There is a unique topology $\mathcal{T}$ on $\mathcal{V}$ such that:

(i) $(\mathcal{V},\mathcal{T})$ is a topological vector space;
(ii) $\mathcal{V}'\subseteq C((\mathcal{V},\mathcal{T}),\mathbb{R})$.

Moreover, $\mathcal{T}=\mathcal{T}_\beta$ for every basis β.

Proof. Fix an arbitrary basis $\beta=(\mathbf{v}_1,\dots,\mathbf{v}_n)$.

(i) *Existence*: put $\mathcal{T}=\mathcal{T}_\beta$. Since $[\cdot]_\beta:(\mathcal{V},\mathcal{T}_\beta)\to\mathbb{R}^n$ is a homeomorphism, the vector operations on $\mathcal{V}$ are transported from the continuous vector operations on $\mathbb{R}^n$.
Also, for $\varphi\in\mathcal{V}'$, $\varphi\circ[\cdot]_\beta^{-1}\in(\mathbb{R}^n)'$, hence is continuous by Proposition 1.6.3. Thus (i)–(ii) hold.

(ii) *Uniqueness*: suppose $\mathcal{T}$ satisfies (i)–(ii).
For $1 \le i \le n$, the coordinate functional $\pi_i \circ [\cdot]_\beta \in \mathcal{V}'$ is continuous, so $[\cdot]_\beta : (\mathcal{V}, \mathcal{T}) \to \mathbb{R}^n$ is continuous.
Its inverse $\Psi(a_1, \dots, a_n) = a_1\mathbf{v}_1 + \cdots + a_n\mathbf{v}_n$ is continuous by (i), since it is built from scalar multiplication and finite sums. Hence $[\cdot]_\beta$ is a homeomorphism, so $\mathcal{T} = \mathcal{T}_\beta$.

Structural Commentary: Functionals give coordinates; operations give inverse coordinates.
Key Insight: Functionals and vector operations determine finite-dimensional coordinates topologically. □

Corollary 1.6.7. *Finite-Dimensional Linear Maps Are Continuous*
Let $\dim_{\mathbb{R}} \mathcal{V}, \dim_{\mathbb{R}} \mathcal{W} < \infty$*, with natural topologies. Then* $\mathcal{L}(\mathcal{V}, \mathcal{W}) \subseteq C(\mathcal{V}, \mathcal{W})$.
If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ *is an isomorphism, then* T *is a homeomorphism.*

Proof. Choose bases β of $\mathcal{V}$ and γ of $\mathcal{W}$. For $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, $[\cdot]_\gamma \circ T \circ [\cdot]_\beta^{-1}(\mathbf{x}) = [T]_{\gamma \leftarrow \beta}\mathbf{x}$, which is continuous by Proposition 1.6.3. Since $[\cdot]_\beta$ and $[\cdot]_\gamma$ are homeomorphisms, T is continuous.
If T is an isomorphism, apply the same argument to T^{-1}.
Structural Commentary: Finite-dimensional linear maps are just matrix maps in natural coordinates.
Key Insight: A basis turns topological questions about linear maps into continuity of matrix multiplication. □

Perspective 1.6.8. *Finite-Dimensional Topology Is Rigid*
For finite-dimensional real vector spaces, topology is algebraically forced.
Note that $\mathcal{T}_\beta = \mathcal{T}_\gamma$; $\mathcal{L}(\mathcal{V}, \mathcal{W}) \subseteq C(\mathcal{V}, \mathcal{W})$; $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ bijective $\Rightarrow$ T is a homeomorphism.
Infinite-dimensional topological vector spaces need not be rigid in this sense.

2 Geometric Linear Algebra

Geometric linear algebra adds inner products, orthogonality, projections, least squares, adjoints, metric bilinear forms, isometries, affine/projective incidence, convex separation, and positive systems to the linear story.

2.1 Inner Product Geometry

Throughout this subsection, $\mathbb{F}$ is $\mathbb{R}$ or $\mathbb{C}$. Inner products turn linear spaces into geometric spaces with norms, orthogonality, orthonormal bases, and Gram matrices.

Definition 2.1.1. *Preliminary Terminology*

(i) For $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$, the *dot product* is $\mathbf{x} \cdot \mathbf{y} := x_1y_1 + \cdots + x_ny_n$.
(ii) For $z = a + bi \in \mathbb{C}$, the *complex conjugate* is $\overline{z} := a - bi$, and the *modulus* is $|z| := \sqrt{a^2 + b^2}$.
(iii) A map $f : \mathcal{V} \to \mathbb{F}$ is *conjugate linear* if $f(\mathbf{u} + \mathbf{v}) = f(\mathbf{u}) + f(\mathbf{v})$ and $f(\lambda\mathbf{v}) = \overline{\lambda} f(\mathbf{v})$.
(iv) A function $B : \mathcal{V} \times \mathcal{V} \to \mathbb{F}$ is *sesquilinear* if it is linear in one slot and conjugate linear in the other.

Motivation 2.1.2. *From Dot Products to Inner Products*
For $\mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$, the Euclidean length is $\|\mathbf{x}\| = \sqrt{x_1^2 + \cdots + x_n^2}$.
The expression under the square root is the dot product of $\mathbf{x}$ with itself. Over $\mathbb{C}$, conjugates are needed to keep squared length real and nonnegative: $\|\mathbf{z}\|^2 = |z_1|^2 + \cdots + |z_n|^2 = z_1\overline{z_1} + \cdots + z_n\overline{z_n}$.
Thus the standard inner product on $\mathbb{F}^n$ is $\langle \mathbf{x}, \mathbf{y} \rangle := x_1\overline{y_1} + \cdots + x_n\overline{y_n}$.

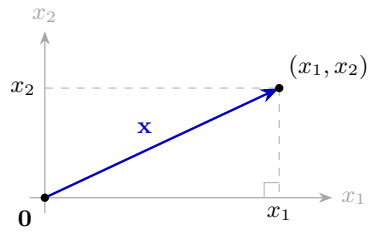


Figure 19: The norm is vector length, with $\|\mathbf{x}\|^2 = \langle \mathbf{x}, \mathbf{x} \rangle$.

Definition 2.1.3. *Inner Product*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$. An *inner product* on $\mathcal{V}$ is a function that assigns to each pair $\mathbf{u}, \mathbf{v} \in \mathcal{V}$ a scalar $\langle \mathbf{u}, \mathbf{v} \rangle \in \mathbb{F}$ satisfying, for all $\mathbf{u}, \mathbf{v}, \mathbf{w} \in \mathcal{V}$ and all $\lambda \in \mathbb{F}$,

i. *positivity*: $\langle \mathbf{v}, \mathbf{v} \rangle \ge 0$;
ii. *definiteness*: $\langle \mathbf{v}, \mathbf{v} \rangle = 0$ if and only if $\mathbf{v} = \mathbf{0}$;
iii. *additivity in the first slot*: $\langle \mathbf{u} + \mathbf{v}, \mathbf{w} \rangle = \langle \mathbf{u}, \mathbf{w} \rangle + \langle \mathbf{v}, \mathbf{w} \rangle$;

iv. *homogeneity in the first slot*: $\langle \lambda \mathbf{u}, \mathbf{v} \rangle = \lambda \langle \mathbf{u}, \mathbf{v} \rangle$;
v. *conjugate symmetry*: $\langle \mathbf{u}, \mathbf{v} \rangle = \overline{\langle \mathbf{v}, \mathbf{u} \rangle}$.

In the positivity axiom, $\langle \mathbf{v}, \mathbf{v} \rangle \geq 0$ means that $\langle \mathbf{v}, \mathbf{v} \rangle$ is a real nonnegative scalar.

Definition 2.1.4. *Inner-Product Space*
An *inner-product space* is a vector space $\mathcal{V}$ together with an inner product on $\mathcal{V}$.

Example 2.1.5. *Standard Inner Product on $\mathbb{F}^n$*
For $\mathbf{x}, \mathbf{y} \in \mathbb{F}^n$, the standard inner product is $\langle \mathbf{x}, \mathbf{y} \rangle = \sum_{j=1}^{n} x_j \overline{y_j}$. When $\mathbb{F} = \mathbb{R}$, this is the usual dot product.

Example 2.1.6. *Function and Matrix Inner Products*
If $\mathcal{V} = C_{\mathbb{F}}[a, b]$, then $\langle f, g \rangle = \int_a^b f(t)\overline{g(t)}\, dt$ is an inner product.
If $\mathcal{V} = \mathrm{Mat}(m, n, \mathbb{F})$, then the *Frobenius inner product* is $\langle A, B \rangle_F := \sum_{i=1}^{m} \sum_{j=1}^{n} a_{ij} \overline{b_{ij}}$.

Example 2.1.7. *Square-Summable Sequences*
Set $\ell^2(\mathbb{F}) := \left\{ \mathbf{x} = (x_j)_{j \geq 1} : \sum_{j=1}^{\infty} |x_j|^2 < \infty \right\}$, where $\langle \mathbf{x}, \mathbf{y} \rangle := \sum_{j=1}^{\infty} x_j \overline{y_j}$. Since $|x_j \overline{y_j}| \leq (|x_j|^2 + |y_j|^2)/2$, the series converges absolutely; termwise operations give an inner product on $\ell^2(\mathbb{F})$, with $\mathbb{F}^{(\mathbb{N})} \subseteq \ell^2(\mathbb{F})$.

Proposition 2.1.8. *Basic Inner-Product Identities*
Let $\mathcal{V}$ be an inner-product space. Then for all $\mathbf{u}, \mathbf{v}, \mathbf{w} \in \mathcal{V}$ and $\lambda \in \mathbb{F}$,

i. for fixed $\mathbf{u}$, the map $\mathbf{v} \mapsto \langle \mathbf{v}, \mathbf{u} \rangle$ is linear;
ii. $\langle \mathbf{0}, \mathbf{u} \rangle = 0$ and $\langle \mathbf{u}, \mathbf{0} \rangle = 0$;
iii. $\langle \mathbf{u}, \mathbf{v} + \mathbf{w} \rangle = \langle \mathbf{u}, \mathbf{v} \rangle + \langle \mathbf{u}, \mathbf{w} \rangle$;
iv. $\langle \mathbf{u}, \lambda \mathbf{v} \rangle = \overline{\lambda} \langle \mathbf{u}, \mathbf{v} \rangle$.

Proof.
(i) This is exactly additivity and homogeneity in the first slot.
(ii) A linear map sends $\mathbf{0}$ to 0, so $\langle \mathbf{0}, \mathbf{u} \rangle = 0$. Conjugate symmetry gives $\langle \mathbf{u}, \mathbf{0} \rangle = \overline{\langle \mathbf{0}, \mathbf{u} \rangle} = 0$.
(iii) By conjugate symmetry and additivity, $\langle \mathbf{u}, \mathbf{v} + \mathbf{w} \rangle = \overline{\langle \mathbf{v} + \mathbf{w}, \mathbf{u} \rangle} = \overline{\langle \mathbf{v}, \mathbf{u} \rangle + \langle \mathbf{w}, \mathbf{u} \rangle} = \langle \mathbf{u}, \mathbf{v} \rangle + \langle \mathbf{u}, \mathbf{w} \rangle$.
(iv) By conjugate symmetry and homogeneity, $\langle \mathbf{u}, \lambda \mathbf{v} \rangle = \overline{\langle \lambda \mathbf{v}, \mathbf{u} \rangle} = \overline{\lambda \langle \mathbf{v}, \mathbf{u} \rangle} = \overline{\lambda} \langle \mathbf{u}, \mathbf{v} \rangle$.
Structural Commentary: First-slot linearity and conjugate symmetry force conjugate linearity in second slot.
Key Insight: Once one slot is linear, the other slot is forced to be conjugate linear. □

Lemma 2.1.9. *Inner Products Separate Vectors*
For $\mathbf{u}, \mathbf{v}, \mathbf{w} \in \mathcal{V}$, $(\forall \mathbf{x} \in \mathcal{V},\ \langle \mathbf{x}, \mathbf{u} \rangle = \langle \mathbf{x}, \mathbf{v} \rangle) \Longrightarrow \mathbf{u} = \mathbf{v}$, and $(\forall \mathbf{x} \in \mathcal{V},\ \langle \mathbf{x}, \mathbf{w} \rangle = 0) \Longleftrightarrow \mathbf{w} = \mathbf{0}$.

Proof. Set $\mathbf{x} = \mathbf{u} - \mathbf{v}$. Then $\|\mathbf{u} - \mathbf{v}\|^2 = \langle \mathbf{u} - \mathbf{v}, \mathbf{u} - \mathbf{v} \rangle = 0$, hence $\mathbf{u} = \mathbf{v}$.
Taking $\mathbf{u} = \mathbf{w}$ and $\mathbf{v} = \mathbf{0}$ gives the second assertion; the reverse implication is immediate.
Structural Commentary: Testing against all vectors detects the vector being tested.
Key Insight: A vector orthogonal to every vector, including itself, must be zero. □

Definition 2.1.10. *Norm Induced by an Inner Product*
For $\mathbf{v} \in \mathcal{V}$, the *norm* of $\mathbf{v}$ is $\|\mathbf{v}\| := \sqrt{\langle \mathbf{v}, \mathbf{v} \rangle}$.

Example 2.1.11. *Norms from Standard Inner Products*
In $\mathbb{F}^n$, the norm from the standard inner product is the *Euclidean norm* $\|(x_1, \ldots, x_n)\|_2 = \sqrt{|x_1|^2 + \cdots + |x_n|^2}$.
In $C_{\mathbb{F}}[a, b]$, with the integral inner product, the derived norm is the *L^2 norm* $\|f\|_2 = \sqrt{\int_a^b |f(t)|^2\, dt}$.
In $\mathrm{Mat}(m, n, \mathbb{F})$, with the Frobenius inner product, the *Frobenius norm* is $\|A\|_F = \sqrt{\langle A, A \rangle_F} = \sqrt{\sum_{i=1}^{m} \sum_{j=1}^{n} |a_{ij}|^2}$, $A = [a_{ij}]$. The Frobenius norm is also called the *Schur norm* or *Hilbert-Schmidt norm*.

Proposition 2.1.12. *Basic Properties of the Derived Norm*
Let $\mathcal{V}$ be an inner-product space with derived norm $\|\cdot\|$. For all $\mathbf{v} \in \mathcal{V}$ and $\lambda \in \mathbb{F}$,

i. $\|\mathbf{v}\|$ is real and nonnegative;
ii. $\|\mathbf{v}\| = 0$ if and only if $\mathbf{v} = \mathbf{0}$;
iii. $\|\lambda \mathbf{v}\| = |\lambda| \, \|\mathbf{v}\|$.

Proof.
(i) By positivity, $\langle \mathbf{v}, \mathbf{v} \rangle$ is real and nonnegative, so its square root is real and nonnegative.

(ii) By definition, $\|\mathbf{v}\| = 0$ if and only if $\langle \mathbf{v}, \mathbf{v} \rangle = 0$, which holds if and only if $\mathbf{v} = \mathbf{0}$ by definiteness.
(iii) Compute $\|\lambda \mathbf{v}\|^2 = \langle \lambda \mathbf{v}, \lambda \mathbf{v} \rangle = \lambda \overline{\lambda} \langle \mathbf{v}, \mathbf{v} \rangle = |\lambda|^2 \|\mathbf{v}\|^2$. Taking nonnegative square roots, $\|\lambda \mathbf{v}\| = |\lambda| \|\mathbf{v}\|$.

Structural Commentary: Positivity gives length; sesquilinearity gives scaling.
Key Insight: Squared length is easier to manipulate than length itself. □

Theorem 2.1.13. *Cauchy-Schwarz Inequality*
For all $\mathbf{u}, \mathbf{v} \in \mathcal{V}$, $|\langle \mathbf{u}, \mathbf{v} \rangle| \leq \|\mathbf{u}\| \|\mathbf{v}\|$. Equality holds if and only if one of $\mathbf{u}, \mathbf{v}$ is a scalar multiple of the other.

Proof. If $\mathbf{v} = \mathbf{0}$, both sides are 0. Assume $\mathbf{v} \neq \mathbf{0}$.
For any $\lambda \in \mathbb{F}$, positivity gives $0 \leq \|\mathbf{u} - \lambda \mathbf{v}\|^2 = \|\mathbf{u}\|^2 - \lambda \langle \mathbf{v}, \mathbf{u} \rangle - \overline{\lambda} \langle \mathbf{u}, \mathbf{v} \rangle + |\lambda|^2 \|\mathbf{v}\|^2$.
Choose $\lambda = \frac{\langle \mathbf{u}, \mathbf{v} \rangle}{\|\mathbf{v}\|^2}$. Since $\langle \mathbf{v}, \mathbf{u} \rangle = \overline{\langle \mathbf{u}, \mathbf{v} \rangle}$, the inequality becomes $0 \leq \|\mathbf{u}\|^2 - \frac{|\langle \mathbf{u}, \mathbf{v} \rangle|^2}{\|\mathbf{v}\|^2}$.
Multiplying by $\|\mathbf{v}\|^2$ and taking square roots proves the inequality. Equality in this argument is equivalent to $\|\mathbf{u} - \lambda \mathbf{v}\| = 0$, hence to $\mathbf{u} = \lambda \mathbf{v}$ for the chosen λ.
Conversely, if $\mathbf{u} = \lambda \mathbf{v}$, then $|\langle \mathbf{u}, \mathbf{v} \rangle| = |\lambda| \langle \mathbf{v}, \mathbf{v} \rangle = \|\lambda \mathbf{v}\| \|\mathbf{v}\| = \|\mathbf{u}\| \|\mathbf{v}\|$, so equality holds.
The case $\mathbf{v} = \mathbf{0}$ is also covered because $\mathbf{v} = 0\mathbf{u}$.
Structural Commentary: Positivity of squared length controls the mixed inner-product term.
Key Insight: The best scalar comparison between two vectors bounds their inner product. □

Theorem 2.1.14. *Triangle Inequality*
For all $\mathbf{u}, \mathbf{v} \in \mathcal{V}$, $\|\mathbf{u} + \mathbf{v}\| \leq \|\mathbf{u}\| + \|\mathbf{v}\|$.
Equality holds if and only if one of $\mathbf{u}, \mathbf{v}$ is a nonnegative scalar multiple of the other.

Proof. We compute

$$\|\mathbf{u} + \mathbf{v}\|^2 = \langle \mathbf{u} + \mathbf{v}, \mathbf{u} + \mathbf{v} \rangle = \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2 + \langle \mathbf{u}, \mathbf{v} \rangle + \langle \mathbf{v}, \mathbf{u} \rangle = \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2 + 2 \operatorname{Re} \langle \mathbf{u}, \mathbf{v} \rangle$$
$$\leq \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2 + 2 |\langle \mathbf{u}, \mathbf{v} \rangle| \leq \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2 + 2 \|\mathbf{u}\| \|\mathbf{v}\| = \big(\|\mathbf{u}\| + \|\mathbf{v}\|\big)^2.$$

Taking square roots gives the inequality. If either vector is $\mathbf{0}$, the equality condition is immediate.
Otherwise, equality in the two displayed inequalities holds exactly when $\operatorname{Re} \langle \mathbf{u}, \mathbf{v} \rangle = |\langle \mathbf{u}, \mathbf{v} \rangle|$ and equality holds in Cauchy-Schwarz. The first condition says $\langle \mathbf{u}, \mathbf{v} \rangle$ is real nonnegative, and the second says $\mathbf{u} = a\mathbf{v}$ for some $a \in \mathbb{F}$. Then $a \|\mathbf{v}\|^2 = \langle \mathbf{u}, \mathbf{v} \rangle$ is real nonnegative, so $a \geq 0$.
Conversely, if $\mathbf{u} = a\mathbf{v}$ with $a \geq 0$, then $\|\mathbf{u} + \mathbf{v}\| = (a + 1) \|\mathbf{v}\| = \|\mathbf{u}\| + \|\mathbf{v}\|$.
Structural Commentary: Square the norm and bound the cross term by Cauchy-Schwarz.
Key Insight: Equality means the two vectors point in the same direction. □

Proposition 2.1.15. *Reverse Triangle Inequality*
For all $\mathbf{u}, \mathbf{v} \in \mathcal{V}$, $|\|\mathbf{u}\| - \|\mathbf{v}\|| \leq \|\mathbf{u} - \mathbf{v}\|$.

Proof. The triangle inequality gives $\|\mathbf{u}\| \leq \|\mathbf{u} - \mathbf{v}\| + \|\mathbf{v}\|$, hence $\|\mathbf{u}\| - \|\mathbf{v}\| \leq \|\mathbf{u} - \mathbf{v}\|$.
Interchanging $\mathbf{u}$ and $\mathbf{v}$ gives $\|\mathbf{v}\| - \|\mathbf{u}\| \leq \|\mathbf{u} - \mathbf{v}\|$. Combining the two inequalities gives the result.
Structural Commentary: Distances control changes in length.
Key Insight: The ordinary triangle inequality applied twice gives the reverse one. □

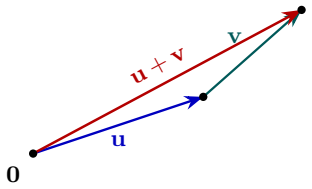


Figure 20: Triangle inequality: $\|\mathbf{u} + \mathbf{v}\| \leq \|\mathbf{u}\| + \|\mathbf{v}\|$.

Definition 2.1.16. *Norm and Normed Vector Space*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$. A *norm* on $\mathcal{V}$ is a function $\|\cdot\| : \mathcal{V} \to [0, \infty)$ such that, $\forall \mathbf{u}, \mathbf{v} \in \mathcal{V}$ and $\lambda \in \mathbb{F}$,

i. $\|\mathbf{u}\| \geq 0$;
ii. $\|\mathbf{u}\| = 0$ if and only if $\mathbf{u} = \mathbf{0}$;
iii. $\|\lambda \mathbf{u}\| = |\lambda| \|\mathbf{u}\|$;
iv. $\|\mathbf{u} + \mathbf{v}\| \leq \|\mathbf{u}\| + \|\mathbf{v}\|$.

A *normed vector space* is a vector space equipped with a norm. Its *unit ball* is $\{\mathbf{v} \in \mathcal{V} : \|\mathbf{v}\| \leq 1\}$.

Definition 2.1.17. *Distance Induced by an Inner Product*
Let $\mathcal{V}$ be an inner-product space. The *distance* between two vectors in $\mathcal{V}$ is $d(\mathbf{u}, \mathbf{v}) := \|\mathbf{u} - \mathbf{v}\|$.

Proposition 2.1.18. *Inner-Product Distance Is a Metric*
The function $d(\mathbf{u},\mathbf{v}) = \|\mathbf{u}-\mathbf{v}\|$ *is a metric on* $\mathcal{V}$*: for all* $\mathbf{u},\mathbf{v},\mathbf{w} \in \mathcal{V}$,

i. $d(\mathbf{u},\mathbf{v}) \geq 0$, *and* $d(\mathbf{u},\mathbf{v}) = 0$ *if and only if* $\mathbf{u} = \mathbf{v}$;
ii. $d(\mathbf{u},\mathbf{v}) = d(\mathbf{v},\mathbf{u})$;
iii. $d(\mathbf{u},\mathbf{v}) \leq d(\mathbf{u},\mathbf{w}) + d(\mathbf{w},\mathbf{v})$.

Proof.

(i) By Proposition 2.1.12, $d(\mathbf{u},\mathbf{v}) = \|\mathbf{u}-\mathbf{v}\| \geq 0$, with equality iff $\mathbf{u}-\mathbf{v} = \mathbf{0}$.
(ii) $d(\mathbf{u},\mathbf{v}) = \|\mathbf{u}-\mathbf{v}\| = \|-(\mathbf{v}-\mathbf{u})\| = \|\mathbf{v}-\mathbf{u}\| = d(\mathbf{v},\mathbf{u})$.
(iii) Since $\mathbf{u}-\mathbf{v} = (\mathbf{u}-\mathbf{w}) + (\mathbf{w}-\mathbf{v})$, Theorem 2.1.14 gives $d(\mathbf{u},\mathbf{v}) \leq d(\mathbf{u},\mathbf{w}) + d(\mathbf{w},\mathbf{v})$.

Structural Commentary: The norm turns displacement vectors into distances.
Key Insight: Metric triangle inequalities are norm triangle inequalities applied to differences. □

Perspective 2.1.19. *Metric Viewpoint*
Define $d(\mathbf{u},\mathbf{v}) = \|\mathbf{u}-\mathbf{v}\|$. Then $\mathbf{v}_n \to \mathbf{v}$ iff $\|\mathbf{v}_n - \mathbf{v}\| \to 0$.
In finite-dimensional inner-product spaces, subspaces are closed, linear maps are continuous, and the ambient space is complete. Infinite-dimensional analogues require additional hypotheses; subsequent projection and basis results therefore state finite-dimensional assumptions explicitly.

Theorem 2.1.20. *Parallelogram Identity*
For all $\mathbf{u},\mathbf{v} \in \mathcal{V}$, $\|\mathbf{u}+\mathbf{v}\|^2 + \|\mathbf{u}-\mathbf{v}\|^2 = 2\big(\|\mathbf{u}\|^2 + \|\mathbf{v}\|^2\big)$.

Proof. Expanding both squared norms gives

$$\|\mathbf{u}+\mathbf{v}\|^2 = \|\mathbf{u}\|^2 + \langle \mathbf{u},\mathbf{v}\rangle + \langle \mathbf{v},\mathbf{u}\rangle + \|\mathbf{v}\|^2, \quad \|\mathbf{u}-\mathbf{v}\|^2 = \|\mathbf{u}\|^2 - \langle \mathbf{u},\mathbf{v}\rangle - \langle \mathbf{v},\mathbf{u}\rangle + \|\mathbf{v}\|^2.$$

Adding the two identities cancels the cross terms and yields the result.
Structural Commentary: Opposite cross terms cancel when the two diagonal norms are added.
Key Insight: Inner-product norms remember parallelogram geometry through cancellation. □

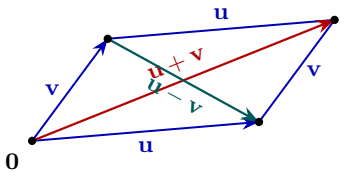


Figure 21: Parallelogram identity: $\|\mathbf{u}+\mathbf{v}\|^2 + \|\mathbf{u}-\mathbf{v}\|^2 = 2(\|\mathbf{u}\|^2 + \|\mathbf{v}\|^2)$.

Proposition 2.1.21. *Coordinate Norms on* $\mathbb{F}^n$
For $\mathbf{x} = (x_1,\ldots,x_n) \in \mathbb{F}^n$, $\|\mathbf{x}\|_1 = \sum_{j=1}^{n} |x_j|$, $\|\mathbf{x}\|_2 = \left(\sum_{j=1}^{n} |x_j|^2\right)^{1/2}$, $\|\mathbf{x}\|_\infty = \max_{1\leq j\leq n} |x_j|$ *are norms on* $\mathbb{F}^n$.

Proof.

(i) For $\|\cdot\|_1$, nonnegativity, definiteness, and homogeneity follow from the corresponding facts for modulus. The triangle inequality follows from $\|\mathbf{x}+\mathbf{y}\|_1 = \sum_{j=1}^{n} |x_j + y_j| \leq \sum_{j=1}^{n} (|x_j| + |y_j|) = \|\mathbf{x}\|_1 + \|\mathbf{y}\|_1$.
(ii) For $\|\cdot\|_2$, the axioms follow from its standard inner product via Proposition 2.1.12 and Theorem 2.1.14.
(iii) For $\|\cdot\|_\infty$, nonnegativity, definiteness, and homogeneity are immediate from the maximum of finitely many moduli. If k is chosen so that $\|\mathbf{x}+\mathbf{y}\|_\infty = |x_k + y_k|$, then $\|\mathbf{x}+\mathbf{y}\|_\infty \leq |x_k| + |y_k| \leq \|\mathbf{x}\|_\infty + \|\mathbf{y}\|_\infty$.

Structural Commentary: The ℓ_1 norm sums coordinate magnitudes, ℓ_∞ keeps only the largest coordinate magnitude, and ℓ_2 comes from the standard inner product.
Key Insight: All three obey the same abstract axioms, even though their unit balls have different shapes. □

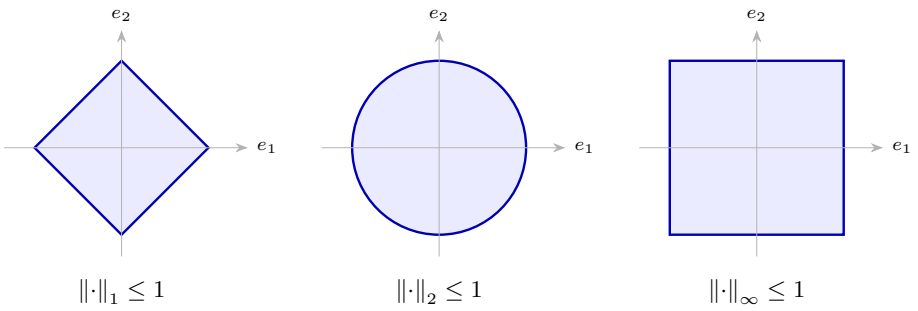


Figure 22: The unit balls of $\|\cdot\|_1$, $\|\cdot\|_2$, and $\|\cdot\|_\infty$ in $\mathbb{R}^2$.

Corollary 2.1.22. *ℓ_1 and ℓ_∞ Are Not Inner-Product Norms*
If $n \geq 2$, then the norms $\|\cdot\|_1$ and $\|\cdot\|_\infty$ on $\mathbb{F}^n$ are not induced by any inner product.

Proof. Let $\mathbf{e}_1, \mathbf{e}_2$ be the first two standard coordinate vectors.
For $\|\cdot\|_1$, $\|\mathbf{e}_1 + \mathbf{e}_2\|_1^2 + \|\mathbf{e}_1 - \mathbf{e}_2\|_1^2 = 4 + 4 = 8$, but $2\|\mathbf{e}_1\|_1^2 + 2\|\mathbf{e}_2\|_1^2 = 4$.
For $\|\cdot\|_\infty$, $\|\mathbf{e}_1 + \mathbf{e}_2\|_\infty^2 + \|\mathbf{e}_1 - \mathbf{e}_2\|_\infty^2 = 1 + 1 = 2$ but $2\|\mathbf{e}_1\|_\infty^2 + 2\|\mathbf{e}_2\|_\infty^2 = 4$.
Both norms violate Theorem 2.1.20, which every inner-product norm must satisfy.
Structural Commentary: The parallelogram identity detects inner-product norms.
Key Insight: Different unit-ball geometry can rule out hidden inner-product structure. □

Theorem 2.1.23. *Polarization Identities*
Let $\mathcal{V}$ be an inner-product space and let $\mathbf{u}, \mathbf{v} \in \mathcal{V}$.

i. If $\mathbb{F} = \mathbb{R}$, then $\langle \mathbf{u}, \mathbf{v} \rangle = \frac{1}{4}\left(\|\mathbf{u} + \mathbf{v}\|^2 - \|\mathbf{u} - \mathbf{v}\|^2\right)$.

ii. If $\mathbb{F} = \mathbb{C}$, then $\langle \mathbf{u}, \mathbf{v} \rangle = \frac{1}{4}\left(\|\mathbf{u} + \mathbf{v}\|^2 - \|\mathbf{u} - \mathbf{v}\|^2 + i\|\mathbf{u} + i\mathbf{v}\|^2 - i\|\mathbf{u} - i\mathbf{v}\|^2\right)$.

Proof. Let $z := \langle \mathbf{u}, \mathbf{v} \rangle$.

(i) If $\mathbb{F} = \mathbb{R}$, then expanding squared norms gives $\|\mathbf{u} + \mathbf{v}\|^2 - \|\mathbf{u} - \mathbf{v}\|^2 = 4 \operatorname{Re} z = 4z$, proving (i).

(ii) If $\mathbb{F} = \mathbb{C}$, the first expansion gives $\|\mathbf{u} + \mathbf{v}\|^2 - \|\mathbf{u} - \mathbf{v}\|^2 = 4 \operatorname{Re} z$. The same expansion with $i\mathbf{v}$ gives $\|\mathbf{u} + i\mathbf{v}\|^2 - \|\mathbf{u} - i\mathbf{v}\|^2 = 4 \operatorname{Im} z$. Combining the real and imaginary parts gives (ii).

Structural Commentary: Two norm differences extract the real and imaginary parts.
Key Insight: The norm determines the inner product because it determines both parts of every cross term. □

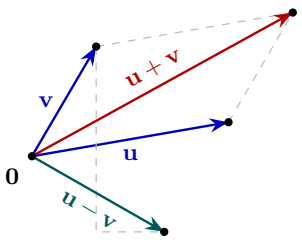


Figure 23: Polarization identity: $\|\mathbf{u} + \mathbf{v}\|^2 - \|\mathbf{u} - \mathbf{v}\|^2 = 4 \operatorname{Re} \langle \mathbf{u}, \mathbf{v} \rangle$.

Proposition 2.1.24. *Operator Zero Tests from Inner Products*
Let $T \in \mathcal{L}(\mathcal{V})$.

i. If $\langle T\mathbf{v}, \mathbf{w} \rangle = 0$ for every $\mathbf{v}, \mathbf{w} \in \mathcal{V}$, then $T = 0$.

ii. If $\mathbb{F} = \mathbb{C}$ and $\langle T\mathbf{v}, \mathbf{v} \rangle = 0$ for every $\mathbf{v} \in \mathcal{V}$, then $T = 0$.

iii. The preceding diagonal test can fail over $\mathbb{R}$.

Proof.

(i) Taking $\mathbf{w} = T\mathbf{v}$ gives $\|T\mathbf{v}\|^2 = 0$ for every $\mathbf{v}$.

(ii) Put $Q(\mathbf{x}) = \langle T\mathbf{x}, \mathbf{x} \rangle$. For $\mathbf{x}, \mathbf{y} \in \mathcal{V}$ and $r \in \mathbb{C}$, the hypothesis applied to $r\mathbf{x} + \mathbf{y}$ gives

$$0 = Q(r\mathbf{x} + \mathbf{y}) = |r|^2 Q(\mathbf{x}) + Q(\mathbf{y}) + r\langle T\mathbf{x}, \mathbf{y} \rangle + \overline{r}\langle T\mathbf{y}, \mathbf{x} \rangle = r\langle T\mathbf{x}, \mathbf{y} \rangle + \overline{r}\langle T\mathbf{y}, \mathbf{x} \rangle .$$

Taking $r = 1, i$ gives $\langle T\mathbf{x}, \mathbf{y} \rangle = 0$ for all $\mathbf{x}, \mathbf{y}$; by (i), $T = 0$.

(iii) On standard $\mathbb{R}^2$, a 90° rotation has $\langle T\mathbf{v}, \mathbf{v} \rangle = 0$ for all $\mathbf{v}$, yet $T \neq 0$.

Structural Commentary: Complex scalar directions recover mixed inner products from diagonal values.
Key Insight: Over $\mathbb{R}$, a skew-symmetric operator can hide on the diagonal; over $\mathbb{C}$, polarization detects it. □

Principle 2.1.25. *Geometry from Algebra*
Most geometric statements in an inner-product space come from expanding $\|\mathbf{u} \pm \mathbf{v}\|^2$ and then interpreting the cross terms $\langle \mathbf{u}, \mathbf{v} \rangle$ and $\langle \mathbf{v}, \mathbf{u} \rangle$.

Definition 2.1.26. *Unit Vector and Normalization*
A vector $\mathbf{u} \in \mathcal{V}$ is a *unit vector* if $\|\mathbf{u}\| = 1$. If $\mathbf{v} \neq \mathbf{0}$, then $\frac{\mathbf{v}}{\|\mathbf{v}\|}$ is called the *normalization* of $\mathbf{v}$.

Proposition 2.1.27. *Normalization Gives a Unit Vector*
If $\mathbf{v} \in \mathcal{V}$ and $\mathbf{v} \neq \mathbf{0}$, then $\frac{\mathbf{v}}{\|\mathbf{v}\|}$ is a unit vector in the same direction as $\mathbf{v}$.

Proof. By Proposition 2.1.12, $\|\mathbf{v}\| > 0$ and $\left\| \frac{\mathbf{v}}{\|\mathbf{v}\|} \right\| = \frac{1}{\|\mathbf{v}\|} \|\mathbf{v}\| = 1$.
The vector $\frac{\mathbf{v}}{\|\mathbf{v}\|}$ is a positive scalar multiple of $\mathbf{v}$, so it lies in the same one-dimensional direction. *Structural Commentary*: Normalization changes scale but preserves direction.
Key Insight: Divide by the length to make the length equal to one. □

Definition 2.1.28. *Orthogonal Vectors*
Vectors $\mathbf{u}, \mathbf{v} \in \mathcal{V}$ are *orthogonal*, written $\mathbf{u} \perp \mathbf{v}$, if $\langle \mathbf{u}, \mathbf{v} \rangle = 0$.

Proposition 2.1.29. *Orthogonality and the Zero Vector*
Let $\mathcal{V}$ be an inner-product space.

i. $\mathbf{u} \perp \mathbf{v}$ *if and only if* $\mathbf{v} \perp \mathbf{u}$.
ii. $\mathbf{0}$ *is orthogonal to every vector in* $\mathcal{V}$.
iii. If $\mathbf{v}$ *is orthogonal to itself, then* $\mathbf{v} = \mathbf{0}$.

Proof.

(i) If $\langle \mathbf{u}, \mathbf{v} \rangle = 0$, then conjugate symmetry gives $\langle \mathbf{v}, \mathbf{u} \rangle = \overline{0} = 0$, and the converse is identical.
(ii) This follows from Proposition 2.1.8, which gives $\langle \mathbf{0}, \mathbf{u} \rangle = 0$ and $\langle \mathbf{u}, \mathbf{0} \rangle = 0$.
(iii) If $\mathbf{v} \perp \mathbf{v}$, then $\langle \mathbf{v}, \mathbf{v} \rangle = 0$, hence $\mathbf{v} = \mathbf{0}$ by definiteness.

Structural Commentary: Conjugate symmetry makes orthogonality symmetric, and definiteness isolates $\mathbf{0}$.
Key Insight: Conjugate symmetry preserves the equation "equals zero." □

Proposition 2.1.30. *Mutually Orthogonal Nonzero Vectors Are Independent*
Mutually orthogonal nonzero vectors $\mathbf{v}_1, \ldots, \mathbf{v}_m$ in an inner-product space are linearly independent.

Proof. If $\sum_{j=1}^{m} c_j \mathbf{v}_j = \mathbf{0}$, then for each k, $0 = \left\langle \sum_{j=1}^{m} c_j \mathbf{v}_j, \mathbf{v}_k \right\rangle = c_k \langle \mathbf{v}_k, \mathbf{v}_k \rangle$.
Since $\langle \mathbf{v}_k, \mathbf{v}_k \rangle > 0$, $c_k = 0$. Thus all coefficients vanish.
Structural Commentary: Orthogonality isolates one coefficient at a time.
Key Insight: Pair a dependence relation with one vector from the list; every cross term disappears. □

Definition 2.1.31. *Orthonormal List*
A list $(\mathbf{e}_1, \ldots, \mathbf{e}_m)$ of vectors in an inner-product space $\mathcal{V}$ is *orthonormal* if

$$\langle \mathbf{e}_j, \mathbf{e}_k \rangle = \begin{cases} 1, & j = k, \\ 0, & j \neq k. \end{cases}$$

Equivalently, every $\mathbf{e}_j$ has norm 1 and distinct vectors in the list are orthogonal.

Definition 2.1.32. *Orthogonal and Orthonormal Sets*
An indexed family $(\mathbf{u}_i)_{i \in I}$ is *orthogonal* if $\langle \mathbf{u}_i, \mathbf{u}_j \rangle = 0$ for $i \neq j$, and *orthonormal* if $\langle \mathbf{u}_i, \mathbf{u}_j \rangle = \delta_{ij}$, $i, j \in I$.
A nonzero orthogonal family becomes orthonormal after replacing each $\mathbf{u}_i$ by $\mathbf{u}_i / \|\mathbf{u}_i\|$.

Proposition 2.1.33. *Norm of an Orthonormal Linear Combination*
If $(\mathbf{e}_1, \ldots, \mathbf{e}_m)$ is an orthonormal list in $\mathcal{V}$, then $\|a_1 \mathbf{e}_1 + \cdots + a_m \mathbf{e}_m\|^2 = |a_1|^2 + \cdots + |a_m|^2$ for all $a_1, \ldots, a_m \in \mathbb{F}$.

Proof. Using orthonormality, $\left\| \sum_{j=1}^{m} a_j \mathbf{e}_j \right\|^2 = \left\langle \sum_{j=1}^{m} a_j \mathbf{e}_j, \sum_{k=1}^{m} a_k \mathbf{e}_k \right\rangle = \sum_{j=1}^{m} \sum_{k=1}^{m} a_j \overline{a_k} \langle \mathbf{e}_j, \mathbf{e}_k \rangle = \sum_{j=1}^{m} |a_j|^2$.
Structural Commentary: Orthonormality deletes every cross term in the squared-norm expansion.
Key Insight: In orthonormal coordinates, length is computed coefficient-by-coefficient. □

Corollary 2.1.34. *Orthonormal Lists Are Linearly Independent*
Every orthonormal list of vectors in an inner-product space is linearly independent.

Proof. If $a_1\mathbf{e}_1 + \cdots + a_m\mathbf{e}_m = \mathbf{0}$, then Proposition 2.1.33 gives $0 = \|\mathbf{0}\|^2 = \|a_1\mathbf{e}_1 + \cdots + a_m\mathbf{e}_m\|^2 = \sum_{j=1}^{m} |a_j|^2$.
Hence each $a_j = 0$, so the list is linearly independent.
Structural Commentary: The norm identity converts a linear relation into a sum of nonnegative scalars.
Key Insight: A zero orthonormal linear combination can only have zero coefficients. □

Definition 2.1.35. *Orthonormal Basis*
An *orthonormal basis* of $\mathcal{V}$ is a basis of $\mathcal{V}$ that is an orthonormal list.

Proposition 2.1.36. *Right-Length Orthonormal Lists Are Bases*
If $\mathcal{V}$ is finite-dimensional and $(\mathbf{e}_1, \ldots, \mathbf{e}_n)$ is an orthonormal list in $\mathcal{V}$ with $n = \dim(\mathcal{V})$, then $(\mathbf{e}_1, \ldots, \mathbf{e}_n)$ is an orthonormal basis of $\mathcal{V}$.

Proof. By Corollary 2.1.34, the list is linearly independent.
Since its length is $\dim(\mathcal{V})$, Proposition 1.1.32 implies that it is a basis.
Structural Commentary: Orthonormality gives independence; the correct length forces spanning.
Key Insight: In finite dimensions, n independent vectors in an n-dimensional space already form a basis. □

Theorem 2.1.37. *Gram–Schmidt Process*
Let $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ be a linearly independent list in an inner-product space $\mathcal{V}$.
Define $\mathbf{u}_1 = \mathbf{v}_1$, $\mathbf{e}_1 = \frac{\mathbf{u}_1}{\|\mathbf{u}_1\|}$ and, for $2 \le j \le m$, $\mathbf{u}_j = \mathbf{v}_j - \sum_{k=1}^{j-1} \langle \mathbf{v}_j, \mathbf{e}_k \rangle \mathbf{e}_k$, $\mathbf{e}_j = \frac{\mathbf{u}_j}{\|\mathbf{u}_j\|}$.
Then $(\mathbf{e}_1, \ldots, \mathbf{e}_m)$ is orthonormal and $\operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_j) = \operatorname{span}(\mathbf{e}_1, \ldots, \mathbf{e}_j)$.

Proof. Proceed by induction on j. The case $j = 1$ is normalization of the nonzero vector $\mathbf{v}_1$.
Assume the result holds through $j-1$. Since $\operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_{j-1}) = \operatorname{span}(\mathbf{e}_1, \ldots, \mathbf{e}_{j-1})$ and the $\mathbf{v}$-list is linearly independent, $\mathbf{v}_j \notin \operatorname{span}(\mathbf{e}_1, \ldots, \mathbf{e}_{j-1})$. Hence $\mathbf{u}_j \neq \mathbf{0}$, because $\mathbf{u}_j = \mathbf{0}$ would put $\mathbf{v}_j$ in that span.
For $k < j$, $\langle \mathbf{u}_j, \mathbf{e}_k \rangle = \langle \mathbf{v}_j, \mathbf{e}_k \rangle - \sum_{i=1}^{j-1} \langle \mathbf{v}_j, \mathbf{e}_i \rangle \langle \mathbf{e}_i, \mathbf{e}_k \rangle = \langle \mathbf{v}_j, \mathbf{e}_k \rangle - \langle \mathbf{v}_j, \mathbf{e}_k \rangle = 0$.
Thus $\mathbf{e}_j$ is a unit vector orthogonal to $\mathbf{e}_1, \ldots, \mathbf{e}_{j-1}$.
Also $\mathbf{e}_j$ lies in $\operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_j)$ because it is a nonzero scalar multiple of $\mathbf{u}_j$, and $\mathbf{u}_j$ is built from $\mathbf{v}_j$ and the previous $\mathbf{e}_k$, whose span agrees with $\operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_{j-1})$.
Conversely, $\mathbf{v}_j = \mathbf{u}_j + \sum_{k=1}^{j-1} \langle \mathbf{v}_j, \mathbf{e}_k \rangle \mathbf{e}_k$ lies in $\operatorname{span}(\mathbf{e}_1, \ldots, \mathbf{e}_j)$.
Combining this with the induction hypothesis gives equality of the two spans at step j.
Structural Commentary: Subtract the previous-span component, then normalize.
Key Insight: Gram–Schmidt changes the basis vectors but preserves every initial span. □

Theorem 2.1.38. *Gram–Schmidt for Arbitrary Finite Lists*
Let $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ be a finite list in an inner-product space.
Define $\mathbf{u}_1 = \mathbf{v}_1$ and, for $2 \le k \le m$, $\mathbf{u}_k = \mathbf{v}_k - \sum_{\substack{1 \le j < k \\ \mathbf{u}_j \neq \mathbf{0}}} \frac{\langle \mathbf{v}_k, \mathbf{u}_j \rangle}{\langle \mathbf{u}_j, \mathbf{u}_j \rangle} \mathbf{u}_j$.
Set $U_k = \operatorname{span}(\mathbf{u}_1, \ldots, \mathbf{u}_k)$ and $V_k = \operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_k)$, with $U_0 = V_0 = \{\mathbf{0}\}$.
Then $(\mathbf{u}_1, \ldots, \mathbf{u}_m)$ is orthogonal, $U_k = V_k$ for $1 \le k \le m$.

Proof. Induct on k. The case $k = 0$ is trivial. Assume $(\mathbf{u}_1, \ldots, \mathbf{u}_{k-1})$ is orthogonal and $U_{k-1} = V_{k-1}$.
For $i < k$ with $\mathbf{u}_i \neq \mathbf{0}$, $\langle \mathbf{u}_k, \mathbf{u}_i \rangle = \langle \mathbf{v}_k, \mathbf{u}_i \rangle - \frac{\langle \mathbf{v}_k, \mathbf{u}_i \rangle}{\langle \mathbf{u}_i, \mathbf{u}_i \rangle} \langle \mathbf{u}_i, \mathbf{u}_i \rangle = 0$, while $\mathbf{u}_i = \mathbf{0}$ is trivial.
Hence the list remains orthogonal. Also $\mathbf{u}_k \in V_k$ and $\mathbf{v}_k \in U_k$, so $U_k = V_k$.
Structural Commentary: Gram–Schmidt preserves every prefix span.
Key Insight: Each step subtracts previous orthogonal directions. □

Corollary 2.1.39. *Zero Remainders Detect Dependence*
Let $(\mathbf{u}_1, \ldots, \mathbf{u}_m)$ be constructed from $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ by Theorem 2.1.38, and set $U_k = \operatorname{span}(\mathbf{u}_1, \ldots, \mathbf{u}_k)$, $V_k = \operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_k)$, with $U_0 = V_0 = \{\mathbf{0}\}$.
Then $\mathbf{u}_k = \mathbf{0} \iff \mathbf{v}_k \in V_{k-1}$, $1 \le k \le m$.

Proof. The case $k = 1$ follows from $\mathbf{u}_1 = \mathbf{v}_1$ and $V_0 = \{\mathbf{0}\}$. Let $2 \le k \le m$.
($\Rightarrow$) If $\mathbf{u}_k = \mathbf{0}$, the recursion gives $\mathbf{v}_k \in U_{k-1} = V_{k-1}$.

($\Leftarrow$) If $\mathbf{v}_k \in V_{k-1} = U_{k-1}$, write $\mathbf{v}_k = \sum_{\substack{1\le j<k \\ \mathbf{u}_j \neq \mathbf{0}}} a_j \mathbf{u}_j$. Orthogonality gives $a_j = \frac{\langle \mathbf{v}_k, \mathbf{u}_j \rangle}{\langle \mathbf{u}_j, \mathbf{u}_j \rangle}$, recursion gives $\mathbf{u}_k = \mathbf{0}$.
Structural Commentary: Zero remainders certify dependence.
Key Insight: A nonzero remainder is a new direction. □

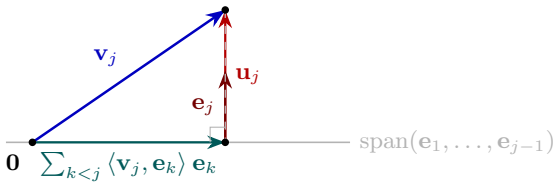


Figure 24: One Gram–Schmidt step: subtract the previous-span component, then normalize.

Example 2.1.40. *Legendre Polynomials from Gram–Schmidt*
Let $\mathbb{R}[x]$ have inner product $\langle p, q \rangle = \int_{-1}^{1} p(x)q(x)\,dx$, and apply Gram–Schmidt to $(1, x, x^2, x^3, \ldots)$.
The first unnormalized vectors are $u_1(x) = 1$, $u_2(x) = x$, $u_3(x) = x^2 - \frac{1}{3}$, $u_4(x) = x^3 - \frac{3}{5}x$.
Thus (u_n) is, up to nonzero scalar multiples, the Legendre sequence.
Structural Commentary: Orthogonal polynomials are orthogonalized monomials.
Key Insight: Changing the inner product changes the polynomial family produced by Gram–Schmidt.

Lemma 2.1.41. *Gram–Schmidt on Orthogonal Lists*
If $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ is a nonzero mutually orthogonal list, then Gram–Schmidt produces $\mathbf{e}_j = \frac{\mathbf{v}_j}{\|\mathbf{v}_j\|}$, $j = 1, \ldots, m$.
In particular, if the original list is orthonormal, Gram–Schmidt leaves it unchanged.

Proof. Assume $\mathbf{e}_i = \mathbf{v}_i / \|\mathbf{v}_i\|$ for $i < j$. Since $\mathbf{v}_j \perp \mathbf{v}_i$, we have $\langle \mathbf{v}_j, \mathbf{e}_i \rangle = 0$ for every $i < j$.
Hence the Gram–Schmidt remainder at step j is $\mathbf{u}_j = \mathbf{v}_j$, so $\mathbf{e}_j = \mathbf{v}_j / \|\mathbf{v}_j\|$.
If the original list is orthonormal, then $\|\mathbf{v}_j\| = 1$ for every j.
Structural Commentary: Orthogonal input has no previous-span component to subtract.
Key Insight: Gram–Schmidt only removes nonorthogonal overlap. □

Corollary 2.1.42. *Existence of Orthonormal Bases*
Every finite-dimensional inner-product space has an orthonormal basis.

Proof. Choose any basis of $\mathcal{V}$ and apply Theorem 2.1.37.
The resulting orthonormal list has the same span as the original basis, hence spans $\mathcal{V}$.
Structural Commentary: Gram–Schmidt orthonormalizes a basis without changing its span.
Key Insight: Finite-dimensional inner-product spaces always admit geometry-friendly coordinates. □

Corollary 2.1.43. *Extending an Orthonormal List*
Every orthonormal list in a finite-dimensional inner-product space can be extended to an orthonormal basis.

Proof. Let $(\mathbf{e}_1, \ldots, \mathbf{e}_m)$ be orthonormal. It is linearly independent by Corollary 2.1.34, so extend it to a basis $(\mathbf{e}_1, \ldots, \mathbf{e}_m, \mathbf{v}_{m+1}, \ldots, \mathbf{v}_n)$ by Theorem 1.1.27.
During Gram–Schmidt, the first m remainders are the original $\mathbf{e}_j$ because the list is already orthonormal; after normalization they remain unchanged. The remaining vectors are orthonormalized by Theorem 2.1.37, and the final list spans $\mathcal{V}$ because Gram–Schmidt preserves the total span.
Structural Commentary: Extend to a basis, then orthonormalize the new part.
Key Insight: Existing orthonormal directions are already cleaned up, so Gram–Schmidt does not move them. □

Perspective 2.1.44. *Hamel and Hilbert Basis Caution*
Here an *orthonormal basis* is a Hamel basis that is orthonormal.
In infinite dimensions, a *Hilbert basis* usually means a maximal orthonormal set, not necessarily a Hamel basis.
For example, the standard $(\mathbf{e}_j)$ in $\ell^2(\mathbb{F})$ is maximal orthonormal, but $\operatorname{span}\{\mathbf{e}_j : j \ge 1\} = \mathbb{F}^{(\mathbb{N})} \subsetneq \ell^2(\mathbb{F})$.

Definition 2.1.45. *Hilbert Basis*

(i) A *Hilbert basis* is a maximal orthonormal set under inclusion.
(ii) A *Hamel basis* is a vector-space basis.
(iii) An *orthonormal basis* is an orthonormal Hamel basis.

Proposition 2.1.46. *Finite Hilbert Bases Are Orthonormal Bases*
Let $\mathcal{O} = \{\mathbf{e}_1, \ldots, \mathbf{e}_m\}$ be orthonormal in $\mathcal{V}$ and set $W = \operatorname{span}(\mathcal{O})$.
Then $\mathcal{O}$ is a Hilbert basis $\iff W = \mathcal{V} \iff \mathcal{O}$ is an orthonormal Hamel basis.

Proof. Since $\mathcal{O}$ is orthonormal, Corollary 2.1.34 gives independence; hence $W = \mathcal{V}$ iff $\mathcal{O}$ is an orthonormal Hamel basis. It remains to prove maximality iff spanning.

($\Rightarrow$) Suppose $\mathcal{O}$ is a Hilbert basis. If $W \neq \mathcal{V}$, choose $\mathbf{v} \notin W$ and set $\mathbf{u} = \mathbf{v} - \sum_j \langle \mathbf{v}, \mathbf{e}_j \rangle \mathbf{e}_j$. Then $\langle \mathbf{u}, \mathbf{e}_i \rangle = 0 \; \forall i$, and $\mathbf{u} \neq \mathbf{0}$ since $\mathbf{v} \notin W$. Thus $\mathcal{O} \cup \{\mathbf{u}/\|\mathbf{u}\|\}$ is orthonormal, contradicting maximality. Hence $W = \mathcal{V}$.

($\Leftarrow$) Suppose $W = \mathcal{V}$. If $\mathbf{w} \perp \mathcal{O}$, write $\mathbf{w} = \sum_j a_j \mathbf{e}_j$. Then $a_j = \langle \mathbf{w}, \mathbf{e}_j \rangle = 0$ for all j, so $\mathbf{w} = \mathbf{0}$. Hence no nonzero vector can be adjoined, and $\mathcal{O}$ is maximal orthonormal.

Structural Commentary: Finite maximal orthonormal sets span.
Key Insight: Missing directions can be orthogonalized and adjoined. □

Example 2.1.47. *A Hilbert Basis Need Not Be a Hamel Basis*
In $\ell^2(\mathbb{F})$, let $\mathbf{e}_j = (0, \ldots, 0, 1, 0, \ldots)$. Then $\{\mathbf{e}_j : j \geq 1\}$ is orthonormal and maximal, since $\mathbf{x} \perp \mathbf{e}_j$ for all j implies $x_j = \langle \mathbf{x}, \mathbf{e}_j \rangle = 0$ for all j. Its finite span is $\mathbb{F}^{(\mathbb{N})} \subsetneq \ell^2(\mathbb{F})$, so it is a Hilbert basis but not a Hamel basis.

Theorem 2.1.48. *Coordinates in an Orthonormal Basis*
Let $\mathcal{V}$ be finite-dimensional and let $\beta = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be an orthonormal basis of $\mathcal{V}$.

Then for every $\mathbf{v} \in \mathcal{V}$, $\mathbf{v} = \sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \mathbf{e}_j$, $\|\mathbf{v}\|^2 = \sum_{j=1}^{n} \left|\langle \mathbf{v}, \mathbf{e}_j \rangle\right|^2$ and $[\mathbf{v}]_\beta = \begin{bmatrix} \langle \mathbf{v}, \mathbf{e}_1 \rangle \\ \vdots \\ \langle \mathbf{v}, \mathbf{e}_n \rangle \end{bmatrix}$.

Proof. Since β is a basis, write $\mathbf{v} = \sum_{j=1}^{n} a_j \mathbf{e}_j$. Taking inner product with $\mathbf{e}_k$ gives $\langle \mathbf{v}, \mathbf{e}_k \rangle = \sum_{j=1}^{n} a_j \langle \mathbf{e}_j, \mathbf{e}_k \rangle = a_k$. Thus the displayed expansion and coordinate column follow.

Applying Proposition 2.1.33 to the same expansion gives $\|\mathbf{v}\|^2 = \sum_{j=1}^{n} |a_j|^2 = \sum_{j=1}^{n} \left|\langle \mathbf{v}, \mathbf{e}_j \rangle\right|^2$.

Structural Commentary: An orthonormal basis makes coordinate extraction an inner-product computation.
Key Insight: Pairing with $\mathbf{e}_k$ isolates the kth coordinate. □

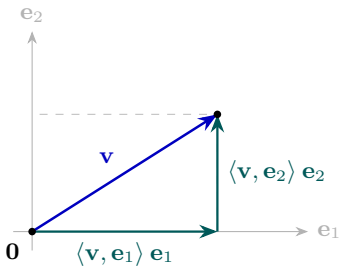


Figure 25: An orthonormal basis resolves a vector into perpendicular coordinate components.

Remark 2.1.49. *Cross Terms*
In an expansion such as $\|\mathbf{u} + \mathbf{v}\|^2$, the terms involving both vectors, $\langle \mathbf{u}, \mathbf{v} \rangle$ and $\langle \mathbf{v}, \mathbf{u} \rangle$, are called *cross terms*. The terms $\|\mathbf{u}\|^2$ and $\|\mathbf{v}\|^2$ are the pure squared-length terms.

Theorem 2.1.50. *Pythagorean Theorem*
If $\mathbf{u}, \mathbf{v} \in \mathcal{V}$ and $\mathbf{u} \perp \mathbf{v}$, then $\|\mathbf{u} + \mathbf{v}\|^2 = \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2$.

Proof. Since $\mathbf{u} \perp \mathbf{v}$, Proposition 2.1.29 gives $\langle \mathbf{v}, \mathbf{u} \rangle = 0$. Thus

$$\|\mathbf{u} + \mathbf{v}\|^2 = \langle \mathbf{u} + \mathbf{v}, \mathbf{u} + \mathbf{v} \rangle = \langle \mathbf{u}, \mathbf{u} \rangle + \langle \mathbf{u}, \mathbf{v} \rangle + \langle \mathbf{v}, \mathbf{u} \rangle + \langle \mathbf{v}, \mathbf{v} \rangle = \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2 .$$

Structural Commentary: Orthogonality cancels the cross terms in the squared-norm expansion.
Key Insight: Perpendicular components add in squared length, not in length. □

Proposition 2.1.51. *Pythagorean Converse over $\mathbb{R}$*
In a real inner-product space, $\mathbf{u} \perp \mathbf{v} \iff \|\mathbf{u} + \mathbf{v}\|^2 = \|\mathbf{u}\|^2 + \|\mathbf{v}\|^2$.

Proof. Use $\|\mathbf{u} + \mathbf{v}\|^2 = \|\mathbf{u}\|^2 + 2\langle \mathbf{u}, \mathbf{v} \rangle + \|\mathbf{v}\|^2$.

($\Rightarrow$) If $\mathbf{u} \perp \mathbf{v}$, then $\langle \mathbf{u}, \mathbf{v} \rangle = 0$, so the identity follows.
($\Leftarrow$) If the identity holds, then $2\langle \mathbf{u}, \mathbf{v} \rangle = 0$, hence $\mathbf{u} \perp \mathbf{v}$.

Structural Commentary: In real inner-product spaces, the cross term has no hidden imaginary part.
Key Insight: The Pythagorean identity detects orthogonality exactly over $\mathbb{R}$. □

Theorem 2.1.52. *Bessel's Inequality*
Let $(\mathbf{e}_1, \ldots, \mathbf{e}_m)$ be an orthonormal list in an inner-product space $\mathcal{V}$.
Then, for every $\mathbf{v} \in \mathcal{V}$, $\sum_{j=1}^{m} \left|\langle \mathbf{v}, \mathbf{e}_j \rangle\right|^2 \leq \|\mathbf{v}\|^2$.

Proof. Set $\mathbf{u} = \sum_{j=1}^{m} \langle \mathbf{v}, \mathbf{e}_j \rangle \mathbf{e}_j$. For each k, orthonormality gives $\langle \mathbf{v} - \mathbf{u}, \mathbf{e}_k \rangle = \langle \mathbf{v}, \mathbf{e}_k \rangle - \sum_{j=1}^{m} \langle \mathbf{v}, \mathbf{e}_j \rangle \langle \mathbf{e}_j, \mathbf{e}_k \rangle = 0$. Thus $\mathbf{v} - \mathbf{u}$ is orthogonal to $\mathbf{u}$. By Pythagorean Theorem 2.1.50 and Proposition 2.1.33,

$$\|\mathbf{v}\|^2 = \|\mathbf{u}\|^2 + \|\mathbf{v} - \mathbf{u}\|^2 \geq \|\mathbf{u}\|^2 = \sum_{j=1}^{m} \left|\langle \mathbf{v}, \mathbf{e}_j \rangle\right|^2.$$

Structural Commentary: The chosen directions give an orthogonal part of $\mathbf{v}$; the residual is perpendicular.
Key Insight: Partial orthonormal coordinates cannot contain more energy than the whole vector. □

Corollary 2.1.53. *Coefficients Against an Orthonormal Sequence Vanish*
Let $(\mathbf{e}_1, \mathbf{e}_2, \ldots)$ be an orthonormal sequence in an inner-product space $\mathcal{V}$, meaning every finite sublist is orthonormal. Then, for every $\mathbf{v} \in \mathcal{V}$, $\lim_{n \to \infty} \langle \mathbf{v}, \mathbf{e}_n \rangle = 0$.

Proof. Fix $\epsilon > 0$. By Theorem 2.1.52, every finite set of indices satisfying $\left|\langle \mathbf{v}, \mathbf{e}_j \rangle\right| \geq \epsilon$ has cardinality at most $\|\mathbf{v}\|^2 / \epsilon^2$. Hence only finitely many j have $\left|\langle \mathbf{v}, \mathbf{e}_j \rangle\right| \geq \epsilon$. This holds for every $\epsilon > 0$, so $\langle \mathbf{v}, \mathbf{e}_n \rangle \to 0$.
Structural Commentary: Bessel bounds the number of large orthonormal coefficients.
Key Insight: Infinitely many large coefficients would force infinite energy. □

Theorem 2.1.54. *Parseval's Identity*
Let $\beta = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be an orthonormal basis of an inner-product space $\mathcal{V}$.
Then, for all $\mathbf{v}, \mathbf{w} \in \mathcal{V}$, $\langle \mathbf{v}, \mathbf{w} \rangle = \sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \overline{\langle \mathbf{w}, \mathbf{e}_j \rangle} = \langle [\mathbf{v}]_\beta, [\mathbf{w}]_\beta \rangle_{\mathbb{F}^n}$.

Proof. By Theorem 2.1.48, $\mathbf{v} = \sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \mathbf{e}_j$ and $\mathbf{w} = \sum_{k=1}^{n} \langle \mathbf{w}, \mathbf{e}_k \rangle \mathbf{e}_k$. By linearity in first slot and conjugate linearity in second, $\langle \mathbf{v}, \mathbf{w} \rangle = \sum_{j=1}^{n} \sum_{k=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \overline{\langle \mathbf{w}, \mathbf{e}_k \rangle} \langle \mathbf{e}_j, \mathbf{e}_k \rangle = \sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \overline{\langle \mathbf{w}, \mathbf{e}_j \rangle}$.
The last expression is exactly the standard inner product of the coordinate columns in $\mathbb{F}^n$.
Structural Commentary: Orthonormal coordinates make inner products coordinatewise.
Key Insight: Orthonormal coordinates preserve angles and lengths, not just linear combinations. □

Theorem 2.1.55. *Matrix Entries from an Orthonormal Codomain Basis*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional vector spaces over $\mathbb{F}$, with $\mathcal{W}$ an inner-product space. Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be a basis of $\mathcal{V}$ and let $\gamma = (\mathbf{w}_1, \ldots, \mathbf{w}_m)$ be an orthonormal basis of $\mathcal{W}$. If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then

$$[T]_{\gamma \leftarrow \beta} = \begin{bmatrix} \langle T\mathbf{v}_1, \mathbf{w}_1 \rangle & \cdots & \langle T\mathbf{v}_n, \mathbf{w}_1 \rangle \\ \vdots & \ddots & \vdots \\ \langle T\mathbf{v}_1, \mathbf{w}_m \rangle & \cdots & \langle T\mathbf{v}_n, \mathbf{w}_m \rangle \end{bmatrix}.$$

Conversely, for each $A \in \mathrm{Mat}(m, n, \mathbb{F})$, there is a unique $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ such that $[T]_{\gamma \leftarrow \beta} = A$.

Proof. The jth column of $[T]_{\gamma \leftarrow \beta}$ is $[T\mathbf{v}_j]_\gamma$. Since γ is orthonormal, Theorem 2.1.48 gives

$$[T\mathbf{v}_j]_\gamma = \begin{bmatrix} \langle T\mathbf{v}_j, \mathbf{w}_1 \rangle \\ \vdots \\ \langle T\mathbf{v}_j, \mathbf{w}_m \rangle \end{bmatrix}.$$

This proves the displayed matrix formula. The converse is exactly the surjectivity and injectivity of the coordinate encoding in Theorem 1.3.89.
Structural Commentary: Columns are basis-image coordinates, read by inner products.
Key Insight: Each entry measures $T\mathbf{v}_j$ against one output direction. □

Corollary 2.1.56. *Inner Products Through an Operator Matrix*
Let $\beta = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be an orthonormal basis of a finite-dimensional inner-product space $\mathcal{V}$, and let $T \in \mathcal{L}(\mathcal{V})$.
Then, for all $\mathbf{v}, \mathbf{w} \in \mathcal{V}$, $\langle T\mathbf{v}, \mathbf{w} \rangle = \langle [T]_{\beta \leftarrow \beta} [\mathbf{v}]_\beta, [\mathbf{w}]_\beta \rangle_{\mathbb{F}^n}$.

Proof. By Proposition 1.3.85, $[T\mathbf{v}]_\beta = [T]_{\beta \leftarrow \beta} [\mathbf{v}]_\beta$. Applying Theorem 2.1.54 to $T\mathbf{v}$ and $\mathbf{w}$ gives the result.
Structural Commentary: Orthonormal coordinates turn $\langle T\mathbf{v}, \mathbf{w} \rangle$ into matrix computation.
Key Insight: Orthonormal coordinates make operator geometry computable by matrices. □

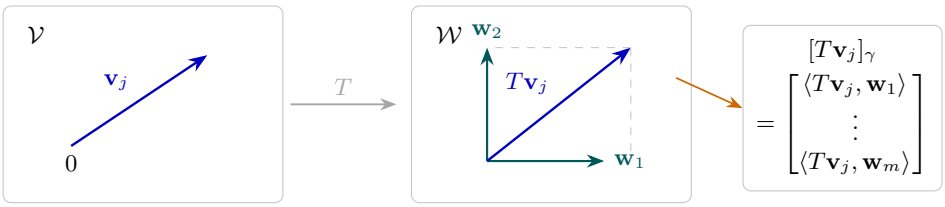


Figure 26: The jth column of $[T]_{\gamma \leftarrow \beta}$ is obtained by resolving $T\mathbf{v}_j$ against the orthonormal basis of $\mathcal{W}$.

Theorem 2.1.57. *Riesz Representation Theorem*
Let $\mathcal{V}$ be a finite-dimensional inner-product space and let $\varphi : \mathcal{V} \to \mathbb{F}$ be a linear functional. Then there is a unique vector $\mathbf{u} \in \mathcal{V}$ such that $\varphi(\mathbf{v}) = \langle \mathbf{v}, \mathbf{u} \rangle$ for every $\mathbf{v} \in \mathcal{V}$.

Proof. Let $(\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be an orthonormal basis of $\mathcal{V}$.
By Theorem 2.1.48, $\varphi(\mathbf{v}) = \varphi\left(\sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \, \mathbf{e}_j\right) = \sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \, \varphi(\mathbf{e}_j)$.
Set $\mathbf{u} = \sum_{j=1}^{n} \overline{\varphi(\mathbf{e}_j)} \, \mathbf{e}_j$. Since inner product is conjugate linear in second slot, $\langle \mathbf{v}, \mathbf{u} \rangle = \sum_{j=1}^{n} \langle \mathbf{v}, \mathbf{e}_j \rangle \, \varphi(\mathbf{e}_j) = \varphi(\mathbf{v})$.
For uniqueness, suppose $\langle \mathbf{v}, \mathbf{u}_1 \rangle = \langle \mathbf{v}, \mathbf{u}_2 \rangle \ \ \forall \mathbf{v} \in \mathcal{V}$. Taking $\mathbf{v} = \mathbf{u}_1 - \mathbf{u}_2$ gives $\|\mathbf{u}_1 - \mathbf{u}_2\|^2 = 0$, hence $\mathbf{u}_1 = \mathbf{u}_2$.
Structural Commentary: An orthonormal basis packages a functional as one representing vector.
Key Insight: Finite-dimensional functionals are inner products with fixed vectors. □

Definition 2.1.58. *Riesz Vector*
Let $\mathcal{V}$ be finite-dimensional and let $\varphi : \mathcal{V} \to \mathbb{F}$ be a linear functional. The unique vector $\mathbf{u} \in \mathcal{V}$ satisfying $\varphi(\mathbf{v}) = \langle \mathbf{v}, \mathbf{u} \rangle$ for every $\mathbf{v} \in \mathcal{V}$ is the *Riesz vector* representing φ.

Theorem 2.1.59. *Riesz Map*
Let $\mathcal{V}$ be finite-dimensional inner-product space, $\tau : \mathcal{V} \to \mathcal{V}'$, $\tau(\mathbf{u}) = \langle \cdot, \mathbf{u} \rangle$. Then τ is bijective, conjugate-linear. Hence $R = \tau^{-1}$ is conjugate-linear, $R(a\varphi + \psi) = \overline{a} \, R(\varphi) + R(\psi)$, $a \in \mathbb{F}$, $\varphi, \psi \in \mathcal{V}'$.
Over $\mathbb{R}$ this is ordinary linearity.

Proof. For $a \in \mathbb{F}$ and $\mathbf{u}, \mathbf{w} \in \mathcal{V}$, $\tau(a\mathbf{u} + \mathbf{w}) = \overline{a}\tau(\mathbf{u}) + \tau(\mathbf{w})$, so τ is conjugate-linear.
If $\tau(\mathbf{u}) = 0$, then $\|\mathbf{u}\|^2 = \tau(\mathbf{u})(\mathbf{u}) = 0$, so τ is injective. Surjectivity is Theorem 2.1.57.
For $a \in \mathbb{F}$, $\varphi, \psi \in \mathcal{V}'$, and $\mathbf{v} \in \mathcal{V}$, $\langle \mathbf{v}, \overline{a}R(\varphi) + R(\psi) \rangle = a \langle \mathbf{v}, R(\varphi) \rangle + \langle \mathbf{v}, R(\psi) \rangle = (a\varphi + \psi)(\mathbf{v})$.
Uniqueness gives $R(a\varphi + \psi) = \overline{a}R(\varphi) + R(\psi)$.
Structural Commentary: Riesz identifies $\mathcal{V}$ with $\mathcal{V}'$ conjugate-linearly.
Key Insight: The representing slot conjugates scalars. □

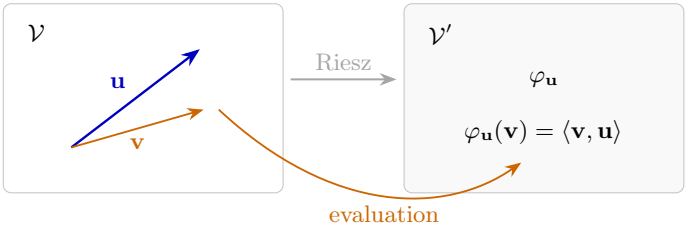


Figure 27: Riesz identifies a functional with its representing vector.

Perspective 2.1.60. *Riesz as Coordinate Duality*
The vector $\sum_{j=1}^{n} \overline{\varphi(\mathbf{e}_j)} \mathbf{e}_j$ seems to depend on the orthonormal basis, but Theorem 2.1.57 shows that the resulting representing vector is intrinsic.

Definition 2.1.61. *Gram Matrix*
For vectors $\mathbf{u}_1, \ldots, \mathbf{u}_n$ in an inner-product space, the *Gram matrix* is

$$\operatorname{Gram}(\mathbf{u}_1, \ldots, \mathbf{u}_n) := \begin{bmatrix} \langle \mathbf{u}_1, \mathbf{u}_1 \rangle & \langle \mathbf{u}_2, \mathbf{u}_1 \rangle & \cdots & \langle \mathbf{u}_n, \mathbf{u}_1 \rangle \\ \langle \mathbf{u}_1, \mathbf{u}_2 \rangle & \langle \mathbf{u}_2, \mathbf{u}_2 \rangle & \cdots & \langle \mathbf{u}_n, \mathbf{u}_2 \rangle \\ \vdots & \vdots & \ddots & \vdots \\ \langle \mathbf{u}_1, \mathbf{u}_n \rangle & \langle \mathbf{u}_2, \mathbf{u}_n \rangle & \cdots & \langle \mathbf{u}_n, \mathbf{u}_n \rangle \end{bmatrix}.$$

Proposition 2.1.62. *Gram Matrix Quadratic Form*

For any vectors $\mathbf{u}_1, \dots, \mathbf{u}_n$*, Gram matrix* $G = \operatorname{Gram}(\mathbf{u}_1, \dots, \mathbf{u}_n)$ *is Hermitian and satisfies* $\mathbf{c}^* G \mathbf{c} = \left\| \sum_{j=1}^{n} c_j \mathbf{u}_j \right\|^2$ *for every* $\mathbf{c} = [c_1 \ \cdots \ c_n]^t \in \mathbb{F}^n$*. Hence the Gram quadratic form is nonnegative. If* $\mathbf{u}_1, \dots, \mathbf{u}_n$ *are linearly independent, then it is positive for every nonzero* $\mathbf{c}$*.*

Proof. Hermitian symmetry follows from conjugate symmetry of the inner product. Expanding the squared norm gives $\left\| \sum_{j=1}^{n} c_j \mathbf{u}_j \right\|^2 = \sum_{i,j} \overline{c_i} \langle \mathbf{u}_j, \mathbf{u}_i \rangle c_j = \mathbf{c}^* G \mathbf{c}$. The right-hand side is a squared norm, so it is nonnegative. If $\mathbf{u}_1, \dots, \mathbf{u}_n$ are linearly independent, the squared norm is 0 only when all $c_j = 0$.
Structural Commentary: A Gram matrix records every pairwise inner product in a list of vectors.
Key Insight: The quadratic form of a Gram matrix is exactly the squared norm of a linear combination. □

2.2 Orthogonal Complements and Projections

Orthogonal complements turn orthogonality into a subspace operation; projections turn it into canonical decompositions and best approximations.

Definition 2.2.1. *Orthogonality, Complements, and Direct Sums*
Let $\mathcal{X}, \mathcal{Y}, \mathcal{U} \subseteq \mathcal{V}$ and let $\mathcal{S}, \mathcal{T} \leq \mathcal{V}$.

i. $\mathcal{X}$ and $\mathcal{Y}$ are *orthogonal*, written $\mathcal{X} \perp \mathcal{Y}$, if $\langle \mathbf{x}, \mathbf{y} \rangle = 0$ for all $\mathbf{x} \in \mathcal{X}$ and $\mathbf{y} \in \mathcal{Y}$.
ii. The *orthogonal complement* of $\mathcal{U}$ is $\mathcal{U}^\perp := \{\mathbf{v} \in \mathcal{V} : \langle \mathbf{v}, \mathbf{u} \rangle = 0 \text{ for all } \mathbf{u} \in \mathcal{U}\}$.
iii. $\mathcal{V}$ is the *orthogonal direct sum* of $\mathcal{S}$ and $\mathcal{T}$, written $\mathcal{V} = \mathcal{S} \odot \mathcal{T}$, if $\mathcal{V} = \mathcal{S} \oplus \mathcal{T}$ and $\mathcal{S} \perp \mathcal{T}$.

Proposition 2.2.2. *Basic Properties of Orthogonal Complements*
For subsets $\mathcal{U}, \mathcal{W} \subseteq \mathcal{V}$*,*

(i) $\mathcal{U}^\perp$ *is a subspace of* $\mathcal{V}$*;*
(ii) $\{\mathbf{0}\}^\perp = \mathcal{V}$*;*
(iii) $\mathcal{V}^\perp = \{\mathbf{0}\}$*;*
(iv) $\mathcal{U} \cap \mathcal{U}^\perp \subseteq \{\mathbf{0}\}$*;*
(v) if $\mathcal{U} \subseteq \mathcal{W}$*, then* $\mathcal{W}^\perp \subseteq \mathcal{U}^\perp$*.*

Proof.

(i) We have $\mathbf{0} \in \mathcal{U}^\perp$. If $\mathbf{v}, \mathbf{w} \in \mathcal{U}^\perp$ and $a \in \mathbb{F}$, then for every $\mathbf{u} \in \mathcal{U}$, $\langle a\mathbf{v} + \mathbf{w}, \mathbf{u} \rangle = a \langle \mathbf{v}, \mathbf{u} \rangle + \langle \mathbf{w}, \mathbf{u} \rangle = 0$.
(ii) For every $\mathbf{v} \in \mathcal{V}$, $\langle \mathbf{v}, \mathbf{0} \rangle = 0$, so $\mathbf{v} \in \{\mathbf{0}\}^\perp$.
(iii) If $\mathbf{v} \in \mathcal{V}^\perp$, then $\langle \mathbf{v}, \mathbf{v} \rangle = 0$, so $\mathbf{v} = \mathbf{0}$.
(iv) If $\mathbf{v} \in \mathcal{U} \cap \mathcal{U}^\perp$, then $\langle \mathbf{v}, \mathbf{v} \rangle = 0$, so $\mathbf{v} = \mathbf{0}$.
(v) If $\mathbf{v} \in \mathcal{W}^\perp$, then $\mathbf{v}$ is orthogonal to every vector in $\mathcal{W}$, hence to every vector in $\mathcal{U}$.

Structural Commentary: Orthogonal complements are solution spaces of homogeneous inner-product equations. Larger input sets impose more equations, so their complements become smaller.
Key Insight: Orthogonality constraints are linear constraints. □

Remark 2.2.3. *Definiteness Matters*
For a bilinear form B, $\mathcal{U}^\perp = \{\mathbf{v} : B(\mathbf{v}, \mathbf{u}) = 0 \ \forall \mathbf{u} \in \mathcal{U}\}$ is still a subspace.
Without positive definiteness, however, one may have $\mathcal{V}^\perp \neq \{\mathbf{0}\}$ or $\mathcal{U} \cap \mathcal{U}^\perp \neq \{\mathbf{0}\}$.
Thus $\mathcal{V} = \mathcal{U} \oplus \mathcal{U}^\perp$ is an inner-product result, not a consequence of bilinearity alone.

Theorem 2.2.4. *Direct Sum with the Orthogonal Complement*
If $\mathcal{U}$ *is a finite-dimensional subspace of* $\mathcal{V}$*, then* $\mathcal{V} = \mathcal{U} \oplus \mathcal{U}^\perp$*.*

Proof. Let $(\mathbf{e}_1, \dots, \mathbf{e}_m)$ be an orthonormal basis of $\mathcal{U}$. For $\mathbf{v} \in \mathcal{V}$, define $\mathbf{u} = \sum_{j=1}^{m} \langle \mathbf{v}, \mathbf{e}_j \rangle \mathbf{e}_j$ and set $\mathbf{w} = \mathbf{v} - \mathbf{u}$.
Clearly $\mathbf{u} \in \mathcal{U}$. For each k, $\langle \mathbf{w}, \mathbf{e}_k \rangle = \langle \mathbf{v}, \mathbf{e}_k \rangle - \sum_{j=1}^{m} \langle \mathbf{v}, \mathbf{e}_j \rangle \langle \mathbf{e}_j, \mathbf{e}_k \rangle = 0$.
If $\mathbf{x} \in \mathcal{U}$, write $\mathbf{x} = \sum_{k=1}^{m} a_k \mathbf{e}_k$. Then $\langle \mathbf{w}, \mathbf{x} \rangle = \sum_{k=1}^{m} \overline{a_k} \langle \mathbf{w}, \mathbf{e}_k \rangle = 0$, so $\mathbf{w} \in \mathcal{U}^\perp$ and $\mathbf{v} = \mathbf{u} + \mathbf{w}$.
Hence $\mathcal{V} = \mathcal{U} + \mathcal{U}^\perp$. The intersection is $\{\mathbf{0}\}$ by Proposition 2.2.2, so the sum is direct.
Structural Commentary: Each vector splits into a $\mathcal{U}$-component and an orthogonal residual.
Key Insight: Projection plus residual gives the canonical direct-sum decomposition. □

Proposition 2.2.5. *Orthogonal Direct Sum Characterization*
For subspaces $\mathcal{S}, \mathcal{T} \subseteq \mathcal{V}$, the following are equivalent:

i. $\mathcal{V} = \mathcal{S} \odot \mathcal{T}$.
ii. $\mathcal{V} = \mathcal{S} \oplus \mathcal{T}$ *and* $\mathcal{T} = \mathcal{S}^\perp$.

Proof.

(*i*) $\Rightarrow$ (*ii*) If $\mathcal{V} = \mathcal{S} \odot \mathcal{T}$, then $\mathcal{V} = \mathcal{S} \oplus \mathcal{T}$ and $\mathcal{T} \subseteq \mathcal{S}^\perp$. For $\mathbf{v} \in \mathcal{S}^\perp$, write $\mathbf{v} = \mathbf{s} + \mathbf{t}$ with $\mathbf{s} \in \mathcal{S}$, $\mathbf{t} \in \mathcal{T}$. Then $\mathbf{s} = \mathbf{v} - \mathbf{t} \in \mathcal{S}^\perp$, so $\mathbf{s} \in \mathcal{S} \cap \mathcal{S}^\perp = \{\mathbf{0}\}$ and $\mathbf{v} \in \mathcal{T}$. Hence $\mathcal{T} = \mathcal{S}^\perp$.
(*ii*) $\Rightarrow$ (*i*) If $\mathcal{V} = \mathcal{S} \oplus \mathcal{T}$ and $\mathcal{T} = \mathcal{S}^\perp$, then $\mathcal{S} \perp \mathcal{T}$, so $\mathcal{V} = \mathcal{S} \odot \mathcal{T}$.

Structural Commentary: Orthogonal complements are maximal orthogonal partners.
Key Insight: A full orthogonal summand equals the complement. □

Corollary 2.2.6. *Dimension of the Orthogonal Complement*
If $\mathcal{V}$ is finite-dimensional and $\mathcal{U}$ is a subspace of $\mathcal{V}$, then $\dim(\mathcal{U}^\perp) = \dim(\mathcal{V}) - \dim(\mathcal{U})$.

Proof. Since $\mathcal{V}$ is finite-dimensional, $\mathcal{U}$ is finite-dimensional by Corollary 1.1.28. Theorem 2.2.4 gives $\mathcal{V} = \mathcal{U} \oplus \mathcal{U}^\perp$. Thus Corollary 1.2.20 gives $\dim(\mathcal{V}) = \dim(\mathcal{U}) + \dim(\mathcal{U}^\perp)$.
Structural Commentary: Dimensions add across the orthogonal split.
Key Insight: Orthogonal complements contain exactly the dimensions not used by $\mathcal{U}$. □

Corollary 2.2.7. *Double Orthogonal Complement*
If $\mathcal{U}$ is a finite-dimensional subspace of $\mathcal{V}$, then $\mathcal{U} = (\mathcal{U}^\perp)^\perp$.

Proof. If $\mathbf{u} \in \mathcal{U}$, then $\mathbf{u}$ is orthogonal to every vector in $\mathcal{U}^\perp$, so $\mathcal{U} \subseteq (\mathcal{U}^\perp)^\perp$. Conversely, let $\mathbf{v} \in (\mathcal{U}^\perp)^\perp$. By Theorem 2.2.4, write $\mathbf{v} = \mathbf{u} + \mathbf{w}$ with $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{U}^\perp$. Since $\mathcal{U} \subseteq (\mathcal{U}^\perp)^\perp$, we have $\mathbf{w} = \mathbf{v} - \mathbf{u} \in (\mathcal{U}^\perp)^\perp$. Hence $\mathbf{w} \in \mathcal{U}^\perp \cap (\mathcal{U}^\perp)^\perp$, so $\mathbf{w} = \mathbf{0}$ and $\mathbf{v} = \mathbf{u} \in \mathcal{U}$.
Structural Commentary: The orthogonal split leaves no extra direction beyond $\mathcal{U}$.
Key Insight: Taking the orthogonal complement twice recovers the original finite-dimensional subspace. □

Corollary 2.2.8. *Double Complement of a Finite Span*
If $\mathcal{X} \subseteq \mathcal{V}$ and $\operatorname{span}(\mathcal{X})$ *is finite-dimensional, then* $(\mathcal{X}^\perp)^\perp = \operatorname{span}(\mathcal{X})$.

Proof. Let $W = \operatorname{span}(\mathcal{X})$. As $\mathcal{X}^\perp = W^\perp$, W is finite-dimensional, by Corollary 2.2.7, $(\mathcal{X}^\perp)^\perp = (W^\perp)^\perp = W$.
Structural Commentary: Orthogonal complements depend only on the span of the tested set.
Key Insight: Double orthogonality recovers the finite-dimensional span, not merely the original subset. □

Perspective 2.2.9. *Orthogonal Versus Complementary*
Orthogonality gives $\mathcal{W} \subseteq \mathcal{U}^\perp$ and $\mathcal{U} \subseteq \mathcal{W}^\perp$, not necessarily equality. In finite dimensions, equality follows once $\dim(\mathcal{U}) + \dim(\mathcal{W}) = \dim(\mathcal{V})$ by Corollary 2.2.6.

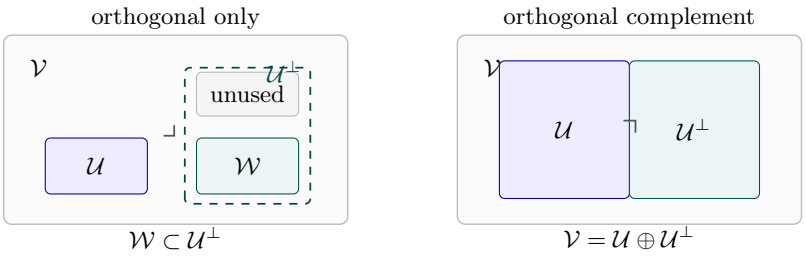


Figure 28: Orthogonal subspaces are complements only when their sum fills $\mathcal{V}$.

Proposition 2.2.10. *Orthogonal Projection onto a Line*
Let $\mathbf{u}, \mathbf{v} \in \mathcal{V}$ with $\mathbf{v} \neq \mathbf{0}$. Define $c := \frac{\langle \mathbf{u}, \mathbf{v} \rangle}{\|\mathbf{v}\|^2}$, $\mathbf{w} := \mathbf{u} - c\mathbf{v}$. Then $\mathbf{w} \perp \mathbf{v}$ and $\mathbf{u} = c\mathbf{v} + \mathbf{w}$.
The vector $\operatorname{proj}_{\mathbf{v}}(\mathbf{u}) := c\mathbf{v} = \left\langle \mathbf{u}, \frac{\mathbf{v}}{\|\mathbf{v}\|} \right\rangle \frac{\mathbf{v}}{\|\mathbf{v}\|}$ *is the orthogonal projection of $\mathbf{u}$ onto* $\operatorname{span}(\mathbf{v})$.

Proof. The decomposition $\mathbf{u} = c\mathbf{v} + \mathbf{w}$ follows from the definition of $\mathbf{w}$.
Moreover, $\langle \mathbf{w}, \mathbf{v} \rangle = \langle \mathbf{u} - c\mathbf{v}, \mathbf{v} \rangle = \langle \mathbf{u}, \mathbf{v} \rangle - c \langle \mathbf{v}, \mathbf{v} \rangle = \langle \mathbf{u}, \mathbf{v} \rangle - \frac{\langle \mathbf{u}, \mathbf{v} \rangle}{\|\mathbf{v}\|^2} \|\mathbf{v}\|^2 = 0$. Thus $\mathbf{w} \perp \mathbf{v}$.
Structural Commentary: Choose c so the remainder $\mathbf{u} - c\mathbf{v}$ has zero inner product with $\mathbf{v}$.
Key Insight: Projection is the part detected by pairing with the direction vector. □

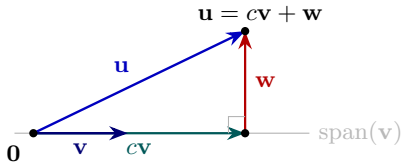


Figure 29: Writing $\mathbf{u}$ as a component parallel to $\mathbf{v}$ plus a component orthogonal to $\mathbf{v}$.

Definition 2.2.11. *Orthogonal Projection onto a Subspace*
Let $\mathcal{U}$ be a finite-dimensional subspace of $\mathcal{V}$.
For $\mathbf{v} \in \mathcal{V}$, write the unique decomposition $\mathbf{v} = \mathbf{u} + \mathbf{w}$, $\mathbf{u} \in \mathcal{U}$, $\mathbf{w} \in \mathcal{U}^\perp$, from Theorem 2.2.4.
The *orthogonal projection* of $\mathcal{V}$ onto $\mathcal{U}$ is the map $P_\mathcal{U} : \mathcal{V} \to \mathcal{V}$ defined by $P_\mathcal{U}\mathbf{v} = \mathbf{u}$.

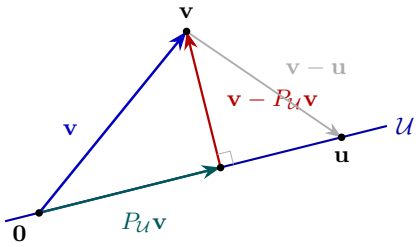


Figure 30: Orthogonal projection decomposes $\mathbf{v}$ into $P_\mathcal{U}\mathbf{v}$ plus an orthogonal residual.

Proposition 2.2.12. *Properties of Orthogonal Projection*
Let $\mathcal{U}$ be a finite-dimensional subspace of $\mathcal{V}$. Then:

i. $P_\mathcal{U} \in \mathcal{L}(\mathcal{V})$;
ii. $P_\mathcal{U}\mathbf{u} = \mathbf{u}$ *for every* $\mathbf{u} \in \mathcal{U}$;
iii. $P_\mathcal{U}\mathbf{w} = \mathbf{0}$ *for every* $\mathbf{w} \in \mathcal{U}^\perp$;
iv. $\operatorname{range}(P_\mathcal{U}) = \mathcal{U}$;
v. $\ker(P_\mathcal{U}) = \mathcal{U}^\perp$, *i.e. the kernel/nullspace of* $P_\mathcal{U}$ *is* $\mathcal{U}^\perp$;
vi. $\mathbf{v} - P_\mathcal{U}\mathbf{v} \in \mathcal{U}^\perp$ *for every* $\mathbf{v} \in \mathcal{V}$;
vii. $P_\mathcal{U}^2 = P_\mathcal{U}$;
viii. $\|P_\mathcal{U}\mathbf{v}\| \le \|\mathbf{v}\|$ *for every* $\mathbf{v} \in \mathcal{V}$;
ix. if $(\mathbf{e}_1, \dots, \mathbf{e}_m)$ *is an orthonormal basis of* $\mathcal{U}$, *then* $P_\mathcal{U}\mathbf{v} = \sum_{j=1}^{m} \langle \mathbf{v}, \mathbf{e}_j \rangle \, \mathbf{e}_j$, $\mathbf{v} \in \mathcal{V}$.

Proof.

(i) For linearity, let $a \in \mathbb{F}$ and write $\mathbf{v}_r = \mathbf{u}_r + \mathbf{w}_r$ with $\mathbf{u}_r \in \mathcal{U}$ and $\mathbf{w}_r \in \mathcal{U}^\perp$ for $r = 1, 2$.
Then $a\mathbf{v}_1 + \mathbf{v}_2 = (a\mathbf{u}_1 + \mathbf{u}_2) + (a\mathbf{w}_1 + \mathbf{w}_2)$, where first summand lies in $\mathcal{U}$ and second in $\mathcal{U}^\perp$.
By uniqueness of the orthogonal decomposition, $P_\mathcal{U}(a\mathbf{v}_1 + \mathbf{v}_2) = aP_\mathcal{U}\mathbf{v}_1 + P_\mathcal{U}\mathbf{v}_2$.
(ii) This follows from $\mathbf{u} = \mathbf{u} + \mathbf{0}$ with $\mathbf{u} \in \mathcal{U}$ and $\mathbf{0} \in \mathcal{U}^\perp$.
(iii) This follows from $\mathbf{w} = \mathbf{0} + \mathbf{w}$ with $\mathbf{0} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{U}^\perp$.
(iv) The definition gives $\operatorname{range}(P_\mathcal{U}) \subseteq \mathcal{U}$, and (ii) gives $\mathcal{U} \subseteq \operatorname{range}(P_\mathcal{U})$.
(v) Part (iii) gives $\mathcal{U}^\perp \subseteq \ker(P_\mathcal{U})$.
Conversely, if $P_\mathcal{U}\mathbf{v} = \mathbf{0}$, then $\mathbf{v} = \mathbf{0} + (\mathbf{v} - P_\mathcal{U}\mathbf{v})$ with $\mathbf{v} - P_\mathcal{U}\mathbf{v} \in \mathcal{U}^\perp$, so $\mathbf{v} \in \mathcal{U}^\perp$.
(vi) This is exactly the residual term in the decomposition $\mathbf{v} = P_\mathcal{U}\mathbf{v} + (\mathbf{v} - P_\mathcal{U}\mathbf{v})$.
(vii) By (ii), $P_\mathcal{U}(P_\mathcal{U}\mathbf{v}) = P_\mathcal{U}\mathbf{v}$ for every $\mathbf{v} \in \mathcal{V}$.
(viii) Write $\mathbf{v} = \mathbf{u} + \mathbf{w}$ with $\mathbf{u} \in \mathcal{U}$, $\mathbf{w} \in \mathcal{U}^\perp$. Then $\mathbf{u} \perp \mathbf{w}$ and $\|\mathbf{v}\|^2 = \|\mathbf{u}\|^2 + \|\mathbf{w}\|^2 \ge \|\mathbf{u}\|^2 = \|P_\mathcal{U}\mathbf{v}\|^2$.
Since norms are nonnegative, $\|P_\mathcal{U}\mathbf{v}\| \le \|\mathbf{v}\|$.
(ix) This is the explicit $\mathcal{U}$-component constructed in the proof of Theorem 2.2.4.

Structural Commentary: Orthogonal projection keeps the $\mathcal{U}$-component and discards the $\mathcal{U}^\perp$-component.
Key Insight: $P_\mathcal{U}$ is the identity on $\mathcal{U}$ and zero on $\mathcal{U}^\perp$. □

Definition 2.2.13. *Rank-One Projection Matrix*
Let $\mathbf{u} \in \mathbb{F}^n$ be a unit vector. The *rank-one projection matrix* onto $\operatorname{span}(\mathbf{u})$ is $P_\mathbf{u} := \mathbf{u}\mathbf{u}^* \in \operatorname{Mat}(n, n, \mathbb{F})$, where * denotes conjugate transpose. For every $\mathbf{x} \in \mathbb{F}^n$, $P_\mathbf{u}\mathbf{x} = \langle \mathbf{x}, \mathbf{u} \rangle_{\mathbb{F}^n} \mathbf{u}$.

Proposition 2.2.14. *Properties of Rank-One Projection Matrices*
Let $\mathbf{u} \in \mathbb{F}^n$ be a unit vector. Then:

i. $\operatorname{range}(P_{\mathbf{u}}) = \operatorname{span}(\mathbf{u})$;
ii. $\ker(P_{\mathbf{u}}) = \operatorname{span}(\mathbf{u})^{\perp}$;
iii. $P_{\mathbf{u}}^* = P_{\mathbf{u}}$;
iv. $P_{\mathbf{u}}^2 = P_{\mathbf{u}}$.

Proof.

(i) The formula $P_{\mathbf{u}}\mathbf{x} = \langle \mathbf{x}, \mathbf{u} \rangle \mathbf{u}$ gives $\operatorname{range}(P_{\mathbf{u}}) \subseteq \operatorname{span}(\mathbf{u})$, and $P_{\mathbf{u}}\mathbf{u} = \mathbf{u}$ gives the reverse inclusion.
(ii) The same formula gives $P_{\mathbf{u}}\mathbf{x} = 0$ exactly when $\langle \mathbf{x}, \mathbf{u} \rangle \mathbf{u} = \mathbf{0}$.
Since $\mathbf{u} \neq \mathbf{0}$, this is equivalent to $\langle \mathbf{x}, \mathbf{u} \rangle = 0$, i.e. $\mathbf{x} \in \operatorname{span}(\mathbf{u})^{\perp}$.
(iii) Since $P_{\mathbf{u}} = \mathbf{u}\mathbf{u}^*$, $P_{\mathbf{u}}^* = (\mathbf{u}\mathbf{u}^*)^* = \mathbf{u}\mathbf{u}^* = P_{\mathbf{u}}$.
(iv) Because $\mathbf{u}^*\mathbf{u} = \|\mathbf{u}\|_2^2 = 1$, $P_{\mathbf{u}}^2 = \mathbf{u}(\mathbf{u}^*\mathbf{u})\mathbf{u}^* = \mathbf{u}\mathbf{u}^* = P_{\mathbf{u}}$.

Structural Commentary: A rank-one projection keeps the $\mathbf{u}$ direction and kills its orthogonal complement.
Key Insight: The outer product $\mathbf{u}\mathbf{u}^*$ encodes projection onto one unit direction. □

Theorem 2.2.15. *Orthogonal Projection Minimizes Distance*
Let $\mathcal{U}$ be a finite-dimensional subspace of $\mathcal{V}$. If $\mathbf{v} \in \mathcal{V}$ and $\mathbf{u} \in \mathcal{U}$, then $\|\mathbf{v} - P_{\mathcal{U}}\mathbf{v}\| \leq \|\mathbf{v} - \mathbf{u}\|$.
Equality holds if and only if $\mathbf{u} = P_{\mathcal{U}}\mathbf{v}$.

Proof. Set $\mathbf{p} = P_{\mathcal{U}}\mathbf{v}$, $\mathbf{r} = \mathbf{v} - \mathbf{p} \in \mathcal{U}^{\perp}$, and $\mathbf{h} = \mathbf{p} - \mathbf{u} \in \mathcal{U}$.
Then $\mathbf{r} \perp \mathbf{h}$, so $\|\mathbf{v} - \mathbf{u}\|^2 = \|\mathbf{r} + \mathbf{h}\|^2 = \|\mathbf{r}\|^2 + \|\mathbf{h}\|^2 \geq \|\mathbf{r}\|^2 = \|\mathbf{v} - P_{\mathcal{U}}\mathbf{v}\|^2$.

($\Rightarrow$) If equality holds, then $\|\mathbf{h}\| = 0$, hence $\mathbf{u} = \mathbf{p} = P_{\mathcal{U}}\mathbf{v}$.
($\Leftarrow$) If $\mathbf{u} = P_{\mathcal{U}}\mathbf{v}$, then equality is immediate.

Structural Commentary: Projection error is orthogonal to every in-subspace correction.
Key Insight: The projection is the unique nearest point. □

Definition 2.2.16. *Fourier Coefficients and Expansion*
Let $\mathcal{O} = (\mathbf{e}_1, \ldots, \mathbf{e}_m)$ be an orthonormal list and $\mathbf{v} \in \mathcal{V}$.
The *Fourier coefficients* of $\mathbf{v}$ relative to $\mathcal{O}$ are $c_j(\mathbf{v}) := \langle \mathbf{v}, \mathbf{e}_j \rangle$, $1 \leq j \leq m$.
The *Fourier expansion* of $\mathbf{v}$ onto $\mathcal{O}$ is $\widehat{\mathbf{v}}_{\mathcal{O}} := \sum_{j=1}^{m} c_j(\mathbf{v})\mathbf{e}_j$.

Theorem 2.2.17. *Fourier Expansion onto a Finite Orthonormal Set*
Let $\mathcal{O} = (\mathbf{e}_1, \ldots, \mathbf{e}_m)$ be orthonormal, let $\mathcal{S} = \operatorname{span}(\mathcal{O})$, and let $\mathbf{v} \in \mathcal{V}$. Then:

i. $\widehat{\mathbf{v}}_{\mathcal{O}}$ *is the unique* $\mathbf{s} \in \mathcal{S}$ *with* $\mathbf{v} - \mathbf{s} \in \mathcal{S}^{\perp}$;
ii. $\widehat{\mathbf{v}}_{\mathcal{O}} = \arg\min_{\mathbf{s} \in \mathcal{S}} \|\mathbf{v} - \mathbf{s}\|$;
iii. $\|\widehat{\mathbf{v}}_{\mathcal{O}}\| \leq \|\mathbf{v}\|$.

Proof.

(i) For $1 \leq k \leq m$, $\langle \mathbf{v} - \widehat{\mathbf{v}}_{\mathcal{O}}, \mathbf{e}_k \rangle = c_k(\mathbf{v}) - \sum_{j=1}^{m} c_j(\mathbf{v})\delta_{jk} = 0$. Hence $\mathbf{v} - \widehat{\mathbf{v}}_{\mathcal{O}} \in \mathcal{S}^{\perp}$.
(ii) If $\mathbf{s} \in \mathcal{S}$ also satisfies $\mathbf{v} - \mathbf{s} \in \mathcal{S}^{\perp}$, then $\mathbf{s} - \widehat{\mathbf{v}}_{\mathcal{O}} \in \mathcal{S} \cap \mathcal{S}^{\perp} = \{\mathbf{0}\}$.
(iii) By (i), $\widehat{\mathbf{v}}_{\mathcal{O}} = P_{\mathcal{S}}\mathbf{v}$; Theorem 2.2.15 gives the minimizer.
(iv) By Theorem 2.1.52, $\|\widehat{\mathbf{v}}_{\mathcal{O}}\|^2 = \sum_{j=1}^{m} |c_j(\mathbf{v})|^2 \leq \|\mathbf{v}\|^2$, hence $\|\widehat{\mathbf{v}}_{\mathcal{O}}\| \leq \|\mathbf{v}\|$.

Structural Commentary: Fourier expansion is orthogonal projection in orthonormal coordinates.
Key Insight: The coefficients are obtained by pairing with the basis vectors because all cross terms vanish. □

Theorem 2.2.18. *Characterizing Finite Orthonormal Bases*
Let $\mathcal{O} = (\mathbf{e}_1, \ldots, \mathbf{e}_m)$ be orthonormal, let $\mathcal{S} = \operatorname{span}(\mathcal{O})$, and write $c_j(\mathbf{v})$ and $\widehat{\mathbf{v}} := \widehat{\mathbf{v}}_{\mathcal{O}}$ as in Definition 2.2.16. The following are equivalent:

i. $\mathcal{S} = \mathcal{V}$.
ii. $\mathcal{S}^{\perp} = \{\mathbf{0}\}$.
iii. $\widehat{\mathbf{v}} = \mathbf{v}$ *for all* $\mathbf{v} \in \mathcal{V}$.
iv. $\|\widehat{\mathbf{v}}\| = \|\mathbf{v}\|$ *for all* $\mathbf{v} \in \mathcal{V}$.
v. For all $\mathbf{v}, \mathbf{w} \in \mathcal{V}$, $\langle \mathbf{v}, \mathbf{w} \rangle = \sum_{j=1}^{m} c_j(\mathbf{v})\overline{c_j(\mathbf{w})}$.

Proof.

(*i*) $\Rightarrow$ (*ii*) If $\mathcal{S} = \mathcal{V}$, then $\mathcal{S}^{\perp} = \mathcal{V}^{\perp} = \{\mathbf{0}\}$.

(*ii*) $\Rightarrow$ (*iii*) By Theorem 2.2.17, $\mathbf{v} - \widehat{\mathbf{v}} \in \mathcal{S}^\perp = \{\mathbf{0}\}$.
(*iii*) $\Rightarrow$ (*i*) Since $\widehat{\mathbf{v}} \in \mathcal{S}$ for all $\mathbf{v}$, condition (*iii*) gives $\mathcal{V} \subseteq \mathcal{S}$, hence $\mathcal{S} = \mathcal{V}$.
(*iii*) $\Rightarrow$ (*iv*) If $\widehat{\mathbf{v}} = \mathbf{v}$, then $\|\widehat{\mathbf{v}}\| = \|\mathbf{v}\|$.
(*iv*) $\Rightarrow$ (*iii*) The decomposition $\mathbf{v} = \widehat{\mathbf{v}} + (\mathbf{v} - \widehat{\mathbf{v}})$ is orthogonal, so $\|\mathbf{v}\|^2 = \|\widehat{\mathbf{v}}\|^2 + \|\mathbf{v} - \widehat{\mathbf{v}}\|^2$. Together with (*iv*), this gives $\mathbf{v} = \widehat{\mathbf{v}}$.
(*iii*) $\Rightarrow$ (*v*) If $\widehat{\mathbf{v}} = \mathbf{v}$ and $\widehat{\mathbf{w}} = \mathbf{w}$, then orthonormality gives $\langle \mathbf{v}, \mathbf{w} \rangle = \sum_{j=1}^m c_j(\mathbf{v})\overline{c_j(\mathbf{w})}$.
(*v*) $\Rightarrow$ (*iv*) Put $\mathbf{w} = \mathbf{v}$ and use $\|\widehat{\mathbf{v}}\|^2 = \sum_{j=1}^m |c_j(\mathbf{v})|^2$ to get $\|\mathbf{v}\| = \|\widehat{\mathbf{v}}\|$.

Structural Commentary: Each condition says projection onto $\mathcal{S}$ loses no component.
Key Insight: Basis means lossless projection. □

Definition 2.2.19. *Distance to a Subspace*
Let $\mathcal{U}$ be a nonempty subset of an inner-product space $\mathcal{V}$.
The *distance* from $\mathbf{v} \in \mathcal{V}$ to $\mathcal{U}$ is $d(\mathbf{v}, \mathcal{U}) := \inf_{\mathbf{u} \in \mathcal{U}} \|\mathbf{v} - \mathbf{u}\|$.
If $\mathcal{U}$ is finite-dimensional, Theorem 2.2.15 gives $d(\mathbf{v}, \mathcal{U}) = \|\mathbf{v} - P_{\mathcal{U}}\mathbf{v}\|$.

Perspective 2.2.20. *Why the Finite-Dimensional Hypothesis Matters*
In $\ell^2(\mathbb{F})$, let $\mathcal{S} = \mathbb{F}^{(\mathbb{N})} = \operatorname{span}\{\mathbf{e}_j : j \geq 1\}$.
If $\mathbf{x} \in \mathcal{S}^\perp$, then $x_j = \langle \mathbf{x}, \mathbf{e}_j \rangle = 0 \quad (\forall j)$, so $\mathcal{S}^\perp = \{\mathbf{0}\}$. Since $\mathcal{S} \subsetneq \ell^2(\mathbb{F})$, $\mathcal{S} \oplus \mathcal{S}^\perp \neq \ell^2(\mathbb{F})$.

2.3 Least Squares and Metric Optimization

Least squares is projection theory as optimization: normal equations, full-rank formulas, and minimum-norm solutions.

Principle 2.3.1. *Least Squares as Orthogonal Projection*
In an inner-product space, the best approximation to $\mathbf{v}$ from a finite-dimensional subspace $\mathcal{U}$ is $P_{\mathcal{U}}\mathbf{v}$. The residual $\mathbf{v} - P_{\mathcal{U}}\mathbf{v}$ is orthogonal to every allowable correction inside $\mathcal{U}$.

Definition 2.3.2. *Matrix Least-Squares Solution*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ and $\mathbf{y} \in \mathbb{F}^m$.
A vector $\mathbf{x}_0 \in \mathbb{F}^n$ is a *least-squares solution* of $A\mathbf{x} = \mathbf{y}$ if $\|A\mathbf{x}_0 - \mathbf{y}\|_2 = \min_{\mathbf{x} \in \mathbb{F}^n} \|A\mathbf{x} - \mathbf{y}\|_2$.
The vector $A\mathbf{x}_0$ is the fitted vector, and $\mathbf{y} - A\mathbf{x}_0$ is the residual.

Theorem 2.3.3. *Normal Equations for Projection*
Let $\mathcal{U}$ be a finite-dimensional subspace of an inner-product space $\mathcal{V}$, and suppose $\mathcal{U} = \operatorname{span}(\mathbf{u}_1, \ldots, \mathbf{u}_n)$. For $\mathbf{v} \in \mathcal{V}$, the projection has the form $P_{\mathcal{U}}\mathbf{v} = \sum_{j=1}^n c_j \mathbf{u}_j$, where $(c_1, \ldots, c_n)$ solves the normal equations

$$\begin{bmatrix} \langle \mathbf{u}_1, \mathbf{u}_1 \rangle & \langle \mathbf{u}_2, \mathbf{u}_1 \rangle & \cdots & \langle \mathbf{u}_n, \mathbf{u}_1 \rangle \\ \langle \mathbf{u}_1, \mathbf{u}_2 \rangle & \langle \mathbf{u}_2, \mathbf{u}_2 \rangle & \cdots & \langle \mathbf{u}_n, \mathbf{u}_2 \rangle \\ \vdots & \vdots & \ddots & \vdots \\ \langle \mathbf{u}_1, \mathbf{u}_n \rangle & \langle \mathbf{u}_2, \mathbf{u}_n \rangle & \cdots & \langle \mathbf{u}_n, \mathbf{u}_n \rangle \end{bmatrix} \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_n \end{bmatrix} = \begin{bmatrix} \langle \mathbf{v}, \mathbf{u}_1 \rangle \\ \langle \mathbf{v}, \mathbf{u}_2 \rangle \\ \vdots \\ \langle \mathbf{v}, \mathbf{u}_n \rangle \end{bmatrix}.$$

This system is consistent. If $\mathbf{u}_1, \ldots, \mathbf{u}_n$ are linearly independent, then the solution is unique.

Proof. Since $P_{\mathcal{U}}\mathbf{v} \in \mathcal{U}$, write $P_{\mathcal{U}}\mathbf{v} = \sum_{j=1}^n c_j \mathbf{u}_j$.
The residual $\mathbf{v} - P_{\mathcal{U}}\mathbf{v}$ lies in $\mathcal{U}^\perp$, so for each i, $0 = \left\langle \mathbf{v} - \sum_{j=1}^n c_j \mathbf{u}_j, \mathbf{u}_i \right\rangle = \langle \mathbf{v}, \mathbf{u}_i \rangle - \sum_{j=1}^n c_j \langle \mathbf{u}_j, \mathbf{u}_i \rangle$.
These equations are exactly the displayed system, so consistency follows from the existence of $P_{\mathcal{U}}\mathbf{v}$.
Conversely, if $(d_1, \ldots, d_n)$ solves displayed system and $\mathbf{q} := \sum_{j=1}^n d_j \mathbf{u}_j$, then $\langle \mathbf{v} - \mathbf{q}, \mathbf{u}_i \rangle = 0$ for every i.
Since $\mathbf{u}_i$ span $\mathcal{U}$, it follows that $\mathbf{v} - \mathbf{q} \in \mathcal{U}^\perp$, hence $\mathbf{q} = P_{\mathcal{U}}\mathbf{v}$ by uniqueness of the decomposition. If $\mathbf{u}_1, \ldots, \mathbf{u}_n$ are linearly independent, they form a basis of $\mathcal{U}$, so the coordinates of this projection in that basis are unique.
Structural Commentary: Normal equations express residual orthogonality.
Key Insight: The projection error is perpendicular to $\mathcal{U}$. □

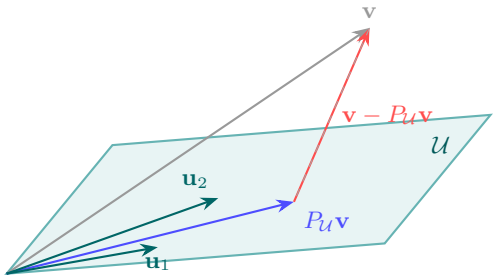


Figure 31: Normal equations impose $\langle \mathbf{v} - P_{\mathcal{U}}\mathbf{v}, \mathbf{u}_i \rangle = 0$ for each spanning vector.

Theorem 2.3.4. *Matrix Least Squares and Normal Equations*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$ and $\mathbf{y} \in \mathbb{F}^m$, where A^* denotes the conjugate transpose of A.
A vector $\mathbf{x}_0 \in \mathbb{F}^n$ satisfies $\|\mathbf{y} - A\mathbf{x}_0\|_2 = \min_{\mathbf{x}\in\mathbb{F}^n} \|\mathbf{y} - A\mathbf{x}\|_2$ if and only if $A^*A\mathbf{x}_0 = A^*\mathbf{y}$.
The normal equations are always consistent, and their solution is unique if and only if $\mathrm{rank}(A) = n$.

Proof. Let $\mathcal{U} = \mathrm{col}(A)$, and write $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n]$. Since $\{A\mathbf{x} : \mathbf{x} \in \mathbb{F}^n\} = \mathcal{U}$, Theorem 2.2.15 shows that $\mathbf{x}_0$ is least-squares optimal exactly when $A\mathbf{x}_0 = P_{\mathcal{U}}\mathbf{y}$, or equivalently when $\mathbf{y} - A\mathbf{x}_0 \in \mathcal{U}^\perp$.

($\Rightarrow$) If $\mathbf{x}_0$ is least-squares optimal, then $\mathbf{y} - A\mathbf{x}_0 \in \mathcal{U}^\perp$. Hence $\langle \mathbf{y} - A\mathbf{x}_0, \mathbf{a}_j \rangle = 0$, $j = 1, \ldots, n$.
These are the entries of $A^*(\mathbf{y} - A\mathbf{x}_0)$, so $A^*(\mathbf{y} - A\mathbf{x}_0) = \mathbf{0}$ and $A^*A\mathbf{x}_0 = A^*\mathbf{y}$.

($\Leftarrow$) If $A^*A\mathbf{x}_0 = A^*\mathbf{y}$, then $A^*(\mathbf{y} - A\mathbf{x}_0) = \mathbf{0}$. Thus $\langle \mathbf{y} - A\mathbf{x}_0, \mathbf{a}_j \rangle = 0$ for every j, so $\mathbf{y} - A\mathbf{x}_0$ is orthogonal to every linear combination of the columns of A. Hence $\mathbf{y} - A\mathbf{x}_0 \in \mathcal{U}^\perp$.
Hence $A\mathbf{x}_0 = P_{\mathcal{U}}\mathbf{y}$, so $\mathbf{x}_0$ is least-squares optimal.

Consistency follows by choosing $\mathbf{x}_0$ such that $A\mathbf{x}_0 = P_{\mathcal{U}}\mathbf{y}$.
If $\mathbf{x}_0, \mathbf{x}_1$ solve the normal equations, then $\|A(\mathbf{x}_0 - \mathbf{x}_1)\|_2^2 = \langle \mathbf{x}_0 - \mathbf{x}_1, A^*A(\mathbf{x}_0 - \mathbf{x}_1) \rangle = 0$, so $\mathbf{x}_0 - \mathbf{x}_1 \in \ker(A)$.
Conversely, if $\mathbf{z} \in \ker(A)$ and $\mathbf{x}_0$ solves the normal equations, then $\mathbf{x}_0 + \mathbf{z}$ also solves them. Thus uniqueness is equivalent to $\ker(A) = \{\mathbf{0}\}$, equivalently $\mathrm{rank}(A) = n$.
Structural Commentary: Least squares is orthogonal projection onto the column space of A.
Key Insight: The normal equations say exactly that the residual is orthogonal to every column of A. □

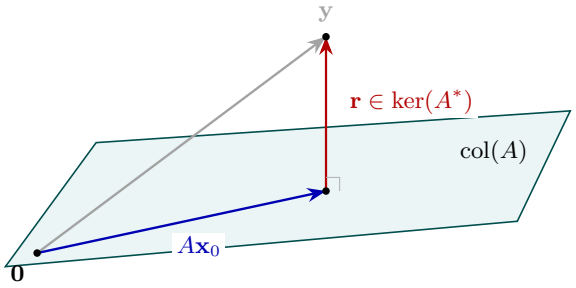


Figure 32: Fitted vector and orthogonal residual: $A\mathbf{x}_0 = P_{\mathrm{col}(A)}\mathbf{y}$ and $\mathbf{r} = \mathbf{y} - A\mathbf{x}_0 \in \mathrm{col}(A)^\perp = \ker(A^*)$.

Proposition 2.3.5. *Least-Squares Solution Set and Minimum Norm*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$, $\mathbf{y} \in \mathbb{F}^m$, and $\mathrm{LS}(A,\mathbf{y}) = \arg\min_{\mathbf{x}\in\mathbb{F}^n} \|A\mathbf{x} - \mathbf{y}\|_2$. If $\mathbf{x}_0 \in \mathrm{LS}(A,\mathbf{y})$, then $\mathrm{LS}(A,\mathbf{y}) = \mathbf{x}_0 + \ker(A)$.
Also $\mathrm{LS}(A,\mathbf{y}) \cap (\ker A)^\perp = \{\mathbf{s}\}$, where $\mathbf{s} = P_{(\ker A)^\perp}\mathbf{x}_0 \in \mathrm{col}(A^*)$, and $\mathbf{s}$ is unique minimizer of $\|\cdot\|_2$ on $\mathrm{LS}(A,\mathbf{y})$.

Proof. Set $\mathbf{b} = P_{\mathrm{col}(A)}\mathbf{y}$. By Theorem 2.3.4, $\mathrm{LS}(A,\mathbf{y}) = \{\mathbf{x} \in \mathbb{F}^n : A\mathbf{x} = \mathbf{b}\}$.
Thus, for $\mathbf{x}_0 \in \mathrm{LS}(A,\mathbf{y})$, $\mathrm{LS}(A,\mathbf{y}) = \{\mathbf{x} : A(\mathbf{x} - \mathbf{x}_0) = \mathbf{0}\} = \mathbf{x}_0 + \ker(A)$.
Write $\mathbf{x}_0 = \mathbf{s} + \mathbf{k}$ with $\mathbf{s} \in (\ker A)^\perp$ and $\mathbf{k} \in \ker(A)$. Since $A\mathbf{s} = A\mathbf{x}_0$, $\mathbf{s} \in \mathrm{LS}(A,\mathbf{y}) \cap (\ker A)^\perp$.
If $\mathbf{u}$ is another element of this intersection, then $\mathbf{u} - \mathbf{s} \in \ker(A) \cap (\ker A)^\perp = \{\mathbf{0}\}$.
Finally, any $\mathbf{x} \in \mathrm{LS}(A,\mathbf{y})$ has $\mathbf{x} = \mathbf{s} + \mathbf{z}$ with $\mathbf{z} \in \ker(A)$, so $\|\mathbf{x}\|_2^2 = \|\mathbf{s}\|_2^2 + \|\mathbf{z}\|_2^2 \geq \|\mathbf{s}\|_2^2$, with equality iff $\mathbf{z} = \mathbf{0}$.
Structural Commentary: Least-squares solutions form an affine translate of the nullspace.
Key Insight: The minimum-norm representative is the one with no nullspace component. □

Corollary 2.3.6. *Projection Formula for Full Column Rank*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$ have rank n.
Then A^*A is invertible and the orthogonal projection onto $\mathrm{col}(A)$ is $P_{\mathrm{col}(A)} = A(A^*A)^{-1}A^*$.

Proof. If $A^*A\mathbf{x} = \mathbf{0}$, then $\|A\mathbf{x}\|_2^2 = \langle \mathbf{x}, A^*A\mathbf{x}\rangle = 0$, so $A\mathbf{x} = \mathbf{0}$.
Since $\operatorname{rank}(A) = n$, we get $\mathbf{x} = \mathbf{0}$. Hence A^*A is invertible. By Theorem 2.3.4, the least-squares solution is $\mathbf{x}_0 = (A^*A)^{-1}A^*\mathbf{y}$. The fitted vector is $P_{\operatorname{col}(A)}\mathbf{y} = A\mathbf{x}_0 = A(A^*A)^{-1}A^*\mathbf{y}$.
Since this holds for every $\mathbf{y}$, the displayed operator formula follows.
Structural Commentary: Full column rank makes the normal-equation coordinates unique.
Key Insight: Solve for coefficients, then map back by A. □

Definition 2.3.7. *Minimum Norm Solution*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ and $\mathbf{y} \in \operatorname{col}(A)$.
A solution $\mathbf{s}$ of $A\mathbf{x} = \mathbf{y}$ is a *minimum norm solution* if $\|\mathbf{s}\|_2 \le \|\mathbf{x}\|_2$ for every solution $\mathbf{x}$ of $A\mathbf{x} = \mathbf{y}$.

Theorem 2.3.8. *Minimum Norm Solution of a Consistent Linear System*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$, $\mathbf{y} \in \mathbb{F}^m$, and suppose $A\mathbf{x} = \mathbf{y}$ is consistent.

i. There is a unique minimum norm solution $\mathbf{s} \in \mathbb{F}^n$, and $\mathbf{s} \in \operatorname{col}(A^*)$.
ii. $\mathbf{s}$ is the only solution of $A\mathbf{x} = \mathbf{y}$ that lies in $\operatorname{col}(A^*)$.
iii. The system $AA^*\mathbf{u} = \mathbf{y}$ is consistent. If $AA^*\mathbf{u}_0 = \mathbf{y}$, then $\mathbf{s} = A^*\mathbf{u}_0$.

Proof. First note that $\operatorname{col}(A^*) = (\ker A)^\perp$. Indeed, if $\mathbf{z} = A^*\mathbf{u}$ and $\mathbf{x} \in \ker(A)$, then $\langle \mathbf{x}, \mathbf{z}\rangle = \langle A\mathbf{x}, \mathbf{u}\rangle = 0$, so $\mathbf{z}$ is orthogonal to every vector in $\ker(A)$. Thus $\operatorname{col}(A^*) \subseteq (\ker A)^\perp$. Also,

$$\dim(\operatorname{col}(A^*)) = \operatorname{rank}(A^*) = \operatorname{rank}(A) = n - \dim(\ker(A)) = \dim((\ker A)^\perp),$$

so the two subspaces are equal. Hence $\mathbb{F}^n = \operatorname{col}(A^*) \oplus \ker(A)$ orthogonally.

(i) Choose a solution $\mathbf{x}_0$ of $A\mathbf{x} = \mathbf{y}$ and write $\mathbf{x}_0 = \mathbf{s} + \mathbf{v}$, with $\mathbf{s} \in \operatorname{col}(A^*)$ and $\mathbf{v} \in \ker(A)$.
Then $A\mathbf{s} = A\mathbf{x}_0 = \mathbf{y}$. If $\mathbf{x}$ is any solution, then $\mathbf{x} - \mathbf{s} \in \ker(A)$, so every solution has the form $\mathbf{s} + \mathbf{w}$ with $\mathbf{w} \in \ker(A)$. Since $\mathbf{s} \perp \mathbf{w}$, $\|\mathbf{s} + \mathbf{w}\|_2^2 = \|\mathbf{s}\|_2^2 + \|\mathbf{w}\|_2^2 \ge \|\mathbf{s}\|_2^2$, with equality only when $\mathbf{w} = \mathbf{0}$.
Thus $\mathbf{s}$ is the unique solution of minimum norm.
(ii) If $\mathbf{s}' \in \operatorname{col}(A^*)$ is another solution, then $\mathbf{s}' - \mathbf{s} \in \ker(A) \cap \operatorname{col}(A^*) = \{\mathbf{0}\}$, so $\mathbf{s}' = \mathbf{s}$.
(iii) Since $\mathbf{s} \in \operatorname{col}(A^*)$, write $\mathbf{s} = A^*\mathbf{z}$. Then $AA^*\mathbf{z} = A\mathbf{s} = \mathbf{y}$, so $AA^*\mathbf{u} = \mathbf{y}$ is consistent.
Conversely, if $AA^*\mathbf{u}_0 = \mathbf{y}$, then $A^*\mathbf{u}_0$ is a solution in $\operatorname{col}(A^*)$, so $A^*\mathbf{u}_0 = \mathbf{s}$ by (ii).

Structural Commentary: Solutions are $\mathbf{s} + \ker(A)$; $\operatorname{col}(A^*)$ meets this set orthogonally.
Key Insight: The minimum norm solution has no nullspace component. □

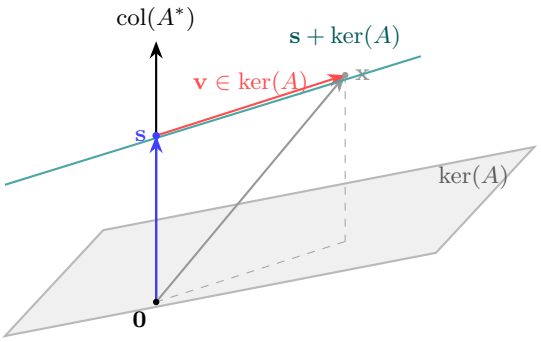


Figure 33: Among all solutions $\mathbf{s} + \ker(A)$, the unique one in $\operatorname{col}(A^*) = (\ker A)^\perp$ has minimum norm.

Perspective 2.3.9. *Roadmap to the Pseudoinverse*
Theorem 2.3.4 projects $\mathbf{y}$ onto $\operatorname{col}(A)$; Proposition 2.3.5 selects the minimum-norm preimage in $\operatorname{col}(A^*)$. Definition 4.3.16 and Theorem 4.3.21 package both choices as $A^\dagger\mathbf{y}$.

2.4 Adjoints and Orthogonal Operators

Adjoints express dual action in inner-product language. They become conjugate transposes in orthonormal coordinates and lead to self-adjoint maps, projections, isometries, and unitary or orthogonal operators.

Definition 2.4.1. *Adjoint of a Linear Map*
Let $\mathcal{V}$ and $\mathcal{W}$ be inner-product spaces and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.
A function $S : \mathcal{W} \to \mathcal{V}$ is an *adjoint* of T if $\langle T\mathbf{v}, \mathbf{w}\rangle_{\mathcal{W}} = \langle \mathbf{v}, S\mathbf{w}\rangle_{\mathcal{V}}$, $\mathbf{v} \in \mathcal{V}$, $\mathbf{w} \in \mathcal{W}$.

Perspective 2.4.2. *Dual Map and Inner-Product Adjoint*
For $T : \mathcal{V} \to \mathcal{W}$, the pullback $\varphi \mapsto \varphi \circ T$ is the *dual map* $T' : \mathcal{W}' \to \mathcal{V}'$. Also called *operator adjoint* and written $T^\times$. It is distinct from the inner-product adjoint T^* of Definition 2.4.1, which requires inner products.

Lemma 2.4.3. *Uniqueness of Adjoints*
A linear map has at most one adjoint.

Proof. Suppose $S_1, S_2 : \mathcal{W} \to \mathcal{V}$ are adjoints of T. For each $\mathbf{w} \in \mathcal{W}$ and every $\mathbf{v} \in \mathcal{V}$, $\langle \mathbf{v}, S_1\mathbf{w} - S_2\mathbf{w}\rangle_{\mathcal{V}} = 0$. Taking $\mathbf{v} = S_1\mathbf{w} - S_2\mathbf{w}$ gives $S_1\mathbf{w} = S_2\mathbf{w}$. Hence $S_1 = S_2$.
Structural Commentary: Same tests force candidate adjoints to agree.
Key Insight: Inner products separate vectors. □

Theorem 2.4.4. *Existence of the Adjoint*
Let $\mathcal{V}$ be a finite-dimensional inner-product space, let $\mathcal{W}$ be an inner-product space, and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. Then T has a unique adjoint $T^* : \mathcal{W} \to \mathcal{V}$, and T^* is linear.

Proof. For each $\mathbf{w} \in \mathcal{W}$, the map $\varphi_{\mathbf{w}} : \mathcal{V} \to \mathbb{F}$ defined by $\varphi_{\mathbf{w}}(\mathbf{v}) = \langle T\mathbf{v}, \mathbf{w}\rangle_{\mathcal{W}}$ is a linear functional. By Theorem 2.1.57, there is a unique vector $T^*\mathbf{w} \in \mathcal{V}$ such that $\langle T\mathbf{v}, \mathbf{w}\rangle_{\mathcal{W}} = \langle \mathbf{v}, T^*\mathbf{w}\rangle_{\mathcal{V}}$ for every $\mathbf{v} \in \mathcal{V}$.
Thus T^* is an adjoint, and it is unique by Lemma 2.4.3. For $\mathbf{w}_1, \mathbf{w}_2 \in \mathcal{W}$ and $a \in \mathbb{F}$,

$$\langle \mathbf{v}, T^*(a\mathbf{w}_1 + \mathbf{w}_2)\rangle_{\mathcal{V}} = \langle T\mathbf{v}, a\mathbf{w}_1 + \mathbf{w}_2\rangle_{\mathcal{W}} = \overline{a}\,\langle T\mathbf{v}, \mathbf{w}_1\rangle_{\mathcal{W}} + \langle T\mathbf{v}, \mathbf{w}_2\rangle_{\mathcal{W}} = \langle \mathbf{v}, aT^*\mathbf{w}_1 + T^*\mathbf{w}_2\rangle_{\mathcal{V}}\,.$$

for every $\mathbf{v} \in \mathcal{V}$. Taking the difference of the two second-slot vectors and then setting $\mathbf{v}$ equal to that difference gives $T^*(a\mathbf{w}_1 + \mathbf{w}_2) = aT^*\mathbf{w}_1 + T^*\mathbf{w}_2$.
Structural Commentary: Each $\mathbf{w}$ gives a Riesz-represented functional on $\mathcal{V}$.
Key Insight: The adjoint is pointwise Riesz representation for T. □

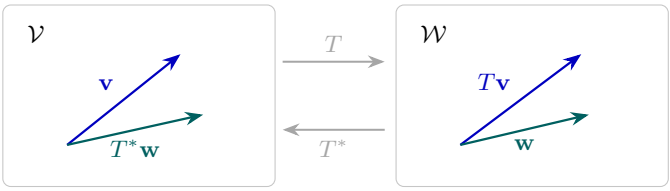


Figure 34: Adjoint relation: $\langle T\mathbf{v}, \mathbf{w}\rangle_{\mathcal{W}} = \langle \mathbf{v}, T^*\mathbf{w}\rangle_{\mathcal{V}}$.

Perspective 2.4.5. *Hilbert Adjoints as Riesz-Transported Dual Maps*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional inner-product spaces, and let $R_{\mathcal{V}} : \mathcal{V}' \to \mathcal{V}$ and $R_{\mathcal{W}} : \mathcal{W}' \to \mathcal{W}$ send each functional to its Riesz vector from Definition 2.1.58. If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then the dual map $T' : \mathcal{W}' \to \mathcal{V}'$ from Definition 1.5.11 and the Hilbert-space adjoint T^* satisfy $T' = \big(R_{\mathcal{V}}\big)^{-1} \circ T^* \circ R_{\mathcal{W}}$.
Indeed, for $\varphi \in \mathcal{W}'$ and $\mathbf{v} \in \mathcal{V}$, $(T'\varphi)(\mathbf{v}) = \varphi(T\mathbf{v}) = \langle T\mathbf{v}, R_{\mathcal{W}}\varphi\rangle_{\mathcal{W}} = \langle \mathbf{v}, T^*R_{\mathcal{W}}\varphi\rangle_{\mathcal{V}}$. Thus the algebraic pullback of functionals is the same operation as moving Riesz vectors by T^* and then returning to functionals.
Structural Commentary: The dual map belongs to the algebraic theory of Definition 1.5.11; the adjoint appears only after the Riesz identification from Theorem 2.1.57.
Key Insight: T' and T^* encode the same reversed action after functionals have been identified with vectors.

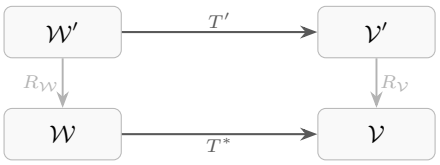


Figure 35: Riesz transports the dual map to the adjoint: $R_{\mathcal{V}}T' = T^*R_{\mathcal{W}}$.

Definition 2.4.6. *Conjugate Transpose*
For $A = [a_{ij}] \in \mathrm{Mat}(m, n, \mathbb{F})$, the *conjugate transpose* of A is $A^* = [\overline{a_{ji}}] \in \mathrm{Mat}(n, m, \mathbb{F})$.
Over $\mathbb{R}$, this is the ordinary transpose.

Theorem 2.4.7. *Matrix of the Adjoint*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional inner-product spaces. Let $\beta = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be an orthonormal basis of $\mathcal{V}$ and let $\gamma = (\mathbf{f}_1, \ldots, \mathbf{f}_m)$ be an orthonormal basis of $\mathcal{W}$. If $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$, then $[T^*]_{\beta\leftarrow\gamma} = \big([T]_{\gamma\leftarrow\beta}\big)^*$.

Proof. Let $A = [T]_{\gamma\leftarrow\beta}$ and $B = [T^*]_{\beta\leftarrow\gamma}$.
By Theorem 2.1.55, $B_{i,j} = \langle T^*\mathbf{f}_j, \mathbf{e}_i\rangle_{\mathcal{V}} = \overline{\langle \mathbf{e}_i, T^*\mathbf{f}_j\rangle_{\mathcal{V}}} = \overline{\langle T\mathbf{e}_i, \mathbf{f}_j\rangle_{\mathcal{W}}} = \overline{A_{j,i}}$. Thus $B = A^*$.
Structural Commentary: Entries swap by moving T across the inner product.
Key Insight: In orthonormal coordinates, taking adjoints is exactly taking conjugate transposes. □

Proposition 2.4.8. *Basic Adjoint Identities*
For finite-dimensional inner-product spaces and composable linear maps:

i. $(aS+T)^* = \overline{a}S^* + T^*$ *for* $S,T \in \mathcal{L}(\mathcal{V},\mathcal{W})$ *and* $a \in \mathbb{F}$;
ii. $(SR)^* = R^*S^*$ *for* $R \in \mathcal{L}(\mathcal{U},\mathcal{V})$ *and* $S \in \mathcal{L}(\mathcal{V},\mathcal{W})$;
iii. $(T^*)^* = T$;
iv. $I_{\mathcal{V}}^* = I_{\mathcal{V}}$.

Proof. Verify the defining identity; uniqueness gives equality.

(i) For $\mathbf{v} \in \mathcal{V}$, $\mathbf{w} \in \mathcal{W}$, $\langle (aS+T)\mathbf{v}, \mathbf{w}\rangle = a\langle S\mathbf{v},\mathbf{w}\rangle + \langle T\mathbf{v},\mathbf{w}\rangle = \langle \mathbf{v}, \overline{a}S^*\mathbf{w} + T^*\mathbf{w}\rangle$, so $(aS+T)^* = \overline{a}S^* + T^*$.
(ii) For $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{W}$, $\langle SR\mathbf{u},\mathbf{w}\rangle = \langle R\mathbf{u}, S^*\mathbf{w}\rangle = \langle \mathbf{u}, R^*S^*\mathbf{w}\rangle$, so $(SR)^* = R^*S^*$.
(iii) Conjugating $\langle T\mathbf{v},\mathbf{w}\rangle = \langle \mathbf{v}, T^*\mathbf{w}\rangle$ gives $\langle T^*\mathbf{w},\mathbf{v}\rangle = \langle \mathbf{w}, T\mathbf{v}\rangle$, so T is the adjoint of T^*.
(iv) Since $\langle I\mathbf{v},\mathbf{w}\rangle = \langle \mathbf{v}, I\mathbf{w}\rangle$, the identity map is its own adjoint.

Structural Commentary: Put both sides in one inner-product slot, then use uniqueness.
Key Insight: Like matrix transposes, adjoints reverse products and conjugate scalars. □

Proposition 2.4.9. *Adjoints of Inverses and Polynomials*
Let $T \in \mathcal{L}(\mathcal{V})$ *on a finite-dimensional inner-product space. For* $p(z) = \sum_{j=0}^{m} a_j z^j \in \mathbb{F}[z]$, *set* $p^{\#}(z) = \sum_{j=0}^{m} \overline{a_j} z^j$.

(i) If T *is invertible, then* T^* *is invertible and* $(T^{-1})^* = (T^*)^{-1}$.
(ii) $p(T)^* = p^{\#}(T^*)$. *In particular,* $p \in \mathbb{R}[z] \Rightarrow p(T)^* = p(T^*)$.

Proof.

(i) Proposition 2.4.8 gives $(T^{-1})^*T^* = (TT^{-1})^* = I_{\mathcal{V}}$ and $T^*(T^{-1})^* = (T^{-1}T)^* = I_{\mathcal{V}}$.
(ii) By conjugate-linearity in scalars and product reversal, $p(T)^* = \left(\sum_{j=0}^{m} a_j T^j\right)^* = \sum_{j=0}^{m} \overline{a_j}(T^*)^j = p^{\#}(T^*)$.

Structural Commentary: Adjoint reverses multiplication and conjugates scalars.
Key Insight: Adjoints send $p(T)$ to $p^{\#}(T^*)$. □

Theorem 2.4.10. *Adjoints and Invariant Orthogonal Complements*
Let $T \in \mathcal{L}(\mathcal{V})$ *on a finite-dimensional inner-product space and let* $\mathcal{U} \leq \mathcal{V}$.

(i) $\mathcal{U} \in \mathrm{Inv}(T) \iff \mathcal{U}^\perp \in \mathrm{Inv}(T^*)$.
(ii) $(\mathcal{U},\mathcal{U}^\perp)$ *reduces* $T \iff \mathcal{U} \in \mathrm{Inv}(T) \cap \mathrm{Inv}(T^*)$.
(iii) If $(\mathcal{U},\mathcal{U}^\perp)$ *reduces* T, *then* $(T|_{\mathcal{U}})^* = (T^*)|_{\mathcal{U}}$.

Proof.

(i) Since $\mathcal{U} = (\mathcal{U}^\perp)^\perp$, $\mathcal{U} \in \mathrm{Inv}(T) \iff \langle T\mathbf{u},\mathbf{z}\rangle = 0 \quad (\mathbf{u} \in \mathcal{U},\ \mathbf{z} \in \mathcal{U}^\perp)$
$\iff \langle \mathbf{u}, T^*\mathbf{z}\rangle = 0$, where $\mathbf{u} \in \mathcal{U},\ \mathbf{z} \in \mathcal{U}^\perp \iff \mathcal{U}^\perp \in \mathrm{Inv}(T^*)$.
(ii) By definition and then (i) applied to T^*, $(\mathcal{U},\mathcal{U}^\perp)$ reduces $T \iff \mathcal{U},\mathcal{U}^\perp \in \mathrm{Inv}(T) \iff \mathcal{U} \in \mathrm{Inv}(T) \cap \mathrm{Inv}(T^*)$.
(iii) By (ii), $(T^*)|_{\mathcal{U}}$ is well-defined. For $\mathbf{u},\mathbf{v} \in \mathcal{U}$,
$\langle T|_{\mathcal{U}}\mathbf{u},\mathbf{v}\rangle_{\mathcal{U}} = \langle T\mathbf{u},\mathbf{v}\rangle_{\mathcal{V}} = \langle \mathbf{u}, T^*\mathbf{v}\rangle_{\mathcal{V}} = \langle \mathbf{u}, (T^*)|_{\mathcal{U}}\mathbf{v}\rangle_{\mathcal{U}}$, so uniqueness of adjoints on $\mathcal{U}$ gives the claim.

Structural Commentary: Adjoints trade invariance with orthogonal-complement invariance.
Key Insight: A subspace reduces T precisely when it is stable under both T and T^*. □

Proposition 2.4.11. *Adjoint Kernel-Range Orthogonality*
Let $\mathcal{V}$ *and* $\mathcal{W}$ *be finite-dimensional inner-product spaces and let* $T \in \mathcal{L}(\mathcal{V},\mathcal{W})$.
Then $\ker(T^*) = (\mathrm{range}(T))^\perp$, $\mathrm{range}(T^*) = (\ker(T))^\perp$, *and* $\ker(T) = (\mathrm{range}(T^*))^\perp$, $\mathrm{range}(T) = (\ker(T^*))^\perp$.
Equivalently, the nullspace of an adjoint is the orthogonal complement of the range of the original map.

Proof. For $\mathbf{w} \in \mathcal{W}$,

$$\begin{aligned}\mathbf{w} \in \ker(T^*) &\iff \langle \mathbf{v}, T^*\mathbf{w}\rangle_{\mathcal{V}} = 0 \quad \text{for every } \mathbf{v} \in \mathcal{V} \\ &\iff \langle T\mathbf{v},\mathbf{w}\rangle_{\mathcal{W}} = 0 \quad \text{for every } \mathbf{v} \in \mathcal{V} \\ &\iff \mathbf{w} \in (\mathrm{range}(T))^\perp.\end{aligned}$$

Applying this identity to T^* gives $\ker(T) = (\mathrm{range}(T^*))^\perp$.
Since all spaces are finite-dimensional, the ranges and kernels are finite-dimensional subspaces.
Taking orthogonal complements and using Corollary 2.2.7 gives

$$(\ker(T^*))^\perp = ((\mathrm{range}(T))^\perp)^\perp = \mathrm{range}(T), \qquad (\ker(T))^\perp = ((\mathrm{range}(T^*))^\perp)^\perp = \mathrm{range}(T^*)$$

Structural Commentary: Adjoints trade ranges for kernels, and orthogonal complements reverse that trade.
Key Insight: A vector is killed by T^* exactly when it is orthogonal to everything T can produce. □

Proposition 2.4.12. *Adjoint Products and Algebraic Projections*
Let $\mathcal{V}$ and $\mathcal{W}$ be finite-dimensional inner-product spaces and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.

(i) $\ker(T^*T) = \ker(T)$ *and* $\ker(TT^*) = \ker(T^*)$.
(ii) $\operatorname{range}(T^*T) = \operatorname{range}(T^*)$ *and* $\operatorname{range}(TT^*) = \operatorname{range}(T)$.
(iii) If $\mathcal{V} = \mathcal{U} \oplus \mathcal{Z}$, then $\left(\rho_{\mathcal{U},\mathcal{Z}}\right)^* = \rho_{\mathcal{Z}^\perp,\mathcal{U}^\perp}$.

Proof.

(i) For $\mathbf{v} \in \mathcal{V}$, $\mathbf{v} \in \ker(T^*T) \iff 0 = \langle T^*T\mathbf{v}, \mathbf{v}\rangle_{\mathcal{V}} = \|T\mathbf{v}\|^2 \iff \mathbf{v} \in \ker(T)$.
Apply this to T^* for $\ker(TT^*) = \ker(T^*)$.

(ii) The inclusions $\operatorname{range}(T^*T) \subseteq \operatorname{range}(T^*)$ and $\operatorname{range}(TT^*) \subseteq \operatorname{range}(T)$ are immediate.
By (i), rank-nullity, and Proposition 2.4.11, $\dim \operatorname{range}(T^*T) = \dim \mathcal{V} - \dim \ker(T) = \dim \operatorname{range}(T^*)$, so equality holds; the TT^* identity follows similarly.

(iii) Put $P = \rho_{\mathcal{U},\mathcal{Z}}$ and $Q = \rho_{\mathcal{Z}^\perp,\mathcal{U}^\perp}$. Since $\mathcal{Z}^\perp \cap \mathcal{U}^\perp = (\mathcal{U}+\mathcal{Z})^\perp = \{\mathbf{0}\}$ and the dimensions add, $\mathcal{V} = \mathcal{Z}^\perp \oplus \mathcal{U}^\perp$.
Write $\mathbf{x} = \mathbf{u} + \mathbf{z}$ and $\mathbf{y} = \mathbf{a} + \mathbf{b}$ with $\mathbf{u} \in \mathcal{U}$, $\mathbf{z} \in \mathcal{Z}$, $\mathbf{a} \in \mathcal{Z}^\perp$, $\mathbf{b} \in \mathcal{U}^\perp$.
Then $\langle P\mathbf{x}, \mathbf{y}\rangle = \langle \mathbf{u}, \mathbf{a}+\mathbf{b}\rangle = \langle \mathbf{u}, \mathbf{a}\rangle = \langle \mathbf{u}+\mathbf{z}, \mathbf{a}\rangle = \langle \mathbf{x}, Q\mathbf{y}\rangle$, so $P^* = Q$.

Structural Commentary: Oblique projection adjoints swap orthogonal-complement summands.
Key Insight: Orthogonal projections are exactly the self-adjoint special case $\mathcal{Z} = \mathcal{U}^\perp$. □

Corollary 2.4.13. *Matrix Fundamental Subspaces as Orthogonal Complements*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ act as $A : \mathbb{F}^n \to \mathbb{F}^m$ with standard inner products.
Then $\ker(A) = \operatorname{col}(A^*)^\perp$ and $\operatorname{col}(A^*) = \ker(A)^\perp$; $\ker(A^*) = \operatorname{col}(A)^\perp$ and $\operatorname{col}(A) = \ker(A^*)^\perp$.
Over $\mathbb{R}$, this says $\ker(A) \perp \operatorname{row}(A)$ in $\mathbb{R}^n$ and $\ker(A^t) \perp \operatorname{col}(A)$ in $\mathbb{R}^m$, with both pairs complementary.

Proof. By Theorem 2.4.7, the adjoint of A is represented by A^*.
Proposition 2.4.11 gives $\ker(A^*) = \operatorname{col}(A)^\perp$ and $\ker(A) = \operatorname{col}(A^*)^\perp$. Taking orthogonal complements and using Corollary 2.2.7 gives the two column-space identities. Over $\mathbb{R}$, $A^* = A^t$ and $\operatorname{col}(A^t) = \operatorname{row}(A)$.
Structural Commentary: The four fundamental subspaces form two orthogonal-complement pairs.
Key Insight: Adjoints turn missing directions into kernels. □

Definition 2.4.14. *Self-Adjoint and Hermitian*
An operator $T \in \mathcal{L}(\mathcal{V})$ on a finite-dimensional inner-product space is *self-adjoint* if $T^* = T$. A square matrix A is *Hermitian* if $A^* = A$. In orthonormal coordinates, these are the same condition by Theorem 2.4.7.

Lemma 2.4.15. *Zero Complex Quadratic Form*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $T \in \mathcal{L}(\mathcal{V})$.
If $\langle T\mathbf{v}, \mathbf{v}\rangle = 0$ for every $\mathbf{v} \in \mathcal{V}$, then $T = 0$.

Proof. Set $Q(\mathbf{x}) = \langle T\mathbf{x}, \mathbf{x}\rangle$.
For $\mathbf{u}, \mathbf{w} \in \mathcal{V}$, expansion gives the polarization identity for the sesquilinear form $(\mathbf{u}, \mathbf{w}) \mapsto \langle T\mathbf{u}, \mathbf{w}\rangle$:

$$\langle T\mathbf{u}, \mathbf{w}\rangle = \frac{Q(\mathbf{u}+\mathbf{w}) - Q(\mathbf{u}-\mathbf{w})}{4} + i\,\frac{Q(\mathbf{u}+i\mathbf{w}) - Q(\mathbf{u}-i\mathbf{w})}{4}.$$

The hypothesis makes every term on the right zero, so $\langle T\mathbf{u}, \mathbf{w}\rangle = 0$ for all $\mathbf{u}, \mathbf{w}$.
Taking $\mathbf{w} = T\mathbf{u}$ gives $\|T\mathbf{u}\|^2 = 0$, hence $T\mathbf{u} = \mathbf{0}$ for every $\mathbf{u}$.
Structural Commentary: Over $\mathbb{C}$, diagonal values of a sesquilinear form determine all mixed values.
Key Insight: Polarization turns the one-vector condition into a two-vector condition. □

Proposition 2.4.16. *Self-Adjointness from Real Quadratic Values*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $T \in \mathcal{L}(\mathcal{V})$.
Then T is self-adjoint if and only if $\langle T\mathbf{v}, \mathbf{v}\rangle \in \mathbb{R}$ for every $\mathbf{v} \in \mathcal{V}$.

Proof. For each $\mathbf{v}$, $\langle (T - T^*)\mathbf{v}, \mathbf{v}\rangle = \langle T\mathbf{v}, \mathbf{v}\rangle - \langle T^*\mathbf{v}, \mathbf{v}\rangle = \langle T\mathbf{v}, \mathbf{v}\rangle - \overline{\langle T\mathbf{v}, \mathbf{v}\rangle}$.
If all values $\langle T\mathbf{v}, \mathbf{v}\rangle$ are real, this expression is zero for every $\mathbf{v}$, and Lemma 2.4.15 gives $T - T^* = 0$.
Conversely, if $T = T^*$, then $\langle T\mathbf{v}, \mathbf{v}\rangle = \langle \mathbf{v}, T\mathbf{v}\rangle = \overline{\langle T\mathbf{v}, \mathbf{v}\rangle}$, so $\langle T\mathbf{v}, \mathbf{v}\rangle$ is real.
Structural Commentary: The imaginary part of $\langle T\mathbf{v}, \mathbf{v}\rangle$ is exactly the diagonal trace of $T - T^*$.
Key Insight: Over $\mathbb{C}$, real quadratic values force the operator to equal its adjoint. □

Proposition 2.4.17. *Zero Quadratic Form for Self-Adjoint Operators*
Let T be a self-adjoint operator on a finite-dimensional inner-product space $\mathcal{V}$.
If $\langle T\mathbf{v}, \mathbf{v}\rangle = 0$ for every $\mathbf{v} \in \mathcal{V}$, then $T = 0$.

Proof. If $\mathbb{F} = \mathbb{C}$, this is Lemma 2.4.15. Assume $\mathbb{F} = \mathbb{R}$.
For $\mathbf{u}, \mathbf{w} \in \mathcal{V}$, self-adjointness gives $\langle T\mathbf{w}, \mathbf{u}\rangle = \langle \mathbf{w}, T\mathbf{u}\rangle = \langle T\mathbf{u}, \mathbf{w}\rangle$, so the real polarization identity gives

$$\langle T\mathbf{u}, \mathbf{w}\rangle = \frac{\langle T(\mathbf{u}+\mathbf{w}), \mathbf{u}+\mathbf{w}\rangle - \langle T(\mathbf{u}-\mathbf{w}), \mathbf{u}-\mathbf{w}\rangle}{4} = 0.$$

Taking $\mathbf{w} = T\mathbf{u}$ gives $T\mathbf{u} = \mathbf{0}$ for every $\mathbf{u}$.
Structural Commentary: Self-adjointness supplies the missing symmetry needed for real polarization.
Key Insight: A self-adjoint operator is determined by its quadratic form. □

Definition 2.4.18. *Orthogonal Projection Matrix*
A square matrix $P \in \mathrm{Mat}(n, n, \mathbb{F})$ is an *orthogonal projection matrix* if $P^2 = P$ and $P^* = P$.

Theorem 2.4.19. *Orthogonal Projections as Self-Adjoint Idempotents*
Let $\mathcal{V}$ be finite-dimensional and let $P \in \mathcal{L}(\mathcal{V})$.
Then P is the orthogonal projection onto $\mathrm{range}(P)$ *if and only if $P^2 = P$ and $P^* = P$.*

Proof.
(⇒) If $P = P_{\mathcal{U}}$ for $\mathcal{U} = \mathrm{range}(P)$, then Proposition 2.2.12 gives $P^2 = P$. For $\mathbf{v} = \mathbf{u} + \mathbf{w}$ and $\mathbf{x} = \mathbf{a} + \mathbf{b}$ with $\mathbf{u}, \mathbf{a} \in \mathcal{U}$ and $\mathbf{w}, \mathbf{b} \in \mathcal{U}^{\perp}$, $\langle P\mathbf{v}, \mathbf{x}\rangle = \langle \mathbf{u}, \mathbf{a}+\mathbf{b}\rangle = \langle \mathbf{u}, \mathbf{a}\rangle = \langle \mathbf{u}+\mathbf{w}, \mathbf{a}\rangle = \langle \mathbf{v}, P\mathbf{x}\rangle$, so $P^* = P$.
(⇐) Suppose $P^2 = P$ and $P^* = P$. For each $\mathbf{v}$, $\mathbf{v} = P\mathbf{v} + (\mathbf{v} - P\mathbf{v})$ with $P\mathbf{v} \in \mathrm{range}(P)$ and $P(\mathbf{v} - P\mathbf{v}) = \mathbf{0}$. Thus the residual lies in $\ker(P)$. Since $P = P^*$, Proposition 2.4.11 gives $\ker(P) = (\mathrm{range}(P))^{\perp}$. Hence every $\mathbf{v}$ has the orthogonal decomposition $\mathbf{v} = P\mathbf{v} + (\mathbf{v} - P\mathbf{v})$ with first part in $\mathrm{range}(P)$, so P is exactly the orthogonal projection onto $\mathrm{range}(P)$.

Structural Commentary: Idempotence splits by range and kernel; self-adjointness makes the split orthogonal.
Key Insight: An idempotent is orthogonal exactly when $\ker(P) = (\mathrm{range}(P))^{\perp}$. □

Theorem 2.4.20. *Unitary Coordinate Form of an Orthogonal Projection*
Let $P \in \mathrm{Mat}(n, n, \mathbb{F})$ be an orthogonal projection matrix and let $r = \mathrm{rank}(P)$. Then P is unitarily similar to $I_r \oplus 0_{n-r}$. Equivalently, there is a unitary matrix U such that

$$U^*PU = \begin{bmatrix} I_r & 0 \\ 0 & 0_{n-r} \end{bmatrix}.$$

Proof. Let $\mathcal{U} = \mathrm{col}(P)$. Identifying P with its induced operator on $\mathbb{F}^n$, Theorem 2.4.19 says P is the orthogonal projection onto $\mathcal{U}$, so P is the identity on $\mathcal{U}$ and zero on $\mathcal{U}^{\perp}$. Choose an orthonormal basis $\mathbf{u}_1, \ldots, \mathbf{u}_r$ of $\mathcal{U}$ and extend it to an orthonormal basis $\mathbf{u}_1, \ldots, \mathbf{u}_n$ of $\mathbb{F}^n$. Let $U = [\mathbf{u}_1 \ \cdots \ \mathbf{u}_n]$.
Then U is unitary. For $j \le r$, $P\mathbf{u}_j = \mathbf{u}_j$, and for $j > r$, $\mathbf{u}_j \in \mathcal{U}^{\perp}$ so $P\mathbf{u}_j = \mathbf{0}$. Therefore the jth column of U^*PU is the coordinate column of $P\mathbf{u}_j$ in this orthonormal basis, giving the displayed block diagonal form.
Structural Commentary: Orthogonal projection splits the space into range and orthogonal kernel.
Key Insight: In a basis adapted to $\mathrm{col}(P) \oplus \mathrm{col}(P)^{\perp}$, projection keeps the first coordinates and kills the rest. □

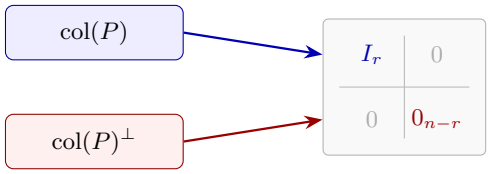


Figure 36: Adapted basis: $U^*PU = I_r \oplus 0_{n-r}$, identity on $\mathrm{col}(P)$ and zero on $\mathrm{col}(P)^{\perp}$.

Definition 2.4.21. *Isometry and Isometric Isomorphism*
Let $\mathcal{V}$ and $\mathcal{W}$ be inner-product spaces and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.

(i) T is an *isometry* if $\|T\mathbf{v}\| = \|\mathbf{v}\|$ for all $\mathbf{v} \in \mathcal{V}$.
(ii) T is an *isometric isomorphism* if T is bijective and isometric.
The spaces $\mathcal{V}$ and $\mathcal{W}$ are *isometrically isomorphic* if such a T exists.

Proposition 2.4.22. *Basic Properties of Isometries*
Let $\mathcal{V}$ be an inner-product space and let $S, T \in \mathcal{L}(\mathcal{V})$ be isometries.

i. ST is an isometry;
ii. T is injective;
iii. if $\mathcal{V}$ is finite-dimensional, then T is invertible and T^{-1} is an isometry.

Proof.

(i) For every $\mathbf{v}$, $\|ST\mathbf{v}\| = \|T\mathbf{v}\| = \|\mathbf{v}\|$.
(ii) If $T\mathbf{v} = \mathbf{0}$, then $\|\mathbf{v}\| = \|T\mathbf{v}\| = 0$, so $\mathbf{v} = \mathbf{0}$. Thus $\ker(T) = \{\mathbf{0}\}$.
(iii) In finite dimensions, injectivity implies invertibility by Theorem 1.3.49.
If $\mathbf{u} = T\mathbf{v}$, then $\left\|T^{-1}\mathbf{u}\right\| = \|\mathbf{v}\| = \|T\mathbf{v}\| = \|\mathbf{u}\|$.

Structural Commentary: Isometries preserve length, so they cannot collapse nonzero vectors.
Key Insight: In finite dimensions, length preservation already forces an invertible geometry-preserving map. □

Example 2.4.23. *An Isometry Need Not Be Onto*
Define $T : \ell^2(\mathbb{F}) \to \ell^2(\mathbb{F})$ by $T(x_1, x_2, \ldots) = (0, x_1, x_2, \ldots)$.
For $\mathbf{x} = (x_j)_{j\geq 1}$, $\|T\mathbf{x}\|_2^2 = \sum_{j\geq 1} |x_j|^2 = \|\mathbf{x}\|_2^2$, so T is an isometry.
Since every vector in $\operatorname{range}(T)$ has first coordinate 0, $(1, 0, 0, \ldots) \notin \operatorname{range}(T)$; hence T is not surjective.
Structural Commentary: Infinite-dimensional isometries can preserve all lengths while missing whole directions.
Key Insight: Injective does not imply surjective outside finite dimensions.

Theorem 2.4.24. *Isometries Preserve Inner Products*
Let $\mathcal{V}$ and $\mathcal{W}$ be inner-product spaces and let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$.
Then T is an isometry if and only if $\langle T\mathbf{u}, T\mathbf{v}\rangle_{\mathcal{W}} = \langle \mathbf{u}, \mathbf{v}\rangle_{\mathcal{V}}$ for every $\mathbf{u}, \mathbf{v} \in \mathcal{V}$.

Proof.

(⇒) Suppose T is an isometry. Then every norm appearing in the polarization identities of Theorem 2.1.23 is unchanged after applying T. Therefore $\langle T\mathbf{u}, T\mathbf{v}\rangle_{\mathcal{W}} = \langle \mathbf{u}, \mathbf{v}\rangle_{\mathcal{V}}$.
(⇐) Suppose T preserves inner products. Then $\|T\mathbf{u}\|^2 = \langle T\mathbf{u}, T\mathbf{u}\rangle_{\mathcal{W}} = \langle \mathbf{u}, \mathbf{u}\rangle_{\mathcal{V}} = \|\mathbf{u}\|^2$, so T is an isometry.

Structural Commentary: Polarization converts preservation of norms into preservation of inner products.
Key Insight: Length preservation for all vectors also preserves angle information. □

Remark 2.4.25. *Isometry Conventions*

(i) For linear maps between inner-product spaces, norm preservation and inner-product preservation are equivalent by Theorem 2.4.24; hence either convention defines the same isometries.
(ii) For general bilinear forms, an isometry is often required to be bijective and form-preserving; if the form is singular, preservation alone need not imply injectivity.

Theorem 2.4.26. *Isometries and the Adjoint Product*
Let $\mathcal{V}$ be finite-dimensional and let $T \in \mathcal{L}(\mathcal{V})$. Then T is an isometry if and only if $T^*T = I_{\mathcal{V}}$.

Proof.

(⇒) Suppose T is an isometry. By Theorem 2.4.24, $\langle T^*T\mathbf{u}, \mathbf{v}\rangle = \langle T\mathbf{u}, T\mathbf{v}\rangle = \langle \mathbf{u}, \mathbf{v}\rangle$ for all $\mathbf{u}, \mathbf{v}$.
Thus $\langle (T^*T - I_{\mathcal{V}})\mathbf{u}, \mathbf{v}\rangle = 0$ for every $\mathbf{v}$, so $T^*T\mathbf{u} = \mathbf{u}$ for every $\mathbf{u}$.
(⇐) Suppose $T^*T = I_{\mathcal{V}}$. Then $\|T\mathbf{u}\|^2 = \langle T\mathbf{u}, T\mathbf{u}\rangle = \langle T^*T\mathbf{u}, \mathbf{u}\rangle = \|\mathbf{u}\|^2$ for every $\mathbf{u}$, so T is an isometry.

Structural Commentary: The operator T^*T records inner products after applying T.
Key Insight: T is an isometry exactly when this Gram operator is the identity. □

Definition 2.4.27. *Unitary and Orthogonal*
An operator $U \in \mathcal{L}(\mathcal{V})$ on a finite-dimensional inner-product space is *unitary* if $U^*U = UU^* = I_{\mathcal{V}}$.
A matrix $U \in \operatorname{Mat}(n, n, \mathbb{F})$ is *unitary* if $U^*U = I_n$. If $\mathbb{F} = \mathbb{R}$, a unitary matrix is called *orthogonal*.

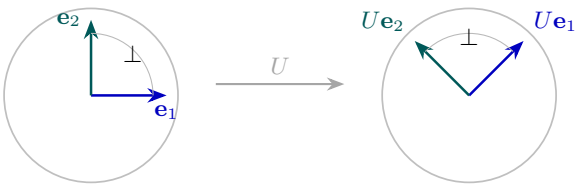


Figure 37: Unitary invariance: $\|U\mathbf{x}\| = \|\mathbf{x}\|$ and $\langle U\mathbf{x}, U\mathbf{y}\rangle = \langle \mathbf{x}, \mathbf{y}\rangle$.

Theorem 2.4.28. *Equivalent Conditions for a Unitary Matrix*
Let $U = [\mathbf{u}_1 \ \cdots \ \mathbf{u}_n] \in \operatorname{Mat}(n, n, \mathbb{F})$. The following are equivalent:

i. U is unitary, i.e. $U^*U = I_n$;
ii. the columns $\mathbf{u}_1, \ldots, \mathbf{u}_n$ are orthonormal in $\mathbb{F}^n$;

iii. $UU^* = I_n$;
iv. U^* *is unitary;*
v. *the rows of* U *are orthonormal;*
vi. U *is invertible and* $U^{-1} = U^*$;
vii. $\|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2$ *for every* $\mathbf{x} \in \mathbb{F}^n$;
viii. $\langle U\mathbf{x}, U\mathbf{y}\rangle_{\mathbb{F}^n} = \langle \mathbf{x}, \mathbf{y}\rangle_{\mathbb{F}^n}$ *for every* $\mathbf{x}, \mathbf{y} \in \mathbb{F}^n$.

Proof.

(*i*) $\Rightarrow$ (*ii*). Since $(U^*U)_{ij} = \mathbf{u}_i^*\mathbf{u}_j = \langle \mathbf{u}_j, \mathbf{u}_i\rangle_{\mathbb{F}^n}$, equation $U^*U = I_n$ states that columns of U are orthonormal.
(*ii*) $\Rightarrow$ (*iii*). Orthonormal columns give $U^*U = I_n$. A left inverse of a square matrix is also a right inverse by Theorem 1.3.87, so $UU^* = I_n$.
(*iii*) $\Rightarrow$ (*iv*). The equation $UU^* = I_n$ is $(U^*)^*U^* = I_n$, so U^* is unitary.
(*iv*) $\Rightarrow$ (*v*). Apply (*i*) $\Rightarrow$ (*ii*) to U^*; the columns of U^* are the conjugates of the rows of U.
(*v*) $\Rightarrow$ (*vi*). Orthonormal rows give $UU^* = I_n$. Thus U^* is the inverse of U, so $U^{-1} = U^*$.
(*vi*) $\Rightarrow$ (*vii*). If $U^{-1} = U^*$, then $U^*U = I_n$. Therefore $\|U\mathbf{x}\|_2^2 = \langle U\mathbf{x}, U\mathbf{x}\rangle_{\mathbb{F}^n} = \langle U^*U\mathbf{x}, \mathbf{x}\rangle_{\mathbb{F}^n} = \|\mathbf{x}\|_2^2$ for every $\mathbf{x} \in \mathbb{F}^n$. Since both norms are nonnegative, $\|U\mathbf{x}\|_2 = \|\mathbf{x}\|_2$.
(*vii*) $\Rightarrow$ (*viii*). Apply Theorem 2.4.24 to $\mathbf{x} \mapsto U\mathbf{x}$.
(*viii*) $\Rightarrow$ (*i*). For every $\mathbf{x}, \mathbf{y} \in \mathbb{F}^n$, $\langle (U^*U - I_n)\mathbf{x}, \mathbf{y}\rangle_{\mathbb{F}^n} = \langle U\mathbf{x}, U\mathbf{y}\rangle_{\mathbb{F}^n} - \langle \mathbf{x}, \mathbf{y}\rangle_{\mathbb{F}^n} = 0$. Taking $\mathbf{y} = (U^*U - I_n)\mathbf{x}$ shows $(U^*U - I_n)\mathbf{x} = \mathbf{0}$ for every $\mathbf{x}$, so $U^*U = I_n$.

Structural Commentary: A unitary matrix is an isometry written in orthonormal coordinates.
Key Insight: Columns, rows, inverses, norms, and inner products are different faces of the same condition. □

Proposition 2.4.29. *Algebra of Unitary Matrices*
Let $U, V \in \mathrm{Mat}(n, n, \mathbb{F})$ *and* $W \in \mathrm{Mat}(m, m, \mathbb{F})$, *where* $V \oplus W$ *denotes the block diagonal matrix with blocks* V *and* W.

i. *If* U *and* V *are unitary, then* UV *is unitary.*
ii. $V \oplus W$ *is unitary if and only if* V *and* W *are unitary.*

Proof.

(i) Compute $(UV)^*(UV) = V^*U^*UV = V^*V = I_n$.
(ii) Since $(V \oplus W)^*(V \oplus W) = V^*V \oplus W^*W$, this product equals I_{n+m} exactly when $V^*V = I_n$ and $W^*W = I_m$.

Structural Commentary: Unitary structure is stable under composition and blockwise orthogonal sums.
Key Insight: The equation $U^*U = I$ makes every verification a short adjoint calculation. □

Theorem 2.4.30. *Extending Isometric Columns to a Unitary Matrix*
If $X \in \mathrm{Mat}(m, n, \mathbb{F})$ *and* $X^*X = I_n$, *then there exists* $X' \in \mathrm{Mat}(m, m-n, \mathbb{F})$ *such that* $[X\ X']$ *is unitary.*

Proof. Write $X = [\mathbf{x}_1\ \cdots\ \mathbf{x}_n]$. Since $(X^*X)_{ij} = \langle \mathbf{x}_j, \mathbf{x}_i\rangle_{\mathbb{F}^m}$, the condition $X^*X = I_n$ says $\mathbf{x}_1, \ldots, \mathbf{x}_n$ are orthonormal in $\mathbb{F}^m$. By Corollary 2.1.43, extend this list to an orthonormal basis $\mathbf{x}_1, \ldots, \mathbf{x}_m$ of $\mathbb{F}^m$.
Let $X' = [\mathbf{x}_{n+1}\ \cdots\ \mathbf{x}_m]$. Then the columns of $[X\ X']$ are orthonormal, so $[X\ X']$ is unitary.
Structural Commentary: Isometric columns are a partial orthonormal basis.
Key Insight: Complete the columns by extending the orthonormal list. □

Definition 2.4.31. *Unitary and Permutation Similarity*
Square matrices $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$ are *unitarily similar* if $A = UBU^*$ for some unitary U. They are *orthogonally similar* if $\mathbb{F} = \mathbb{R}$ and U can be chosen real orthogonal.
A *permutation matrix* is a square matrix with exactly one entry equal to 1 in each row and each column, and all other entries equal to 0. Matrices are *permutation similar* if $A = PBP^*$ for some permutation matrix P.

Proposition 2.4.32. *Permutation Matrices Are Orthogonal*
If $P = [\mathbf{e}_{\sigma(1)}\ \cdots\ \mathbf{e}_{\sigma(n)}]$ *is a permutation matrix, then* $P^*P = PP^* = I_n$. *Hence* P *is real orthogonal.*

Proof. The columns of P are a rearrangement of the standard orthonormal basis of $\mathbb{R}^n$, hence they are orthonormal. Theorem 2.4.28 gives $P^*P = PP^* = I_n$.
Structural Commentary: A permutation matrix only reorders the standard basis.
Key Insight: Reordering orthonormal vectors preserves orthonormality. □

Definition 2.4.33. *Householder Matrix*
Let $\mathbf{w} \in \mathbb{F}^n$ be nonzero and set $\mathbf{u} = \mathbf{w}/\|\mathbf{w}\|_2$.
The *Householder matrix* associated to $\mathbf{w}$ is $U_{\mathbf{w}} := I_n - 2P_{\mathbf{u}} = I_n - 2\frac{\mathbf{w}\mathbf{w}^*}{\|\mathbf{w}\|_2^2}$.
It acts by $\mathbf{x} \mapsto U_{\mathbf{w}}\mathbf{x} = \mathbf{x} - 2\langle \mathbf{x}, \mathbf{u}\rangle_{\mathbb{F}^n}\mathbf{u}$.

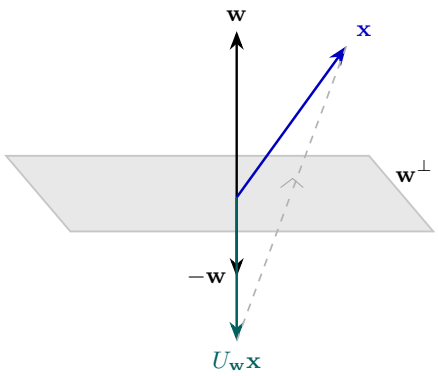


Figure 38: A Householder matrix fixes $\mathbf{w}^\perp$ and flips the span($\mathbf{w}$) component.

Theorem 2.4.34. *Householder Matrices Are Unitary Reflections*
Every Householder matrix $U_\mathbf{w}$ is unitary, Hermitian, and involutive: $U_\mathbf{w}^* U_\mathbf{w} = I_n$, $U_\mathbf{w}^* = U_\mathbf{w}$, $U_\mathbf{w}^2 = I_n$.
If $\mathbf{w} \in \mathbb{R}^n$, then $U_\mathbf{w}$ is real orthogonal and symmetric.

Proof. Let $\mathbf{u} = \mathbf{w}/\|\mathbf{w}\|_2$. By Proposition 2.2.14, $P_\mathbf{u}^* = P_\mathbf{u}$ and $P_\mathbf{u}^2 = P_\mathbf{u}$.
Therefore $U_\mathbf{w}^* = I_n - 2P_\mathbf{u} = U_\mathbf{w}$, $U_\mathbf{w}^2 = (I_n - 2P_\mathbf{u})^2 = I_n - 4P_\mathbf{u} + 4P_\mathbf{u}^2 = I_n$.
Thus $U_\mathbf{w}^* U_\mathbf{w} = U_\mathbf{w}^2 = I_n$, so $U_\mathbf{w}$ is unitary. The real case is the same statement with conjugation removed.
Structural Commentary: A Householder matrix reflects across an orthogonal hyperplane.
Key Insight: The identity $P_\mathbf{u}^2 = P_\mathbf{u}$ makes $(I - 2P_\mathbf{u})^2 = I$. □

Theorem 2.4.35. *Householder Map Between Equal-Norm Vectors*
Let $\mathbf{x}, \mathbf{y} \in \mathbb{F}^n$ satisfy $\|\mathbf{x}\|_2 = \|\mathbf{y}\|_2 \neq 0$. Define $\sigma = \begin{cases} 1, & \langle \mathbf{x}, \mathbf{y} \rangle_{\mathbb{F}^n} = 0, \\ -\dfrac{\overline{\langle \mathbf{x}, \mathbf{y} \rangle_{\mathbb{F}^n}}}{|\langle \mathbf{x}, \mathbf{y} \rangle_{\mathbb{F}^n}|}, & \langle \mathbf{x}, \mathbf{y} \rangle_{\mathbb{F}^n} \neq 0, \end{cases}$ where $\mathbf{w} = \mathbf{y} - \sigma\mathbf{x}$.
Then $\mathbf{w} \neq \mathbf{0}$, $|\sigma| = 1$, and $\sigma U_\mathbf{w} \mathbf{x} = \mathbf{y}$.

Proof. Clearly $|\sigma| = 1$. Let $z = \langle \mathbf{x}, \mathbf{y} \rangle$ and $r = \|\mathbf{x}\|_2 = \|\mathbf{y}\|_2$. By construction, $\sigma z = -|z|$, with the convention that this equals 0 when $z = 0$. Hence $\|\mathbf{w}\|_2^2 = \|\mathbf{y} - \sigma\mathbf{x}\|_2^2 = 2r^2 - \sigma z - \overline{\sigma z} = 2\left(r^2 + |z|\right) > 0$, so $\mathbf{w} \neq \mathbf{0}$.
Since $\mathbf{y} = \mathbf{w} + \sigma\mathbf{x}$, $\sigma U_\mathbf{w} \mathbf{x} = \sigma\mathbf{x} - 2\sigma \frac{\langle \mathbf{x}, \mathbf{w} \rangle}{\|\mathbf{w}\|_2^2} \mathbf{w} = \mathbf{y} - \left(1 + 2\sigma \frac{\langle \mathbf{x}, \mathbf{w} \rangle}{\|\mathbf{w}\|_2^2}\right) \mathbf{w}$.
But $\langle \mathbf{x}, \mathbf{w} \rangle = z - \overline{\sigma} r^2$, so $2\sigma \langle \mathbf{x}, \mathbf{w} \rangle = 2\sigma z - 2r^2 = -2(r^2 + |z|) = -\|\mathbf{w}\|_2^2$.
Coefficient of $\mathbf{w}$ is hence zero, proving $\sigma U_\mathbf{w} \mathbf{x} = \mathbf{y}$.
Structural Commentary: The scalar σ chooses the phase so that $\mathbf{y} - \sigma\mathbf{x}$ gives a stable reflection direction.
Key Insight: Householder reflections connect equal-norm vectors up to a unit scalar. □

Corollary 2.4.36. *Householder Unitary Annihilator*
Let $\mathbf{x} = [x_1 \ \cdots \ x_n]^t \in \mathbb{F}^n$ be nonzero. Define $\sigma = \begin{cases} 1, & x_1 = 0, \\ -\dfrac{\overline{x_1}}{|x_1|}, & x_1 \neq 0, \end{cases}$ where $\mathbf{w} = \|\mathbf{x}\|_2 \mathbf{e}_1 - \sigma\mathbf{x}$.
Then $\sigma U_\mathbf{w} \mathbf{x} = \|\mathbf{x}\|_2 \mathbf{e}_1$. If $\mathbf{x}$ is a unit vector, then $\overline{\sigma} U_\mathbf{w}$ is unitary and has $\mathbf{x}$ as its first column.

Proof. Apply Theorem 2.4.35 with $\mathbf{y} = \|\mathbf{x}\|_2 \mathbf{e}_1$, noting that $\langle \mathbf{x}, \mathbf{y} \rangle = x_1 \|\mathbf{x}\|_2$. If $\|\mathbf{x}\|_2 = 1$, then $(\sigma U_\mathbf{w})\mathbf{x} = \mathbf{e}_1$, so $(\sigma U_\mathbf{w})^* \mathbf{e}_1 = \mathbf{x}$. Since $(\sigma U_\mathbf{w})^* = \overline{\sigma} U_\mathbf{w}$, this says the first column of $\overline{\sigma} U_\mathbf{w}$ is $\mathbf{x}$.
Structural Commentary: The reflector concentrates all mass of $\mathbf{x}$ into the first coordinate.
Key Insight: Householder matrices produce explicit unitary matrices with a prescribed first column. □

2.5 Metric Vector Spaces and Bilinear Forms

Metric vector spaces keep the algebraic part of inner-product geometry and drop positivity. Throughout this subsection $\mathbb{F}$ is an arbitrary field and vector spaces are finite-dimensional unless explicitly stated otherwise. The bilinear form is written $\langle \mathbf{x}, \mathbf{y} \rangle$ when it is fixed.

Definition 2.5.1. *Bilinear Pairings and Their Radicals*
Let $\mathcal{U}, \mathcal{W}$ be vector spaces over $\mathbb{F}$.

(i) A *bilinear pairing* is a map $B : \mathcal{U} \times \mathcal{W} \to \mathbb{F}$ linear in each variable.
(ii) The *left radical* is $\operatorname{rad}_L(B) := \{\mathbf{u} \in \mathcal{U} : B(\mathbf{u}, \mathbf{w}) = 0 \ \forall \mathbf{w} \in \mathcal{W}\}$.
The *right radical* is $\operatorname{rad}_R(B) := \{\mathbf{w} \in \mathcal{W} : B(\mathbf{u}, \mathbf{w}) = 0 \ \forall \mathbf{u} \in \mathcal{U}\}$.

(iii) The *right orthogonal* of $S \subseteq \mathcal{U}$ is $S^{\perp_R} := \{\mathbf{w} \in \mathcal{W} : B(\mathbf{s}, \mathbf{w}) = 0 \ \forall \mathbf{s} \in S\}$.
The *left orthogonal* of $T \subseteq \mathcal{W}$ is ${}^{\perp_L}T := \{\mathbf{u} \in \mathcal{U} : B(\mathbf{u}, \mathbf{t}) = 0 \ \forall \mathbf{t} \in T\}$.
(iv) B is *left nonsingular* if $\operatorname{rad}_L(B) = \{\mathbf{0}\}$, *right nonsingular* if $\operatorname{rad}_R(B) = \{\mathbf{0}\}$, and *perfect* if both hold.

Proposition 2.5.2. *Matrices of Bilinear Pairings*
Let $\alpha = (\mathbf{u}_1, \dots, \mathbf{u}_m)$, $\beta = (\mathbf{w}_1, \dots, \mathbf{w}_n)$ be bases of $\mathcal{U}$, $\mathcal{W}$. Let $M = M_{\alpha,\beta} := [B(\mathbf{u}_i, \mathbf{w}_j)]_{i,j} \in \operatorname{Mat}(m, n, \mathbb{F})$.

(i) $B(\mathbf{u}, \mathbf{w}) = [\mathbf{u}]_\alpha^t M [\mathbf{w}]_\beta$.
(ii) $\{[\mathbf{u}]_\alpha : \mathbf{u} \in \operatorname{rad}_L(B)\} = \ker(M^t)$ *and* $\{[\mathbf{w}]_\beta : \mathbf{w} \in \operatorname{rad}_R(B)\} = \ker(M)$.
(iii) $\dim \operatorname{rad}_L(B) = m - \operatorname{rank}(M)$ *and* $\dim \operatorname{rad}_R(B) = n - \operatorname{rank}(M)$.
(iv) If B is perfect, then $m = n$; when $m = n$, B is perfect if and only if M is invertible.

Proof.

(i) If $\mathbf{u} = \sum_i x_i \mathbf{u}_i$ and $\mathbf{w} = \sum_j y_j \mathbf{w}_j$, then $B(\mathbf{u}, \mathbf{w}) = \sum_{i,j} x_i M_{ij} y_j = [\mathbf{u}]_\alpha^t M [\mathbf{w}]_\beta$.
(ii) For $\mathbf{x} = [\mathbf{u}]_\alpha$ and $\mathbf{y} = [\mathbf{w}]_\beta$, $\mathbf{u} \in \operatorname{rad}_L(B) \iff \mathbf{x}^t M \mathbf{y} = 0 \ \forall \mathbf{y} \iff M^t \mathbf{x} = \mathbf{0}$.
Similarly, $\mathbf{w} \in \operatorname{rad}_R(B) \iff M\mathbf{y} = \mathbf{0}$.
(iii) By (ii) and rank-nullity, $\dim \operatorname{rad}_L(B) = \dim \ker(M^t) = m - \operatorname{rank}(M^t) = m - \operatorname{rank}(M)$,
$\dim \operatorname{rad}_R(B) = \dim \ker(M) = n - \operatorname{rank}(M)$.
(iv) By (iii), perfection is equivalent to $m - \operatorname{rank}(M) = 0$ and $n - \operatorname{rank}(M) = 0$. Hence $m = n$; when $m = n$, this is equivalent to M being invertible.

Structural Commentary: A pairing matrix has one radical for each input slot.
Key Insight: Nonsingularity of a pairing is ordinary full-rank behavior of its representing matrix. □

Proposition 2.5.3. *Two-Sided Basis Change for Pairings*
With notation as above, let α' and β' be bases, and set $P = C_{\alpha \leftarrow \alpha'}$, $Q = C_{\beta \leftarrow \beta'}$. Then $M_{\alpha',\beta'} = P^t M_{\alpha,\beta} Q$. For square forms this is congruence; for rectangular pairings it is two-sided equivalence.

Proof. By Proposition 2.5.2, $[\mathbf{u}]_{\alpha'}^t M_{\alpha',\beta'} [\mathbf{w}]_{\beta'} = B(\mathbf{u}, \mathbf{w}) = (P[\mathbf{u}]_{\alpha'})^t M_{\alpha,\beta} (Q[\mathbf{w}]_{\beta'}) = [\mathbf{u}]_{\alpha'}^t P^t M_{\alpha,\beta} Q [\mathbf{w}]_{\beta'}$.
Thus $M_{\alpha',\beta'} = P^t M_{\alpha,\beta} Q$.
Structural Commentary: Pairing matrices change independently in each input slot.
Key Insight: The number of input slots determines how basis changes act on the matrix. □

Perspective 2.5.4. *Square Pairings*
Henceforth $B : \mathcal{V} \times \mathcal{V} \to \mathbb{F}$ is a square pairing, usually symmetric, skew-symmetric, or alternating.
In these cases $\operatorname{rad}_L(B) = \operatorname{rad}_R(B)$, so one radical and one orthogonality notation suffice.

Definition 2.5.5. *Symmetric, Skew-Symmetric, and Alternating Forms*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$ and let $B : \mathcal{V} \times \mathcal{V} \to \mathbb{F}$ be bilinear.

(i) B is a *bilinear form* on $\mathcal{V}$; when fixed, write $B(\mathbf{x}, \mathbf{y}) = \langle \mathbf{x}, \mathbf{y} \rangle$.
(ii) $\operatorname{Sym}(B)$ means $B(\mathbf{x}, \mathbf{y}) = B(\mathbf{y}, \mathbf{x})$ for all $\mathbf{x}, \mathbf{y} \in \mathcal{V}$; then B is *symmetric*.
(iii) $\operatorname{Skew}(B)$ means $B(\mathbf{x}, \mathbf{y}) = -B(\mathbf{y}, \mathbf{x})$ for all $\mathbf{x}, \mathbf{y} \in \mathcal{V}$; then B is *skew-symmetric*.
(iv) $\operatorname{Alt}(B)$ means $B(\mathbf{x}, \mathbf{x}) = 0$ for all $\mathbf{x} \in \mathcal{V}$; then B is *alternating*.
(v) A *metric vector space* is a pair $(\mathcal{V}, B)$ with $\operatorname{Sym}(B)$, $\operatorname{Skew}(B)$, or $\operatorname{Alt}(B)$.
The symmetric case is an *orthogonal geometry*; the alternating case is a *symplectic geometry*.

Proposition 2.5.6. *Characteristic Controls the Form Types*
Let $B(\mathbf{x}, \mathbf{y}) = \langle \mathbf{x}, \mathbf{y} \rangle$ be a bilinear form on $\mathcal{V}$.

(i) $\operatorname{Alt}(B) \Rightarrow \operatorname{Skew}(B)$.
(ii) If $\operatorname{char}(\mathbb{F}) \neq 2$, then $\operatorname{Alt}(B) \iff \operatorname{Skew}(B)$, and $\operatorname{Sym}(B) \wedge \operatorname{Alt}(B) \Rightarrow B = 0$.
(iii) If $\operatorname{char}(\mathbb{F}) = 2$, then $\operatorname{Sym}(B) \iff \operatorname{Skew}(B)$ and $\operatorname{Alt}(B) \Rightarrow \operatorname{Sym}(B)$.

Proof.

(i) If $\operatorname{Alt}(B)$, then $0 = \langle \mathbf{x} + \mathbf{y}, \mathbf{x} + \mathbf{y} \rangle = \langle \mathbf{x}, \mathbf{y} \rangle + \langle \mathbf{y}, \mathbf{x} \rangle$, so $\langle \mathbf{x}, \mathbf{y} \rangle = -\langle \mathbf{y}, \mathbf{x} \rangle$.
(ii) Assume $\operatorname{char}(\mathbb{F}) \neq 2$. By (i), $\operatorname{Alt}(B) \Rightarrow \operatorname{Skew}(B)$. Conversely, $\operatorname{Skew}(B)$ gives $\langle \mathbf{x}, \mathbf{x} \rangle = -\langle \mathbf{x}, \mathbf{x} \rangle$, hence $2\langle \mathbf{x}, \mathbf{x} \rangle = 0$ and $\operatorname{Alt}(B)$. If also $\operatorname{Sym}(B)$ and $\operatorname{Alt}(B)$, then $\langle \mathbf{x}, \mathbf{y} \rangle = -\langle \mathbf{x}, \mathbf{y} \rangle$, so $\langle \mathbf{x}, \mathbf{y} \rangle = 0$ for all $\mathbf{x}, \mathbf{y}$.
(iii) If $\operatorname{char}(\mathbb{F}) = 2$, then $-1 = 1$, so $\operatorname{Sym}(B) \iff \operatorname{Skew}(B)$.
Moreover $\operatorname{Alt}(B) \Rightarrow \operatorname{Skew}(B)$ by (i), hence $\operatorname{Alt}(B) \Rightarrow \operatorname{Sym}(B)$.

Structural Commentary: Skew-symmetry collapses by characteristic.
Key Insight: The equation $\langle \mathbf{x} + \mathbf{y}, \mathbf{x} + \mathbf{y} \rangle = 0$ polarizes alternation. □

Definition 2.5.7. *Minkowski Space*
Minkowski space is the orthogonal geometry $\mathbb{R}^{3,1} := \mathbb{R}^4$ with standard basis $(\mathbf{e}_1, \dots, \mathbf{e}_4)$ and bilinear form $\langle \mathbf{x}, \mathbf{y} \rangle = x_1y_1 + x_2y_2 + x_3y_3 - x_4y_4$.
Equivalently, its Gram matrix is $\operatorname{diag}(1, 1, 1, -1)$. It is nonsingular but not positive definite.

Example 2.5.8. *Null Vectors in Minkowski Space*
In $\mathbb{R}^{3,1}$, set $q(\mathbf{x}) := \langle \mathbf{x}, \mathbf{x} \rangle = x_1^2 + x_2^2 + x_3^2 - x_4^2$.
Null cone is $\mathcal{N} := \{\mathbf{x} \neq \mathbf{0} : q(\mathbf{x}) = 0\} = \{\mathbf{x} \neq \mathbf{0} : x_1^2 + x_2^2 + x_3^2 = x_4^2\}$. Every $\mathbf{x} \in \mathcal{N}$ is isotropic; i.e., $\mathbf{e}_1 + \mathbf{e}_4 \in \mathcal{N}$.
As $\operatorname{diag}(1, 1, 1, -1)$, I_4 have different real inertias, $(3, 1, 0)$ and $(4, 0, 0)$, $\mathbb{R}^{3,1}$ is not isometric to Euclidean $\mathbb{R}^4$.

Definition 2.5.9. *Matrix, Rank, and Discriminant of a Form*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$ be an ordered basis of a metric vector space $\mathcal{V}$.

(i) The *matrix of the form* in β is $M_\beta := [\langle \mathbf{v}_i, \mathbf{v}_j \rangle]_{i,j}$.
(ii) The *rank* of the form is $\operatorname{rank}(M_\beta)$, independent of β.
(iii) The *discriminant set* is $\Delta(\mathcal{V}) := \{r^2 \det(M_\beta) : r \in \mathbb{F}^\times\}$.
(iv) If the form is nonsingular, then $\Delta(\mathcal{V})$ is the square class of $\det(M_\beta)$ in $\mathbb{F}^\times/(\mathbb{F}^\times)^2$.

Proposition 2.5.10. *Matrices Represent Forms up to Congruence*
Let β, γ be ordered bases of $\mathcal{V}$ and put $C = C_{\beta \leftarrow \gamma}$.

(i) $\langle \mathbf{x}, \mathbf{y} \rangle = [\mathbf{x}]_\beta^t M_\beta [\mathbf{y}]_\beta$.
(ii) $M_\gamma = C^t M_\beta C$.
(iii) Matrices A, N represent the same form on $\mathcal{V}$ if and only if $N = P^t A P$ for some invertible P.

Proof.

(i) If $\mathbf{x} = \sum_i a_i \mathbf{v}_i$ and $\mathbf{y} = \sum_j b_j \mathbf{v}_j$, then $\langle \mathbf{x}, \mathbf{y} \rangle = \sum_{i,j} a_i b_j \langle \mathbf{v}_i, \mathbf{v}_j \rangle = [\mathbf{x}]_\beta^t M_\beta [\mathbf{y}]_\beta$.
(ii) Since $[\mathbf{x}]_\beta = C[\mathbf{x}]_\gamma$ and $[\mathbf{y}]_\beta = C[\mathbf{y}]_\gamma$, $[\mathbf{x}]_\gamma^t M_\gamma [\mathbf{y}]_\gamma = \langle \mathbf{x}, \mathbf{y} \rangle = [\mathbf{x}]_\gamma^t C^t M_\beta C [\mathbf{y}]_\gamma$, so $M_\gamma = C^t M_\beta C$.
(iii) If $A = M_\beta$ and $N = M_\gamma$ represent one form, then (ii) gives $N = C_{\beta \leftarrow \gamma}^t A C_{\beta \leftarrow \gamma}$.
Conversely, suppose $N = P^t A P$ with $P = [p_{ij}]$ invertible and $A = M_\beta$.
Define $\gamma = (\mathbf{w}_1, \dots, \mathbf{w}_n)$ by $\mathbf{w}_j = \sum_i p_{ij} \mathbf{v}_i$. Then $C_{\beta \leftarrow \gamma} = P$, so (ii) gives $M_\gamma = P^t A P = N$.

Structural Commentary: Congruence is basis change for forms.
Key Insight: Basis change acts in both slots: $M \mapsto P^t M P$. □

Definition 2.5.11. *Quadratic Forms and Polar Forms*
A *quadratic form* on $\mathcal{V}$ is a function $Q : \mathcal{V} \to \mathbb{F}$ such that:

(i) $Q(a\mathbf{v}) = a^2 Q(\mathbf{v})$ for all $a \in \mathbb{F}$ and $\mathbf{v} \in \mathcal{V}$;
(ii) the *polar form* $B_Q(\mathbf{u}, \mathbf{v}) := Q(\mathbf{u} + \mathbf{v}) - Q(\mathbf{u}) - Q(\mathbf{v})$ is bilinear.

Proposition 2.5.12. *Quadratic Forms and Symmetric Bilinear Forms*
For every quadratic form Q, its polar form B_Q is symmetric.
If $\operatorname{char}(\mathbb{F}) \neq 2$, *maps* $\{B : \operatorname{Sym}(B)\} \underset{Q \mapsto B_Q}{\overset{B \mapsto Q_B}{\rightleftarrows}} \{Q : Q \text{ quadratic}\}$ *and* $Q_B(\mathbf{x}) := \frac{1}{2} B(\mathbf{x}, \mathbf{x})$ *are inverse correspondences between symmetric bilinear forms and quadratic forms.*

Proof. First, $B_Q(\mathbf{u}, \mathbf{v}) = Q(\mathbf{u} + \mathbf{v}) - Q(\mathbf{u}) - Q(\mathbf{v}) = B_Q(\mathbf{v}, \mathbf{u})$, so B_Q is symmetric.
($\Rightarrow$) If $\operatorname{Sym}(B)$, then $Q_B(a\mathbf{x}) = a^2 Q_B(\mathbf{x})$ and $B_{Q_B}(\mathbf{u}, \mathbf{v}) = \frac{1}{2}\big(B(\mathbf{u}+\mathbf{v}, \mathbf{u}+\mathbf{v}) - B(\mathbf{u}, \mathbf{u}) - B(\mathbf{v}, \mathbf{v})\big) = B(\mathbf{u}, \mathbf{v})$.
($\Leftarrow$) If Q is quadratic, then $Q_{B_Q}(\mathbf{x}) = \frac{1}{2} B_Q(\mathbf{x}, \mathbf{x}) = \frac{1}{2}\big(Q(2\mathbf{x}) - 2Q(\mathbf{x})\big) = Q(\mathbf{x})$.

Structural Commentary: Away from characteristic 2, quadratic and symmetric bilinear data coincide.
Key Insight: Dividing by 2 recovers diagonal data from polar data. □

Proposition 2.5.13. *Polar Forms in Characteristic Two*
If $\operatorname{char}(\mathbb{F}) = 2$, *then* $\operatorname{Alt}(B_Q)$ *for every quadratic form Q on $\mathcal{V}$.*

Proof. For $\mathbf{x} \in \mathcal{V}$, $B_Q(\mathbf{x}, \mathbf{x}) = Q(2\mathbf{x}) - Q(\mathbf{x}) - Q(\mathbf{x}) = Q(\mathbf{0}) = 0$. Thus $\operatorname{Alt}(B_Q)$.
Structural Commentary: In characteristic 2, every polar form is alternating.
Key Insight: The diagonal of the polar form vanishes automatically because both $2\mathbf{x}$ and $2Q(\mathbf{x})$ vanish. □

Example 2.5.14. *Characteristic Two Loses Diagonal Data*
Assume $\operatorname{char}(\mathbb{F}) = 2$. On $\mathbb{F}^2$, let $Q_0(x, y) = xy$ and $Q_1(x, y) = xy + x^2$.
Then $B_{Q_0}((x, y), (u, v)) = xv + uy = B_{Q_1}((x, y), (u, v))$, and $Q_0(1, 0) = 0 \neq 1 = Q_1(1, 0)$.
Indeed, the added square has zero polar contribution, where $(x + u)^2 - x^2 - u^2 = 2xu = 0$.
More generally, for any linear $L : \mathcal{V} \to \mathbb{F}$, Q and $Q + L^2$ have the same polar form.

Structural Commentary: In characteristic 2, squares are invisible to polarization.
Key Insight: Characteristic 2 quadratic forms need data beyond their polar forms.

Example 2.5.15. *Alternating versus Nonalternating Symmetric Forms*
Assume $\text{char}(\mathbb{F}) = 2$.
The symmetric matrix $H = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$ defines a nonsingular alternating form on $\mathbb{F}^2$, since $\begin{bmatrix} x & y \end{bmatrix} H \begin{bmatrix} x \\ y \end{bmatrix} = 2xy = 0$.
Thus H is not congruent to any nonzero diagonal form: alternating means all self-pairings vanish, while a nonzero diagonal form has some basis vector with nonzero self-pairing. By contrast, [1] is symmetric and nonalternating.
Structural Commentary: In characteristic 2, symmetric forms split into alternating and nonalternating parts.
Key Insight: Nonsingular alternating blocks are invisible on the diagonal.

Proposition 2.5.16. *Isometries from Quadratic Values*
Let $\mathcal{V}$ be an orthogonal geometry with $\text{char}(\mathbb{F}) \neq 2$, let $q(\mathbf{x}) := \langle \mathbf{x}, \mathbf{x} \rangle$, and let $T : \mathcal{V} \to \mathcal{V}$ be bijective. Then T is an isometry if and only if $q(T\mathbf{x}) = q(\mathbf{x})$ for all $\mathbf{x} \in \mathcal{V}$.

Proof.
($\Rightarrow$) If T is an isometry, then $q(T\mathbf{x}) = \langle T\mathbf{x}, T\mathbf{x} \rangle = \langle \mathbf{x}, \mathbf{x} \rangle = q(\mathbf{x})$.
($\Leftarrow$) If $q \circ T = q$, then $\langle T\mathbf{x}, T\mathbf{y} \rangle = \frac{q(T\mathbf{x}+T\mathbf{y})-q(T\mathbf{x})-q(T\mathbf{y})}{2} = \frac{q(\mathbf{x}+\mathbf{y})-q(\mathbf{x})-q(\mathbf{y})}{2} = \langle \mathbf{x}, \mathbf{y} \rangle$.
Thus T preserves the form, so it is an isometry.

Structural Commentary: In odd characteristic, q determines the orthogonal form.
Key Insight: Polarization turns preservation of lengths into preservation of inner products. □

Definition 2.5.17. *Orthogonality, Isotropy, and Radical*
Let $\mathcal{V}$ be a metric vector space.

(i) $\mathbf{x} \perp \mathbf{y}$ means $\langle \mathbf{x}, \mathbf{y} \rangle = 0$; then $\mathbf{x}, \mathbf{y}$ are *orthogonal*. For $S \subseteq \mathcal{V}$, $S^\perp := \{\mathbf{v} \in \mathcal{V} : \langle \mathbf{v}, \mathbf{s} \rangle = 0 \ \forall \mathbf{s} \in S\}$
(ii) $\mathcal{S} \leq \mathcal{V}$ is *totally isotropic* or *totally degenerate* if $\mathcal{S} \subseteq \mathcal{S}^\perp$, and *nonsingular as a subspace* if $\mathcal{S} \cap \mathcal{S}^\perp = \{\mathbf{0}\}$.
(iii) $\mathbf{x} \neq \mathbf{0}$ is *isotropic* or *null* if $\langle \mathbf{x}, \mathbf{x} \rangle = 0$, and *nonisotropic* otherwise.
The space $\mathcal{V}$ is *isotropic* if such an $\mathbf{x}$ exists, and *anisotropic* otherwise.
(iv) $\mathbf{x}$ is *degenerate* if $\mathbf{x} \perp \mathcal{V}$. The *radical* is $\text{rad}(\mathcal{V}) := \mathcal{V}^\perp$
(v) $\mathcal{V}$ is *nonsingular* or *nondegenerate* if $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$, *singular* otherwise.
(vi) $\mathcal{V}$ is *totally singular* or *totally degenerate* if $\text{rad}(\mathcal{V}) = \mathcal{V}$.

Proposition 2.5.18. *Metric Radical as a Bilinear Radical*
Let $B(\mathbf{x}, \mathbf{y}) = \langle \mathbf{x}, \mathbf{y} \rangle$ be the form of a metric vector space $\mathcal{V}$. Then $\text{rad}_L(B) = \text{rad}_R(B) = \text{rad}(\mathcal{V})$.
Thus $\mathcal{V}$ is nonsingular if and only if B is perfect.

Proof.
(i) By definition, $\text{rad}(\mathcal{V}) = \{\mathbf{x} \in \mathcal{V} : \langle \mathbf{x}, \mathbf{y} \rangle = 0 \ \forall \mathbf{y} \in \mathcal{V}\} = \text{rad}_L(B)$
(ii) If $\mathbf{x} \in \text{rad}_L(B)$, then $\langle \mathbf{x}, \mathbf{y} \rangle = 0$ for all $\mathbf{y}$. Since B is symmetric or skew-symmetric, $\langle \mathbf{y}, \mathbf{x} \rangle = \pm \langle \mathbf{x}, \mathbf{y} \rangle = 0$, so $\mathbf{x} \in \text{rad}_R(B)$. The reverse inclusion is identical.
(iii) B is perfect if and only if $\text{rad}_L(B) = \text{rad}_R(B) = \{\mathbf{0}\}$, if and only if $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$.

Structural Commentary: Metric forms are the square bilinear forms whose two null directions agree.
Key Insight: Metric geometry uses the common left-right radical. □

Proposition 2.5.19. *Radical and Representing Matrices*
Let $\beta = (\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be an ordered basis of $\mathcal{V}$.
Then $\{[\mathbf{x}]_\beta : \mathbf{x} \in \text{rad}(\mathcal{V})\} = \ker(M_\beta^t)$, and $\dim \text{rad}(\mathcal{V}) = \dim(\mathcal{V}) - \text{rank}(M_\beta)$.

Proof. For $\mathbf{c} = [\mathbf{x}]_\beta$, $\mathbf{x} \in \text{rad}(\mathcal{V}) \iff \langle \mathbf{x}, \mathbf{v}_j \rangle = 0 \ \forall j \iff \mathbf{c}^t M_\beta = 0 \iff \mathbf{c} \in \ker(M_\beta^t)$.
Rank-nullity gives $\dim \text{rad}(\mathcal{V}) = n - \text{rank}(M_\beta^t) = n - \text{rank}(M_\beta)$.
Structural Commentary: Degenerate vectors are precisely the coordinate null vectors of the representing matrix.
Key Insight: The radical is the coordinate nullspace of M_β^t. □

Proposition 2.5.20. *Matrix Criterion for Nonsingularity*
Let β be an ordered basis of $\mathcal{V}$. Then $\mathcal{V}$ is nonsingular if and only if M_β is invertible.

Proof.
($\Rightarrow$) If $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$, then $\ker(M_\beta^t) = \{\mathbf{0}\}$. Since M_β is square, M_β is invertible.
($\Leftarrow$) If M_β is invertible, then $\ker(M_\beta^t) = \{\mathbf{0}\}$, so $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$.

Structural Commentary: Zero radical is matrix invertibility in coordinates.
Key Insight: Nonsingularity of a form is ordinary matrix nonsingularity in any basis. □

Remark 2.5.21. *Subspace Radicals*
For $\mathcal{S} \leq \mathcal{V}$, compute $\text{rad}(\mathcal{S})$ using $\langle \cdot, \cdot \rangle|_{\mathcal{S}\times\mathcal{S}}$, where $\text{rad}(\mathcal{S}) = \{\mathbf{s} \in \mathcal{S} : \langle \mathbf{s}, \mathbf{t} \rangle = 0 \ \forall \mathbf{t} \in \mathcal{S}\} = \mathcal{S} \cap \mathcal{S}^\perp$.
In general, $\text{rad}(\mathcal{S}) \neq \mathcal{S}^\perp$; the latter is ambient.

Proposition 2.5.22. *Riesz Map Isomorphism Criterion*
Let $\mathcal{V}$ be a finite-dimensional metric vector space and define $\tau : \mathcal{V} \to \mathcal{V}'$ where $\tau(\mathbf{x})(\mathbf{v}) = \langle \mathbf{v}, \mathbf{x} \rangle$.
Then $\ker(\tau) = \text{rad}(\mathcal{V})$, and τ is an isomorphism $\iff \text{rad}(\mathcal{V}) = \{\mathbf{0}\}$.

Proof. Bilinearity gives $\tau \in \mathcal{L}(\mathcal{V}, \mathcal{V}')$, and
$\mathbf{x} \in \ker(\tau) \iff \langle \mathbf{v}, \mathbf{x} \rangle = 0 \ \forall \mathbf{v} \in \mathcal{V} \iff \langle \mathbf{x}, \mathbf{v} \rangle = 0 \ \forall \mathbf{v} \in \mathcal{V} \iff \mathbf{x} \in \mathcal{V}^\perp = \text{rad}(\mathcal{V})$
($\Rightarrow$) If τ is an isomorphism, then $\ker(\tau) = \{\mathbf{0}\}$, hence $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$.
($\Leftarrow$) If $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$, then $\ker(\tau) = \{\mathbf{0}\}$. Since $\dim \mathcal{V} = \dim \mathcal{V}'$, τ is an isomorphism.
Structural Commentary: The Riesz map is invertible exactly off the radical.
Key Insight: Nonsingularity is injectivity of the Riesz map. □

Theorem 2.5.23. *Riesz Map for Nonsingular Metric Spaces*
Let $\mathcal{V}$ be a finite-dimensional nonsingular metric vector space. Then $\forall f \in \mathcal{V}' \ \exists! \, \mathbf{x} \in \mathcal{V}$, $f(\mathbf{v}) = \langle \mathbf{v}, \mathbf{x} \rangle \ \ \forall \mathbf{v} \in \mathcal{V}$.

Proof. Let $\tau : \mathcal{V} \to \mathcal{V}'$ be given by $\tau(\mathbf{x})(\mathbf{v}) = \langle \mathbf{v}, \mathbf{x} \rangle$. By Proposition 2.5.22, τ is an isomorphism.
Hence each $f \in \mathcal{V}'$ has a unique $\mathbf{x} \in \mathcal{V}$ with $\tau(\mathbf{x}) = f$.
Structural Commentary: Nonsingularity identifies vectors with covectors; positivity is irrelevant.
Key Insight: Riesz representation is fundamentally nonsingularity, not order or length. □

Definition 2.5.24. *Adjoint for a Nonsingular Bilinear Form*
Let $\mathcal{V}$ be nonsingular.
For $T \in \mathcal{L}(\mathcal{V})$, the *metric adjoint* is the unique $T^* \in \mathcal{L}(\mathcal{V})$ such that $\langle T\mathbf{x}, \mathbf{y} \rangle = \langle \mathbf{x}, T^*\mathbf{y} \rangle \ \forall \mathbf{x}, \mathbf{y} \in \mathcal{V}$.

Proposition 2.5.25. *Metric Adjoints and Isometries*
Let $\mathcal{V}$ be nonsingular and let $T \in \mathcal{L}(\mathcal{V})$.

(i) $\exists! \, T^* \in \mathcal{L}(\mathcal{V})$.
(ii) If β is ordered, $M = M_\beta$, and $A = [T]_\beta$, then $[T^*]_\beta = M^{-1}A^tM$.
(iii) T is an isometry $\iff$ T is bijective and $T^*T = I$. For bijective T, this is equivalent to $T^{-1} = T^*$.

Proof.
(i) For $\mathbf{y} \in \mathcal{V}$, set $f_\mathbf{y}(\mathbf{x}) = \langle T\mathbf{x}, \mathbf{y} \rangle$. Theorem 2.5.23 gives a unique $T^*\mathbf{y}$ such that $f_\mathbf{y}(\mathbf{x}) = \langle \mathbf{x}, T^*\mathbf{y} \rangle$ for all $\mathbf{x}$. Uniqueness gives linearity of $\mathbf{y} \mapsto T^*\mathbf{y}$.
(ii) With $A = [T]_\beta$ and $B = [T^*]_\beta$, the adjoint identity gives $A^tM = MB$, so $B = M^{-1}A^tM$.
(iii) $\forall \mathbf{x}, \mathbf{y} \in \mathcal{V}$, $\langle T\mathbf{x}, T\mathbf{y} \rangle = \langle \mathbf{x}, T^*T\mathbf{y} \rangle$. so form preservation $\iff \langle \mathbf{x}, (T^*T - I)\mathbf{y} \rangle = 0 \ \forall \mathbf{x}, \mathbf{y} \iff T^*T = I$. This gives the stated bijective criterion; if T is bijective, then $T^*T = I \iff T^* = T^{-1}$.
Structural Commentary: Adjoint theory needs nonsingularity, not positivity.
Key Insight: Isometry means the adjoint is the inverse. □

Theorem 2.5.26. *Dimensions of Orthogonal Complements*
Let $\mathcal{S} \leq \mathcal{V}$. If $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$ or $\text{rad}(\mathcal{S}) = \{\mathbf{0}\}$, then $\dim(\mathcal{S}) + \dim(\mathcal{S}^\perp) = \dim(\mathcal{V})$.
If $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$, then $\mathcal{S}^{\perp\perp} = \mathcal{S}$, $\text{rad}(\mathcal{S}) = \{\mathbf{0}\} \iff \text{rad}(\mathcal{S}^\perp) = \{\mathbf{0}\}$.

Proof.
i. Define $\rho : \mathcal{V} \to \mathcal{S}'$ by $\rho(\mathbf{x})(\mathbf{s}) = \langle \mathbf{s}, \mathbf{x} \rangle$. Then $\ker(\rho) = \mathcal{S}^\perp$.
ii. If $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$, extend $f \in \mathcal{S}'$ to $\widetilde{f} \in \mathcal{V}'$ and use Theorem 2.5.23. If $\text{rad}(\mathcal{S}) = \{\mathbf{0}\}$, represent f inside $\mathcal{S}$. Thus ρ is onto, so rank-nullity gives $\dim \mathcal{V} = \dim \mathcal{S}^\perp + \dim \mathcal{S}$.
iii. If $\text{rad}(\mathcal{V}) = \{\mathbf{0}\}$, then $\mathcal{S} \subseteq \mathcal{S}^{\perp\perp}$ and $\dim \mathcal{S}^{\perp\perp} = \dim \mathcal{V} - \dim \mathcal{S}^\perp = \dim \mathcal{S}$, hence $\mathcal{S}^{\perp\perp} = \mathcal{S}$.
($\Rightarrow$) If $\text{rad}(\mathcal{S}) = \{\mathbf{0}\}$, then $\text{rad}(\mathcal{S}^\perp) = \mathcal{S}^\perp \cap \mathcal{S}^{\perp\perp} = \mathcal{S}^\perp \cap \mathcal{S}$, so $\text{rad}(\mathcal{S}^\perp) = \{\mathbf{0}\}$.
($\Leftarrow$) If $\text{rad}(\mathcal{S}^\perp) = \{\mathbf{0}\}$, then $\mathcal{S}^\perp \cap \mathcal{S} = \{\mathbf{0}\}$, so $\text{rad}(\mathcal{S}) = \mathcal{S} \cap \mathcal{S}^\perp = \{\mathbf{0}\}$.
Structural Commentary: Complement dimensions come from surjectivity of the restriction-pairing map.
Key Insight: Orthogonal complement dimensions are rank-nullity for the restriction-of-pairing map. □

Definition 2.5.27. *Orthogonal Direct Sum*
For subspaces $\mathcal{S}, \mathcal{T} \leq \mathcal{V}$, write $\mathcal{V} = \mathcal{S} \odot \mathcal{T} \iff \mathcal{V} = \mathcal{S} \oplus \mathcal{T}$ and $\mathcal{S} \perp \mathcal{T}$.
Then $\mathcal{V}$ is the *orthogonal direct sum* of $\mathcal{S}$ and $\mathcal{T}$.

Theorem 2.5.28. *Splitting off the Radical*
Let $R = \text{rad}(\mathcal{V})$. Then $\exists \mathcal{S} \leq \mathcal{V}$ such that $\mathcal{V} = R \odot \mathcal{S}$, $\text{rad}(\mathcal{S}) = \{\mathbf{0}\}$, and R is totally singular.

Proof. Choose $\mathcal{S}$ with $\mathcal{V} = R \oplus \mathcal{S}$. Since $R \perp \mathcal{V}$, $\mathcal{V} = R \odot \mathcal{S}$ and $R \subseteq R^\perp$, so R is totally singular. If $\mathbf{s} \in \operatorname{rad}(\mathcal{S})$, then $\mathbf{s} \perp \mathcal{S}$ and $\mathbf{s} \perp R$, hence $\mathbf{s} \in R \cap \mathcal{S} = \{\mathbf{0}\}$.
Structural Commentary: Singularity splits as an invisible summand.
Key Insight: Once the radical is removed, the remaining form is nonsingular. □

Proposition 2.5.29. *Metric Quotients*
Let $\mathcal{U} \leq \mathcal{V}$. The rule $\langle \mathbf{x}+\mathcal{U}, \mathbf{y}+\mathcal{U}\rangle_{\mathcal{V}/\mathcal{U}} := \langle \mathbf{x}, \mathbf{y}\rangle$ is well-defined if and only if $\mathcal{U} \leq \operatorname{rad}(\mathcal{V})$.
Then $\operatorname{rad}(\mathcal{V}/\mathcal{U}) = \operatorname{rad}(\mathcal{V})/\mathcal{U}$. In particular, $\mathcal{V}/\operatorname{rad}(\mathcal{V})$ is nonsingular.

Proof.
($\Rightarrow$) If the rule is well-defined, then $0 = \langle \mathbf{u}+\mathcal{U}, \mathbf{y}+\mathcal{U}\rangle_{\mathcal{V}/\mathcal{U}} = \langle \mathbf{u}, \mathbf{y}\rangle \ \forall\, \mathbf{u} \in \mathcal{U},\ \mathbf{y} \in \mathcal{V}$, so $\mathcal{U} \leq \operatorname{rad}(\mathcal{V})$.
($\Leftarrow$) If $\mathcal{U} \leq \operatorname{rad}(\mathcal{V})$, replacing $\mathbf{x}, \mathbf{y}$ by $\mathbf{x}+\mathbf{u}, \mathbf{y}+\mathbf{v}$ changes $\langle \mathbf{x}, \mathbf{y}\rangle$ by $\langle \mathbf{u}, \mathbf{y}\rangle + \langle \mathbf{x}, \mathbf{v}\rangle + \langle \mathbf{u}, \mathbf{v}\rangle = 0$, rule well-defined.
If $\mathcal{U} \leq \operatorname{rad}(\mathcal{V})$, then $\mathbf{x}+\mathcal{U} \in \operatorname{rad}(\mathcal{V}/\mathcal{U}) \iff \langle \mathbf{x}, \mathbf{y}\rangle = 0\ \forall\, \mathbf{y} \in \mathcal{V} \iff \mathbf{x} \in \operatorname{rad}(\mathcal{V})$, hence $\operatorname{rad}(\mathcal{V}/\mathcal{U}) = \operatorname{rad}(\mathcal{V})/\mathcal{U}$.
Structural Commentary: Quotients preserve the form exactly along invisible directions.
Key Insight: The radical is the largest subspace that can be collapsed without making the form ambiguous. □

Proposition 2.5.30. *Radicals and Orthogonal Direct Sums*
If $\mathcal{V} = \mathcal{S} \odot \mathcal{T}$, then $\operatorname{rad}(\mathcal{V}) = \operatorname{rad}(\mathcal{S}) \odot \operatorname{rad}(\mathcal{T})$; $\mathcal{V}/\operatorname{rad}(\mathcal{V}) \cong \mathcal{S}/\operatorname{rad}(\mathcal{S}) \odot \mathcal{T}/\operatorname{rad}(\mathcal{T})$ as metric vector spaces.
Hence $\mathcal{V}$ is nonsingular if and only if $\mathcal{S}$ and $\mathcal{T}$ are nonsingular.

Proof. For $\mathbf{v} = \mathbf{s}+\mathbf{t}$, $\mathbf{v} \in \operatorname{rad}(\mathcal{V}) \iff \mathbf{v} \perp \mathcal{V} \iff \mathbf{s} \perp \mathcal{S}$ and $\mathbf{t} \perp \mathcal{T} \iff \mathbf{s} \in \operatorname{rad}(\mathcal{S}),\ \mathbf{t} \in \operatorname{rad}(\mathcal{T})$.
Hence $\operatorname{rad}(\mathcal{V}) = \operatorname{rad}(\mathcal{S}) \odot \operatorname{rad}(\mathcal{T})$. The isomorphism is induced by $\Phi((\mathbf{s}+\mathbf{t})+\operatorname{rad}(\mathcal{V})) = (\mathbf{s}+\operatorname{rad}(\mathcal{S}), \mathbf{t}+\operatorname{rad}(\mathcal{T}))$.
It is well-defined and bijective by the radical formula, and it preserves the form since cross terms vanish.
($\Rightarrow$) If $\mathcal{V}$ is nonsingular, then $\operatorname{rad}(\mathcal{S}) \odot \operatorname{rad}(\mathcal{T}) = \{\mathbf{0}\}$, so $\mathcal{S}$ and $\mathcal{T}$ are nonsingular.
($\Leftarrow$) If $\mathcal{S}$ and $\mathcal{T}$ are nonsingular, then $\operatorname{rad}(\mathcal{V}) = \operatorname{rad}(\mathcal{S}) \odot \operatorname{rad}(\mathcal{T}) = \{\mathbf{0}\}$.
Structural Commentary: Orthogonal direct sums split both visible and invisible parts.
Key Insight: The regular quotient of an orthogonal sum is the orthogonal sum of the regular quotients. □

Corollary 2.5.31. *Recovering a Singular Space from Its Regular Quotient*
If $\dim \mathcal{V} = \dim \mathcal{W}$ and $\mathcal{V}/\operatorname{rad}(\mathcal{V}) \cong \mathcal{W}/\operatorname{rad}(\mathcal{W})$, then $\mathcal{V} \cong \mathcal{W}$ as metric vector spaces.

Proof. By Theorem 2.5.28, choose $\mathcal{V} = \operatorname{rad}(\mathcal{V}) \odot \mathcal{S}$, $\mathcal{W} = \operatorname{rad}(\mathcal{W}) \odot \mathcal{T}$, with $\mathcal{S}, \mathcal{T}$ nonsingular.
The quotient isometry gives $\mathcal{S} \cong \mathcal{T}$. Since $\dim \mathcal{V} = \dim \mathcal{W}$, also $\dim \operatorname{rad}(\mathcal{V}) = \dim \operatorname{rad}(\mathcal{W})$; choose any vector-space isomorphism between the radicals, which is an isometry because both radicals are totally singular.
The orthogonal direct sum of these two isometries gives $\mathcal{V} \cong \mathcal{W}$.
Structural Commentary: A singular metric space is its regular quotient plus invisible dimension.
Key Insight: Once the regular quotient and total dimension are fixed, the radical has no further metric data. □

Definition 2.5.32. *Hyperbolic Pairs and Hyperbolic Spaces*

(i) A *hyperbolic pair* is $(\mathbf{u}, \mathbf{v})$ with $\langle \mathbf{u}, \mathbf{u}\rangle = \langle \mathbf{v}, \mathbf{v}\rangle = 0$ and $\langle \mathbf{u}, \mathbf{v}\rangle = 1$.
(ii) The plane $\operatorname{span}(\mathbf{u}, \mathbf{v})$ of a hyperbolic pair is a *hyperbolic plane*.
(iii) A *hyperbolic space* is an orthogonal direct sum of hyperbolic planes.
(iv) A basis $(\mathbf{u}_1, \mathbf{v}_1, \ldots, \mathbf{u}_k, \mathbf{v}_k)$ is a *hyperbolic basis* if each $(\mathbf{u}_i, \mathbf{v}_i)$ is hyperbolic and the planes $\operatorname{span}(\mathbf{u}_i, \mathbf{v}_i)$ are mutually orthogonal. In the symplectic case it is a *symplectic basis*.

Theorem 2.5.33. *Capturing Isotropic Vectors*
Let $\mathcal{V}$ be nonsingular. If $\mathcal{V}$ is orthogonal, assume $\operatorname{char}(\mathbb{F}) \neq 2$.
For every isotropic $\mathbf{u} \in \mathcal{V}$, $\exists\, \mathbf{v} \in \mathcal{V}$ such that $(\mathbf{u}, \mathbf{v})$ is hyperbolic.

Proof. Choose $\mathbf{x}$ with $c := \langle \mathbf{u}, \mathbf{x}\rangle \neq 0$. If the form is alternating, take $\mathbf{v} = c^{-1}\mathbf{x}$. Then $\langle \mathbf{u}, \mathbf{v}\rangle = 1$ and $\langle \mathbf{v}, \mathbf{v}\rangle = 0$.
If the form is symmetric, set $\mathbf{v} = c^{-1}\mathbf{x} - \frac{\langle \mathbf{x}, \mathbf{x}\rangle}{2c^2}\mathbf{u}$. Using $\langle \mathbf{u}, \mathbf{u}\rangle = 0$ gives $\langle \mathbf{u}, \mathbf{v}\rangle = 1$ and $\langle \mathbf{v}, \mathbf{v}\rangle = 0$.
Structural Commentary: Nonsingularity supplies a partner; one adjustment makes it isotropic.
Key Insight: Isotropic directions in nonsingular spaces come in hyperbolic pairs. □

Proposition 2.5.34. *Recognizing Orthogonal Hyperbolic Planes*
Let $\mathcal{H}$ be a 2-dimensional nonsingular orthogonal geometry with $\operatorname{char}(\mathbb{F}) \neq 2$. The following are equivalent:

(i) $\mathcal{H}$ is a hyperbolic plane;
(ii) $\exists\, 0 \neq \mathbf{u} \in \mathcal{H}$ with $\langle \mathbf{u}, \mathbf{u}\rangle = 0$;
(iii) $\Delta(\mathcal{H}) = (\mathbb{F}^\times)^2(-1)$ as a square class.

Proof.
(i) $\Rightarrow$ *(ii)* A hyperbolic pair consists of nonzero isotropic vectors.

(*ii*) $\Rightarrow$ (*i*) By Theorem 2.5.33, choose $\mathbf{v}$ with $(\mathbf{u}, \mathbf{v})$ hyperbolic. Since $\dim \mathcal{H} = 2$, $\mathcal{H} = \operatorname{span}(\mathbf{u}, \mathbf{v})$.

(*i*) $\Rightarrow$ (*iii*) In a hyperbolic basis, $M = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$, so $\det M = -1$ and $\Delta(\mathcal{H}) = (\mathbb{F}^\times)^2(-1)$.

(*iii*) $\Rightarrow$ (*ii*) Diagonalize $M = \operatorname{diag}(a, b)$ with $a, b \neq 0$. If $\Delta(\mathcal{H}) = (\mathbb{F}^\times)^2(-1)$, then $ab = -r^2$ for some $r \neq 0$. Hence $a(r/a)^2 + b = 0$. Thus $(r/a, 1)$ is isotropic in this basis.

Structural Commentary: In dimension 2, isotropy and discriminant detect hyperbolicity.
Key Insight: Hyperbolic means the form represents 0 nontrivially. □

Theorem 2.5.35. *Nonsingular Completions*
Let $\mathcal{U}$ be a subspace of a nonsingular metric vector space $\mathcal{V}$, and assume $\operatorname{char}(\mathbb{F}) \neq 2$ *in the orthogonal case. Write $\mathcal{U} = \operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_r) \odot \mathcal{W}$, where $(\mathbf{u}_1, \dots, \mathbf{u}_r)$ is a basis of* $\operatorname{rad}(\mathcal{U})$ *and $\mathcal{W}$ is nonsingular. Then $\exists\, \mathbf{v}_1, \dots, \mathbf{v}_r$ such that $\widehat{\mathcal{U}} = \mathcal{H}_r \odot \mathcal{W}$; $\mathcal{H}_r = \operatorname{span}(\mathbf{u}_1, \mathbf{v}_1, \dots, \mathbf{u}_r, \mathbf{v}_r)$ is nonsingular, contains $\mathcal{U}$, and* $\dim \widehat{\mathcal{U}} = \dim \mathcal{U} + \dim \operatorname{rad}(\mathcal{U})$.
The minimal nonsingular extensions of $\mathcal{U}$ are exactly such hyperbolic extensions, up to isometry.

Proof.

(i) Since $\mathcal{W}$ is nonsingular, Theorem 2.5.26 gives $\mathcal{V} = \mathcal{W} \odot \mathcal{W}^\perp$. Thus $\mathbf{u}_i \in \mathcal{W}^\perp$ for all i.

(ii) Construct $\mathbf{v}_j$ inductively. Suppose $\mathcal{H}_{j-1}$ is built and set $\mathcal{S}_j = \operatorname{span}(\mathbf{u}_{j+1}, \dots, \mathbf{u}_r) \odot \mathcal{W} \odot \mathcal{H}_{j-1}$. Then $\mathbf{u}_j \perp \mathcal{S}_j$ and $\mathbf{u}_j \notin \mathcal{S}_j = \mathcal{S}_j^{\perp\perp}$, so $\exists\, \mathbf{x} \in \mathcal{S}_j^\perp$ with $c := \langle \mathbf{u}_j, \mathbf{x} \rangle \neq 0$. Put

$$\mathbf{v}_j = \begin{cases} c^{-1}\mathbf{x}, & \text{alternating case,} \\ c^{-1}\mathbf{x} - \dfrac{\langle \mathbf{x}, \mathbf{x} \rangle}{2c^2}\mathbf{u}_j, & \text{orthogonal case.} \end{cases}$$

Then $\mathbf{v}_j \in \mathcal{S}_j^\perp$ and $(\mathbf{u}_j, \mathbf{v}_j)$ is hyperbolic; set $\mathcal{H}_j = \mathcal{H}_{j-1} \odot \operatorname{span}(\mathbf{u}_j, \mathbf{v}_j)$.

(iii) After r steps, $\widehat{\mathcal{U}} = \mathcal{H}_r \odot \mathcal{W}$ is nonsingular and $\dim \widehat{\mathcal{U}} = \dim \mathcal{U} + r$.

(iv) Let $\mathcal{X} \supseteq \mathcal{U}$ be nonsingular and choose $\mathcal{X} = \mathcal{U} \oplus \mathcal{Y}$. If $0 \neq \mathbf{u} \in \operatorname{rad}(\mathcal{U})$ paired trivially with $\mathcal{Y}$, then $\mathbf{u} \in \operatorname{rad}(\mathcal{X})$, impossible. Hence the pairing $\operatorname{rad}(\mathcal{U}) \times \mathcal{Y} \to \mathbb{F}$ has zero left radical, so $\dim \mathcal{Y} \geq r$ and $\dim \mathcal{X} \geq \dim \mathcal{U} + r$. If equality holds, the construction inside $\mathcal{X}$ gives $\widehat{\mathcal{U}} \subseteq \mathcal{X}$ with equal dimension, hence $\widehat{\mathcal{U}} = \mathcal{X}$. Any two minimal completions are isometric since each is $\mathcal{W} \odot$ a hyperbolic r-space.

Structural Commentary: Add one hyperbolic partner per radical direction.
Key Insight: Minimal nonsingular completion is radical repair by hyperbolic planes. □

Principle 2.5.36. *Witt Extension and Cancellation*
Let $\mathcal{V} \cong \mathcal{V}'$ be nonsingular metric vector spaces of same kind, either symplectic or orthogonal with $\operatorname{char}(\mathbb{F}) \neq 2$. Then the following equivalent principles hold:

(i) *Witt extension*: every isometry $\tau : \mathcal{S} \to \mathcal{T}$ between subspaces $\mathcal{S} \leq \mathcal{V}$ and $\mathcal{T} \leq \mathcal{V}'$ extends to an isometry $\overline{\tau} : \mathcal{V} \to \mathcal{V}'$ with $\overline{\tau}|_{\mathcal{S}} = \tau$;

(ii) *Witt cancellation*: for nonsingular subspaces $\mathcal{S} \leq \mathcal{V}$ and $\mathcal{T} \leq \mathcal{V}'$, $\mathcal{S} \cong \mathcal{T} \Rightarrow \mathcal{S}^\perp \cong \mathcal{T}^\perp$.

Proposition 2.5.37. *Witt Extension–Cancellation Bridge*
Let $\mathcal{V} \cong \mathcal{V}'$ be nonsingular metric vector spaces of same kind, either symplectic or orthogonal with $\operatorname{char}(\mathbb{F}) \neq 2$. *Using Theorem 2.5.35, Witt extension $\iff$ Witt cancellation.*

Proof.

($\Rightarrow$) Let $\mathcal{V} = \mathcal{S} \odot \mathcal{S}^\perp$, $\mathcal{V}' = \mathcal{T} \odot \mathcal{T}^\perp$, and let $\sigma : \mathcal{S} \to \mathcal{T}$ be an isometry. Extend σ to $\widehat{\sigma} : \mathcal{V} \to \mathcal{V}'$. Then $\widehat{\sigma}(\mathcal{S}^\perp) = \mathcal{T}^\perp$, hence $\mathcal{S}^\perp \cong \mathcal{T}^\perp$.

($\Leftarrow$) Let $\sigma : \mathcal{S} \to \mathcal{T}$ be an isometry. By Theorem 2.5.35, choose minimal nonsingular completions $\widehat{\mathcal{S}} \supseteq \mathcal{S}$ and $\widehat{\mathcal{T}} \supseteq \mathcal{T}$, and extend σ to $\widehat{\sigma} : \widehat{\mathcal{S}} \to \widehat{\mathcal{T}}$ by matching hyperbolic partners of $\operatorname{rad}(\mathcal{S})$ and $\operatorname{rad}(\mathcal{T})$. Then $\mathcal{V} = \widehat{\mathcal{S}} \odot \widehat{\mathcal{S}}^\perp$, $\mathcal{V}' = \widehat{\mathcal{T}} \odot \widehat{\mathcal{T}}^\perp$. Cancellation gives $\eta : \widehat{\mathcal{S}}^\perp \to \widehat{\mathcal{T}}^\perp$, so $\widehat{\sigma} \odot \eta : \mathcal{V} \to \mathcal{V}'$ extends σ.

Structural Commentary: Complete first, then cancel complements.
Key Insight: Complete singular subspaces first; after that, the complement is the only missing piece. □

Theorem 2.5.38. *Witt Extension Theorem*
Let $\mathcal{V} \cong \mathcal{W}$ be nonsingular metric vector spaces of same kind, either symplectic or orthogonal with $\operatorname{char}(\mathbb{F}) \neq 2$. *For every subspace isometry $\sigma : \mathcal{S} \to \mathcal{T}$, $\mathcal{S} \leq \mathcal{V}$, $\mathcal{T} \leq \mathcal{W}$, there exists an isometry $\widehat{\sigma} : \mathcal{V} \to \mathcal{W}$ with $\widehat{\sigma}|_{\mathcal{S}} = \sigma$. Thus, if $\mathcal{V} = \mathcal{W}$, every subspace isometry is the restriction of an ambient isometry.*

Proof. Induct simultaneously on $(\dim \mathcal{V}, \dim \mathcal{S})$. If $\mathcal{S} = 0$, take any fixed isometry $\mathcal{V} \cong \mathcal{W}$.
Let $\mathcal{R} = \operatorname{rad}(\mathcal{S})$, $\mathcal{R}' = \sigma\mathcal{R} = \operatorname{rad}(\mathcal{T})$. By Theorem 2.5.35, choose minimal nonsingular completions $\widehat{\mathcal{S}} \supseteq \mathcal{S}$ and

$\widehat{\mathcal{T}} \supseteq \mathcal{T}$; matching hyperbolic partners over $\mathcal{R}$ and $\mathcal{R}'$ extends σ to $\widehat{\sigma} : \widehat{\mathcal{S}} \to \widehat{\mathcal{T}}$. Hence assume $\mathcal{S}$, $\mathcal{T}$ nonsingular. Choose a nonsingular summand $\mathcal{A} \leq \mathcal{S}$, where

$$\mathcal{A} = \begin{cases} \operatorname{span}(\mathbf{s}, \mathbf{u}) \text{ hyperbolic}, & \text{symplectic case}, \\ \operatorname{span}(\mathbf{s}),\ \langle \mathbf{s}, \mathbf{s} \rangle \neq 0, & \text{orthogonal case}. \end{cases}$$

Set $\mathcal{B} = \sigma\mathcal{A}$, $\mathcal{S}_0 = \mathcal{S} \cap \mathcal{A}^\perp$, and $\mathcal{T}_0 = \mathcal{T} \cap \mathcal{B}^\perp$. One-summand case extends $\sigma|_{\mathcal{A}}$ to $\alpha : \mathcal{V} \to \mathcal{W}$; hence $\alpha(\mathcal{A}^\perp) = \mathcal{B}^\perp$ and $\mathcal{V} = \mathcal{A} \odot \mathcal{A}^\perp$, $\mathcal{W} = \mathcal{B} \odot \mathcal{B}^\perp$, $\mathcal{S} = \mathcal{A} \odot \mathcal{S}_0$, and $\mathcal{T} = \mathcal{B} \odot \mathcal{T}_0$. Since $\dim \mathcal{A}^\perp < \dim \mathcal{V}$, induction gives an isometry $\eta : \mathcal{A}^\perp \to \mathcal{B}^\perp$ with $\eta|_{\mathcal{S}_0} = \sigma|_{\mathcal{S}_0}$. Then $(\sigma|_{\mathcal{A}}) \odot \eta : \mathcal{V} \to \mathcal{W}$ extends σ.
Structural Commentary: Subspace isometries extend ambiently.
Key Insight: Complete radicals, then peel matched summands. □

Corollary 2.5.39. *Witt Cancellation*
Let $\mathcal{V} = \mathcal{S} \odot \mathcal{A}$, $\mathcal{W} = \mathcal{T} \odot \mathcal{B}$ be nonsingular metric vector spaces of the same kind, either symplectic or orthogonal with $\operatorname{char}(\mathbb{F}) \neq 2$. If $\mathcal{V} \cong \mathcal{W}$ and $\mathcal{S} \cong \mathcal{T}$, then $\mathcal{A} \cong \mathcal{B}$.

Proof. Let $\sigma : \mathcal{S} \to \mathcal{T}$ be an isometry. By Theorem 2.5.38, $\exists\, \widehat{\sigma} : \mathcal{V} \to \mathcal{W}$ with $\widehat{\sigma}|_{\mathcal{S}} = \sigma$.
Since $\mathcal{A} = \mathcal{S}^\perp$ and $\mathcal{B} = \mathcal{T}^\perp$, $\widehat{\sigma}(\mathcal{A}) = \widehat{\sigma}(\mathcal{S}^\perp) = (\widehat{\sigma}\mathcal{S})^\perp = \mathcal{T}^\perp = \mathcal{B}$. Thus $\widehat{\sigma}|_{\mathcal{A}} : \mathcal{A} \to \mathcal{B}$ is an isometry.
Structural Commentary: Matched summands cancel.
Key Insight: Extend, then restrict to complements. □

Perspective 2.5.40. *The Classification Problem*
Metric vector spaces are isometric exactly when their representing matrices are congruent.
Thus classification up to isometry is equivalent to classification of the associated matrices up to congruence, i.e. to specifying canonical representatives or complete congruence invariants.

Theorem 2.5.41. *Classification of Symplectic Geometries*
Let $\mathcal{V}$ be a symplectic geometry, $R = \operatorname{rad}(\mathcal{V})$, and $n = \dim(\mathcal{V})$.

(i) If $R = 0$, then $\mathcal{V}$ is hyperbolic; hence n is even.
(ii) $\mathcal{V} = R \odot \mathcal{H}$, where R is totally singular and $\mathcal{H}$ is hyperbolic.
(iii) If $r = \operatorname{rank}(\mathcal{V}) = 2k$, then some basis has representing matrix $M_{\mathcal{V}} = X_{2k,n-2k} := \operatorname{diag}(J_{2k}, 0_{n-2k})$, $J_{2k} := \bigoplus_{i=1}^{k} \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$. For symplectic geometries, $\mathcal{V}, \mathcal{W}$, $\mathcal{V} \cong \mathcal{W} \iff (\dim \mathcal{V}, \operatorname{rank} \mathcal{V}) = (\dim \mathcal{W}, \operatorname{rank} \mathcal{W})$.

Proof.

(i) The case $\mathcal{V} = 0$ is clear. If $R = 0$ and $\mathcal{V} \neq 0$, then every $0 \neq x \in \mathcal{V}$ is isotropic.
By Theorem 2.5.33, x lies in a hyperbolic plane $\mathcal{H}_1 \subset \mathcal{V}$. Since $\mathcal{H}_1$ is nonsingular, Theorem 2.5.26 gives $\mathcal{V} = \mathcal{H}_1 \odot \mathcal{H}_1^\perp$. Induction on n gives the claim.
(ii) By Theorem 2.5.28, $\mathcal{V} = R \odot \mathcal{H}$ with $\operatorname{rad}(\mathcal{H}) = 0$.
Hence $\mathcal{H}$ is hyperbolic by (i), and R is totally singular since $R \perp \mathcal{V}$.
(iii) In (ii), $\dim \mathcal{H} = r = 2k$ and $\dim R = n - 2k$.
A symplectic basis of $\mathcal{H}$ followed by any basis of R gives $M_{\mathcal{V}} = X_{2k,n-2k}$. Dimension and rank are isometry invariants; conversely, equal (n, r) gives the same normal matrix, hence an isometry.

Structural Commentary: Alternating forms have only one nonsingular building block: the hyperbolic plane.
Key Insight: Symplectic classification is rank plus dimension. □

Theorem 2.5.42. *Witt Theorems for Symplectic Geometries*
Let $\mathcal{V} \cong \mathcal{V}'$ be nonsingular symplectic geometries. Then Witt extension and cancellation hold for all subspaces.

Proof. By Proposition 2.5.37, it suffices to prove cancellation for nonsingular $\mathcal{S} \subset \mathcal{V}$ and $\mathcal{T} \subset \mathcal{V}'$ with $\mathcal{S} \cong \mathcal{T}$.
Let $\mathcal{V} = \mathcal{S} \odot \mathcal{S}^\perp$, $\mathcal{V}' = \mathcal{T} \odot \mathcal{T}^\perp$. Then $\dim \mathcal{S}^\perp = \dim \mathcal{V} - \dim \mathcal{S} = \dim \mathcal{V}' - \dim \mathcal{T} = \dim \mathcal{T}^\perp$.
By Theorem 2.5.26, $\mathcal{S}^\perp$ and $\mathcal{T}^\perp$ are nonsingular symplectic geometries; by Theorem 2.5.41, equal dimension gives $\mathcal{S}^\perp \cong \mathcal{T}^\perp$. Thus cancellation holds, and Proposition 2.5.37 gives extension.
Structural Commentary: Cancellation reduces to matching hyperbolic summand counts.
Key Insight: For symplectic forms, dimension alone identifies nonsingular complements. □

Definition 2.5.43. *Symplectic Transvection*
Let $\mathcal{V}$ be a nonsingular symplectic geometry.
For $\mathbf{v} \in \mathcal{V} \setminus \{\mathbf{0}\}$, $a \in \mathbb{F}$, *symplectic transvection* with direction $\mathbf{v}$, parameter a is $\tau_{\mathbf{v},a} : \mathcal{V} \to \mathcal{V}$, $\mathbf{x} \mapsto \mathbf{x} + a \langle \mathbf{x}, \mathbf{v} \rangle \mathbf{v}$.

Proposition 2.5.44. *Basic Properties of Symplectic Transvections*
Let $\mathcal{V}$ be a nonsingular symplectic geometry, $0 \neq \mathbf{v} \in \mathcal{V}$, $a, b \in \mathbb{F}$, $c \in \mathbb{F}^\times$, and $\sigma \in \operatorname{Sp}(\mathcal{V})$. Then:

(i) $\tau_{\mathbf{v},a} \in \operatorname{Sp}(\mathcal{V})$;

(ii) $\tau_{\mathbf{v},a} = I$ *if and only if* $a = 0$;
(iii) $\tau_{\mathbf{v},a}\tau_{\mathbf{v},b} = \tau_{\mathbf{v},a+b}$ *and* $\tau_{\mathbf{v},a}^{-1} = \tau_{\mathbf{v},-a}$;
(iv) $\sigma\tau_{\mathbf{v},a}\sigma^{-1} = \tau_{\sigma\mathbf{v},a}$;
(v) $\tau_{c\mathbf{v},a} = \tau_{\mathbf{v},ac^2}$.

Proof.

(i) For $\mathbf{x}, \mathbf{y} \in \mathcal{V}$, alternation and skew-symmetry give
$\langle \tau_{\mathbf{v},a}\mathbf{x}, \tau_{\mathbf{v},a}\mathbf{y}\rangle = \langle \mathbf{x}, \mathbf{y}\rangle + a\langle \mathbf{x}, \mathbf{v}\rangle\langle \mathbf{v}, \mathbf{y}\rangle + a\langle \mathbf{y}, \mathbf{v}\rangle\langle \mathbf{x}, \mathbf{v}\rangle + a^2\langle \mathbf{x}, \mathbf{v}\rangle\langle \mathbf{y}, \mathbf{v}\rangle\langle \mathbf{v}, \mathbf{v}\rangle = \langle \mathbf{x}, \mathbf{y}\rangle$.
Also $\tau_{\mathbf{v},-a}\tau_{\mathbf{v},a}(\mathbf{x}) = \mathbf{x} = \tau_{\mathbf{v},a}\tau_{\mathbf{v},-a}(\mathbf{x})$, so $\tau_{\mathbf{v},a}$ is invertible and hence symplectic.
(ii) If $a = 0$, then $\tau_{\mathbf{v},a} = I$. If $a \neq 0$, nonsingularity gives $\mathbf{x} \in \mathcal{V}$ with $\langle \mathbf{x}, \mathbf{v}\rangle \neq 0$, so $\tau_{\mathbf{v},a}\mathbf{x} - \mathbf{x} = a\langle \mathbf{x}, \mathbf{v}\rangle \mathbf{v} \neq \mathbf{0}$.
(iii) For $\mathbf{x} \in \mathcal{V}$, since $\langle \mathbf{v}, \mathbf{v}\rangle = 0$, $\tau_{\mathbf{v},a}\tau_{\mathbf{v},b}(\mathbf{x}) = \mathbf{x} + (a+b)\langle \mathbf{x}, \mathbf{v}\rangle \mathbf{v} = \tau_{\mathbf{v},a+b}(\mathbf{x})$. Taking $b = -a$ gives $\tau_{\mathbf{v},a}^{-1} = \tau_{\mathbf{v},-a}$.
(iv) For $\mathbf{x} \in \mathcal{V}$, $\sigma\tau_{\mathbf{v},a}\sigma^{-1}(\mathbf{x}) = \mathbf{x} + a\langle \sigma^{-1}\mathbf{x}, \mathbf{v}\rangle \sigma\mathbf{v} = \mathbf{x} + a\langle \mathbf{x}, \sigma\mathbf{v}\rangle \sigma\mathbf{v} = \tau_{\sigma\mathbf{v},a}(\mathbf{x})$.
(v) For $\mathbf{x} \in \mathcal{V}$, $\tau_{c\mathbf{v},a}(\mathbf{x}) = \mathbf{x} + a\langle \mathbf{x}, c\mathbf{v}\rangle c\mathbf{v} = \mathbf{x} + ac^2\langle \mathbf{x}, \mathbf{v}\rangle \mathbf{v} = \tau_{\mathbf{v},ac^2}(\mathbf{x})$

Structural Commentary: Fix $\mathbf{v}^\perp$; shear along $\mathbb{F}\mathbf{v}$.
Key Insight: Alternation cancels the possible quadratic error term. □

Theorem 2.5.45. *Symplectic Transvections Generate*
For a nonsingular symplectic geometry $\mathcal{V}$, $\mathrm{Sp}(\mathcal{V}) = \langle \tau_{\mathbf{v},a} : 0 \neq \mathbf{v} \in \mathcal{V},\ a \in \mathbb{F}\rangle$.

Proof. Let G denote the right-hand side. By Proposition 2.5.44, $G \leq \mathrm{Sp}(\mathcal{V})$.
If $\mathbf{x}, \mathbf{y} \neq \mathbf{0}$, $\langle \mathbf{x}, \mathbf{y}\rangle \neq 0$, then $\tau_{\mathbf{y}-\mathbf{x},\,\langle \mathbf{x},\mathbf{y}\rangle^{-1}}(\mathbf{x}) = \mathbf{y}$. If $\langle \mathbf{x}, \mathbf{y}\rangle = 0$, choose $\mathbf{z}$ with $\langle \mathbf{x}, \mathbf{z}\rangle\langle \mathbf{z}, \mathbf{y}\rangle \neq 0$, $\mathbf{x} \mapsto \mathbf{z} \mapsto \mathbf{y}$.
Hence G is transitive on nonzero vectors; the same argument in $\mathbf{u}^\perp/\mathbb{F}\mathbf{u}$ gives transitivity on hyperbolic pairs.
Let $\sigma \in \mathrm{Sp}(\mathcal{V})$ and induct on $n = \dim \mathcal{V}$. For $n = 2$, hyperbolic-pair transitivity gives $\sigma \in G$.
For $n > 2$, choose a hyperbolic plane $\mathcal{H} = \mathrm{span}(\mathbf{u}, \mathbf{v})$. Pair transitivity gives $\gamma \in G$ with $\gamma\sigma|_{\mathcal{H}} = I_{\mathcal{H}}$.
Then $\gamma\sigma$ restricts to an element of $\mathrm{Sp}(\mathcal{H}^\perp)$, which is a product of transvections by induction; extended by the identity on $\mathcal{H}$, these are transvections of $\mathcal{V}$. Thus $\sigma \in G$.
Structural Commentary: Symplectic transformations are built from elementary shears along null directions.
Key Insight: Hyperbolic pairs are the movable frames of symplectic geometry. □

Proposition 2.5.46. *Symplectic Transformations Have Determinant One*
If $\mathcal{V}$ *is nonsingular symplectic, then* $\mathrm{Sp}(\mathcal{V}) \subseteq \mathrm{SL}(\mathcal{V})$.

Proof. By Theorem 2.5.45, it suffices to show $\det(\tau_{\mathbf{v},a}) = 1$. Let $\tau_{\mathbf{v},a} = I + a\,\mathbf{v} \otimes \lambda_{\mathbf{v}}$, $\lambda_{\mathbf{v}}(\mathbf{x}) = \langle \mathbf{x}, \mathbf{v}\rangle$.
Since $\lambda_{\mathbf{v}}(\mathbf{v}) = \langle \mathbf{v}, \mathbf{v}\rangle = 0$, $\tau_{\mathbf{v},a}$ fixes $\mathbb{F}\mathbf{v}$ and induces I on $\mathcal{V}/\mathbb{F}\mathbf{v}$. Hence its matrix in any basis starting with $\mathbf{v}$ is unipotent triangular, so $\det(\tau_{\mathbf{v},a}) = 1$.
Structural Commentary: Generators rule out $\det = -1$.
Key Insight: Symplectic transvections have determinant 1, and they generate the symplectic group. □

Theorem 2.5.47. *Orthogonal Bases and Diagonal Congruence*
Let $(\mathcal{V}, B)$ *be an orthogonal geometry,* $q(\mathbf{x}) := B(\mathbf{x}, \mathbf{x})$, $n = \dim \mathcal{V}$, *and* $r = \mathrm{rank}\, B$.
Then $\exists$ *basis* β *such that* $M_\beta = \mathrm{diag}(a_1, \ldots, a_r, 0_{n-r})$, $\quad a_i \in \mathbb{F}^\times \iff \neg\big(\mathrm{char}(\mathbb{F}) = 2 \ \wedge\ q \equiv 0 \ \wedge\ B \neq 0\big)$.
Equivalently, a symmetric matrix is congruent to a diagonal matrix iff it is not a nonzero alternating matrix in characteristic 2.

Proof.

(⇒) If M_β is diagonal and $\mathrm{char}(\mathbb{F}) = 2$ with $q \equiv 0$, then all diagonal entries of M_β vanish. Hence $M_\beta = 0$, so $B = 0$.
(⇐) If $B = 0$, any basis works. Otherwise, the hypothesis and Proposition 2.5.6 give $q \not\equiv 0$.
Peel off nonisotropic lines, where $\mathcal{V} = \mathbb{F}\mathbf{u}_1 \odot \cdots \odot \mathbb{F}\mathbf{u}_s \odot \mathcal{W}$, $q(\mathbf{u}_i) = a_i \in \mathbb{F}^\times$, $q|_{\mathcal{W}} \equiv 0$.
If $\mathrm{char}(\mathbb{F}) \neq 2$, Proposition 2.5.6 gives $B|_{\mathcal{W}} = 0$, so append any basis of $\mathcal{W}$.
If $\mathrm{char}(\mathbb{F}) = 2$, then $\mathcal{W} = R \odot \mathcal{H}_1 \odot \cdots \odot \mathcal{H}_t$ by Theorem 2.5.41.
For each hyperbolic plane $\mathcal{H}_j = \mathrm{span}(\mathbf{e}, \mathbf{f})$, rescale so $B(\mathbf{e}, \mathbf{f}) = a = q(\mathbf{u})$ for some current nonisotropic $\mathbf{u}$.
Then $\mathbb{F}\mathbf{u} \odot \mathcal{H}_j = \mathbb{F}(\mathbf{u} + \mathbf{e}) \odot \mathbb{F}(\mathbf{u} + \mathbf{f}) \odot \mathbb{F}(\mathbf{u} + \mathbf{e} + \mathbf{f})$ is orthogonal with all three norms a.
Absorb all $\mathcal{H}_j$ and append any basis of R.

Structural Commentary: Only nonzero alternating characteristic-2 blocks obstruct diagonals.
Key Insight: Peel off nonisotropic lines until only the radical or characteristic-2 alternating residue remains. □

Example 2.5.48. *Characteristic Two Orthogonal-Symplectic Plane*
Assume $\mathrm{char}(\mathbb{F}) = 2$. Let $\mathcal{H} = \mathbb{F}\mathbf{e} \oplus \mathbb{F}\mathbf{f}$ with Gram matrix $\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$.
Then B is symmetric, alternating, and nonsingular. For every $a, b \in \mathbb{F}$, $B(a\mathbf{e} + b\mathbf{f}, a\mathbf{e} + b\mathbf{f}) = 2ab = 0$, so $\mathcal{H}$ is

both orthogonal and symplectic. If $0 \neq \mathbf{x} \in \mathcal{H}$, then $\{\mathbf{x}\}^\perp = \mathbb{F}\mathbf{x}$; hence no two independent nonzero vectors are orthogonal. Thus $\mathcal{H}$ has no orthogonal basis, the obstruction excluded in Theorem 2.5.47.

Method 2.5.49. *Lagrange Orthogonalization*
Let $\mathcal{V}$ be an orthogonal geometry with $\operatorname{char}(\mathbb{F}) \neq 2$, and set $q(\mathbf{x}) = \langle \mathbf{x}, \mathbf{x} \rangle$.
From a generating list $L = (\mathbf{v}_1, \dots, \mathbf{v}_n)$, construct an orthogonal basis as follows.

(i) In the current summand with generators $(\mathbf{x}_i)$, choose a pivot $\mathbf{u}$ with $q(\mathbf{u}) \neq 0$; if none exists, stop with a totally degenerate remainder. If all $\mathbf{x}_i$ are isotropic but the summand is not totally degenerate, take $\mathbf{u} = \mathbf{x}_i + \mathbf{x}_j$ with $\langle \mathbf{x}_i, \mathbf{x}_j \rangle \neq 0$.
(ii) Given mutually orthogonal pivots $\mathbf{u}_1, \dots, \mathbf{u}_k$, let $S_k = \operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_k)$.
Replace each generator $\mathbf{x}_i$ by $\mathbf{w}_i = \mathbf{x}_i - \sum_{j=1}^{k} \frac{\langle \mathbf{x}_i, \mathbf{u}_j \rangle}{q(\mathbf{u}_j)} \mathbf{u}_j$. Then the nonzero $\mathbf{w}_i$ generate $S_k^\perp$.
(iii) Iterate inside $S_k^\perp$ until the remaining summand R is totally degenerate.
Output $\beta = (\mathbf{u}_1, \dots, \mathbf{u}_r, \mathbf{z}_1, \dots, \mathbf{z}_m)$, where $(\mathbf{z}_1, \dots, \mathbf{z}_m)$ is any basis of R.

Only nonisotropic pivots are allowed, since the projection step divides by $q(\mathbf{u}_j) \neq 0$.

Theorem 2.5.50. *Canonical Orthogonal Forms over Common Fields*
Let $M = M^t \in M_n(\mathbb{F})$, let $r = \operatorname{rank} M$, and write $M \sim N$ for congruence. If $\operatorname{char}(\mathbb{F}) = 2$, assume M is not nonzero alternating.

(i) If $\mathbb{F}$ is algebraically closed, then $M \sim I_r \oplus 0_{n-r}$; r is complete.
(ii) If $\mathbb{F} = \mathbb{R}$, then $M \sim I_p \oplus (-I_m) \oplus 0_k$ with $p + m + k = n$; (p, m) is complete.
(iii) If $\mathbb{F} = \mathbb{F}_q$ and $\operatorname{char}(\mathbb{F}) = 2$, then $M \sim I_r \oplus 0_{n-r}$.
(iv) If $\mathbb{F} = \mathbb{F}_q$ has odd characteristic and $d \in \mathbb{F}_q^\times$ is a fixed nonsquare, then for $r > 0$, $M \sim I_r \oplus 0_{n-r}$ or $M \sim I_{r-1} \oplus [d] \oplus 0_{n-r}$ with exactly one alternative; for $r = 0$, $M \sim 0_n$.

Proof.

(i) By Theorem 2.5.47, write $M \sim \operatorname{diag}(a_1, \dots, a_r, 0_{n-r})$ with $a_i \neq 0$.
If $\mathbb{F}$ is algebraically closed, $a_i = s_i^2$, so scaling by s_i^{-1} gives $M \sim I_r \oplus 0_{n-r}$.
(ii) Write $M \sim \operatorname{diag}(a_1, \dots, a_r, 0_k)$ over $\mathbb{R}$. Rescaling gives $a_i = \pm 1$, hence $M \sim I_p \oplus (-I_m) \oplus 0_k$.
Uniqueness of (p, m) is Sylvester's inertia theorem, Theorem 4.1.75.
(iii) Write $M \sim \operatorname{diag}(a_1, \dots, a_r, 0_{n-r})$ over $\mathbb{F}_q$ with $\operatorname{char}(\mathbb{F}_q) = 2$.
Since $x \mapsto x^2$ is an automorphism, each a_i is a square; hence $M \sim I_r \oplus 0_{n-r}$.
(iv) Write $M \sim \operatorname{diag}(a_1, \dots, a_r, 0_{n-r})$ over odd $\mathbb{F}_q$. Since $\mathbb{F}_q^\times / (\mathbb{F}_q^\times)^2 = \{1, d\}$, the nonsingular diagonal part is determined by its discriminant square class $\delta(M) := a_1 \cdots a_r \pmod{(\mathbb{F}_q^\times)^2}$, giving I_r when $\delta(M) = 1$ and $I_{r-1} \oplus [d]$ when $\delta(M) = d$. Appending the zero block gives the stated forms, and δ separates them.

Structural Commentary: Once diagonalized, congruence classification becomes square-class bookkeeping.
Key Insight: The base field determines how many diagonal coefficients can be normalized to 1. □

Definition 2.5.51. *Universal Orthogonal Geometry*
An orthogonal geometry $\mathcal{V}$ over $\mathbb{F}$ is *universal* if its norm map $q(\mathbf{x}) := \langle \mathbf{x}, \mathbf{x} \rangle$ represents every scalar $q(\mathcal{V}) = \mathbb{F}$.

Proposition 2.5.52. *Finite-Field Square Classes and Universality*
Let $\mathbb{F}_q$ be a finite field and let $S = (\mathbb{F}_q^\times)^2$.

(i) If $\operatorname{char}(\mathbb{F}_q) = 2$, then every element of $\mathbb{F}_q$ is a square.
(ii) If $\operatorname{char}(\mathbb{F}_q) \neq 2$, then $|S| = (q-1)/2$ and $\mathbb{F}_q^\times \setminus S = dS$ for any nonsquare d.
(iii) If $\mathcal{V}$ is orthogonal over odd $\mathbb{F}_q$ and contains a nonsingular subspace of dimension ≥ 2, then $\mathcal{V}$ is universal.

Proof.

(i) The map $\varphi : \mathbb{F}_q \to \mathbb{F}_q$, $\varphi(x) = x^2$, is the Frobenius automorphism in characteristic 2.
(ii) For $\operatorname{char}(\mathbb{F}_q) \neq 2$, the homomorphism $\varphi : \mathbb{F}_q^\times \to \mathbb{F}_q^\times$, $\varphi(x) = x^2$, has kernel $\{\pm 1\}$.
Hence $|S| = (q-1)/2$, and the nonsquares are the other coset dS.
(iii) Let $q(\mathbf{x}) = \langle \mathbf{x}, \mathbf{x} \rangle$. By Theorem 2.5.47, the nonsingular subspace contains orthogonal nonisotropic $\mathbf{u}, \mathbf{v}$; set $a = q(\mathbf{u})$, $b = q(\mathbf{v})$. For fixed $c \in \mathbb{F}_q$, the sets $A = a\{\alpha^2 : \alpha \in \mathbb{F}_q\}$ and $B = c - b\{\beta^2 : \beta \in \mathbb{F}_q\}$ both have $(q+1)/2$ elements, so $A \cap B \neq \varnothing$. Thus $c = a\alpha^2 + b\beta^2 = q(\alpha\mathbf{u} + \beta\mathbf{v})$ for some $\alpha, \beta \in \mathbb{F}_q$.

Structural Commentary: Over finite fields, cardinality replaces order arguments.
Key Insight: Two half-sized subsets of a finite field must overlap. □

Example 2.5.53. *Two Rank-Two Orthogonal Planes over $\mathbb{F}_5$*
In $\mathbb{F}_5$, $(\mathbb{F}_5^\times)^2 = \{1, 4\}$ and the nonsquare class is $\{2, 3\}$. Let $q_1(x, y) = x^2 + y^2$, $q_2(x, y) = x^2 + 2y^2$.
Then q_1 has Gram matrix I_2 and is hyperbolic, since $q_1(1, 2) = 1 + 4 = 0$.

The form q_2 has Gram matrix $\mathrm{diag}(1,2)$ and is anisotropic: if $q_2(x,y)=0$ with $y\neq 0$, then $\left(\frac{x}{y}\right)^2=-2=3$, impossible since 3 is nonsquare; if $y=0$, then $x=0$.
Thus I_2 and $\mathrm{diag}(1,2)$ realize the two nonsingular rank-2 classes in Theorem 2.5.50.
Structural Commentary: Odd finite-field classes are square-class classes.
Key Insight: In rank 2, the square class of the determinant decides whether the plane is hyperbolic or anisotropic.

Lemma 2.5.54. *Odd Finite-Field Isotropy Bound*
Let $\mathbb{F}_q$ have odd characteristic, $\mathcal{V}$ be nonsingular orthogonal with $q(\mathbf{x}):=\langle \mathbf{x},\mathbf{x}\rangle$. Let $\mathcal{H}$ be hyperbolic plane.

(i) If $\dim \mathcal{V}\geq 3$, then $\mathcal{V}$ is isotropic.
(ii) $\mathcal{V}\cong \mathcal{H}^{\odot s}\odot\mathcal{R}$, where $\mathcal{R}$ is nonsingular anisotropic and $\dim\mathcal{R}\leq 2$.

Proof.

(i) By Theorem 2.5.47, choose a nonsingular plane $\mathcal{P}\subset\mathcal{V}$ and $0\neq\mathbf{z}\in\mathcal{P}^{\perp}$. If $q(\mathbf{z})=0$, then $\mathbf{z}$ is isotropic. Otherwise Proposition 2.5.52 gives $\mathbf{p}\in\mathcal{P}$ with $q(\mathbf{p})=-q(\mathbf{z})$, so $\mathbf{p}+\mathbf{z}$ is isotropic.
(ii) Induct on $\dim\mathcal{V}$. If $\mathcal{V}$ is anisotropic, take $s=0$ and $\mathcal{R}=\mathcal{V}$; then (i) gives $\dim\mathcal{R}\leq 2$. Otherwise Theorem 2.5.33 gives $\mathcal{V}\cong\mathcal{H}\odot\mathcal{W}$ with $\mathcal{W}$ nonsingular and $\dim\mathcal{W}=\dim\mathcal{V}-2$, so induction gives the decomposition.

Structural Commentary: Odd finite fields force isotropy in dimension ≥ 3.
Key Insight: Hyperbolic splitting leaves only a residue of dimension at most 2. □

Theorem 2.5.55. *Odd Finite-Field Orthogonal Decomposition*
Let $\mathbb{F}_q$ have odd characteristic, fix a nonsquare $d\in\mathbb{F}_q^{\times}$, and let $\mathcal{V}$ be nonsingular orthogonal. Write $\mathcal{H}$ for a hyperbolic plane, $\langle a\rangle$ for a line of value a, and $\mathcal{A}_d:=\langle 1\rangle\odot\langle -d\rangle$.

(i) If $\dim\mathcal{V}=2m+1$, then $\exists\, a\in\mathbb{F}_q^{\times}$ such that $\mathcal{V}\cong\mathcal{H}^{\odot m}\odot\langle a\rangle$, and $a(\mathbb{F}_q^{\times})^2=(-1)^m\Delta(\mathcal{V})$.
(ii) If $\dim\mathcal{V}=2m$, then exactly one holds: $\mathcal{V}\cong\mathcal{H}^{\odot m}\iff\Delta(\mathcal{V})=(-1)^m(\mathbb{F}_q^{\times})^2$ or $\mathcal{V}\cong\mathcal{H}^{\odot(m-1)}\odot\mathcal{A}_d\iff\Delta(\mathcal{V})=(-1)^m d(\mathbb{F}_q^{\times})^2$.

Proof.

(i) By Lemma 2.5.54, $\mathcal{V}\cong\mathcal{H}^{\odot m}\odot\mathcal{R}$ with $\dim\mathcal{R}=1$, so $\mathcal{R}\cong\langle a\rangle$.
Since $\Delta(\mathcal{H})=-(\mathbb{F}_q^{\times})^2$, $\Delta(\mathcal{V})=(-1)^m a(\mathbb{F}_q^{\times})^2$, hence $a(\mathbb{F}_q^{\times})^2=(-1)^m\Delta(\mathcal{V})$.
(ii) By Lemma 2.5.54, $\mathcal{V}\cong\mathcal{H}^{\odot s}\odot\mathcal{R}$ with $\dim\mathcal{R}=0$ or 2. If $\dim\mathcal{R}=0$, then $s=m$ and $\Delta(\mathcal{V})=(-1)^m(\mathbb{F}_q^{\times})^2$. If $\dim\mathcal{R}=2$, then $s=m-1$ and Proposition 2.5.34 shows that the anisotropic residue is the nonsquare class $\mathcal{A}_d$, with $\Delta(\mathcal{A}_d)=-d(\mathbb{F}_q^{\times})^2$; hence $\Delta(\mathcal{V})=(-1)^m d(\mathbb{F}_q^{\times})^2$.

Structural Commentary: The anisotropic residue has dimension ≤ 2.
Key Insight: Parity and discriminant determine the residue. □

Example 2.5.56. *Quadratic Refinements over $\mathbb{F}_2$*
On $V=\mathbb{F}_2^2$, set $B((x,y),(u,v))=xv+yu$, $Q_0(x,y)=xy$, $Q_1(x,y)=x^2+xy+y^2$. Then $B_{Q_0}=B_{Q_1}=B$, the unique nonsingular alternating plane over $\mathbb{F}_2$. However, $Q_0^{-1}(0)\setminus\{\mathbf{0}\}=\{(1,0),(0,1)\}$, $Q_1^{-1}(0)=\{\mathbf{0}\}$.
Hence $Q_0\ncong Q_1$ although they have the same polar form.
Structural Commentary: One polar form, two refinements.
Key Insight: Track Q, not only B_Q.

Definition 2.5.57. *Arf Invariant over $\mathbb{F}_2$*
Let $Q:\mathcal{V}\to\mathbb{F}_2$ be quadratic.

(i) Its *polar form* is $B_Q(\mathbf{x},\mathbf{y})=Q(\mathbf{x}+\mathbf{y})+Q(\mathbf{x})+Q(\mathbf{y})$ and is assumed nonsingular.
(ii) For a *symplectic basis* $\beta=(\mathbf{u}_1,\mathbf{v}_1,\dots,\mathbf{u}_m,\mathbf{v}_m)$ of B_Q,
the *Arf invariant* is $\mathrm{Arf}(Q):=\sum_{i=1}^{m}Q(\mathbf{u}_i)Q(\mathbf{v}_i)\in\mathbb{F}_2$, independent of β.

Theorem 2.5.58. *Quadratic Refinements over $\mathbb{F}_2$*
Let $Q:\mathcal{V}\to\mathbb{F}_2$ be quadratic with nonsingular polar form B_Q. Then $\dim\mathcal{V}=2m$, and $Q\cong Q_{\mathrm{Arf}(Q)}$.
Here $Q_0(x_1,y_1,\dots,x_m,y_m):=\sum_{i=1}^{m}x_iy_i$ and, for $m\geq 1$, $Q_1(x_1,y_1,\dots,x_m,y_m):=x_1^2+x_1y_1+y_1^2+\sum_{i=2}^{m}x_iy_i$.
Q_0 has Arf invariant 0, and Q_1 has Arf invariant 1.

Proof. As B_Q nonsingular alternating, Theorem 2.5.41 gives symplectic basis $(\mathbf{u}_1,\mathbf{v}_1,\dots,\mathbf{u}_m,\mathbf{v}_m)$, $\dim\mathcal{V}=2m$.
On each hyperbolic pair $(\mathbf{u},\mathbf{v})$, if $Q(\mathbf{u})Q(\mathbf{v})=0$, a change of pair gives the refined plane xy.
If $Q(\mathbf{u})=Q(\mathbf{v})=1$, the plane is x^2+xy+y^2.
Thus Q is an orthogonal sum of refined planes of types $H_Q:=xy$ and $A_Q:=x^2+xy+y^2$.
Since $A_Q\odot A_Q\cong H_Q\odot H_Q$, all pairs of anisotropic refinements cancel.

Remaining parity is $\operatorname{Arf}(Q) = \sum_{i=1}^{m} Q(\mathbf{u}_i)Q(\mathbf{v}_i) \in \mathbb{F}_2$, so $Q \cong Q_0$ when $\operatorname{Arf}(Q) = 0$ and $Q \cong Q_1$ when $\operatorname{Arf}(Q) = 1$.
Structural Commentary: In characteristic 2, the same symplectic polar form supports two quadratic geometries.
Key Insight: The Arf invariant is the missing one-bit datum that the polar bilinear form cannot see. □

Definition 2.5.59. *Isometries and Classical Groups*
Let $\mathcal{V}, \mathcal{W}$ be metric vector spaces.

(i) A bijective linear map $\sigma : \mathcal{V} \to \mathcal{W}$ is an *isometry* if $\langle \sigma\mathbf{x}, \sigma\mathbf{y}\rangle = \langle \mathbf{x}, \mathbf{y}\rangle \ \ \forall \mathbf{x}, \mathbf{y} \in \mathcal{V}$.
(ii) The group of isometries $\mathcal{V} \to \mathcal{V}$ is denoted $\operatorname{Iso}(\mathcal{V})$. If $\mathcal{V}$ is nonsingular orthogonal, this is the *orthogonal group* $\mathcal{O}(\mathcal{V})$; if $\mathcal{V}$ is nonsingular symplectic, this is the *symplectic group* $\operatorname{Sp}(\mathcal{V})$.

Proposition 2.5.60. *Matrix Criterion for Isometries*
Let β be a basis of a metric vector space $\mathcal{V}$, let $M = M_\beta$, and let $A = [\sigma]_{\beta \leftarrow \beta}$.
Then $\sigma \in \operatorname{Iso}(\mathcal{V}) \iff \sigma$ is bijective and $A^t M A = M$. If M is nonsingular, then $\det A = \pm 1$.

Proof.

($\Rightarrow$) If σ is an isometry, then it is bijective and, for all coordinate vectors X, Y, $X^t A^t M A Y = (AX)^t M (AY) = X^t M Y$. Hence $A^t M A = M$.
($\Leftarrow$) If σ is bijective and $A^t M A = M$, then for all $\mathbf{x}, \mathbf{y} \in \mathcal{V}$ with coordinate vectors X, Y, $\langle \sigma\mathbf{x}, \sigma\mathbf{y}\rangle = (AX)^t M (AY) = X^t M Y = \langle \mathbf{x}, \mathbf{y}\rangle$, so σ is an isometry.

If M is invertible, taking determinants in $A^t M A = M$ gives $(\det A)^2 = 1$, hence $\det A = \pm 1$.
Structural Commentary: Isometries are exactly the coordinate changes that stabilize the representing matrix.
Key Insight: Orthogonal and symplectic groups are stabilizer groups under congruence. □

Proposition 2.5.61. *Fixed Spaces of Isometries*
Let $\mathcal{V}$ be nonsingular and let $\tau \in \operatorname{Iso}(\mathcal{V})$. Then $\ker(I - \tau) = \operatorname{im}(I - \tau)^\perp$, and $\dim \ker(I - \tau) = \dim \operatorname{im}(I - \tau)^\perp$.

Proof. If $\mathbf{x}$ is fixed by τ, then for every $\mathbf{y}$, $\langle \mathbf{x}, (I - \tau)\mathbf{y}\rangle = \langle \mathbf{x}, \mathbf{y}\rangle - \langle \tau\mathbf{x}, \tau\mathbf{y}\rangle = 0$, so $\ker(I - \tau) \subseteq \operatorname{im}(I - \tau)^\perp$.
The dimensions agree by rank-nullity and Theorem 2.5.26, hence equality holds.
Structural Commentary: The fixed space of an isometry is the orthogonal complement of its displacement space.
Key Insight: Fixed vectors are exactly vectors orthogonal to all displacements. □

Proposition 2.5.62. *Symmetries and Fixed Hyperplanes*
Let $\mathcal{V}$ be nonsingular orthogonal with $\operatorname{char}(\mathbb{F}) \neq 2$, and let $\tau \in \mathcal{O}(\mathcal{V})$ with $\tau \neq I$. Then $\tau = \sigma_{\mathbf{u}}$ for some nonisotropic $\mathbf{u} \iff \ker(I - \tau)$ is a nonsingular hyperplane. In this case $\ker(I - \sigma_{\mathbf{u}}) = \operatorname{span}(\mathbf{u})^\perp$.

Proof.

($\Rightarrow$) The formula for $\sigma_{\mathbf{u}}$ gives $\sigma_{\mathbf{u}}\mathbf{u} = -\mathbf{u}$ and fixes $\operatorname{span}(\mathbf{u})^\perp$ pointwise. Hence $\ker(I - \sigma_{\mathbf{u}}) = \operatorname{span}(\mathbf{u})^\perp$, a nonsingular hyperplane.
($\Leftarrow$) Let $\mathcal{H} = \ker(I - \tau)$ be a nonsingular hyperplane and write $\mathcal{V} = \mathcal{H} \odot \mathcal{L}$ with $\dim \mathcal{L} = 1$. Since τ fixes $\mathcal{H}$ and is an isometry, $\tau\mathcal{L} = \mathcal{L}$. Thus $\tau|_{\mathcal{L}} = \lambda I$ with $\lambda^2 = 1$; as $\tau \neq I$ and $\operatorname{char}(\mathbb{F}) \neq 2$, $\lambda = -1$. For $0 \neq \mathbf{u} \in \mathcal{L}$, $\tau = \sigma_{\mathbf{u}}$.

Structural Commentary: Symmetries are exactly fixed-hyperplane isometries.
Key Insight: A symmetry fixes a nonsingular hyperplane and flips its orthogonal line. □

Proposition 2.5.63. *No Unipotent Isometries in Anisotropic Spaces*
Let $\mathcal{V}$ be nonsingular anisotropic. If $\tau \in \operatorname{Iso}(\mathcal{V})$ is unipotent, then $\tau = I$.

Proof. Put $N = \tau - I$. Then N is nilpotent and, by Proposition 2.5.25, $N^* = \tau^* - I = \tau^{-1} - I \in N\mathbb{F}[N]$.
If $N \neq 0$, choose maximal $k \geq 1$ with $N^k \neq 0$ and choose $\mathbf{x}$ with $\mathbf{u} = N^k\mathbf{x} \neq 0$.
Since $N^{k+1} = 0$ and $(N^*)^k \in N^k\mathbb{F}[N]$, $\langle \mathbf{u}, \mathbf{u}\rangle = \langle N^k\mathbf{x}, N^k\mathbf{x}\rangle = \langle \mathbf{x}, (N^*)^k N^k \mathbf{x}\rangle = 0$, contradicting anisotropy.
Hence $N = 0$ and $\tau = I$.
Structural Commentary: Top nilpotent displacements are isotropic.
Key Insight: An anisotropic form leaves no room for unipotent shear. □

Definition 2.5.64. *Rotations, Reflections, and Symmetries*
Let $\mathcal{V}$ be a nonsingular orthogonal geometry.

(i) For $\sigma \in \mathcal{O}(\mathcal{V})$, $\det \sigma = 1$ means σ is a *rotation*, and $\det \sigma = -1$ means σ is a *reflection*.
(ii) If $\operatorname{char}(\mathbb{F}) \neq 2$ and $q(\mathbf{u}) = \langle \mathbf{u}, \mathbf{u}\rangle \neq 0$, the *symmetry* in the line $\mathbb{F}\mathbf{u}$ is $\sigma_{\mathbf{u}}(\mathbf{x}) = \mathbf{x} - 2\frac{\langle \mathbf{x}, \mathbf{u}\rangle}{\langle \mathbf{u}, \mathbf{u}\rangle}\mathbf{u}$.
It satisfies $\sigma_{\mathbf{u}}\mathbf{u} = -\mathbf{u}$ and $\sigma_{\mathbf{u}}|_{\operatorname{span}(\mathbf{u})^\perp} = I$.

Proposition 2.5.65. *Matching Equal-Length Vectors by a Symmetry*
Let $\mathcal{V}$ be nonsingular orthogonal with $\mathrm{char}(\mathbb{F}) \neq 2$, and set $q(\mathbf{x}) = \langle \mathbf{x}, \mathbf{x} \rangle$.
If $q(\mathbf{u}) = q(\mathbf{v}) \neq 0$, then $\exists$ a symmetry σ with $\sigma\mathbf{u} = \mathbf{v}$ or $\sigma\mathbf{u} = -\mathbf{v}$.

Proof. Let $\mathbf{a} = \mathbf{u} - \mathbf{v}$ and $\mathbf{b} = \mathbf{u} + \mathbf{v}$. Then $\langle \mathbf{a}, \mathbf{b} \rangle = q(\mathbf{u}) - q(\mathbf{v}) = 0$.
They cannot both be isotropic: if $q(\mathbf{a}) = q(\mathbf{b}) = 0$, then $q(2\mathbf{u}) = 0$, contradicting $q(\mathbf{u}) \neq 0$.
If $q(\mathbf{a}) \neq 0$, then $\sigma_{\mathbf{a}}\mathbf{u} = \mathbf{v}$; if $q(\mathbf{b}) \neq 0$, then $\sigma_{\mathbf{b}}\mathbf{u} = -\mathbf{v}$.
Structural Commentary: One symmetry matches equal-length lines.
Key Insight: Use $\mathbf{u} \pm \mathbf{v}$ as the reflecting direction. □

Theorem 2.5.66. *Orthogonal Transformations Are Products of Symmetries*
Let $\mathcal{V}$ be nonsingular orthogonal with $\mathrm{char}(\mathbb{F}) \neq 2$, and set $q(\mathbf{x}) := \langle \mathbf{x}, \mathbf{x} \rangle$. Then $\mathcal{O}(\mathcal{V}) = \langle \sigma_{\mathbf{u}} : q(\mathbf{u}) \neq 0 \rangle$.

Proof. Induct on $n = \dim \mathcal{V}$.
If $n = 0$, this is trivial; if $n = 1$, any $\tau \in \mathcal{O}(\mathcal{V})$ acts on a nonisotropic line by ± 1, hence is I or the symmetry of that line. For $n > 1$, choose $\mathbf{u}$ with $q(\mathbf{u}) \neq 0$. Since $q(\tau\mathbf{u}) = q(\mathbf{u})$, Proposition 2.5.65 gives a symmetry σ with $\sigma\tau\mathbf{u} = \mathbf{u}$ or $-\mathbf{u}$. Composing further with $\sigma_{\mathbf{u}}$ if necessary, obtain a product ρ of symmetries such that $\rho\tau$ fixes $\mathbf{u}$. Hence $\rho\tau$ preserves $\mathrm{span}(\mathbf{u})^{\perp}$, and its restriction lies in $\mathcal{O}(\mathrm{span}(\mathbf{u})^{\perp})$.
By induction, this restriction is a product of symmetries in $\mathrm{span}(\mathbf{u})^{\perp}$, extended by I on $\mathbb{F}\mathbf{u}$; these are symmetries of $\mathcal{V}$. Thus τ is a product of symmetries.
Structural Commentary: Symmetries are the orthogonal analogue of elementary generators.
Key Insight: Match one nonisotropic line, then reduce to its orthogonal complement. □

Theorem 2.5.67. *Witt Theorems for Orthogonal Geometries*
Let $\mathcal{V} \cong \mathcal{V}'$ be nonsingular orthogonal geometries over $\mathbb{F}$ with $\mathrm{char}(\mathbb{F}) \neq 2$, and write $q(\mathbf{x}) := \langle \mathbf{x}, \mathbf{x} \rangle$.
If $\mathcal{S} \subseteq \mathcal{V}$ and $\mathcal{T} \subseteq \mathcal{V}'$ are nonsingular and $\tau : \mathcal{S} \to \mathcal{T}$ is an isometry, then for $\mathcal{S}^{\perp} \cong \mathcal{T}^{\perp}$ $\exists$ Φ such that $\mathcal{V} \to \mathcal{V}'$ isometric with $\Phi|_{\mathcal{S}} = \tau$.

Proof. By Proposition 2.5.37, it suffices to prove cancellation. Transport along a fixed isometry $\mathcal{V} \cong \mathcal{V}'$ and assume $\mathcal{V}' = \mathcal{V}$. Thus $\mathcal{V} = \mathcal{S} \odot \mathcal{S}^{\perp} = \mathcal{T} \odot \mathcal{T}^{\perp}$, with $\tau : \mathcal{S} \to \mathcal{T}$ an isometry.
Induct on $r = \dim \mathcal{S}$. If $r = 0$, the claim is trivial. If $r = 1$, write $\mathcal{S} = \mathbb{F}\mathbf{s}$ with $q(\mathbf{s}) \neq 0$. Since $q(\tau\mathbf{s}) = q(\mathbf{s})$, Proposition 2.5.65 gives $\sigma \in \mathcal{O}(\mathcal{V})$ with $\sigma\mathcal{S} = \mathcal{T}$; hence $\sigma(\mathcal{S}^{\perp}) = \mathcal{T}^{\perp}$.
For $r > 1$, choose $\mathbf{s} \in \mathcal{S}$ with $q(\mathbf{s}) \neq 0$ and write $\mathcal{S} = \mathbb{F}\mathbf{s} \odot \mathcal{U}$. By case $r = 1$, choose $\sigma \in \mathcal{O}(\mathcal{V})$ with $\sigma(\mathbb{F}\mathbf{s}) = \mathbb{F}\tau\mathbf{s}$. Inside $\mathcal{W} := (\mathbb{F}\tau\mathbf{s})^{\perp}$, $\mathcal{W} = \sigma\mathcal{U} \odot \sigma(\mathcal{S}^{\perp}) = \tau\mathcal{U} \odot \mathcal{T}^{\perp}$, $\sigma\mathcal{U} \cong \tau\mathcal{U}$. Since $\dim \mathcal{U} < r$, induction in $\mathcal{W}$ gives $\sigma(\mathcal{S}^{\perp}) \cong \mathcal{T}^{\perp}$. Thus $\mathcal{S}^{\perp} \cong \mathcal{T}^{\perp}$. Cancellation follows, and extension follows from Proposition 2.5.37.
Structural Commentary: Orthogonal Witt cancellation reduces one nonisotropic line at a time.
Key Insight: Symmetries align the first line; induction handles the remaining orthogonal complement. □

Definition 2.5.68. *Witt Index*
Let $\mathcal{V}$ be nonsingular orthogonal with $\mathrm{char}(\mathbb{F}) \neq 2$. The *Witt index* is $w(\mathcal{V}) := \max\{\dim \mathcal{L} : \mathcal{L} \leq \mathcal{V},\ \mathcal{L} \subseteq \mathcal{L}^{\perp}\}$.

Theorem 2.5.69. *Witt Index and Anisotropic Decomposition*
Let $\mathcal{V}$ be orthogonal over $\mathbb{F}$ with $\mathrm{char}(\mathbb{F}) \neq 2$. Then:

(i) if $\mathcal{V}$ is nonsingular, every maximal totally isotropic subspace has dimension $w(\mathcal{V})$;
(ii) if $\mathcal{V}$ is nonsingular, every maximal hyperbolic subspace has dimension $2w(\mathcal{V})$;
(iii) for $R = \mathrm{rad}(\mathcal{V})$ and a nonsingular complement $\mathcal{U}$ of R, there exist $\mathcal{H} \leq \mathcal{U}$ hyperbolic and $\mathcal{S} \leq \mathcal{U}$ anisotropic such that $\mathcal{V} = \mathcal{S} \odot \mathcal{H} \odot R$, $\dim \mathcal{H} = 2w(\mathcal{U})$.

Proof.

(i) Let $\mathcal{L}$ be maximal totally isotropic and choose $\mathcal{M}$ totally isotropic with $\dim \mathcal{M} = w(\mathcal{V})$.
Since $\dim \mathcal{L} \leq w(\mathcal{V})$, choose $\mathcal{M}_0 \leq \mathcal{M}$ with $\dim \mathcal{M}_0 = \dim \mathcal{L}$.
Any isomorphism $\mathcal{L} \to \mathcal{M}_0$ is an isometry, so Theorem 2.5.38 extends it to $\Phi \in \mathcal{O}(\mathcal{V})$.
Then $\mathcal{L} \subseteq \Phi^{-1}\mathcal{M}$; maximality gives $\mathcal{L} = \Phi^{-1}\mathcal{M}$, hence $\dim \mathcal{L} = w(\mathcal{V})$.

(ii) If $\mathcal{H}$ has k hyperbolic planes, then $\mathcal{H}$ contains a k-dimensional totally isotropic subspace, so $k \leq w(\mathcal{V})$.
Conversely, Theorem 2.5.35 completes a $w(\mathcal{V})$-dimensional totally isotropic subspace to a hyperbolic subspace $\mathcal{H}_0$ of dimension $2w(\mathcal{V})$. If a maximal hyperbolic $\mathcal{H}$ had dimension $2k < 2w(\mathcal{V})$, choose hyperbolic $\mathcal{H}_1 \leq \mathcal{H}_0$ with $\dim \mathcal{H}_1 = 2k$ and extend an isometry $\mathcal{H} \to \mathcal{H}_1$ to $\Phi \in \mathcal{O}(\mathcal{V})$; then $\Phi^{-1}\mathcal{H}_0$ is a larger hyperbolic subspace containing $\mathcal{H}$, contradiction.

(iii) By Theorem 2.5.28, $\mathcal{V} = R \odot \mathcal{U}$ with $\mathcal{U}$ nonsingular. Choose maximal hyperbolic $\mathcal{H} \leq \mathcal{U}$ and set $\mathcal{S} = \mathcal{H}^{\perp} \cap \mathcal{U}$.
Then $\dim \mathcal{H} = 2w(\mathcal{U})$ by (ii), and $\mathcal{U} = \mathcal{H} \odot \mathcal{S}$. If $\mathcal{S}$ had a nonzero isotropic vector, Theorem 2.5.33 would give a hyperbolic plane in $\mathcal{S}$, enlarging $\mathcal{H}$; hence $\mathcal{S}$ is anisotropic.

Structural Commentary: Orthogonal spaces split into radical, hyperbolic part, and anisotropic residue.
Key Insight: The Witt index counts how many hyperbolic planes can be split off. □

2.6 Affine and Projective Geometry

Affine geometry studies cosets of subspaces as geometric objects in their own right. Throughout this subsection, vector spaces are finite-dimensional over $\mathbb{F}$ unless explicitly stated otherwise.

Definition 2.6.1. *Flats* Let $\mathcal{S} \leq \mathcal{V}$ and $\mathbf{x} \in \mathcal{V}$, and set $\mathbf{x}+\mathcal{S} = \{\mathbf{x}+\mathbf{s} : \mathbf{s} \in \mathcal{S}\}$.

(i) The affine subset $\mathbf{x}+\mathcal{S}$ from Definition 1.2.26 is a *flat* in $\mathcal{V}$; $\mathcal{S}$ is its *base*, and $\mathbf{x}$ is a *flat representative*.
(ii) Let $\mathcal{A}(\mathcal{V})$ denote the set of all flats in $\mathcal{V}$, with $\dim(\mathcal{A}(\mathcal{V})) := \dim(\mathcal{V})$.
(iii) The *dimension* of $\mathbf{x}+\mathcal{S}$ is $\dim(\mathcal{S})$.
A 0-flat is a *point*, a 1-flat is a *line*, a 2-flat is a *plane*, and a codimension-one flat is a *hyperplane*.
(iv) Two flats $X = \mathbf{x}+\mathcal{S}$ and $Y = \mathbf{y}+\mathcal{T}$ are *parallel*, written $X \parallel Y$, if $\mathcal{S} \subseteq \mathcal{T}$ or $\mathcal{T} \subseteq \mathcal{S}$.

Structural Commentary: A flat has many representatives but one base: $\mathbf{x}+\mathcal{S} = \mathbf{y}+\mathcal{T} \Rightarrow \mathcal{S} = \mathcal{T}$.

Proposition 2.6.2. *Basic Incidence of Flats*
Let $X = \mathbf{x}+\mathcal{S}$ and $Y = \mathbf{y}+\mathcal{T}$ be flats in $\mathcal{V}$, and write $\mathbf{w}+X := \{\mathbf{w}+\mathbf{u} : \mathbf{u} \in X\}$.

(i) $(\exists\, \mathbf{w} \in \mathcal{V})(\mathbf{w}+X \subseteq Y) \iff \mathcal{S} \subseteq \mathcal{T}$.
(ii) $(\exists\, \mathbf{w} \in \mathcal{V})(\mathbf{w}+X = Y) \iff \mathcal{S} = \mathcal{T}$.
(iii) If $X \cap Y \neq \emptyset$, then $X \subseteq Y \iff \mathcal{S} \subseteq \mathcal{T}$.
(iv) If $X \cap Y \neq \emptyset$, then $X = Y \iff \mathcal{S} = \mathcal{T}$.
(v) If $X \parallel Y$, then $X \subseteq Y$, or $Y \subseteq X$, or $X \cap Y = \emptyset$.
(vi) $X \parallel Y \iff [(\exists\, \mathbf{w} \in \mathcal{V})(\mathbf{w}+X \subseteq Y)] \vee [(\exists\, \mathbf{w} \in \mathcal{V})(\mathbf{w}+Y \subseteq X)]$.

Proof.

(i) Since $\mathbf{w}+X = (\mathbf{w}+\mathbf{x})+\mathcal{S}$, every translate of X has base $\mathcal{S}$. If $\mathbf{w}+X \subseteq Y$, choose $\mathbf{z} \in \mathbf{w}+X$; then $\mathbf{w}+X = \mathbf{z}+\mathcal{S}$ and $Y = \mathbf{z}+\mathcal{T}$, hence $\mathcal{S} \subseteq \mathcal{T}$. Conversely, if $\mathcal{S} \subseteq \mathcal{T}$, then $(\mathbf{y}-\mathbf{x})+X = \mathbf{y}+\mathcal{S} \subseteq Y$.
(ii) If $\mathbf{w}+X = Y$, then the bases are equal. Conversely, if $\mathcal{S} = \mathcal{T}$, then $(\mathbf{y}-\mathbf{x})+X = Y$.
(iii) Choose $\mathbf{z} \in X \cap Y$. Then $X = \mathbf{z}+\mathcal{S}$ and $Y = \mathbf{z}+\mathcal{T}$, so translating by $-\mathbf{z}$ gives $X \subseteq Y \iff \mathcal{S} \subseteq \mathcal{T}$.
(iv) With the same $\mathbf{z}$, translate by $-\mathbf{z}$ to get $X = Y \iff \mathcal{S} = \mathcal{T}$.
(v) If $X \parallel Y$, then $\mathcal{S} \subseteq \mathcal{T}$ or $\mathcal{T} \subseteq \mathcal{S}$. If $X \cap Y \neq \emptyset$, part (iii) gives $X \subseteq Y$ or $Y \subseteq X$; otherwise $X \cap Y = \emptyset$.
(vi) If $X \parallel Y$, then either $\mathcal{S} \subseteq \mathcal{T}$ or $\mathcal{T} \subseteq \mathcal{S}$, so part (i) gives a translate $X' \subseteq Y$ or a translate $Y' \subseteq X$. Conversely, either containment gives $\mathcal{S} \subseteq \mathcal{T}$ or $\mathcal{T} \subseteq \mathcal{S}$ by part (i), hence $X \parallel Y$.

Structural Commentary: Intersecting flats reduce affine incidence to subspace incidence.
Key Insight: The representative may move, but the base controls containment, equality, and parallelism. □

Definition 2.6.3. *Affine Combinations*

(i) For $\mathbf{x}_1, \ldots, \mathbf{x}_n \in \mathcal{V}$ and $r_1, \ldots, r_n \in \mathbb{F}$, the sum $\sum_{j=1}^n r_j \mathbf{x}_j$ is an *affine combination* if $\sum_{j=1}^n r_j = 1$.
(ii) A nonempty $E \subseteq \mathcal{V}$ is *affine closed* if $\sum_{j=1}^n r_j \mathbf{x}_j \in E$ whenever $\mathbf{x}_j \in E$, $r_j \in \mathbb{F}$, and $\sum_{j=1}^n r_j = 1$.
(iii) A nonempty $E \subseteq \mathcal{V}$ is *two-affine closed* if $r\mathbf{x}+(1-r)\mathbf{y} \in E$ for all $\mathbf{x}, \mathbf{y} \in E$ and $r \in \mathbb{F}$.
(iv) For $\mathbf{x} \neq \mathbf{y}$, the affine line through them is $\overline{\mathbf{xy}} := \{r\mathbf{x}+(1-r)\mathbf{y} : r \in \mathbb{F}\} = \mathbf{y}+\operatorname{span}(\mathbf{x}-\mathbf{y})$.
Equivalently, a nonempty $E \subseteq \mathcal{V}$ is two-affine closed iff $\overline{\mathbf{xy}} \subseteq E$ for all distinct $\mathbf{x}, \mathbf{y} \in E$.

Theorem 2.6.4. *Two-Point Affine Closure*
Assume $\operatorname{char}(\mathbb{F}) \neq 2$. For every $\emptyset \neq E \subseteq \mathcal{V}$, E is affine closed $\iff$ E is two-affine closed.

Proof.

($\Rightarrow$) A two-point affine combination is an affine combination, so affine closure implies two-affine closure.
($\Leftarrow$) Assume E is two-affine closed.

Prove by induction on n that every affine combination $\sum_{j=1}^n r_j \mathbf{x}_j$, with $\mathbf{x}_j \in E$ and $\sum_{j=1}^n r_j = 1$, lies in E.

The cases $n = 1, 2$ are immediate. Let $n \geq 3$ and set $\mathbf{z} := \sum_{j=1}^n r_j \mathbf{x}_j$, $\sum_{j=1}^n r_j = 1$.

If some $r_j \neq 1$, relabel so $r_1 \neq 1$ and put $\mathbf{u} := \sum_{j=2}^n \frac{r_j}{1-r_1} \mathbf{x}_j$. Then $\sum_{j=2}^n \frac{r_j}{1-r_1} = 1$, so $\mathbf{u} \in E$ by induction, and $\mathbf{z} = r_1\mathbf{x}_1 + (1-r_1)\mathbf{u} \in E$ by two-affine closure.
Otherwise all $r_j = 1$; since 2 is invertible, $\mathbf{z} = 2\left(\frac{1}{2}\mathbf{x}_1 + \frac{1}{2}\mathbf{x}_2\right) + r_3\mathbf{x}_3 + \cdots + r_n\mathbf{x}_n$, where $\frac{1}{2}\mathbf{x}_1 + \frac{1}{2}\mathbf{x}_2 \in E$ by two-affine closure; this is an $(n-1)$-term affine combination, hence lies in E by induction.

Structural Commentary: For $\operatorname{char}(\mathbb{F}) \neq 2$, two-point closure generates finite affine closure.
Key Insight: The only exceptional induction case is handled by averaging two coefficient-1 terms. □

Remark 2.6.5. *Characteristic Two Warning*
Condition $\operatorname{char}(\mathbb{F}) \neq 2$ is necessary: in $\mathbb{F}_2^2$, $E := \{(0,0),(1,0),(0,1)\}$ is two-affine closed, but $\operatorname{aff}(E) = \mathbb{F}_2^2 \neq E$.

Theorem 2.6.6. *Flats Are the Affine-Closed Sets*
For every $\emptyset \neq E \subseteq \mathcal{V}$, $E \in \mathcal{A}(\mathcal{V}) \iff$ *E is affine closed.*

Proof.

(i) Suppose $E = \mathbf{x} + \mathcal{S}$ with $\mathcal{S} \leq \mathcal{V}$. If $\mathbf{x}_j = \mathbf{x} + \mathbf{s}_j \in E$ and $\sum_j r_j = 1$,
then $\sum_j r_j \mathbf{x}_j = \sum_j r_j(\mathbf{x} + \mathbf{s}_j) = \mathbf{x} + \sum_j r_j \mathbf{s}_j \in \mathbf{x} + \mathcal{S}$. Thus E is affine closed.

(ii) Suppose E is affine closed. Choose $\mathbf{x} \in E$ and set $\mathcal{S} := E - \mathbf{x} = \{\mathbf{e} - \mathbf{x} : \mathbf{e} \in E\}$.
For $\mathbf{s}_1, \mathbf{s}_2 \in \mathcal{S}$, write $\mathbf{s}_j = \mathbf{x}_j - \mathbf{x}$ with $\mathbf{x}_j \in E$.
For $r_1, r_2 \in \mathbb{F}$, $r_1\mathbf{s}_1 + r_2\mathbf{s}_2 + \mathbf{x} = r_1\mathbf{x}_1 + r_2\mathbf{x}_2 + (1 - r_1 - r_2)\mathbf{x} \in E$, so $r_1\mathbf{s}_1 + r_2\mathbf{s}_2 \in \mathcal{S}$.
Hence $\mathcal{S} \leq \mathcal{V}$ by Theorem 1.1.8, and $E = \mathbf{x} + \mathcal{S}$ is a flat.

Structural Commentary: An affine-closed set becomes a subspace after choosing one point as an origin.
Key Insight: Flats are precisely translated subspaces. □

Theorem 2.6.7. *Flats Are the Two-Affine-Closed Sets*
Assume $\operatorname{char}(\mathbb{F}) \neq 2$. *For every* $\emptyset \neq E \subseteq \mathcal{V}$, $E \in \mathcal{A}(\mathcal{V}) \iff$ *E is two-affine closed.*

Proof.

(i) If $E \in \mathcal{A}(\mathcal{V})$, then E is affine closed by Theorem 2.6.6, hence two-affine closed by Theorem 2.6.4.

(ii) If E is two-affine closed, then E is affine closed by Theorem 2.6.4, hence $E \in \mathcal{A}(\mathcal{V})$ by Theorem 2.6.6.

Structural Commentary: In characteristic not 2, two-point closure detects flats.
Key Insight: Two-affine closure equals finite affine closure, hence flatness. □

Definition 2.6.8. *Affine Hull and Affine Span*
Let $\emptyset \neq E \subseteq \mathcal{V}$.

(i) The *affine hull* $\operatorname{affhull}(E)$ is the least flat $X \in \mathcal{A}(\mathcal{V})$ with $E \subseteq X$.

(ii) The *affine span* $\operatorname{affspan}(E)$ is the set of all finite affine combinations of points of E.

Theorem 2.6.9. *Affine Hull Formula*
Let $\emptyset \neq E \subseteq \mathcal{V}$, $\mathbf{x} \in E$, *and* $\mathcal{S}_\mathbf{x} := \operatorname{span}(E - \mathbf{x})$.
Then $\operatorname{affhull}(E) = \operatorname{affspan}(E) = \mathbf{x} + \mathcal{S}_\mathbf{x}$, *and hence* $\dim(\operatorname{affhull}(E)) = \dim(\mathcal{S}_\mathbf{x})$.

Proof. If $\mathbf{z} = \sum_{j=1}^{n} r_j \mathbf{e}_j$ with $\mathbf{e}_j \in E$ and $\sum_j r_j = 1$, then $\mathbf{z} = \mathbf{x} + \sum_{j=1}^{n} r_j(\mathbf{e}_j - \mathbf{x}) \in \mathbf{x} + \mathcal{S}_\mathbf{x}$, so $\operatorname{affspan}(E) \subseteq \mathbf{x} + \mathcal{S}_\mathbf{x}$.
Conversely, for $\sum_{j=1}^{n} a_j(\mathbf{e}_j - \mathbf{x}) \in \mathcal{S}_\mathbf{x}$, $\mathbf{x} + \sum_{j=1}^{n} a_j(\mathbf{e}_j - \mathbf{x}) = \left(1 - \sum_{j=1}^{n} a_j\right)\mathbf{x} + \sum_{j=1}^{n} a_j \mathbf{e}_j \in \operatorname{affspan}(E)$, hence $\operatorname{affspan}(E) = \mathbf{x} + \mathcal{S}_\mathbf{x}$. This set is a flat containing E; any flat containing E is affine closed by Theorem 2.6.6, hence contains $\operatorname{affspan}(E)$. Therefore it is $\operatorname{affhull}(E)$. The dimension formula follows from Definition 2.6.1.
Structural Commentary: Affine span is ordinary span after temporarily choosing one point as the origin.
Key Insight: Differences from a base point record all directions in the affine hull. □

Lemma 2.6.10. *Intersections of Flats*
Let $X_i = \mathbf{x}_i + \mathcal{S}_i$ *be flats indexed by* $i \in I$. *If* $\bigcap_{i \in I} X_i \neq \emptyset$ *and* $\mathbf{z} \in \bigcap_{i \in I} X_i$, *then* $\bigcap_{i \in I} X_i = \mathbf{z} + \bigcap_{i \in I} \mathcal{S}_i$.

Proof. Since $\mathbf{z} \in X_i$, each flat has the form $X_i = \mathbf{z} + \mathcal{S}_i$. Hence $\mathbf{u} \in \bigcap_{i \in I} X_i \iff \mathbf{u} - \mathbf{z} \in \bigcap_{i \in I} \mathcal{S}_i$, QED.
Structural Commentary: A common point turns affine intersection into subspace intersection.
Key Insight: Translate to the common point; only the bases remain. □

Theorem 2.6.11. *Lattice of Flats*
Let $\mathcal{A}_0(\mathcal{V}) := \mathcal{A}(\mathcal{V}) \cup \{\emptyset\}$, *ordered by inclusion. Then* $\mathcal{A}_0(\mathcal{V})$ *is a complete lattice.*
For every family $\mathcal{C} \subseteq \mathcal{A}_0(\mathcal{V})$, $\bigwedge \mathcal{C} = \bigcap_{X \in \mathcal{C}} X$, $\bigvee \mathcal{C} = \begin{cases} \operatorname{affhull}\left(\bigcup_{X \in \mathcal{C}} X\right), & \bigcup_{X \in \mathcal{C}} X \neq \emptyset, \\ \emptyset, & \bigcup_{X \in \mathcal{C}} X = \emptyset. \end{cases}$

Proof.

(i) Arbitrary intersections of elements of $\mathcal{A}_0(\mathcal{V})$ lie in $\mathcal{A}_0(\mathcal{V})$: they are either empty, equal to $\mathcal{V}$ for the empty family, or flats by Lemma 2.6.10.

(ii) If $\bigcup_{X\in\mathcal{C}} X = \emptyset$, then $\emptyset$ is the least upper bound of $\mathcal{C}$.
Otherwise, affhull$(\bigcup_{X\in\mathcal{C}} X)$ is the least flat containing every $X \in \mathcal{C}$, so it is the join.

(iii) Since arbitrary meets and joins exist, $\mathcal{A}_0(\mathcal{V})$ is complete.

Structural Commentary: Adding the empty set repairs the failure of affine flats to be closed under intersection.
Key Insight: Meet is common incidence; join is the affine hull forced by union. □

Remark 2.6.12. *Empty Affine Hull Convention*
Within $\mathcal{A}_0(\mathcal{V})$, one may set affhull$(\emptyset) := \emptyset$ to write the join formula uniformly.

Corollary 2.6.13. *Join of Two Flats*
Let $X = \mathbf{x} + \mathcal{S}$ and $Y = \mathbf{y} + \mathcal{T}$ be flats. Then $X \vee Y = \mathbf{x} + \operatorname{span}\big(\{\mathbf{x}-\mathbf{y}\} \cup \mathcal{S} \cup \mathcal{T}\big)$.
If $X \cap Y \neq \emptyset$, then $X \vee Y = \mathbf{x} + (\mathcal{S} + \mathcal{T})$.

Proof. By Theorem 2.6.11, $X \vee Y = \text{affhull}(X \cup Y)$. Apply Theorem 2.6.9 to $X \cup Y$ with base point $\mathbf{x}$, then $(X \cup Y) - \mathbf{x} = \mathcal{S} \cup \big((\mathbf{y} - \mathbf{x}) + \mathcal{T}\big)$, whose span is $\operatorname{span}(\{\mathbf{x}-\mathbf{y}\} \cup \mathcal{S} \cup \mathcal{T})$.
If $X \cap Y \neq \emptyset$, then $\mathbf{x} - \mathbf{y} \in \mathcal{S} + \mathcal{T}$, so this span is $\mathcal{S} + \mathcal{T}$.
Structural Commentary: The join adds both bases and the displacement between representatives.
Key Insight: The displacement direction disappears exactly when the flats meet. □

Corollary 2.6.14. *Dimension of the Join of Two Flats*
Let $X = \mathbf{x} + \mathcal{S}$ and $Y = \mathbf{y} + \mathcal{T}$ be flats.
(i) If $X \cap Y \neq \emptyset$, then $\dim(X \vee Y) = \dim(X) + \dim(Y) - \dim(X \cap Y)$.
(ii) If $X \cap Y = \emptyset$, then $\dim(X \vee Y) = \dim(\mathcal{S} + \mathcal{T}) + 1$.

Proof.

(i) If $X \cap Y \neq \emptyset$, choose $\mathbf{z} \in X \cap Y$.
By Corollary 2.6.13 and Lemma 2.6.10, $X \vee Y = \mathbf{x} + (\mathcal{S} + \mathcal{T})$; $X \cap Y = \mathbf{z} + (\mathcal{S} \cap \mathcal{T})$. Thus the dimension formula for subspace sums gives $\dim(X \vee Y) = \dim(\mathcal{S} + \mathcal{T}) = \dim(\mathcal{S}) + \dim(\mathcal{T}) - \dim(\mathcal{S} \cap \mathcal{T})$.
Since $\dim X = \dim \mathcal{S}$, $\dim Y = \dim \mathcal{T}$, and $\dim(X \cap Y) = \dim(\mathcal{S} \cap \mathcal{T})$, this is (i).

(ii) If $X \cap Y = \emptyset$, Corollary 2.6.13 gives base $\mathcal{U} := \operatorname{span}(\{\mathbf{x}-\mathbf{y}\} \cup \mathcal{S} \cup \mathcal{T})$.
Since $\mathbf{x} - \mathbf{y} \notin \mathcal{S} + \mathcal{T}$, we have $\dim \mathcal{U} = \dim(\mathcal{S} + \mathcal{T}) + 1$, hence $\dim(X \vee Y) = \dim(\mathcal{S} + \mathcal{T}) + 1$.

Structural Commentary: Intersections subtract overlap; disjoint joins add one connecting direction.
Key Insight: The vector between representatives is new exactly when the flats are disjoint. □

Definition 2.6.15. *Affinely Independent Sets*
A nonempty set $E \subseteq \mathcal{V}$ is *affinely independent* if, for every $\mathbf{x} \in E$, the set $(E - \mathbf{x}) \setminus \{\mathbf{0}\}$ is linearly independent.

Theorem 2.6.16. *Affine Independence Criteria*
For $\emptyset \neq E \subseteq \mathcal{V}$, conditions (i)–(iv) are equivalent. If E is finite, they are also equivalent to (v).
(i) E is affinely independent.
(ii) For every $\mathbf{x} \in E$, $\mathbf{x} \notin \text{affhull}(E \setminus \{\mathbf{x}\})$, with affhull$(\emptyset) = \emptyset$ in the singleton case.
(iii) $\forall$ distinct $\mathbf{x}_1, \ldots, \mathbf{x}_n \in E$ and $r_1, \ldots, r_n \in \mathbb{F}$, $\sum_{j=1}^{n} r_j \mathbf{x}_j = \mathbf{0}$ and $\sum_{j=1}^{n} r_j = 0$ imply $r_1 = \cdots = r_n = 0$.
(iv) $\forall$ distinct $\mathbf{x}_1, \ldots, \mathbf{x}_n \in E$ and affine coefficient vectors $\mathbf{r}, \mathbf{s} \in \mathbb{F}^n$,
$\sum_{j=1}^{n} r_j \mathbf{x}_j = \sum_{j=1}^{n} s_j \mathbf{x}_j$ and $\sum_{j=1}^{n} r_j = \sum_{j=1}^{n} s_j = 1$ imply $\mathbf{r} = \mathbf{s}$.
(v) If $E = \{\mathbf{x}_1, \ldots, \mathbf{x}_n\}$ is finite, then $\dim(\text{affhull}(E)) = n - 1$.

Proof.

(i) $\Rightarrow$ (iii) Let $\sum_{j=1}^{n} r_j \mathbf{x}_j = \mathbf{0}$ with distinct $\mathbf{x}_j \in E$ and $\sum_j r_j = 0$. Then $\sum_{j=1}^{n-1} r_j(\mathbf{x}_j - \mathbf{x}_n) = \mathbf{0}$.
By (i), $r_1 = \cdots = r_{n-1} = 0$, and then $r_n = 0$.

(iii) $\Rightarrow$ (i) Fix $\mathbf{x} \in E$. Any relation $\sum_{j=1}^{m} a_j(\mathbf{y}_j - \mathbf{x}) = \mathbf{0}$, $\mathbf{y}_j \in E \setminus \{\mathbf{x}\}$ is the zero-sum relation
$\sum_{j=1}^{m} a_j \mathbf{y}_j - (\sum_j a_j)\mathbf{x} = \mathbf{0}$, so (iii) gives $a_1 = \cdots = a_m = 0$.

(iii) $\Rightarrow$ (ii) If $\mathbf{x}_k \in \text{affhull}(E \setminus \{\mathbf{x}_k\})$, then $\mathbf{x}_k = \sum_j r_j \mathbf{x}_j$ for $\mathbf{x}_j \neq \mathbf{x}_k$ and $\sum_j r_j = 1$,
giving a nontrivial zero-sum relation, contradiction.

(*ii*) $\Rightarrow$ (*iii*) If a zero-sum relation is nontrivial, choose $r_k \neq 0$.
Then $\mathbf{x}_k = \sum_{j\neq k} \left(-\frac{r_j}{r_k}\right)\mathbf{x}_j$, $\sum_{j\neq k} -\frac{r_j}{r_k} = 1$, so $\mathbf{x}_k \in \operatorname{affhull}(E \setminus \{\mathbf{x}_k\})$, contradicting (ii).
(*iii*) $\Rightarrow$ (*iv*) Subtract the two affine representations; the coefficient vector $\mathbf{r}-\mathbf{s}$ has sum 0, so (iii) gives $\mathbf{r}=\mathbf{s}$.
(*iv*) $\Rightarrow$ (*iii*) Let $\sum_j r_j\mathbf{x}_j = \mathbf{0}$ and $\sum_j r_j = 0$. Choose $\mathbf{a} \in \mathbb{F}^n$ with $\sum_j a_j = 1$.
Then $\mathbf{a}$ and $\mathbf{a}+\mathbf{r}$ are affine coefficient vectors for the same point. By (iv), $\mathbf{r}=\mathbf{0}$.
(*i*) $\Leftrightarrow$ (*v*) If $E = \{\mathbf{x}_1,\dots,\mathbf{x}_n\}$ and $\mathbf{x} \in E$, then Theorem 2.6.9 gives $\dim(\operatorname{affhull}(E)) = \dim(\operatorname{span}(E-\mathbf{x}))$.
Since $(E-\mathbf{x})\setminus\{\mathbf{0}\}$ has $n-1$ vectors, this dimension is $n-1$ iff those vectors are linearly independent.

Structural Commentary: Affine independence is linear independence after forgetting one chosen origin.
Key Insight: Uniqueness is absence of nontrivial zero-sum affine relations. □

Definition 2.6.17. *Affine Bases and Barycentric Coordinates* Let $A = \mathbf{x}+\mathcal{S}$ be an n-flat.

(i) A subset $\mathcal{B} \subseteq A$ is an *affine basis* of A if $\mathcal{B}$ is affinely independent and $\operatorname{affspan}(\mathcal{B}) = A$; an *ordered affine basis* is an ordering of such a basis.
(ii) If $(\mathbf{s}_1,\dots,\mathbf{s}_n)$ is an ordered basis of $\mathcal{S}$, then $(\mathbf{x}+\mathbf{s}_1,\dots,\mathbf{x}+\mathbf{s}_n,\mathbf{x})$ is an ordered affine basis of A.
(iii) For an ordered affine basis $(\mathbf{p}_1,\dots,\mathbf{p}_{n+1})$, every $\mathbf{v} \in A$ has a unique expression $\mathbf{v} = \sum_{j=1}^{n+1} r_j\mathbf{p}_j$, $\sum_{j=1}^{n+1} r_j = 1$.
The scalars $(r_1,\dots,r_{n+1})$ are the *barycentric coordinates* of $\mathbf{v}$ in this basis.

Example 2.6.18. *Barycentric Coordinates on a Plane*
Let $A = \mathbf{x}+\operatorname{span}(\mathbf{v}_1,\mathbf{v}_2) \subseteq \mathbb{R}^3$, with $(\mathbf{v}_1,\mathbf{v}_2)$ linearly independent.
Set $\mathbf{p}_1 = \mathbf{x}+\mathbf{v}_1$, $\mathbf{p}_2 = \mathbf{x}+\mathbf{v}_2$, and $\mathbf{p}_3 = \mathbf{x}$. Then $(\mathbf{p}_1,\mathbf{p}_2,\mathbf{p}_3)$ is an ordered affine basis of A, and every $\mathbf{a} \in A$ has a unique barycentric expansion $\mathbf{a} = r_1\mathbf{p}_1 + r_2\mathbf{p}_2 + r_3\mathbf{p}_3$, $r_1+r_2+r_3 = 1$, with $r_1,r_2,r_3 \in \mathbb{R}$.

Definition 2.6.19. *Affine Transformation*
A map $f : \mathcal{V} \to \mathcal{V}$ is an *affine transformation* if it preserves finite affine combinations:
for all $\mathbf{x}_j \in \mathcal{V}$ and $r_j \in \mathbb{F}$ with $\sum_j r_j = 1$, $f\left(\sum_j r_j\mathbf{x}_j\right) = \sum_j r_j f(\mathbf{x}_j)$.

Theorem 2.6.20. *Two-Point Test for Affine Maps*
Assume $\operatorname{char}(\mathbb{F}) \neq 2$.
A map $f : \mathcal{V} \to \mathcal{V}$ is affine if and only if $f(r\mathbf{x}+(1-r)\mathbf{y}) = rf(\mathbf{x})+(1-r)f(\mathbf{y})$ for all $\mathbf{x},\mathbf{y} \in \mathcal{V}$ and $r \in \mathbb{F}$.

Proof.
($\Rightarrow$) Since $r+(1-r)=1$, affinity gives $f(r\mathbf{x}+(1-r)\mathbf{y}) = rf(\mathbf{x})+(1-r)f(\mathbf{y})$.
($\Leftarrow$) Assume displayed identity. Prove by induction on n that $f\left(\sum_{j=1}^n r_j\mathbf{x}_j\right) = \sum_{j=1}^n r_j f(\mathbf{x}_j)$; $\sum_{j=1}^n r_j = 1$.
The cases $n = 1, 2$ are immediate. For $n \geq 3$, set $\mathbf{z} := \sum_{j=1}^n r_j\mathbf{x}_j$. If some $r_j \neq 1$, relabel so $r_1 \neq 1$ and put $\mathbf{u} := \sum_{j=2}^n \frac{r_j}{1-r_1}\mathbf{x}_j$. Then $f(\mathbf{z}) = f(r_1\mathbf{x}_1+(1-r_1)\mathbf{u}) = r_1f(\mathbf{x}_1)+(1-r_1)f(\mathbf{u}) = \sum_{j=1}^n r_j f(\mathbf{x}_j)$, by the two-point identity and induction. Otherwise all $r_j = 1$; since 2 is invertible, $\mathbf{z} = 2\left(\frac{1}{2}\mathbf{x}_1+\frac{1}{2}\mathbf{x}_2\right)+r_3\mathbf{x}_3+\cdots+r_n\mathbf{x}_n$, an $(n-1)$-term affine combination, so induction plus two-point identity gives same formula for $f(\mathbf{z})$.

Structural Commentary: In characteristic not 2, line preservation gives affine preservation.
Key Insight: The same two-point assembly that builds affine-closed sets also tests affine maps. □

Definition 2.6.21. *Translations*
For $\mathbf{a} \in \mathcal{V}$, the *translation by* $\mathbf{a}$ is the map $T_\mathbf{a} : \mathcal{V} \to \mathcal{V}$ defined by $T_\mathbf{a}(\mathbf{x}) = \mathbf{x}+\mathbf{a}$.

Theorem 2.6.22. *Affine Maps Are Linear Maps Followed by Translations*
A map $f : \mathcal{V} \to \mathcal{V}$ is affine iff $f = T_\mathbf{a} \circ \tau$ for unique $\mathbf{a} \in \mathcal{V}$ and $\tau \in \mathcal{L}(\mathcal{V})$, namely $\mathbf{a} = f(\mathbf{0})$ and $\tau = T_{-\mathbf{a}} \circ f$.

Proof.
($\Rightarrow$) Set $\mathbf{a} := f(\mathbf{0})$ and $\tau(\mathbf{x}) := f(\mathbf{x})-\mathbf{a}$. For $r,s \in \mathbb{F}$ and $\mathbf{x},\mathbf{y} \in \mathcal{V}$,
$\tau(r\mathbf{x}+s\mathbf{y}) = f(r\mathbf{x}+s\mathbf{y}+(1-r-s)\mathbf{0})-\mathbf{a} = r\tau(\mathbf{x})+s\tau(\mathbf{y})$, so $\tau \in \mathcal{L}(\mathcal{V})$ and $f = T_\mathbf{a} \circ \tau$.
Uniqueness follows from $\mathbf{a} = f(\mathbf{0})$ and $\tau(\mathbf{x}) = f(\mathbf{x})-\mathbf{a}$.
($\Leftarrow$) If $f = T_\mathbf{a} \circ \tau$ with τ linear,
then for $\sum_j r_j = 1$, $f\left(\sum_j r_j\mathbf{x}_j\right) = \mathbf{a}+\tau\left(\sum_j r_j\mathbf{x}_j\right) = \sum_j r_j(\mathbf{a}+\tau\mathbf{x}_j) = \sum_j r_j f(\mathbf{x}_j)$, hence f is affine.

Structural Commentary: An affine map is linear after moving its image of the origin back to the origin.
Key Insight: The failure of an affine map to be linear is exactly a translation. □

Definition 2.6.23. *Affine Group Data*

(i) The *general linear group* is $\mathrm{GL}(\mathcal{V}) := \{\tau \in \mathcal{L}(\mathcal{V}) : \tau \text{ is invertible}\}$.
(ii) The *affine group* is $\mathrm{Aff}(\mathcal{V}) := \{T_{\mathbf{a}} \circ \tau : \mathbf{a} \in \mathcal{V}, \tau \in \mathrm{GL}(\mathcal{V})\}$.
(iii) The *translation group* is $\mathrm{Trans}(\mathcal{V}) := \{T_{\mathbf{a}} : \mathbf{a} \in \mathcal{V}\}$.
(iv) The *linear-part map* is $\Lambda : \mathrm{Aff}(\mathcal{V}) \to \mathrm{GL}(\mathcal{V})$; $\Lambda(T_{\mathbf{a}} \circ \tau) = \tau$, by Theorem 2.6.22.

Corollary 2.6.24. *Bijective Affine Maps*
For $\mathbf{a} \in \mathcal{V}$ and $\tau \in \mathcal{L}(\mathcal{V})$, $T_{\mathbf{a}} \circ \tau$ is bijective $\iff$ $\tau \in \mathrm{GL}(\mathcal{V})$.

Proof.

($\Rightarrow$) If $T_{\mathbf{a}} \circ \tau$ is bijective, then $\tau = T_{-\mathbf{a}} \circ (T_{\mathbf{a}} \circ \tau)$ is bijective, hence $\tau \in \mathrm{GL}(\mathcal{V})$.
($\Leftarrow$) If $\tau \in \mathrm{GL}(\mathcal{V})$, then $T_{\mathbf{a}} \circ \tau$ is bijective with inverse $\tau^{-1} \circ T_{-\mathbf{a}}$.

Structural Commentary: Translation never affects invertibility; only the linear part does.
Key Insight: Bijective affine maps are exactly translations of linear isomorphisms. □

Corollary 2.6.25. *Affine Group and Translations*

(i) For $\mathbf{a}, \mathbf{b} \in \mathcal{V}$ and $\tau, \sigma \in \mathcal{L}(\mathcal{V})$, $(T_{\mathbf{a}} \circ \tau) \circ (T_{\mathbf{b}} \circ \sigma) = T_{\mathbf{a}+\tau\mathbf{b}} \circ (\tau\sigma)$.
(ii) $\mathrm{Aff}(\mathcal{V})$ is a group, $\mathrm{Trans}(\mathcal{V}) \trianglelefteq \mathrm{Aff}(\mathcal{V})$, and Λ is a surjective group homomorphism with $\ker \Lambda = \mathrm{Trans}(\mathcal{V})$.
(iii) Consequently, $\mathrm{Aff}(\mathcal{V})/\mathrm{Trans}(\mathcal{V}) \cong \mathrm{GL}(\mathcal{V})$.

Proof.

i. For $\mathbf{x} \in \mathcal{V}$, $((T_{\mathbf{a}} \circ \tau) \circ (T_{\mathbf{b}} \circ \sigma))(\mathbf{x}) = \mathbf{a} + \tau\mathbf{b} + (\tau\sigma)\mathbf{x}$.
ii. For $\tau, \sigma \in \mathrm{GL}(\mathcal{V})$, (i) gives closure in $\mathrm{Aff}(\mathcal{V})$; the identity is $T_{\mathbf{0}} \circ I_{\mathcal{V}}$, and $(T_{\mathbf{a}} \circ \tau)^{-1} = T_{-\tau^{-1}\mathbf{a}} \circ \tau^{-1}$. Thus $\mathrm{Aff}(\mathcal{V})$ is a group. Also (i) gives $\Lambda(gh) = \Lambda(g)\Lambda(h)$, $\ker \Lambda = \{T_{\mathbf{a}} \circ I_{\mathcal{V}} : \mathbf{a} \in \mathcal{V}\} = \mathrm{Trans}(\mathcal{V})$, and Λ is surjective since $\tau = T_{\mathbf{0}} \circ \tau$ for all $\tau \in \mathrm{GL}(\mathcal{V})$. Hence $\mathrm{Trans}(\mathcal{V}) \trianglelefteq \mathrm{Aff}(\mathcal{V})$.
iii. Since Λ is surjective and $\ker \Lambda = \mathrm{Trans}(\mathcal{V})$, the first isomorphism theorem gives the induced isomorphism $\overline{\Lambda} : \mathrm{Aff}(\mathcal{V})/\mathrm{Trans}(\mathcal{V}) \to \mathrm{GL}(\mathcal{V})$; $\overline{\Lambda}(g\ \mathrm{Trans}(\mathcal{V})) = \Lambda(g)$.

Structural Commentary: The affine group is the translation part together with the invertible linear part.
Key Insight: Forgetting the translation records a homomorphism onto the linear automorphism group. □

Definition 2.6.26. *Projective Geometry of a Vector Space*

(i) Let $\mathrm{Sub}(\mathcal{V})$ be the set of subspaces of $\mathcal{V}$.
The *projective geometry* of $\mathcal{V}$ is $\mathrm{Sub}(\mathcal{V})$ with $\mathrm{pdim}(\mathcal{U}) := \dim(\mathcal{U}) - 1$.
(ii) A *projective point* is a 1-dimensional subspace; a *projective line* is a 2-dimensional subspace.
(iii) If $H = \mathbf{x} + \mathcal{K}$ is an affine hyperplane with $\mathbf{0} \notin H$ and $X \subseteq H$ is an affine flat, then X is represented projectively by $\mathrm{span}(X)$, with $\mathrm{pdim}(\mathrm{span}(X)) = \dim(X)$.

Perspective 2.6.27. *Projective Completion of an Affine Hyperplane*
Let $H = \mathbf{x} + \mathcal{K}$ be a hyperplane with $\mathbf{0} \notin H$. The map $X \subseteq H \mapsto \mathrm{span}(X) \leq \mathcal{V}$ adjoins the missing directions in $\mathcal{K}$. Thus, when $\dim H = 2$, parallel affine lines in H acquire a common projective point at infinity.

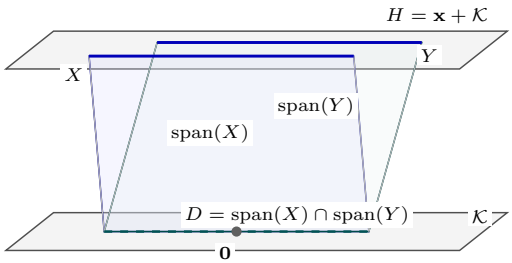


Figure 39: In the projective completion of $H = \mathbf{x} + \mathcal{K}$, an affine flat $X \subseteq H$ is represented by $\mathrm{span}(X)$; parallel flats meet in their common direction $D \leq \mathcal{K}$.

Theorem 2.6.28. *Linear Span Map from Affine to Projective Geometry*
Let $H = \mathbf{x} + \mathcal{K}$ be an affine hyperplane with $\mathbf{0} \notin H$, and set $\mathcal{A}(H) := \{X \in \mathcal{A}(\mathcal{V}) : X \subseteq H\}$.
The map $P : \mathcal{A}(H) \to \mathrm{Sub}(\mathcal{V})$, $P(X) = \mathrm{span}(X)$ satisfies:

(i) P is injective and $P^{-1}(U) = U \cap H$ for $U \in \mathrm{im}\, P$.

(ii) $\operatorname{im} P = \{\mathcal{U} \le \mathcal{V} : \mathcal{U} \not\subseteq \mathcal{K}\}$.
(iii) $X \subseteq Y$ *if and only if* $P(X) \subseteq P(Y)$.
(iv) *If* $\bigcap_{i\in I} X_i \neq \emptyset$, *then* $P\left(\bigcap_{i\in I} X_i\right) = \bigcap_{i\in I} P(X_i)$.
(v) *For any family* $\{X_i : i \in I\} \subseteq \mathcal{A}(H)$, $P\left(\bigvee_{i\in I} X_i\right) = \bigvee_{i\in I} P(X_i)$, *where* $\bigvee_{i\in I} P(X_i)$ *is computed in* $\operatorname{Sub}(\mathcal{V})$, *i.e.* $\bigvee_{i\in I} P(X_i) = \sum_{i\in I} P(X_i)$.
(vi) $\operatorname{pdim}(P(X)) = \dim(X)$.
(vii) $X \parallel Y$ *if and only if one of* $P(X) \cap \mathcal{K}$ *and* $P(Y) \cap \mathcal{K}$ *is contained in the other.*

Proof.

(i) Let $X = \mathbf{u} + \mathcal{S} \subseteq H$. Then $\mathcal{S} \le \mathcal{K}$, $\mathbf{u} \notin \mathcal{K}$, and $P(X) = \mathbb{F}\mathbf{u} \oplus \mathcal{S}$.
If $\mathbf{z} \in P(X) \cap H$, then $\mathbf{z} = r\mathbf{u} + \mathbf{s} = \mathbf{u} + \mathbf{k}$ with $\mathbf{s} \in \mathcal{S}$ and $\mathbf{k} \in \mathcal{K}$. Hence $(1-r)\mathbf{u} \in \mathcal{K}$, so $r = 1$, and $P(X) \cap H = X$. Thus P is injective and $P^{-1}(U) = U \cap H$ on $\operatorname{im} P$.
(ii) If $\mathcal{U} \not\subseteq \mathcal{K}$, choose $\mathbf{u} \in \mathcal{U} \setminus \mathcal{K}$. Since $\mathcal{V}/\mathcal{K}$ is 1-dimensional and $H = \mathbf{x} + \mathcal{K}$, there is $r \neq 0$ with $\mathbf{h} := r\mathbf{u} \in H$. Then $\mathcal{U} \cap H = \mathbf{h} + (\mathcal{U} \cap \mathcal{K})$; $P(\mathcal{U} \cap H) = \mathbb{F}\mathbf{h} + (\mathcal{U} \cap \mathcal{K}) = \mathcal{U}$. Thus $\operatorname{im} P = \{\mathcal{U} \le \mathcal{V} : \mathcal{U} \not\subseteq \mathcal{K}\}$.
(iii) If $X \subseteq Y$, then $P(X) \subseteq P(Y)$. Conversely, if $P(X) \subseteq P(Y)$, then $X = P(X) \cap H \subseteq P(Y) \cap H = Y$.
(iv) If $\mathbf{z} \in \bigcap_i X_i$ and $X_i = \mathbf{z} + \mathcal{S}_i$, then Lemma 2.6.10 gives $\bigcap_i X_i = \mathbf{z} + \bigcap_i \mathcal{S}_i$.
Since $P(X_i) = \mathbb{F}\mathbf{z} \oplus \mathcal{S}_i$, $\bigcap_i P(X_i) = \mathbb{F}\mathbf{z} \oplus \bigcap_i \mathcal{S}_i = P\left(\bigcap_i X_i\right)$.
(v) By Theorem 2.6.11, $P\left(\bigvee_i X_i\right) = \operatorname{span}\left(\operatorname{affhull}\left(\bigcup_i X_i\right)\right) = \operatorname{span}\left(\bigcup_i X_i\right) = \bigvee_i P(X_i)$.
(vi) For $X = \mathbf{u} + \mathcal{S} \subseteq H$, $\operatorname{pdim}(P(X)) = \dim(\mathbb{F}\mathbf{u} \oplus \mathcal{S}) - 1 = \dim \mathcal{S} = \dim X$.
(vii) If $X = \mathbf{u} + \mathcal{S}$ and $Y = \mathbf{v} + \mathcal{T}$, then $P(X) \cap \mathcal{K} = \mathcal{S}$ and $P(Y) \cap \mathcal{K} = \mathcal{T}$.
Hence $X \parallel Y \iff \mathcal{S} \subseteq \mathcal{T}$ or $\mathcal{T} \subseteq \mathcal{S} \iff P(X) \cap \mathcal{K} \subseteq P(Y) \cap \mathcal{K}$ or $P(Y) \cap \mathcal{K} \subseteq P(X) \cap \mathcal{K}$.

Structural Commentary: H selects affine representatives for projective subspaces outside $\mathcal{K}$.
Key Insight: Intersecting with H undoes the linear span exactly because H does not pass through the origin. □

2.7 Positive Systems, Convexity, and Separation

This subsection uses finite-dimensional real inner-product geometry to decide when linear systems have positive solutions. The guiding theme is that a missing positive solution is often detected by a separating hyperplane, hence by a dual vector certificate.

Definition 2.7.1. *Positive Orthants*
For $\mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$, use the following order and orthant notation:

(i) $\mathbf{x} \ge \mathbf{0}$ iff $\mathbf{x}$ is *coordinatewise nonnegative*: $x_i \ge 0$ for every i.
(ii) $\mathbf{x} > \mathbf{0}$ iff $\mathbf{x}$ is *nonzero nonnegative*: $\mathbf{x} \ge \mathbf{0}$ and $\mathbf{x} \neq \mathbf{0}$.
(iii) $\mathbf{x} \gg \mathbf{0}$ iff $\mathbf{x}$ is *coordinatewise positive*: $x_i > 0$ for every i.
(iv) $\mathbb{R}^n_+ := \{\mathbf{x} \in \mathbb{R}^n : \mathbf{x} > \mathbf{0}\}$ is the *nonzero nonnegative orthant*.
(v) $\mathbb{R}^n_{++} := \{\mathbf{x} \in \mathbb{R}^n : \mathbf{x} \gg \mathbf{0}\}$ is the *coordinatewise positive interior*.

Motivation 2.7.2. *Positive Null Vectors as Probabilities*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \operatorname{Mat}(m, n, \mathbb{R})$ and let $\mathbf{x} = (x_1, \dots, x_n) > \mathbf{0}$ satisfy $A\mathbf{x} = \mathbf{0}$.

(i) Set $s := x_1 + \cdots + x_n > 0$ and $\boldsymbol{\pi} := s^{-1}\mathbf{x}$.
(ii) Then $\pi_i \ge 0$ and $\sum_{i=1}^n \pi_i = 1$, so $\boldsymbol{\pi}$ is a probability vector.
(iii) Since $A\boldsymbol{\pi} = s^{-1}A\mathbf{x} = \mathbf{0}$, equivalently $\sum_{i=1}^n \pi_i \mathbf{a}_i = \mathbf{0}$.
(iv) If $\mathbf{x} \gg \mathbf{0}$, then $\pi_i > 0$ for every i.

Thus a positive null vector normalizes to probabilities under which the columns of A have zero expected value.

Principle 2.7.3. *Positive Solutions Are Orthant Intersections*
For $A \in \operatorname{Mat}(m, n, \mathbb{R})$, $\ker(A) = \operatorname{row}(A)^\perp$ by Corollary 2.4.13.
Hence $\{\mathbf{x} \in \mathbb{R}^n : A\mathbf{x} = \mathbf{0},\ \mathbf{x} > \mathbf{0}\} = \operatorname{row}(A)^\perp \cap \mathbb{R}^n_+$; and $\{\mathbf{x} \in \mathbb{R}^n : A\mathbf{x} = \mathbf{0},\ \mathbf{x} \gg \mathbf{0}\} = \operatorname{row}(A)^\perp \cap \mathbb{R}^n_{++}$.
Thus *positive null vectors* are precisely orthant intersections with $\operatorname{row}(A)^\perp$.

Definition 2.7.4. *Convex Combinations, Cones, and Hulls*
Let $\mathbf{x}_1, \dots, \mathbf{x}_k \in \mathbb{R}^n$.

i. A *convex combination* of these vectors is a vector $t_1\mathbf{x}_1 + \cdots + t_k\mathbf{x}_k$, $t_i \ge 0$, $\sum_{i=1}^k t_i = 1$.
ii. A set $C \subseteq \mathbb{R}^n$ is *convex* if the line segment between any two of its points lies in C.
iii. A set $K \subseteq \mathbb{R}^n$ is a *cone* if $\mathbf{x} \in K$ and $a \ge 0$ imply $a\mathbf{x} \in K$.

iv. *Convex hull* of finite set $S = \{\mathbf{x}_1, \dots, \mathbf{x}_k\}$ is smallest convex subset of $\mathbb{R}^n$ containing S, denoted $\operatorname{conv}(S)$.

Closedness and compactness use the Euclidean metric conventions from Definitions 2.1.14 and 2.2.1.

Proposition 2.7.5. *Finite Convex Hulls*

Let $S = \{\mathbf{x}_1, \dots, \mathbf{x}_k\} \subseteq \mathbb{R}^n$ *and* $\Delta_{k-1} := \left\{(t_1, \dots, t_k) \in \mathbb{R}^k : t_i \geq 0,\ \sum_{i=1}^{k} t_i = 1\right\}$.

Then $\operatorname{conv}(S) = \left\{\sum_{i=1}^{k} t_i \mathbf{x}_i : (t_1, \dots, t_k) \in \Delta_{k-1}\right\}$. *Moreover,* $\operatorname{conv}(S)$ *is compact.*

Proof. Let $T : \Delta_{k-1} \to \mathbb{R}^n$ be $T(t_1, \dots, t_k) = \sum_{i=1}^{k} t_i \mathbf{x}_i$, and set $C := T(\Delta_{k-1})$. Taking $t_i = 1$ and $t_j = 0$ for $j \neq i$ gives $S \subseteq C$. If $K \subseteq \mathbb{R}^n$ is convex and $S \subseteq K$, then $C \subseteq K$, hence $C \subseteq \operatorname{conv}(S)$.
Conversely, Δ_{k-1} is convex and T is affine, so C is convex; since $S \subseteq C$, $\operatorname{conv}(S) \subseteq C$. Thus $\operatorname{conv}(S) = C$.
The simplex Δ_{k-1} is closed and bounded, hence compact by Heine–Borel, Theorem 2.2.39.
Since T is continuous, $\operatorname{conv}(S) = T(\Delta_{k-1})$ is compact by Theorem 2.3.33.
Structural Commentary: A finite convex hull is the continuous image of a compact coefficient simplex.
Key Insight: Convexity is carried by the coefficients t_i, not by the ambient coordinates. □

Proposition 2.7.6. *Sums of Convex and Closed Sets*

For $X, Y \subseteq \mathbb{R}^n$, *set* $X + Y := \{\mathbf{x} + \mathbf{y} : \mathbf{x} \in X,\ \mathbf{y} \in Y\}$. *Then:*

(i) X, Y *convex* $\implies$ $X + Y$ *convex.*
(ii) X *compact and* Y *closed* $\implies$ $X + Y$ *closed.*

Proof.

(i) Assume X, Y are convex. Let $\mathbf{z}_r = \mathbf{x}_r + \mathbf{y}_r \in X + Y$ for $r = 0, 1$, and let $0 \leq t \leq 1$. Then $t\mathbf{z}_0 + (1-t)\mathbf{z}_1 = \big(t\mathbf{x}_0 + (1-t)\mathbf{x}_1\big) + \big(t\mathbf{y}_0 + (1-t)\mathbf{y}_1\big) \in X + Y$. Thus $X + Y$ is convex.
(ii) Assume X is compact and Y is closed. Let $\mathbf{z}_j \in X + Y$ and $\mathbf{z}_j \to \mathbf{z}$.
Let $\mathbf{z}_j = \mathbf{x}_j + \mathbf{y}_j$ with $\mathbf{x}_j \in X$ and $\mathbf{y}_j \in Y$. Since X is compact, a subsequence satisfies $\mathbf{x}_{j_\ell} \to \mathbf{x} \in X$.
Then $\mathbf{y}_{j_\ell} = \mathbf{z}_{j_\ell} - \mathbf{x}_{j_\ell} \to \mathbf{z} - \mathbf{x} \in Y$, as Y is closed. Hence $\mathbf{z} = \mathbf{x} + (\mathbf{z} - \mathbf{x}) \in X + Y$, so $X + Y$ is closed.

Structural Commentary: Compactness lets one hold the X-part still enough for closedness of Y to apply.
Key Insight: A limit in a sum can be split by passing to a convergent compact subsequence. □

Definition 2.7.7. *Hyperplanes, Half-Spaces, and Separation*

Let $\mathbf{N} \in \mathbb{R}^n \setminus \{\mathbf{0}\}$ and $b \in \mathbb{R}$.

(i) The *affine hyperplane* with normal $\mathbf{N}$ and level b is $\mathcal{H}(\mathbf{N}, b) := \{\mathbf{x} \in \mathbb{R}^n : \langle \mathbf{N}, \mathbf{x} \rangle = b\}$.
(ii) The case $b = 0$ gives the *linear hyperplane* $\mathcal{H}(\mathbf{N}, 0) = \operatorname{span}(\mathbf{N})^{\perp}$.
(iii) The *closed half-spaces* determined by $\mathcal{H}(\mathbf{N}, b)$ are
$\mathcal{H}_+(\mathbf{N}, b) := \{\mathbf{x} \in \mathbb{R}^n : \langle \mathbf{N}, \mathbf{x} \rangle \geq b\}$ and $\mathcal{H}_-(\mathbf{N}, b) := \{\mathbf{x} \in \mathbb{R}^n : \langle \mathbf{N}, \mathbf{x} \rangle \leq b\}$.
The *open half-spaces* are obtained by replacing $\geq, \leq$ with $>, <$.
(iv) For $X \subseteq \mathbb{R}^n$ and $\star \in \{<, \leq, >, \geq\}$, use the *set-inequality convention* $\langle \mathbf{N}, X \rangle \star b \iff \langle \mathbf{N}, \mathbf{x} \rangle \star b\ \forall \mathbf{x} \in X$.
(v) An *orientation* of $\mathcal{H}(\mathbf{N}, b)$ is $\sigma(\mathbf{N}, b) := (\sigma\mathbf{N}, \sigma b)$ for $\sigma \in \{\pm 1\}$.
(vi) Sets $X, Y \subseteq \mathbb{R}^n$ are *strictly separated* by $\mathcal{H}(\mathbf{N}, b)$ if $\exists \sigma \in \{\pm 1\}$ such that $\langle \sigma\mathbf{N}, X \rangle < \sigma b < \langle \sigma\mathbf{N}, Y \rangle$
(vii) Sets $X, Y \subseteq \mathbb{R}^n$ are *strongly separated* by $\mathcal{H}(\mathbf{N}, b)$ if $\exists \sigma \in \{\pm 1\}$ and $\varepsilon > 0$ such that
$\langle \sigma\mathbf{N}, X \rangle < \sigma b - \varepsilon < \sigma b + \varepsilon < \langle \sigma\mathbf{N}, Y \rangle$.

Remark 2.7.8. *Hyperplane Geometry*

The half-spaces satisfy $\mathcal{H}_+(\mathbf{N}, b) \cap \mathcal{H}_-(\mathbf{N}, b) = \mathcal{H}(\mathbf{N}, b)$, and two open half-spaces with $\mathcal{H}(\mathbf{N}, b)$ partition $\mathbb{R}^n$.
Moreover, $\mathbf{N} \in \mathcal{H}(\mathbf{N}, \|\mathbf{N}\|^2)$, by Cauchy–Schwarz (Theorem 2.1.13), $\mathbf{N}$ is closest point of this hyperplane to $\mathbf{0}$.

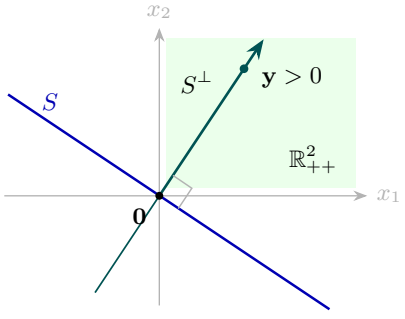


Figure 40: In $\mathbb{R}^2$, a negative-slope line S has $S^{\perp}$ meeting the strongly positive orthant.

Theorem 2.7.9. *Nearest Points and Separating a Point*
Let $C \subseteq \mathbb{R}^n$ be nonempty, closed, and convex.

(i) There exists a unique $\mathbf{N}_C \in C$ such that $\|\mathbf{N}_C\| = \min_{\mathbf{x}\in C} \|\mathbf{x}\|$.
(ii) If $\mathbf{b} \notin C$, set $D := C - \mathbf{b}$ and let $\mathbf{N} \in D$ satisfy $\|\mathbf{N}\| = \min_{\mathbf{z}\in D} \|\mathbf{z}\|$.
Then $\mathbf{N} \neq \mathbf{0}$ and $\langle \mathbf{N}, C\rangle \geq \|\mathbf{N}\|^2 + \langle \mathbf{N}, \mathbf{b}\rangle > \langle \mathbf{N}, \mathbf{b}\rangle$.
With $\alpha := \langle \mathbf{N}, \mathbf{b}\rangle + \frac{1}{2}\|\mathbf{N}\|^2$, the hyperplane $\mathcal{H}(\mathbf{N}, \alpha)$ strongly separates $\{\mathbf{b}\}$ and C.

Proof.

(i) Choose $\mathbf{x}_0 \in C$ and $s \geq \|\mathbf{x}_0\|$. Then $K := C \cap \overline{B}_s(\mathbf{0})$ is nonempty and compact by Theorem 2.2.39. Since $\|\cdot\|$ is continuous, Theorem 2.3.35 gives $\mathbf{N}_C \in K$ minimizing $\|\cdot\|$ on K. If $\mathbf{x} \in C$ and $\|\mathbf{x}\| < \|\mathbf{N}_C\|$, then $\mathbf{x} \in K$, a contradiction; hence $\mathbf{N}_C$ minimizes on C. If $\mathbf{u} \neq \mathbf{v}$ are minimizers of norm m, then $\mathbf{w} := (\mathbf{u}+\mathbf{v})/2 \in C$ and, by the parallelogram identity, Theorem 2.1.20, $\|\mathbf{w}\|^2 = m^2 - \frac{1}{4}\|\mathbf{u}-\mathbf{v}\|^2 < m^2$, a contradiction.

(ii) The set $D = C - \mathbf{b}$ is nonempty, closed, and convex, so (i) gives its unique minimum-norm vector $\mathbf{N}$. Since $\mathbf{b} \notin C$, $\mathbf{0} \notin D$, hence $\mathbf{N} \neq \mathbf{0}$. For $\mathbf{z} \in D$, define $\phi(t) := \|\mathbf{N} + t(\mathbf{z} - \mathbf{N})\|^2$, $0 \leq t \leq 1$. Minimality gives $\phi'_+(0) \geq 0$, hence $\langle \mathbf{N}, \mathbf{z}\rangle \geq \|\mathbf{N}\|^2$. Taking $\mathbf{z} = \mathbf{x} - \mathbf{b}$ yields $\langle \mathbf{N}, \mathbf{x}\rangle \geq \|\mathbf{N}\|^2 + \langle \mathbf{N}, \mathbf{b}\rangle$ for all $\mathbf{x} \in C$. For $\alpha := \langle \mathbf{N}, \mathbf{b}\rangle + \|\mathbf{N}\|^2/2$ and $\varepsilon := \|\mathbf{N}\|^2/4$, then $\langle \mathbf{N}, \mathbf{b}\rangle < \alpha - \varepsilon < \alpha + \varepsilon < \langle \mathbf{N}, C\rangle$, so $\mathcal{H}(\mathbf{N}, \alpha)$ strongly separates $\{\mathbf{b}\}$ and C.

Structural Commentary: The nearest point in $C - \mathbf{b}$ supplies the normal vector to a separating hyperplane.
Key Insight: Minimal norm turns a metric problem into the linear inequality $\langle \mathbf{N}, \mathbf{z}\rangle \geq \|\mathbf{N}\|^2$. □

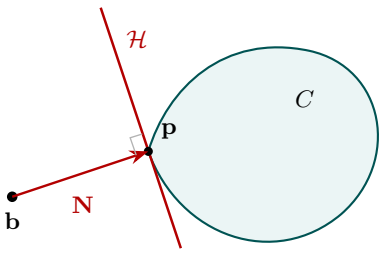


Figure 41: A nearest vector in $C - \mathbf{b}$ gives a normal direction and a strongly separating hyperplane.

Theorem 2.7.10. *Separating a Compact Convex Set from a Subspace*
Let $C \subseteq \mathbb{R}^n$ be nonempty, compact, and convex, and let $S \leq \mathbb{R}^n$ satisfy $C \cap S = \emptyset$.
Then there exists $\mathbf{N} \in S^\perp \setminus \{\mathbf{0}\}$ such that $\langle \mathbf{N}, C\rangle \geq \|\mathbf{N}\|^2 > 0$.
Consequently, $\mathcal{H}(\mathbf{N}, \|\mathbf{N}\|^2/2)$ strongly separates S and C.

Proof.

(i) Set $D := S + C$. Since S is closed in finite dimension, Proposition 2.7.6 gives D closed and convex. Also $\mathbf{0} \notin D$, for $\mathbf{0} = \mathbf{s} + \mathbf{c}$ would imply $\mathbf{c} = -\mathbf{s} \in C \cap S$.

(ii) Apply Theorem 2.7.9 to D and $\mathbf{0}$. Then $\exists \mathbf{N} \neq \mathbf{0}$ such that $\langle \mathbf{N}, \mathbf{s} + \mathbf{c}\rangle \geq \|\mathbf{N}\|^2$, $\mathbf{s} \in S$, $\mathbf{c} \in C$. Fix $\mathbf{c}_0 \in C$. If $a := \langle \mathbf{N}, \mathbf{s}\rangle \neq 0$ for some $\mathbf{s} \in S$, then $\langle \mathbf{N}, \lambda\mathbf{s} + \mathbf{c}_0\rangle = \lambda a + \langle \mathbf{N}, \mathbf{c}_0\rangle \to -\infty$ for λ of opposite sign and $|\lambda| \to \infty$, a contradiction. Hence $\mathbf{N} \in S^\perp$. Taking $\mathbf{s} = \mathbf{0}$ gives $\langle \mathbf{N}, C\rangle \geq \|\mathbf{N}\|^2 > 0$. Since $\langle \mathbf{N}, \mathbf{s}\rangle = 0$ for all $\mathbf{s} \in S$, the choice $\varepsilon := \|\mathbf{N}\|^2/4$ gives strong separation by $\mathcal{H}(\mathbf{N}, \|\mathbf{N}\|^2/2)$.

Structural Commentary: The sum $S + C$ is a closed convex set disjoint from $\mathbf{0}$, so point separation applies.
Key Insight: Since the separating inequality holds for all translates by S, scaling inside S forces $\mathbf{N}$ to annihilate S. □

Theorem 2.7.11. *Positive Orthant Alternatives for Subspaces*
Let $S \leq \mathbb{R}^n$. Then:

(i) $S \cap \mathbb{R}^n_+ = \emptyset \iff S^\perp \cap \mathbb{R}^n_{++} \neq \emptyset$.
(ii) $S \cap \mathbb{R}^n_{++} = \emptyset \iff S^\perp \cap \mathbb{R}^n_+ \neq \emptyset$.

Proof.

(i) ($\Rightarrow$) Let $\Delta := \operatorname{conv}\{\mathbf{e}_1, \ldots, \mathbf{e}_n\}$. By Proposition 2.7.5, Δ is compact and convex, and $\Delta \subseteq \mathbb{R}^n_+$. Hence $\Delta \cap S = \emptyset$. Theorem 2.7.10 gives $\mathbf{N} \in S^\perp \setminus \{\mathbf{0}\}$ with $\langle \mathbf{N}, \mathbf{e}_i\rangle \geq \|\mathbf{N}\|^2 > 0$, $1 \leq i \leq n$. QED.

(i) ($\Leftarrow$) If $\mathbf{x} \in S \cap \mathbb{R}^n_+$ and $\mathbf{N} \in S^\perp \cap \mathbb{R}^n_{++}$, then $\langle \mathbf{N}, \mathbf{x}\rangle > 0$, contradicting $\mathbf{N} \perp S$.

(*ii*) ($\Rightarrow$) If $S = \{\mathbf{0}\}$, then $\mathbf{e}_1 \in S^\perp \cap \mathbb{R}^n_+$. Assume $S \neq \{\mathbf{0}\}$.
Let $\mathbf{B}_1, \dots, \mathbf{B}_k$ be a basis of S and set $M = [\mathbf{B}_1 \ \cdots \ \mathbf{B}_k] \in \mathrm{Mat}(n, k, \mathbb{R})$, $\mathrm{rows}(M) = \{\mathbf{R}_1, \dots, \mathbf{R}_n\} \subseteq \mathbb{R}^k$.
For $\mathbf{a} = (a_1, \dots, a_n)$, $\mathbf{a} \in S^\perp \iff \sum_{i=1}^n a_i \mathbf{R}_i = \mathbf{0}$.
Set $K := \mathrm{conv}\{\mathbf{R}_1, \dots, \mathbf{R}_n\}$. Hence $S^\perp \cap \mathbb{R}^n_+ \neq \emptyset \iff \mathbf{0} \in K$. If $\mathbf{0} \notin K$, Theorem 2.7.9 gives $\mathbf{u} \in \mathbb{R}^k$ with $\langle \mathbf{u}, \mathbf{R}_i \rangle > 0$ for all i. Then $\mathbf{v} := M\mathbf{u} \in S$ has coordinates $v_i = \langle \mathbf{u}, \mathbf{R}_i \rangle > 0$, so $\mathbf{v} \in S \cap \mathbb{R}^n_{++}$, a contradiction. Thus $\mathbf{0} \in K$, and $S^\perp \cap \mathbb{R}^n_+ \neq \emptyset$.

Structural Commentary: Use the simplex and row hull as compact witnesses.
Key Insight: Missing positivity yields an orthogonal positive certificate. □

Theorem 2.7.12. *Homogeneous Positive Alternatives*
Let $A \in \mathrm{Mat}(m, n, \mathbb{R})$. Exactly one alternative holds in each pair:

(i) $\ker(A) \cap \mathbb{R}^n_{++} \neq \emptyset$ *or* $\mathrm{col}(A^t) \cap \mathbb{R}^n_+ \neq \emptyset$.
(ii) $\ker(A) \cap \mathbb{R}^n_+ \neq \emptyset$ *or* $\mathrm{col}(A^t) \cap \mathbb{R}^n_{++} \neq \emptyset$.

Proof.

(i) Let $S := \mathrm{row}(A) = \mathrm{col}(A^t) \subseteq \mathbb{R}^n$. By Principle 2.7.3, $\ker(A) = S^\perp$.
Moreover, $A^t\mathbf{v} \in \mathrm{col}(A^t) = S$ for every $\mathbf{v} \in \mathbb{R}^m$, and every vector of S is of this form.
Hence the first pair is equivalent to $S^\perp \cap \mathbb{R}^n_{++} \neq \emptyset$ or $S \cap \mathbb{R}^n_+ \neq \emptyset$.
Theorem 2.7.11(i) says exactly that these two alternatives are mutually exclusive and exhaustive.
(ii) With the same S, one has $\ker(A) \cap \mathbb{R}^n_+ = S^\perp \cap \mathbb{R}^n_+$, and $\{A^t\mathbf{v} : \mathbf{v} \in \mathbb{R}^m\} \cap \mathbb{R}^n_{++} = S \cap \mathbb{R}^n_{++}$.
Thus the second pair is equivalent to $S^\perp \cap \mathbb{R}^n_+ \neq \emptyset$ or $S \cap \mathbb{R}^n_{++} \neq \emptyset$.
Theorem 2.7.11(ii) says exactly that these two alternatives are mutually exclusive and exhaustive.

Structural Commentary: The row space gives the certificate side; the nullspace gives the solution side.
Key Insight: Positive solvability is an orthogonal-complement alternative. □

Method 2.7.13. *Certifying Positive Homogeneous Systems*
Let $A \in \mathrm{Mat}(m, n, \mathbb{R})$ and $S := \mathrm{row}(A) = \mathrm{col}(A^t)$.

(i) $\ker(A) \cap \mathbb{R}^n_{++} = \emptyset$ is certified by $\exists\, \mathbf{v} \in \mathbb{R}^m$ with $A^t\mathbf{v} > \mathbf{0}$.
(ii) $\ker(A) \cap \mathbb{R}^n_+ = \emptyset$ is certified by $\exists\, \mathbf{v} \in \mathbb{R}^m$ with $A^t\mathbf{v} \gg \mathbf{0}$.
(iii) Existence follows from row-space avoidance,
$S \cap \mathbb{R}^n_+ = \emptyset \Rightarrow \ker(A) \cap \mathbb{R}^n_{++} \neq \emptyset$; or $S \cap \mathbb{R}^n_{++} = \emptyset \Rightarrow \ker(A) \cap \mathbb{R}^n_+ \neq \emptyset$.

Lemma 2.7.14. *Nonnegative Matrix Cone*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(m, n, \mathbb{R})$ and $C_A := A(\mathbb{R}^n_{\geq 0}) = \left\{ \sum_{i=1}^n t_i \mathbf{a}_i : t_i \geq 0 \right\}$.
Then C_A is a closed convex cone in $\mathbb{R}^m$.

Proof. For $\mathbf{x}_r = A\mathbf{y}_r \in C_A$ with $\mathbf{y}_r \geq \mathbf{0}$.
For $0 \leq t \leq 1$, $a \geq 0$, $t\mathbf{x}_1 + (1-t)\mathbf{x}_2 = A(t\mathbf{y}_1 + (1-t)\mathbf{y}_2) \in C_A$, where $a\mathbf{x}_1 = A(a\mathbf{y}_1) \in C_A$, C_A is convex cone.
For closedness, choose for $\mathbf{x} \in C_A$ a representation $\mathbf{x} = \sum_{i \in I} t_i \mathbf{a}_i, \qquad t_i > 0$ with I minimal.
If $\{\mathbf{a}_i : i \in I\}$ were dependent, choose $(d_i)_{i \in I} \neq 0$ with $\sum_{i \in I} d_i \mathbf{a}_i = \mathbf{0}$.
Replacing d by $-d$ if necessary, some $d_i < 0$; then $\lambda := \min_{d_i < 0} \frac{t_i}{-d_i}$ makes $t_i + \lambda d_i \geq 0$ for all i and zero for at least one i, contradicting minimality. Hence the active columns are independent.
Thus $C_A = \bigcup_{\substack{I \subseteq \{1,\dots,n\} \\ \{\mathbf{a}_i : i \in I\} \text{ independent}}} K_I$, where $K_I := \left\{ \sum_{i \in I} t_i \mathbf{a}_i : t_i \geq 0 \right\}$.
For such I, the map $T_I : \mathbb{R}^I \to \mathrm{span}\{\mathbf{a}_i : i \in I\}$, where $T_I((t_i)) = \sum_{i \in I} t_i \mathbf{a}_i$, is a linear isomorphism, so $K_I = T_I(\mathbb{R}^I_{\geq 0})$ is closed in its finite-dimensional range, and the range is closed in $\mathbb{R}^m$.
Hence each K_I is closed in $\mathbb{R}^m$, and the finite union above is closed.
Structural Commentary: Reduce C_A to finitely many independent-column cones.
Key Insight: Minimal support turns generators independent, making each cone closed. □

Theorem 2.7.15. *Farkas's Lemma*
Let $A \in \mathrm{Mat}(m, n, \mathbb{R})$, $\mathbf{b} \in \mathbb{R}^m \setminus \{\mathbf{0}\}$, and $C_A := \{A\mathbf{y} : \mathbf{y} \in \mathbb{R}^n,\ \mathbf{y} \geq \mathbf{0}\}$.
Exactly one of the following holds:

(i) $\mathbf{b} \in C_A$, *equivalently* $\exists\, \mathbf{u} > \mathbf{0}$ *with* $A\mathbf{u} = \mathbf{b}$.
(ii) $\exists\, \mathbf{v} \in \mathbb{R}^m$ *such that* $A^t\mathbf{v} \leq \mathbf{0}$ *and* $\langle \mathbf{v}, \mathbf{b} \rangle > 0$.

Proof.

(i) If $\mathbf{b} = A\mathbf{u}$ with $\mathbf{u} \geq \mathbf{0}$ and $A^t\mathbf{v} \leq \mathbf{0}$, then $\langle \mathbf{v}, \mathbf{b} \rangle = \langle \mathbf{v}, A\mathbf{u} \rangle = \langle A^t\mathbf{v}, \mathbf{u} \rangle \leq 0$.

(ii) Suppose $\mathbf{b} \notin C_A$. By Lemma 2.7.14, C_A is closed and convex. Since $\mathbf{0} \in C_A$, Theorem 2.7.9, after changing sign if necessary, gives $\mathbf{v} \in \mathbb{R}^m$ and $\alpha \in \mathbb{R}$ such that $\langle \mathbf{v}, C_A \rangle < \alpha < \langle \mathbf{v}, \mathbf{b} \rangle$.
Thus $\alpha > 0$ and $\langle \mathbf{v}, \mathbf{b} \rangle > 0$. For every $\mathbf{y} \geq \mathbf{0}$, $\langle A^t\mathbf{v}, \mathbf{y} \rangle = \langle \mathbf{v}, A\mathbf{y} \rangle < \alpha$.
If $(A^t\mathbf{v})_i > 0$, taking $\mathbf{y} = \lambda \mathbf{e}_i$ and $\lambda \to \infty$ contradicts the bounded inequality. Hence $A^t\mathbf{v} \leq \mathbf{0}$.

Structural Commentary: Nonrepresentability means $\mathbf{b} \notin C_A$.
Key Insight: Separation turns nonmembership in a cone into the dual inequality $A^t\mathbf{v} \leq 0$. □

Example 2.7.16. *Why the Weak Inequality Matters*

The conclusion $A^t\mathbf{v} \leq \mathbf{0}$ cannot generally be strengthened to $A^t\mathbf{v} \ll \mathbf{0}$. Let $A = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & -1 \end{bmatrix}$ and $\mathbf{b} = \begin{bmatrix} -1 \\ 0 \end{bmatrix}$.

If $\mathbf{u} \geq \mathbf{0}$, then $(A\mathbf{u})_1 = u_1 \geq 0$, so $A\mathbf{u} \neq \mathbf{b}$.

The certificate $\mathbf{v} = \begin{bmatrix} -1 \\ 0 \end{bmatrix}$ satisfies $A^t\mathbf{v} = (-1, 0, 0)^t \leq \mathbf{0}$ and $\langle \mathbf{v}, \mathbf{b} \rangle = 1 > 0$.

For arbitrary $\mathbf{w} = (r, s)^t$, $A^t\mathbf{w} = (r, s, -s)^t$, so $A^t\mathbf{w} \ll \mathbf{0}$ would force $s < 0 < -s$, impossible.
Structural Commentary: Opposite columns force a zero in every certificate.
Key Insight: Boundary separation is sometimes the best possible separation.

3 Spectral and Operator Structure

Spectral and operator structure studies linear maps through invariant subspaces, polynomial relations, canonical forms, spectral theorems, and operator functions.

3.1 Eigenstructure and Invariant Subspaces

Spectral linear algebra understands an operator through subspaces it preserves. Eigenstructure, triangularization, diagonalization, and commuting families all refine this invariant-subspace viewpoint.

Definition 3.1.1. *Eigenvalue and Eigenvector*
Let $T \in \mathcal{L}(\mathcal{V})$. A scalar $\lambda \in \mathbb{F}$ is an *eigenvalue* of T if there exists $\mathbf{v} \in \mathcal{V}$ with $\mathbf{v} \neq \mathbf{0}$ and $T\mathbf{v} = \lambda\mathbf{v}$. Such a vector $\mathbf{v}$ is an *eigenvector* of T corresponding to λ.

Definition 3.1.2. *Eigenspace*
Let $T \in \mathcal{L}(\mathcal{V})$ and $\lambda \in \mathbb{F}$. The *eigenspace* of T corresponding to λ is $E_\lambda(T) := \ker(T - \lambda I_\mathcal{V})$. Thus λ is an eigenvalue of T if and only if $E_\lambda(T) \neq \{\mathbf{0}\}$. Equivalently, $E_\lambda(T)$ is the nullspace of the shifted operator $T - \lambda I_\mathcal{V}$.

Definition 3.1.3. *Spectrum and Geometric Multiplicity*
Let $T \in \mathcal{L}(\mathcal{V})$. The *spectrum* of T is $\operatorname{spec}(T) := \{\lambda \in \mathbb{F} : \lambda \text{ is an eigenvalue of } T\}$.
The *geometric multiplicity* of $\lambda \in \mathbb{F}$ is $\dim E_\lambda(T)$. For a square matrix $A \in \operatorname{Mat}(n, n, \mathbb{F})$, $E_\lambda(A) := \ker(A - \lambda I_n)$ and $\operatorname{spec}(A)$ denotes the spectrum of the induced operator on $\mathbb{F}^n$.

Proposition 3.1.4. *Eigenvectors as Kernel Vectors*
Let $T \in \mathcal{L}(\mathcal{V})$, $\lambda \in \mathbb{F}$. Vector $\mathbf{v}$ is eigenvector of T corresponding to λ if and only if $\mathbf{v} \in \ker(T - \lambda I_\mathcal{V})$, $\mathbf{v} \neq \mathbf{0}$.

Proof.

($\Rightarrow$) If $\mathbf{v}$ is an eigenvector corresponding to λ, then $\mathbf{v} \neq \mathbf{0}$ and $T\mathbf{v} = \lambda\mathbf{v}$, so $(T - \lambda I_\mathcal{V})\mathbf{v} = \mathbf{0}$.
($\Leftarrow$) If $\mathbf{v} \in \ker(T - \lambda I_\mathcal{V})$ and $\mathbf{v} \neq \mathbf{0}$, then $(T - \lambda I_\mathcal{V})\mathbf{v} = \mathbf{0}$, so $T\mathbf{v} = \lambda\mathbf{v}$.

Structural Commentary: Eigenvectors are kernels of shifted operators.
Key Insight: Subtract $\lambda I_\mathcal{V}$ to convert $T\mathbf{v} = \lambda\mathbf{v}$ into homogeneous form. □

Theorem 3.1.5. *Equivalent Conditions for an Eigenvalue*
Let $\mathcal{V}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V})$, and let $\lambda \in \mathbb{F}$. The following are equivalent:

(i) λ is an eigenvalue of T;
(ii) $T - \lambda I_\mathcal{V}$ is not injective;
(iii) $T - \lambda I_\mathcal{V}$ is not surjective;
(iv) $T - \lambda I_\mathcal{V}$ is not invertible.

Proof.

(*i*) $\Rightarrow$ (*ii*) If λ is an eigenvalue, Proposition 3.1.4 gives nonzero vector in $\ker(T - \lambda I_\mathcal{V})$, so $T - \lambda I_\mathcal{V}$ not injective.
(*ii*) $\Rightarrow$ (*i*) If $T - \lambda I_\mathcal{V}$ not injective, then $\ker(T - \lambda I_\mathcal{V})$ contains nonzero vector, so λ is eigenvalue (Prop 3.1.4).

(*ii*) $\Rightarrow$ (*iii*) If $T - \lambda I_{\mathcal{V}}$ were surjective, Theorem 1.3.49 would make it injective, a contradiction.
(*iii*) $\Rightarrow$ (*iv*) If $T - \lambda I_{\mathcal{V}}$ were invertible, then it would be surjective, a contradiction.
(*iv*) $\Rightarrow$ (*ii*) If $T - \lambda I_{\mathcal{V}}$ were injective, Theorem 1.3.49 would make it invertible, a contradiction.

Structural Commentary: In finite dimensions, noninvertibility has equivalent forms.
Key Insight: A scalar λ is spectral precisely when $T - \lambda I_{\mathcal{V}}$ has a nontrivial kernel. □

Theorem 3.1.6. *Eigenvectors for Distinct Eigenvalues Are Independent*
Let $T \in \mathcal{L}(\mathcal{V})$. Suppose $\lambda_1, \ldots, \lambda_m$ are distinct eigenvalues of T, and let $\mathbf{v}_1, \ldots, \mathbf{v}_m$ be corresponding eigenvectors. Then $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ is linearly independent.

Proof. Suppose the list is linearly dependent. Let k be smallest positive integer s.t. $\mathbf{v}_k \in \operatorname{span}(\mathbf{v}_1, \ldots, \mathbf{v}_{k-1})$. Then $k \geq 2$, $(\mathbf{v}_1, \ldots, \mathbf{v}_{k-1})$ is linearly independent, and $\exists$ scalars $a_1, \ldots, a_{k-1}$ s.t. $\mathbf{v}_k = a_1\mathbf{v}_1 + \cdots + a_{k-1}\mathbf{v}_{k-1}$. Applying T gives $\lambda_k\mathbf{v}_k = a_1\lambda_1\mathbf{v}_1 + \cdots + a_{k-1}\lambda_{k-1}\mathbf{v}_{k-1}$. Multiplying first displayed equation by λ_k and subtracting gives $0 = a_1(\lambda_k - \lambda_1)\mathbf{v}_1 + \cdots + a_{k-1}(\lambda_k - \lambda_{k-1})\mathbf{v}_{k-1}$. Since $\mathbf{v}_1, \ldots, \mathbf{v}_{k-1}$ are linearly independent and $\lambda_k \neq \lambda_i$ for $i < k$, all $a_i = 0$. Hence $\mathbf{v}_k = \mathbf{0}$, contradicting that $\mathbf{v}_k$ is an eigenvector.
Structural Commentary: A smallest dependence isolates the first new eigenvector that depends on earlier ones. Applying T and subtracting separates the distinct eigenvalue factors.
Key Insight: Distinct eigenvalues make the same dependence relation impossible after applying T. □

Corollary 3.1.7. *Number of Distinct Eigenvalues*
If $\mathcal{V}$ is finite-dimensional and $T \in \mathcal{L}(\mathcal{V})$, then T has at most $\dim(\mathcal{V})$ distinct eigenvalues.

Proof. If T had more than $\dim(\mathcal{V})$ distinct eigenvalues, choose $m > \dim(\mathcal{V})$ of them, one eigenvector for each. By Theorem 3.1.6, these eigenvectors form linearly independent list of length m, contradicting Theorem 1.1.24.
Structural Commentary: Distinct eigenvalues give independent directions.
Key Insight: Count one eigenvector per eigenspace. □

Corollary 3.1.8. *Invertibility and the Zero Eigenvalue*
Let $\mathcal{V}$ be finite-dimensional and let $T \in \mathcal{L}(\mathcal{V})$. Then T is invertible if and only if $0 \notin \operatorname{spec}(T)$.

Proof. Apply Theorem 3.1.5 with $\lambda = 0$. Then 0 is an eigenvalue exactly when $T - 0I_{\mathcal{V}} = T$ is not invertible.
Structural Commentary: The zero eigenvalue is the spectral form of singularity.
Key Insight: Invertibility is equivalent to having trivial kernel, and $E_0(T) = \ker T$. □

Corollary 3.1.9. *Matrix Invertibility and the Spectrum*
Let $A \in \operatorname{Mat}(n, n, \mathbb{F})$. Then A is invertible if and only if $0 \notin \operatorname{spec}(A)$.

Proof.
($\Rightarrow$) If A is invertible, then $\ker(A) = \{\mathbf{0}\}$, so 0 is not an eigenvalue of A.
($\Leftarrow$) If $0 \notin \operatorname{spec}(A)$, then $\ker(A) = \{\mathbf{0}\}$. Corollary 1.3.73 gives full rank, hence A is invertible.

Structural Commentary: Matrix invertibility is the coordinate form of trivial kernel.
Key Insight: A zero eigenvalue is exactly a nonzero solution of $A\mathbf{x} = \mathbf{0}$. □

Definition 3.1.10. *Invariant Subspace*
Let $T \in \mathcal{L}(\mathcal{V})$. A subspace $\mathcal{U}$ of $\mathcal{V}$ is *invariant under T* if $T\mathbf{u} \in \mathcal{U}$ for every $\mathbf{u} \in \mathcal{U}$.
The set of all T-invariant subspaces is $\operatorname{Inv}(T) := \{\mathcal{U} \leq \mathcal{V} : T\mathcal{U} \subseteq \mathcal{U}\}$.
Key Insight: $\mathcal{U}$ is invariant under T iff $T|_{\mathcal{U}} \in \mathcal{L}(\mathcal{U})$.

Proposition 3.1.11. *One-Dimensional Invariant Subspaces and Eigenvectors*
Let $T \in \mathcal{L}(\mathcal{V})$, $\mathbf{v} \neq \mathbf{0}$. Then $\operatorname{span}(\mathbf{v})$ is invariant under T if and only if there exists $\lambda \in \mathbb{F}$ such that $T\mathbf{v} = \lambda\mathbf{v}$.

Proof.
($\Rightarrow$) If $\operatorname{span}(\mathbf{v})$ is invariant under T, then $T\mathbf{v} \in \operatorname{span}(\mathbf{v})$, so $T\mathbf{v} = \lambda\mathbf{v}$ for some $\lambda \in \mathbb{F}$.
($\Leftarrow$) If $T\mathbf{v} = \lambda\mathbf{v}$ and $\mathbf{u} = a\mathbf{v} \in \operatorname{span}(\mathbf{v})$, then $T\mathbf{u} = aT\mathbf{v} = a\lambda\mathbf{v} \in \operatorname{span}(\mathbf{v})$.

Structural Commentary: One-dimensional invariance means rescaling one direction.
Key Insight: In dimension one, staying inside the subspace means being a scalar multiple of the same vector. □

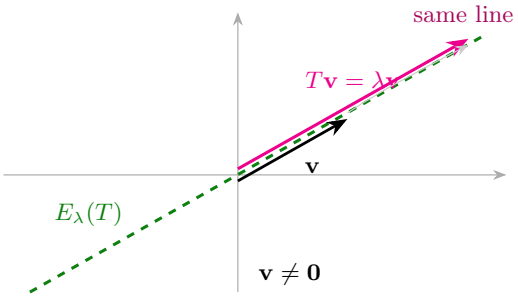


Figure 42: An eigendirection is a one-dimensional invariant subspace.

Proposition 3.1.12. *Invariant Subspaces and Block Form*
Let $T \in \mathcal{L}(\mathcal{V})$ and suppose $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$.

(i) If $\mathcal{U}$ is invariant under T, then in a basis adapted to $\mathcal{U} \oplus \mathcal{W}$, matrix of T has upper block-triangular form

$$\begin{bmatrix} A & B \\ 0 & D \end{bmatrix}.$$

(ii) If both $\mathcal{U}$ and $\mathcal{W}$ are invariant under T, then the matrix of T has the block-diagonal form

$$\begin{bmatrix} A & 0 \\ 0 & D \end{bmatrix}.$$

Proof. Choose basis $(\mathbf{u}_1, \ldots, \mathbf{u}_p)$ of $\mathcal{U}$ and basis $(\mathbf{w}_1, \ldots, \mathbf{w}_q)$ of $\mathcal{W}$, and use concatenation as basis of $\mathcal{V}$.

(i) If $\mathcal{U}$ is invariant, then each $T\mathbf{u}_j$ lies in $\mathcal{U}$, so its coordinates along $\mathbf{w}_1, \ldots, \mathbf{w}_q$ are zero. These entries form the lower-left block.

(ii) If $\mathcal{W}$ is also invariant, then each $T\mathbf{w}_j$ lies in $\mathcal{W}$, so its coordinates along $\mathbf{u}_1, \ldots, \mathbf{u}_p$ are zero. These entries form the upper-right block.

Structural Commentary: Block form records how an operator moves vectors between chosen summands. Invariant summands force the corresponding off-diagonal movement to vanish.
Key Insight: Zero blocks are coordinate evidence of invariant subspaces. □

Definition 3.1.13. *Reducing Pair*
Let $T \in \mathcal{L}(\mathcal{V})$. A pair of subspaces $(\mathcal{U}, \mathcal{W})$ *reduces* T if $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$; $T(\mathcal{U}) \subseteq \mathcal{U}$; and $T(\mathcal{W}) \subseteq \mathcal{W}$.
Then $T = T|_{\mathcal{U}} \oplus T|_{\mathcal{W}}$ with respect to $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$.

Proposition 3.1.14. *Reducing Pairs and Projections*
Let $T \in \mathcal{L}(\mathcal{V})$, let $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$, and let $\rho = \rho_{\mathcal{U},\mathcal{W}}$. Then $(\mathcal{U}, \mathcal{W})$ reduces $T \iff T\rho = \rho T$.

Proof.

($\Rightarrow$) Let $\mathbf{v} = \mathbf{u} + \mathbf{w}$ with $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{W}$.
Since both summands are invariant, $T\rho\mathbf{v} = T\mathbf{u}$ and $\rho T\mathbf{v} = \rho(T\mathbf{u} + T\mathbf{w}) = T\mathbf{u}$. Thus $T\rho = \rho T$.

($\Leftarrow$) Suppose $T\rho = \rho T$. If $\mathbf{u} \in \mathcal{U}$, then $\rho(T\mathbf{u}) = T(\rho\mathbf{u}) = T\mathbf{u}$, so $T\mathbf{u} \in \operatorname{range}(\rho) = \mathcal{U}$.
If $\mathbf{w} \in \mathcal{W}$, then $\rho(T\mathbf{w}) = T(\rho\mathbf{w}) = \mathbf{0}$, so $T\mathbf{w} \in \ker(\rho) = \mathcal{W}$. Hence $(\mathcal{U}, \mathcal{W})$ reduces T.

Structural Commentary: Commutation with ρ means T preserves the two projection components.
Key Insight: Reducing pairs are exactly decompositions whose projection commutes with T. □

Proposition 3.1.15. *Invariant Subspaces and Algebraic Projections*
Let $T \in \mathcal{L}(\mathcal{V})$, let $\mathcal{V} = \mathcal{U} \oplus \mathcal{W}$, and let $\rho = \rho_{\mathcal{U},\mathcal{W}}$. Then $\mathcal{U} \in \operatorname{Inv}(T) \iff \rho T\rho = T\rho$.
Consequently, if $\mathcal{U} \in \operatorname{Inv}(T)$, then $\rho_{\mathcal{U},\mathcal{Z}} T \rho_{\mathcal{U},\mathcal{Z}} = T\rho_{\mathcal{U},\mathcal{Z}}$ for every complement $\mathcal{Z}$ of $\mathcal{U}$.

Proof. For $\mathbf{v} = \mathbf{u} + \mathbf{w}$ with $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{W}$, $\mathcal{U} \in \operatorname{Inv}(T)$ gives $(\rho T\rho - T\rho)\mathbf{v} = \rho T\mathbf{u} - T\mathbf{u} = \mathbf{0}$.
Thus $\rho T\rho = T\rho$. Conversely, if $\rho T\rho = T\rho$ and $\mathbf{u} \in \mathcal{U}$, then $\rho T\mathbf{u} = \rho T\rho\mathbf{u} = T\rho\mathbf{u} = T\mathbf{u}$.
Hence $T\mathbf{u} \in \operatorname{range}(\rho) = \mathcal{U}$, so $\mathcal{U} \in \operatorname{Inv}(T)$.
Structural Commentary: A projection onto $\mathcal{U}$ tests whether the $\mathcal{W}$-component of $T\mathbf{u}$ vanishes.
Key Insight: Invariance says applying T after projecting into $\mathcal{U}$ needs no further correction by ρ. □

Perspective 3.1.16. *Matrix Viewpoint on Structure*
Basic block matrices already capture important geometry. Upper block-triangular form records the presence of an invariant subspace; block diagonal form records a genuine decomposition into independent invariant pieces.

Definition 3.1.17. *Diagonal and Upper-Triangular Matrices*
The *diagonal* of a square matrix is its upper-left to lower-right entries. A square matrix is *diagonal* if all off-diagonal entries are zero, and *upper triangular* if all entries below the diagonal are zero.

Theorem 3.1.18. *Invariant Subspaces, Blocks, and Projections*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$, let $\mathcal{U}$ be an r-dimensional subspace of $\mathbb{F}^n$ with $1 \le r \le n-1$, and let $Q = [Q_1\ Q_2]$ be unitary, where the columns of Q_1 and Q_2 are orthonormal bases of $\mathcal{U}$ and $\mathcal{U}^\perp$, respectively. Let $P = Q_1 Q_1^*$ be the orthogonal projection onto $\mathcal{U}$. The following are equivalent:

(i) $\mathcal{U}$ *is invariant under* A*;*
(ii) $Q^*AQ = \begin{bmatrix} B & X \\ 0 & C \end{bmatrix}$ *for some compatible matrices* B, X, C*;*
(iii) $PAP = AP$.

Proof. Let $M = Q^*AQ = \begin{bmatrix} B & X \\ Y & C \end{bmatrix}$ and $E = Q^*PQ = I_r \oplus 0_{n-r}$. Since Q is unitary and $P = Q_1Q_1^*$, the matrix of P in the Q-coordinates is E. Also $Y = Q_2^*AQ_1$, and $\ker(Q_2^*) = \mathrm{col}(Q_1) = \mathcal{U}$.

(i) $\Rightarrow$ *(ii)* If $A\mathcal{U} \subseteq \mathcal{U}$, then each column of AQ_1 lies in $\mathcal{U} = \ker(Q_2^*)$, so $Y = Q_2^*AQ_1 = 0$.

(ii) $\Rightarrow$ *(iii)* If $Y = 0$, then $EME = \begin{bmatrix} B & 0 \\ 0 & 0 \end{bmatrix}$ and $ME = \begin{bmatrix} B & 0 \\ 0 & 0 \end{bmatrix}$.
Thus $Q^*(PAP)Q = EME = ME = Q^*(AP)Q$, so $PAP = AP$.

(iii) $\Rightarrow$ *(i)* If $PAP = AP$, then $EME = ME$. Since $EME = \begin{bmatrix} B & 0 \\ 0 & 0 \end{bmatrix}$ and $ME = \begin{bmatrix} B & 0 \\ Y & 0 \end{bmatrix}$, we get $Y = 0$.
Hence the columns of AQ_1 lie in $\ker(Q_2^*) = \mathcal{U}$. Since the columns of Q_1 span $\mathcal{U}$, linearity gives $A\mathcal{U} \subseteq \mathcal{U}$.

Structural Commentary: In adapted coordinates, the lower-left block is the component of $A\mathcal{U}$ in $\mathcal{U}^\perp$.
Key Insight: $PAP = AP$ exactly kills this component. □

Corollary 3.1.19. *Reducing Subspaces and Commuting Projections*
Maintain the notation of Theorem 3.1.18. The following are equivalent:

(i) $\mathcal{U}$ *is invariant under both* A *and* A^**;*
(ii) $Q^*AQ = B \oplus C$ *for some* $B \in \mathrm{Mat}(r,r,\mathbb{F})$ *and* $C \in \mathrm{Mat}(n-r, n-r, \mathbb{F})$*;*
(iii) $PA = AP$.

Proof. Let $M = Q^*AQ = \begin{bmatrix} B & X \\ Y & C \end{bmatrix}$ and $E = Q^*PQ = I_r \oplus 0_{n-r}$.
Then $Q^*A^*Q = M^*$, $Q^*(PA)Q = EM$, and $Q^*(AP)Q = ME$.

(i) $\Rightarrow$ *(ii)* By Theorem 3.1.18, A-invariance gives $Y = 0$. Applying the same theorem to A^* gives that the lower-left block of M^* is zero, namely $X^* = 0$, so $X = 0$. Thus $M = B \oplus C$.

(ii) $\Rightarrow$ *(iii)* If $M = B \oplus C$, then $EM = ME$, hence $PA = AP$.

(iii) $\Rightarrow$ *(i)* If $PA = AP$, then $EM = ME$. Since $EM = \begin{bmatrix} B & X \\ 0 & 0 \end{bmatrix}$ and $ME = \begin{bmatrix} B & 0 \\ Y & 0 \end{bmatrix}$, we get $X = Y = 0$.
The equation $Y = 0$ gives A-invariance by Theorem 3.1.18. Since the lower-left block of M^* is $X^* = 0$, applying the same theorem to A^* gives A^*-invariance.

Structural Commentary: Reducing subspaces are exactly block diagonal decompositions in adapted coordinates.
Key Insight: Commutation with P kills both off-diagonal blocks. □

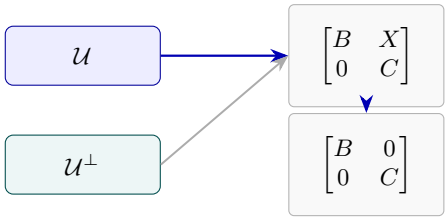


Figure 43: Adapted coordinates: invariance gives upper block form; $PA = AP$ gives block diagonal form.

Definition 3.1.20. *Restriction and Quotient Operators*
Let $T \in \mathcal{L}(\mathcal{V})$ and let $\mathcal{U}$ be invariant under T.

(i) The *restriction operator* $T|_{\mathcal{U}} \in \mathcal{L}(\mathcal{U})$ is defined by $T|_{\mathcal{U}}(\mathbf{u}) = T\mathbf{u}$ for $\mathbf{u} \in \mathcal{U}$.
(ii) The *quotient operator* $T/\mathcal{U} \in \mathcal{L}(\mathcal{V}/\mathcal{U})$ is defined by $(T/\mathcal{U})(\mathbf{v} + \mathcal{U}) = T\mathbf{v} + \mathcal{U}$, $\mathbf{v} \in \mathcal{V}$.

Proposition 3.1.21. *Restriction and Quotient Operators Are Well-Defined*
Let $T \in \mathcal{L}(\mathcal{V})$ and let $\mathcal{U}$ be invariant under T. Then $T|_{\mathcal{U}} \in \mathcal{L}(\mathcal{U})$ and $T/\mathcal{U} \in \mathcal{L}(\mathcal{V}/\mathcal{U})$.

Proof. Since $\mathcal{U}$ is invariant under T, $T\mathbf{u} \in \mathcal{U}$ whenever $\mathbf{u} \in \mathcal{U}$, so $T|_{\mathcal{U}}$ maps $\mathcal{U}$ into itself. Its linearity is inherited from T. For the quotient operator, suppose $\mathbf{v} + \mathcal{U} = \mathbf{w} + \mathcal{U}$. Then $\mathbf{v} - \mathbf{w} \in \mathcal{U}$.
Invariance gives $T(\mathbf{v} - \mathbf{w}) \in \mathcal{U}$, so $T\mathbf{v} - T\mathbf{w} \in \mathcal{U}$ and hence $T\mathbf{v} + \mathcal{U} = T\mathbf{w} + \mathcal{U}$. Thus $T/\mathcal{U}$ is well-defined.
For $a \in \mathbb{F}$, $(T/\mathcal{U})(a(\mathbf{v} + \mathcal{U}) + (\mathbf{w} + \mathcal{U})) = T(a\mathbf{v} + \mathbf{w}) + \mathcal{U} = a(T\mathbf{v} + \mathcal{U}) + (T\mathbf{w} + \mathcal{U})$, so $T/\mathcal{U}$ is linear.
Structural Commentary: Restriction stays inside $\mathcal{U}$; quotient collapses $\mathcal{U}$.
Key Insight: Invariance makes both operators well-defined. □

Perspective 3.1.22. *Two Smaller Views of an Operator*
When $\mathcal{U}$ is invariant under T, the operators $T|_{\mathcal{U}}$ and $T/\mathcal{U}$ split attention between what happens inside $\mathcal{U}$ and what remains after $\mathcal{U}$ is collapsed. These two views often simplify an operator, even though they need not determine the whole operator by themselves.

Theorem 3.1.23. *Complex Operators Have Eigenvalues*
Every operator on a finite-dimensional, nonzero complex vector space has an eigenvalue.

Proof. Let $T \in \mathcal{L}(\mathcal{V})$. By Proposition 3.2.4, T has a nonconstant annihilating polynomial. Since $\mathcal{V}$ is complex and nonzero, Theorem 3.2.36 implies that one root of this polynomial is an eigenvalue of T.
Structural Commentary: Finite dimension gives an annihilating polynomial; over $\mathbb{C}$, it factors linearly.
Key Insight: One noninjective factor produces an eigenvector. □

Definition 3.1.24. *Commutant*
For $\emptyset \neq \mathcal{F} \subseteq \mathrm{Mat}(n, n, \mathbb{F})$, the *commutant* of $\mathcal{F}$ is $\mathrm{Comm}(\mathcal{F}) := \{X \in \mathrm{Mat}(n, n, \mathbb{F}) : XA = AX \ \forall A \in \mathcal{F}\}$.
For a single matrix A, write $\mathrm{Comm}(A)$.

Lemma 3.1.25. *Commuting with a Block Scalar Diagonal Matrix*
Let $\lambda_1, \dots, \lambda_k \in \mathbb{F}$ and let $\Lambda = \lambda_1 I_{n_1} \oplus \cdots \oplus \lambda_k I_{n_k}$.
Partition $A = [A_{ij}] \in \mathrm{Mat}(n, n, \mathbb{F})$ conformally with Λ, so $A_{ij} \in \mathrm{Mat}(n_i, n_j, \mathbb{F})$.

(i) If A is block diagonal, then $A\Lambda = \Lambda A$.
(ii) If $A\Lambda = \Lambda A$ and $\lambda_1, \dots, \lambda_k$ are distinct, then A is block diagonal.

Proof. Block multiplication gives $(A\Lambda)_{ij} = \lambda_j A_{ij}$ and $(\Lambda A)_{ij} = \lambda_i A_{ij}$. If A is block diagonal, then $A_{ij} = 0$ for $i \neq j$, and these two expressions agree for every block. Conversely, if $A\Lambda = \Lambda A$, then $(\lambda_j - \lambda_i)A_{ij} = 0$ for every i, j. Distinctness of the λ_i forces $A_{ij} = 0$ whenever $i \neq j$, so A is block diagonal.
Structural Commentary: Distinct scalar blocks force commuting matrices to respect the block decomposition.
Key Insight: Distinct block scalars make every off-diagonal block visible and force it to vanish. □

Example 3.1.26. *Distinctness Is Necessary*
The distinctness hypothesis in Lemma 3.1.25 cannot be dropped. If $\lambda_1 = \cdots = \lambda_k = 1$, then $\Lambda = I_n$ commutes with every matrix in $\mathrm{Mat}(n, n, \mathbb{F})$, including matrices that are not block diagonal for the chosen partition.

Theorem 3.1.27. *Common Eigenvector for a Commuting Family*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$, and let $\mathcal{F} \subseteq \mathcal{L}(\mathcal{V})$ be a nonempty commuting family: $ST = TS$ for all $S, T \in \mathcal{F}$. Then there exists $\mathbf{v} \neq \mathbf{0}$ that is an eigenvector of every operator in $\mathcal{F}$.

Proof. First suppose $\mathcal{F} = \{T_1, \dots, T_m\}$ is finite, and induct on m. The case $m = 1$ is Theorem 3.1.23.
For $m > 1$, choose $\lambda \in \mathrm{spec}(T_1)$ and set $\mathcal{U} = E_\lambda(T_1)$. Then $\mathcal{U} \neq \{\mathbf{0}\}$.
If $\mathbf{u} \in \mathcal{U}$ and $j \geq 2$, commutation gives $(T_1 - \lambda I_{\mathcal{V}})T_j\mathbf{u} = T_j(T_1 - \lambda I_{\mathcal{V}})\mathbf{u} = \mathbf{0}$, so $\mathcal{U}$ is invariant under $T_2, \dots, T_m$.
Their restrictions to $\mathcal{U}$ still commute, and by induction have a common eigenvector $\mathbf{v} \in \mathcal{U}$, $\mathbf{v} \neq \mathbf{0}$.
Since $\mathbf{v} \in E_\lambda(T_1)$, it is also an eigenvector of T_1. For general $\mathcal{F}$, the subspace $\mathrm{span}(\mathcal{F})$ of $\mathcal{L}(\mathcal{V})$ is finite-dimensional. Choose a basis $T_1, \dots, T_m$ of this span from the spanning set $\mathcal{F}$.
The finite case gives a nonzero $\mathbf{v}$ and scalars μ_j with $T_j\mathbf{v} = \mu_j\mathbf{v}$. If $S \in \mathcal{F}$, write $S = \sum_j a_j T_j$.
Then $S\mathbf{v} = (\sum_j a_j\mu_j)\mathbf{v}$, so $\mathbf{v}$ is an eigenvector of every member of $\mathcal{F}$.
Structural Commentary: Commutation makes one eigenspace invariant, reducing the family by one.
Key Insight: Find the common eigenvector inside an eigenspace. □

Theorem 3.1.28. *Criteria for an Upper-Triangular Operator Matrix*
Let $T \in \mathcal{L}(\mathcal{V})$ and let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$ be a basis of $\mathcal{V}$. The following are equivalent:

(i) $[T]_{\beta \leftarrow \beta}$ is upper triangular;
(ii) $T\mathbf{v}_j \in \mathrm{span}(\mathbf{v}_1, \dots, \mathbf{v}_j)$ for each $j = 1, \dots, n$;
(iii) $\mathrm{span}(\mathbf{v}_1, \dots, \mathbf{v}_j)$ is invariant under T for each $j = 1, \dots, n$.

Proof.

(*i*) $\Rightarrow$ (*ii*) The jth column of $[T]_{\beta\leftarrow\beta}$ is the coordinate column of $T\mathbf{v}_j$ in the basis β. Since the matrix is upper triangular, this column has zero entries below position j, so $T\mathbf{v}_j \in \operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j)$.

(*ii*) $\Rightarrow$ (*i*) If $T\mathbf{v}_j \in \operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j)$ for every j, then the jth coordinate column of $T\mathbf{v}_j$ has zero entries below position j. Hence $[T]_{\beta\leftarrow\beta}$ is upper triangular.

(*ii*) $\Rightarrow$ (*iii*) If $\mathbf{u} \in \operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j)$, then linearity gives $T\mathbf{u}$ as a linear combination of $T\mathbf{v}_1,\dots,T\mathbf{v}_j$, all of which lie in $\operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j)$. Hence this span is invariant.

(*iii*) $\Rightarrow$ (*ii*) Apply invariance of $\operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j)$ to $\mathbf{v}_j$.

Structural Commentary: Upper triangularity records an invariant flag.
Key Insight: Columns below the diagonal vanish exactly when $T\mathbf{v}_j$ uses only earlier basis vectors and itself. □

Corollary 3.1.29. *Upper-Triangular Form in an Orthonormal Basis*
Let $\mathcal{V}$ be a finite-dimensional inner-product space and let $T \in \mathcal{L}(\mathcal{V})$. If T has an upper-triangular matrix with respect to some basis of $\mathcal{V}$, then T has an upper-triangular matrix with respect to some orthonormal basis of $\mathcal{V}$.

Proof. Let $\beta = (\mathbf{v}_1,\dots,\mathbf{v}_n)$ be a basis for which $[T]_{\beta\leftarrow\beta}$ is upper triangular. By Theorem 3.1.28, each subspace

$$\mathcal{U}_j = \operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j), \qquad j = 1,\dots,n,$$

is invariant under T. Apply Gram–Schmidt to β and obtain an orthonormal basis $\gamma = (\mathbf{e}_1,\dots,\mathbf{e}_n)$ with

$$\operatorname{span}(\mathbf{e}_1,\dots,\mathbf{e}_j) = \operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j) = \mathcal{U}_j.$$

Thus each $\operatorname{span}(\mathbf{e}_1,\dots,\mathbf{e}_j)$ is invariant under T, and Theorem 3.1.28 shows that $[T]_{\gamma\leftarrow\gamma}$ is upper triangular.
Structural Commentary: Upper triangularity is preservation of a nested flag of subspaces. Gram–Schmidt preserves that flag while making the basis orthonormal.
Key Insight: Triangular form depends on the flag, not on the particular basis used inside the flag. □

Theorem 3.1.30. *Complex Upper Triangularisation*
Let $\mathcal{V}$ be a finite-dimensional complex vector space and let $T \in \mathcal{L}(\mathcal{V})$. Then T has an upper-triangular matrix with respect to some basis of $\mathcal{V}$.

Proof. Induct on $n = \dim(\mathcal{V})$. If $n \le 1$, the result is immediate.
Suppose $n > 1$ and the result holds in smaller dimensions. Choose an eigenvector $\mathbf{v}_1$ of T by Theorem 3.1.23, and set $\mathcal{U} = \operatorname{span}(\mathbf{v}_1)$. Then $\mathcal{U}$ is T-invariant, so the quotient operator $T/\mathcal{U}$ is defined on $\mathcal{V}/\mathcal{U}$, which has dimension $n-1$. By induction, choose a basis $(\mathbf{v}_2+\mathcal{U},\dots,\mathbf{v}_n+\mathcal{U})$ of $\mathcal{V}/\mathcal{U}$ with respect to which $T/\mathcal{U}$ is upper triangular, and choose representatives $\mathbf{v}_2,\dots,\mathbf{v}_n \in \mathcal{V}$.
The lifted list $(\mathbf{v}_1,\dots,\mathbf{v}_n)$ is a basis of $\mathcal{V}$: any relation descends to the quotient basis, forcing all coefficients of $\mathbf{v}_2,\dots,\mathbf{v}_n$ to be zero, and then the coefficient of $\mathbf{v}_1$ is zero. Now $T\mathbf{v}_1 \in \operatorname{span}(\mathbf{v}_1)$.
For $j \ge 2$, upper triangularity of $T/\mathcal{U}$ gives $T\mathbf{v}_j + \mathcal{U} = (T/\mathcal{U})(\mathbf{v}_j+\mathcal{U}) \in \operatorname{span}(\mathbf{v}_2+\mathcal{U},\dots,\mathbf{v}_j+\mathcal{U})$, so $T\mathbf{v}_j \in \operatorname{span}(\mathbf{v}_1,\dots,\mathbf{v}_j)$. Hence Theorem 3.1.28 gives an upper-triangular matrix for T in the basis $(\mathbf{v}_1,\dots,\mathbf{v}_n)$.
Structural Commentary: Lift an upper-triangular quotient basis above an eigenline.
Key Insight: The quotient builds the remaining invariant flag. □

Definition 3.1.31. *Almost Upper-Triangular Real Matrices*
Let $A \in \operatorname{Mat}(n,n,\mathbb{R})$. The matrix A is *almost upper triangular* if there exist $d_1,\dots,d_m \in \{1,2\}$ with $\sum_j d_j = n$ such that, in the corresponding block form $A = [A_{ij}]_{i,j=1}^{m}$,

(i) $A_{ij} = 0$ for $i > j$;

(ii) each diagonal block is either $A_{jj} = [\lambda_j]$, $\lambda_j \in \mathbb{R}$, or $A_{jj} = \begin{bmatrix} a_j & -b_j \\ b_j & a_j \end{bmatrix}$, $a_j, b_j \in \mathbb{R}$, $b_j \neq 0$.

The matrix A is in *real Schur block form* if only (i) holds for some such block sizes $d_j \in \{1,2\}$.

Theorem 3.1.32. *Real Almost Upper Triangularisation*
Let $\dim_{\mathbb{R}} \mathcal{V} < \infty$ and let $T \in \mathcal{L}(\mathcal{V})$. Then $\exists$ a basis β of $\mathcal{V}$ such that $[T]_{\beta\leftarrow\beta}$ is almost upper triangular.

Proof. Induct on $n = \dim_{\mathbb{R}} \mathcal{V}$. If $n \le 1$, clear.
For $n \ge 2$, Theorem 3.1.23 gives an eigenpair $T_{\mathbb{C}}(\mathbf{u}+i\mathbf{v}) = (a+bi)(\mathbf{u}+i\mathbf{v})$, where $\mathbf{u}+i\mathbf{v} \neq \mathbf{0}$.
Thus $T\mathbf{u} = a\mathbf{u} - b\mathbf{v}$ and $T\mathbf{v} = b\mathbf{u} + a\mathbf{v}$.
If $b = 0$, choose $\mathbf{0} \neq \mathbf{x} \in \{\mathbf{u},\mathbf{v}\}$, set $\mathcal{U} = \operatorname{span}(\mathbf{x})$, $\alpha = (\mathbf{x})$, so $[T|_{\mathcal{U}}]_{\alpha\leftarrow\alpha} = [a]$.
If $b \neq 0$, set $\mathcal{U} = \operatorname{span}(\mathbf{u},\mathbf{v})$, $\alpha = (\mathbf{u},-\mathbf{v})$; then $\mathbf{u},\mathbf{v}$ are independent and $[T|_{\mathcal{U}}]_{\alpha\leftarrow\alpha} = \begin{bmatrix} a & -b \\ b & a \end{bmatrix}$.
In either case, $T(\mathcal{U}) \subseteq \mathcal{U}$, $\dim\mathcal{U} \in \{1,2\}$, and the first diagonal block is admissible.
By induction, $T/\mathcal{U}$ has an almost upper-triangular matrix in some quotient basis $\overline{\gamma}$.
Lift $\overline{\gamma}$ to $\gamma \subseteq \mathcal{V}$ and set $\beta = (\alpha,\gamma)$. Then $[T]_{\beta\leftarrow\beta}$ is block upper triangular with first block $[T|_{\mathcal{U}}]_{\alpha\leftarrow\alpha}$ and quotient block $[T/\mathcal{U}]_{\overline{\gamma}\leftarrow\overline{\gamma}}$, hence is almost upper triangular.

Structural Commentary: Real triangularisation starts with an invariant line or plane.
Key Insight: Complex eigenvectors split into real and imaginary parts; those parts form the first real block. □

Corollary 3.1.33. *Orthonormal Real Schur Block Form*
Let $\dim_{\mathbb{R}} \mathcal{V} < \infty$ *and let* $T \in \mathcal{L}(\mathcal{V})$. *Then* $\exists$ *an orthonormal basis* γ *of* $\mathcal{V}$ *s.t.* $[T]_{\gamma \leftarrow \gamma}$ *is in real Schur block form.*

Proof. By Theorem 3.1.32, choose $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$ with block sizes $d_j \in \{1, 2\}$.
Put $s_k = d_1 + \cdots + d_k$ and $\mathcal{U}_k = \operatorname{span}(\mathbf{v}_1, \dots, \mathbf{v}_{s_k})$. Then $T(\mathcal{U}_k) \subseteq \mathcal{U}_k$ for all k.
Apply Gram–Schmidt to β and obtain an orthonormal basis $\gamma = (\mathbf{e}_1, \dots, \mathbf{e}_n)$ with $\operatorname{span}(\mathbf{e}_1, \dots, \mathbf{e}_{s_k}) = \mathcal{U}_k$, $1 \le k \le m$. Thus $[T]_{\gamma \leftarrow \gamma}$ is block upper triangular for the same $d_j \in \{1, 2\}$, hence is in real Schur block form.
Structural Commentary: Gram–Schmidt preserves the invariant block flag.
Key Insight: Orthonormalization changes blocks, not the flag. □

Proposition 3.1.34. *Invertibility from an Upper-Triangular Matrix*
Suppose $T \in \mathcal{L}(\mathcal{V})$ *has an upper-triangular matrix with diagonal entries* $\lambda_1, \dots, \lambda_n$ *with respect to some basis of* $\mathcal{V}$. *Then* T *is invertible if and only if* $\lambda_j \neq 0$ *for every* j.

Proof. Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$ be such a basis.

($\Leftarrow$) If every $\lambda_j \neq 0$, then upper-triangular matrix $[T]_{\beta \leftarrow \beta}$ is invertible by Theorem 1.4.10, so T is invertible.
($\Rightarrow$) Suppose T is invertible. If $\lambda_j = 0$ for some j, set $\mathcal{U}_j = \operatorname{span}(\mathbf{v}_1, \dots, \mathbf{v}_j)$ and $\mathcal{U}_0 = \{\mathbf{0}\}$. Upper triangularity gives $T(\mathcal{U}_j) \subseteq \mathcal{U}_{j-1}$ because the coefficient of $\mathbf{v}_j$ in $T\mathbf{v}_j$ is $\lambda_j = 0$ and $T\mathbf{v}_i \in \mathcal{U}_{j-1}$ for $i < j$.
Since $\dim(\mathcal{U}_j) > \dim(\mathcal{U}_{j-1})$, the map $\mathcal{U}_j \to \mathcal{U}_{j-1}$ induced by T is not injective. Thus there is a nonzero $\mathbf{u} \in \mathcal{U}_j$ with $T\mathbf{u} = \mathbf{0}$, contradicting invertibility. Hence every $\lambda_j \neq 0$.

Structural Commentary: Diagonal entries control each flag layer.
Key Insight: A zero diagonal entry forces a kernel vector. □

Proposition 3.1.35. *Eigenvalues from an Upper-Triangular Matrix*
Suppose $T \in \mathcal{L}(\mathcal{V})$ *has an upper-triangular matrix with diagonal entries* $\lambda_1, \dots, \lambda_n$ *with respect to some basis of* $\mathcal{V}$. *Then the eigenvalues of* T *are precisely* $\lambda_1, \dots, \lambda_n$.

Proof. For $\lambda \in \mathbb{F}$, the matrix of $T - \lambda I_{\mathcal{V}}$ in the same basis is upper triangular with diagonal entries $\lambda_1 - \lambda, \dots, \lambda_n - \lambda$. By Proposition 3.1.34, $T - \lambda I_{\mathcal{V}}$ is not invertible if and only if $\lambda = \lambda_j$ for some j. By Theorem 3.1.5, this is equivalent to λ being an eigenvalue of T.
Structural Commentary: Scalar shifts preserve triangular form.
Key Insight: Eigenvalues occur exactly when one shifted diagonal entry becomes zero. □

Perspective 3.1.36. *Why Triangular Form Matters*
Upper-triangular form is the first useful simplification beyond finding one eigenvector. It preserves enough of the operator to read eigenvalues from the diagonal while still allowing nontrivial coupling above the diagonal.

Definition 3.1.37. *Diagonalizable Operator*
An operator $T \in \mathcal{L}(\mathcal{V})$ is *diagonalizable* if T has a diagonal matrix with respect to some basis of $\mathcal{V}$.

Proposition 3.1.38. *Eigenspaces for Distinct Eigenvalues Form a Direct Sum*
Let $\mathcal{V}$ *be finite-dimensional, let* $T \in \mathcal{L}(\mathcal{V})$, *and suppose* $\lambda_1, \dots, \lambda_m$ *are distinct eigenvalues of* T.
Then $E_{\lambda_1}(T) + \cdots + E_{\lambda_m}(T)$ *is a direct sum. Consequently,* $\dim E_{\lambda_1}(T) + \cdots + \dim E_{\lambda_m}(T) \le \dim(\mathcal{V})$.

Proof. Suppose $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$ with $\mathbf{u}_j \in E_{\lambda_j}(T)$. If some $\mathbf{u}_j$ were nonzero, then the nonzero vectors among $\mathbf{u}_1, \dots, \mathbf{u}_m$ would be eigenvectors corresponding to distinct eigenvalues, and Theorem 3.1.6 would make them linearly independent. This contradicts their sum being $\mathbf{0}$. Hence every $\mathbf{u}_j = \mathbf{0}$, so the sum is direct.
By Theorem 1.2.21, $\dim(E_{\lambda_1}(T) + \cdots + E_{\lambda_m}(T)) = \dim E_{\lambda_1}(T) + \cdots + \dim E_{\lambda_m}(T)$.
The left side is the dimension of a subspace of $\mathcal{V}$, so Proposition 1.1.31 gives the stated inequality.
Structural Commentary: Distinct eigenspaces form independent pieces.
Key Insight: A zero sum across them would contradict eigenvector independence. □

Theorem 3.1.39. *Conditions Equivalent to Diagonalizability*
Let $\mathcal{V}$ *be finite-dimensional, let* $T \in \mathcal{L}(\mathcal{V})$, *and let* $\lambda_1, \dots, \lambda_m$ *denote the distinct eigenvalues of* T. *Set* $n = \dim(\mathcal{V})$. *The following are equivalent:*

(i) T *is diagonalizable;*
(ii) $\mathcal{V}$ *has a basis consisting of eigenvectors of* T*;*
(iii) $\exists$ *one-dimensional subspaces* $\mathcal{U}_1, \dots, \mathcal{U}_n$ *of* $\mathcal{V}$, *each invariant under* T, *such that* $\mathcal{V} = \mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_n$*;*
(iv) $\mathcal{V} = E_{\lambda_1}(T) \oplus \cdots \oplus E_{\lambda_m}(T)$*;*
(v) $\dim(\mathcal{V}) = \dim E_{\lambda_1}(T) + \cdots + \dim E_{\lambda_m}(T)$.

Proof.

(*i*) $\Rightarrow$ (*ii*) If T is diagonalizable, then T has a diagonal matrix with respect to some basis $(\mathbf{v}_1, \dots, \mathbf{v}_n)$. Hence $T\mathbf{v}_j = a_j\mathbf{v}_j$ for each j, so this basis consists of eigenvectors of T.

(*ii*) $\Rightarrow$ (*i*) If $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ is a basis of eigenvectors of T, then $T\mathbf{v}_j = a_j\mathbf{v}_j$ for each j. Thus the matrix of T with respect to this basis is diagonal.

(*ii*) $\Rightarrow$ (*iii*) Let $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ be an eigenbasis, and put $\mathcal{U}_j = \operatorname{span}(\mathbf{v}_j)$. Each $\mathcal{U}_j$ is one-dimensional and invariant under T by Proposition 3.1.11; the basis decomposition gives $\mathcal{V} = \mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_n$.

(*iii*) $\Rightarrow$ (*ii*) Choose nonzero $\mathbf{v}_j \in \mathcal{U}_j$. Proposition 3.1.11 makes each $\mathbf{v}_j$ an eigenvector of T, and the direct-sum decomposition makes $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ a basis of $\mathcal{V}$.

(*ii*) $\Rightarrow$ (*iv*) Every vector in an eigenbasis lies in one of $E_{\lambda_1}(T), \dots, E_{\lambda_m}(T)$, so these eigenspaces span $\mathcal{V}$. Proposition 3.1.38 shows their sum is direct.

(*iv*) $\Rightarrow$ (*v*) If $\mathcal{V} = E_{\lambda_1}(T) \oplus \cdots \oplus E_{\lambda_m}(T)$, then dimensions add by Theorem 1.2.21.

(*v*) $\Rightarrow$ (*ii*) For each j, choose a basis β_j of $E_{\lambda_j}(T)$, and concatenate these bases into a list β. The length of β is $\dim(\mathcal{V})$. If a linear combination of vectors in β is zero, grouping terms by eigenspace gives $\mathbf{u}_1 + \cdots + \mathbf{u}_m = \mathbf{0}$ with $\mathbf{u}_j \in E_{\lambda_j}(T)$. Proposition 3.1.38 gives $\mathbf{u}_j = \mathbf{0}$ for every j; since each β_j is independent, all coefficients are zero. Thus β is independent, Proposition 1.1.32 makes β a basis of $\mathcal{V}$ consisting of eigenvectors of T.

Structural Commentary: Diagonalizability means the eigenspaces exhaust $\mathcal{V}$. Directness is automatic for distinct eigenspaces, so the real test is spanning, or equivalently the dimension count.
Key Insight: A diagonal basis is exactly an eigenbasis. □

Perspective 3.1.40. *Diagonalization as Exhaustion by Eigenspaces*
Distinct eigenspaces are already separated by Proposition 3.1.38. Thus diagonalizability is the stronger statement that these separated pieces account for every direction in $\mathcal{V}$.

Theorem 3.1.41. *Algebraic Spectral Resolution*
Let $0 \neq \mathcal{V}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V})$, and let $\Lambda = \{\lambda_1, \dots, \lambda_m\} \subseteq \mathbb{F}$ with $\lambda_i \neq \lambda_j$ for $i \neq j$. The following are equivalent:

(i). T is diagonalizable and $\operatorname{spec}(T) = \Lambda$.

(ii). $\exists$ nonzero projections $\rho_1, \dots, \rho_m \in \mathcal{L}(\mathcal{V})$ s.t. $\rho_i\rho_j = 0$ $(i \neq j)$, $\sum_{i=1}^{m} \rho_i = I_{\mathcal{V}}$, $T = \sum_{i=1}^{m} \lambda_i\rho_i$.

In this case, for each i, $\operatorname{range}(\rho_i) = E_{\lambda_i}(T)$; $\ker(\rho_i) = \bigoplus_{j\neq i} E_{\lambda_j}(T)$; and $\rho_i = \prod_{j\neq i} \frac{T-\lambda_j I_{\mathcal{V}}}{\lambda_i-\lambda_j}$

Proof.

(*i*) $\Rightarrow$ (*ii*) By Theorem 3.1.39, $\mathcal{V} = \bigoplus_{i=1}^{m} E_{\lambda_i}(T)$. Let $\rho_i = \rho_{E_{\lambda_i}(T), \bigoplus_{j\neq i} E_{\lambda_j}(T)}$.

Then $\rho_i\rho_j = 0$ for $i \neq j$, $\sum_{i=1}^{m} \rho_i = I_{\mathcal{V}}$, and $T = \sum_{i=1}^{m} \lambda_i\rho_i$. Also $\operatorname{range}(\rho_i)$ and $\ker(\rho_i)$ are as displayed.

For $\ell_i(z) = \prod_{j\neq i} \frac{z-\lambda_j}{\lambda_i-\lambda_j}$, we have $\ell_i(\lambda_k) = \delta_{ik}$, hence $\ell_i(T) = \rho_i$.

(*ii*) $\Rightarrow$ (*i*) Put $\mathcal{V}_i = \operatorname{range}(\rho_i)$. By Theorem 1.3.19, $\mathcal{V} = \bigoplus_{i=1}^{m} \mathcal{V}_i$.

If $\mathbf{x} \in \mathcal{V}_i$, then $\rho_i\mathbf{x} = \mathbf{x}$ and $\rho_j\mathbf{x} = \mathbf{0}$ for $j \neq i$, so $T\mathbf{x} = \lambda_i\mathbf{x}$. Thus $\mathcal{V}_i \subseteq E_{\lambda_i}(T)$.

Conversely, if $T\mathbf{x} = \lambda_i\mathbf{x}$ and $\mathbf{x} = \sum_{j=1}^{m} \mathbf{x}_j$ with $\mathbf{x}_j \in \mathcal{V}_j$, then $\sum_{j=1}^{m} (\lambda_j - \lambda_i)\mathbf{x}_j = \mathbf{0}$, so $\mathbf{x}_j = \mathbf{0}$ for $j \neq i$.

Hence $\mathcal{V}_i = E_{\lambda_i}(T)$. Since each $\mathcal{V}_i \neq 0$, each $\lambda_i \in \operatorname{spec}(T)$; concatenating bases of the $\mathcal{V}_i$ gives an eigenbasis, so T is diagonalizable and $\operatorname{spec}(T) = \Lambda$.

The kernel and polynomial formulas follow from $\mathcal{V} = \bigoplus_{i=1}^{m} \mathcal{V}_i$ and $\ell_i(\lambda_k) = \delta_{ik}$.

Structural Commentary: A diagonalizable operator is a weighted sum of the projections onto its eigenspaces.
Key Insight: Spectral projections are Lagrange polynomials in the operator. □

Lemma 3.1.42. *Diagonalizable Restrictions to Invariant Subspaces*
Let $T \in \mathcal{L}(\mathcal{V})$ be diagonalizable and let $\mathcal{U} \in \operatorname{Inv}(T)$. For $\lambda \in \operatorname{spec}(T)$, set $\mathcal{U}_\lambda = \mathcal{U} \cap E_\lambda(T)$. Then $\mathcal{U} = \bigoplus_{\lambda\in\operatorname{spec}(T)} \mathcal{U}_\lambda$, where $T|_{\mathcal{U}}$ is diagonalizable.

Proof. Let $\rho_\lambda = p_\lambda(T)$ be the spectral projection from Theorem 3.1.41. Since $\mathcal{U} \in \operatorname{Inv}(T)$, $\rho_\lambda(\mathcal{U}) \subseteq \mathcal{U}$. Hence, for $\mathbf{u} \in \mathcal{U}$, $\mathbf{u} = \sum_{\lambda\in\operatorname{spec}(T)} \rho_\lambda\mathbf{u}$, where $\rho_\lambda\mathbf{u} \in \mathcal{U}_\lambda$. Thus $\mathcal{U} = \sum_{\lambda\in\operatorname{spec}(T)} \mathcal{U}_\lambda$, and the sum is direct since $\mathcal{U}_\lambda \subseteq E_\lambda(T)$. Bases of the nonzero $\mathcal{U}_\lambda$ concatenate to an eigenbasis for $T|_{\mathcal{U}}$.

Structural Commentary: Invariant subspaces inherit spectral projections.
Key Insight: Polynomial spectral projections cannot move vectors out of a T-invariant subspace. □

Theorem 3.1.43. *Simultaneous Diagonalisation of Commuting Diagonalizable Operators*
Let $\dim \mathcal{V} < \infty$, *and let* $\emptyset \neq \mathcal{F} \subseteq \mathcal{L}(\mathcal{V})$ *be finite with every* $T \in \mathcal{F}$ *diagonalizable.*
Then $\forall S, T \in \mathcal{F},\ ST = TS \iff \exists$ *a basis* β *of* $\mathcal{V}$ *such that* $[T]_{\beta\leftarrow\beta}$ *is diagonal* $\forall T \in \mathcal{F}$.

Proof.

($\Rightarrow$) Induct on $n = \dim \mathcal{V}$. If every $T \in \mathcal{F}$ is scalar, any basis works. Otherwise choose nonscalar $S \in \mathcal{F}$.
By Theorem 3.1.41, $\mathcal{V} = \bigoplus_{\lambda \in \operatorname{spec}(S)} E_\lambda(S)$, with at least two nonzero summands of dimension $< n$.
If $T \in \mathcal{F}$ and $\mathbf{x} \in E_\lambda(S)$, then $S(T\mathbf{x}) = T(S\mathbf{x}) = \lambda T\mathbf{x}$, so $E_\lambda(S) \in \operatorname{Inv}(T)$.
By Lemma 3.1.42, each $T|_{E_\lambda(S)}$ is diagonalizable; the restricted family still commutes.
Induction gives a common eigenbasis of each $E_\lambda(S)$; concatenate over λ.

($\Leftarrow$) For $S, T \in \mathcal{F}$, set $D_S = [S]_{\beta\leftarrow\beta}$ and $D_T = [T]_{\beta\leftarrow\beta}$.
Since D_S, D_T are diagonal, $[ST]_{\beta\leftarrow\beta} = D_S D_T = D_T D_S = [TS]_{\beta\leftarrow\beta}$. Thus $ST = TS$.

Structural Commentary: Commutation preserves eigenspaces.
Key Insight: Diagonalize one operator, then restrict. □

Corollary 3.1.44. *Enough Eigenvalues Imply Diagonalizability*
If $\mathcal{V}$ *is finite-dimensional and* $T \in \mathcal{L}(\mathcal{V})$ *has* $\dim(\mathcal{V})$ *distinct eigenvalues, then* T *is diagonalizable.*

Proof. Let $n = \dim(\mathcal{V})$, and let $\lambda_1, \ldots, \lambda_n$ be distinct eigenvalues of T. For each j, choose an eigenvector $\mathbf{v}_j$ corresponding to λ_j. By Theorem 3.1.6, $(\mathbf{v}_1, \ldots, \mathbf{v}_n)$ is linearly independent. Since this list has length n, Proposition 1.1.32 makes it a basis of $\mathcal{V}$. Thus $\mathcal{V}$ has an eigenbasis, so T is diagonalizable by Theorem 3.1.39.
Structural Commentary: One eigenvector from each distinct eigenvalue already gives a full-size independent list.
Key Insight: Once the independent eigenvectors reach $\dim(\mathcal{V})$, they must form a basis. □

Theorem 3.1.45. *Commutant of a Diagonalizable Matrix*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$ *be diagonalizable. Suppose its distinct eigenvalues are* $\mu_1, \ldots, \mu_d$, *and set* $n_i = \dim E_{\mu_i}(A)$.
Write $A = S(\mu_1 I_{n_1} \oplus \cdots \oplus \mu_d I_{n_d})S^{-1}$. *Then* $\operatorname{Comm}(A) = \{S(X_1 \oplus \cdots \oplus X_d)S^{-1} : X_i \in \operatorname{Mat}(n_i, n_i, \mathbb{C})\}$.
In particular, $\dim \operatorname{Comm}(A) = n_1^2 + \cdots + n_d^2$.

Proof. Let $\Lambda = \mu_1 I_{n_1} \oplus \cdots \oplus \mu_d I_{n_d}$. Matrix B commutes with A if and only if $S^{-1}BS$ commutes with Λ.
By Lemma 3.1.25, this happens exactly when $S^{-1}BS = X_1 \oplus \cdots \oplus X_d$ with $X_i \in \operatorname{Mat}(n_i, n_i, \mathbb{C})$.
Dimension count follows because conjugation by S is a vector-space isomorphism and $\dim \operatorname{Mat}(n_i, n_i, \mathbb{C}) = n_i^2$.
Structural Commentary: Diagonalizing A reduces commutation to respecting equal-eigenvalue blocks.
Key Insight: Distinct eigenvalue blocks cannot mix under a commuting matrix. □

Theorem 3.1.46. *Simultaneous Schur Triangularisation*
Let $\emptyset \neq \mathcal{F} \subseteq \operatorname{Mat}(n, n, \mathbb{C})$ *be a commuting family.*
Then there is a unitary $U \in \operatorname{Mat}(n, n, \mathbb{C})$ *such that* U^*AU *is upper triangular for every* $A \in \mathcal{F}$.
If every matrix in $\mathcal{F}$ *is real and has only real eigenvalues, then* U *can be chosen real orthogonal.*

Proof. Induct on n. The case $n = 1$ is immediate.
For $n > 1$, Theorem 3.1.27 gives a common eigenvector; normalize it to a unit vector $\mathbf{x}$ and extend to an orthonormal basis. Let $V = [\mathbf{x}\ V_2]$ be the corresponding unitary matrix.
For every $A \in \mathcal{F}$, since $A\mathbf{x} = \lambda_A \mathbf{x}$, we have $V^*AV = \begin{bmatrix} \lambda_A & * \\ 0 & A' \end{bmatrix}$.
If $A, B \in \mathcal{F}$, then $AB = BA$. The lower-right block of $V^*ABV = (V^*AV)(V^*BV)$ is $A'B'$, while the lower-right block of V^*BAV is $B'A'$. Hence $A'B' = B'A'$, so the compressed family is still commuting.
By induction, choose a unitary W that simultaneously upper triangularises all compressed blocks.
Then $U = V(1 \oplus W)$ is unitary and U^*AU is upper triangular for every $A \in \mathcal{F}$.
For the real case, run the same induction over $\mathbb{R}$. A common real eigenvector exists as follows. If all matrices are scalar, any real unit vector works. Otherwise choose a nonscalar $A_0 \in \mathcal{F}$.
Since A_0 has only real eigenvalues, some real eigenspace $\mathcal{U} = E_\lambda(A_0)$ is nonzero; since A_0 is not scalar, $\mathcal{U} \neq \mathbb{R}^n$.
Commutation gives $(A_0 - \lambda I)B\mathbf{u} = B(A_0 - \lambda I)\mathbf{u} = \mathbf{0}$ for $B \in \mathcal{F}$ and $\mathbf{u} \in \mathcal{U}$, so $\mathcal{U}$ is invariant under every B.
The restricted family on $\mathcal{U}$ is commuting and still has only real eigenvalues: in a basis adapted to $\mathcal{U}$ each B has block form $\begin{bmatrix} B|_{\mathcal{U}} & * \\ 0 & * \end{bmatrix}$, so the characteristic polynomial of $B|_{\mathcal{U}}$ divides that of B. Induction on $\dim(\mathcal{U})$ gives a common real eigenvector. Extending by real orthonormal bases gives a real orthogonal U.
Structural Commentary: A common eigenline creates the first block; commutation descends to the compression.
Key Insight: Simultaneous triangularisation is Schur induction with a common eigenvector at every step. □

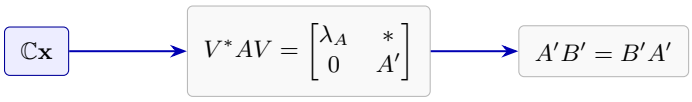


Figure 44: A shared eigenline gives a simultaneous first block; lower-right compressions still commute.

Corollary 3.1.47. *Eigenvalues of Sums and Products of Commuting Matrices*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{C})$ and $AB = BA$. Their eigenvalues can be ordered as $\lambda_1, \dots, \lambda_n$ and $\mu_1, \dots, \mu_n$ so that $A + B$ has eigenvalues $\lambda_i + \mu_i$ and AB has eigenvalues $\lambda_i \mu_i$, counted with multiplicity.

Proof. Choose simultaneous Schur forms $T = U^*AU$ and $R = U^*BU$. Set $\lambda_i = t_{ii}$ and $\mu_i = r_{ii}$. Then $T + R$ and TR are upper triangular with diagonals $\lambda_i + \mu_i$ and $\lambda_i \mu_i$, and are similar to $A + B$ and AB.
Structural Commentary: Simultaneous Schur form aligns both spectra on one diagonal.
Key Insight: In common triangular coordinates, diagonals add and multiply. □

3.2 Polynomial Structure of Operators

Polynomial structure records the equations an operator satisfies and turns them into invariant pieces through characteristic and minimal polynomials, Cayley–Hamilton, generalized eigenspaces, and nilpotent parts.

Definition 3.2.1. *Operator Powers and Polynomial Evaluation*
Let $T \in \mathcal{L}(\mathcal{V})$.

(i) The *powers* of T are $T^m := \underbrace{T \cdots T}_{m \text{ times}}$ for $m \in \mathbb{N}$, with $T^0 := I_{\mathcal{V}}$. If T is invertible, set $T^{-m} := (T^{-1})^m$.

(ii) For $p(z) = \sum_{j=0}^{m} a_j z^j \in \mathbb{F}[z]$, the *polynomial p applied to T* is $p(T) := \sum_{j=0}^{m} a_j T^j$.

(iii) $p \in \mathbb{F}[z]$ is an *annihilating polynomial* for T if $p(T) = 0_{\mathcal{L}(\mathcal{V})}$; equivalently, p *annihilates* T.

Definition 3.2.2. *Product of Polynomials*
If $p, q \in \mathbb{F}[z]$, their product $pq \in \mathbb{F}[z]$ is defined by $(pq)(z) = p(z)q(z)$ for $z \in \mathbb{F}$.

Proposition 3.2.3. *Multiplicative Properties of Operator Polynomials*
Let $p, q \in \mathbb{F}[z]$ and $T \in \mathcal{L}(\mathcal{V})$. Then
(i) $(pq)(T) = p(T)q(T)$;
(ii) $p(T)q(T) = q(T)p(T)$.

Proof.

(i) Write $p(z) = \sum_{j=0}^{m} a_j z^j$ and $q(z) = \sum_{k=0}^{n} b_k z^k$. Then $(pq)(z) = \sum_{j=0}^{m} \sum_{k=0}^{n} a_j b_k z^{j+k}$.
Substituting T gives $(pq)(T) = \sum_{j=0}^{m} \sum_{k=0}^{n} a_j b_k T^{j+k} = \left(\sum_{j=0}^{m} a_j T^j \right) \left(\sum_{k=0}^{n} b_k T^k \right) = p(T)q(T)$.

(ii) Since polynomial multiplication is commutative, $pq = qp$.
By part (i), $p(T)q(T) = (pq)(T) = (qp)(T) = q(T)p(T)$.

Structural Commentary: Polynomial evaluation preserves products.
Key Insight: Expanding a polynomial product works the same with T as with z. □

Proposition 3.2.4. *Existence of Annihilating Polynomials*
Let $\mathcal{V}$ be finite-dimensional with $n = \dim(\mathcal{V}) > 0$, and let $T \in \mathcal{L}(\mathcal{V})$.
Then T has a nonconstant annihilating polynomial of degree at most n^2.

Proof. By Theorem 1.3.89, the vector space $\mathcal{L}(\mathcal{V})$ has dimension n^2, so the $n^2 + 1$ operators $I_{\mathcal{V}}, T, T^2, \dots, T^{n^2}$ are linearly dependent. Hence there are scalars $a_0, \dots, a_{n^2}$, not all zero, such that $\sum_{j=0}^{n^2} a_j T^j = 0$.
If $a_1 = \cdots = a_{n^2} = 0$, then $a_0 I_{\mathcal{V}} = 0$, forcing $a_0 = 0$ because $\mathcal{V} \neq \{\mathbf{0}\}$, a contradiction. Thus $p(z) = \sum_{j=0}^{n^2} a_j z^j$ is a nonconstant annihilating polynomial.
Structural Commentary: Finite-dimensionality makes the powers of one operator linearly dependent inside $\mathcal{L}(\mathcal{V})$.
Key Insight: A dependence among $I_{\mathcal{V}}, T, \dots, T^{n^2}$ is exactly a polynomial equation satisfied by T. □

Proposition 3.2.5. *Polynomial Evaluation on Eigenvectors*
Let $T \in \mathcal{L}(\mathcal{V})$, let $p \in \mathbb{F}[z]$, and suppose $T\mathbf{v} = \lambda \mathbf{v}$ with $\mathbf{v} \neq \mathbf{0}$. Then $p(T)\mathbf{v} = p(\lambda)\mathbf{v}$.

Proof. Induct on j to get $T^j\mathbf{v} = \lambda^j\mathbf{v}$ for every $j \geq 0$.
If $p(z) = \sum_{j=0}^{m} a_j z^j$, then $p(T)\mathbf{v} = \sum_{j=0}^{m} a_j T^j\mathbf{v} = \sum_{j=0}^{m} a_j\lambda^j\mathbf{v} = p(\lambda)\mathbf{v}$.
Structural Commentary: On a one-dimensional eigenspace, T acts like the scalar λ.
Key Insight: Polynomial expressions in T reduce to the same polynomial evaluated at λ. □

Motivation 3.2.6. *Two Spectral Polynomials*
The characteristic polynomial records the whole primary decomposition with dimensions.
The minimal polynomial records the shortest polynomial relation satisfied by the operator.

Definition 3.2.7. *Characteristic Polynomial*
For $A \in \mathrm{Mat}(n, n, \mathbb{F})$, the *characteristic polynomial* of A is $p_A(z) := \det(zI_n - A)$.

Theorem 3.2.8. *Basic Form of the Characteristic Polynomial*
Let $A = [a_{ij}] \in \mathrm{Mat}(n, n, \mathbb{F})$. Then p_A is monic of degree n and $p_A(z) = z^n - (\operatorname{tr} A)z^{n-1} + \cdots + (-1)^n \det A$. Moreover, for $\lambda \in \mathbb{F}$, $\lambda \in \operatorname{spec}(A)$ if and only if $p_A(\lambda) = 0$.

Proof. The only determinant term of degree greater than $n-2$ comes from the identity permutation:

$$(z - a_{11})(z - a_{22}) \cdots (z - a_{nn}) = z^n - (a_{11} + \cdots + a_{nn})z^{n-1} + \cdots .$$

Thus p_A is monic and the coefficient of z^{n-1} is $-\operatorname{tr} A$; also $p_A(0) = \det(-A) = (-1)^n \det A$. Finally, $p_A(\lambda) = 0$ if and only if $\lambda I_n - A$ is singular, equivalently $\ker(A - \lambda I_n) \neq \{\mathbf{0}\}$.
Structural Commentary: The diagonal product determines the top two coefficients.
Key Insight: Eigenvalues are precisely the roots of $\det(zI_n - A)$. □

Definition 3.2.9. *Algebraic Multiplicity*
For $\lambda \in \mathbb{F}$, let $m_A(\lambda)$ be the multiplicity of λ as a root of p_A, with $m_A(\lambda) = 0$ if $\lambda \notin \operatorname{spec}(A)$.
For $\lambda \in \operatorname{spec}(A)$, this is the *algebraic multiplicity* of λ. Unqualified multiplicity means algebraic multiplicity.

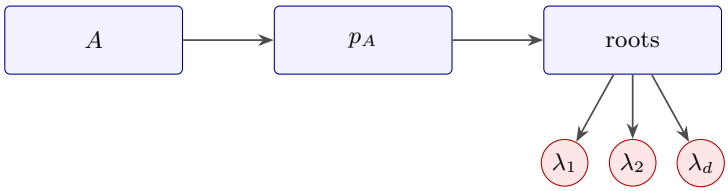


Figure 45: Form $p_A(z) = \det(zI_n - A)$ and factor it; each root λ is counted by $m_A(\lambda)$.

Method 3.2.10. *Characteristic Equation and Eigenspace Computation*
To compute eigenvalues and eigenspaces of $A \in \mathrm{Mat}(n, n, \mathbb{F})$:

(i) Form the characteristic polynomial $p_A(z) = \det(zI_n - A)$ from Definition 3.2.7.
(ii) Solve the characteristic equation $p_A(\lambda) = 0$. By Theorem 3.2.8, the solutions are exactly the eigenvalues in $\mathbb{F}$.
(iii) For each eigenvalue λ, row-reduce $A - \lambda I_n$ and compute

$$E_\lambda(A) = \ker(A - \lambda I_n).$$

(iv) Choose a basis of each eigenspace. The dimension of this basis is the geometric multiplicity of λ.

Algebraic multiplicity comes from the factorisation of p_A, while geometric multiplicity comes from the nullspace computation.

Proposition 3.2.11. *Upper-Triangular Eigenvalues*
If $A = [a_{ij}] \in \mathrm{Mat}(n, n, \mathbb{F})$ is upper triangular, then $p_A(z) = \prod_{j=1}^{n} (z - a_{jj})$.
Thus $\operatorname{spec}(A) = \{a_{jj} : 1 \leq j \leq n\}$, and for $\lambda \in \operatorname{spec}(A)$, $m_A(\lambda) = \#\{j : a_{jj} = \lambda\}$.

Proof. Since $zI_n - A$ is upper triangular with diagonal entries $z - a_{jj}$, $p_A(z) = \det(zI_n - A) = \prod_{j=1}^{n} (z - a_{jj})$.
Thus the roots of p_A are exactly the diagonal entries, and the exponent of $(z - \lambda)$ is $\#\{j : a_{jj} = \lambda\}$.
Structural Commentary: Triangular form removes all off-diagonal determinant contributions.
Key Insight: Once a matrix is triangular, the spectrum is visible on the diagonal. □

Theorem 3.2.12. *Similarity Preserves Characteristic Data*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$ and suppose $B = SAS^{-1}$ for some invertible S. Then $p_B = p_A$. In particular, for each $\lambda \in \operatorname{spec}(A)$, $m_B(\lambda) = m_A(\lambda)$ and $\dim E_\lambda(B) = \dim E_\lambda(A)$.

Proof. For every z, $p_B(z) = \det(zI_n - SAS^{-1}) = \det(S(zI_n - A)S^{-1}) = \det(S)\det(zI_n - A)\det(S)^{-1} = p_A(z)$. Thus roots and root multiplicities agree, so $m_B(\lambda) = m_A(\lambda)$. Also $B - \lambda I_n = S(A - \lambda I_n)S^{-1}$.
Hence $\mathbf{x} \in \ker(A - \lambda I_n)$ if and only if $S\mathbf{x} \in \ker(B - \lambda I_n)$, S restricts to isomorphism between two eigenspaces.
Structural Commentary: Similarity changes coordinates, not the shifted determinant or shifted kernel dimension.
Key Insight: Conjugation inserts S and S^{-1}, which cancel in determinants and identify kernels. □

Proposition 3.2.13. *Block Upper-Triangular Characteristic Polynomial*
Let $A = \begin{bmatrix} B & C \\ 0 & D \end{bmatrix}$, *$B$ and D are square. Then* $p_A(z) = p_B(z)p_D(z)$, *so* $m_A(\lambda) = m_B(\lambda) + m_D(\lambda)$ $\forall \lambda \in \mathbb{F}$.

Proof. The shifted matrix is block upper triangular, $zI - A = \begin{bmatrix} zI - B & -C \\ 0 & zI - D \end{bmatrix}$.
Theorem 4.1.11 gives $p_A(z) = \det(zI - B)\det(zI - D) = p_B(z)p_D(z)$.
Roots of a product are the roots of the factors, with multiplicity added.
Structural Commentary: The off-diagonal block couples vectors but not determinant factors.
Key Insight: In block upper-triangular form, spectral data lives on the diagonal blocks. □

Theorem 3.2.14. *Geometric Multiplicity Is at Most Algebraic Multiplicity*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ *and let* $\lambda \in \mathrm{spec}(A)$. *Then* $\dim E_\lambda(A) \le m_A(\lambda)$.

Proof. Let $k = \dim E_\lambda(A)$ and choose a basis $\mathbf{x}_1, \dots, \mathbf{x}_k$ of $E_\lambda(A)$. Extend it to a basis of $\mathbb{F}^n$, and let S be the matrix whose columns are this basis. In the new basis,

$$S^{-1}AS = \begin{bmatrix} \lambda I_k & * \\ 0 & C \end{bmatrix}.$$

By Theorem 3.2.12, Proposition 3.2.13, $p_A(z) = p_{S^{-1}AS}(z) = (z - \lambda)^k p_C(z)$. Thus $m_A(\lambda) \ge k = \dim E_\lambda(A)$.
Structural Commentary: An eigenspace-first basis exposes a λI_k block.
Key Insight: Each independent eigenvector forces one visible factor $z - \lambda$ in the characteristic polynomial. □

Theorem 3.2.15. *Diagonalizability and Multiplicity*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. *Then A is diagonalizable if and only if* $\dim E_\lambda(A) = m_A(\lambda)$ *for every* $\lambda \in \mathrm{spec}(A)$.

Proof.

($\Rightarrow$) If A is diagonalizable, then A is similar to a diagonal matrix D. Theorem 3.2.12 gives $m_A(\lambda) = m_D(\lambda)$ and $\dim E_\lambda(A) = \dim E_\lambda(D)$. For diagonal D, both numbers count diagonal entries equal to λ.

($\Leftarrow$) Assume $\dim E_\lambda(A) = m_A(\lambda)$ for every eigenvalue.

If $\mu_1, \dots, \mu_d$ are distinct eigenvalues, then $\sum_{i=1}^{d} m_A(\mu_i) = n$. Hence $\dim E_{\mu_1}(A) + \cdots + \dim E_{\mu_d}(A) = n$.

Theorem 3.1.39 gives an eigenbasis, so A is diagonalizable.

Structural Commentary: Diagonalizability means root counts match eigenspace dimensions.
Key Insight: No missing eigendirections means an eigenbasis. □

Definition 3.2.16. *Characteristic Polynomial of an Operator*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$, and let $T \in \mathcal{L}(\mathcal{V})$. If $\lambda_1, \dots, \lambda_m$ are the distinct eigenvalues of T and $d_j = \dim \mathcal{G}_{\lambda_j}(T)$, then the *characteristic polynomial* of T is $q_T(z) := \prod_{j=1}^{m} (z - \lambda_j)^{d_j}$.

Proposition 3.2.17. *Degree and Roots of q_T*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ *be finite-dimensional over* $\mathbb{C}$, *let* $T \in \mathcal{L}(\mathcal{V})$, *and let q_T be the characteristic polynomial of T.*

(i) $\deg q_T = \dim(\mathcal{V})$.
(ii) The roots of q_T are precisely the eigenvalues of T.

Proof. Write $q_T(z) = \prod_{j=1}^{m} (z - \lambda_j)^{d_j}$, where $\lambda_1, \dots, \lambda_m$ are the distinct eigenvalues and $d_j = \dim \mathcal{G}_{\lambda_j}(T)$.

(i) Proposition 3.2.60 gives $\deg q_T = \sum_j d_j = \dim(\mathcal{V})$.
(ii) The displayed factorization has exactly the roots $\lambda_1, \dots, \lambda_m$.

Structural Commentary: Each primary summand contributes one root factor.
Key Insight: Multiplicity is dimension. □

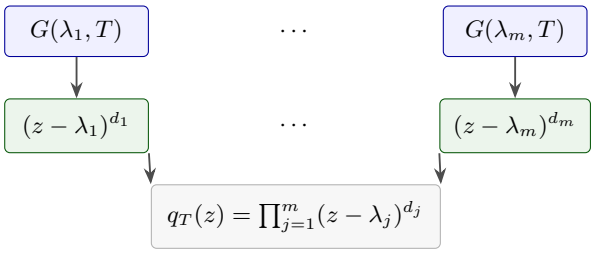


Figure 46: Each primary summand contributes $(z-\lambda_j)^{d_j}$, where $d_j = \dim \mathcal{G}_{\lambda_j}(T)$.

Definition 3.2.18. *Monic Polynomial*
A nonzero polynomial is *monic* if its highest-degree coefficient is 1.

Theorem 3.2.19. *Existence of the Minimal Polynomial*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dim over $\mathbb{F}$ and let $T \in \mathcal{L}(\mathcal{V})$. Then T has a unique least-degree monic annihilator μ_T.

Proof. Proposition 3.2.4 gives a nonzero annihilator; scaling makes it monic, so a least-degree monic annihilator exists. If p and q are two of them, then $(p-q)(T) = 0$ and $\deg(p-q) < \deg p$.
If $p \neq q$, scaling $p - q$ gives a smaller monic annihilator, impossible. Hence $p = q$.
Structural Commentary: Existence comes from finite-dimensional dependence; uniqueness from minimality.
Key Insight: Two least monic annihilators cannot differ. □

Definition 3.2.20. *Minimal Polynomial of an Operator*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional and let $T \in \mathcal{L}(\mathcal{V})$. The *minimal polynomial* of T, denoted μ_T, is the unique monic polynomial of smallest degree such that $\mu_T(T) = 0$.

Definition 3.2.21. *Minimal Polynomial of a Matrix*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$. The *minimal polynomial* of A, denoted μ_A, is the unique monic polynomial of smallest positive degree such that $\mu_A(A) = 0$.

Theorem 3.2.22. *Cayley–Hamilton Theorem for Operators*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$, q_T be characteristic polynomial of $T \in \mathcal{L}(\mathcal{V})$. Then $q_T(T) = 0$.

Proof. Let $\lambda_1, \ldots, \lambda_m$ be the distinct eigenvalues and set $d_j = \dim \mathcal{G}_{\lambda_j}(T)$.
It suffices, by primary decomposition, to show $q_T(T)$ kills each $\mathcal{G}_{\lambda_j}(T)$.
On $\mathcal{G}_{\lambda_j}(T)$, the operator $T - \lambda_j I_{\mathcal{V}}$ is nilpotent, so Corollary 3.2.52 gives $(T-\lambda_j I_{\mathcal{V}})^{d_j} = 0$ on $\mathcal{G}_{\lambda_j}(T)$.
Since the factors of $q_T(T)$ commute, this factor kills the whole product on $\mathcal{G}_{\lambda_j}(T)$. Hence $q_T(T) = 0$.
Structural Commentary: Check the polynomial on each primary summand.
Key Insight: One factor of $q_T(T)$ kills each summand. □

Theorem 3.2.23. *Cayley–Hamilton Theorem*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ have characteristic polynomial $p_A(z) = \det(zI_n - A)$. Then $p_A(A) = 0$.

Proof. Choose a Schur form $A = UTU^*$ with diagonal $\lambda_1, \ldots, \lambda_n$.
Since $p_A(A) = Up_A(T)U^*$ and $p_A(z) = \prod_{i=1}^{n}(z-\lambda_i)$, it suffices to prove $(T-\lambda_1 I_n)(T-\lambda_2 I_n)\cdots(T-\lambda_n I_n) = 0$.
Let $Q_j := \prod_{i=1}^{j}(T-\lambda_i I_n)$. We show that the first j columns of Q_j are zero. For $j = 1$, the first column of $T-\lambda_1 I_n$ is zero as T is upper triangular with first diagonal entry λ_1. Assume the first j columns of Q_j are zero.
Since $T-\lambda_{j+1} I_n$ is upper triangular and its $(j+1, j+1)$ entry is zero, its first $j+1$ columns lie in $\mathrm{span}(\mathbf{e}_1, \ldots, \mathbf{e}_j)$.
By the induction hypothesis, Q_j kills this span, so $Q_{j+1} = Q_j(T - \lambda_{j+1} I_n)$ has first $j+1$ zero columns.
At $j = n$, all columns of Q_n are zero, so $Q_n = 0$.
Structural Commentary: Schur form reduces the theorem to a triangular column-killing induction.
Key Insight: Each factor $T - \lambda_i I_n$ kills the next leading column. □

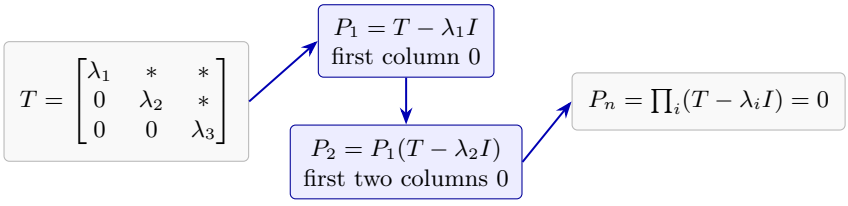


Figure 47: In Schur form, the factors $T - \lambda_i I$ kill the leading columns one layer at a time.

Theorem 3.2.24. *Basic Properties of the Minimal Polynomial*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$, and let $\lambda_1, \ldots, \lambda_d$ be the distinct eigenvalues of A.

(i) μ_A exists, is unique, and $\deg \mu_A \leq n$.
(ii) If a nonzero polynomial f annihilates A, then μ_A divides f. In particular, μ_A divides p_A.
(iii) There are integers q_i satisfying $1 \leq q_i \leq m_A(\lambda_i)$ such that $\mu_A(z) = \prod_{i=1}^{d} (z - \lambda_i)^{q_i}$.

Proof.

(i) Cayley–Hamilton supplies the monic annihilator p_A, so a least-degree monic annihilator μ_A exists and has degree at most n. If g is another least-degree monic annihilator, write $g = h\mu_A + r$, where $r = 0$ or $\deg r < \deg \mu_A$. Evaluating at A gives $r(A) = 0$. If $r \neq 0$, scaling r to be monic would give a smaller-degree monic annihilator, a contradiction. Hence $r = 0$, so $g = h\mu_A$.
Since g and μ_A are monic of the same degree, $h = 1$ and $g = \mu_A$.

(ii) If $f(A) = 0$, divide $f = q\mu_A + r$ with $\deg r < \deg \mu_A$. Evaluating at A gives $r(A) = 0$, hence $r = 0$ and $\mu_A \mid f$. Taking $f = p_A$ gives $\mu_A \mid p_A$.

(iii) Since $\mu_A \mid p_A = \prod_{i=1}^{d} (z - \lambda_i)^{m_A(\lambda_i)}$, its roots are eigenvalues with exponents at most the algebraic multiplicities. If $A\mathbf{v} = \lambda_i \mathbf{v}$ and $\mathbf{v} \neq \mathbf{0}$, then $\mu_A(\lambda_i)\mathbf{v} = \mu_A(A)\mathbf{v} = 0$, so every λ_i occurs.

Structural Commentary: Polynomial division makes μ_A the least monic generator of annihilators.
Key Insight: Cayley–Hamilton gives existence; eigenvectors force the roots. □

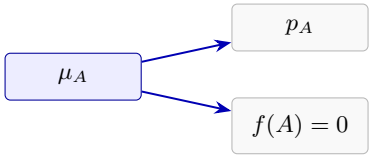


Figure 48: μ_A is the least monic annihilator; arrows indicate that μ_A divides p_A and every annihilator f.

Theorem 3.2.25. *Minimal Polynomial of a Spectrally Disjoint Direct Sum*
Let $A = A_1 \oplus \cdots \oplus A_d$, where the spectra of $A_1, \ldots, A_d$ are pairwise disjoint. Then $\mu_A = \mu_{A_1} \mu_{A_2} \cdots \mu_{A_d}$.

Proof. Let $f = \prod_i \mu_{A_i}$. Since $f(A_i) = 0$ for all i, blockwise evaluation gives $f(A) = 0$, and therefore $\mu_A \mid f$.
Conversely, $\mu_A(A_i) = 0$ for each i, so $\mu_{A_i} \mid \mu_A$ for each i. The roots of μ_{A_i} are exactly $\mathrm{spec}(A_i)$ by Theorem 3.2.33, so the disjointness of the spectra makes the polynomials μ_{A_i} pairwise relatively prime.
Hence their product f divides μ_A. Since $\mu_A \mid f$ and $f \mid \mu_A$, and both are monic, $\mu_A = f$.
Structural Commentary: Direct sums are annihilated blockwise; disjoint spectra make the block factors coprime.
Key Insight: Each spectral component contributes its own minimal-polynomial factor. □

Theorem 3.2.26. *Similarity Preserves the Minimal Polynomial*
If $A, B \in \mathrm{Mat}(n, n, \mathbb{C})$ are similar, then $p_A = p_B$ and $\mu_A = \mu_B$.

Proof. The identity $p_A = p_B$ is Theorem 3.2.12. Let $A = SBS^{-1}$, then $f(A) = Sf(B)S^{-1}$ for every polynomial f, so A and B have the same annihilating polynomials. Hence $\mu_A = \mu_B$.
Structural Commentary: Similarity preserves polynomial equations.
Key Insight: Conjugation commutes with polynomial evaluation. □

Theorem 3.2.27. *Diagonalizability and the Minimal Polynomial*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ have distinct eigenvalues $\lambda_1, \ldots, \lambda_d$.
Then A is diagonalizable if and only if $\mu_A(z) = (z - \lambda_1) \cdots (z - \lambda_d)$.
Equivalently, A is diagonalizable if and only if μ_A is square-free.

Proof. Set $p(z) := \prod_{i=1}^{d} (z - \lambda_i)$.

($\Rightarrow$) If $A = S\Lambda S^{-1}$ is diagonalizable, then $p(\Lambda) = 0$, so $p(A) = 0$.
Since μ_A contains every factor $z - \lambda_i$ and divides p, we have $\mu_A = p$.

($\Leftarrow$) Suppose $\mu_A = p$. On the primary summand $\mathcal{G}_{\lambda_i}(A)$, the restriction of $A - \lambda_j I_n$ for $j \neq i$ is $(A - \lambda_i I_n) + (\lambda_i - \lambda_j)I$, a nonzero scalar plus a nilpotent operator. By Lemma 3.2.53, these restricted factors are invertible. Since $0 = p(A)|_{\mathcal{G}_{\lambda_i}(A)} = \left(\prod_{j \neq i} (A - \lambda_j I_n) \right) |_{\mathcal{G}_{\lambda_i}(A)} (A - \lambda_i I_n)|_{\mathcal{G}_{\lambda_i}(A)}$, the invertible product of the $j \neq i$ factors can be cancelled, giving $(A - \lambda_i I_n) = 0$ on $\mathcal{G}_{\lambda_i}(A)$.
Thus every primary summand is an eigenspace, so A is diagonalizable.

Structural Commentary: Repeated factors of μ_A record nilpotent parts in primary blocks.
Key Insight: Square-free μ_A kills each primary block in one step. □

Corollary 3.2.28. *Square-Free Annihilator Criterion*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. If a polynomial f with no repeated roots satisfies $f(A) = 0$, then A is diagonalizable.

Proof. Since $\mu_A \mid f$, the minimal polynomial is square-free; Theorem 3.2.27 applies.
Structural Commentary: A square-free annihilator forces square-free μ_A.
Key Insight: Square-free μ_A means diagonalizable. □

Method 3.2.29. *Computing μ_T from a Matrix*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ be the matrix of T in some basis.

(i) For $m = 1, 2, \ldots$, test the monic candidate $p_m(z) = a_0 + a_1 z + \cdots + a_{m-1} z^{m-1} + z^m$.
(ii) Solve linear system by comparing matrix entries in $a_0 I + a_1 A + \cdots + a_{m-1} A^{m-1} = -A^m$ for $a_0, \ldots, a_{m-1}$.
(iii) If no solution exists, increase m. The first solvable candidate is μ_T.

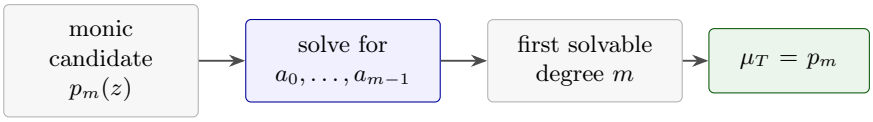


Figure 49: Solve $a_0 I + \cdots + a_{m-1} A^{m-1} = -A^m$; the first solvable monic candidate is μ_T.

Method 3.2.30. *Testing a Minimal Polynomial from p_A*
Suppose $p_A(z) = \prod_i (z - \lambda_i)^{a_i}$ is known.

i. Write each candidate as $m_{\mathbf{q}}(z) = \prod_i (z - \lambda_i)^{q_i}$, where $1 \le q_i \le a_i$.
ii. Test exponent vectors $\mathbf{q} = (q_i)$ in increasing total degree $\sum_i q_i$.
iii. For each $\mathbf{q}$, compute $m_{\mathbf{q}}(A) = \prod_i (A - \lambda_i I)^{q_i}$.
iv. The first candidate with $m_{\mathbf{q}}(A) = 0$ is μ_A, by Theorem 3.2.24.

Theorem 3.2.31. *Annihilators Are Multiples of μ_T*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dim, $T \in \mathcal{L}(\mathcal{V})$, and $f \in \mathbb{F}[z]$ be nonzero. Then $f(T) = 0$ if and only if μ_T divides f.

Proof.
(⇒) Suppose $f(T) = 0$. Divide $f = s\mu_T + r$, where $r = 0$ or $\deg r < \deg \mu_T$.
Then $r(T) = 0$, so minimality forces $r = 0$. Hence μ_T divides f.
(⇐) If $f = \mu_T s$, then $f(T) = \mu_T(T) s(T) = 0$.

Structural Commentary: Polynomial division reduces every annihilator to its remainder modulo μ_T.
Key Insight: Minimality forces that remainder to be zero. □

Corollary 3.2.32. *q_T Is a Multiple of μ_T*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$ and let $T \in \mathcal{L}(\mathcal{V})$.
Then the characteristic polynomial q_T is a polynomial multiple of the minimal polynomial μ_T.

Proof. Theorem 3.2.22 gives $q_T(T) = 0$. Theorem 3.2.31 then implies that μ_T divides q_T.
Structural Commentary: Cayley–Hamilton turns the characteristic polynomial into an annihilator.
Key Insight: The minimal polynomial divides every annihilator, hence divides q_T. □

Theorem 3.2.33. *Roots of μ_T Are Eigenvalues*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional and let $T \in \mathcal{L}(\mathcal{V})$. Then the roots of μ_T are precisely the eigenvalues of T.

Proof.
(⇒) If $\mu_T(\lambda) = 0$, write $\mu_T(z) = (z - \lambda) g(z)$. Since $\deg g < \deg \mu_T$, minimality gives $g(T) \neq 0$; choose $\mathbf{v}$ with $\mathbf{u} = g(T)\mathbf{v} \neq \mathbf{0}$. Then $(T - \lambda I_{\mathcal{V}})\mathbf{u} = \mu_T(T)\mathbf{v} = \mathbf{0}$.
(⇐) If $T\mathbf{v} = \lambda \mathbf{v}$ with $\mathbf{v} \neq \mathbf{0}$, then $0 = \mu_T(T)\mathbf{v} = \mu_T(\lambda)\mathbf{v}$, so $\mu_T(\lambda) = 0$.

Structural Commentary: A root splits off one shifted operator.
Key Insight: The quotient cannot already annihilate T. □

Corollary 3.2.34. *Minimal and Characteristic Polynomials Have the Same Roots*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$ and let $T \in \mathcal{L}(\mathcal{V})$.
Then $\{\lambda \in \mathbb{C} : \mu_T(\lambda) = 0\} = \{\lambda \in \mathbb{C} : q_T(\lambda) = 0\} = \mathrm{spec}(T)$.

Proof. By Theorem 3.2.33, $\{\lambda \in \mathbb{C} : \mu_T(\lambda) = 0\} = \operatorname{spec}(T)$.
By Proposition 3.2.17, $\{\lambda \in \mathbb{C} : q_T(\lambda) = 0\} = \operatorname{spec}(T)$.
Structural Commentary: Different multiplicities, same spectral scalars.
Key Insight: Minimal polynomial roots and characteristic polynomial roots are both just eigenvalues. □

Perspective 3.2.35. *Characteristic vs. Minimal Polynomial*
The characteristic polynomial records every primary block with its dimension. The minimal polynomial records the least polynomial needed to annihilate all blocks at once.

Theorem 3.2.36. *Annihilating Polynomials Control the Spectrum*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V})$, and let p be a nonconstant annihilating polynomial for T.

(i) Every eigenvalue of T is a root of p.
(ii) If $\mathbb{F} = \mathbb{C}$, then at least one root of p is an eigenvalue of T.

Proof.

(i) If $T\mathbf{v} = \lambda\mathbf{v}$ with $\mathbf{v} \neq \mathbf{0}$, then Proposition 3.2.5 gives $0 = p(T)\mathbf{v} = p(\lambda)\mathbf{v}$. Since $\mathbf{v} \neq \mathbf{0}$, $p(\lambda) = 0$.
(ii) Factor $p(z) = c(z - \lambda_1)\cdots(z - \lambda_m)$ over $\mathbb{C}$, with $c \neq 0$. Then $0 = p(T) = c(T - \lambda_1 I_{\mathcal{V}})\cdots(T - \lambda_m I_{\mathcal{V}})$. At least one factor $T - \lambda_j I_{\mathcal{V}}$ is not injective; otherwise product would be injective. Hence λ_j is an eigenvalue.

Structural Commentary: Annihilating polynomial restricts spectrum to its roots; over $\mathbb{C}$, one root must appear.
Key Insight: Eigenvectors turn $p(T)$ into $p(\lambda)$. □

Proposition 3.2.37. *Kernel and Range of a Polynomial in an Operator Are Invariant*
Let $T \in \mathcal{L}(\mathcal{V})$ and $p \in \mathbb{F}[z]$. Then $\ker(p(T))$ and $\operatorname{range}(p(T))$ are invariant under T.

Proof. If $\mathbf{v} \in \ker(p(T))$, then $p(T)\mathbf{v} = \mathbf{0}$.
Since $p(T)$ commutes with T by Proposition 3.2.3, $p(T)(T\mathbf{v}) = T(p(T)\mathbf{v}) = T\mathbf{0} = \mathbf{0}$.
If $\mathbf{v} \in \operatorname{range}(p(T))$, write $\mathbf{v} = p(T)\mathbf{u}$. Then $T\mathbf{v} = Tp(T)\mathbf{u} = p(T)(T\mathbf{u})$, so $T\mathbf{v} \in \operatorname{range}(p(T))$.
Structural Commentary: Polynomial kernels and ranges are stable under the original operator.
Key Insight: Since T commutes with $p(T)$, it preserves both $\ker p(T)$ and $\operatorname{range} p(T)$. □

Theorem 3.2.38. *Sylvester Linear Matrix Equation*
Let $A \in \operatorname{Mat}(m, m, \mathbb{C})$ and $B \in \operatorname{Mat}(n, n, \mathbb{C})$. If $\operatorname{spec}(A) \cap \operatorname{spec}(B) = \emptyset$, then for every $C \in \operatorname{Mat}(m, n, \mathbb{C})$ the equation $AX - XB = C$ has a unique solution $X \in \operatorname{Mat}(m, n, \mathbb{C})$.

Proof. Let $L(X) = AX - XB$. Since L has equal finite-dimensional domain and codomain, it suffices to show $\ker L = \{0\}$. If $AX = XB$, then $p(A)X = Xp(B)$ for every polynomial p.
With $p = p_B$, Cayley–Hamilton gives $p_B(A)X = Xp_B(B) = 0$. For every $\lambda \in \operatorname{spec}(A)$, disjointness gives $p_B(\lambda) \neq 0$; hence $p_B(A)$ is invertible by Corollary 3.5.17. Thus $X = 0$, so L is bijective.
Structural Commentary: Polynomial calculus pass through homogeneous solutions; disjoint spectra invert $p_B(A)$.
Key Insight: Spectral separation turns an intertwining relation into the zero relation. □

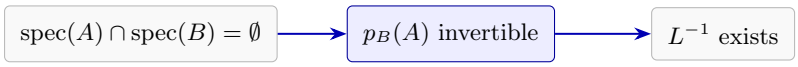


Figure 50: Disjoint spectra make $p_B(A)$ invertible, so the Sylvester map $L(X) = AX - XB$ is invertible.

Corollary 3.2.39. *Removing a Spectrally Separated Off-Diagonal Block*
Let $M = \begin{bmatrix} A & C \\ 0 & B \end{bmatrix}$ and $\operatorname{spec}(A) \cap \operatorname{spec}(B) = \emptyset$. Then M is similar to $A \oplus B$.

Proof. By Theorem 3.2.38, choose X such that $AX - XB = -C$. Then

$$\begin{bmatrix} I & -X \\ 0 & I \end{bmatrix} \begin{bmatrix} A & C \\ 0 & B \end{bmatrix} \begin{bmatrix} I & X \\ 0 & I \end{bmatrix} = \begin{bmatrix} A & AX + C - XB \\ 0 & B \end{bmatrix} = \begin{bmatrix} A & 0 \\ 0 & B \end{bmatrix}.$$

Structural Commentary: A block shear cancels the off-diagonal block via Sylvester's equation.
Key Insight: Spectrally separated diagonal blocks cannot carry essential upper-right coupling. □

Theorem 3.2.40. *Schur Form Grouped by Eigenvalues*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ have distinct eigenvalues $\lambda_1, \ldots, \lambda_d$ with algebraic multiplicities $n_1, \ldots, n_d$. Then A is unitarily similar to a block upper-triangular matrix

$$T = \begin{bmatrix} T_{11} & T_{12} & \cdots & T_{1d} \\ 0 & T_{22} & \cdots & T_{2d} \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & T_{dd} \end{bmatrix}, \qquad T_{ii} \in \mathrm{Mat}(n_i, n_i, \mathbb{C}),$$

where each T_{ii} is upper triangular with all diagonal entries equal to λ_i. Moreover, A is similar to $T_{11} \oplus \cdots \oplus T_{dd}$.

Proof. Order the Schur diagonal as $\lambda_1, \ldots, \lambda_1, \lambda_2, \ldots, \lambda_2, \ldots, \lambda_d, \ldots, \lambda_d$.
Theorem 3.4.10 gives the displayed block Schur form. Since $\mathrm{spec}(T_{11})$ is disjoint from the trailing block spectrum, Corollary 3.2.39 splits off T_{11}. Iterating splits off every diagonal block.
Structural Commentary: Grouped Schur form isolates spectral blocks; Sylvester shears split disjoint groups.
Key Insight: Cross-group upper blocks are coordinate artefacts. □

Corollary 3.2.41. *Commuting with Disjoint Blocks*
Let $A = A_{11} \oplus \cdots \oplus A_{dd}$, where the spectra of $A_{11}, \ldots, A_{dd}$ are pairwise disjoint. Partition $B = [B_{ij}]$ conformally with A. If $AB = BA$, then $B_{ij} = 0$ for all $i \neq j$.

Proof. Equating (i, j) blocks in $AB = BA$ gives $A_{ii}B_{ij} - B_{ij}A_{jj} = 0$.
For $i \neq j$, the spectra of A_{ii} and A_{jj} are disjoint, so Theorem 3.2.38 with $C = 0$ gives $B_{ij} = 0$.
Structural Commentary: A commuting matrix cannot mix direct-sum blocks with disjoint spectra.
Key Insight: The off-diagonal block solves a homogeneous Sylvester equation, whose only solution is zero. □

Generalized eigenvectors repair the main defect of ordinary eigenvectors: an operator may not have enough eigendirections to form a basis. The replacement is to study vectors killed by powers of shifted operators.

Definition 3.2.42. *Generalized Eigenvectors and Generalized Eigenspaces*
Let $T \in \mathcal{L}(\mathcal{V})$, $\lambda \in \mathbb{F}$. Nonzero vector $\mathbf{v} \in \mathcal{V}$ is a *generalized eigenvector* of T to λ if $(T - \lambda I_{\mathcal{V}})^j \mathbf{v} = \mathbf{0}$ for some positive integer j. The *generalized eigenspace* corresponding to λ is

$$\mathcal{G}_\lambda(T) := \{\mathbf{0}\} \cup \{\mathbf{v} \neq \mathbf{0} : \mathbf{v} \text{ is a generalized eigenvector of } T \text{ corresponding to } \lambda\}.$$

Every eigenvector is a generalized eigenvector by taking $j = 1$. The point is that higher powers of $T - \lambda I_{\mathcal{V}}$ may reveal directions that are not eigenvectors but are eventually pushed into the eigenspace.

Proposition 3.2.43. *Kernels of Powers Increase*
Let $T \in \mathcal{L}(\mathcal{V})$. Then $\{\mathbf{0}\} = \ker(T^0) \subseteq \ker(T) \subseteq \ker(T^2) \subseteq \cdots \subseteq \ker(T^k) \subseteq \ker(T^{k+1}) \subseteq \cdots$.

Proof. If $\mathbf{v} \in \ker(T^k)$, then $T^k\mathbf{v} = \mathbf{0}$, hence $T^{k+1}\mathbf{v} = T(T^k\mathbf{v}) = \mathbf{0}$. Thus $\ker(T^k) \subseteq \ker(T^{k+1})$ for every $k \geq 0$.
Structural Commentary: Powers of T form an increasing chain of kernels.
Key Insight: Once a vector is killed by T^k, applying one more T cannot revive it. □

Proposition 3.2.44. *Equality Forces Stabilization*
Let $T \in \mathcal{L}(\mathcal{V})$ and $m \geq 0$. If $\ker(T^m) = \ker(T^{m+1})$, then $\ker(T^m) = \ker(T^{m+1}) = \ker(T^{m+2}) = \cdots$.

Proof. By Proposition 3.2.43, it suffices to prove $\ker(T^{m+k+1}) \subseteq \ker(T^{m+k})$ for every $k \geq 0$.
Take $\mathbf{v} \in \ker(T^{m+k+1})$. Then $T^{m+1}(T^k\mathbf{v}) = T^{m+k+1}\mathbf{v} = \mathbf{0}$, so $T^k\mathbf{v} \in \ker(T^{m+1}) = \ker(T^m)$.
Hence $T^{m+k}\mathbf{v} = T^m(T^k\mathbf{v}) = \mathbf{0}$, proving the reverse inclusion.
Structural Commentary: Stabilization at one step propagates through the whole tail.
Key Insight: A later kernel vector can be shifted down into the already-stable kernel. □

Theorem 3.2.45. *Kernels Stop Growing by $\dim(\mathcal{V})$*
If $\mathcal{V}$ is finite-dimensional, $T \in \mathcal{L}(\mathcal{V})$, and $n = \dim(\mathcal{V})$, then $\ker(T^n) = \ker(T^{n+1}) = \ker(T^{n+2}) = \cdots$.

Proof. If $\ker(T^n) \neq \ker(T^{n+1})$, Prop 3.2.44 implies all inclusions $\{\mathbf{0}\} = \ker(T^0) \subsetneq \ker(T) \subsetneq \cdots \subsetneq \ker(T^{n+1})$ are strict. Each strict inclusion raises dimension by at least 1, so $\dim \ker(T^{n+1}) \geq n + 1$, impossible because $\ker(T^{n+1}) \subseteq \mathcal{V}$. Thus $\ker(T^n) = \ker(T^{n+1})$, and Proposition 3.2.44 gives the tail equality.
Structural Commentary: Strict kernel growth is dimension-limited.
Key Insight: Dimension bounds force the kernel chain to stabilize no later than step $\dim(\mathcal{V})$. □

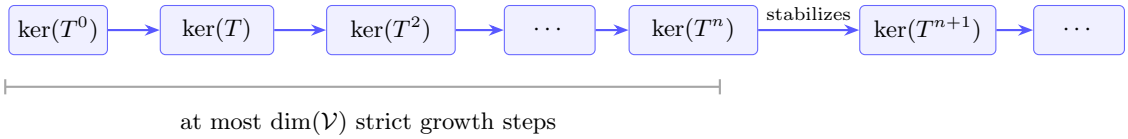


Figure 51: The kernel filtration grows until it stabilizes; after equality appears once, the tail is constant.

Theorem 3.2.46. *Ranks of Matrix Powers Stabilize*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$. Then:

(i) $\mathrm{rank}(A^k) \geq \mathrm{rank}(A^{k+1})$ *for* $k = 1, 2, \ldots$;
(ii) if $\mathrm{rank}(A^k) = \mathrm{rank}(A^{k+1})$, *then* $\mathrm{col}(A^k) = \mathrm{col}(A^{k+p})$ *and* $\mathrm{rank}(A^k) = \mathrm{rank}(A^{k+p})$ *for every* $p = 1, 2, \ldots$;
(iii) if A is not invertible, then there is a least positive integer $q \leq n$ such that $\mathrm{rank}(A^q) = \mathrm{rank}(A^{q+1})$.

Proof.

(i) Since $A^{k+1} = A^k A$, each column of A^{k+1} is a linear combination of the columns of A^k. Hence $\mathrm{col}(A^{k+1}) \subseteq \mathrm{col}(A^k)$, so $\mathrm{rank}(A^{k+1}) \leq \mathrm{rank}(A^k)$.
(ii) The inclusion $\mathrm{col}(A^{k+1}) \subseteq \mathrm{col}(A^k)$ and equality of dimensions give $\mathrm{col}(A^{k+1}) = \mathrm{col}(A^k)$. Let $U = \mathrm{col}(A^k)$. Then multiplication by A maps U onto U, so induction gives $\mathrm{col}(A^{k+p}) = U$ for every $p \geq 1$.
(iii) If A is not invertible, then Corollary 1.3.95 gives $\mathrm{rank}(A) \leq n-1$. If all inequalities $\mathrm{rank}(A^k) > \mathrm{rank}(A^{k+1})$ held for $k = 1, \ldots, n$, then $\mathrm{rank}(A^{n+1}) \leq \mathrm{rank}(A) - n < 0$, impossible. Hence equality occurs for some $q \leq n$; choose the least such q.

Structural Commentary: The column spaces of powers form a decreasing chain. In finite dimensions, a decreasing chain cannot keep losing dimension forever.
Key Insight: Once two consecutive ranks agree, the image has stabilized under multiplication by A. □

Definition 3.2.47. *Index of a Matrix*
If $A \in \mathrm{Mat}(n, n, \mathbb{F})$ is invertible, define its *index* to be 0. If A is not invertible, its index is the least positive integer q such that $\mathrm{rank}(A^q) = \mathrm{rank}(A^{q+1})$.

Theorem 3.2.48. *Stable Kernel-Range Decomposition*
Let $\mathcal{V}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V})$, and set $n = \dim(\mathcal{V})$. Then $\mathcal{V} = \ker(T^n) \oplus \mathrm{range}(T^n)$.

Proof. First show the intersection is trivial. If $\mathbf{v} \in \ker(T^n) \cap \mathrm{range}(T^n)$, then $\mathbf{v} = T^n\mathbf{u}$ and $T^n\mathbf{v} = \mathbf{0}$. Hence $T^{2n}\mathbf{u} = \mathbf{0}$, so $\mathbf{u} \in \ker(T^{2n}) = \ker(T^n)$ by Theorem 3.2.45. Thus $\mathbf{v} = T^n\mathbf{u} = \mathbf{0}$.
By Rank-Nullity Theorem 1.3.27, $\dim\ker(T^n) + \dim\mathrm{range}(T^n) = \dim(\mathcal{V})$.
The zero intersection and the dimension count imply the sum is direct and equals $\mathcal{V}$.
Structural Commentary: At the stable power, killed directions and surviving image directions separate cleanly.
Key Insight: Stabilization turns the usual rank-nullity count into an actual direct-sum decomposition. □

Theorem 3.2.49. *Description of Generalized Eigenspaces*
Let $\mathcal{V}$ be finite-dimensional, $T \in \mathcal{L}(\mathcal{V})$, $\lambda \in \mathbb{F}$, and $n = \dim(\mathcal{V})$. Then $\mathcal{G}_\lambda(T) = \ker\big((T - \lambda I_\mathcal{V})^n\big)$.
In particular, $\mathcal{G}_\lambda(T)$ is a subspace of $\mathcal{V}$.

Proof. Set $S = T - \lambda I_\mathcal{V}$.

($\Rightarrow$) If $\mathbf{v} \in \mathcal{G}_\lambda(T)$, then $S^j\mathbf{v} = \mathbf{0}$ for some $j \geq 1$. If $j \leq n$, Proposition 3.2.43 gives $\mathbf{v} \in \ker(S^n)$; if $j > n$, Theorem 3.2.45 gives $\ker(S^j) = \ker(S^n)$.
($\Leftarrow$) If $\mathbf{v} \in \ker(S^n)$, then either $\mathbf{v} = \mathbf{0}$ or $\mathbf{v}$ is a generalized eigenvector of T corresponding to λ.

Thus $\mathcal{G}_\lambda(T) = \ker(S^n)$, which is a subspace.
Structural Commentary: Generalized eigenspaces are stable shifted kernels.
Key Insight: Finite dimension makes exponent $\dim(\mathcal{V})$ enough. □

Proposition 3.2.50. *Generalized Eigenspaces Are Invariant*
Let $\mathcal{V}$ be finite-dimensional and $T \in \mathcal{L}(\mathcal{V})$. Then $\mathcal{G}_\lambda(T)$ is invariant under T for every $\lambda \in \mathbb{F}$.

Proof. By Theorem 3.2.49, $\mathcal{G}_\lambda(T) = \ker((T - \lambda I_\mathcal{V})^n)$, where $n = \dim(\mathcal{V})$.
This is the kernel of a polynomial in T, so it is invariant under T by Proposition 3.2.37.
Structural Commentary: A generalized eigenspace is cut out by a polynomial equation in T.
Key Insight: Polynomial kernels commute with the operator that defines them. □

Definition 3.2.51. *Nilpotent Operator*
An operator $N \in \mathcal{L}(\mathcal{V})$ is *nilpotent* if $N^m = 0$ for some positive integer m.

Corollary 3.2.52. *Nilpotent Power Bound*
If $\mathcal{V}$ is finite-dimensional, $N \in \mathcal{L}(\mathcal{V})$ is nilpotent, and $n = \dim(\mathcal{V})$, then $N^n = 0$.

Proof. Since N is nilpotent, $\ker(N^m) = \mathcal{V}$ for some m.
The kernel chain has stabilized by the nth step by Theorem 3.2.45, so $\ker(N^n) = \mathcal{V}$. Hence $N^n = 0$.
Structural Commentary: Nilpotence says some kernel in the chain is already the whole space.
Key Insight: The stabilization bound pulls that full-kernel conclusion back to exponent $\dim(\mathcal{V})$. □

Lemma 3.2.53. *Nonzero Scalar Plus Nilpotent Is Invertible*
Let $N \in \mathcal{L}(\mathcal{V})$ be nilpotent and let $a \in \mathbb{F}$ with $a \neq 0$. Then $aI_{\mathcal{V}} + N$ is invertible.

Proof. Choose m such that $N^m = 0$. Then $(aI_{\mathcal{V}} + N)\left(a^{-1}\sum_{r=0}^{m-1}(-a^{-1}N)^r\right) = I_{\mathcal{V}} - (-a^{-1}N)^m = I_{\mathcal{V}}$, and the same computation in the reverse order gives a two-sided inverse.
Structural Commentary: A nilpotent perturbation has a finite geometric-series inverse.
Key Insight: Nilpotence truncates the Neumann series exactly. □

Theorem 3.2.54. *Generalized Eigenvectors for Distinct Eigenvalues Are Independent*
Let $\mathcal{V}$ be finite-dimensional and $T \in \mathcal{L}(\mathcal{V})$. Suppose $\lambda_1, \ldots, \lambda_m$ are distinct eigenvalues of T, and let $\mathbf{v}_1, \ldots, \mathbf{v}_m$ be corresponding generalized eigenvectors. Then $(\mathbf{v}_1, \ldots, \mathbf{v}_m)$ is linearly independent.

Proof. Let $n = \dim(\mathcal{V})$ and suppose $a_1\mathbf{v}_1 + \cdots + a_m\mathbf{v}_m = \mathbf{0}$.
Fix $r \in \{1, \ldots, m\}$ and apply $Q_r(T) := \prod_{i \neq r}(T - \lambda_i I_{\mathcal{V}})^n$ to the relation.
For $i \neq r$, Theorem 3.2.49 gives $Q_r(T)\mathbf{v}_i = \mathbf{0}$, so $a_r Q_r(T)\mathbf{v}_r = \mathbf{0}$.
On $\mathcal{G}_{\lambda_r}(T)$, the operator $T - \lambda_i I_{\mathcal{V}}$ equals $(T - \lambda_r I_{\mathcal{V}}) + (\lambda_r - \lambda_i)I_{\mathcal{V}}$. The first term is nilpotent on $\mathcal{G}_{\lambda_r}(T)$ and $\lambda_r - \lambda_i \neq 0$, so Lemma 3.2.53 makes each factor invertible on $\mathcal{G}_{\lambda_r}(T)$. Thus $Q_r(T)\mathbf{v}_r \neq \mathbf{0}$, forcing $a_r = 0$. Since r was arbitrary, all coefficients vanish.
Structural Commentary: Polynomial filters isolate one generalized eigenspace at a time.
Key Insight: Distinct eigenvalues let polynomial filters kill the other components and preserve the target one. □

Theorem 3.2.55. *Nilpotent Operators and Strictly Upper-Triangular Matrices*
Let $\mathcal{V}$ be finite-dimensional and $N \in \mathcal{L}(\mathcal{V})$. Then N is nilpotent if and only if N has a strictly upper-triangular matrix in some basis.

Proof. Let $n = \dim(\mathcal{V})$. By Corollary 3.2.52, $\{\mathbf{0}\} \subseteq \ker(N) \subseteq \ker(N^2) \subseteq \cdots \subseteq \ker(N^n) = \mathcal{V}$.

(⇒) Choose a basis of $\ker(N)$, extend it to a basis of $\ker(N^2)$, and continue until obtaining a basis of $\mathcal{V}$. If $\mathbf{v}$ first appears in $\ker(N^k)$, then $N\mathbf{v} \in \ker(N^{k-1})$, which is spanned by earlier basis vectors. Hence the matrix of N is strictly upper triangular.

(⇐) If the matrix of N is $n \times n$ strictly upper triangular, then N sends each basis vector into the span of earlier basis vectors. After n applications every basis vector is sent to $\mathbf{0}$, so $N^n = 0$.

Structural Commentary: Kernel filtration gives adapted basis; triangular form records motion to earlier layers.
Key Insight: Nilpotence is strict upper-triangularity in a suitable basis. □

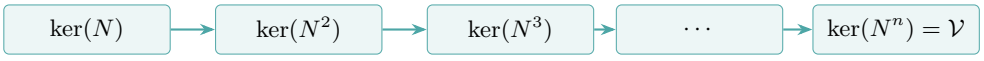


Figure 52: Extending bases through the kernel filtration makes N move vectors to earlier layers.

Perspective 3.2.56. *Eigenvalue Plus Nilpotent Part*
On $\mathcal{G}_\lambda(T)$, the operator T behaves as $\lambda I_{\mathcal{V}}$ plus the nilpotent operator $T - \lambda I_{\mathcal{V}}$. Generalized eigenvectors are therefore ordinary eigenvectors enriched by nilpotent chain data.

Theorem 3.2.57. *Primary Decomposition over $\mathbb{C}$*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dim over $\mathbb{C}$, let $T \in \mathcal{L}(\mathcal{V})$, and let $\lambda_1, \ldots, \lambda_m$ be the distinct eigenvalues of T. Then

(i) $\mathcal{V} = \mathcal{G}_{\lambda_1}(T) \oplus \cdots \oplus \mathcal{G}_{\lambda_m}(T)$;
(ii) each $(T - \lambda_j I_{\mathcal{V}})|_{\mathcal{G}_{\lambda_j}(T)}$ *is nilpotent.*

Proof. Let $n = \dim(\mathcal{V})$.

(i) Proceed by induction on n. If $n = 1$, then T has single eigenvalue λ_1 and $\mathcal{V} = \mathcal{G}_{\lambda_1}(T)$.
Assume $n > 1$. Let $S = T - \lambda_1 I_{\mathcal{V}}$, $\mathcal{U} := \operatorname{range}(S^n)$.
By Theorems 3.2.48 and 3.2.49, $\mathcal{V} = \ker(S^n) \oplus \operatorname{range}(S^n) = \mathcal{G}_{\lambda_1}(T) \oplus \mathcal{U}$. The subspace $\mathcal{U}$ is T-invariant by Proposition 3.2.37, and $\dim(\mathcal{U}) < n$ because $\mathcal{G}_{\lambda_1}(T) \neq \{\mathbf{0}\}$. If $\mathcal{U} = \{\mathbf{0}\}$, then (i) follows.
Otherwise distinct eigenvalues of $T|_{\mathcal{U}}$ are $\lambda_2, \ldots, \lambda_m$: every eigenvalue of $T|_{\mathcal{U}}$ is an eigenvalue of T but not λ_1, while an eigenvector $\mathbf{v}$ of T for λ_j, $j \geq 2$, satisfies $S^n\mathbf{v} = (\lambda_j - \lambda_1)^n\mathbf{v}$ and hence lies in $\mathcal{U}$.
Induction gives $\mathcal{U} = \mathcal{G}_{\lambda_2}(T|_{\mathcal{U}}) \oplus \cdots \oplus \mathcal{G}_{\lambda_m}(T|_{\mathcal{U}})$. For $j \geq 2$, $\mathcal{G}_{\lambda_j}(T|_{\mathcal{U}}) = \mathcal{G}_{\lambda_j}(T)$.
The inclusion $\subseteq$ follows from the definition. Conversely, if $\mathbf{v} \in \mathcal{G}_{\lambda_j}(T)$, write $\mathbf{v} = \mathbf{g} + \mathbf{u}$ with $\mathbf{g} \in \mathcal{G}_{\lambda_1}(T)$ and $\mathbf{u} \in \mathcal{U}$. For some r, $(T - \lambda_j I_{\mathcal{V}})^r\mathbf{v} = \mathbf{0}$. Since both summands are T-invariant, the two terms $(T - \lambda_j I_{\mathcal{V}})^r\mathbf{g}$ and $(T - \lambda_j I_{\mathcal{V}})^r\mathbf{u}$ remain in the two direct summands and have sum zero, so both are zero. Thus $\mathbf{g} \in \mathcal{G}_{\lambda_1}(T) \cap \mathcal{G}_{\lambda_j}(T) = \{\mathbf{0}\}$ by Theorem 3.2.54, and $\mathbf{v} = \mathbf{u} \in \mathcal{G}_{\lambda_j}(T|_{\mathcal{U}})$. Hence (i) follows.

(ii) By Proposition 3.2.50, the restriction is an operator on $\mathcal{G}_{\lambda_j}(T)$. By Theorem 3.2.49, $(T - \lambda_j I_{\mathcal{V}})^n\mathbf{v} = \mathbf{0}$ $\forall \mathbf{v} \in \mathcal{G}_{\lambda_j}(T)$. Hence nth power of $(T - \lambda_j I_{\mathcal{V}})|_{\mathcal{G}_{\lambda_j}(T)}$ is zero.

Structural Commentary: Peel off one generalized eigenspace, then induct on the invariant range.
Key Insight: Stable kernel-range splitting isolates one eigenvalue at a time. □

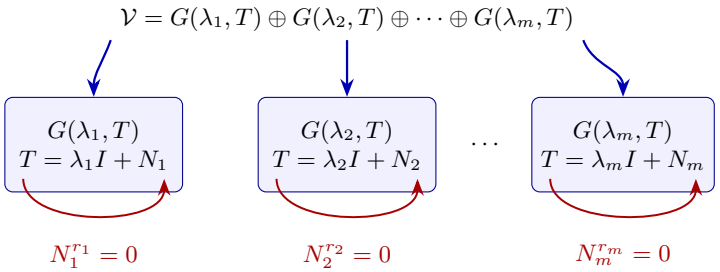


Figure 53: Primary decomposition: on $\mathcal{G}_{\lambda_j}(T)$, $T = \lambda_j I$ plus nilpotent.

Corollary 3.2.58. *Basis of Generalized Eigenvectors*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$ and $T \in \mathcal{L}(\mathcal{V})$. Then $\mathcal{V}$ has a generalized-eigenvector basis for T.

Proof. Let $\lambda_1, \ldots, \lambda_m$ be the distinct eigenvalues of T. By Theorem 3.2.57, $\mathcal{V} = \mathcal{G}_{\lambda_1}(T) \oplus \cdots \oplus \mathcal{G}_{\lambda_m}(T)$. Choose a basis of each generalized eigenspace and concatenate the bases. Directness makes the concatenated list a basis of $\mathcal{V}$, and every nonzero vector in each list is a generalized eigenvector.
Structural Commentary: A direct sum basis is obtained by choosing bases in each summand.
Key Insight: Bases of the direct-sum pieces concatenate. □

Definition 3.2.59. *Multiplicity of an Eigenvalue*
Let $\mathcal{V}$ be finite-dimensional and let $T \in \mathcal{L}(\mathcal{V})$.
The *multiplicity* of an eigenvalue λ of T is $\dim \mathcal{G}_\lambda(T) = \dim \ker\big((T - \lambda I_{\mathcal{V}})^{\dim(\mathcal{V})}\big)$.
This is the algebraic multiplicity in generalized-eigenspace language.

Proposition 3.2.60. *Sum of Multiplicities*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be a finite-dimensional complex vector space and let $T \in \mathcal{L}(\mathcal{V})$. Then the sum of the multiplicities of all eigenvalues of T equals $\dim(\mathcal{V})$.

Proof. Let $\lambda_1, \ldots, \lambda_m$ be the distinct eigenvalues of T. By Theorem 3.2.57, $\mathcal{V} = \mathcal{G}_{\lambda_1}(T) \oplus \cdots \oplus \mathcal{G}_{\lambda_m}(T)$.
Taking dimensions, by Corollary 1.2.20, $\dim(\mathcal{V}) = \sum_{j=1}^{m} \dim \mathcal{G}_{\lambda_j}(T)$, exactly the sum of the multiplicities.
Structural Commentary: Multiplicity is dimension of the generalized-eigenspace summand.
Key Insight: Multiplicities are dimensions of the direct-sum pieces. □

Theorem 3.2.61. *Block Diagonal Form with Upper-Triangular Primary Blocks*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$, and let $\lambda_1, \ldots, \lambda_m$ be the distinct eigenvalues of $T \in \mathcal{L}(\mathcal{V})$, with multiplicities $d_1, \ldots, d_m$. Then T has a matrix of the form

$$\begin{bmatrix} A_1 & & 0 \\ & \ddots & \\ 0 & & A_m \end{bmatrix},$$

where each $A_j \in \operatorname{Mat}(d_j, d_j, \mathbb{C})$ is upper triangular with diagonal $\lambda_j, \ldots, \lambda_j$.

Proof. Set $\mathcal{U}_j = \mathcal{G}_{\lambda_j}(T)$ and $N_j = (T - \lambda_j I_{\mathcal{V}})|_{\mathcal{U}_j}$.
By Theorem 3.2.57, the $\mathcal{U}_j$ form a T-invariant direct-sum decomposition and each N_j is nilpotent.
Choose a basis of each $\mathcal{U}_j$ in which N_j is strictly upper triangular. Then $T|_{\mathcal{U}_j} = \lambda_j I_{\mathcal{U}_j} + N_j$ has an upper triangular matrix with diagonal $\lambda_j, \ldots, \lambda_j$. Concatenating these bases gives the desired block diagonal matrix.
Structural Commentary: Primary decomposition gives the blocks; nilpotent triangularisation shapes them.
Key Insight: On each generalized eigenspace, T is eigenvalue plus nilpotent. □

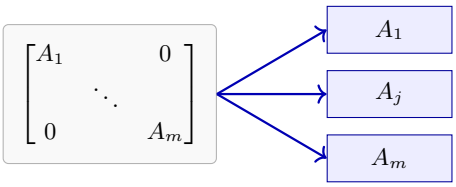


Figure 54: Primary blocks: $A_j = \lambda_j I + N_j$, with N_j strictly upper triangular.

3.3 Jordan and Canonical Forms

Jordan and canonical forms turn primary decomposition into explicit blocks. Jordan chains, complexification, and real block forms organize the nilpotent and real spectral structure. Jordan blocks are elementary divisors $(z - \lambda)^k$ after the relevant polynomial factors split into linear factors.

Definition 3.3.1. *Jordan Chains, Blocks, and Matrices*
For $k \geq 1$, the *nilpotent Jordan block* $J_k(0)$ is the $k \times k$ matrix with 1 on the first superdiagonal and 0 elsewhere:

$$J_k(0) = \begin{bmatrix} 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \ddots & \vdots \\ \vdots & \ddots & \ddots & \ddots & 0 \\ 0 & \cdots & 0 & 0 & 1 \\ 0 & \cdots & \cdots & 0 & 0 \end{bmatrix}.$$

For $\lambda \in \mathbb{C}$, set $J_k(\lambda) := \lambda I_k + J_k(0)$. A *Jordan matrix* is a direct sum of Jordan blocks. A basis is a *Jordan basis* for an operator T if the matrix of T in that basis is a Jordan matrix.

Perspective 3.3.2. *Jordan Strings as Columns*
A length-k Jordan string for A at λ is $(\mathbf{x}_1, \ldots, \mathbf{x}_k)$ with $(A - \lambda I)\mathbf{x}_1 = \mathbf{0}$; $(A - \lambda I)\mathbf{x}_j = \mathbf{x}_{j-1}$, $2 \leq j \leq k$.
For $X = [\mathbf{x}_1 \ \cdots \ \mathbf{x}_k]$, $AX = XJ_k(\lambda)$.
If all strings form the columns of a basis matrix M, then $AM = MJ$ and $M^{-1}AM = J$.

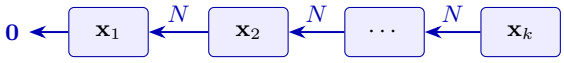


Figure 55: Jordan chain for $N = A - \lambda I$; columns give $AX = XJ_k(\lambda)$.

Method 3.3.3. *Filippov Column-Space Jordan Strings*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$, let $\lambda \in \mathrm{spec}(A)$, and set $N = A - \lambda I_n$ on the generalized eigenspace $\mathcal{G}_\lambda(A) = \ker(N^n)$.

(i) Work first inside $\mathrm{range}(N) \cap \mathcal{G}_\lambda(A)$, where N is again nilpotent.
(ii) Build Jordan strings for this smaller nilpotent operator.
(iii) For each string ending in $\mathbf{u}$, solve $N\mathbf{v} = \mathbf{u}$ in $\mathcal{G}_\lambda(A)$ and lift the string by appending $\mathbf{v}$.
(iv) Add enough vectors from $\ker(N)$ to complete a basis of $\mathcal{G}_\lambda(A)$; these added vectors are length-one strings.
(v) Repeat the process recursively until all of $\mathcal{G}_\lambda(A)$ is covered, then concatenate the strings for all eigenvalues.

Theorem 3.3.5 justifies the lifts; Perspective 3.3.2 turns the resulting columns into $AM = MJ$.

Proposition 3.3.4. *Powers of a Jordan Block*
For $p \geq 0$, $J_k(\lambda)^p = \sum_{r=0}^{\min(p,k-1)} \binom{p}{r} \lambda^{p-r} J_k(0)^r$. Thus $J_k(\lambda)^p \to 0$ if and only if $|\lambda| < 1$.

Proof.

($\Rightarrow$) If $J_k(\lambda)^p \to 0$, then its diagonal entries λ^p tend to 0, so $|\lambda| < 1$.
($\Leftarrow$) If $|\lambda| < 1$, then for each fixed $0 \leq r < k$ the scalar $\binom{p}{r}\lambda^{p-r}$ tends to 0, because a polynomial in p is dominated by the exponential decay of $|\lambda|^p$.

Structural Commentary: Jordan powers are exponential-polynomial terms.
Key Insight: Polynomial growth cannot defeat exponential decay when $|\lambda| < 1$. □

Theorem 3.3.5. *Nilpotent Jordan Chain Basis*
Let $\mathcal{V}$ be finite-dim, $N \in \mathcal{L}(\mathcal{V})$ be nilpotent. Then $\exists$ vectors $\mathbf{v}_1, \dots, \mathbf{v}_r \in \mathcal{V}$ and nonnegative integers $m_1, \dots, m_r$ such that $\mathcal{B}_N = \big(N^{m_1}\mathbf{v}_1, \dots, N\mathbf{v}_1, \mathbf{v}_1, \dots, N^{m_r}\mathbf{v}_r, \dots, N\mathbf{v}_r, \mathbf{v}_r\big)$ is a basis of $\mathcal{V}$ and $N^{m_j+1}\mathbf{v}_j = \mathbf{0}$ for every j.

Proof. Induct on $n = \dim(\mathcal{V})$. If $N = 0$, any basis works with all $m_j = 0$.
Now assume $N \neq 0$. Then N is not injective, so $\operatorname{range}(N)$ is a proper subspace, and $N|_{\operatorname{range}(N)}$ is nilpotent.
By induction, $\operatorname{range}(N)$ has a chain basis $\big(N^{m_1}\mathbf{u}_1, \dots, N\mathbf{u}_1, \mathbf{u}_1, \dots, N^{m_r}\mathbf{u}_r, \dots, N\mathbf{u}_r, \mathbf{u}_r\big)$ with $N^{m_j+1}\mathbf{u}_j = \mathbf{0}$.
Choose $\mathbf{v}_j \in \mathcal{V}$ such that $N\mathbf{v}_j = \mathbf{u}_j$, and lift the chains to $\mathcal{L} = \big(N^{m_1+1}\mathbf{v}_1, \dots, N\mathbf{v}_1, \mathbf{v}_1, \dots, N^{m_r+1}\mathbf{v}_r, \dots, N\mathbf{v}_r, \mathbf{v}_r\big)$.
The list $\mathcal{L}$ is linearly independent. Indeed, applying N to a dependence among vectors in $\mathcal{L}$ gives a dependence among the range-chain basis, so all coefficients except possibly those on the leading vectors $N^{m_j+1}\mathbf{v}_j$ vanish.
The remaining relation is a relation among $N^{m_j+1}\mathbf{v}_j = N^{m_j}\mathbf{u}_j$, which are members of the range-chain basis, so those coefficients vanish too. Also $N(\operatorname{span}(\mathcal{L})) = \operatorname{range}(N)$: each range-chain basis vector $N^q\mathbf{u}_j$ is $N(N^q\mathbf{v}_j)$ for the corresponding lifted vector. Extend $\mathcal{L}$ to a basis by adding $\mathbf{w}_1, \dots, \mathbf{w}_s$. For each i, choose $\mathbf{x}_i \in \operatorname{span}(\mathcal{L})$ such that $N\mathbf{x}_i = N\mathbf{w}_i$, and set $\mathbf{z}_i = \mathbf{w}_i - \mathbf{x}_i$. Then $N\mathbf{z}_i = \mathbf{0}$. Since each $\mathbf{x}_i$ lies in $\operatorname{span}(\mathcal{L})$, replacing $\mathbf{w}_i$ by $\mathbf{z}_i$ preserves the span and linear independence of the full basis. Thus $\mathcal{L}, \mathbf{z}_1, \dots, \mathbf{z}_s$ is the desired chain basis.
Structural Commentary: Build chains in $\operatorname{range}(N)$, lift them once, then add kernel chains.
Key Insight: Applying N turns lifted-chain dependences into known range-chain dependences. □

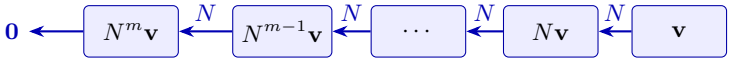


Figure 56: A nilpotent chain shifts left under N; in this basis order it becomes one Jordan block.

Theorem 3.3.6. *Jordan Form over $\mathbb{C}$*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$ and let $T \in \mathcal{L}(\mathcal{V})$. Then $\mathcal{V}$ has a Jordan basis for T.

Proof. Let $\lambda_1, \dots, \lambda_m$ be the distinct eigenvalues of T. By Theorem 3.2.57, $\mathcal{V} = \mathcal{G}_{\lambda_1}(T) \oplus \cdots \oplus \mathcal{G}_{\lambda_m}(T)$, and each $N_j := (T - \lambda_j I_{\mathcal{V}})|_{\mathcal{G}_{\lambda_j}(T)}$ is nilpotent. Apply Theorem 3.3.5 to N_j on each generalized eigenspace.
The resulting chain basis of $\mathcal{G}_{\lambda_j}(T)$ is a basis of that summand, and the direct-sum decomposition makes the concatenation of these bases a basis of $\mathcal{V}$. In a chain of length r for N_j, the matrix of N_j is $J_r(0)$, so the matrix of $T|_{\mathcal{G}_{\lambda_j}(T)} = \lambda_j I + N_j$ on the same chain is $J_r(\lambda_j)$. Therefore the concatenated basis is a Jordan basis for T.
Structural Commentary: Primary pieces decompose into nilpotent chains.
Key Insight: Jordan form is "eigenvalue plus nilpotent chains" on every generalized eigenspace. □

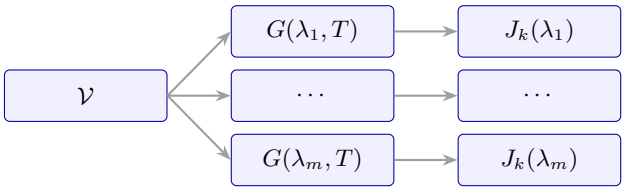


Figure 57: Primary pieces $\mathcal{G}_{\lambda_j}(T)$ give nilpotent chains, hence Jordan blocks.

Corollary 3.3.7. *Diagonalizability in Jordan Form*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$ have Jordan form J. For each eigenvalue λ, $\dim E_\lambda(A)$ equals the number of λ-blocks in J. Hence A is diagonalizable if and only if every Jordan block has size 1.

Proof. Write $A = SJS^{-1}$. Then $S : E_\lambda(J) \to E_\lambda(A)$ is an isomorphism. A block $J_k(\mu)$ contributes to $E_\lambda(J)$ only when $\mu = \lambda$, and then $\dim\ker(J_k(\lambda) - \lambda I_k) = \dim\ker J_k(0) = 1$.
Thus $\dim E_\lambda(A)$ is the number of λ-blocks. By Theorem 3.2.15, diagonalizability means geometric and algebraic multiplicities agree for every λ; here this holds exactly when every block has size 1.
Structural Commentary: Each Jordan block contributes one eigenvector direction.
Key Insight: Missing eigenvectors are measured by block sizes greater than 1. □

Lemma 3.3.8. *Powers of a Nilpotent Jordan Block*
For $k \geq 1$ and $p \geq 0$, $\operatorname{rank}\big(J_k(0)^p\big) = \max(k - p, 0)$ and $\dim\ker\big(J_k(0)^p\big) = \min(p, k)$.

Proof. Let $N = J_k(0)$. In the standard basis, $N\mathbf{e}_1 = \mathbf{0}$ and $N\mathbf{e}_j = \mathbf{e}_{j-1}$ for $2 \le j \le k$. Iterating gives

$$N^p \mathbf{e}_j = \begin{cases} \mathbf{0}, & j \le p, \\ \mathbf{e}_{j-p}, & j > p. \end{cases}$$

If $p < k$, the image is $\operatorname{span}(\mathbf{e}_1, \dots, \mathbf{e}_{k-p})$, so the rank is $k - p$; if $p \ge k$, the image is $\{\mathbf{0}\}$, so the rank is 0. Hence $\operatorname{rank}(N^p) = \max(k-p, 0)$. Rank-nullity gives $\dim \ker(N^p) = k - \max(k-p, 0) = \min(p, k)$.
Structural Commentary: A nilpotent Jordan block is a left-shift on the standard basis.
Key Insight: Each power deletes one more leading layer of the chain. □

Definition 3.3.9. *Weyr Characteristic*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$ and $\lambda \in \operatorname{spec}(A)$.
For $p \ge 1$, define $w_p(\lambda, A) := \operatorname{rank}\big((A - \lambda I_n)^{p-1}\big) - \operatorname{rank}\big((A - \lambda I_n)^p\big)$.
The positive part of the sequence $w_1, w_2, \dots$ is the *Weyr characteristic* of A associated with λ.

Theorem 3.3.10. *Jordan Block Sizes from Rank Drops*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$, $\lambda \in \operatorname{spec}(A)$, and $p \ge 1$. Then $w_p(\lambda, A)$ is the number of λ-Jordan blocks of size at least p; hence $w_p(\lambda, A) - w_{p+1}(\lambda, A)$ is the number of λ-blocks of size exactly p.

Proof. Rank is similarity-invariant, so assume A is in Jordan form.
Blocks $J_k(\mu)$ with $\mu \ne \lambda$ contribute 0 to the rank drop because $J_k(\mu) - \lambda I_k$ is invertible.
A block $J_k(\lambda)$ contributes $\operatorname{rank}\big(J_k(0)^{p-1}\big) - \operatorname{rank}\big(J_k(0)^p\big)$, which is 1 exactly when $p \le k$ by Lemma 3.3.8.
Summing over blocks proves the first claim; subtracting the $p+1$ count gives the exact-size count.
Structural Commentary: Rank drops count surviving chain lengths.
Key Insight: A length-k block contributes to drops $1, \dots, k$. □

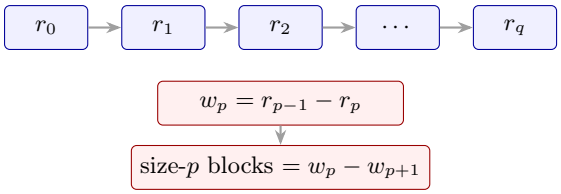


Figure 58: For $r_p = \operatorname{rank}((A - \lambda I)^p)$, drops count surviving chains; $w_p - w_{p+1}$ counts size-p blocks.

Method 3.3.11. *Computing Jordan Blocks by Rank Drops*
For each eigenvalue λ of $A \in \operatorname{Mat}(n, n, \mathbb{C})$:

(i) Compute $r_p = \operatorname{rank}\big((A - \lambda I_n)^p\big)$ for $p = 0, 1, 2, \dots$ until the differences stabilize at 0.
(ii) Set $w_p = r_{p-1} - r_p$. Then w_p is the number of λ-blocks of size at least p.
(iii) The number of blocks $J_p(\lambda)$ is $w_p - w_{p+1}$.
(iv) Assemble the Jordan form by taking the direct sum over all eigenvalues.

Theorem 3.3.12. *Nonzero Eigenvalues of AB and BA*
Let $A \in \operatorname{Mat}(m, n, \mathbb{C})$, $B \in \operatorname{Mat}(n, m, \mathbb{C})$, and suppose $n \ge m$. Then $p_{BA}(z) = z^{n-m} p_{AB}(z)$.
Consequently, for $\lambda \ne 0$, $m_{AB}(\lambda) = m_{BA}(\lambda)$. If $m_{AB}(0) = k$, then $m_{BA}(0) = k + n - m$.

Proof. For $z \ne 0$, Sylvester's determinant identity gives

$$\det(zI_n - BA) = z^n \det(I_n - z^{-1}BA) = z^n \det(I_m - z^{-1}AB) = z^{n-m} \det(zI_m - AB).$$

Both sides are polynomials in z, so the identity holds for all z. The claimed multiplicities follow by factoring p_{AB} and comparing with $p_{BA} = z^{n-m} p_{AB}$.
Structural Commentary: Sylvester's determinant identity moves the rectangular product across the determinant.
Key Insight: Passing from AB to BA only adds the necessary extra zero roots when the sizes differ. □

Theorem 3.3.13. *Nonzero Jordan Blocks of AB and BA*
Let $A, B \in \operatorname{Mat}(n, n, \mathbb{C})$. For every $\lambda \ne 0$, the Jordan blocks of AB and BA at λ coincide.

Proof. Set $X = \begin{bmatrix} AB & A \\ 0 & 0 \end{bmatrix}$ and $Y = \begin{bmatrix} 0 & A \\ 0 & BA \end{bmatrix}$. Then $\begin{bmatrix} I_n & 0 \\ B & I_n \end{bmatrix} X \begin{bmatrix} I_n & 0 \\ -B & I_n \end{bmatrix} = Y$, so X and Y are similar.
For $\lambda \ne 0$ and $p \ge 1$, the shifted powers have block triangular forms

$$(X - \lambda I_{2n})^p = \begin{bmatrix} (AB - \lambda I_n)^p & * \\ 0 & (-\lambda)^p I_n \end{bmatrix}, \qquad (Y - \lambda I_{2n})^p = \begin{bmatrix} (-\lambda)^p I_n & * \\ 0 & (BA - \lambda I_n)^p \end{bmatrix}$$

The diagonal block $(-\lambda)^p I_n$ is invertible.
Block Gaussian elimination by invertible row and column operations removes the starred block, so

$$\operatorname{rank}\big((X-\lambda I_{2n})^p\big) = n + \operatorname{rank}\big((AB-\lambda I_n)^p\big),$$
$$\operatorname{rank}\big((Y-\lambda I_{2n})^p\big) = n + \operatorname{rank}\big((BA-\lambda I_n)^p\big).$$

Since X and Y are similar, these left sides agree for all p.
Hence the shifted rank drops of AB and BA at λ agree, so Theorem 3.3.10 gives the same Jordan blocks.
Structural Commentary: Similar block embeddings make the nonzero shifted ranks match.
Key Insight: At $\lambda \neq 0$, the extra block is invertible and splits off. □

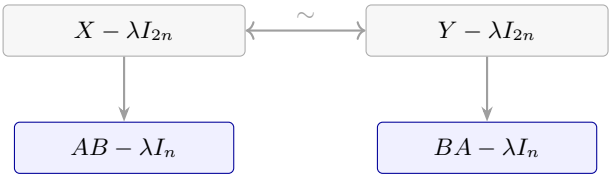


Figure 59: For $\lambda \neq 0$, the auxiliary similarity reduces the comparison to shifted ranks of AB and BA.

Corollary 3.3.14. *Uniqueness of Jordan Form*
If J and J' are Jordan matrices similar to $A \in \operatorname{Mat}(n,n,\mathbb{C})$, they differ only by a permutation of Jordan blocks.

Proof. Similarity preserves spectra and the ranks of $(X-\lambda I_n)^p$. Thus J and J' have the same Weyr characteristic for each eigenvalue. Theorem 3.3.10 then gives the same block multiplicities, so only block order may differ.
Structural Commentary: Weyr data is similarity-invariant and fixes block multiplicities.
Key Insight: Jordan form is unique modulo block order. □

Corollary 3.3.15. *Similarity Criterion for AB and BA*
Let $A, B \in \operatorname{Mat}(n,n,\mathbb{C})$. AB is similar to BA if and only if $\operatorname{rank}\big((AB)^p\big) = \operatorname{rank}\big((BA)^p\big)$, $p = 1,\dots,n$.

Proof.

(⇒) If $AB \sim BA$, then $(AB)^p \sim (BA)^p$ for all p, so their ranks are equal.
(⇐) Assume the rank equalities hold for $p = 1,\dots,n$. Theorem 3.3.13 gives the same nonzero Jordan blocks. Together with the equality at $p = 0$, the rank equalities give the same Weyr characteristic at $\lambda = 0$ by Theorem 3.3.10. Hence the nilpotent blocks also agree, so the Jordan forms coincide up to block order.

Structural Commentary: Nonzero blocks agree automatically; only the nilpotent blocks must be tested.
Key Insight: Powers of AB and BA measure exactly the zero-Jordan-block data. □

Method 3.3.16. *Checking Whether AB and BA Are Similar*
For $A, B \in \operatorname{Mat}(n,n,\mathbb{C})$:

(i) The nonzero Jordan blocks of AB and BA already match.
(ii) Compute $\operatorname{rank}\big((AB)^p\big)$ and $\operatorname{rank}\big((BA)^p\big)$ for $p = 1,\dots,n$.
(iii) If all ranks agree, then $AB \sim BA$; if some rank differs, then only the nilpotent Jordan structure differs.

Definition 3.3.17. *Nonderogatory Matrix*
A matrix $A \in \operatorname{Mat}(n,n,\mathbb{C})$ is *nonderogatory* if its Jordan form has exactly one Jordan block for each distinct eigenvalue. Equivalently, $\dim E_\lambda(A) = 1$ for every $\lambda \in \operatorname{spec}(A)$.

Lemma 3.3.18. *Commutant of One Jordan Block*
If $J = J_m(\lambda)$ and $XJ = JX$, then $X = p(J)$ for some polynomial p with $\deg p < m$.

Proof. It suffices to treat $N = J_m(0)$, since $J_m(\lambda) = \lambda I_m + N$.
The vector $\mathbf{e}_m$ is cyclic for N: $(\mathbf{e}_m, N\mathbf{e}_m, \dots, N^{m-1}\mathbf{e}_m)$ is the standard basis in reverse order.
Choose the unique polynomial p with $\deg p < m$ such that $p(N)\mathbf{e}_m = X\mathbf{e}_m$.
For every $r = 0,\dots,m-1$, $p(N)N^r\mathbf{e}_m = N^r p(N)\mathbf{e}_m = N^r X\mathbf{e}_m = XN^r\mathbf{e}_m$.
Thus $p(N)$ and X agree on a basis, so $X = p(N) = q(J)$ for $q(z) = p(z-\lambda)$.
Structural Commentary: A single Jordan block is generated by one cyclic vector.
Key Insight: The cyclic vector determines the commutant. □

Theorem 3.3.19. *Nonderogatory Commutants Are Polynomials*
If $A \in \operatorname{Mat}(n,n,\mathbb{C})$ is nonderogatory, then $\operatorname{Comm}(A) = \{p(A) : p \in \mathbb{C}[z]\}$.

Proof. The inclusion $\{p(A) : p \in \mathbb{C}[z]\} \subseteq \mathrm{Comm}(A)$ follows from Proposition 3.2.3.
For reverse inclusion, first replace A by its Jordan form $J = J_{m_1}(\lambda_1) \oplus \cdots \oplus J_{m_d}(\lambda_d)$, where λ_i are distinct.
If $XJ = JX$ and $X = [X_{ij}]$ is partitioned conformally with J, the block equation is $X_{ij}J_{m_j}(\lambda_j) = J_{m_i}(\lambda_i)X_{ij}$.
Corollary 3.2.41 gives $X_{ij} = 0$ for $i \neq j$. Each diagonal block X_{ii} commutes with $J_{m_i}(\lambda_i)$, so Lemma 3.3.18 gives $X_{ii} = p_i(J_{m_i}(\lambda_i))$ for some polynomial p_i with $\deg p_i < m_i$. The ideals generated by $(z-\lambda_i)^{m_i}$ are pairwise coprime, so Hermite interpolation gives a polynomial p with $p \equiv p_i \pmod{(z-\lambda_i)^{m_i}}$ for every i.
On the ith block, $p - p_i = (z-\lambda_i)^{m_i} q_i$ for some q_i, and $(J_{m_i}(\lambda_i) - \lambda_i I)^{m_i} = 0$. Hence $p(J_{m_i}(\lambda_i)) = p_i(J_{m_i}(\lambda_i)) = X_{ii}$ for every i, so $X = p(J)$. Conjugating back to A proves the claim.
Structural Commentary: Disjoint eigenvalue blocks force block diagonality; one Jordan block per eigenvalue makes each diagonal commutant polynomial.
Key Insight: Hermite interpolation glues the separate block polynomials into one polynomial in A. □

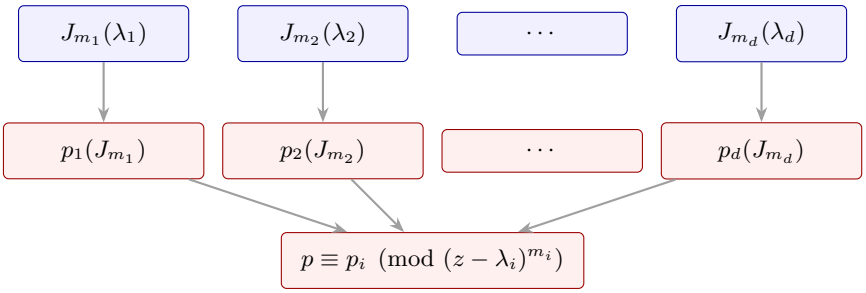


Figure 60: Blockwise commutant polynomials glue by interpolation.

Theorem 3.3.20. *Conjugate Similarity Criterion*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. Then A is similar to $\overline{A}$ if and only if each $J_k(\lambda)$ occurs in the Jordan form of A with the same multiplicity as $J_k(\overline{\lambda})$.

Proof. Let J be a Jordan form of A. Then $\overline{J}$ is a Jordan form of $\overline{A}$, obtained by replacing each $J_k(\lambda)$ by $J_k(\overline{\lambda})$.
($\Rightarrow$) If A is similar to $\overline{A}$, then J and $\overline{J}$ are similar.
By Corollary 3.3.14, their Jordan blocks have same multiplicities.
($\Leftarrow$) If the multiplicities agree, then J and $\overline{J}$ differ only by a permutation of blocks. Hence A is similar to $\overline{A}$.
Structural Commentary: Conjugation changes eigenvalue labels but not block sizes.
Key Insight: Non-real Jordan blocks must occur in conjugate pairs. □

Corollary 3.3.21. *Hermitian Products and Conjugate Products*
Let $A, H, K \in \mathrm{Mat}(n, n, \mathbb{C})$, and suppose H and K are Hermitian.
Then HK is similar to KH, and $A\overline{A}$ is similar to $\overline{A}A$.

Proof. For each $p = 1, \ldots, n$, adjoints preserves rank, $((HK)^p)^* = (KH)^p$. Hence $\mathrm{rank}\big((HK)^p\big) = \mathrm{rank}\big((KH)^p\big)$.
Similarly, entrywise conjugation preserves rank and gives $\overline{(A\overline{A})^p} = (\overline{A}A)^p$, so $\mathrm{rank}\big((A\overline{A})^p\big) = \mathrm{rank}\big((\overline{A}A)^p\big)$.
Apply Corollary 3.3.15 in both cases.
Structural Commentary: Adjoint and conjugation preserve rank and reverse product order.
Key Insight: The rank-power test turns order reversal into similarity. □

Definition 3.3.22. *Reversal Matrix*
For $k \geq 1$, *reversal matrix* $R_k \in \mathrm{Mat}(k, k, \mathbb{C})$ defined $(R_k)_{ij} = \begin{cases} 1, & i+j = k+1, \\ 0, & i+j \neq k+1 \end{cases}$, thus $R_k^t = R_k$, $R_k^2 = I_k$.

Lemma 3.3.23. *Reversal Conjugates a Jordan Block to Its Transpose*
For every $k \geq 1$ and $\lambda \in \mathbb{C}$, $J_k(\lambda)^t = R_k J_k(\lambda) R_k$. Moreover, $R_k J_k(\lambda)$ and $J_k(\lambda)R_k$ are symmetric.

Proof. Let $N = J_k(0)$. Since R_k reverses standard basis order, $R_k N R_k = N^t$.
Adding scalar gives $R_k J_k(\lambda) R_k = J_k(\lambda)^t$. As $R_k = R_k^t = R_k^{-1}$, this identity gives $J_k(\lambda)^t R_k = R_k J_k(\lambda)$ and $R_k J_k(\lambda)^t = J_k(\lambda) R_k$, so $R_k J_k(\lambda)$ and $J_k(\lambda)R_k$ are symmetric.
Structural Commentary: Reversal changes the superdiagonal shift into the subdiagonal shift.
Key Insight: Transposition reads the same Jordan chain backwards. □

Theorem 3.3.24. *Transpose Similarity and Symmetric Factorisations*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. Then:
(i) A is similar to A^t via a symmetric similarity matrix.

(ii) $A = BC = DE$, *where* B, C, D, E *are symmetric and* B, E *are invertible.*

Proof. Choose a Jordan form $A = SJS^{-1}$, and let R be the direct sum of the reversal matrices matching the Jordan blocks of J. By Lemma 3.3.23, $J^t = RJR$, $R^t = R$, $R^2 = I_n$, and both RJ and JR are symmetric.

(i) Set $P = (S^{-1})^t RS^{-1}$. Then $P^t = P$, $P^{-1} = SRS^t$, and $PAP^{-1} = (S^{-1})^t RJRS^t = (S^{-1})^t J^t S^t = A^t$.

(ii) Using $R^2 = I_n$, $A = (SRS^t)\big((S^{-1})^t RJS^{-1}\big)$ and $A = (SJRS^t)\big((S^{-1})^t RS^{-1}\big)$. The products equal SJS^{-1} because $S^t(S^{-1})^t = I_n$ and $R^2 = I_n$. All four factors are symmetric because R, RJ, JR are symmetric; SRS^t and $(S^{-1})^t RS^{-1}$ are invertible.

Structural Commentary: Reduce to Jordan blocks, reverse chains, then conjugate back.
Key Insight: Symmetric involutions produce both transpose similarity and symmetric factors. □

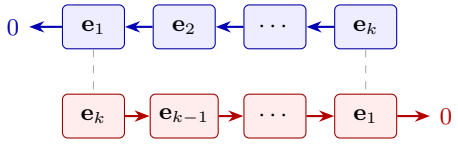


Figure 61: Reversal gives $R_k J_k(0) R_k = J_k(0)^t$: the same chain read backwards.

Theorem 3.3.25. *Conjugate Similarity and Hermitian Products*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. *The following are equivalent:*

(i) A *is similar to* $\overline{A}$.
(ii) A *is similar to* A^*.
(iii) A *is similar to* A^* *via a Hermitian similarity matrix.*
(iv) $A = HK = LM$, *where* H, K, L, M *are Hermitian, and* H, M *are invertible.*
(v) $A = HK$ *for some Hermitian matrices* H, K.

Proof.

(i) $\Rightarrow$ *(ii)* Theorem 3.3.24 applied to $\overline{A}$ gives $\overline{A} \sim (\overline{A})^t = A^*$. Thus $A \sim A^*$.

(ii) $\Rightarrow$ *(iii)* Choose invertible S such that $A = SA^*S^{-1}$, so $AS = SA^*$. For $\theta \in \mathbb{R}$, put $S_\theta = e^{i\theta}S$. Then $AS_\theta = S_\theta A^*$ and, by taking adjoints, $AS_\theta^* = S_\theta^* A^*$. Hence $A(S_\theta + S_\theta^*) = (S_\theta + S_\theta^*)A^*$. The matrix $H_\theta = S_\theta + S_\theta^*$ is Hermitian. Also $S_\theta^{-1} H_\theta = I_n + e^{-2i\theta} S^{-1}S^*$. Thus H_θ is singular exactly when $-e^{2i\theta}$ is an eigenvalue of $S^{-1}S^*$. There are only finitely many such values, so choose θ for which H_θ is invertible. Then $A = H_\theta A^* H_\theta^{-1}$.

(iii) $\Rightarrow$ *(iv)* Let R be invertible Hermitian and satisfy $A = RA^*R^{-1}$. Then $AR = RA^*$ and $R^{-1}A = A^*R^{-1}$. Thus A^*R^{-1} and RA^* are Hermitian, and $A = R(A^*R^{-1}) = (RA^*)R^{-1}$. This gives the two asserted Hermitian factorizations, with the first factor in the first product and the second factor in the second product invertible.

(iv) $\Rightarrow$ *(v)* The first factorization in (iv) already has the form $A = HK$ with H and K Hermitian.

(v) $\Rightarrow$ *(i)* If $A = HK$ with H, K Hermitian, then Corollary 3.3.21 gives $HK \sim KH$. By Theorem 3.3.24, $KH \sim (KH)^t = H^t K^t = \overline{H}\,\overline{K} = \overline{A}$. Hence $A \sim \overline{A}$.

Structural Commentary: The equivalences move between conjugation, adjoints, and Hermitian products by combining transpose similarity with a Hermitianized similarity matrix.
Key Insight: A phase-rotated average $e^{i\theta}S + e^{-i\theta}S^*$ turns an arbitrary similarity to A^* into a Hermitian one. □

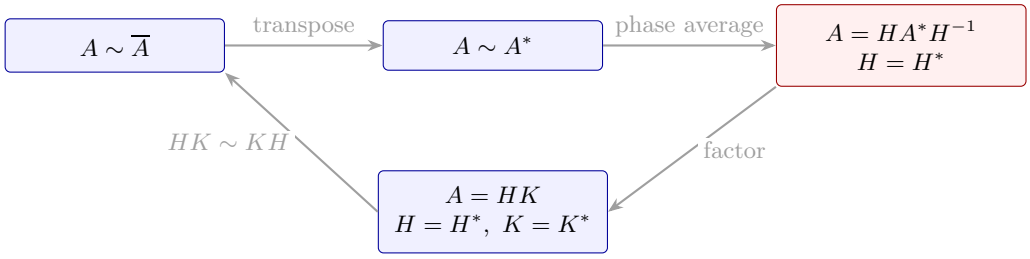


Figure 62: Conjugate similarity, adjoint similarity, and Hermitian product factorisations are equivalent.

3.4 Spectral Theory of Structured Operators

Structured spectral theory studies operators whose geometry and spectrum align. Normal, self-adjoint, positive, and unitary maps admit orthogonal decompositions and spectral resolutions.

Definition 3.4.1. *Normal Operator*
An operator $T \in \mathcal{L}(\mathcal{V})$ on a finite-dim inner-product space is *normal* if it commutes with its adjoint: $TT^* = T^*T$.
A square matrix A is *normal* if $A^*A = AA^*$.
Every self-adjoint operator, every Hermitian matrix, and every unitary matrix is normal.

Theorem 3.4.2. *Normal Operators and Equal Adjoint Norms*
An operator $T \in \mathcal{L}(\mathcal{V})$ is normal if and only if $\|T\mathbf{v}\| = \|T^*\mathbf{v}\|$ for every $\mathbf{v} \in \mathcal{V}$.

Proof. The operator $A = T^*T - TT^*$ is self-adjoint.
For every $\mathbf{v} \in \mathcal{V}$, $\langle A\mathbf{v}, \mathbf{v}\rangle = \langle T^*T\mathbf{v}, \mathbf{v}\rangle - \langle TT^*\mathbf{v}, \mathbf{v}\rangle = \|T\mathbf{v}\|^2 - \|T^*\mathbf{v}\|^2$.
Thus T is normal exactly when $A = 0$, which by Proposition 2.4.17 is equivalent to the displayed quadratic form being zero for every $\mathbf{v}$. This is equivalent to $\|T\mathbf{v}\| = \|T^*\mathbf{v}\|$ for every $\mathbf{v}$.
Structural Commentary: Normality is measured by the self-adjoint commutator $T^*T - TT^*$.
Key Insight: A normal operator and its adjoint move every vector the same distance. □

Theorem 3.4.3. *Normal Operators Share Eigenvectors with Their Adjoints*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal. If $T\mathbf{v} = \lambda\mathbf{v}$ for some $\mathbf{v} \neq \mathbf{0}$, then $T^*\mathbf{v} = \overline{\lambda}\mathbf{v}$.

Proof. The operator $T - \lambda I_{\mathcal{V}}$ is normal by Proposition 2.4.8 and scalar multiples of $I_{\mathcal{V}}$ commute with T and T^*. Applying Theorem 3.4.2 to $T - \lambda I_{\mathcal{V}}$ gives $0 = \|(T - \lambda I_{\mathcal{V}})\mathbf{v}\| = \|(T - \lambda I_{\mathcal{V}})^*\mathbf{v}\| = \|(T^* - \overline{\lambda} I_{\mathcal{V}})\mathbf{v}\|$.
Structural Commentary: For normal operators, kernels of $T - \lambda I$ and $T^* - \overline{\lambda} I$ coincide.
Key Insight: Normality makes eigenvector information visible from both T and T^*. □

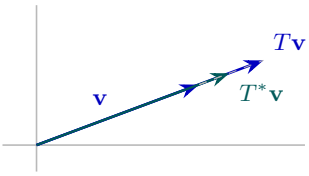


Figure 63: For normal T, $T\mathbf{v} = \lambda\mathbf{v}$ implies $T^*\mathbf{v} = \overline{\lambda}\mathbf{v}$.

Theorem 3.4.4. *Orthogonal Eigenvectors for Normal Operators*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal. Eigenvectors of T corresponding to distinct eigenvalues are orthogonal.

Proof. Suppose $T\mathbf{u} = \alpha\mathbf{u}$ and $T\mathbf{v} = \beta\mathbf{v}$ with $\alpha \neq \beta$. By Theorem 3.4.3, $T^*\mathbf{v} = \overline{\beta}\mathbf{v}$.
Hence $\alpha\langle \mathbf{u}, \mathbf{v}\rangle = \langle T\mathbf{u}, \mathbf{v}\rangle = \langle \mathbf{u}, T^*\mathbf{v}\rangle = \langle \mathbf{u}, \overline{\beta}\mathbf{v}\rangle = \beta\langle \mathbf{u}, \mathbf{v}\rangle$. Since $\alpha - \beta \neq 0$, $\langle \mathbf{u}, \mathbf{v}\rangle = 0$.
Structural Commentary: Normality lets the operator move from one eigenvector to the other through the adjoint.
Key Insight: Distinct eigenvalue scalings are incompatible unless the inner product is zero. □

Theorem 3.4.5. *Eigenvalues of Self-Adjoint Operators Are Real*
Every eigenvalue of a self-adjoint operator on a finite-dimensional inner-product space is real.

Proof. Let $T = T^*$ and $T\mathbf{v} = \lambda\mathbf{v}$ with $\mathbf{v} \neq \mathbf{0}$. Then $\lambda\|\mathbf{v}\|^2 = \langle \lambda\mathbf{v}, \mathbf{v}\rangle = \langle T\mathbf{v}, \mathbf{v}\rangle = \langle \mathbf{v}, T\mathbf{v}\rangle = \langle \mathbf{v}, \lambda\mathbf{v}\rangle = \overline{\lambda}\|\mathbf{v}\|^2$.
Since $\|\mathbf{v}\|^2 > 0$, $\lambda = \overline{\lambda}$, so $\lambda \in \mathbb{R}$.
Structural Commentary: Self-adjointness moves T between inner-product slots without changing the value.
Key Insight: An eigenvalue must equal its conjugate because the eigenvector has positive squared norm. □

Proposition 3.4.6. *Invertible Quadratic Expressions*
Let $T \in \mathcal{L}(\mathcal{V})$ be self-adjoint on a finite-dimensional inner-product space.
If $b, c \in \mathbb{R}$ and $b^2 < 4c$, then $T^2 + bT + cI_{\mathcal{V}}$ is invertible.

Proof. Let $S = T^2 + bT + cI_{\mathcal{V}}$. For $\mathbf{v} \neq \mathbf{0}$,

$$\begin{aligned}\langle S\mathbf{v}, \mathbf{v}\rangle &= \|T\mathbf{v}\|^2 + b\langle T\mathbf{v}, \mathbf{v}\rangle + c\|\mathbf{v}\|^2 \geq \|T\mathbf{v}\|^2 - |b|\,\|T\mathbf{v}\|\,\|\mathbf{v}\| + c\|\mathbf{v}\|^2 \\ &= \left(\|T\mathbf{v}\| - \frac{|b|}{2}\|\mathbf{v}\|\right)^2 + \left(c - \frac{b^2}{4}\right)\|\mathbf{v}\|^2 > 0\end{aligned}$$

Thus $S\mathbf{v} \neq \mathbf{0}$ for every nonzero $\mathbf{v}$, so $\ker(S) = \{\mathbf{0}\}$. In finite dimensions, injectivity implies invertibility.
Structural Commentary: Negative discriminant keeps this self-adjoint quadratic positive.
Key Insight: Self-adjointness converts $\langle T^2\mathbf{v}, \mathbf{v}\rangle$ into $\|T\mathbf{v}\|^2$, so completing the square controls the kernel. □

Lemma 3.4.7. *Self-Adjoint Operators Have Eigenvalues*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be a finite-dimensional real inner-product space. Every self-adjoint $T \in \mathcal{L}(\mathcal{V})$ has an eigenvalue.

Proof. Choose $\mathbf{v} \neq \mathbf{0}$ and set $n = \dim(\mathcal{V})$.
The list $\mathbf{v}, T\mathbf{v}, \dots, T^n\mathbf{v}$ is linearly dependent, so some nonzero real polynomial p satisfies $p(T)\mathbf{v} = \mathbf{0}$.
Factor p over $\mathbb{R}$ as $p(x) = a \prod_{j=1}^{M} (x^2 + b_j x + c_j) \prod_{k=1}^{m} (x - \lambda_k)$, where $a \neq 0$, $b_j^2 < 4c_j$, $\lambda_k \in \mathbb{R}$.
By Proposition 3.4.6, each quadratic factor $Q_j(T) := T^2 + b_j T + c_j I_{\mathcal{V}}$ is invertible. Let $R(T) = \prod_j Q_j(T)$ and $L(T) = \prod_k (T - \lambda_k I_{\mathcal{V}})$. The factors commute because they are polynomials in T, and $R(T)$ is invertible.
If $m = 0$, then $p(T) = aR(T)$ would be invertible, contradicting $p(T)\mathbf{v} = \mathbf{0}$ with $\mathbf{v} \neq \mathbf{0}$. Hence $m \geq 1$.
Since $\mathbf{0} = p(T)\mathbf{v} = aR(T)L(T)\mathbf{v}$ and $aR(T)$ is invertible, $L(T)\mathbf{v} = \mathbf{0}$. Thus $L(T)$ has nontrivial kernel.
If every $T - \lambda_k I_{\mathcal{V}}$ were injective, their product $L(T)$ would be injective, a contradiction. Therefore some $T - \lambda_k I_{\mathcal{V}}$ is not injective, and λ_k is an eigenvalue of T.
Structural Commentary: An annihilating polynomial remains; quadratic factors are invertible.
Key Insight: The nontrivial kernel must come from a real linear factor, hence from an eigenvalue. □

Proposition 3.4.8. *Self-Adjoint Operators and Invariant Subspaces*
Let $T \in \mathcal{L}(\mathcal{V})$ be self-adjoint, and let $\mathcal{U}$ be a subspace of $\mathcal{V}$ invariant under T. Then:

i. $\mathcal{U}^\perp$ is invariant under T;
ii. $T|_{\mathcal{U}} \in \mathcal{L}(\mathcal{U})$ is self-adjoint;
iii. $T|_{\mathcal{U}^\perp} \in \mathcal{L}(\mathcal{U}^\perp)$ is self-adjoint.

Proof.

(i) If $\mathbf{v} \in \mathcal{U}^\perp$ and $\mathbf{u} \in \mathcal{U}$, then $\langle T\mathbf{v}, \mathbf{u}\rangle = \langle \mathbf{v}, T\mathbf{u}\rangle = 0$ because $T\mathbf{u} \in \mathcal{U}$. Hence $T\mathbf{v} \in \mathcal{U}^\perp$.
(ii) For $\mathbf{u}, \mathbf{v} \in \mathcal{U}$, $\langle (T|_{\mathcal{U}})\mathbf{u}, \mathbf{v}\rangle = \langle T\mathbf{u}, \mathbf{v}\rangle = \langle \mathbf{u}, T\mathbf{v}\rangle = \langle \mathbf{u}, (T|_{\mathcal{U}})\mathbf{v}\rangle$. Thus $T|_{\mathcal{U}}$ is self-adjoint.
(iii) By (i), $\mathcal{U}^\perp$ is invariant, so (ii)'s argument applies to $T|_{\mathcal{U}^\perp}$.

Structural Commentary: Self-adjointness preserves orthogonal complements of invariant subspaces.
Key Insight: Move T across the inner product; invariance of $\mathcal{U}$ then forces orthogonality to persist. □

Theorem 3.4.9. *Schur Triangularisation*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $T \in \mathcal{L}(\mathcal{V})$. Then T has an upper-triangular matrix with respect to some orthonormal basis of $\mathcal{V}$.

Proof. Proceed by induction on $n = \dim(\mathcal{V})$. The cases $n = 0$ and $n = 1$ are immediate.
Suppose $n \geq 2$ and the result holds for smaller dimensions. By Theorem 3.1.23, choose a unit eigenvector $\mathbf{e}_1$ of T, say $T\mathbf{e}_1 = \lambda \mathbf{e}_1$. Let $\mathcal{U} = \operatorname{span}(\mathbf{e}_1)$ and define $S \in \mathcal{L}(\mathcal{U}^\perp)$ by $S\mathbf{w} = P_{\mathcal{U}^\perp}(T\mathbf{w})$.
By induction, choose an orthonormal basis $(\mathbf{e}_2, \dots, \mathbf{e}_n)$ of $\mathcal{U}^\perp$ with respect to which S has an upper-triangular matrix. Then $(\mathbf{e}_1, \dots, \mathbf{e}_n)$ is an orthonormal basis of $\mathcal{V}$.
For $j \geq 2$, the coordinates of $T\mathbf{e}_j$ along $\mathbf{e}_2, \dots, \mathbf{e}_n$ are exactly the coordinates of $S\mathbf{e}_j$, hence have no entries below row j. Also $T\mathbf{e}_1 = \lambda \mathbf{e}_1$. Therefore $[T]_{\gamma \leftarrow \gamma}$ is upper triangular for $\gamma = (\mathbf{e}_1, \dots, \mathbf{e}_n)$.
Structural Commentary: Schur triangularisation builds an eigenline first, then triangularises the projected action on its orthogonal complement.
Key Insight: The lower-right block is controlled by the compression $P_{\mathcal{U}^\perp} T|_{\mathcal{U}^\perp}$. □

Theorem 3.4.10. *Matrix Schur Triangularisation with Ordered Diagonal*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$ have eigenvalues $\lambda_1, \dots, \lambda_n$ in a prescribed order. If $A\mathbf{x} = \lambda_1 \mathbf{x}$ and $\|\mathbf{x}\|_2 = 1$, then there is a unitary $U = [\mathbf{x}\ U_2]$ such that $U^* A U$ is upper triangular with diagonal $\lambda_1, \dots, \lambda_n$.
If A is real, all $\lambda_i \in \mathbb{R}$, and $\mathbf{x} \in \mathbb{R}^n$, then U may be chosen real orthogonal.

Proof. Induct on n. The case $n = 1$ is immediate. Extend $\mathbf{x}$ to an orthonormal basis and set $V = [\mathbf{x}\ V_2]$.
Since $A\mathbf{x} = \lambda_1 \mathbf{x}$, the first column of $V^* A V$ is $\lambda_1 \mathbf{e}_1$, so $V^* A V = \begin{bmatrix} \lambda_1 & * \\ 0 & A' \end{bmatrix}$.
Block triangularity gives $p_{V^*AV}(z) = (z - \lambda_1) p_{A'}(z)$, so A' has eigenvalues $\lambda_2, \dots, \lambda_n$, counted with algebraic multiplicity. By induction, choose unitary W with $W^* A' W = T'$ upper triangular and diagonal $\lambda_2, \dots, \lambda_n$.
Then $U := V(I_1 \oplus W)$ has first column $\mathbf{x}$ and $U^* A U = \begin{bmatrix} \lambda_1 & * \\ 0 & T' \end{bmatrix}$.
The real case uses real orthonormal extensions; the block A' is real, and the same characteristic-polynomial factorization leaves the real eigenvalue list $\lambda_2, \dots, \lambda_n$ for the induction step.
Structural Commentary: Isolate the chosen eigenline, then repeat on the lower-right block.
Key Insight: The diagonal order is controlled by which eigenline is peeled off first. □

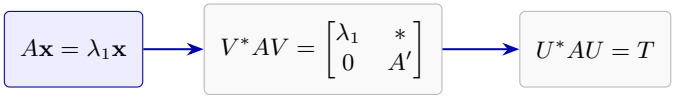


Figure 64: Schur induction: isolate an eigenline, then triangularise the compressed block.

Corollary 3.4.11. *Eigenvalues of the Inverse*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ be invertible with eigenvalues $\lambda_1, \ldots, \lambda_n$, counted with algebraic multiplicity. Then A^{-1} has eigenvalues $\lambda_1^{-1}, \ldots, \lambda_n^{-1}$, counted with algebraic multiplicity.

Proof. Take a Schur form $A = UTU^*$ with diagonal $\lambda_1, \ldots, \lambda_n$.
Then $A^{-1} = UT^{-1}U^*$, and T^{-1} is upper triangular with diagonal $\lambda_1^{-1}, \ldots, \lambda_n^{-1}$.
Structural Commentary: Invert the Schur form and read the diagonal.
Key Insight: The inverse spectrum is the reciprocal spectrum. □

Theorem 3.4.12. *Complex Spectral Theorem*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $T \in \mathcal{L}(\mathcal{V})$. The following are equivalent:

i. T is normal;
ii. $\mathcal{V}$ has an orthonormal basis consisting of eigenvectors of T;
iii. T has a diagonal matrix with respect to some orthonormal basis of $\mathcal{V}$.

Proof.

(i) $\Rightarrow$ *(ii)* By Schur triangularisation (Theorem 3.4.9), choose an orthonormal basis $\gamma = (\mathbf{e}_1, \ldots, \mathbf{e}_n)$ in which $A = [T]_{\gamma \leftarrow \gamma} = [a_{ij}]$ is upper triangular. Show by induction on k that the kth row has no entries to the right of the diagonal. Assume this has been proved for rows $1, \ldots, k-1$.
Upper triangularity gives $T\mathbf{e}_k = \sum_{i \le k} a_{ik}\mathbf{e}_i$ and $T^*\mathbf{e}_k = \sum_{j \ge k} \overline{a_{kj}}\mathbf{e}_j$.
The induction hypothesis gives $a_{ik} = 0$ for $i < k$, so $\|T\mathbf{e}_k\|^2 = |a_{kk}|^2$.
Since T is normal, Theorem 3.4.2 gives $|a_{kk}|^2 = \|T\mathbf{e}_k\|^2 = \|T^*\mathbf{e}_k\|^2 = |a_{kk}|^2 + \sum_{j>k} |a_{kj}|^2$.
Thus each a_{kj} with $j > k$ is zero. Induction makes A diagonal, so γ is an orthonormal basis of eigenvectors.

(ii) $\Rightarrow$ *(iii)* In an orthonormal eigenbasis, the matrix of T is diagonal.

(iii) $\Rightarrow$ *(i)* If $[T]$ is diagonal in an orthonormal basis, then $[T^*] = [T]^*$ is diagonal as well.
Diagonal matrices commute, so $TT^* = T^*T$; hence T is normal.

Structural Commentary: Schur gives triangular form; normality kills the upper entries.
Key Insight: Equal norms of $T\mathbf{e}_j$ and $T^*\mathbf{e}_j$ force each row tail to vanish. □

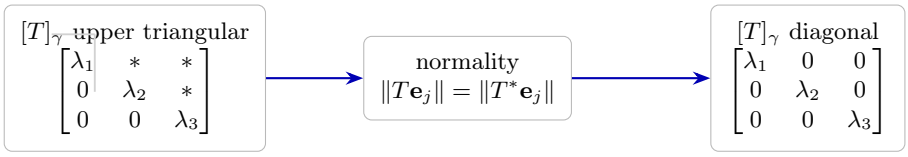


Figure 65: Normality turns Schur triangular form diagonal by forcing row tails to vanish.

Theorem 3.4.13. *Real Spectral Theorem*
Let $\mathcal{V}$ be a finite-dimensional real inner-product space and let $T \in \mathcal{L}(\mathcal{V})$. The following are equivalent:

i. T is self-adjoint;
ii. $\mathcal{V}$ has an orthonormal basis consisting of eigenvectors of T;
iii. T has a diagonal matrix with respect to some orthonormal basis of $\mathcal{V}$.

Proof.

(i) $\Rightarrow$ *(ii)* Assume $T = T^*$ and induct on $n = \dim(\mathcal{V})$. For $n = 0$ there is nothing to prove, and for $n = 1$ every nonzero vector is an eigenvector after normalization.
Let $n \ge 2$. Lemma 3.4.7 gives a unit eigenvector $\mathbf{u}$, and set $\mathcal{U} = \mathrm{span}(\mathbf{u})$. The subspace $\mathcal{U}$ is T-invariant.
By Proposition 3.4.8, $\mathcal{U}^\perp$ is T-invariant and the restriction $T|_{\mathcal{U}^\perp}$ is self-adjoint.
Since $\dim(\mathcal{U}^\perp) = n - 1$, the induction hypothesis gives an orthonormal eigenbasis of $\mathcal{U}^\perp$.
Adjoining $\mathbf{u}$ gives an orthonormal eigenbasis of $\mathcal{V}$ by Theorem 2.2.4.

(ii) $\Rightarrow$ *(iii)* In an orthonormal eigenbasis, the matrix of T is diagonal.

(iii) $\Rightarrow$ *(i)* In an orthonormal basis where $[T]$ is real diagonal, $[T^*] = [T]^t = [T]$. Hence $T = T^*$.

Structural Commentary: Split off one eigendirection, then restrict to its orthogonal complement.
Key Insight: One eigenvector reduces the dimension. □

Perspective 3.4.14. *Normal Means Diagonal in Orthogonal Coordinates*
Complex spectral theorem states that normal operators are precisely the operators whose full structure is visible in an orthonormal eigenbasis. This is why normality is the natural coordinate-free version of being diagonal.

Perspective 3.4.15. *Real Normal Operators Use Planes*
Over $\mathbb{R}$, a normal operator need not be orthogonally diagonalizable. Its real spectral form is an orthogonal sum of invariant lines and planes, with each plane represented by $\begin{bmatrix} a & -b \\ b & a \end{bmatrix}$. See Theorem 4.5.11.

Proposition 3.4.16. *Basic Spectral Properties of Normal Operators*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space, let $T \in \mathcal{L}(\mathcal{V})$ be normal, and let $p \in \mathbb{C}[z]$.

(i) $p(T)$ is normal and commutes with T.
(ii) T^* is normal and $E_\lambda(T) = E_{\overline{\lambda}}(T^*)$ for every $\lambda \in \mathbb{C}$.
(iii) $\ker T = \ker T^*$ and $\operatorname{range}(T) = \operatorname{range}(T^*)$.

Proof. Choose an orthonormal eigenbasis $\mathbf{e}_1, \dots, \mathbf{e}_n$ with $T\mathbf{e}_j = \mu_j \mathbf{e}_j$.

(i) $p(T)\mathbf{e}_j = p(\mu_j)\mathbf{e}_j$, $p(T)$ is diagonal in this basis, hence normal. It commutes with T by Proposition 3.2.3.
(ii) Since $T^*\mathbf{e}_j = \overline{\mu_j}\mathbf{e}_j$, T^* is diagonal in the same basis, hence normal. Also $T\mathbf{v} = \lambda\mathbf{v}$ iff $T^*\mathbf{v} = \overline{\lambda}\mathbf{v}$.
(iii) Part (ii) with $\lambda = 0$ gives $\ker T = \ker T^*$. Both ranges are $\operatorname{span}\{\mathbf{e}_j : \mu_j \neq 0\}$.

Structural Commentary: Normality gives one orthonormal basis that diagonalises both T and T^*.
Key Insight: Kernel and range are read from the nonzero diagonal entries in that basis. □

Corollary 3.4.17. *Normal Operators Have No Nilpotent Primary Part*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $T \in \mathcal{L}(\mathcal{V})$ be normal.

(i) If $S = p(T)$ for some $p \in \mathbb{C}[z]$, then $\ker(S^k) = \ker(S)$ for every $k \geq 1$.
(ii) The minimal polynomial of T is square-free, where $\mu_T(z) = \prod_{\lambda \in \operatorname{spec}(T)} (z - \lambda)$.

Proof. Let $\mathbf{e}_1, \dots, \mathbf{e}_n$ be an orthonormal eigenbasis with $T\mathbf{e}_j = \lambda_j \mathbf{e}_j$.

(i) For $S = p(T)$, $k \geq 1$, $S^k\mathbf{e}_j = p(\lambda_j)^k \mathbf{e}_j$; $\ker(S^k) = \operatorname{span}\{\mathbf{e}_j : p(\lambda_j)^k = 0\} = \operatorname{span}\{\mathbf{e}_j : p(\lambda_j) = 0\} = \ker(S)$.
(ii) Let $m(z) = \prod_{\lambda \in \operatorname{spec}(T)}(z - \lambda)$. Since $m(T)\mathbf{e}_j = 0$ for all j, $m(T) = 0$, so $\mu_T \mid m$. Conversely, every $\lambda \in \operatorname{spec}(T)$ is a root of μ_T by Theorem 3.2.33, hence $m \mid \mu_T$. Since both polynomials are monic, $\mu_T = m$.

Structural Commentary: Normality makes every primary component an eigenspace.
Key Insight: A diagonal operator has no hidden nilpotent part. □

Theorem 3.4.18. *Spectral Resolution of a Normal Operator*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $\mu_1, \dots, \mu_d$ be distinct scalars. For $T \in \mathcal{L}(\mathcal{V})$, consider:

(i) T is normal and $\operatorname{spec}(T) \subseteq \{\mu_1, \dots, \mu_d\}$.
(ii) $\exists$ pairwise orthogonal projections $P_1, \dots, P_d \in \mathcal{L}(\mathcal{V})$ such that $P_i P_j = 0$ for $i \neq j$, $P_1 + \cdots + P_d = I_\mathcal{V}$, and $T = \mu_1 P_1 + \cdots + \mu_d P_d$
(iii) Any projections satisfying (ii) obey $\operatorname{range}(P_j) = E_{\mu_j}(T)$ whenever $\mu_j \in \operatorname{spec}(T)$.

Then (i) and (ii) are equivalent, and any projections satisfying (ii) satisfy (iii).

Proof.

(i) $\Rightarrow$ (ii). By Theorem 3.4.12, $\mathcal{V} = \bigoplus_{\lambda \in \operatorname{spec}(T)} E_\lambda(T)$ as an orthogonal direct sum.

Let P_j be the orthogonal projection onto $E_{\mu_j}(T)$, and set $P_j = 0$ when $\mu_j \notin \operatorname{spec}(T)$. Since $\operatorname{spec}(T) \subseteq \{\mu_1, \dots, \mu_d\}$, these projections are pairwise orthogonal, sum to $I_\mathcal{V}$, and satisfy $T = \sum_j \mu_j P_j$.

(ii) $\Rightarrow$ (i). Since $P_j^* = P_j$ and $P_i P_j = 0$ for $i \neq j$, $T^* = \sum_{j=1}^{d} \overline{\mu_j} P_j$ and $T^*T = \sum_{j=1}^{d} |\mu_j|^2 P_j = TT^*$.

Thus T is normal. If $\lambda \in \operatorname{spec}(T)$ and $\mathbf{0} \neq \mathbf{x} \in E_\lambda(T)$, then $0 = (T - \lambda I_\mathcal{V})\mathbf{x} = \sum_{j=1}^{d} (\mu_j - \lambda) P_j \mathbf{x}$.

The summands lie in the pairwise orthogonal ranges $\operatorname{range}(P_j)$, hence $0 = \sum_{j=1}^{d} |\mu_j - \lambda|^2 \|P_j \mathbf{x}\|^2$.

Because $\sum_j P_j = I_\mathcal{V}$, at least one $P_k \mathbf{x}$ is nonzero, so $\mu_k = \lambda$. Thus $\operatorname{spec}(T) \subseteq \{\mu_1, \dots, \mu_d\}$.

(*ii*) $\Rightarrow$ (*iii*). For $\mathbf{x} \in \mathcal{V}$, $\mathbf{x} = \sum_{i=1}^{d} P_i\mathbf{x}$ and $(T - \mu_j I_\mathcal{V})\mathbf{x} = \sum_{i \neq j} (\mu_i - \mu_j) P_i\mathbf{x}$.

The summands are orthogonal, so $\|(T - \mu_j I_\mathcal{V})\mathbf{x}\|^2 = \sum_{i \neq j} |\mu_i - \mu_j|^2 \|P_i\mathbf{x}\|^2$, which vanishes exactly when $P_i\mathbf{x} = 0$ for all $i \neq j$, equivalently $\mathbf{x} = P_j\mathbf{x}$. Thus $E_{\mu_j}(T) = \operatorname{range}(P_j)$.

Structural Commentary: A normal operator decomposes into scalar actions on orthogonal eigenspaces.
Key Insight: Spectral projections isolate eigenspaces; the operator is the weighted sum of those projections. □



Figure 66: Spectral resolution: $\mathcal{V} = \bigoplus_j \operatorname{range}(P_j)$, $P_iP_j = 0$ for $i \neq j$, $\sum_j P_j = I_\mathcal{V}$, and $T = \sum_j \mu_j P_j$.

Perspective 3.4.19. *Spectral Resolutions Are Orthogonal Resolutions*
For the projections P_j in Theorem 3.4.18, $P_j = P_j^* = P_j^2$; $P_iP_j = 0$ $(i \neq j)$; $\sum_j P_j = I_\mathcal{V}$.
Hence $\mathcal{V} = \bigoplus_j^\perp \operatorname{range}(P_j)$ is the orthogonal decomposition into T-invariant spectral subspaces.
Structural Commentary: Each P_j is forced by its eigenspace: it is the orthogonal projection onto it.
Key Insight: Normality turns diagonalization into a projection-valued decomposition.

Corollary 3.4.20. *Functional Calculus from Spectral Resolution*
Let $T = \sum_{j=1}^d \mu_j P_j$ be the spectral resolution of a normal operator and let $f \in \mathbb{C}[z]$. Then $f(T) = \sum_{j=1}^{d} f(\mu_j) P_j$.
In particular, for each j there is a polynomial p_j such that $p_j(T) = P_j$.

Proof. Since $P_iP_j = 0$ for $i \neq j$ and $P_j^2 = P_j$, induction gives $T^m = \sum_j \mu_j^m P_j$ for every $m \geq 0$.
Linearity then gives the displayed formula for every polynomial f.
For the final assertion, choose the Lagrange polynomial p_j satisfying $p_j(\mu_i) = \delta_{ij}$. Then $p_j(T) = P_j$.
Structural Commentary: Orthogonal spectral projections multiply like coordinate idempotents.
Key Insight: Polynomial functions of a normal operator act by scalar functions on each eigenspace. □

Lemma 3.4.21. *Adjoints of Normal Matrices as One Polynomial*
Let $A_1, \ldots, A_k$ be normal complex matrices. Then $\exists$ a polynomial $p \in \mathbb{C}[z]$ s.t. $p(A_j) = A_j^*$ $\forall j = 1, \ldots, k$.

Proof. Let $\Omega = \bigcup_{j=1}^{k} \operatorname{spec}(A_j)$. Since Ω is finite, choose a Lagrange interpolation polynomial p satisfying $p(\lambda) = \overline{\lambda}$ for every $\lambda \in \Omega$. For each A_j, Corollary 3.4.20 gives $p(A_j) = A_j^*$.
Structural Commentary: One polynomial interpolates conjugation on all spectra.
Key Insight: For normal matrices, A^* is $p(A)$ with $p(\lambda) = \overline{\lambda}$. □

Theorem 3.4.22. *Fuglede–Putnam Theorem*
Let $A \in \operatorname{Mat}(m, m, \mathbb{C})$ and $B \in \operatorname{Mat}(n, n, \mathbb{C})$ be normal, and let $X \in \operatorname{Mat}(m, n, \mathbb{C})$.
Then $AX = XB \iff A^*X = XB^*$.

Proof.
($\Rightarrow$) Assume $AX = XB$. Then $A^rX = XB^r$ for every $r \geq 0$, so $p(A)X = Xp(B)$ for every polynomial p. By Lemma 3.4.21, choose p with $p(A) = A^*$ and $p(B) = B^*$. Thus $A^*X = XB^*$.
($\Leftarrow$) Assume $A^*X = XB^*$. Apply the same argument to the normal matrices A^* and B^*, choosing q with $q(A^*) = A$ and $q(B^*) = B$. Hence $AX = XB$.

Structural Commentary: An intertwining relation passes through every polynomial in the two normal matrices.
Key Insight: Normality makes adjoints polynomial, so intertwining transfers to adjoints. □

Method 3.4.23. *Finding a Spectral Decomposition of a Normal Matrix*
For a normal matrix $A \in \operatorname{Mat}(n, n, \mathbb{C})$:

i. Find the distinct eigenvalues $\mu_1, \ldots, \mu_d$.
ii. For each j, solve $(A - \mu_j I_n)\mathbf{x} = 0$ to get $E_{\mu_j}(A)$.
iii. Orthonormalise a basis of each eigenspace.
iv. Concatenate these orthonormal bases as the columns of U; then U is unitary and $A = U\Lambda U^*$, where Λ repeats each μ_j according to $\dim E_{\mu_j}(A)$.

Theorem 3.4.24. *Spectral Classification of Normal Operators*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $T \in \mathcal{L}(\mathcal{V})$ be normal.

(i) T is unitary if and only if $|\lambda| = 1$ for every $\lambda \in \operatorname{spec}(T)$.
(ii) T is self-adjoint if and only if $\operatorname{spec}(T) \subseteq \mathbb{R}$.
(iii) T is an orthogonal projection if and only if $\operatorname{spec}(T) \subseteq \{0, 1\}$.
(iv) T is skew-adjoint, meaning $T^* = -T$, if and only if every $\lambda \in \operatorname{spec}(T)$ is purely imaginary.

Proof. Write the spectral resolution $T = \sum_j \mu_j P_j$, where the μ_j are the distinct eigenvalues and $P_j \neq 0$.
Then $I_{\mathcal{V}} = \sum_j P_j$, $T^* = \sum_j \overline{\mu_j} P_j$, $T^*T = \sum_j |\mu_j|^2 P_j$, and $T^2 = \sum_j \mu_j^2 P_j$.
Since ranges of P_j are nonzero and independent, an identity $\sum_j \alpha_j P_j = 0$ holds only when every $\alpha_j = 0$.
(i) T is unitary iff $T^*T = I_{\mathcal{V}} = \sum_j P_j$, iff $|\mu_j| = 1$ for every j.
(ii) T is self-adjoint iff $T = T^*$, iff $\mu_j = \overline{\mu_j}$ for every j.
(iii) T is an orthogonal projection iff $T = T^*$ and $T^2 = T$, iff $\mu_j = \overline{\mu_j}$ and $\mu_j^2 = \mu_j$ $\forall j$, iff $\mu_j \in \{0, 1\}$ $\forall j$.
(iv) $T^* = -T$ iff $\overline{\mu_j} = -\mu_j$ for every j, iff every μ_j is purely imaginary.
Structural Commentary: For normal operators, structural classes are scalar conditions on spectral blocks.
Key Insight: For normal T, these properties reduce to scalar tests. □

Proposition 3.4.25. *Cartesian Decomposition and Normality*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space.
For every $T \in \mathcal{L}(\mathcal{V})$, there are unique self-adjoint operators $H, K \in \mathcal{L}(\mathcal{V})$ such that $T = H + iK$.
They are $H = \frac{1}{2}(T + T^*)$ and $K = \frac{1}{2i}(T - T^*)$. Moreover, T is normal if and only if $HK = KH$.

Proof. The displayed formulas give self-adjoint H, K and $T = H + iK$. If $T = X + iY$ with X, Y self-adjoint, then $T^* = X - iY$, so $X = \frac{1}{2}(T + T^*)$ and $Y = \frac{1}{2i}(T - T^*)$; uniqueness follows.
For $T = H + iK$, $T^*T = (H - iK)(H + iK) = H^2 + K^2 + i(HK - KH)$,
whereas $TT^* = (H + iK)(H - iK) = H^2 + K^2 - i(HK - KH)$. Thus $T^*T = TT^*$ if and only if $HK = KH$.
Structural Commentary: Every complex operator splits into Hermitian and skew-Hermitian parts.
Key Insight: Normality says the two Hermitian coordinates of T commute. □

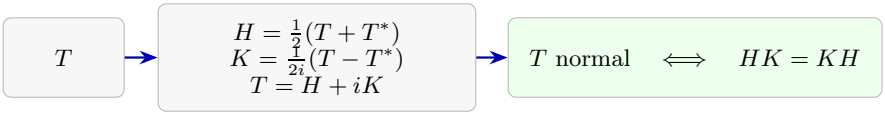


Figure 67: Cartesian decomposition: T is normal exactly when $HK = KH$.

Lemma 3.4.26. *Upper-Triangular Normal Matrices Are Diagonal*
If $A \in \operatorname{Mat}(n, n, \mathbb{C})$ is upper triangular and normal, then A is diagonal.

Proof. Induct on n. The case $n = 1$ is immediate.
Since A is upper triangular, $A\mathbf{e}_1 = a_{11}\mathbf{e}_1$, while $A^*\mathbf{e}_1 = \overline{a_{11}}\mathbf{e}_1 + \overline{a_{12}}\mathbf{e}_2 + \cdots + \overline{a_{1n}}\mathbf{e}_n$.
Normality gives $\|A\mathbf{e}_1\|_2 = \|A^*\mathbf{e}_1\|_2$, hence $a_{12} = \cdots = a_{1n} = 0$. Thus $A = a_{11} \oplus B$ with B upper triangular.
The block diagonal normality equation gives $B^*B = BB^*$, so B is normal. By induction, B is diagonal.
Structural Commentary: Normality forbids upper-triangular row tails.
Key Insight: Compare the lengths of $A\mathbf{e}_1$ and $A^*\mathbf{e}_1$ to kill the first row, then induct. □

Proposition 3.4.27. *Normal Block Upper-Triangular Matrices Collapse*
Let $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$. Let $A = \begin{bmatrix} B & X \\ 0 & C \end{bmatrix} \in \operatorname{Mat}(n, n, \mathbb{F})$, where B and C are square.
Then A is normal if and only if $X = 0$ and both B and C are normal.

Proof.
(⇒) Suppose A is normal. Comparing the $(1,1)$ blocks of A^*A and AA^* gives $B^*B = BB^* + XX^*$.
Taking traces and using cyclicity, $\operatorname{tr}(XX^*) = \operatorname{tr}(B^*B) - \operatorname{tr}(BB^*) = 0$. Since $\operatorname{tr}(XX^*) = \|X\|_F^2$, we get $X = 0$. Thus $A = B \oplus C$, and $A^*A = AA^*$ reduces blockwise to $B^*B = BB^*$ and $C^*C = CC^*$.
(⇐) If $X = 0$ and B, C are normal, then $A = B \oplus C$, so blockwise multiplication gives $A^*A = AA^*$.
Structural Commentary: A normal operator cannot have one-way coupling between orthogonal blocks.
Key Insight: The off-diagonal block contributes the nonnegative trace $\|X\|_F^2$, so normality forces it to vanish. □

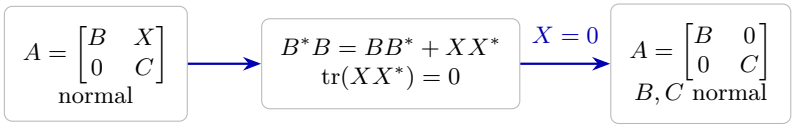


Figure 68: Normality removes the upper-right block in a block upper-triangular matrix.

Theorem 3.4.28. *Invariant Subspaces of Normal Matrices Reduce*
Let $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ be normal, and let $\mathcal{U} \subseteq \mathbb{F}^n$ be invariant under A. Then $\mathcal{U}$ is invariant under A^*. If P is the orthogonal projection onto $\mathcal{U}$, then $AP = PA$.

Proof. The cases $\mathcal{U} = \{0\}$ and $\mathcal{U} = \mathbb{F}^n$ are immediate. Otherwise choose a unitary $Q = [Q_1\ Q_2]$ whose columns form orthonormal bases of $\mathcal{U}$ and $\mathcal{U}^\perp$. Invariance of $\mathcal{U}$ says AQ_1 has columns in $\mathcal{U} = \mathrm{col}(Q_1)$, hence $Q_2^* A Q_1 = 0$. Therefore $Q^*AQ = \begin{bmatrix} B & X \\ 0 & C \end{bmatrix}$. This matrix is normal because it is unitarily similar to A.
Proposition 3.4.27 gives $X = 0$, so $Q^*AQ = B \oplus C$. Hence $Q^*A^*Q = B^* \oplus C^*$, which shows $\mathcal{U}$ is A^*-invariant. Also $Q^*PQ = I_{\dim(\mathcal{U})} \oplus 0$, so Q^*AQ commutes with Q^*PQ; conjugating back gives $AP = PA$.
Structural Commentary: In adapted coordinates, normal invariance becomes block diagonality.
Key Insight: Normal invariant subspaces are automatically reducing subspaces. □

Corollary 3.4.29. *Invariant Subspaces of Normal Operators Reduce*
Let $\mathcal{V}$ be a finite-dimensional inner-product space, let $T \in \mathcal{L}(\mathcal{V})$ be normal, let $\mathcal{U}$ be invariant under T. Then:

(i) $\mathcal{U}^\perp$ is invariant under T;
(ii) $\mathcal{U}$ is invariant under T^*;
(iii) $(T|_{\mathcal{U}})^* = (T^*)|_{\mathcal{U}}$;
(iv) $T|_{\mathcal{U}}$ and $T|_{\mathcal{U}^\perp}$ are normal.

Proof. Choose an orthonormal basis adapted to $\mathcal{U} \oplus \mathcal{U}^\perp$.
In these coordinates, Theorem 3.4.28 makes the matrix of T block diagonal, say $A = B \oplus C$.
The block diagonal form shows $\mathcal{U}^\perp$ is invariant under T and $\mathcal{U}$ is invariant under T^*. It also gives $(T|_{\mathcal{U}})^* = (T^*)|_{\mathcal{U}}$ because the adjoint block on $\mathcal{U}$ is B^*. Finally, normality of A reduces blockwise to normality of B and C.
Structural Commentary: Normal operator turns every invariant subspace into an orthogonal reducing summand.
Key Insight: Normality kills the one-way coupling between $\mathcal{U}$ and $\mathcal{U}^\perp$. □

Proposition 3.4.30. *Coprime Normal Primary Pieces Are Orthogonal*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal on a finite-dimensional complex inner-product space, and let $\mathcal{S}, \mathcal{R} \leq \mathcal{V}$ be T-invariant. If $p, q \in \mathbb{C}[z]$, $(p, q) = 1$, $p(T)|_{\mathcal{S}} = 0$, and $q(T)|_{\mathcal{R}} = 0$, then $\mathcal{S} \perp \mathcal{R}$.

Proof. Choose $a, b \in \mathbb{C}[z]$ with $ap + bq = 1$. For $\mathbf{s} \in \mathcal{S}$, $\mathbf{s} = a(T)p(T)\mathbf{s} + b(T)q(T)\mathbf{s} = b(T)q(T)\mathbf{s}$.
By Corollary 3.4.29, $\mathcal{R}$ reduces T. Hence $b(T)^*\mathcal{R} \subseteq \mathcal{R}$ and $q(T)^*|_{\mathcal{R}} = (q(T)|_{\mathcal{R}})^* = 0$.
Thus, for $\mathbf{r} \in \mathcal{R}$, $\langle \mathbf{s}, \mathbf{r} \rangle = \langle b(T)q(T)\mathbf{s}, \mathbf{r} \rangle = \langle \mathbf{s}, q(T)^* b(T)^* \mathbf{r} \rangle = 0$.
Structural Commentary: Coprime annihilators orthogonally separate normal invariant pieces.
Key Insight: Normality makes invariant subspaces reduce, so annihilation also holds after taking adjoints. □

Theorem 3.4.31. *Simultaneous Unitary Diagonalisation of Commuting Normal Operators*
Let $\emptyset \neq \mathcal{F} \subseteq \mathcal{L}(\mathcal{V})$ be a normal family on a finite-dimensional complex inner-product space. Then $\mathcal{F}$ is pairwise commuting if and only if $\mathcal{V}$ has an orthonormal basis of common eigenvectors for $\mathcal{F}$.

Proof.
($\Leftarrow$) In a common orthonormal eigenbasis, every $A \in \mathcal{F}$ is diagonal. Hence all operators in $\mathcal{F}$ commute.
($\Rightarrow$) Induct on $\dim(\mathcal{V})$. The zero-dimensional case is vacuous. If every $A \in \mathcal{F}$ is scalar, any orthonormal basis works. Otherwise choose a nonscalar $A_0 \in \mathcal{F}$ and write its spectral resolution as $A_0 = \sum_j \mu_j P_j$.
Since A_0 is nonscalar, at least two nonzero ranges $\mathrm{range}(P_j)$ occur, and each has dimension strictly smaller than $\dim(\mathcal{V})$. By Corollary 3.4.20, each P_j is a polynomial in A_0.
If $B \in \mathcal{F}$, then $BA_0 = A_0B$, so B commutes with every polynomial in A_0, in particular with P_j. Hence $B(\mathrm{range}(P_j)) \subseteq \mathrm{range}(P_j)$ because $P_jB\mathbf{x} = BP_j\mathbf{x} = B\mathbf{x}$ whenever $P_j\mathbf{x} = \mathbf{x}$.
The same commutation preserves $\ker(P_j) = \mathrm{range}(P_j)^\perp$.
In the decomposition $\mathcal{V} = \mathrm{range}(P_j) \oplus \ker(P_j)$, the matrix of B is therefore block diagonal; since B is normal, its diagonal restrictions are normal block by block.
Thus, on each nonzero $\mathrm{range}(P_j)$, the restricted family is again pairwise commuting and normal.
By induction it has a common orthonormal eigenbasis. Concatenating these bases over all j gives a common orthonormal eigenbasis of $\mathcal{V}$.

Structural Commentary: Spectral projections reduce the whole commuting family.
Key Insight: Diagonalise one nonscalar member, then induct on its eigenspaces. □

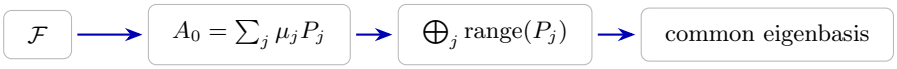


Figure 69: For $A_0 = \sum_j \mu_j P_j$, commutativity gives $BP_j = P_j B$; the ranges reduce $\mathcal{F}$, and induction gives a common orthonormal eigenbasis.

Corollary 3.4.32. *Sums and Products of Commuting Normal Operators*
Let $S, T \in \mathcal{L}(\mathcal{V})$ be commuting normal operators on a finite-dimensional complex inner-product space. Then $S+T$ and ST are normal.

Proof. By Theorem 3.4.31, S and T are diagonal in a common orthonormal basis.
Then $S+T$ and ST are also diagonal in that basis, hence normal.
Structural Commentary: Commuting normal operators share orthonormal coordinates.
Key Insight: In common diagonal form, sums and products remain diagonal. □

Theorem 3.4.33. *Normality Transfers from AB to BA*
Let $A, B \in \operatorname{Mat}(n, n, \mathbb{C})$. If A, B, and AB are normal, then BA is normal.

Proof. By Theorem 3.3.12, AB and BA have the same eigenvalues counted with algebraic multiplicity. Also,

$$\begin{aligned}\operatorname{tr}\big((BA)^*(BA)\big) &= \operatorname{tr}(A^*B^*BA) = \operatorname{tr}(AA^*B^*B)\\ &= \operatorname{tr}(A^*AB^*B) = \operatorname{tr}(A^*ABB^*) = \operatorname{tr}(B^*A^*AB) = \operatorname{tr}\big((AB)^*(AB)\big),\end{aligned}$$

using cyclicity of trace and normality of A and B. Thus $\|BA\|_F^2 = \|AB\|_F^2$, so $\Delta(BA) = \Delta(AB)$.
Since AB is normal, $\Delta(AB) = 0$; hence $\Delta(BA) = 0$, and Theorem 4.3.9 gives that BA is normal.
Structural Commentary: Same spectrum and Frobenius norm give the same Schur defect.
Key Insight: Schur defect turns normality of AB into normality of BA. □

Definition 3.4.34. *Circulant Matrix*
On the 0-indexed basis of $\mathbb{C}^n$, define the *cyclic shift* $P_n \mathbf{e}_j = \mathbf{e}_{(j-1) \bmod n}$.
A matrix $C \in \operatorname{Mat}(n, n, \mathbb{C})$ is *circulant* if, for some $a_0, \dots, a_{n-1} \in \mathbb{C}$,

$$\begin{aligned}C &= a_0 I_n + a_1 P_n + \cdots + a_{n-1} P_n^{n-1} = p(P_n),\\ p(z) &= a_0 + a_1 z + \cdots + a_{n-1} z^{n-1}\end{aligned}$$

Theorem 3.4.35. *Fourier Basis for Cyclic Shifts and Circulants*
Let $n \geq 1$, let $\omega_n := e^{2\pi i/n}$, and let P_n be the cyclic shift.
For $0 \leq k < n$, set $\mathbf{f}_k := (\omega_n^{jk})_{j=0}^{n-1}$; $F_n := [\mathbf{f}_0 \ \cdots \ \mathbf{f}_{n-1}]$; and $\Omega_n := \operatorname{diag}(1, \omega_n, \dots, \omega_n^{n-1})$.
Then $F_n^* F_n = n I_n$, so $Q_n := n^{-1/2} F_n$ is unitary, and $P_n F_n = F_n \Omega_n$.
If $C = p(P_n)$ is circulant, then $Q_n^* C Q_n = p(\Omega_n) = \operatorname{diag}\big(p(1), p(\omega_n), \dots, p(\omega_n^{n-1})\big)$.

Proof. For $0 \leq k, \ell \leq n-1$, $\langle \mathbf{f}_k, \mathbf{f}_\ell \rangle = \sum_{r=0}^{n-1} (\omega_n^{k-\ell})^r = n\delta_{k\ell}$. For $k = \ell$ the sum is n; for $k \neq \ell$, $W = \omega_n^{k-\ell}$ has $W \neq 1$ and $W^n = 1$, so $(1-W)\sum_{r=0}^{n-1} W^r = 0$. Hence $F_n^* F_n = n I_n$, and Theorem 2.4.28 makes Q_n unitary.
Also $P_n \mathbf{f}_k = \omega_n^k \mathbf{f}_k$, so $Q_n^* P_n Q_n = \Omega_n$; applying p gives the formula.
Structural Commentary: Fourier coordinates diagonalize every circulant.
Key Insight: Orthogonality of roots of unity turns a cyclic matrix problem into scalar evaluation. □

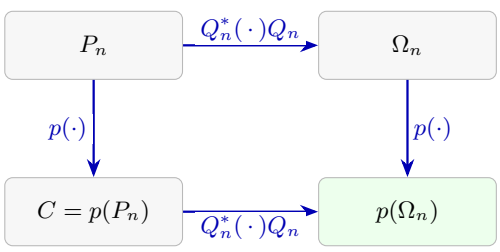


Figure 70: Fourier diagonalizes P_n; polynomial evaluation diagonalizes $C = p(P_n)$.

Proposition 3.4.36. *Permutation Similarity Reorders Diagonal Entries*
Let $\Lambda = \operatorname{diag}(\lambda_1, \dots, \lambda_n)$ and let $P = [\mathbf{e}_{\sigma(1)} \ \cdots \ \mathbf{e}_{\sigma(n)}]$ be a permutation matrix. Then $P^*\Lambda P = \operatorname{diag}(\lambda_{\sigma(1)}, \dots, \lambda_{\sigma(n)})$.

Proof. The (i,j) entry of $P^*\Lambda P$ is $\mathbf{e}_{\sigma(i)}^*\Lambda\mathbf{e}_{\sigma(j)} = \lambda_{\sigma(j)}\mathbf{e}_{\sigma(i)}^*\mathbf{e}_{\sigma(j)} = \lambda_{\sigma(j)}\delta_{ij}$.
Thus all off-diagonal entries are zero, and the jth diagonal entry is $\lambda_{\sigma(j)}$.
Structural Commentary: Permutation similarity applies the same basis reordering to rows and columns.
Key Insight: Diagonal entries move exactly as the standard basis vectors are rearranged. □

Corollary 3.4.37. *Normal Matrices with the Same Spectrum Are Unitarily Similar*
If $A, B \in \operatorname{Mat}(n, n, \mathbb{C})$ are normal and have the same eigenvalues counted with algebraic multiplicity, then A and B are unitarily similar.

Proof. By Theorem 3.4.12, write $A = U\Lambda U^*$ and $B = V\Gamma V^*$ with U, V unitary and Λ, Γ diagonal.
Since eigenvalues agree with multiplicity, Proposition 3.4.36 gives a permutation matrix P such that $\Gamma = P^*\Lambda P$.
Hence $B = (VP^*U^*)A(UPV^*)$, and UPV^* is unitary.
Structural Commentary: Normal matrices are classified up to unitary similarity by their eigenvalue multiset.
Key Insight: After diagonalising both matrices, only a permutation of diagonal entries remains. □

Theorem 3.4.38. *Description of Complex Isometries*
Let $\mathcal{V}$ be a finite-dimensional complex inner-product space and let $S \in \mathcal{L}(\mathcal{V})$. Then S is an isometry if and only if $\mathcal{V}$ has an orthonormal basis consisting of eigenvectors of S whose eigenvalues all have absolute value 1.

Proof. If $S\mathbf{e}_j = \lambda_j\mathbf{e}_j$ in an orthonormal basis with $|\lambda_j| = 1$, then for $\mathbf{v} = \sum_j c_j\mathbf{e}_j$,
$\|S\mathbf{v}\|^2 = \sum_j |\lambda_j c_j|^2 = \sum_j |c_j|^2 = \|\mathbf{v}\|^2$, so S is an isometry. Conversely, if S is an isometry, then $S^*S = I_{\mathcal{V}}$ by Theorem 2.4.26. Finite-dimensional invertibility gives $SS^* = I_{\mathcal{V}}$, so S is normal. By Complex Spectral Theorem, choose orthonormal eigenbasis $S\mathbf{e}_j = \lambda_j\mathbf{e}_j$. Since $\mathbf{e}_j$ is unit and S preserves norms, $|\lambda_j| = \|S\mathbf{e}_j\| = 1$.
Structural Commentary: Complex isometries are normal with unit-circle spectrum.
Key Insight: In an eigenbasis, isometry means $|\lambda_j| = 1$. □

Definition 3.4.39. *Positive Operator*
$T \in \mathcal{L}(\mathcal{V})$ is *positive* if $T = T^*$ and $\langle T\mathbf{v}, \mathbf{v}\rangle \geq 0$ for all $\mathbf{v} \in \mathcal{V}$.

Definition 3.4.40. *Square Root of an Operator*
An operator $R \in \mathcal{L}(\mathcal{V})$ is a *square root* of $T \in \mathcal{L}(\mathcal{V})$ if $R^2 = T$.

Theorem 3.4.41. *Characterization of Positive Operators*
For $T \in \mathcal{L}(\mathcal{V})$, the following are equivalent:

i. *T is positive;*
ii. *T is self-adjoint and every eigenvalue of T is nonnegative;*
iii. *T has a positive square root;*
iv. *T has a self-adjoint square root;*
v. *there exists $R \in \mathcal{L}(\mathcal{V})$ such that $T = R^*R$.*

Proof.

(*i*) $\Rightarrow$ (*ii*) If $T\mathbf{v} = \lambda\mathbf{v}$ and $\mathbf{v} \neq \mathbf{0}$, then $\lambda \in \mathbb{R}$ by Theorem 3.4.5, and $0 \leq \langle T\mathbf{v}, \mathbf{v}\rangle = \lambda\|\mathbf{v}\|^2$, so $\lambda \geq 0$.
(*ii*) $\Rightarrow$ (*iii*) By the real or complex spectral theorem, choose an orthonormal eigenbasis $(\mathbf{e}_1, \dots, \mathbf{e}_n)$ with $T\mathbf{e}_j = \lambda_j\mathbf{e}_j$ and $\lambda_j \geq 0$. Define $S\mathbf{e}_j = \sqrt{\lambda_j}\mathbf{e}_j$ and extend linearly.
Then S is diagonal with nonnegative real diagonal entries in an orthonormal basis, hence self-adjoint and positive, and $S^2\mathbf{e}_j = \lambda_j\mathbf{e}_j = T\mathbf{e}_j$ for every j. Thus $S^2 = T$.
(*iii*) $\Rightarrow$ (*iv*) A positive square root is self-adjoint.
(*iv*) $\Rightarrow$ (*v*) If $T = R^2$ with $R = R^*$, then $T = R^*R$.
(*v*) $\Rightarrow$ (*i*) If $T = R^*R$, then $T^* = T$ and $\langle T\mathbf{v}, \mathbf{v}\rangle = \langle R\mathbf{v}, R\mathbf{v}\rangle = \|R\mathbf{v}\|^2 \geq 0$ for every $\mathbf{v}$, so T is positive.

Structural Commentary: Positivity equals nonnegative spectrum and square-root factorization.
Key Insight: In an eigenbasis, the proof is scalar: $\lambda \geq 0$ iff $\sqrt{\lambda}$ exists. □

Proposition 3.4.42. *Unique Positive Square Root*
Every positive operator on $\mathcal{V}$ has a unique positive square root.

Proof. Existence follows from Theorem 3.4.41. For uniqueness, decompose $\mathcal{V} = E_{\lambda_1}(T) \oplus \cdots \oplus E_{\lambda_m}(T)$.
Let R be a positive square root of T. Since $R^2 = T$, we have $RT = R^3 = TR$, so each $E_{\lambda_j}(T)$ is invariant under R. On $E_{\lambda_j}(T)$, the restriction $R_j := R|_{E_{\lambda_j}(T)}$ is positive and satisfies $R_j^2 = \lambda_j I$.
By the spectral theorem for the positive operator R_j, every eigenvalue ρ of R_j satisfies $\rho \geq 0$ and $\rho^2 = \lambda_j$,

hence $\rho = \sqrt{\lambda_j}$. Therefore $R_j = \sqrt{\lambda_j} I$ on $E_{\lambda_j}(T)$.
Thus every positive square root agrees with the same scalar operator on each eigenspace $E_{\lambda_j}(T)$, so it is unique.
Structural Commentary: Each eigenspace determines its own scalar root.
Key Insight: Positivity selects the nonnegative root. □

Definition 3.4.43. *Positive Square Root Notation*
If T is positive, then $\sqrt{T}$ denotes the unique positive square root of T.

3.5 Operator Functions

Operator functions are defined by applying scalar functions to spectral data.
Corollary 3.4.20 gives $f(T) = \sum_j f(\mu_j) P_j$ for normal operators; this subsection extends that rule to diagonalizable and Jordan forms, with nilpotent terms where needed.

Theorem 3.5.1. *Polynomial Functional Calculus for Diagonalizable Matrices*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ be diagonalizable and write $A = S\Lambda S^{-1}$, where $\Lambda = \mathrm{diag}(\lambda_1, \ldots, \lambda_n)$.
If $p \in \mathbb{F}[z]$, then $p(A) = S\,\mathrm{diag}(p(\lambda_1), \ldots, p(\lambda_n))S^{-1}$.

Proof. For every $j \geq 0$, $A^j = S\Lambda^j S^{-1}$.

If $p(z) = \sum_j c_j z^j$, then $p(A) = \sum_j c_j A^j = S\left(\sum_j c_j \Lambda^j\right) S^{-1} = Sp(\Lambda)S^{-1}$.

Since Λ is diagonal, $p(\Lambda) = \mathrm{diag}(p(\lambda_1), \ldots, p(\lambda_n))$.
Structural Commentary: Diagonalization reduces $p(A)$ to scalar evaluation on eigenvalues.
Key Insight: Compute functions of A by applying the function to the eigenvalues. □

Definition 3.5.2. *Spectral Resolution Notation*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal on a finite-dimensional complex inner-product space.
Define $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$, where $\Omega_T = \mathrm{spec}(T)$.

(i) *Spectral Set*: $\Omega_T = \mathrm{spec}(T)$.
(ii) *Spectral Projection*: P_λ is the orthogonal projection onto $E_\lambda(T)$.
(iii) *Spectral Resolution*: $\{P_\lambda\}_{\lambda \in \Omega_T}$, equivalently $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$.

Theorem 3.5.3. *Finite Spectral Functional Calculus for Normal Operators*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal, with spectral resolution $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$.
The finite spectral functional calculus is $\mathbb{C}^{\Omega_T} \to \mathcal{L}(\mathcal{V})$, where $f \mapsto f(T) := \sum_{\lambda \in \Omega_T} f(\lambda) P_\lambda$.
If $p \in \mathbb{C}[z]$ and $p|_{\Omega_T} = f$, then $p(T) = f(T)$. Hence every $f \in \mathbb{C}^{\Omega_T}$ has a polynomial representative in T.

Proof. For $p|_{\Omega_T} = f$, Corollary 3.4.20 gives $p(T) = \sum_{\lambda \in \Omega_T} p(\lambda) P_\lambda = \sum_{\lambda \in \Omega_T} f(\lambda) P_\lambda = f(T)$.
Since Ω_T is finite, Lagrange interpolation gives $p \in \mathbb{C}[z]$ with $p|_{\Omega_T} = f$.
Structural Commentary: Spectral sums reduce finite functions to polynomial evaluation.
Key Insight: Finite spectral data is interpolated by one polynomial. □

Theorem 3.5.4. *Rules for Normal Spectral Functions*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal, with spectral resolution $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$. For $f, g : \Omega_T \to \mathbb{C}$:

(i) $(f + g)(T) = f(T) + g(T)$ and $(fg)(T) = f(T)g(T)$;
(ii) $f(T)^* = \overline{f}(T)$, where $\overline{f}(\lambda) = \overline{f(\lambda)}$;
(iii) $f(T)$ is normal and $\mathrm{spec}(f(T)) = f(\Omega_T)$.

Proof.

(i) Since $P_\lambda P_\mu = \delta_{\lambda\mu} P_\lambda$, $f(T)g(T) = \sum_{\lambda \in \Omega_T} f(\lambda)g(\lambda) P_\lambda = (fg)(T)$; and $(f + g)(T) = f(T) + g(T)$.

(ii) Since $P_\lambda^* = P_\lambda$, $f(T)^* = \sum_{\lambda \in \Omega_T} \overline{f(\lambda)} P_\lambda = \overline{f}(T)$.

(iii) For each $\lambda \in \Omega_T$, $f(T)|_{\mathrm{range}(P_\lambda)} = f(\lambda) I$.
Thus $f(T)$ is diagonal in an orthonormal eigenbasis, hence normal, and $\mathrm{spec}(f(T)) = f(\Omega_T)$.

Structural Commentary: Spectral projections make operator rules pointwise.
Key Insight: Spectral functions act as scalars on spectral subspaces. □

Example 3.5.5. *Basic Functions from the Spectral Calculus*
Let $T \in \mathcal{L}(\mathcal{V})$ be normal, with spectral resolution $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$. Then $T^* = \sum_{\lambda \in \Omega_T} \overline{\lambda} P_\lambda$.
If $0 \notin \Omega_T$, then $T^{-1} = \sum_{\lambda \in \Omega_T} \lambda^{-1} P_\lambda$. If $T \geq 0$, then $T^{1/2} = \sum_{\lambda \in \Omega_T} \lambda^{1/2} P_\lambda$, where $\lambda^{1/2} \geq 0$.

Theorem 3.5.6. *Commutativity via Spectral Projections*
Let $T, S \in \mathcal{L}(\mathcal{V})$ be normal, with spectral resolutions $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$ and $S = \sum_{\mu \in \Omega_S} \mu Q_\mu$.

(i) For $R \in \mathcal{L}(\mathcal{V})$, $RT = TR$ if and only if $RP_\lambda = P_\lambda R$ for all $\lambda \in \Omega_T$.
(ii) $TS = ST$ if and only if $P_\lambda Q_\mu = Q_\mu P_\lambda$ for all $\lambda \in \Omega_T$ and $\mu \in \Omega_S$.
(iii) If $f : \Omega_T \to \mathbb{C}$ and $g : \Omega_S \to \mathbb{C}$ are injective, then $f(T)g(S) = g(S)f(T)$ if and only if $TS = ST$.

Proof.

(i). If $RT = TR$, then $Rp(T) = p(T)R$ for all $p \in \mathbb{C}[z]$. Since $P_\lambda \in \mathbb{C}[T]$ by Corollary 3.4.20, $RP_\lambda = P_\lambda R$. Conversely, if $RP_\lambda = P_\lambda R$ for all $\lambda \in \Omega_T$, then $RT = \sum_{\lambda \in \Omega_T} \lambda R P_\lambda = \sum_{\lambda \in \Omega_T} \lambda P_\lambda R = TR$.

(ii). If $TS = ST$, then (i) with $R = S$ gives $SP_\lambda = P_\lambda S$. Since $Q_\mu \in \mathbb{C}[S]$, $P_\lambda Q_\mu = Q_\mu P_\lambda$ for all λ, μ. Conversely, if $P_\lambda Q_\mu = Q_\mu P_\lambda$ for all λ, μ, then $TS = \sum_{\lambda,\mu} \lambda\mu P_\lambda Q_\mu = \sum_{\lambda,\mu} \lambda\mu Q_\mu P_\lambda = ST$

(iii). Suppose f and g are injective. If $f(T)g(S) = g(S)f(T)$, choose $h : f(\Omega_T) \to \Omega_T$ and $k : g(\Omega_S) \to \Omega_S$ with $h \circ f = \mathrm{id}_{\Omega_T}$ and $k \circ g = \mathrm{id}_{\Omega_S}$. By Theorem 3.5.3, $T = h(f(T))$ and $S = k(g(S))$, so $TS = ST$. Conversely, if $TS = ST$, then (ii) gives $P_\lambda Q_\mu = Q_\mu P_\lambda$,
and $f(T)g(S) = \sum_{\lambda,\mu} f(\lambda)g(\mu)P_\lambda Q_\mu = \sum_{\lambda,\mu} f(\lambda)g(\mu)Q_\mu P_\lambda = g(S)f(T)$.

Structural Commentary: Spectral projections are the coordinate idempotents of a normal operator.
Key Insight: To commute with a normal operator is to preserve each of its eigenspaces. □

Corollary 3.5.7. *A Commuting Normal Pair Comes from One Normal Operator*
Let $T, S \in \mathcal{L}(\mathcal{V})$ be normal on a finite-dimensional complex inner-product space.
Then $TS = ST \iff \exists\, H \in \mathcal{L}(\mathcal{V})$ normal, $p, q \in \mathbb{C}[z]$ such that $T = p(H)$, $S = q(H)$.
If $TS = ST$, then H may be chosen in $\mathbb{C}[T, S]$.

Proof.

($\Rightarrow$) Suppose $TS = ST$. Write $T = \sum_{\lambda \in \Omega_T} \lambda P_\lambda$ and $S = \sum_{\mu \in \Omega_S} \mu Q_\mu$.
By Theorem 3.5.6, $P_\lambda Q_\mu = Q_\mu P_\lambda$. Let $\mathcal{I} = \{(\lambda, \mu) \in \Omega_T \times \Omega_S : P_\lambda Q_\mu \neq 0\}$.
The projections $P_\lambda Q_\mu$, $(\lambda, \mu) \in \mathcal{I}$, are pairwise orthogonal and sum to $I_\mathcal{V}$.
Choose distinct scalars $\alpha_{\lambda\mu}$ for $(\lambda, \mu) \in \mathcal{I}$ and set $H = \sum_{(\lambda,\mu) \in \mathcal{I}} \alpha_{\lambda\mu} P_\lambda Q_\mu$.
Then H is normal, and $H \in \mathbb{C}[T, S]$ since $P_\lambda \in \mathbb{C}[T]$ and $Q_\mu \in \mathbb{C}[S]$.
Choose $p, q \in \mathbb{C}[z]$ with $p(\alpha_{\lambda\mu}) = \lambda$ and $q(\alpha_{\lambda\mu}) = \mu$ on $\mathcal{I}$.
Hence $p(H) = \sum_{(\lambda,\mu) \in \mathcal{I}} \lambda P_\lambda Q_\mu = T$, and $q(H) = \sum_{(\lambda,\mu) \in \mathcal{I}} \mu P_\lambda Q_\mu = S$.

($\Leftarrow$) If $T = p(H)$ and $S = q(H)$, then $TS = p(H)q(H) = q(H)p(H) = ST$.

Structural Commentary: One scalar labels each common spectral block.
Key Insight: Joint spectral projections reduce a pair of commuting normal operators to finite scalar data. □

Definition 3.5.8. *Matrix Exponential*
For $A \in \mathrm{Mat}(n, n, \mathbb{C})$ and $t \in \mathbb{C}$, define $e^{tA} := \sum_{k=0}^{\infty} \frac{t^k A^k}{k!}$.

Proposition 3.5.9. *Pure Modes for Powers and Exponentials*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ be diagonalizable with $A = S\Lambda S^{-1}$, where $\Lambda = \mathrm{diag}(\lambda_1, \ldots, \lambda_n)$ and the columns of S are $\mathbf{v}_1, \ldots, \mathbf{v}_n$. If $\mathbf{x}_0 = \sum_{j=1}^{n} c_j \mathbf{v}_j$, then $A^k \mathbf{x}_0 = \sum_{j=1}^{n} c_j \lambda_j^k \mathbf{v}_j$ for $k \geq 0$; and $e^{tA}\mathbf{x}_0 = \sum_{j=1}^{n} c_j e^{\lambda_j t} \mathbf{v}_j$ for $t \in \mathbb{C}$.

Proof. Let $c = (c_1, \ldots, c_n)^t$, so $\mathbf{x}_0 = Sc$.
Then $A^k = S\Lambda^k S^{-1}$, and Definition 3.5.8 gives $e^{tA} = S\left(\sum_{r=0}^{\infty} \frac{t^r \Lambda^r}{r!}\right) S^{-1} = S\,\mathrm{diag}(e^{\lambda_1 t}, \ldots, e^{\lambda_n t}) S^{-1}$.
Applying these two formulas to Sc gives the stated identities.
Structural Commentary: Eigenbasis coordinates decouple powers and exponentials.
Key Insight: Diagonalization turns coupled evolution into independent scalar modes. □

Corollary 3.5.10. *Simple Spectrum Commutants Are Polynomials*
If $A \in \mathrm{Mat}(n,n,\mathbb{C})$ has n distinct eigenvalues, then $\mathrm{Comm}(A) = \{p(A) : p \in \mathbb{C}[z]\}$, $\dim \mathrm{Comm}(A) = n$.

Proof. Write $A = S\,\mathrm{diag}(\lambda_1,\dots,\lambda_n)S^{-1}$ with distinct λ_i. Theorem 3.1.45 says every $B \in \mathrm{Comm}(A)$ has the form $S\,\mathrm{diag}(\xi_1,\dots,\xi_n)S^{-1}$. Choose the Lagrange interpolation polynomial $p(z) = \sum_{i=1}^{n} \xi_i \prod_{j\neq i} \frac{z-\lambda_j}{\lambda_i-\lambda_j}$.
Then $p(\lambda_i) = \xi_i$ for every i, so Theorem 3.5.1 gives $B = p(A)$.
The dimension statement is the case $n_1 = \cdots = n_n = 1$ of Theorem 3.1.45.
Structural Commentary: Simple spectrum leaves one scalar choice per eigendirection.
Key Insight: Interpolation converts arbitrary diagonal data into a polynomial in A. □

Corollary 3.5.11. *Reduction Modulo the Characteristic Polynomial*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$ and let $f \in \mathbb{C}[z]$. If $f = p_A q + r$ with $\deg r < n$, then $f(A) = r(A)$.

Proof. Evaluate division identity at A: $f(A) = p_A(A)q(A) + r(A) = 0q(A) + r(A) = r(A)$.
Theorem 3.2.23 gives $p_A(A) = 0$.
Structural Commentary: Cayley–Hamilton makes polynomial expressions in A reducible modulo p_A.
Key Insight: Every high power of A can be rewritten using powers below n. □

Corollary 3.5.12. *Powers Span at Most n Dimensions*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$. Then $\mathrm{span}(I_n, A, A^2, \dots) = \mathrm{span}(I_n, A, \dots, A^{n-1})$, so this has dimension at most n.

Proof. For each $k \geq 0$, Corollary 3.5.11 applied to $f(z) = z^{n+k}$ gives a polynomial r_k of degree at most $n-1$ such that $A^{n+k} = r_k(A)$. Thus every power of A lies in $\mathrm{span}(I_n, A, \dots, A^{n-1})$.
Structural Commentary: Polynomial reduction collapses the infinite power list to its first n terms.
Key Insight: Cayley–Hamilton gives the recurrence relation for powers of A. □

Corollary 3.5.13. *Inverse from the Characteristic Polynomial*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$ have $p_A(z) = z^n + c_{n-1}z^{n-1} + \cdots + c_1 z + c_0$.
If A is invertible, then $A^{-1} = -c_0^{-1}(A^{n-1} + c_{n-1}A^{n-2} + \cdots + c_2 A + c_1 I_n)$.

Proof. Since A is invertible, $0 \notin \mathrm{spec}(A)$, so $c_0 = p_A(0) = (-1)^n \det A \neq 0$.
Cayley–Hamilton gives $A^n + c_{n-1}A^{n-1} + \cdots + c_1 A + c_0 I_n = 0$.
Move $c_0 I_n$ to the other side and multiply by $-c_0^{-1}A^{-1}$.
Structural Commentary: The characteristic equation contains an identity term exactly when A is invertible.
Key Insight: Solve Cayley–Hamilton for A^{-1} after factoring out one A. □

Proposition 3.5.14. *Eigenvalues of a Polynomial in a Matrix*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$, let $p \in \mathbb{C}[z]$, and let $\lambda_1,\dots,\lambda_n$ be the eigenvalues of A, counted with algebraic multiplicity. Then the eigenvalues of $p(A)$ are $p(\lambda_1),\dots,p(\lambda_n)$, counted with algebraic multiplicity.

Proof. Choose a Schur form $U^*AU = T$ with diagonal $\lambda_1,\dots,\lambda_n$. Then $p(A) = Up(T)U^*$, and $p(T)$ is upper triangular with diagonal $p(\lambda_1),\dots,p(\lambda_n)$. Hence $p(A)$ has precisely these eigenvalues.
Structural Commentary: Schur form reduces the claim to the diagonal.
Key Insight: A polynomial preserves triangular form and applies p to each diagonal entry. □

Corollary 3.5.15. *Polynomial Spectral Mapping*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$ and $p \in \mathbb{C}[z]$. Then $\mathrm{spec}(p(A)) = \{p(\lambda) : \lambda \in \mathrm{spec}(A)\}$.

Proof. Let $\lambda_1,\dots,\lambda_n$ be the eigenvalues of A, counted with algebraic multiplicity.
By Proposition 3.5.14, the eigenvalues of $p(A)$, counted with algebraic multiplicity, are $p(\lambda_1),\dots,p(\lambda_n)$. Taking underlying sets gives the claim.
Structural Commentary: Polynomial evaluation pushes the spectrum forward through the scalar polynomial.
Key Insight: Once Schur form has exposed the diagonal, p acts entry by entry there. □

Corollary 3.5.16. *Polynomial Spectral Mapping for Operators*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{C}$, let $T \in \mathcal{L}(\mathcal{V})$, and let $p \in \mathbb{C}[z]$.
Then $\mathrm{spec}(p(T)) = \{p(\lambda) : \lambda \in \mathrm{spec}(T)\}$.

Proof. Choose a basis β and set $A = [T]_{\beta\leftarrow\beta}$.
Then $[p(T)]_{\beta\leftarrow\beta} = p(A)$, and coordinate isomorphism gives $\mathrm{spec}(T) = \mathrm{spec}(A)$ and $\mathrm{spec}(p(T)) = \mathrm{spec}(p(A))$.
Corollary 3.5.15 gives $\mathrm{spec}(p(T)) = \mathrm{spec}(p(A)) = \{p(\lambda) : \lambda \in \mathrm{spec}(A)\} = \{p(\lambda) : \lambda \in \mathrm{spec}(T)\}$.
Structural Commentary: Coordinates reduce the operator statement to the matrix statement.
Key Insight: Polynomial spectral mapping is an operator fact seen through any convenient basis. □

Corollary 3.5.17. *Polynomial Invertibility Test*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$ and $p \in \mathbb{C}[z]$. If $p(\lambda) \neq 0$ for every $\lambda \in \mathrm{spec}(A)$, then $p(A)$ is invertible.

Proof. By Proposition 3.5.14, $0 \notin \mathrm{spec}(p(A))$. Corollary 3.1.9 gives invertibility.
Structural Commentary: Invertibility of $p(A)$ is checked on spectral values of A.
Key Insight: Avoiding zero on the spectrum prevents $p(A)$ from having a kernel. □

Proposition 3.5.18. *Exponential of a Jordan Block*
For $k \geq 1$, $\lambda \in \mathbb{C}$, and $t \in \mathbb{C}$, $e^{tJ_k(\lambda)} = e^{\lambda t} \sum_{r=0}^{k-1} \frac{t^r}{r!} J_k(0)^r$.

Proof. Since $J_k(\lambda) = \lambda I_k + J_k(0)$ and the two summands commute, $e^{tJ_k(\lambda)} = e^{\lambda t I_k} e^{tJ_k(0)}$.
The first factor is $e^{\lambda t} I_k$, and the second exponential truncates because $J_k(0)^k = 0$.
Structural Commentary: The scalar part supplies $e^{\lambda t}$; the nilpotent part supplies a finite polynomial in t.
Key Insight: Nilpotence turns the exponential series into a finite sum. □

Lemma 3.5.19. *Identity Plus Nilpotent Has a Square Root*
Let $\mathcal{V}$ be a complex vector space and let $N \in \mathcal{L}(\mathcal{V})$ be nilpotent. Then $I_\mathcal{V} + N$ has a square root.

Proof. Choose $m \geq 1$ such that $N^m = 0$. If $m = 1$, then $N = 0$ and $I_\mathcal{V}$ is a square root.
Otherwise seek $R = I_\mathcal{V} + a_1 N + \cdots + a_{m-1} N^{m-1}$. With $a_0 = 1$, choose $a_1, \ldots, a_{m-1}$ recursively so the coefficients of R^2 match those of $I_\mathcal{V} + N$: $2a_1 = 1$, and $2a_k + \sum_{r=1}^{k-1} a_r a_{k-r} = 0 \quad (2 \leq k \leq m-1)$
Terms of degree $\geq m$ vanish because $N^m = 0$, so $R^2 = I_\mathcal{V} + N$.
Structural Commentary: Nilpotence reduces the square-root series to a finite polynomial.
Key Insight: Match coefficients recursively. □



Figure 71: Nilpotence truncates the square-root series; coefficients are chosen so $R^2 = I + N$.

Theorem 3.5.20. *Invertible Complex Operators Have Square Roots*
Let $\mathcal{V}$ be finite-dimensional over $\mathbb{C}$ and let $T \in \mathcal{L}(\mathcal{V})$ be invertible. Then T has a square root.

Proof. Let $\mathcal{U}_j = \mathcal{G}_{\lambda_j}(T)$ for the distinct eigenvalues of T. Primary decomposition gives $\mathcal{V} = \mathcal{U}_1 \oplus \cdots \oplus \mathcal{U}_m$, and invertibility gives $\lambda_j \neq 0$. Put $N_j = (T - \lambda_j I_\mathcal{V})|_{\mathcal{U}_j}$.
Then $T|_{\mathcal{U}_j} = \lambda_j \left(I_{\mathcal{U}_j} + \lambda_j^{-1} N_j\right)$ with $\lambda_j^{-1} N_j$ nilpotent. By Lemma 3.5.19, choose $B_j^2 = I_{\mathcal{U}_j} + \lambda_j^{-1} N_j$; choose $\alpha_j^2 = \lambda_j$ and set $R_j = \alpha_j B_j$. Then $R_j^2 = T|_{\mathcal{U}_j}$. Define $R = R_1 \oplus \cdots \oplus R_m$ on the direct sum. Then $R^2 = T$.
Structural Commentary: Take square roots on primary blocks, then recombine.
Key Insight: Invertibility makes every scalar block nonzero. □

3.6 The Umbral Calculus

The umbral calculus is the formal linear algebra of polynomial sequences. It identifies a formal power series with three related objects: an algebra element, a linear functional on $\mathbb{C}[x]$, and an operator on $\mathbb{C}[x]$. Refines duality of Definition 1.5.11 and polynomial operator calculus of Definition 3.2.1, without invoking convergence.

Definition 3.6.1. *Formal Power Series and Order*
Let $\mathscr{F} := \mathbb{C}[[t]]$.

(i) $f \in \mathscr{F}$ are *formal power series* $f(t) = \sum_{k \geq 0} a_k t^k$, w. coefficientwise addition, Cauchy product multiplication.
(ii) The *order* of f is $o(f) := \min\{k : a_k \neq 0\}$, where $o(0) := +\infty$.
(iii) A series is *invertible* if it has a multiplicative inverse in $\mathscr{F}$, and is a *delta series* if $o(f) = 1$.
(iv) If $o(f) = 1$, its *compositional inverse* is denoted $\overline{f}$ and satisfies $f(\overline{f}(t)) = \overline{f}(f(t)) = t$.
(v) The *lower factorial polynomials* are $x)_0 := 1$ and $(x)_n := x(x-1)\cdots(x-n+1)$, for $n \geq 1$

Proposition 3.6.2. *Formal Order Calculus*
For $f, g \in \mathscr{F}$:

(i) $f \in \mathscr{F}^\times \iff o(f) = 0$;
(ii) $o(fg) = o(f) + o(g)$ *and* $o(f+g) \geq \min\{o(f), o(g)\}$;
(iii) if $g(t) = \sum_{k \geq 0} b_k t^k$ *and* $o(f) \geq 1$, *then* $g \circ f := g(f(t)) := \sum_{k \geq 0} b_k f(t)^k \in \mathscr{F}$;

(iv) *if* $o(f) \geq 1$ *and* $g \neq 0$, *then* $o(g \circ f) = o(g)o(f)$;
(v) *if* $o(f) = 1$, *then for every* $g \in \mathscr{F}$ *there are unique* $c_k \in \mathbb{C}$ *such that* $g(t) = \sum_{k\geq 0} c_k f(t)^k$.

Proof.

(i) Let $[t^n] : \mathscr{F} \to \mathbb{C}$ denote coefficient extraction. Then $f \in \mathscr{F}^\times \iff [t^0]f \in \mathbb{C}^\times \iff o(f) = 0$.
(ii) Let $r = o(f)$ and $s = o(g)$. Then $[t^{r+s}]fg = [t^r]f\,[t^s]g \neq 0$, and all lower coefficients vanish; hence $o(fg) = r + s$. Also $[t^m](f + g) = 0$ for $m < \min\{r, s\}$, giving the stated lower bound.
(iii) Since $o(f) \geq 1$, $o(f^k) \to \infty$; hence $\forall n$, coefficient $[t^n] \sum_{k\geq 0} b_k f^k$ has only finitely many nonzero terms.
(iv) Let $s = o(g)$. Then $g \circ f = f^s \left(b_s + \sum_{k>s} b_k f^{k-s} \right)$, where second factor has order 0; hence $o(g \circ f) = s\, o(f)$.
(v) If $o(f) = 1$ and $f(t) = a_1 t + \cdots$ with $a_1 \neq 0$, then $f^k = a_1^k t^k +$ higher terms. Thus $\{t^k\}_{k\geq 0} \mapsto \{f^k\}_{k\geq 0}$ is triangular with nonzero diagonal, so recursive coefficient matching gives existence and uniqueness.

Structural Commentary: Delta series behave like a formal change of coordinate at $t = 0$.
Key Insight: Order replaces convergence as the bookkeeping device. □

Definition 3.6.3. *Derivation*
Let $\mathcal{A}$ be an algebra over a field. Linear operator $\partial : \mathcal{A} \to \mathcal{A}$ is a *derivation* if $\partial(ab) = (\partial a)b + a(\partial b)$ $\forall a, b \in \mathcal{A}$.

Proposition 3.6.4. *Formal Differentiation*
Define $\partial_t : \mathscr{F} \to \mathscr{F}$ *by* $\partial_t \left(\sum_{k\geq 0} a_k t^k \right) := \sum_{k\geq 1} k a_k t^{k-1}$.
Then $\partial_t \in \mathrm{Der}_{\mathbb{C}}(\mathscr{F})$; *equivalently, for all* $f, g \in \mathscr{F}$, $\partial_t(fg) = (\partial_t f)g + f(\partial_t g)$.

Proof. Linearity is immediate. Let $f(t) = \sum_{k\geq 0} a_k t^k$ and $g(t) = \sum_{k\geq 0} b_k t^k$.

For $n \geq 0$, $[t^n]\partial_t(fg) = (n+1) \sum_{j=0}^{n+1} a_j b_{n+1-j}$, while $[t^n]\big((\partial_t f)g + f(\partial_t g)\big) = \sum_{j=1}^{n+1} j a_j b_{n+1-j} + \sum_{j=0}^{n} (n+1-j) a_j b_{n+1-j}$.
The last two sums combine to the preceding coefficient, so the Leibniz identity holds.
Structural Commentary: Formal differentiation is the canonical derivation of $\mathbb{C}[[t]]$.
Key Insight: The Leibniz rule is a Cauchy-product coefficient identity. □

Definition 3.6.5. *The Umbral Algebra*

(i) *Polynomial algebra*: set $\mathscr{P} := \mathbb{C}[x]$ and let $\mathscr{P}' := \mathrm{Hom}_{\mathbb{C}}(\mathscr{P}, \mathbb{C})$.
(ii) *Umbral pairing*: for $L \in \mathscr{P}'$ and $p(x) \in \mathscr{P}$, write $\langle L \mid p(x) \rangle := L(p(x))$.
(iii) *Exponential normalization*: write each $f \in \mathscr{F}$ as $f(t) = \sum_{k\geq 0} \frac{a_k}{k!} t^k$.
(iv) *Associated functional*: set $\langle f(t) \mid x^n \rangle := a_n$ for $n \geq 0$, extending $\mathbb{C}$-linearly.

Theorem 3.6.6. *Formal Series Represent All Algebraic Functionals*
The $\mathbb{C}$*-linear map* $\Phi : \mathscr{P}' \xrightarrow{\sim} \mathscr{F}$; $L \longmapsto \sum_{k\geq 0} \frac{\langle L | x^k \rangle}{k!} t^k$ *is an isomorphism with inverse*

$$\Phi^{-1}\left(\sum_{k\geq 0} \frac{a_k}{k!} t^k \right)\left(\sum_{k=0}^{d} c_k x^k \right) := \sum_{k=0}^{d} c_k a_k.$$

Thus, under Φ^{-1}, $\langle t^k \mid x^n \rangle = n!\delta_{n,k}$ *and* $\langle t^k \mid p \rangle = p^{(k)}(0)$, $(p \in \mathscr{P})$.
Thus $t^0 = \mathrm{ev}_0$ *and* $t^k = D^k|_0$.

Proof. Since $\{x^k\}_{k\geq 0}$ is a Hamel basis of $\mathscr{P}$, $\mathscr{P}' \cong \mathbb{C}^{\mathbb{N}}$ via $L \mapsto (\langle L \mid x^k \rangle)_{k\geq 0}$.
The inverse sends the coefficient sequence (a_k) to the unique functional with $L(x^k) = a_k$, so Φ is bijective.
Moreover, $t^k = k!\, t^k/k!$ gives $\langle t^k \mid x^n \rangle = n!\delta_{n,k}$, and linearity gives $\langle t^k \mid p \rangle = p^{(k)}(0)$.
Structural Commentary: $\mathscr{P}'$ is the sequence space encoded by exponentially normalised formal series.
Key Insight: A formal series is a compact record of the functional's values on all monomials. □

Definition 3.6.7. *Delta and Invertible Functionals*
Under the identification $\mathscr{F} \simeq \mathscr{P}'$:

(i) *Delta functional*: a functional represented by $f \in \mathscr{F}$ with $o(f) = 1$.
(ii) *Invertible functional*: a functional represented by $f \in \mathscr{F}^\times$, equivalently $o(f) = 0$.

Proposition 3.6.8. *Basic Umbral Identities*
For $f, g \in \mathscr{F}$, $p \in \mathscr{P}$, *and* $n \geq 0$:

(i) $f(t) = \sum_{k\geq 0} \frac{\langle f(t) \mid x^k \rangle}{k!} t^k$;

(ii) $p(x) = \sum_{k=0}^{\deg p} \frac{\langle t^k \mid p\rangle}{k!} x^k$;

(iii) $\langle fg \mid x^n\rangle = \sum_{k=0}^{n} \binom{n}{k} \langle f \mid x^k\rangle \langle g \mid x^{n-k}\rangle$;

(iv) *if* $o(f) > \deg p$, *then* $\langle f \mid p\rangle = 0$;

(v) *if* $o(f_k) = k$ *for all* $k \geq 0$, *then* $\left\langle \sum_{k\geq 0} a_k f_k \middle| p \right\rangle = \sum_{k=0}^{\deg p} a_k \langle f_k \mid p\rangle$, *with* $\sum_{k\geq 0} a_k f_k$ *well-defined in* $\mathscr{F}$.

Proof.

(i) If $f(t) = \sum_{k\geq 0} a_k t^k/k!$, then $a_k = \langle f \mid x^k\rangle$ by definition.

(ii) Since $\langle t^k \mid p\rangle = p^{(k)}(0)$, this is Taylor's formula for the polynomial p.

(iii) Write $f = \sum a_k t^k/k!$ and $g = \sum b_k t^k/k!$.

The Cauchy product is $fg = \sum_{n\geq 0} \left(\sum_{k=0}^{n} \binom{n}{k} a_k b_{n-k} \right) \frac{t^n}{n!}$, and $a_k = \langle f \mid x^k\rangle$, $b_k = \langle g \mid x^k\rangle$.

(iv) If $p = \sum_{j=0}^{d} c_j x^j$ and $o(f) > d$, then $\langle f \mid x^j\rangle = 0$ for $0 \leq j \leq d$.

(v) Since $o(f_k) = k$, each coefficient of $\sum_{k\geq 0} a_k f_k$ has only finitely many contributions.

With $d = \deg p$, (iv) gives $\langle f_k \mid p\rangle = 0$ for $k > d$, proving the finite formula.

Structural Commentary: The umbral pairing dualizes coefficient calculus.
Key Insight: Order makes infinite formal sums act finitely on each polynomial. □

Theorem 3.6.9. *Multiplication by x Moves to Formal Differentiation*
For all $f \in \mathscr{F}$ and $p \in \mathscr{P}$, $\langle f \mid xp\rangle = \langle \partial_t f \mid p\rangle$.
Equivalently, multiplication by x on $\mathscr{P}$ has adjoint ∂_t on $\mathscr{F}$.

Proof. By linearity it suffices to take $p = x^n$.
For $f = \sum_{k\geq 0} a_k t^k/k!$, $\partial_t f = \sum_{k\geq 0} a_{k+1} \frac{t^k}{k!}$, where $\langle \partial_t f \mid x^n\rangle = a_{n+1} = \langle f \mid x^{n+1}\rangle$.
Structural Commentary: Under the umbral pairing, $x^* = \partial_t$.
Key Insight: The factorial normalisation shifts coefficients by one without extra constants. □

Definition 3.6.10. *Basic Umbral Functionals*
For $a \in \mathbb{C}$ and $p \in \mathscr{P}$:

(i) *Evaluation*: $\epsilon_a = e^{at}$ and $\langle e^{at} \mid p\rangle = p(a)$; hence $\epsilon_a \epsilon_b = \epsilon_{a+b}$.
(ii) *Forward difference*: $\Delta_a = e^{at} - 1$ and $\langle \Delta_a \mid p\rangle = p(a) - p(0)$.
(iii) *Abel functional*: te^{at} and $\langle te^{at} \mid p\rangle = p'(a)$.
(iv) *Gamma functional*: $(1-t)^{-1}$ and $\left\langle (1-t)^{-1} \middle| p \right\rangle = \int_0^\infty p(u) e^{-u}\, du$.
(v) *Integral functional*: $\left\langle \frac{e^{at}-1}{t} \middle| p \right\rangle = \int_0^a p(u)\, du$
(vi) *Bernoulli numbers*: for $a = 1$, $B_n = \left\langle \frac{t}{e^t - 1} \middle| x^n \right\rangle$.

Definition 3.6.11. *Formal Series as Polynomial Operators*
For $f(t) = \sum_{k\geq 0} a_k t^k/k! \in \mathscr{F}$:

(i) *Operator action*: define $f(t) \in \mathrm{End}_{\mathbb{C}}(\mathscr{P})$ by $f(t)p(x) := \sum_{k\geq 0} \frac{a_k}{k!} p^{(k)}(x)$, a finite sum for each $p \in \mathscr{P}$.

(ii) *Monomial action*: for $n \geq 0$, $f(t)x^n = \sum_{k=0}^{n} \binom{n}{k} a_k x^{n-k}$.

Theorem 3.6.12. *Functional-Operator Compatibility*
For $f, g \in \mathscr{F}$ and $p \in \mathscr{P}$, $\langle fg \mid p\rangle = \langle f \mid gp\rangle$. In particular, $\langle f \mid p\rangle = \langle t^0 \mid fp\rangle$.

Proof. By linearity take $p = x^n$. Write $f = \sum_{i\geq 0} a_i t^i/i!$ and $g = \sum_{j\geq 0} b_j t^j/j!$.
Then $\langle fg \mid x^n\rangle = \sum_{j=0}^{n} \binom{n}{j} a_{n-j} b_j$. Also $gx^n = \sum_{j=0}^{n} \binom{n}{j} b_j x^{n-j}$ and $\langle f \mid gx^n\rangle = \sum_{j=0}^{n} \binom{n}{j} b_j a_{n-j}$.
Thus $\langle fg \mid x^n\rangle = \langle f \mid gx^n\rangle$. Taking $(f, g) = (t^0, f)$ in the first identity gives $\langle f \mid p\rangle = \langle t^0 \mid fp\rangle$.
Structural Commentary: Product in $\mathscr{F}$ is adjoint to operator action.
Key Insight: The same series acts on either side of the umbral pairing. □

Definition 3.6.13. *Basic Umbral Operators*
For $a \in \mathbb{C}$ and $p \in \mathscr{P}$:

(i) *Translation*: $e^{at}p(x) = p(x+a)$.
(ii) *Forward difference*: $(e^{at}-1)p(x) = p(x+a) - p(x)$.
(iii) *Abel operator*: $te^{at}p(x) = p'(x+a)$.
(iv) *Gamma operator*: $(1-t)^{-1}p(x) = \int_0^\infty p(x+u)e^{-u}\,du$
(v) *Integral operator*: $\frac{e^{at}-1}{t}p(x) = \int_x^{x+a} p(u)\,du$

Theorem 3.6.14. *Shift-Invariant Operators*
For $T \in \mathrm{End}_{\mathbb{C}}(\mathscr{P})$, the following are equivalent:

(i) $T = f(t)$ *for some* $f \in \mathscr{F}$;
(ii) $[T,t] = 0$;
(iii) $[T, h(t)] = 0$ *for every delta series* h;
(iv) $[T, e^{at}] = 0$ *for all* $a \in \mathbb{C}$.

Proof.

(i) $\Rightarrow$ *(ii), (iii), (iv)* If $T = f(t)$, then $\forall u(t) \in \mathscr{F}$, $[T, u(t)] = [f(t), u(t)] = 0$, since $\mathscr{F}$ is commutative.
(iii) $\Rightarrow$ *(ii)* Since $o(t) = 1$, $(\forall h \in \mathscr{F},\ o(h) = 1 \Rightarrow [T, h(t)] = 0) \Rightarrow [T,t] = 0$.
(iv) $\Rightarrow$ *(ii)* For each $p \in \mathscr{P}$, $0 = [a^1]\,[T, e^{at}]p = [T,t]p$, so $[T,t] = 0$.
(ii) $\Rightarrow$ *(i)* Set $g(t) := \sum_{k\geq 0} \frac{\langle t^0 | Tx^k\rangle}{k!} t^k$. Let $a_k := \langle t^0 \mid Tx^k\rangle$.

Since $[T,t] = 0$, for $m, k \geq 0$, $\langle t^m \mid Tx^k\rangle = \langle t^0 \mid t^m Tx^k\rangle = \langle t^0 \mid Tt^m x^k\rangle = (k)_m a_{k-m}$, with the convention $a_{k-m} = 0$ if $m > k$. The same computation from Definition 3.6.11 gives $\langle t^m \mid g(t)x^k\rangle = (k)_m a_{k-m}$.
Thus $\langle t^m \mid Tx^k - g(t)x^k\rangle = 0$ for all m, k. Since $\{t^m : m \geq 0\}$ separates $\mathscr{P}$, $T = g(t)$.

Structural Commentary: The commutant of t on $\mathscr{P}$ is $\mathscr{F}$.
Key Insight: The derivative is enough to recover every coefficient of the operator. □

Definition 3.6.15. *Sheffer, Associated, and Appell Sequences*
Let f be a delta series, g an invertible series, and $\deg s_n = n$.

(i) *Sheffer sequence*: $(s_n)_{n\geq 0}$ is Sheffer for (g,f) if $\langle g(t)f(t)^k \mid s_n(x)\rangle = n!\delta_{n,k}$ for $n, k \geq 0$.
(ii) *Associated sequence*: a Sheffer sequence for $(1, f)$.
(iii) *Appell sequence*: a Sheffer sequence for (g, t).

Theorem 3.6.16. *Existence and Expansion Theorems*
Let f be a delta series, g an invertible series, and $q_k := g(t)f(t)^k$.

(i) There is a unique sequence $(s_n)_{n\geq 0}$ with $\deg s_n = n$ and $\langle q_k \mid s_n\rangle = n!\delta_{n,k}$.
(ii) If $p_n := g(t)s_n$, then (s_n) is Sheffer for (g,f) iff (p_n) is associated with f.
(iii) For all $h \in \mathscr{F}$, $h = \sum_{k\geq 0} \frac{\langle h | s_k\rangle}{k!} q_k$.
(iv) For all $p \in \mathscr{P}$, $p = \sum_{k=0}^{\deg p} \frac{\langle q_k | p\rangle}{k!} s_k$.

Proof.

(i) Since $o(q_k) = k$, the matrix $A_n = (\langle q_i \mid x^j\rangle)_{0\leq i,j\leq n}$ is triangular with nonzero diagonal.
Hence the equations $\langle q_k \mid s_n\rangle = n!\delta_{n,k}$ have a unique solution $s_n \in \mathscr{P}_{\leq n}$, necessarily with $\deg s_n = n$.
(ii) By Theorem 3.6.12, $\langle f(t)^k \mid p_n\rangle = \langle f(t)^k \mid g(t)s_n\rangle = \langle q_k \mid s_n\rangle$.
(iii) Since $o(q_k) = k$, every $h \in \mathscr{F}$ has a unique expansion $h = \sum c_k q_k$.
Pairing with s_n gives $\langle h \mid s_n\rangle = c_n n!$, so $c_n = \langle h \mid s_n\rangle / n!$.
(iv) Let $d = \deg p$ and $r := p - \sum_{k=0}^{d} \langle q_k \mid p\rangle s_k/k!$.

Then $\langle q_j \mid r\rangle = 0$ for all $j \leq d$, and also for $j > d$ by $o(q_j) > d$. Since the q_j separate $\mathscr{P}$, $r = 0$.

Structural Commentary: (s_n) is dual to $q_k = gf^k$.
Key Insight: The order filtration makes the orthogonality equations triangular. □

Theorem 3.6.17. *Generating and Operator Characterizations*
Let f be a delta series with compositional inverse $\overline{f}$, let $g \in \mathscr{F}^\times$, and let $(p_n), (s_n)$ be polynomial sequences with $\deg p_n = \deg s_n = n$. Set $P(y,u) := \sum_{n\geq 0} \frac{p_n(y)}{n!}u^n$ and $S(y,u) := \sum_{n\geq 0} \frac{s_n(y)}{n!}u^n$.

(i) (p_n) is associated with f iff $P(y,u) = e^{y\overline{f}(u)}$

(ii) (s_n) *is Sheffer for* (g,f) *iff* $S(y,u)=g(\overline{f}(u))^{-1}e^{y\overline{f}(u)}$
(iii) (p_n) *is associated with* f *iff* $p_n(0)=\delta_{n,0}$, $f(t)p_n=np_{n-1}$, *where* $p_{-1}:=0$.
(iv) (s_n) *is Sheffer for* (γ,f) *for some* $\gamma\in\mathscr{F}^\times$ *iff* $f(t)s_n=ns_{n-1}$, *where* $s_{-1}:=0$.
(v) If (p_n) *is associated with* f, *then* $p_n(x)=\sum_{k=0}^{n}\frac{\langle\overline{f}(t)^k|x^n\rangle}{k!}\,x^k$
(vi) If (s_n) *is Sheffer for* (g,f), *then* $s_n(x)=\sum_{k=0}^{n}\frac{1}{k!}\left\langle g(\overline{f}(t))^{-1}\overline{f}(t)^k\middle|x^n\right\rangle x^k$.

Proof.

(i) ($\Rightarrow$) By Theorem 3.6.16, $e^{yt}=\sum_{n\ge0}\frac{p_n(y)}{n!}f(t)^n$. Substitute $t=\overline{f}(u)$ to obtain $P(y,u)=e^{y\overline{f}(u)}$.

(i) ($\Leftarrow$) If $P(x,u)=e^{x\overline{f}(u)}$, then for all $k\ge0$, $\sum_{n\ge0}\frac{\langle f(t)^k|p_n\rangle}{n!}u^n=\langle f(t)^k\mid P(x,u)\rangle=f(\overline{f}(u))^k=u^k$.
Thus $\langle f(t)^k\mid p_n\rangle=n!\delta_{n,k}$.

(ii) ($\Rightarrow$) By Theorem 3.6.16, $e^{yt}=\sum_{n\ge0}\frac{s_n(y)}{n!}g(t)f(t)^n$. Substitute $t=\overline{f}(u)$ and divide by $g(\overline{f}(u))$.

(ii) ($\Leftarrow$) If $S(x,u)=g(\overline{f}(u))^{-1}e^{x\overline{f}(u)}$, then $\sum_{n\ge0}\frac{\langle g(t)f(t)^k|s_n\rangle}{n!}u^n=\langle g(t)f(t)^k\mid S(x,u)\rangle=u^k$,
so $\langle g(t)f(t)^k\mid s_n\rangle=n!\delta_{n,k}$.

(iii) ($\Rightarrow$) If (p_n) is associated, then $p_n(0)=\langle1\mid p_n\rangle=\delta_{n,0}$ and, for all k, $\langle f(t)^k\mid f(t)p_n-np_{n-1}\rangle=0$.
Since $\{f(t)^k:k\ge0\}$ separates $\mathscr{P}$, $f(t)p_n=np_{n-1}$.

(iii) ($\Leftarrow$) The lowering relation gives $f(t)^kp_n=(n)_kp_{n-k}$.
Hence $\langle f(t)^k\mid p_n\rangle=\langle1\mid f(t)^kp_n\rangle=(n)_kp_{n-k}(0)=n!\delta_{n,k}$.

(iv) ($\Rightarrow$) If (s_n) is Sheffer for (γ,f), then for all k, $\langle\gamma(t)f(t)^k\mid f(t)s_n-ns_{n-1}\rangle=0$.
Since $\{\gamma(t)f(t)^k:k\ge0\}$ separates $\mathscr{P}$, $f(t)s_n=ns_{n-1}$.

(iv) ($\Leftarrow$) Define $L\in\mathscr{P}'$ by $L(s_0)=1$ and $L(s_n)=0$ for $n>0$, and put $\gamma:=\Phi(L)$. Then $\langle\gamma\mid1\rangle=L(1)\neq0$, so $\gamma\in\mathscr{F}^\times$, and $\langle\gamma(t)f(t)^k\mid s_n\rangle=\langle\gamma(t)\mid f(t)^ks_n\rangle=(n)_kL(s_{n-k})=n!\delta_{n,k}$.

(v) From (i), $e^{x\overline{f}(u)}=\sum_{k\ge0}\frac{\overline{f}(u)^k}{k!}x^k$. Applying $n![u^n]$ gives the stated formula for p_n.

(vi) From (ii), $g(\overline{f}(u))^{-1}e^{x\overline{f}(u)}=\sum_{k\ge0}\frac{g(\overline{f}(u))^{-1}\overline{f}(u)^k}{k!}x^k$. Applying $n![u^n]$ gives the stated formula for s_n.

Structural Commentary: $\overline{f}$ converts f-lowering to exponential generating functions.
Key Insight: Generating functions package the whole triangular duality into one formal identity. □

Method 3.6.18. *Building Sheffer Sequences*
To construct the Sheffer sequence for (g,f):

(i) Verify $o(f)=1$ and $o(g)=0$.
(ii) If $\overline{f}$ is explicit, compute $S(y,t):=g(\overline{f}(t))^{-1}e^{y\overline{f}(t)}$, where $s_n(y)=n![t^n]S(y,t)$.
(iii) If $\overline{f}$ is unavailable, solve recursively from $f(t)s_n=ns_{n-1}$, where $\langle g(t)f(t)^k\mid s_n\rangle=n!\delta_{n,k}$.
(iv) For associated sequences, set $g=1$ and use Theorem 3.6.30.

Definition 3.6.19. *Binomial Type and Sheffer Identity*
Let (p_n) and (s_n) be polynomial sequences.

(i) *Binomial type*: (p_n) is of binomial type if $p_n(x+y)=\sum_{k=0}^{n}\binom{n}{k}p_k(y)p_{n-k}(x)$.

(ii) *Sheffer identity*: (s_n) satisfies the Sheffer identity relative to (p_n) if $s_n(x+y)=\sum_{k=0}^{n}\binom{n}{k}p_k(y)s_{n-k}(x)$.

Theorem 3.6.20. *Action and Binomial Identities*
Let (s_n) *be Sheffer for* (g,f), *let* (p_n) *be associated with* f, *and write* $E_y:=e^{yt}$. *Then:*

(i) for $h\in\mathscr{F}$ *and* $n\ge0$, $h(t)s_n=\sum_{k=0}^{n}\binom{n}{k}\langle h\mid s_k\rangle p_{n-k}$;
(ii) (p_n) *is associated with a delta series iff* (p_n) *is of binomial type;*
(iii) (s_n) *is Sheffer for some* $g\in\mathscr{F}^\times$ *and same* f *iff* (s_n) *satisfies the Sheffer identity relative to* (p_n).

Proof.

(i) By Theorem 3.6.16, $h=\sum_{k\ge0}\frac{\langle h|s_k\rangle}{k!}g(t)f(t)^k$.
Since $g(t)s_n=p_n$ and $f(t)^kp_n=(n)_kp_{n-k}$, $h(t)s_n=\sum_{k=0}^{n}\frac{\langle h|s_k\rangle}{k!}(n)_kp_{n-k}$.

(ii) If (p_n) is associated, apply (i) with $s_n = p_n$, $g = 1$, and $h = E_y$ to get the binomial identity.
Conversely, if the binomial identity holds, then $p_n(0) = \delta_{n,0}$ and the operator $Qp_n := np_{n-1}$ commutes with every E_y; hence $Q = f(t)$ by Theorem 3.6.14.
Since Q lowers degree by one, $o(f) = 1$, and Theorem 3.6.17 gives that (p_n) is associated with f.
(iii) If (s_n) is Sheffer, apply (i) with $h = E_y$ to obtain the Sheffer identity.
Conversely, applying $f(t)$ in y-variable to Sheffer identity and then setting $y = 0$ gives $f(t)s_n = ns_{n-1}$; Theorem 3.6.17 yields a $g \in \mathscr{F}^\times$ such that (s_n) is Sheffer for (g, f).

Structural Commentary: Translation turns the operator formula into addition laws for polynomial sequences.
Key Insight: Binomial identities are the visible polynomial shadow of a delta operator. □

Example 3.6.21. *Standard Sheffer Examples*

(i) *Monomials*: $p_n(x) = x^n$ is associated with $f(t) = t$; the binomial identity is the usual binomial theorem.
(ii) *Lower factorials*: $(x)_n$ is associated with $f(t) = e^t - 1$, with $(1+t)^y = e^{y\log(1+t)} = \sum_{k\geq 0} \frac{(y)_k}{k!} t^k$,
and Vandermonde identity $(x+y)_n = \sum_{k=0}^{n} \binom{n}{k} (x)_k (y)_{n-k}$.
(iii) *Abel polynomials*: associated with $f(t) = te^{at}$, $A_n(x; a) = x(x - an)^{n-1}$.
(iv) *Hermite polynomials*: for $g(t) = e^{t^2/2}$, the polynomial is $H_n(x) = e^{-t^2/2} x^n$, where $e^{yt - t^2/2} = \sum_{k\geq 0} \frac{H_k(y)}{k!} t^k$.
(v) *Laguerre polynomials*: order α, Sheffer for $g(t) = (1-t)^{-\alpha-1}$, $f(t) = \frac{t}{t-1}$, $\frac{1}{(1-t)^{\alpha+1}} e^{yt/(t-1)} = \sum_{k\geq 0} \frac{L_k^{(\alpha)}(y)}{k!} t^k$.

Definition 3.6.22. *Umbral Operators, Continuity, and Adjoints*
Let (s_n) be Sheffer for (g, f), and let (p_n) be associated with f.
(i) *Sheffer operator*: $\lambda_{g,f} \in \mathcal{L}(\mathscr{P})$ is defined by $\lambda_{g,f}(x^n) = s_n(x)$.
(ii) *Umbral operator*: $\lambda_f \in \mathcal{L}(\mathscr{P})$ is defined by $\lambda_f(x^n) = p_n(x)$.
(iii) *Sheffer/umbral shift*: $\theta_{g,f} s_n = s_{n+1}$; in the associated case, $\theta_f p_n = p_{n+1}$.
(iv) *Continuous operator*: $T \in \mathcal{L}(\mathscr{F})$ is continuous iff $o(f_k) \to \infty \Longrightarrow o(Tf_k) \to \infty$.
Equivalently, for every defined formal sum with $o(f_k) \to \infty$, $T\left(\sum_{k\geq 0} a_k f_k\right) = \sum_{k\geq 0} a_k T(f_k)$
(v) *Algebraic adjoint*: for $\tau \in \mathcal{L}(\mathscr{P})$, $\tau^\times \in \mathcal{L}(\mathscr{F})$ is defined by $\langle \tau^\times h \mid p\rangle = \langle h \mid \tau p\rangle$ for $h \in \mathscr{F}$, $p \in \mathscr{P}$.
This is the dual-map adjoint of Definition 1.5.11, not an inner-product adjoint.

Theorem 3.6.23. *Continuous Operators Are Exactly Operator Adjoints*
Let $\mathcal{L}_{\mathrm{cont}}(\mathscr{F})$ denote the continuous linear operators on $\mathscr{F}$.
Then $(-)^\times : \mathcal{L}(\mathscr{P}) \xrightarrow{\sim} \mathcal{L}_{\mathrm{cont}}(\mathscr{F})$, $\tau \longmapsto \tau^\times$ is a vector-space isomorphism.

Proof.
(⇒) Let $\tau \in \mathcal{L}(\mathscr{P})$ and $o(h_j) \to \infty$. For fixed n, $\tau x^n \in \mathscr{P}$, hence $\langle \tau^\times h_j \mid x^n\rangle = \langle h_j \mid \tau x^n\rangle = 0$ for all sufficiently large j. Thus $o(\tau^\times h_j) \to \infty$, so $\tau^\times \in \mathcal{L}_{\mathrm{cont}}(\mathscr{F})$.
If $\tau^\times = 0$, then $\langle h \mid \tau p\rangle = 0$ for all h, p, hence $\tau = 0$; the map is injective.
(⇐) Let $T \in \mathcal{L}_{\mathrm{cont}}(\mathscr{F})$. Define $\tau_T x^n := \sum_{k\geq 0} \frac{\langle Tt^k \mid x^n\rangle}{k!} x^k$. Since $o(Tt^k) \to \infty$, the sum is finite for each n.
For all m, n, $\langle \tau_T^\times t^m \mid x^n\rangle = \langle t^m \mid \tau_T x^n\rangle = \langle Tt^m \mid x^n\rangle$. Thus $\tau_T^\times t^m = Tt^m$ for every m.
By continuity, $\tau_T^\times = T$, so the map is surjective.

Structural Commentary: Continuity is the finiteness condition for dualising polynomial operators.
Key Insight: Finiteness on each polynomial degree is the hidden topology. □

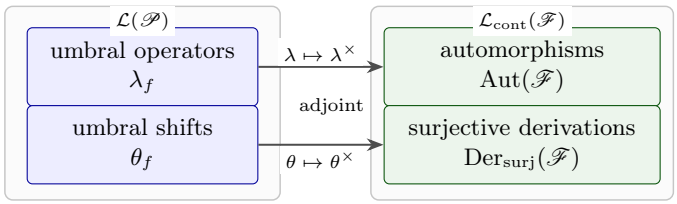


Figure 72: Adjoints identify $\mathcal{L}(\mathscr{P})$ with $\mathcal{L}_{\mathrm{cont}}(\mathscr{F})$ and map the umbral classes shown.

Theorem 3.6.24. *Umbral Operators and Adjoint Automorphisms*
Let f be a delta series with associated sequence (p_n) and umbral operator $\lambda_f(x^n) = p_n(x)$.
Let $C_{\overline{f}} : \mathscr{F} \to \mathscr{F}$ be substitution by $\overline{f}$. Then $\lambda_f^\times = C_{\overline{f}}$, $C_{\overline{f}}(h) = h(\overline{f})$, so $\lambda_f^\times \in \mathrm{Aut}_{\mathbb{C}\text{-}alg}(\mathscr{F})$ and $\lambda_f^\times f = t$.
Moreover, $\lambda \in \mathcal{L}(\mathscr{P})$ is umbral iff $\lambda^\times \in \mathrm{Aut}_{\mathbb{C}\text{-}alg}(\mathscr{F})$.

Proof.
($\Rightarrow$) If $\lambda = \lambda_f$, then for $k, n \geq 0$, $\langle \lambda_f^\times f^k \mid x^n \rangle = \langle f^k \mid p_n \rangle = n!\delta_{n,k} = \langle t^k \mid x^n \rangle$.
Thus $\lambda_f^\times f^k = t^k$. For $h = q(f)$, equivalently $q = h(\overline{f})$, continuity gives $\lambda_f^\times h = q(t) = h(\overline{f}(t))$.
Hence $\lambda_f^\times = C_{\overline{f}} \in \mathrm{Aut}_{\mathbb{C}\text{-alg}}(\mathscr{F})$.
($\Leftarrow$) Let $A := \lambda^\times \in \mathrm{Aut}_{\mathbb{C}\text{-alg}}(\mathscr{F})$ and set $f := A^{-1}t$. Since A^{-1} preserves (t) and $(t)^2$, $o(f) = 1$.
With $p_n := \lambda x^n$, $\langle f^k \mid p_n \rangle = \langle Af^k \mid x^n \rangle = \langle t^k \mid x^n \rangle = n!\delta_{n,k}$.
Thus (p_n) is associated with the delta series f, so $\lambda = \lambda_f$.

Structural Commentary: Umbral operators are formal coordinate changes on $\mathscr{F}$.
Key Insight: The adjoint remembers the delta series as the coordinate sent to t. □

Definition 3.6.25. *Umbral Composition*

For polynomial sequences (p_n) and (q_n) with $p_n(x) = \sum_{k=0}^n p_{n,k}x^k$, define $(p \circ_U q)_n(x) := \sum_{k=0}^n p_{n,k}q_k(x)$.

Theorem 3.6.26. *Group Law for Umbral Operators*
If f, g are delta series, then $\lambda_g \circ \lambda_f = \lambda_{f\circ g}$ and $\lambda_f^{-1} = \lambda_{\overline{f}}$.
Equivalently, associated sequences form a group under umbral composition.

Proof. Using $(AB)^\times = B^\times A^\times$ and Theorem 3.6.24, $(\lambda_g\lambda_f)^\times h = \lambda_f^\times \lambda_g^\times h = h(\overline{g}(\overline{f})) = h(\overline{f \circ g})$,
so $\lambda_g\lambda_f = \lambda_{f\circ g}$. Taking $g = \overline{f}$ gives $\lambda_f^{-1} = \lambda_{\overline{f}}$. If (p_n) and (q_n) are associated with f and g,
then $(\lambda_g\lambda_f)x^n = \lambda_g p_n = \sum_{k=0}^n p_{n,k}q_k = (p \circ_U q)_n$, so umbral composition corresponds to $f \circ g$.
Structural Commentary: The umbral group is the substitution group of delta series.
Key Insight: Composition of associated sequences is composition of delta series in reverse adjoint form. □

Proposition 3.6.27. *Sheffer Operators*
Let f, k be delta series and let $g, h \in \mathscr{F}^\times$.
If (s_n) is Sheffer for (g, f) and (p_n) is associated with f, then $s_n = g(t)^{-1}p_n$, and $\lambda_{g,f} = g(t)^{-1}\lambda_f$.
Moreover, $\lambda_{g,f} \circ \lambda_{h,k} = \lambda_{g\cdot(h\circ f),\, k\circ f}$ and $\lambda_{g,f}^{-1} = \lambda_{g^{-1}(\overline{f}),\, \overline{f}}$.

Proof. Theorem 3.6.16 gives $g(t)s_n = p_n$, hence $s_n = g(t)^{-1}p_n$ and $\lambda_{g,f} = g(t)^{-1}\lambda_f$.
For $a \in \mathscr{F}$, the adjoint substitution identity gives $\lambda_f a(t) = a(f(t))\lambda_f$.
Therefore $\lambda_{g,f}\lambda_{h,k} = g(t)^{-1}\lambda_f h(t)^{-1}\lambda_k = g(t)^{-1}h(f(t))^{-1}\lambda_{k\circ f}$.
This is $\lambda_{g\cdot(h\circ f),\, k\circ f}$. Taking $(h, k) = (g^{-1}(\overline{f}), \overline{f})$ gives the inverse formula.
Structural Commentary: Sheffer operators are invertible weights followed by umbral substitutions.
Key Insight: Their group law is the semidirect product law $(g, f)(h, k) = (g(h \circ f), k \circ f)$. □

Theorem 3.6.28. *Umbral Shift Derivations*
Let f be a delta series and let θ_f be its umbral shift.
Define the continuous derivation $\partial_f : \mathscr{F} \to \mathscr{F}$ by $\partial_f(f^k) = kf^{k-1}$, $k \geq 0$. Then $\theta_f^\times = \partial_f$, $\partial_f f = 1$.
If f, g are delta series, then $\partial_g = (\partial_g f)\partial_f$.

Proof. Let (p_n) be associated with f.
For $k, n \geq 0$, $\langle \theta_f^\times f^k \mid p_n \rangle = \langle f^k \mid p_{n+1} \rangle = (n+1)!\delta_{n+1,k} = \langle kf^{k-1} \mid p_n \rangle$.
Hence $\theta_f^\times f^k = kf^{k-1}$ for all k, so $\theta_f^\times = \partial_f$ by continuity.
For the chain rule, test on the coordinate powers f^k, which is $(\partial_g f)\partial_f(f^k) = k(\partial_g f)f^{k-1} = \partial_g(f^k)$.
Continuity gives the identity on $\mathscr{F}$.
Structural Commentary: An umbral shift is differentiation in its delta coordinate.
Key Insight: The associated basis turns shifting into formal differentiation. □

Theorem 3.6.29. *Umbral Shifts and Surjective Derivations*
For $\theta \in \mathcal{L}(\mathscr{P})$, θ is an umbral shift $\iff$ $\theta^\times$ is a surjective derivation of $\mathscr{F}$.

Proof.
($\Rightarrow$) If $\theta = \theta_f$, then $\theta^\times = \partial_f$ by Theorem 3.6.28. Since $\partial_f\left(\frac{f^{k+1}}{k+1}\right) = f^k$, $k \geq 0$, ∂_f is a surjective derivation.

($\Leftarrow$) Let $D := \theta^\times$ be a surjective derivation. Choose $u \in \mathscr{F}$ with $Du = 1$ and set $f := u - [t^0]u$.
Then $Df = 1$ and, since D lowers order by one, $o(f) = 1$. Let (p_n) be associated with f.
For $k, n \geq 0$, $\langle f^k \mid \theta p_n\rangle = \langle Df^k \mid p_n\rangle = \langle kf^{k-1} \mid p_n\rangle = \langle f^k \mid p_{n+1}\rangle$.
The functionals f^k separate $\mathscr{P}$, so $\theta p_n = p_{n+1}$ and $\theta = \theta_f$.

Structural Commentary: Surjective derivations are derivatives in delta coordinates.
Key Insight: Surjectivity chooses the coordinate whose derivative is 1. □

Theorem 3.6.30. *Recurrence and Transfer Formulas*
Let (p_n) be associated with the delta series f, and set $u(t) := f(t)/t \in \mathscr{F}^\times$.
Then $p_{n+1} = x[f'(t)]^{-1}p_n$, and, for $n \geq 1$, $p_n = f'(t)u(t)^{-n-1}x^n = xu(t)^{-n}x^{n-1}$.

Proof. By the chain rule for umbral shifts with coordinate t, $\theta_t = (\partial_t f)\theta_f = f'(t)\theta_f$.
Since $\theta_t = x$ and $\theta_f p_n = p_{n+1}$, $p_{n+1} = x[f'(t)]^{-1}p_n$. The transfer identity follows by applying this recurrence inductively after writing $f = tu$, hence $p_n = f'(t)u(t)^{-n-1}x^n = xu(t)^{-n}x^{n-1}$.
Structural Commentary: The recurrence generates (p_n) from $f'(t)^{-1}$.
Key Insight: Transfer formulas convert the difficult inverse coordinate problem into formal operator algebra. □

Corollary 3.6.31. *Sheffer Shift Formula*
If (s_n) is Sheffer for (g, f), then $\theta_{g,f} = g(t)^{-1}\theta_f g(t) = \left[x - \frac{g'(t)}{g(t)}\right]\frac{1}{f'(t)}$, and hence $s_{n+1} = \left(x - \frac{g'}{g}\right)(f')^{-1}s_n$.

Proof. Since $g(t)s_n$ is associated with f, $\theta_{g,f} = g^{-1}\theta_f g$. By Theorem 3.6.30, $\theta_f = x(f')^{-1}$.
Using the Pincherle commutator $g'(t) = g(t)x - xg(t)$, we get $g^{-1}x(f')^{-1}g = g^{-1}xg(f')^{-1} = \left(x - \frac{g'}{g}\right)(f')^{-1}$.
Applying this operator to s_n gives the recurrence.
Structural Commentary: The Sheffer shift is a conjugated associated shift.
Key Insight: The logarithmic derivative g'/g is the correction term. □

Definition 3.6.32. *Stirling and Exponential Polynomials*

(i) *Stirling numbers*: $S(n,k)$ are defined by $x^n = \sum_{k=0}^{n} S(n,k)(x)_k$.
(ii) *Exponential polynomials*: define $\phi_n(x) := \sum_{k=0}^{n} S(n,k)x^k$.
(iii) *Boundary convention*: $S(n,-1) = 0$ and $S(n,k) = 0$ for $k > n$.

Example 3.6.33. *Lower Factorials and Stirling Numbers*
For $f(t) = e^t - 1$, $f'(t) = e^t$, hence $p_{n+1}(x) = xe^{-t}p_n(x) = xp_n(x-1)$. With $p_0 = 1$, this gives $p_n = (x)_n$.
For $f(t) = \log(1+t)$, $\overline{f}(t) = e^t - 1$, and the associated sequence is (ϕ_n).
The inverse umbral relation is $x^n = \sum_{k=0}^{n} S(n,k)(x)_k$, and $\phi_n(x) = \sum_{k=0}^{n} S(n,k)x^k$.
The recurrence formula gives $\phi_{n+1}(x) = x(1+t)\phi_n(x) = x\big(\phi_n(x) + \phi_n'(x)\big)$,
so coefficient comparison yields $S(n+1,k) = S(n,k-1) + kS(n,k)$.

Definition 3.6.34. *q- and Logarithmic Umbral Variants*

(i) *Gaussian coefficient*: for $0 \leq k \leq n$, $\binom{n}{k}_q := \prod_{j=1}^{n}(1-q^j)\Big/\left(\prod_{j=1}^{k}(1-q^j)\prod_{j=1}^{n-k}(1-q^j)\right)$.
(ii) *q-umbral calculus*: the deformation $\binom{n}{k} \rightsquigarrow \binom{n}{k}_q$.
(iii) *Logarithmic umbral calculus*: the logarithmic-type variant based on harmonic logarithms.

Remark 3.6.35. *Further Directions*
Classical umbral methods extend to many polynomial families and operations:

(i) *Families*: lower factorial, Stirling, exponential, Gould, central factorial, Abel, Mittag–Leffler, Bessel, Bell, Hermite, Bernoulli, Euler, Laguerre, Poisson–Charlier, actuarial, Meixner, Pidduck, Narumi, Boole, Peters, Mahler, and Mott polynomials.
(ii) *Operations*: connection constants, duplication formulas, Lagrange inversion, cross sequences, Steffensen sequences, operational formulas, inverse relations, recurrence solving, and binomial convolution.
(iii) *Variants*: q-umbral and logarithmic umbral calculi replace or modify the binomial structure through the objects in Definition 3.6.34.

4 Matrix Analysis and Numerical Structure

Matrix analysis turns linear structure into quantitative matrix tools: invariants, decompositions, norms, conditioning, perturbation bounds, and numerical algorithms.

4.1 Matrix Classes and Invariants

Matrix classes package operator structure in coordinates. Trace, determinant, rank, positivity, idempotents, nilpotents, spectral radius, and structured products give computable invariants.

Definition 4.1.1. *Trace and Determinant of a Matrix*
For $A = [A_{j,k}] \in \mathrm{Mat}(n, n, \mathbb{F})$, the *trace* of A is $\mathrm{tr}(A) := A_{1,1} + \cdots + A_{n,n}$.
The *determinant* of A is $\det A := \sum_{\sigma \in \mathfrak{S}_n} \mathrm{sign}(\sigma) A_{\sigma(1),1} \cdots A_{\sigma(n),n}$.

Definition 4.1.2. *Gram Determinant*
Vectors $\mathbf{u}_1, \ldots, \mathbf{u}_n$ in inner-product space has *Gram determinant* $\mathrm{GramDet}(\mathbf{u}_1, \ldots, \mathbf{u}_n) := \det \mathrm{Gram}(\mathbf{u}_1, \ldots, \mathbf{u}_n)$.

Example 4.1.3. *Trace Functional and the Frobenius Inner Product*
On $\mathrm{Mat}(n, n, \mathbb{F})$ with the Frobenius inner product from Example 2.1.6, the Riesz vector of the trace functional $A \mapsto \mathrm{tr}(A)$ is I_n, since $\langle A, I_n \rangle_F = \sum_{j=1}^{n} a_{jj} = \mathrm{tr}(A)$.

Proposition 4.1.4. *Trace Is Cyclic for Two Factors*
If $A \in \mathrm{Mat}(m, n, \mathbb{F})$ *and* $B \in \mathrm{Mat}(n, m, \mathbb{F})$, *then* $\mathrm{tr}(AB) = \mathrm{tr}(BA)$.

Proof. Expand both traces by entries: $\mathrm{tr}(AB) = \sum_{j=1}^{m} \sum_{k=1}^{n} A_{j,k} B_{k,j} = \sum_{k=1}^{n} \sum_{j=1}^{m} B_{k,j} A_{j,k} = \mathrm{tr}(BA)$.
Structural Commentary: Cyclicity is just reindexing the same double sum.
Key Insight: The two double sums contain exactly the same scalar terms. □

Proposition 4.1.5. *Trace Arithmetic*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$ *and* $r \in \mathbb{F}$. *Then*
(i) $\mathrm{tr}(rA) = r\,\mathrm{tr}(A)$;
(ii) $\mathrm{tr}(A + B) = \mathrm{tr}(A) + \mathrm{tr}(B)$.

Proof. We prove the two identities separately.
(i) Summing the diagonal entries of rA gives $\mathrm{tr}(rA) = \sum_{j=1}^{n} r a_{jj} = r \sum_{j=1}^{n} a_{jj} = r\,\mathrm{tr}(A)$.
(ii) Summing the diagonal entries of $A+B$ gives $\mathrm{tr}(A+B) = \sum_{j=1}^{n} (a_{jj}+b_{jj}) = \sum_{j=1}^{n} a_{jj} + \sum_{j=1}^{n} b_{jj} = \mathrm{tr}(A)+\mathrm{tr}(B)$.
Structural Commentary: Trace is linear because it is the sum of diagonal entries.
Key Insight: Scalar multiplication and addition pass through the diagonal sum. □

Proposition 4.1.6. *Cyclic Products*
Let $m \geq 2$, $A_j \in \mathrm{Mat}(n_j, n_{j+1}, \mathbb{F})$ *for* $j = 1, \ldots, m$, $n_{m+1} = n_1$, $P_{p,q} := A_p A_{p+1} \cdots A_q$.
Then $\mathrm{tr}(P_{1,m}) = \mathrm{tr}(A_m P_{1,m-1})$. *Consequently, for* $A, B, C \in \mathrm{Mat}(n, n, \mathbb{F})$, $\mathrm{tr}(ABC) = \mathrm{tr}(BCA) = \mathrm{tr}(CAB)$

Proof. Since $P_{1,m-1} \in \mathrm{Mat}(n_1, n_m, \mathbb{F})$ and $A_m \in \mathrm{Mat}(n_m, n_1, \mathbb{F})$,
Proposition 4.1.4 gives $\mathrm{tr}(P_{1,m}) = \mathrm{tr}(P_{1,m-1} A_m) = \mathrm{tr}(A_m P_{1,m-1})$.
Applying this cyclic move twice to ABC gives the three-factor identity.
Structural Commentary: Trace preserves cyclic order, not arbitrary order.
Key Insight: One cyclic move is $\mathrm{tr}(XY) = \mathrm{tr}(YX)$ with $X = P_{1,m-1}$ and $Y = A_m$. □

Example 4.1.7. *Cyclicity Is Not Arbitrary Commutativity*
Let $A = E_{1,2}$, $B = E_{2,1}$, and $C = E_{1,1}$ in $\mathrm{Mat}(2, 2, \mathbb{F})$. Then $\mathrm{tr}(ABC) = 1$, while $\mathrm{tr}(ACB) = 0$.
Thus trace permits cyclic rotations of factors, but not arbitrary rearrangements.

Theorem 4.1.8. *Basic Determinant Laws*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$.
(i) $\det I_n = 1$.
(ii) $\det(A)$ *is linear as a function of each column of* A *when the other columns are fixed.*
(iii) Swapping two columns negates the determinant; equal columns force $\det A = 0$.
(iv) If A *is triangular, then* $\det A = A_{1,1} \cdots A_{n,n}$.
(v) $\det(AB) = \det(A) \det(B)$.
(vi) A *is invertible if and only if* $\det A \neq 0$. *In that case* $\det(A^{-1}) = (\det A)^{-1}$.
(vii) $\det(A^t) = \det(A)$.

Proof.

(i) $\det I_n = \sum_{\sigma\in\mathfrak{S}_n} \operatorname{sign}(\sigma)(I_n)_{\sigma(1),1}\cdots(I_n)_{\sigma(n),n} = 1$ as product is nonzero only when $\sigma(k)=k\ \forall k$.

(ii) For fixed k, $\det(A_{\cdot,1},\dots,\mathbf{x},\dots,A_{\cdot,n}) = \sum_{\sigma\in\mathfrak{S}_n} \operatorname{sign}(\sigma)x_{\sigma(k)}\prod_{\ell\neq k} A_{\sigma(\ell),\ell}$, which is linear in $\mathbf{x}$.

(iii) If $B_{\cdot,k} = A_{\cdot,\tau(k)}$ for transposition τ, then $\det B = \sum_{\rho\in\mathfrak{S}_n} \operatorname{sign}(\rho\tau)\prod_{k=1}^{n} A_{\rho(k),k} = -\det A$.
If the swapped columns are equal, then $B=A$, so $\det A = -\det A$ and $\det A = 0$.

(iv) If A is upper triangular and $\prod_{k=1}^{n} A_{\sigma(k),k}\neq 0$, then $\sigma(k)\le k$ for all k, hence $\sigma(k)=k$ for all k.
Thus $\det A = A_{1,1}\cdots A_{n,n}$. If A is lower triangular, nonzero products require $\sigma(k)\ge k$ for all k, again forcing $\sigma(k)=k$ for all k, so the same diagonal product formula holds.

(v) If F is alternating and multilinear in columns, then expanding each column in the standard basis gives $F(A_{\cdot,1},\dots,A_{\cdot,n}) = \sum_{i_1,\dots,i_n} A_{i_1,1}\cdots A_{i_n,n}F(\mathbf{e}_{i_1},\dots,\mathbf{e}_{i_n}) = \sum_{\sigma\in\mathfrak{S}_n}\operatorname{sign}(\sigma)F(I_n)A_{\sigma(1),1}\cdots A_{\sigma(n),n}$.
The terms with repeated i_j vanish by alternation, and the surviving ordered bases differ from $(\mathbf{e}_1,\dots,\mathbf{e}_n)$ by sign of permutation. For $H(B)=\det(AB)$, map H is alternating multilinear and $H(I_n)=\det A$.
Hence $\det(AB)=H(B)=H(I_n)\det B=\det(A)\det(B)$.

(vi) If A is invertible, then $1=\det I_n=\det(AA^{-1})=\det(A)\det(A^{-1})$.
If $A_{\cdot,k}=\sum_{\ell\neq k} c_\ell A_{\cdot,\ell}$, then $\det A=\sum_{\ell\neq k} c_\ell\det(A_{\cdot,1},\dots,A_{\cdot,\ell},\dots,A_{\cdot,\ell},\dots,A_{\cdot,n})=0$.
Thus $\det A\neq 0$ iff the columns of A are independent, iff A is invertible.

(vii) Using the permutation definition and replacing σ by σ^{-1},
$\det(A^t)=\sum_{\sigma\in\mathfrak{S}_n}\operatorname{sign}(\sigma)A_{1,\sigma(1)}\cdots A_{n,\sigma(n)}=\sum_{\tau\in\mathfrak{S}_n}\operatorname{sign}(\tau)A_{\tau(1),1}\cdots A_{\tau(n),n}=\det(A)$; $\operatorname{sign}(\tau^{-1})=\operatorname{sign}(\tau)$.

Structural Commentary: The determinant is the normalized alternating multilinear measure of columns.
Key Insight: Multilinearity expands columns; alternation kills repeated directions. □

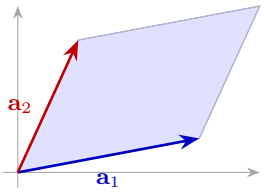


Figure 73: Determinant is signed column-volume; swaps reverse orientation and dependence gives zero volume.

Method 4.1.9. *Determinant by Structure*
To compute or simplify a determinant:

(i) first look for triangular or block triangular form;
(ii) use column linearity to split complicated columns only when it reduces structure;
(iii) use column swaps to move useful pivots into position, tracking a factor of -1 for each swap;
(iv) detect zero determinants from dependent columns or repeated columns;
(v) for products or changes of basis, use multiplicativity instead of expanding permutations.

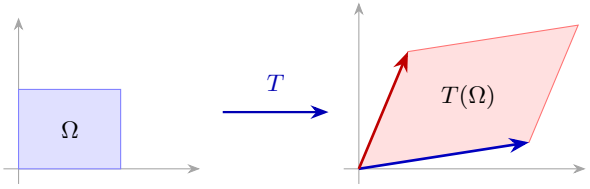


Figure 74: A linear map scales oriented volume by its determinant; singular maps collapse volume to zero.

Lemma 4.1.10. *Identity Blocks Do Not Change Determinant*
Let $A\in\operatorname{Mat}(m,m,\mathbb{F})$ *and let* $n\ge 0$. *Then* $\det(I_n\oplus A)=\det(A)=\det(A\oplus I_n)$.

Proof. In the permutation formula for $I_n\oplus A$, a nonzero term must satisfy $\sigma(k)=k$ for $1\le k\le n$. Remaining choices are exactly permutations of $\{n{+}1,\dots,n{+}m\}$, after relabeling, gives $\det(A)$. Hence $\det(I_n\oplus A)=\det(A)$.
The same argument applied to the last n coordinates gives $\det(A\oplus I_n)=\det(A)$.
The case $n=0$ is the empty identity block convention.

Structural Commentary: Identity blocks add directions without changing volume.
Key Insight: In the determinant expansion, the identity block forces the matching coordinates to be fixed. □

Theorem 4.1.11. *Determinant of an Upper Block-Triangular Matrix*
Let $A \in \mathrm{Mat}(r,r,\mathbb{F})$, $B \in \mathrm{Mat}(r,s,\mathbb{F})$, *and* $D \in \mathrm{Mat}(s,s,\mathbb{F})$. *Then* $\det \begin{bmatrix} A & B \\ 0 & D \end{bmatrix} = \det(A)\det(D)$.

Proof. Factor $\begin{bmatrix} A & B \\ 0 & D \end{bmatrix} = \begin{bmatrix} I_r & 0 \\ 0 & D \end{bmatrix} \begin{bmatrix} I_r & B \\ 0 & I_s \end{bmatrix} \begin{bmatrix} A & 0 \\ 0 & I_s \end{bmatrix}$.
The first and last factors have determinants $\det(D)$ and $\det(A)$ by Lemma 4.1.10.
The middle factor is upper triangular with every diagonal entry equal to 1, so its determinant is 1.
Multiplicativity of determinant gives the formula.
Structural Commentary: Block upper triangular matrices separate determinant contribution into diagonal blocks. The off-diagonal block does not affect the determinant.
Key Insight: Factor the matrix into two block diagonal pieces and one unit upper triangular piece. □

Proposition 4.1.12. *Schur Complement Determinant Formula*
Let $M = \begin{bmatrix} A & B \\ C & D \end{bmatrix}$, $A \in \mathrm{Mat}(r,r,\mathbb{F})$, $D \in \mathrm{Mat}(s,s,\mathbb{F})$.
If A *is invertible, then the Schur complement of* A *in* M *is* $M/A := D - CA^{-1}B$, *and* $\det(M) = \det(A)\det(M/A)$.

Proof. Block Gaussian elimination gives $\begin{bmatrix} I_r & 0 \\ -CA^{-1} & I_s \end{bmatrix} \begin{bmatrix} A & B \\ C & D \end{bmatrix} = \begin{bmatrix} A & B \\ 0 & D - CA^{-1}B \end{bmatrix}$.
The elimination matrix is lower block-triangular with identity diagonal blocks, so its determinant is 1. Take determinants, apply Theorem 4.1.11 to right-hand block triangular matrix, $\det(M) = \det(A)\det(D - CA^{-1}B)$.
Structural Commentary: Schur complementation is block Gaussian elimination.
Key Insight: Eliminating C splits the determinant into $\det(A)$ times the remaining block determinant. □

Proposition 4.1.13. *Schur Complement of the Lower-Right Block*
Let $M = \begin{bmatrix} A & B \\ C & D \end{bmatrix}$, $A \in \mathrm{Mat}(r,r,\mathbb{F})$, $D \in \mathrm{Mat}(s,s,\mathbb{F})$.
If D *is invertible, then the Schur complement of* D *in* M *is* $M/D := A - BD^{-1}C$, *and* $\det(M) = \det(D)\det(M/D)$.

Proof. Block Gaussian elimination gives $\begin{bmatrix} I_r & -BD^{-1} \\ 0 & I_s \end{bmatrix} \begin{bmatrix} A & B \\ C & D \end{bmatrix} = \begin{bmatrix} A - BD^{-1}C & 0 \\ C & D \end{bmatrix}$.
The elimination matrix has determinant 1. The transpose of the matrix on the right is upper block-triangular, and $\det(X^t) = \det(X)$ by Theorem 4.1.8. Therefore Theorem 4.1.11 gives $\det(M) = \det(A - BD^{-1}C)\det(D)$.
Structural Commentary: This is the lower-right version of Schur complementation. Eliminating the upper-right block isolates the determinant contribution of D.
Key Insight: Eliminate by blocks, then read the determinant from the diagonal blocks. □

Theorem 4.1.14. *Sylvester Determinant Identity*
Let $A \in \mathrm{Mat}(n,m,\mathbb{F})$ *and* $B \in \mathrm{Mat}(m,n,\mathbb{F})$. *Then* $\det(I_n + AB) = \det(I_m + BA)$.

Proof. Let $M = \begin{bmatrix} I_n & -A \\ B & I_m \end{bmatrix}$.
Using Proposition 4.1.13 with lower-right block I_m gives $\det(M) = \det(I_n + AB)$. Using Proposition 4.1.12 with upper-left block I_n gives $\det(M) = \det(I_m + BA)$. Hence the two determinants are equal.
Structural Commentary: Two Schur complements of one block matrix give Sylvester's identity.
Key Insight: Compare the same determinant after eliminating different identity blocks. □

Theorem 4.1.15. *Cauchy–Binet Formula*
Let $A \in \mathrm{Mat}(k,n,\mathbb{F})$, $B \in \mathrm{Mat}(n,k,\mathbb{F})$, *and* $k \le n$. *For* $J = \{j_1 < \cdots < j_k\} \subseteq \{1,\dots,n\}$, *let* A^J *use the columns of* A *in* J *and* B_J *use the rows of* B *in* J.
Then $\det(AB) = \sum_J \det(A^J)\det(B_J)$, *where* J *ranges over all* k*-subsets.*

Proof. Put $S(A,B) := \sum_J \det(A^J)\det(B_J)$. For an elementary $n \times n$ matrix E, determinant multilinearity and row/column operation rules give $S(A,B) = S(AE, E^{-1}B)$, and $AB = (AE)(E^{-1}B)$.
If $\operatorname{rank} A < k$, both sides vanish. Otherwise reduce A by elementary column operations to $A' = [C\ 0]$ with C invertible, and transform B oppositely to B'.
Then $AB = A'B'$ and $\det(A'B') = \det(C)\det(B'_{J_0}) = S(A',B')$, $J_0 = \{1,\dots,k\}$, as every other k-minor of A' is zero. Hence $\det(AB) = S(A,B)$.

Structural Commentary: Cauchy–Binet is determinant multiplicativity for rectangular middle dimensions.
Key Insight: The middle space contributes one k-coordinate shadow at a time. □

Proposition 4.1.16. *Bordered Determinant Formulas*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$, let $\mathbf{x},\mathbf{y} \in \mathbb{F}^n$, and let $c \in \mathbb{F}$.

(i) If $c \neq 0$, then $\det \begin{bmatrix} c & \mathbf{x}^t \\ \mathbf{y} & A \end{bmatrix} = c \det\left(A - \frac{1}{c}\mathbf{y}\mathbf{x}^t\right)$.

(ii) If A is invertible, then $\det \begin{bmatrix} A & \mathbf{x} \\ \mathbf{y}^t & c \end{bmatrix} = \left(c - \mathbf{y}^t A^{-1}\mathbf{x}\right)\det(A)$

Proof.

(i) This is Proposition 4.1.12 with the scalar block c in the upper-left corner.
(ii) This is Proposition 4.1.12 with A as the upper-left block and $D = c$.

Structural Commentary: Bordered determinants are Schur complements with scalar borders.
Key Insight: Treat the border as a 1×1 block and apply the Schur complement formula. □

Definition 4.1.17. *Minors, Cofactors, and Adjugate*
For $A \in \mathrm{Mat}(n,n,\mathbb{F})$, let M_{ij} delete row i and column j from A; set $\det(\emptyset_{0\times 0}) = 1$. Define $C_{ij} := (-1)^{i+j}\det(M_{ij})$. Then $C = [C_{ij}]$ is the *cofactor matrix*, and $\mathrm{adj}(A) := C^t$ is the *adjugate*.

Lemma 4.1.18. *Cofactor Expansion*
For $A = [a_{ij}] \in \mathrm{Mat}(n,n,\mathbb{F})$, $\det A = \sum_{k=1}^{n} a_{ik}C_{ik}$ *and* $\det A = \sum_{k=1}^{n} a_{kj}C_{kj}$.

Proof. Using $\det A = \sum_{\sigma\in S_n} \mathrm{sgn}(\sigma) \prod_{r=1}^{n} a_{r,\sigma(r)}$, group terms by $\sigma(i) = k$.
For fixed k, the remaining permutation sum equals $(-1)^{i+k}\det(M_{ik}) = C_{ik}$, giving the row formula. Applying this to A^t and using Theorem 4.1.8 gives the column formula.
Structural Commentary: Cofactor expansion is the big determinant formula collected around one row or column.
Key Insight: Fix one entry; the rest is its minor. □

Remark 4.1.19. *Adjugate Form of the Bordered Formula*
Proposition 4.1.16 gives the adjugate form $\det \begin{bmatrix} A & \mathbf{x} \\ \mathbf{y}^t & c \end{bmatrix} = c\det(A) - \mathbf{y}^t \mathrm{adj}(A)\mathbf{x}$. No invertibility of A is required.

Theorem 4.1.20. *Adjugate Identity and Cramer's Rule*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$ and let $\mathbf{b} \in \mathbb{F}^n$. Then $A\,\mathrm{adj}(A) = \mathrm{adj}(A)A = (\det A)I_n$. If $\det A \neq 0$, then, for $A\mathbf{x} = \mathbf{b}$, $A^{-1} = \frac{1}{\det A}\mathrm{adj}(A)$ and $x_j = \frac{\det A_j(\mathbf{b})}{\det A}$ for $1 \le j \le n$, where $A_j(\mathbf{b})$ replaces column j of A by $\mathbf{b}$.

Proof. Since $(\mathrm{adj}\,A)_{kj} = C_{jk}$, $(A\,\mathrm{adj}\,A)_{ij} = \sum_k a_{ik}C_{jk}$.
For $i = j$ this is $\det A$ by Lemma 4.1.18; for $i \neq j$ it is the cofactor expansion of a matrix with rows i and j equal, hence 0. Thus $A\,\mathrm{adj}(A) = (\det A)I_n$. The same argument applied to A^t gives $\mathrm{adj}(A)A = (\det A)I_n$. If $\det A \neq 0$, division gives the inverse formula. Finally, $(\mathrm{adj}\,A\,\mathbf{b})_j = \sum_i C_{ij}b_i = \det A_j(\mathbf{b})$, by expansion down column j, so $x = A^{-1}\mathbf{b}$ gives Cramer's rule.
Structural Commentary: Cofactors assemble all one-column replacements of A into a single inverse formula.
Key Insight: Off-diagonal products vanish because they are determinants with two equal rows or columns. □

Theorem 4.1.21. *Trace and Determinant Are Similarity Invariants*
If $A, B \in \mathrm{Mat}(n,n,\mathbb{F})$ are similar, then $\mathrm{tr}(A) = \mathrm{tr}(B)$, $\det(A) = \det(B)$.

Proof. Write $B = PAP^{-1}$ with P invertible.
Proposition 4.1.4 gives $\mathrm{tr}(B) = \mathrm{tr}\big(P(AP^{-1})\big) = \mathrm{tr}\big((AP^{-1})P\big) = \mathrm{tr}(A)$.
Theorem 4.1.8 gives $\det(B) = \det(P)\det(A)\det(P^{-1}) = \det(A)$.
Structural Commentary: Trace and determinant ignore the coordinate system used to write an operator.
Key Insight: Conjugating by P inserts P and P^{-1}, which cancel inside trace and determinant. □

Principle 4.1.22. *Trace and Determinant Are Coordinate-Free*
Trace and determinant are first computed from a matrix, but their real structural meaning is operator-level: changing basis conjugates the matrix, and conjugation preserves both quantities.

Theorem 4.1.23. *Trace and Determinant of an Operator Are Basis-Independent*
Let $\mathcal{V}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V})$, and let β, γ be bases of $\mathcal{V}$.
Then $\mathrm{tr}\big([T]_{\beta\leftarrow\beta}\big) = \mathrm{tr}\big([T]_{\gamma\leftarrow\gamma}\big)$, $\det\big([T]_{\beta\leftarrow\beta}\big) = \det\big([T]_{\gamma\leftarrow\gamma}\big)$.

Proof. By Theorem 1.3.105, the two matrices are similar: $[T]_{\gamma\leftarrow\gamma} = C_{\gamma\leftarrow\beta}[T]_{\beta\leftarrow\beta}C_{\beta\leftarrow\gamma}$.
Theorem 4.1.21 now gives equality of traces and determinants.
Structural Commentary: Basis changes conjugate matrices; trace and determinant remain operator invariants.
Key Insight: Similar matrices have the same trace and determinant. □

Definition 4.1.24. *Trace and Determinant of an Operator*
Let $\mathcal{V}$ be finite-dim, let $T \in \mathcal{L}(\mathcal{V})$. Choose any basis β of $\mathcal{V}$. Define $\operatorname{tr} T := \operatorname{tr}\big([T]_{\beta\leftarrow\beta}\big)$, $\det T := \det\big([T]_{\beta\leftarrow\beta}\big)$.
Theorem 4.1.23 shows these definitions do not depend on chosen basis.

Proposition 4.1.25. *Unitary Matrices and Isometries Have Determinant of Modulus One*
Let $U \in \operatorname{Mat}(n,n,\mathbb{F})$ be unitary. Then $|\det U| = 1$. More generally, if $S \in \mathcal{L}(\mathcal{V})$ is an isometry on a finite-dimensional inner-product space over $\mathbb{F} \in \{\mathbb{R},\mathbb{C}\}$, then $|\det S| = 1$.

Proof. The permutation definition gives $\det(U^*) = \overline{\det U}$. By Theorem 4.1.8, taking determinants in $U^*U = I_n$ gives $1 = \det(U^*)\det(U) = \overline{\det U}\det U = |\det U|^2$. Hence $|\det U| = 1$.
For an isometry S, choose an orthonormal basis and let U be the matrix of S in that basis. Then U is unitary by Theorem 2.4.26, and $\det S = \det U$.
Structural Commentary: In orthonormal coordinates, isometries are unitary or orthogonal, so the determinant calculation is the matrix one.
Key Insight: Determinants turn $U^*U = I$ into $|\det U|^2 = 1$. □

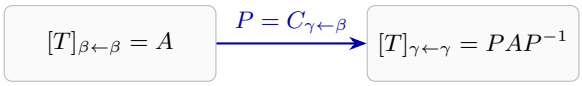


Figure 75: Basis changes conjugate matrices; trace and determinant are unchanged.

Proposition 4.1.26. *Trace and Determinant from q_T*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dim over $\mathbb{F} \in \{\mathbb{R},\mathbb{C}\}$, let $T \in \mathcal{L}(\mathcal{V})$, $q_T(z) = z^n + c_{n-1}z^{n-1} + \cdots + c_1 z + c_0$, $n = \dim \mathcal{V}$. Then $\operatorname{tr} T = -c_{n-1}$ and $\det T = (-1)^n c_0$. If $\mathbb{F} = \mathbb{C}$ and $\lambda_1,\ldots,\lambda_n$ list the eigenvalues of T with algebraic multiplicity, then $\operatorname{tr} T = \lambda_1 + \cdots + \lambda_n$, $\det T = \lambda_1 \cdots \lambda_n$.

Proof. Choose a basis β and set $A = [T]_{\beta\leftarrow\beta}$.
Then $q_T = p_A$, and by Definition 4.1.24, $\operatorname{tr} T = \operatorname{tr} A$ and $\det T = \det A$.
Theorem 3.2.8 gives $\operatorname{tr} T = -c_{n-1}$ and $\det T = (-1)^n c_0$.
If $\mathbb{F} = \mathbb{C}$, then comparing $q_T(z) = (z-\lambda_1)\cdots(z-\lambda_n)$ with the displayed form of q_T gives the formulas.
Structural Commentary: Characteristic polynomial stores trace and determinant in first and last coefficients.
Key Insight: Compare coefficients of $\det(zI - T)$. □

Proposition 4.1.27. *Trace and Determinant from Eigenvalues*
Let $\lambda_1,\ldots,\lambda_n$ list the eigenvalues of $A \in \operatorname{Mat}(n,n,\mathbb{C})$ according to m_A. Then $\operatorname{tr} A = \sum_{j=1}^{n} \lambda_j$ and $\det A = \prod_{j=1}^{n} \lambda_j$.

Proof. By Fundamental Theorem of Algebra, $p_A(z) = \prod_{j=1}^{n}(z-\lambda_j) = z^n - \left(\sum_{j=1}^{n}\lambda_j\right) z^{n-1} + \cdots + (-1)^n \prod_{j=1}^{n}\lambda_j$.
Theorem 3.2.8 gives $p_A(z) = z^n - (\operatorname{tr} A)z^{n-1} + \cdots + (-1)^n \det A$. Uniqueness of coefficients gives the identities.
Structural Commentary: Coefficients of p_A encode symmetric functions of the eigenvalues.
Key Insight: Trace is the first root sum and determinant is the full root product. □

Proposition 4.1.28. *Transpose, Conjugate, and Adjoint Spectra*
Let $A \in \operatorname{Mat}(n,n,\mathbb{C})$. Then $p_{A^t}(z) = p_A(z)$, $p_{\overline{A}}(z) = \overline{p_A(\overline{z})}$, and $p_{A^*}(z) = \overline{p_A(\overline{z})}$.
Consequently, if $\lambda \in \operatorname{spec}(A)$, then $m_{A^t}(\lambda) = m_A(\lambda)$ and $m_{\overline{A}}(\overline{\lambda}) = m_{A^*}(\overline{\lambda}) = m_A(\lambda)$.

Proof. Use determinant identities:

$$p_{A^t}(z) = \det(zI_n - A^t) = \det((zI_n - A)^t) = p_A(z),$$
$$p_{\overline{A}}(z) = \det(zI_n - \overline{A}) = \det(\overline{\overline{z}I_n - A}) = \overline{\det(\overline{z}I_n - A)} = \overline{p_A(\overline{z})}.$$

Since $A^* = \overline{A}^{\,t}$, the first identity applied to $\overline{A}$ gives the formula for A^*.
Multiplicity is preserved because conjugating a polynomial sends a factor $(z-\lambda)^k$ to $(z-\overline{\lambda})^k$.
Structural Commentary: Transpose preserves p_A; conjugation conjugates its roots.
Key Insight: The adjoint combines conjugation with transposition, so its spectrum is conjugated. □

Proposition 4.1.29. *Transpose Preserves Spectrum and Geometric Multiplicity*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$ *and* $\lambda \in \mathbb{F}$. *Then* $\lambda \in \mathrm{spec}(A) \iff \lambda \in \mathrm{spec}(A^t)$, *and* $\dim E_\lambda(A) = \dim E_\lambda(A^t)$.

Proof. Since $(A-\lambda I_n)^t = A^t - \lambda I_n$, Theorem 1.3.74 gives $\mathrm{rank}(A-\lambda I_n) = \mathrm{rank}(A^t - \lambda I_n)$.
Matrix rank-nullity then gives $\dim E_\lambda(A) = n - \mathrm{rank}(A-\lambda I_n) = n - \mathrm{rank}(A^t - \lambda I_n) = \dim E_\lambda(A^t)$.
($\Rightarrow$) If $\lambda \in \mathrm{spec}(A)$, then $E_\lambda(A) \neq \{\mathbf{0}\}$, so $E_\lambda(A^t) \neq \{\mathbf{0}\}$ and $\lambda \in \mathrm{spec}(A^t)$.
($\Leftarrow$) The same argument with A and A^t interchanged gives the converse.

Structural Commentary: Transposition preserves rank, so shifted kernels have the same dimension.
Key Insight: Spectrum and geometric multiplicity are kernel facts for $A - \lambda I_n$. □

Corollary 4.1.30. *Square Matrix Algebraic Invertibility Dictionary*
Let $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n] \in \mathrm{Mat}(n,n,\mathbb{F})$. *The following are equivalent:*

(i) A *is invertible;*
(ii) $\mathbf{a}_1, \ldots, \mathbf{a}_n$ *form a basis of* $\mathbb{F}^n$*;*
(iii) $\mathbf{a}_1, \ldots, \mathbf{a}_n$ *are linearly independent;*
(iv) $\mathbf{a}_1, \ldots, \mathbf{a}_n$ *span* $\mathbb{F}^n$*;*
(v) for every $\mathbf{b} \in \mathbb{F}^n$, *the system* $A\mathbf{x} = \mathbf{b}$ *has a unique solution;*
(vi) the homogeneous system $A\mathbf{x} = \mathbf{0}$ *has only the zero solution.*

Proof.

(*i*) $\Rightarrow$ (*ii*) By Theorem 1.3.86, the columns of an invertible matrix form a basis.
(*ii*) $\Rightarrow$ (*iii*) A basis is linearly independent.
(*iii*) $\Rightarrow$ (*iv*) In $\mathbb{F}^n$, n linearly independent vectors form a basis, hence span.
(*iv*) $\Rightarrow$ (*v*) In $\mathbb{F}^n$, n spanning vectors form a basis, so every $\mathbf{b}$ has unique coordinates $\mathbf{x}$ in $A\mathbf{x} = \mathbf{b}$.
(*v*) $\Rightarrow$ (*vi*) Apply (v) to $\mathbf{b} = \mathbf{0}$; since $\mathbf{x} = \mathbf{0}$ is a solution, it is the only one.
(*vi*) $\Rightarrow$ (*i*) Condition (vi) says the columns are independent.
Hence they form a basis of $\mathbb{F}^n$, and Corollary 1.3.93 gives invertibility.

Structural Commentary: Square invertibility is column basishood.
Key Insight: Invertible columns give unique coordinates. □

Corollary 4.1.31. *Nonsingularity by Rank, Determinant, and Spectrum*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$. *The following are equivalent:*

(i) A *is invertible;*
(ii) $\mathrm{rank}(A) = n$*;*
(iii) $\det A \neq 0$*;*
(iv) $0 \notin \mathrm{spec}(A)$.

Over $\mathbb{C}$, *this is also equivalent to all eigenvalues of* A *being nonzero.*

Proof.

(*i*) $\Rightarrow$ (*ii*) This is Corollary 1.3.95.
(*ii*) $\Rightarrow$ (*iii*) Full rank gives invertibility by Corollary 1.3.95, hence $\det A \neq 0$ by Theorem 4.1.8.
(*iii*) $\Rightarrow$ (*iv*) If $\det A \neq 0$, then A is invertible by Theorem 4.1.8; hence $0 \notin \mathrm{spec}(A)$ by Corollary 3.1.9.
(*iv*) $\Rightarrow$ (*i*) This is Corollary 3.1.9.

Over $\mathbb{C}$, Theorem 3.2.8 identifies $\lambda_1, \ldots, \lambda_n$ as the roots of p_A, so (iv) is equivalent to $\lambda_j \neq 0$ for all j.
Structural Commentary: Rank, determinant, and spectrum detect singularity.
Key Insight: Singularity can be seen as rank loss, zero determinant, or the zero eigenvalue. □

Method 4.1.32. *Trace–Determinant Invariant Method*
To get eigenvalue data without solving for eigenvalues:

(i) choose a basis in which $[T]_{\beta\leftarrow\beta}$ is easiest to compute;
(ii) compute $\mathrm{tr}\, T$ and $\det T$ from this matrix;
(iii) if the eigenvalues are $\lambda_1, \ldots, \lambda_n$, conclude $\lambda_1 + \cdots + \lambda_n = \mathrm{tr}\, T$, and $\lambda_1 \cdots \lambda_n = \det T$;
(iv) for equations $ST - TS = A$, first test $\mathrm{tr}\, A = 0$; otherwise no solution exists.

Theorem 4.1.33. *Rank Bounds for Products and Concatenations*
Let $A \in \mathrm{Mat}(m,k,\mathbb{F})$, $B \in \mathrm{Mat}(k,n,\mathbb{F})$, *and* $C \in \mathrm{Mat}(m,p,\mathbb{F})$. *Then* $\mathrm{rank}(AB) \leq \min(\mathrm{rank}(A), \mathrm{rank}(B))$ *and* $\mathrm{rank}(AB) = \mathrm{rank}(A)$ *if and only if* $\mathrm{col}(AB) = \mathrm{col}(A)$.
Also, $\max(\mathrm{rank}(A), \mathrm{rank}(C)) \leq \mathrm{rank}\begin{bmatrix} A & C \end{bmatrix}$ *and* $\mathrm{rank}\begin{bmatrix} A & C \end{bmatrix} = \mathrm{rank}(A)$ *if and only if* $\mathrm{col}(A) = \mathrm{col}(A) + \mathrm{col}(C)$.

Proof. Each column of AB is a linear combination of the columns of A, so $\mathrm{col}(AB) \subseteq \mathrm{col}(A)$. Therefore $\mathrm{rank}(AB) \leq \mathrm{rank}(A)$; since both spaces are finite-dimensional, equality holds exactly when $\mathrm{col}(AB) = \mathrm{col}(A)$. Applying this same argument to $(AB)^t = B^t A^t$ and using Theorem 1.3.74 gives $\mathrm{rank}(AB) \leq \mathrm{rank}(B)$.
For the block matrix, its columns are precisely the columns of A and C, hence $\mathrm{col}\begin{bmatrix} A & C \end{bmatrix} = \mathrm{col}(A) + \mathrm{col}(C)$. This space contains both $\mathrm{col}(A)$ and $\mathrm{col}(C)$, giving the lower bound. Equality with $\mathrm{rank}(A)$ occurs exactly when the containing subspace has the same dimension as $\mathrm{col}(A)$, equivalently when $\mathrm{col}(A) + \mathrm{col}(C) = \mathrm{col}(A)$.
Structural Commentary: Products can only reduce accessible column directions, while horizontal concatenation can only add column directions.
Key Insight: Rank tracks whether a matrix operation preserves, loses, or adds column-space dimension. □

Theorem 4.1.34. *Sylvester Rank Inequality*
Let $A \in \mathrm{Mat}(m, k, \mathbb{F})$ and $B \in \mathrm{Mat}(k, n, \mathbb{F})$. Then $\mathrm{rank}(A) + \mathrm{rank}(B) - k \leq \mathrm{rank}(AB)$.

Proof. Let $R = \mathrm{col}(B) \subseteq \mathbb{F}^k$. The columns of AB are A applied to the columns of B, so $\mathrm{col}(AB) = A(R)$.
Apply rank-nullity to the restriction $T_A|_R : R \to \mathbb{F}^m$. Its range is $A(R) = \mathrm{col}(AB)$ and its kernel is $R \cap \ker(A)$.
Hence $\mathrm{rank}(AB) = \dim(A(R)) = \dim(R) - \dim(R \cap \ker(A)) \geq \mathrm{rank}(B) - \dim(\ker(A))$.
Since A has k columns, Corollary 1.3.73 gives $\dim(\ker(A)) = k - \mathrm{rank}(A)$. Substitution gives desired inequality.
Structural Commentary: Multiplication by A can lose only the part of $\mathrm{col}(B)$ lying in $\ker(A)$. Sylvester's inequality is this loss estimate in dimensions.
Key Insight: The shared middle dimension k bounds how much rank can disappear in the product. □

Theorem 4.1.35. *Frobenius Rank Inequality*
Let $A \in \mathrm{Mat}(m, k, \mathbb{F})$, $B \in \mathrm{Mat}(k, p, \mathbb{F})$, and $C \in \mathrm{Mat}(p, n, \mathbb{F})$.
Then $\mathrm{rank}(AB) + \mathrm{rank}(BC) \leq \mathrm{rank}(B) + \mathrm{rank}(ABC)$.

Proof. Let $R = \mathrm{col}(B)$ and $S = \mathrm{col}(BC)$. Since every column of BC is a linear combination of columns of B, we have $S \subseteq R$. The induced map T_A maps R onto $\mathrm{col}(AB)$ and maps S onto $\mathrm{col}(ABC)$.
Rank-nullity on the restrictions $T_A|_R$ and $T_A|_S$ gives

$$\mathrm{rank}(AB) = \dim R - \dim(R \cap \ker(A)), \quad \mathrm{rank}(ABC) = \dim S - \dim(S \cap \ker(A)).$$

Since $S \cap \ker(A) \subseteq R \cap \ker(A)$, we have $\dim(S \cap \ker(A)) \leq \dim(R \cap \ker(A))$.
Thus $\dim R - \dim(R \cap \ker(A)) + \dim S \leq \dim R + \dim S - \dim(S \cap \ker(A))$.
Substituting $\dim R = \mathrm{rank}(B)$ and $\dim S = \mathrm{rank}(BC)$ gives the stated inequality.
Structural Commentary: Frobenius tracks A on nested column spaces.
Key Insight: The smaller column space cannot lose more dimensions to $\ker(A)$ than the larger one. □

Corollary 4.1.36. *Rank Inequalities from Frobenius*
The following are special cases of Theorem 4.1.35:

(i) $\mathrm{rank}(AB) \leq \min(\mathrm{rank}(A), \mathrm{rank}(B))$;
(ii) $\mathrm{rank}(A) + \mathrm{rank}(B) - k \leq \mathrm{rank}(AB)$.

Proof.

(i) Apply Frobenius to conformable factors X, Y, Z. Taking $Z = 0$ gives $\mathrm{rank}(XY) \leq \mathrm{rank}(Y)$, because $\mathrm{rank}(Y0) = \mathrm{rank}(XY0) = 0$. Taking $X = 0$ gives $\mathrm{rank}(YZ) \leq \mathrm{rank}(Y)$.
Renaming the factors gives $\mathrm{rank}(AB) \leq \mathrm{rank}(B)$ and $\mathrm{rank}(AB) \leq \mathrm{rank}(A)$.
(ii) Taking the middle matrix to be I_k gives $\mathrm{rank}(A) + \mathrm{rank}(B) \leq k + \mathrm{rank}(AB)$, which is Sylvester's inequality.

Structural Commentary: Frobenius is the three-factor parent inequality. The usual upper and lower product-rank bounds appear by making one factor zero or identity.
Key Insight: Zero factors isolate upper bounds; an identity middle factor yields the Sylvester lower bound. □

Proposition 4.1.37. *Sylvester's Law of Nullity*
For a matrix X, let $\mathrm{nullity}(X) = \dim(\ker(X))$; this is the dimension of its nullspace.

(i) If $A \in \mathrm{Mat}(m, k, \mathbb{F})$ and $B \in \mathrm{Mat}(k, n, \mathbb{F})$, then $\mathrm{nullity}(AB) \leq \mathrm{nullity}(A) + \mathrm{nullity}(B)$.
(ii) If $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$, then $\max(\mathrm{nullity}(A), \mathrm{nullity}(B)) \leq \mathrm{nullity}(AB)$.

Proof.

(i) By rank-nullity and Theorem 4.1.34,

$$\begin{aligned} \mathrm{nullity}(AB) &= n - \mathrm{rank}(AB) \leq n - \mathrm{rank}(A) - \mathrm{rank}(B) + k \\ &= (k - \mathrm{rank}(A)) + (n - \mathrm{rank}(B)) = \mathrm{nullity}(A) + \mathrm{nullity}(B). \end{aligned}$$

(ii) For square A, B, $\mathrm{rank}(AB) \leq \mathrm{rank}(A)$, $\mathrm{rank}(AB) \leq \mathrm{rank}(B)$. Subtract from n gives desired lower bounds.

Structural Commentary: Nullity is rank loss from the domain.
Key Insight: Use rank-nullity to translate rank bounds into kernel bounds. □

Theorem 4.1.38. *Rank Triangle Inequality*
If $A, B \in \mathrm{Mat}(m, n, \mathbb{F})$, then $|\mathrm{rank}(A) - \mathrm{rank}(B)| \leq \mathrm{rank}(A + B) \leq \mathrm{rank}(A) + \mathrm{rank}(B)$.

Proof. Since $\mathrm{col}(A + B) \subseteq \mathrm{col}(A) + \mathrm{col}(B)$, Theorem 1.2.19 gives $\mathrm{rank}(A + B) \leq \dim(\mathrm{col}(A) + \mathrm{col}(B)) \leq \mathrm{rank}(A) + \mathrm{rank}(B)$. Applying this upper bound to $A = (A + B) - B$ gives $\mathrm{rank}(A) \leq \mathrm{rank}(A + B) + \mathrm{rank}(B)$, so $\mathrm{rank}(A) - \mathrm{rank}(B) \leq \mathrm{rank}(A + B)$. Applying it to $B = (A + B) - A$ gives $\mathrm{rank}(B) - \mathrm{rank}(A) \leq \mathrm{rank}(A + B)$. Together these are exactly $|\mathrm{rank}(A) - \mathrm{rank}(B)| \leq \mathrm{rank}(A + B)$.
Structural Commentary: The rank of a sum is controlled by the sum of the two column spaces. The reverse bounds come by solving for either summand.
Key Insight: Rank behaves like a dimension-valued size under addition. □

Definition 4.1.39. *Commutator*
For $S, T \in \mathcal{L}(\mathcal{V})$, the *commutator* of S and T is $[S, T] := ST - TS$.
For $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$, the *matrix commutator* is $[A, B] := AB - BA$.

Proposition 4.1.40. *Trace Obstruction for Commutators*
If $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$, then $\mathrm{tr}(AB - BA) = 0$. Hence any matrix similar to a commutator has trace zero.

Proof. By linearity of trace and Proposition 4.1.4, $\mathrm{tr}(AB - BA) = \mathrm{tr}(AB) - \mathrm{tr}(BA) = 0$.
If C is similar to $AB - BA$, then Theorem 4.1.21 gives $\mathrm{tr}(C) = \mathrm{tr}(AB - BA) = 0$.
Structural Commentary: Trace kills commutators because it is cyclic.
Key Insight: AB and BA have the same trace, so their difference has trace zero. □

Lemma 4.1.41. *Scalar Matrices and One-Step Orbits*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ with $n \geq 2$.
Then $(\mathbf{x}, A\mathbf{x})$ is linearly dependent for every $\mathbf{x} \in \mathbb{F}^n$ if and only if A is a scalar matrix.

Proof.

($\Rightarrow$) Suppose $A = cI_n$. Then $A\mathbf{x} = c\mathbf{x}$ for every $\mathbf{x}$, so $(\mathbf{x}, A\mathbf{x})$ is linearly dependent.
($\Leftarrow$) Suppose $(\mathbf{x}, A\mathbf{x})$ is linearly dependent for every $\mathbf{x}$.
For each standard basis vector $\mathbf{e}_i$, $\exists a_i \in \mathbb{F}$ such that $A\mathbf{e}_i = a_i\mathbf{e}_i$. For $i \geq 2$, dependence of $\mathbf{e}_1 + \mathbf{e}_i$ and $A(\mathbf{e}_1 + \mathbf{e}_i) = a_1\mathbf{e}_1 + a_i\mathbf{e}_i$ means $a_1\mathbf{e}_1 + a_i\mathbf{e}_i = b(\mathbf{e}_1 + \mathbf{e}_i)$ for some $b \in \mathbb{F}$. Comparing the independent vectors $\mathbf{e}_1, \mathbf{e}_i$ gives $a_1 = b = a_i$. Hence every a_i equals a_1, so $A = a_1 I_n$.

Structural Commentary: A scalar matrix preserves every one-dimensional coordinate direction and every diagonal line between two coordinates. These diagonal lines force all coordinate scalars to agree.
Key Insight: If every vector is an eigenvector, then there can be only one eigenvalue. □

Lemma 4.1.42. *Nonscalar Matrices Have a Zero Diagonal Entry After Similarity*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ with $n \geq 2$.
If A is not a scalar matrix, then A is similar over $\mathbb{F}$ to a matrix whose $(1, 1)$ entry is 0.

Proof. By Lemma 4.1.41, there is $\mathbf{x} \in \mathbb{F}^n$ such that $(\mathbf{x}, A\mathbf{x})$ is linearly independent. Extend $(\mathbf{x}, A\mathbf{x})$ to a basis β of $\mathbb{F}^n$. In the basis β, the first column of the matrix of A is the coordinate vector of $A\mathbf{x}$, namely $(0, 1, 0, \ldots, 0)^t$. Thus the $(1, 1)$ entry is 0, and this matrix is similar to A.
Structural Commentary: Nonscalarity produces a vector whose first iterate points in a genuinely new direction. Using these two vectors as the start of a basis forces a zero in the first diagonal position.
Key Insight: Put $\mathbf{x}$ first and $A\mathbf{x}$ second in the basis. □

Theorem 4.1.43. *Similarity to a Constant-Diagonal Matrix*
Let $n \geq 1$, let $A \in \mathrm{Mat}(n, n, \mathbb{F})$, and assume $k \cdot 1_{\mathbb{F}} \neq 0$ for every $1 \leq k \leq n$. Then A is similar over $\mathbb{F}$ to a matrix whose diagonal entries are all equal to $\frac{1}{n}\mathrm{tr}(A)$.

Proof. Set $\tau = \frac{1}{n}\mathrm{tr}(A)$ and $B = A - \tau I_n$. Then $\mathrm{tr}(B) = 0$. It is enough to show that every trace-zero matrix is similar to a matrix with zero diagonal, because Theorem 1.3.113 then adds τI_n back to the diagonal.
To prove the trace-zero claim by induction on n. If $n = 1$, a trace-zero matrix is $[0]$. Let $n \geq 2$ and assume the claim holds for smaller sizes. Let $\mathrm{tr}(B) = 0$. If B is scalar, then $B = cI_n$ and $0 = \mathrm{tr}(B) = nc$.
Since $n \cdot 1_{\mathbb{F}} \neq 0$, we get $c = 0$, so B already has zero diagonal.
If B is not scalar, Lemma 4.1.42 makes B similar to $\begin{bmatrix} 0 & * \\ * & B_1 \end{bmatrix}$, and $\mathrm{tr}(B_1) = 0$.
By induction, choose an invertible S_1 such that $S_1^{-1}B_1S_1$ has zero diagonal. Conjugating the displayed block matrix by $1 \oplus S_1$ leaves the first diagonal entry equal to 0 and replaces the lower-right block by $S_1^{-1}B_1S_1$.

Hence the resulting similar matrix has every diagonal entry equal to 0.
Structural Commentary: The proof removes one zero diagonal entry at a time.
Key Insight: Similarity creates the first zero; trace zero handles the smaller block. □

Theorem 4.1.44. *Shoda's Theorem in the Distinct-Diagonal Case*
Let $n \geq 1$ and assume $k \cdot 1_{\mathbb{F}} \neq 0$ for every $1 \leq k \leq n$. For $A \in \mathrm{Mat}(n, n, \mathbb{F})$, the matrix A is a commutator in $\mathrm{Mat}(n, n, \mathbb{F})$ if and only if $\mathrm{tr}(A) = 0$. In particular, no nonzero scalar matrix is a commutator.

Proof.

(⇒) Proposition 4.1.40 gives $\mathrm{tr}(A) = 0$.
Thus cI_n can be a commutator only when $nc = \mathrm{tr}(cI_n) = 0$, hence only when $c = 0$.

(⇐) Suppose $\mathrm{tr}(A) = 0$. By Theorem 4.1.43, there is an invertible S such that $B = S^{-1}AS = [b_{ij}]$ has zero diagonal. Let $X = \mathrm{diag}(1, 2, \ldots, n)$, viewing the integers as scalars in $\mathbb{F}$, and define $Y = [y_{ij}]$ by

$$y_{ij} = \begin{cases} (i-j)^{-1}b_{ij}, & i \neq j, \\ 0, & i = j. \end{cases}$$

The hypothesis ensures that $i - j \neq 0$ in $\mathbb{F}$ whenever $i \neq j$. Thus $XY - YX = [(i-j)y_{ij}] = B$.
Hence $A = (SXS^{-1})(SYS^{-1}) - (SYS^{-1})(SXS^{-1})$, so A is a commutator.

Structural Commentary: Trace is the obstruction; similarity exposes zero diagonal form.
Key Insight: Cyclicity forces trace zero; distinct diagonals realize it. □

Remark 4.1.45. *Field Hypothesis in the Diagonal Proof*
The proof above uses that the integers $1, \ldots, n$ are distinct and that each smaller block size remains nonzero in $\mathbb{F}$. In small characteristic, this diagonal proof must be replaced by a more delicate argument.

Corollary 4.1.46. *The Identity Is Not a Commutator*
Assume $\dim \mathcal{V} \cdot 1_{\mathbb{F}} \neq 0$. There do not exist $S, T \in \mathcal{L}(\mathcal{V})$ such that $ST - TS = I_{\mathcal{V}}$.

Proof. If $ST - TS = I_{\mathcal{V}}$, choose a basis and let $A = [S]$ and $B = [T]$ in that basis. Then $AB - BA = I_n$.
Taking traces and using Proposition 4.1.4, $0 = \mathrm{tr}(AB) - \mathrm{tr}(BA) = \mathrm{tr}(I_n) = n \cdot 1_{\mathbb{F}}$, contradicting the hypothesis.
Structural Commentary: Commutators have zero trace, while the identity has dimension as its trace.
Key Insight: Trace detects a scalar obstruction invisible from multiplication alone. □

Definition 4.1.47. *Idempotent and Nilpotent Matrices*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$.

(i) A is *idempotent* if $A^2 = A$.
(ii) A is *nilpotent* if $A^k = 0$ for some positive integer k.

An idempotent matrix is an algebraic projection: it fixes its range and kills its kernel.

Proposition 4.1.48. *Basic Invariants of Idempotents*
If $P \in \mathrm{Mat}(n, n, \mathbb{F})$ satisfies $P^2 = P$, then $\mathbb{F}^n = \mathrm{range}(P) \oplus \ker(P)$, P is diagonalizable, and its only possible eigenvalues are 0 and 1. Moreover $\mathrm{tr}(P) = \mathrm{rank}(P)$.
If $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, then P is the orthogonal projection onto $\mathrm{range}(P)$ if and only if $P = P^*$.

Proof. For every $\mathbf{x}$, write $\mathbf{x} = P\mathbf{x} + (\mathbf{x} - P\mathbf{x})$. The first term lies in $\mathrm{range}(P)$, while $P(\mathbf{x} - P\mathbf{x}) = P\mathbf{x} - P^2\mathbf{x} = 0$, so the second term lies in $\ker(P)$. The intersection is zero because $\mathbf{x} = P\mathbf{y}$ and $P\mathbf{x} = 0$ imply $\mathbf{x} = P^2\mathbf{y} = P\mathbf{x} = 0$.
In a basis adapted to $\mathrm{range}(P) \oplus \ker(P)$, the matrix of P is $I_r \oplus 0$, where $r = \mathrm{rank}(P)$.
Hence P is diagonalizable, its eigenvalues lie in $\{0, 1\}$, and $\mathrm{tr}(P) = r$.
For the final claim, self-adjoint idempotents are exactly orthogonal projections by Theorem 2.4.19.
Structural Commentary: Idempotents split the space into fixed directions and killed directions.
Key Insight: In adapted coordinates, every idempotent is just $I_r \oplus 0$. □

Proposition 4.1.49. *Basic Invariants of Nilpotents*
For $N \in \mathrm{Mat}(n, n, \mathbb{F})$, the following are equivalent:

(i) N is nilpotent;
(ii) the characteristic polynomial of N is t^n;
(iii) $N^n = 0$.

In this case, $\mathrm{tr}(N) = 0$, $\det(N) = 0$ for $n \geq 1$, and the spectral radius of N over $\mathbb{C}$ is 0.

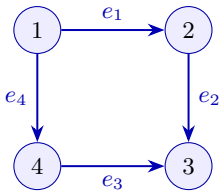
Proof. If $N^k = 0$, then every eigenvalue λ of N over an algebraic closure satisfies $\lambda^k = 0$, hence $\lambda = 0$. Thus the characteristic polynomial is t^n. If the characteristic polynomial is t^n, Cayley–Hamilton gives $N^n = 0$. The implication $N^n = 0 \Rightarrow$ nilpotent is immediate. The trace, determinant, and spectral-radius statements follow from the eigenvalue data and from Definitions 4.1.1 and 4.1.92.
Structural Commentary: Nilpotent matrices have no spectral mass away from zero.
Key Insight: Cayley–Hamilton turns the characteristic polynomial t^n into the equation $N^n = 0$. □

Definition 4.1.50. *Positive Semidefinite and Positive Definite Matrices*
Let $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$ and let $A \in \mathrm{Mat}(n, n, \mathbb{F})$.
The matrix A is *positive semidefinite*, written $A \succeq 0$, if $A = A^*$ and $\mathbf{x}^* A\mathbf{x} \geq 0$ for every $\mathbf{x} \in \mathbb{F}^n$.
It is *positive definite*, written $A \succ 0$, if $A = A^*$ and $\mathbf{x}^* A\mathbf{x} > 0$ for every nonzero $\mathbf{x} \in \mathbb{F}^n$.
Negative semidefinite and negative definite mean $-A \succeq 0$ and $-A \succ 0$, respectively.

Proposition 4.1.51. *Rectangular Gram Products Are Positive Semidefinite*
Let $R \in \mathrm{Mat}(m, n, \mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$. Then $R^*R \succeq 0$, $RR^* \succeq 0$, and $(R^*R)_{j,j} = \|R_{\cdot,j}\|_2^2$; $(RR^*)_{i,i} = \|R_{i,\cdot}\|_2^2$. Full column rank gives $R^*R \succ 0$; full row rank gives $RR^* \succ 0$.

Proof. The matrices are self-adjoint by Proposition 2.4.8.
For $\mathbf{x} \in \mathbb{F}^n$ and $\mathbf{y} \in \mathbb{F}^m$, $\mathbf{x}^* R^* R\mathbf{x} = \|R\mathbf{x}\|_2^2 \geq 0$ and $\mathbf{y}^* RR^*\mathbf{y} = \|R^*\mathbf{y}\|_2^2 \geq 0$. Hence both products are positive semidefinite. Taking $\mathbf{x} = \mathbf{e}_j$ and $\mathbf{y} = \mathbf{e}_i$ gives the diagonal formulas. If R has full column rank, then $R\mathbf{x} \neq 0$ for $\mathbf{x} \neq 0$, so $R^*R \succ 0$. If R has full row rank, apply the same argument to R^*.
Structural Commentary: Multiplying a matrix by its adjoint produces a Gram matrix of columns or rows.
Key Insight: The quadratic form is a squared norm, so no cancellation can make it negative. □

Definition 4.1.52. *Directed Incidence Matrix*
For edges $e = u \to v$ on vertices $1, \ldots, n$, $B \in \mathrm{Mat}(m, n, \mathbb{R})$ has row e equal to $-\mathbf{e}_u^t + \mathbf{e}_v^t$.



$$B = \begin{bmatrix} -1 & 1 & 0 & 0 \\ 0 & -1 & 1 & 0 \\ 0 & 0 & 1 & -1 \\ -1 & 0 & 0 & 1 \end{bmatrix}$$

Figure 76: A directed incidence matrix records oriented edge differences from vertices to edges.

Example 4.1.53. *A Directed Incidence Matrix*
For Figure 76, $B = \begin{bmatrix} -1 & 1 & 0 & 0 \\ 0 & -1 & 1 & 0 \\ 0 & 0 & 1 & -1 \\ -1 & 0 & 0 & 1 \end{bmatrix}$. For node potentials $\mathbf{x} = (x_1, x_2, x_3, x_4)^t$, $B\mathbf{x} = \begin{bmatrix} x_2 - x_1 \\ x_3 - x_2 \\ x_3 - x_4 \\ x_4 - x_1 \end{bmatrix}$.
Hence $\ker(B) = \mathrm{span}\{(1, 1, 1, 1)^t\}$, and $\mathrm{rank}(B) = 3$. The loop vector $\mathbf{y} = (1, 1, -1, -1)^t$ spans $\ker(B^t)$.
Thus $\mathbf{b} \in \mathrm{col}(B)$ if and only if $\mathbf{y}^t\mathbf{b} = b_1 + b_2 - b_3 - b_4 = 0$. This is Kirchhoff loop compatibility.
Structural Commentary: Incidence subspaces encode graph constraints.
Key Insight: A sparse matrix can encode both the algebra of a linear map and the combinatorics of a graph.

Perspective 4.1.54. *Incidence Data and Least Squares*
For an incidence matrix B as in Example 4.1.53, $B\mathbf{x} = \mathbf{b}$ seeks node potentials matching edge differences. If loop consistency fails, least squares projects $\mathbf{b}$ onto $\mathrm{col}(B)$, where $B^*(B\mathbf{x}_0 - \mathbf{b}) = \mathbf{0}$ (Theorem 2.3.4 on B). The constant-vector kernel means potentials are unique only after fixing one node or choosing minimum norm.

Proposition 4.1.55. *Incidence Laplacians Are Positive Semidefinite*
Let B be a directed incidence matrix and set $L = B^tB$. Then $L \succeq 0$ and $\mathbf{x}^t L\mathbf{x} = \sum_{e=(u\to v)} (x_v - x_u)^2$, with $\ker(L)$ equal to the vectors constant on each connected component.

Proof. Since $L = B^tB$, Proposition 4.1.51 gives $L \succeq 0$.
For $e = (u \to v)$, $(B\mathbf{x})_e = x_v - x_u$, so $\mathbf{x}^t L\mathbf{x} = \|B\mathbf{x}\|_2^2 = \sum_{e=(u\to v)} (x_v - x_u)^2$. Thus $\mathbf{x} \in \ker(L)$ if and only if $B\mathbf{x} = \mathbf{0}$, equivalently $x_u = x_v$ on every edge, which means $\mathbf{x}$ is constant on each connected component.
Structural Commentary: Laplacians sum squared edge differences.
Key Insight: A graph Laplacian is a Gram product, so positivity is automatic. □

Theorem 4.1.56. *Equivalent Forms of Positive Semidefiniteness*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R},\mathbb{C}\}$. The following are equivalent:

(i) $A \succeq 0$.
(ii) *A is Hermitian and all eigenvalues of A are nonnegative.*
(iii) $A = B^*B$ *for some* $B \in \mathrm{Mat}(m,n,\mathbb{F})$.
(iv) *A has a positive semidefinite square root.*

In this case the positive semidefinite square root is unique.
If $A = U\Lambda U^*$ with $\Lambda = \mathrm{diag}(\lambda_1,\dots,\lambda_n)$ and $\lambda_i \geq 0$, then $A^{1/2} = U\,\mathrm{diag}(\sqrt{\lambda_1},\dots,\sqrt{\lambda_n})U^*$.

Proof.

(i) $\Rightarrow$ *(ii)* This is the matrix form of Theorem 3.4.41.
(ii) $\Rightarrow$ *(iv)* By the spectral theorem, write $A = U\Lambda U^*$, where U is unitary and $\Lambda = \mathrm{diag}(\lambda_1,\dots,\lambda_n)$ with $\lambda_i \geq 0$. Then $A^{1/2} = U\,\mathrm{diag}(\sqrt{\lambda_1},\dots,\sqrt{\lambda_n})U^*$ is positive semidefinite and satisfies $(A^{1/2})^2 = A$.
(iv) $\Rightarrow$ *(iii)* If $C \succeq 0$ and $C^2 = A$, then $C = C^*$ and $A = C^*C$, so (iii) holds with $m = n$.
(iii) $\Rightarrow$ *(i)* If $A = B^*B$, then $A^* = A$ and $\mathbf{x}^*A\mathbf{x} = \mathbf{x}^*B^*B\mathbf{x} = \|B\mathbf{x}\|_2^2 \geq 0$, so $A \succeq 0$.

The positive semidefinite square root is unique by Proposition 3.4.42.
Structural Commentary: Positive semidefinite matrices are Gram-type matrices with nonnegative spectrum.
Key Insight: Positivity is best checked either on eigenvalues or by finding a factor B with $A = B^*B$. □

Method 4.1.57. *PSD Spectral Factorisation Method*
To analyse a Hermitian matrix A:

(i) diagonalise $A = U\Lambda U^*$;
(ii) check the signs of the diagonal entries of Λ;
(iii) if $\Lambda \geq 0$, set $A^{1/2} = U\Lambda^{1/2}U^*$ and use $A = (A^{1/2})^*A^{1/2}$;
(iv) translate kernel, rank, and zero-quadratic-form questions into the factor $A^{1/2}$.

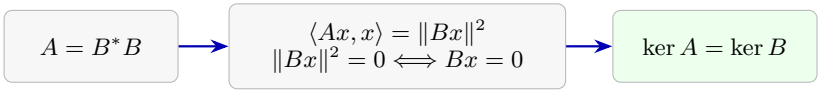


Figure 77: For $A = B^*B$, the squared-norm identity gives $\ker A = \ker B$.

Theorem 4.1.58. *Equivalent Forms of Positive Definiteness*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R},\mathbb{C}\}$. The following are equivalent:

(i) $A \succ 0$.
(ii) *A is Hermitian and all eigenvalues of A are positive.*
(iii) $A \succeq 0$ *and A is invertible.*
(iv) $A = B^*B$ *for some invertible* $B \in \mathrm{Mat}(n,n,\mathbb{F})$.
(v) *A is invertible and $A^{-1} \succ 0$.*

Proof.

(i) $\Rightarrow$ *(ii)* By definition A is Hermitian. If $(\lambda,\mathbf{v})$ is an eigenpair with $\mathbf{v} \neq \mathbf{0}$, then $\lambda\|\mathbf{v}\|_2^2 = \mathbf{v}^*A\mathbf{v} > 0$, so every eigenvalue is positive.
(ii) $\Rightarrow$ *(iii)* Positive eigenvalues imply $A \succeq 0$ by Theorem 4.1.56; they also exclude 0, so A is invertible.
(iii) $\Rightarrow$ *(iv)* By Theorem 4.1.56, A has a positive semidefinite square root $A^{1/2}$.
Since A is invertible, $A^{1/2}$ is invertible, and $A = (A^{1/2})^*A^{1/2}$.
(iv) $\Rightarrow$ *(i)* If $A = B^*B$ and B is invertible, then for every nonzero $\mathbf{x}$, $\mathbf{x}^*A\mathbf{x} = \|B\mathbf{x}\|_2^2 > 0$.
(ii) $\Leftrightarrow$ *(v)* Hermitian-ness is preserved by inversion, and eigenvalues of A^{-1} are reciprocal eigenvalues of A. Thus positive eigenvalues of A are equivalent to positive eigenvalues of A^{-1}.

Structural Commentary: Positive definite matrices are the invertible part of the positive semidefinite cone.
Key Insight: Positive definite means every eigenvalue is strictly on the positive side of zero. □

Proposition 4.1.59. *Trace and Determinant of Positive Semidefinite Matrices*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$ be positive semidefinite, where $\mathbb{F} \in \{\mathbb{R},\mathbb{C}\}$. Then $\mathrm{tr}\,A \geq 0$, with equality if and only if $A = 0$, and $\det A \geq 0$, with strict inequality if and only if A is positive definite.

Proof. List the eigenvalues of A as $\lambda_1,\dots,\lambda_n$. By Theorem 4.1.56, each $\lambda_j \geq 0$.
Hence $\mathrm{tr}\,A = \sum_{j=1}^{n} \lambda_j \geq 0$ and $\det A = \prod_{j=1}^{n} \lambda_j \geq 0$.

The trace is zero exactly when every $\lambda_j = 0$, which for a Hermitian matrix means $A = 0$. The determinant is positive exactly when every $\lambda_j > 0$, which is positive definiteness by Theorem 4.1.58.
Structural Commentary: Trace and determinant aggregate the nonnegative eigenvalues by sum and product.
Key Insight: Equality is controlled by which eigenvalues vanish. □

Theorem 4.1.60. *Kernel, Rank, and Columns of B^*B*
Let $B \in \mathrm{Mat}(m, n, \mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, and set $A = B^*B$. Then:

(i) $A \succeq 0$;
(ii) $\ker(A) = \ker(B)$;
(iii) $\mathrm{rank}(A) = \mathrm{rank}(B)$;
(iv) $\mathrm{col}(A) = \mathrm{col}(B^*)$;
(v) $A = 0$ *if and only if* $B = 0$.

Proof.

(i) Positive semidefiniteness follows from Theorem 4.1.56.
(ii) If $B\mathbf{x} = 0$, then $A\mathbf{x} = 0$. Conversely, if $A\mathbf{x} = 0$, then $0 = \mathbf{x}^*A\mathbf{x} = \|B\mathbf{x}\|_2^2$ so $B\mathbf{x} = 0$.
(iii) Since A and B both have n columns, rank-nullity and (ii) give $\mathrm{rank}(A) = \mathrm{rank}(B)$.
(iv) Since $A = A^*$, Proposition 2.4.11 and (ii) give $\mathrm{col}(A) = \ker(A)^\perp = \ker(B)^\perp = \mathrm{col}(B^*)$.
(v) By (iii), $\mathrm{rank}(A) = 0$ if and only if $\mathrm{rank}(B) = 0$.
A matrix has rank zero exactly when it is the zero matrix, so $A = 0$ if and only if $B = 0$.

Structural Commentary: The Gram product B^*B preserves the kernel of B.
Key Insight: The identity $\mathbf{x}^*B^*B\mathbf{x} = \|B\mathbf{x}\|_2^2$ is the whole mechanism. □

Corollary 4.1.61. *Principal Blocks and Column Inclusion*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$ be PSD, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, and partition $A = \begin{bmatrix} A_{11} & A_{12} \\ A_{12}^* & A_{22} \end{bmatrix}$, $A_{11} \in \mathrm{Mat}(k, k, \mathbb{F})$.
Then $A_{11} \succeq 0$, $\mathrm{col}(A_{12}) \subseteq \mathrm{col}(A_{11})$, and $\mathrm{rank}(A_{12}) \leq \mathrm{rank}(A_{11})$. If $A \succ 0$, then $A_{11} \succ 0$.

Proof. Write $A = B^*B$ and partition $B = [B_1 \; B_2]$ conformally. Then $A_{11} = B_1^*B_1$ and $A_{12} = B_1^*B_2$.
Hence $A_{11} \succeq 0$ and $\mathrm{col}(A_{12}) = \mathrm{col}(B_1^*B_2) \subseteq \mathrm{col}(B_1^*) = \mathrm{col}(B_1^*B_1) = \mathrm{col}(A_{11})$, equality $\mathrm{col}(B_1^*) = \mathrm{col}(B_1^*B_1)$ follows from Theorem 4.1.60. The rank inequality follows by taking dimensions
If $A \succ 0$ and $\mathbf{x} \neq 0$ in $\mathbb{F}^k$, then $\mathbf{x}^*A_{11}\mathbf{x} = \begin{bmatrix} \mathbf{x} \\ 0 \end{bmatrix}^* A \begin{bmatrix} \mathbf{x} \\ 0 \end{bmatrix} > 0$, so $A_{11} \succ 0$.

Structural Commentary: In a positive block matrix, cross-block columns lie in principal block's column space.
Key Insight: Factor first; then column inclusion is immediate from $A_{12} = B_1^*B_2$. □

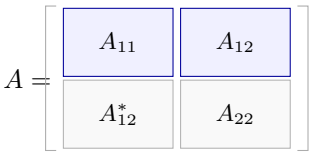


Figure 78: For $A \succeq 0$, factorisation forces $\mathrm{col}(A_{12}) \subseteq \mathrm{col}(A_{11})$.

Definition 4.1.62. *Principal and Leading Principal Minors*
For $A \in \mathrm{Mat}(n, n, \mathbb{C})$, a *principal minor* is the determinant of a principal submatrix of A.
A *leading principal minor* is the determinant of the leading $k \times k$ principal submatrix for some $k = 1, \ldots, n$.

Theorem 4.1.63. *Sylvester's Criterion*
A Hermitian matrix $A \in \mathrm{Mat}(n, n, \mathbb{C})$ is positive definite if and only if all leading principal minors are positive.

Proof.

($\Rightarrow$) If $A \succ 0$, then each leading principal submatrix is positive definite by Corollary 4.1.61. Its eigenvalues are positive, so its determinant is positive.
($\Leftarrow$) Induct on n. If $n = 1$, the single leading principal minor is the only diagonal entry, so positivity of that minor is exactly positive definiteness. Assume $n \geq 2$ and write $A = \begin{bmatrix} B & \mathbf{y} \\ \mathbf{y}^* & a \end{bmatrix}$, $B \in \mathrm{Mat}(n-1, n-1, \mathbb{C})$
The leading principal minors of B are positive, so $B \succ 0$.
Since $\det(B) > 0$ and $\det(A) > 0$, Proposition 4.1.12 gives $a - \mathbf{y}^*B^{-1}\mathbf{y} = \frac{\det(A)}{\det(B)} > 0$.
For $\mathbf{u} \in \mathbb{C}^{n-1}$ and $t \in \mathbb{C}$, $\begin{bmatrix} \mathbf{u} \\ t \end{bmatrix}^* A \begin{bmatrix} \mathbf{u} \\ t \end{bmatrix} = (\mathbf{u} + B^{-1}\mathbf{y}t)^*B(\mathbf{u} + B^{-1}\mathbf{y}t) + \left(a - \mathbf{y}^*B^{-1}\mathbf{y}\right)|t|^2$.

First term is nonnegative as $B \succ 0$, and second coefficient is positive. If displayed vector is nonzero, then either $t \neq 0$ and second term is positive, or $t = 0$ and then $\mathbf{u} \neq 0$, making first term positive. Thus $A \succ 0$.

Structural Commentary: Leading minors make each leading block positive and final Schur complement positive.
Key Insight: Complete the square with the leading block. □

Method 4.1.64. *2×2 Quadratic Form Test*
Let $A = \begin{bmatrix} a & b \\ b & c \end{bmatrix}$ be real symmetric, $\Delta = ac - b^2$, and $q(x,y) = \begin{bmatrix} x & y \end{bmatrix} A \begin{bmatrix} x \\ y \end{bmatrix} = ax^2 + 2bxy + cy^2$.
For $a \neq 0$, $q(x,y) = a\left(x + \frac{b}{a}y\right)^2 + \frac{\Delta}{a}y^2$.
If $a = 0$, swap variables when $c \neq 0$; if $a = c = 0$, use $q = 2bxy$. Thus:

(i) $q \succ 0$ if and only if $a > 0$ and $\Delta > 0$;
(ii) $q \prec 0$ if and only if $a < 0$ and $\Delta > 0$;
(iii) q is indefinite if and only if $\Delta < 0$;
(iv) if $\Delta = 0$, then $q \succeq 0$ iff $a, c \geq 0$, and $q \preceq 0$ iff $a, c \leq 0$.

This is the 2×2 form of Sylvester's criterion and inertia.
Structural Commentary: Completing the square is the 2×2 shadow of the LDL^* and Schur-complement tests.
Key Insight: Δ and the leading sign classify the form.

Corollary 4.1.65. *Zero Quadratic Form Detects the Kernel*
If $A \in \mathrm{Mat}(n, n, \mathbb{F})$ is positive semidefinite, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, then $\mathbf{x}^* A\mathbf{x} = 0$ if and only if $A\mathbf{x} = 0$.

Proof. Write $A = B^*B$.
($\Rightarrow$) If $\mathbf{x}^*A\mathbf{x} = 0$, then $\|B\mathbf{x}\|_2^2 = 0$, so $B\mathbf{x} = 0$ and hence $A\mathbf{x} = B^*B\mathbf{x} = 0$.
($\Leftarrow$) If $A\mathbf{x} = 0$, then $\mathbf{x}^*A\mathbf{x} = 0$.

Structural Commentary: A zero PSD quadratic value has no hidden cancellation.
Key Insight: Zero length-square means the factored vector itself is zero. □

Theorem 4.1.66. *Stability Rules for Positive Semidefinite Matrices*
Let $A, B \succeq 0$ over $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$. Whenever the expressions are defined:

(i) $aA + bB \succeq 0$ for $a, b \geq 0$; if $a, b > 0$ and one of A, B is positive definite, then $aA + bB \succ 0$.
(ii) $S^*AS \succeq 0$; moreover, $S^*AS \succ 0$ if and only if $S\mathbf{x} \notin \ker(A)$ for all $\mathbf{x} \neq \mathbf{0}$.
(iii) $A \oplus B \succeq 0$, with $A \oplus B \succ 0$ if and only if $A \succ 0$ and $B \succ 0$.

Proof.
(i) $\mathbf{x}^*(aA + bB)\mathbf{x} = a\mathbf{x}^*A\mathbf{x} + b\mathbf{x}^*B\mathbf{x} \geq 0$. If, for example, $A \succ 0$ and $a > 0$, then for every nonzero $\mathbf{x}$ the first summand is positive and the second is nonnegative, so $aA + bB \succ 0$; the case $B \succ 0$ is identical.
(ii) $\mathbf{x}^*S^*AS\mathbf{x} = (S\mathbf{x})^*A(S\mathbf{x}) \geq 0$. By Corollary 4.1.65, equality occurs exactly when $S\mathbf{x} \in \ker(A)$.
(iii) $(\mathbf{x} \oplus \mathbf{y})^*(A \oplus B)(\mathbf{x} \oplus \mathbf{y}) = \mathbf{x}^*A\mathbf{x} + \mathbf{y}^*B\mathbf{y} \geq 0$. This expression is positive for every nonzero $\mathbf{x} \oplus \mathbf{y}$ exactly when each summand is positive on its own nonzero vectors, namely when $A \succ 0$ and $B \succ 0$.

Structural Commentary: The PSD cone is stable under sums, pullbacks, and direct sums.
Key Insight: Test the quadratic form after each construction. □

Proposition 4.1.67. *Explicit Square Root for a 2×2 Positive Semidefinite Matrix*
Let $A \in \mathrm{Mat}(2, 2, \mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, be positive semidefinite and nonzero. Set $\tau = \left(\operatorname{tr} A + 2\sqrt{\det A}\right)^{1/2}$. Then $\tau > 0$ and $A^{1/2} = \frac{1}{\tau}\left(A + \sqrt{\det A}\, I_2\right)$.

Proof. Let $C = \frac{1}{\tau}\left(A + \sqrt{\det A}\, I_2\right)$. By Proposition 4.1.59, $\det A \geq 0$ and $\operatorname{tr} A > 0$, hence $\tau > 0$; by Theorem 4.1.66, $C \succeq 0$. Cayley–Hamilton gives $A^2 - (\operatorname{tr} A)A + (\det A)I_2 = 0$, so

$$C^2 = \frac{1}{\tau^2}\left(A^2 + 2\sqrt{\det A}\, A + (\det A)I_2\right) = \frac{1}{\tau^2}\left((\operatorname{tr} A)A + 2\sqrt{\det A}\, A\right) = A.$$

Thus $C = A^{1/2}$ by uniqueness of the positive semidefinite square root.
Structural Commentary: In dimension 2, Cayley–Hamilton makes $A^{1/2}$ linear in A.
Key Insight: Square the linear candidate and use Cayley–Hamilton. □

Theorem 4.1.68. *Simultaneous Congruence with One Positive Definite Form*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$ be Hermitian, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, and let $A \succ 0$.
Then, for some invertible S and real diagonal Λ, $A = SS^*$ and $B = S\Lambda S^*$.
The diagonal entries of Λ are the eigenvalues of $A^{-1}B$.

Proof. Let $C = A^{-1/2}BA^{-1/2} = U\Lambda U^*$ be spectral decomposition, $S = A^{1/2}U$. Then $SS^* = A$ and $S\Lambda S^* = B$. Also $A^{-1}B = A^{-1/2}CA^{1/2}$, so $A^{-1}B$ is similar to C and has eigenvalues given by the diagonal entries of Λ.
Structural Commentary: The positive definite form A normalizes B to a Hermitian matrix.
Key Insight: Normalize, diagonalise, then transform back. □

Corollary 4.1.69. *Products with One Positive Definite Factor*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, and let $A \succ 0$.

(i) If B is Hermitian, then AB is diagonalizable with real spectrum.
(ii) If $B \succeq 0$, then AB is diagonalizable with nonnegative real spectrum; if $B \neq 0$, one eigenvalue is positive.
(iii) If $B \succ 0$, then AB is diagonalizable with positive real spectrum.

Proof. Let $C = A^{1/2}BA^{1/2}$. Since $A^{-1/2}(AB)A^{1/2} = C$, the matrices AB and C are similar.

(i) If B is Hermitian, then C is Hermitian, hence diagonalizable with real spectrum.
(ii) If $B \succeq 0$, then $C \succeq 0$ by Theorem 4.1.66, so its spectrum is nonnegative.
If $B \neq 0$, then $C \neq 0$, hence some eigenvalue is positive.
(iii) If $B \succ 0$, then $C \succ 0$ by Theorem 4.1.66, so all eigenvalues are positive.

Structural Commentary: A positive definite factor makes AB similar to a Hermitian congruence.
Key Insight: Use $AB \sim A^{1/2}BA^{1/2}$. □

Theorem 4.1.70. *Principal Axes for Real Quadratic Forms*
Let $A \in \mathrm{Mat}(n, n, \mathbb{R})$ be symmetric. Then:

(i) there are a real orthogonal Q and real $\lambda_1, \dots, \lambda_n$ such that, with $\mathbf{y} = Q^t\mathbf{x}$, $\mathbf{x}^tA\mathbf{x} = \sum_{j=1}^{n} \lambda_j y_j^2$.
(ii) $A \succeq 0$ if and only if every $\lambda_j \geq 0$;
(iii) $A \succ 0$ if and only if every $\lambda_j > 0$;
(iv) A is indefinite if and only if some $\lambda_j > 0$ and some $\lambda_k < 0$.

Proof.

(i) By the real spectral theorem, write $A = Q\,\mathrm{diag}(\lambda_1, \dots, \lambda_n)Q^t$.
Then $\mathbf{x}^tA\mathbf{x} = \mathbf{x}^tQ\,\mathrm{diag}(\lambda_j)Q^t\mathbf{x} = \mathbf{y}^t\,\mathrm{diag}(\lambda_j)\mathbf{y} = \sum_{j=1}^{n} \lambda_j y_j^2$.
(ii) The displayed sum is nonnegative for all $\mathbf{y}$ exactly when each coefficient λ_j is nonnegative.
(iii) The displayed sum is positive for all nonzero $\mathbf{y}$ exactly when each coefficient λ_j is positive.
(iv) The diagonal form takes both signs exactly when its coefficients include both signs.

Structural Commentary: Principal axes turn the form into signed squares.
Key Insight: Eigenvalue signs classify the quadratic form. □

Method 4.1.71. *Quadratic Form Principal-Axes Method*
To analyse a real quadratic form:

(i) write it as $\mathbf{x}^tA\mathbf{x}$ with $A = A^t$; split each cross coefficient equally between a_{ij} and a_{ji};
(ii) diagonalise $A = Q\Lambda Q^t$;
(iii) change variables by $\mathbf{y} = Q^t\mathbf{x}$;
(iv) classify the form from the signs of the eigenvalues in Λ.

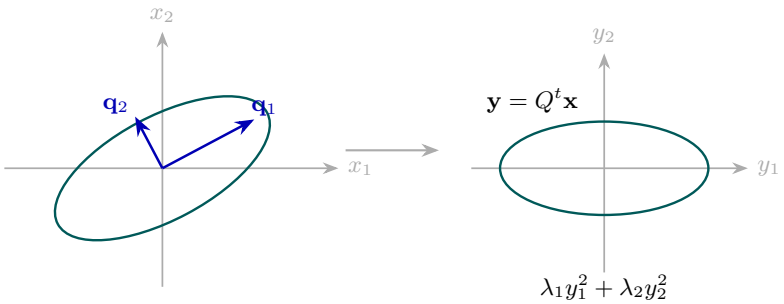


Figure 79: Orthogonal coordinates reveal the principal axes of a quadratic form.

Definition 4.1.72. *$*$-Congruence and Hermitian Inertia*
Two matrices $A, B \in \mathrm{Mat}(n, n, \mathbb{C})$ are *$*$-congruent* if $A = SBS^*$ for some invertible $S \in \mathrm{Mat}(n, n, \mathbb{C})$.
For Hermitian A, its *inertia* is $\mathrm{in}(A) = (n_+(A), n_-(A), n_0(A))$, where these are the numbers of positive, negative, and zero eigenvalues of A, counted with multiplicity.

Perspective 4.1.73. *Congruence versus Similarity*
Similarity $B = P^{-1}AP$ is change of basis for an operator. Congruence $N = P^tMP$ is change of basis for a bilinear form, with the same basis change in both slots. Over $\mathbb{R}$ this is the real case of $*$-congruence. Hence inertia classifies forms, not similarity classes of operators.

Lemma 4.1.74. *Positive Subspaces and Inertia*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ be Hermitian. The maximal dimensions of subspaces on which $\mathbf{x}^*A\mathbf{x}$ is positive or negative for every nonzero $\mathbf{x}$ are $n_+(A)$ and $n_-(A)$, respectively.

Proof. Let P_+, $P_{\leq 0}$ be positive, nonpositive spectral subspaces. Then $\dim P_+ = n_+(A)$, P_+ is positive for A.
If $\dim \mathcal{U} > n_+(A)$, then $\dim \mathcal{U} + \dim P_{\leq 0} > n_+(A) + n_-(A) + n_0(A) = n$, so $\mathcal{U} \cap P_{\leq 0}$ contains a nonzero vector.
On this vector quadratic form is nonpositive, so $\mathcal{U}$ cannot be positive. Apply this to $-A$ for negative case.
Structural Commentary: The positive spectral subspace is maximal; the nonpositive part blocks enlargement.
Key Insight: Inertia counts maximal sign-definite dimensions. □

Theorem 4.1.75. *Sylvester's Inertia Theorem*
Hermitian matrices A, B of the same size are $*$-congruent if and only if $n_+(A) = n_+(B)$ and $n_-(A) = n_-(B)$.

Proof.
($\Rightarrow$) Suppose $A = SBS^*$ with S invertible. Then $\mathbf{x}^*A\mathbf{x} = (S^*\mathbf{x})^*B(S^*\mathbf{x})$.
Thus S^* transports positive and negative subspaces, so Lemma 4.1.74 gives $n_\pm(A) = n_\pm(B)$.
($\Leftarrow$) Let the common inertia be (p, q, r) and set $D = I_p \oplus (-I_q) \oplus 0_r$.
By the spectral theorem, each of A and B is $*$-congruent to D, hence they are $*$-congruent to each other.
Structural Commentary: $*$-congruence is a coordinate change for Hermitian quadratic forms.
Key Insight: Spectral decomposition reduces the form to signed squares and zero directions. □

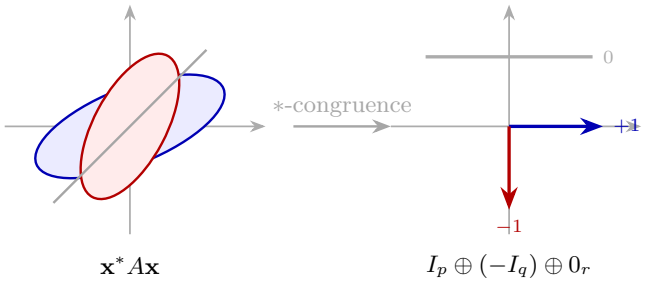


Figure 80: $*$-congruence reduces a Hermitian form to signed coordinate axes.

Definition 4.1.76. *Phase Inertia Matrix*
A *phase inertia matrix* is a diagonal complex matrix whose nonzero diagonal entries have modulus 1.

Proposition 4.1.77. *Phase Inertia Form for Normal Matrices*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ be normal.
Then A is $*$-congruent to a phase inertia matrix D, whose entries record the spectral rays of A.

Proof. Write $A = U \operatorname{diag}(\lambda_j)U^*$. Write $\lambda_j = r_j d_j$, where $r_j > 0$ and $|d_j| = 1$ if $\lambda_j \neq 0$, while $r_j = 1$ and $d_j = 0$ if $\lambda_j = 0$. Set $D = \operatorname{diag}(d_j)$, $S = U \operatorname{diag}(\sqrt{r_j})$. Then S is invertible and $A = SDS^*$.
The entries $d_j = e^{i\theta}$ count exactly the eigenvalues on the ray $\{\rho e^{i\theta} : \rho > 0\}$.
Structural Commentary: Normal diagonalization separates modulus from phase; $*$-congruence absorbs moduli.
Key Insight: Only spectral phases survive. □

Perspective 4.1.78. *Hermitian Inertia as Real Phase Inertia*
For Hermitian matrices, the only nonzero spectral rays are the positive and negative real axes.
Sylvester's inertia theorem is therefore the real-ray version of phase inertia.

Definition 4.1.79. *Kronecker Product*
Let $A = [a_{ij}] \in \mathrm{Mat}(m, n, \mathbb{F})$ and $B \in \mathrm{Mat}(p, q, \mathbb{F})$. The *Kronecker product* of A and B is the block matrix

$$A \otimes B := \begin{bmatrix} a_{11}B & a_{12}B & \cdots & a_{1n}B \\ a_{21}B & a_{22}B & \cdots & a_{2n}B \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1}B & a_{m2}B & \cdots & a_{mn}B \end{bmatrix} \in \mathrm{Mat}(mp, nq, \mathbb{F}).$$

The Kronecker product is also called the tensor product of matrices.

Proposition 4.1.80. *Basic Kronecker Identities*
Whenever the expressions are defined,

(i) $c(A \otimes B) = (cA) \otimes B = A \otimes (cB)$;
(ii) $(A + B) \otimes C = A \otimes C + B \otimes C$ *and* $A \otimes (B + C) = A \otimes B + A \otimes C$;
(iii) $(A \otimes B) \otimes C = A \otimes (B \otimes C)$;
(iv) $(A \otimes B)^t = A^t \otimes B^t$;
(v) $I_m \otimes I_n = I_{mn}$.

If $\mathbb{F} = \mathbb{C}$, *then also* $\overline{A \otimes B} = \overline{A} \otimes \overline{B}$ *and* $(A \otimes B)^* = A^* \otimes B^*$.

Proof.

(i) The (i, j) block of $c(A \otimes B)$ is $c(a_{ij}B) = (ca_{ij})B = a_{ij}(cB)$, which is (i, j) block of $(cA) \otimes B$ and $A \otimes (cB)$.
(ii) The (i, j) block of $(A + B) \otimes C$ is $(a_{ij} + b_{ij})C = a_{ij}C + b_{ij}C$, so $(A + B) \otimes C = A \otimes C + B \otimes C$. The proof of $A \otimes (B + C) = A \otimes B + A \otimes C$ is identical inside each block.
(iii) If $B = [b_{rs}]$, then both $(A \otimes B) \otimes C$ and $A \otimes (B \otimes C)$ have blocks $a_{ij}b_{rs}C$ in the same relative positions.
(iv) The (j, i) block of $(A \otimes B)^t$ is $(a_{ij}B)^t = a_{ij}B^t$, which is the (j, i) block of $A^t \otimes B^t$.
(v) The block matrix $I_m \otimes I_n$ has I_n on its diagonal blocks and zero elsewhere, hence equals I_{mn}.

If $\mathbb{F} = \mathbb{C}$, conjugation is checked blockwise, adjoint identity follows from conjugation together with part (iv).
Structural Commentary: Kronecker products replace scalar entries by whole scaled blocks. Standard algebraic rules survive because each block is still controlled by scalar distributivity and transposition.
Key Insight: Check identities block-by-block; each block reduces to an ordinary scalar-matrix identity. □

Theorem 4.1.81. *Mixed Product Property*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$, $B \in \mathrm{Mat}(p, q, \mathbb{F})$, $C \in \mathrm{Mat}(n, r, \mathbb{F})$, $D \in \mathrm{Mat}(q, s, \mathbb{F})$. $(A \otimes B)(C \otimes D) = (AC) \otimes (BD)$.

Proof. Write $A = [a_{ik}]$ and $C = [c_{kj}]$. The (i, j) block of $(A \otimes B)(C \otimes D)$ is
$\sum_{k=1}^{n} (a_{ik}B)(c_{kj}D) = \left(\sum_{k=1}^{n} a_{ik}c_{kj} \right) BD = (AC)_{ij}BD$, which is exactly the (i, j) block of $(AC) \otimes (BD)$.
Structural Commentary: The outer indices multiply as AC, while the inner block action multiplies as BD.
Key Insight: Kronecker multiplication separates into ordinary multiplication in each factor. □

Corollary 4.1.82. *Inverse of a Kronecker Product*
If $A \in \mathrm{Mat}(m, m, \mathbb{F})$ *and* $B \in \mathrm{Mat}(n, n, \mathbb{F})$ *are invertible, then* $A \otimes B$ *is invertible and* $(A \otimes B)^{-1} = A^{-1} \otimes B^{-1}$.

Proof. By Theorem 4.1.81, $(A \otimes B)(A^{-1} \otimes B^{-1}) = I_m \otimes I_n = I_{mn}$ and the reverse product is identical.
Structural Commentary: Inverting a Kronecker product inverts each independent factor.
Key Insight: The mixed product property turns the inverse check into two ordinary inverse checks. □

Proposition 4.1.83. *Kronecker Product and Sum Spectra in the Diagonalizable Case*
Let $A = P\Lambda P^{-1}$, $B = QMQ^{-1}$, *where* $\Lambda = \mathrm{diag}(\lambda_1, \ldots, \lambda_n)$ *and* $M = \mathrm{diag}(\mu_1, \ldots, \mu_q)$.
Then $A \otimes B \sim \Lambda \otimes M$; *and* $A \otimes I_q + I_n \otimes B \sim \Lambda \otimes I_q + I_n \otimes M$.
Their diagonal entries are respectively $\lambda_i\mu_j$ *and* $\lambda_i + \mu_j$ *for* $1 \le i \le n$, $1 \le j \le q$.

Proof. By Theorem 4.1.81 and $(P \otimes Q)^{-1} = P^{-1} \otimes Q^{-1}$, $A \otimes B = (P \otimes Q)(\Lambda \otimes M)(P^{-1} \otimes Q^{-1})$.
Similarly, $A \otimes I_q + I_n \otimes B = (P \otimes Q)(\Lambda \otimes I_q + I_n \otimes M)(P^{-1} \otimes Q^{-1})$.
The middle matrices are diagonal with the stated (i, j)-indexed entries.
Structural Commentary: Eigenvectors tensor together, and Kronecker sums add the two scalar eigenvalues.
Key Insight: Diagonalize each factor first; the Kronecker operation becomes entrywise on eigenvalue pairs. □

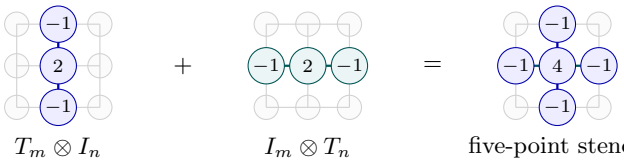


Figure 81: Vertical plus horizontal second differences give the five-point stencil.

Definition 4.1.84. *Vectorisation*
Let $X = [\mathbf{x}_1\ \mathbf{x}_2\ \cdots\ \mathbf{x}_n] \in \mathrm{Mat}(m,n,\mathbb{F})$. The *vectorisation* of X is

$$\mathrm{vec}(X) := \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_n \end{bmatrix} \in \mathbb{F}^{mn}.$$

Thus vec stacks the columns of X vertically.

Theorem 4.1.85. *Vectorisation of a Triple Product*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$, $X \in \mathrm{Mat}(n,p,\mathbb{F})$, and $B \in \mathrm{Mat}(p,q,\mathbb{F})$. Then $\mathrm{vec}(AXB) = (B^t \otimes A)\,\mathrm{vec}(X)$.

Proof. Write $B = [\mathbf{b}_1\ \cdots\ \mathbf{b}_q]$ and $X = [\mathbf{x}_1\ \cdots\ \mathbf{x}_p]$.
The kth column of AXB is $AX\mathbf{b}_k = A\sum_{i=1}^{p} b_{ik}\mathbf{x}_i = \sum_{i=1}^{p} b_{ik}A\mathbf{x}_i = (\mathbf{b}_k^t \otimes A)\,\mathrm{vec}(X)$.
Stacking these q columns gives $(B^t \otimes A)\,\mathrm{vec}(X)$.
Structural Commentary: Vectorisation rewrites two-sided multiplication as one matrix action.
Key Insight: Left multiplication contributes A; right multiplication contributes B^t. □

Perspective 4.1.86. *Vectorising a Minimal-Polynomial Test*
For $A \in \mathrm{Mat}(n,n,\mathbb{F})$ and a monic candidate $p_m(z) = a_0 + a_1 z + \cdots + a_{m-1}z^{m-1} + z^m$,
the matrix equation $a_0 I + a_1 A + \cdots + a_{m-1}A^{m-1} = -A^m$ can be written as the ordinary linear system

$$\begin{bmatrix} \mathrm{vec}(I) & \mathrm{vec}(A) & \cdots & \mathrm{vec}(A^{m-1}) \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{m-1} \end{bmatrix} = -\,\mathrm{vec}(A^m)$$

Structural Commentary: Vectorisation turns matrix identities into linear systems.
Key Insight: Polynomial membership becomes a coefficient system in vectorized powers.

Definition 4.1.87. *Hadamard Product*
Let $A = [a_{ij}], B = [b_{ij}] \in \mathrm{Mat}(m,n,\mathbb{F})$.
The *Hadamard product*, or *entrywise product*, of A and B is $A \circ B := [a_{ij}b_{ij}] \in \mathrm{Mat}(m,n,\mathbb{F})$.

Proposition 4.1.88. *Basic Hadamard Identities*
For matrices of compatible sizes:

(i) $A \circ B = B \circ A$;
(ii) $A \circ (\gamma B + C) = \gamma(A \circ B) + A \circ C$;
(iii) $(A \circ B)^* = A^* \circ B^*$;
(iv) if E is the all-ones matrix, then $A \circ E = A$;
(v) if $D = \mathrm{diag}(d_1,\ldots,d_n)$, then $A \circ D = \mathrm{diag}(a_{11}d_1,\ldots,a_{nn}d_n)$.

Proof.

(i) $(A \circ B)_{ij} = a_{ij}b_{ij} = b_{ij}a_{ij} = (B \circ A)_{ij}$.
(ii) $\big(A \circ (\gamma B + C)\big)_{ij} = a_{ij}(\gamma b_{ij} + c_{ij}) = \gamma a_{ij}b_{ij} + a_{ij}c_{ij}$.
(iii) $\big((A \circ B)^*\big)_{ij} = \overline{a_{ji}b_{ji}} = \overline{a_{ji}}\,\overline{b_{ji}} = (A^* \circ B^*)_{ij}$.
(iv) $(A \circ E)_{ij} = a_{ij}$.
(v) Since D has zero off-diagonal entries, $A \circ D$ is diagonal with entries $a_{ii}d_i$.

Structural Commentary: Hadamard multiplication is entrywise, with no row-column summation.
Key Insight: Prove Hadamard identities entry by entry. □

Theorem 4.1.89. *Schur Product Theorem*
If $A, B \in \mathrm{Mat}(n,n,\mathbb{F})$ are positive semidefinite, where $\mathbb{F} \in \{\mathbb{R},\mathbb{C}\}$, then $A \circ B$ is positive semidefinite.

Proof. By Theorem 4.1.56, regard A and B as Gram matrices: $a_{ij} = \langle \mathbf{u}_j, \mathbf{u}_i \rangle$, $b_{ij} = \langle \mathbf{v}_j, \mathbf{v}_i \rangle$ for suitable vectors $\mathbf{u}_1, \dots, \mathbf{u}_n$ and $\mathbf{v}_1, \dots, \mathbf{v}_n$. Then $(A \circ B)_{ij} = \langle \mathbf{u}_j, \mathbf{u}_i \rangle \langle \mathbf{v}_j, \mathbf{v}_i \rangle = \langle \mathbf{u}_j \otimes \mathbf{v}_j, \mathbf{u}_i \otimes \mathbf{v}_i \rangle$.
Thus $A \circ B$ is the Gram matrix of the vectors $\mathbf{u}_j \otimes \mathbf{v}_j$, hence is positive semidefinite by Proposition 2.1.62.
Structural Commentary: The Schur product is the Gram matrix of paired tensor vectors.
Key Insight: Pair the jth vector from each Gram representation and take tensor products. □

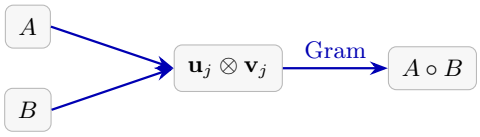


Figure 82: For Gram matrices, entrywise products become tensor inner products, so $A \circ B$ is Gram.

Corollary 4.1.90. *Positive Definite Hadamard Products*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$ be positive semidefinite, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$.

(i) If A is positive definite and every diagonal entry of B is positive, then $A \circ B$ is positive definite.
(ii) If A and B are positive definite, then $A \circ B$ is positive definite.

Proof.

(i) Let $\lambda_{\min} > 0$ be the smallest eigenvalue of A. Then $A - \lambda_{\min} I_n \succeq 0$, so $(A - \lambda_{\min} I_n) \circ B \succeq 0$ by Theorem 4.1.89. Thus, for $\mathbf{x} \neq \mathbf{0}$, $\mathbf{x}^*(A \circ B)\mathbf{x} = \mathbf{x}^*((A - \lambda_{\min} I_n) \circ B)\mathbf{x} \lambda_{\min} \sum_{j=1}^{n} b_{jj} |x_j|^2 > 0$.

(ii) If $B \succ 0$, then $b_{jj} = \mathbf{e}_j^* B \mathbf{e}_j > 0$ for every j, so (i) applies.

Structural Commentary: Subtract the spectral floor of A; the diagonal of B gives strictness.
Key Insight: Use $I_n \circ B = \mathrm{diag}(b_{11}, \dots, b_{nn})$. □

Method 4.1.91. *Schur Product Positivity Method*
To prove an entrywise product is positive semidefinite:

(i) factor or interpret each matrix as a Gram matrix;
(ii) pair the corresponding vectors;
(iii) take tensor products of the paired vectors;
(iv) identify the Hadamard product as the Gram matrix of those tensor products.

Definition 4.1.92. *Spectral Radius and Convergent Matrix*
For $A \in \mathrm{Mat}(n, n, \mathbb{C})$, the *spectral radius* is $\rho(A) := \max\{|\lambda| : \lambda \in \mathrm{spec}(A)\}$.
We say $A^p \to 0$ if every entry of A^p converges to 0 as $p \to \infty$.

Theorem 4.1.93. *Matrix Powers Converge to Zero*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. Then $A^p \to 0$ if and only if $\rho(A) < 1$.

Proof. Choose a Jordan form $A = SJS^{-1}$.
Since S and S^{-1} are fixed, $A^p \to 0$ if and only if every Jordan block power in J^p converges to 0.
($\Rightarrow$) Each block power tends to 0, so Proposition 3.3.4 gives $|\lambda| < 1$ for every eigenvalue. Thus $\rho(A) < 1$.
($\Leftarrow$) If $\rho(A) < 1$, then every eigenvalue satisfies $|\lambda| < 1$. Proposition 3.3.4 makes each block power tend to 0.
Structural Commentary: Similarity reduces matrix-power convergence to Jordan block-power convergence.
Key Insight: The largest eigenvalue modulus controls the asymptotic size of all blocks. □

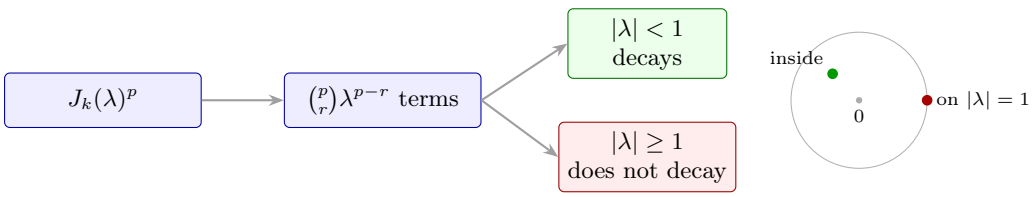


Figure 83: The gray circle is $|\lambda| = 1$: inside it, powers decay; on or outside it, they do not.

Definition 4.1.94. *Power Bounded Matrix*
A matrix $A \in \mathrm{Mat}(n, n, \mathbb{C})$ is *power bounded* if for each entry position (i, j), the sequence of entries of A^p in that position is bounded for $p = 1, 2, \dots$.

Theorem 4.1.95. *Power Bounded Jordan Criterion*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$. The following are equivalent:

(i) *A is power bounded.*
(ii) *$\rho(A) \leq 1$, and every eigenvalue λ with $|\lambda| = 1$ has only 1×1 Jordan blocks.*
(iii) *$\rho(A) \leq 1$, and $\dim E_\lambda(A) = m_A(\lambda)$ for every $\lambda \in \mathrm{spec}(A)$ with $|\lambda| = 1$.*

Proof. Work in Jordan form.

(i) $\Rightarrow$ *(ii)*: If $|\lambda| > 1$, the diagonal entries λ^p are unbounded.
If $|\lambda| = 1$ and $J_k(\lambda)$ has $k \geq 2$, then the first superdiagonal contains $p\lambda^{p-1}$, also unbounded.
(ii) $\Rightarrow$ *(i)*: Blocks with $|\lambda| < 1$ converge to 0 by Proposition 3.3.4; 1×1 blocks with $|\lambda| = 1$ are bounded.
Finite direct sums and fixed similarities preserve boundedness.
(ii) $\Leftrightarrow$ *(iii)*: For fixed λ, $\dim E_\lambda(A)$ counts Jordan blocks, while $m_A(\lambda)$ counts their total size.
Equality holds exactly when every block has size 1.

Structural Commentary: Unit-circle eigenvalues are harmless only when their nilpotent parts vanish.
Key Insight: Nontrivial Jordan blocks on the unit circle create polynomial growth. □

4.2 Matrix Decompositions

Decompositions expose usable structure in a matrix. LU, QR, Cholesky, SVD, compact SVD, and polar decomposition separate elimination, geometry, positivity, rank, and stretching.

Perspective 4.2.1. *Matrix Factorisation Roadmap*
Factorisations expose structure by task: elimination gives LU and LDL^t (Theorem 4.2.5, Corollary 4.2.6); positivity gives Cholesky (Theorem 4.2.12); geometry gives QR (Theorem 4.2.8); singular directions give SVD, compact SVD, and the pseudoinverse (Theorems 4.2.22, 4.2.26, Definition 4.3.16); polar decomposition separates motion from stretch (Theorem 4.2.31). Spectral factorisations appeared earlier: diagonalisation, spectral theorems, Jordan form, and Schur form (Theorems 3.1.39, 3.4.12, 3.4.13, 3.3.6, 3.4.10).
Key Insight: Choose the factorisation that exposes the needed structure.

Definition 4.2.2. *Full-Rank Factorization*
A factorization $A = XY$ with X full column rank and Y full row rank is called a *full-rank factorization* of A.

Theorem 4.2.3. *Full-Rank Factorization*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$ have rank $r \geq 1$, and let the columns of $X \in \mathrm{Mat}(m, r, \mathbb{F})$ be a basis of $\mathrm{col}(A)$. Then there is a unique $Y \in \mathrm{Mat}(r, n, \mathbb{F})$ such that $A = XY$, and $\mathrm{rank}(X) = \mathrm{rank}(Y) = r$.

Proof. Write $A = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_n]$. Since the columns of X form a basis of $\mathrm{col}(A)$, for each j there is a unique $\mathbf{y}_j \in \mathbb{F}^r$ such that $\mathbf{a}_j = X\mathbf{y}_j$. Set $Y = [\mathbf{y}_1 \ \cdots \ \mathbf{y}_n]$. Then $A = [X\mathbf{y}_1 \ \cdots \ X\mathbf{y}_n] = XY$ and uniqueness of the coordinates $\mathbf{y}_j$ gives uniqueness of Y. The columns of X are independent, so $\mathrm{rank}(X) = r$.
By Theorem 4.1.33, $r = \mathrm{rank}(A) = \mathrm{rank}(XY) \leq \mathrm{rank}(Y) \leq r$ because Y has only r rows. Hence $\mathrm{rank}(Y) = r$.
Structural Commentary: A matrix factors through a basis of its column space.
Key Insight: Full-rank factorization separates column directions from their coordinates. □

Definition 4.2.4. *LU Factorisation*
For $A \in \mathrm{Mat}(n, n, \mathbb{F})$, an *LU factorisation* is a product $A = LU$, where L is lower triangular and U is upper triangular. It is *unit lower triangular* if every diagonal entry of L is 1.

Theorem 4.2.5. *LU Factorisation from Nonzero Leading Principal Minors*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$. If every leading principal minor of A is nonzero, then $A = LU$ for a unique unit lower triangular L and a unique upper triangular U.

Proof. Induct on n. The case $n = 1$ is immediate. Write $A = \begin{bmatrix} a & \mathbf{r} \\ \mathbf{c} & B \end{bmatrix}$, $a \neq 0$.
Set $\boldsymbol{\ell} = a^{-1}\mathbf{c}$ and $S = B - a^{-1}\mathbf{c}\mathbf{r}$. For each leading $k \times k$ principal block of A, the block determinant formula gives determinant a times the corresponding leading $(k-1) \times (k-1)$ principal minor of S. Hence all leading principal minors of S are nonzero. By induction, $S = L_1U_1$ with L_1 unit lower triangular and U_1 upper triangular.
Then $A = \begin{bmatrix} 1 & 0 \\ \boldsymbol{\ell} & L_1 \end{bmatrix} \begin{bmatrix} a & \mathbf{r} \\ 0 & U_1 \end{bmatrix}$. For uniqueness, comparing the first row and first column of LU determines a, $\mathbf{r}$, and $\boldsymbol{\ell}$, then reduces uniqueness to the lower-right Schur complement.
Structural Commentary: Gaussian elimination peels off one pivot and passes to a Schur complement.
Key Insight: Nonzero leading principal minors are exactly the no-row-swap pivot condition. □

Corollary 4.2.6. *Symmetric LU Gives an LDL^t Factorisation*
Let $A = A^t \in \mathrm{Mat}(n, n, \mathbb{R})$ have nonzero leading principal minors.
Then $A = LDL^t$, with L unit lower triangular and D diagonal invertible.

Proof. By Theorem 4.2.5, write $A = LU$ uniquely with L unit lower triangular and U upper triangular. Since the leading principal minors are nonzero, each $u_{ii} \neq 0$. Put $D = \mathrm{diag}(u_{11}, \ldots, u_{nn})$ and $U_0 = D^{-1}U$; then U_0 is unit upper triangular and $A = LDU_0$. Since $A = A^t$, also $A = U_0^t DL^t$. This is another LU factorisation with unit lower factor U_0^t, so uniqueness gives $L = U_0^t$. Hence $A = LDL^t$.
Structural Commentary: Symmetry makes the upper triangular unit factor the transpose of the lower one.
Key Insight: Once the pivots are separated into D, the two triangular halves mirror each other. □

Definition 4.2.7. *Wide and Narrow QR Factorisation*
For $A \in \mathrm{Mat}(m, n, \mathbb{F})$ with $m \geq n$, factorisation $A = V \begin{bmatrix} R \\ 0 \end{bmatrix}$, $V \in \mathrm{Mat}(m, m, \mathbb{F})$ unitary, $R \in \mathrm{Mat}(n, n, \mathbb{F})$ upper triangular is called a *wide QR factorisation*. Writing $V = [Q\ Q']$ gives the *narrow QR factorisation* $A = QR$.

Theorem 4.2.8. *QR Factorisation*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$ and suppose $m \geq n$.

(i) $\exists$ a unitary V and an upper triangular R with real nonnegative diagonal entries such that $A = V \begin{bmatrix} R \\ 0 \end{bmatrix}$.
Writing $V = [Q\ Q']$ with Q the first n columns gives $A = QR$, with Q orthonormal columns.
(ii) If $\mathrm{rank}(A) = n$, then narrow $A = QR$ is unique when R has positive diagonal entries.

Proof.

(i) At step k, after the first $k-1$ columns have been triangularized, apply Corollary 2.4.36 to the tail of the kth column. This gives a unitary $I_{k-1} \oplus U_k$ that sends that tail to $r_{kk}\mathbf{e}_1$ with $r_{kk} \geq 0$.
Since the first $k-1$ entries of every earlier column already contain all possible nonzero entries, acting only on rows $k, \ldots, m$ preserves the earlier triangular columns.
After n steps, $U_n \cdots U_1 A = \begin{bmatrix} R \\ 0 \end{bmatrix}$, where R is upper triangular with nonnegative diagonal entries.
Let $V = (U_n \cdots U_1)^*$. Then $A = V \begin{bmatrix} R \\ 0 \end{bmatrix}$, and $V = [Q\ Q']$ gives $A = QR$ with Q orthonormal columns.
If $\mathrm{rank}(A) = n$, then $n = \mathrm{rank}(A) = \mathrm{rank}(R)$, so the triangular matrix R is invertible and its nonnegative diagonal entries are positive.

(ii) Suppose $A = Q_1R_1 = Q_2R_2$, where $Q_i^*Q_i = I_n$ and the R_i are upper triangular with positive diagonal entries. Then $A^*A = R_1^*R_1 = R_2^*R_2$. Let $D = R_1R_2^{-1}$. Substituting $R_1 = DR_2$ gives $D^*D = I_n$.
Thus the columns of D are orthonormal. Since D is upper triangular, the first column has only one possible nonzero entry; orthogonality with it forces the rest of the first row to vanish.
Inducting on the lower-right principal submatrix shows D is diagonal, and $D^*D = I_n$ gives diagonal entries of modulus 1. These diagonal entries are ratios of positive diagonal entries of R_1 and R_2, so they are positive. Hence $D = I_n$. Therefore $R_1 = R_2$, and then $Q_1 = Q_2$.

Structural Commentary: Householder steps triangularise columns; full rank makes the diagonal positive.
Key Insight: Positive diagonals remove the sign/phase ambiguity. □

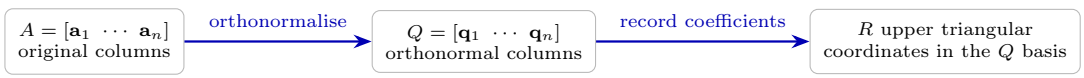


Figure 84: QR writes $A = QR$ while preserving prefix column spaces: $\mathrm{col}(A_k) = \mathrm{col}(Q_k)$.

Perspective 4.2.9. *Gram–Schmidt and Householder QR*
Let $A = [\mathbf{a}_1\ \cdots\ \mathbf{a}_n]$ have full column rank.
Gram–Schmidt gives orthonormal columns $\mathbf{q}_1, \ldots, \mathbf{q}_n$ s.t. $\mathrm{span}(\mathbf{a}_1, \ldots, \mathbf{a}_k) = \mathrm{span}(\mathbf{q}_1, \ldots, \mathbf{q}_k)$, $k = 1, \ldots, n$.
Hence $A = QR$ with $Q^*Q = I$ and R upper triangular. The Householder construction gives the same factorization via unitary transformations, which is the numerically stable form.

Perspective 4.2.10. *Rank-Deficient QR from Gram–Schmidt*
Let $A = [\mathbf{a}_1\ \cdots\ \mathbf{a}_n] \in \mathrm{Mat}(m, n, \mathbb{F})$, $m \geq n$, and let $\mathbf{u}_j$ be the arbitrary-list Gram–Schmidt remainders from Theorem 2.1.38. Set $\mathcal{J} = \{j : \mathbf{u}_j \neq \mathbf{0}\}$. For $j \in \mathcal{J}$ put $\mathbf{q}_j = \mathbf{u}_j / \|\mathbf{u}_j\|$; for $j \notin \mathcal{J}$ choose $\mathbf{q}_j \in \mathrm{col}(A)^\perp$ so that $Q = [\mathbf{q}_1\ \cdots\ \mathbf{q}_n]$ has orthonormal columns. With $R = Q^*A$, one has $A = QR$, R upper triangular, and $r_{jj} = 0$ exactly when $\mathbf{u}_j = \mathbf{0}$. Thus uniqueness fails unless $\mathcal{J} = \{1, \ldots, n\}$ and $r_{jj} > 0$ for all j.

Structural Commentary: Rank defect becomes zero QR pivots.
Key Insight: The full-rank hypothesis in uniqueness is a statement about nonzero orthogonal remainders.

Corollary 4.2.11. *Prefix Column Spaces in QR Factorisation*
Let $A = QR$ be the narrow QR factorisation of $A \in \mathrm{Mat}(m, n, \mathbb{F})$ with $\mathrm{rank}(A) = n$ and positive diagonal entries in R. For $k = 1, \ldots, n$, let A_k and Q_k be the first k columns of A and Q, respectively. Then $\mathrm{col}(A_k) = \mathrm{col}(Q_k)$. Thus the columns of Q_k form an orthonormal basis of $\mathrm{col}(A_k)$.

Proof. Let R_k be the leading $k \times k$ principal submatrix of R. Since R is upper triangular with positive diagonal entries, R_k is invertible. The first k columns of $A = QR$ satisfy $A_k = Q_k R_k$, because upper triangularity prevents later columns of Q from contributing. Hence $\mathrm{col}(A_k) \subseteq \mathrm{col}(Q_k)$ and $Q_k = A_k R_k^{-1}$ gives the reverse inclusion. The columns of Q_k are orthonormal because Q has orthonormal columns.
Structural Commentary: QR preserves prefix column spaces while orthonormalising them.
Key Insight: Invertible R_k only changes coordinates within that space. □

Theorem 4.2.12. *Cholesky Factorisation*
Let $A \in \mathrm{Mat}(n, n, \mathbb{F})$, where $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$. If $A \succeq 0$, then $A = LL^*$ for some lower triangular L with nonnegative diagonal entries. If $A \succ 0$, then L is unique and has positive diagonal entries.

Proof. Let $B = A^{1/2}$ and take a QR factorisation $B = QR$ with R upper triangular and diagonal entries nonnegative. Then $A = B^*B = R^*R$, so $L = R^*$ works.
If $A \succ 0$, then L is invertible, so its diagonal entries are positive. For uniqueness, suppose $A = L_1 L_1^* = L_2 L_2^*$ with L_i lower triangular and positive diagonal. Set $R_i = L_i^*$. Then R_i is upper triangular with positive diagonal and $B^2 = A = R_i^* R_i$. Define $Q_i = B R_i^{-1}$. Then $Q_i^* Q_i = (R_i^{-1})^* B^2 R_i^{-1} = (R_i^{-1})^* R_i^* R_i R_i^{-1} = I_n$, so $B = Q_i R_i$ are QR factorisations with positive diagonal. Uniqueness of QR gives $R_1 = R_2$, hence $L_1 = L_2$.
Structural Commentary: Cholesky is QR applied to $A^{1/2}$.
Key Insight: Triangularise the square root, then set $L = R^*$. □

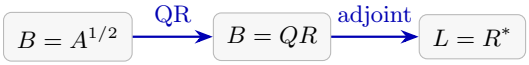


Figure 85: Apply QR to $B = A^{1/2}$; then $A = R^*R = LL^*$, with uniqueness in the positive definite case.

Theorem 4.2.13. *Polar Decomposition*
For every $T \in \mathcal{L}(\mathcal{V})$ on a finite-dimensional inner-product space, $\exists$ an isometry $S \in \mathcal{L}(\mathcal{V})$ s.t. $T = S\sqrt{T^*T}$.

Proof. Let $P = \sqrt{T^*T}$. Since T^*T is positive, P is well defined by Proposition 3.4.42.
For every $\mathbf{v} \in \mathcal{V}$, $\|T\mathbf{v}\|^2 = \langle T^*T\mathbf{v}, \mathbf{v}\rangle = \langle P^2\mathbf{v}, \mathbf{v}\rangle = \|P\mathbf{v}\|^2$.
Define $S_1 : \mathrm{range}(P) \to \mathrm{range}(T)$ by $S_1(P\mathbf{v}) = T\mathbf{v}$. The norm identity shows this is well defined: if $P\mathbf{v}_1 = P\mathbf{v}_2$, then $\|T(\mathbf{v}_1 - \mathbf{v}_2)\| = \|P(\mathbf{v}_1 - \mathbf{v}_2)\| = 0$. It also shows S_1 is an isometry onto $\mathrm{range}(T)$.
The same norm identity gives $\ker(P) = \ker(T)$, so $\mathrm{rank}(P) = \mathrm{rank}(T)$ by rank-nullity. Hence $\mathrm{range}(P)^\perp$ and $\mathrm{range}(T)^\perp$ have the same dimension. Choose any isometry between these complements and combine it orthogonally with S_1 to obtain an isometry S on all of $\mathcal{V}$.
Then $SP\mathbf{v} = S_1(P\mathbf{v}) = T\mathbf{v}$ for every $\mathbf{v}$, so $T = S\sqrt{T^*T}$.
Structural Commentary: The positive factor $\sqrt{T^*T}$ reproduces exactly the lengths produced by T.
Key Insight: Once the metric stretching has been extracted, the remaining map is forced to be an isometry. □

Corollary 4.2.14. *Uniqueness in Operator Polar Decomposition*
Let $T \in \mathcal{L}(\mathcal{V})$ on a finite-dimensional inner-product space, and put $|T| := \sqrt{T^*T}$ and $|T^*| := \sqrt{TT^*}$.

(i) If $T = SP$, $S^*S = I_{\mathcal{V}}$, and $P \succeq 0$, then $P = |T|$. Hence the positive right polar factor is unique.
(ii) $\exists$ left polar decomposition $T = |T^*|\, R$ with R an isometry, and $|T^*|$ is unique positive left polar factor.
(iii) If T is invertible, then both isometric factors are unique: $S = T\,|T|^{-1}$ and $R = |T^*|^{-1}\,T$.

Proof.

(i) If $T = SP$, then $T^*T = PS^*SP = P^2$. Since $P \succeq 0$, Proposition 3.4.42 gives $P = \sqrt{T^*T} = |T|$.
(ii) Apply Theorem 4.2.13 to T^*. There is an isometry S_0 such that $T^* = S_0\,|T^*|$.
Taking adjoints gives $T = |T^*|\, S_0^*$. Set $R = S_0^*$. If also $T = QR$ with $Q \succeq 0$ and R an isometry, then $RR^* = I_{\mathcal{V}}$ in finite dimensions, so $TT^* = QRR^*Q = Q^2$. Hence $Q = |T^*|$.
(iii) If T is invertible, then $|T|$ and $|T^*|$ are positive definite, hence invertible.
Therefore any right and left polar decompositions force $S = T\,|T|^{-1}$, $R = |T^*|^{-1}\,T$.

Structural Commentary: The positive polar factor is canonical.
Key Insight: Invertibility removes the kernel complement, so the whole factorisation becomes unique. □

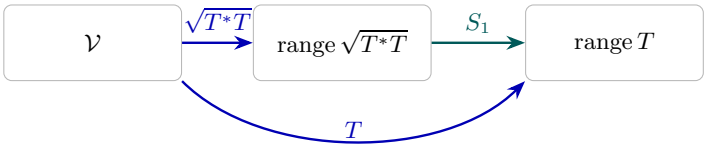


Figure 86: Polar decomposition builds S_1 using $\|T\mathbf{v}\| = \left\|\sqrt{T^*T}\mathbf{v}\right\|$.

Perspective 4.2.15. *Polar Decomposition Separates Stretch and Direction*
In $T = S\sqrt{T^*T}$, the positive factor $\sqrt{T^*T}$ determines the metric stretch and the isometry S determines the direction. Thus polar decomposition is the operator analogue of modulus times phase.
Structural Commentary: For normal T, the polar factors commute and share spectral data.
Key Insight: Polar decomposition factors an operator into positive geometry and isometric motion.

Definition 4.2.16. *Singular Values*
For finite-dimensional inner-product spaces $\mathcal{U}, \mathcal{V}$ and $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$, the *singular values* of T are the eigenvalues of $\sqrt{T^*T} \in \mathcal{L}(\mathcal{U})$, where each eigenvalue λ is repeated $\dim E_\lambda(\sqrt{T^*T})$.

Theorem 4.2.17. *Singular Value Decomposition*
Let $T \in \mathcal{L}(\mathcal{V})$ have singular values $s_1, \dots, s_n$. Then there exist orthonormal bases $(\mathbf{e}_1, \dots, \mathbf{e}_n)$ and $(\mathbf{f}_1, \dots, \mathbf{f}_n)$ of $\mathcal{V}$ such that $T\mathbf{v} = s_1 \langle \mathbf{v}, \mathbf{e}_1 \rangle \mathbf{f}_1 + \cdots + s_n \langle \mathbf{v}, \mathbf{e}_n \rangle \mathbf{f}_n$ for every $\mathbf{v} \in \mathcal{V}$.

Proof. Let $|T| = \sqrt{T^*T}$. Since $|T|$ is positive, the real or complex spectral theorem gives an orthonormal basis $(\mathbf{e}_1, \dots, \mathbf{e}_n)$ with $|T|\,\mathbf{e}_j = s_j \mathbf{e}_j$. Thus $|T|\,\mathbf{v} = \sum_{j=1}^{n} s_j \langle \mathbf{v}, \mathbf{e}_j \rangle \mathbf{e}_j$.
By Theorem 4.2.13, $T = S\,|T|$ for some isometry S. Put $\mathbf{f}_j = S\mathbf{e}_j$. The vectors $\mathbf{f}_1, \dots, \mathbf{f}_n$ form an orthonormal basis, and applying S to the displayed expansion gives the desired formula.
Structural Commentary: SVD diagonalises $|T|$, then transports its eigenvectors by S.
Key Insight: The s_j are stretches; the $\mathbf{f}_j$ are their landing directions. □

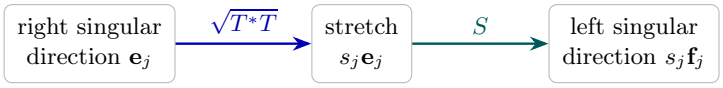


Figure 87: Since $T = S\sqrt{T^*T}$, each $\mathbf{e}_j$ maps to $s_j\mathbf{f}_j$.

Proposition 4.2.18. *Singular Values from T^*T*
The singular values of $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ are the nonnegative square roots of the eigenvalues of T^*T, with the same eigenspace multiplicities.

Proof. By the spectral theorem, choose an orthonormal basis $\mathbf{e}_1, \dots, \mathbf{e}_n$ and nonnegative scalars $\lambda_1, \dots, \lambda_n$ such that $T^*T\mathbf{e}_j = \lambda_j \mathbf{e}_j$. The unique positive square root of T^*T sends $\mathbf{e}_j$ to $\sqrt{\lambda_j}\mathbf{e}_j$.
Thus the singular values are $\sqrt{\lambda_j}$, with the same eigenspace multiplicities.
Structural Commentary: The positive square root keeps eigenspaces and sends λ to $\sqrt{\lambda}$.
Key Insight: Singular values can be read from T^*T directly. □

Theorem 4.2.19. *Singular Vectors for a Linear Map*
Let $\mathcal{U}, \mathcal{V}$ be finite-dimensional inner-product spaces over $\mathbb{R}$ or $\mathbb{C}$, and let $T \in \mathcal{L}(\mathcal{U}, \mathcal{V})$ have rank r.
Put $n = \dim \mathcal{U}$, $m = \dim \mathcal{V}$, and $\mathcal{U}_+ = \operatorname{range}(T^*)$, $\mathcal{U}_0 = \ker(T)$; $\mathcal{V}_+ = \operatorname{range}(T)$, $\mathcal{V}_0 = \ker(T^*)$.
Then $\exists$ ordered orthonormal bases $\beta = (\mathbf{u}_1, \dots, \mathbf{u}_r, \mathbf{u}_{r+1}, \dots, \mathbf{u}_n)$; $\gamma = (\mathbf{v}_1, \dots, \mathbf{v}_r, \mathbf{v}_{r+1}, \dots, \mathbf{v}_m)$
with $\operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_r) = \mathcal{U}_+$; $\operatorname{span}(\mathbf{u}_{r+1}, \dots, \mathbf{u}_n) = \mathcal{U}_0$; $\operatorname{span}(\mathbf{v}_1, \dots, \mathbf{v}_r) = \mathcal{V}_+$; $\operatorname{span}(\mathbf{v}_{r+1}, \dots, \mathbf{v}_m) = \mathcal{V}_0$,
and numbers $s_1, \dots, s_r > 0$ such that $T\mathbf{u}_k = s_k \mathbf{v}_k$; $T^*\mathbf{v}_k = s_k \mathbf{u}_k$, for $1 \le k \le r$.
The s_k are the positive singular values; $\mathbf{u}_k$ and $\mathbf{v}_k$ are right and left singular vectors.

Proof. Let $A = T^*T$. Then $A \succeq 0$, $A = A^*$, and $\ker(A) = \ker(T)$. By spectral theorem, choose rthonormal eigenbasis $\mathbf{u}_1, \dots, \mathbf{u}_n$ s.t. $A\mathbf{u}_k = s_k^2 \mathbf{u}_k$ for $1 \le k \le r$; $A\mathbf{u}_k = 0$ for $k > r$; $s_k > 0$.
Thus $\operatorname{span}(\mathbf{u}_{r+1}, \dots, \mathbf{u}_n) = \mathcal{U}_0$, and Proposition 2.4.11 gives $\operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_r) = \mathcal{U}_+$. For $1 \le k \le r$,

set $\mathbf{v}_k = s_k^{-1} T\mathbf{u}_k$. Then $\langle \mathbf{v}_i, \mathbf{v}_j \rangle = \frac{1}{s_i s_j} \langle T\mathbf{u}_i, T\mathbf{u}_j \rangle = \frac{1}{s_i s_j} \langle T^*T\mathbf{u}_i, \mathbf{u}_j \rangle = \delta_{ij}$; $T^*\mathbf{v}_k = s_k^{-1} A\mathbf{u}_k = s_k \mathbf{u}_k$.
Hence $\mathbf{v}_1, \dots, \mathbf{v}_r$ are orthonormal and span $\mathcal{V}_+ = T(\mathcal{U}_+)$. Complete them by an orthonormal basis of $\mathcal{V}_0 = \mathcal{V}_+^{\perp}$.
Structural Commentary: T^*T gives right directions; T gives left directions.
Key Insight: Away from the two kernels, T and T^* act symmetrically by the same positive scalars. □

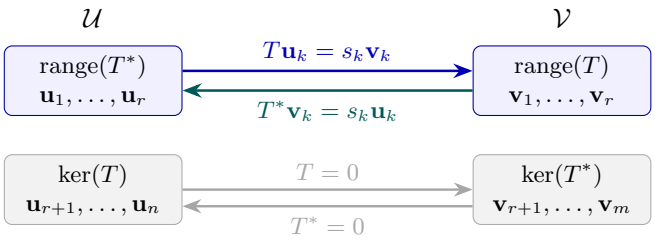


Figure 88: The singular-vector bases split the domain and codomain into active singular subspaces and kernels.

Perspective 4.2.20. *Operator-to-Matrix SVD Bridge*
In the bases from Theorem 4.2.19, the matrix of T is rectangular diagonal with diagonal entries $s_1, \dots, s_r$ and zeros elsewhere. Returning to standard orthonormal bases gives the matrix form $A = V\Sigma W^*$ in Theorem 4.2.22.

Definition 4.2.21. *Matrix Singular Values*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$, let $q = \min(m, n)$, and let $r = \mathrm{rank}(A)$. The *positive singular values* of A are the positive eigenvalues of $|A| = (A^*A)^{1/2}$, written $\sigma_1 \geq \sigma_2 \geq \cdots \geq \sigma_r > 0$.
The full list of singular values is $\sigma_1, \dots, \sigma_r, \sigma_{r+1}, \dots, \sigma_q$, where $\sigma_{r+1} = \cdots = \sigma_q = 0$.

Theorem 4.2.22. *Matrix Singular Value Decomposition*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$ have singular values $\sigma_1 \geq \cdots \geq \sigma_q \geq 0$, where $q = \min(m, n)$. Then $\exists$ unitary matrices $V \in \mathrm{Mat}(m, m, \mathbb{F})$, $W \in \mathrm{Mat}(n, n, \mathbb{F})$ such that $A = V\Sigma W^*$, $\Sigma = \begin{bmatrix} \mathrm{diag}(\sigma_1, \dots, \sigma_q) & 0 \\ 0 & 0 \end{bmatrix} \in \mathrm{Mat}(m, n, \mathbb{R})$ with the evident deletion of empty blocks.

Proof. If $A = 0$, take $\Sigma = 0$ and choose any unitary V, W. Assume $r = \mathrm{rank}(A) \geq 1$.
Diagonalise A^*A as $W^*A^*AW = \mathrm{diag}(\sigma_1^2, \dots, \sigma_r^2, 0, \dots, 0)$, where $\mathbf{w}_1, \dots, \mathbf{w}_n$ are the columns of W.
For $1 \leq j \leq r$, set $\mathbf{v}_j = \sigma_j^{-1} A\mathbf{w}_j$. Then $\mathbf{v}_i^*\mathbf{v}_j = \frac{1}{\sigma_i \sigma_j} \mathbf{w}_i^* A^* A \mathbf{w}_j = \delta_{ij}$ so $\mathbf{v}_1, \dots, \mathbf{v}_r$ are orthonormal.
Complete them to an orthonormal basis of $\mathbb{F}^m$ and let V be the matrix with these columns.
For $j > r$, $\|A\mathbf{w}_j\|_2^2 = \mathbf{w}_j^* A^* A \mathbf{w}_j = 0$, hence $A\mathbf{w}_j = 0$. Thus $AW = V\Sigma$, equivalently $A = V\Sigma W^*$.
Structural Commentary: A^*A determines the right singular directions and singular values.
Key Insight: Apply A and normalize to get the left singular directions. □

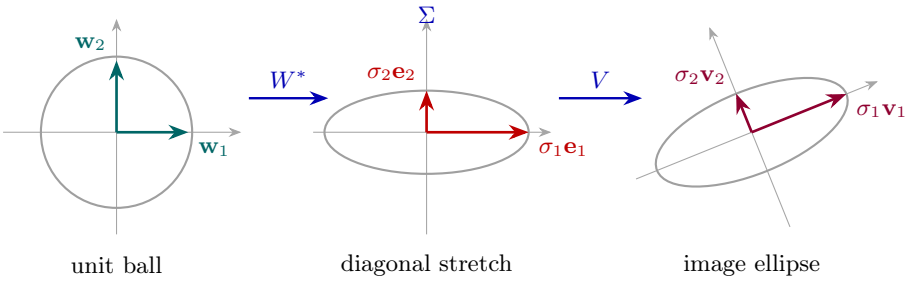


Figure 89: SVD: right-unitary motion, diagonal stretch, left-unitary motion.

Proposition 4.2.23. *Basic Singular Value Identities*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$ have singular values $\sigma_1 \geq \cdots \geq \sigma_q \geq 0$, where $q = \min(m, n)$ and $r = \mathrm{rank}(A)$. Let $c \in \mathbb{F}$.

(i) The positive eigenvalues of A^*A and AA^* are $\sigma_1^2, \dots, \sigma_r^2$.
(ii) $\sum_{j=1}^{q} \sigma_j^2 = \mathrm{tr}(A^*A) = \mathrm{tr}(AA^*) = \|A\|_F^2$, where $\|A\|_F^2 := \mathrm{tr}(A^*A)$.
(iii) A and A^* have the same singular values. Over $\mathbb{C}$, so do A^t and $\overline{A}$.
(iv) The singular values of cA are $|c|\,\sigma_1, \dots, |c|\,\sigma_q$.

Proof. Let $A = V\Sigma W^*$ be an SVD. Then $A^*A = W(\Sigma^*\Sigma)W^*$ and $AA^* = V(\Sigma\Sigma^*)V^*$.
(i) The nonzero diagonal entries of $\Sigma^*\Sigma$ and $\Sigma\Sigma^*$ are $\sigma_1^2, \dots, \sigma_r^2$.

(ii) Since A^*A and $\Sigma^*\Sigma$ are unitarily similar, and AA^* and $\Sigma\Sigma^*$ are unitarily similar, $\operatorname{tr}(A^*A) = \operatorname{tr}(\Sigma^*\Sigma) = \sum_{j=1}^{q} \sigma_j^2 = \operatorname{tr}(\Sigma\Sigma^*) = \operatorname{tr}(AA^*)$. The final equality is the definition of $\|A\|_F$.

(iii) For A^*, the two Gram matrices are interchanged. Over $\mathbb{C}$, $(A^t)^*A^t = (AA^*)^t$ and $(\overline{A})^*\overline{A} = (A^*A)^t$. Since a matrix and its transpose have the same characteristic polynomial, the spectra agree.

(iv) Since $(cA)^*(cA) = |c|^2 A^*A$, singular values scale by $|c|$.

Structural Commentary: Unitary factors do not change Gram spectra; Σ contains all singular data.
Key Insight: Reduce each claim to $\Sigma^*\Sigma$ and $\Sigma\Sigma^*$. □

Corollary 4.2.24. *Square Invertibility and Singular Values*
Let $A \in \operatorname{Mat}(n, n, \mathbb{F})$, $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, have singular values $\sigma_1 \geq \cdots \geq \sigma_n \geq 0$. The following are equivalent:

(i) A is invertible;
(ii) $\operatorname{rank}(A) = n$;
(iii) $\sigma_n > 0$;
(iv) $A^*A \succ 0$.

If $r = \operatorname{rank}(A) < n$, then $\sigma_1, \ldots, \sigma_r > 0$ and $\sigma_{r+1} = \cdots = \sigma_n = 0$.

Proof.

(i) ⇒ *(ii)* By Corollary 1.3.95, invertibility gives full rank.
(ii) ⇒ *(iii)* By Definition 4.2.21, the number of positive singular values is $\operatorname{rank}(A)$.
(iii) ⇒ *(iv)* If $\sigma_n > 0$, then A has full column rank, so Proposition 4.1.51 gives $A^*A \succ 0$.
(iv) ⇒ *(i)* If $A\mathbf{x} = \mathbf{0}$, then $\mathbf{x}^*A^*A\mathbf{x} = 0$; PSD gives $\mathbf{x} = \mathbf{0}$. Thus $\ker(A) = \{\mathbf{0}\}$, so A is invertible.

The rank-defect statement follows again from Definition 4.2.21.
Structural Commentary: SVD sees nonsingularity as the absence of zero stretch directions.
Key Insight: A zero singular value is exactly a nonzero vector sent to zero. □

Definition 4.2.25. *Compact Singular Value Decomposition*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ have rank $r \geq 1$.
A factorisation $A = X\Sigma_r Y^*$ is a *compact singular value decomposition* if $X \in \operatorname{Mat}(m, r, \mathbb{F})$ and $Y \in \operatorname{Mat}(n, r, \mathbb{F})$ have orthonormal columns and $\Sigma_r = \operatorname{diag}(\sigma_1, \ldots, \sigma_r)$ with $\sigma_1 \geq \cdots \geq \sigma_r > 0$.

Theorem 4.2.26. *Compact Singular Value Decomposition*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ have rank $r \geq 1$.

(i) If $A = V\Sigma W^*$ is an SVD and X, Y are the first r columns of V, W, then $A = X\Sigma_r Y^*$.
(ii) Conversely, if X, Y have orthonormal columns and $\Sigma_r = \operatorname{diag}(\sigma_1, \ldots, \sigma_r)$ with $\sigma_1 \geq \cdots \geq \sigma_r > 0$, then $X\Sigma_r Y^*$ has positive singular values $\sigma_1, \ldots, \sigma_r$.

Proof.

(i) Partition $V = [X\ X_0]$, $W = [Y\ Y_0]$, and $\Sigma = \begin{bmatrix} \Sigma_r & 0 \\ 0 & 0 \end{bmatrix}$. Multiplication gives $A = V\Sigma W^* = X\Sigma_r Y^*$.

(ii) Let $B = X\Sigma_r Y^*$. Since $X^*X = Y^*Y = I_r$, $B^*B = Y\Sigma_r^2 Y^*$.
Completing Y to a unitary matrix $\widetilde{Y} = [Y\ Y_0]$ gives $\widetilde{Y}^*B^*B\widetilde{Y} = \Sigma_r^2 \oplus 0$.
Thus the positive eigenvalues of B^*B are $\sigma_1^2, \ldots, \sigma_r^2$, so the positive singular values of B are $\sigma_1, \ldots, \sigma_r$.

Structural Commentary: Compact SVD keeps only the nonzero singular directions.
Key Insight: The rank-r data X, Σ_r, Y reconstructs A. □

Theorem 4.2.27. *Column and Kernel Spaces from Compact SVD*
Let $A = X\Sigma_r Y^*$ be a compact SVD. Then $\operatorname{col}(A) = \operatorname{col}(X) = (\ker A^*)^\perp$, and $\ker(A) = \operatorname{col}(Y)^\perp = (\operatorname{col}(A^*))^\perp$.

Proof. Since $A = X\Sigma_r Y^*$ and $AY\Sigma_r^{-1} = X$, $\operatorname{col}(A) = \operatorname{col}(X)$.
For $\mathbf{z} \in \mathbb{F}^n$, $A\mathbf{z} = 0 \iff X\Sigma_r Y^*\mathbf{z} = 0 \iff Y^*\mathbf{z} = 0 \iff \mathbf{z} \perp \operatorname{col}(Y)$, because X has independent columns and Σ_r is invertible. The orthogonal-complement identities follow from Proposition 2.4.11.
Structural Commentary: Compact SVD exposes image directions X and null directions $\operatorname{col}(Y)^\perp$.
Key Insight: A acts only on the right singular subspace $\operatorname{col}(Y)$. □

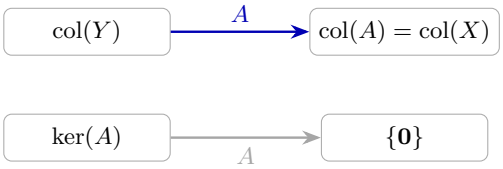


Figure 90: Compact SVD: A maps $\mathrm{col}(Y)$ onto $\mathrm{col}(X)$ and kills $\ker(A) = \mathrm{col}(Y)^{\perp}$.

Theorem 4.2.28. *Uniqueness of Compact SVD up to Singular Subspaces*
Let $A = X_1\Sigma_r Y_1^* = X_2\Sigma_r Y_2^*$ be compact SVDs with the same Σ_r. Then $X_1 = X_2U$, $Y_1 = Y_2U$ for some unitary U satisfying $U\Sigma_r = \Sigma_r U$. Hence $U = U_1 \oplus \cdots \oplus U_d$ along the equal singular-value blocks of Σ_r.

Proof. Theorem 4.2.27 gives $\mathrm{col}(X_1) = \mathrm{col}(X_2)$ and $\mathrm{col}(Y_1) = \mathrm{col}(Y_2)$. Since each X_i and Y_i has orthonormal columns, there are unitary $U, V \in \mathrm{Mat}(r, r, \mathbb{F})$ such that $X_1 = X_2U$ and $Y_1 = Y_2V$. Substitution and multiplication on the left by X_2^* and on the right by Y_2 give $\Sigma_r = U\Sigma_r V^*$, hence $U\Sigma_r = \Sigma_r V$.
Therefore $U\Sigma_r^2 U^* = \Sigma_r^2$, so U commutes with Σ_r^2. Lemma 3.1.25 makes U block diagonal on equal singular-value blocks. Hence $U\Sigma_r = \Sigma_r U$, and $U\Sigma_r = \Sigma_r V$ forces $V = U$ because Σ_r is invertible.
Structural Commentary: Compact SVD is unique up to unitary changes inside equal singular-value blocks.
Key Insight: Distinct singular values fix subspaces; repeated ones only fix the subspace, not its basis. □

Definition 4.2.29. *Matrix Modulus*
For $A \in \mathrm{Mat}(m, n, \mathbb{F})$, the *modulus* of A is $|A| := (A^*A)^{1/2} \in \mathrm{Mat}(n, n, \mathbb{F})$.
Thus $|A^*| = (AA^*)^{1/2} \in \mathrm{Mat}(m, m, \mathbb{F})$.

Proposition 4.2.30. *Modulus from SVD*
Let $A = V\Sigma W^*$ be an SVD. Then $|A| = W(\Sigma^*\Sigma)^{1/2}W^*$ and $|A^*| = V(\Sigma\Sigma^*)^{1/2}V^*$.

Proof. Since $A^*A = W(\Sigma^*\Sigma)W^*$ and $AA^* = V(\Sigma\Sigma^*)V^*$, the displayed matrices are positive semidefinite square roots of A^*A and AA^*, respectively. Uniqueness of the positive semidefinite square root gives the claim.
Structural Commentary: The modulus keeps only the singular-value stretch.
Key Insight: In an SVD, $|A|$ is the right singular-vector version of the diagonal stretch. □

Theorem 4.2.31. *Rectangular Polar Decomposition*
Let $A = V\Sigma W^* \in \mathrm{Mat}(m, n, \mathbb{F})$ be an SVD.

(i) If $m \geq n$, let V_n be the first n columns of V and set $U = V_nW^*$. Then $U^*U = I_n$ and $A = U|A|$.
(ii) If $m \leq n$, let W_m be the first m columns of W and set $U = VW_m^*$. Then $UU^* = I_m$ and $A = |A^*|U$.
(iii) If $m = n$, then $U = VW^*$ is unitary and $A = U|A| = |A^*|U$.

Proof.

(i) Proposition 4.2.30 gives $|A| = W(\Sigma^*\Sigma)^{1/2}W^*$.
Since $V_n(\Sigma^*\Sigma)^{1/2} = V\Sigma$, $U|A| = V_nW^*W(\Sigma^*\Sigma)^{1/2}W^* = V\Sigma W^* = A$. Also $U^*U = WV_n^*V_nW^* = I_n$.
(ii) Apply (i) to A^* with the SVD $A^* = W\Sigma^*V^*$, obtaining $A^* = U_0|A^*|$ with $U_0 = W_mV^*$ and $U_0^*U_0 = I_m$.
Taking adjoints gives $A = |A^*|U_0^* = |A^*|VW_m^*$, and $U_0^*(U_0^*)^* = U_0^*U_0 = I_m$. Thus $U = VW_m^*$ works.
(iii) This is the common case of (i) and (ii), with $U = VW^*$ unitary.

Structural Commentary: Polar decomposition compresses the two SVD unitary factors into one partial isometry.
Key Insight: SVD records every singular direction; polar form records stretch plus motion. □

Method 4.2.32. *SVD and Polar Computation Method*
To compute an SVD and the associated polar factors of $A \in \mathrm{Mat}(m, n, \mathbb{F})$:

(i) Diagonalise A^*A; the square roots of its positive eigenvalues are the positive singular values.
(ii) Choose orthonormal right singular vectors $\mathbf{w}_j$ for the positive eigenvalues and set $\mathbf{v}_j = \sigma_j^{-1}A\mathbf{w}_j$.
(iii) Complete the $\mathbf{v}_j$ and $\mathbf{w}_j$ to unitary matrices V, W to obtain $A = V\Sigma W^*$.
Omit zero-singular directions for compact SVD.
(iv) Use $|A| = W(\Sigma^*\Sigma)^{1/2}W^*$ and merge V, W as in Theorem 4.2.31 to obtain the polar factor.

Theorem 4.2.33. *Uniqueness of the Positive Polar Factor*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$.

(i) If $m \geq n$ and $A = UB$ with $U^*U = I_n$ and $B \succeq 0$, then $B = |A|$.
If $\mathrm{rank}(A) = n$, then $U = A|A|^{-1}$ is unique.
(ii) If $m \leq n$ and $A = CU$ with $UU^* = I_m$ and $C \succeq 0$, then $C = |A^*|$.
If $\mathrm{rank}(A) = m$, then $U = |A^*|^{-1}A$ is unique.

Proof.

(i) $A^*A = BU^*UB = B^2$, $B = |A|$ by uniqueness of positive square root. If $\operatorname{rank}(A) = n$, then $\ker(A) = \{\mathbf{0}\}$. Since $\mathbf{x}^*A^*A\mathbf{x} = \|A\mathbf{x}\|_2^2$, this gives $\ker(A^*A) = \{\mathbf{0}\}$ and $|A|$ is invertible. Thus $U = A\,|A|^{-1}$.

(ii) $AA^* = CUU^*C = C^2$, so $C = |A^*|$. If $\operatorname{rank}(A) = m$, then A^* has full column rank, so $\ker(A^*) = \{\mathbf{0}\}$. Since $\mathbf{y}^*AA^*\mathbf{y} = \|A^*\mathbf{y}\|_2^2$, this gives $\ker(AA^*) = \{\mathbf{0}\}$ and $|A^*|$ is invertible. Thus $U = |A^*|^{-1}\,A$.

Structural Commentary: The positive factor is fixed by the Gram matrix.
Key Insight: Full rank makes the partial isometry unique. □

Theorem 4.2.34. *Normality and Polar Commutation*
Let $A \in \operatorname{Mat}(n, n, \mathbb{F})$ and suppose $A = U\,|A|$ with U unitary. Then A is normal if and only if $U\,|A| = |A|\,U$.

Proof.

($\Rightarrow$) If A is normal, then $|A|^2 = A^*A = AA^* = U\,|A|^2\,U^*$. Thus $U\,|A|^2\,U^* = |A|^2$. The matrix $U\,|A|\,U^*$ is PSD and squares to $|A|^2$, so uniqueness of positive square root gives $U\,|A|\,U^* = |A|$. Hence $U\,|A| = |A|\,U$.

($\Leftarrow$) If $U\,|A| = |A|\,U$, then $AA^* = U\,|A|^2\,U^* = |A|^2 = A^*A$.

Structural Commentary: Normality says that the polar motion preserves the stretching eigenspaces.
Key Insight: Normality is commutation of polar motion and stretch. □

Proposition 4.2.35. *Absolute Determinant Comes from the Positive Polar Factor*
For $T \in \mathcal{L}(\mathcal{V})$ on a finite-dimensional inner-product space, $|\det T| = \det\sqrt{T^*T}$.

Proof. By Theorem 4.2.13, $T = S\sqrt{T^*T}$ for some isometry S. Hence $|\det T| = |\det S|\left|\det\sqrt{T^*T}\right| = \det\sqrt{T^*T}$, where Proposition 4.1.25 gives $|\det S| = 1$, and positivity of $\sqrt{T^*T}$ gives nonnegative determinant.
Structural Commentary: Polar decomposition isolates stretch; only the positive factor changes volume.
Key Insight: Determinant sign or phase belongs to the isometric part; volume belongs to the positive part. □

4.3 Matrix Norms and Conditioning

Matrix norms measure stretch, energy, approximation quality, and sensitivity. Singular values organize norms, low-rank approximation, pseudoinverses, conditioning, and error bounds.

Definition 4.3.1. *Spectral Norm and Extreme Singular Values*
For $A \in \operatorname{Mat}(m, n, \mathbb{F})$, define the *spectral norm* $\|A\|_2 := \max_{\|\mathbf{x}\|_2 = 1} \|A\mathbf{x}\|_2$.

If the singular values of A are $\sigma_1 \geq \cdots \geq \sigma_q \geq 0$, $q = \min(m, n)$, set $\sigma_{\max}(A) := \sigma_1$.
When A is square, $\sigma_{\min}(A) := \sigma_n$.

Theorem 4.3.2. *Spectral Norm is the Largest Singular Value*
Let $A \in \operatorname{Mat}(m, n, \mathbb{F})$ have singular values $\sigma_1 \geq \cdots \geq \sigma_q \geq 0$. Then $\|A\|_2 = \sigma_{\max}(A)$.
Moreover, if $A = V\Sigma W^*$ is an SVD and $\mathbf{w}_1$ is the first column of W, then $\|A\mathbf{w}_1\|_2 = \sigma_1$.

Proof. Let $\mathbf{w}_1, \ldots, \mathbf{w}_n$ be the columns of W. For $\mathbf{x} = \sum_{j=1}^n c_j\mathbf{w}_j$ with $\|\mathbf{x}\|_2 = 1$, unitary invariance of the Euclidean norm gives $\|A\mathbf{x}\|_2^2 = \|\Sigma W^*\mathbf{x}\|_2^2 = \sum_{j=1}^{q} \sigma_j^2\,|c_j|^2 \leq \sigma_1^2 \sum_{j=1}^{n} |c_j|^2 = \sigma_1^2$.
Thus $\|A\|_2 \leq \sigma_1$. Taking $\mathbf{x} = \mathbf{w}_1$ gives equality.
Structural Commentary: The right singular basis diagonalises the stretching computation.
Key Insight: The largest stretch occurs in a top right singular direction. □

Proposition 4.3.3. *Basic Properties of the Spectral Norm*
For compatible matrices A, B and scalar c:

(i) $\|A\|_2 \geq 0$, *and* $\|A\|_2 = 0$ *if and only if* $A = 0$;
(ii) $\|cA\|_2 = |c|\,\|A\|_2$;
(iii) $\|A + B\|_2 \leq \|A\|_2 + \|B\|_2$;
(iv) $\|AB\|_2 \leq \|A\|_2\,\|B\|_2$.

Proof.

(i) Since $\|A\mathbf{x}\|_2 \geq 0$, $\|A\|_2 \geq 0$. Also $\|A\|_2 = 0$ iff $A\mathbf{x} = 0$ for every unit $\mathbf{x}$, iff $A = 0$.
(ii) $\|cA\|_2 = \max_{\|\mathbf{x}\|_2=1} \|cA\mathbf{x}\|_2 = |c|\,\|A\|_2$.
(iii) For $\|\mathbf{x}\|_2 = 1$, $\|(A + B)\mathbf{x}\|_2 \leq \|A\mathbf{x}\|_2 + \|B\mathbf{x}\|_2 \leq \|A\|_2 + \|B\|_2$. Maximize over unit $\mathbf{x}$.
(iv) For $\|\mathbf{x}\|_2 = 1$, $\|AB\mathbf{x}\|_2 \leq \|A\|_2\,\|B\mathbf{x}\|_2 \leq \|A\|_2\,\|B\|_2$. Maximize over unit $\mathbf{x}$.

Structural Commentary: The spectral norm is the operator norm induced by Euclidean geometry.
Key Insight: Bound each output vector first, then take the worst case. □

Definition 4.3.4. *Operator Norm on $\mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$*
For $A \in \mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$, define $\|A\|_{\text{op}} := \sup_{\|\mathbf{h}\|_2 \le 1} \|A\mathbf{h}\|_2$.
Equivalently, after standard bases identify A with its matrix, $\|A\|_{\text{op}}$ is the spectral norm from Definition 4.3.1.
Thus $\|A\mathbf{h}\|_2 \le \|A\|_{\text{op}} \|\mathbf{h}\|_2$, and sums and products obey the norm estimates in Proposition 4.3.3.
The induced metric on $\mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$ is $d(A, B) := \|A - B\|_{\text{op}}$; this is the metric used for continuity of derivative maps in Definition 4.2.12.

Proposition 4.3.5. *Frobenius and Spectral Norm Comparison*
Let $A \in \text{Mat}(m, n, \mathbb{F})$ have singular values $\sigma_1 \ge \cdots \ge \sigma_q \ge 0$. Then $\|A\|_F^2 = \sum_{j=1}^q \sigma_j^2$ and $\|A\|_2 \le \|A\|_F$.
If $A \neq 0$, equality $\|A\|_2 = \|A\|_F$ holds if and only if $\text{rank}(A) = 1$.

Proof. The Frobenius identity is Proposition 4.2.23. Hence $\|A\|_2^2 = \sigma_1^2 \le \sum_{j=1}^q \sigma_j^2 = \|A\|_F^2$. Assume $A \neq 0$.
($\Rightarrow$) If $\|A\|_2 = \|A\|_F$, then $\sum_{j=2}^q \sigma_j^2 = 0$, so only σ_1 is nonzero and $\text{rank}(A) = 1$.
($\Leftarrow$) If $\text{rank}(A) = 1$, then $\sigma_2 = \cdots = \sigma_q = 0$, so $\|A\|_2 = \|A\|_F$.
Structural Commentary: Frobenius counts all singular energy; spectral keeps the largest part.
Key Insight: Equality means exactly one nonzero singular value. □

Theorem 4.3.6. *Submultiplicativity of Frobenius and Spectral Norms*
Let $A \in \text{Mat}(m, n, \mathbb{F})$, $B \in \text{Mat}(n, k, \mathbb{F})$. Then $\|AB\|_F \le \|A\|_2 \|B\|_F \le \|A\|_F \|B\|_F$ and $\|AB\|_2 \le \|A\|_2 \|B\|_2$.

Proof. Write $B = [\mathbf{b}_1 \ \cdots \ \mathbf{b}_k]$. Then $\|AB\|_F^2 = \sum_{j=1}^k \|A\mathbf{b}_j\|_2^2 \le \|A\|_2^2 \sum_{j=1}^k \|\mathbf{b}_j\|_2^2 = \|A\|_2^2 \|B\|_F^2$.
Proposition 4.3.5 gives $\|A\|_2 \le \|A\|_F$. The spectral-norm inequality is Proposition 4.3.3.
Structural Commentary: Frobenius norm is column energy; each column of AB is A applied to a column of B.
Key Insight: Control every column by the spectral norm of A. □

Proposition 4.3.7. *Unitary Invariance of Singular Values*
Let $A \in \text{Mat}(m, n, \mathbb{F})$, and let $U \in \text{Mat}(m, m, \mathbb{F})$ and $V \in \text{Mat}(n, n, \mathbb{F})$ be unitary. Then A and UAV have the same singular values. In particular, $\|UAV\|_2 = \|A\|_2$ and $\|UAV\|_F = \|A\|_F$.

Proof. Since $(UAV)^*(UAV) = V^*A^*AV$, the Gram matrices of A and UAV are unitarily similar, hence have the same eigenvalues. Singular values are the nonnegative square roots of these eigenvalues. The two norm identities follow from Theorem 4.3.2 and Proposition 4.3.5.
Structural Commentary: Left and right unitary factors change coordinates, not singular-value data.
Key Insight: Singular values are spectral data of A^*A, and unitary conjugation preserves that spectrum. □

Example 4.3.8. *Similarity Need Not Preserve Singular Values*
Let $A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$ and $B = \begin{bmatrix} 1 & 2 \\ 0 & 1 \end{bmatrix}$. Then $B = DAD^{-1}$ with $D = \text{diag}(2, 1)$, so A and B are similar.
However $A^*A = \begin{bmatrix} 1 & 1 \\ 1 & 2 \end{bmatrix}$, $B^*B = \begin{bmatrix} 1 & 2 \\ 2 & 5 \end{bmatrix}$ have different eigenvalues, so A and B have different singular values.
Singular values are preserved by unitary left-right equivalence, not by arbitrary similarity.

Theorem 4.3.9. *Schur's Inequality and Normality*
Let $A \in \text{Mat}(n, n, \mathbb{C})$ have eigenvalues $\lambda_1, \ldots, \lambda_n$, counted algebraically.
Then $\sum_{j=1}^n |\lambda_j|^2 \le \|A\|_F^2$ with equality if and only if A is normal.

Proof. Let $A = UTU^*$ be a Schur form, where $T = [t_{ij}]$ is upper triangular and $t_{jj} = \lambda_j$.
Then $\|A\|_F^2 = \text{tr}(T^*T) = \sum_{j=1}^n |\lambda_j|^2 + \sum_{i<j} |t_{ij}|^2$. Inequality holds. Equality holds exactly when T is diagonal.
($\Rightarrow$) If equality holds, then T is diagonal, so A is unitarily diagonalizable and hence normal.
($\Leftarrow$) If A is normal, then T is normal and upper triangular, hence diagonal by Lemma 3.4.26.
Structural Commentary: Schur form splits Frobenius norm into eigenvalue mass plus off-diagonal mass.
Key Insight: Normality means the Schur off-diagonal part vanishes. □

Definition 4.3.10. *Defect from Normality*
Let $A \in \text{Mat}(n, n, \mathbb{C})$ have eigenvalues $\lambda_1, \ldots, \lambda_n$, counted with algebraic multiplicity.
The *defect from normality* of A is $\Delta(A) := \|A\|_F^2 - \sum_{j=1}^n |\lambda_j|^2$.
By Theorem 4.3.9, $\Delta(A) \ge 0$, with equality if and only if A is normal.

Theorem 4.3.11. *Eigenvalues Versus Singular Values*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$ have eigenvalues $\lambda_1,\dots,\lambda_n$ ordered so that $|\lambda_1| \geq \cdots \geq |\lambda_n|$, and singular values $\sigma_1 \geq \cdots \geq \sigma_n \geq 0$. Then:

(i) $|\det A| = \prod_{j=1}^n |\lambda_j| = \prod_{j=1}^n \sigma_j$;
(ii) $\rho(A) \leq \|A\|_2$;
(iii) $\sigma_n \leq |\lambda_n|$;
(iv) A *is normal if and only if* $|\lambda_j| = \sigma_j$ *for every* j;
(v) A *is positive semidefinite if and only if* $\lambda_j = \sigma_j$ *for every* j.

Proof.

(i) The eigenvalue product gives $|\det A| = \prod_j |\lambda_j|$.
If $A = V\Sigma W^*$ is an SVD, then $|\det V| = |\det W| = 1$, so $|\det A| = |\det \Sigma| = \prod_j \sigma_j$.
(ii) If $A\mathbf{x} = \lambda\mathbf{x}$ with $\|\mathbf{x}\|_2 = 1$, then $|\lambda| = \|A\mathbf{x}\|_2 \leq \|A\|_2$. Maximize over eigenvalues.
(iii) If A is singular, both sides are 0 or the inequality is immediate.
If A is invertible, apply (ii) to A^{-1}: $\frac{1}{|\lambda_n|} = \rho(A^{-1}) \leq \|A^{-1}\|_2 = \frac{1}{\sigma_n}$.
(iv) If A is normal, the spectral theorem gives $\sigma_j = |\lambda_j|$. Conversely, equality gives $\sum_j |\lambda_j|^2 = \sum_j \sigma_j^2 = \|A\|_F^2$, so Theorem 4.3.9 gives normality.
(v) If $A \succeq 0$, then A is normal and $\lambda_j \geq 0$, so (iv) gives $\lambda_j = \sigma_j$. Conversely, equality gives normality with nonnegative real eigenvalues, hence $A \succeq 0$ by Theorem 4.1.56.

Structural Commentary: Normality aligns eigenvalue moduli with singular values; positivity removes the moduli.
Key Insight: Normality is the bridge that turns spectral data into singular-value data. □

Theorem 4.3.12. *Trace Bound by Singular Values*
If $A \in \mathrm{Mat}(n,n,\mathbb{C})$, then $|\operatorname{tr} A| \leq \sum_{j=1}^n \sigma_j(A)$. Consequently, $\sum_{j=1}^n |\lambda_j(A)| \leq \sum_{j=1}^n \sigma_j(A)$ and $\sum_{j=1}^n |a_{jj}| \leq \sum_{j=1}^n \sigma_j(A)$.

Proof. Let $A = V\Sigma W^*$ be an SVD.
Writing $\mathbf{v}_k, \mathbf{w}_k$ for the columns of V, W, $\operatorname{tr} A = \sum_{k=1}^n \sigma_k \mathbf{w}_k^* \mathbf{v}_k$, $|\mathbf{w}_k^* \mathbf{v}_k| \leq 1$, so $|\operatorname{tr} A| \leq \sum_k \sigma_k$.
For eigenvalues, take a Schur form T of A and a diagonal unitary D with $\operatorname{tr}(DT) = \sum_j |\lambda_j|$.
For diagonal entries, take diagonal unitary D with $\operatorname{tr}(DA) = \sum_j |a_{jj}|$.
In both cases, unitary invariance of singular values and the trace bound give the result.
Structural Commentary: Phase rotations turn trace estimates into eigenvalue and diagonal estimates.
Key Insight: Prove $|\operatorname{tr} A| \leq \sum_j \sigma_j$ once, then rotate phases. □

Proposition 4.3.13. *Rank-One SVD Terms and Truncation Error*
Let $A = X\Sigma_r Y^* = \sum_{j=1}^r \sigma_j \mathbf{x}_j \mathbf{y}_j^*$ be a compact SVD, with $\sigma_1 \geq \cdots \geq \sigma_r > 0$.
For $0 \leq k \leq r$, set $A_k := \sum_{j=1}^k \sigma_j \mathbf{x}_j \mathbf{y}_j^*$, $A_0 := 0$.
Then the rank-one summands are mutually orthogonal in the Frobenius inner product, and

$$\|A\|_F^2 = \sum_{j=1}^r \sigma_j^2, \qquad \|A - A_k\|_F^2 = \sum_{j=k+1}^r \sigma_j^2, \qquad \|A - A_k\|_2 = \begin{cases} \sigma_{k+1}, & k < r, \\ 0, & k = r. \end{cases}$$

Proof. For $i,j \in \{1,\dots,r\}$, $\langle \sigma_i \mathbf{x}_i \mathbf{y}_i^*, \sigma_j \mathbf{x}_j \mathbf{y}_j^* \rangle_F = \sigma_i \sigma_j \operatorname{tr}(\mathbf{y}_i \mathbf{x}_i^* \mathbf{x}_j \mathbf{y}_j^*) = \sigma_i \sigma_j \delta_{ij}$.
Pythagoras in the Frobenius inner product gives the two Frobenius formulas. The residual $A - A_k$ has compact SVD $\sum_{j=k+1}^r \sigma_j \mathbf{x}_j \mathbf{y}_j^*$, so Theorem 4.3.2 gives the spectral-norm formula.
Structural Commentary: Compact SVD decomposes a matrix into orthogonal rank-one energy packets.
Key Insight: Truncation error is exactly the energy left in the discarded singular directions. □

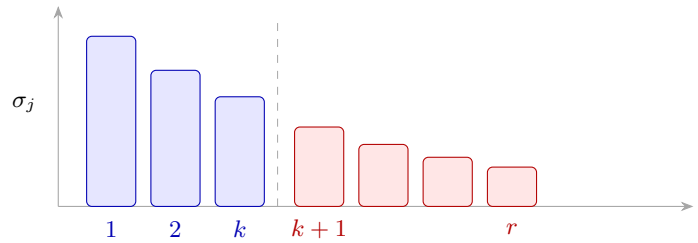


Figure 91: Blue bars form A_k; red bars are discarded, with $\|A - A_k\|_F^2 = \sum_{j>k} \sigma_j^2$.

Theorem 4.3.14. *Eckart–Young Low-Rank Approximation*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$ have singular values $\sigma_1 \geq \cdots \geq \sigma_q \geq 0$, where $q = \min(m, n)$ and $\sigma_j = 0$ for $j > q$. If A_k is the rank-k truncated SVD, $0 \leq k < q$, and $B \in \mathrm{Mat}(m, n, \mathbb{F})$ has $\mathrm{rank}(B) \leq k$, then

$$\|A - A_k\|_2 \leq \|A - B\|_2 , \qquad \|A - A_k\|_F \leq \|A - B\|_F .$$

Proof. Let $\mathbf{w}_1, \ldots, \mathbf{w}_n$ be right singular vectors for A.

i. Since $\mathrm{rank}(B) \leq k$, rank-nullity gives $\dim \ker(B) \geq n - k$. The subspaces $\ker(B)$ and $\mathrm{span}(\mathbf{w}_1, \ldots, \mathbf{w}_{k+1})$ in $\mathbb{F}^n$ have dimensions summing at least $n + 1$, so their intersection contains a nonzero vector.
Choose unit $\mathbf{x}$ in this intersection. Let $\mathbf{x} = \sum_{i=1}^{k+1} c_i \mathbf{w}_i$ gives $\sum_i |c_i|^2 = 1$ and $\|A\mathbf{x}\|_2^2 = \sum_{i=1}^{k+1} \sigma_i^2 |c_i|^2 \geq \sigma_{k+1}^2$.
Thus $\|A - B\|_2 \geq \|(A - B)\mathbf{x}\|_2 = \|A\mathbf{x}\|_2 \geq \sigma_{k+1} = \|A - A_k\|_2$.

ii. Choose an $(n-k)$-dimensional subspace $\mathcal{U} \subseteq \ker(B)$, with projection P and orthonormal basis $\mathbf{z}_1, \ldots, \mathbf{z}_{n-k}$.
Since $B\mathbf{z}_j = 0$, $\|A - B\|_F^2 \geq \sum_{j=1}^{n-k} \|(A - B)\mathbf{z}_j\|_2^2 = \sum_{j=1}^{n-k} \|A\mathbf{z}_j\|_2^2 = \sum_{i=1}^{n} \sigma_i^2 \|P\mathbf{w}_i\|_2^2$.
Put $\alpha_i = \|P\mathbf{w}_i\|_2^2$. Then $0 \leq \alpha_i \leq 1$ and $\sum_i \alpha_i = \mathrm{tr}(P) = n - k$.
Among all such weights, the sum $\sum_i \sigma_i^2 \alpha_i$ is minimized by placing weight 1 on the $n - k$ smallest values of the decreasing sequence σ_i^2, hence $\sum_i \sigma_i^2 \alpha_i \geq \sum_{i=k+1}^{n} \sigma_i^2 = \|A - A_k\|_F^2$.

Structural Commentary: Rank-k competitors vanish on at least $n - k$ directions.
Key Insight: The optimal loss is precisely the singular-value tail. □

Method 4.3.15. *Truncated SVD Approximation Method*
To compress or approximate A by rank k:

(i) compute or estimate the leading singular triples $(\sigma_j, \mathbf{x}_j, \mathbf{y}_j)$;
(ii) keep $A_k = \sum_{j=1}^{k} \sigma_j \mathbf{x}_j \mathbf{y}_j^*$;
(iii) use $\|A - A_k\|_F^2 = \sum_{j>k} \sigma_j^2$ and $\|A - A_k\|_2 = \sigma_{k+1}$ to quantify the lost information.

Definition 4.3.16. *Moore–Penrose Pseudoinverse*
Let $A \in \mathrm{Mat}(m, n, \mathbb{F})$. If $A = 0$, define $A^\dagger = 0_{n,m}$.
If $\mathrm{rank}(A) = r \geq 1$ and $A = X\Sigma_r Y^*$ is a compact SVD, define $A^\dagger := Y\Sigma_r^{-1} X^* \in \mathrm{Mat}(n, m, \mathbb{F})$.
This is the *pseudoinverse* or *Moore–Penrose pseudoinverse* of A.

Proposition 4.3.17. *Pseudoinverse is Well Defined*
The definition of $A^\dagger$ does not depend on the chosen compact SVD of A.

Proof. Let $A = X_1\Sigma_r Y_1^* = X_2\Sigma_r Y_2^*$. By Theorem 4.2.28, $X_1 = X_2 U$, $Y_1 = Y_2 U$, and $U\Sigma_r^{-1}U^* = \Sigma_r^{-1}$.
Hence $Y_1\Sigma_r^{-1}X_1^* = Y_2 U\Sigma_r^{-1}U^*X_2^* = Y_2\Sigma_r^{-1}X_2^*$.
Structural Commentary: Compact SVD ambiguity cancels in $Y\Sigma_r^{-1}X^*$.
Key Insight: Inverting the singular values preserves the same block symmetries. □

Corollary 4.3.18. *Pseudoinverse from a Full SVD*
Let $A = V\Sigma W^*$ be a full SVD of $A \in \mathrm{Mat}(m, n, \mathbb{F})$, with rank r and positive singular values $\sigma_1, \ldots, \sigma_r$.
Write $V = [X \ X_0]$, $W = [Y \ Y_0]$, and $\Sigma_r = \mathrm{diag}(\sigma_1, \ldots, \sigma_r)$.
In conformal blocks, $\Sigma = \begin{bmatrix} \Sigma_r & 0 \\ 0 & 0 \end{bmatrix}$, $\Sigma^\dagger := \begin{bmatrix} \Sigma_r^{-1} & 0 \\ 0 & 0 \end{bmatrix} \in \mathrm{Mat}(n, m, \mathbb{F})$. Then $A^\dagger = W\Sigma^\dagger V^*$.

Proof. If $r = 0$, then $A = 0$ and $W\Sigma^\dagger V^* = 0$.
If $r > 0$, the block form gives $A = X\Sigma_r Y^*$; $A^\dagger = Y\Sigma_r^{-1}X^*$; and $W\Sigma^\dagger V^* = Y\Sigma_r^{-1}X^*$. Thus $A^\dagger = W\Sigma^\dagger V^*$.
Structural Commentary: Full SVD pads compact SVD by zero blocks.
Key Insight: The pseudoinverse flips V and W and inverts only the positive singular values. □

Theorem 4.3.19. *Basic Properties of the Pseudoinverse*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$. Then:

(i) $\mathrm{col}(A^\dagger) = \mathrm{col}(A^*)$ *and* $\ker(A^\dagger) = \ker(A^*)$;
(ii) $AA^\dagger$ *is the orthogonal projection of* $\mathbb{F}^m$ *onto* $\mathrm{col}(A)$;
(iii) $A^\dagger A$ *is the orthogonal projection of* $\mathbb{F}^n$ *onto* $\mathrm{col}(A^*)$;
(iv) $AA^\dagger A = A$ *and* $A^\dagger AA^\dagger = A^\dagger$;
(v) $(A^\dagger)^\dagger = A$;
(vi) if $m = n$ *and* A *is invertible, then* $A^\dagger = A^{-1}$.

Proof. The zero case is immediate. Let $A = X\Sigma_r Y^*$ be a compact SVD; then $A^\dagger = Y\Sigma_r^{-1}X^*$.

(i) $\mathrm{col}(A^\dagger) = \mathrm{col}(Y) = \mathrm{col}(A^*)$ and $\ker(A^\dagger) = \ker(X^*) = \ker(A^*)$.
(ii) $AA^\dagger = XX^*$, the orthogonal projection onto $\mathrm{col}(X) = \mathrm{col}(A)$.
(iii) $A^\dagger A = YY^*$, the orthogonal projection onto $\mathrm{col}(Y) = \mathrm{col}(A^*)$.
(iv) These follow from $XX^*X = X$ and $YY^*Y = Y$.
(v) The compact SVD of $A^\dagger$ is $Y\Sigma_r^{-1}X^*$, so $(A^\dagger)^\dagger = A$.
(vi) If A is square and invertible, then $AA^\dagger = A^\dagger A = I_n$, hence $A^\dagger = A^{-1}$.

Structural Commentary: The pseudoinverse inverts nonzero singular directions and projects away null directions.
Key Insight: Replace each $\sigma_j > 0$ by σ_j^{-1} on the same singular subspaces. □

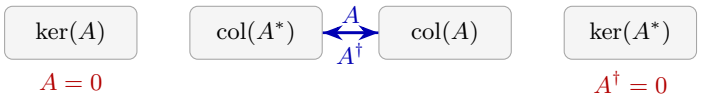


Figure 92: A and $A^\dagger$ invert the nonzero singular subspaces; they vanish on $\ker(A)$ and $\ker(A^*)$, respectively.

Theorem 4.3.20. *Penrose Characterization of the Pseudoinverse*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$, $\mathcal{C} = \mathrm{col}(A)$, and $\mathcal{D} = \mathrm{col}(A^*) = (\ker A)^\perp$. Then $A^\dagger$ is unique $B \in \mathrm{Mat}(n,m,\mathbb{F})$ satisfying

(i) $ABA = A$;
(ii) $BAB = B$;
(iii) AB *is Hermitian;*
(iv) BA *is Hermitian.*

Proof.

(i) Theorem 4.3.19 gives $AA^\dagger A = A$; $A^\dagger AA^\dagger = A^\dagger$, and $AA^\dagger$, $A^\dagger A$ are orthogonal projections. Hence $A^\dagger$ satisfies the four equations.
(ii) Let B satisfy the four equations. Then $(AB)^2 = ABAB = AB$; $(AB)^* = AB$. Also $\mathrm{range}(AB) = \mathcal{C}$, since $\mathrm{range}(AB) \subseteq \mathcal{C}$ and $AB(A\mathbf{x}) = ABA\mathbf{x} = A\mathbf{x}$. By Theorem 2.4.19, $AB = P_{\mathcal{C}}$.
(iii) $(BA)^2 = BABA = BA$; $(BA)^* = BA$. Also $\ker(BA) = \ker(A)$, as $BA\mathbf{x} = 0 \Rightarrow A\mathbf{x} = ABA\mathbf{x} = 0$ and the reverse inclusion is immediate. By Theorem 2.4.19, $BA = P_{\mathcal{D}}$.
(iv) For $\mathbf{y} \in \mathbb{F}^m$, write $\mathbf{y} = \mathbf{p} + \mathbf{z}$ with $\mathbf{p} \in \mathcal{C}$ and $\mathbf{z} \in \mathcal{C}^\perp = \ker(A^*)$. Since $AB\mathbf{z} = 0$, $B\mathbf{z} = BAB\mathbf{z} = 0$. If $\mathbf{p} = A\mathbf{x}$, then $B\mathbf{p} = BA\mathbf{x} = P_{\mathcal{D}}\mathbf{x}$, independent of the choice of $\mathbf{x}$ because $\ker(P_{\mathcal{D}}) = \mathcal{D}^\perp = \ker(A)$. Thus the four equations determine B uniquely. Since $A^\dagger$ satisfies them, $B = A^\dagger$.

Structural Commentary: Penrose equations force the two projection factors.
Key Insight: Once AB and BA are projections, B is forced. □

Theorem 4.3.21. *Pseudoinverse and Minimum-Norm Least Squares*
Let $A \in \mathrm{Mat}(m,n,\mathbb{F})$, let $\mathbf{y} \in \mathbb{F}^m$, and set $\mathbf{x}_0 = A^\dagger\mathbf{y}$.

(i) If $A\mathbf{x} = \mathbf{y}$ is consistent, then $A\mathbf{x}_0 = \mathbf{y}$ and $\mathbf{x}_0$ is the unique solution of smallest Euclidean norm.
(ii) For arbitrary $\mathbf{y}$, $\mathbf{x}_0$ is a least-squares solution of $A\mathbf{x} \approx \mathbf{y}$, and it is the unique least-squares solution of smallest Euclidean norm.

Proof. Let $P = AA^\dagger = P_{\mathrm{col}(A)}$ and note that $\mathbf{x}_0 \in \mathrm{col}(A^\dagger) = \mathrm{col}(A^*) = (\ker A)^\perp$.

(i) If $\mathbf{y} \in \mathrm{col}(A)$, then $A\mathbf{x}_0 = P\mathbf{y} = \mathbf{y}$. For any solution $\mathbf{x}$, $\mathbf{x} - \mathbf{x}_0 \in \ker(A)$, so $\|\mathbf{x}\|_2^2 = \|\mathbf{x}_0 + (\mathbf{x} - \mathbf{x}_0)\|_2^2 = \|\mathbf{x}_0\|_2^2 + \|\mathbf{x} - \mathbf{x}_0\|_2^2$ with equality only when $\mathbf{x} = \mathbf{x}_0$.
(ii) For arbitrary $\mathbf{y}$, $A\mathbf{x}_0 = P\mathbf{y}$, so $\mathbf{y} - A\mathbf{x}_0 \perp \mathrm{col}(A)$ and $\mathbf{x}_0$ is least-squares optimal by Theorem 2.3.4. Any least-squares solution $\mathbf{x}$ satisfies $A\mathbf{x} = P\mathbf{y} = A\mathbf{x}_0$, hence the same orthogonal decomposition gives $\|\mathbf{x}\|_2 \geq \|\mathbf{x}_0\|_2$, with equality only for $\mathbf{x} = \mathbf{x}_0$.

Structural Commentary: Least squares projects onto $\operatorname{col}(A)$, then removes the $\ker(A)$ component.
Key Insight: $A^\dagger \mathbf{y}$ is the fitted solution orthogonal to $\ker(A)$. □

Method 4.3.22. *Pseudoinverse Solution Method*
For a system $A\mathbf{x} \approx \mathbf{y}$:

(i) compute a compact SVD $A = X\Sigma_r Y^*$;
(ii) form $A^\dagger = Y\Sigma_r^{-1}X^*$;
(iii) set $\mathbf{x}_0 = A^\dagger \mathbf{y}$;
(iv) interpret $A\mathbf{x}_0 = P_{\operatorname{col}(A)}\mathbf{y}$ and $\mathbf{x}_0 \in (\ker A)^\perp$.

Definition 4.3.23. *Spectral Condition Number*
If $A \in \operatorname{Mat}(n,n,\mathbb{F})$ is invertible, its *spectral condition number* is $\kappa_2(A) := \|A\|_2 \left\|A^{-1}\right\|_2$.
Equivalently, $\kappa_2(A) = \frac{\sigma_{\max}(A)}{\sigma_{\min}(A)}$. If A is singular, $\kappa_2(A)$ is undefined.

Proposition 4.3.24. *Properties of the Spectral Condition Number*
Let $A, B \in \operatorname{Mat}(n,n,\mathbb{F})$ be invertible.

(i) $\kappa_2(AB) \le \kappa_2(A)\kappa_2(B)$;
(ii) $\kappa_2(A) \ge 1$, *with equality if and only if A is a nonzero scalar multiple of a unitary matrix;*
(iii) $\kappa_2(cA) = \kappa_2(A)$ *for every nonzero scalar c;*
(iv) $\kappa_2(A) = \kappa_2(A^*) = \kappa_2(\overline{A}) = \kappa_2(A^t) = \kappa_2(A^{-1})$ *over $\mathbb{C}$, with the evident real special cases;*
(v) $\kappa_2(A^*A) = \kappa_2(A)^2$;
(vi) if U, V are unitary, then $\kappa_2(UAV) = \kappa_2(A)$.

Proof.

(i) Use submultiplicativity: $\kappa_2(AB) = \|AB\|_2 \left\|B^{-1}A^{-1}\right\|_2 \le \|A\|_2 \|B\|_2 \left\|B^{-1}\right\|_2 \left\|A^{-1}\right\|_2$.
(ii) Since $I_n = AA^{-1}$, $1 = \|I_n\|_2 \le \kappa_2(A)$. Equality means $\sigma_{\max}(A) = \sigma_{\min}(A)$, equivalently every singular value is the same positive number s. If $A = V\Sigma W^*$, then $\Sigma = sI_n$, so $A = sVW^*$ is a nonzero scalar multiple of a unitary matrix. The converse is immediate.
(iii) This follows from $\|cA\|_2 = |c|\,\|A\|_2$ and $(cA)^{-1} = c^{-1}A^{-1}$.
(iv) These matrices have the same singular values as A, except A^{-1} has singular values $\sigma_n^{-1} \ge \cdots \ge \sigma_1^{-1}$.
(v) The singular values of A^*A are $\sigma_1(A)^2, \ldots, \sigma_n(A)^2$.
(vi) This follows from Proposition 4.3.7.

Structural Commentary: Condition number compares largest and smallest singular stretches.
Key Insight: A well-conditioned matrix stretches every direction by nearly the same factor. □

Corollary 4.3.25. *Positive Definite Condition Number by Eigenvalues*
Let $A \succ 0$ have eigenvalues $0 < \lambda_1 \le \cdots \le \lambda_n$. Then $\|A\|_2 = \lambda_n$; $\left\|A^{-1}\right\|_2 = \frac{1}{\lambda_1}$; and $\kappa_2(A) = \frac{\lambda_n}{\lambda_1}$.

Proof. By Theorems 4.1.58 and 4.3.11, the singular values of A are $\lambda_1, \ldots, \lambda_n$.
Hence Theorem 4.3.2 gives $\|A\|_2 = \lambda_n$. Since $\operatorname{spec}(A^{-1}) = \{\lambda_1^{-1}, \ldots, \lambda_n^{-1}\}$, same argument gives $\left\|A^{-1}\right\|_2 = \lambda_1^{-1}$, and multiplying yields $\kappa_2(A) = \lambda_n/\lambda_1$.
Structural Commentary: For $A \succ 0$, singular and eigenvalue stretches coincide.
Key Insight: Conditioning compares the largest and smallest positive eigendirections. □

Perspective 4.3.26. *Determinant is Not Conditioning*
Determinants measure volume, not sensitivity.
$\kappa_2(cA) = \kappa_2(A)$ for $c \neq 0$ (Proposition 4.3.24), but $\det(cA) = c^n \det(A)$.

Theorem 4.3.27. *Condition Number and Linear-System Error*
Let $A \in \operatorname{Mat}(n,n,\mathbb{F})$ be invertible, let $\mathbf{y} \neq \mathbf{0}$, and let $A\mathbf{x} = \mathbf{y}$.
If $A(\Delta\mathbf{x}) = \Delta\mathbf{y}$, then $\frac{\|\Delta\mathbf{x}\|_2}{\|\mathbf{x}\|_2} \le \kappa_2(A)\frac{\|\Delta\mathbf{y}\|_2}{\|\mathbf{y}\|_2}$

Proof. From $\mathbf{y} = A\mathbf{x}$, we have $\|\mathbf{y}\|_2 \le \|A\|_2 \|\mathbf{x}\|_2$, hence $1/\|\mathbf{x}\|_2 \le \|A\|_2 / \|\mathbf{y}\|_2$. Also $\Delta\mathbf{x} = A^{-1}\Delta\mathbf{y}$, so $\|\Delta\mathbf{x}\|_2 \le \left\|A^{-1}\right\|_2 \|\Delta\mathbf{y}\|_2$. Multiplying the two estimates gives the displayed bound.
Structural Commentary: Solution error is controlled by the ratio of extreme singular stretches.
Key Insight: $\kappa_2(A)$ measures worst-case error magnification. □

Theorem 4.3.28. *Condition Number and Matrix Perturbation*
Let $A \in \operatorname{Mat}(n,n,\mathbb{F})$ be invertible. Suppose $A\mathbf{x} = \mathbf{y}$; $(A + \Delta A)(\mathbf{x} + \Delta\mathbf{x}) = \mathbf{y}$, and $\mathbf{x} + \Delta\mathbf{x} \neq \mathbf{0}$.
Then $\frac{\|\Delta\mathbf{x}\|_2}{\|\mathbf{x}+\Delta\mathbf{x}\|_2} \le \kappa_2(A)\frac{\|\Delta A\|_2}{\|A\|_2}$.

Proof. Subtracting gives $A\Delta\mathbf{x} + \Delta A(\mathbf{x}+\Delta\mathbf{x}) = \mathbf{0}$, so $\Delta\mathbf{x} = -A^{-1}\Delta A(\mathbf{x}+\Delta\mathbf{x})$.
Thus $\|\Delta\mathbf{x}\|_2 \le \|A^{-1}\|_2 \|\Delta A\|_2 \|\mathbf{x}+\Delta\mathbf{x}\|_2 = \kappa_2(A)\frac{\|\Delta A\|_2}{\|A\|_2}\|\mathbf{x}+\Delta\mathbf{x}\|_2$. Divide by $\|\mathbf{x}+\Delta\mathbf{x}\|_2 \neq 0$.
Structural Commentary: Perturbing the matrix changes the operator applied to the perturbed solution.
Key Insight: The same condition number controls matrix errors and right-hand-side errors. □

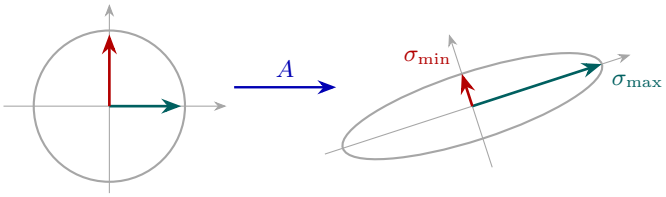


Figure 93: A maps a unit error ball to an ellipse; $\kappa_2(A) = \sigma_{\max}/\sigma_{\min}$ is the long/short semiaxis ratio.

4.4 Perturbation and Stability

Perturbation theory studies how matrix information changes under small errors. Bauer–Fike, Gersgorin, interlacing, Weyl inequalities, and pseudospectra measure robustness.

Lemma 4.4.1. *Small Spectral-Norm Perturbations of the Identity*
If $E \in \mathrm{Mat}(n,n,\mathbb{C})$ and $\|E\|_2 < 1$, then $I_n + E$ is invertible.

Proof. Theorem 4.3.11 gives $\rho(E) \le \|E\|_2 < 1$, so $-1 \notin \mathrm{spec}(E)$.
Hence $0 \notin \mathrm{spec}(I_n + E)$, and $I_n + E$ is invertible.
Structural Commentary: Spectral-norm smallness traps every eigenvalue inside the unit disk.
Key Insight: If -1 is not an eigenvalue of E, then $I + E$ has no kernel. □

Theorem 4.4.2. *Bauer–Fike Eigenvalue Perturbation Bound*
Let $A \in \mathrm{Mat}(n,n,\mathbb{C})$ be diagonalizable, with $A = S\Lambda S^{-1}$ and $\Lambda = \mathrm{diag}(\lambda_1,\dots,\lambda_n)$.
If $\Delta A \in \mathrm{Mat}(n,n,\mathbb{C})$ and $\mu \in \mathrm{spec}(A+\Delta A)$, then $\min_{1\le j\le n}|\mu-\lambda_j| \le \kappa_2(S)\|\Delta A\|_2$.

Proof. If $\mu = \lambda_j$ for some j, the claim is immediate. Otherwise $\Lambda - \mu I_n$ is invertible.
Since μ is an eigenvalue of $A+\Delta A$, the matrix $A + \Delta A - \mu I_n$ is singular.
Conjugating gives $S^{-1}(A+\Delta A-\mu I_n)S = (\Lambda-\mu I_n)\left(I_n + (\Lambda-\mu I_n)^{-1}S^{-1}\Delta A S\right)$ so the second factor is singular.
Lemma 4.4.1 implies $1 \le \|(\Lambda-\mu I_n)^{-1}S^{-1}\Delta A S\|_2 \le \frac{\|S^{-1}\|_2\|\Delta A\|_2\|S\|_2}{\min_j|\mu-\lambda_j|}$. Rearranging gives the bound.
Structural Commentary: Diagonalize, then test how perturbation breaks shifted invertibility.
Key Insight: Eigenvalue sensitivity is controlled by how ill-conditioned the eigenvector matrix S is. □

Method 4.4.3. *Bauer–Fike Perturbation Method*
To bound eigenvalue movement for a diagonalizable matrix:

(i) diagonalize $A = S\Lambda S^{-1}$;
(ii) estimate $\|\Delta A\|_2$;
(iii) compute or bound $\kappa_2(S)$;
(iv) place every eigenvalue of $A+\Delta A$ inside the union of disks $|z-\lambda_j| \le \kappa_2(S)\|\Delta A\|_2$.

Perspective 4.4.4. *Pseudospectral Intuition*
For a matrix A, the ε-pseudospectrum records complex numbers z for which $A - zI$ is singular or nearly singular. Equivalently, it asks where a perturbation of size at most ε can move the spectrum. Bauer–Fike gives a clean enclosure for diagonalizable matrices, while Gersgorin and norm bounds give cruder but often cheaper enclosures. The main lesson is that eigenvalues of nonnormal matrices can be much more sensitive than their locations alone suggest.

Definition 4.4.5. *Gersgorin Disks*
Let $A = [a_{ij}] \in \mathrm{Mat}(n,n,\mathbb{C})$. The *deleted absolute row sum* in row k is $R_k^{\mathrm{off}}(A) := \sum_{j\neq k}|a_{kj}|$.

The *kth Gersgorin disk* of A is $\mathrm{Ger}_k(A) := \{z \in \mathbb{C} : |z-a_{kk}| \le R_k^{\mathrm{off}}(A)\}$, where $\mathrm{Ger}(A) := \bigcup_{k=1}^{n}\mathrm{Ger}_k(A)$.

Theorem 4.4.6. *Gersgorin Disk Theorem*
Let $A = [a_{ij}] \in \mathrm{Mat}(n,n,\mathbb{C})$. Then $\mathrm{spec}(A) \subseteq \mathrm{Ger}(A) = \bigcup_{k=1}^{n}\mathrm{Ger}_k(A)$.

Proof. Let $\lambda \in \operatorname{spec}(A)$ and choose an eigenvector $\mathbf{x} = [x_1 \ \cdots \ x_n]^t$.
Pick k such that $|x_k| = \max_j |x_j|$; then $x_k \neq 0$. The kth row of $A\mathbf{x} = \lambda\mathbf{x}$ gives $(\lambda - a_{kk})x_k = \sum_{j\neq k} a_{kj}x_j$.

Divide by x_k and use $|x_j/x_k| \leq 1$: $|\lambda - a_{kk}| = \left|\sum_{j\neq k} a_{kj}\frac{x_j}{x_k}\right| \leq \sum_{j\neq k} |a_{kj}|\left|\frac{x_j}{x_k}\right| \leq R_k^{\text{off}}(A)$.

Thus $\lambda \in \operatorname{Ger}_k(A) \subseteq \operatorname{Ger}(A)$.
Structural Commentary: Normalize by a largest component; all row ratios have modulus at most one.
Key Insight: Row k traps λ in the disk centered at a_{kk} with radius $R_k^{\text{off}}(A)$. □

Method 4.4.7. *Gersgorin Localisation Method*
For $A = [a_{ij}]$, Gersgorin localisation gives a disk enclosure for the spectrum.

(i) For each row k, place a disk at the diagonal entry a_{kk} with radius $R_k^{\text{off}}(A) = \sum_{j\neq k} |a_{kj}|$.
The row enclosure is $\operatorname{Ger}(A) = \bigcup_{k=1}^n \operatorname{Ger}_k(A)$.
(ii) Theorem 4.4.6 guarantees that every eigenvalue of A lies in $\operatorname{Ger}(A)$.
(iii) For a sharper enclosure, repeat the same construction for A^t. The centers are still a_{kk}, while the radii becomes deleted column sums of A. Then use $\operatorname{spec}(A) \subseteq \operatorname{Ger}(A) \cap \operatorname{Ger}(A^t)$. This is the intersection of two unions of disks; an eigenvalue need not lie in the overlap of each individual row-column disk pair.

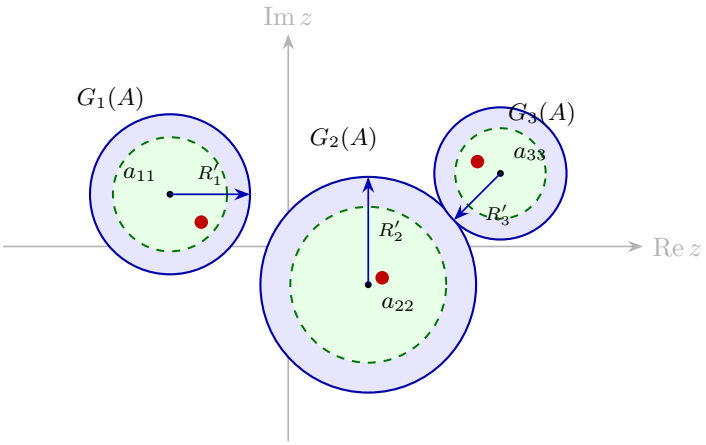


Figure 94: Blue row disks contain the spectrum; green dashed column disks give the sharper row-column enclosure $\operatorname{Ger}(A) \cap \operatorname{Ger}(A^t)$. Red points illustrate possible eigenvalue locations.

Corollary 4.4.8. *Row-Column Gersgorin Refinement*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$. *Then* $\operatorname{spec}(A) \subseteq \operatorname{Ger}(A) \cap \operatorname{Ger}(A^t)$.
If $R_k(A) = \sum_{j=1}^n |a_{kj}|$ *and* $R_{\max}(A) = \max_{1\leq k\leq n} R_k(A)$, *then* $\operatorname{spec}(A) \subseteq \{z \in \mathbb{C} : |z| \leq \min(R_{\max}(A), R_{\max}(A^t))\}$.

Proof. Theorem 4.4.6 gives $\operatorname{spec}(A) \subseteq \operatorname{Ger}(A)$.
Proposition 4.1.29 gives $\operatorname{spec}(A) = \operatorname{spec}(A^t)$, so applying Gersgorin to A^t gives $\operatorname{spec}(A) \subseteq \operatorname{Ger}(A^t)$.
Also $\operatorname{Ger}_k(A) \subseteq \{z : |z| \leq |a_{kk}| + R_k^{\text{off}}(A) = R_k(A)\}$, and similarly for A^t.
Structural Commentary: Row disks and column disks are both valid because A and A^t have the same spectrum.
Key Insight: Intersecting independent enclosures usually sharpens the eigenvalue search region. □

Definition 4.4.9. *Diagonal Dominance*
Let $A = [a_{ij}] \in \operatorname{Mat}(n, n, \mathbb{C})$. Matrix A is *diagonally dominant* if $|a_{kk}| \geq R_k^{\text{off}}(A)$ for every k, and *strictly diagonally dominant* if $|a_{kk}| > R_k^{\text{off}}(A)$ for every k.

Corollary 4.4.10. *Strict Diagonal Dominance Implies Invertibility*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$.

(i) If A is strictly diagonally dominant, then A is invertible.
(ii) If A has real spectrum and nonnegative diagonal entries, then diagonal dominance gives $\operatorname{spec}(A) \subseteq [0, \infty)$*; strict diagonal dominance gives* $\operatorname{spec}(A) \subseteq (0, \infty)$.

Proof.

(i) Strict dominance gives $0 \notin \operatorname{Ger}_k(A)$ for every k, since $|a_{kk}| > R_k^{\text{off}}(A)$.
Thus $0 \notin \operatorname{spec}(A)$ by Gersgorin, so A is invertible by Corollary 3.1.9.

(ii) If $a_{kk} \geq 0$ and $a_{kk} \geq R_k^{\text{off}}(A)$, then each $\text{Ger}_k(A)$ lies in the closed right half-plane. In the strict case, each disk lies in the open right half-plane. Gersgorin and the real-spectrum hypothesis give the claim.

Structural Commentary: Gersgorin converts diagonal dominance into disk placement.
Key Insight: The disks control whether eigenvalues can reach 0 or the left half-plane. □

Corollary 4.4.11. *Diagonal Dominance and Positive Semidefiniteness*
Let $A = [a_{ij}] \in \text{Mat}(n, n, \mathbb{C})$ be Hermitian with nonnegative diagonal entries.

(i) If A is diagonally dominant, then A is positive semidefinite.
(ii) If A is diagonally dominant and invertible, then A is positive definite.
(iii) If A is strictly diagonally dominant, then A is positive definite.

Proof. Hermitian matrices have real spectrum.

(i) Corollary 4.4.10 gives $\text{spec}(A) \subseteq [0, \infty)$; hence $A \succeq 0$ by Theorem 4.1.56.
(ii) If A is invertible, then $0 \notin \text{spec}(A)$, so every eigenvalue is positive; hence $A \succ 0$.
(iii) Strict diagonal dominance implies invertibility by Corollary 4.4.10; apply (ii).

Structural Commentary: Gersgorin traps the real spectrum in $[0, \infty)$; invertibility removes 0.
Key Insight: For Hermitian matrices, positivity is spectral location. □

Theorem 4.4.12. *Weak Diagonal Dominance Criterion*
Let $A = [a_{ij}] \in \text{Mat}(n, n, \mathbb{C})$ be diagonally dominant with no zero entries.
If $|a_{kk}| > R_k^{\text{off}}(A)$ for some k, then A is invertible.

Proof. Assume $A\mathbf{x} = 0$ with $\mathbf{x} \neq \mathbf{0}$, and choose k with $|x_k| = \|\mathbf{x}\|_\infty$.

Row k gives $|a_{kk}| \|\mathbf{x}\|_\infty = \left| -\sum_{j \neq k} a_{kj} x_j \right| \leq \sum_{j \neq k} |a_{kj}| |x_j| \leq R_k^{\text{off}}(A) \|\mathbf{x}\|_\infty \leq |a_{kk}| \|\mathbf{x}\|_\infty$.

Hence all inequalities are equalities. Since no a_{kj} vanishes, $|x_j| = \|\mathbf{x}\|_\infty$ for every j. Applying the same chain to any row i gives $|a_{ii}| = R_i'(A)$ for every i, contradicting the strict row. Thus $\ker(A) = \{\mathbf{0}\}$, so A is invertible.
Structural Commentary: A zero vector in the kernel would force equality in every Gersgorin row.
Key Insight: One strict row breaks that equality pattern. □

Definition 4.4.13. *Increasing Hermitian Eigenvalue Order*
For a Hermitian matrix $A \in \text{Mat}(n, n, \mathbb{C})$, write $\lambda_1(A) \leq \lambda_2(A) \leq \cdots \leq \lambda_n(A)$ for its eigenvalues counted with algebraic multiplicity. If only one Hermitian matrix is under discussion, we write λ_i instead of $\lambda_i(A)$.

Theorem 4.4.14. *Rayleigh Quotient on a Spectral Interval*
Let $A \in \text{Mat}(n, n, \mathbb{C})$ be Hermitian with increasing eigenvalues $\lambda_1 \leq \cdots \leq \lambda_n$ and corresponding orthonormal eigenvectors $\mathbf{u}_1, \ldots, \mathbf{u}_n$. For a unit vector, write $\text{Ray}_A(\mathbf{x}) := \mathbf{x}^* A \mathbf{x}$. If $1 \leq p \leq q \leq n$ and $\mathbf{x} \in \text{span}(\mathbf{u}_p, \ldots, \mathbf{u}_q)$ is a unit vector, then $\lambda_p \leq \text{Ray}_A(\mathbf{x}) \leq \lambda_q$.

Proof. Write $\mathbf{x} = \sum_{i=p}^q c_i \mathbf{u}_i$. Orthonormality gives $\sum_{i=p}^q |c_i|^2 = 1$, and $\mathbf{x}^* A \mathbf{x} = \sum_{i=p}^{q} \lambda_i |c_i|^2$.

This is a convex combination of $\lambda_p, \ldots, \lambda_q$, so it lies between the smallest and largest of them.
Structural Commentary: In an eigenbasis, the Rayleigh quotient is an eigenvalue average.
Key Insight: Spectral coordinates turn a quadratic form into weighted scalar data. □

Corollary 4.4.15. *Rayleigh Extrema and Diagonal Entry Bounds*
Let $A \in \text{Mat}(n, n, \mathbb{C})$ be Hermitian.

(i) If $\mathcal{U}_{p,q} = \text{span}(\mathbf{u}_p, \ldots, \mathbf{u}_q)$, then $\min_{\substack{\mathbf{x} \in \mathcal{U}_{p,q} \\ \|\mathbf{x}\|_2 = 1}} \mathbf{x}^* A \mathbf{x} = \lambda_p(A)$, and $\max_{\substack{\mathbf{x} \in \mathcal{U}_{p,q} \\ \|\mathbf{x}\|_2 = 1}} \mathbf{x}^* A \mathbf{x} = \lambda_q(A)$.
(ii) For every unit vector $\mathbf{x}$, $\lambda_1(A) \leq \mathbf{x}^* A \mathbf{x} \leq \lambda_n(A)$.
(iii) If $A = [a_{ij}]$, then $\lambda_1(A) \leq a_{ii} \leq \lambda_n(A)$ for $i = 1, \ldots, n$.

Proof.

(i) Theorem 4.4.14 gives the inequalities on $\mathcal{U}_{p,q}$, with equality at $\mathbf{u}_p$ and $\mathbf{u}_q$.
(ii) This is (i) with $p = 1$ and $q = n$.
(iii) Apply (ii) to $\mathbf{x} = \mathbf{e}_i$, since $\mathbf{e}_i^* A \mathbf{e}_i = a_{ii}$.

Structural Commentary: Rayleigh extrema on extreme eigendirections; diagonal entries are coordinate quotients.
Key Insight: The whole numerical range of a Hermitian matrix sits inside its spectral interval. □

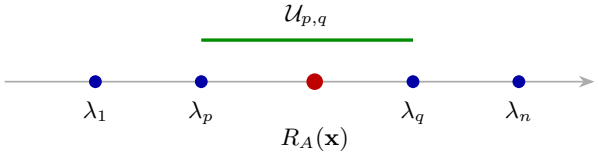


Figure 95: For $\mathbf{x} \in \mathcal{U}_{p,q}$, $\mathrm{Ray}_A(\mathbf{x})$ is a convex combination of $\lambda_p, \dots, \lambda_q$.

Theorem 4.4.16. *Rayleigh Minimax Principle*
Let $A = A^* \in \mathrm{Mat}(n, n, \mathbb{C})$ have eigenvalues $\lambda_1 \le \cdots \le \lambda_n$, and set $R_A(\mathbf{x}) = \mathbf{x}^* A\mathbf{x}/(\mathbf{x}^*\mathbf{x})$ for $\mathbf{x} \neq \mathbf{0}$.
For $1 \le k \le n$, $\lambda_k = \min\limits_{\substack{\mathcal{S} \subseteq \mathbb{C}^n \\ \dim \mathcal{S} = k}} \max\limits_{\mathbf{0} \neq \mathbf{x} \in \mathcal{S}} R_A(\mathbf{x}) = \max\limits_{\substack{\mathcal{S} \subseteq \mathbb{C}^n \\ \dim \mathcal{S} = n-k+1}} \min\limits_{\mathbf{0} \neq \mathbf{x} \in \mathcal{S}} R_A(\mathbf{x})$.

Proof. Let $\mathbf{u}_1, \dots, \mathbf{u}_n$ be an orthonormal eigenbasis in increasing order.
Set $E_k = \mathrm{span}(\mathbf{u}_1, \dots, \mathbf{u}_k)$; $F_k = \mathrm{span}(\mathbf{u}_k, \dots, \mathbf{u}_n)$.
By Theorem 4.4.14, $\max_{\mathbf{0} \neq \mathbf{x} \in E_k} R_A(\mathbf{x}) = \lambda_k$ and $\min_{\mathbf{0} \neq \mathbf{x} \in F_k} R_A(\mathbf{x}) = \lambda_k$.
If $\dim \mathcal{S} = k$, then $\mathcal{S} \cap F_k \neq \{\mathbf{0}\}$, so $\max_{\mathbf{0} \neq \mathbf{x} \in \mathcal{S}} R_A(\mathbf{x}) \ge \lambda_k$; equality is attained by E_k.
If $\dim \mathcal{T} = n - k + 1$, then $\mathcal{T} \cap E_k \neq \{\mathbf{0}\}$, so $\min_{\mathbf{0} \neq \mathbf{x} \in \mathcal{T}} R_A(\mathbf{x}) \le \lambda_k$; equality is attained by F_k.
Structural Commentary: Dimension forces intersections with spectral flags.
Key Insight: Minimax is Rayleigh's quotient plus dimension counting. □

Method 4.4.17. *Rayleigh Eigenbasis Method*
To bound a Hermitian quadratic form:

(i) diagonalise A with an orthonormal eigenbasis;
(ii) expand the unit vector in that basis;
(iii) rewrite $\mathbf{x}^* A\mathbf{x}$ as $\sum_i \lambda_i |c_i|^2$;
(iv) use the eigenvalue range allowed by the support of $\mathbf{x}$.

Theorem 4.4.18. *Rank-One Positive Perturbation Interlacing*
Let $A, E \in \mathrm{Mat}(n, n, \mathbb{C})$ be Hermitian, with $E \succeq 0$ and $\mathrm{rank}(E) \le 1$.
Then $\lambda_i(A) \le \lambda_i(A + E) \le \lambda_{i+1}(A) \quad (1 \le i < n)$, $\lambda_n(A) \le \lambda_n(A + E)$
Equivalently, $\lambda_i(A - E) \le \lambda_i(A) \le \lambda_{i+1}(A - E) \quad (1 \le i < n)$, $\lambda_n(A - E) \le \lambda_n(A)$.

Proof. Let $\mathbf{u}_j$ and $\mathbf{v}_j$ be orthonormal eigenvectors of A and $A + E$, respectively, in increasing eigenvalue order.
For the lower bound, set $\mathcal{U} = \mathrm{span}(\mathbf{u}_i, \dots, \mathbf{u}_n)$, $\mathcal{V} = \mathrm{span}(\mathbf{v}_1, \dots, \mathbf{v}_i)$. Since $\dim \mathcal{U} + \dim \mathcal{V} = n + 1$, choose a unit $\mathbf{x} \in \mathcal{U} \cap \mathcal{V}$. Then $\lambda_i(A) \le \mathbf{x}^* A\mathbf{x} \le \mathbf{x}^*(A + E)\mathbf{x} \le \lambda_i(A + E)$ by Theorem 4.4.14 and $E \succeq 0$.
For the upper bound, set $\mathcal{U} = \mathrm{span}(\mathbf{u}_1, \dots, \mathbf{u}_{i+1})$, $\mathcal{V} = \mathrm{span}(\mathbf{v}_i, \dots, \mathbf{v}_n)$, $\mathcal{W} = \ker(E)$. Here $\dim(\mathcal{U} \cap \mathcal{V}) \ge 2$ and $\dim \mathcal{W} \ge n - 1$, so choose a unit $\mathbf{x} \in \mathcal{U} \cap \mathcal{V} \cap \mathcal{W}$. Thus $\lambda_i(A + E) \le \mathbf{x}^*(A + E)\mathbf{x} = \mathbf{x}^* A\mathbf{x} \le \lambda_{i+1}(A)$.
The statements for $A - E$ follow by applying the proved result to $A - E$ and E.
Structural Commentary: Interlacing comes from test vectors in overlapping spectral subspaces.
Key Insight: Rank one leaves a large kernel, so the upper-bound vector can avoid the perturbation. □

Corollary 4.4.19. *Rank-Two Indefinite Interlacing*
Let $A, H \in \mathrm{Mat}(n, n, \mathbb{C})$ be Hermitian, $n \ge 2$. Suppose $\mathrm{rank}(H) \le 2$ and H has at most one positive and one negative eigenvalue. Then

$$\lambda_1(A + H) \le \lambda_2(A), \qquad \lambda_{i-1}(A) \le \lambda_i(A + H) \le \lambda_{i+1}(A) \quad (2 \le i < n), \qquad \lambda_{n-1}(A) \le \lambda_n(A + H).$$

Proof. Write $H = E_+ - E_-$ with $E_\pm \succeq 0$ and $\mathrm{rank}(E_\pm) \le 1$. Set $B = A - E_-$; then $A = B + E_-$ and $A + H = B + E_+$. For $2 \le i < n$, rank-one interlacing gives $\lambda_{i-1}(A) \le \lambda_i(B) \le \lambda_i(A + H) \le \lambda_{i+1}(B) \le \lambda_{i+1}(A)$.
The endpoints follow similarly: $\lambda_1(A + H) \le \lambda_2(B) \le \lambda_2(A)$, and $\lambda_{n-1}(A) \le \lambda_n(B) \le \lambda_n(A + H)$.
Structural Commentary: Such an H is a difference of two rank-one PSD perturbations.
Key Insight: Interlace once with E_-, then once with E_+. □

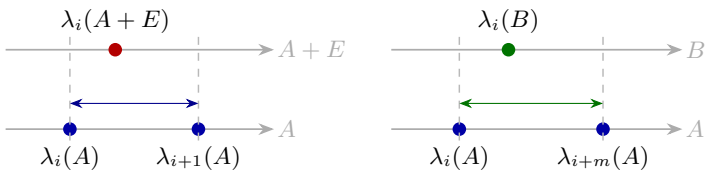


Figure 96: Rank-one and principal-submatrix interlacing place new eigenvalues between old ones.

Theorem 4.4.20. *Cauchy Interlacing: One Principal Deletion*
Let $A = \begin{bmatrix} B & \mathbf{y} \\ \mathbf{y}^* & c \end{bmatrix} \in \mathrm{Mat}(n+1, n+1, \mathbb{C})$ *be Hermitian, with* $B \in \mathrm{Mat}(n, n, \mathbb{C})$.
Then, for $1 \le i \le n$, $\lambda_i(A) \le \lambda_i(B) \le \lambda_{i+1}(A)$.

Proof. Let $\mathbf{u}_j$ and $\mathbf{v}_j$ be orthonormal eigenvectors of A and B, respectively, in increasing eigenvalue order. Identify $\mathbb{C}^n$ with $\{(\mathbf{z}, 0)\} \subseteq \mathbb{C}^{n+1}$, so $(\mathbf{z}, 0)^* A (\mathbf{z}, 0) = \mathbf{z}^* B \mathbf{z}$.
For the lower bound, set $\mathcal{S} = \mathrm{span}((\mathbf{v}_1, 0), \ldots, (\mathbf{v}_i, 0))$ and $\mathcal{T} = \mathrm{span}(\mathbf{u}_i, \ldots, \mathbf{u}_{n+1})$. Their dimensions are i and $n + 2 - i$, so $\dim \mathcal{S} + \dim \mathcal{T} > n + 1$ and $\mathcal{S} \cap \mathcal{T}$ contains a nonzero vector.
Choose a unit $\mathbf{x} = (\mathbf{z}, 0)$ in the intersection. The spectral expansion in the $\mathbf{u}_j$ gives $\mathbf{x}^* A \mathbf{x} \ge \lambda_i(A)$, while the expansion in the $\mathbf{v}_j$ gives $\mathbf{z}^* B \mathbf{z} \le \lambda_i(B)$. Hence $\lambda_i(A) \le \lambda_i(B)$.
For the upper bound, take $\mathcal{S}' = \mathrm{span}((\mathbf{v}_i, 0), \ldots, (\mathbf{v}_n, 0))$ and $\mathcal{T}' = \mathrm{span}(\mathbf{u}_1, \ldots, \mathbf{u}_{i+1})$. Again the dimensions sum to $n + 2$, so choose a unit $\mathbf{x} = (\mathbf{z}, 0) \in \mathcal{S}' \cap \mathcal{T}'$.
Now $\mathbf{z}^* B \mathbf{z} \ge \lambda_i(B)$ and $\mathbf{x}^* A \mathbf{x} \le \lambda_{i+1}(A)$, giving $\lambda_i(B) \le \lambda_{i+1}(A)$.
Structural Commentary: A principal submatrix is a coordinate slice of the larger quadratic form.
Key Insight: Deleting one coordinate costs one dimension, so eigenvalues shift by at most one index. □

Corollary 4.4.21. *Singular Values Interlace After Deleting a Row or Column*
Let B *be obtained from* $A \in \mathrm{Mat}(n, n, \mathbb{C})$ *by deleting one row or column.*
With singular values ordered decreasingly, $\sigma_i(A) \ge \sigma_i(B) \ge \sigma_{i+1}(A) \quad (1 \le i < n)$.

Proof. For a row deletion, permute rows so that $PA = \begin{bmatrix} B \\ \mathbf{y}^* \end{bmatrix}$.
Then PA has the same singular values as A, and BB^* is a principal submatrix of $(PA)(PA)^*$.
Cauchy interlacing applied to these Gram matrices gives the result after taking square roots and reversing the increasing eigenvalue order. A column deletion follows by applying the row case to A^*.
Structural Commentary: Row deletion makes BB^* a principal submatrix of a permuted AA^*.
Key Insight: Singular-value interlacing is Cauchy interlacing applied to Gram matrices. □

Corollary 4.4.22. *Cauchy Interlacing for Principal Submatrices*
Let $B \in \mathrm{Mat}(n, n, \mathbb{C})$ *be a principal submatrix of Hermitian* $A \in \mathrm{Mat}(n+m, n+m, \mathbb{C})$.
Then, for $1 \le i \le n$, $\lambda_i(A) \le \lambda_i(B) \le \lambda_{i+m}(A)$.

Proof. After a permutation similarity, B is leading.
Delete the other m coordinates one at a time and apply Theorem 4.4.20 at each step.
Structural Commentary: Principal submatrices come from repeated coordinate deletion.
Key Insight: Each deletion shifts eigenvalue indices by at most one. □

Theorem 4.4.23. *Schur Majorization of the Diagonal*
Let $A \in \mathrm{Mat}(n, n, \mathbb{C})$ *be Hermitian, with increasing eigenvalues* $\lambda_i(A)$ *and increasing diagonal entries* d_i.
Then $\sum_{i=1}^{k} \lambda_i(A) \le \sum_{i=1}^{k} d_i$, $(1 \le k \le n)$ *with equality for* $k = n$.

Proof. Permute coordinates so that $a_{ii} = d_i$, and let A_k be the leading $k \times k$ principal submatrix.
Cauchy interlacing gives $\lambda_i(A) \le \lambda_i(A_k)$ for $1 \le i \le k$, hence $\sum_{i=1}^{k} \lambda_i(A) \le \sum_{i=1}^{k} \lambda_i(A_k) = \mathrm{tr}(A_k) = \sum_{i=1}^{k} d_i$.
For $k = n$, this is $\mathrm{tr}(A) = \mathrm{tr}(A)$.
Structural Commentary: Ordered diagonal sums are traces of leading principal blocks.
Key Insight: Diagonal majorization is Cauchy interlacing plus trace. □

Method 4.4.24. *Eigenvalue Interlacing Method*
To prove or use an interlacing estimate:

(i) order Hermitian eigenvalues increasingly;
(ii) identify dimension lost: rank-one perturbation, one deleted coordinate, m-coordinate principal submatrix;
(iii) choose test vectors by intersecting spectral subspaces with the unaffected subspace;
(iv) convert the Rayleigh quotient bounds into eigenvalue inequalities.

Theorem 4.4.25. *Weyl Eigenvalue Inequalities*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{C})$ *be Hermitian. For* $1 \le i \le n$,

$$\lambda_i(A + B) \le \lambda_{i+j}(A) + \lambda_{n-j}(B), \qquad 0 \le j \le n - i,$$
$$\lambda_{i-j+1}(A) + \lambda_j(B) \le \lambda_i(A + B), \qquad 1 \le j \le i.$$

Proof. Let $\mathbf{u}_k, \mathbf{v}_k, \mathbf{w}_k$ be increasing eigenbases for $A, A+B, B$.
Fix i and $0 \le j \le n-i$, and set $\mathcal{U} = \operatorname{span}(\mathbf{u}_1, \dots, \mathbf{u}_{i+j})$, $\mathcal{V} = \operatorname{span}(\mathbf{v}_i, \dots, \mathbf{v}_n)$, $\mathcal{W} = \operatorname{span}(\mathbf{w}_1, \dots, \mathbf{w}_{n-j})$.
Their dimensions sum to $2n+1$, so choose a unit $\mathbf{x} \in \mathcal{U} \cap \mathcal{V} \cap \mathcal{W}$.
Then $\lambda_i(A+B) \le \mathbf{x}^*(A+B)\mathbf{x} = \mathbf{x}^*A\mathbf{x} + \mathbf{x}^*B\mathbf{x} \le \lambda_{i+j}(A) + \lambda_{n-j}(B)$.
The lower bound follows by applying this upper bound to $-A$ and $-B$ and using $\lambda_k(-C) = -\lambda_{n-k+1}(C)$.
Structural Commentary: Weyl bounds come from intersecting the spectral flags of A, B, and $A+B$.
Key Insight: One shared test vector controls all three Rayleigh quotients. □

Method 4.4.26. *Weyl Bound Method*
To bound eigenvalues of $A+B$:

(i) choose the target index i;
(ii) use the upper Weyl family with $j = 0, \dots, n-i$ to get upper bounds;
(iii) use the lower Weyl family with $j = 1, \dots, i$ to get lower bounds;
(iv) for positive semidefinite or low-rank B, insert the known zero or nonnegative eigenvalues to recover monotonicity or interlacing.

Theorem 4.4.27. *Brauer Rank-One Eigenvalue Shift*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$, let $A\mathbf{x} = \lambda\mathbf{x}$ with $\mathbf{x} \neq \mathbf{0}$, and let $\mathbf{y} \in \mathbb{C}^n$. If $\lambda, \lambda_2, \dots, \lambda_n$ are the eigenvalues of A with multiplicity, then the eigenvalues of $A + \mathbf{x}\mathbf{y}^*$ are $\lambda + \mathbf{y}^*\mathbf{x}, \lambda_2, \dots, \lambda_n$.
Moreover, $\mathbf{x}$ is an eigenvector of $A + \mathbf{x}\mathbf{y}^*$ corresponding to $\lambda + \mathbf{y}^*\mathbf{x}$.

Proof. Put $\mathbf{u} = \mathbf{x}/\|\mathbf{x}\|_2$. By Schur triangularisation, choose a unitary $U = [\mathbf{u}\ U_2]$ such that $U^*AU = \begin{bmatrix} \lambda & * \\ 0 & T' \end{bmatrix}$, where the diagonal entries of T' are $\lambda_2, \dots, \lambda_n$. Also $U^*(\mathbf{x}\mathbf{y}^*)U = \begin{bmatrix} \mathbf{y}^*\mathbf{x} & * \\ 0 & 0 \end{bmatrix}$. Therefore $U^*(A + \mathbf{x}\mathbf{y}^*)U$ is upper triangular with diagonal entries $\lambda + \mathbf{y}^*\mathbf{x}, \lambda_2, \dots, \lambda_n$. Finally, $(A + \mathbf{x}\mathbf{y}^*)\mathbf{x} = A\mathbf{x} + \mathbf{x}(\mathbf{y}^*\mathbf{x}) = (\lambda + \mathbf{y}^*\mathbf{x})\mathbf{x}$.
Structural Commentary: A rank-one eigenvector update changes only the first Schur diagonal entry.
Key Insight: Align the update with the first coordinate; only one diagonal term shifts. □

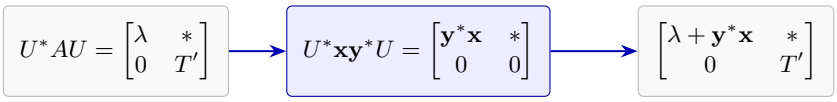


Figure 97: Brauer's update shifts one Schur diagonal entry.

Corollary 4.4.28. *Damped Rank-One Eigenvalue Reset*
Let $A \in \operatorname{Mat}(n, n, \mathbb{C})$ have eigenpair $(\lambda, \mathbf{x})$, let $\mathbf{y}^*\mathbf{x} = 1$, and let $\tau \in \mathbb{C}$. If $\lambda, \lambda_2, \dots, \lambda_n$ are the eigenvalues of A, then $A_\tau := \tau A + (1-\tau)\lambda\mathbf{x}\mathbf{y}^*$ has eigenvalues $\lambda, \tau\lambda_2, \dots, \tau\lambda_n$.

Proof. The eigenvalues of τA are $\tau\lambda, \tau\lambda_2, \dots, \tau\lambda_n$. Since $A_\tau = \tau A + \mathbf{x}((1-\overline{\tau})\overline{\lambda}\mathbf{y})^*$, Theorem 4.4.27 shifts only $\tau\lambda$ by $(1-\tau)\lambda\mathbf{y}^*\mathbf{x} = (1-\tau)\lambda$. Thus $\tau\lambda$ becomes λ, while the other eigenvalues stay $\tau\lambda_i$.
Structural Commentary: Damping scales the spectrum; Brauer's update restores one eigenvalue.
Key Insight: Brauer's theorem changes one spectral coordinate without moving the rest. □

4.5 Numerical Linear Algebra

Numerical linear algebra turns matrix structure into algorithms for direct solves, least squares, eigenvalues, SVD, structured transforms, Krylov methods, iteration, and preconditioning.

Perspective 4.5.1. *Solve with Factors Before Forming Inverses*
For $A\mathbf{x} = \mathbf{b}$, use available factors such as LU, QR, or Cholesky to solve triangular or orthogonal subproblems. This avoids forming A^{-1} and preserves sparse or triangular structure.

Perspective 4.5.2. *Iterative Refinement by Residual Correction*
Given an approximate solution $\mathbf{x}_k$ of $A\mathbf{x} = \mathbf{b}$, set $\mathbf{r}_k = \mathbf{b} - A\mathbf{x}_k$.
Solve $A\mathbf{d}_k = \mathbf{r}_k$ with the same factors or preconditioner, then update $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{d}_k$.
Exact arithmetic corrects in one step; floating point helps when $\mathbf{r}_k$ is accurate and conditioning is moderate.

Method 4.5.3. *Augmented Inverse Computation Preview*
For $A \in \operatorname{Mat}(n, n, \mathbb{F})$:

(i) Form $[\,A\ I_n\,]$.
(ii) Apply Gauss–Jordan reduction to the left block.

(iii) If $[\,A\ I_n\,] \sim [\,I_n\ B\,]$, then $B = A^{-1}$.
(iv) Otherwise A is not invertible, by Corollary 1.3.95.

Proposition 4.5.4. *QR Route for Least Squares*
Let $A = QR \in \mathrm{Mat}(m,n,\mathbb{F})$ be a full-column-rank narrow QR factorisation. For $\mathbf{y} \in \mathbb{F}^m$, the least-squares solution solves $R\mathbf{x} = Q^*\mathbf{y}$. Moreover $P_{\mathrm{col}(A)} = QQ^*$.

Proof. The normal equations reduce by $A = QR$ and $Q^*Q = I_n$ to $R^*R\mathbf{x} = R^*Q^*\mathbf{y}$. Since R is invertible, this is equivalent to $R\mathbf{x} = Q^*\mathbf{y}$. Also $\mathrm{col}(A) = \mathrm{col}(Q)$, and the columns of Q are orthonormal, so $P_{\mathrm{col}(A)} = QQ^*$.
Structural Commentary: QR replaces $\mathrm{col}(A)$ by an orthonormal basis Q.
Key Insight: Project onto Q, then solve triangularly. □

Perspective 4.5.5. *Why Normal Equations Can Be Numerically Fragile*
For full-column-rank least squares, solving $A^*A\mathbf{x} = A^*\mathbf{y}$ squares the spectral condition number because $\kappa_2(A^*A) = \kappa_2(A)^2$. QR-based methods avoid this squaring and therefore preserve conditioning better.

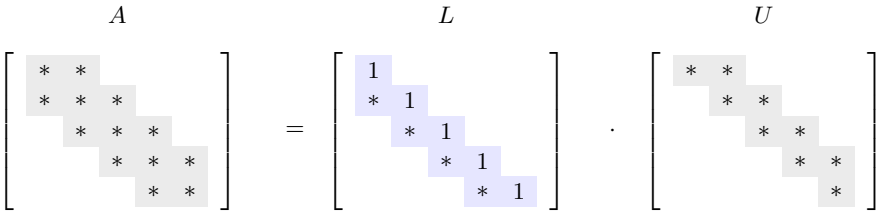


Figure 98: A tridiagonal matrix can factor as bidiagonal L times bidiagonal U, retaining local band structure.

Definition 4.5.6. *Band and Tridiagonal Matrices*
For $A = [a_{ij}] \in \mathrm{Mat}(n,n,\mathbb{F})$ and $1 \le w \le n$, say that A has *half-bandwidth* w if $a_{ij} = 0$ whenever $|i-j| \ge w$. Thus $w = 1$ gives a diagonal matrix, while $w = 2$ gives a *tridiagonal* matrix.

Method 4.5.7. *Banded Elimination Preview*
Let $A \in \mathrm{Mat}(n,n,\mathbb{F})$ have half-bandwidth w.

(i) Eliminate without row exchanges that leave the band.
(ii) At pivot k, update only entries with $|i-k| < w$ and $|j-k| < w$.
(iii) Obtain $A = LU$ with L and U of comparable bandwidth.

Cost: $O(nw^2)$ operations; for tridiagonal matrices ($w = 2$), $O(n)$.

Perspective 4.5.8. *Five-Point Stencils Are Banded but Less Local*
The 1D second-difference pattern $-1, 2, -1$ is tridiagonal; the row-ordered $N \times N$ five-point stencil is still sparse, but vertical neighbors create diagonals offset by about N. Thus bandwidth grows with grid width, and incomplete LU keeps a prescribed sparse pattern for a cheap approximate solve or preconditioner.

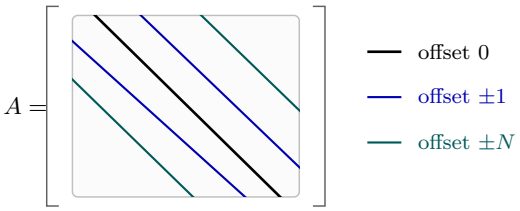


Figure 99: Row ordering gives matrix bands at offsets $0, \pm 1, \pm N$.

Definition 4.5.9. *Upper Hessenberg Matrix*
Let $A = [a_{ij}] \in \mathrm{Mat}(n,n,\mathbb{F})$. Then A is *upper Hessenberg* if $a_{ij} = 0$ whenever $i > j+1$.
Equivalently, all entries below the first subdiagonal are zero.

Numerical linear algebra also replaces a matrix by a unitarily similar one with many forced zero entries. Upper Hessenberg form is the standard preliminary reduction before QR-type eigenvalue algorithms.

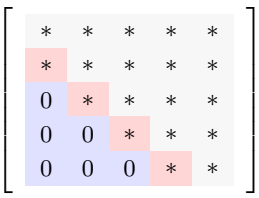


Figure 100: Upper Hessenberg pattern: red is the first subdiagonal; blue entries below it are forced zero.

Theorem 4.5.10. *Unitary Upper Hessenberg Reduction*
Every $A \in \mathrm{Mat}(n, n, \mathbb{F})$ is unitarily similar to upper Hessenberg matrix with real nonnegative subdiagonal entries.

Proof. Set $A_0 = A$. For $k = 1, \ldots, n-2$, suppose columns $1, \ldots, k-1$ of A_{k-1} already have zeros below the first subdiagonal. Let $\mathbf{w}$ be the tail of column k in rows $k+1, \ldots, n$.
By Corollary 2.4.36, choose a unitary V_k with $V_k^* \mathbf{w} = \alpha \mathbf{e}_1$, $\alpha \geq 0$.
Put $U_k = I_k \oplus V_k$ and $A_k = U_k^* A_{k-1} U_k$. Right multiplication by U_k changes only columns $k+1, \ldots, n$, so columns $1, \ldots, k-1$ keep their existing Hessenberg zeros. Left multiplication by U_k^* acts only on rows $k+1, \ldots, n$; those rows are already zero in columns $1, \ldots, k-1$, so the zeros remain. In column k, the tail becomes $V_k^* \mathbf{w} = \alpha \mathbf{e}_1$, so entries below row $k+1$ are zero and the new subdiagonal entry is real nonnegative.
After $n-2$ steps, $A_{n-2} = U^* A U$ is upper Hessenberg, where $U = U_1 \cdots U_{n-2}$ is unitary.
Structural Commentary: Each trailing unitary creates one Hessenberg column and preserves earlier ones.
Key Insight: Kill one column tail at a time. □

Perspective 4.5.11. *Why Hessenberg Form Matters*
Hessenberg form preserves spectral data by unitary similarity, making QR-type eigenvalue algorithms cheaper.

Method 4.5.12. *QR Eigenvalue Algorithm Preview*
For a square matrix A_0, the basic QR eigenvalue iteration is:

(i) In practice, first reduce A_0 to Hessenberg form by Theorem 4.5.10.
(ii) Factor $A_k = Q_k R_k$, with Q_k unitary and R_k upper triangular.
(iii) Set $A_{k+1} = R_k Q_k = Q_k^* A_k Q_k$, using shifts to speed convergence.

Each step is a unitary similarity, so the spectrum is preserved.

Method 4.5.13. *Hessenberg QR Step by Plane Rotations*
Let $H \in \mathrm{Mat}(n, n, \mathbb{C})$ be upper Hessenberg and let $\mu \in \mathbb{C}$.

(i) Set $B_0 = H - \mu I_n$.
(ii) For $j = 1, \ldots, n-1$, choose a unitary plane rotation G_j acting on coordinates $j, j+1$, so that $(G_j^* B_{j-1})_{j+1,j} = 0$ and $B_j = G_j^* B_{j-1}$.
(iii) Put $Q = G_1 \cdots G_{n-1}$ and $R = B_{n-1}$; then $H - \mu I_n = QR$.
(iv) Define the shifted QR update $H_+ = RQ + \mu I_n = Q^* H Q$

Adjacent rotations implement the Hessenberg step by local bulge chasing.

Theorem 4.5.14. *Power Method Dominant-Eigenvector Preview*
Let $A\mathbf{v}_j = \lambda_j \mathbf{v}_j$, where $\mathbf{v}_1, \ldots, \mathbf{v}_n$ is a basis of $\mathbb{C}^n$, and assume $|\lambda_1| > \max_{j \geq 2} |\lambda_j|$.
If $\mathbf{u}_0 = \sum_{j=1}^{n} c_j \mathbf{v}_j$ with $c_1 \neq 0$, then $\lambda_1^{-k} A^k \mathbf{u}_0 \longrightarrow c_1 \mathbf{v}_1$. Thus the directions of $A^k \mathbf{u}_0$ converge to $\mathrm{span}\{\mathbf{v}_1\}$.

Proof. Set $\rho = \max_{j \geq 2} |\lambda_j / \lambda_1| < 1$. Since $A^k \mathbf{v}_j = \lambda_j^k \mathbf{v}_j$, $\lambda_1^{-k} A^k \mathbf{u}_0 = c_1 \mathbf{v}_1 + \sum_{j=2}^{n} c_j \left(\frac{\lambda_j}{\lambda_1}\right)^k \mathbf{v}_j$.

The remainder has norm at most $\rho^k \sum_{j=2}^{n} |c_j| \, \|\mathbf{v}_j\|$, hence tends to 0.

Structural Commentary: The power method is spectral filtering by repeated multiplication.
Key Insight: A spectral gap makes every non-dominant component decay relative to the dominant one. □

Method 4.5.15. *Inverse and Shifted Power Iteration Preview*
Let $\mathcal{N}(\mathbf{z}) := \mathbf{z} / \|\mathbf{z}\|_2$ for $\mathbf{z} \neq \mathbf{0}$.

(i) *Inverse*: if $0 \notin \mathrm{spec}(A)$, solve $A\mathbf{y}_{k+1} = \mathbf{x}_k$ and set $\mathbf{x}_{k+1} = \mathcal{N}(\mathbf{y}_{k+1})$.
This is power iteration on A^{-1} and targets minimal $|\lambda|$.
(ii) *Shift-invert*: for $\sigma \notin \mathrm{spec}(A)$, $(A - \sigma I)\mathbf{y}_{k+1} = \mathbf{x}_k$, $\mathbf{x}_{k+1} = \mathcal{N}(\mathbf{y}_{k+1})$.
Eigenvalues become $(\lambda - \sigma)^{-1}$, targeting λ nearest σ.
(iii) *Shifted power*: $\mathbf{x}_{k+1} = \mathcal{N}\big((A - \sigma I)\mathbf{x}_k\big)$; eigenvalues become $\lambda - \sigma$, emphasizing large $|\lambda - \sigma|$.

Use factors of A or $A - \sigma I$ for solves; do not form explicit inverses.

Method 4.5.16. *SVD Numerics Preview*
Stable numerical SVD computation proceeds as:

(i) Avoid forming A^*A when possible, since it can square the condition number.
(ii) Use orthogonal or unitary transformations to reduce A to bidiagonal form.
(iii) Compute singular values from the bidiagonal problem, then use them for low-rank approximation, pseudoinverse truncation, and least-squares diagnostics through Theorem 4.3.14 and Definition 4.3.23.

Definition 4.5.17. *Discrete Fourier Matrix*
For $n \geq 1$, set $\omega_n = e^{2\pi i/n}$ and $F_n := [\omega_n^{jk}]_{j,k=0}^{n-1}$.
Then $F_n\mathbf{c}$ is the unnormalised DFT of $\mathbf{c} \in \mathbb{C}^n$; $Q_n := n^{-1/2}F_n$ is the normalized Fourier matrix.

Proposition 4.5.18. *Inverse of the Fourier Matrix*
The Fourier matrix satisfies $F_n^*F_n = nI_n$, and therefore $F_n^{-1} = \frac{1}{n}F_n^*$.

Proof. Theorem 3.4.35 gives $Q_n^*Q_n = I_n$ for $Q_n = n^{-1/2}F_n$. Multiplying by n gives $F_n^*F_n = nI_n$, so $n^{-1}F_n^*$ is a left inverse of F_n. Since F_n is square, Theorem 1.3.87 makes it the inverse.
Structural Commentary: Fourier inversion is conjugate transpose plus one scalar normalization.
Key Insight: The inverse is simple because the Fourier columns are orthogonal. □

Method 4.5.19. *Radix-Two FFT Split*
Let $n = 2m$ and $\mathbf{y} = F_n\mathbf{c}$.

(i) Split $\mathbf{c}' := (c_{2r})_{r=0}^{m-1}$ and $\mathbf{c}'' := (c_{2r+1})_{r=0}^{m-1}$.
(ii) Compute $\mathbf{y}' := F_m\mathbf{c}'$ and $\mathbf{y}'' := F_m\mathbf{c}''$.
(iii) With $d_j := \omega_n^j$, and using $\omega_n^2 = \omega_m$ and $\omega_n^m = -1$, set $y_j = y'_j + d_j y''_j$; and $y_{j+m} = y'_j - d_j y''_j$ for $0 \leq j < m$.
(iv) For $n = 2^\ell$, repeat (i)–(iii) for ℓ stages.
n^2 dense products become about $\frac{1}{2}n\log_2 n$ twiddle multiplications plus additions.

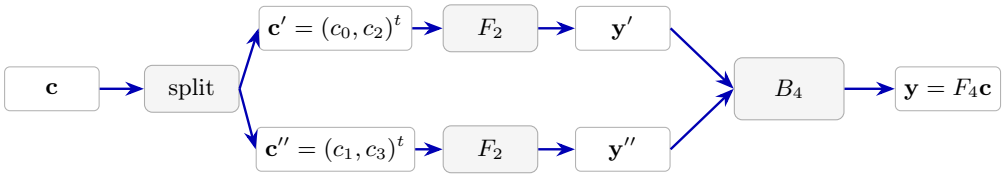


Figure 101: For $n = 4$: split, apply two F_2 transforms, then $B_4\binom{\mathbf{y}'}{\mathbf{y}''} = \binom{\mathbf{y}'+D\mathbf{y}''}{\mathbf{y}'-D\mathbf{y}''}$ with $D = \mathrm{diag}(1, i)$.

Definition 4.5.20. *Krylov Subspace*
For $A \in \mathrm{Mat}(n, n, \mathbb{F})$ and a vector $\mathbf{b} \in \mathbb{F}^n$, the mth *Krylov subspace* generated by $(A, \mathbf{b})$ is
$\mathcal{K}_m(A, \mathbf{b}) := \mathrm{span}\{\mathbf{b}, A\mathbf{b}, \ldots, A^{m-1}\mathbf{b}\}$.

Perspective 4.5.21. *Krylov Methods Preview*
Krylov methods work in $\mathcal{K}_m(A, \mathbf{b})$: Arnoldi and Lanczos build bases, while CG and GMRES impose projection or residual-minimization conditions using only matrix-vector products.

Method 4.5.22. *Arnoldi, Lanczos, and CG Preview*
Let $\mathcal{K}_m = \mathcal{K}_m(A, \mathbf{b})$ and let $Q_m = [\mathbf{q}_1 \ \cdots \ \mathbf{q}_m]$ be an orthonormal basis.

(i) *Arnoldi*: $AQ_m = Q_mH_m + h_{m+1,m}\mathbf{q}_{m+1}\mathbf{e}_m^t$, with H_m upper Hessenberg.
(ii) *Lanczos*: if $A = A^*$, then H_m is tridiagonal and $A\mathbf{q}_j = \beta_{j-1}\mathbf{q}_{j-1} + \alpha_j\mathbf{q}_j + \beta_j\mathbf{q}_{j+1}$.
(iii) *CG*: if $A \succ 0$, choose $\mathbf{x}_k \in \mathbf{x}_0 + \mathcal{K}_k(A, \mathbf{r}_0)$ with $\mathbf{r}_k \perp \mathcal{K}_k(A, \mathbf{r}_0)$; search directions satisfy $\mathbf{p}_i^*A\mathbf{p}_j = 0$, $i \neq j$.

Each step uses one matrix-vector product.

Definition 4.5.23. *Matrix Splitting Iteration*
Let $A = S - T$ with S invertible, and set $B := S^{-1}T$.
The *splitting iteration* for $A\mathbf{x} = \mathbf{b}$ is $S\mathbf{x}_{k+1} = T\mathbf{x}_k + \mathbf{b}$, and $\mathbf{x}_{k+1} = B\mathbf{x}_k + S^{-1}\mathbf{b}$.
For a solution $\mathbf{x}_*$, the error $\mathbf{e}_k := \mathbf{x}_k - \mathbf{x}_*$ satisfies $\mathbf{e}_{k+1} = B\mathbf{e}_k$.

Theorem 4.5.24. *Splitting Convergence Criterion*
Let $A = S - T$, with S invertible, and set $B = S^{-1}T$. For $\mathbf{x}_{k+1} = B\mathbf{x}_k + S^{-1}\mathbf{b}$:

(i) if A is invertible, then $\mathbf{x}_k \to A^{-1}\mathbf{b}$ for all $\mathbf{x}_0, \mathbf{b}$ if and only if $\rho(B) < 1$;

(ii) $\rho(B) < 1$ *implies* A *is invertible.*

Proof.

(i) Let $\mathbf{x}_* = A^{-1}\mathbf{b}$ and $\mathbf{e}_k = \mathbf{x}_k - \mathbf{x}_*$. Then $\mathbf{e}_k = B^k\mathbf{e}_0$, so convergence for all $\mathbf{x}_0, \mathbf{b}$ is equivalent to $B^k \to 0$, equivalently $\rho(B) < 1$ by Theorem 4.1.93.

(ii) If $\rho(B) < 1$, then $1 \notin \operatorname{spec}(B)$, so $I_n - B$ is invertible. Since $A = S - T = S(I_n - B)$, A is invertible.

Structural Commentary: A stationary method is a matrix-power problem in disguise.
Key Insight: The iteration matrix $B = S^{-1}T$ controls the errors. □

Corollary 4.5.25. *Singular Matrices Obstruct Splitting Convergence*
Let $A = S - T$*, with* S *invertible. If* A *is singular, then* $1 \in \operatorname{spec}(S^{-1}T)$.
Hence $\rho(S^{-1}T) \geq 1$*, so* $\rho(S^{-1}T) < 1$ *cannot hold.*

Proof. Choose $\mathbf{0} \neq \mathbf{x} \in \ker A$. Then $(S - T)\mathbf{x} = \mathbf{0}$, so $S\mathbf{x} = T\mathbf{x}$ and $S^{-1}T\mathbf{x} = \mathbf{x}$. Thus 1 is eigenvalue of $S^{-1}T$.
Structural Commentary: Singular A forces an undamped error mode.
Key Insight: A null vector of A becomes a fixed error mode. □

Method 4.5.26. *Stationary Iterative Method*
For $A\mathbf{x} = \mathbf{b}$, a stationary iteration is:

(i) Choose fixed M and $\mathbf{c}$.
(ii) Update $\mathbf{x}_{k+1} = M\mathbf{x}_k + \mathbf{c}$.
(iii) If $\mathbf{x}_*$ exists, use $\mathbf{x}_{k+1} - \mathbf{x}_* = M(\mathbf{x}_k - \mathbf{x}_*)$ and Theorem 4.1.93 for convergence.

Method 4.5.27. *Jacobi, Gauss–Seidel, and SOR Splittings*
Let $A = D + L + U$, with D diagonal, L strictly lower triangular, and U strictly upper triangular. For a splitting $A = S - T$, use $S\mathbf{x}_{k+1} = T\mathbf{x}_k + \mathbf{b}$ and $B = S^{-1}T$.

(i) *Jacobi*: $S = D$, $T = -(L + U)$, $B_J = -D^{-1}(L + U)$.
(ii) *Gauss–Seidel*: $S = D + L$, $T = -U$, $B_{\mathrm{GS}} = -(D + L)^{-1}U$.
(iii) *SOR*: for $\omega \neq 0$, $S_\omega = \frac{1}{\omega}D + L$; $T_\omega = \left(\frac{1}{\omega} - 1\right)D - U$; and $B_\omega = S_\omega^{-1}T_\omega$.

Convergence is controlled by $\rho(B)$; $\omega = 1$ gives Gauss–Seidel.

Corollary 4.5.28. *Strict Diagonal Dominance Gives Jacobi Convergence*
If $A = [a_{ij}] \in \mathrm{Mat}(n, n, \mathbb{C})$ *is strictly row diagonally dominant, then Jacobi converges for every right-hand side.*

Proof. Here $a_{ii} \neq 0$ and $B_J = -D^{-1}(L+U) = I_n - D^{-1}A$. Thus $(B_J)_{ii} = 0$ and $\sum_{j \neq i} |(B_J)_{ij}| = \sum_{j \neq i} |a_{ij}| \,/\, |a_{ii}| < 1$.
Gersgorin gives $\operatorname{spec}(B_J) \subset \{z : |z| < 1\}$, so $\rho(B_J) < 1$; Theorem 4.5.24 applies.
Structural Commentary: Diagonal dominance puts the Jacobi spectrum in the unit disk.
Key Insight: Small off-diagonal row sums force the error matrix to have spectral radius below 1. □

Example 4.5.29. *Two-by-Two Jacobi and Gauss–Seidel*
For $A = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$, $\mathbf{b} \in \mathbb{R}^2$, write $A = D + L + U$.
The Jacobi iteration matrix is $B_J = -D^{-1}(L + U) = \begin{bmatrix} 0 & -\frac{1}{2} \\ -\frac{1}{2} & 0 \end{bmatrix}$, so $\rho(B_J) = \frac{1}{2}$.
The Gauss–Seidel iteration matrix is $B_{\mathrm{GS}} = -(D + L)^{-1}U = \begin{bmatrix} 0 & -\frac{1}{2} \\ 0 & \frac{1}{4} \end{bmatrix}$, so $\rho(B_{\mathrm{GS}}) = \frac{1}{4}$.
Both converge by Theorem 4.5.24, and smaller spectral radius predicts faster asymptotic decay for Gauss–Seidel.

Example 4.5.30. *Poisson Tridiagonal Stationary Rates*
For $A_n = \operatorname{tridiag}(-1, 2, -1)$, Jacobi gives $B_J = I_n - \frac{1}{2}A_n$; $\lambda_k(B_J) = \cos\frac{k\pi}{n+1}$; and $\rho(B_J) = \cos\frac{\pi}{n+1}$.
For the same ordering, Gauss–Seidel satisfies $\rho(B_{\mathrm{GS}}) = \rho(B_J)^2 = \cos^2\frac{\pi}{n+1}$. As $n \to \infty$, both radii tend to 1; fine grids need acceleration or preconditioning.

Perspective 4.5.31. *Preconditioning Preview*
A preconditioner replaces $A\mathbf{x} = \mathbf{b}$ by an equivalent system such as $M^{-1}A\mathbf{x} = M^{-1}\mathbf{b}$ or $AM^{-1}\mathbf{y} = \mathbf{b}$, with $\mathbf{x} = M^{-1}\mathbf{y}$. It should improve conditioning or convergence while being much cheaper to apply than A^{-1}.

Part III
Algebraic Structures

1 Algebraic Language and Categories

1.1 Sets, Maps, Relations, and Universal Properties

The elementary set, order, choice, function, and bijection vocabulary used here is collected in the foundational prerequisites. This subsection continues with set algebra, universal properties, and equivalence-based structure.

Theorem 1.1.1. *Universal Property of Quotients*
Let $\pi : \mathcal{V} \to \mathcal{V}/\mathcal{U}$ be the quotient map.
If $L \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ and $\mathcal{U} \subseteq \ker(L)$, then $\exists!\, \overline{L} \in \mathcal{L}(\mathcal{V}/\mathcal{U}, \mathcal{W})$ such that $\overline{L}\pi = L$, $\overline{L}(\mathbf{v}+\mathcal{U}) = L\mathbf{v}$.
Moreover, $\ker(\overline{L}) = \ker(L)/\mathcal{U}$ and $\operatorname{range}(\overline{L}) = \operatorname{range}(L)$.

Proof. If $\mathbf{v}+\mathcal{U} = \mathbf{w}+\mathcal{U}$, then $\mathbf{v}-\mathbf{w} \in \mathcal{U} \subseteq \ker(L)$, so $L\mathbf{v} = L\mathbf{w}$. Hence the formula defines a linear map $\overline{L}$ and gives $\overline{L}\pi = L$. If $M\pi = L$, then $M(\mathbf{v}+\mathcal{U}) = M\pi(\mathbf{v}) = L\mathbf{v} = \overline{L}(\mathbf{v}+\mathcal{U})$, so $M = \overline{L}$.
Finally, $\mathbf{v}+\mathcal{U} \in \ker(\overline{L}) \iff L\mathbf{v} = \mathbf{0} \iff \mathbf{v} \in \ker(L)$, giving $\ker(\overline{L}) = \ker(L)/\mathcal{U}$.
The range equality follows from $\overline{L}(\mathbf{v}+\mathcal{U}) = L\mathbf{v}$.
Structural Commentary: Maps killing $\mathcal{U}$ descend uniquely to $\mathcal{V}/\mathcal{U}$.
Key Insight: Quotients are the natural domain for maps whose ambiguity contains $\mathcal{U}$. □

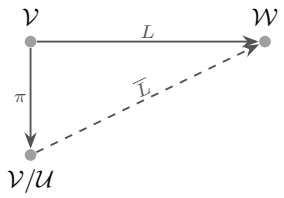


Figure 102: If $\mathcal{U} \subseteq \ker(L)$, then $L = \overline{L}\pi$ uniquely.

Definition 1.1.2. *Equivalence Relations, Classes, Invariants, and Canonical Forms*
Let $\sim$ be a relation on a set X.

(i) $\sim$ is an *equivalence relation* if, for all $x, y, z \in X$, $x \sim x$, $x \sim y \Rightarrow y \sim x$, and $(x \sim y$ and $y \sim z) \Rightarrow x \sim z$.
(ii) If $\sim$ is an equivalence relation, the *equivalence class* of $x \in X$ is $[x] := \{y \in X : y \sim x\}$.
(iii) A function $f : X \to Y$ is an *invariant* of $\sim$ if $x \sim y \Rightarrow f(x) = f(y)$, and a *complete invariant* if $x \sim y \iff f(x) = f(y)$.
(iv) A subset $\mathcal{C} \subseteq X$ is an *set of canonical forms* for $\sim$ if each equivalence class meets $\mathcal{C}$ in exactly one element.

Perspective 1.1.3. *Similarity Invariants and Canonical Forms*
Similarity is equality up to change of basis. Trace, determinant, characteristic polynomials, and minimal polynomials are similarity invariants by Theorems 4.1.21, 3.2.12, and 3.2.26, but they are not generally complete. Canonical forms, such as the Jordan form of Theorem 3.3.6, choose one structured representative of each similarity class when possible.

1.2 Categories and Functors

Definition 1.2.1. *Morphism Terminology for Linear Maps*
For $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$:

(i) T is a *homomorphism*.
(ii) If $\mathcal{V} = \mathcal{W}$, then T is an *endomorphism*.
(iii) If T is injective, then T is a *monomorphism*, or *embedding*.
(iv) If T is surjective, then T is an *epimorphism*.
(v) If T is bijective, then T is an *isomorphism*.
(vi) If T is a bijective endomorphism, then T is an *automorphism*.

1.3 Algebraic Structures as Objects

Definition 1.3.1. *Fields and Ordered Fields*
A *field* is a set F with addition and multiplication satisfying closure, commutativity, associativity, identities $0, 1$ with $1 \neq 0$, additive inverses, multiplicative inverses for nonzero elements, and distributivity.
An *ordered field* is a field such that $y < z$ implies $x + y < x + z$, and $x > 0$ and $y > 0$ imply $xy > 0$.

Remark 1.3.2. *Fields as Objects and Scalars*
Unlike their scalar role in Definition 1.1.1, fields are now Analysis objects: $\mathbb{Q}$ is ordered but incomplete, while $\mathbb{R}$ has the least-upper-bound property.

Proposition 1.3.3. *Basic Consequences of the Field Axioms*
Every field has cancellation, unique identities and inverses, $0x = 0$, no zero divisors, and the sign rules

$$(-x)y = -(xy) = x(-y), \qquad (-x)(-y) = xy.$$

Proof. Cancellation comes from adding or multiplying by the relevant inverse; uniqueness follows from cancellation. Also $0x + 0x = (0 + 0)x = 0x$, so $0x = 0$. If $x, y \neq 0$ and $xy = 0$, multiplying by $x^{-1}y^{-1}$ gives $1 = 0$, impossible. The sign rules follow because the displayed negatives are additive inverses of xy.
Structural Commentary: The field axioms license ordinary algebra.
Key Insight: Inverses force cancellation and signs. □

Proposition 1.3.4. *Basic Inequality Rules in Ordered Fields*
In every ordered field,

(i) $x > 0$ *iff* $-x < 0$;
(ii) multiplying an inequality by a positive element preserves the inequality;
(iii) multiplying an inequality by a negative element reverses the inequality;
(iv) if $x \neq 0$, *then* $x^2 > 0$, *so in particular* $1 > 0$;
(v) if $0 < x < y$, *then* $0 < 1/y < 1/x$.

Proof.

(i) If $x > 0$, then $0 = -x + x > -x + 0$, so $-x < 0$; the converse is the same with $-x$.
(ii) If $y < z$ and $x > 0$, then $z - y > 0$, so $x(z - y) > 0$ and hence $xy < xz$.
(iii) Apply (ii) to $-x > 0$.
(iv) Either $x > 0$ or $-x > 0$, so $x^2 > 0$; take $x = 1$ for $1 > 0$.
(v) Since $x, y > 0$, multiply $x < y$ by $(1/x)(1/y) > 0$ to get $1/y < 1/x$.

Structural Commentary: Ordered fields make algebraic signs compatible with order.
Key Insight: Positivity decides whether multiplication preserves order. □

Perspective 1.3.5. *Finite-Support Coordinate Model*
Proposition 1.3.67 identifies an n-dimensional vector space with $\mathbb{F}^n$, and Theorem 1.3.89 identifies linear maps with matrices after bases are chosen. The space $\mathbb{F}^{(\mathbb{N})}$ is the countable standard-coordinate model with finite support. This is the field analogue of the free-module notation $(R^B)_0$.

1.4 Exactness and Diagrams

Theorem 1.4.1. *First Isomorphism Theorem*
Let $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$. Then $\widehat{T} : \mathcal{V}/\ker(T) \to \operatorname{range}(T)$; $\widehat{T}(\mathbf{v} + \ker(T)) = T\mathbf{v}$ is a well-defined linear isomorphism.

Proof. If $\mathbf{v} + \ker(T) = \mathbf{w} + \ker(T)$, then $\mathbf{v} - \mathbf{w} \in \ker(T)$, so $T\mathbf{v} = T\mathbf{w}$; hence $\widehat{T}$ is well-defined.
Linearity follows from the quotient operations, $\widehat{T}(a(\mathbf{v} + \ker(T)) + (\mathbf{w} + \ker(T))) = aT\mathbf{v} + T\mathbf{w}$.
Also, $\widehat{T}(\mathbf{v} + \ker(T)) = \mathbf{0} \iff T\mathbf{v} = \mathbf{0} \iff \mathbf{v} \in \ker(T) \iff \mathbf{v} + \ker(T) = \ker(T)$, so $\widehat{T}$ is injective.
If $\mathbf{y} \in \operatorname{range}(T)$, then $\mathbf{y} = T\mathbf{v} = \widehat{T}(\mathbf{v} + \ker(T))$ for some $\mathbf{v} \in \mathcal{V}$, so $\widehat{T}$ is surjective. Thus $\widehat{T}$ is linear isomorphism.
Structural Commentary: The quotient by the kernel is the part of the domain actually seen by T.
Key Insight: Image is domain modulo kernel. □

Theorem 1.4.2. *Second Isomorphism Theorem*
Let $\mathcal{S}$ and $\mathcal{T}$ be subspaces of $\mathcal{V}$. Then $(\mathcal{S} + \mathcal{T})/\mathcal{T} \cong \mathcal{S}/(\mathcal{S} \cap \mathcal{T})$.

Proof. Define $L : \mathcal{S} \to (\mathcal{S} + \mathcal{T})/\mathcal{T}$ by $L\mathbf{s} = \mathbf{s} + \mathcal{T}$. Then L is linear by the quotient operations.
Moreover, $\ker(L) = \{\mathbf{s} \in \mathcal{S} : \mathbf{s} + \mathcal{T} = \mathcal{T}\} = \mathcal{S} \cap \mathcal{T}$. If $(\mathbf{s} + \mathbf{t}) + \mathcal{T} \in (\mathcal{S} + \mathcal{T})/\mathcal{T}$, with $\mathbf{s} \in \mathcal{S}$ and $\mathbf{t} \in \mathcal{T}$, then $(\mathbf{s} + \mathbf{t}) + \mathcal{T} = \mathbf{s} + \mathcal{T} = L\mathbf{s}$, so L is surjective. Theorem 1.4.1 gives the stated isomorphism.
Structural Commentary: Modding out by $\mathcal{T}$ leaves only the part of $\mathcal{S}$ not already inside $\mathcal{T}$.
Key Insight: In a quotient by $\mathcal{T}$, adding $\mathcal{T}$ to $\mathcal{S}$ costs nothing. □

Theorem 1.4.3. *Third Isomorphism Theorem*
If $\mathcal{U} \subseteq \mathcal{T} \subseteq \mathcal{V}$ are subspaces, then $(\mathcal{V}/\mathcal{U})/(\mathcal{T}/\mathcal{U}) \cong \mathcal{V}/\mathcal{T}$.

Proof. Define $L : \mathcal{V}/\mathcal{U} \to \mathcal{V}/\mathcal{T}$ by $L(\mathbf{v}+\mathcal{U}) = \mathbf{v}+\mathcal{T}$.
If $\mathbf{v}+\mathcal{U} = \mathbf{w}+\mathcal{U}$, then $\mathbf{v}-\mathbf{w} \in \mathcal{U} \subseteq \mathcal{T}$, so L is well-defined.
It is linear by the quotient operations, and it is surjective since $L(\mathbf{v}+\mathcal{U}) = \mathbf{v}+\mathcal{T}$ for every coset $\mathbf{v}+\mathcal{T}$.
Moreover, $\ker(L) = \{\mathbf{v}+\mathcal{U} : \mathbf{v}+\mathcal{T} = \mathcal{T}\} = \{\mathbf{v}+\mathcal{U} : \mathbf{v} \in \mathcal{T}\} = \mathcal{T}/\mathcal{U}$. Theorem 1.4.1 gives $(\mathcal{V}/\mathcal{U})/(\mathcal{T}/\mathcal{U}) \cong \mathcal{V}/\mathcal{T}$.
Structural Commentary: Collapsing $\mathcal{U}$ and then collapsing $\mathcal{T}/\mathcal{U}$ is the same as collapsing $\mathcal{T}$ at once.
Key Insight: Quotients of quotients cancel only when the nesting is explicit. □

2 Group Theory

2.1 Groups, Subgroups, and Homomorphisms

Groups are the algebraic language of reversible operations. They appear here as the additive structure underlying rings and modules, and later as symmetry groups acting on mathematical objects.

Definition 2.1.1. *Groups*
A *group* is a nonempty set G with a binary operation $(a,b) \mapsto ab$ such that:

(i) $(ab)c = a(bc)$ for all $a,b,c \in G$;
(ii) there is an element $e \in G$ with $eg = ge = g$ for all $g \in G$;
(iii) for each $g \in G$, there is $g^{-1} \in G$ with $gg^{-1} = g^{-1}g = e$.

The group is *abelian* if $ab = ba$ for all $a,b \in G$.
In an abelian group the operation is often written additively, with identity 0 and inverse $-a$.

Example 2.1.2. *Basic Groups*

(i) The bijections $S \to S$ form a group under composition; it is usually nonabelian.
(ii) $\mathrm{Mat}(m,n,\mathbb{F})$ is an abelian group under matrix addition.
(iii) $\mathrm{GL}_n(\mathbb{F})$, the invertible $n \times n$ matrices, is a group under multiplication.
(iv) $\mathbb{Z}/n\mathbb{Z}$ is a finite abelian group under addition modulo n.

Definition 2.1.3. *Subgroup*
A *subgroup* of a group G is a nonempty subset $H \subseteq G$ that is itself a group under operation inherited from G.

Proposition 2.1.4. *Subgroup Test*
Let G be a group and let $H \subseteq G$ be nonempty. Then $H \leq G \iff ab^{-1} \in H \ \forall a,b \in H$.

Proof.

($\Rightarrow$) If $H \leq G$, then $b^{-1} \in H$ and $ab^{-1} \in H$ for all $a,b \in H$.
($\Leftarrow$) Assume $ab^{-1} \in H$ for all $a,b \in H$. Choose $h \in H$. Then $e_G = hh^{-1} \in H$. For $b \in H$, $b^{-1} = e_G b^{-1} \in H$. For $a,c \in H$, since $c^{-1} \in H$, we have $ac = a(c^{-1})^{-1} \in H$. Hence $H \leq G$.

Structural Commentary: Closure under one expression packages identity, inverses, and products.
Key Insight: In a group, subtraction-style closure is enough. □

Definition 2.1.5. *Generated Groups and Homomorphisms*
Let G, H be groups.

(i) For $a \in G$, the *subgroup generated by a* is $\langle a \rangle := \{a^k : k \in \mathbb{Z}\} \leq G$.
(ii) G is *cyclic* if $G = \langle a \rangle$ for some $a \in G$.
If it exists, $\mathrm{ord}(a) := \min\{n \in \mathbb{N}_{>0} : a^n = e_G\}$; then $\langle a \rangle = \{e_G, a, \ldots, a^{\mathrm{ord}(a)-1}\}$.
(iii) A *homomorphism* $\varphi : G \to H$ satisfies $\varphi(ab) = \varphi(a)\varphi(b)$ for all $a,b \in G$.
(iv) For a homomorphism $\varphi : G \to H$, $\ker\varphi := \{g \in G : \varphi(g) = e_H\}$; and $\mathrm{im}\,\varphi := \varphi(G)$.
An *isomorphism* is a bijective homomorphism.

Proposition 2.1.6. *Basic Homomorphism Properties*
Let $\varphi : G \to H$ be a group homomorphism. Then:

(i) $\varphi(e_G) = e_H$;
(ii) $\varphi(g^{-1}) = \varphi(g)^{-1}$ *for all* $g \in G$;
(iii) $\ker\varphi \leq G$;
(iv) $\mathrm{im}\,\varphi \leq H$.

Proof.

(i) Since $\varphi(e_G) = \varphi(e_G e_G) = \varphi(e_G)^2$, cancellation gives $\varphi(e_G) = e_H$.
(ii) For $g \in G$, $\varphi(g)\varphi(g^{-1}) = \varphi(gg^{-1}) = e_H$, hence $\varphi(g^{-1}) = \varphi(g)^{-1}$.
(iii) If $a, b \in \ker\varphi$, then $\varphi(ab^{-1}) = \varphi(a)\varphi(b)^{-1} = e_H$, so $\ker\varphi \leq G$ by Proposition 2.1.4.
(iv) If $x = \varphi(a)$ and $y = \varphi(b)$ lie in $\operatorname{im}\varphi$, then $xy^{-1} = \varphi(ab^{-1}) \in \operatorname{im}\varphi$, so $\operatorname{im}\varphi \leq H$ by Proposition 2.1.4.

Structural Commentary: Homomorphisms preserve all group structure once they preserve multiplication.
Key Insight: $\varphi(ab) = \varphi(a)\varphi(b)$ preserves e, inverses, kernels, and images. □

Proposition 2.1.7. *Kernel Criterion for Injectivity*
Let $\varphi : G \to H$ be a group homomorphism. Then φ injective $\iff$ $\ker\varphi = \{e_G\}$.

Proof.

($\Rightarrow$) If φ injective, $g \in \ker\varphi$, then $\varphi(g) = e_H = \varphi(e_G)$ by Proposition 2.1.6, so $g = e_G$. Hence $\ker\varphi = \{e_G\}$.
($\Leftarrow$) If $\ker\varphi = \{e_G\}$ and $\varphi(a) = \varphi(b)$, then $\varphi(ab^{-1}) = \varphi(a)\varphi(b)^{-1} = e_H$, so $ab^{-1} = e_G$ and $a = b$.

Structural Commentary: Injectivity is detected by the identity fiber $\ker\varphi$.
Key Insight: The kernel measures exactly the failure of injectivity. □

2.2 Group Actions

A group action packages the idea that elements of a group act as reversible motions of a set. The conjugation action is the finite-group ingredient used later in Wedderburn's theorem.

Definition 2.2.1. *Actions, Orbits, and Stabilizers*
Let G be a group and let $\mathcal{X}$ be a set.

(i) *Left action* of G on $\mathcal{X}$ is map $G \times \mathcal{X} \to \mathcal{X}$, $(g, x) \mapsto gx$, s.t. $e_G x = x$ and $(gh)x = g(hx)$ $\forall g, h \in G$, $x \in \mathcal{X}$.
(ii) The *orbit* of $x \in \mathcal{X}$ is $Gx := \{gx : g \in G\}$.
(iii) The *stabilizer* of $x \in \mathcal{X}$ is $G_x := \{g \in G : gx = x\}$.

Theorem 2.2.2. *Orbit-Stabilizer*
Let a finite group G act on $\mathcal{X}$. For each $x \in \mathcal{X}$, $|Gx| = [G : G_x] = \frac{|G|}{|G_x|}$.

Proof. Define $\Phi : G/G_x \to Gx$, $\Phi(gG_x) = gx$.
For $g, h \in G$, $gG_x = hG_x \iff h^{-1}g \in G_x \iff h^{-1}gx = x \iff gx = hx$. Thus Φ is well defined and injective. It is surjective by definition of Gx, hence bijective, so $|Gx| = |G/G_x| = [G : G_x]$.
Structural Commentary: Orbits are parametrized by cosets of the stabilizer.
Key Insight: Same orbit point $\iff$ same stabilizer coset. □

2.3 Normality and Structure Theorems

2.4 Sylow Theory and Finite Groups

Definition 2.4.1. *Conjugacy Classes and Centralizers*
Let G be a group.

(i) The *conjugation action* of G on itself is $g \cdot x := gxg^{-1}$.
(ii) The *conjugacy class* of $x \in G$ is $\{x\}^G := \{gxg^{-1} : g \in G\}$.
(iii) The *centralizer* of $x \in G$ is $C_G(x) := \{g \in G : gx = xg\}$.
(iv) The *center* of G is $Z(G) := \{g \in G : gx = xg \text{ for all } x \in G\}$.

Theorem 2.4.2. *Class Equation*
Let G be finite, $S \subseteq G$ contain one representative of each noncentral conjugacy class. Then $|G| = |Z(G)| + \sum_{x \in S} [G : C_G(x)]$.

Proof. The conjugacy classes form a partition $G = Z(G) \sqcup \bigsqcup_{x \in S} \{x\}^G$.
Indeed, $\left|\{x\}^G\right| = 1 \iff gxg^{-1} = x$ for all $g \in G \iff x \in Z(G)$. For $x \in S$, Theorem 2.2.2 applied to conjugation gives $\left|\{x\}^G\right| = [G : C_G(x)]$. Taking cardinalities yields the formula.
Structural Commentary: G decomposes into central singleton classes and noncentral conjugacy classes.
Key Insight: Count conjugacy classes by orbit-stabilizer. □

2.5 Representation-Theoretic Preview

Definition 2.5.1. *Group Algebra*
For a group G and field $\mathbb{F}$, the *group algebra* is $\mathbb{F}[G] := \{\sum_{g\in G} r_g g : r_g \in \mathbb{F},\ |\{g : r_g \neq 0\}| < \infty\}$.
It has coefficientwise vector-space operations and product $(\sum_g r_g g)(\sum_h s_h h) := \sum_{g,h} r_g s_h (gh)$.

Example 2.5.2. *Group Algebras*
In $\mathbb{F}[G]$, G is an $\mathbb{F}$-basis, e_G is the algebra identity, and associativity follows from $(gh)k = g(hk)$ by bilinearity.
If G is finite, then $\dim_{\mathbb{F}} \mathbb{F}[G] = |G|$. Moreover, $\mathbb{F}[G]$ is commutative iff G is abelian.

Example 2.5.3. *Cyclic Group Algebras and Circulant Shifts*
Let $G = \langle\alpha\rangle$, $|G| = n$, and let $\mathcal{B} = (1, \alpha, \dots, \alpha^{n-1})$ be the ordered $\mathbb{F}$-basis of $\mathbb{F}[G]$.
For $0 \le j, k < n$, $L_{\alpha^k}(\alpha^j) = \alpha^{j+k}$ with exponents read modulo n, so L_{α^k} is the k-fold cyclic shift of $\mathcal{B}$.
Let $P := [L_\alpha]_{\mathcal{B}}$, one has $[L_{\alpha^k}]_{\mathcal{B}} = P^k$, where

$$P = \begin{bmatrix} 0 & 0 & \cdots & 0 & 1 \\ 1 & 0 & \cdots & 0 & 0 \\ 0 & 1 & \ddots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 1 & 0 \end{bmatrix}.$$

Hence $[L_{\sum_{k=0}^{n-1} c_k \alpha^k}]_{\mathcal{B}} = \sum_{k=0}^{n-1} c_k P^k$, precisely the circulant matrices.

3 Rings, Ideals, and Modules

3.1 Rings and Ideals

Rings add multiplication to an abelian group. Ideals are the subsets by which quotient rings can be formed, just as subspaces are the subsets by which quotient vector spaces can be formed.

Definition 3.1.1. *Rings, Identities, and Fields*
Let R be a set with operations $+$ and $\cdot$.

(i) R is a *ring* if $(R, +)$ is an abelian group, multiplication is associative, and $a(b+c) = ab + ac$, $(a+b)c = ac + bc$ for all $a, b, c \in R$.
(ii) A ring R is *commutative* if $ab = ba$ for all $a, b \in R$.
(iii) A ring R has *identity* if there is $1_R \in R$ such that $1_R a = a 1_R = a$ for all $a \in R$.
(iv) A *field* is a commutative ring with identity $1_R \neq 0$ such that every $a \in R \setminus \{0\}$ has $a^{-1} \in R$.

Example 3.1.2. *Basic Rings*

(i) $(\mathbb{Z}, +, \cdot)$ and $(\mathbb{F}[x], +, \cdot)$ are commutative rings with identity 1.
(ii) $\mathrm{Mat}(n, n, \mathbb{F})$ is a unital ring under matrix addition and multiplication, with identity I_n; it is noncommutative for $n \ge 2$.
(iii) $2\mathbb{Z}$ is a commutative ring under inherited addition and multiplication, but has no identity.
(iv) $\mathbb{Z}/n\mathbb{Z}$ is a commutative ring under addition and multiplication modulo n, with identity $1 + n\mathbb{Z}$.

Definition 3.1.3. *Ring Homomorphisms and Subrings*
Let R, S be rings with identity.

(i) A *ring homomorphism* $\varphi : R \to S$ preserves $+$, $\cdot$, and 1, where $\varphi(a+b) = \varphi(a) + \varphi(b)$, $\varphi(ab) = \varphi(a)\varphi(b)$, $\varphi(1_R) = 1_S$ for all $a, b \in R$.
(ii) A *subring* of R is a subset $T \subseteq R$ that is a ring under the inherited operations and has identity 1_R.

Proposition 3.1.4. *Subring Test*
Let R be a ring with identity and let $S \subseteq R$. Then $S \le R \iff 1_R \in S$, $a - b \in S$, $ab \in S$ $\forall a, b \in S$.

Proof.
($\Rightarrow$) If $S \le R$, then $1_R \in S$ and S is closed under subtraction and multiplication.
($\Leftarrow$) Assume $1_R \in S$, $a - b \in S$, and $ab \in S$ for all $a, b \in S$. Then S is nonempty, $(S, +) \le (R, +)$ by Proposition 2.1.4, and multiplication restricts to S. All ring laws and the identity 1_R are inherited from R, so $S \le R$.
Structural Commentary: A subring is a subset stable under the operations and containing the ambient unit.
Key Insight: Subtraction handles the additive group part in one stroke. □

Definition 3.1.5. *Ideals and Generated Ideals*
Let R be a ring.

(i) A nonempty subset $\mathcal{I} \subseteq R$ is a *two-sided ideal* if $a-b \in \mathcal{I}\ \forall a,b \in \mathcal{I}$, and $ra, ar \in \mathcal{I}\ \forall a \in \mathcal{I}$, $r \in R$.

(ii) If R is commutative with identity, $S \subseteq R$, *ideal generated by S* is $\langle S\rangle := \left\{ \sum_{j=1}^{n} r_j s_j : n \geq 1,\ r_j \in R,\ s_j \in S \right\}$.

(iii) *Finite generated ideal notation*: if $S = \{s_1, \ldots, s_n\}$, write $\langle s_1, \ldots, s_n\rangle := \langle S\rangle$.

(iv) In $\mathbb{F}[x]$, $\langle p(x)\rangle = \{q(x)p(x) : q(x) \in \mathbb{F}[x]\}$ is the *principal* ideal generated by $p(x)$.

Proposition 3.1.6. *Generated Ideals Are Smallest*
Let R be a commutative ring with identity and let $S \subseteq R$.
Then $\langle S\rangle$ is an ideal, $S \subseteq \langle S\rangle$, and for every ideal $\mathcal{J} \subseteq R$, $S \subseteq \mathcal{J} \implies \langle S\rangle \subseteq \mathcal{J}$.

Proof. Let $x = \sum_{j=1}^{m} r_j s_j$ and $y = \sum_{k=1}^{n} t_k u_k$ lie in $\langle S\rangle$.
Then $x - y = \sum_{j=1}^{m} r_j s_j + \sum_{k=1}^{n} (-t_k) u_k \in \langle S\rangle$; $ax = \sum_{j=1}^{m} (ar_j) s_j \in \langle S\rangle$ for all $a \in R$, so $\langle S\rangle$ is an ideal.
Since $s = 1_R s \in \langle S\rangle$, $S \subseteq \langle S\rangle$. If $\mathcal{J}$ is an ideal with $S \subseteq \mathcal{J}$, then $r_j s_j \in \mathcal{J}$ and finite sums remain in $\mathcal{J}$, QED.
Structural Commentary: Generated ideals close a set under exactly the operations ideals must allow.
Key Insight: To force S into an ideal, one must include all ring-linear combinations of elements of S. □

Proposition 3.1.7. *Intersections and Chain Unions of Ideals*
Let R be a ring.

(i) If $\{\mathcal{I}_\alpha\}_{\alpha\in A}$ is a nonempty family of ideals of R, then $\bigcap_{\alpha\in A} \mathcal{I}_\alpha$ is an ideal of R.
(ii) If $\{\mathcal{I}_\alpha\}_{\alpha\in A}$ is a chain of ideals of R under inclusion, then $\bigcup_{\alpha\in A} \mathcal{I}_\alpha$ is an ideal of R.

Proof.

(i) Let $\mathcal{I} = \bigcap_{\alpha\in A} \mathcal{I}_\alpha$. If $a, b \in \mathcal{I}$ and $r \in R$, then $a-b, ra, ar \in \mathcal{I}_\alpha$ for every $\alpha \in A$, hence $a-b, ra, ar \in \mathcal{I}$.

(ii) Let $\mathcal{U} = \bigcup_{\alpha\in A} \mathcal{I}_\alpha$. If $a, b \in \mathcal{U}$, then $a \in \mathcal{I}_\alpha$ and $b \in \mathcal{I}_\beta$ for some α, β. Since the family is a chain, one of $\mathcal{I}_\alpha, \mathcal{I}_\beta$ contains both a, b; hence $a-b \in \mathcal{U}$. If $a \in \mathcal{U}$ and $r \in R$, then $a \in \mathcal{I}_\alpha$ for some α, so $ra, ar \in \mathcal{I}_\alpha \subseteq \mathcal{U}$.

Structural Commentary: Chains avoid the usual failure of unions to preserve closure.
Key Insight: Comparability places any finite test inside one ideal. □

Definition 3.1.8. *Quotient Rings*
Let R be a ring and let $\mathcal{I} \trianglelefteq R$ be a two-sided ideal.

(i) The underlying set is the additive quotient group $R/\mathcal{I} = \{a + \mathcal{I} : a \in R\}$.

(ii) Addition and multiplication are $(a+\mathcal{I}) + (b+\mathcal{I}) = (a+b) + \mathcal{I}$ and $(a+\mathcal{I})(b+\mathcal{I}) = ab + \mathcal{I}$.

(iii) With these operations, $R/\mathcal{I}$ is the *quotient ring* of R by $\mathcal{I}$.

Proposition 3.1.9. *Ideals Make Quotient Multiplication Well Defined*
Let $\mathcal{I} \leq (R, +)$. The rule $(a+\mathcal{I})(b+\mathcal{I}) := ab + \mathcal{I}$ is well defined on $R/\mathcal{I}$ iff $\mathcal{I} \trianglelefteq R$.

Proof.

($\Rightarrow$) If the rule is well defined and $i \in \mathcal{I}$, then $i + \mathcal{I} = 0 + \mathcal{I}$.
For $r \in R$, $ri + \mathcal{I} = r0 + \mathcal{I}$, and $ir + \mathcal{I} = 0r + \mathcal{I}$, so $ri, ir \in \mathcal{I}$. Thus $\mathcal{I} \trianglelefteq R$.

($\Leftarrow$) If $\mathcal{I} \trianglelefteq R$ and $a' = a + i$, $b' = b + j$ with $i, j \in \mathcal{I}$, then $a'b' - ab = aj + ib + ij \in \mathcal{I}$.
Hence $a'b' + \mathcal{I} = ab + \mathcal{I}$, so the product is independent of representatives.

Structural Commentary: Ideals are exactly the additive subgroups compatible with multiplication of cosets.
Key Insight: Coset multiplication is well defined exactly for ideals. □

Definition 3.1.10. *Maximal Ideals*
Let R be a ring and let $\mathcal{M} \trianglelefteq R$.

(i) $\mathcal{M}$ is *proper* if $\mathcal{M} \neq R$.

(ii) $\mathcal{M}$ is *maximal* if $\mathcal{M} \neq R$ and, for every ideal $\mathcal{J} \trianglelefteq R$, $\mathcal{M} \subseteq \mathcal{J} \subseteq R \implies \mathcal{J} = \mathcal{M}$ or $\mathcal{J} = R$.

Theorem 3.1.11. *Maximal Ideals and Fields*
Let R be a commutative ring with identity and let $\mathcal{I} \trianglelefteq R$. Then $R/\mathcal{I}$ is a field $\iff$ $\mathcal{I}$ is maximal

Proof.

($\Rightarrow$) If $R/\mathcal{I}$ is a field and $\mathcal{I} \subsetneq \mathcal{J} \trianglelefteq R$, then $\mathcal{J}/\mathcal{I}$ is a nonzero ideal of $R/\mathcal{I}$.
Hence $\mathcal{J}/\mathcal{I} = R/\mathcal{I}$, so $\mathcal{J} = R$. Thus $\mathcal{I}$ is maximal.

($\Leftarrow$) If $\mathcal{I}$ is maximal and $0 \neq a+\mathcal{I} \in R/\mathcal{I}$, then $a \notin \mathcal{I}$.
The ideal $\mathcal{I}+\langle a\rangle$ strictly contains $\mathcal{I}$, so maximality gives $\mathcal{I}+\langle a\rangle = R$. Thus $1 = i + ra$ for some $i \in \mathcal{I}$, $r \in R$, and $(r+\mathcal{I})(a+\mathcal{I}) = 1+\mathcal{I}$. Hence every nonzero coset is invertible.

Structural Commentary: Maximal ideals are exactly kernels whose quotients have no residual ideal structure.
Key Insight: Maximality makes every nonzero coset invertible. □

Theorem 3.1.12. *Existence of Maximal Ideals*
If R is a commutative ring with identity and $R \neq 0$, then R has a maximal ideal.

Proof. Let $\mathcal{P} := \{\mathcal{I} \trianglelefteq R : \mathcal{I} \neq R\}$, ordered by inclusion. Since $R \neq 0$, $\{0\} \in \mathcal{P}$.
If $\mathcal{C} \subseteq \mathcal{P}$ is a chain, then $\mathcal{U} := \bigcup_{\mathcal{I}\in\mathcal{C}} \mathcal{I}$ is an ideal by Proposition 3.1.7.
Also $\mathcal{U} \neq R$: if $1_R \in \mathcal{U}$, then $1_R \in \mathcal{I}$ for some $\mathcal{I} \in \mathcal{C}$, forcing $\mathcal{I} = R$, contradiction. Thus $\mathcal{U} \in \mathcal{P}$ is an upper bound for $\mathcal{C}$. By Zorn's lemma, Principle 6.5.2, $\mathcal{P}$ has a maximal element, i.e. R has a maximal ideal.
Structural Commentary: Maximal ideals are produced by maximality, not by an explicit formula in general.
Key Insight: Chain unions keep properness because containing 1 is a finite obstruction. □

Definition 3.1.13. *Zero Divisors and Integral Domains*
Let R be a ring.

(i) A *zero divisor* is a nonzero $a \in R$ such that $\exists\, 0 \neq b \in R$ with $ab = 0$ or $ba = 0$.
(ii) An *integral domain* is a commutative ring with identity $1_R \neq 0$ and no zero divisors.

Proposition 3.1.14. *Cancellation in Integral Domains*
Let R be a commutative ring with identity $1_R \neq 0$.
Then R is an integral domain $\iff$ $(\forall r, x, y \in R)(r \neq 0 \wedge rx = ry \Rightarrow x = y)$.

Proof.

($\Rightarrow$) If $rx = ry$ and $r \neq 0$, then $r(x-y) = 0$; no zero divisors gives $x - y = 0$.
($\Leftarrow$) If $ab = 0$ and $a \neq 0$, then $ab = a0$; cancellation gives $b = 0$, so no zero divisors exist.

Structural Commentary: Domains are exactly cancellative commutative rings.
Key Insight: Zero divisors obstruct cancellation. □

Definition 3.1.15. *Field of Quotients*
Let R be an integral domain. On $R \times (R \setminus \{0\})$, set $(a,b) \sim (c,d) \iff ad = bc$.
The *field of quotients* of R is $(R \times (R \setminus \{0\}))/\sim$; write $[(a,b)] = a/b$. For $b, d \neq 0$, $\frac{a}{b} + \frac{c}{d} = \frac{ad+bc}{bd}$; $\frac{a}{b}\frac{c}{d} = \frac{ac}{bd}$.

Proposition 3.1.16. *Domains Embed in Fields*
Let R be an integral domain with field of quotients K.
Then $\iota : R \hookrightarrow K$; $\iota(a) = a/1$ is an injective ring homomorphism.

Proof. By Proposition 3.1.14, the relation and operations in Definition 3.1.15 are well defined.
For $a, b \in R$, $\iota(a+b) = \frac{a+b}{1} = \frac{a}{1} + \frac{b}{1}$; $\iota(ab) = \frac{ab}{1} = \frac{a}{1}\frac{b}{1}$; and $\iota(1_R) = \frac{1_R}{1}$..
Thus ι is a ring homomorphism. If $\iota(a) = \iota(b)$, then $a/1 = b/1$, so $a = b$; hence ι is injective.
Structural Commentary: The field of quotients formally adds denominators to a domain.
Key Insight: Integral domains are exactly the setting where fraction arithmetic is consistent. □

Definition 3.1.17. *Characteristic*
Let R be a ring with identity. Its *characteristic* is $\operatorname{char}(R) := \begin{cases} \min\{n \in \mathbb{N}_{>0} : n1_R = 0_R\}, & \text{if this set is nonempty,} \\ 0, & \text{otherwise.} \end{cases}$.

Proposition 3.1.18. *Finite Domains Have Prime Characteristic*
Let R be a finite ring with identity. Then $\operatorname{char}(R) \neq 0$. If R is an integral domain, then $\operatorname{char}(R)$ is prime.

Proof. Since R is finite, the map $n \mapsto n1_R$ from $\mathbb{N}$ to R is not injective.
Thus $i1_R = j1_R$ for some $i < j$, so $(j-i)1_R = 0_R$ and $\operatorname{char}(R) > 0$.
Assume R is a domain and set $c = \operatorname{char}(R)$. Since $1_R \neq 0_R$, $c > 1$.
If $c = ab$ with $1 < a, b < c$, then $(a1_R)(b1_R) = ab1_R = c1_R = 0_R$.
By minimality of c, $a1_R \neq 0_R$ and $b1_R \neq 0_R$, contradicting that R has no zero divisors. Hence c is prime.
Structural Commentary: Characteristic records the additive order of the unit.
Key Insight: In a domain, a composite characteristic would factor zero nontrivially. □

Definition 3.1.19. *Algebras over a Field*
Let $\mathbb{F}$ be a field.
An *$\mathbb{F}$-algebra* is an $\mathbb{F}$-vector space $\mathcal{A}$ with a ring product satisfying $\lambda(ab) = (\lambda a)b = a(\lambda b)$ $\forall \lambda \in \mathbb{F}$ and $a, b \in \mathcal{A}$.

Example 3.1.20. *Matrix and Polynomial Algebras*
$\operatorname{Mat}(n, n, \mathbb{F})$ and $\mathbb{F}[x]$ are algebras over $\mathbb{F}$. The first is generally noncommutative; the second is commutative.

3.2 Polynomial Rings and Factorization

Polynomial rings over fields behave like the integers in the ways linear algebra needs: division with remainder, gcds, Bezout identities, irreducible factorization, and principal ideals.

Definition 3.2.1. *Polynomials over a Field*
Let $\mathbb{F}$ be a field.

(i) $\mathbb{F}[x]$ denotes the ring of polynomials in one variable over $\mathbb{F}$.
(ii) If $0 \neq p \in \mathbb{F}[x]$ and $p(x) = a_0 + a_1x + \cdots + a_nx^n$, $a_n \neq 0$, then $\operatorname{lc}(p) := a_n$, $\deg p := n$. Also, p is *monic* iff $\operatorname{lc}(p) = 1$. Set $\deg 0 = -\infty$.
(iii) A nonzero $p \in \mathbb{F}[x]$ *splits over* $\mathbb{F}$ iff $p = \alpha \prod_{i=1}^{m} (x - \lambda_i)$ for some $m \geq 0$, $\alpha \in \mathbb{F}^\times$, and $\lambda_i \in \mathbb{F}$.

Theorem 3.2.2. *Polynomial Division Algorithm*
Let $f, g \in \mathbb{F}[x]$ with $g \neq 0$. Then $\exists!\, (q, r) \in \mathbb{F}[x]^2$ such that $f = qg + r$, $r = 0$ or $\deg r < \deg g$.

Proof. Let $d = \deg g$ and $\ell = \operatorname{lc}(g)$. Existence follows by induction on $\deg f$. If $f = 0$ or $\deg f < d$, take $(q, r) = (0, f)$. Otherwise, let $n = \deg f$, $a = \operatorname{lc}(f)$, and $f_1 = f - a\ell^{-1}x^{n-d}g$. Then $\deg f_1 < n$.
By induction, $f_1 = q_1g + r$ with $r = 0$ or $\deg r < d$, hence $f = (q_1 + a\ell^{-1}x^{n-d})g + r$.
For uniqueness, if $qg + r = q'g + r'$ with admissible remainders, then $(q - q')g = r' - r$.
If $q \neq q'$, then $\deg((q - q')g) \geq d$, while $\deg(r' - r) < d$, contradiction. Hence $q = q'$ and then $r = r'$.
Structural Commentary: Division works because the leading coefficient of g is invertible in the field.
Key Insight: Cancel leading terms until the remainder is too small to divide again. □

Definition 3.2.3. *Divisibility, GCDs, and Coprimality*
For $p, q \in \mathbb{F}[x]$, write $p \mid q$ if $q = hp$ for some $h \in \mathbb{F}[x]$.
A *greatest common divisor* of $f, g \in \mathbb{F}[x]$, not both zero, is the unique monic polynomial $\gcd(f, g)$ such that:

(i) $\gcd(f, g) \mid f$ and $\gcd(f, g) \mid g$;
(ii) every common divisor of f and g divides $\gcd(f, g)$.

The polynomials f, g are *relatively prime* if $\gcd(f, g) = 1$.

Theorem 3.2.4. *Polynomial GCD and Bezout Identity*
For $f, g \in \mathbb{F}[x]$, not both zero, $\exists$ a unique monic $d \in \mathbb{F}[x]$ s.t. $d \mid f$, $d \mid g$, $h \mid f$, $h \mid g \implies h \mid d$
Moreover, $d = \gcd(f, g)$ satisfies $d = af + bg$ for some $a, b \in \mathbb{F}[x]$.
Hence $\gcd(f, g) = 1$ iff $\exists\, a, b \in \mathbb{F}[x]$ with $af + bg = 1$.

Proof. If one of f, g is 0, the claim is immediate from the monic associate of the other.
Assume $f, g \neq 0$ and run the Euclidean algorithm: $r_{i-1} = q_ir_i + r_{i+1}$, where $r_{i+1} < \deg r_i$, starting with $(r_0, r_1) = (f, g)$ and ending with $r_N \neq 0$, $r_{N+1} = 0$. Then $r_N \mid r_{N-1}$, and upward substitution gives $r_N \mid f, g$.
Conversely, any common divisor of f, g divides every r_i, hence divides r_N. Thus $d := \operatorname{lc}(r_N)^{-1}r_N$ is a monic gcd; uniqueness follows since two monic gcds divide each other.
Back-substitution gives $d = af + bg$ for some $a, b \in \mathbb{F}[x]$. Finally, $d = 1$ gives $af + bg = 1$; conversely, if $af + bg = 1$, every common divisor divides 1, so $d = 1$.
Structural Commentary: The Euclidean algorithm produces both the gcd and its Bezout certificate.
Key Insight: GCD data is ideal data in $\mathbb{F}[x]$. □

Definition 3.2.5. *Irreducible Polynomials*
A nonconstant polynomial $p \in \mathbb{F}[x]$ is *irreducible over* $\mathbb{F}$ if every factorization $p = ab$ in $\mathbb{F}[x]$ has a or b constant.

Proposition 3.2.6. *Irreducible Polynomials Are Prime in $\mathbb{F}[x]$*
Let $p \in \mathbb{F}[x]$ be nonconstant. Then p is irreducible over $\mathbb{F}$ iff $\forall f, g \in \mathbb{F}[x]$, $p \mid fg \implies p \mid f$ or $p \mid g$.

Proof.

($\Rightarrow$) Assume p irreducible and $p \mid fg$. If $p \mid f$, done. Otherwise $\gcd(p, f) = 1$, so Theorem 3.2.4 gives $a, b \in \mathbb{F}[x]$ with $ap + bf = 1$. Multiplying by g gives $apg + bfg = g$; since $p \mid apg$ and $p \mid bfg$, we get $p \mid g$.
($\Leftarrow$) Assume the divisibility property and write $p = ab$. Since $p \mid ab$, either $p \mid a$ or $p \mid b$. If $p \mid a$, then $a = ph$, so $p = phb$ and cancellation gives $1 = hb$, hence $b \in \mathbb{F}^\times$. Similarly, $p \mid b$ forces $a \in \mathbb{F}^\times$. Thus p is irreducible.

Structural Commentary: Over $\mathbb{F}[x]$, irreducible and prime are the same notion.
Key Insight: Bezout turns coprimality into a divisibility test. □

Theorem 3.2.7. *Unique Factorization in $\mathbb{F}[x]$*

For every nonconstant $f \in \mathbb{F}[x]$, $\exists\, c \in \mathbb{F}^\times$ and monic irreducibles $p_1, \ldots, p_m \in \mathbb{F}[x]$ such that $f = c \prod_{i=1}^{m} p_i$.

If $f = d \prod_{j=1}^{n} q_j$ is another such factorization, then $m = n$, $c = d$, and $p_i = q_{\sigma(i)}$ for some $\sigma \in S_m$.

Proof.

(i) *Existence*: Induct on $\deg f$. If f is irreducible, write $f = \mathrm{lc}(f)\big(\mathrm{lc}(f)^{-1}f\big)$.
If $f = uv$ with $0 < \deg u, \deg v < \deg f$, apply induction to u and v and concatenate the factors.

(ii) *Uniqueness*: Suppose $c\prod_{i=1}^{m} p_i = d\prod_{j=1}^{n} q_j$ with all p_i, q_j monic irreducible. Since $p_1 \mid \prod_{j=1}^{n} q_j$, iterating Proposition 3.2.6 gives $p_1 \mid q_j$ for some j. As p_1, q_j are monic irreducible, $p_1 = q_j$.
Cancel and continue by induction on m; the final scalar equality follows from leading coefficients.

Structural Commentary: Polynomial factorization over a field mirrors prime factorization in $\mathbb{Z}$.
Key Insight: The prime property of irreducibles drives uniqueness. □

Definition 3.2.8. *Units, Associates, Prime and Irreducible Elements*
Let R be an integral domain.

(i) A *unit* is an element $u \in R$ with $uv = 1$ for some $v \in R$.
(ii) Nonzero elements $a, b \in R$ are *associates*, written $a \sim b$, if $a = ub$ for some unit u.
(iii) A nonzero nonunit $p \in R$ is *prime* if $p \mid ab$ implies $p \mid a$ or $p \mid b$.
(iv) A nonzero nonunit $r \in R$ is *irreducible* if $r = ab$ implies a or b is a unit.

Proposition 3.2.9. *Divisibility and Principal Ideals*
Let R be an integral domain and let $r, s \in R$ be nonzero.

(i) $r \in R^\times \iff \langle r\rangle = R$.
(ii) $r \sim s \iff \langle r\rangle = \langle s\rangle$.
(iii) $r \mid s \iff \langle s\rangle \subseteq \langle r\rangle$.

Proof.

(i) ($\Rightarrow$) If $rv = 1$, then $1 \in \langle r\rangle$, so $\langle r\rangle = R$.
(i) ($\Leftarrow$) If $\langle r\rangle = R$, then $1 = ar$ for some $a \in R$, hence $r \in R^\times$.
(ii) ($\Rightarrow$) If $r = us$ with $u \in R^\times$, then $r \in \langle s\rangle$ and $s = u^{-1}r \in \langle r\rangle$, so $\langle r\rangle = \langle s\rangle$.
(ii) ($\Leftarrow$) If $\langle r\rangle = \langle s\rangle$, then $r = as$ and $s = br$ for some $a, b \in R$.
Thus $r = abr$, so $ab = 1$ by cancellation; hence $a \in R^\times$ and $r \sim s$.
(iii) ($\Rightarrow$) If $s = ar$, then every element of $\langle s\rangle$ is of the form $xs = xar \in \langle r\rangle$.
(iii) ($\Leftarrow$) If $\langle s\rangle \subseteq \langle r\rangle$, then $s \in \langle r\rangle$, so $s = ar$ for some $a \in R$; hence $r \mid s$.

Structural Commentary: Divisibility reverses inclusion of principal ideals.
Key Insight: Principal ideals translate multiplicative divisibility into order by containment. □

Definition 3.2.10. *PIDs and UFDs*
Let R be an integral domain.

(i) R is a *principal ideal domain* if every ideal $\mathcal{I} \trianglelefteq R$ has $\mathcal{I} = \langle a\rangle$ for some $a \in R$.
(ii) R is a *unique factorization domain* if every $0 \neq a \in R \setminus R^\times$ factors into irreducibles, uniquely up to reordering and associates.

Theorem 3.2.11. *$\mathbb{Z}$ and $\mathbb{F}[x]$ Are Principal Ideal Domains*

(i) If $0 \neq \mathcal{I} \trianglelefteq \mathbb{Z}$, then $\mathcal{I} = \langle d\rangle$ for $d = \min(\mathcal{I} \cap \mathbb{N}_{>0})$.
(ii) If $0 \neq \mathcal{I} \trianglelefteq \mathbb{F}[x]$, then $\mathcal{I} = \langle m\rangle$ for the unique monic $m \in \mathcal{I}$ of minimal degree.

Proof.

(i) If $\mathcal{I} = 0$, then $\mathcal{I} = \langle 0\rangle$. Otherwise choose $d = \min(\mathcal{I} \cap \mathbb{N}_{>0})$. For $a \in \mathcal{I}$, write $a = qd + r$ with $0 \leq r < d$. Since $r = a - qd \in \mathcal{I}$, minimality gives $r = 0$, so $\mathcal{I} = \langle d\rangle$.

(ii) If $\mathcal{I} = 0$, then $\mathcal{I} = \langle 0\rangle$. Otherwise choose monic $m \in \mathcal{I}$ of minimal degree. For $p \in \mathcal{I}$, divide $p = qm + r$ with $r = 0$ or $\deg r < \deg m$. Since $r = p - qm \in \mathcal{I}$, minimality gives $r = 0$, so $\mathcal{I} = \langle m\rangle$.
If m' is another monic minimal generator, then $m \mid m'$ and $m' \mid m$, hence $m = m'$.

Structural Commentary: Euclidean division turns arbitrary ideals into one-generator ideals.
Key Insight: Choose the smallest positive integer or smallest-degree monic polynomial and divide by it. □

Corollary 3.2.12. *GCDs Generate Polynomial Ideals*
For $0 \neq p_1, \ldots, p_n \in \mathbb{F}[x]$ and monic $m = \gcd(p_1, \ldots, p_n)$, $\langle p_1, \ldots, p_n\rangle = \langle m\rangle$.

Proof. By Theorem 3.2.11, $\langle p_1, \ldots, p_n\rangle = \langle g\rangle$ for a unique monic g. Since $p_i \in \langle g\rangle$, $g \mid p_i$ for all i.
If $q \mid p_i$ for all i, then $\langle p_1, \ldots, p_n\rangle \subseteq \langle q\rangle$, so $g \in \langle q\rangle$ and $q \mid g$. Hence $g = m$, and the claim follows.
Structural Commentary: In $\mathbb{F}[x]$, gcds and finitely generated ideals are the same data.
Key Insight: The gcd is the monic generator of the ideal generated by the polynomials. □

Example 3.2.13. *A Polynomial Ring That Is Not a PID*
In $\mathbb{F}[x,y]$, the ideal $\langle x,y\rangle$ is not principal. If $\langle x,y\rangle = \langle p\rangle$, then $p \mid x$ and $p \mid y$, so $p \in \mathbb{F}^\times$. Hence $\langle x,y\rangle = \mathbb{F}[x,y]$, contradicting $1 \notin \langle x,y\rangle$ since every element of $\langle x,y\rangle$ has zero constant term.

Proposition 3.2.14. *PIDs Satisfy the Ascending Chain Condition*
If R is a PID and $\mathcal{I}_1 \subseteq \mathcal{I}_2 \subseteq \cdots$ is a chain of ideals, then $\exists N$ such that

$$\mathcal{I}_j = \mathcal{I}_N \qquad \forall j \geq N.$$

Proof. Let $\mathcal{I} = \bigcup_{j\geq 1} \mathcal{I}_j$. By Proposition 3.1.7, $\mathcal{I}$ is an ideal. Since R is a PID, $\mathcal{I} = \langle a\rangle$. Choose N with $a \in \mathcal{I}_N$. Then $\mathcal{I} = \langle a\rangle \subseteq \mathcal{I}_N \subseteq \mathcal{I}$, so $\mathcal{I}_j = \mathcal{I}_N$ for all $j \geq N$.
Structural Commentary: The generator appears at a finite stage.
Key Insight: Principal generation converts infinite ideal growth into a finite containment. □

Theorem 3.2.15. *Prime, Irreducible, and Bezout in a PID*
Let R be a principal ideal domain.

(i) For $0 \neq r \in R \setminus R^\times$, r is irreducible iff $\langle r\rangle$ is maximal among proper ideals.
(ii) Every irreducible $r \in R$ is prime.
(iii) For $a,b \in R$, $\gcd(a,b) = 1$ iff $\exists x,y \in R$ such that $xa + yb = 1$.

Proof.

(i) If r is irreducible and $\langle r\rangle \subseteq \mathcal{I} \subseteq R$, write $\mathcal{I} = \langle a\rangle$. Then $r = xa$, so irreducibility gives $\mathcal{I} = \langle r\rangle$ or $\mathcal{I} = R$. Conversely, if $r = ab$ with $a,b \notin R^\times$, then $\langle r\rangle \subsetneq \langle a\rangle \subsetneq R$.
(ii) Let r be irreducible and $r \mid ab$. By (i), $\langle r,a\rangle = \langle r\rangle$ or R. In the first case $r \mid a$. In the second, $1 = xr + ya$, so $b = xrb + yab$ and $r \mid b$.
(iii) If $xa + yb = 1$, every common divisor of a,b is a unit. Conversely, let $\langle a,b\rangle = \langle d\rangle$. If $\gcd(a,b) = 1$, then $d \in R^\times$, so $\langle a,b\rangle = R$ and $1 = xa + yb$ for some $x,y \in R$.

Structural Commentary: In a PID, factorization statements become ideal-containment statements.
Key Insight: Maximal principal ideals are the bridge from irreducible to prime. □

Theorem 3.2.16. *Every PID Is a UFD*
If R is a PID, then R is a UFD.

Proof. Let $0 \neq r \in R \setminus R^\times$. If factorization into nonunits never terminates, obtain $s_1, s_2, \ldots$ with $s_{j+1} \mid s_j$ properly. By Proposition 3.2.9,

$$\langle s_1\rangle \subsetneq \langle s_2\rangle \subsetneq \cdots,$$

contradicting Proposition 3.2.14; hence irreducible factorizations exist.
For uniqueness, let $p_1 \cdots p_m = q_1 \cdots q_n$ with all p_i, q_j irreducible. By Theorem 3.2.15, p_1 is prime, so $p_1 \mid q_j$ for some j; irreducibility gives $p_1 \sim q_j$. Cancel associates and iterate. Thus $m = n$ and the factors agree up to order and associates.
Structural Commentary: ACC gives existence; irreducible-is-prime gives uniqueness.
Key Insight: PIDs have enough ideal control to force unique factorization. □

Definition 3.2.17. *Roots of Unity and Cyclotomic Polynomials*
For $n \geq 1$:

(i) The complex nth roots of unity are $U_n := \{z \in \mathbb{C} : z^n = 1\}$.
(ii) A root $\omega \in U_n$ is *primitive* iff $\operatorname{ord}(\omega) = n$; set $\Omega_n := \{\omega \in U_n : \operatorname{ord}(\omega) = n\}$.
(iii) The nth *cyclotomic polynomial* is $\Phi_n(x) := \prod_{\omega\in\Omega_n} (x - \omega)$.
(iv) The number of primitive nth roots is $\varphi(n) := |\Omega_n|$.

Proposition 3.2.18. *Roots of Unity Factorization*
Let $\zeta_n = e^{2\pi i/n}$. Then $U_n = \langle \zeta_n\rangle$; $|U_n| = n$; $\Omega_n = \{\zeta_n^k : 1 \leq k \leq n,\ \gcd(k,n) = 1\}$; $|\Omega_n| = \varphi(n)$.
Moreover, $U_n = \coprod_{d\mid n} \Omega_d$ and $x^n - 1 = \prod_{d\mid n} \Phi_d(x)$.

Proof. The roots of $x^n - 1$ are precisely ζ_n^k for $0 \leq k < n$, hence $U_n = \langle \zeta_n\rangle$ and $|U_n| = n$.
In a cyclic group, $\operatorname{ord}(\zeta_n^k) = n/\gcd(k,n)$, so ζ_n^k is primitive iff $\gcd(k,n) = 1$; thus $|\Omega_n| = \varphi(n)$.
Each $z \in U_n$ has unique order $d \mid n$, giving $U_n = \coprod_{d\mid n} \Omega_d$.
Therefore $x^n - 1 = \prod_{z\in U_n} (x - z) = \prod_{d\mid n} \prod_{\omega\in\Omega_d} (x - \omega) = \prod_{d\mid n} \Phi_d(x)$ *Structural Commentary*: Orders partition roots of unity by divisors.
Key Insight: Cyclicity turns root counting into order counting. □

Proposition 3.2.19. *Cyclotomic Polynomials Have Integer Coefficients*
For every $n \geq 1$, $\Phi_n(x) \in \mathbb{Z}[x]$ is monic.

Proof. Induct on n. For $n = 1$, $\Phi_1(x) = x - 1$.
Assume $\Phi_d \in \mathbb{Z}[x]$ for every proper divisor $d \mid n$, and set $P_n(x) := \prod_{\substack{d \mid n \\ d < n}} \Phi_d(x) \in \mathbb{Z}[x]$.
Then P_n is monic, and Proposition 3.2.18 gives $x^n - 1 = \Phi_n(x)P_n(x)$.
Divide $x^n - 1$ by P_n in $\mathbb{Z}[x]$: $x^n - 1 = hP_n + r$, where $h, r \in \mathbb{Z}[x]$ and $r = 0$ or $\deg r < \deg P_n$.
Uniqueness of division over $\mathbb{C}[x]$ gives $h = \Phi_n$ and $r = 0$. Hence $\Phi_n \in \mathbb{Z}[x]$, and Φ_n is monic by definition.
Structural Commentary: Proper divisors isolate the next cyclotomic factor.
Key Insight: Monic division keeps the quotient in $\mathbb{Z}[x]$. □

Lemma 3.2.20. *A Cyclotomic Size Estimate*
For $n > 1$ and $z \in \mathbb{Z}_{\geq 2}$, $|\Phi_n(z)| > z - 1$.

Proof. For $\omega \in \Omega_n$, $\omega \neq 1$ and $|\omega| = 1$, hence $|z - \omega| > z - 1$. Thus $|\Phi_n(z)| = \prod_{\omega \in \Omega_n} |z - \omega| > (z-1)^{\varphi(n)} \geq z - 1$.
Structural Commentary: Primitive factors all lie outside the closest point 1.
Key Insight: Each primitive root contributes a factor larger than $z - 1$. □

3.3 Modules

Throughout this subsection, R denotes a nonzero commutative ring with identity 1.
Modules are vector spaces with the scalar field replaced by a ring; the price is that nonzero scalars may fail to be invertible and may even kill nonzero vectors.

Perspective 3.3.1. *Polynomial Scalars from One Operator*
Let $\mathcal{V}$ be a vector space over $\mathbb{F}$ and let $T \in \mathcal{L}(\mathcal{V})$. Every polynomial $p(z) \in \mathbb{F}[z]$ acts on $\mathcal{V}$ by $p \cdot \mathbf{v} := p(T)\mathbf{v}$.
Thus one operator turns $\mathcal{V}$ into an object with scalar ring $\mathbb{F}[z]$.

Definition 3.3.2. *R-Module*
An *R-module* is an abelian group $(M, +)$ equipped with a scalar action $R \times M \to M$, $(r, u) \mapsto ru$, such that for all $r, s \in R$ and $u, v \in M$,

(i) $r(u + v) = ru + rv$.
(ii) $(r + s)u = ru + su$.
(iii) $(rs)u = r(su)$.
(iv) $1u = u$.

The ring R is the *base ring* of M.

Example 3.3.3. *Basic Module Models*

(i) R^n is an R-module under componentwise addition and scalar multiplication.
(ii) $\mathrm{Mat}(m, n, R)$ is an R-module under entrywise operations.
(iii) R is an R-module over itself by ring multiplication.
(iv) Every vector space over a field $\mathbb{F}$ is an $\mathbb{F}$-module.

Structural Commentary: Modules are abelian groups with compatible scalar actions.
Key Insight: The new behavior comes from the scalar ring, not from changing addition.

Perspective 3.3.4. *The Base Ring Controls the Geometry*
The structure of an R-module is governed by the arithmetic of R. When R is a field, every nonzero scalar is invertible, giving the usual theory of bases, dimension, and complements. For a general ring, nonunits and zero divisors obstruct division and cancellation, so these notions need not retain their vector-space properties.

Definition 3.3.5. *Submodule*
A *submodule* of an R-module M is a nonempty subset $N \subseteq M$ that is an R-module under the operations inherited from M, written $N \leq M$.

Theorem 3.3.6. *Submodule Test*
Let M be an R-module and let $\emptyset \neq N \subseteq M$. Then $N \leq M \iff \forall r, s \in R$, $\forall u, v \in N$, $ru + sv \in N$.

Proof.
($\Rightarrow$) If $N \leq M$, then $ru, sv \in N$, hence $ru + sv \in N$.

($\Leftarrow$) Assume the displayed condition. Choose $n \in N$. Then $0 = 0n + 0n \in N$.
For $x, y \in N$ and $r \in R$, $x + y = 1x + 1y \in N$, $-x = (-1)x + 0x \in N$, $rx = rx + 0x \in N$.
Thus N is an additive subgroup of M closed under R-scalar multiplication, so the remaining module axioms are inherited from M and $N \leq M$.

Structural Commentary: One closure test packages addition, additive inverses, and scalar multiplication.
Key Insight: Submodules are exactly the subsets closed under all two-term R-linear combinations. □

Proposition 3.3.7. *Intersections and Sums of Submodules*
Let M be an R-module.
If $A, B \leq M$, then $A \cap B \leq M$; $A + B := \{a + b : a \in A,\ b \in B\} \leq M$, and $A + B = \bigcap_{\substack{N \leq M \\ A \cup B \subseteq N}} N$.

Proof. Since $0 \in A \cap B$ and $rx + sy \in A \cap B$ for all $x, y \in A \cap B$ and $r, s \in R$, the submodule test gives $A \cap B \leq M$.
If $x = a + b$ and $y = a' + b'$ with $a, a' \in A$ and $b, b' \in B$, then $rx + sy = (ra + sa') + (rb + sb') \in A + B$.
Thus $A + B \leq M$. Moreover, $A \cup B \subseteq A + B$, and if $N \leq M$ contains $A \cup B$, then $a + b \in N$ for all $a \in A$, $b \in B$; hence $A + B \subseteq N$. This proves the displayed intersection formula.
Structural Commentary: $A \cap B$ is the meet; $A + B$ is the join.
Key Insight: Sums are forced by closure under addition. □

Definition 3.3.8. *Ideal*
An *ideal* of R is an additive subgroup $\mathcal{I} \leq (R, +)$ such that $R\mathcal{I} := \{ri : r \in R,\ i \in \mathcal{I}\} \subseteq \mathcal{I}$, written $\mathcal{I} \trianglelefteq R$.

Proposition 3.3.9. *Ideals Are Submodules of the Regular Module*
View R as the regular R-module ${}_RR$. For every $N \subseteq R$,

$$N \leq {}_RR \iff N \trianglelefteq R.$$

Proof. For $N \subseteq R$,

$$N \leq {}_RR \iff N \leq (R, +) \text{ and } RN \subseteq N \iff N \trianglelefteq R,$$

by Definition 3.3.8.
Structural Commentary: Ideals are the module-theoretic subobjects of the ring acting on itself.
Key Insight: Ring multiplication becomes scalar multiplication. □

Definition 3.3.10. *Generated Submodule*

(i) For $S \subseteq M$, *submodule generated by S* is $\langle\langle S \rangle\rangle_R := \left\{ \sum_{j=1}^{n} r_j s_j : n \geq 1,\ r_j \in R,\ s_j \in S \right\}$; $\langle\langle \emptyset \rangle\rangle_R = \{0\}$.
(ii) A subset $S \subseteq M$ *spans* or *generates* M if $\langle\langle S \rangle\rangle_R = M$.

Proposition 3.3.11. *Generated Submodule Is Smallest*
For every $S \subseteq M$, $\langle\langle S \rangle\rangle_R = \bigcap_{\substack{N \leq M \\ S \subseteq N}} N$.

Proof. Let $G := \langle\langle S \rangle\rangle_R$. By construction, $S \subseteq G$ and the submodule test gives $G \leq M$.
If $N \leq M$ and $S \subseteq N$, then $\sum_{j=1}^{n} r_j s_j \in N$, where $r_j \in R,\ s_j \in S$, so $G \subseteq N$.
Structural Commentary: Generation means closing a set under the operations the module axioms require.
Key Insight: The generated submodule adds no elements except those forced by finite R-linear combination. □

Definition 3.3.12. *Cyclic and Finitely Generated Modules*

(i) For $v \in M$, the *cyclic submodule generated by v* is $\langle\langle v \rangle\rangle_R = Rv = \{rv : r \in R\}$.
(ii) M is *finitely generated* if $M = \langle\langle v_1, \ldots, v_n \rangle\rangle_R$ for some $v_1, \ldots, v_n \in M$.
(iii) M is *cyclic* if $M = Rv$ for some $v \in M$.

Example 3.3.13. *A Finitely Generated Module with a Non-Finitely Generated Submodule*
Let $R = \mathbb{F}[x_1, x_2, \ldots]$. Then $R = \langle\langle 1 \rangle\rangle_R$ as an R-module.
Set $I := (x_1, x_2, \ldots) = \{f \in R : f(0, 0, \ldots) = 0\}$. Then $I \trianglelefteq R$, hence $I \leq R$, but I is not finitely generated.
If $I = \langle\langle p_1, \ldots, p_n \rangle\rangle_R$, choose k such that x_k occurs in no p_i. Let $\varphi_k : R \to \mathbb{F}[x_k]$ send $x_k \mapsto x_k$ and $x_j \mapsto 0$ for $j \neq k$. Since $\varphi_k(p_i) = 0$ for all i, every element of $\langle\langle p_1, \ldots, p_n \rangle\rangle_R$ maps to 0, but $\varphi_k(x_k) = x_k \neq 0$, contradiction.

Definition 3.3.14. *Linear Independence in a Module*

(i) Subset $S \subseteq M$ is *linearly independent over R* if,
$\forall$ distinct $v_1, \ldots, v_n \in S$, $r_1 v_1 + \cdots + r_n v_n = 0 \implies r_1 = \cdots = r_n = 0$.
(ii) A subset $S \subseteq M$ is *linearly dependent* if it is not linearly independent.

Example 3.3.15. *Dependence Without Solving for One Vector*
In the $\mathbb{Z}$-module $\mathbb{Z}$, $3 \cdot 2 - 2 \cdot 3 = 0$. Thus $\{2, 3\}$ is linearly dependent.
However, $2 \notin 3\mathbb{Z}$; $3 \notin 2\mathbb{Z}$, so dependence need not express one element as a scalar multiple of another.

Proposition 3.3.16. *Scaling Independent Sets over a Domain*
Assume R is an integral domain. If $\{v_1, \dots, v_n\}$ is R-linearly independent and $0 \neq r \in R$, then $\{rv_1, \dots, rv_n\}$ is R-linearly independent.

Proof. Let $a_1, \dots, a_n \in R$ and suppose $a_1(rv_1) + \cdots + a_n(rv_n) = 0$. Then $(ra_1)v_1 + \cdots + (ra_n)v_n = 0$.
Independence gives $ra_i = 0$ for all i. Since R is a domain and $r \neq 0$, $a_i = 0$ for all i.
Structural Commentary: Domains preserve independence under nonzero scaling.
Key Insight: The failure of this statement is exactly zero-divisor behavior. □

Definition 3.3.17. *Torsion*

(i) A nonzero element $v \in M$ is a *torsion element* if $\exists\, 0 \neq r \in R$ such that $rv = 0$.
(ii) M is *torsion-free* if it has no torsion elements.
(iii) M is a *torsion module* if every nonzero element of M is torsion.
(iv) The *torsion set* of M is $M_{\text{tor}} := \{0\} \cup \{v \in M : v \text{ is torsion}\}$

Example 3.3.18. *$\mathbb{Z}/n\mathbb{Z}$ as a Torsion Module*
For $n > 1$, $\mathbb{Z}/n\mathbb{Z}$ is a $\mathbb{Z}$-module via $k \cdot \overline{a} = \overline{ka}$. For every $\overline{a} \in \mathbb{Z}/n\mathbb{Z}$, $n\overline{a} = \overline{0}$.
Hence $\mathbb{Z}/n\mathbb{Z}$ is a torsion module, and $\{\overline{a}\}$ is linearly dependent whenever $\overline{a} \neq \overline{0}$.

Proposition 3.3.19. *Torsion Submodule over an Integral Domain*
If R is an integral domain, then $M_{\text{tor}} \leq M$, $(M/M_{\text{tor}})_{\text{tor}} = \{0\}$.

Proof. Let $x, y \in M_{\text{tor}}$ and $r, s \in R$.
Choose $a, b \in R \setminus \{0\}$ with $ax = 0$ and $by = 0$, taking $a = 1$ if $x = 0$ and $b = 1$ if $y = 0$.
Since R is a domain, $ab \neq 0$, and $ab(rx + sy) = rb(ax) + sa(by) = 0$. Thus $rx + sy \in M_{\text{tor}}$, so $M_{\text{tor}} \leq M$.
Now let $0 \neq r \in R$ and suppose $r(m + M_{\text{tor}}) = M_{\text{tor}}$. Then $rm \in M_{\text{tor}}$, so $srm = 0$ for some $0 \neq s \in R$; if $rm = 0$, take $s = 1$. Since $sr \neq 0$, $m \in M_{\text{tor}}$, hence $m + M_{\text{tor}} = M_{\text{tor}}$.
Structural Commentary: Domains let annihilators be multiplied without becoming zero.
Key Insight: Domains make torsion closed and quotients torsion-free. □

Definition 3.3.20. *Annihilators*

(i) For $v \in M$, the *annihilator* of v is $\text{Ann}_R(v) := \{r \in R : rv = 0\}$.
(ii) For $N \leq M$, the *annihilator* of N is $\text{Ann}_R(N) := \{r \in R : rN = \{0\}\}$.

Proposition 3.3.21. *Annihilator Properties*
For $v \in M$ and $N \leq M$:

(i) $\text{Ann}_R(v) \trianglelefteq R$ *and* $\text{Ann}_R(N) \trianglelefteq R$.
(ii) v *is torsion* $\iff v \neq 0$ *and* $\text{Ann}_R(v) \neq \{0\}$.
(iii) $A \leq B \leq M \implies \text{Ann}_R(B) \subseteq \text{Ann}_R(A)$.

Proof.

(i) Since $0 \in \text{Ann}_R(v)$ and, for $a, b \in \text{Ann}_R(v)$ and $r, s \in R$, $(ra + sb)v = r(av) + s(bv) = 0$, $\text{Ann}_R(v) \trianglelefteq R$. The same computation for every $n \in N$ gives $\text{Ann}_R(N) \trianglelefteq R$.
(ii) ($\Rightarrow$) If v is torsion, then $v \neq 0$ and $\exists\, 0 \neq r \in R$ with $rv = 0$, so $\text{Ann}_R(v) \neq \{0\}$.
(ii) ($\Leftarrow$) If $v \neq 0$ and $\text{Ann}_R(v) \neq \{0\}$, choose $0 \neq r \in \text{Ann}_R(v)$. Then $rv = 0$, so v is torsion.
(iii) If $A \leq B \leq M$ and $r \in \text{Ann}_R(B)$, then $rA = \{0\}$; hence $r \in \text{Ann}_R(A)$.

Structural Commentary: Annihilators reverse inclusion.
Key Insight: Killing more vectors gives fewer scalars. □

Proposition 3.3.22. *Torsion Generators Give a Nonzero Annihilator over a Domain*
Assume R is an integral domain. If $M = \langle\langle u_1, \dots, u_n \rangle\rangle_R$, $u_i \in M_{\text{tor}}$ $\forall i$, then $\text{Ann}_R(M) \neq \{0\}$.

Proof. Choose $0 \neq a_i \in R$ with $a_i u_i = 0$, and set $a := \prod_{i=1}^{n} a_i$. Since R is a domain, $a \neq 0$.
For each i, $au_i = 0$, hence $aM = \{0\}$ because the u_i generate M. Thus $0 \neq a \in \text{Ann}_R(M)$.
Structural Commentary: Finite generation lets many annihilators be compressed into one product.
Key Insight: A finite torsion generating set gives a global order for the module. □

Definition 3.3.23. *Free Module and Basis*

(i) A subset $\mathcal{B} \subseteq M$ is a *basis* of M if $\mathcal{B}$ is R-linearly independent and $\langle\langle \mathcal{B} \rangle\rangle_R = M$.
(ii) An R-module M is *free* if $M = \{0\}$ or M has a basis.
(iii) If $\mathcal{B}$ is a basis of M, then M is *free on* $\mathcal{B}$.

Theorem 3.3.24. *Basis Means Unique Finite Coordinates*
A subset $\mathcal{B} \subseteq M$ is a basis if and only if $\forall v \in M,\ \exists!\, (F, (r_b)_{b\in F})$, with $F \subseteq_{\text{fin}} \mathcal{B}$, $r_b \in R \setminus \{0\}$, $v = \sum_{b\in F} r_b b$.

The case $F = \emptyset$ gives the empty sum 0.

Proof.

(⇒) Spanning gives a finite expression; collecting equal basis elements and deleting zero coefficients gives the displayed form. If $\sum_{b\in F} r_b b = \sum_{b\in G} s_b b$, then $\sum_{b\in F\cup G} (r_b - s_b)b = 0$, where missing coefficients are 0. Independence gives $r_b = s_b$ for all b, hence $F = G$ and the coefficients agree.
(⇐) Existence gives $\langle\langle \mathcal{B} \rangle\rangle_R = M$. If $\sum_{b\in F} r_b b = 0$ and $F_0 := \{b \in F : r_b \neq 0\}$ is nonempty, then $0 = \sum_{b\in F_0} r_b b$ contradicts uniqueness of the empty expression. Thus all $r_b = 0$, so $\mathcal{B}$ is independent.

Structural Commentary: Bases give unique finite coordinates.
Key Insight: Existence is spanning; uniqueness is independence. □

Proposition 3.3.25. *Bases Are Minimal Spanning and Maximal Independent*
If $\mathcal{B}$ is a basis of M, then $\forall b \in \mathcal{B}$, $\langle\langle \mathcal{B} \setminus \{b\} \rangle\rangle_R \neq M$, and $\forall v \in M \setminus \mathcal{B}$, $\mathcal{B} \cup \{v\}$ is R-linearly dependent.

Proof. If $b \in \mathcal{B}$ and $b \in \langle\langle \mathcal{B} \setminus \{b\} \rangle\rangle_R$, then b has two finite coordinate expressions, contradicting Theorem 3.3.24. Hence deleting b destroys spanning. If $v \notin \mathcal{B}$, write $v = \sum_{b\in F} r_b b$ with $F \subseteq_{\text{fin}} \mathcal{B}$.
Then $v - \sum_{b\in F} r_b b = 0$ is a nontrivial relation in $\mathcal{B} \cup \{v\}$, so $\mathcal{B} \cup \{v\}$ is dependent.
Structural Commentary: Basis still implies the familiar minimal/maximal properties.
Key Insight: The converses are vector-space privileges, not general module facts. □

Example 3.3.26. *A Free Module with a Nonfree Submodule*
Let $R = \mathbb{Z} \times \mathbb{Z}$ with componentwise multiplication. As an R-module, $R = R(1,1)$, so R is free on $\{(1,1)\}$.
Set $N := \mathbb{Z} \times \{0\} \leq R$. For every $0 \neq (a,0) \in N$, $(0,1)(a,0) = 0$.
Thus every nonzero element of N is torsion, so N has no nonempty independent subset and is not free.

Perspective 3.3.27. *Dimension as the Field Case of Rank*
An $\mathbb{F}$-vector space is a free $\mathbb{F}$-module. By Theorem 1.1.29, finite bases have invariant cardinality; over general rings, rank invariance requires additional hypotheses.

Theorem 3.3.28. *Maps Out of a Free Module*
Let M be free on $\mathcal{B}$ and let N be an R-module.
For every function $f : \mathcal{B} \to N$, there exists a unique $\widetilde{f} \in \text{Hom}_R(M, N)$ such that $\widetilde{f}|_{\mathcal{B}} = f$.

Proof. For $v = \sum_{b\in F} r_b b$ with $F \subseteq_{\text{fin}} \mathcal{B}$, define $\widetilde{f}(v) := \sum_{b\in F} r_b f(b)$.
Theorem 3.3.24 makes the expression of v unique, so $\widetilde{f}$ is well-defined.
For $x = \sum a_b b$ and $y = \sum c_b b$, combining supports gives $\widetilde{f}(rx+sy) = r\widetilde{f}(x) + s\widetilde{f}(y)$, so $\widetilde{f}$ is R-linear and clearly extends f. If $\phi \in \text{Hom}_R(M,N)$ also extends f, then $\phi\left(\sum_{b\in F} r_b b\right) = \sum_{b\in F} r_b f(b) = \widetilde{f}\left(\sum_{b\in F} r_b b\right)$, hence $\phi = \widetilde{f}$.
Structural Commentary: A free module is controlled entirely by its basis values.
Key Insight: Freeness is the universal permission to choose images of basis elements arbitrarily. □

Definition 3.3.29. *Module Homomorphism*
Let M, N be R-modules and let $\tau : M \to N$.

(i) τ is *R-linear* if, $\forall r, s \in R$ and $u, v \in M$, $\tau(ru+sv) = r\tau(u)+s\tau(v)$, and is then called an *R-homomorphism*.
(ii) The set of all R-homomorphisms $M \to N$ is $\text{Hom}_R(M, N)$.
(iii) An R-homomorphism $M \to M$ is an *endomorphism*.
(iv) An injective, surjective, or bijective R-homomorphism is a *monomorphism*, *epimorphism*, or *isomorphism*.

Proposition 3.3.30. *Hom, Kernel, and Image*
Let M, N be R-modules and $\tau \in \text{Hom}_R(M, N)$.

(i) With $(a\sigma + b\rho)(m) := a\sigma(m) + b\rho(m)$, $\text{Hom}_R(M, N)$ is an R-module.
(ii) $\ker \tau := \tau^{-1}(\{0\}) \leq M$.

(iii) $\operatorname{im}\tau := \tau(M) \leq N$.
(iv) τ *is a monomorphism* $\iff \ker\tau = \{0\}$.

Proof. By Theorem 3.3.6, for $A \subseteq P$ in an R-module P, $A \leq P \iff A \neq \emptyset$ and $\forall a,b \in R, \forall x,y \in A, ax+by \in A$.

(i) For $\sigma,\rho \in \operatorname{Hom}_R(M,N)$ and $a,b,r,s \in R$, $(a\sigma + b\rho)(ru+sv) = r(a\sigma+b\rho)(u) + s(a\sigma+b\rho)(v)$, where $u,v \in M$, using commutativity of R. Since $0 \in \operatorname{Hom}_R(M,N)$, the displayed criterion gives the claim inside the pointwise module of maps $M \to N$.
(ii) Since $0 \in \ker\tau$, and $x,y \in \ker\tau$, $a,b \in R$ imply $\tau(ax+by) = a\tau(x) + b\tau(y) = 0$, so the displayed criterion gives $\ker\tau \leq M$.
(iii) Since $0 = \tau(0) \in \operatorname{im}\tau$, and $y_i = \tau(x_i) \in \operatorname{im}\tau$ $(i = 1,2)$, $a,b \in R$ imply $ay_1 + by_2 = \tau(ax_1 + bx_2) \in \operatorname{im}\tau$, so the displayed criterion gives $\operatorname{im}\tau \leq N$.
(iv) If τ is injective, then $\tau(m) = 0 = \tau(0)$ implies $m = 0$, so $\ker\tau = \{0\}$. Conversely, if $\ker\tau = \{0\}$ and $\tau(m) = \tau(m')$, then $\tau(m-m') = 0$, hence $m = m'$.

Structural Commentary: Kernels and images are the canonical submodules attached to a module map.
Key Insight: Homomorphisms preserve precisely the finite R-linear combinations that define submodules. □

Definition 3.3.31. *Quotient Module*
Let $S \leq M$. Define $u \equiv v \pmod S \iff u - v \in S$.
The *quotient module* is $M/S := \{u + S : u \in M\}$; $u + S := \{u + s : s \in S\}$,
with operations $(u+S)+(v+S) := (u+v)+S$; $r(u+S) := ru + S$.

Proposition 3.3.32. *Quotient Operations Are Well-Defined*
Let $S \leq M$. The operations of Definition 3.3.31 are well-defined and make M/S an R-module. Moreover, $0_{M/S} = S$, $\pi : M \twoheadrightarrow M/S$, $\pi(u) := u + S$, $\ker\pi = S$..

Proof. If $u \equiv u' \pmod S$ and $v \equiv v' \pmod S$, then $(u+v)-(u'+v') = (u-u')+(v-v') \in S$ and $ru - ru' = r(u-u') \in S$. Thus $+$ and R-scaling on cosets are representative-independent. The module axioms descend from M. Finally, $0_{M/S} = 0 + S = S$, π is surjective, and $\ker\pi = \{u \in M : u + S = S\} = \{u \in M : u \in S\} = S$. *Structural Commentary*: Quotients collapse exactly the submodule S to zero.
Key Insight: Two representatives define the same coset precisely when their difference is invisible modulo S. □

Example 3.3.33. *A Quotient of a Free Module Need Not Be Free*
As a $\mathbb{Z}$-module, $\mathbb{Z}$ is free on $\{1\}$. For $n > 1$, the submodule $n\mathbb{Z}$ is cyclic, but the quotient $\mathbb{Z}/n\mathbb{Z}$ is torsion by Example 3.3.18. Hence $\mathbb{Z}/n\mathbb{Z}$ is not free as a $\mathbb{Z}$-module.

Theorem 3.3.34. *Correspondence Theorem for Modules*
Let $S \leq M$ and let $\pi : M \twoheadrightarrow M/S$ be the quotient map. Then $\Phi : \{T \leq M : S \leq T\} \longrightarrow \{U \leq M/S\}$, where $\Phi(T) := T/S$ is an order isomorphism, with inverse $\Psi(U) := \pi^{-1}(U)$.

Proof. If $S \leq T \leq M$, then $\Phi(T) = \pi(T) \leq M/S$.
If $U \leq M/S$, then $\Psi(U) = \pi^{-1}(U) \leq M$ and $S = \ker\pi \leq \Psi(U)$.
Moreover, $\Psi\Phi(T) = \pi^{-1}(T/S) = T$ and $\Phi\Psi(U) = \pi(\pi^{-1}(U)) = U$, where the second equality uses surjectivity of π. Finally, $T_1 \leq T_2 \Longrightarrow T_1/S \leq T_2/S$, and $U_1 \leq U_2 \Longrightarrow \pi^{-1}(U_1) \leq \pi^{-1}(U_2)$. Thus Φ is an order isomorphism.
Structural Commentary: Submodules above S are exactly the submodules visible after S is collapsed.
Key Insight: Pull back along the quotient map to recover the original submodule. □

Theorem 3.3.35. *Module Isomorphism Theorems*
The following canonical maps are R-module isomorphisms.

(i) If $\tau \in \operatorname{Hom}_R(M,N)$, then $\overline{\tau} : M/\ker\tau \xrightarrow{\sim} \operatorname{im}\tau$ and $\overline{\tau}(m + \ker\tau) := \tau(m)$.
(ii) If $S,T \leq M$, then $\theta : S/(S\cap T) \xrightarrow{\sim} (S+T)/T$ and $\theta\big(s + (S\cap T)\big) := s + T$.
(iii) If $S \leq T \leq M$, then $\eta : (M/S)/(T/S) \xrightarrow{\sim} M/T$ and $\eta\big((m+S)+T/S\big) := m + T$.

Proof.

(i) If $m - m' \in \ker\tau$, then $\tau(m) = \tau(m')$, so $\overline{\tau}$ is well-defined and R-linear. Also $\operatorname{im}\overline{\tau} = \operatorname{im}\tau$ and $\ker\overline{\tau} = \{m + \ker\tau : \tau(m) = 0\} = \{0_{M/\ker\tau}\}$.
(ii) Apply (i) to $\varphi \in \operatorname{Hom}_R(S,(S+T)/T)$, $\varphi(s) := s + T$. Then $\operatorname{im}\varphi = (S+T)/T$ and $\ker\varphi = \{s \in S : s \in T\} = S \cap T$.
(iii) Apply (i) to $\psi \in \operatorname{Hom}_R(M/S, M/T)$, $\psi(m+S) := m + T$. Then $\operatorname{im}\psi = M/T$ and $\ker\psi = \{m + S : m \in T\} = T/S$.

Structural Commentary: Each theorem identifies a quotient by the exact ambiguity in a natural map.
Key Insight: Image is domain modulo kernel, repeatedly. □

Definition 3.3.36. *Direct Sums and Direct Summands*

(i) For R-modules $M_1, \dots, M_n$, their *external direct sum* is $\bigoplus_{i=1}^{n} M_i := \{(m_i)_{i=1}^n : m_i \in M_i\}$ with componentwise R-module operations.

(ii) For submodules $\{S_i : i \in I\}$ of M, their *internal direct sum* is denoted $M = \bigoplus_{i\in I} S_i$, which means $M = \sum_{i\in I} S_i$ and $\forall i \in I,\ S_i \cap \sum_{j\neq i} S_j = \{0\}$.

(iii) A submodule $S \leq M$ is a *direct summand* if $\exists\, T \leq M$ such that $M = S \oplus T$. Such a T is a *complement* of S in M.

Proposition 3.3.37. *Direct Sum Characterizations*
For a family $\{S_i : i \in I\}$ of distinct submodules of M, the following conditions are equivalent:

(i) $\sum_{i\in I} S_i$ *is direct, i.e.* $\forall k \in I,\ S_k \cap \sum_{i\neq k} S_i = \{0\}$.

(ii) For every $F \subseteq_{\text{fin}} I$ *and* $x_i \in S_i$ $(i \in F)$, $\sum_{i\in F} x_i = 0 \Longrightarrow x_i = 0\ \forall i \in F$.

(iii) Every $0 \neq x \in \sum_{i\in I} S_i$ *has a unique finite nonzero decomposition,* $\exists!\,(F, (x_i)_{i\in F})$ *where* $F \subseteq_{\text{fin}} I$; $0 \neq x_i \in S_i$; $x = \sum_{i\in F} x_i$.

Proof.

(*i*) $\Rightarrow$ (*ii*). If (i) holds and $\sum_{i\in F} x_i = 0$, then any $x_k \neq 0$ gives $x_k = -\sum_{i\in F\setminus\{k\}} x_i \in S_k \cap \sum_{i\neq k} S_i$, a contradiction.

(*ii*) $\Rightarrow$ (*i*). If $x \in S_k \cap \sum_{i\neq k} S_i$, write $x = \sum_{i\in F} x_i$ with $F \subseteq_{\text{fin}} I \setminus \{k\}$ and $x_i \in S_i$. Then $x - \sum_{i\in F} x_i = 0$, so (ii) gives $x = 0$.

(*ii*) $\Rightarrow$ (*iii*). Existence follows from the definition of $\sum_i S_i$ after deleting zero terms. If $\sum_{i\in F} x_i = \sum_{j\in G} y_j$, $0 \neq x_i \in S_i$, $0 \neq y_j \in S_j$, then, over $H = F \cup G$, missing terms set to 0, $\sum_{h\in H} (x_h - y_h) = 0$. By (ii), $x_h = y_h$ for all $h \in H$, hence $F = G$ and the terms agree.

(*iii*) $\Rightarrow$ (*ii*). If $\sum_{i\in F} x_i = 0$, set $F_0 := \{i \in F : x_i \neq 0\}$. If $F_0 \neq \emptyset$, choose $k \in F_0$. If $F_0 = \{k\}$, then $x_k = 0$, contradiction; otherwise $0 \neq x_k = -\sum_{i\in F_0\setminus\{k\}} x_i$, giving two finite nonzero decompositions of x_k, contradiction. Hence $F_0 = \emptyset$.

Structural Commentary: Directness is the absence of hidden cancellation among different summands.
Key Insight: Unique decomposition is controlled by the zero vector. □

Example 3.3.38. *Most Submodules of $\mathbb{Z}$ Are Not Complemented*
As a $\mathbb{Z}$-module, every submodule of $\mathbb{Z}$ is $n\mathbb{Z}$ for a unique $n \geq 0$. For $m, n > 0$, $m\mathbb{Z} \cap n\mathbb{Z} = \operatorname{lcm}(m,n)\mathbb{Z} \neq \{0\}$. Thus $m\mathbb{Z} \oplus n\mathbb{Z} \neq \mathbb{Z}$ for all $m, n > 0$, so the only complemented submodules of $\mathbb{Z}$ are $\{0\}$ and $\mathbb{Z}$.

Theorem 3.3.39. *Complements and Quotients*
If $M = S \oplus T$, then $\pi|_T : T \xrightarrow{\sim} M/S$, $t \mapsto t + S$, where $\pi : M \twoheadrightarrow M/S$ is the quotient map. Hence any two complements of S are isomorphic.

Proof. $\ker(\pi|_T) = T \cap S = \{0\}$. Also, for $m = s + t$ with $s \in S$, $t \in T$, one has $m + S = t + S$; hence $\operatorname{im}(\pi|_T) = M/S$. Thus $\pi|_T$ is an isomorphism, and all complements are isomorphic to M/S.
Structural Commentary: A complement gives a concrete copy of the quotient inside the original module.
Key Insight: Quotienting by S remembers exactly the complementary coordinate. □

Theorem 3.3.40. *Direct Summands Extend Epimorphisms*
Let $N \leq M$ and let M_1 be an R-module.

(i) If $M = N \oplus H$ and $\sigma : N \twoheadrightarrow M_1$, then $\exists$ a unique $\overline{\sigma} \in \operatorname{Hom}_R(M, M_1)$ such that $\overline{\sigma}(n + h) := \sigma(n)$, $n \in N$, $h \in H$, $\overline{\sigma}|_N = \sigma$, $\overline{\sigma} : M \twoheadrightarrow M_1$, and $\ker \overline{\sigma} = \ker \sigma \oplus H$.

(ii) If $\sigma : N \xrightarrow{\sim} M_1$, then $\mathcal{E}_\sigma := \{\overline{\sigma} : M \twoheadrightarrow M_1 : \overline{\sigma}|_N = \sigma\} \longleftrightarrow \mathcal{C}_N(M) := \{H \leq M : M = N \oplus H\}$ via $\overline{\sigma} \mapsto \ker \overline{\sigma}$ and $H \mapsto \overline{\sigma}_H$, where $\overline{\sigma}_H(n + h) := \sigma(n)$.

Proof.

(i) Direct-sum uniqueness makes $\overline{\sigma}$ well-defined, and componentwise computation gives R-linearity. Moreover, $\operatorname{im} \overline{\sigma} = \operatorname{im} \sigma = M_1$; $\ker \overline{\sigma} = \{n + h : \sigma(n) = 0\} = \ker \sigma \oplus H$. This shows uniqueness.

(ii) If $H \in \mathcal{C}_N(M)$, then (i) gives $\overline{\sigma}_H \in \mathcal{E}_\sigma$ and, since $\ker \sigma = \{0\}$, $\ker \overline{\sigma}_H = H$. Conversely, if $\overline{\sigma} \in \mathcal{E}_\sigma$, then $N \cap \ker \overline{\sigma} = \{0\}$, and for each $m \in M$, $n := \sigma^{-1}(\overline{\sigma}(m)) \in N$, $m - n \in \ker \overline{\sigma}$. Thus $M = N \oplus \ker \overline{\sigma}$, so the two displayed assignments are inverse.

Structural Commentary: A complement supplies the directions on which an extension may vanish.
Key Insight: Extending an isomorphism is equivalent to choosing its kernel complement. □

Definition 3.3.41. *Retract*
Let $N \leq M$.

(i) A *retraction* onto N is a map $\rho \in \mathrm{Hom}_R(M,N)$ such that $\rho|_N = \mathrm{id}_N$.
(ii) The submodule N is a *retract* of M iff $\exists\, \rho \in \mathrm{Hom}_R(M,N)$ with $\rho|_N = \mathrm{id}_N$.

Corollary 3.3.42. *Retracts Are Direct Summands*
For $N \leq M$, N is a retract of $M \iff \exists\, H \leq M$ such that $M = N \oplus H$.

Proof.

($\Rightarrow$) Let $\rho : M \to N$ satisfy $\rho|_N = \mathrm{id}_N$. For every $m \in M$, $m = \rho(m)+(m-\rho(m))$, $\rho(m) \in N$, $m-\rho(m) \in \ker\rho$. Moreover $N \cap \ker\rho = \{0\}$, so $M = N \oplus \ker\rho$.
($\Leftarrow$) If $M = N \oplus H$, define $\rho : M \to N$ by $\rho(n+h) := n$, $n \in N$, $h \in H$. Then $\rho \in \mathrm{Hom}_R(M,N)$ and $\rho|_N = \mathrm{id}_N$, so N is a retract.

Structural Commentary: Retractions are projections onto submodules.
Key Insight: A submodule is split exactly when it can be recovered by an idempotent map. □

Definition 3.3.43. *One-Sided Inverses*
Let $\tau \in \mathrm{Hom}_R(A,B)$.

(i) A *left inverse* of τ is $\lambda \in \mathrm{Hom}_R(B,A)$ such that $\lambda\tau = \mathrm{id}_A$.
(ii) A *right inverse* of τ is $\rho \in \mathrm{Hom}_R(B,A)$ such that $\tau\rho = \mathrm{id}_B$.
(iii) A *two-sided inverse* of τ is $\upsilon \in \mathrm{Hom}_R(B,A)$ such that $\upsilon\tau = \mathrm{id}_A$ and $\tau\upsilon = \mathrm{id}_B$.

Theorem 3.3.44. *One-Sided Inverses and Complements*
Let $\sigma : M \to M_1$ be an R-homomorphism.
For $S \leq N$, set $\mathcal{C}_S(N) := \{H \leq N : N = S \oplus H\}$, $\mathcal{L}_\sigma := \{\lambda \in \mathrm{Hom}_R(M_1,M) : \lambda\sigma = \mathrm{id}_M\}$,
and $\mathcal{R}_\sigma := \{\rho \in \mathrm{Hom}_R(M_1,M) : \sigma\rho = \mathrm{id}_{M_1}\}$.

(i) If σ is injective, then $\mathcal{L}_\sigma \xrightarrow{\sim} \mathcal{C}_{\mathrm{im}\,\sigma}(M_1)$ by $\lambda \mapsto \ker\lambda$.
(ii) If σ is surjective, then $\mathcal{R}_\sigma \xrightarrow{\sim} \mathcal{C}_{\ker\sigma}(M)$ by $\rho \mapsto \mathrm{im}\,\rho$.

Proof.

(i) For $H \in \mathcal{C}_{\mathrm{im}\,\sigma}(M_1)$, define $\lambda_H(\sigma m + h) := m$. Directness and injectivity give well-definedness; $\lambda_H\sigma = \mathrm{id}_M$ and $\ker\lambda_H = H$. Conversely, if $\lambda \in \mathcal{L}_\sigma$, then $\mathrm{im}\,\sigma \cap \ker\lambda = \{0\}$, since $\sigma m \in \ker\lambda \Rightarrow m = 0$, and $y = \sigma\lambda y + (y - \sigma\lambda y)$ with $y - \sigma\lambda y \in \ker\lambda$. Thus $M_1 = \mathrm{im}\,\sigma \oplus \ker\lambda$, and the constructions are inverse.
(ii) For $H \in \mathcal{C}_{\ker\sigma}(M)$, $\sigma|_H : H \to M_1$ is an isomorphism; let $\rho_H = (\sigma|_H)^{-1}$ followed by $H \hookrightarrow M$. Then $\rho_H \in \mathcal{R}_\sigma$ and $\mathrm{im}\,\rho_H = H$. Conversely, if $\rho \in \mathcal{R}_\sigma$, then ρ is injective, $\mathrm{im}\,\rho \cap \ker\sigma = \{0\}$, and $m = (m - \rho\sigma m) + \rho\sigma m$ with $m - \rho\sigma m \in \ker\sigma$, $\rho\sigma m \in \mathrm{im}\,\rho$. Thus $M = \ker\sigma \oplus \mathrm{im}\,\rho$, and the constructions are inverse.

Structural Commentary: One-sided inverses exist exactly when the relevant image or kernel splits.
Key Insight: Injective maps need a complemented image; surjective maps need a complemented kernel. □

Proposition 3.3.45. *Split Epimorphisms Split the Domain*
Let $\sigma : M \twoheadrightarrow N$ be an R-homomorphism and set $\mathcal{R}_\sigma := \{\rho \in \mathrm{Hom}_R(N,M) : \sigma\rho = \mathrm{id}_N\}$. These are equivalent:

(i) $\mathcal{R}_\sigma \neq \emptyset$.
(ii) $\exists\, H \leq M$ such that $M = \ker\sigma \oplus H$.

Moreover, $\rho \in \mathcal{R}_\sigma \Longrightarrow M = \ker\sigma \oplus \mathrm{im}\,\rho$.

Proof.

($\Rightarrow$) Let $\rho \in \mathcal{R}_\sigma$. If $\rho n = 0$, then $n = \sigma\rho n = 0$, so ρ is injective.
If $\rho n \in \ker\sigma$, then $n = \sigma\rho n = 0$, hence $\mathrm{im}\,\rho \cap \ker\sigma = \{0\}$.
For $m \in M$, $m = (m - \rho\sigma m) + \rho\sigma m$, with $m - \rho\sigma m \in \ker\sigma$ and $\rho\sigma m \in \mathrm{im}\,\rho$. Thus $M = \ker\sigma \oplus \mathrm{im}\,\rho$.
($\Leftarrow$) Let $M = \ker\sigma \oplus H$. Then $\ker(\sigma|_H) = H \cap \ker\sigma = \{0\}$, and $\mathrm{im}(\sigma|_H) = N$ since $n = \sigma m$, $m = k + h$, gives $n = \sigma h$. Thus $\sigma|_H : H \xrightarrow{\sim} N$, and $\rho := \iota_H(\sigma|_H)^{-1} \in \mathcal{R}_\sigma$, where $\iota_H : H \hookrightarrow M$.

Structural Commentary: A split epimorphism is a quotient map with complemented kernel.
Key Insight: The right inverse chooses canonical representatives for the cosets of $\ker\sigma$. □

Principle 3.3.46. *Modules Are Not Just Vector Spaces with New Letters*
For a general ring R, the following vector-space implications fail:

(i) $N \leq M \not\Rightarrow \exists\, H \leq M$ with $M = N \oplus H$.

(ii) M finitely generated and $N \leq M \not\Rightarrow N$ finitely generated.
(iii) $M \neq 0 \not\Rightarrow \exists \emptyset \neq S \subseteq M$ linearly independent.
(iv) S minimal spanning or maximal independent $\not\Rightarrow S$ is a basis.
(v) M free and $N \leq M \not\Rightarrow N$ free.
(vi) F free and $S \leq F \not\Rightarrow F/S$ free.

Over a field, these obstructions vanish because $R \setminus \{0\} \subseteq R^\times$.

3.4 Structure Theorems

This subsection begins with the structure facts that make free modules and finite generation usable: rank, splitting maps onto free modules, Noetherian modules, and Hilbert's basis theorem. It then develops the classification of finitely generated modules over a principal ideal domain. The guiding model is the abelian-group theorem over $\mathbb{Z}$ and, later, the polynomial-module classification behind rational canonical form.

Theorem 3.4.1. *Rank Is Well Defined fo r Free Modules*
Let A benonzero commutative ring with identity, M be free A-module, $\mathrm{Bas}_A(M)$ denote set of A-bases of M.

(i) $B, B' \in \mathrm{Bas}_A(M) \Longrightarrow |B| = |B'|$.
(ii) $B \in \mathrm{Bas}_A(M)$, $S \subseteq M$, *and* $\langle\langle S \rangle\rangle_A = M \Longrightarrow |B| \leq |S|$.

Proof.

(i) Choose a maximal ideal $\mathcal{I} \lhd A$ by Theorem 3.1.12.
Put $K := A/\mathcal{I}$, a field by Theorem 3.1.11, $\mathcal{I}M := \langle\langle im : i \in \mathcal{I},\ m \in M \rangle\rangle_A$, and $V := M/\mathcal{I}M$.
Since $\mathcal{I}V = 0$, V is a K-vector space via $(a+\mathcal{I})(m+\mathcal{I}M) := am + \mathcal{I}M$.

(ii) Let $B \in \mathrm{Bas}_A(M)$ and set $\overline{B} := \{\overline{b} := b + \mathcal{I}M : b \in B\}$.
In B-coordinates, $\mathcal{I}M = \{\sum_{b \in F} i_b b : F \subseteq_{\mathrm{fin}} B,\ i_b \in \mathcal{I}\}$; hence $\overline{B}$ spans V.
If $\sum_{j=1}^{n} (a_j + \mathcal{I})\overline{b_j} = 0$ with distinct $b_j \in B$, then $\sum_{j=1}^{n} a_j b_j \in \mathcal{I}M$, so $a_j \in \mathcal{I}\ \forall j$; hence $\overline{B}$ is K-independent.
If $\overline{b} = \overline{b'}$ with $b \neq b'$, then $b - b' \in \mathcal{I}M$, forcing the B-coefficient $1 \in \mathcal{I}$, impossible. Thus $B \xrightarrow{\sim} \overline{B}$.

(iii) Therefore $|B| = |\overline{B}| = \dim_K V$ for every $B \in \mathrm{Bas}_A(M)$, so (i) follows from Theorem 1.1.36.
If $\langle\langle S \rangle\rangle_A = M$, then $\overline{S} := \{s + \mathcal{I}M : s \in S\}$ spans V, so $|B| = \dim_K V \leq |\overline{S}| \leq |S|$.

Structural Commentary: Modulo a maximal ideal, free module becomes a vector space with same basis size.
Key Insight: The field quotient $A/\mathcal{I}$ lets vector-space dimension certify module rank. □

Definition 3.4.2. *Rank of a Free Module*
Let A be a nonzero commutative ring with identity. The *rank* of a free A-module M, denoted $\mathrm{rk}_A(M)$, is the cardinality of any basis of M. The zero module has rank 0.

Example 3.4.3. *Why Commutativity Matters for Rank*
Let $V = \bigoplus_{k \geq 0} \mathbb{F} e_k$ and $E = \mathcal{L}(V)$. As a left E-module, $E = E\,\mathrm{id}_V$, so $\{\mathrm{id}_V\}$ is a basis.

For each $n \geq 1$, define $\beta_s \in E$, $0 \leq s < n$, by $\beta_s(e_{kn+t}) = \begin{cases} e_k, & t = s, \\ 0, & t \neq s, \end{cases}$ where $k \geq 0$, $0 \leq t < n$.

If $\sum_{s=0}^{n-1} \alpha_s \beta_s = 0$, then evaluating at e_{kn+t} gives $\alpha_t(e_k) = 0$ for all k, t, hence $\alpha_t = 0$ for all t.

Conversely, for $h \in E$, define $\alpha_s(e_k) := h(e_{kn+s})$; then $h = \sum_{s=0}^{n-1} \alpha_s \beta_s$.

Thus $\{\beta_0, \ldots, \beta_{n-1}\}$ is an E-basis of E, so E has bases of all positive finite cardinalities.
Structural Commentary: Noncommutative scalar rings can make a free module have incompatible basis sizes.
Key Insight: The rank theorem is a commutative-ring theorem, not a formal consequence of freeness alone.

Theorem 3.4.4. *Finite-Support Coordinate Modules*
Let A be a commutative ring with identity and let B be a set. Put $A^{(B)} := \{f \in A^B : |\mathrm{supp}(f)| < \infty\}$.
For $b \in B$, let $\delta_b \in A^{(B)}$ be given by $\delta_b(b) = 1$ and $\delta_b(x) = 0$ for $x \neq b$.
Then $A^{(B)} = \bigoplus_{b \in B} A\delta_b$, so $\{\delta_b : b \in B\}$ is an A-basis. If $A \neq 0$, then $\mathrm{rk}_A(A^{(B)}) = |B|$.

Proof. For $f, g \in A^{(B)}$ and $a \in A$, $\mathrm{supp}(f+g) \subseteq \mathrm{supp}(f) \cup \mathrm{supp}(g)$ and $\mathrm{supp}(af) \subseteq \mathrm{supp}(f)$, hence $A^{(B)} \leq A^B$.
For $f \in A^{(B)}$, $f = \sum_{b \in \mathrm{supp}(f)} f(b)\delta_b$, so the δ_b span. If $\sum_{j=1}^{m} a_j \delta_{b_j} = 0$ with $b_1, \ldots, b_m$ distinct, then evaluation at b_i gives $a_i = 0$ for every i. Thus the δ_b are independent, and the rank statement follows from the displayed basis.
Structural Commentary: Free modules are finite-support coordinate modules on their bases.
Key Insight: Finite support is the module-theoretic version of finite linear combination. □

Proposition 3.4.5. *Free Modules Are Coordinate Modules*
Let A be a commutative ring with identity and let M be free on a basis B.
Then the coordinate map $\Phi_B : M \to A^{(B)}$ is an isomorphism.
If $A \neq 0$ and M, N are free A-modules, then $M \cong N \iff \mathrm{rk}_A(M) = \mathrm{rk}_A(N)$.

Proof. For $m = \sum_{b \in F} a_b b$, set $\Phi_B(m) = f_m$, with $f_m(b) = a_b$ for $b \in F$ and $f_m(b) = 0$ for $b \notin F$.
Unique coordinates give $\Phi_B : M \xrightarrow{\sim} A^{(B)}$, with inverse $f \mapsto \sum_{b \in \mathrm{supp}(f)} f(b)b$.

($\Rightarrow$) If $\psi : M \xrightarrow{\sim} N$ and C is a basis of M, then $\psi(C)$ is a basis of N, so $\mathrm{rk}_A(M) = |C| = |\psi(C)| = \mathrm{rk}_A(N)$.
($\Leftarrow$) If $\mathrm{rk}_A(M) = \mathrm{rk}_A(N)$, choose bases C, D with $|C| = |D|$.
Any bijection $C \to D$ extends, by Theorem 3.3.28, to $M \xrightarrow{\sim} N$.

Structural Commentary: A basis identifies a free module with a standard coordinate model.
Key Insight: Over a nonzero commutative ring, rank classifies free modules up to isomorphism. □

Theorem 3.4.6. *Independent Sets Are Bounded by Rank over a Domain*
Let A be integral domain with field of quotients Q, M free over A, $S \subseteq M$ A-independent. Then $|S| \leq \mathrm{rk}_A(M)$.

Proof. Choose an A-basis B and identify $M = A^{(B)} \subseteq Q^{(B)}$.
If $\sum_{i=1}^{n} q_i s_i = 0$ with $q_i \in Q$ and distinct $s_i \in S$, choose $d \in A \setminus \{0\}$ with $dq_i \in A$ for all i.
Then $\sum_{i=1}^{n} (dq_i)s_i = 0$ in M, so $dq_i = 0$ by A-independence, hence $q_i = 0$ for all i.
Thus S is Q-independent in $Q^{(B)}$, and $|S| \leq \dim_Q Q^{(B)} = |B| = \mathrm{rk}_A(M)$.
Structural Commentary: The field of quotients lets ordinary dimension bound module independence.
Key Insight: Domains allow denominator clearing without creating false zero relations. □

Theorem 3.4.7. *Free Quotients Split*
Let A be a commutative ring with identity.

(i) If $\sigma : M \twoheadrightarrow F$ and F is free, then $\exists N \leq M$ such that $M = \ker(\sigma) \oplus N$ and $N \cong F$.
(ii) If $S \leq M$ and M/S is free, then $\exists T \leq M$ such that $M = S \oplus T$ and $T \cong M/S$; hence $M \cong S \oplus M/S$.
(iii) If $A \neq 0$ and $M, S, M/S$ are free, then $\mathrm{rk}_A(M) = \mathrm{rk}_A(S) + \mathrm{rk}_A(M/S)$, with cardinal addition; in finite rank, $\mathrm{rk}_A(M/S) = \mathrm{rk}_A(M) - \mathrm{rk}_A(S)$.

Proof.

(i) Let B be a basis of F. For $b \in B$, choose $m_b \in M$ with $\sigma(m_b) = b$. By Theorem 3.3.28, $b \mapsto m_b$ extends to $\tau : F \to M$. Since $\sigma\tau = \mathrm{id}_F$, Theorem 3.3.44 gives $M = \ker(\sigma) \oplus \mathrm{im}(\tau)$ and $\mathrm{im}(\tau) \cong F$.
(ii) Apply (i) to $\pi : M \twoheadrightarrow M/S$. Then $M = S \oplus T$ for some $T \cong M/S$, so $M \cong S \oplus M/S$.
(iii) From (ii), $M = S \oplus T$ with $T \cong M/S$. If B_S, B_T are bases of S, T, then $B_S \sqcup B_T$ is a basis of M, hence $\mathrm{rk}_A(M) = |B_S| + |B_T| = \mathrm{rk}_A(S) + \mathrm{rk}_A(M/S)$. The finite formula follows.

Structural Commentary: Free modules are projective: maps onto them have linear sections.
Key Insight: Choosing representatives for basis vectors of the quotient is enough to split the quotient map. □

Definition 3.4.8. *Noetherian Modules and Rings*
Let A be a commutative ring with identity.

(i) An A-module M satisfies *ACC* if every chain $M_1 \leq M_2 \leq \cdots$ of submodules satisfies $\exists n_0, \forall n \geq n_0, M_n = M_{n_0}$. Such M is *Noetherian*.
(ii) The ring A is *Noetherian* if the regular A-module A is Noetherian, equivalently if every ideal-chain $I_1 \subseteq I_2 \subseteq \cdots$ stabilizes.

Theorem 3.4.9. *Noetherian Means Finitely Generated Submodules*
Let M be an A-module.

(i) M is Noetherian $\iff$ every $N \leq M$ is f.g.
(ii) If A is commutative with identity, then A is Noetherian $\iff$ every $\mathcal{I} \lhd A$ is f.g.

Proof.

(i) ($\Rightarrow$) Let $N \leq M$. If N is not f.g., choose $u_{n+1} \in N \setminus \langle\langle u_1, \ldots, u_n \rangle\rangle_A$ inductively. Then $\langle\langle u_1 \rangle\rangle_A < \langle\langle u_1, u_2 \rangle\rangle_A < \cdots$, contradicting ACC.
(i) ($\Leftarrow$) Let $M_1 \leq M_2 \leq \cdots$ and put $U = \bigcup_{n \geq 1} M_n$. Then $U \leq M$, so $U = \langle\langle u_1, \ldots, u_r \rangle\rangle_A$. Choose k with $u_1, \ldots, u_r \in M_k$. Hence $U = M_k$ and $M_n = M_k$ for all $n \geq k$.
(ii) Apply (i) to the regular A-module A: its submodules are exactly the ideals of A.

Structural Commentary: ACC and finite generation are two faces of the same finiteness condition.
Key Insight: Failure of finite generation builds an infinite increasing chain one new generator at a time. □

Theorem 3.4.10. *Finitely Generated Modules over Noetherian Rings*
Let A be a commutative ring with identity. Then A is Noetherian $\iff$ every f.g. A-module is Noetherian.

Proof.
($\Rightarrow$) Assume A Noetherian. First prove by induction on n that every $S \leq A^n$ is f.g.
The case $n = 1$ is Theorem 3.4.9. For $n > 1$, let $\pi : S \to A$ be last-coordinate projection.
Then $\pi(S) \lhd A$ is f.g. and $\ker(\pi) \leq A^{n-1}$ is f.g. by induction. Lifts of generators of $\pi(S)$, together with generators of $\ker(\pi)$, generate S. Thus A^n is Noetherian. If M is f.g., choose $A^n \twoheadrightarrow M$; for $N \leq M$, its preimage in A^n is f.g., hence N is f.g. Therefore M is Noetherian.
($\Leftarrow$) Since $A = \langle\!\langle 1 \rangle\!\rangle_A$, the regular A-module A is f.g.; by hypothesis it is Noetherian. Hence A is Noetherian.
Structural Commentary: Noetherianity passes from rings to finite free modules to quotients.
Key Insight: Finitely generated modules are quotients of finite free modules. □

Theorem 3.4.11. *Hilbert Basis Theorem*
Let A be commutative with identity. If A is Noetherian, then $A[x]$ is Noetherian.

Proof. It suffices to show every ideal $\mathcal{J} \lhd A[x]$ is f.g. If $\mathcal{J} = 0$, done.
For $g \neq 0$, write $\mathrm{lc}(g)$ and $\mathrm{lt}(g)$. Set $L = \langle \mathrm{lc}(g) : g \in \mathcal{J}, g \neq 0 \rangle \lhd A$. Since A is Noetherian, choose $a_1, \ldots, a_m$ generating L and $f_i \in \mathcal{J}$ with $\mathrm{lc}(f_i) = a_i$. Replacing f_i by powers of x times f_i, assume $\deg f_i = d$ for all i.
For $0 \leq k < d$, set $L_k = \langle \mathrm{lc}(g) : g \in \mathcal{J}, \deg g = k \rangle$ and choose finite $P_k \subseteq \mathcal{J}$ with $\deg p = k$ for $p \in P_k$ and $\{\mathrm{lc}(p) : p \in P_k\}$ generating L_k. Put $G = \{f_1, \ldots, f_m\} \cup \bigcup_{k=0}^{d-1} P_k$ and $\mathcal{A} = \langle G \rangle \lhd A[x]$. Then $\mathcal{A} \subseteq \mathcal{J}$.
For $g \in \mathcal{J} \setminus \mathcal{A}$ of minimal degree, set $r = \deg g$. If $r < d$, write $\mathrm{lc}(g) = \sum_{p \in P_r} c_p \,\mathrm{lc}(p)$ and set $h = \sum_{p \in P_r} c_p p$.
If $r \geq d$, write $\mathrm{lc}(g) = \sum_{i=1}^m c_i a_i$ and set $h = \sum_{i=1}^m c_i x^{r-d} f_i$. In either case $h \in \mathcal{A}$ and $\mathrm{lt}(h) = \mathrm{lt}(g)$.
Then $g - h \in \mathcal{J} \setminus \mathcal{A}$ and $\deg(g - h) < \deg g$, contradiction. Thus $\mathcal{J} = \mathcal{A}$ is f.g.
Structural Commentary: Ideals in $A[x]$ are controlled by finitely many leading coefficients in A.
Key Insight: Match the leading term, subtract, and induct on degree. □

Definition 3.4.12. *Standing PID Notation*
In the PID-structure results, R denotes a principal ideal domain unless otherwise stated.
For $r \in R$, write $\langle r \rangle = Rr$; for an R-module M and $v \in M$, write $\langle\!\langle v \rangle\!\rangle_R = Rv$.

Perspective 3.4.13. *PID Arithmetic Controls Module Structure*
Fields have $F \setminus \{0\} \subseteq F^\times$. PIDs replace invertibility by factorization, gcds, and lcms.
This divisibility data classifies finitely generated modules.

Definition 3.4.14. *Order of a Submodule or Element*
Let M be an R-module and $N \leq M$. Since $\mathrm{Ann}_R(N) = \{r \in R : rN = 0\} = \langle \alpha \rangle$, any such α is an *order* of N; write $o(N) = \alpha$, up to associates. For $v \in M$, set $o(v) := o(\langle\!\langle v \rangle\!\rangle_R)$.

Proposition 3.4.15. *Orders and Direct Sums*
Let $M = A \oplus B$. If $o(A) = \alpha$ and $o(B) = \beta$, then $o(M) \sim \mathrm{lcm}(\alpha, \beta)$.
More generally, $o(\bigoplus_{i=1}^n M_i) \sim \mathrm{lcm}(o(M_1), \ldots, o(M_n))$.

Proof. $r(A \oplus B) = 0 \iff rA = 0$ and $rB = 0$.
Hence $\mathrm{Ann}_R(A \oplus B) = \mathrm{Ann}_R(A) \cap \mathrm{Ann}_R(B) = \langle \alpha \rangle \cap \langle \beta \rangle = \langle \mathrm{lcm}(\alpha, \beta) \rangle$.
Thus $o(A \oplus B) \sim \mathrm{lcm}(\alpha, \beta)$, and induction gives the finite case.
Structural Commentary: Direct sums intersect annihilator conditions.
Key Insight: Lcms appear because one scalar must kill every summand at once. □

Theorem 3.4.16. *Cyclic Modules over a PID*
Let M be an R-module and let $v \in M$.
(i) If $\mathrm{Ann}_R(v) = \langle \alpha \rangle$, then $R/\langle \alpha \rangle \xrightarrow{\sim} \langle\!\langle v \rangle\!\rangle_R$, $r + \langle \alpha \rangle \mapsto rv$.
(ii) If $C = \langle\!\langle v \rangle\!\rangle_R$ and $S \leq C$, then $S = \langle\!\langle sv \rangle\!\rangle_R$ for some $s \in R$.
(iii) If $o(v) = \alpha$ and $\beta \in R$, then $o(\beta v) \sim \alpha / \gcd(\alpha, \beta)$, and $\langle\!\langle \beta v \rangle\!\rangle_R = \langle\!\langle v \rangle\!\rangle_R \iff \gcd(\alpha, \beta) \sim 1$.

Proof.
(i) The map $\phi : R \to \langle\!\langle v \rangle\!\rangle_R$, $\phi(r) = rv$, is onto and $\ker \phi = \mathrm{Ann}_R(v) = \langle \alpha \rangle$. Hence $R/\langle \alpha \rangle \xrightarrow{\sim} \langle\!\langle v \rangle\!\rangle_R$.
(ii) For $S \leq C$, set $I = \{r \in R : rv \in S\} \lhd R$.
Since R is a PID, $I = \langle s \rangle$ for some $s \in R$, and $S = \{rv : r \in I\} = \langle\!\langle sv \rangle\!\rangle_R$.

(iii) Put $d = \gcd(\alpha, \beta)$, $\alpha = d\alpha'$, and $\beta = d\beta'$ with $\gcd(\alpha', \beta') \sim 1$.
For $r \in R$, $r\beta v = 0 \iff \alpha \mid r\beta \iff \alpha' \mid r\beta' \iff \alpha' \mid r$, the last step by Bezout.
Thus $\mathrm{Ann}_R(\beta v) = \langle \alpha' \rangle = \langle \alpha / \gcd(\alpha, \beta) \rangle$. Finally, $\langle\!\langle \beta v \rangle\!\rangle_R = \langle\!\langle v \rangle\!\rangle_R \iff (\exists c \in R)(v = c\beta v)$
$\iff (\exists c \in R)(\alpha \mid 1 - c\beta) \iff \gcd(\alpha, \beta) \sim 1$.

Structural Commentary: A cyclic module is the base ring modulo the first relation satisfied by its generator.
Key Insight: Multiplying a cyclic generator divides its order by the gcd of the multiplier and the old order. □

Theorem 3.4.17. *Composing and Decomposing Cyclic Modules*
Let M be an R-module.

(i) If $o(u_i) = \alpha_i \neq 0$ and $\gcd(\alpha_i, \alpha_j) \sim 1$ for $i \neq j$, then $o(\sum_{i=1}^n u_i) \sim \prod_{i=1}^n \alpha_i$ and $\bigoplus_{i=1}^n \langle\!\langle u_i \rangle\!\rangle_R = \langle\!\langle \sum_{i=1}^n u_i \rangle\!\rangle_R$.

(ii) If $o(v) = \alpha = \prod_{i=1}^n \alpha_i \neq 0$ and $\gcd(\alpha_i, \alpha_j) \sim 1$ for $i \neq j$,
then $\exists v_1, \ldots, v_n \in M$ such that $v = \sum_{i=1}^n v_i$, $o(v_i) \sim \alpha_i$, and $\langle\!\langle v \rangle\!\rangle_R = \bigoplus_{i=1}^n \langle\!\langle v_i \rangle\!\rangle_R$.

Proof.

(i) It suffices to prove $n = 2$ and iterate. Let $o(u) = \alpha$, $o(w) = \beta$, and $\gcd(\alpha, \beta) \sim 1$.
If $x \in \langle\!\langle u \rangle\!\rangle_R \cap \langle\!\langle w \rangle\!\rangle_R$, then $\alpha x = \beta x = 0$, so Bezout gives $x = 0$. If $r(u + w) = 0$, then $ru = -rw$ lies in this intersection; hence $ru = rw = 0$, so $\alpha \mid r$ and $\beta \mid r$, hence $\alpha\beta \mid r$. Since $\alpha\beta(u+w) = 0$, $o(u+w) \sim \alpha\beta$.
Choose $a, b \in R$ with $a\alpha + b\beta = 1$. Then $u = b\beta(u+w)$ and $w = a\alpha(u+w)$, so $\langle\!\langle u \rangle\!\rangle_R \oplus \langle\!\langle w \rangle\!\rangle_R = \langle\!\langle u+w \rangle\!\rangle_R$.

(ii) Put $\beta_i = \alpha/\alpha_i$. Since $\gcd(\beta_1, \ldots, \beta_n) \sim 1$, choose $c_i \in R$ with $\sum_{i=1}^n c_i\beta_i = 1$, and set $v_i = c_i\beta_i v$.
Then $v = \sum_{i=1}^n v_i$. Also $\gcd(c_i, \alpha_i) \sim 1$ and $\gcd(\alpha, \beta_i) \sim \beta_i$, so Theorem 3.4.16 gives $o(v_i) \sim \alpha_i$.
Part (i) gives $\langle\!\langle v \rangle\!\rangle_R = \bigoplus_{i=1}^n \langle\!\langle v_i \rangle\!\rangle_R$.

Structural Commentary: Relatively prime cyclic pieces merge into one cyclic piece.
Key Insight: Bezout identities recover each component from the single sum. □

Perspective 3.4.18. *Noetherian Modules over a PID*
Over a PID, $N \leq R^n$ implies N free with $\mathrm{rk}_R(N) \leq n$; Noetherianity is the f.g. shadow.

Theorem 3.4.19. *Submodules of Finite Free PID Modules*
Let F be free over R, $\mathrm{rk}_R(F) = n < \infty$, and $N \leq F$. Then $\exists m \leq n$ such that $N \cong R^m$.
Consequently, for $\mu_R(M) := \min\{|X| : \langle\!\langle X \rangle\!\rangle_R = M\}$, $\mu_R(M) \leq n$ implies $S \leq M \Longrightarrow \mu_R(S) \leq n$.

Proof. Identify $F \simeq R^n$. For $1 \leq k \leq n$, set $I_k := \{x_k : x = (x_1, \ldots, x_k, 0, \ldots, 0) \in N\} \lhd R$ and write $I_k = \langle r_k \rangle$..
If $r_k \neq 0$, choose $u_k \in N \cap (R^k \times \{0\}^{n-k})$ with $(u_k)_k = r_k$. Let $K = \{k : r_k \neq 0\}$.
If $\sum_{k \in K} c_k u_k = 0$ and $j = \max\{k : c_k \neq 0\}$, then the jth coordinate gives $c_j r_j = 0$, hence $c_j = 0$, contradiction.
Thus $\{u_k : k \in K\}$ is R-independent. Set $\ell(0) = 0$ and $\ell(x) = \max \mathrm{supp}(x)$ for $0 \neq x \in N$.
If $j = \ell(x) > 0$, then $x_j \in I_j = \langle r_j \rangle$, say $x_j = cr_j$, and $\ell(x - cu_j) < j$.
Induction on $\ell(x)$ gives $N = \bigoplus_{k \in K} Ru_k \cong R^{|K|}$, so $\mathrm{rk}_R(N) = |K| \leq n$.
If $\mu_R(M) \leq n$, choose $\pi : R^n \twoheadrightarrow M$. For $S \leq M$, $P := \pi^{-1}(S) \leq R^n$ has $P = \langle\!\langle v_1, \ldots, v_m \rangle\!\rangle_R$ with $m \leq n$.
Hence $S = \langle\!\langle \pi(v_1), \ldots, \pi(v_m) \rangle\!\rangle_R$ and $\mu_R(S) \leq m \leq n$.
Structural Commentary: PIDs let one choose pivot generators in a submodule of R^n.
Key Insight: Pivot ideals I_k yield triangular free generators. □

Corollary 3.4.20. *Finitely Generated PID Modules Are Noetherian*
If $\mu_R(M) < \infty$, then M is Noetherian.

Proof. Let $n = \mu_R(M)$. By Theorem 3.4.19, $N \leq M \Longrightarrow \mu_R(N) \leq n < \infty$.
Thus every submodule of M is f.g.; Theorem 3.4.9 gives M Noetherian.
Structural Commentary: Finite generation of every submodule prevents infinite strict growth.
Key Insight: Generators of the union appear at a finite stage. □

Definition 3.4.21. *Rank of a Finite Free Module*
If F is a finite free R-module and $F \cong R^n$, the integer n is the *rank* of F, denoted $\mathrm{rk}(F)$.

Proposition 3.4.22. *Finite Free Rank Is Well-Defined over a Domain*
Let R be an integral domain and $m, n \geq 0$. Then $R^m \cong_R R^n \Longrightarrow m = n$.

Proof. Let $K = \mathrm{Frac}(R)$ and $\phi : R^m \xrightarrow{\sim} R^n$. Then $K \otimes_R \phi : K \otimes_R R^m \xrightarrow{\sim} K \otimes_R R^n$, i.e. $K^m \xrightarrow{\sim} K^n$.
Hence $m = \dim_K K^m = \dim_K K^n = n$.
Structural Commentary: Passing to the fraction field recovers ordinary dimension.
Key Insight: Rank over a domain is detected after denominators are allowed. □

Proposition 3.4.23. *A Full-Size Spanning Set Is a Basis*
Let $\mathrm{rk}_R(F) = n < \infty$. *If* $F = \langle\langle s_1, \dots, s_n \rangle\rangle_R$, *then* $(s_1, \dots, s_n)$ *is an* R*-basis of* F.

Proof. Choose a basis $(b_i)_{i=1}^n$ and define $\tau \in \mathrm{End}_R(F)$ by $\tau(b_i) = s_i$. Then τ is onto.
Since F is free, $\exists \rho : F \to F$ with $\tau\rho = \mathrm{id}_F$; hence $F = \ker\tau \oplus \mathrm{im}\,\rho$ and $\mathrm{im}\,\rho \cong F$.
By Theorem 3.4.19, $\ker\tau \cong R^r$ for some $r \le n$. Thus $R^n \cong F \cong \ker\tau \oplus F \cong R^{r+n}$, so $r = 0$ by Proposition 3.4.22. Hence $\ker\tau = 0$, τ is an isomorphism, and $(s_i) = \tau(b_i)$ is a basis.
Structural Commentary: Surjective finite-free endomorphisms over a PID are injective.
Key Insight: Rank leaves no room for a split kernel. □

Perspective 3.4.24. *Smith Form as Diagonalized Relations*
For a presentation $R^n \xrightarrow{A} R^m \to M \to 0$, Smith form replaces A by a diagonal relation matrix, decoupling finite relations into cyclic constraints.

Theorem 3.4.25. *Smith Normal Form over a PID*
Let $A \in \mathrm{Mat}(m, n, R)$, *viewed as* $R^n \to R^m$.
$\exists$ *invertible* $P \in \mathrm{Mat}(m, m, R)$ *and* $Q \in \mathrm{Mat}(n, n, R)$ *such that* $PAQ = \mathrm{diag}(d_1, \dots, d_r, 0, \dots, 0)$,
where $d_i \in R \setminus \{0\}$, $d_1 \mid \cdots \mid d_r$, *and* $r = \mathrm{rk}_R(\mathrm{im}\,A)$. *The ideals* $\langle d_i \rangle$ *are unique.*

Proof.

(i) *Existence*: If $A = 0$, done. Set $I(A) = \langle a_{ij} \rangle = \langle d \rangle$. Elementary operations move a Bezout combination equal to d to position $(1,1)$. If a pivot $p \nmid a$, Bezout on (p, a) replaces p by $\gcd(p, a)$, so $\langle p \rangle$ strictly increases; ACC in a PID forces termination. Then p divides every entry, so clear its row and column. Induct on the complementary submatrix. Finally replace each bad adjacent pair (a, b) by $(\gcd(a, b), \mathrm{lcm}(a, b))$ until $d_1 \mid \cdots \mid d_r$.

(ii) *Uniqueness*: Let $\Delta_k(A)$ be the ideal of $k \times k$ minors and $\Delta_0(A) = R$. Matrix equivalence preserves each Δ_k. In Smith form, $\Delta_k(A) = \langle d_1 \cdots d_k \rangle$ for $1 \le k \le r$, and $\Delta_k(A) = 0$ for $k > r$. Since R is a domain, $(\Delta_k(A) : \Delta_{k-1}(A)) = \langle d_k \rangle$, so d_k is determined up to associates.

Structural Commentary: Smith form preserves rank plus divisibility data.
Key Insight: Determinantal ideals encode invariant factors. □

Theorem 3.4.26. *Adapted Bases for Submodules of Free Modules*
Let $N \le F$, $\mathrm{rk}_R(F) = n$, *and* $\mathrm{rk}_R(N) = k$.
Then $\exists (v_i)_{i=1}^n \in \mathrm{Bas}_R(F)$ *and* $r_1, \dots, r_k \in R \setminus \{0\}$ *such that* $r_1 \mid \cdots \mid r_k$ *and* $(r_i v_i)_{i=1}^k \in \mathrm{Bas}_R(N)$.

Proof. Choose bases of N and F, and let A represent $\iota : N \hookrightarrow F$.
Smith form gives $PAQ = \mathrm{diag}(r_1, \dots, r_k, 0, \dots, 0)$ with $r_i \neq 0$ and $r_1 \mid \cdots \mid r_k$, since ι is injective.
The row operations give a basis $(v_i)_{i=1}^n$ of F and the column operations give a basis $(w_i)_{i=1}^k$ of N with $\iota(w_i) = r_i v_i$. Hence $(r_i v_i)_{i=1}^k = (w_i)_{i=1}^k$ is a basis of N.
Structural Commentary: Inclusions diagonalize after choosing bases.
Key Insight: Diagonal entries are the required rescalings. □

Corollary 3.4.27. *Smith Form Presents Finitely Generated PID Modules*
Let $R^n \xrightarrow{A} R^m \to M \to 0$ *be a finite presentation.*
If $A \sim \mathrm{diag}(d_1, \dots, d_r, 0, \dots, 0)$ *in Smith form.* $I = \{i : 1 \le i \le r,\ d_i \notin R^\times\}$, *then* $M \cong R^{m-r} \oplus \bigoplus_{i \in I} R/\langle d_i \rangle$.

Proof. Equivalence $A \sim PAQ$ changes bases in R^m, R^n, so $\mathrm{coker}\,A \cong \mathrm{coker}(PAQ)$.
For the diagonal map, $d_i \in R^\times$ gives 0, $d_i \notin R^\times$ gives $R/\langle d_i \rangle$, and each zero row gives R.
Structural Commentary: Smith form turns presentations into direct sums.
Key Insight: Diagonal relations create cyclic torsion summands, and zero rows create free summands. □

Method 3.4.28. *Reading a Module from Smith Form*
For finite presentation $R^n \xrightarrow{A} R^m \to M \to 0$ over PID, put $A \sim D = \mathrm{diag}(d_1, \dots, d_r, 0, \dots, 0)$, $d_1 \mid \cdots \mid d_r$:

(i) $d_i \in R^\times \implies$ no summand;
(ii) $d_i \notin R^\times \implies R/\langle d_i \rangle$;
(iii) each zero target coordinate contributes one R;
(iv) factor nonunit d_i into prime powers for elementary divisors.

Key Insight: Smith form changes arbitrary relations into independent one-coordinate relations.

Example 3.4.29. *Smith Form for a Finite Abelian Group*
Let $M = \mathbb{Z}^2/\langle\!\langle (6,0),(0,10),(2,4)\rangle\!\rangle_{\mathbb{Z}}$, with relation matrix $A = \begin{bmatrix} 6 & 0 & 2 \\ 0 & 10 & 4 \end{bmatrix}$. Then $\Delta_1(A) = \langle 6, 10, 2, 4\rangle = \langle 2\rangle$ and $\Delta_2(A) = \langle 60, 24, -20\rangle = \langle 4\rangle$. Hence $d_1 \sim 2$, $d_1 d_2 \sim 4$, so $d_2 \sim 2$ and $M \cong \mathbb{Z}/2\mathbb{Z} \oplus \mathbb{Z}/2\mathbb{Z}$.
Key Insight: Determinantal ideals read off cyclic summands.

Example 3.4.30. *Recombining Elementary Divisors*
Let $G = \mathbb{Z}/12\mathbb{Z} \oplus \mathbb{Z}/18\mathbb{Z}$. Then $\mathrm{ElemDiv}(G) = \{2, 2^2, 3, 3^2\}$. Recombining equal-height primary blocks gives $2 \cdot 3 = 6$ and $2^2 \cdot 3^2 = 36$; hence $G \cong \mathbb{Z}/6\mathbb{Z} \oplus \mathbb{Z}/36\mathbb{Z}$, with $6 \mid 36$.
Structural Commentary: Elementary divisors split by primes; invariant factors recombine them.
Key Insight: Prime-wise data equals one divisibility chain.

Theorem 3.4.31. *Finitely Generated Torsion-Free Modules Are Free*
For a f.g. R-module M, M is free $\iff$ M is torsion-free.

Proof.
($\Rightarrow$) If $M \cong R^n$, then $am = 0$ with $a \neq 0$ forces every coordinate of m to be 0, since R is a domain.
($\Leftarrow$) Let M be torsion-free and $G = \{u_1, \ldots, u_n\}$ generate M.
Choose maximal R-independent $S = \{v_1, \ldots, v_k\} \subseteq G$. For $w \in G \setminus S$, dependence of $S \cup \{w\}$ gives $a_w \neq 0$ and $b_{w,i} \in R$ with $a_w w + \sum_i b_{w,i} v_i = 0$. Put $a = \prod_{w \in G\setminus S} a_w$. Then $aM \subseteq \langle\!\langle S\rangle\!\rangle_R$, so aM is free by Theorem 3.4.19. Since M is torsion-free, $M \xrightarrow{\sim} aM$, $m \mapsto am$; hence M is free.
Structural Commentary: Clear denominators into a free submodule.
Key Insight: One scalar embeds M into a free sp an. □

Theorem 3.4.32. *Free Plus Torsion Decomposition*
For f.g. M, $M = M_{\mathrm{tor}} \oplus F$ with F f.g. free. Moreover M_{tor} is intrinsic and $F \cong M/M_{\mathrm{tor}}$.

Proof. By Proposition 3.3.19, $T := M_{\mathrm{tor}} \leq M$ and $Q := M/T$ is f.g. torsion-free.
Thus Q is free by Theorem 3.4.31. Since Q is free, the quotient map $\pi : M \twoheadrightarrow Q$ splits; hence $M = T \oplus \operatorname{im}\sigma$ for any section $\sigma : Q \to M$, and $\operatorname{im}\sigma \cong Q$. If $M = T' \oplus G$ with T' torsion and G free, then $m = t' + g \in M_{\mathrm{tor}}$ forces $g = 0$, so $T' = M_{\mathrm{tor}}$. Every complement is isomorphic to M/M_{tor}.
Structural Commentary: F.g. PID modules split into torsion plus free.
Key Insight: The quotient by torsion is free, and freeness makes the quotient map split. □

Definition 3.4.33. *Primary Module*
Let $p \in R$ be prime. A torsion R-module M is *p-primary* if $o(M) \sim p^e$ for some $e \geq 1$.

Theorem 3.4.34. *Primary Decomposition*
Let M be f.g. torsion with $o(M) = \mu \sim \prod_{i=1}^{n} p_i^{e_i}$, where the p_i are distinct nonassociate primes.

Put $q_i := p_i^{e_i}$, $\mu_i := \mu/q_i$, and $M_i := \mu_i M$. Then $M = \bigoplus_{i=1}^{n} M_i$, $M_i = \{v \in M : q_i v = 0\}$, and $o(M_i) \sim q_i$.

The decomposition is unique up to ordering.

Proof. Since $\gcd(\mu_1, \ldots, \mu_n) \sim 1$, choose $a_i \in R$ with $\sum_i a_i \mu_i = 1$. Hence $x = \sum_i a_i \mu_i x$, so $M = \sum_i M_i$.
Also $q_i M_i = \mu M = 0$. Conversely, if $q_i v = 0$, choose $a, b \in R$ with $a\mu_i + bq_i = 1$; then $v = a\mu_i v \in M_i$.
Thus $M_i = \{v : q_i v = 0\}$. If $\sum_i x_i = 0$ with $x_i \in M_i$, applying μ_i gives $\mu_i x_i = 0$ as $q_j \mid \mu_i$ for $j \neq i$.
Since $q_i x_i = 0$ and $\gcd(q_i, \mu_i) \sim 1$, Bezout gives $x_i = 0$; hence the sum is direct.
Now $o(M_i) \mid q_i$ and $\operatorname{lcm}_i o(M_i) \sim o(M) \sim \prod_i q_i$, so $o(M_i) \sim q_i$.
Finally each summand is the intrinsic set $\{v : q_i v = 0\}$, where q_i is the p_i-part of $o(M)$.
Structural Commentary: Primary decomposition separates a torsion module prime by prime.
Key Insight: Bezout idempotents built from the coprime factors of the order project onto the primary pieces. □

Lemma 3.4.35. *The First p-Torsion Layer*
Let $p \in R$ be prime and $M[p] := \ker(p : M \to M) = \{v \in M : pv = 0\}$.
Then $M[p]$ is an $R/\langle p\rangle$-vector space. If $M = S \oplus T$, then $M[p] = S[p] \oplus T[p]$.

Proof. $\langle p\rangle$ is maximal by Theorem 3.2.15 and Theorem 3.1.11; hence $\kappa := R/\langle p\rangle$ is a field. Define $(r + \langle p\rangle)v := rv$.
If $r - r' \in \langle p\rangle$ and $v \in M[p]$, then $(r - r')v = 0$, so the action is well-defined; the axioms are inherited from M.
If $v = s + t \in M[p]$ under $M = S \oplus T$, then $0 = pv = ps + pt$ gives $ps = pt = 0$ by directness.
Thus $v \in S[p] \oplus T[p]$, and the reverse inclusion is immediate.
Structural Commentary: Killing by p turns scalar multiplication modulo p into field scalar multiplication.
Key Insight: The layer $M[p]$ counts the number of cyclic primary summands. □

Theorem 3.4.36. *Cyclic Decomposition of a Primary Module*
Let M be f.g. p-primary torsion with $o(M) \sim p^e$.
Then $M = \bigoplus_{i=1}^{m} \langle\!\langle v_i \rangle\!\rangle_R$ with $\mathrm{Ann}_R(v_i) = \langle p^{e_i} \rangle$ and $e = e_1 \geq e_2 \geq \cdots \geq e_m \geq 1$.
The annihilator chain $\langle p^{e_1} \rangle \subseteq \cdots \subseteq \langle p^{e_m} \rangle$ is intrinsic.

Proof.

(i) *Existence*: Choose a finite presentation of M.
Corollary 3.4.27 gives $M \cong R^s \oplus \bigoplus_{\ell=1}^{t} R/\langle d_\ell \rangle$, with d_ℓ nonunit. Since M is torsion, $s = 0$.
Since M is p-primary, $d_\ell \sim p^{e_\ell}$ with $e_\ell \geq 1$. Reorder so $e_1 \geq \cdots \geq e_m$.
Proposition 3.4.15 gives $o(M) \sim p^{e_1}$, hence $e_1 = e$.

(ii) *Uniqueness*: Put $\kappa = R/\langle p \rangle$. By Lemma 3.4.35, each summand $R/\langle p^{e_i} \rangle$ contributes one κ-dimension to $M[p]$, so $\dim_\kappa M[p] = m$. Passing from M to pM replaces each $e_i > 1$ by $e_i - 1$ and removes each $e_i = 1$. Induction on e recovers the exponent multiset, hence the annihilator chain.

Structural Commentary: A primary module splits into cyclic strings whose lengths are the exponents e_i.
Key Insight: The dimensions of the layers $M[p], pM[p], p^2M[p], \ldots$ determine the string lengths. □

Proposition 3.4.37. *Layer Counts Recover Primary Exponents*
Let $\kappa = R/\langle p \rangle$ and $\lambda_k(M) := \dim_\kappa(p^{k-1}M/p^kM)$ for $k \geq 1$.
If $M \cong \bigoplus_{i=1}^{m} R/\langle p^{e_i} \rangle$ with $e_1 \geq \cdots \geq e_m \geq 1$, then $\lambda_k(M) = \#\{i : e_i \geq k\}$.
Hence the number of summands of exact exponent k is $\lambda_k(M) - \lambda_{k+1}(M)$.

Proof. For $C = R/\langle p^e \rangle$, $p^{k-1}C/p^kC$ is 0 if $k > e$ and is κ if $k \leq e$, generated by the class of p^{k-1}.
Since direct sums commute with $p^{k-1}M$ and p^kM, dimensions add, giving $\lambda_k(M) = \#\{i : e_i \geq k\}$.
Thus $\lambda_k - \lambda_{k+1}$ counts exactly the summands ending at exponent k.
Structural Commentary: Successive p-layers determine the exponent list.
Key Insight: A cyclic summand of order p^e contributes one dimension to exactly the first e layers. □

Definition 3.4.38. *Elementary Divisors*
For finitely generated torsion M, if its primary cyclic summands have annihilators $\langle p_i^{e_{i,j}} \rangle$,
set $\mathrm{ElemDiv}(M) := \{p_i^{e_{i,j}}\}_{i,j}$ as a multiset, up to associates.

Theorem 3.4.39. *Primary Cyclic Decomposition*
Every f.g. torsion R-module M admits $M = \bigoplus_i \bigoplus_{j=1}^{k_i} \langle\!\langle v_{i,j} \rangle\!\rangle_R$, where $\mathrm{Ann}_R(v_{i,j}) = \langle p_i^{e_{i,j}} \rangle$, p_i are distinct nonassociate primes, and $e_{i,1} \geq \cdots \geq e_{i,k_i} \geq 1$. For f.g. torsion M, N, $M \cong N \iff \mathrm{ElemDiv}(M) = \mathrm{ElemDiv}(N)$.

Proof.

($\Rightarrow$) Theorem 3.4.34 gives $M = \bigoplus_i M_{p_i}$ with each M_{p_i} f.g. p_i-primary. Theorem 3.4.36 gives $M_{p_i} = \bigoplus_{j=1}^{k_i} \langle\!\langle v_{i,j} \rangle\!\rangle_R$ with $\mathrm{Ann}_R(v_{i,j}) = \langle p_i^{e_{i,j}} \rangle$. If $M \cong N$, the same uniqueness statements identify the primary pieces and their exponent chains; hence $\mathrm{ElemDiv}(M) = \mathrm{ElemDiv}(N)$.

($\Leftarrow$) If $\mathrm{ElemDiv}(M) = \mathrm{ElemDiv}(N)$, then both modules are isomorphic to the same direct sum $\bigoplus_{q \in \mathrm{ElemDiv}(M)} R/\langle q \rangle$, since a cyclic summand with annihilator $\langle q \rangle$ is $R/\langle q \rangle$ by Theorem 3.4.16. Thus $M \cong N$.

Structural Commentary: A torsion module is a direct sum of prime-power cyclic blocks.
Key Insight: Elementary divisors are the atoms of the PID module classification. □

Proposition 3.4.40. *Elementary Divisors of a Direct Sum*
For f.g. torsion A, B, $\mathrm{ElemDiv}(A \oplus B) = \mathrm{ElemDiv}(A) \uplus \mathrm{ElemDiv}(B)$, where $\uplus$ is multiset union.

Proof. If $A = \bigoplus_i A_i$ and $B = \bigoplus_j B_j$ are primary cyclic decompositions, then $A \oplus B = (\bigoplus_i A_i) \oplus (\bigoplus_j B_j)$ is one for $A \oplus B$. Thus the prime-power annihilator multiset is the multiset union.
Structural Commentary: Direct sums concatenate cyclic primary blocks.
Key Insight: Elementary divisors count blocks with multiplicity. □

Definition 3.4.41. *Invariant Factors*
For f.g. torsion M, write $\mathrm{ElemDiv}(M) = \{p_i^{e_{i,j}} : 1 \leq j \leq k_i\}_i$ with $e_{i,1} \geq \cdots \geq e_{i,k_i} \geq 1$.
Let $m = \max_i k_i$, set $e_{i,j} = 0$ for $j > k_i$, and define $d_j := \prod_i p_i^{e_{i,j}}$.
The nonunit d_j, up to associates, are the *invariant factors*; they satisfy $d_m \mid \cdots \mid d_1$.

Method 3.4.42. *From Elementary Divisors to Invariant Factors*

(i) Group ElemDiv(M) by prime: $\{p_i^{e_{i,j}}\}_j$.
(ii) Reorder each list so $e_{i,1} \geq e_{i,2} \geq \cdots$ and pad by $e_{i,j} = 0$ to common length.
(iii) Form $d_j = \prod_i p_i^{e_{i,j}}$.
(iv) Discard unit d_j; the rest are the invariant factors from Definition 3.4.41.

Example 3.4.43. *Packing Elementary Divisors into Invariant Factors*
Over $R = \mathbb{F}[z]$, let ElemDiv$(M) = \{(z-1)^3, z-1, (z^2+1)^2, z^2+1\}$. The exponent rows are $(3,1)$ for $z-1$ and $(2,1)$ for z^2+1, so $d_1 = (z-1)^3(z^2+1)^2$ and $d_2 = (z-1)(z^2+1)$, with $d_2 \mid d_1$. Hence $M \cong R/\langle d_1\rangle \oplus R/\langle d_2\rangle$.
Structural Commentary: Elementary divisors separate primes; invariant factors pack rows.
Key Insight: Line up equal-height primary strings and multiply across the row.

Theorem 3.4.44. *Invariant Factor Decomposition*
For f.g. torsion M, $M \cong \bigoplus_{j=1}^m R/\langle d_j\rangle$ with $d_j \notin R^\times$ and $d_m \mid \cdots \mid d_1$.
Writing Inv$(M) := (d_m, \ldots, d_1)$ *up to associates, for f.g. torsion M, N, $M \cong N \iff$* Inv$(M) =$ Inv(N).

Proof. By Theorem 3.4.39, $M \cong \bigoplus_{q \in \text{ElemDiv}(M)} R/\langle q\rangle$.
Let ElemDiv$(M) = \{p_i^{e_{i,j}} : 1 \leq j \leq k_i\}_i$ with $e_{i,1} \geq \cdots \geq e_{i,k_i}$, pad by $e_{i,j} = 0$, and set $d_j = \prod_i p_i^{e_{i,j}}$.
For fixed j, the nonunit $p_i^{e_{i,j}}$ are pairwise coprime, so Theorem 3.4.17 gives $\bigoplus_{e_{i,j}>0} R/\langle p_i^{e_{i,j}}\rangle \cong R/\langle d_j\rangle$.
Hence $M \cong \bigoplus_j R/\langle d_j\rangle$ with $d_m \mid \cdots \mid d_1$.
($\Rightarrow$) If $M \cong N$, then ElemDiv$(M) =$ ElemDiv(N) by Theorem 3.4.39.
Grouping by height gives Inv$(M) =$ Inv(N).
($\Leftarrow$) If Inv$(M) =$ Inv(N), by factorization of each d_j, ElemDiv$(M) =$ ElemDiv(N); QED (Theorem 3.4.39).
Structural Commentary: Invariant factors pack prime-power blocks.
Key Insight: Same data, different packaging. □

Corollary 3.4.45. *Classification of Finitely Generated PID Modules*
For f.g. M, $M \cong R^r \oplus \bigoplus_{j=1}^m R/\langle d_j\rangle$, where $r \geq 0$, $d_j \notin R^\times$, and $d_m \mid \cdots \mid d_1$.
The pair $(r,$ Inv$(M_{\text{tor}}))$*, equivalently $(r,$* ElemDiv$(M_{\text{tor}}))$*, is a complete invariant.*

Proof. Let $T := M_{\text{tor}}$. Theorem 3.4.32 gives $M = T \oplus F$ with $F \cong R^r$.
Theorem 3.4.44 gives $T \cong \bigoplus_{j=1}^m R/\langle d_j\rangle$ with $d_m \mid \cdots \mid d_1$. Since T is intrinsic and $F \cong M/T$, uniqueness and completeness follow from rank uniqueness and the invariant-factor theorem.
Structural Commentary: F.g. PID modules are free plus cyclic torsion.
Key Insight: The free rank records the zero-annihilator part; invariant factors record the torsion part. □

Perspective 3.4.46. *Largest Invariant Factor and Product*
If $M \cong \bigoplus_{j=1}^m R/\langle d_j\rangle$ with $d_m \mid \cdots \mid d_1$, then $\text{Ann}_R(M) = \langle d_1\rangle$ and $\prod_{j=1}^m d_j \sim \prod_{q \in \text{ElemDiv}(M)} q$.
For $R = \mathbb{F}[z]$ and $M = (\mathbb{F}^n)_A$, after monic normalisation, $d_1 = m_A$ and $\prod_j d_j = \chi_A$.

Theorem 3.4.47. *Characterizing Cyclic Torsion Modules*
Let M be f.g. torsion with order $\mu \sim \prod_{i=1}^n p_i^{e_i}$, p_i distinct nonassociate primes. These are equivalent:

(i) M is cyclic.
(ii) $\forall i$, the primary component M_{p_i} is cyclic.
(iii) ElemDiv$(M) = \{p_i^{e_i} : 1 \leq i \leq n\}$.
(iv) $M \cong R/\langle\mu\rangle$.

Proof.

(i) $\Rightarrow$ (ii) If $M = \langle\!\langle x\rangle\!\rangle_R$, Theorem 3.4.17 gives $M = \bigoplus_i \langle\!\langle x_i\rangle\!\rangle_R$ with $o(x_i) \sim p_i^{e_i}$; hence $M_{p_i} = \langle\!\langle x_i\rangle\!\rangle_R$.
(ii) $\Rightarrow$ (iii) By Theorem 3.4.39, cyclic M_{p_i} of order $p_i^{e_i}$ has sole elementary divisor $p_i^{e_i}$; use Proposition 3.4.40.
(iii) $\Rightarrow$ (iv) Method 3.4.42 packs the displayed elementary divisors into one invariant factor $d_1 \sim \prod_i p_i^{e_i} \sim \mu$.
(iv) $\Rightarrow$ (i) Since $R/\langle\mu\rangle = \langle\!\langle 1 + \langle\mu\rangle\rangle\!\rangle_R$, any isomorphism $R/\langle\mu\rangle \xrightarrow{\sim} M$ sends $1 + \langle\mu\rangle$ to a generator of M.

Structural Commentary: Cyclic torsion means one string per prime.
Key Insight: Multiple elementary divisors for the same prime are the obstruction to one generator. □

Definition 3.4.48. *Indecomposable Module*
A nonzero R-module M is *indecomposable* if $M = A \oplus B \Longrightarrow A = 0$ or $B = 0$.

Theorem 3.4.49. *Indecomposable Finitely Generated Torsion Modules*
Let $0 \neq M$ be f.g. torsion. The following are equivalent:

(i) M is indecomposable.

(ii) M *is primary cyclic.*
(iii) $\mathrm{ElemDiv}(M) = \{p^e\}$ *for some prime* p *and* $e \geq 1$.

Proof.

(*i*) $\Rightarrow$ (*ii*) By Theorem 3.4.39, write $M \cong \bigoplus_{\alpha\in A} C_\alpha$, where $0 \neq C_\alpha \cong R/\langle p_\alpha^{e_\alpha}\rangle$.
If $|A| \geq 2$, then $M \cong C_{\alpha_0} \oplus \bigoplus_{\alpha\neq\alpha_0} C_\alpha$ is a nontrivial decomposition, contradiction.
Since $M \neq 0$, $|A| = 1$; hence M is primary cyclic.
(*ii*) $\Rightarrow$ (*iii*) If $M \cong R/\langle p^e\rangle$, then $\mathrm{ElemDiv}(M) = \{p^e\}$.
(*iii*) $\Rightarrow$ (*i*) If $M = A \oplus B$, Proposition 3.4.40 gives $\{p^e\} = \mathrm{ElemDiv}(A) \sqcup \mathrm{ElemDiv}(B)$.
Hence one multiset is empty, so one summand is 0.

Structural Commentary: Indecomposable torsion modules are precisely the single primary cyclic blocks.
Key Insight: Elementary divisors are the indecomposable pieces. □

Proposition 3.4.50. *Prime Divisors Give Prime-Order Cyclic Submodules*
Let M *be f.g. torsion with order* μ. *If* $p \mid \mu$ *is prime, then* $\exists w \in M$ *with* $o(w) \sim p$; *equivalently* $Rw \cong R/\langle p\rangle$.

Proof. Write $\mu = pq$. Since $q \notin \mathrm{Ann}_R(M) = \langle\mu\rangle$, choose $v \in M$ with $w := qv \neq 0$.
Then $pw = \mu v = 0$, so $o(w) \mid p$. Since $w \neq 0$, $o(w)$ is nonunit; as p is prime, $o(w) \sim p$.
Structural Commentary: Every prime in the global order appears in some one-dimensional primary layer.
Key Insight: Divide the order by one prime and apply the remaining scalar to an element it does not kill. □

3.5 Linear Operators as Modules

Throughout this subsection, $\mathcal{V}$ is finite-dimensional over $\mathbb{F}$ unless explicitly stated otherwise.

Perspective 3.5.1. *An Operator as a Polynomial Action*
Fix $T \in \mathcal{L}(\mathcal{V})$. Evaluation at T defines an $\mathbb{F}[z]$-action on $\mathcal{V}$ by $p \cdot \mathbf{v} := p(T)\mathbf{v}$.
Thus the orbit $\mathbf{v}, T\mathbf{v}, T^2\mathbf{v}, \ldots$ is encoded by the cyclic subspace $\{p(T)\mathbf{v} : p \in \mathbb{F}[z]\} = \mathrm{span}\{T^j\mathbf{v} : j \geq 0\}$.

Definition 3.5.2. *The Module Attached to an Operator*
Let $T \in \mathcal{L}(\mathcal{V})$. The *operator module* attached to T is the $\mathbb{F}[z]$-module $\mathcal{V}_T$ with underlying additive group $\mathcal{V}$ and action $p(z) \cdot \mathbf{v} := p(T)\mathbf{v}$. For $A \in \mathrm{Mat}(n, n, \mathbb{F})$, define $(\mathbb{F}^n)_A$ analogously by $p(z) \cdot \mathbf{x} := p(A)\mathbf{x}$.

Proposition 3.5.3. *Operator Modules Are Finitely Generated Torsion Modules*
Let $0 < \dim \mathcal{V} < \infty$ *and* $T \in \mathcal{L}(\mathcal{V})$. *Then* $\mathcal{V}_T$ *is a finitely generated torsion* $\mathbb{F}[z]$*-module.*

Proof. Let $(\mathbf{v}_1, \ldots, \mathbf{v}_n)$ be an $\mathbb{F}$-basis of $\mathcal{V}$. Since $\mathbb{F} \subseteq \mathbb{F}[z]$, this basis generates $\mathcal{V}_T$.
By Proposition 3.2.4, there exists $0 \neq p \in \mathbb{F}[z]$ with $p(T) = 0$. Thus $p \cdot \mathbf{v} = p(T)\mathbf{v} = \mathbf{0}$ for all $\mathbf{v} \in \mathcal{V}$, so $\mathcal{V}_T$ is torsion. Finally, $\mathbb{F}[z]$ is a PID by Theorem 3.2.11.
Structural Commentary: Finite-dimensionality gives a global polynomial relation for the operator.
Key Insight: Rational canonical form is the coordinate form of the structure theorem for $\mathcal{V}_T$. □

Theorem 3.5.4. *Submodules Are Invariant Subspaces*
Let $T \in \mathcal{L}(\mathcal{V})$ *and* $\mathcal{S} \subseteq \mathcal{V}$. *Then* $\mathcal{S} \leq_{\mathbb{F}[z]} \mathcal{V}_T$ *if and only if* $\mathcal{S} \leq_{\mathbb{F}} \mathcal{V}$ *and* $T\mathcal{S} \subseteq \mathcal{S}$.

Proof.

($\Rightarrow$) If $\mathcal{S} \leq_{\mathbb{F}[z]} \mathcal{V}_T$, then $\mathbb{F} \subseteq \mathbb{F}[z]$ gives $\mathcal{S} \leq_{\mathbb{F}} \mathcal{V}$. Since $z \cdot \mathbf{s} = T\mathbf{s}$ for $\mathbf{s} \in \mathcal{S}$, we have $T\mathcal{S} \subseteq \mathcal{S}$.
($\Leftarrow$) If $\mathcal{S} \leq_{\mathbb{F}} \mathcal{V}$ and $T\mathcal{S} \subseteq \mathcal{S}$, then $T^j\mathcal{S} \subseteq \mathcal{S}$ for all $j \geq 0$.
Hence for $p(z) = \sum_j a_j z^j$ and $\mathbf{s} \in \mathcal{S}$, $p \cdot \mathbf{s} = p(T)\mathbf{s} = \sum_j a_j T^j\mathbf{s} \in \mathcal{S}$, so $\mathcal{S} \leq_{\mathbb{F}[z]} \mathcal{V}_T$.

Structural Commentary: The single scalar z remembers the operator.
Key Insight: Invariance under T is exactly closure under all polynomial scalars. □

Theorem 3.5.5. *Similarity Is Module Isomorphism*
Let $T, S \in \mathcal{L}(\mathcal{V})$. *Then* $\mathcal{V}_T \cong_{\mathbb{F}[z]} \mathcal{V}_S \iff T \sim S$. *More precisely, for* $\Phi \in \mathrm{GL}(\mathcal{V})$, $\Phi : \mathcal{V}_T \to \mathcal{V}_S$ *is* $\mathbb{F}[z]$*-linear if and only if* $\Phi T = S\Phi$, *equivalently* $S = \Phi T \Phi^{-1}$.

Proof.

($\Rightarrow$) If $\Phi : \mathcal{V}_T \to \mathcal{V}_S$ is an $\mathbb{F}[z]$-isomorphism, then $\Phi \in \mathrm{GL}(\mathcal{V})$, $\Phi(T\mathbf{v}) = \Phi(z \cdot \mathbf{v}) = z \cdot \Phi(\mathbf{v}) = S\Phi(\mathbf{v})$ $\forall \mathbf{v} \in \mathcal{V}$.
Hence $\Phi T = S\Phi$, so $S = \Phi T \Phi^{-1}$.
($\Leftarrow$) If $\Phi \in \mathrm{GL}(\mathcal{V})$ and $S = \Phi T \Phi^{-1}$, then $\Phi T = S\Phi$, hence $\Phi T^j = S^j\Phi$ for all $j \geq 0$.
Thus for $p(z) = \sum_j a_j z^j$, $\Phi(p(T)\mathbf{v}) = p(S)\Phi(\mathbf{v})$, so $\Phi : \mathcal{V}_T \to \mathcal{V}_S$ is an $\mathbb{F}[z]$-isomorphism.

Structural Commentary: A module isomorphism is precisely an intertwiner for the action of z.
Key Insight: Similarity is isomorphism of the associated $\mathbb{F}[z]$-modules. □

Corollary 3.5.6. *Matrix Similarity Is Matrix-Module Isomorphism*
Let $A, B \in \mathrm{Mat}(n,n,\mathbb{F})$. Then $A \sim B \iff (\mathbb{F}^n)_A \cong_{\mathbb{F}[z]} (\mathbb{F}^n)_B$. More precisely, for $X \in \mathrm{GL}_n(\mathbb{F})$, $X : (\mathbb{F}^n)_A \to (\mathbb{F}^n)_B$ is $\mathbb{F}[z]$-linear if and only if $XA = BX$, equivalently $B = XAX^{-1}$.

Proof.
($\Rightarrow$) If $X : (\mathbb{F}^n)_A \to (\mathbb{F}^n)_B$ is an $\mathbb{F}[z]$-isomorphism, then $X \in \mathrm{GL}_n(\mathbb{F})$ and $X(A\mathbf{x}) = X(z \cdot \mathbf{x}) = z \cdot X\mathbf{x} = BX\mathbf{x}$ for all $\mathbf{x} \in \mathbb{F}^n$. Hence $XA = BX$, so $B = XAX^{-1}$.
($\Leftarrow$) If $X \in \mathrm{GL}_n(\mathbb{F})$ and $B = XAX^{-1}$, then $XA = BX$, hence $XA^j = B^jX$ for all $j \geq 0$. Thus $Xp(A) = p(B)X$ for all $p \in \mathbb{F}[z]$, so $X : (\mathbb{F}^n)_A \to (\mathbb{F}^n)_B$ is an $\mathbb{F}[z]$-isomorphism.

Structural Commentary: Similarity is invertible module intertwining.
Key Insight: The matrix is remembered by the scalar action of z. □

Definition 3.5.7. *Module Annihilator and Minimal Polynomial*
(i) The *annihilator* of $\mathcal{V}_T$ is $\mathrm{Ann}_{\mathbb{F}[z]}(\mathcal{V}_T) := \{p \in \mathbb{F}[z] : p(T)\mathcal{V} = \{\mathbf{0}\}\}$.
(ii) *Minimal polynomial* μ_T is unique monic generator of this nonzero principal ideal, so $\mathrm{Ann}_{\mathbb{F}[z]}(\mathcal{V}_T) = \langle \mu_T \rangle$.
(iii) For $A \in \mathrm{Mat}(n,n,\mathbb{F})$, define $\mu_A := \mu_{T_A}$, where $T_A(\mathbf{x}) = A\mathbf{x}$.

Proposition 3.5.8. *Minimal Polynomial as Module Order*
Let $T \in \mathcal{L}(\mathcal{V})$. Then μ_T is the unique monic polynomial of minimal degree in $\{p \in \mathbb{F}[z] : p(T) = 0\}$. If $S \sim T$, then $\mu_S = \mu_T$; hence similar matrices have the same minimal polynomial.

Proof. By Definition 3.5.7, $\mathrm{Ann}_{\mathbb{F}[z]}(\mathcal{V}_T) = \langle \mu_T \rangle$ and $p \in \mathrm{Ann}_{\mathbb{F}[z]}(\mathcal{V}_T)$ iff $p(T) = 0$. Thus μ_T is the unique monic annihilator of minimal degree. If $S = \Phi T \Phi^{-1}$, then $p(S) = \Phi p(T) \Phi^{-1}$ for all $p \in \mathbb{F}[z]$, so $\mathrm{Ann}_{\mathbb{F}[z]}(\mathcal{V}_S) = \mathrm{Ann}_{\mathbb{F}[z]}(\mathcal{V}_T)$ and $\mu_S = \mu_T$. The matrix case follows by applying this to the associated operators.
Structural Commentary: The minimal polynomial is the order of the whole operator module.
Key Insight: Similarity preserves the module and therefore preserves its annihilator ideal. □

Definition 3.5.9. *Cyclic Subspace and Vector Annihilator*
Let $T \in \mathcal{L}(\mathcal{V})$ and $\mathbf{v} \in \mathcal{V}$.
(i) The *T-cyclic subspace generated by $\mathbf{v}$* is $\mathcal{C}_T(\mathbf{v}) := \mathrm{span}\{T^j\mathbf{v} : j \geq 0\}$
(ii) The vector $\mathbf{v}$ is a *cyclic vector* for T if $\mathcal{C}_T(\mathbf{v}) = \mathcal{V}$.
(iii) The *annihilator of $\mathbf{v}$ for T* is $\mathrm{ann}_T(\mathbf{v}) := \{p \in \mathbb{F}[z] : p(T)\mathbf{v} = \mathbf{0}\}$

Proposition 3.5.10. *Cyclic Subspaces Are Smallest Invariant Hosts*
Let $T \in \mathcal{L}(\mathcal{V})$ and $\mathbf{v} \in \mathcal{V}$. Then $\mathcal{C}_T(\mathbf{v}) = \{p(T)\mathbf{v} : p \in \mathbb{F}[z]\}$.
Moreover, $\mathcal{C}_T(\mathbf{v})$ is the smallest T-invariant subspace containing $\mathbf{v}$, and $\mathrm{ann}_T(\mathbf{v})$ is an ideal of $\mathbb{F}[z]$.

Proof. Equality follows as $p(T)\mathbf{v}$ is a finite linear combination of the orbit vectors $T^j\mathbf{v}$, and every such finite linear combination equals $p(T)\mathbf{v}$ for some $p \in \mathbb{F}[z]$. Also $T(p(T)\mathbf{v}) = (zp)(T)\mathbf{v}$, so $\mathcal{C}_T(\mathbf{v})$ is T-invariant.
If $\mathcal{U}$ is T-invariant and contains $\mathbf{v}$, then $T^j\mathbf{v} \in \mathcal{U}$ for every $j \geq 0$, hence $\mathcal{C}_T(\mathbf{v}) \subseteq \mathcal{U}$. Finally, if $p, q \in \mathrm{ann}_T(\mathbf{v})$ and $r \in \mathbb{F}[z]$, then $(p+q)(T)\mathbf{v} = \mathbf{0}$ and $(rp)(T)\mathbf{v} = r(T)p(T)\mathbf{v} = \mathbf{0}$ by Proposition 3.2.3. $\mathrm{ann}_T(\mathbf{v})$ is an ideal.
Structural Commentary: A cyclic subspace is the span of the orbit generated by repeatedly applying T.
Key Insight: Invariance is exactly what lets the whole orbit stay inside the same host subspace. □

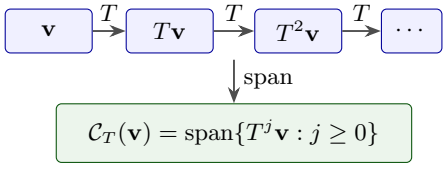


Figure 103: The T-orbit of $\mathbf{v}$ spans $\mathcal{C}_T(\mathbf{v})$.

Definition 3.5.11. *Order of a Vector under an Operator*
Let $\mathcal{V}$ be finite-dimensional, let $T \in \mathcal{L}(\mathcal{V})$, and let $\mathbf{v} \neq \mathbf{0}$. The *order of $\mathbf{v}$ under T* is the unique monic polynomial $m_{\mathbf{v}} \in \mathrm{ann}_T(\mathbf{v})$ of least degree; equivalently, $\mathrm{ann}_T(\mathbf{v}) = (m_{\mathbf{v}})$

Definition 3.5.12. *Companion Matrix*
Let $f(z) = z^n + c_{n-1}z^{n-1} + \cdots + c_1 z + c_0$ be monic. For $n \geq 2$, the *companion matrix* of f is

$$C_f := \begin{bmatrix} 0 & 0 & \cdots & 0 & -c_0 \\ 1 & 0 & \cdots & 0 & -c_1 \\ 0 & 1 & \cdots & 0 & -c_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 1 & -c_{n-1} \end{bmatrix}.$$

For $n = 1$, set $C_f = [-c_0]$.

Theorem 3.5.13. *Companion Matrix Realises a Polynomial*
Let $f(z) = z^n + \sum_{i=0}^{n-1} c_i z^i$ be monic and let C_f be its companion matrix. Then $p_{C_f} = \mu_{C_f} = f$.

Proof. Let $(\mathbf{e}_1, \ldots, \mathbf{e}_n)$ be the standard basis.

For $0 \leq j < n$, $C_f^j \mathbf{e}_1 = \mathbf{e}_{j+1}$, and last column gives $C_f^n \mathbf{e}_1 = -\sum_{i=0}^{n-1} c_i C_f^i \mathbf{e}_1$. Hence $f(C_f)\mathbf{e}_1 = 0$.
Since $\mathbf{e}_j = C_f^{j-1}\mathbf{e}_1$ and $f(C_f)C_f = C_f f(C_f)$, $f(C_f)\mathbf{e}_j = 0\ \forall j$, so $f(C_f) = 0$.
If $0 \neq g(z) = \sum_{i=0}^{m} b_i z^i$ with $m < n$, then $g(C_f)\mathbf{e}_1 = \sum_{i=0}^{m} b_i \mathbf{e}_{i+1} \neq \mathbf{0}$.
Thus no nonzero polynomial of degree $< n$ annihilates C_f, so $\deg \mu_{C_f} = n$.
Since $\mu_{C_f} \mid f$ and both are monic of degree n, $\mu_{C_f} = f$.
By Cayley–Hamilton, $\mu_{C_f} \mid p_{C_f}$; as p_{C_f} and f are monic of degree n, $p_{C_f} = f$.
Structural Commentary: The cyclic vector $\mathbf{e}_1$ generates the basis, and last column imposes relation $f(C_f) = 0$.
Key Insight: The companion matrix makes f the first possible annihilating relation. □

Theorem 3.5.14. *Cyclic Subspaces as Polynomial Quotients*
Let $\dim \mathcal{V} < \infty$, $T \in \mathcal{L}(\mathcal{V})$, $0 \neq \mathbf{v} \in \mathcal{V}$, and let $m_{\mathbf{v}}(z) = z^d + \sum_{i=0}^{d-1} c_i z^i$ be the order of $\mathbf{v}$.
Define $\Phi_{\mathbf{v}} : \mathbb{F}[z] \to \mathcal{C}_T(\mathbf{v})$ by $\Phi_{\mathbf{v}}(p) = p(T)\mathbf{v}$ and $\mathcal{B}_{\mathbf{v}} := (\mathbf{v}, T\mathbf{v}, \ldots, T^{d-1}\mathbf{v})$. Then
(i) $\operatorname{im} \Phi_{\mathbf{v}} = \mathcal{C}_T(\mathbf{v})$, $\ker \Phi_{\mathbf{v}} = \langle m_{\mathbf{v}} \rangle$, and $\mathcal{C}_T(\mathbf{v}) \cong \mathbb{F}[z]/\langle m_{\mathbf{v}} \rangle$;
(ii) $\mathcal{B}_{\mathbf{v}}$ is a basis of $\mathcal{C}_T(\mathbf{v})$;
(iii) $[T|_{\mathcal{C}_T(\mathbf{v})}]_{\mathcal{B}_{\mathbf{v}}} = C_{m_{\mathbf{v}}}$.

Proof.
(i) By Proposition 3.5.10, $\operatorname{im} \Phi_{\mathbf{v}} = \mathcal{C}_T(\mathbf{v})$. $\ker \Phi_{\mathbf{v}} = \operatorname{ann}_T(\mathbf{v}) = \langle m_{\mathbf{v}} \rangle$, by Theorem 1.3.41 $\mathcal{C}_T(\mathbf{v}) \cong \mathbb{F}[z]/\langle m_{\mathbf{v}} \rangle$.
(ii) If $\sum_{j=0}^{d-1} a_j T^j \mathbf{v} = \mathbf{0}$, then $q(z) = \sum_{j=0}^{d-1} a_j z^j \in \operatorname{ann}_T(\mathbf{v})$ and $\deg q < d$, hence $q = 0$; thus $\mathcal{B}_{\mathbf{v}}$ is independent.
For $p = hm_{\mathbf{v}} + r$ with $\deg r < d$ or $r = 0$, $p(T)\mathbf{v} = r(T)\mathbf{v} \in \operatorname{span} \mathcal{B}_{\mathbf{v}}$, so $\mathcal{B}_{\mathbf{v}}$ spans.
(iii) For $0 \leq j < d-1$, $T(T^j\mathbf{v}) = T^{j+1}\mathbf{v}$. Since $m_{\mathbf{v}}(T)\mathbf{v} = \mathbf{0}$, $T^d\mathbf{v} = -\sum_{i=0}^{d-1} c_i T^i \mathbf{v}$, last-column relation of $C_{m_{\mathbf{v}}}$.

Structural Commentary: The first relation cuts out the quotient and companion block.
Key Insight: One orbit carries the whole cyclic subspace. □

Perspective 3.5.15. *The Translation Dictionary*
The module vocabulary for $\mathcal{V}_T$ and the vector-space vocabulary for T match as follows:

$\mathbb{F}[z]$-module $\mathcal{V}_T$	$\mathbb{F}$-vector space $\mathcal{V}$
$p(z) \cdot \mathbf{v}$	$p(T)\mathbf{v}$
submodule	T-invariant subspace
$\operatorname{Ann}_{\mathbb{F}[z]}(\mathcal{V}_T) = \langle \mu_T \rangle$	minimal polynomial of T
$\langle\langle \mathbf{v} \rangle\rangle_{\mathbb{F}[z]}$	$\mathcal{C}_T(\mathbf{v})$
$\mathbb{F}[z]/\langle m_{\mathbf{v}} \rangle$	cyclic block with companion matrix $C_{m_{\mathbf{v}}}$

Key Insight: Once z is interpreted as T, module structure becomes operator structure.

Method 3.5.16. *Moving Between Module and Spectral Data*
For $T \in \mathcal{L}(\mathcal{V})$, translate the $\mathbb{F}[z]$-module data of $\mathcal{V}_T$ into spectral data as follows.
(i) Factor $\mu_T = \prod_{i=1}^{s} p_i^{e_i}$ over $\mathbb{F}$; the primary components are $\ker p_i(T)^{e_i}$.
(ii) Decompose each $\ker p_i(T)^{e_i}$ into cyclic summands; their orders $p_i^{e_{i,j}}$ are the elementary divisors.
(iii) Over $\mathbb{F}$, each elementary divisor $p_i^{e_{i,j}}$ gives a companion block $C_{p_i^{e_{i,j}}}$ in rational canonical form.
(iv) Over splitting field for p_i, same primary cyclic summands refine into Jordan blocks (Perspective 3.5.25).

(v) T is diagonalizable exactly when every elementary divisor is linear, i.e., Corollary 3.5.38 applies.

Key Insight: Rational and Jordan forms encode the same elementary-divisor data over different fields.

Definition 3.5.17. *Invariant Factors and Elementary Divisors*
Let $T \in \mathcal{L}(\mathcal{V})$.

(i) A *cyclic decomposition* is $\mathcal{V} = \bigoplus_{j=1}^{r} \mathcal{C}_j$, where each $\mathcal{C}_j$ is nonzero, T-invariant, and cyclic.
(ii) $\mathrm{InvFact}(T) := (f_1, \dots, f_r)$, where f_j are monic nonconstant invariant factors of $\mathcal{V}_T$; $f_r \mid f_{r-1} \mid \cdots \mid f_1$.
(iii) $\mathrm{ElemDiv}(T)$ is the multiset of monic prime powers in the f_j, counted with multiplicity.
(iv) *Elementary divisor basis* is $\mathcal{B}_{\mathrm{ED}} := \bigsqcup_\alpha (\mathbf{v}_\alpha, T\mathbf{v}_\alpha, \dots, T^{d_\alpha - 1}\mathbf{v}_\alpha)$, where $d_\alpha = \deg m_{\mathbf{v}_\alpha}$; $\{m_{\mathbf{v}_\alpha}\}_\alpha = \mathrm{ElemDiv}(T)$.
(v) For $A \in \mathrm{Mat}(n, n, \mathbb{F})$, set $\mathrm{InvFact}(A) := \mathrm{InvFact}(T_A)$ and $\mathrm{ElemDiv}(A) := \mathrm{ElemDiv}(T_A)$.

Theorem 3.5.18. *Primary Cyclic Decomposition for an Operator*
Let $\mu_T = \prod_{i=1}^{s} p_i^{e_i}$ with $p_i \in \mathbb{F}[z]$ distinct monic irreducible, and set $\mathcal{V}_{p_i} := \ker p_i(T)^{e_i}$. Then

(i) $\mathcal{V}_T = \bigoplus_{i=1}^{s} \mathcal{V}_{p_i}$;
(ii) $\mathcal{V}_{p_i} = \bigoplus_{j=1}^{k_i} \mathcal{C}_T(\mathbf{v}_{i,j})$, with $m_{\mathbf{v}_{i,j}} = p_i^{e_{i,j}}$ and $e_i = e_{i,1} \geq e_{i,2} \geq \cdots \geq e_{i,k_i}$;
(iii) $\mathrm{ElemDiv}(T) = \{p_i^{e_{i,j}} : 1 \leq i \leq s,\ 1 \leq j \leq k_i\}$, corresponding cyclic bases form elementary divisor basis.

Proof.

(i) By Theorem 3.4.39, $\mathcal{V}_T$, $\mathcal{V}_T = \bigoplus_i \mathcal{V}_T[p_i^\infty]$; $\mathcal{V}_T[p_i^\infty] := \{\mathbf{x} \in \mathcal{V} : \exists N \geq 1,\ p_i(T)^N \mathbf{x} = 0\} = \ker p_i(T)^{e_i} = \mathcal{V}_{p_i}$.
(ii) By Theorem 3.4.39, $\mathcal{V}_{p_i} = \bigoplus_{j=1}^{k_i} \mathcal{C}_T(\mathbf{v}_{i,j})$, $m_{\mathbf{v}_{i,j}} = p_i^{e_{i,j}}$ and $e_i = e_{i,1} \geq \cdots \geq e_{i,k_i}$.
(iii) Definition 3.5.17 gives $\mathrm{ElemDiv}(T) = \{p_i^{e_{i,j}} : 1 \leq i \leq s,\ 1 \leq j \leq k_i\}$, and concatenating the bases of $\mathcal{C}_T(\mathbf{v}_{i,j})$ gives an elementary divisor basis.

Structural Commentary: Primary module pieces are polynomial-kernel pieces for the operator.
Key Insight: Elementary divisors are exactly the orders of the cyclic primary summands. □

Theorem 3.5.19. *Invariant-Factor Skeleton for Operators*
Let $0 < \dim \mathcal{V} < \infty$, $T \in \mathcal{L}(\mathcal{V})$, and $\mathrm{InvFact}(T) = (f_1, \dots, f_r)$.
Then $f_r \mid \cdots \mid f_1$ and there are T-cyclic subspaces $\mathcal{C}_j$ such that $\mathcal{V} = \bigoplus_{j=1}^{r} \mathcal{C}_j$, $\mathcal{C}_j \cong \mathbb{F}[z]/\langle f_j \rangle$, $[T] = \bigoplus_{j=1}^{r} C_{f_j}$ in the concatenated cyclic bases, $\mu_T = f_1$, and $p_T = \prod_{j=1}^{r} f_j$.

Proof. By Proposition 3.5.3, $\mathcal{V}_T$ is finitely generated torsion over $\mathbb{F}[z]$.
Theorem 3.4.44 gives $\mathcal{V}_T \cong \bigoplus_{j=1}^{r} \mathbb{F}[z]/\langle f_j \rangle$ with $f_r \mid \cdots \mid f_1$, uniquely.
Transport cyclic generators to T-cyclic subspaces $\mathcal{C}_j$; Theorem 3.5.14 gives $\mathcal{C}_j \cong \mathbb{F}[z]/\langle f_j \rangle$ and block C_{f_j}.
Since $\mathrm{Ann}(\mathcal{V}_T) = \langle \mathrm{lcm}(f_1, \dots, f_r) \rangle = \langle f_1 \rangle = \langle \mu_T \rangle$, $f_1 = \mu_T$.
Finally, Theorem 3.5.13 gives $p_{C_{f_j}} = f_j$, so $p_T = \prod_{j=1}^{r} f_j$.
Structural Commentary: Torsion module decomposition produces companion blocks.
Key Insight: Invariant factors are the polynomial-module version of choosing canonical blocks for similarity. □

Theorem 3.5.20. *Reducing Decompositions Combine the Data*
Let $\mathcal{V} = \mathcal{A} \oplus \mathcal{B}$ with $\mathcal{A}, \mathcal{B}$ T-invariant, and set $T_\mathcal{A} := T|_\mathcal{A}$, $T_\mathcal{B} := T|_\mathcal{B}$.
Then $\mu_T = \mathrm{lcm}(\mu_{T_\mathcal{A}}, \mu_{T_\mathcal{B}})$ and $\mathrm{ElemDiv}(T) = \mathrm{ElemDiv}(T_\mathcal{A}) \sqcup \mathrm{ElemDiv}(T_\mathcal{B})$, where $\sqcup$ is multiset union.
The invariant factors are obtained by recombining the elementary divisors into a divisibility chain.

Proof. By Theorem 3.5.4, $\mathcal{V}_T = \mathcal{A}_{T_\mathcal{A}} \oplus \mathcal{B}_{T_\mathcal{B}}$ as $\mathbb{F}[z]$-modules.
Hence $\mathrm{Ann}(\mathcal{V}_T) = \mathrm{Ann}(\mathcal{A}_{T_\mathcal{A}}) \cap \mathrm{Ann}(\mathcal{B}_{T_\mathcal{B}}) = \langle \mathrm{lcm}(\mu_{T_\mathcal{A}}, \mu_{T_\mathcal{B}}) \rangle$.
Proposition 3.4.40 gives elementary-divisor identity.
Structural Commentary: Reducing decompositions split the operator module.
Key Insight: Direct sums take lcm on global orders and multiset union on elementary divisors. □

Definition 3.5.21. *Characteristic Polynomial from Elementary Divisors*
For $T \in \mathcal{L}(\mathcal{V})$, define $\chi_T(z) := \prod_{q \in \mathrm{ElemDiv}(T)} q(z)$, with product over the multiset $\mathrm{ElemDiv}(T)$.
For $A \in \mathrm{Mat}(n, n, \mathbb{F})$, set $\chi_A := \chi_{T_A}$.

Theorem 3.5.22. *Minimal and Characteristic Polynomials from Elementary Divisors*
Let $\mathrm{ElemDiv}(T) = \{p_i^{e_{i,j}} : 1 \leq i \leq s,\ 1 \leq j \leq k_i\}$.
Then $\mu_T = \prod_i p_i^{\max_j e_{i,j}}$, $\chi_T = \prod_{i,j} p_i^{e_{i,j}}$, and $\deg \chi_T = \dim \mathcal{V}$.
Hence $\mu_T \mid \chi_T$, and μ_T, χ_T have the same irreducible factors.

Proof. By Theorem 3.5.18, $\mathcal{V}_T = \bigoplus_{i,j} \mathcal{C}_T(\mathbf{v}_{i,j})$ with $m_{\mathbf{v}_{i,j}} = p_i^{e_{i,j}}$.
Hence $\operatorname{Ann}(\mathcal{V}_T) = \langle \operatorname{lcm}_{i,j} p_i^{e_{i,j}} \rangle = \langle \prod_i p_i^{\max_j e_{i,j}} \rangle$, so $\mu_T = \prod_i p_i^{\max_j e_{i,j}}$. Definition 3.5.21 gives $\chi_T = \prod_{i,j} p_i^{e_{i,j}}$.
By Theorem 3.5.14, $\dim \mathcal{C}_T(\mathbf{v}_{i,j}) = \deg p_i^{e_{i,j}}$; summing gives $\deg \chi_T = \dim \mathcal{V}$.
Divisibility and common irreducible factors are immediate.
Structural Commentary: μ_T records maximal exponents; χ_T records all exponents.
Key Insight: The pair (μ_T, χ_T) forgets the block distribution. □

Definition 3.5.23. *Nonderogatory Operators*
$T \in \mathcal{L}(\mathcal{V})$ is *nonderogatory* if $\mu_T = \chi_T$, equivalently $\deg \mu_T = \dim \mathcal{V}$.
For $A \in \operatorname{Mat}(n, n, \mathbb{F})$, A is nonderogatory if T_A is.

Theorem 3.5.24. *Cyclic and Nonderogatory Are the Same*
For $T \in \mathcal{L}(\mathcal{V})$, the following are equivalent:

(i) $\mathcal{V}_T$ *is cyclic;*
(ii) T *has a cyclic vector;*
(iii) $\operatorname{ElemDiv}(T) = \{p_1^{e_1}, \ldots, p_s^{e_s}\}$, *one elementary divisor for each irreducible factor of* μ_T*;*
(iv) T *is nonderogatory.*

Proof.

(i) $\Rightarrow$ *(ii)* If $\mathcal{V}_T = \mathbb{F}[z]\mathbf{v}$, then $\mathcal{V} = \{p(T)\mathbf{v} : p \in \mathbb{F}[z]\} = \mathcal{C}_T(\mathbf{v})$, so $\mathbf{v}$ is cyclic.
(ii) $\Rightarrow$ *(i)* If $\mathcal{C}_T(\mathbf{v}) = \mathcal{V}$, then $\mathcal{V}_T = \mathbb{F}[z]\mathbf{v}$.
(i) $\Leftrightarrow$ *(iii)* Theorem 3.4.47 on $\mathcal{V}_T$. Cyclic torsion module has one cyclic summand for each primary factor.
(iii) $\Rightarrow$ *(iv)* If $\operatorname{ElemDiv}(T) = \{p_i^{e_i}\}_{i=1}^s$, then $\chi_T = \prod_i p_i^{e_i} = \mu_T$, so T is nonderogatory.
(iv) $\Rightarrow$ *(iii)* Let $\operatorname{ElemDiv}(T) = \{p_i^{e_{i,j}}\}$. If $\mu_T = \chi_T$, then $\prod_i p_i^{\max_j e_{i,j}} = \prod_{i,j} p_i^{e_{i,j}}$, so $k_i = 1$ for all i.

Structural Commentary: Nonderogatory means all primary information can be generated by one vector.
Key Insight: Equality $\mu_T = \chi_T$ says that product and lcm of the cyclic orders coincide. □

Perspective 3.5.25. *Elementary Divisors and Jordan Blocks*
Over $\mathbb{C}$, elementary divisors are exactly $(z-\lambda)^k$, and $\mathbb{C}[z]/\langle (z-\lambda)^k \rangle$ corresponds to the Jordan block $J_k(\lambda)$.
Thus $\operatorname{ElemDiv}(T)$ is the multiset of Jordan blocks.

Perspective 3.5.26. *Crosswalk with the Spectral Presentation*
The module and spectral presentations encode the same data at different splitting levels:

(i) diagonalizable $\Longleftrightarrow$ $\operatorname{ElemDiv}(T)$ consists of linear factors, by Theorem 3.2.27;
(ii) $(z-\lambda)^k \Longleftrightarrow J_k(\lambda)$ over an algebraically closed field, by Perspective 3.5.25 and Theorem 3.3.6;
(iii) rank drops recover the exponents k via Proposition 3.4.37, with $p = z - \lambda$ and $M = \mathcal{V}_T$.

Key Insight: Rational form is Jordan form before irreducibles split.

Perspective 3.5.27. *Operator Modules and Similarity*
For $T \in \mathcal{L}(\mathcal{V})$, $\mathcal{V}_T$ has action $p \cdot \mathbf{v} = p(T)\mathbf{v}$. Then $S \sim T \Longleftrightarrow \mathcal{V}_S \cong_{\mathbb{F}[z]} \mathcal{V}_T$; $\mathbb{F}[z]/\langle f \rangle \Longleftrightarrow C_f$ (Theorem 3.5.14).

Theorem 3.5.28. *Complete Similarity Invariants*
Let $T, S \in \mathcal{L}(\mathcal{V})$ and $A, B \in \operatorname{Mat}(n, n, \mathbb{F})$.

(i) $T \sim S \Longleftrightarrow \mathcal{V}_T \cong_{\mathbb{F}[z]} \mathcal{V}_S \Longleftrightarrow \operatorname{ElemDiv}(T) = \operatorname{ElemDiv}(S) \Longleftrightarrow \operatorname{InvFact}(T) = \operatorname{InvFact}(S)$.
(ii) $A \sim B \Longleftrightarrow (\mathbb{F}^n)_A \cong_{\mathbb{F}[z]} (\mathbb{F}^n)_B \Longleftrightarrow \operatorname{ElemDiv}(A) = \operatorname{ElemDiv}(B) \Longleftrightarrow \operatorname{InvFact}(A) = \operatorname{InvFact}(B)$.
(iii) If $A = [T]_{\mathcal{B}}$ *for some basis* $\mathcal{B}$ *of* $\mathcal{V}$, *then* $\operatorname{ElemDiv}(A) = \operatorname{ElemDiv}(T)$ *and* $\operatorname{InvFact}(A) = \operatorname{InvFact}(T)$.

Proof.

(i) $(\Rightarrow)$ If $T \sim S$, then $\mathcal{V}_T \cong_{\mathbb{F}[z]} \mathcal{V}_S$ by Theorem 3.5.5.
Module isomorphism preserves the PID decomposition, hence preserves ElemDiv and InvFact.
(i) $(\Leftarrow)$ Equality of either ElemDiv or InvFact gives $\mathcal{V}_T \cong_{\mathbb{F}[z]} \mathcal{V}_S$ by uniqueness of the PID decomposition.
Thus $T \sim S$ by Theorem 3.5.5.
(ii) $(\Rightarrow)$ If $A \sim B$, then $(\mathbb{F}^n)_A \cong_{\mathbb{F}[z]} (\mathbb{F}^n)_B$ by Corollary 3.5.6; hence the divisor data agree.
(ii) $(\Leftarrow)$ Equality of either matrix divisor datum gives $(\mathbb{F}^n)_A \cong_{\mathbb{F}[z]} (\mathbb{F}^n)_B$, so $A \sim B$ by Corollary 3.5.6.
(iii) If $A = [T]_{\mathcal{B}}$, then $\Phi(\mathbf{v}) = [\mathbf{v}]_{\mathcal{B}}$ satisfies $\Phi(p(T)\mathbf{v}) = p(A)\Phi(\mathbf{v})$.
Hence $\Phi : \mathcal{V}_T \to (\mathbb{F}^n)_A$ is an $\mathbb{F}[z]$-isomorphism, so the divisor data agree.

Structural Commentary: Similarity is module isomorphism plus PID data.
Key Insight: Elementary divisors classify operator modules. □

Lemma 3.5.29. *Determinant of a Companion Block*
Let $f(z) = z^n + \sum_{i=0}^{n-1} c_i z^i$ be monic. Then $\det(zI_n - C_f) = f(z)$.

Proof. Set $D_n(c_0,\dots,c_{n-1}) := \det(zI_n - C_f)$. For $n=1$, $D_1(c_0) = z + c_0$.
For $n \geq 2$, first-row expansion gives $D_n(c_0,\dots,c_{n-1}) = zD_{n-1}(c_1,\dots,c_{n-1}) + c_0$, since the $(1,1)$-minor is again companion-shaped and the last first-row cofactor contributes c_0. Induction gives $D_n = z^n + \sum_{i=0}^{n-1} c_i z^i = f(z)$.
Structural Commentary: The last column of C_f stores the coefficients of the polynomial relation.
Key Insight: The shifted determinant reads that stored relation back as f. □

Definition 3.5.30. *Rational Canonical Block Forms*

(i) *Invariant-factor RCF*: $A = \operatorname{diag}(C_{s_1},\dots,C_{s_r})$, where $s_j \in \mathbb{F}[z]$ are monic nonconstant and $s_r \mid \cdots \mid s_1$.
(ii) *Elementary-divisor RCF*: $A = \operatorname{diag}(C_{q_1},\dots,C_{q_m})$, where each q_j is power of monic irreducible polynomial.

Theorem 3.5.31. *Rational Canonical Form Skeleton*
For $A \in \operatorname{Mat}(n,n,\mathbb{F})$, there is a unique tuple $\operatorname{InvFact}(A) = (s_1,\dots,s_r)$ with $s_j \in \mathbb{F}[z]$ monic nonconstant, $s_r \mid \cdots \mid s_1$, and $A \sim_{\mathbb{F}} R_A := \operatorname{diag}(C_{s_1},\dots,C_{s_r})$. Moreover, $\mu_A = s_1$ and $p_A = \prod_{j=1}^r s_j$.
Equivalently, for $A, B \in \operatorname{Mat}(n,n,\mathbb{F})$, $A \sim_{\mathbb{F}} B \iff \operatorname{ElemDiv}(A) = \operatorname{ElemDiv}(B)$ as multisets.

Proof. Let $T_A : \mathbb{F}^n \to \mathbb{F}^n$, $\mathbf{x} \mapsto A\mathbf{x}$.
Theorem 3.5.19 gives $\mathbb{F}^n = \bigoplus_{j=1}^r \mathcal{C}_j$, $\mathcal{C}_j \cong_{\mathbb{F}[z]} \mathbb{F}[z]/\langle s_j\rangle$, $s_r \mid \cdots \mid s_1$, $[T_A] = R_A$, $\mu_{T_A} = s_1$, and $p_{T_A} = \prod_{j=1}^r s_j$.
Hence $A \sim_{\mathbb{F}} R_A$, $\mu_A = s_1$, and $p_A = \prod_{j=1}^r s_j$; uniqueness of $\operatorname{InvFact}(A)$ is uniqueness in the same theorem.
($\Rightarrow$) If $A \sim_{\mathbb{F}} B$, then Theorem 3.5.28 gives $\operatorname{ElemDiv}(A) = \operatorname{ElemDiv}(B)$.
($\Leftarrow$) If $\operatorname{ElemDiv}(A) = \operatorname{ElemDiv}(B)$, then Theorem 3.5.28 gives $A \sim_{\mathbb{F}} B$.
Structural Commentary: Companion blocks classify similarity without requiring p_A to split.
Key Insight: Similarity classes can be represented by polynomial data rather than by choosing roots. □

Theorem 3.5.32. *Rational Canonical Form: Two Versions*
Let $T \in \mathcal{L}(\mathcal{V})$.

(i) If $\operatorname{ElemDiv}(T) = \{q_1,\dots,q_m\}$, then $\exists$ basis $\mathcal{R}$ with $[T]_{\mathcal{R}} = \operatorname{diag}(C_{q_1},\dots,C_{q_m})$. The form is unique up to block permutation.
(ii) If $\operatorname{InvFact}(T) = (s_1,\dots,s_r)$ and $s_r \mid \cdots \mid s_1$, then $\exists$ basis $\mathcal{B}$ with $[T]_{\mathcal{B}} = \operatorname{diag}(C_{s_1},\dots,C_{s_r})$. The form is unique in this order.
(iii) If $A \in \operatorname{Mat}(n,n,\mathbb{F})$, then the similarity class $[A]_{\sim_{\mathbb{F}}}$ contains both forms. Elementary-divisor form is unique up to block order, invariant-factor form by its divisibility chain.

Proof.

(i) Theorem 3.5.18 gives $\mathcal{V} = \bigoplus_{\alpha=1}^m \mathcal{C}_\alpha$ with orders q_α; Theorem 3.5.14 gives cyclic bases and blocks C_{q_α}. Permuting summands conjugates by a permutation matrix.
(ii) Theorem 3.5.19 gives $\mathcal{V} = \bigoplus_{j=1}^r \mathcal{C}_j$, $\mathcal{C}_j \cong_{\mathbb{F}[z]} \mathbb{F}[z]/\langle s_j\rangle$, and cyclic bases with block matrix $\operatorname{diag}(C_{s_1},\dots,C_{s_r})$.
(iii) For A, apply (i),(ii) to $T_A : \mathbf{x} \mapsto A\mathbf{x}$; any chosen basis gives a similarity matrix. Uniqueness follows from Theorem 3.5.28.

Structural Commentary: Elementary divisors split primary blocks; invariant factors recombine them.
Key Insight: Both forms classify the same similarity classes, but they expose different polynomial data. □

Method 3.5.33. *Constructing an Operator from Elementary Divisors*
Let $Q = \{q_1,\dots,q_m\}$ be a finite multiset of monic prime powers in $\mathbb{F}[z]$, and set $n := \sum_{j=1}^m \deg q_j$.

(i) Set $R_Q := \operatorname{diag}(C_{q_1},\dots,C_{q_m}) \in \operatorname{Mat}(n,n,\mathbb{F})$.
(ii) Then $\operatorname{ElemDiv}(T_{R_Q}) = Q$; hence every such Q is realised on $\mathbb{F}^n$.
(iii) Apply Method 3.4.42 to recombine Q into $\operatorname{InvFact}(T_{R_Q}) = (s_1,\dots,s_r)$, then $\operatorname{diag}(C_{s_1},\dots,C_{s_r}) \sim_{\mathbb{F}} R_Q$.

Structural Commentary: RCF constructs operators from divisor data.
Key Insight: Companion blocks realise cyclic summands; block sums realise direct sums.

Example 3.5.34. *Rational Canonical Possibilities from Degree Counts*
Let $T \in \mathcal{L}(\mathbb{R}^7)$ and $\mu_T = (z-1)(z^2+1)^2$. For $Q = \operatorname{ElemDiv}(T)$, $\operatorname{lcm} Q = \mu_T$ and $\sum_{q\in Q} \deg q = 7$.
Hence $Q_1 = \{z-1, (z^2+1)^2, z^2+1\}$ or $Q_2 = \{z-1, z-1, z-1, (z^2+1)^2\}$.
Thus the elementary-divisor RCF is $R_i = \bigoplus_{q\in Q_i} C_q$ for $i = 1, 2$, up to block order.

Theorem 3.5.35. *Determinant Form of the Characteristic Polynomial*
For $A \in \operatorname{Mat}(n,n,\mathbb{F})$, $\chi_A(z) = \det(zI_n - A)$. If $A = [T]_{\mathcal{B}}$, then $\chi_T = \chi_A$.

Proof. Let $Q := \operatorname{ElemDiv}(A)$ and choose elementary-divisor RCF $R = \bigoplus_{q\in Q} C_q \sim_{\mathbb{F}} A$.
By Lemma 3.5.29, $\det(zI - R) = \prod_{q\in Q} \det(zI - C_q) = \prod_{q\in Q} q = \chi_R$.
Since similarity preserves both Q and $\det(zI - \cdot)$, $\chi_A = \chi_R = \det(zI - R) = \det(zI - A)$.
If $A = [T]_{\mathcal{B}}$, then $\operatorname{ElemDiv}(A) = \operatorname{ElemDiv}(T)$, hence $\chi_A = \chi_T$.

Structural Commentary: Rational canonical form reduces the determinant computation to companion blocks.
Key Insight: The abstract product of elementary divisors is the concrete shifted determinant. □

Theorem 3.5.36. *Changing the Base Field*
Let $\mathbb{F} \subseteq \mathbb{K}$, $A \in \mathrm{Mat}(n,n,\mathbb{F})$, *and* $\mathrm{ElemDiv}_{\mathbb{F}}(A) = \{p_i^{e_{i,j}}\}_{i,j}$ *with* $p_i \in \mathbb{F}[z]$ *monic irreducible.*
If $p_i = \prod_{\ell=1}^{m_i} a_{i,\ell}^{d_{i,\ell}}$ *in* $\mathbb{K}[z]$, *with* $a_{i,\ell}$ *pairwise distinct monic irreducible over* $\mathbb{K}$, *then* $\mathrm{ElemDiv}_{\mathbb{K}}(A) = \{a_{i,\ell}^{d_{i,\ell}e_{i,j}}\}_{i,j,\ell}$.
Moreover, $\mathrm{InvFact}_{\mathbb{K}}(A) = \mathrm{InvFact}_{\mathbb{F}}(A)$.

Proof. Take elementary-divisor RCF over $\mathbb{F}$: $R = \bigoplus_{i,j} C_{p_i^{e_{i,j}}} \sim_{\mathbb{F}} A$.
Since $R \in \mathrm{Mat}(n,n,\mathbb{K})$ and $p_i^{e_{i,j}} = \prod_{\ell=1}^{m_i} a_{i,\ell}^{d_{i,\ell}e_{i,j}}$ with coprime factors in $\mathbb{K}[z]$,
Theorem 3.4.17 gives $C_{p_i^{e_{i,j}}} \sim_{\mathbb{K}} \bigoplus_{\ell=1}^{m_i} C_{a_{i,\ell}^{d_{i,\ell}e_{i,j}}}$. Summing over (i,j) yields $\mathrm{ElemDiv}_{\mathbb{K}}(A)$.
Recombination into the divisibility chain gives the original invariant factors.
Structural Commentary: Field extension refines elementary divisors.
Key Insight: Elementary divisors split; invariant factors persist. □

Corollary 3.5.37. *Nonderogatory Means One Companion Block*
For $A \in \mathrm{Mat}(n,n,\mathbb{F})$, *the following are equivalent:*

(i) A *is nonderogatory, i.e.* $\mu_A = \chi_A$;
(ii) $(\mathbb{F}^n)_A$ *is cyclic;*
(iii) $A \sim_{\mathbb{F}} C_{\chi_A}$.

Proof.

$(i) \Rightarrow (ii)$ This is Theorem 3.5.24.
$(ii) \Rightarrow (iii)$ Cyclicity gives $\mathrm{InvFact}(A) = (s)$. By Theorem 3.5.19, $s = \mu_A = \chi_A$; hence $A \sim_{\mathbb{F}} C_{\chi_A}$.
$(iii) \Rightarrow (i)$ Since $\mu_{C_f} = f$ by Theorem 3.5.13, similarity gives $\mu_A = \mu_{C_{\chi_A}} = \chi_A$.

Structural Commentary: Nonderogatory means one cyclic block.
Key Insight: One cyclic generator gives one companion block. □

Corollary 3.5.38. *When the Elementary-Divisor Form Is Diagonal*
Let $T \in \mathcal{L}(\mathcal{V})$, $Q = \mathrm{ElemDiv}(T)$, *and* $R_{\mathrm{ed}} = \bigoplus_{q \in Q} C_q$.
Then R_{ed} *is diagonal iff* $\deg q = 1$ *for all* $q \in Q$, *iff* $\mu_T = \prod_{\lambda \in \Lambda}(z-\lambda)$ *for some finite* $\Lambda \subseteq \mathbb{F}$.

Proof.

($\Rightarrow$) If R_{ed} is diagonal, each block C_q has size 1, hence $\deg q = 1$ for all $q \in Q$.
Thus $\mu_T = \mathrm{lcm}\, Q$ is a product of distinct linear factors.
($\Leftarrow$) If $\mu_T = \prod_{\lambda \in \Lambda}(z-\lambda)$, then each prime power $q \in Q$ divides μ_T, so $q = z - \lambda$ for some $\lambda \in \Lambda$.
Hence each C_q is 1×1, and R_{ed} is diagonal.

Structural Commentary: Diagonal RCF means all companion blocks are 1×1.
Key Insight: Nonlinear or repeated factors obstruct diagonal RCF. □

Proposition 3.5.39. *AB and BA Have the Same Characteristic Polynomial*
For $A, B \in \mathrm{Mat}(n,n,\mathbb{F})$, $\chi_{AB} = \chi_{BA}$.
If $A \in \mathrm{GL}(n,\mathbb{F})$ *or* $B \in \mathrm{GL}(n,\mathbb{F})$, *then* $AB \sim_{\mathbb{F}} BA$; *otherwise not necessarily.*

Proof. In $\mathbb{F}(z)$, by Sylvester's identitym $\det(zI-AB) = z^n \det(I - z^{-1}AB) = z^n \det(I - z^{-1}BA) = \det(zI - BA)$, hence $\chi_{AB} = \chi_{BA}$. If A is invertible, $BA = A^{-1}(AB)A$; if B is invertible, $AB = B^{-1}(BA)B$.
Without invertibility, take $A = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}$ and $B = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$.
Then $AB = B$, $BA = 0$, $\chi_{AB} = \chi_{BA} = z^2$, but $\mathrm{rank}(AB) = 1 \neq 0 = \mathrm{rank}(BA)$, so $AB \not\sim_{\mathbb{F}} BA$.
Structural Commentary: Characteristic polynomials see less than similarity.
Key Insight: AB and BA share determinant data; invertibility yields similarity. □

Example 3.5.40. *The Same Polynomial over Different Fields*
Let $\dim V = 4$ and $\mu_T = (z^2+1)(z^2-2)$. Since $\deg \mu_T = 4$, T is nonderogatory and $\mathrm{InvFact}_K(T) = (\mu_T)$ for $K = \mathbb{Q}, \mathbb{R}, \mathbb{C}$; hence $R_{\mathrm{inv},K} = C_{\mu_T}$. Meanwhile

K	$\mathrm{ElemDiv}_K(T)$
$\mathbb{Q}$	$\{z^2+1,\ z^2-2\}$
$\mathbb{R}$	$\{z^2+1,\ z-\sqrt{2},\ z+\sqrt{2}\}$
$\mathbb{C}$	$\{z-i,\ z+i,\ z-\sqrt{2},\ z+\sqrt{2}\}$

Thus $R_{\mathrm{ed},K} = \bigoplus_{q \in \mathrm{ElemDiv}_K(T)} C_q$: scalar extension refines elementary divisors while preserving the invariant factor. Over $\mathbb{C}$, $R_{\mathrm{ed},\mathbb{C}}$ is diagonal.

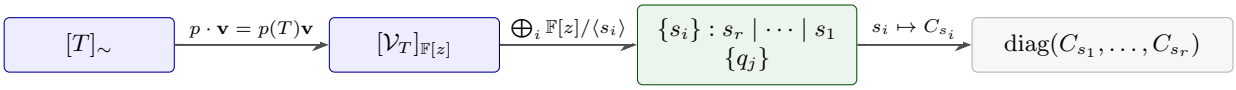


Figure 104: Rational-canonical dictionary: $[T]_\sim$, $[\mathcal{V}_T]_{\mathbb{F}[z]}$, $\{s_i\}$ or $\{q_j\}$, and $\mathrm{diag}(C_{s_i})$ are equivalent data.

3.6 Associative Algebras

The module viewpoint in Definition 3.5.2 uses one operator at a time. An algebra keeps track of a whole vector space of operators or products at once. The guiding example is $\mathcal{L}(\mathcal{V})$, where addition and scalar multiplication are pointwise and multiplication is composition.

Definition 3.6.1. *Unital Associative Algebra*
Let $\mathbb{F}$ be a field.

(i) An *associative $\mathbb{F}$-algebra* is an $\mathbb{F}$-vector space $\mathcal{A}$ and unital ring $(\mathcal{A}, +, \cdot, 1_{\mathcal{A}})$ such that $r(ab) = (ra)b = a(rb)$ for all $r \in \mathbb{F}$ and $a, b \in \mathcal{A}$.
(ii) $\mathcal{A}$ is *finite-dimensional* iff $\dim_{\mathbb{F}} \mathcal{A} < \infty$, and *commutative* iff $ab = ba$ for all $a, b \in \mathcal{A}$.
(iii) $a \in \mathcal{A}$ is *invertible* iff $\exists b \in \mathcal{A}$ such that $ab = ba = 1_{\mathcal{A}}$.
(iv) In this subsection, *algebra* means unital associative $\mathbb{F}$-algebra unless stated otherwise.

Remark 3.6.2. *Convention*
Relative to Definition 3.1.19, assume multiplication is associative and unital.
Thus, *algebra* means unital associative $\mathbb{F}$-algebra.

Definition 3.6.3. *Center and Central Algebras*

(i) The *center* of an $\mathbb{F}$-algebra $\mathcal{A}$ is $Z(\mathcal{A}) := \{a \in \mathcal{A} : ax = xa \ \forall x \in \mathcal{A}\}$.
(ii) $\mathcal{A}$ is *central over $\mathbb{F}$* iff $Z(\mathcal{A}) = \mathbb{F}1_{\mathcal{A}} := \{r1_{\mathcal{A}} : r \in \mathbb{F}\}$.

Proposition 3.6.4. *The Scalar Field Sits in the Center*
If $\mathcal{A}$ is an $\mathbb{F}$-algebra with $1_{\mathcal{A}} \neq 0$, then $\iota_{\mathcal{A}} : \mathbb{F} \hookrightarrow Z(\mathcal{A})$, $\quad r \mapsto r1_{\mathcal{A}}$, is an injective unital ring homomorphism. Conversely, if R is a unital ring and $\mathbb{F} \leq Z(R)$ is a subfield, then R is an $\mathbb{F}$-algebra via $r \cdot x := rx$.

Proof.

(⇒) Write $\iota = \iota_{\mathcal{A}}$. For $r, s \in \mathbb{F}$, $\iota(r+s) = (r+s)1_{\mathcal{A}} = \iota(r) + \iota(s)$, $\iota(rs) = (rs)1_{\mathcal{A}} = (r1_{\mathcal{A}})(s1_{\mathcal{A}})$, $\iota(1) = 1_{\mathcal{A}}$. For $r \in \mathbb{F}$ and $a \in \mathcal{A}$, $\iota(r)a = (r1_{\mathcal{A}})a = r(1_{\mathcal{A}}a) = ra = r(a1_{\mathcal{A}}) = a(r1_{\mathcal{A}}) = a\iota(r)$, so $\operatorname{im}\iota \subseteq Z(\mathcal{A})$. If $\iota(r) = 0$ and $r \neq 0$, then $1_{\mathcal{A}} = \iota(r^{-1})\iota(r) = 0$, contradiction; hence $\ker\iota = \{0\}$.
(⇐) Define $r \cdot x = rx$. The field laws in $\mathbb{F} \subseteq R$ give the vector-space axioms, and $1_{\mathbb{F}} \cdot x = x$. For $r \in \mathbb{F}$ and $x, y \in R$, centrality gives $r \cdot (xy) = rxy = (rx)y = x(ry) = x(r \cdot y)$, so the scalar action is compatible with the ring product.

Structural Commentary: An algebra structure is a choice of central scalar copy $\mathbb{F}1_{\mathcal{A}} \subseteq Z(\mathcal{A})$.
Key Insight: Scalars are precisely central multipliers of the unit. □

Method 3.6.5. *Building Multiplication from Structure Constants*
Let $V = \bigoplus_{i \in I} \mathbb{F}b_i$. Choose $c_{ij}^k \in \mathbb{F}$ with $\{k : c_{ij}^k \neq 0\}$ finite for each (i, j).

(i) Define $b_i b_j := \sum_{k \in I} c_{ij}^k b_k$ and extend bilinearly: $(\sum_i r_i b_i)(\sum_j s_j b_j) := \sum_{i,j,k} r_i s_j c_{ij}^k b_k$.
(ii) The product is associative iff $\sum_{\ell \in I} c_{ij}^\ell c_{\ell k}^m = \sum_{\ell \in I} c_{jk}^\ell c_{i\ell}^m$ for all $i, j, k, m \in I$.
(iii) The product is commutative iff $c_{ij}^k = c_{ji}^k$ for all $i, j, k \in I$.
(iv) A basis vector b_e is an identity iff $c_{ej}^m = \delta_j^m = c_{je}^m$ for all $j, m \in I$.

Example 3.6.6. *Basic Algebras*

(i) If $\mathbb{F} \subseteq \mathbb{E}$ are fields, then $\mathbb{E}$ is an $\mathbb{F}$-algebra by $r \cdot x = rx$ and multiplication in $\mathbb{E}$.
(ii) $\mathbb{F}[x]$ is a commutative $\mathbb{F}$-algebra with unit 1 and scalar embedding $r \mapsto r$.
(iii) $\mathrm{Mat}(n, n, \mathbb{F})$ is an $\mathbb{F}$-algebra with unit I_n and product AB.
(iv) For an $\mathbb{F}$-vector space $\mathcal{V}$, $\mathcal{L}(\mathcal{V}) = \mathrm{End}_{\mathbb{F}}(\mathcal{V})$ is an $\mathbb{F}$-algebra with unit $I_{\mathcal{V}}$ and product $S \circ T$.

Proposition 3.6.7. *Endomorphisms and Matrices Are the Same Algebra after Choosing a Basis*
Let $\dim \mathcal{V} = n$ and fix an ordered basis $\mathcal{B}$ of $\mathcal{V}$.
Then $\mu_{\mathcal{B}} : \mathcal{L}(\mathcal{V}) \to \mathrm{Mat}(n, n, \mathbb{F})$, $T \mapsto [T]_{\mathcal{B}}$, is an $\mathbb{F}$-algebra isomorphism.

Proof. By Theorem 1.3.89, $\mu_{\mathcal{B}}$ is an $\mathbb{F}$-linear bijection.
For $S, T \in \mathcal{L}(\mathcal{V})$ and $x \in \mathcal{V}$, $[STx]_{\mathcal{B}} = [S]_{\mathcal{B}}[Tx]_{\mathcal{B}} = [S]_{\mathcal{B}}[T]_{\mathcal{B}}[x]_{\mathcal{B}}$, hence $\mu_{\mathcal{B}}(ST) = \mu_{\mathcal{B}}(S)\mu_{\mathcal{B}}(T)$.
Also $\mu_{\mathcal{B}}(I_{\mathcal{V}}) = I_n$; therefore $\mu_{\mathcal{B}}$ is an $\mathbb{F}$-algebra isomorphism.
Structural Commentary: A basis is a coordinate bridge from abstract endomorphisms to concrete matrices.
Key Insight: Composition becomes matrix multiplication after taking coordinate columns. □

Definition 3.6.8. *Subalgebras, Ideals, Quotients, Homomorphisms, and Regular Representation*
Let $\mathcal{A}$ be an $\mathbb{F}$-algebra.

(i) A *subalgebra* is a subset $\mathcal{B} \subseteq \mathcal{A}$ such that $1_{\mathcal{A}} \in \mathcal{B}$, $\mathcal{B} \leq_{\mathbb{F}} \mathcal{A}$, and $b, b' \in \mathcal{B} \Rightarrow bb' \in \mathcal{B}$.
(ii) For $X \subseteq \mathcal{A}$, the *subalgebra generated by X* is $\langle X \rangle_{\text{subalg}} := \bigcap\{\mathcal{B} : \mathcal{B} \text{ is a subalgebra of } \mathcal{A},\ X \subseteq \mathcal{B}\}$.
Equivalently, $\langle X \rangle_{\text{subalg}} = \text{span}_{\mathbb{F}}\{1_{\mathcal{A}},\ x_1 \cdots x_m : m \geq 1,\ x_i \in X\}$.
(iii) A *two-sided ideal* is a subset $\mathcal{I} \subseteq \mathcal{A}$ such that $\mathcal{I} \neq \emptyset$, $u, v \in \mathcal{I} \Rightarrow u - v \in \mathcal{I}$, and $a \in \mathcal{A},\ u \in \mathcal{I} \Rightarrow au, ua \in \mathcal{I}$.
(iv) For nonempty $X \subseteq \mathcal{A}$, the *ideal generated by X* is $\langle X \rangle_{\text{ideal}} := \{\sum_{j=1}^{n} a_j x_j b_j : n \geq 1,\ x_j \in X,\ a_j, b_j \in \mathcal{A}\}$.
(v) $\mathcal{A}$ is *simple* iff $\mathcal{A}^2 \neq 0$ and its only two-sided ideals are 0 and $\mathcal{A}$.
(vi) If $\mathcal{I} \triangleleft \mathcal{A}$, then $\mathcal{A}/\mathcal{I}$ is an $\mathbb{F}$-algebra by $(a + \mathcal{I})(b + \mathcal{I}) := ab + \mathcal{I}$ and $r(a + \mathcal{I}) := ra + \mathcal{I}$.
(vii) For $\mathbb{F}$-algebras $\mathcal{A}, \mathcal{B}$, an *algebra homomorphism* is a map $\sigma : \mathcal{A} \to \mathcal{B}$ such that, for all $a, a' \in \mathcal{A}$ and $r \in \mathbb{F}$, $\sigma(a + a') = \sigma(a) + \sigma(a')$, $\sigma(aa') = \sigma(a)\sigma(a')$, $\sigma(1_{\mathcal{A}}) = 1_{\mathcal{B}}$, and $\sigma(ra) = r\sigma(a)$.
(viii) For $a \in \mathcal{A}$, the *left multiplication operator* is $L_a : \mathcal{A} \to \mathcal{A}$, $x \mapsto ax$.
The map $L_{\mathcal{A}} : \mathcal{A} \to \mathcal{L}(\mathcal{A})$, $a \mapsto L_a$, is the *left regular representation* of $\mathcal{A}$.

Proposition 3.6.9. *Kernel and First Isomorphism Theorem for Algebras*
If $\sigma : \mathcal{A} \to \mathcal{B}$ is an $\mathbb{F}$-algebra homomorphism,
then $\ker\sigma \triangleleft \mathcal{A}$ and $\overline{\sigma} : \mathcal{A}/\ker\sigma \to \text{im}\,\sigma$, $a + \ker\sigma \mapsto \sigma(a)$, is an $\mathbb{F}$-algebra isomorphism.

Proof. Set $K := \ker\sigma$. Linearity gives $K \leq_{\mathbb{F}} \mathcal{A}$.
If $u \in K$ and $a \in \mathcal{A}$, then $\sigma(au) = \sigma(a)\sigma(u) = 0 = \sigma(u)\sigma(a) = \sigma(ua)$, hence $au, ua \in K$ and $K \triangleleft \mathcal{A}$.
Define $\overline{\sigma}(a + K) := \sigma(a)$. Since $a + K = a' + K \iff a - a' \in K \iff \sigma(a) = \sigma(a')$, $\overline{\sigma}$ is well defined and injective; it is surjective onto $\text{im}\,\sigma$.
For $a, a' \in \mathcal{A}$ and $r \in \mathbb{F}$, $\overline{\sigma}((a + K)(a' + K)) = \sigma(aa') = \sigma(a)\sigma(a')$, $\overline{\sigma}(r(a + K)) = r\sigma(a)$, and $\overline{\sigma}(1_{\mathcal{A}} + K) = 1_{\mathcal{B}}$.
Thus $\overline{\sigma}$ is an $\mathbb{F}$-algebra isomorphism.
Structural Commentary: Quotienting by the kernel removes exactly the relations invisible to σ.
Key Insight: The image is the faithful algebra obtained after collapsing $\ker\sigma$. □

Theorem 3.6.10. *Left Regular Representation*
For every $\mathbb{F}$-algebra $\mathcal{A}$, $L_{\mathcal{A}} : \mathcal{A} \hookrightarrow \mathcal{L}(\mathcal{A})$, $a \mapsto L_a$, is an $\mathbb{F}$-algebra monomorphism.

Proof. For $a, b \in \mathcal{A}$, $r \in \mathbb{F}$, and $x \in \mathcal{A}$, $L_{a+b}(x) = (a + b)x = L_a(x) + L_b(x)$ and $L_{ra}(x) = (ra)x = rL_a(x)$, so $L_{\mathcal{A}}$ is $\mathbb{F}$-linear. Also $L_{ab}(x) = (ab)x = a(bx) = (L_aL_b)(x)$ and $L_{1_{\mathcal{A}}} = I_{\mathcal{A}}$, so products and units are preserved.
If $L_a = 0$, then $a = L_a(1_{\mathcal{A}}) = 0$, so $L_{\mathcal{A}}$ is injective.
Structural Commentary: Left multiplication realises $\mathcal{A}$ inside its own operator algebra.
Key Insight: The unit detects elements, since $L_a(1_{\mathcal{A}}) = a$. □

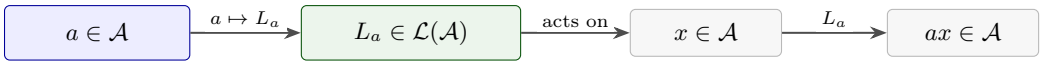


Figure 105: The left regular representation sends $a \in \mathcal{A}$ to $L_a : x \mapsto ax$; $L_{ab} = L_aL_b$ and $L_a(1_{\mathcal{A}}) = a$ give multiplicativity and injectivity.

Theorem 3.6.11. *Center and Ideals of an Endomorphism Algebra*
Let $0 \neq \mathcal{V}$ be an $\mathbb{F}$-vector space, and set $\mathcal{F}(\mathcal{V}) := \{T \in \mathcal{L}(\mathcal{V}) : \dim\text{im}\,T < \infty\}$.

(i) $Z(\mathcal{L}(\mathcal{V})) = \mathbb{F}I_{\mathcal{V}}$, so $\mathcal{L}(\mathcal{V})$ is central.
(ii) $\mathcal{F}(\mathcal{V}) \triangleleft \mathcal{L}(\mathcal{V})$, and $0 \neq \mathcal{J} \triangleleft \mathcal{L}(\mathcal{V}) \Rightarrow \mathcal{F}(\mathcal{V}) \subseteq \mathcal{J}$.
(iii) $\mathcal{L}(\mathcal{V})$ is simple iff $\dim\mathcal{V} < \infty$.

Proof.

(i) Let $T \in Z(\mathcal{L}(\mathcal{V}))$. For $0 \neq v \in \mathcal{V}$ and $\varphi \in \mathcal{V}^*$ with $\varphi(v) = 1$, commutation with $v \otimes \varphi : x \mapsto \varphi(x)v$ gives $Tv = T(v \otimes \varphi)(v) = (v \otimes \varphi)T(v) = \varphi(Tv)v$, so every line is T-stable. Thus $Tv = \lambda_v v$ for $v \neq 0$, and applying this to independent u, v and to $u + v$ gives $\lambda_u = \lambda_v$; hence $T = rI_{\mathcal{V}}$.
Scalars commute with all endomorphisms.

(ii) If $R, S \in \mathcal{L}(\mathcal{V})$ and $T \in \mathcal{F}(\mathcal{V})$, then $\operatorname{im}(RTS) \subseteq R(\operatorname{im} T)$, so $RTS \in \mathcal{F}(\mathcal{V})$ and $\mathcal{F}(\mathcal{V}) \triangleleft \mathcal{L}(\mathcal{V})$. If $0 \neq \mathcal{J} \triangleleft \mathcal{L}(\mathcal{V})$, choose $0 \neq S \in \mathcal{J}$, e with $Se \neq 0$, and $\eta \in \mathcal{V}^*$ with $\eta(Se) = 1$. For any $v \in \mathcal{V}$ and $\varphi \in \mathcal{V}^*$, $v \otimes \varphi = (v \otimes \eta)S(e \otimes \varphi) \in \mathcal{J}$, so $\mathcal{J}$ contains all rank-one maps, hence all finite-rank maps.
(iii) If $\dim \mathcal{V} < \infty$, then $\mathcal{F}(\mathcal{V}) = \mathcal{L}(\mathcal{V})$, so (ii) makes every nonzero two-sided ideal equal to $\mathcal{L}(\mathcal{V})$. If $\dim \mathcal{V} = \infty$, then $0 \neq \mathcal{F}(\mathcal{V}) \neq \mathcal{L}(\mathcal{V})$, so $\mathcal{L}(\mathcal{V})$ is not simple.

Structural Commentary: Rank-one operators are the atoms controlling both centrality and ideals in $\mathcal{L}(\mathcal{V})$.
Key Insight: Sandwich one nonzero map to obtain every rank-one map. □

Definition 3.6.12. *Algebraic Elements and Minimal Polynomials*
Let $\mathcal{A}$ be an $\mathbb{F}$-algebra and $a \in \mathcal{A}$.

(i) The *evaluation map at a* is $\operatorname{ev}_a : \mathbb{F}[x] \to \mathcal{A}$, $p \mapsto p(a)$.
(ii) a is *algebraic over $\mathbb{F}$* iff $\ker \operatorname{ev}_a \neq 0$.
(iii) If a is algebraic, its *minimal polynomial* m_a is the unique monic generator of $\ker \operatorname{ev}_a$, equivalently the monic polynomial of least degree with $m_a(a) = 0$.

Theorem 3.6.13. *Finite-Dimensional Algebras Are Algebraic*
If $\dim_{\mathbb{F}} \mathcal{A} < \infty$, then every $a \in \mathcal{A}$ is algebraic over $\mathbb{F}$. Moreover, the algebraic closure $\overline{\mathbb{Q}}_{\mathbb{C}}$ of $\mathbb{Q}$ inside $\mathbb{C}$ is an infinite-dimensional $\mathbb{Q}$-algebra whose elements are all algebraic.

Proof. Let $\dim_{\mathbb{F}} \mathcal{A} = n$ and $a \in \mathcal{A}$. The $n+1$ vectors $1, a, \dots, a^n$ are linearly dependent, so $\sum_{j=0}^{n} c_j a^j = 0$ for some $(c_0, \dots, c_n) \neq 0$; hence $p(x) = \sum_{j=0}^{n} c_j x^j \neq 0$ satisfies $p(a) = 0$.
If $\alpha, \beta \in \overline{\mathbb{Q}}_{\mathbb{C}}$, then $\mathbb{Q}(\alpha, \beta)/\mathbb{Q}$ is finite, so $\alpha + \beta$, $\alpha\beta$, and $r\alpha$ for $r \in \mathbb{Q}$ are algebraic; hence $\overline{\mathbb{Q}}_{\mathbb{C}}$ is a $\mathbb{Q}$-algebra.
For each $n \geq 1$, $x^n - 2$ is Eisenstein at 2, so it has a root of degree n in $\overline{\mathbb{Q}}_{\mathbb{C}}$; therefore $\dim_{\mathbb{Q}} \overline{\mathbb{Q}}_{\mathbb{C}} = \infty$.
Structural Commentary: Finite-dimensionality forces powers of each element into a linear relation.
Key Insight: Algebraic need not mean finite-dimensional as an algebra. □

Proposition 3.6.14. *The Algebra Generated by One Algebraic Element*
Let $a \in \mathcal{A}$ be algebraic with minimal polynomial m_a. Then $\mathbb{F}[a] := \operatorname{im} \operatorname{ev}_a = \{p(a) : \deg p < \deg m_a\}$ and $\mathbb{F}[a] \cong \mathbb{F}[x]/\langle m_a \rangle$.

Proof. $\operatorname{ev}_a : \mathbb{F}[x] \to \mathcal{A}$ is an $\mathbb{F}$-algebra homomorphism with image $\mathbb{F}[a]$.
Since $\ker \operatorname{ev}_a = \{p : p(a) = 0\}$ and m_a is the monic generator of this kernel, $\ker \operatorname{ev}_a = \langle m_a \rangle$.
Proposition 3.6.9 gives $\mathbb{F}[a] \cong \mathbb{F}[x]/\langle m_a \rangle$. $\forall p \in \mathbb{F}[x]$, by division $p = qm_a + r$, $\deg r < \deg m_a$, hence $p(a) = r(a)$.
Structural Commentary: One algebraic element gives one polynomial quotient.
Key Insight: Evaluation turns polynomial division modulo m_a into normal forms in $\mathbb{F}[a]$. □

Theorem 3.6.15. *Minimal Polynomial, Annihilator, and Invertibility*
Let $a \in \mathcal{A}$ be algebraic and write $m_a(x) = x^n + \sum_{j=0}^{n-1} \alpha_j x^j$.
Then $\operatorname{ann}_{\mathbb{F}[x]}(a) := \{f \in \mathbb{F}[x] : f(a) = 0\} = \langle m_a \rangle$, and $a \in \mathcal{A}^\times \iff \alpha_0 \neq 0$.

Proof. By Definition 3.6.12, $\operatorname{ann}_{\mathbb{F}[x]}(a) = \ker \operatorname{ev}_a = \langle m_a \rangle$.
($\Rightarrow$) If $a \in \mathcal{A}^\times$ and $\alpha_0 = 0$, then $m_a = xq$ with $\deg q < \deg m_a$, and $0 = m_a(a) = aq(a)$ gives $q(a) = 0$, contradicting minimality of m_a; hence $\alpha_0 \neq 0$.
($\Leftarrow$) If $\alpha_0 \neq 0$, then $0 = m_a(a) = a(a^{n-1} + \sum_{j=1}^{n-1} \alpha_j a^{j-1}) + \alpha_0$, so $b := -\alpha_0^{-1}(a^{n-1} + \sum_{j=1}^{n-1} \alpha_j a^{j-1})$ satisfies $ab = 1_{\mathcal{A}}$. Since $b \in \mathbb{F}[a]$, $ab = ba$, so $ba = 1_{\mathcal{A}}$ and $a \in \mathcal{A}^\times$.
Structural Commentary: Invertibility is read from the constant term of the minimal polynomial.
Key Insight: A nonzero constant term solves for a^{-1} as a polynomial in a. □

Definition 3.6.16. *Zero Divisors and Spectrum*
Let $\mathcal{A}$ be an $\mathbb{F}$-algebra and $0 \neq a \in \mathcal{A}$.

(i) a is a *left zero divisor* iff $\exists b \in \mathcal{A} \setminus \{0\}$ with $ab = 0$.
(ii) a is a *right zero divisor* iff $\exists c \in \mathcal{A} \setminus \{0\}$ with $ca = 0$.
(iii) If a is algebraic, *spectrum* is $\operatorname{spec}_{\mathcal{A}}(a) := \{r \in \mathbb{F} : m_a(r) = 0\}$; elements are algebraic *eigenvalues* of a.

Theorem 3.6.17. *Invertible if and only if Not a Zero Divisor*
Let $a \in \mathcal{A}$ be algebraic. Then $a \in \mathcal{A}^\times$ iff a is neither a left nor a right zero divisor.

Proof.
($\Rightarrow$) If $a \in \mathcal{A}^\times$ and $ab = 0$, then $b = a^{-1}ab = 0$; if $ca = 0$, then $c = caa^{-1} = 0$.
Hence a is neither a left nor a right zero divisor.
($\Leftarrow$) Prove the contrapositive. If $a \notin \mathcal{A}^\times$, Theorem 3.6.15 gives $m_a(x) = xp(x)$ with $p \neq 0$ and $\deg p < \deg m_a$.
Then $0 = m_a(a) = ap(a) = p(a)a$, while minimality gives $p(a) \neq 0$. Thus a is both left, right zero divisor.

Structural Commentary: For algebraic elements, noninvertibility forces a polynomial kernel witness.
Key Insight: Factoring m_a by x manufactures the zero-divisor witness $p(a)$. □

Corollary 3.6.18. *One-Sided Inverses Become Two-Sided*
Let $a \in \mathcal{A}$ be algebraic. If $\exists b \in \mathcal{A}$ with $ba = 1_{\mathcal{A}}$ or $ab = 1_{\mathcal{A}}$, then $a \in \mathcal{A}^\times$.

Proof. Assume $ba = 1_{\mathcal{A}}$. If $ac = 0$, then $c = bac = 0$; if $ax = 0$, then $x = bax = 0$. Hence a is neither a right nor a left zero divisor, so Theorem 3.6.17 gives $a \in \mathcal{A}^\times$. The case $ab = 1_{\mathcal{A}}$ is symmetric.
Structural Commentary: Algebraicity upgrades one-sided cancellation to two-sided invertibility.
Key Insight: A one-sided inverse rules out both zero-divisor obstructions. □

Theorem 3.6.19. *Roots Detected by Invertibility*
Let $a \in \mathcal{A}$ be algebraic and $r \in \mathbb{F}$. Then $m_a(r) = 0 \iff a - r1_{\mathcal{A}} \notin \mathcal{A}^\times$. In particular, $a \in \mathcal{A}^\times \iff 0 \notin \operatorname{spec}_{\mathcal{A}}(a)$.

Proof.
($\Rightarrow$) If $m_a(r) = 0$, write $m_a = (x - r)q$ with $\deg q < \deg m_a$.
Then $0 = m_a(a) = (a - r1_{\mathcal{A}})q(a) = q(a)(a - r1_{\mathcal{A}})$ and $q(a) \neq 0$ by minimality, so $a - r1_{\mathcal{A}} \notin \mathcal{A}^\times$.
($\Leftarrow$) If $m_a(r) \neq 0$, then $\gcd(x - r, m_a) = 1$, so $\exists u, v \in \mathbb{F}[x]$ with $u(x)(x - r) + v(x)m_a(x) = 1$.
Evaluating at a gives $u(a)(a - r1_{\mathcal{A}}) = 1_{\mathcal{A}} = (a - r1_{\mathcal{A}})u(a)$, hence $a - r1_{\mathcal{A}} \in \mathcal{A}^\times$, contradiction.
Taking $r = 0$ gives $a \in \mathcal{A}^\times \iff m_a(0) \neq 0 \iff 0 \notin \operatorname{spec}_{\mathcal{A}}(a)$.
Structural Commentary: Roots of m_a are precisely the nonunit scalar shifts of a.
Key Insight: Bézout turns $\gcd(x - r, m_a) = 1$ into an explicit inverse. □

Theorem 3.6.20. *Polynomial Spectral Mapping for Algebraic Elements*
Let $\mathbb{F}$ be algebraically closed, $a \in \mathcal{A}$ algebraic, and $p \in \mathbb{F}[x]$.
Then $p(a)$ is algebraic and $\operatorname{spec}_{\mathcal{A}}(p(a)) = p(\operatorname{spec}_{\mathcal{A}}(a)) := \{p(r) : r \in \operatorname{spec}_{\mathcal{A}}(a)\}$.

Proof. Since $\mathbb{F}[p(a)] \subseteq \mathbb{F}[a]$ and $\dim_{\mathbb{F}} \mathbb{F}[a] < \infty$, $p(a)$ is algebraic. If p is constant, the claim is immediate.
If $\lambda \in \operatorname{spec}_{\mathcal{A}}(p(a))$, then $p(a) - \lambda 1_{\mathcal{A}} \notin \mathcal{A}^\times$. Factor $p(x) - \lambda = c\prod_{j=1}^{m}(x - r_j)^{e_j}$ over $\mathbb{F}$.
Since a product of units is a unit, $a - r_j 1_{\mathcal{A}} \notin \mathcal{A}^\times$ for some j. Theorem 3.6.19 gives $r_j \in \operatorname{spec}_{\mathcal{A}}(a)$ and $\lambda = p(r_j)$.
Conversely, if $r \in \operatorname{spec}_{\mathcal{A}}(a)$, then $p(x) - p(r) = (x - r)q(x)$, so $p(a) - p(r)1_{\mathcal{A}} = (a - r1_{\mathcal{A}})q(a)$ is noninvertible; otherwise $a - r1_{\mathcal{A}}$ has a one-sided inverse and Corollary 3.6.18 makes it invertible. Thus $p(r) \in \operatorname{spec}_{\mathcal{A}}(p(a))$.
Structural Commentary: Polynomial evaluation carries spectra by applying the same polynomial to roots.
Key Insight: Factor $p(x) - \lambda$ and detect a nonunit scalar shift of a. □

Proposition 3.6.21. *Spectra of ab and ba*
If $\dim_{\mathbb{F}} \mathcal{A} < \infty$, $\mathbb{F}$ is algebraically closed, and $a, b \in \mathcal{A}$, then $\operatorname{spec}_{\mathcal{A}}(ab) = \operatorname{spec}_{\mathcal{A}}(ba)$.

Proof. All elements are algebraic by Theorem 3.6.13. Let $r \in \mathbb{F}^\times$.
If $ba - r1_{\mathcal{A}} \in \mathcal{A}^\times$, set $u := a(ba - r1_{\mathcal{A}})^{-1}b - 1_{\mathcal{A}}$; then $(ab - r1_{\mathcal{A}})u = r1_{\mathcal{A}}$, so $ab - r1_{\mathcal{A}}$ has a right inverse, hence is invertible by Corollary 3.6.18. Thus $r \notin \operatorname{spec}_{\mathcal{A}}(ba) \Rightarrow r \notin \operatorname{spec}_{\mathcal{A}}(ab)$, and symmetry gives equality off 0.
If $ba \in \mathcal{A}^\times$ with inverse y, then $b(ay) = 1_{\mathcal{A}}$ and $(yb)a = 1_{\mathcal{A}}$, so $b, a \in \mathcal{A}^\times$ by Corollary 3.6.18; hence $ab \in \mathcal{A}^\times$.
The converse is symmetric, so 0 belongs to both spectra or neither.
Structural Commentary: $\operatorname{spec}(ab) \setminus \{0\} = \operatorname{spec}(ba) \setminus \{0\}$, and algebraicity resolves the 0-case.
Key Insight: The resolvent identity transfers invertibility between $ba - r1_{\mathcal{A}}$ and $ab - r1_{\mathcal{A}}$ for $r \neq 0$. □

Definition 3.6.22. *Division Algebras*
An associative $\mathbb{F}$-algebra $\mathcal{D}$ is a *division algebra* iff $\mathcal{D} \setminus \{0\} \subseteq \mathcal{D}^\times$.

Example 3.6.23. *Hamilton's Quaternions*
Hamilton's quaternion algebra is $\mathbb{H} := \mathbb{R}1 \oplus \mathbb{R}i \oplus \mathbb{R}j \oplus \mathbb{R}k$ with unit 1,
relations $i^2 = j^2 = k^2 = -1$, $ij = k = -ji$, $jk = i = -kj$, and $ki = j = -ik$, extended $\mathbb{R}$-bilinearly.
It is noncommutative since $ij = -ji$. For $q = r_0 + r_1 i + r_2 j + r_3 k$, define $\overline{q} := r_0 - r_1 i - r_2 j - r_3 k$ and $N(q) := q\overline{q} = r_0^2 + r_1^2 + r_2^2 + r_3^2$. If $q \neq 0$, then $N(q) > 0$ and $q^{-1} = N(q)^{-1}\overline{q}$, so $\mathbb{H}$ is a real division algebra.

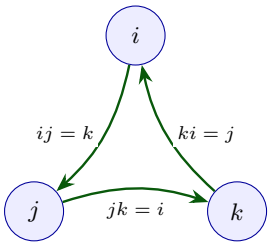


Figure 106: Quaternion basis products: the cycle $i \to j \to k \to i$ gives $ij = k$, $jk = i$, and $ki = j$; reversed order negates the product. Also $i^2 = j^2 = k^2 = -1$, and 1 is the unit.

Theorem 3.6.24. *Division Algebras over Algebraically Closed Fields*
If $\dim_{\mathbb{F}} \mathcal{D} < \infty$, $\mathcal{D}$ is a division $\mathbb{F}$-algebra, and $\mathbb{F}$ is algebraically closed, then $\mathcal{D} = \mathbb{F}1_{\mathcal{D}} \cong \mathbb{F}$.

Proof. Let $a \in \mathcal{D}$. By Theorem 3.6.13, a is algebraic with minimal polynomial m_a.
If $m_a = pq$ with $0 < \deg p, \deg q < \deg m_a$, then $0 = m_a(a) = p(a)q(a)$, while minimality gives $p(a), q(a) \neq 0$; this contradicts $\mathcal{D} \setminus \{0\} \subseteq \mathcal{D}^\times$. Thus m_a is irreducible, hence linear over algebraically closed $\mathbb{F}$: $m_a = x - r$.
Therefore $a = r1_{\mathcal{D}}$, so $\mathcal{D} \subseteq \mathbb{F}1_{\mathcal{D}}$, and the reverse inclusion is the scalar copy.
Structural Commentary: Algebraic closure leaves only scalars.
Key Insight: Irreducible minimal polynomials collapse to linear polynomials. □

Theorem 3.6.25. *Wedderburn's Theorem*
Every finite division algebra $\mathcal{D}$ is commutative; equivalently $\mathcal{D} = Z(\mathcal{D})$, hence $\mathcal{D}$ is a finite field.

Proof. Let $F := Z(\mathcal{D})$, $|F| = z$, $[\mathcal{D} : F] = n$, and $G := \mathcal{D}^\times$.
Suppose $\mathcal{D}$ is noncommutative, and let B be representatives of the noncentral conjugacy classes of G.
For $\beta \in B$, $C_{\mathcal{D}}(\beta)$ is a proper division subalgebra containing F, so $|C_{\mathcal{D}}(\beta)| = z^{k_\beta}$ with $1 < k_\beta < n$ and $k_\beta \mid n$.
By the class equation, Theorem 2.4.2, $z^n - 1 = z - 1 + \sum_{\beta \in B} (z^n - 1)/(z^{k_\beta} - 1)$.
Since $k_\beta \mid n$ and $k_\beta < n$, the primitive nth roots occur in $(x^n - 1)/(x^{k_\beta} - 1)$, so Propositions 3.2.18 and 3.2.19 give $Q_n(z) \mid (z^n - 1)/(z^{k_\beta} - 1)$ for each β, and also $Q_n(z) \mid z^n - 1$.
The class equation therefore gives $Q_n(z) \mid z - 1$, contradicting Lemma 3.2.20, which gives $|Q_n(z)| > z - 1$.
Thus G has no noncentral conjugacy classes, so G is abelian and $\mathcal{D}$ is commutative.
Structural Commentary: The class equation converts noncommutativity into cyclotomic divisibility.
Key Insight: The cyclotomic lower bound is too large to divide $z - 1$. □

Theorem 3.6.26. *Frobenius Classification over $\mathbb{R}$*
If $\dim_{\mathbb{R}} \mathcal{D} < \infty$ and $\mathcal{D}$ is a real division algebra, then $\mathcal{D} \cong \mathbb{R}$, $\mathcal{D} \cong \mathbb{C}$, or $\mathcal{D} \cong \mathbb{H}$.

Proof. For each $d \in \mathcal{D}$, m_d is linear or irreducible quadratic over $\mathbb{R}$; completing the square gives $d = r + u$ with $r \in \mathbb{R}$ and $u^2 \in \mathbb{R}_{\leq 0}$. Set $\mathcal{D}_0 := \{u \in \mathcal{D} : u^2 \in \mathbb{R}_{\leq 0}\}$; the preceding decomposition and uniqueness give $\mathcal{D} = \mathbb{R} \oplus \mathcal{D}_0$. If $\mathcal{D}_0 = 0$, then $\mathcal{D} \cong \mathbb{R}$. If $\dim \mathcal{D}_0 = 1$, choose $i \in \mathcal{D}_0$ with $i^2 = -1$, so $\mathcal{D} = \mathbb{R} \oplus \mathbb{R}i \cong \mathbb{C}$.
If $\dim \mathcal{D}_0 > 1$, use $\langle u, v \rangle := -\frac{1}{2}(uv + vu)$ on $\mathcal{D}_0$ to choose orthonormal $i, j \in \mathcal{D}_0$ with $i^2 = j^2 = -1$; $ij + ji = 0$.
Then $k := ij$ satisfies $k^2 = -1$, $ij = k$, $jk = i$, $ki = j$, and reversed products change sign.
If $\ell \in \mathcal{D}_0$ is orthogonal to i, j, k, the same identities force $\ell = 0$, so $\mathcal{D} = \mathbb{R} \oplus \mathbb{R}i \oplus \mathbb{R}j \oplus \mathbb{R}k \cong \mathbb{H}$.
Structural Commentary: A real division algebra splits into scalars plus a Euclidean pure part.
Key Insight: Orthogonal imaginary units generate exactly $\mathbb{C}$ or $\mathbb{H}$. □

Perspective 3.6.27. *Operator Connection*
$T \in \mathcal{L}(\mathcal{V})$ gives an $\mathbb{F}[x]$-module by $x \cdot v = T(v)$. For any algebra $\mathcal{A}$, the left regular map $L_{\mathcal{A}} : \mathcal{A} \hookrightarrow \mathcal{L}(\mathcal{A})$ realizes $\mathcal{A}$ as a faithful endomorphism algebra.

4 Fields and Galois Theory

4.1 Field Extensions

4.2 Splitting Fields and Normality

4.3 Galois Groups

4.4 Classical Applications

Perspective 4.4.1. *Finite Fields and Characteristic*
If $\operatorname{char}(F) = 2$, then $(\forall a \in F)\ 2a = 0$, so $a = -a$. A finite field has $|F| = p^n$ with p prime and $n \geq 1$; conversely,

each prime power $q = p^n$ determines a unique field $\mathbb{F}_q$ up to isomorphism.
Thus the classification is field-theoretic, while the characteristic constraint is ring-theoretic.

4.5 Scalar Extension in Linear Algebra

Definition 4.5.1. *Complexification of a Real Vector Space*
Let $\mathcal{V}$ be a real vector space.

(i) The *complexification* of $\mathcal{V}$ is $\mathcal{V}_{\mathbb{C}} := \mathcal{V} \times \mathcal{V}$, with $(\mathbf{u}, \mathbf{v})$ denoted $\mathbf{u} + i\mathbf{v}$.
(ii) Addition and $\mathbb{C}$-scalar multiplication are

$$(\mathbf{u}_1 + i\mathbf{v}_1) + (\mathbf{u}_2 + i\mathbf{v}_2) = (\mathbf{u}_1 + \mathbf{u}_2) + i(\mathbf{v}_1 + \mathbf{v}_2),$$
$$(a + bi)(\mathbf{u} + i\mathbf{v}) = (a\mathbf{u} - b\mathbf{v}) + i(b\mathbf{u} + a\mathbf{v}), \qquad a, b \in \mathbb{R}.$$

(iii) *Complexification map* is real-linear embedding $\operatorname{cpx} : \mathcal{V} \hookrightarrow \mathcal{V}_{\mathbb{C}}$; $\operatorname{cpx}(\mathbf{u}) = \mathbf{u} + i\mathbf{0}$, where $\{\mathbf{u} + i\mathbf{0} : \mathbf{u} \in \mathcal{V}\}$.
(iv) For $\mathbf{z} = \mathbf{u} + i\mathbf{v} \in \mathcal{V}_{\mathbb{C}}$, define $\operatorname{Re}\mathbf{z} := \mathbf{u}$ and $\operatorname{Im}\mathbf{z} := \mathbf{v}$. Hence $\mathbf{z} = \mathbf{0}$ iff $\operatorname{Re}\mathbf{z} = \operatorname{Im}\mathbf{z} = \mathbf{0}$.

Proposition 4.5.2. *Complexification Preserves Bases*
Let $\mathcal{V}$ be real. Then $\mathcal{V}_{\mathbb{C}}$ is complex, and every real basis $(\mathbf{v}_1, \dots, \mathbf{v}_n)$ of $\mathcal{V}$ is a complex basis of $\mathcal{V}_{\mathbb{C}}$.
Thus $\dim_{\mathbb{C}}(\mathcal{V}_{\mathbb{C}}) = \dim_{\mathbb{R}}(\mathcal{V})$.

Proof. The vector-space axioms follow componentwise from those of $\mathcal{V}$.
For $\mathbf{u}, \mathbf{v} \in \mathcal{V}$, write $\mathbf{u} = \sum_{j=1}^{n} a_j\mathbf{v}_j$ and $\mathbf{v} = \sum_{j=1}^{n} b_j\mathbf{v}_j$, $a_j, b_j \in \mathbb{R}$. Then $\mathbf{u} + i\mathbf{v} = \sum_{j=1}^{n} (a_j + ib_j)\mathbf{v}_j$, so the list spans.
If $\sum_{j=1}^{n} \lambda_j\mathbf{v}_j = \mathbf{0}$ and $\lambda_j = a_j + ib_j$, then $\sum_{j=1}^{n} a_j\mathbf{v}_j = \mathbf{0}$ and $\sum_{j=1}^{n} b_j\mathbf{v}_j = \mathbf{0}$.
Real linear independence gives $a_j = b_j = 0$ for all j, hence $\lambda_j = 0$ for all j.
Structural Commentary: Complexification keeps the basis length unchanged.
Key Insight: Separate real and imaginary parts. □

Corollary 4.5.3. *Coordinates Survive Complexification*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_n)$ be a real basis of $\mathcal{V}$ and set $\beta_{\mathbb{C}} := \operatorname{cpx}(\beta)$. Then, for all $\mathbf{v} \in \mathcal{V}$, $[\operatorname{cpx}(\mathbf{v})]_{\beta_{\mathbb{C}}} = [\mathbf{v}]_{\beta} \in \mathbb{R}^n \subseteq \mathbb{C}^n$.

Proof. If $[\mathbf{v}]_{\beta} = (a_1, \dots, a_n)^t$, then $a_j \in \mathbb{R}$ and $\operatorname{cpx}(\mathbf{v}) = \mathbf{v} + i\mathbf{0} = \sum_{j=1}^{n} a_j(\mathbf{v}_j + i\mathbf{0}) = \sum_{j=1}^{n} a_j \operatorname{cpx}(\mathbf{v}_j)$.
Thus $[\operatorname{cpx}(\mathbf{v})]_{\beta_{\mathbb{C}}} = (a_1, \dots, a_n)^t = [\mathbf{v}]_{\beta}$.
Structural Commentary: The embedded real copy keeps exactly the same coordinate columns.
Key Insight: Complexification enlarges the scalar field, not the coordinates of real vectors. □

Proposition 4.5.4. *Complexification as a Real Space*
If $\mathcal{V}$ has real basis $(\mathbf{v}_1, \dots, \mathbf{v}_n)$, then $\mathcal{V}_{\mathbb{C}}$, as a real vector space, has basis $(\mathbf{v}_1, \dots, \mathbf{v}_n, i\mathbf{v}_1, \dots, i\mathbf{v}_n)$.
Hence $\dim_{\mathbb{R}}(\mathcal{V}_{\mathbb{C}}) = 2\dim_{\mathbb{R}}(\mathcal{V})$, *and* $\operatorname{cpx} : \mathcal{V} \to \mathcal{V}_{\mathbb{C}}$, $\operatorname{cpx}(\mathbf{v}) = \mathbf{v} + i\mathbf{0}$, *is real-linear and injective.*

Proof. For $\mathbf{u} = \sum a_j\mathbf{v}_j$ and $\mathbf{w} = \sum b_j\mathbf{v}_j$, $\mathbf{u} + i\mathbf{w} = \sum a_j\mathbf{v}_j + \sum b_j(i\mathbf{v}_j)$, so the displayed list spans $\mathcal{V}_{\mathbb{C}}$ over $\mathbb{R}$.
If $\sum a_j\mathbf{v}_j + \sum b_j(i\mathbf{v}_j) = \mathbf{0}$, then Re and Im give $\sum a_j\mathbf{v}_j = \sum b_j\mathbf{v}_j = \mathbf{0}$; all a_j, b_j vanish. Thus this is real basis.
For $r \in \mathbb{R}$, $\operatorname{cpx}(r\mathbf{v} + \mathbf{w}) = r\operatorname{cpx}(\mathbf{v}) + \operatorname{cpx}(\mathbf{w})$, and $\ker(\operatorname{cpx}) = \{\mathbf{0}\}$, so cpx is real-linear and injective.
Structural Commentary: Complexification is two real copies of $\mathcal{V}$.
Key Insight: The real and imaginary parts are independent real copies of the same basis. □

Definition 4.5.5. *Complexification of a Real Operator*
Let $\mathcal{V}$ be a real vector space and let $T \in \mathcal{L}(\mathcal{V})$.
The *complexification* of T is the operator $T_{\mathbb{C}} \in \mathcal{L}(\mathcal{V}_{\mathbb{C}})$ defined by $T_{\mathbb{C}}(\mathbf{u} + i\mathbf{v}) := T\mathbf{u} + iT\mathbf{v}$.

Proposition 4.5.6. *Complexification Preserves Operator Data*
Let $\mathcal{V}$ be a real vector space and let $T \in \mathcal{L}(\mathcal{V})$.

(i) $T_{\mathbb{C}} \in \mathcal{L}(\mathcal{V}_{\mathbb{C}})$.
(ii) $(T_{\mathbb{C}})^k(\mathbf{u} + i\mathbf{v}) = T^k\mathbf{u} + iT^k\mathbf{v}$ *for every integer* $k \geq 0$.
(iii) If β is a real basis of $\mathcal{V}$, then $[T_{\mathbb{C}}]_{\beta \leftarrow \beta} = [T]_{\beta \leftarrow \beta}$, now regarded as a complex matrix.

Proof.

(i) Additivity is immediate. For $\alpha = a + bi \in \mathbb{C}$,

$$T_{\mathbb{C}}\big(\alpha(\mathbf{u} + i\mathbf{v})\big) = T(a\mathbf{u} - b\mathbf{v}) + iT(b\mathbf{u} + a\mathbf{v}) = (aT\mathbf{u} - bT\mathbf{v}) + i(bT\mathbf{u} + aT\mathbf{v}) = \alpha T_{\mathbb{C}}(\mathbf{u} + i\mathbf{v}),$$

so $T_{\mathbb{C}}$ is complex-linear.

(ii) For $k = 0$, the claim is the identity. If it holds for k, then

$$(T_{\mathbb{C}})^{k+1}(\mathbf{u} + i\mathbf{v}) = T_{\mathbb{C}}(T^k\mathbf{u} + iT^k\mathbf{v}) = T^{k+1}\mathbf{u} + iT^{k+1}\mathbf{v}.$$

Thus the formula holds for all $k \geq 0$.

(iii) If $T\mathbf{v}_j = \sum_i a_{ij}\mathbf{v}_i$ with $a_{ij} \in \mathbb{R}$, then $T_{\mathbb{C}}\mathbf{v}_j = \sum_i a_{ij}\mathbf{v}_i$. Hence the coordinate matrix is unchanged.

Structural Commentary: Complexification keeps the coordinate matrix unchanged.
Key Insight: Same real entries; larger scalar field. □



Figure 107: Complexification packages two real vectors: $T_{\mathbb{C}}(\mathbf{u} + i\mathbf{v}) = T\mathbf{u} + iT\mathbf{v}$.

Method 4.5.7. *Complexifying a Real Matrix*
To use complex eigenvalues for a real matrix $A \in \mathrm{Mat}(n, n, \mathbb{R})$:

(i) Regard the same real matrix A as a complex matrix on $\mathbb{C}^n$.
(ii) Find complex eigenpairs or invariant data for this complexified operator.
(iii) If $A(\mathbf{u} + i\mathbf{v}) = (a + bi)(\mathbf{u} + i\mathbf{v})$, separate parts: $A\mathbf{u} = a\mathbf{u} - b\mathbf{v}$, $A\mathbf{v} = b\mathbf{u} + a\mathbf{v}$
(iv) Conclude that $\mathrm{span}(\mathbf{u}, \mathbf{v}) \subseteq \mathbb{R}^n$ is an invariant line or plane.

Theorem 4.5.8. *Invariant Line or Plane over $\mathbb{R}$*
Every operator on a nonzero finite-dimensional real vector space has an invariant subspace of dimension 1 *or* 2.

Proof. Let $\mathcal{V}$ be a nonzero finite-dimensional real vector space and let $T \in \mathcal{L}(\mathcal{V})$.
By Theorem 3.1.23, $T_{\mathbb{C}}$ has an eigenvalue $\lambda = a + bi$ with eigenvector $\mathbf{u} + i\mathbf{v} \neq \mathbf{0}$.
The equation $T_{\mathbb{C}}(\mathbf{u} + i\mathbf{v}) = (a + bi)(\mathbf{u} + i\mathbf{v})$ gives $T\mathbf{u} = a\mathbf{u} - b\mathbf{v}$ and $T\mathbf{v} = b\mathbf{u} + a\mathbf{v}$.
Thus $\mathcal{U} = \mathrm{span}(\mathbf{u}, \mathbf{v})$ is nonzero, has dimension 1 or 2, and is invariant under T.
Structural Commentary: A complex eigenvector decomposes into two real directions that are closed under T.
Key Insight: Nonreal eigenvalues become real invariant planes; real eigenvalues become invariant lines. □

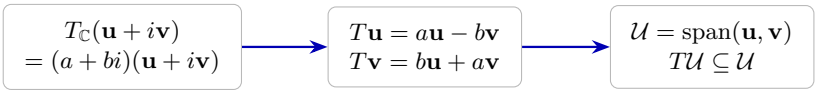


Figure 108: Separate a complex eigenpair into a real invariant plane.

Perspective 4.5.9. *Real Normal Operators Use Planes*
Over $\mathbb{R}$, normal operators need not have enough eigenvectors. The replacement for a missing conjugate eigenpair is a two-dimensional invariant plane carrying a rotation-dilation block.

Proposition 4.5.10. *Two-Dimensional Real Normal Blocks*
Let $\dim(\mathcal{V}) = 2$ *over* $\mathbb{R}$ *and let* $T \in \mathcal{L}(\mathcal{V})$.
Then T *is normal but not self-adjoint if and only if, in every orthonormal basis,* $[T] = \begin{bmatrix} a & -b \\ b & a \end{bmatrix}$, $b \neq 0$.
Moreover, some orthonormal basis makes $b > 0$.

Proof.

($\Rightarrow$) In an orthonormal basis, write $A = [T] = \begin{bmatrix} a & c \\ b & d \end{bmatrix}$.
Normality gives $A^tA = AA^t$. Comparing the $(1,1)$ and $(2,2)$ entries gives $b^2 = c^2$. If $c = b$, then $A = A^t$, impossible because T is not self-adjoint. Thus $c = -b$ and $b \neq 0$. Comparing off-diagonal entries gives $-ab + bd = ab - bd$, so $d = a$. Since the original orthonormal basis was arbitrary, every orthonormal basis has the displayed form.

($\Leftarrow$) A matrix $\begin{bmatrix} a & -b \\ b & a \end{bmatrix}$ with $b \neq 0$ satisfies $A^tA = AA^t$ and is not symmetric. Hence T is normal but not self-adjoint. If $b < 0$, replace the second basis vector by its negative to make $b > 0$.

Structural Commentary: In real dimension 2, non-self-adjoint normality is exactly a scalar plus a rotation block.
Key Insight: Normality forces $c = -b$ and $d = a$. □

Theorem 4.5.11. *Real Normal Operator Block Form*
Let $\mathcal{V}$ be finite-dimensional over $\mathbb{R}$ and let $T \in \mathcal{L}(\mathcal{V})$. Then T is normal if and only if some orthonormal basis makes $[T]$ block diagonal, with each diagonal block either 1-by-1 or of the form $\begin{bmatrix} a & -b \\ b & a \end{bmatrix}$, $b > 0$.

Proof.

($\Rightarrow$) Induct on $n = \dim(\mathcal{V})$. The cases $n \leq 1$ are immediate.
For $n \geq 2$, Theorem 4.5.8 gives an invariant subspace $\mathcal{U}$ with $\dim(\mathcal{U}) \in \{1, 2\}$.
By Corollary 3.4.29, $\mathcal{U}^{\perp}$ is invariant under T and both restrictions are normal.
If $\dim(\mathcal{U}) = 1$, take a unit vector in $\mathcal{U}$. If $\dim(\mathcal{U}) = 2$, use the Real Spectral Theorem 3.4.13 when $T|_{\mathcal{U}}$ is self-adjoint, giving scalar 1-by-1 blocks on $\mathcal{U}$, and use Proposition 4.5.10 otherwise.
By induction, choose the required orthonormal block basis on $\mathcal{U}^{\perp}$. Concatenating it with the basis chosen on $\mathcal{U}$ gives an orthonormal basis of $\mathcal{V}$ because $\mathcal{V} = \mathcal{U} \oplus \mathcal{U}^{\perp}$.

($\Leftarrow$) Each allowed block is normal. If $[T]$ is the block diagonal sum of these blocks, then $[T]^*[T]$ and $[T][T]^*$ are the corresponding block diagonal sums of B^*B and BB^*, hence are equal block by block.

Structural Commentary: Normality lets invariant lines or planes split orthogonally.
Key Insight: Real normal operators split into scalar lines and rotation-dilation planes. □

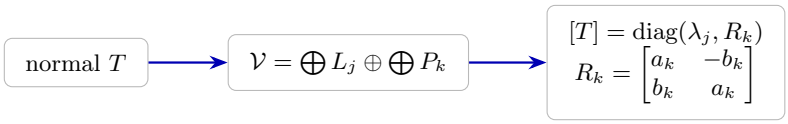


Figure 109: Orthogonal line/plane split into scalar and rotation blocks.

Example 4.5.12. *Plane Rotations*
For $\theta \in \mathbb{R}$, the real matrix $R_\theta = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix}$ is an isometry of $\mathbb{R}^2$. If θ is not an integer multiple of π, then R_θ has no real eigenvalue, which is why real isometries require two-dimensional rotation blocks.

Theorem 4.5.13. *Real Isometry Block Form*
Let $\mathcal{V}$ be finite-dimensional over $\mathbb{R}$ and let $S \in \mathcal{L}(\mathcal{V})$. Then S is an isometry if and only if some orthonormal basis makes $[S]$ block diagonal, with each diagonal block one of $[1]$, $[-1]$, $\begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix}$ with $0 < \theta < \pi$.

Proof.

($\Rightarrow$) Since S is an isometry, $S^*S = I_{\mathcal{V}}$ by Theorem 2.4.26. Finite-dimensional invertibility gives $SS^* = I_{\mathcal{V}}$, so S is normal. Apply Theorem 4.5.11. A 1-by-1 block $[\lambda]$ has $\lambda^2 = 1$, so $\lambda = \pm 1$.
A 2-by-2 block $\begin{bmatrix} a & -b \\ b & a \end{bmatrix}$ has $a^2 + b^2 = 1$ and $b > 0$, so $a = \cos\theta$ and $b = \sin\theta$ for a unique $\theta \in (0, \pi)$.

($\Leftarrow$) Each displayed block B satisfies $B^tB = I$ on its coordinate subspace. Therefore the block diagonal matrix satisfies $[S]^t[S] = I$, so S preserves norms and is an isometry.

Structural Commentary: Isometries are the unit-length special case of real normal blocks.
Key Insight: Scalar blocks reduce to ± 1, and planar normal blocks reduce to rotations. □

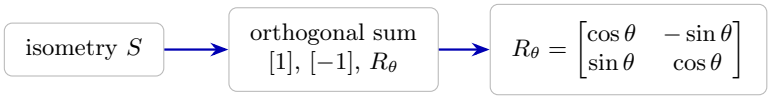


Figure 110: Normal form: fixed/flipped lines and rotation planes.

Proposition 4.5.14. *Real Eigenvalues of a Complexification*
Let $\mathcal{V}$ be real, let $T \in \mathcal{L}(\mathcal{V})$, and let $\lambda \in \mathbb{R}$. Then λ is an eigenvalue of T if and only if it is an eigenvalue of $T_{\mathbb{C}}$.

Proof.

($\Rightarrow$) If $T\mathbf{v} = \lambda\mathbf{v}$ with $\mathbf{v} \neq \mathbf{0}$, then $T_{\mathbb{C}}(\mathbf{v} + i\mathbf{0}) = \lambda(\mathbf{v} + i\mathbf{0})$.

($\Leftarrow$) If $T_{\mathbb{C}}(\mathbf{u} + i\mathbf{v}) = \lambda(\mathbf{u} + i\mathbf{v})$ and $\mathbf{u} + i\mathbf{v} \neq \mathbf{0}$, then $T\mathbf{u} = \lambda\mathbf{u}$ and $T\mathbf{v} = \lambda\mathbf{v}$.
Since $\mathbf{u}$ or $\mathbf{v}$ is nonzero, T has eigenvalue λ.

Structural Commentary: Real eigenvalues survive complexification exactly.
Key Insight: For real λ, real and imaginary parts decouple. □

Lemma 4.5.15. *Conjugation Transfers Generalized Eigenspaces*
Let $\mathcal{V}$ be real, let $T \in \mathcal{L}(\mathcal{V})$, let $\lambda \in \mathbb{C}$, and let $k \geq 0$. Then $(T_{\mathbb{C}} - \lambda I)^k(\mathbf{u}+i\mathbf{v}) = \mathbf{0} \iff (T_{\mathbb{C}} - \overline{\lambda} I)^k(\mathbf{u}-i\mathbf{v}) = \mathbf{0}$.

Proof. Let $\Gamma(\mathbf{u}+i\mathbf{v}) = \mathbf{u} - i\mathbf{v}$.
Then Γ is a conjugate-linear bijection, $\Gamma T_{\mathbb{C}} = T_{\mathbb{C}}\Gamma$, and $\Gamma\big((T_{\mathbb{C}} - \lambda I)^k\mathbf{x}\big) = (T_{\mathbb{C}} - \overline{\lambda} I)^k\Gamma\mathbf{x}$.
($\Rightarrow$) Apply the displayed identity to $\mathbf{x} = \mathbf{u} + i\mathbf{v}$.
($\Leftarrow$) Apply the first direction with $\overline{\lambda}$ in place of λ and $\mathbf{u} - i\mathbf{v}$ in place of $\mathbf{u} + i\mathbf{v}$.
Structural Commentary: Conjugation commutes with $T_{\mathbb{C}}$ and conjugates scalars.
Key Insight: Conjugation pairs generalized eigenspaces. □

Corollary 4.5.16. *Conjugate Eigenvalues Have Equal Multiplicity*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{R}$ and let $T \in \mathcal{L}(\mathcal{V})$.
If λ is an eigenvalue of $T_{\mathbb{C}}$, then $\overline{\lambda}$ is an eigenvalue of $T_{\mathbb{C}}$, and $\dim_{\mathbb{C}} \mathcal{G}_{\lambda}(T_{\mathbb{C}}) = \dim_{\mathbb{C}} \mathcal{G}_{\overline{\lambda}}(T_{\mathbb{C}})$.

Proof. Lemma 4.5.15 with $k = 1$ sends a nonzero eigenvector for λ to a nonzero eigenvector for $\overline{\lambda}$. Thus $\overline{\lambda}$ is an eigenvalue. With $k = \dim_{\mathbb{C}} \mathcal{V}_{\mathbb{C}}$, Theorem 3.2.49 identifies generalized eigenspaces with the corresponding stabilized kernels. By Lemma 4.5.15, $\Gamma(\mathbf{u}+i\mathbf{v}) = \mathbf{u}-i\mathbf{v}$ restricts to conjugate-linear bijection $\mathcal{G}_{\lambda}(T_{\mathbb{C}}) \longrightarrow \mathcal{G}_{\overline{\lambda}}(T_{\mathbb{C}})$. Conjugate-linear bijections preserve complex dimension, so the multiplicities agree.
Structural Commentary: Conjugation mirrors generalized eigenspaces.
Key Insight: Nonreal spectral data of a real operator arrives in equal-sized conjugate pairs. □

Corollary 4.5.17. *Odd-Dimensional Real Operators Have Eigenvalues*
Every operator on a nonzero odd-dimensional real vector space has a real eigenvalue.

Proof. Let $T \in \mathcal{L}(\mathcal{V})$, where $\dim_{\mathbb{R}} \mathcal{V}$ is odd. By Proposition 4.5.2, $\dim_{\mathbb{C}} \mathcal{V}_{\mathbb{C}}$ is odd.
The sum of the multiplicities of all eigenvalues of $T_{\mathbb{C}}$ equals this odd dimension by Proposition 3.2.60. Nonreal eigenvalues contribute in conjugate pairs of equal multiplicity by Corollary 4.5.16, hence contribute an even total. Therefore $T_{\mathbb{C}}$ has at least one real eigenvalue, which is an eigenvalue of T by Proposition 4.5.14.
Structural Commentary: Odd dimension forces a real eigenvalue.
Key Insight: Parity forces at least one real spectral direction. □

Definition 4.5.18. *Characteristic Polynomial of a Real Operator*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be a finite-dimensional real vector space and let $T \in \mathcal{L}(\mathcal{V})$.
The *characteristic polynomial* of T is $q_T := q_{T_{\mathbb{C}}}$, the characteristic polynomial of its complexification.

Theorem 4.5.19. *Characteristic Polynomial over $\mathbb{R}$*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be a finite-dimensional real vector space, let $T \in \mathcal{L}(\mathcal{V})$, and let q_T be its characteristic polynomial.
(i) $q_T \in \mathbb{R}[z]$.
(ii) $\deg q_T = \dim_{\mathbb{R}} \mathcal{V}$.
(iii) The real zeros of q_T are precisely the eigenvalues of T.
(iv) $q_T(T) = 0$.

Proof.
(i) Factor $q_{T_{\mathbb{C}}}$ over $\mathbb{C}$. Real eigenvalues give real linear factors. If $\lambda \notin \mathbb{R}$ has multiplicity m, then $\overline{\lambda}$ has multiplicity m by Corollary 4.5.16, and $(z-\lambda)^m(z-\overline{\lambda})^m = \big(z^2 - 2\operatorname{Re}(\lambda)z + |\lambda|^2\big)^m$, which is real.
(ii) Proposition 3.2.17 and Proposition 4.5.2 give $\deg q_T = \dim_{\mathbb{R}} \mathcal{V}$.
(iii) Real zeros of $q_T = q_{T_{\mathbb{C}}}$ are real eigenvalues of $T_{\mathbb{C}}$ (Prop 3.2.17), hence eigenvalues of T (Prop 4.5.14).
(iv) Cayley–Hamilton over $\mathbb{C}$ gives $q_T(T_{\mathbb{C}}) = 0$. By (i) and Proposition 4.5.6, $q_T(T_{\mathbb{C}}) = (q_T(T))_{\mathbb{C}}$.
If $q_T(T)$ were nonzero on some $\mathbf{u} \in \mathcal{V}$, then $(q_T(T))_{\mathbb{C}}(\mathbf{u} + i\mathbf{0}) = q_T(T)\mathbf{u} + i\mathbf{0}$ would be nonzero, a contradiction. Hence $q_T(T) = 0$.
Structural Commentary: Nonreal factors pair into real quadratics.
Key Insight: Complexify, then pair conjugates. □

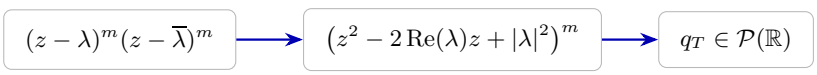


Figure 111: Conjugate factors pair into real quadratics.

Corollary 4.5.20. *q_T Is a Multiple of μ_T over $\mathbb{R}$*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be a finite-dimensional real vector space and let $T \in \mathcal{L}(\mathcal{V})$.
Then $\deg \mu_T \leq \dim_{\mathbb{R}} \mathcal{V}$, and q_T is a polynomial multiple of μ_T in $\mathbb{R}[z]$.

Proof. Theorem 4.5.19 gives $q_T(T) = 0$.
By Theorem 3.2.31, μ_T divides q_T in $\mathbb{R}[z]$. The degree bound follows from $\deg q_T = \dim_{\mathbb{R}} \mathcal{V}$.
Structural Commentary: Cayley–Hamilton makes q_T a real annihilator of T.
Key Insight: The minimal polynomial divides every real annihilator. □

Proposition 4.5.21. *Minimal Polynomial Survives Complexification*
Let $\mathcal{V} \neq \{\mathbf{0}\}$ be finite-dimensional over $\mathbb{R}$ and let $T \in \mathcal{L}(\mathcal{V})$. Then $\mu_{T_{\mathbb{C}}} = \mu_T$ in $\mathbb{C}[z]$.

Proof.
($\Rightarrow$) Since $\mu_T(T) = 0$, Proposition 4.5.6 gives $\mu_T(T_{\mathbb{C}}) = (\mu_T(T))_{\mathbb{C}} = 0$. Hence $\mu_{T_{\mathbb{C}}}$ divides μ_T in $\mathbb{C}[z]$.
($\Leftarrow$) Let $q = \mu_{T_{\mathbb{C}}} = r + is$, with $r, s \in \mathbb{R}[z]$. Since q is monic, r is monic and $\deg r = \deg q$.
For every $\mathbf{u} \in \mathcal{V}$, $0 = q(T_{\mathbb{C}})(\mathbf{u} + i\mathbf{0}) = r(T)\mathbf{u} + is(T)\mathbf{u}$. Thus $r(T) = 0$, so $\deg \mu_T \leq \deg r = \deg q$. The first part gives opposite degree inequality. As q divides μ_T and both are monic with equal degree, $q = \mu_T$.
Structural Commentary: Annihilation passes to complexification and back by real-part extraction.
Key Insight: Evaluate $q(T_{\mathbb{C}})$ on $\mathbf{u} + i\mathbf{0}$. □

5 Tensor, Exterior, and Homological Algebra

5.1 Tensor Products and Universal Properties

Tensor products are the device that turns bilinear data into linear data. The guiding rule is the same one already used for quotient spaces in Theorem 1.1.1: a universal object is recognised by the unique way every compatible map factors through it.

Definition 5.1.1. *Finite-Coordinate Tensor Product Model*
Let $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_m)$ and $\gamma = (\mathbf{w}_1, \dots, \mathbf{w}_n)$ be ordered bases of $\mathcal{V}$ and $\mathcal{W}$.
Define the *coordinate tensor model* $\mathcal{V} \otimes_{\beta,\gamma} \mathcal{W} := \mathrm{Mat}(m, n, \mathbb{F})$ by $\mathbf{v}_i \otimes_{\beta,\gamma} \mathbf{w}_j := E_{i,j}$.
If $\mathbf{v} = \sum_{i=1}^{m} x_i \mathbf{v}_i$ and $\mathbf{w} = \sum_{j=1}^{n} y_j \mathbf{w}_j$, then $\mathbf{v} \otimes_{\beta,\gamma} \mathbf{w} := \sum_{i=1}^{m} \sum_{j=1}^{n} x_i y_j E_{i,j} \in \mathrm{Mat}(m, n, \mathbb{F})$.

Proposition 5.1.2. *Tensor Model Basis and Bilinear Evaluation*
In the coordinate tensor model associated to bases $\beta = (\mathbf{v}_1, \dots, \mathbf{v}_m)$ and $\gamma = (\mathbf{w}_1, \dots, \mathbf{w}_n)$:

(i) the matrices $E_{i,j} = \mathbf{v}_i \otimes_{\beta,\gamma} \mathbf{w}_j$ form a basis;
(ii) the map $(\mathbf{v}, \mathbf{w}) \mapsto \mathbf{v} \otimes_{\beta,\gamma} \mathbf{w}$ is bilinear;
(iii) every bilinear form $B : \mathcal{V} \times \mathcal{W} \to \mathbb{F}$ determines a unique linear functional Θ_B on $\mathrm{Mat}(m, n, \mathbb{F})$ such that $\Theta_B(\mathbf{v} \otimes_{\beta,\gamma} \mathbf{w}) = B(\mathbf{v}, \mathbf{w})$ for all $\mathbf{v} \in \mathcal{V}$ and $\mathbf{w} \in \mathcal{W}$.

Proof.
(i) This is the standard entry-matrix basis of $\mathrm{Mat}(m, n, \mathbb{F})$ from Proposition 1.3.55.
(ii) The coordinates of $a\mathbf{v}_1 + \mathbf{v}_2$ are the same linear combination of the coordinates of $\mathbf{v}_1$ and $\mathbf{v}_2$, and the displayed formula is linear in those coordinates. The same argument applies in the second variable.
(iii) Define $\Theta_B(E_{i,j}) = B(\mathbf{v}_i, \mathbf{w}_j)$ on the basis from (i), and extend linearly by Proposition 1.3.11. If $\mathbf{v} = \sum_i x_i \mathbf{v}_i$ and $\mathbf{w} = \sum_j y_j \mathbf{w}_j$, then $\Theta_B(\mathbf{v} \otimes_{\beta,\gamma} \mathbf{w}) = \sum_{i=1}^{m} \sum_{j=1}^{n} x_i y_j B(\mathbf{v}_i, \mathbf{w}_j) = B(\mathbf{v}, \mathbf{w})$, by bilinearity of B.
Uniqueness follows since the values on basis matrices $E_{i,j}$ are forced by $\mathbf{v}_i, \mathbf{w}_j$.
Structural Commentary: The coordinate tensor model stores each pair of basis directions in its own matrix slot, so a bilinear form can be evaluated by a single linear functional on those slots.
Key Insight: Tensor coordinates turn products of coordinates into linear coordinates. □

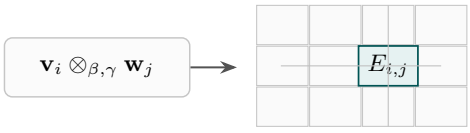


Figure 112: A basis pair occupies the (i, j) matrix slot.

Definition 5.1.3. *Bilinear Maps*
Let $\mathcal{U}, \mathcal{V}, \mathcal{W}$ be $\mathbb{F}$-vector spaces and $f : \mathcal{U} \times \mathcal{V} \to \mathcal{W}$.

(i) f is *bilinear* if $\forall r, s \in \mathbb{F}$, $\mathbf{u}, \mathbf{u}' \in \mathcal{U}$, $\mathbf{v}, \mathbf{v}' \in \mathcal{V}$, $f(r\mathbf{u} + s\mathbf{u}', \mathbf{v}) = rf(\mathbf{u}, \mathbf{v}) + sf(\mathbf{u}', \mathbf{v})$ and $f(\mathbf{u}, r\mathbf{v} + s\mathbf{v}') = rf(\mathbf{u}, \mathbf{v}) + sf(\mathbf{u}, \mathbf{v}')$.

(ii) $\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{W}) := \{f : \mathcal{U}\times\mathcal{V}\to\mathcal{W} : f \text{ bilinear}\}$.
(iii) If $\mathcal{W}=\mathbb{F}$, then $\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathbb{F})$ is the space of *bilinear forms* from Definition 1.5.30.

Example 5.1.4. *Basic Bilinear Maps*

(i) If A is an $\mathbb{F}$-algebra, then $\mu : A\times A\to A$, $\mu(a,b) := ab$, is bilinear.
(ii) Evaluation $\mathrm{ev} : \mathcal{L}(\mathcal{V},\mathcal{W})\times\mathcal{V}\to\mathcal{W}$, $\mathrm{ev}(T,\mathbf{v}) := T\mathbf{v}$, is bilinear.
(iii) If $\varphi\in\mathcal{V}'$ and $\psi\in\mathcal{W}'$, then $B_{\varphi,\psi} : \mathcal{V}\times\mathcal{W}\to\mathbb{F}$, $B_{\varphi,\psi}(\mathbf{v},\mathbf{w}) := \varphi(\mathbf{v})\psi(\mathbf{w})$, is bilinear.

Definition 5.1.5. *Tensor Product by Universal Property*
Let $\mathcal{U},\mathcal{V}$ be $\mathbb{F}$-vector spaces.

(i) A pair $(\mathcal{T},t)$ is a *tensor product* of $\mathcal{U}$ and $\mathcal{V}$ if $\mathcal{T}$ is an $\mathbb{F}$-vector space, $t\in\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{T})$, and for every $\mathbb{F}$-vector space $\mathcal{W}$ and $f\in\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{W})$, $\exists!\widetilde{f}\in\mathcal{L}(\mathcal{T},\mathcal{W})$ such that $f=\widetilde{f}\circ t$.
(ii) The map $\widetilde{f}$ is the *mediating map* for f.
(iii) Write $\mathcal{T}=\mathcal{U}\otimes\mathcal{V}$ and $\mathbf{u}\otimes\mathbf{v} := t(\mathbf{u},\mathbf{v})$; such tensors are *decomposable*.

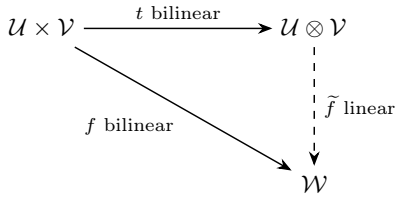


Figure 113: A tensor product is universal for bilinear maps out of $\mathcal{U}\times\mathcal{V}$.

Theorem 5.1.6. *Coordinate-Free Quotient Construction*
Let $\mathcal{U},\mathcal{V}$ be $\mathbb{F}$-vector spaces and let $\mathbb{F}^{(\mathcal{U}\times\mathcal{V})}$ have basis $\{[\mathbf{u},\mathbf{v}] : (\mathbf{u},\mathbf{v})\in\mathcal{U}\times\mathcal{V}\}$.
Let $\mathcal{R}$ be the span of all $r[\mathbf{u}_1,\mathbf{v}]+s[\mathbf{u}_2,\mathbf{v}]-[r\mathbf{u}_1+s\mathbf{u}_2,\mathbf{v}]$; and $r[\mathbf{u},\mathbf{v}_1]+s[\mathbf{u},\mathbf{v}_2]-[\mathbf{u},r\mathbf{v}_1+s\mathbf{v}_2]$, where $r,s\in\mathbb{F}$, $\mathbf{u},\mathbf{u}_1,\mathbf{u}_2\in\mathcal{U}$, and $\mathbf{v},\mathbf{v}_1,\mathbf{v}_2\in\mathcal{V}$. Then $\mathcal{U}\otimes\mathcal{V} := \mathbb{F}^{(\mathcal{U}\times\mathcal{V})}/\mathcal{R}$, $t(\mathbf{u},\mathbf{v})=[\mathbf{u},\mathbf{v}]+\mathcal{R}$ is a tensor product of $\mathcal{U}$ and $\mathcal{V}$.

Proof. Let $\pi : \mathbb{F}^{(\mathcal{U}\times\mathcal{V})}\to\mathcal{U}\otimes\mathcal{V}$ be the quotient map.
Since the generators of $\mathcal{R}$ vanish in the quotient, $t(\mathbf{u},\mathbf{v})=\pi[\mathbf{u},\mathbf{v}]$ is bilinear.
Let $f\in\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{W})$ and define $\alpha : \mathbb{F}^{(\mathcal{U}\times\mathcal{V})}\to\mathcal{W}$ by $\alpha[\mathbf{u},\mathbf{v}] := f(\mathbf{u},\mathbf{v})$.
Bilinearity gives $\mathcal{R}\subseteq\ker\alpha$, so Theorem 1.1.1 gives a unique $\widetilde{f}\in\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathcal{W})$ with $\widetilde{f}\circ\pi=\alpha$. Hence $\widetilde{f}\circ t=f$.
If $L\circ t=f$, then L and $\widetilde{f}$ agree on the spanning set $\{[\mathbf{u},\mathbf{v}]+\mathcal{R}\}$, so $L=\widetilde{f}$.
Structural Commentary: The quotient kills precisely the failure of the raw ordered-pair basis to be bilinear.
Key Insight: A tensor product is a free bilinear map, followed by no extra identifications. □

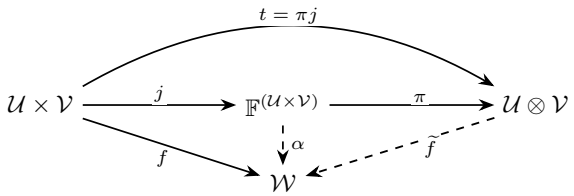


Figure 114: The free linear map α kills $\mathcal{R}$, so it descends uniquely to $\widetilde{f}:\mathcal{U}\otimes\mathcal{V}\to\mathcal{W}$.

Corollary 5.1.7. *Tensor Products Are Unique up to Isomorphism*
If $(\mathcal{T},t)$ and $(\mathcal{S},s)$ are tensor products of $\mathcal{U}$ and $\mathcal{V}$,
then there is a unique isomorphism $\Phi : \mathcal{T}\to\mathcal{S}$ satisfying $\Phi(t(\mathbf{u},\mathbf{v}))=s(\mathbf{u},\mathbf{v})$, where $\mathbf{u}\in\mathcal{U}$, $\mathbf{v}\in\mathcal{V}$.

Proof. Apply universal property of $(\mathcal{T},t)$ to bilinear map s and universal property of $(\mathcal{S},s)$ to bilinear map t.
This gives linear maps $\Phi : \mathcal{T}\to\mathcal{S}$ and $\Psi : \mathcal{S}\to\mathcal{T}$. The composites $\Psi\Phi$ and $I_{\mathcal{T}}$ are both mediating maps for t, while $\Phi\Psi$ and $I_{\mathcal{S}}$ are both mediating maps for s. Uniqueness forces $\Psi\Phi=I_{\mathcal{T}}$ and $\Phi\Psi=I_{\mathcal{S}}$.
Structural Commentary: Universal objects are unique up to canonical isomorphism.
Key Insight: Two universal factorisations force the comparison maps to be inverse. □

Theorem 5.1.8. *Tensor Bases, Dimension, and Coordinate Rank*
Let $\mathcal{U}$ and $\mathcal{V}$ be finite-dimensional with bases $\beta=(\mathbf{u}_1,\dots,\mathbf{u}_m)$ and $\gamma=(\mathbf{v}_1,\dots,\mathbf{v}_n)$.

(i) The tensors $\mathbf{u}_i\otimes\mathbf{v}_j$ form a basis of $\mathcal{U}\otimes\mathcal{V}$.

(ii) $\dim(\mathcal{U}\otimes\mathcal{V}) = \dim(\mathcal{U})\dim(\mathcal{V})$.

(iii) *Every* $\mathbf{z}\in\mathcal{U}\otimes\mathcal{V}$ *has a unique matrix* $R_{\beta,\gamma}(\mathbf{z}) = [r_{ij}] \in \mathrm{Mat}(m,n,\mathbb{F})$ *such that* $\mathbf{z} = \sum_{i=1}^{m}\sum_{j=1}^{n} r_{ij}(\mathbf{u}_i\otimes\mathbf{v}_j)$.
If β',γ' *are other bases, then* $\operatorname{rank} R_{\beta',\gamma'}(\mathbf{z}) = \operatorname{rank} R_{\beta,\gamma}(\mathbf{z})$.

Proof.

(i) The coordinate tensor model of Definition 5.1.1 has the universal property in Definition 5.1.5. Hence Corollary 5.1.7 identifies $\mathcal{U}\otimes\mathcal{V}\cong\mathrm{Mat}(m,n,\mathbb{F})$ with $\mathbf{u}_i\otimes\mathbf{v}_j\mapsto E_{i,j}$, and $\{E_{i,j}\}$ is a basis.
(ii) Counting the mn basis tensors gives $\dim(\mathcal{U}\otimes\mathcal{V}) = mn = \dim(\mathcal{U})\dim(\mathcal{V})$.
(iii) The basis in (i) gives unique coefficients r_{ij}, hence $R_{\beta,\gamma}(\mathbf{z})$. If β',γ' are other bases, then $R_{\beta',\gamma'}(\mathbf{z}) = PR_{\beta,\gamma}(\mathbf{z})Q$ for invertible change-of-basis matrices P,Q, so $\operatorname{rank} R_{\beta',\gamma'}(\mathbf{z}) = \operatorname{rank} R_{\beta,\gamma}(\mathbf{z})$.

Structural Commentary: Tensor rank is coordinate-free matrix rank.
Key Insight: Changing bases performs row and column changes, which preserve matrix rank. □

Definition 5.1.9. *Tensor Coordinate Matrix and Rank*
Let $\mathcal{U},\mathcal{V}$ be finite-dimensional, with bases $\beta = (\mathbf{u}_1,\dots,\mathbf{u}_m)$ and $\gamma = (\mathbf{v}_1,\dots,\mathbf{v}_n)$.

(i) For $\mathbf{z} = \sum_{i=1}^{m}\sum_{j=1}^{n} r_{ij}(\mathbf{u}_i\otimes\mathbf{v}_j)$, define the *coordinate matrix* $R_{\beta,\gamma}(\mathbf{z}) := [r_{ij}]$.
(ii) The *rank of the tensor* $\mathbf{z}$ is $\operatorname{rank}(\mathbf{z}) := \operatorname{rank} R_{\beta,\gamma}(\mathbf{z})$, independent of β,γ by Theorem 5.1.8.

Theorem 5.1.10. *Zero Tests and Decomposable Rank*
Let $\mathcal{U}$ *and* $\mathcal{V}$ *be vector spaces.*

(i) $\forall\mathbf{u}\in\mathcal{U},\mathbf{v}\in\mathcal{V}$, $\mathbf{0}\otimes\mathbf{v} = \mathbf{u}\otimes\mathbf{0} = \mathbf{0}$.

(ii) *If* $\mathbf{u}_1,\dots,\mathbf{u}_r$ *are linearly independent and* $\sum_{i=1}^{r}\mathbf{u}_i\otimes\mathbf{v}_i = \mathbf{0}$, *then* $\mathbf{v}_1 = \cdots = \mathbf{v}_r = \mathbf{0}$.

(iii) *If* $\dim\mathcal{U},\dim\mathcal{V} < \infty$, *then* $\mathbf{z}\neq\mathbf{0}$ *is decomposable* $\iff \operatorname{rank}(\mathbf{z}) = 1$.
More generally, $\operatorname{rank}(\mathbf{z}) = \min\left\{r\geq 0 : \mathbf{z} = \sum_{i=1}^{r}\mathbf{x}_i\otimes\mathbf{y}_i,\ \mathbf{x}_i\in\mathcal{U},\mathbf{y}_i\in\mathcal{V}\right\}$.

Proof.

(i) The tensor map $t:\mathcal{U}\times\mathcal{V}\to\mathcal{U}\otimes\mathcal{V}$ is bilinear, hence linear in each coordinate; therefore $t(\mathbf{0},\mathbf{v}) = t(\mathbf{u},\mathbf{0}) = \mathbf{0}$.
(ii) For each k, extend the dual basis functional on $\operatorname{span}(\mathbf{u}_1,\dots,\mathbf{u}_r)$ to $\varphi_k\in\mathcal{U}'$ with $\varphi_k(\mathbf{u}_i) = \delta_{ki}$. The bilinear map $(\mathbf{u},\mathbf{v})\mapsto\varphi_k(\mathbf{u})\mathbf{v}$ induces $T_k\in\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathcal{V})$.
Applying T_k gives $\mathbf{0} = T_k\left(\sum_{i=1}^{r}\mathbf{u}_i\otimes\mathbf{v}_i\right) = \sum_{i=1}^{r}\delta_{ki}\mathbf{v}_i = \mathbf{v}_k$.
(iii) Choose bases β,γ. If $\mathbf{z} = \mathbf{x}\otimes\mathbf{y}\neq\mathbf{0}$, then $R_{\beta,\gamma}(\mathbf{z}) = [\mathbf{x}]_\beta[\mathbf{y}]_\gamma^{\mathsf{T}}$, so $\operatorname{rank}(\mathbf{z}) = 1$.
Conversely, if $\operatorname{rank} R_{\beta,\gamma}(\mathbf{z}) = r$, write $R_{\beta,\gamma}(\mathbf{z}) = \sum_{i=1}^{r}a_ib_i^{\mathsf{T}}$ with columns a_i,b_i.
Translating a_i,b_i back to vectors $\mathbf{x}_i,\mathbf{y}_i$ gives $\mathbf{z} = \sum_{i=1}^{r}\mathbf{x}_i\otimes\mathbf{y}_i$ as a sum of r decomposable tensors.
Conversely, any sum of s decomposable tensors gives a sum of s rank-one coordinate matrices, so $\operatorname{rank}(\mathbf{z})\leq s$; hence the displayed minimum equals $\operatorname{rank}(\mathbf{z})$.

Structural Commentary: Linear functionals can test one tensor coordinate at a time.
Key Insight: Tensor rank is matrix rank after choosing bases, but the number itself survives changing bases. □

Theorem 5.1.11. *Bilinearity Equals Linearity on a Tensor Product*
For $\mathbb{F}$*-vector spaces* $\mathcal{U},\mathcal{V},\mathcal{W}$, *the map* $\Phi:\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{W})\to\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathcal{W})$, $\Phi(f) = \widetilde{f}$, *where* $\widetilde{f}(\mathbf{u}\otimes\mathbf{v}) = f(\mathbf{u},\mathbf{v})$, *is an isomorphism.*

Proof. The universal property gives existence and uniqueness of $\widetilde{f}$.
Define $\Psi:\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathcal{W})\to\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{W})$ by $\Psi(L) := L\circ t$.
Then $\Psi(\Phi(f)) = \widetilde{f}\circ t = f$, and $\Phi(\Psi(L)) = L$ by uniqueness of the mediating map for $L\circ t$. Hence Φ is bijective.
For $a\in\mathbb{F}$ and f,g bilinear, $\Phi(af+g)(\mathbf{u}\otimes\mathbf{v}) = af(\mathbf{u},\mathbf{v}) + g(\mathbf{u},\mathbf{v}) = (a\Phi(f)+\Phi(g))(\mathbf{u}\otimes\mathbf{v})$, so Φ is linear.
Structural Commentary: Tensor products package bilinear maps as ordinary linear maps.
Key Insight: To define a linear map out of $\mathcal{U}\otimes\mathcal{V}$, define a bilinear rule on $\mathcal{U}\times\mathcal{V}$. □

Method 5.1.12. *Defining Maps on Tensor Products*
To define $L\in\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathcal{W})$:

(i) define a candidate rule $f:\mathcal{U}\times\mathcal{V}\to\mathcal{W}$ by $f(\mathbf{u},\mathbf{v})$;
(ii) prove $f\in\mathrm{Bil}_{\mathbb{F}}(\mathcal{U},\mathcal{V};\mathcal{W})$;

(iii) set $L := \Phi(f)$ from Theorem 5.1.11, so $L(\mathbf{u}\otimes\mathbf{v}) = f(\mathbf{u},\mathbf{v})$.

Do not assign arbitrary values on decomposable tensors: they span $\mathcal{U}\otimes\mathcal{V}$ but are usually linearly dependent.

Theorem 5.1.13. *Tensor Product of Linear Maps*
Let $A\in\mathcal{L}(\mathcal{U},\mathcal{U}')$ and $B\in\mathcal{L}(\mathcal{V},\mathcal{V}')$.
There exists a unique $A\otimes B:\mathcal{U}\otimes\mathcal{V}\to\mathcal{U}'\otimes\mathcal{V}'$ such that $(A\otimes B)(\mathbf{u}\otimes\mathbf{v}) = A\mathbf{u}\otimes B\mathbf{v}$.
Moreover, $(A_2A_1)\otimes(B_2B_1) = (A_2\otimes B_2)(A_1\otimes B_1)$, $I_{\mathcal{U}}\otimes I_{\mathcal{V}} = I_{\mathcal{U}\otimes\mathcal{V}}$.
In finite dimensions, $[A\otimes B]^{\beta'\otimes\gamma'}_{\beta\otimes\gamma} = [A]^{\beta'}_{\beta}\otimes[B]^{\gamma'}_{\gamma}$ for lexicographically ordered tensor-product bases, where the right-hand $\otimes$ is the matrix Kronecker product.

Proof. The map $(\mathbf{u},\mathbf{v})\mapsto A\mathbf{u}\otimes B\mathbf{v}$ is bilinear, so Theorem 5.1.11 gives the unique $A\otimes B$.
The two functorial identities hold on decomposable tensors $((A_2\otimes B_2)(A_1\otimes B_1))(\mathbf{u}\otimes\mathbf{v}) = A_2A_1\mathbf{u}\otimes B_2B_1\mathbf{v}$ and $(I_{\mathcal{U}}\otimes I_{\mathcal{V}})(\mathbf{u}\otimes\mathbf{v}) = \mathbf{u}\otimes\mathbf{v}$, and decomposable tensors span. If $A\mathbf{u}_j = \sum_i a_{ij}\mathbf{u}'_i$ and $B\mathbf{v}_\ell = \sum_k b_{k\ell}\mathbf{v}'_k$, then $(A\otimes B)(\mathbf{u}_j\otimes\mathbf{v}_\ell) = \sum_{i,k} a_{ij}b_{k\ell}(\mathbf{u}'_i\otimes\mathbf{v}'_k)$, which is the Kronecker block-entry rule.
Structural Commentary: Tensoring maps acts independently in the two tensor coordinates.
Key Insight: The mixed product property for Kronecker products is the coordinate shadow of functoriality. □

Theorem 5.1.14. *Tensor Product of Spaces of Linear Maps*
Let $\mathcal{U},\mathcal{V},\mathcal{X},\mathcal{Y}$ be $\mathbb{F}$-vector spaces. The assignment $(A,B)\mapsto A\otimes B$ induces a natural injective map $\Theta:\mathcal{L}(\mathcal{U},\mathcal{X})\otimes\mathcal{L}(\mathcal{V},\mathcal{Y})\to\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathcal{X}\otimes\mathcal{Y})$ with $\Theta(A\otimes B)(\mathbf{u}\otimes\mathbf{v}) = A\mathbf{u}\otimes B\mathbf{v}$.
If $\mathcal{U},\mathcal{V},\mathcal{X},\mathcal{Y}$ are finite-dimensional, then Θ is an isomorphism.
In particular, $\mathcal{U}'\otimes\mathcal{V}'\hookrightarrow(\mathcal{U}\otimes\mathcal{V})'$; $(\varphi\otimes\psi)(\mathbf{u}\otimes\mathbf{v}) = \varphi(\mathbf{u})\psi(\mathbf{v})$ is an isomorphism when $\dim\mathcal{U},\dim\mathcal{V}<\infty$.

Proof.

(i) By Theorem 5.1.13, $(A,B)\mapsto A\otimes B$ is bilinear, hence induces Θ.

(ii) Let $0\neq Z = \sum_{i=1}^{r} A_i\otimes B_i$ with r minimal. Then $A_i\neq 0$ and (B_i) is linearly independent.
Choose $\mathbf{u}$ with some $A_i\mathbf{u}\neq 0$, after reindexing let $A_1\mathbf{u},\dots,A_m\mathbf{u}$ be basis of $\operatorname{span}\{A_i\mathbf{u} : 1\le i\le r\}$.
Absorb remaining coefficients into B_i gives $C_1,\dots,C_m$, not all zero, s.t. $\Theta(Z)(\mathbf{u}\otimes\mathbf{v}) = \sum_{i=1}^{m}(A_i\mathbf{u})\otimes C_i\mathbf{v}$.
If $\Theta(Z)=0$, Theorem 5.1.10 gives $C_i\mathbf{v}=0$ for all $i,\mathbf{v}$, contradiction. Thus $\ker\Theta = 0$.

(iii) If $\mathcal{U},\mathcal{V},\mathcal{X},\mathcal{Y}$ are finite-dimensional, then the domain and codomain of Θ have equal dimension by Theorem 5.1.8, so injectivity implies bijectivity. Taking $\mathcal{X}=\mathcal{Y}=\mathbb{F}$, $\mathcal{L}(\mathcal{U},\mathbb{F}) = \mathcal{U}'$, $\mathcal{L}(\mathcal{U}\otimes\mathcal{V},\mathbb{F}) = (\mathcal{U}\otimes\mathcal{V})'$ gives the dual embedding and finite-dimensional isomorphism.

Structural Commentary: Tensoring maps yields maps on tensor products.
Key Insight: The rule on decomposable tensors is enough, and finite dimensions make the embedding onto. □

Proposition 5.1.15. *Scalar Extension by Tensor Product*
Let $K/\mathbb{F}$ be a field extension and let $\mathcal{V}_{\mathbb{F}}$ be an $\mathbb{F}$-vector space.
Then $K\otimes_{\mathbb{F}}\mathcal{V}_{\mathbb{F}}$ is a K-vector space via $\alpha\cdot(\beta\otimes\mathbf{v}) := (\alpha\beta)\otimes\mathbf{v}$. With $\mu:\mathcal{V}_{\mathbb{F}}\to K\otimes_{\mathbb{F}}\mathcal{V}_{\mathbb{F}}$, $\mu(\mathbf{v}) := 1\otimes\mathbf{v}$, every $\mathbb{F}$-linear $f:\mathcal{V}_{\mathbb{F}}\to\mathcal{Y}$ into a K-vector space $\mathcal{Y}$ admits a unique K-linear $\widetilde{f}:K\otimes_{\mathbb{F}}\mathcal{V}_{\mathbb{F}}\to\mathcal{Y}$ such that $\widetilde{f}\circ\mu = f$.

Proof. For $\alpha\in K$, $(\beta,\mathbf{v})\mapsto(\alpha\beta)\otimes\mathbf{v}$ is $\mathbb{F}$-bilinear, hence induces $M_\alpha\in\mathcal{L}_{\mathbb{F}}(K\otimes_{\mathbb{F}}\mathcal{V}_{\mathbb{F}})$.
On decomposable tensors, $M_\alpha M_\lambda = M_{\alpha\lambda}$, $M_1 = I$, and distributivity holds; since decomposable tensors span, these define a K-vector-space structure.
For f, set $g(\beta,\mathbf{v}) := \beta f(\mathbf{v})$. Then g is $\mathbb{F}$-bilinear, so it induces $\widetilde{f}$ with $\widetilde{f}(\beta\otimes\mathbf{v}) = \beta f(\mathbf{v})$. This is K-linear and gives $\widetilde{f}\circ\mu = f$. If h is K-linear and $h\circ\mu = f$, then $h(\beta\otimes\mathbf{v}) = \beta h(1\otimes\mathbf{v}) = \beta f(\mathbf{v})$, so $h=\widetilde{f}$.
Structural Commentary: Scalar extension tensors in the larger field.
Key Insight: Tensoring with K upgrades which scalars are allowed to act. □

5.2 Symmetric and Exterior Algebras

The previous subsection shows that tensor products are the natural place where bilinear maps become linear. The same idea extends to maps with more variables, and then to tensor spaces, symmetric tensors, exterior tensors, and the determinant.

Definition 5.2.1. *Multilinear Maps and Iterated Tensor Products*
Let $\mathcal{V}_1,\dots,\mathcal{V}_n,\mathcal{W}$ be $\mathbb{F}$-vector spaces.

(i) A map $f:\mathcal{V}_1\times\cdots\times\mathcal{V}_n\to\mathcal{W}$ is *multilinear*, or *n-linear*, if for each k, the map $\mathbf{v}_k\mapsto f(\mathbf{v}_1,\dots,\mathbf{v}_k,\dots,\mathbf{v}_n)$ is linear with the other entries fixed.

(ii) $\operatorname{Mult}_{\mathbb{F}}(\mathcal{V}_1,\dots,\mathcal{V}_n;\mathcal{W}) := \{f:\mathcal{V}_1\times\cdots\times\mathcal{V}_n\to\mathcal{W} : f \text{ multilinear}\}$.

(iii) If $\mathcal{W}=\mathbb{F}$, elements of $\operatorname{Mult}_{\mathbb{F}}(\mathcal{V}_1,\dots,\mathcal{V}_n;\mathbb{F})$ are *multilinear forms*.

(iv) A pair $(\mathcal{T}, t)$ is an *iterated tensor product* of $\mathcal{V}_1, \dots, \mathcal{V}_n$ if $t \in \mathrm{Mult}_{\mathbb{F}}(\mathcal{V}_1, \dots, \mathcal{V}_n; \mathcal{T})$ and, for every $\mathbb{F}$-vector space $\mathcal{Y}$ and $f \in \mathrm{Mult}_{\mathbb{F}}(\mathcal{V}_1, \dots, \mathcal{V}_n; \mathcal{Y})$, $\exists! \widetilde{f} \in \mathcal{L}(\mathcal{T}, \mathcal{Y})$ such that $f = \widetilde{f} \circ t$.
(v) Write $\mathcal{T} = \mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n$ and $\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_n := t(\mathbf{v}_1, \dots, \mathbf{v}_n)$.

Example 5.2.2. *Basic Multilinear Maps*

(i) If A is an associative $\mathbb{F}$-algebra, then $m_n : A^n \to A$, $m_n(a_1, \dots, a_n) := a_1 \cdots a_n$, is n-linear.
(ii) If $\dim \mathcal{V} = n$, then $\det : \mathcal{V}^n \to \mathbb{F}$, applied to coordinate columns in any basis, is alternating n-linear form.

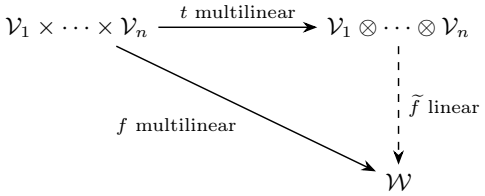


Figure 115: The iterated tensor product is universal for multilinear maps.

Theorem 5.2.3. *Basic Laws for Iterated Tensor Products*
Let $\mathcal{V}_1, \dots, \mathcal{V}_n, \mathcal{V}, \mathcal{W}$ be $\mathbb{F}$-vector spaces.

(i) $\mathrm{Mult}_{\mathbb{F}}(\mathcal{V}_1, \dots, \mathcal{V}_n; \mathcal{W}) \cong \mathcal{L}(\mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n, \mathcal{W})$ *naturally.*
(ii) For $1 \le r < n$, the canonical α_r is $\alpha_r : (\mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_r) \otimes (\mathcal{V}_{r+1} \otimes \cdots \otimes \mathcal{V}_n) \cong \mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n$, where $\alpha_r((\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_r) \otimes (\mathbf{v}_{r+1} \otimes \cdots \otimes \mathbf{v}_n)) = \mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_n$.
(iii) For $\pi \in S_n$, the canonical τ_π is $\tau_\pi : \mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n \cong \mathcal{V}_{\pi(1)} \otimes \cdots \otimes \mathcal{V}_{\pi(n)}$, where $\tau_\pi(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_n) = \mathbf{v}_{\pi(1)} \otimes \cdots \otimes \mathbf{v}_{\pi(n)}$.
(iv) $\lambda : \mathbb{F} \otimes \mathcal{V} \cong \mathcal{V}$, $\lambda(r \otimes \mathbf{v}) = r\mathbf{v}$, and $\rho : \mathcal{V} \otimes \mathbb{F} \cong \mathcal{V}$, $\rho(\mathbf{v} \otimes r) = r\mathbf{v}$.
(v) If $\dim \mathcal{V}_i < \infty$ for all i, then $\dim(\mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n) = \prod_{i=1}^n \dim(\mathcal{V}_i)$.

Proof.

(i) Let $t : \mathcal{V}_1 \times \cdots \times \mathcal{V}_n \to \mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n$ be canonical. Define $\Phi(f) = \widetilde{f}$, where $\widetilde{f}$ is the unique linear map satisfying $\widetilde{f}(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_n) = f(\mathbf{v}_1, \dots, \mathbf{v}_n)$; conversely define $\Psi(T) = T \circ t$. Then $\Psi \circ \Phi$ and $\Phi \circ \Psi$ are identity maps, so Φ is a linear bijection.
(ii) Set $A_r = (\mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_r) \otimes (\mathcal{V}_{r+1} \otimes \cdots \otimes \mathcal{V}_n)$ and $B = \mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n$. The rules defining $\alpha_r : A_r \to B$ and $\beta_r : B \to A_r$, $\beta_r(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_n) = (\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_r) \otimes (\mathbf{v}_{r+1} \otimes \cdots \otimes \mathbf{v}_n)$, are induced by multilinearity; $\alpha_r \circ \beta_r$ and $\beta_r \circ \alpha_r$ fix the spanning decomposable tensors, hence are identities.
(iii) Set $C_\pi = \mathcal{V}_{\pi(1)} \otimes \cdots \otimes \mathcal{V}_{\pi(n)}$. The map $p_\pi : \mathcal{V}_1 \times \cdots \times \mathcal{V}_n \to C_\pi$, $p_\pi(\mathbf{v}_1, \dots, \mathbf{v}_n) = \mathbf{v}_{\pi(1)} \otimes \cdots \otimes \mathbf{v}_{\pi(n)}$, is multilinear, so the universal property induces the unique τ_π with $\tau_\pi \circ t = p_\pi$.
The map $\sigma_\pi : C_\pi \to \mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n$, $\sigma_\pi(\mathbf{v}_{\pi(1)} \otimes \cdots \otimes \mathbf{v}_{\pi(n)}) = \mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_n$, is induced similarly; $\sigma_\pi \tau_\pi$ and $\tau_\pi \sigma_\pi$ fix decomposable spanning tensors, hence are identities.
(iv) The maps $m_L : \mathbb{F} \times \mathcal{V} \to \mathcal{V}$, $m_L(r, \mathbf{v}) = r\mathbf{v}$, and $m_R : \mathcal{V} \times \mathbb{F} \to \mathcal{V}$, $m_R(\mathbf{v}, r) = r\mathbf{v}$, are bilinear, hence induce λ and ρ. Since $\lambda(1 \otimes \mathbf{v}) = \mathbf{v}$ and $1 \otimes r\mathbf{v} = r \otimes \mathbf{v}$, $\lambda^{-1}(\mathbf{v}) = 1 \otimes \mathbf{v}$.
Since $\rho(\mathbf{v} \otimes 1) = \mathbf{v}$ and $r\mathbf{v} \otimes 1 = \mathbf{v} \otimes r$, $\rho^{-1}(\mathbf{v}) = \mathbf{v} \otimes 1$.
(v) Choose finite bases B_i of $\mathcal{V}_i$ and set $\mathcal{B} = \{\mathbf{b}_1 \otimes \cdots \otimes \mathbf{b}_n : \mathbf{b}_i \in B_i\}$. By repeated application of Theorem 5.1.8, $\mathcal{B}$ is a basis; hence $\dim(\mathcal{V}_1 \otimes \cdots \otimes \mathcal{V}_n) = \#\mathcal{B} = \prod_{i=1}^n \#B_i = \prod_{i=1}^n \dim \mathcal{V}_i$.

Structural Commentary: Tensor parenthesising and ordering are canonical bookkeeping.
Key Insight: Universal properties extend pure-tensor rules uniquely. □

Definition 5.2.4. *Tensor Spaces of Type (p, q)*
Let $\dim \mathcal{V} < \infty$.

(i) For $(p, q) \in \mathbb{Z}^2_{\ge 0} \setminus \{(0, 0)\}$, with empty tensor powers equal to $\mathbb{F}$, $\mathcal{T}^p_q(\mathcal{V}) := \mathcal{V}^{\otimes p} \otimes (\mathcal{V}')^{\otimes q}$.
(ii) $\mathcal{T}^p_q(\mathcal{V})$ is the *tensor space of type (p, q)*; p is the *contravariant type*, q the *covariant type*.
(iii) $\mathcal{T}^0_0(\mathcal{V}) := \mathbb{F}$.

Proposition 5.2.5. *Tensor Spaces as Multilinear Functionals*
Let $\dim \mathcal{V} = n < \infty$ and $p, q \in \mathbb{Z}_{\ge 0}$.

(i) $\dim \mathcal{T}^p_q(\mathcal{V}) = n^{p+q}$.
(ii) Let $\iota : \mathcal{V} \to \mathcal{V}''$, $\iota(\mathbf{v})(\eta) = \eta(\mathbf{v})$. Via ι, natural isomorphism is $\Theta_{p,q} : \mathcal{T}^p_q(\mathcal{V}) \xrightarrow{\sim} \mathrm{Mult}_{\mathbb{F}}((\mathcal{V}')^{\times p} \times \mathcal{V}^{\times q}; \mathbb{F})$; where $\Theta_{p,q}(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_p \otimes \varphi_1 \otimes \cdots \otimes \varphi_q)(\eta_1, \dots, \eta_p, \mathbf{x}_1, \dots, \mathbf{x}_q) = \prod_{i=1}^p \eta_i(\mathbf{v}_i) \prod_{j=1}^q \varphi_j(\mathbf{x}_j)$, with empty products equal to 1.

Proof.

(i) By Theorem 5.2.3, $\dim \mathcal{V}' = n$, $\dim \mathcal{T}_q^p(\mathcal{V}) = \dim(\mathcal{V}^{\otimes p}) \dim((\mathcal{V}')^{\otimes q}) = n^p n^q = n^{p+q}$, with $n^0 = 1$.

(ii) Set $X = (\mathcal{V}')^{\otimes p} \otimes \mathcal{V}^{\otimes q}$. Since $\iota : \mathcal{V} \cong \mathcal{V}''$ and all factors are finite-dimensional, $\mathcal{T}_q^p(\mathcal{V}) = \mathcal{V}^{\otimes p} \otimes (\mathcal{V}')^{\otimes q} \cong ((\mathcal{V}')')^{\otimes p} \otimes \mathcal{V}'^{\otimes q} \cong X' \cong \mathrm{Mult}_{\mathbb{F}}((\mathcal{V}')^{\times p} \times \mathcal{V}^{\times q}; \mathbb{F})$. On decomposables, the composite is $\Theta_{p,q}$; spanning gives uniqueness.

Structural Commentary: Tensor elements are multilinear functionals after finite-dimensional duality.
Key Insight: Choosing which factors are active turns tensor factors into input slots. □

Definition 5.2.6. *Contraction*
Let $p, q \geq 1$. The *contraction* is the unique linear map $C : \mathcal{T}_q^p(\mathcal{V}) \to \mathcal{T}_{q-1}^{p-1}(\mathcal{V})$, where $C(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_p \otimes \varphi_1 \otimes \cdots \otimes \varphi_q) = \varphi_1(\mathbf{v}_1)\mathbf{v}_2 \otimes \cdots \otimes \mathbf{v}_p \otimes \varphi_2 \otimes \cdots \otimes \varphi_q$.
For (i, j), C_{ij} is C conjugated by the factor permutation placing $\mathbf{v}_i, \varphi_j$ first.

Example 5.2.7. *Rank-One Operators and Trace*
Under $\mathcal{V} \otimes \mathcal{V}' \cong \mathcal{L}(\mathcal{V})$, $\mathbf{v} \otimes \varphi \mapsto T_{\mathbf{v},\varphi}$, where $T_{\mathbf{v},\varphi}(\mathbf{x}) = \varphi(\mathbf{x})\mathbf{v}$. Hence $C(\mathbf{v} \otimes \varphi) = \varphi(\mathbf{v}) = \mathrm{tr}(T_{\mathbf{v},\varphi})$.

Definition 5.2.8. *Symmetric, Skew-Symmetric, and Alternating Maps*
Let $f : \mathcal{V}^{\times n} \to \mathcal{W}$ be multilinear.

(i) f is *symmetric* if interchanging any two coordinate positions does not change its value.
(ii) f is *skew-symmetric* if interchanging any two coordinate positions multiplies its value by -1.
(iii) f is *alternating* if $f(\mathbf{v}_1, \ldots, \mathbf{v}_n) = 0$ whenever two coordinate entries are equal.

If $\mathrm{char}(\mathbb{F}) \neq 2$, alternating and skew-symmetric are equivalent.
If $\mathrm{char}(\mathbb{F}) = 2$, skew-symmetric is the same as symmetric, and alternating is the stronger condition.

Definition 5.2.9. *Permutation Action and Symmetric/Exterior Tensor Spaces*
Let $\dim \mathcal{V} < \infty$, $p \in \mathbb{Z}_{\geq 0}$, $\mathcal{V}^{\otimes 0} = \mathbb{F}$, $T_p := \mathcal{V}^{\otimes p}$, and $G_p := \mathfrak{S}_p$.

(i) For $\sigma \in G_p$, the *permutation action* is $\lambda_\sigma : T_p \to T_p$, determined by $\lambda_\sigma(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_p) = \mathbf{v}_{\sigma(1)} \otimes \cdots \otimes \mathbf{v}_{\sigma(p)}$.
(ii) The *symmetric tensors* are $\mathrm{ST}^p(\mathcal{V}) := \{\mathbf{z} \in T_p : \forall \sigma \in G_p, \lambda_\sigma \mathbf{z} = \mathbf{z}\}$.
(iii) The *antisymmetric tensors* are $\mathrm{AT}^p(\mathcal{V}) := \{\mathbf{z} \in T_p : \forall \sigma \in G_p, \lambda_\sigma \mathbf{z} = \mathrm{sign}(\sigma)\mathbf{z}\}$.
(iv) Set $R_\vee := \mathrm{span}\{\lambda_\sigma \mathbf{z} - \mathbf{z} : \sigma \in G_p, \mathbf{z} \in T_p\}$.
(v) The *symmetric power* is $\mathrm{Sym}^p(\mathcal{V}) := T_p / R_\vee$.
(vi) Set $A_p := \mathrm{span}\{\lambda_\sigma \mathbf{z} - \mathrm{sign}(\sigma)\mathbf{z} : \sigma \in G_p, \mathbf{z} \in T_p\}$ and $D_p := \mathrm{span}\{\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_p : \exists i < j, \mathbf{v}_i = \mathbf{v}_j\}$. The *exterior power* is $\bigwedge^p \mathcal{V} := T_p / (A_p + D_p)$.
(vii) For quotient map $q_\vee$, the *symmetric product* is $q_\vee(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_p) = \mathbf{v}_1 \vee \cdots \vee \mathbf{v}_p$
(viii) For quotient map $q_\wedge$, the *exterior product* is $q_\wedge(\mathbf{v}_1 \otimes \cdots \otimes \mathbf{v}_p) = \mathbf{v}_1 \wedge \cdots \wedge \mathbf{v}_p$.

Theorem 5.2.10. *Universal Properties of Symmetric and Exterior Powers*
Let $\mathcal{U}$ be a vector space, $M = \mathrm{Mult}_{\mathbb{F}}(\mathcal{V}, \ldots, \mathcal{V}; \mathcal{U})$, and $f^\sigma(\mathbf{v}_1, \ldots, \mathbf{v}_p) := f(\mathbf{v}_{\sigma(1)}, \ldots, \mathbf{v}_{\sigma(p)})$.

(i) $\mathcal{L}(\mathrm{Sym}^p(\mathcal{V}), \mathcal{U}) \cong M_\vee$, where $M_\vee := \{f \in M : \forall \sigma \in \mathfrak{S}_p,\ f^\sigma = f\}$.
(ii) $\mathcal{L}(\bigwedge^p \mathcal{V}, \mathcal{U}) \cong M_\wedge$, where $M_\wedge := \{f \in M : \forall \sigma \in \mathfrak{S}_p,\ f^\sigma = \mathrm{sign}(\sigma) f,\ \forall (\mathbf{v}_1, \ldots, \mathbf{v}_p), (\exists i < j, \mathbf{v}_i = \mathbf{v}_j) \Rightarrow f(\mathbf{v}_1, \ldots, \mathbf{v}_p) = 0\}$.

Proof.

(i) Let $t : \mathcal{V}^{\times p} \to \mathcal{V}^{\otimes p}$ and $q_\vee : \mathcal{V}^{\otimes p} \to \mathrm{Sym}^p(\mathcal{V})$ be canonical. The map $\Phi_\vee(T) = T q_\vee t$ sends $\mathcal{L}(\mathrm{Sym}^p(\mathcal{V}), \mathcal{U})$ into $M_\vee$. Conversely, for $f \in M_\vee$, there is a unique $\widetilde{f} : \mathcal{V}^{\otimes p} \to \mathcal{U}$ with $f = \widetilde{f} t$; since $\widetilde{f}(R_\vee) = 0$, it factors uniquely through $q_\vee$ as $\widetilde{f} = \overline{f} q_\vee$, where $\overline{f} : \mathrm{Sym}^p(\mathcal{V}) \to \mathcal{U}$. Thus $f \mapsto \overline{f}$ is $\Phi_\vee^{-1}$.

(ii) Let $t : \mathcal{V}^{\times p} \to \mathcal{V}^{\otimes p}$ and $q_\wedge : \mathcal{V}^{\otimes p} \to \bigwedge^p \mathcal{V}$ be canonical. The map $\Phi_\wedge(T) = T q_\wedge t$ sends $\mathcal{L}(\bigwedge^p \mathcal{V}, \mathcal{U})$ into $M_\wedge$. Conversely, for $f \in M_\wedge$, the induced $\widetilde{f} : \mathcal{V}^{\otimes p} \to \mathcal{U}$ satisfies $\widetilde{f}(R_\wedge) = 0$, hence factors uniquely as $\widetilde{f} = \overline{f} q_\wedge$; thus $f \mapsto \overline{f}$ is $\Phi_\wedge^{-1}$.

Structural Commentary: Symmetric and exterior powers are tensor quotients by permutation relations.
Key Insight: Symmetric and alternating multilinear maps are ordinary linear maps out of the right quotient. □

Theorem 5.2.11. *Dimensions of Symmetric and Exterior Powers*
Let $\dim \mathcal{V} = n < \infty$ and $p \in \mathbb{Z}_{\geq 0}$. Set $I_\vee = \{1 \leq i_1 \leq \cdots \leq i_p \leq n\}$ and $I_\wedge = \{1 \leq i_1 < \cdots < i_p \leq n\}$.

(i) $\dim \mathrm{Sym}^p(\mathcal{V}) = \#I_\vee = \binom{n+p-1}{p}$.
(ii) $\dim \bigwedge^p \mathcal{V} = \#I_\wedge = \binom{n}{p}$, with $\binom{n}{p} = 0$ for $p > n$.

Proof. Let $E = (\mathbf{e}_1, \dots, \mathbf{e}_n)$ be a basis. For $I = (i_1, \dots, i_p)$, set $\mathbf{e}_I^\vee = \mathbf{e}_{i_1} \vee \cdots \vee \mathbf{e}_{i_p}$ and $\mathbf{e}_I^\wedge = \mathbf{e}_{i_1} \wedge \cdots \wedge \mathbf{e}_{i_p}$.

(i) $R_\vee$ identifies each word with its nondecreasing representative. Thus $B_\vee = \{\mathbf{e}_I^\vee : I \in I_\vee\}$ is the quotient basis, and $\dim \mathrm{Sym}^p(\mathcal{V}) = \#B_\vee = \#I_\vee = \binom{n+p-1}{p}$.

(ii) $R_\wedge$ kills repeated words and identifies each remaining word with its increasing representative up to sign. Thus $B_\wedge = \{\mathbf{e}_I^\wedge : I \in I_\wedge\}$ is the quotient basis, and $\dim \bigwedge^p \mathcal{V} = \#B_\wedge = \#I_\wedge = \binom{n}{p}$.

Structural Commentary: Symmetric powers count repetitions; exterior powers forbid repetitions.
Key Insight: Multisets give symmetric dimensions, subsets give exterior dimensions. □

Proposition 5.2.12. *Symmetrization in Characteristic Zero*
Assume $\mathrm{char}(\mathbb{F}) = 0$. *Define* $S, A \in \mathcal{L}(\mathcal{V}^{\otimes p})$ *by* $S := \frac{1}{p!} \sum_{\sigma \in \mathfrak{S}_p} \lambda_\sigma$ *and* $A := \frac{1}{p!} \sum_{\sigma \in \mathfrak{S}_p} \mathrm{sign}(\sigma) \lambda_\sigma$.

(i) $S^2 = S$ *and* $\operatorname{im} S = \mathrm{ST}^p(\mathcal{V})$.
(ii) $A^2 = A$ *and* $\operatorname{im} A = \mathrm{AT}^p(\mathcal{V})$.

Proof.

(i) For $\tau \in \mathfrak{S}_p$, $\lambda_\tau S = p!^{-1} \sum_\sigma \lambda_{\tau\sigma} = S$, hence $\operatorname{im} S \subseteq \mathrm{ST}^p(\mathcal{V})$.
If $\mathbf{z} \in \mathrm{ST}^p(\mathcal{V})$, then $S\mathbf{z} = \mathbf{z}$, so $\operatorname{im} S = \mathrm{ST}^p(\mathcal{V})$ and $S^2 = S$.

(ii) For $\tau \in \mathfrak{S}_p$, $\lambda_\tau A = p!^{-1} \sum_\sigma \mathrm{sign}(\sigma) \lambda_{\tau\sigma} = \mathrm{sign}(\tau) A$, hence $\operatorname{im} A \subseteq \mathrm{AT}^p(\mathcal{V})$.
If $\mathbf{z} \in \mathrm{AT}^p(\mathcal{V})$, then $A\mathbf{z} = p!^{-1} \sum_\sigma \mathrm{sign}(\sigma)^2 \mathbf{z} = \mathbf{z}$, so $\operatorname{im} A = \mathrm{AT}^p(\mathcal{V})$ and $A^2 = A$.

Structural Commentary: Averaging gives symmetry projections.
Key Insight: The factor $1/p!$ makes the averaging map fix the tensors it projects onto. □

Definition 5.2.13. *Tensor, Symmetric, and Exterior Algebras*
Let $\dim \mathcal{V} < \infty$.

(i) The *tensor algebra* is $\mathcal{T}(\mathcal{V}) := \bigoplus_{p \geq 0} \mathcal{V}^{\otimes p}$, with $\mathcal{V}^{\otimes 0} := \mathbb{F}$ and product induced by tensor concatenation.
(ii) The *symmetric algebra* is $\mathrm{Sym}(\mathcal{V}) := \bigoplus_{p \geq 0} \mathrm{Sym}^p(\mathcal{V})$, with product induced by $\vee$.
(iii) The *exterior algebra* is $\bigwedge \mathcal{V} := \bigoplus_{0 \leq p \leq \dim \mathcal{V}} \bigwedge^p \mathcal{V}$, with product induced by $\wedge$.
(iv) Each is *graded*: degree p times degree q lies in degree $p + q$.

Corollary 5.2.14. *Determinant from the Exterior Universal Property*
Let $\dim \mathcal{V} = n$, $E = (\mathbf{e}_1, \dots, \mathbf{e}_n)$, $\Omega_E := \mathbf{e}_1 \wedge \cdots \wedge \mathbf{e}_n$, *and* $A_E(\mathbf{v}_1, \dots, \mathbf{v}_n) = ([\mathbf{v}_j]_i)_{i,j}$.

(i) There is a unique $\ell_E \in (\bigwedge^n \mathcal{V})'$ *with* $\ell_E(\Omega_E) = 1$*; hence* $d_E(\mathbf{v}_1, \dots, \mathbf{v}_n) := \ell_E(\mathbf{v}_1 \wedge \cdots \wedge \mathbf{v}_n)$ *is the unique alternating form satisfying* $d_E(\mathbf{e}_1, \dots, \mathbf{e}_n) = 1$.

(ii) For all $\mathbf{v}_1, \dots, \mathbf{v}_n \in \mathcal{V}$, $d_E(\mathbf{v}_1, \dots, \mathbf{v}_n) = \sum_{\sigma \in \mathfrak{S}_n} \mathrm{sign}(\sigma) \prod_{j=1}^n [\mathbf{v}_j]_{\sigma(j)} = \det A_E(\mathbf{v}_1, \dots, \mathbf{v}_n)$.

Proof.

(i) By Theorem 5.2.11, $\bigwedge^n \mathcal{V} = \mathbb{F}\Omega_E$, so $\ell_E(\Omega_E) = 1$ fixes a unique $\ell_E \in (\bigwedge^n \mathcal{V})'$; via Theorem 5.2.10, ℓ_E corresponds to $d_E(\mathbf{v}_1, \dots, \mathbf{v}_n) = \ell_E(\mathbf{v}_1 \wedge \cdots \wedge \mathbf{v}_n)$, hence d_E is unique.

(ii) Let $\mathbf{v}_j = \sum_{i=1}^n a_{ij} \mathbf{e}_i$, $a_{ij} = [\mathbf{v}_j]_i$. Repeated index tuples have zero wedge; distinct tuples are $(\sigma(1), \dots, \sigma(n))$, with wedge $\mathrm{sign}(\sigma)\Omega_E$. Hence $\mathbf{v}_1 \wedge \cdots \wedge \mathbf{v}_n = \sum_{\sigma \in \mathfrak{S}_n} \mathrm{sign}(\sigma) \prod_{j=1}^n a_{\sigma(j),j} \Omega_E$, and

$$d_E(\mathbf{v}_1, \dots, \mathbf{v}_n) = \ell_E(\mathbf{v}_1 \wedge \cdots \wedge \mathbf{v}_n) = \sum_{\sigma \in \mathfrak{S}_n} \mathrm{sign}(\sigma) \prod_{j=1}^n a_{\sigma(j),j}$$
$$= \det(a_{ij}) = \det A_E(\mathbf{v}_1, \dots, \mathbf{v}_n).$$

Structural Commentary: Determinant is the basis-normalised alternating form.
Key Insight: Top exterior dimension one turns normalisation into determinant. □

Corollary 5.2.15. *Determinant Properties from Exterior Powers*
Let $A, B \in \mathrm{Mat}(n, n, \mathbb{F})$, *and set* $d(A) := d_E(A_{\bullet 1}, \dots, A_{\bullet n})$ *for* E *the standard basis of* $\mathbb{F}^n$.

(i) $d(A^t) = d(A)$.
(ii) $d(AB) = d(A)d(B)$.
(iii) $A \in \mathrm{GL}(n, \mathbb{F}) \iff d(A) \neq 0$.

Proof. Let $\omega = \mathbf{e}_1 \wedge \cdots \wedge \mathbf{e}_n$ and let T_A denote the linear map with matrix A.

(i) By Corollary 5.2.14, $d(A^t) = \sum_\sigma \mathrm{sign}(\sigma) \prod_{j=1}^n A_{j,\sigma(j)}$. With $\tau = \sigma^{-1}$, $\sum_\tau \mathrm{sign}(\tau) \prod_{i=1}^n A_{\tau(i),i} = d(A)$.

(ii) Since $(\bigwedge^n T_A)\omega = d(A)\omega$ and $T_{AB} = T_A \circ T_B$, applying $\bigwedge^n (T_A \circ T_B) = (\bigwedge^n T_A)(\bigwedge^n T_B)$ to ω gives $d(AB)\omega = d(A)d(B)\omega$, hence $d(AB) = d(A)d(B)$.

(iii) ($\Rightarrow$) If $A \in \mathrm{GL}(n, \mathbb{F})$, then $I = AA^{-1}$, so $1 = d(I) = d(A)d(A^{-1})$ by (ii); hence $d(A) \neq 0$.

(iii) ($\Leftarrow$) Let $C_j = A_{\bullet j}$. Suppose $\sum_{j=1}^n c_j C_j = 0$ and $c_k \neq 0$. Then $C_k = -c_k^{-1} \sum_{j\neq k} c_j C_j$, so multilinearity in the kth slot gives $d(A) = d(C_1, \ldots, C_k, \ldots, C_n) = -c_k^{-1} \sum_{j\neq k} c_j d(C_1, \ldots, C_{k-1}, C_j, C_{k+1}, \ldots, C_n) = 0$, since each summand has equal jth and kth entries.
Thus $d(A) \neq 0$ implies $C_1, \ldots, C_n$ are independent, hence $A \in \mathrm{GL}(n, \mathbb{F})$.

Structural Commentary: Determinant laws are top-exterior scalar laws.
Key Insight: Top exterior dimension one makes determinants multiplicative. □

5.3 Multilinear Algebra Revisited

Definition 5.3.1. *Double Dual Space*
The *double dual space* of $\mathcal{V}$ is $\mathcal{V}'' := (\mathcal{V}')'$. Define $\Lambda_{\mathcal{V}} : \mathcal{V} \to \mathcal{V}''$ by $(\Lambda_{\mathcal{V}}\mathbf{v})(\varphi) = \varphi(\mathbf{v}), \mathbf{v} \in \mathcal{V},\ \varphi \in \mathcal{V}'$.

Theorem 5.3.2. *Natural Map to the Double Dual*
The map $\Lambda_{\mathcal{V}} : \mathcal{V} \to \mathcal{V}''$ satisfies:

(i) $\Lambda_{\mathcal{V}}$ is linear;
(ii) if $T \in \mathcal{L}(\mathcal{V}, \mathcal{W})$ and $T'' = (T')'$, then $T''\Lambda_{\mathcal{V}} = \Lambda_{\mathcal{W}} T$;
(iii) if $\mathcal{V}$ is finite-dimensional, then $\Lambda_{\mathcal{V}}$ is an isomorphism.

Proof.

(i) For fixed $\mathbf{v} \in \mathcal{V}$, the map $\varphi \mapsto \varphi(\mathbf{v})$ is linear on $\mathcal{V}'$, so $\Lambda_{\mathcal{V}}\mathbf{v} \in \mathcal{V}''$.
For $\mathbf{u}, \mathbf{v} \in \mathcal{V}$, $a \in \mathbb{F}$, and $\varphi \in \mathcal{V}'$, $(\Lambda_{\mathcal{V}}(a\mathbf{u} + \mathbf{v}))(\varphi) = \varphi(a\mathbf{u} + \mathbf{v}) = a(\Lambda_{\mathcal{V}}\mathbf{u})(\varphi) + (\Lambda_{\mathcal{V}}\mathbf{v})(\varphi)$, so $\Lambda_{\mathcal{V}}$ is linear.
(ii) If $\psi \in \mathcal{W}'$, then $\left(T''(\Lambda_{\mathcal{V}}\mathbf{v})\right)(\psi) = (\Lambda_{\mathcal{V}}\mathbf{v})(T'\psi) = (T'\psi)(\mathbf{v}) = \psi(T\mathbf{v}) = \left(\Lambda_{\mathcal{W}}(T\mathbf{v})\right)(\psi)$; $T''\Lambda_{\mathcal{V}} = \Lambda_{\mathcal{W}} T$.
(iii) Suppose $\mathcal{V}$ is finite-dimensional. If $\mathbf{v} \neq \mathbf{0}$, extend $(\mathbf{v})$ to a basis of $\mathcal{V}$ and let φ be the first functional of the corresponding dual basis. Then $\varphi(\mathbf{v}) = 1$, so $(\Lambda_{\mathcal{V}}\mathbf{v})(\varphi) = 1$ and $\Lambda_{\mathcal{V}}\mathbf{v} \neq 0$. Hence $\ker(\Lambda_{\mathcal{V}}) = \{\mathbf{0}\}$ and $\Lambda_{\mathcal{V}}$ is injective. Since $\dim(\mathcal{V}'') = \dim(\mathcal{V}') = \dim(\mathcal{V})$, Corollary 1.3.30 shows that $\Lambda_{\mathcal{V}}$ is an isomorphism.

Structural Commentary: The double-dual map sends a vector to its evaluation action on functionals. It is natural because it commutes with every linear map.
Key Insight: Finite-dimensional vectors are completely recovered by how all linear functionals evaluate them. □

Perspective 5.3.3. *Why the Double Dual Is Natural*
A finite-dimensional space $\mathcal{V}$ is isomorphic to $\mathcal{V}'$, but such an isomorphism usually requires choosing a basis. The map $\Lambda_{\mathcal{V}} : \mathcal{V} \to \mathcal{V}''$ requires no choice of basis.

5.4 Chain Complexes and Homology

5.5 Derived Functors Preview

Part IV
Algebraic Geometry and Commutative Algebra

1 Commutative Algebra

1.1 Prime Ideals and Spectra
1.2 Localization and Local Rings
1.3 Noetherian Rings and Modules
1.4 Integral Dependence and Dimension
1.5 Regularity and Singularities Preview

2 Algebraic Geometry

2.1 Affine Algebraic Sets
2.2 Morphisms and Coordinate Rings
2.3 Projective Geometry
2.4 Schemes
2.5 Sheaves and Cohomology Preview
2.6 Smoothness, Tangent Spaces, and Dimension
2.7 Divisors, Line Bundles, and Curves

Part V
Analysis and Continuous Structure

1 Foundations of Analysis

Foundations of Analysis builds on the global mathematical language prerequisites, then fixes number systems, cardinality tools, and real sequences before topology begins.

1.1 Number Systems and Completeness

Number systems supply the ordered and algebraic background for Analysis: $\mathbb{Q}$ has gaps, $\mathbb{R}$ repairs them by completeness, and $\mathbb{C}$ keeps field algebra without order.

Proposition 1.1.1. *The Rational Gap at $\sqrt{2}$*
There is no rational p such that $p^2 = 2$.

Proof. Let $p = m/n$. Then $m^2 = 2n^2$, so m is even; $m = 2k$. Hence $n^2 = 2k^2$, so n is even, a contradiction.
Structural Commentary: This exhibits a gap in $\mathbb{Q}$.
Key Insight: $\sqrt{2}$ is not rational. □

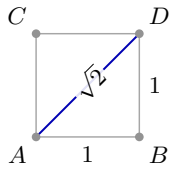


Figure 116: The diagonal of a unit square has length $\sqrt{2}$, so the rational gap is geometric as well as algebraic.

Example 1.1.2. *No Largest Lower Side and No Smallest Upper Side*
Let $A := \{p \in \mathbb{Q} : p > 0,\ p^2 < 2\}$, and $B := \{p \in \mathbb{Q} : p > 0,\ p^2 > 2\}$.
For every $p > 0$ define $q := p - \frac{p^2-2}{p+2} = \frac{2p+2}{p+2}$. Then $q^2 - 2 = \frac{2(p^2-2)}{(p+2)^2}$. Hence if $p \in A$, then $q \in A$ and $q > p$; if $p \in B$, then $q \in B$ and $0 < q < p$. Thus A has no largest element in $\mathbb{Q}$, and B has no smallest element in $\mathbb{Q}$.

Remark 1.1.3. *Density Does Not Mean Completeness*
Between two rationals $r < s$ there is another rational, namely $(r + s)/2$. Nevertheless, $\mathbb{Q}$ is not complete.

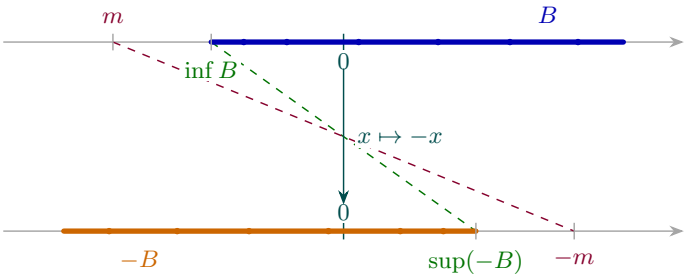


Figure 117: Reflecting B through 0 turns lower bounds of B into upper bounds of $-B$, so $\inf B = -\sup(-B)$.

Corollary 1.1.4. *Infimum by Reflection*
If $B \subseteq \mathbb{R}$ is nonempty and bounded below, then $\inf B = -\sup(-B)$ and $-B := \{-b : b \in B\}$.

Proof. Since B is nonempty and any lower bound m of B gives an upper bound $-m$ of $-B$, let $s = \sup(-B)$. For $b \in B$, $-b \le s$, so $-s \le b$; hence $-s$ is a lower bound of B. If t is any lower bound of B, then $-t$ is an upper bound of $-B$, so $s \le -t$ and $t \le -s$. Thus $-s = \inf B$.
Structural Commentary: Reflection swaps lower bounds of B with upper bounds of $-B$.
Key Insight: The infimum of B is the negative supremum of its reflection. □

Definition 1.1.5. *Real Absolute Value and Distance*
For $x \in \mathbb{R}$, define $|x| := \begin{cases} x, & x \geq 0, \\ -x, & x < 0. \end{cases}$.
For $a, b \in \mathbb{R}$, the distance between a and b on the real line is $|a - b|$.

Proposition 1.1.6. *Real Absolute-Value Laws*
For $x, y, a, b, c \in \mathbb{R}$,

(i) $|x| \geq 0$, $|x| = 0$ *iff* $x = 0$, *and* $-|x| \leq x \leq |x|$;
(ii) $|xy| = |x|\,|y|$;
(iii) $|x + y| \leq |x| + |y|$;
(iv) $|a - b| \leq |a - c| + |c - b|$.

Proof.

(i) Immediate from Definition 1.1.5.
(ii) Check the four sign cases for x and y.
(iii) Add $-|x| \leq x \leq |x|$ and $-|y| \leq y \leq |y|$, then $-(|x| + |y|) \leq x + y \leq |x| + |y|$.
Since $|x| + |y| \geq 0$, this gives $|x + y| \leq |x| + |y|$.
(iv) Apply (iii) to $a - b = (a - c) + (c - b)$.

Structural Commentary: Absolute value turns order bounds into distance estimates.
Key Insight: The triangle inequality on $\mathbb{R}$ is the inequality $|(a - c) + (c - b)| \leq |a - c| + |c - b|$. □

Perspective 1.1.7. *Completeness Becomes a Theorem*
Dedekind cuts build completeness from $\mathbb{Q}$: real numbers are rational cuts, rational cuts embed $\mathbb{Q}$, and least upper bounds arise by taking unions of cuts.

Definition 1.1.8. *Dedekind Cuts and Their Order*
A *cut* is a set $\alpha \subseteq \mathbb{Q}$ such that

(i) α is nonempty and $\alpha \neq \mathbb{Q}$;
(ii) if $p \in \alpha$, $q \in \mathbb{Q}$, and $q < p$, then $q \in \alpha$;
(iii) if $p \in \alpha$, then some $r \in \alpha$ satisfies $p < r$.

Let $\mathscr{D}$ be the set of all cuts. For $\alpha, \beta \in \mathscr{D}$, define $\alpha < \beta$ to mean $\alpha \subsetneq \beta$.
Rational letters $p, q, r, \ldots$ denote elements of $\mathbb{Q}$; Greek letters $\alpha, \beta, \gamma, \ldots$ denote cuts.
A cut is the rational left side of a real number: a proper initial segment with no last point.

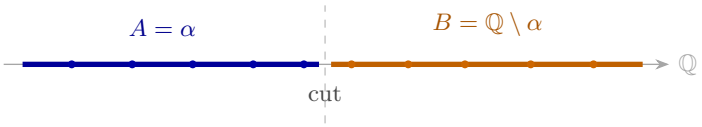


Figure 118: A Dedekind cut splits $\mathbb{Q}$ into α and $\mathbb{Q} \setminus \alpha$, with no last rational in α.

Remark 1.1.9. *Left Cuts Versus Two-Sided Cuts*
Two-sided notation writes a cut as $A|B$, with $A = \alpha$ and $B = \mathbb{Q} \setminus \alpha$; Definition 1.1.8 records only A.
Thus "no largest element" is the same left-side condition.

Definition 1.1.10. *Arithmetic of Cuts*
For cuts α, β, define $\alpha + \beta := \{r + s : r \in \alpha,\ s \in \beta\}$ and $0^* := \{q \in \mathbb{Q} : q < 0\}$.
Let candidate additive inverse of α be $-\alpha := \{p \in \mathbb{Q} : \exists r \in \mathbb{Q},\ r > 0,\ -p - r \notin \alpha\}$. Let $\mathscr{D}^+ := \{\alpha \in \mathscr{D} : \alpha > 0^*\}$.
For $\alpha, \beta \in \mathscr{D}^+$, define $\alpha\beta := \{p \in \mathbb{Q} : \exists r \in \alpha,\ \exists s \in \beta,\ r > 0,\ s > 0,\ p \leq rs\}$ and $1^* := \{q \in \mathbb{Q} : q < 1\}$.
Extend multiplication to all cuts by setting $\alpha 0^* = 0^* \alpha = 0^*$ and, for nonzero signs,

$$\alpha\beta := \begin{cases} (-\alpha)(-\beta), & \alpha < 0^*,\ \beta < 0^*, \\ -\big((-\alpha)\beta\big), & \alpha < 0^*,\ \beta > 0^*, \\ -\big(\alpha(-\beta)\big), & \alpha > 0^*,\ \beta < 0^*. \end{cases}$$

Motivation: arithmetic uses rational approximations inside cuts; signs come after positive products.

Definition 1.1.11. *Rational Cuts*
For $r \in \mathbb{Q}$, define the *rational cut*, $r^* := \{p \in \mathbb{Q} : p < r\}$.
The map $r \mapsto r^*$ is intended to identify $\mathbb{Q}$ with a subfield of the cut field.

Theorem 1.1.12. *Existence of the Real Field*
There exists an ordered field $\mathbb{R}$ with the least-upper-bound property. Moreover, $\mathbb{R}$ contains $\mathbb{Q}$ as a subfield.

Proof. Construct $\mathbb{R}$ from $\mathbb{Q}$ using Definitions 1.1.8–1.1.11.

(i) *Order.* The relation $\alpha < \beta$ is transitive, and at most one of $\alpha < \beta$, $\alpha = \beta$, and $\beta < \alpha$ can hold. For trichotomy, suppose $\alpha \not\subseteq \beta$ and $\alpha \neq \beta$. Choose $p \in \alpha$ with $p \notin \beta$. If $q \in \beta$, then $q < p$, since $p \notin \beta$ and β is downward closed. Hence $q \in \alpha$, so $\beta \subseteq \alpha$. Since $\alpha \neq \beta$, $\beta < \alpha$. Thus $\mathscr{D}$ is ordered.
Commentary: One missing rational decides containment.

(ii) *Least Upper Bounds.* Let $A \subseteq \mathscr{D}$ be nonempty and bounded above by a cut β. Define $\gamma := \bigcup_{\alpha \in A} \alpha$.
Since some $\alpha_0 \in A$ is nonempty, γ is nonempty. Since every $\alpha \in A$ lies in β, $\gamma \subseteq \beta$, so $\gamma \neq \mathbb{Q}$. If $p \in \gamma$ and $q < p$, then $p \in \alpha_1$ for some $\alpha_1 \in A$, hence $q \in \alpha_1 \subseteq \gamma$. If $p \in \gamma$, choose $\alpha_1 \in A$ with $p \in \alpha_1$ and then choose $r \in \alpha_1$ with $p < r$; then $r \in \gamma$. Thus γ is a cut.
Every $\alpha \in A$ satisfies $\alpha \leq \gamma$. If $\delta < \gamma$, choose $s \in \gamma \setminus \delta$. Then $s \in \alpha$ for some $\alpha \in A$, so $\delta < \alpha$ and δ is not an upper bound of A. Hence $\gamma = \sup A$.
Motivation: Supremum is union.
Commentary: Bounded unions stay proper and give the smallest upper bound.

(iii) *Addition.* The set $\alpha + \beta$ is a cut. It is nonempty by choosing one element from each summand.
If $r' \notin \alpha$ and $s' \notin \beta$, then every $r \in \alpha$ and $s \in \beta$ satisfies $r < r'$ and $s < s'$, so $r + s < r' + s'$; hence $r' + s' \notin \alpha + \beta$, and $\alpha + \beta$ is proper. If $p = r + s \in \alpha + \beta$ and $q < p$, then $q = r + (s + q - p)$ with $s+q-p < s$, so $s+q-p \in \beta$ and $q \in \alpha+\beta$. If $p = r+s$, choose $r_1 \in \alpha$ with $r < r_1$; then $p < r_1+s \in \alpha+\beta$. Thus $\alpha + \beta$ is a cut. Commutativity and associativity follow by rewriting representatives in $\mathbb{Q}$.
The cut 0^* is the additive identity: $\alpha + 0^* \subseteq \alpha$ by downward closure, while if $p \in \alpha$, choose $r \in \alpha$ with $p < r$ and write $p = r + (p - r)$ with $p - r < 0$.
The inverse $-\alpha$ is a cut, and $\alpha + (-\alpha) = 0^*$. The inclusion $\alpha + (-\alpha) \subseteq 0^*$ follows because $s \in -\alpha$ forces $-s \notin \alpha$, so any $r \in \alpha$ satisfies $r < -s$. Conversely, if $v < 0$, put $w = -v/2 > 0$ and choose an integer n with $nw \in \alpha$ but $(n+1)w \notin \alpha$; then $p = -(n+2)w$ lies in $-\alpha$ and $v = nw + p \in \alpha + (-\alpha)$.
Commentary: Addition inherits its laws from $\mathbb{Q}$ once closure and inverses are checked.

(iv) *Order Compatibility with Addition.* The additive laws make $\mathscr{D}$ an additive group.
If $\beta < \gamma$, then $\alpha+\beta \subseteq \alpha+\gamma$. Equality would imply $\beta = \gamma$ after adding $-\alpha$ to both sides, so $\alpha+\beta < \alpha+\gamma$. Consequently, $\alpha > 0^*$ if and only if $-\alpha < 0^*$.
Commentary: Cancellation turns inclusion into strict inclusion.

(v) *Positive Multiplication.* If $\alpha, \beta \in \mathscr{D}^+$, then $\alpha\beta$ is a cut. Since each positive cut contains some positive rational, $\alpha\beta$ is nonempty. Choose positive $r' \notin \alpha$ and $s' \notin \beta$; then every positive $r \in \alpha$ and $s \in \beta$ satisfies $rs < r's'$, so $r's' \notin \alpha\beta$, and the product is proper. Downward closure is built into the condition $p \leq rs$. If $p \leq rs$ with $r, s > 0$, then either choose q with $p < q < rs$, or, when needed, choose $r_1 \in \alpha$ with $r < r_1$ and take q with $p < q \leq r_1 s$; hence the product has no largest element.
Commutativity follows from $rs = sr$. Associativity and distributivity are proved by unfolding membership and using the rational identities $(rs)t = r(st)$ and $r(s+t) = rs+rt$, with a positive rational slack inserted whenever a representative must be made strictly larger inside a cut. The identity is 1^*: if $p \leq rs$ with $r \in \alpha$, $s < 1$, and $r, s > 0$, then $p < r$, so $p \in \alpha$; conversely, if $p \in \alpha$, choose $r \in \alpha$ with $p < r$ and then choose $s \in \mathbb{Q}$ with $0 < s < 1$ and $p \leq rs$. In particular, if $\alpha > 0^*$ and $\beta > 0^*$, then $\alpha\beta > 0^*$.
Commentary: Positivity avoids sign reversals.

(vi) *Signed Multiplication.* The definitions extend multiplication to all cuts by reducing to positive factors and additive inverses. Since $\delta = -(-\delta)$ and the sign rules for addition have already been proved, commutativity, associativity, and inverses reduce to the positive case.
For distributivity, split into signs of $\alpha, \beta, \gamma, \beta + \gamma$. For example, if $\alpha > 0^*$, $\beta < 0^*$, and $\beta + \gamma > 0^*$, then $\gamma = (\beta + \gamma) + (-\beta)$, so positive distributivity gives $\alpha\gamma = \alpha(\beta + \gamma) + \alpha(-\beta)$. Because $\alpha(-\beta) = -(\alpha\beta)$, adding $\alpha\beta$ to both sides gives $\alpha\beta + \alpha\gamma = \alpha(\beta + \gamma)$. The other sign cases are identical reductions after replacing negative terms by their additive inverses.
Commentary: Distributivity reduces to positive cases plus signs.

(vii) *Embedding $\mathbb{Q}$.* Each r^* is a cut. For $r, s \in \mathbb{Q}$, $r^* + s^* = (r + s)^*$, $r^*s^* = (rs)^*$, and $r^* < s^* \iff r < s$.
For addition, if $p = u+v$ with $u < r$ and $v < s$, then $p < r+s$. Conversely, if $p < r+s$, choose $t > 0$ with $2t = r + s - p$ and put $r' = r - t$, $s' = s - t$; then $r' < r$, $s' < s$, and $p = r' + s'$. For positive products, if $p \leq uv$ with $0 < u < r$ and $0 < v < s$, then $p < rs$. Conversely, if $0 < p < rs$, choose positive $u < r$ and $v < s$ close enough to r and s that $p \leq uv$; if $p \leq 0$, use any positive $u < r$, $v < s$. The signed product cases follow from the definitions by additive inverses. The order equivalence follows because $r < s$ puts r in s^* but not in r^*, while $r^* < s^*$ gives some $p < s$ with $p \notin r^*$, hence $r \leq p < s$.
Commentary: The map $r \mapsto r^*$ preserves arithmetic and order.

Identify $\mathbb{Q}$ with the subfield $\mathbb{Q}^* := \{r^* : r \in \mathbb{Q}\} \subseteq \mathscr{D}$, and write $\mathbb{R}$ for $\mathscr{D}$. Items (i), (iii)–(vi) give an ordered field; item (ii) gives the least-upper-bound property; item (vii) embeds $\mathbb{Q}$ as a subfield.

Any two ordered fields with the least-upper-bound property are isomorphic, so this construction characterizes the real field up to ordered-field structure. □

Perspective 1.1.13. *Number Systems from Representatives*
Build by representatives, then pass to classes:

$$\mathbb{Z}: \ (a,b) \mapsto [(a,b)], \qquad \mathbb{Q}: \ (a,b),\ b \neq 0 \mapsto [(a,b)], \qquad \mathbb{R}: \ \text{cuts or Cauchy data} \mapsto [\cdot].$$

The pattern is: choose $\sim$, prove $+, \cdot, <, \ldots$ are well defined on $X/\sim$, then use the classes as the number system; compare Definition 4.1.1 and Proposition 1.2.31.
Key Insight: Representatives are scaffolding; the quotient laws are the interface.

Perspective 1.1.14. *Cauchy Construction as a Parallel Model*
Cuts build $\mathbb{R}$ here; Cauchy representatives give parallel model:

$$(a_n) \sim (b_n) \iff |a_n - b_n| \to 0, \qquad \mathbb{R} \simeq \{\text{rational Cauchy sequences}\}/\sim$$

This is the representative pattern from Perspective 1.1.13; compare the completion skeleton in Perspective 2.2.22. Definition 4.1.1 is the first well-definedness checkpoint.

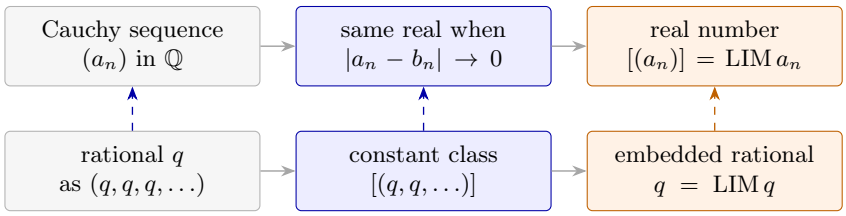


Figure 119: Real numbers as equivalence classes of rational Cauchy representatives.

Theorem 1.1.15. *Archimedean Property and Density of $\mathbb{Q}$*

(i) If $x, y \in \mathbb{R}$ and $x > 0$, then there is a positive integer n such that $nx > y$.
(ii) If $x, y \in \mathbb{R}$ and $x < y$, then there exists $p \in \mathbb{Q}$ such that $x < p < y$.

Proof.

(i) Let $A = \{nx : n \in \mathbb{N}\}$. If no nx exceeded y, then A would be bounded above. By Theorem 1.1.12, let $\alpha = \sup A$. Since $x > 0$, $\alpha - x < \alpha$, so $\alpha - x$ is not an upper bound.
Hence $\alpha - x < mx$ for some $m \in \mathbb{N}$, and $\alpha < (m+1)x \in A$, a contradiction.

(ii) Since $y - x > 0$, part (i) gives $n \in \mathbb{N}$ with $n(y-x) > 1$. Apply part (i) again to get $m_1, m_2 \in \mathbb{N}$ with $m_1 > nx$ and $m_2 > -nx$. Then some $m \in \mathbb{Z}$ satisfies $m - 1 \leq nx < m$.
Thus $nx < m \leq nx + 1 < ny$, so $x < m/n < y$ and $m/n \in \mathbb{Q}$.

Structural Commentary: Completeness makes $\{nx : n \in \mathbb{N}\}$ unbounded.
Key Insight: Archimedean unboundedness gives rational density. □

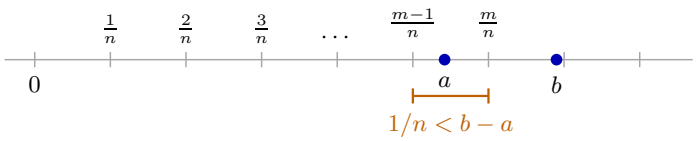


Figure 120: Choose n so the spacing $1/n$ is smaller than $b - a$; then one grid point m/n falls inside (a, b).

Theorem 1.1.16. *Existence and Uniqueness of Positive nth Roots*
For every $x \in \mathbb{R}$ with $x > 0$ and every positive integer n, $\exists$ a unique positive real number y such that $y^n = x$.

Proof.

(i) *Uniqueness.* If $0 < y_1 < y_2$, then $y_1^n < y_2^n$.

(ii) *Existence.* Let $E := \{t \in \mathbb{R} : t > 0,\ t^n < x\}$. $E \neq \varnothing$: for $t = x/(1+x)$, $0 < t < 1$ and $t^n \leq t < x$.
Also E is bounded above by $1 + x$, since $t > 1 + x$ implies $t^n \geq t > x$. By Theorem 1.1.12, let $y = \sup E$.
For $0 < a < b$, $b^n - a^n = (b-a)(b^{n-1} + b^{n-2}a + \cdots + a^{n-1}) < (b-a)nb^{n-1}$.
If $y^n < x$, choose h with $0 < h < 1$ and $h < \frac{x - y^n}{n(y+1)^{n-1}}$.
Then $(y+h)^n - y^n < hn(y+h)^{n-1} < hn(y+1)^{n-1} < x - y^n$.

Thus $y+h \in E$, contradicting that y is an upper bound. If $y^n > x$, set $k := \frac{y^n - x}{ny^{n-1}}$. Then $0 < k < y$.
For $t \geq y-k$, $y^n - t^n \leq y^n - (y-k)^n < kny^{n-1} = y^n - x$, so $t^n > x$ and $t \notin E$.
Hence $y-k$ is an upper bound for E, contradicting $y = \sup E$. Therefore $y^n = x$.

Structural Commentary: The root is $\sup\{t > 0 : t^n < x\}$.
Key Insight: Completeness supplies the missing root. □

Corollary 1.1.17. *Density of Rational and Irrational Points*
If $x < y$ in $\mathbb{R}$, then there exist $r \in \mathbb{Q}$ and $s \in \mathbb{R} \setminus \mathbb{Q}$ such that $x < r < y$ and $x < s < y$.

Proof. The rational point r is supplied by Theorem 1.1.15. Let $\eta = \sqrt{2}$.
Choose $n \in \mathbb{N}$ with $0 < \eta/n < y - x$, and choose $q \in \mathbb{Q}$ with $x - \eta/n < q < y - \eta/n$.
Then $s := q + \eta/n$ lies in (x, y) and is irrational, since $s \in \mathbb{Q}$ would imply $\eta = n(s-q) \in \mathbb{Q}$.
Structural Commentary: Irrational density is rational density shifted by a small irrational amount.
Key Insight: Once one irrational is known, every interval contains irrational translates of rational points. □

Corollary 1.1.18. *Product Rule for Positive Roots*
If $a, b > 0$ and n is a positive integer, then $(ab)^{1/n} = a^{1/n}b^{1/n}$.

Proof. Let $\alpha = a^{1/n}$ and $\beta = b^{1/n}$. Then $ab = \alpha^n\beta^n = (\alpha\beta)^n$ by commutativity of multiplication.
The uniqueness part of Theorem 1.1.16 gives the identity.
Structural Commentary: Use uniqueness after checking nth powers.
Key Insight: Root identities are uniqueness arguments. □

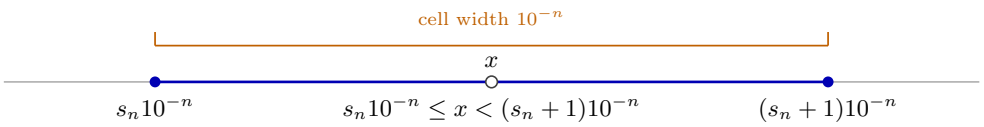


Figure 121: At precision 10^{-n}, the truncation $s_n 10^{-n}$ traps x in a cell of width 10^{-n}.

Proposition 1.1.19. *Decimal Truncations Squeeze a Real*
Let $x \geq 0$. For $n \geq 0$, set $s_n := \max\{k \in \mathbb{Z}_{\geq 0} : k10^{-n} \leq x\}$.
Then $s_n 10^{-n} \leq x < (s_n+1)10^{-n}$; and $0 \leq x - s_n 10^{-n} < 10^{-n}$; so $s_n 10^{-n} \to x$.
Moreover $s_{n+1} = 10s_n + a_{n+1}$ for some $a_{n+1} \in \{0, \ldots, 9\}$.

Proof. The maximum exists by the Archimedean property, Theorem 1.1.15, and well-ordering, Principle 2.6.2.
Maximality gives the displayed inequalities. The squeeze and Theorem 1.3.25 give $s_n 10^{-n} \to x$.
At levels n and $n+1$, $10s_n \leq s_{n+1} < 10s_n + 10$, so $a_{n+1} := s_{n+1} - 10s_n \in \{0, \ldots, 9\}$.
Structural Commentary: Decimal existence is an Archimedean grid argument followed by a squeeze.
Key Insight: Refining the mesh by a factor of 10 adds exactly one digit. □

Definition 1.1.20. *Extended Real Number System*
The *extended real number system* is $\overline{\mathbb{R}} := \mathbb{R} \cup \{-\infty, +\infty\}$.
The order on $\mathbb{R}$ is preserved, and $-\infty < x < +\infty$, $(x \in \mathbb{R})$.
Every nonempty subset of $\overline{\mathbb{R}}$ has least upper bound; if $E \subseteq \mathbb{R}$ is nonempty and unbounded above in $\mathbb{R}$, then $\sup E = +\infty$. Same convention applies to greatest lower bounds.

Remark 1.1.21. *Extended Arithmetic Conventions*
$\overline{\mathbb{R}}$ is not a field. For real x, $x + \infty = +\infty$, $x - \infty = -\infty$, and $\frac{x}{+\infty} = \frac{x}{-\infty} = 0$.
If $x > 0$, then $x(+\infty) = +\infty$ and $x(-\infty) = -\infty$; if $x < 0$, the signs reverse.
Real numbers are called *finite* when one wants to distinguish them from $\pm\infty$.

Definition 1.1.22. *Complex Numbers*
A *complex number* is an ordered pair (a, b) of real numbers. For $x = (a, b)$ and $y = (c, d)$, define equality componentwise and set $x + y := (a + c, b + d)$ and $xy := (ac - bd, ad + bc)$

Theorem 1.1.23. *The Complex Numbers Form a Field*
With the operations in Definition 1.1.22, the set of complex numbers is a field.
Its additive and multiplicative identities are $(0, 0)$ and $(1, 0)$.

Proof. The field axioms are checked directly from the two coordinate formulas and the field structure of $\mathbb{R}$. If $x = (a, b) \neq (0, 0)$, then $a^2 + b^2 > 0$ by Proposition 1.3.4, and $x^{-1} = \left(\frac{a}{a^2+b^2}, \frac{-b}{a^2+b^2}\right)$ satisfies $xx^{-1} = (1, 0)$.
Structural Commentary: Inverses use conjugate-over-square.
Key Insight: Positivity of $a^2 + b^2$ supplies the denominator for every nonzero complex number. □

Proposition 1.1.24. *Real Numbers Inside Complex Numbers*
For all real a, b, $(a, 0) + (b, 0) = (a + b, 0)$, and $(a, 0)(b, 0) = (ab, 0)$.
Thus complex numbers of form $(a, 0)$ have same arithmetic as real numbers, and identify $(a, 0)$ with a.
The complex unit is then $i := (0, 1)$.

Proposition 1.1.25. *Coordinate Form and i^2*
We have $i^2 = -1$. Moreover, if $a, b \in \mathbb{R}$, then $(a, b) = a + bi$.

Proof. First, $i^2 = (0, 1)(0, 1) = (-1, 0) = -1$. Also, $a + bi = (a, 0) + (b, 0)(0, 1) = (a, 0) + (0, b) = (a, b)$.
Structural Commentary: The ordered-pair construction recovers the usual notation after $i = (0, 1)$ is named.
Key Insight: The symbol i is not assumed; it is produced by the pair model. □

Definition 1.1.26. *Conjugate, Real Part, and Imaginary Part*
If $z = a + bi$ with $a, b \in \mathbb{R}$, then the *complex conjugate* of z is $\overline{z} := a - bi$.
The numbers a and b are the *real part* and *imaginary part*, written $a = \operatorname{Re}(z)$ and $b = \operatorname{Im}(z)$.

Proposition 1.1.27. *Basic Conjugation Identities*
If $z, w \in \mathbb{C}$, then
(i) $\overline{z + w} = \overline{z} + \overline{w}$;
(ii) $\overline{zw} = \overline{z}\,\overline{w}$;
(iii) $z + \overline{z} = 2\operatorname{Re}(z)$ *and* $z - \overline{z} = 2i\operatorname{Im}(z)$;
(iv) $z\overline{z}$ *is real and positive unless* $z = 0$.

Proof. Write $z = a + bi$ and $w = c + di$.
(i) $\overline{z + w} = (a + c) - (b + d)i = \overline{z} + \overline{w}$.
(ii) $\overline{zw} = (ac - bd) - (ad + bc)i = (a - bi)(c - di) = \overline{z}\,\overline{w}$.
(iii) $z + \overline{z} = 2a = 2\operatorname{Re}(z)$ and $z - \overline{z} = 2bi = 2i\operatorname{Im}(z)$.
(iv) $z\overline{z} = a^2 + b^2$, positive unless $a = b = 0$, i.e. unless $z = 0$.
Structural Commentary: Conjugation fixes Re and negates Im.
Key Insight: $z\overline{z} = a^2 + b^2 \geq 0$. □

Definition 1.1.28. *Absolute Value of a Complex Number*
For $z \in \mathbb{C}$, define $|z| := (z\overline{z})^{1/2}$, the nonnegative square root supplied by Theorem 1.1.16 and Proposition 1.1.27.

Theorem 1.1.29. *Absolute Value Identities*
For $z, w \in \mathbb{C}$,
(i) $|z| > 0$ *unless* $z = 0$, *and* $|0| = 0$;
(ii) $|\overline{z}| = |z|$;
(iii) $|zw| = |z|\,|w|$;
(iv) $|\operatorname{Re}(z)| \leq |z|$;
(v) $|z + w| \leq |z| + |w|$.

Proof.
(i) If $z = a + bi$, then $|z| = (a^2 + b^2)^{1/2}$; this is positive unless $a = b = 0$, and $|0| = 0$.
(ii) $|\overline{z}|^2 = \overline{z}z = z\overline{z} = |z|^2$, so nonnegative square roots give $|\overline{z}| = |z|$.
(iii) $|zw|^2 = zw\overline{zw} = z\overline{z}w\overline{w} = |z|^2\,|w|^2$; take nonnegative square roots.
(iv) If $z = a + bi$, then $|\operatorname{Re}(z)| = |a| \leq (a^2 + b^2)^{1/2} = |z|$.
(v) $z\overline{w}$ and $\overline{z}w$ are conjugates, $|z + w|^2 = |z|^2 + 2\operatorname{Re}(z\overline{w}) + |w|^2 \leq |z|^2 + 2\,|z\overline{w}| + |w|^2 = (|z| + |w|)^2$, using (iii) and (iv). Take nonnegative square roots.
Structural Commentary: The complex absolute value behaves like length because $z\overline{z}$ is a squared length.
Key Insight: Triangle inequality is a real-part estimate plus multiplicativity. □

Theorem 1.1.30. *Finite Schwarz Inequality for Complex Scalars*
If $a_1, \ldots, a_n$ and $b_1, \ldots, b_n$ are complex numbers, then $\left|\sum_{j=1}^{n} a_j\overline{b_j}\right|^2 \leq \left(\sum_{j=1}^{n} |a_j|^2\right)\left(\sum_{j=1}^{n} |b_j|^2\right)$.

Proof. This is Theorem 2.1.13 in $\mathbb{C}^n$ with the standard inner product.
Equivalently, set $A = \sum |a_j|^2$, $B = \sum |b_j|^2$, and $C = \sum a_j \overline{b_j}$; when $B > 0$, $\sum |Ba_j - Cb_j|^2 = B(AB - |C|^2) \geq 0$.
Structural Commentary: This scalar inequality is the coordinate model for geometric Cauchy–Schwarz inequality.
Key Insight: A nonnegative square sum controls the mixed term. □

1.2 Cardinality and Infinite Sets

Cardinality counts sets by matching them with other sets, using the earlier language of functions, power sets, and bijections. The foundational prerequisites collect same cardinality, finite/countable vocabulary, Cantor's theorem, Schroeder–Bernstein, and basic cardinal arithmetic; this subsection develops countability, diagonal arguments, finite-set summation, and the continuum.

1.3 Real Sequences and Completeness Methods

Sequences are the first testing ground for limits, order, and completeness. The main tools are subsequences, Cauchy criteria, monotone convergence, limsup/liminf, and Bolzano–Weierstrass.

Definition 1.3.1. *Convergent Real and Complex Sequences*
A real or complex sequence (s_n) *converges* to s if, $\forall \varepsilon > 0$, $\exists N \in \mathbb{N}$ s.t. $n \geq N$ implies $|s_n - s| < \varepsilon$.
Notation is $s_n \to s$ or $\lim_{n\to\infty} s_n = s$. If no such s exists, the sequence *diverges*.

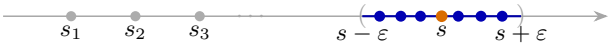


Figure 122: For all $n \geq N$, the terms s_n stay inside $(s - \varepsilon, s + \varepsilon)$.

Definition 1.3.2. *Range and Boundedness of a Real Sequence*
The *range* of (s_n) is set $\{s_n : n \in \mathbb{N}\}$. Real sequence is *bounded* if its range is bounded above and below in $\mathbb{R}$.

Theorem 1.3.3. *Limit Algebra for Complex Sequences*
Suppose (s_n) and (t_n) are complex sequences with $s_n \to s$ and $t_n \to t$. Then, for every $c \in \mathbb{C}$,

(i) $s_n + t_n \to s + t$;
(ii) $cs_n \to cs$ *and* $c + s_n \to c + s$;
(iii) $s_n t_n \to st$;
(iv) if $s \neq 0$ *and* $s_n \neq 0$ *for all* n, *then* $1/s_n \to 1/s$.

Proof.

(i) Given $\varepsilon > 0$, choose N such that $|s_n - s| < \varepsilon/2$, $|t_n - t| < \varepsilon/2$ for $n \geq N$.
Then $|(s_n + t_n) - (s + t)| \leq |s_n - s| + |t_n - t| < \varepsilon$.
(ii) Scalar multiplication uses $|cs_n - cs| = |c|\,|s_n - s|$; translation is (i) with constant sequence c.
(iii) Since $s_n \to s$, choose $N \geq 2$ with $|s_n - s| < 1$ for $n \geq N$ and set $M := \max\{|s_1|, \ldots, |s_{N-1}|, |s| + 1\}$.
Then $|s_n t_n - st| \leq |s_n|\,|t_n - t| + |t|\,|s_n - s| \leq M\,|t_n - t| + |t|\,|s_n - s| \to 0$.
(iv) If $s \neq 0$, then for all sufficiently large n, $|s_n| > |s|/2$. Hence $\left|\frac{1}{s_n} - \frac{1}{s}\right| = \frac{|s_n - s|}{|s_n||s|} \leq \frac{2}{|s|^2}\,|s_n - s| \to 0$.

Structural Commentary: Boundedness controls products; nonzero limits control reciprocals.
Key Insight: Product and reciprocal rules are estimates after the easy sum rule. □

Theorem 1.3.4. *Coordinatewise Convergence in $\mathbb{R}^k$*
Let $\mathbf{x}_n = (\alpha_{j,n})_{j=1}^k$ and $\mathbf{x} = (\alpha_j)_{j=1}^k$ in $\mathbb{R}^k$. Then $\mathbf{x}_n \to \mathbf{x} \iff \alpha_{j,n} \to \alpha_j$ $(1 \leq j \leq k)$.
Moreover, if $\mathbf{x}_n \to \mathbf{x}$, $\mathbf{y}_n \to \mathbf{y}$, and $\beta_n \to \beta$ in $\mathbb{R}$, then $\mathbf{x}_n + \mathbf{y}_n \to \mathbf{x} + \mathbf{y}$; $\mathbf{x}_n \cdot \mathbf{y}_n \to \mathbf{x} \cdot \mathbf{y}$; $\beta_n \mathbf{x}_n \to \beta \mathbf{x}$.

Proof.

($\Rightarrow$) Since $|\alpha_{j,n} - \alpha_j| \leq \|\mathbf{x}_n - \mathbf{x}\|_2$ for each j, each coordinate converges.
($\Leftarrow$) Given $\varepsilon > 0$, choose N so that $|\alpha_{j,n} - \alpha_j| < \varepsilon/\sqrt{k}$ for all j and $n \geq N$. Then $\|\mathbf{x}_n - \mathbf{x}\|_2 < \varepsilon$.

The algebraic conclusions follow coordinatewise from Theorem 1.3.3.
Structural Commentary: Euclidean convergence is exactly simultaneous coordinate convergence.
Key Insight: The norm packages finitely many coordinate estimates into one estimate. □

Remark 1.3.5. *Coordinate Metrics Give the Same Limits*
For $k \in \mathbb{N}$ and $x, y \in \mathbb{R}^k$, set $d_1(x,y) = \sum_{j=1}^k |x_j - y_j|$, $d_2(x,y) = \left(\sum_{j=1}^k |x_j - y_j|^2\right)^{1/2}$, $d_\infty(x,y) = \max_{1\leq j\leq k} |x_j - y_j|$.
Since $d_\infty \leq d_2 \leq d_1 \leq k d_\infty$, these have same convergent sequences, equivalently same coordinatewise criterion.

Definition 1.3.6. *Subsequences and Subsequential Limits*
Let (p_n) be a sequence. If $n_1 < n_2 < n_3 < \cdots$ is a strictly increasing sequence of positive integers, then (p_{n_k}) is a *subsequence* of (p_n). If $p_{n_k} \to p$, then p is a *subsequential limit* of (p_n).

Remark 1.3.7. *Subsequences as Selection Maps*
Equivalently, a subsequence is the composite of the original sequence, viewed as a function on $\mathbb{N}$ as in Definition 6.1.4, with a strictly increasing map $\sigma : \mathbb{N} \to \mathbb{N}$. In this notation the subsequence is $k \mapsto p_{\sigma(k)}$.
Key Insight: A subsequence keeps the original order but is allowed to skip terms.

Remark 1.3.8. *Subsequences of Convergent Sequences*
If $p_n \to p$, then every subsequence of (p_n) also converges to p. Conversely, if every subsequence of (p_n) converges to p, then (p_n) itself converges to p, since (p_n) is one of its subsequences.

Definition 1.3.9. *Monotone Real Sequence*
A real sequence (s_n) is *monotonically increasing* if $s_n \le s_{n+1}$ for all n, and *monotonically decreasing* if $s_n \ge s_{n+1}$ for all n. A *monotonic* sequence is either increasing or decreasing.

Definition 1.3.10. *Dominant Index*
For a real sequence (s_n), an index n is *dominant* if $s_n \ge s_m$ for every $m > n$.

Theorem 1.3.11. *Monotone Subsequence Theorem*
Every real sequence has a monotone subsequence.

Proof. If the dominant indices $n_1 < n_2 < \cdots$ are infinite, then $s_{n_1} \ge s_{n_2} \ge \cdots$, giving a decreasing subsequence. Otherwise choose N beyond all dominant indices and set $n_1 > N$.
Given $n_k > N$, nondominance gives $n_{k+1} > n_k$ with $s_{n_{k+1}} > s_{n_k}$; hence (s_{n_k}) is increasing.
Structural Commentary: Dominant indices give decreasing case; nondominance recursively gives increasing case.
Key Insight: The proof organizes arbitrary real behavior into one ordered thread. □

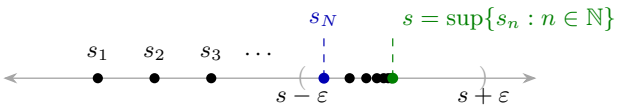


Figure 123: For an increasing bounded sequence, $s = \sup\{s_n\}$ is the limit.

Theorem 1.3.12. *Monotone Convergence Theorem for Sequences*
A monotonic real sequence converges if and only if it is bounded.

Proof.
($\Rightarrow$) If $s_n \to s$, choose N with $|s_n - s| < 1$ for $n \ge N$. The tail lies in $(s-1, s+1)$ and the initial set $\{s_1, \ldots, s_{N-1}\}$ is finite, so (s_n) is bounded.
($\Leftarrow$) Let $E := \{s_n : n \in \mathbb{N}\}$. If (s_n) is increasing, set $s := \sup E$. Given $\varepsilon > 0$, choose N with $s - \varepsilon < s_N \le s$; then $s - \varepsilon < s_n \le s$ for $n \ge N$, so $s_n \to s$. If (s_n) is decreasing, set $s := \inf E$. Given $\varepsilon > 0$, choose N with $s \le s_N < s + \varepsilon$; then $s \le s_n \le s_N < s + \varepsilon$ for $n \ge N$, so $s_n \to s$.
Structural Commentary: Monotonicity traps bounded tails at $\sup E$ or $\inf E$.
Key Insight: The least upper bound is forced to be the limit. □

Theorem 1.3.13. *Bolzano–Weierstrass for Real Sequences*
Every bounded real sequence has a convergent subsequence.

Proof. By Theorem 1.3.11, choose a monotone subsequence. It is bounded because the original sequence is bounded, so Theorem 5.5.1 gives convergence.
Structural Commentary: Boundedness prevents escape; monotone subsequences force one ordered thread.
Key Insight: Monotone subsequence plus monotone convergence is the real-line compactness mechanism. □

Definition 1.3.14. *Cauchy Real Sequence*
A real sequence (s_n) is *Cauchy* if, for every $\varepsilon > 0$, there is $N \in \mathbb{N}$ such that $m, n \ge N$ implies $|s_n - s_m| < \varepsilon$.

Lemma 1.3.15. *Cauchy Plus One Convergent Subsequence*
If a real Cauchy sequence has a subsequence converging to $s \in \mathbb{R}$, then the whole sequence converges to s.

Proof. Given $\varepsilon > 0$, choose N_0 so that $m, n \geq N_0$ implies $|s_m - s_n| < \varepsilon/2$. Choose j with $n_j \geq N_0$ and $|s_{n_j} - s| < \varepsilon/2$. Then for every $n \geq N_0$, $|s_n - s| \leq |s_n - s_{n_j}| + |s_{n_j} - s| < \varepsilon$.
Structural Commentary: Cauchy tails follow any convergent thread.
Key Insight: Mutual tail closeness turns one subsequential limit into the full limit. □

Theorem 1.3.16. *Real Cauchy Criterion*
A real sequence converges in $\mathbb{R}$ if and only if it is Cauchy.

Proof. If $s_n \to s$, choose N such that $|s_n - s| < \varepsilon/2$ for $n \geq N$; then $|s_n - s_m| < \varepsilon$ for $m, n \geq N$. Conversely, suppose (s_n) is Cauchy. Choose N so that $m, n \geq N$ implies $|s_m - s_n| < 1$. Then the tail lies in $(s_N - 1, s_N + 1)$, and the finitely many initial terms are bounded, so (s_n) is bounded. By Theorem 1.3.13, some subsequence converges; Lemma 1.3.15 makes the whole sequence converge.
Structural Commentary: Completeness of $\mathbb{R}$ means Cauchy behavior has a real endpoint.
Key Insight: Cauchy gives boundedness, Bolzano–Weierstrass gives a subsequential endpoint, and the Cauchy property pulls the whole sequence there. □

Definition 1.3.17. *Divergence to $\pm\infty$*
For a real sequence (s_n),

$$s_n \to +\infty \iff \forall M \in \mathbb{R},\ \exists N \in \mathbb{N} :\ n \geq N \Rightarrow s_n \geq M,$$
$$s_n \to -\infty \iff \forall M \in \mathbb{R},\ \exists N \in \mathbb{N} :\ n \geq N \Rightarrow s_n \leq M.$$

This extends $\to$ notation for divergent real sequences; Definition 1.3.1 is unchanged.

Proposition 1.3.18. *Unbounded Real Sequences Have Infinite Subsequences*
A real sequence is unbounded iff some subsequence tends to $+\infty$ or $-\infty$.

Proof. Set $n_0 = 0$. If (s_n) is unbounded above, every tail is unbounded above, so choose $n_k > n_{k-1}$ with $s_{n_k} \geq k$; then $s_{n_k} \to +\infty$. The unbounded-below case similarly gives $s_{n_k} \leq -k$, hence $s_{n_k} \to -\infty$.
Conversely, a subsequence tending to $\pm\infty$ is unbounded, so (s_n) is unbounded.
Structural Commentary: Unbounded tails supply the thresholds $\pm k$.
Key Insight: Large late terms form a divergent subsequence. □

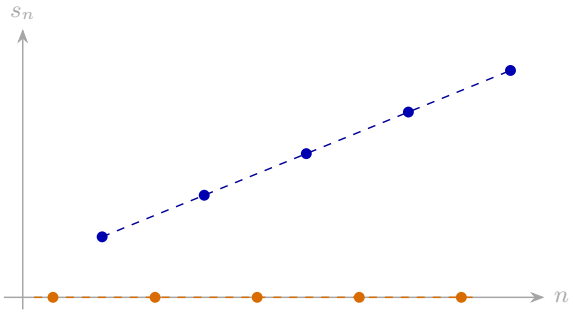


Figure 124: For $s_n = [(-1)^n + 1]n$, odd terms (orange) stay at 0, while even terms (blue) diverge to $+\infty$.

Definition 1.3.19. *Upper and Lower Limits*
For real sequence (s_n), let $E := \{x \in \overline{\mathbb{R}} : \exists\, n_1 < n_2 < \cdots,\ s_{n_k} \to x\}$, allowing $x = \pm\infty$.
Set $s^* := \sup E$ and $s_* := \inf E$; equivalently, $\limsup_{n\to\infty} s_n = s^*$ and $\liminf_{n\to\infty} s_n = s_*$.

Theorem 1.3.20. *Characterization of the Upper Limit*
Let (s_n) be a real sequence, and let E, s^*, s_* be as in Definition 1.3.19. Then:

(i) $s^* \in E$.
(ii) If $x > s^*$, then $s_n < x$ eventually.
(iii) These two properties characterize s^* uniquely.

Dually, s_* is uniquely characterized by $s_* \in E$ and $x < s_*$ implies $s_n > x$ eventually.

Proof.

(i) Since $s^* = \sup E$: if $s^* = +\infty$, use Proposition 1.3.18; if $s^* \in \mathbb{R}$, use Theorem 2.2.4. If $s^* = -\infty$ and $s_n \not\to -\infty$, infinitely many terms are bounded below; they are either unbounded above or have a bounded subsequence, giving an element of E above $-\infty$, a contradiction.

(ii) If $s_n \geq x > s^*$ infinitely often, that subsequence is unbounded above or has a finite subsequential limit $y \geq x$; both contradict $s^* = \sup E < x$. Hence $s_n < x$ eventually.

(iii) If $p < q$ both satisfy (i)–(ii), choose $p < x < q$. Then $s_n < x$ eventually, contradicting $q \in E$.

Apply the same argument to $(-s_n)$ for s_*.

Structural Commentary: Limsup is the maximal subsequential limit and eventual ceiling.

Key Insight: Limsup combines subsequences with eventual upper bounds. □

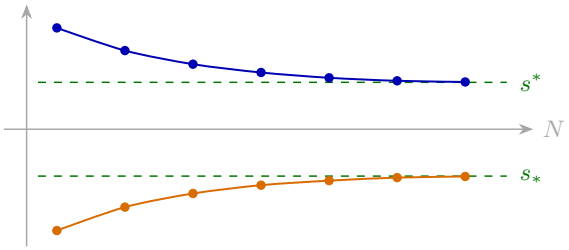


Figure 125: Blue tail suprema u_N decrease to $s^* = \limsup s_n$; orange tail infima ℓ_N increase to $s_* = \liminf s_n$.

Proposition 1.3.21. *Tail Suprema Formula*

For a real sequence (s_n), put $T_N := \{s_n : n \geq N\}$ and, in $\overline{\mathbb{R}}$, $u_N := \sup T_N$, $\ell_N := \inf T_N$. Then $u_N \downarrow$, $\ell_N \uparrow$, and $\limsup_{n\to\infty} s_n = \inf_N u_N$ and $\liminf_{n\to\infty} s_n = \sup_N \ell_N$.

Proof. Since $T_{N+1} \subseteq T_N$, $u_N \downarrow$ and $\ell_N \uparrow$. Let $s^* = \limsup s_n$ and $u = \inf_N u_N$.

If $x > s^*$, then $s_n < x$ eventually, so $u \leq u_N \leq x$ for some N. If $y < s^*$, then $s^* \in E$ gives a subsequence tending to s^*; hence every tail meets (y, ∞), so $u_N \geq y$ for all N and $u \geq y$. Thus $u = s^*$.

Apply the same formula to $(-s_n)$ and use Corollary 1.1.4.

Structural Commentary: Tail envelopes converge to lim sup and lim inf.

Key Insight: Shrinking tails compute the extremal subsequential limits. □

Theorem 1.3.22. *Convergence and Equal Upper/Lower Limits*

A real sequence (s_n) converges to a finite number $s \in \mathbb{R}$ if and only if $\limsup_{n\to\infty} s_n = \liminf_{n\to\infty} s_n = s$.

Proof.

($\Rightarrow$) If $s_n \to s$, every subsequence converges to s, and the tail is bounded, so no subsequence tends to $\pm\infty$. Thus $E = \{s\}$, hence $\limsup s_n = \liminf s_n = s$.

($\Leftarrow$) Suppose $\limsup s_n = \liminf s_n = s$. Given $\varepsilon > 0$, Theorem 1.3.20 gives $s_n < s + \varepsilon$ and $s_n > s - \varepsilon$ eventually. Hence $|s_n - s| < \varepsilon$ eventually, so $s_n \to s$.

Structural Commentary: Equality of the two extremal subsequential limits traps the whole tail.

Key Insight: Limsup supplies an eventual ceiling; liminf supplies an eventual floor. □

Theorem 1.3.23. *Order Comparison for Upper and Lower Limits*

If $s_n \leq t_n$ for all $n \geq N$, where N is fixed, then:

(i) $\limsup_{n\to\infty} s_n \leq \limsup_{n\to\infty} t_n$.

(ii) $\liminf_{n\to\infty} s_n \leq \liminf_{n\to\infty} t_n$.

Proof. Discard finitely many terms, so $s_n \leq t_n$ for all n.

(i) Put $a = \limsup s_n$ and $b = \limsup t_n$. If $x < a$, then $s_n > x$ infinitely often, hence $t_n > x$ infinitely often; Theorem 1.3.20 gives $x \leq b$. Thus $a \leq b$.

(ii) Apply (i) to $-t_n \leq -s_n$ to get $\limsup(-t_n) \leq \limsup(-s_n)$, equivalently $\liminf s_n \leq \liminf t_n$.

Structural Commentary: Eventual order is inherited by the extremal subsequential limits.

Key Insight: Limsup and liminf ignore finitely many initial terms. □

Theorem 1.3.24. *Pinching Theorem for Real Sequences*

Let (a_n), (s_n), and (b_n) be real sequences with $a_n \leq s_n \leq b_n$ for all sufficiently large n. If $a_n \to s$ and $b_n \to s$, then $s_n \to s$.

Proof. Let $\varepsilon > 0$. Choose N_0 s.t. $a_n \leq s_n \leq b_n$ for $n \geq N_0$; N_1 s.t. $s - \varepsilon < a_n$ for $n \geq N_1$; and N_2 s.t. $b_n < s + \varepsilon$ for $n \geq N_2$. If $N = \max\{N_0, N_1, N_2\}$, then $s - \varepsilon < s_n < s + \varepsilon$ for all $n \geq N$. Hence $s_n \to s$.

Structural Commentary: Two convergent bounds squeeze the middle sequence to the same limit.

Key Insight: Once both walls are inside the ε-neighborhood, the middle term has no room to escape. □

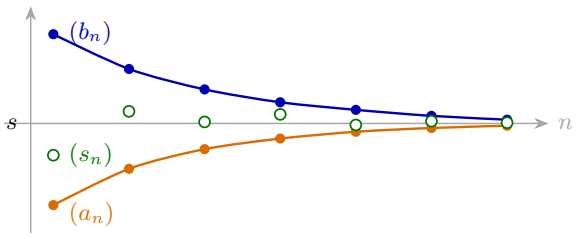


Figure 126: With $a_n \le s_n \le b_n$ and $a_n, b_n \to s$, the middle sequence also tends to s.

Theorem 1.3.25. *Special Sequence Limits*

(i) If $p > 0$, then $\lim_{n\to\infty} n^{-p} = 0$.
(ii) If $p > 0$, then $\lim_{n\to\infty} \sqrt[n]{p} = 1$.
(iii) $\lim_{n\to\infty} \sqrt[n]{n} = 1$.
(iv) If $p > 0$ and $\alpha \in \mathbb{R}$, then $\lim_{n\to\infty} \frac{n^\alpha}{(1+p)^n} = 0$.
(v) If $|x| < 1$, then $\lim_{n\to\infty} x^n = 0$.

Proof.

(i) Given $\varepsilon > 0$, choose $N > (1/\varepsilon)^{1/p}$ by Theorem 1.1.15. Then $n \ge N$ gives $n^{-p} < \varepsilon$.
(ii) If $p > 1$, set $x_n := \sqrt[n]{p} - 1$. Then $x_n > 0$ and $1 + nx_n \le (1 + x_n)^n = p$, so $0 < x_n \le (p-1)/n \to 0$. By Theorem 1.3.24, $x_n \to 0$. If $p = 1$, then $\sqrt[n]{p} = 1$ for every n. If $0 < p < 1$, then $1/p > 1$, so $\sqrt[n]{1/p} \to 1$ by the previous case. Since $\sqrt[n]{p} = 1/\sqrt[n]{1/p}$ and limits preserve reciprocals away from 0, $\sqrt[n]{p} \to 1$.
(iii) Put $x_n := \sqrt[n]{n} - 1$. Then $x_n \ge 0$, and for $n \ge 2$ the binomial theorem gives $n = (1 + x_n)^n \ge \frac{n(n-1)}{2} x_n^2$. Thus $0 \le x_n \le \sqrt{2/(n-1)}$, so Theorem 1.3.24 gives $x_n \to 0$.
(iv) Choose an integer $k > \alpha$ with $k > 0$. For $n > 2k$, $(1+p)^n \ge \binom{n}{k} p^k = \frac{n(n-1)\cdots(n-k+1)}{k!} p^k > \frac{n^k p^k}{2^k k!}$. Hence $0 < \frac{n^\alpha}{(1+p)^n} < \frac{2^k k!}{p^k} n^{\alpha-k}$, and $n^{\alpha-k} \to 0$ by (i). Apply Theorem 1.3.24.
(v) If $x \ne 0$, write $|x| = (1+p)^{-1}$ with $p > 0$ and apply (iv) with $\alpha = 0$ to get $|x|^n \to 0$. Since $-|x|^n \le x^n \le |x|^n$, Theorem 1.3.24 gives $x^n \to 0$; $x = 0$ is immediate.

Structural Commentary: Each limit is reduced to an elementary comparison sequence.
Key Insight: The binomial theorem supplies the right upper bounds for roots and exponential domination. □

Method 1.3.26. *Sequence Proof Patterns*

(i) *Convergence*: fix $\varepsilon > 0$, then find N so that all later terms lie in the relevant ε-ball.
(ii) *Uniqueness*: compare two proposed limits with $\varepsilon/2$ and use the triangle inequality.
(iii) *Subsequences*: choose $n_k \uparrow$ recursively with distances $< 1/k$.
(iv) *Cauchy arguments*: estimate the diameter of tails rather than guessing the limit first.
(v) *Monotone sequences*: use sup or inf and then prove eventual trapping.

Key Insight: Most sequence proofs are bookkeeping for quantifiers, indices, and one controlling inequality.

2 Metric and Topological Structure

Metric and Topological Structure abstracts limits, continuity, compactness, completeness, and connectedness, ending with only the general topology needed later.

2.1 Metric Spaces and Euclidean Topology

Metric topology studies sets through balls: openness, closure, and limit points come from distance. Euclidean estimates rely on norms, Cauchy–Schwarz, and the triangle inequality.

Definition 2.1.1. *Metric Space*
A *metric space* is a set X equipped with a function $d : X \times X \to \mathbb{R}$ such that, for all $p, q, r \in X$,

i. $d(p, q) > 0$ if $p \ne q$, and $d(p, p) = 0$;
ii. $d(p, q) = d(q, p)$;
iii. $d(p, q) \le d(p, r) + d(r, q)$.

The function d is called a *distance function* or *metric*.

Perspective 2.1.2. *ε-Closeness Becomes Metric Language*
On $\mathbb{R}$, ε-close means $|x-y|<\varepsilon$, i.e. $d(x,y)<\varepsilon$ with $d(x,y)=|x-y|$; may be abstracted to metric balls.
Key Insight: Absolute value is distance on $\mathbb{R}$.

Example 2.1.3. *Euclidean Metric and Subspaces*
The Euclidean space $\mathbb{R}^k$ is a metric space with $d(\mathbf{x},\mathbf{y}) := \|\mathbf{x}-\mathbf{y}\|_2, \qquad \mathbf{x},\mathbf{y}\in\mathbb{R}^k$.
Every subset Y of a metric space X is a metric space by restricting the same distance function to points of Y.

Example 2.1.4. *Discrete Metric*
On any set X, define $d(x,y)=\begin{cases}0, & x=y,\\ 1, & x\neq y.\end{cases}$.
This is the *discrete metric*: distinct points are all distance 1 apart.

Example 2.1.5. *Hamming Distance*
For $x,y\in\{0,1\}^n$, define $h(x,y) := |\{i : x_i\neq y_i\}|$. Then h is a metric.

Proof. $h\geq 0$, $h(x,y)=h(y,x)$, and $h(x,y)=0 \iff x_i=y_i\ \forall i \iff x=y$.
For $x,y,z\in\{0,1\}^n$, $\{i : x_i\neq y_i\}\subseteq\{i : x_i\neq z_i\}\cup\{i : z_i\neq y_i\}$, so $h(x,y)\leq h(x,z)+h(z,y)$.
Structural Commentary: Hamming distance counts coordinate disagreements.
Key Insight: Each disagreement between x and y must be paid for through z in at least one leg. □

Example 2.1.6. *Coordinate Norm Metrics on $\mathbb{R}^k$*
By Proposition 2.1.21 and Theorem 2.1.13, $d_1(\mathbf{x},\mathbf{y}) := \|\mathbf{x}-\mathbf{y}\|_1$; $d_2(\mathbf{x},\mathbf{y}) := \|\mathbf{x}-\mathbf{y}\|_2$; $d_\infty(\mathbf{x},\mathbf{y}) := \|\mathbf{x}-\mathbf{y}\|_\infty$ are metrics on $\mathbb{R}^k$, and $\forall \mathbf{x},\mathbf{y}\in\mathbb{R}^k$, $d_\infty(\mathbf{x},\mathbf{y})\leq d_2(\mathbf{x},\mathbf{y})\leq d_1(\mathbf{x},\mathbf{y})\leq \sqrt{k}\, d_2(\mathbf{x},\mathbf{y})$; $d_2(\mathbf{x},\mathbf{y})\leq \sqrt{k}\, d_\infty(\mathbf{x},\mathbf{y})$.
Hence d_1, d_2, d_∞ have the same convergent sequences.

Theorem 2.1.7. *Hölder and Minkowski for Sequence Spaces*
Let $\mathbb{F}\in\{\mathbb{R},\mathbb{C}\}$.

(i) If $p,q>1$ and $p^{-1}+q^{-1}=1$, then $\mathbf{x}\in\ell^p(\mathbb{F})$, $\mathbf{y}\in\ell^q(\mathbb{F})$ imply $\sum_{n\geq 1}|x_n y_n|\leq \|\mathbf{x}\|_p \|\mathbf{y}\|_q$.
(ii) If $1\leq p<\infty$ and $\mathbf{x},\mathbf{y}\in\ell^p(\mathbb{F})$, then $\mathbf{x}+\mathbf{y}\in\ell^p(\mathbb{F})$ and $\|\mathbf{x}+\mathbf{y}\|_p\leq \|\mathbf{x}\|_p+\|\mathbf{y}\|_p$.

Proof.

(i) If $\|\mathbf{x}\|_p\|\mathbf{y}\|_q=0$, the claim is immediate.
Otherwise set $X := \|\mathbf{x}\|_p$, $Y := \|\mathbf{y}\|_q$, $u_n := |x_n|/X$, and $v_n := |y_n|/Y$. Young's inequality gives $u_n v_n\leq u_n^p/p+v_n^q/q$; summing over $1\leq n\leq N$ and letting $N\to\infty$ gives $\sum_{n\geq 1}|x_n y_n|\leq XY$.

(ii) For $p=1$, apply $|x_n+y_n|\leq|x_n|+|y_n|$ and sum. For $p>1$, put $s_N := (\sum_{n=1}^N |x_n+y_n|^p)^{1/p}$. If $s_N=0$, the bound is trivial; if $s_N>0$, then $\sum_{n=1}^N|x_n+y_n|^p\leq\sum_{n=1}^N(|x_n|+|y_n|)\,|x_n+y_n|^{p-1}\leq(\|\mathbf{x}\|_p+\|\mathbf{y}\|_p)s_N^{p-1}$ by (i), so $s_N\leq\|\mathbf{x}\|_p+\|\mathbf{y}\|_p$. Let $N\to\infty$.

Structural Commentary: Hölder controls products; Minkowski controls sums.
Key Insight: The metric property of ℓ^p is exactly Minkowski's inequality. □

Example 2.1.8. *Function and Sequence Metric Spaces*
Let $\mathbb{F}\in\{\mathbb{R},\mathbb{C}\}$.

(i) On $C_{\mathbb{F}}[a,b]$, $d_\infty(f,g) := \sup_{t\in[a,b]}|f(t)-g(t)|$ and $d_1(f,g) := \int_a^b |f(t)-g(t)|\,dt$ are metrics.
Definiteness of d_1 uses continuity.
(ii) On $\ell^\infty(\mathbb{F})$, $d_\infty(\mathbf{x},\mathbf{y}) := \sup_{n\geq 1}|x_n-y_n|$ is a metric.
(iii) For $1\leq p<\infty$, $d_p(\mathbf{x},\mathbf{y}) := \|\mathbf{x}-\mathbf{y}\|_p = (\sum_{n\geq 1}|x_n-y_n|^p)^{1/p}$ is a metric on $\ell^p(\mathbb{F})$ by Theorem 2.1.7.

Structural Commentary: Norm estimates produce metrics on function and sequence spaces.
Key Insight: Infinite-dimensional examples require estimates, not just coordinate formulas.

Definition 2.1.9. *Euclidean Intervals, Cells, Balls, and Convexity*

(i) In $\mathbb{R}$, $(a,b) := \{x : a<x<b\}$, $[a,b] := \{x : a\leq x\leq b\}$, with $[a,b)$ and $(a,b]$ defined by mixed inequalities.
(ii) For $a_i<b_i$, a *k-cell* is $\prod_{i=1}^k[a_i,b_i]=\{\mathbf{x}\in\mathbb{R}^k : a_i\leq x_i\leq b_i,\ 1\leq i\leq k\}$.
(iii) For $\mathbf{x}\in\mathbb{R}^k$ and $r>0$, $B_r(\mathbf{x}) := \{\mathbf{y} : \|\mathbf{y}-\mathbf{x}\|_2<r\}$ and $\overline{B}_r(\mathbf{x}) := \{\mathbf{y} : \|\mathbf{y}-\mathbf{x}\|_2\leq r\}$.
(iv) $E\subseteq\mathbb{R}^k$ is *convex* iff $\lambda\mathbf{x}+(1-\lambda)\mathbf{y}\in E$ for all $\mathbf{x},\mathbf{y}\in E$ and $0\leq\lambda\leq 1$.
(v) The *integer lattice* is $\mathbb{Z}^k$, and the *first orthant* is $\{\mathbf{x}\in\mathbb{R}^k : x_i\geq 0\ \forall i\}$.
(vi) The *unit ball* is $B^k := \{\mathbf{x}\in\mathbb{R}^k : \|\mathbf{x}\|_2\leq 1\}$, and the *unit sphere* is $S^{k-1} := \{\mathbf{x}\in\mathbb{R}^k : \|\mathbf{x}\|_2=1\}$.
(vii) A *convex combination* of $\mathbf{x},\mathbf{y}$ is $\lambda\mathbf{x}+(1-\lambda)\mathbf{y}$ with $0\leq\lambda\leq 1$; these points form the line segment $[\mathbf{x},\mathbf{y}]$.

Proposition 2.1.10. *Balls and Cells Are Convex*
Open balls, closed balls, and k-cells in $\mathbb{R}^k$ are convex.

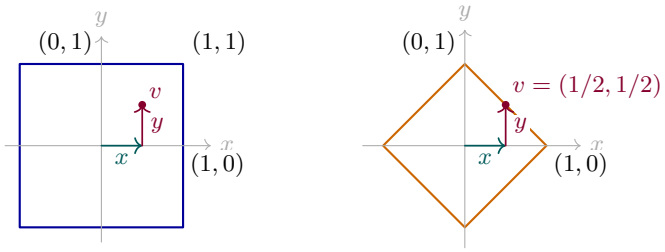


Figure 127: Unit balls in $\mathbb{R}^2$: $\max\{|x|, |y|\} < 1$ and $|x| + |y| < 1$.

Proof. Let $\mathbf{y}, \mathbf{z} \in B_r(\mathbf{x})$ and $0 < \lambda < 1$.
Then $\|\lambda\mathbf{y} + (1-\lambda)\mathbf{z} - \mathbf{x}\|_2 = \|\lambda(\mathbf{y} - \mathbf{x}) + (1-\lambda)(\mathbf{z} - \mathbf{x})\|_2 \leq \lambda \|\mathbf{y} - \mathbf{x}\|_2 + (1-\lambda) \|\mathbf{z} - \mathbf{x}\|_2 < r$.
The same proof with $\leq r$ gives closed balls. For a k-cell, each coordinate of $\lambda\mathbf{y} + (1-\lambda)\mathbf{z}$ lies between the corresponding coordinates of $\mathbf{y}$ and $\mathbf{z}$, hence remains in the same interval $[a_i, b_i]$.
Structural Commentary: Convexity is checked by following the line segment between two points.
Key Insight: Balls use the triangle inequality; cells use coordinate inequalities. □

Definition 2.1.11. *Product Metrics*
For metric spaces (X, d_X) and (Y, d_Y), and $p = (x, y)$, $p' = (x', y')$ in $X \times Y$, define

$$d_E(p, p') := \sqrt{d_X(x, x')^2 + d_Y(y, y')^2},$$
$$d_{\max}(p, p') := \max\{d_X(x, x'), d_Y(y, y')\},$$
$$d_{\text{sum}}(p, p') := d_X(x, x') + d_Y(y, y')$$

These are the Euclidean, maximum, and sum product metrics; the sum metric is often simplest for estimates.

Proposition 2.1.12. *Comparing Product Metrics*
On $X \times Y$, $d_{\max} \leq d_E \leq d_{\text{sum}} \leq 2d_{\max}$. Hence the three product metrics have same convergent sequences.

Proof. Let $a = d_X(x, x')$ and $b = d_Y(y, y')$.
Then $\max\{a, b\} \leq \sqrt{a^2 + b^2} \leq a + b \leq 2\max\{a, b\}$; convergence follows by squeezing.
Structural Commentary: The three product metrics differ by controlled constants.
Key Insight: Equivalent metric estimates preserve convergence. □

Proposition 2.1.13. *Bounded Compression of a Metric*
If (X, d) is a metric space, then $\rho(x, y) := d(x, y)/(1 + d(x, y))$ is a bounded metric on X, and $\tau_\rho = \tau_d$.

Proof. $0 \leq \rho < 1$, and definiteness and symmetry follow from d.
For $a, b \geq 0$, $(a + b)/(1 + a + b) \leq a/(1 + a) + b/(1 + b)$.
Since $t/(1 + t)$ is increasing, $d(x, z) \leq d(x, y) + d(y, z)$ gives $\rho(x, z) \leq \rho(x, y) + \rho(y, z)$.
For balls, $B_d(x, \eta) \subseteq B_\rho(x, \eta)$ for $0 < \eta < 1$, and $B_\rho(x, \eta) \subseteq B_d(x, \eta/(1 - \eta))$; hence $\tau_\rho = \tau_d$.
Structural Commentary: Compressing distances into $[0, 1)$ changes boundedness but not local neighborhoods.
Key Insight: Topology depends on small-scale control, not on the large-scale size of distances. □

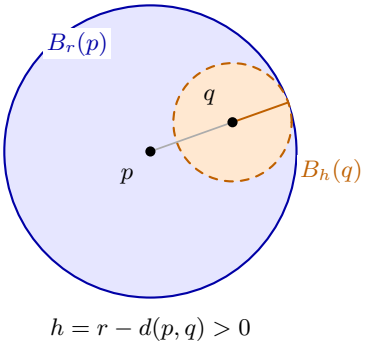


Figure 128: Every point inside a metric ball has a smaller ball still contained in the original ball.

Definition 2.1.14. *Metric Topology Terminology*
Let X be a metric space, and let all points and sets below be understood as points of X and subsets of X.

(i) A *neighborhood* of p with radius $r > 0$ is $B_r(p) := \{q \in X : d(p,q) < r\}$.
The corresponding *closed ball* is $\overline{B}_r(p) := \{q \in X : d(p,q) \leq r\}$, and the *sphere* of radius r about p is $\{q \in X : d(p,q) = r\}$.
(ii) A point p is a *limit point* of E if every neighborhood of p contains a point $q \neq p$ with $q \in E$.
(iii) If $p \in E$ and p is not a limit point of E, then p is an *isolated point* of E.
(iv) A set E is *closed* if every limit point of E is a point of E.
(v) A point p is an *interior point* of E if some neighborhood of p is contained in E.
(vi) A set E is *open* if every point of E is an interior point of E.
(vii) The *complement* of E is $E^c := X \setminus E$.
(viii) A set E is *dense* in X if every point of X is either a limit point of E, a point of E, or both.

For $X = \mathbb{R}$, this gives $B_\varepsilon(a) = (a - \varepsilon, a + \varepsilon)$.
Structural Commentary: Openness and closedness are ambient notions, not opposites.

Example 2.1.15. *Open, Closed, Both, or Neither*
In $\mathbb{R}$, $\emptyset, \mathbb{R}$ are open and closed, while $[a,b)$, $a < b$, is neither. If d is discrete on X, then $B_{1/2}(x) = \{x\}$ for all $x \in X$, so every $E \subseteq X$ and E^c are open; hence every $E \subseteq X$ is closed.

Theorem 2.1.16. *Every Neighborhood Is Open*
Every neighborhood in a metric space is an open set.

Proof. Let $E = B_r(p)$, and choose $q \in E$. Put $h := r - d(p,q) > 0$.
If $s \in B_h(q)$, then $d(p,s) \leq d(p,q) + d(q,s) < d(p,q) + h = r$.
Thus $B_h(q) \subseteq E$, so q is an interior point of E. Since q was arbitrary, E is open.
Structural Commentary: Around any point inside a ball, the leftover distance to boundary gives a smaller ball.
Key Insight: The triangle inequality turns radius slack into openness. □

Corollary 2.1.17. *Closed Balls Are Closed*
For every $p \in X$ and $r > 0$, the closed ball $\overline{B}_r(p)$ is closed.

Proof. Let $q \notin \overline{B}_r(p)$, so $d(p,q) > r$. Put $h := (d(p,q) - r)/2$.
If $s \in B_h(q)$, then $d(p,s) \geq d(p,q) - d(q,s) > d(p,q) - h > r$, so $s \notin \overline{B}_r(p)$.
Hence the complement of $\overline{B}_r(p)$ is open, and $\overline{B}_r(p)$ is closed.
Structural Commentary: Outside a closed ball, positive distance remains to the boundary.
Key Insight: Reverse triangle estimates certify closedness. □

Theorem 2.1.18. *Limit Points Force Infinitely Many Nearby Points*
If p is a limit point of E, then every neighborhood of p contains infinitely many points of E.

Proof. Suppose some neighborhood N of p contains only finitely many points of E distinct from p, say $q_1, \ldots, q_n$.
If there are none, then N itself contradicts that p is a limit point. Otherwise put $\rho := \min_{1 \leq m \leq n} d(p, q_m) > 0$.
The neighborhood $B_\rho(p)$ contains no point $q \in E$ with $q \neq p$, contradicting the definition of limit point.
Structural Commentary: A finite list of nearby exceptions has a smallest positive distance from p.
Key Insight: Limit points cannot be witnessed by only finitely many neighboring points. □

Corollary 2.1.19. *Finite Point Sets Have No Limit Points*
A finite point set in a metric space has no limit points.

Proof. Let $E \subseteq X$ be finite. If p were a limit point of E, then by Theorem 2.1.18, every neighborhood of p would contain infinitely many distinct points of E, impossible as E is finite. Hence E has no limit points.
Structural Commentary: This corollary is the finite obstruction to accumulation.
Key Insight: A limit point requires infinitely many nearby witnesses. □

Theorem 2.1.20. *Open Sets and Closed Complements*
A set E is open if and only if its complement E^c is closed.

Proof.

(⇒) Suppose E is open. If x is a limit point of E^c and $x \in E$, then some neighborhood N of x satisfies $N \subseteq E$, contradicting $N \cap (E^c \setminus \{x\}) \neq \emptyset$. Hence $x \in E^c$, so E^c is closed.
(⇐) Suppose E^c is closed. If $x \in E$, then $x \notin E^c$, hence x is not a limit point of E^c. Thus some neighborhood N of x has $N \cap E^c = \emptyset$, i.e. $N \subseteq E$. Hence E is open.

Structural Commentary: Openness of a set is equivalent to closure of the outside.
Key Insight: Interior points of E cannot be limit points of E^c. □

Corollary 2.1.21. *Closed Sets and Open Complements*
A set F is closed if and only if F^c is open.

Proof. Apply Theorem 2.1.20 with $E = F^c$. Then F^c is open if and only if $(F^c)^c$ is closed. Since $(F^c)^c = F$, this is exactly the assertion that F^c is open if and only if F is closed.
Structural Commentary: This is the same open–closed duality after taking complements.
Key Insight: Closedness of F is openness of the outside F^c. □

Theorem 2.1.22. *Basic Topological Calculus*
In a metric space:

(i) any union of open sets is open;
(ii) any intersection of closed sets is closed;
(iii) any finite intersection of open sets is open;
(iv) any finite union of closed sets is closed.

Proof.

(i) Let $G := \bigcup_\alpha G_\alpha$. If $x \in G$, then $x \in G_\alpha$ for some α.
Since G_α is open, some neighborhood of x lies in $G_\alpha \subseteq G$. Hence G is open.

(ii) Let $F := \bigcap_\alpha F_\alpha$. By Theorem 1.4.3, $F^c = \bigcup_\alpha F_\alpha^c$.
Each F_α^c is open by Corollary 2.1.21, so F^c is open by (i). Hence F is closed.

(iii) Let $H := \bigcap_{i=1}^n G_i$. If $x \in H$, let $r_i > 0$, $B_{r_i}(x) \subseteq G_i$, $r := \min_{1\le i\le n} r_i$. Then $B_r(x) \subseteq H$, so H is open.

(iv) Let $F := \bigcup_{i=1}^n F_i$. By Theorem 1.4.3, $F^c = \bigcap_{i=1}^n F_i^c$.
Each F_i^c is open by Corollary 2.1.21, so F^c is open by (iii). Hence F is closed.

Structural Commentary: Arbitrary unions are easy for open sets; finite intersections require a minimum radius.
Key Insight: The finite minimum is the reason infinite intersections of open sets need not be open. □

Definition 2.1.23. *Closure*
If $E \subseteq X$ and E^{acc} denotes the set of all limit points of E in X, then the *closure* of E is $\overline{E} := E \cup E^{\text{acc}}$.

Theorem 2.1.24. *Closure Is the Smallest Closed Superset*
If X is a metric space and $E \subseteq X$, then:

(i) $\overline{E}$ is closed;
(ii) $\overline{E} = E$ if and only if E is closed;
(iii) $\overline{E} \subseteq F$ for every closed set $F \subseteq X$ such that $E \subseteq F$.

Thus $\overline{E}$ is the smallest closed subset of X that contains E.

Proof.

(i) Let $p \notin \overline{E}$. Then $p \notin E$ and p is not a limit point of E, so $B_r(p) \cap E = \emptyset$ for some $r > 0$. If $q \in B_r(p)$, then $q \notin E$ and $B_{r-d(p,q)}(q) \subseteq B_r(p)$, so $q \notin E^{\text{acc}}$.
Hence $B_r(p) \subseteq (\overline{E})^c$, and $\overline{E}$ is closed by Corollary 2.1.21.

(ii) If $\overline{E} = E$, then E is closed by (i). Conversely, if E is closed, then $E^{\text{acc}} \subseteq E$, so $\overline{E} = E \cup E^{\text{acc}} = E$.

(iii) If F is closed and $E \subseteq F$, then $E^{\text{acc}} \subseteq F^{\text{acc}} \subseteq F$. Thus $\overline{E} = E \cup E^{\text{acc}} \subseteq F$.

Structural Commentary: Closure adds exactly the limit points forced by E.
Key Insight: Closed supersets must contain all limit points, so they must contain $\overline{E}$. □

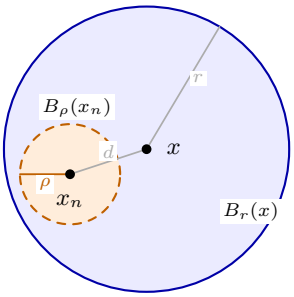


Figure 129: If $d = d(x, x_n) < r$, then any $0 < \rho < r - d$ gives $B_\rho(x_n) \subseteq B_r(x)$.

Definition 2.1.25. *Interior, Boundary, Adherence, and Distance*
Let (X, d) be a metric space and $E, F \subseteq X$.

(i) $E^\circ := \{p \in E : \exists r > 0,\ B_r(p) \subseteq E\}$ is the *interior*
(ii) $\text{ext}(E) := (E^c)^\circ = \{p \in X : \exists r > 0,\ B_r(p) \cap E = \emptyset\}$ is the *exterior*.
(iii) $\text{bd}(E) := \overline{E} \setminus E^\circ = \{p \in X : \forall r > 0,\ B_r(p) \cap E \neq \emptyset,\ B_r(p) \cap E^c \neq \emptyset\}$ is the *boundary*.
(iv) $p \in X$ is *adherent* to E iff $\forall r > 0,\ B_r(p) \cap E \neq \emptyset$, equivalently $p \in \overline{E}$.
(v) For $p \in X$, nonempty E, F, $\text{dist}(p, E) := \inf\{d(p, q) : q \in E\}$' $\text{dist}(E, F) := \inf\{d(p, q) : p \in E,\ q \in F\}$.

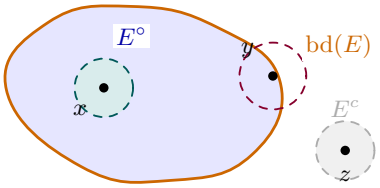


Figure 130: Interior, exterior, and boundary points are distinguished by how small balls meet E and E^c.

Proposition 2.1.26. *Distance-to-Set Closure Criterion*
Let $\emptyset \neq E \subseteq X$. For $p \in X$, $p \in \overline{E} \iff \text{dist}(p, E) = 0$.

Proof.

($\Rightarrow$) If $p \in \overline{E}$, then $\forall \varepsilon > 0$, $B_\varepsilon(p) \cap E \neq \emptyset$; choose $q_\varepsilon \in E$ with $d(p, q_\varepsilon) < \varepsilon$.
Thus $0 \leq \text{dist}(p, E) \leq \varepsilon$ for all $\varepsilon > 0$, so $\text{dist}(p, E) = 0$.
($\Leftarrow$) If $\text{dist}(p, E) = 0$, then $\forall r > 0$, $\exists q \in E$ with $d(p, q) < r$.
Hence $B_r(p) \cap E \neq \emptyset$ for all $r > 0$, so p is adherent to E, equivalently $p \in \overline{E}$.

Structural Commentary: Closure can be detected by vanishing distance to the set.
Key Insight: The infimum packages all small-ball tests into one number. □

Remark 2.1.27. *Set Distance Is Not a Metric on Subsets*
The function $\text{dist}(E, F)$ is useful, but it is not a metric on nonempty subsets.
In $\mathbb{R}$, for example, $\text{dist}((0, 1), (1, 2)) = 0$ although the sets are distinct and disjoint.



Figure 131: Adherent endpoints need not lie in E; isolated points are adherent but not limit points.

Definition 2.1.28. *Bounded Sets and Metric Sequences*
Let (X, d) be a metric space, $E \subseteq X$, and (p_n) a sequence in X.

(i) E is *bounded* iff $\exists q \in X, \exists M > 0$ such that $E \subseteq B_M(q)$
(ii) E is *perfect* iff E is closed and every $p \in E$ is a limit point of E.
(iii) (p_n) *converges* to $p \in X$ iff $\forall \varepsilon > 0, \exists N \in \mathbb{N}$ such that $n \geq N \Longrightarrow d(p_n, p) < \varepsilon$; write $p_n \to p$ or $\lim_{n\to\infty} p_n = p$.
If no such p exists, (p_n) *diverges*.
(iv) The *range* of (p_n) is $\text{ran}(p_n) := \{p_n : n \in \mathbb{N}\}$; (p_n) is *bounded* iff $\text{ran}(p_n)$ is bounded.
(v) (p_n) is *eventually in E* iff $\exists N \in \mathbb{N}$ such that $n \geq N \Longrightarrow p_n \in E$

(vi) (p_n) is *frequently in E* iff $\forall N \in \mathbb{N}, \exists n \geq N$ such that $p_n \in E$.

Structural Commentary: Boundedness, compactness, and completeness are intrinsic to restricted metrics; openness and closedness are ambient.
Key Insight: Convergence is tail control; frequent membership is arbitrarily late recurrence.

Proposition 2.1.29. *Limit Points Are Subsequence Limits*
For $(p_n) \subseteq X$ and $p \in X$, the following are equivalent:

(i) $\forall r > 0, \forall N \in \mathbb{N}, \exists n \geq N$ *such that* $p_n \in B_r(p)$.
(ii) $\exists n_1 < n_2 < \cdots$ *such that* $p_{n_k} \to p$.

Proof.

(i) $\Rightarrow$ *(ii)* Set $n_0 = 0$. By (i), recurse $n_k > n_{k-1}$ with $p_{n_k} \in B_{1/k}(p)$; then $d(p_{n_k}, p) < 1/k$, so $p_{n_k} \to p$.
(ii) $\Rightarrow$ *(i)* If $p_{n_k} \to p$, then for any $r > 0$ and $N \in \mathbb{N}$, choose k with $n_k \geq N$ and $d(p_{n_k}, p) < r$.
Thus $p_{n_k} \in B_r(p)$, so (i) holds.

Structural Commentary: Frequent visits are exactly enough to extract a convergent subsequence.
Key Insight: Limits require all sufficiently late terms; limit points require arbitrarily late witnesses. □

Theorem 2.1.30. *Basic Properties of Metric Convergence*
Let X be a metric space and $(p_n) \subseteq X$.

(i) $p_n \to p \iff \forall r > 0, \exists N \in \mathbb{N}$ *such that* $n \geq N \implies p_n \in B_r(p)$.
(ii) If $p_n \to p$ *and* $p_n \to p'$, *then* $p = p'$.
(iii) If $E \subseteq X$ *and* $p \in E^{\text{acc}}$, *then* $\exists (q_n) \subseteq E \setminus \{p\}$ *such that* $q_n \to p$.

Proof.

(i) $(\Rightarrow)$ If $p_n \to p$ and $r > 0$, take $\varepsilon = r$ in the definition.
(i) $(\Leftarrow)$ If the ball condition holds, then for every $\varepsilon > 0$ take $r = \varepsilon$, giving $d(p_n, p) < \varepsilon$ eventually.
(ii) Given $\varepsilon > 0$, choose N, N' with $n \geq N \implies d(p_n, p) < \varepsilon/2$ and $n \geq N' \implies d(p_n, p') < \varepsilon/2$.
If $n \geq \max\{N, N'\}$, then $d(p, p') \leq d(p, p_n) + d(p_n, p') < \varepsilon$, so $d(p, p') = 0$.
(iii) Since $p \in E^{\text{acc}}$, for each $n \in \mathbb{N}$ choose $q_n \in E \setminus \{p\}$ with $d(q_n, p) < 1/n$. Then $q_n \to p$.

Structural Commentary: Eventual containment gives convergence; triangle inequality gives uniqueness.
Key Insight: A limit point can always be approached by a sequence of points from the set. □

Corollary 2.1.31. *Subsequences Preserve Metric Limits*
If $p_n \to p$ in a metric space, then every subsequence (p_{n_k}) also converges to p.

Proof. Given $\varepsilon > 0$, choose N such that $n \geq N$ implies $d(p_n, p) < \varepsilon$.
Since $n_k \to \infty$, all sufficiently large k satisfy $n_k \geq N$, hence $d(p_{n_k}, p) < \varepsilon$.
Structural Commentary: A subsequence is only a delayed tail of the same sequence.
Key Insight: Eventual estimates survive deleting terms. □

Proposition 2.1.32. *Discrete-Metric Convergence*
Let X carry the discrete metric $d(x, y) = 0$ if $x = y$ and $d(x, y) = 1$ if $x \neq y$.
For $(p_n) \subset X$ and $p \in X$, $p_n \to p$ if and only if $\exists N \in \mathbb{N}$ such that $n \geq N \implies p_n = p$.

Proof.

$(\Rightarrow)$ Suppose $p_n \to p$. For $\varepsilon = 1/2$, choose $N \in \mathbb{N}$ such that $n \geq N \implies d(p_n, p) < 1/2$.
Since d takes only the values $0, 1$, $d(p_n, p) = 0$, hence $p_n = p$ for all $n \geq N$.
$(\Leftarrow)$ Suppose $\exists N \in \mathbb{N}$ s.t. $n \geq N \implies p_n = p$. Then $\forall \varepsilon > 0$, $n \geq N$ gives $d(p_n, p) = 0 < \varepsilon$, so $p_n \to p$.

Structural Commentary: The discrete metric has no room for distinct points to approach one another.
Key Insight: In the discrete metric, $d(x, y) < 1 \iff x = y$. □

Proposition 2.1.33. *Sequential Closure Criterion*
Let $E \subseteq X$, where X is a metric space, and let $p \in X$. Then $p \in \overline{E} \iff \exists (p_n)_{n \geq 1}$ with $p_n \in E$ $\forall n$, $p_n \to p$.

Proof.

$(\Rightarrow)$ If $p \in \overline{E}$, choose $p_n \in E \cap B_{1/n}(p)$ for $n \geq 1$ by Principle 6.5.1. Then $d(p_n, p) < 1/n$, so $p_n \to p$.
$(\Leftarrow)$ If $p_n \in E$ and $p_n \to p$, then every ball $B_r(p)$ contains some p_n.
Thus p is adherent to E by Definition 2.1.25, so $p \in \overline{E}$.

Structural Commentary: Closure can be read either by neighborhoods or by sequences.
Key Insight: Adherence supplies one witness at every scale. □

Theorem 2.1.34. *Closed Sets and Limit Points via Sequences*
Let X be a metric space and $E \subseteq X$.

(i) $p \in E^{\text{acc}} \iff \exists (p_n) \subseteq E \setminus \{p\}$ *such that* $p_n \to p$.
(ii) E *is closed* $\iff$ *every* $(p_n) \subseteq E$ *with* $p_n \to p \in X$ *satisfies* $p \in E$.

Proof.

(*i*) ($\Rightarrow$) If $p \in E^{\text{acc}}$, choose $p_n \in (E \setminus \{p\}) \cap B_{1/n}(p)$ for each n. Then $d(p_n, p) < 1/n$, so $p_n \to p$.
(*i*) ($\Leftarrow$) If $(p_n) \subseteq E \setminus \{p\}$ and $p_n \to p$, then every $B_r(p)$ contains some $p_n \neq p$. Hence $p \in E^{\text{acc}}$.
(*ii*) ($\Rightarrow$) If E is closed, $(p_n) \subseteq E$, and $p_n \to p$, then Proposition 2.1.33 gives $p \in \overline{E} = E$.
(*ii*) ($\Leftarrow$) If every convergent sequence in E has limit in E, then Proposition 2.1.33 gives $\overline{E} \subseteq E$. QED.

Structural Commentary: In metric spaces, closedness can be tested by sequences.
Key Insight: Balls at radii $1/n$ convert local information into a sequence. □

Proposition 2.1.35. *Density and Closure*
Let X be a metric space and let $E \subseteq X$. Then E is dense in X if and only if $\overline{E} = X$.

Proof.

($\Rightarrow$) Suppose E is dense in X. Let $p \in X$. By Definition 2.1.14, p is either a point of E or a limit point of E. Hence $p \in E \cup E^{\text{acc}} = \overline{E}$ by Definition 2.1.23. Thus $X \subseteq \overline{E}$. Since $\overline{E} \subseteq X$, $\overline{E} = X$.
($\Leftarrow$) Suppose $\overline{E} = X$. Let $p \in X$. Then $p \in \overline{E} = E \cup E^{\text{acc}}$, so p is either a point of E or a limit point of E. Therefore E is dense in X by Definition 2.1.14.

Structural Commentary: Density says every point is already in E or is forced by nearby points of E.
Key Insight: The closure packages exactly those two possibilities. □

Proposition 2.1.36. *Real-Line Density Interval Test*
Let $E \subseteq \mathbb{R}$. Then E is dense in $\mathbb{R}$ if and only if every nonempty open interval (a, b) satisfies $(a, b) \cap E \neq \emptyset$.

Proof.

($\Rightarrow$) Let $(a, b) \neq \emptyset$ and choose $p \in (a, b)$, $r := \frac{1}{2}\min\{p - a, b - p\}$. If $p \in E$, done. If $p \notin E$, density gives $p \in E^{\text{acc}}$, so choose $q \in (E \setminus \{p\}) \cap B_r(p)$. Since $B_r(p) \subseteq (a, b)$, $q \in (a, b) \cap E$.
($\Leftarrow$) Suppose $(a, b) \cap E \neq \emptyset$ for every nonempty open interval (a, b). For $p \in \mathbb{R}$, either $p \in E$, or $p \notin E$ and $\forall \varepsilon > 0$, choose $q \in (p - \varepsilon, p + \varepsilon) \cap E$. Then $q \neq p$ and $q \in B_\varepsilon(p)$, so $p \in E^{\text{acc}}$. Hence $\mathbb{R} = E \cup E^{\text{acc}} = \overline{E}$, so E is dense in $\mathbb{R}$.

Structural Commentary: On $\mathbb{R}$, balls and open intervals test the same local condition.
Key Insight: Density means no interval can avoid the set. □

Definition 2.1.37. *Separable Metric Space*
A metric space X is *separable* if it contains a countable dense subset.

Example 2.1.38. *Separable and Nonseparable Metric Spaces*

(i) $\mathbb{R}$ is separable: $\mathbb{Q}$ is countable and dense. Hence $\mathbb{R}^k$ is separable via $\mathbb{Q}^k$, and $\mathbb{C}^k$ via $\mathbb{C}^k \simeq \mathbb{R}^{2k}$.
(ii) If X is discrete, then X is separable $\iff$ X is countable,
since $B_{1/2}(x) = \{x\}$ forces every dense $D \subseteq X$ to contain x.
(iii) $\ell^\infty(\mathbf{F})$ is not separable. The set $\{0, 1\}^{\mathbb{N}}$ is uncountable, and $s \neq t$ gives $d_\infty(s, t) = 1$; the balls $B_{1/3}(s)$ are pairwise disjoint, so no countable dense set can meet all of them.
(iv) For $1 \leq p < \infty$, $\ell^p(\mathbf{F})$ is separable. Let D be the countable set of finite-support sequences with rational coordinates, using rational real and imaginary parts if $\mathbf{F} = \mathbb{C}$.
For $x \in \ell^p$ and $\varepsilon > 0$, choose N with $\sum_{n>N} |x_n|^p < (\varepsilon/2)^p$, then choose $y \in D$ with support in $\{1, \ldots, N\}$ and $\sum_{n=1}^{N} |x_n - y_n|^p < (\varepsilon/2)^p$. Thus $\|x - y\|_p < \varepsilon$.

Structural Commentary: Separability is the existence of a countable set meeting every nonempty open ball.
Key Insight: Finite rational data are dense in ℓ^p, while ℓ^∞ contains an uncountable 1-separated set.

Theorem 2.1.39. *Suprema Belong to Closures*
Let $\emptyset \neq E \subseteq \mathbb{R}$ be bounded above and $y := \sup E$. Then $y \in \overline{E}$; if E is closed, then $y \in E$.

Proof. If $y \in E$, then $y \in \overline{E}$. If $y \notin E$, then $\forall h > 0$, $y - h$ is not an upper bound; choose $x_h \in E$ with $y - h < x_h \leq y$. Thus $x_h \neq y$, $|x_h - y| < h$, and $B_h(y) \cap (E \setminus \{y\}) \neq \emptyset$, so $y \in E^{\text{acc}} \subseteq \overline{E}$.
If E is closed, $\overline{E} = E$ by Theorem 2.1.24(ii).
Structural Commentary: LUB gives points of E arbitrarily close below $\sup E$.
Key Insight: A supremum is either attained or approached. □

Definition 2.1.40. *Relative Openness*
Suppose $E \subseteq Y \subseteq X$, where X is a metric space. The set E is *open relative to Y* if for each $p \in E$ there is an $r > 0$ such that $q \in E$ whenever $q \in Y$ and $d(p,q) < r$.

Theorem 2.1.41. *Relative Opens Are Intersections*
Suppose $Y \subseteq X$. A subset E of Y is open relative to Y if and only if $E = Y \cap G$ for some open subset G of X.

Proof.
($\Rightarrow$) For each $p \in E$, choose $r_p > 0$ with $B_{r_p}(p) \cap Y \subseteq E$, and set $G := \bigcup_{p \in E} B_{r_p}(p)$. Then G is open in X, and $E \subseteq G \cap Y$. If $q \in G \cap Y$, then $q \in B_{r_p}(p) \cap Y$ for some $p \in E$, so $q \in E$. Thus $E = G \cap Y$.
($\Leftarrow$) If $E = G \cap Y$ with G open in X, then for each $p \in E$ choose $r > 0$ with $B_r(p) \subseteq G$. Hence $B_r(p) \cap Y \subseteq G \cap Y = E$, so E is open relative to Y.

Structural Commentary: Relative openness is ordinary openness viewed through the ambient subset Y.
Key Insight: Intersect open sets in X with Y to get exactly the open sets of Y. □

Example 2.1.42. *Relative Topology Depends on the Ambient Space*
Let $X = \{(x,0) : x \in \mathbb{R}\} \subseteq \mathbb{R}^2$ with the metric inherited from $\mathbb{R}^2$. The set $E = \{(x,0) : -1 < x < 1\}$ is not open as a subset of $\mathbb{R}^2$, but it is open relative to X, since $E = X \cap \big((-1,1) \times (-1,1)\big)$.
Likewise, $[0,1)$ is not closed in $\mathbb{R}$, but it is closed relative to $(-1,1)$ because $[0,1) = (-1,1) \cap [0,\infty)$.

Method 2.1.43. *Metric Proof Patterns*
(i) To prove G is open, fix $x \in G$ and find $r > 0$ with $B_r(x) \subseteq G$.
(ii) To prove F is closed, show $F^{\text{acc}} \subseteq F$, or equivalently show F^c is open.
(iii) To prove $C = \overline{E}$, show $E \subseteq C$, C is closed, and $C \subseteq F$ for every closed $F \supseteq E$.
(iv) To prove $p \in \overline{E}$, show $\operatorname{dist}(p,E) = 0$, or build a sequence in E converging to p.

2.2 Compactness and Completeness

Compactness extracts finite or convergent structure from infinite behavior; completeness asks whether Cauchy approximations converge internally.

Definition 2.2.1. *Open Covers and Compact Sets*
Let $E \subseteq X$, where X is a metric space.
An *open cover* of E is a collection $\{G_\alpha\}_{\alpha \in A}$ of open subsets of X such that $E \subseteq \bigcup_{\alpha \in A} G_\alpha$. A set $K \subseteq X$ is *compact* if every open cover of K has a finite subcover: $K \subseteq G_{\alpha_1} \cup \cdots \cup G_{\alpha_n}$ for some finite subcollection.



Figure 132: Compactness extracts a finite subcover.

Definition 2.2.2. *Sequential Compactness*
A metric space is *sequentially compact* if every sequence has a convergent subsequence.

Theorem 2.2.3. *Compactness Produces Convergent Subsequences*
(i) If (p_n) is a sequence in a compact metric space X, then $p_{n_i} \to p$ for some subsequence and some $p \in X$.
(ii) If $(\mathbf{x}_n)$ is bounded in $\mathbb{R}^k$, then it has a convergent subsequence.

Proof.
(i) Let $E = \{p_n : n \in \mathbb{N}\}$. If E is finite, choose a value p attained infinitely often; the corresponding constant subsequence converges. If E is infinite, Theorem 2.2.34 gives a limit point $p \in X$ of E. By Theorem 2.1.18, choose $n_i > n_{i-1}$ with $d(p_{n_i}, p) < 1/i$. Then $p_{n_i} \to p$.
(ii) The range of $(\mathbf{x}_n)$ has compact closure by Theorem 2.2.39; apply (i) in that closure.

Structural Commentary: Compactness turns arbitrary infinite behavior into an accumulating subsequence.
Key Insight: In $\mathbb{R}^k$, boundedness is enough because bounded sets sit inside compact closed cells. □

Theorem 2.2.4. *Subsequential Limits Form a Closed Set*
For a sequence (p_n) in a metric space X, the set of all subsequential limits is closed.

Proof. Let E^* be the set of subsequential limits and let q be a limit point of E^*. Set $n_0 = 0$.
For each i, choose $x_i \in E^*$ with $d(x_i, q) < 2^{-i}$. Since x_i is a subsequential limit, choose $n_i > n_{i-1}$ such that $d(p_{n_i}, x_i) < 2^{-i}$. Then $d(p_{n_i}, q) < 2^{1-i}$, so $p_{n_i} \to q$, and $q \in E^*$.
Structural Commentary: A limit of subsequential limits is reached by diagonalizing through those subsequences.
Key Insight: Choose later and closer terms at the same time. □

Definition 2.2.5. *Cauchy Sequence*
A sequence (p_n) in a metric space X is *Cauchy* if $\forall \varepsilon > 0$ $\exists$ integer N s.t. $n \geq N,\ m \geq N \implies d(p_n, p_m) < \varepsilon$.

Definition 2.2.6. *Diameter*
Let E be a nonempty subset of a metric space X. The *diameter* of E is $\operatorname{diam} E := \sup\{d(p,q) : p, q \in E\}$, with the supremum interpreted in the extended real sense when the set of distances is unbounded.

Proposition 2.2.7. *Basic Diameter Properties*
Let $\emptyset \neq E, F \subseteq X$.

(i) E is bounded $\iff \operatorname{diam} E < \infty \iff \exists q \in X, \exists M > 0$ such that $E \subseteq B_M(q)$.
(ii) $\operatorname{diam} E = 0 \iff |E| = 1$.
(iii) $E \subseteq F \implies \operatorname{diam} E \leq \operatorname{diam} F$.
(iv) E, F bounded $\implies E \cup F$ bounded.

Proof.

(i) If $E \subseteq B_M(q)$, then $d(x,y) \leq d(x,q) + d(q,y) < 2M$ for $x, y \in E$, so $\operatorname{diam} E < \infty$.
If $\operatorname{diam} E < \infty$ and $q \in E$, then $E \subseteq B_{\operatorname{diam} E + 1}(q)$.
(ii) $\operatorname{diam} E = 0 \iff d(x,y) = 0$ for all $x, y \in E \iff x = y$ for all $x, y \in E$; since $E \neq \emptyset$, this is $|E| = 1$.
(iii) $\{d(x,y) : x, y \in E\} \subseteq \{d(x,y) : x, y \in F\}$, so monotonicity of sup gives $\operatorname{diam} E \leq \operatorname{diam} F$.
(iv) Choose $e \in E$, $f \in F$, and set $D := \operatorname{diam} E + d(e,f) + \operatorname{diam} F < \infty$. For $x, y \in E \cup F$, either x, y lie in the same set, or after swapping, $x \in E$, $y \in F$, and $d(x,y) \leq d(x,e) + d(e,f) + d(f,y) \leq D$.
Thus $\operatorname{diam}(E \cup F) < \infty$, so $E \cup F$ is bounded by (i).

Structural Commentary: Diameter is a single-number form of boundedness for nonempty sets.
Key Insight: Cross-distances between two bounded sets are controlled by one bridge distance. □

Proposition 2.2.8. *Cauchy Sequences and Tail Diameters*
For (p_n) in a metric space X, set $E_N := \{p_j : j \geq N\}$. Then (p_n) is Cauchy if and only if $\operatorname{diam} E_N \to 0$.

Proof.

($\Rightarrow$) Given $\varepsilon > 0$, choose N_0 so that $m, n \geq N_0$ implies $d(p_m, p_n) < \varepsilon/2$.
If $N \geq N_0$, then every pair in E_N has distance $< \varepsilon/2$, so $\operatorname{diam} E_N < \varepsilon$.
($\Leftarrow$) Given $\varepsilon > 0$, choose N_0 with $\operatorname{diam} E_{N_0} < \varepsilon$.
If $m, n \geq N_0$, then $p_m, p_n \in E_{N_0}$, hence $d(p_m, p_n) \leq \operatorname{diam} E_{N_0} < \varepsilon$.

Structural Commentary: Cauchy behavior is shrinking diameter of sequence tails.
Key Insight: The limit is not named; only mutual closeness is measured. □

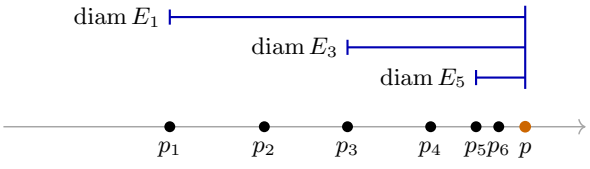


Figure 133: Tail diameters $\operatorname{diam} E_N$ shrink toward 0.

Theorem 2.2.9. *Nested Compact Sets with Shrinking Diameter*
Let X be a metric space.

(i) If $\emptyset \neq E \subseteq X$, then $\operatorname{diam} \overline{E} = \operatorname{diam} E$.
(ii) If $K_1 \supseteq K_2 \supseteq \cdots$ are nonempty compact subsets of X and $\operatorname{diam} K_n \to 0$, then $\left|\bigcap_{n=1}^{\infty} K_n\right| = 1$.

Proof.

(i) Since $E \subseteq \overline{E}$, $\operatorname{diam} E \leq \operatorname{diam} \overline{E}$. For $p, q \in \overline{E}$ and $\varepsilon > 0$, choose $p', q' \in E$ with $d(p,p'), d(q,q') < \varepsilon$; then $d(p,q) \leq \operatorname{diam} E + 2\varepsilon$. Hence $\operatorname{diam} \overline{E} \leq \operatorname{diam} E$.

(ii) Set $I := \bigcap_{n\geq 1} K_n$. Corollary 2.2.33 gives $I \neq \emptyset$. If $x, y \in I$, then $d(x,y) \leq \operatorname{diam} K_n$ for all n, so $d(x,y) = 0$ as $\operatorname{diam} K_n \to 0$. Thus $x = y$, hence $|I| = 1$.

Structural Commentary: Closure preserves diam; $K_n \downarrow$ with $\operatorname{diam} K_n \to 0$ forces a singleton.
Key Insight: Compactness gives existence; shrinking diameter gives uniqueness. □

Theorem 2.2.10. *Cauchy Criterion and Completeness Results*

(i) In any metric space, every convergent sequence is Cauchy.
(ii) If X is a compact metric space, then every Cauchy sequence in X converges to a point of X.
(iii) In $\mathbb{R}^k$, every Cauchy sequence converges.

Proof.

(i) If $p_n \to p$, choose N with $d(p_n, p) < \varepsilon/2$ for $n \geq N$. For $m, n \geq N$, $d(p_n, p_m) < \varepsilon$.
(ii) Let $E_N := \{p_j : j \geq N\}$ and $F_N := \overline{E_N}$. Then $\operatorname{diam} F_N \to 0$, the F_N are compact and nested, and Theorem 2.2.9 gives $\bigcap_N F_N = \{p\}$. If $\operatorname{diam} F_N < \varepsilon$, then $p, p_n \in F_N$ for $n \geq N$, so $d(p_n, p) < \varepsilon$.
(iii) For Cauchy $(\mathbf{x}_n) \subseteq \mathbb{R}^k$, choose N with $\operatorname{diam}\{\mathbf{x}_j : j \geq N\} < 1$. Its range is bounded, hence has compact closure by Theorem 2.2.39; apply (ii).

Structural Commentary: Cauchy controls $\operatorname{diam} E_N$; compactness turns $\bigcap_N F_N$ into the limit.
Key Insight: In $\mathbb{R}^k$, the Cauchy criterion is equivalent to convergence. □

Lemma 2.2.11. *A Cauchy Sequence with a Convergent Subsequence Converges*
Let (p_n) be a Cauchy sequence in a metric space X. If some subsequence (p_{n_j}) converges to $p \in X$, then $p_n \to p$.

Proof. Given $\varepsilon > 0$, choose N_0 so that $m, n \geq N_0$ implies $d(p_m, p_n) < \varepsilon/2$.
Choose j with $n_j \geq N_0$ and $d(p_{n_j}, p) < \varepsilon/2$. Then for every $n \geq N_0$, $d(p_n, p) \leq d(p_n, p_{n_j}) + d(p_{n_j}, p) < \varepsilon$.
Structural Commentary: The Cauchy tail follows any convergent subsequence to the same limit.
Key Insight: Mutual closeness plus one convergent thread forces the whole tail to converge. □

Proposition 2.2.12. *Boundedness and Distances of Cauchy Sequences*
Let $(p_n), (q_n)$ be Cauchy sequences in a metric space X.

(i) Cauchy boundedness: the range of $n \mapsto p_n$, $\operatorname{ran}(p_n) := \{p_n : n \in \mathbb{N}\}$ is bounded.
(ii) distance convergence: the sequence $n \mapsto d(p_n, q_n)$ converges in $\mathbb{R}$.
(iii) nonconverse: there exists a bounded sequence in $\mathbb{R}$ that is not Cauchy.

Proof.

(i) Choose N with $m, n \geq N \Longrightarrow d(p_m, p_n) < 1$.
Then $\{p_n : n \geq N\} \subseteq B_1(p_N)$, and $\{p_1, \ldots, p_{N-1}\}$ is finite, so $\operatorname{ran}(p_n)$ is bounded.
(ii) For $m, n \in \mathbb{N}$, $|d(p_n, q_n) - d(p_m, q_m)| \leq d(p_n, p_m) + d(q_n, q_m)$. Hence $(d(p_n, q_n))$ Cauchy in $\mathbb{R}$, converges.
(iii) Let $a_n := (-1)^n$. Then $|a_n| \leq 1$ for all n, but with $\varepsilon = 1$, for every N there exist $m, n \geq N$ with $|a_m - a_n| = 2$; hence (a_n) is bounded but not Cauchy.

Structural Commentary: Cauchy tails are bounded; Cauchy-pair distances converge.
Key Insight: The four-point triangle estimate controls distance sequences. □

Definition 2.2.13. *Complete Metric Space*
A metric space X is *complete* if every Cauchy sequence in X converges to a point of X.
Compact metric spaces and $\mathbb{R}^k$ are complete (by Theorem 2.2.10).

Proposition 2.2.14. *Completeness and Closed Subspaces*
Let (X, d) be a metric space, $E \subseteq X$, and $d_E := d|_{E\times E}$.

(i) If (E, d_E) is complete, then E is closed in X, equivalently $E = \overline{E}^X$.
(ii) If (X, d) is complete, then (E, d_E) is complete if and only if E is closed in X, equivalently $E = \overline{E}^X$.

Proof.

(i) Assume (E, d_E) complete. Let $p \in \overline{E}^X$. Choose $(p_n) \subseteq E$ with $d(p_n, p) < 1/n$.
Then $p_n \to p$ in X, hence (p_n) is Cauchy in (E, d_E). Thus $p_n \to q \in E$ in (E, d_E), hence also in X.
Uniqueness in X gives $p = q \in E$. Therefore $\overline{E}^X \subseteq E$, so $E = \overline{E}^X$.
(ii) Assume (X, d) complete. By (i), (E, d_E) complete $\Rightarrow E = \overline{E}^X$.
Conversely, if $E = \overline{E}^X$ and $(p_n) \subseteq E$ is Cauchy in (E, d_E), then (p_n) is Cauchy in X, so $p_n \to p \in X$.
Since $p \in \overline{E}^X = E$, every Cauchy sequence in (E, d_E) converges in (E, d_E).

Structural Commentary: Completeness gives the limit in X; closedness keeps it in E.
Key Insight: In a complete ambient space, E is complete exactly when $\overline{E} = E$. □

Example 2.2.15. *Finite-Dimensional Complete Metric Spaces*
For $k \in \mathbb{N}$, $(\mathbb{R}^k, d_2)$ and $(\mathbb{C}^k, d_2)$ are complete; in particular, so are $\mathbb{R}$ and $\mathbb{C}$.

Proof. Completeness of $\mathbb{R}$ is Theorem 1.3.16.
If $z_n = x_n + iy_n$ is Cauchy in $\mathbb{C}$, then (x_n) and (y_n) are Cauchy in $\mathbb{R}$, so $z_n \to x + iy \in \mathbb{C}$.
Theorem 2.2.10 gives completeness of $\mathbb{R}^k$, and the isometry $\mathbb{C}^k \cong \mathbb{R}^{2k}$ gives completeness of $\mathbb{C}^k$.
Structural Commentary: Finite-dimensional completeness is coordinatewise completeness.
Key Insight: A Cauchy sequence in finitely many coordinates converges coordinate by coordinate. □

Example 2.2.16. *Two Metrics on $C[0,1]$*
For $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, $C_{\mathbb{F}} := C([0,1], \mathbb{F})$ is complete under $d_\infty(f,g) := \sup_{x\in[0,1]} |f(x) - g(x)|$ by Theorem 3.3.1, but not under $d_1(f,g) := \int_0^1 |f(x) - g(x)|\, dx$.

Proof. For $n \geq 2$, define $f_n \in C_{\mathbb{F}}$ by

$$f_n(x) := \begin{cases} 0, & 0 \leq x \leq \frac{1}{2}, \\ n(x - \frac{1}{2}), & \frac{1}{2} \leq x \leq \frac{1}{2} + \frac{1}{n}, \\ 1, & \frac{1}{2} + \frac{1}{n} \leq x \leq 1. \end{cases}$$

Let $s := \mathbf{1}_{(1/2,1]}$. Then $d_1(f_n, s) = 1/(2n)$, hence (f_n) is d_1-Cauchy. If $f_n \to g \in C_{\mathbb{F}}$ in d_1, then $d_1(g,s) = 0$, so $g = s$ a.e. Continuity forces $g = 0$ on $[0, 1/2]$ and $g = 1$ on $(1/2, 1]$, impossible at $1/2$.
Structural Commentary: The integral metric permits a continuous ramp to collapse to a jump.
Key Insight: Supremum control preserves continuity; integral control can miss jumps. □

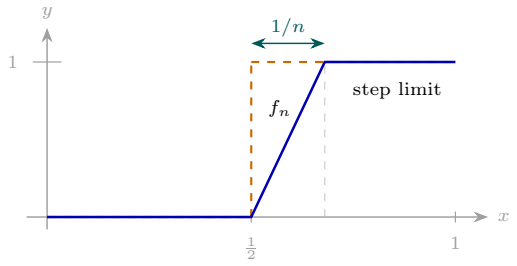


Figure 134: The continuous ramps f_n converge in the integral metric to a discontinuous step function.

Example 2.2.17. *Completeness of ℓ^∞ and ℓ^p*
For $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$, $(\ell^\infty(\mathbb{F}), \|\cdot\|_\infty)$ and $(\ell^p(\mathbb{F}), \|\cdot\|_p)$, $1 \leq p < \infty$, are complete.

Proof. For ℓ^∞, let $\mathbf{x}^{(n)} = (x_j^{(n)})_j$ be Cauchy. Then $x_j^{(n)} \to y_j \in \mathbb{F}$ for each j.
If $\left\|\mathbf{x}^{(n)} - \mathbf{x}^{(m)}\right\|_\infty < \varepsilon$ for $m, n \geq N$, then $\left|x_j^{(n)} - y_j\right| \leq \varepsilon$ for all j, so $\left\|\mathbf{x}^{(n)} - \mathbf{y}\right\|_\infty \leq \varepsilon$.
Taking $n = N$ gives $\mathbf{y} \in \ell^\infty$, hence $\mathbf{x}^{(n)} \to \mathbf{y}$.
For ℓ^p, let $\mathbf{x}^{(n)}$ be Cauchy and define $\mathbf{y} = (y_j)$ by coordinatewise limits.
If $\left\|\mathbf{x}^{(n)} - \mathbf{x}^{(m)}\right\|_p < \varepsilon$ for $m, n \geq N$, then for each r, $\sum_{j=1}^{r} \left|x_j^{(n)} - y_j\right|^p = \lim_{m\to\infty} \sum_{j=1}^{r} \left|x_j^{(n)} - x_j^{(m)}\right|^p \leq \varepsilon^p$.
Letting $r \to \infty$ gives $\left\|\mathbf{x}^{(n)} - \mathbf{y}\right\|_p \leq \varepsilon$ for $n \geq N$. Taking $n = N$ gives $\mathbf{y} \in \ell^p$, hence $\mathbf{x}^{(n)} \to \mathbf{y}$.
Structural Commentary: Completeness of scalar coordinates combines with uniform tail control.
Key Insight: Norm-Cauchy control keeps coordinatewise limits inside the space. □

Definition 2.2.18. *Contraction*
Let X be a metric space with metric d.
A map $\varphi : X \to X$ is a *contraction* if $\exists$ a constant $0 \leq c < 1$ such that $d(\varphi(x), \varphi(y)) \leq c\, d(x,y)$ for $x, y \in X$.
Structural Commentary: Bound $c < 1$ excludes nonexpansive maps and is the hypothesis in Theorem 2.2.19.

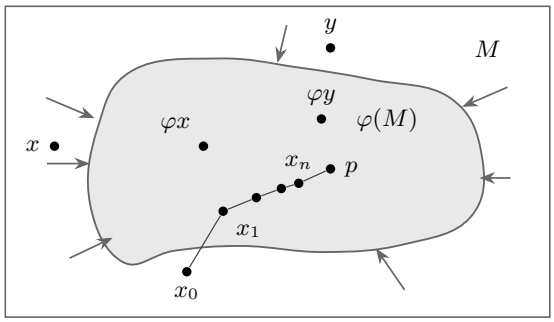


Figure 135: A contraction pulls iterates toward the unique fixed point.

Theorem 2.2.19. *Contraction Principle*
If X is a complete metric space and $\varphi : X \to X$ is a contraction, then $\exists$ a unique $x \in X$ such that $\varphi(x) = x$.

Proof. Choose $x_0 \in X$ and define $x_{n+1} = \varphi(x_n)$.
If $c < 1$ is a contraction constant, induction gives $d(x_{n+1}, x_n) \leq c^n d(x_1, x_0)$.
For $n < m$, the triangle inequality and the finite geometric sum give $d(x_n, x_m) \leq \sum_{j=n}^{m-1} d(x_j, x_{j+1}) \leq \frac{c^n}{1-c} d(x_1, x_0)$.
Thus (x_n) is Cauchy, so completeness gives $x_n \to x \in X$.
Since a contraction is continuous, $\varphi(x) = \lim \varphi(x_n) = \lim x_{n+1} = x$.
If $\varphi(x) = x$ and $\varphi(y) = y$, then $d(x, y) \leq c\, d(x, y)$, hence $d(x, y) = 0$ and $x = y$.
Structural Commentary: Iteration turns the contraction estimate into a geometric Cauchy estimate.
Key Insight: Completeness supplies the limiting point; the strict constant $c < 1$ forces uniqueness. □

Perspective 2.2.20. *Fixed Points as Existence Proofs*
A contraction proof solves $x = \varphi(x)$ by iteration in a complete metric space. This is the template behind the inverse function theorem, Theorem 4.3.4.

Example 2.2.21. *A Noncomplete Metric Space*
The rational numbers with $d(x, y) = |x - y|$ are not complete.
Let (r_n) be the decimal truncation sequence $1.4,\ 1.41,\ 1.414,\ 1.4142, \ldots$, for $\sqrt{2}$.
This is a concrete instance of the truncation squeeze in Proposition 1.1.19. If $k, n \geq N$, then $|r_k - r_n| \leq 10^{-N}$, so (r_n) is Cauchy in $\mathbb{Q}$. As a sequence in $\mathbb{R}$, it converges to $\sqrt{2}$.
If it converged to some $r \in \mathbb{Q}$ in the inherited metric, uniqueness of limits in $\mathbb{R}$ would force $r = \sqrt{2}$, contradicting Proposition 1.1.1. Thus $\mathbb{Q}$ is not complete in the inherited metric. *Structural Commentary*: The decimal approximants are Cauchy before their missing limit is available inside $\mathbb{Q}$.
Key Insight: Completeness asks whether every internally Cauchy approximation has an internal endpoint.

Perspective 2.2.22. *Completion by Cauchy Sequences*
Complete by identifying Cauchy sequences with $d(p_n, q_n) \to 0$ and setting $D([(p_n)], [(q_n)]) = \lim_{n\to\infty} d(p_n, q_n)$.
Constant sequences embed the original space; triangle inequalities and diagonals give a complete metric.

Theorem 2.2.23. *Metric Completion*
For every metric space (X, d), $\exists$ complete metric space $(\widehat{X}, \widehat{d})$, isometry $\iota : X \to \widehat{X}$ with $\overline{\iota(X)}^{\widehat{X}} = \widehat{X}$.
If (Y, d_Y, τ), (Z, d_Z, σ) are two such completions, then $\exists!$ bijective isometry $\Phi : Y \to Z$ s.t. $\Phi \circ \tau = \sigma$.

Proof.

i. *Existence*. Let $\mathrm{CS}(X)$ be the Cauchy sequences in X.
For $a = (a_n), b = (b_n) \in \mathrm{CS}(X)$, $|d(a_n, b_n) - d(a_m, b_m)| \leq d(a_n, a_m) + d(b_n, b_m)$, so $D(a, b) := \lim_{n\to\infty} d(a_n, b_n)$ exists. Put $a \sim b \Leftrightarrow D(a, b) = 0$, $\widehat{X} := \mathrm{CS}(X)/\!\sim$, and $\widehat{d}([a], [b]) := D(a, b)$.
The estimate gives well-definedness; metric axioms pass to limits. Map $\iota(x) := [(x, x, \ldots)]$ is an isometry.
If $a = (a_n)$ and $d(a_n, a_N) < \varepsilon$ for $n \geq N$, then $\widehat{d}(\iota(a_N), [a]) \leq \varepsilon$, hence $\overline{\iota(X)}^{\widehat{X}} = \widehat{X}$.
Let $([a^{(k)}])$ be Cauchy in $\widehat{X}$. Choose $x_k \in X$ with $\widehat{d}(\iota(x_k), [a^{(k)}]) < 1/k$.
Since $d(x_k, x_l) \leq 1/k + \widehat{d}([a^{(k)}], [a^{(l)}]) + 1/l$, (x_k) is Cauchy in X,
and $\widehat{d}([a^{(k)}], [(x_j)]) \leq 1/k + \limsup_{j\to\infty} d(x_k, x_j) \to 0$. Thus $([a^{(k)}]) \to [(x_j)]$, so $(\widehat{X}, \widehat{d})$ is complete.

ii. *Uniqueness*. Let (Y, d_Y, τ) and (Z, d_Z, σ) be completions.
Define $\rho_0 : \tau(X) \to \sigma(X)$ by $\rho_0(\tau x) := \sigma x$; it is an isometry.
For $y \in Y$, choose $y_n \in \tau(X)$ with $y_n \to y$. Since $\rho_0(y_n)$ is Cauchy in Z, define $\Phi(y) := \lim_{n\to\infty} \rho_0(y_n)$.
If $z_n \in \tau(X)$ and $z_n \to y$, then $d_Z(\rho_0(y_n), \rho_0(z_n)) = d_Y(y_n, z_n) \to 0$, so Φ is well defined.

If $y'_n \to y'$, then $d_Z(\Phi(y), \Phi(y')) = \lim_{n\to\infty} d_Z(\rho_0(y_n), \rho_0(y'_n)) = \lim_{n\to\infty} d_Y(y_n, y'_n) = d_Y(y, y')$.
Hence Φ is an isometry extending ρ_0. Its image is complete, hence closed in Z, and contains the dense set $\sigma(X)$; therefore $\Phi(Y) = Z$. Uniqueness follows from continuity and density of $\tau(X)$.

Structural Commentary: Missing limits are installed as equivalence classes of Cauchy approximations.
Key Insight: Completion adds Cauchy limits and is unique over the dense copy of X. □

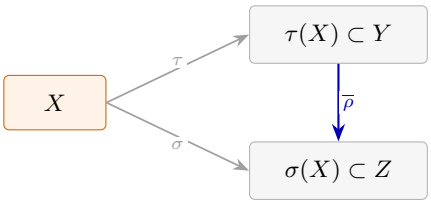


Figure 136: For completions Y, Z, the dense-copy isometry $\sigma\tau^{-1} : \tau(X) \to \sigma(X)$ extends uniquely to $\overline{\rho} : Y \cong Z$.

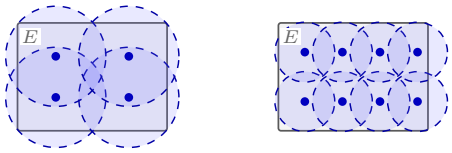


Figure 137: Finite ε- and δ-nets cover E, with $0 < \delta < \varepsilon$.

Definition 2.2.24. *Total Boundedness*
Let (X, d) be a metric space. It is *totally bounded* if, $\forall \varepsilon > 0$, $\exists$ a finite set $F \subseteq X$ such that $X \subseteq \bigcup_{x\in F} B_\varepsilon(x)$.
F is a finite ε-net. $E \subseteq X$ totally bounded if same condition holds with finite $F \subseteq E$ and X.

Theorem 2.2.25. *Total Boundedness Implies Boundedness*
Let X be a metric space. If X is totally bounded, then X is bounded.

Proof. If $X = \emptyset$, this is trivial. Otherwise cover X by finitely many 1-balls $B_1(x_i)$ and fix x_1.
Then $d(x, x_1) < 1 + \max_i d(x_i, x_1)$ for every $x \in X$.
Structural Commentary: One finite cover at scale 1 gives one global bound.
Key Insight: Total boundedness is boundedness with finite control at every scale. □

Theorem 2.2.26. *Compactness as Complete Finite Resolution*
Let X be a metric space. Then, X is compact $\iff$ X is complete and totally bounded.

Proof.
($\Rightarrow$) Let (p_k) be Cauchy in X. By Theorem 2.2.10, compactness forces (p_k) to converge, so X is complete.
Given $\varepsilon > 0$, cover $\{B_\varepsilon(x) : x \in X\}$ has finite subcover $B_\varepsilon(y_1), \ldots, B_\varepsilon(y_m)$; hence X is totally bounded.
($\Leftarrow$) Let $(x_n) \subseteq X$. Choose nested subsequences $(x_k^{(j)})_{k\geq 1}$ and points $a_j \in X$ with $x_k^{(j)} \in B_{1/j}(a_j)$ $\forall k$.
Diagonal $y_j := x_j^{(j)}$ satisfies $y_k, y_\ell \in B_{1/j}(a_j)$ for $k, \ell \geq j$, so it is Cauchy and converges by completeness; hence X is sequentially compact. Now let $\{U_\alpha\}$ cover X. If no $\delta > 0$ has every $B_\delta(x)$ contained in some U_α, choose x_n with $B_{1/n}(x_n) \nsubseteq U_\alpha$ $\forall \alpha$. Passing to $x_{n_k} \to x \in U_{\alpha_0}$, and openness gives $B_{1/n_k}(x_{n_k}) \subseteq U_{\alpha_0}$ for large k, a contradiction. Thus such δ exists.
Choose a finite $\delta/2$-net $\{z_i\}_{i=1}^m$ and α_i with $B_\delta(z_i) \subseteq U_{\alpha_i}$; then $\{U_{\alpha_i}\}_{i=1}^m$ covers X.
Structural Commentary: Finite nets plus completeness force compactness.
Key Insight: Compact metric spaces are exactly complete spaces visible at every finite scale. □

Example 2.2.27. *An Open Cover Without a Finite Subcover*
For $x \in (0, 1)$, set $O_x := (x/2, 1)$. These sets cover $(0, 1)$, since $y \in O_x$ whenever $0 < x < y < 1$.
No finite subcollection covers $(0, 1)$: if $x' = \min\{x_1, \ldots, x_n\}$, then $(0, x'/2]$ is missed by $O_{x_1} \cup \cdots \cup O_{x_n}$.

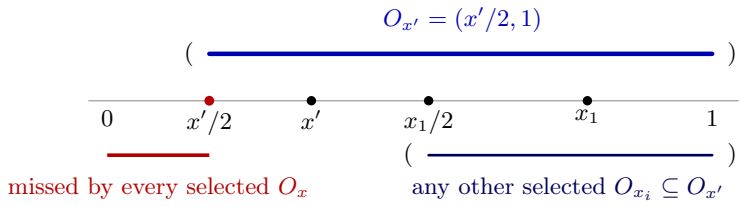


Figure 138: For finite choice, $x' = \min_i x_i$ makes $O_{x'}$ largest selected interval, but $(0, x'/2]$ remains uncovered.

Theorem 2.2.28. *Compactness Is Independent of the Ambient Space*
Suppose $K \subseteq Y \subseteq X$, where X is a metric space.
Then K is compact relative to X if and only if K is compact relative to Y.

Proof.
($\Rightarrow$) Let $\{V_\alpha\}_{\alpha \in A}$ be a relative-open cover of K in Y.
By Theorem 2.1.41, write $V_\alpha = Y \cap G_\alpha$ with G_α open in X. Since $\{G_\alpha\}_{\alpha \in A}$ covers K, compactness in X gives a finite $F \subseteq A$ such that $K \subseteq \bigcup_{\alpha \in F} G_\alpha$. As $K \subseteq Y$, $K \subseteq \bigcup_{\alpha \in F} V_\alpha$.
($\Leftarrow$) Let $\{G_\alpha\}_{\alpha \in A}$ be an open cover of K in X. Then $\{Y \cap G_\alpha\}_{\alpha \in A}$ is a relative-open cover of K in Y.
Compactness in Y gives a finite $F \subseteq A$ such that $K \subseteq \bigcup_{\alpha \in F}(Y \cap G_\alpha) \subseteq \bigcup_{\alpha \in F} G_\alpha$.
Structural Commentary: Relative covers are ambient covers intersected with Y.
Key Insight: Openness depends on the ambient space, but compactness does not. □

Theorem 2.2.29. *Compact Sets Are Closed*
Every compact subset of a metric space is closed.

Proof. Let K be compact in X. To prove K^c is open.
Fix $p \in K^c$. For each $q \in K$, put $V_q := B_{\frac{1}{2}d(p,q)}(p)$, $W_q := B_{\frac{1}{2}d(p,q)}(q)$.
The sets W_q form an open cover of K, so compactness gives $q_1, \ldots, q_n \in K$ with $K \subseteq W_{q_1} \cup \cdots \cup W_{q_n}$.
Let $V := V_{q_1} \cap \cdots \cap V_{q_n}$, which is open by Theorem 2.1.22.
No point of V lies in any W_{q_i}, by the triangle inequality. Hence $V \cap K = \emptyset$, so p is an interior point of K^c.
Structural Commentary: Compactness makes the separation finite.
Key Insight: A finite subcover lets local point-separations become one open neighborhood of the complement. □

Theorem 2.2.30. *Closed Subsets of Compact Sets Are Compact*
Suppose $F \subseteq K \subseteq X$, where F is closed in X and K is compact. Then F is compact.

Proof. Let $\{V_\alpha\}_{\alpha \in A}$ cover F. Since F^c is open, $\{V_\alpha\}_{\alpha \in A} \cup \{F^c\}$ covers K.
Thus $K \subseteq F^c \cup \bigcup_{\alpha \in F_0} V_\alpha$ for some finite $F_0 \subseteq A$, so $F \subseteq \bigcup_{\alpha \in F_0} V_\alpha$.
Structural Commentary: The open complement of a closed set covers exactly the part of K not under discussion.
Key Insight: Cover K, then discard the complement piece. □

Corollary 2.2.31. *Closed Intersections with Compact Sets*
If F is closed in X and K is compact in X, then $F \cap K$ is compact.

Proof. By Theorem 2.2.29, K is closed in X. Since F is closed, Theorem 2.1.22 gives that $F \cap K$ is closed in X. Also $F \cap K \subseteq K$, so Theorem 2.2.30 implies $F \cap K$ is compact.
Structural Commentary: The intersection is closed inside a compact set.
Key Insight: Closed subsets of compact sets remain compact. □

Theorem 2.2.32. *Finite Intersection Property for Compact Sets*
Let $\{K_\alpha\}_{\alpha \in A}$ be compact subsets of a metric space X, with $A \neq \emptyset$.
If $\bigcap_{\alpha \in F} K_\alpha \neq \emptyset$ for every finite nonempty $F \subseteq A$, then $\bigcap_{\alpha \in A} K_\alpha \neq \emptyset$.

Proof. Choose $\alpha_0 \in A$ and suppose $\bigcap_{\alpha \in A} K_\alpha = \emptyset$. Then $K_{\alpha_0} \subseteq \bigcup_{\alpha \in A} K_\alpha^c$. Since each K_α^c is open by Theorem 2.2.29, compactness of K_{α_0} gives a finite $F \subseteq A$ with $K_{\alpha_0} \subseteq \bigcup_{\alpha \in F} K_\alpha^c$. Hence $\bigcap_{\alpha \in F \cup \{\alpha_0\}} K_\alpha = \emptyset$, a contradiction.
Structural Commentary: Complements turn failure into a finite cover.
Key Insight: For compact sets, failure of intersection is already witnessed finitely. □

Corollary 2.2.33. *Nested Nonempty Compact Sets Intersect*
If $K_1 \supseteq K_2 \supseteq K_3 \supseteq \cdots$ is a sequence of nonempty compact sets, then $\bigcap_{n=1}^{\infty} K_n \neq \emptyset$.

Proof. Let $\emptyset \neq F \subseteq \mathbb{N}$ be finite and put $m := \max F$. Since the sequence is nested, $\bigcap_{n\in F} K_n = K_m \neq \emptyset$.

Thus $\{K_n\}_{n\in\mathbb{N}}$ has the finite intersection property, so Theorem 2.2.32 gives $\bigcap_{n=1}^{\infty} K_n \neq \emptyset$.

Structural Commentary: Nestedness reduces finite intersections to one term.
Key Insight: A nested compact sequence has the finite intersection property. □

Theorem 2.2.34. *Infinite Subsets of Compact Sets Have Limit Points*
Let $K \subseteq X$ be compact in a metric space X. If $E \subseteq K$ is infinite, then $E^{\text{acc}} \cap K \neq \emptyset$.

Proof. Suppose $E^{\text{acc}} \cap K = \emptyset$. For each $q \in K$, choose $r_q > 0$ such that $B_{r_q}(q) \cap (E \setminus \{q\}) = \emptyset$, so $\#(B_{r_q}(q) \cap E) \leq 1$.
Since $\{B_{r_q}(q) : q \in K\}$ covers K, compactness gives $q_1, \ldots, q_n \in K$ with $K \subseteq \bigcup_{i=1}^{n} B_{r_{q_i}}(q_i)$.

As $E \subseteq K$, $E \subseteq \bigcup_{i=1}^{n} B_{r_{q_i}}(q_i)$, finite union of sets meeting E in at most one point. Thus E is finite, contradiction.

Structural Commentary: Compactness makes the local "at most one point" obstruction finite.
Key Insight: Infinite subsets of compact sets must accumulate somewhere inside the compact set. □

Theorem 2.2.35. *Nested Closed Intervals Intersect*
If $I_1 \supseteq I_2 \supseteq I_3 \supseteq \cdots$ is a sequence of closed intervals in $\mathbb{R}$, then $\bigcap_{n=1}^{\infty} I_n \neq \emptyset$.

Proof. Let $I_n = [a_n, b_n]$. The set $\{a_n : n \in \mathbb{N}\}$ is nonempty and bounded above by each b_m.
Let $x = \sup\{a_n : n \in \mathbb{N}\}$. Then $a_m \leq x \leq b_m$ for every m, so $x \in I_m$ for all m.
Structural Commentary: Completeness traps a common point.
Key Insight: The least upper bound of the left endpoints lies in every interval. □

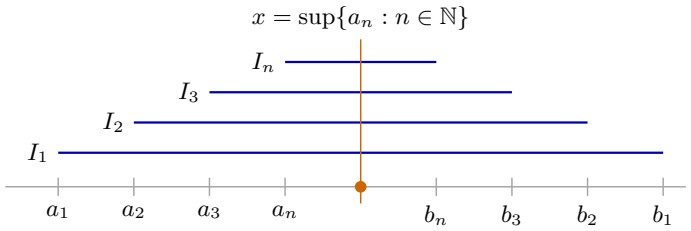


Figure 139: For nested closed intervals $[a_n, b_n]$, the supremum of the left endpoints lies between every a_n and b_n.

Theorem 2.2.36. *Nested k-Cells Intersect*
If $I_1 \supseteq I_2 \supseteq I_3 \supseteq \cdots$ is a sequence of k-cells in $\mathbb{R}^k$, then $\bigcap_{n=1}^{\infty} I_n \neq \emptyset$.

Proof. Let $I_n = \{\mathbf{x} \in \mathbb{R}^k : a_{n,j} \leq x_j \leq b_{n,j},\ 1 \leq j \leq k\}$. For each coordinate j, the intervals $[a_{n,j}, b_{n,j}]$ are nested, so Theorem 2.2.35 gives a point x_j^* lying in all of them. Then $\mathbf{x}^* = (x_1^*, \ldots, x_k^*)$ lies in every I_n.
Structural Commentary: A nested cell problem reduces coordinatewise to nested intervals.
Key Insight: Choose one persistent coordinate at a time. □

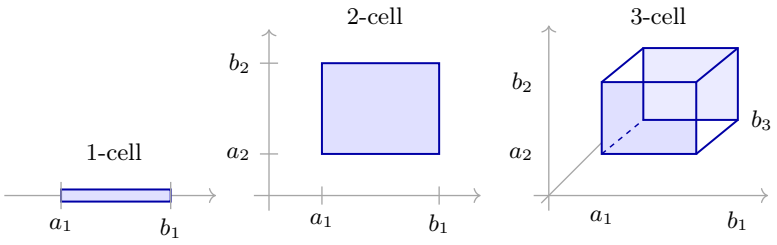


Figure 140: A 1-cell is a closed interval, a 2-cell is a coordinate rectangle, and a 3-cell is a coordinate box.

Theorem 2.2.37. *k-Cells Are Compact*
Every k-cell in $\mathbb{R}^k$ is compact.

Proof. Let $I = \{\mathbf{x} : a_j \le x_j \le b_j,\ 1 \le j \le k\}$ and set $\delta := \left(\sum_{j=1}^{k} (b_j - a_j)^2\right)^{1/2}$. Suppose an open cover $\{G_\alpha\}$ of I has no finite subcover. Bisect each coordinate interval. Among the resulting 2^k subcells, at least one cannot be finitely covered by the G_α; call it I_1. Continue inductively, obtaining nested cells $I \supseteq I_1 \supseteq I_2 \supseteq \cdots$ such that no I_n is finitely covered and every two points of I_n are within $2^{-n}\delta$ of each other.
By Theorem 2.2.36, choose $\mathbf{x}^* \in \bigcap_n I_n$. Since $\{G_\alpha\}$ covers I, some G_α contains $\mathbf{x}^*$. Because G_α is open, $B_r(\mathbf{x}^*) \subseteq G_\alpha$ for some $r > 0$. By the Archimedean property in Theorem 1.1.15, choose n with $n > \delta/r$; since $2^n \ge n$, we have $2^{-n}\delta < r$. Then $I_n \subseteq G_\alpha$, contradicting the construction of I_n.
Structural Commentary: Repeated bisection traps a point in cells whose diameters shrink to zero.
Key Insight: An open set around the trapped point eventually contains the whole small cell. □

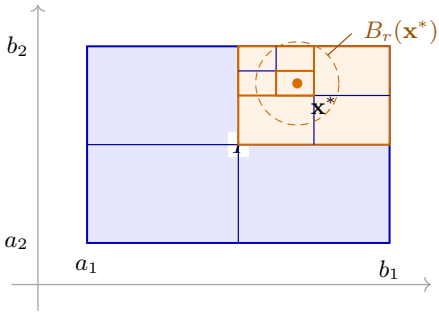


Figure 141: Repeated bisection traps a point.

Lemma 2.2.38. *Bounded Euclidean Sets Fit in k-Cells*
Every bounded subset $E \subseteq \mathbb{R}^k$ is contained in some k-cell.

Proof. If $E = \emptyset$, any k-cell contains E. Suppose $E \neq \emptyset$. By Definition 2.1.28, choose $\mathbf{q} = (q_1, \ldots, q_k) \in \mathbb{R}^k$ and $M > 0$ such that $\|\mathbf{p} - \mathbf{q}\|_2 < M$ for all $\mathbf{p} \in E$. Let $I := \{\mathbf{x} = (x_1, \ldots, x_k) : q_j - M \le x_j \le q_j + M,\ 1 \le j \le k\}$.
For $\mathbf{p} = (p_1, \ldots, p_k) \in E$, $|p_j - q_j| \le \|\mathbf{p} - \mathbf{q}\|_2 < M$ for each j, so $\mathbf{p} \in I$.
Thus $E \subseteq I$, and I is a k-cell by Definition 2.1.9.
Structural Commentary: A Euclidean ball centered at $\mathbf{q}$ sits inside the coordinate box with side half-length M.
Key Insight: Boundedness supplies one radius; coordinates turn it into interval bounds. □

Theorem 2.2.39. *Heine–Borel in $\mathbb{R}^k$*
For $E \subseteq \mathbb{R}^k$, any one of the following three properties implies the other two:

(i) E is closed and bounded;
(ii) E is compact;
(iii) every infinite subset of E has a limit point in E.

Proof.

(*i*) $\Rightarrow$ (*ii*): Since E is bounded, Lemma 2.2.38 gives $E \subseteq I$ for some k-cell I. By Theorem 2.2.37, I is compact; since E is closed, Theorem 2.2.30 gives that E is compact.
(*ii*) $\Rightarrow$ (*iii*): This is Theorem 2.2.34.
(*iii*) $\Rightarrow$ (*i*): If E is unbounded, choose $\mathbf{x}_n \in E$ with $\|\mathbf{x}_n\|_2 > n$ and set $S := \{\mathbf{x}_n : n \in \mathbb{N}\}$. If $\mathbf{x}_n \in B_R(\mathbf{y})$, then $n < \|\mathbf{y}\|_2 + R$; hence S has no limit point, contradicting (iii). Thus E is bounded. If E is not closed, choose $\mathbf{x}_0 \in E^{\text{acc}} \setminus E$ and $\mathbf{x}_n \in E$ with $0 < \|\mathbf{x}_n - \mathbf{x}_0\|_2 < 1/n$. Then $S := \{\mathbf{x}_n : n \in \mathbb{N}\}$ is infinite and $S^{\text{acc}} \subseteq \{\mathbf{x}_0\} \subseteq E^c$, contradicting (iii). Thus E is closed.

Structural Commentary: In $\mathbb{R}^k$, the three tests coincide.
Key Insight: Heine–Borel is the bridge between open-cover compactness and Euclidean bounded geometry. □

Corollary 2.2.40. *Heine–Borel on the Real Line*
For $E \subseteq \mathbb{R}$, the following are equivalent:

(i) E is closed and bounded;
(ii) E is compact;
(iii) every sequence in E has a subsequence converging to a point of E.

Proof.

(*i*) $\Rightarrow$ (*ii*) Since $E \subseteq \mathbb{R} = \mathbb{R}^1$ is closed and bounded, Theorem 2.2.39 with $k = 1$ implies E is compact.

(*ii*) $\Rightarrow$ (*iii*) Compactness and Theorem 2.2.26 give convergent subsequence; compactness keeps limit in E.
(*iii*) $\Rightarrow$ (*i*) Let $F \subseteq E$ be infinite. Choose distinct $x_n \in F$ using Principle 6.5.1. By (iii), some $x_{n_j} \to x \in E$. Since the x_n are distinct, x is a limit point of F in the sense of Definition 2.1.14. Hence every infinite subset of E has a limit point in E, so Theorem 2.2.39 gives that E is closed and bounded.

Structural Commentary: Real-line compactness equals closed-bounded sequential compactness.
Key Insight: Closed bounded sets keep sequences from escaping. □

Theorem 2.2.41. *Weierstrass Limit-Point Theorem*
Every bounded infinite subset of $\mathbb{R}^k$ has a limit point in $\mathbb{R}^k$.

Proof. By Lemma 2.2.38, a bounded set E lies in some k-cell I. By Theorem 2.2.37, I is compact. Since E is an infinite subset of I, Theorem 2.2.34 gives a limit point of E in $I \subseteq \mathbb{R}^k$.
Structural Commentary: Boundedness places the set inside a compact cell.
Key Insight: Compact cells force bounded infinite sets to accumulate. □

Theorem 2.2.42. *Perfect Sets Are Uncountable*
Every nonempty perfect subset of $\mathbb{R}^k$ is uncountable.

Proof. Assume $P = \{\mathbf{x}_1, \mathbf{x}_2, \ldots\}$. Choose a neighborhood V_1 with $V_1 \cap P \neq \emptyset$. Inductively, since $P \cap V_n$ contains infinitely many points by Theorem 2.1.18, choose $\mathbf{y}_n \in P \cap V_n$ with $\mathbf{y}_n \neq \mathbf{x}_n$ and then a neighborhood V_{n+1} of $\mathbf{y}_n$ such that $\overline{V}_{n+1} \subseteq V_n$ and $\mathbf{x}_n \notin \overline{V}_{n+1}$.

Put $K_n := \overline{V}_n \cap P$. Then K_n is nonempty, compact, and $K_{n+1} \subseteq K_n$. By Corollary 2.2.33, choose $\mathbf{z} \in \bigcap_{n=1}^{\infty} K_n$. For each n, $\mathbf{z} \in K_{n+1}$ but $\mathbf{x}_n \notin K_{n+1}$, so $\mathbf{z} \neq \mathbf{x}_n$ for all n, contradicting the enumeration.
Structural Commentary: Nested compact sets dodge the list.
Key Insight: Perfectness gives new points; compactness keeps one. □

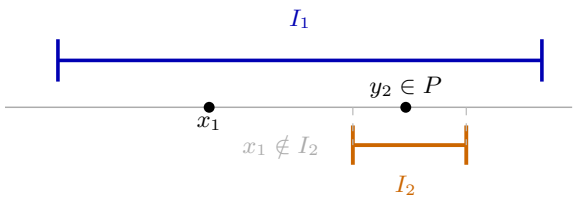


Figure 142: Choose $y_2 \in P \cap I_1$ with $y_2 \neq x_1$, then take a smaller interval $I_2 \subseteq I_1$ around y_2 that avoids x_1.

Corollary 2.2.43. *Intervals and the Real Line Are Uncountable*
Every interval $[a, b] \subseteq \mathbb{R}$ with $a < b$ is uncountable. In particular, $\mathbb{R}$ is uncountable.

Proof. The interval $I := [a, b]$ is a nonempty 1-cell, hence compact by Theorem 2.2.37 and closed by Theorem 2.2.29. For every $x \in I$ and $\varepsilon > 0$, $I \cap \big(B_\varepsilon(x) \setminus \{x\}\big) \neq \emptyset$, so I is perfect. Theorem 2.2.42 gives that I is uncountable. If $\mathbb{R}$ were countable, then $I \subseteq \mathbb{R}$ would be countable, a contradiction.
Structural Commentary: Closed intervals are perfect compact sets.
Key Insight: One closed interval already forces $\mathbb{R}$ to be uncountable. □

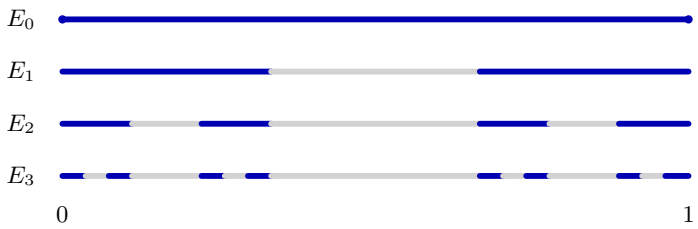


Figure 143: The first stages E_0, E_1, E_2, E_3 in the Cantor-set construction.

Example 2.2.44. *The Cantor Set*
Let $E_0 = [0, 1]$. Remove the open middle third $(1/3, 2/3)$ and call the remaining set E_1. Continue by removing the open middle third from each interval component, so

$$E_1 = [0, 1/3] \cup [2/3, 1],$$
$$E_2 = [0, 1/9] \cup [2/9, 3/9] \cup [6/9, 7/9] \cup [8/9, 1].$$

Then E_n is a union of 2^n compact intervals of length 3^{-n} and $E_1 \supseteq E_2 \supseteq E_3 \supseteq \cdots$.
The *Cantor set* is $P := \bigcap_{n=1}^{\infty} E_n$. By Corollary 2.2.33, $P \neq \emptyset$. Since each E_n is closed and $P \subseteq E_1$, P is a closed subset of the compact set E_1, hence compact.
If $J \subseteq [0,1]$ is nonempty open interval, for some $m \in \mathbb{N}$ and $k \geq 0$, $R_{m,k} := \left(\frac{3k+1}{3^m}, \frac{3k+2}{3^m}\right) \subseteq J$ and $R_{m,k} \cap P = \emptyset$. Thus P contains no nontrivial interval. Finally, let $x \in P$ and let U be an interval neighborhood of x. If $I_n(x)$ is the component of E_n containing x, choose n with $I_n(x) \subseteq U$. An endpoint $y \neq x$ of $I_n(x)$ lies in P, so $U \cap (P \setminus \{x\}) \neq \emptyset$. Hence P is perfect.
Structural Commentary: The Cantor set is built by nested compact deletion, not by listing points.
Key Insight: It is compact and perfect, but contains no interval.

Definition 2.2.45. *F_σ, G_δ, and Nowhere-Dense Sets*
For subsets of a metric space X:

(i) A is an *F_σ set* if $A = \bigcup_{n=1}^{\infty} F_n$ for closed sets F_n.
(ii) B is a *G_δ set* if $B = \bigcap_{n=1}^{\infty} G_n$ for open sets G_n.
(iii) E is *nowhere dense* if $(\overline{E})^\circ = \emptyset$; equivalently, no nonempty open set is contained in $\overline{E}$.

For $X = \mathbb{R}$, this says exactly that $\overline{E}$ contains no nonempty open interval.

Definition 2.2.46. *First Category*
A set $E \subseteq X$ is *of first category*, *meager*, or *thin*, if covered by countably many nowhere-dense sets: $E \subseteq \bigcup_{n=1}^{\infty} E_n$.
The Baire Category Theorem states that a nonempty complete metric space is not first category in itself.

Definition 2.2.47. *Residual, Thick, and Generic Sets*
A set $G \subseteq X$ is *residual*, or *thick*, if it contains a countable intersection $\bigcap_{n=1}^{\infty} O_n$ of dense open subsets of X.
A property is *generic* in X if the set of points satisfying it is residual.
By Theorem 2.2.49, residual subsets of a complete metric space are dense.
In this sense, meager sets are topologically small and residual sets are topologically large.

Theorem 2.2.48. *Dense Open Intersections in $\mathbb{R}$*
If $G_1, G_2, G_3, \ldots$ are dense open subsets of $\mathbb{R}$, then $\bigcap_{n=1}^{\infty} G_n \neq \emptyset$.

Proof. We construct nested closed intervals. Since G_1 is dense and open, $G_1 \cap (0,1)$ contains a nonempty open interval; choose a closed interval $I_1 \subseteq G_1 \cap (0,1)$ with nonempty interior.
Given $I_n = [a_n, b_n]$, density of G_{n+1} gives a point of G_{n+1} in (a_n, b_n). Since G_{n+1} is open, choose a closed interval $I_{n+1} \subseteq G_{n+1} \cap (a_n, b_n)$ with nonempty interior and length $< 1/(n+1)$.
Then $I_1 \supseteq I_2 \supseteq I_3 \supseteq \cdots$, so Theorem 2.2.35 gives $x \in \bigcap_n I_n$. Since $I_n \subseteq G_n$ for every n, this x lies in $\bigcap_n G_n$.
Structural Commentary: Dense openness builds nested intervals; completeness keeps a limit point.
Key Insight: A countable dense-open intersection cannot be destroyed one interval at a time. □

Theorem 2.2.49. *Dense Open Intersections in Complete Metric Spaces*
Let (X,d) be complete, let $O_1, O_2, \ldots$ be dense open subsets of X, and let $U \subseteq X$ be nonempty and open. Then $U \cap \bigcap_{n=1}^{\infty} O_n \neq \emptyset$. Consequently, $\bigcap_{n=1}^{\infty} O_n$ is dense in X.

Proof. Choose $x_1 \in U \cap O_1$ and $0 < r_1 < 1$ with $C_1 := \overline{B}_{r_1}(x_1) \subseteq U \cap O_1$.
Inductively, if $C_n = \overline{B}_{r_n}(x_n) \subseteq U \cap O_1 \cap \cdots \cap O_n$ and $0 < r_n < 2^{-n}$, density and openness of O_{n+1} give $x_{n+1} \in B_{r_n}(x_n) \cap O_{n+1}$ and $0 < r_{n+1} < 2^{-(n+1)}$ such that $C_{n+1} := \overline{B}_{r_{n+1}}(x_{n+1}) \subseteq B_{r_n}(x_n) \cap O_{n+1}$.
Then $C_{n+1} \subseteq C_n$. If $m > n$, then $x_m \in B_{r_n}(x_n)$, so $d(x_m, x_n) < 2^{-n}$; hence (x_n) is Cauchy. Let $x_n \to x$.
Since each C_n is closed and contains all sufficiently late x_m, we have $x \in C_n$ for every n. Thus $x \in U \cap \bigcap_n O_n$.
Since every nonempty open U meets $\bigcap_n O_n$, the intersection is dense.
Structural Commentary: The interval proof becomes a nested closed-ball proof.
Key Insight: Dense open sets choose the next ball; completeness supplies the limiting point. □

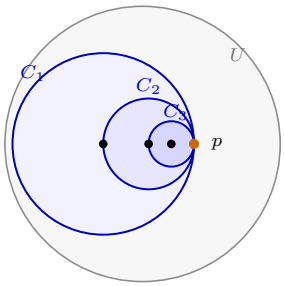


Figure 144: Nested closed balls shrink to p.

Theorem 2.2.50. *Baire Category Theorem for Complete Metric Spaces*
A nonempty complete metric space is not the union of a countable collection of nowhere-dense sets.

Proof. Suppose $X = \bigcup_n E_n$ with each E_n nowhere dense, and set $O_n := (\overline{E_n})^c$.
Then each O_n is dense open: it is open, and every nonempty open set meets it because no such set lies in $\overline{E_n}$.
By Theorem 2.2.49, choose $x \in \bigcap_n O_n$. Then $x \notin E_n$ for all n, contradicting $X = \bigcup_n E_n$.
Structural Commentary: Countably many nowhere-dense sets have dense open complements.
Key Insight: Completeness forbids countable assembly from thin pieces. □

Perspective 2.2.51. *Category in Function Spaces*
Baire category is a theorem about complete metric spaces, not just $\mathbb{R}$.
It will be reused for complete function spaces under the supremum metric.

Theorem 2.2.52. *Baire's Theorem for $\mathbb{R}$*
The real line $\mathbb{R}$ cannot be written as a countable union of nowhere-dense sets.

Proof. Assume $\mathbb{R} = \bigcup_n E_n$ with each E_n nowhere dense, and set $U_n := (\overline{E_n})^c$.
Each U_n is dense open: it is open, and every nonempty interval meets it because no such interval lies in $\overline{E_n}$.
By Theorem 2.2.48, choose $x \in \bigcap_n U_n$. Then $x \notin E_n$ for all n, contradicting $\mathbb{R} = \bigcup_n E_n$.
Structural Commentary: Nowhere-dense sets have dense open complements.
Key Insight: Countably many thin sets cannot exhaust the real line. □

Corollary 2.2.53. *Rationals and Irrationals by Category*
The set $\mathbb{Q}$ is F_σ but not G_δ. The set $\mathbb{R} \setminus \mathbb{Q}$ is G_δ but not F_σ.

Proof. Write $\mathbb{Q} = \{q_1, q_2, \ldots\}$. Since singletons are closed, $\mathbb{Q} = \bigcup_n \{q_n\}$ is F_σ, and $\mathbb{R} \setminus \mathbb{Q} = \bigcap_n (\{q_n\})^c$ is G_δ.
If $\mathbb{Q} = \bigcap_n G_n$ with each G_n open, then each G_n is dense because it contains the dense set $\mathbb{Q}$. Set $H_n := G_n \setminus \{q_n\}$.
Each H_n is dense open, so Theorem 2.2.48 gives a point of $\bigcap_n H_n$. But $\bigcap_n H_n = \emptyset$, since $\bigcap_n G_n = \mathbb{Q}$ and q_n is removed at stage n. Hence $\mathbb{Q}$ is not G_δ. Therefore $\mathbb{R} \setminus \mathbb{Q}$ is not F_σ, or else its complement $\mathbb{Q}$ would be G_δ.
Structural Commentary: Category separates the familiar countable dense set from its dense complement.
Key Insight: Dense is not the same as large in the Baire-category sense. □

Method 2.2.54. *Baire-Category Proof Patterns*

(i) *Dense-open intersection*: build nested closed intervals $I_n \subseteq G_n$ and apply Theorem 2.2.35.
(ii) *Complete metric version*: build nested closed balls $C_{n+1} \subseteq C_n$ with $r_n \to 0$ and use completeness.
(iii) *Nowhere dense*: replace E by the dense open set $(\overline{E})^c$.
(iv) *First category*: show the exceptional set is contained in $\bigcup_n E_n$ with each E_n nowhere dense.
(v) *Residual/generic*: show the desired property contains $\bigcap_n O_n$ with each O_n dense open.
(vi) *F_σ/G_δ complements*: pass between countable unions and intersections by De Morgan's laws.

2.3 Connectedness and Continuity

Connectedness prevents a space from splitting into separated pieces; continuity transports topological structure through maps. This subsection also records IVT/EVT, uniform continuity, and real-line refinements.

Definition 2.3.1. *Separated and Connected Sets*
Let X be a metric space.

(i) Subsets $A, B \subseteq X$ are *separated* if $A \cap \overline{B} = \overline{A} \cap B = \emptyset$.
(ii) A subset $E \subseteq X$ is *connected* if no nonempty separated subsets $A, B \subseteq E$ satisfy $E = A \cup B$.
(iii) A subset $U \subseteq X$ is *clopen* if it is both open and closed.

Structural Commentary: Separated sets avoid each other's closures.
Key Insight: Disconnected means a nontrivial clopen split; continuous preimages preserve such splits.

Theorem 2.3.2. *Connected Subsets of the Real Line*
A set $E \subseteq \mathbb{R}$ is connected if and only if $x, y \in E,\ x < y \implies [x, y] \subseteq E$.

Proof.

($\Rightarrow$) If the interval property fails, choose $x, y \in E$ and $x < z < y$ with $z \notin E$. Then $A := E \cap (-\infty, z)$ and $B := E \cap (z, \infty)$ are nonempty separated sets with $E = A \cup B$, contradicting connectedness.
($\Leftarrow$) Suppose $E = A \cup B$ with A, B nonempty separated. Choose $x \in A$, $y \in B$, $x < y$, and set $S := A \cap [x, y]$, $z := \sup S$. By Theorem 2.1.39, $z \in \overline{A}$; hence $z \notin B$ and $z < y$. If $z \notin A$, then $z \in [x, y] \setminus E$.
If $z \in A$, then $z \notin \overline{B}$, so choose $z_1 \in (z, y)$ with $z_1 \notin B$; by maximality of z, also $z_1 \notin A$.
Thus $z_1 \in [x, y] \setminus E$, contradicting $[x, y] \subseteq E$.

Structural Commentary: On the real line, connectedness is exactly having no gaps between any two of its points.
Key Insight: A separation creates a missing point by taking a supremum at the boundary of one side. □

Method 2.3.3. *Compactness and Connectedness Proof Patterns*

(i) *Open-cover compactness*: for every cover $\{G_\alpha\}_{\alpha \in A}$, find finite $F \subseteq A$ with $K \subseteq \bigcup_{\alpha \in F} G_\alpha$.
(ii) *Euclidean closed-bounded compactness*: if $E \subseteq \mathbb{R}^k$ is closed and bounded, choose a compact k-cell $I \supseteq E$ and apply Theorem 2.2.30.
(iii) *Connectedness on $\mathbb{R}$*: replace connectedness of $E \subseteq \mathbb{R}$ by $x, y \in E, x < y \implies [x, y] \subseteq E$.

Definition 2.3.4. *Limit of a Function Between Metric Spaces*
Let X, Y be metric spaces, $E \subseteq X$, $f : E \to Y$, and let p be a limit point of E. For $q \in Y$, $\lim_{x \to p} f(x) = q$ means that, for every $\varepsilon > 0$, there exists $\delta > 0$ such that $x \in E$, $0 < d_X(x, p) < \delta \implies d_Y(f(x), q) < \varepsilon$.
The point p need not belong to E, and $f(p)$ is irrelevant when it does.
Key Insight: Domain E controls approach: $J \setminus \{a\}$ gives two-sided limit, (a, b) or (c, a) gives one-sided limits.

Remark 2.3.5. *Adherent-Point Limit Convention*
The adherent-point convention takes $p \in \overline{E}$ and allows $x = p$ in the δ-test; Definition 2.3.4 instead assumes p is a limit point and requires $0 < d_X(x, p)$. If $p \in E$, the unpunctured convention also tests $f(p)$, forcing $q = f(p)$; in the punctured convention this is exactly continuity at p by Theorem 2.3.12.
Key Insight: Punctured limits study nearby values; continuity adds compatibility with the value at the point.

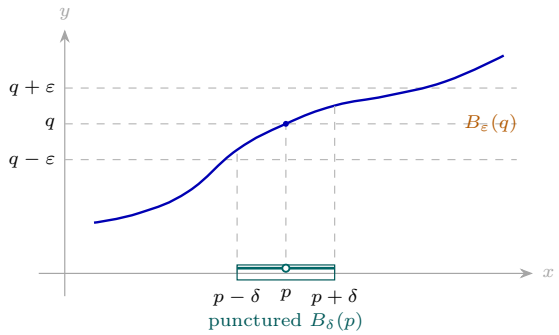


Figure 145: The epsilon–delta geometry for $\lim_{x \to p} f(x) = q$.

Theorem 2.3.6. *Sequential Characterization of Function Limits*
Let X, Y be metric spaces, $E \subseteq X$, $f : E \to Y$, and let p be a limit point of E. For $q \in Y$, $\lim_{x \to p} f(x) = q$ if and only if $f(p_n) \to q$ for every sequence $(p_n) \subseteq E \setminus \{p\}$ with $p_n \to p$.

Proof.

($\Rightarrow$) If $p_n \to p$ with $p_n \in E \setminus \{p\}$, then eventually $0 < d_X(p_n, p) < \delta$; hence $d_Y(f(p_n), q) < \varepsilon$, so $f(p_n) \to q$.
($\Leftarrow$) If the limit fails, then for some $\varepsilon > 0$ and each n there is $x_n \in E$ with $0 < d_X(x_n, p) < 1/n$ and $d_Y(f(x_n), q) \geq \varepsilon$. Then $x_n \to p$, $x_n \neq p$, but $f(x_n) \not\to q$.

Structural Commentary: Function limits can be tested by all sequences approaching the input point.
Key Insight: The negation of an epsilon–delta limit manufactures a bad sequence. □

Principle 2.3.7. *Challenge–Response Order*
Limits and continuity are controlled by quantifier order. In Definition 2.3.4, the ε-challenge at the range point must be answered by a δ-restriction near the domain point. Uniform continuity, Definition 2.3.37, strengthens the same pattern by requiring one δ to work for all relevant domain points at once.

Corollary 2.3.8. *Uniqueness of Function Limits*
If a function has a limit at p, then that limit is unique.

Proof. Let q and r be limits of f at p. Choose $p_n \in E \setminus \{p\}$ with $d_X(p_n, p) < 1/n$. Then $p_n \to p$, so Theorem 2.3.6 gives $f(p_n) \to q$ and $f(p_n) \to r$. By Theorem 2.1.30, $q = r$.
Structural Commentary: Limit uniqueness is inherited from metric sequence uniqueness.
Key Insight: One approaching sequence cannot have two different limits. □

Theorem 2.3.9. *Algebra of Function Limits*
Let X be a metric space, $E \subseteq X$, p be limit point of E.
For $f, g : E \to \mathbb{C}$, assume $\lim_{x\to p} f(x) = A$ and $\lim_{x\to p} g(x) = B$. Then:

(i) $\lim\limits_{x\to p}(f+g)(x) = A + B$;
(ii) $\lim\limits_{x\to p}(fg)(x) = AB$;
(iii) if $B \neq 0$ and f/g is defined near p, then $\lim\limits_{x\to p}(f/g)(x) = A/B$.

For $f, g : E \to \mathbb{R}^k$ with limits a, b, the corresponding limits are $f + g \to a + b$, $\lambda f \to \lambda a$, and $f \cdot g \to a \cdot b$.

Proof.

(i) For any $p_n \in E \setminus \{p\}$ with $p_n \to p$, Theorem 2.3.6 gives $f(p_n) \to A$ and $g(p_n) \to B$. Hence $(f+g)(p_n) \to A + B$ by Theorem 1.3.3, so Theorem 2.3.6 gives (i).
(ii) The same sequence gives $(fg)(p_n) \to AB$ by Theorem 1.3.3; hence (ii).
(iii) If $B \neq 0$ and f/g is defined near p, then every such sequence is eventually in the domain of f/g, and $(f/g)(p_n) \to A/B$ by Theorem 1.3.3; hence (iii).

The vector laws follow by the same sequential characterization and Theorem 1.3.4.
Structural Commentary: Limit algebra for functions is pulled back to sequence algebra.
Key Insight: Sequential characterization lets previous sequence theorems do the work. □

Proposition 2.3.10. *Function Limits Are Local*
Let X, Y be metric spaces, let $E \subseteq X$, let $f : E \to Y$, let p be a limit point of E, and let $q \in Y$.
For every $r > 0$, $\lim\limits_{x\to p;\, x\in E} f(x) = q \iff \lim\limits_{x\to p;\, x\in E\cap B_r(p)} f(x) = q$.

Proof.

($\Rightarrow$) Since $E \cap B_r(p) \subseteq E$, any δ witnessing the limit on E also witnesses the limit on $E \cap B_r(p)$.
($\Leftarrow$) Assume the restricted limit. Given $\varepsilon > 0$, choose $\delta_0 > 0$ such that $x \in E \cap B_r(p)$ and $0 < d_X(x, p) < \delta_0$ imply $d_Y(f(x), q) < \varepsilon$. Let $\delta := \min\{\delta_0, r\}$.
If $x \in E$ and $0 < d_X(x, p) < \delta$, then $x \in E \cap B_r(p)$, so $d_Y(f(x), q) < \varepsilon$.

Structural Commentary: Values away from p cannot affect the limit at p.
Key Insight: Shrink δ until every tested point lies in the chosen local neighborhood. □

Definition 2.3.11. *Continuity at a Point and on a Set*
Let $(X, d_X), (Y, d_Y)$ be metric spaces, $E \subseteq X$, and $f : E \to Y$.

(i) For $p \in E$, f is *continuous at p* if $\forall \varepsilon > 0\ \exists \delta > 0\ \forall x \in E$, $d_X(x, p) < \delta \implies d_Y(f(x), f(p)) < \varepsilon$.
(ii) f is *continuous on E* if $\forall p \in E$, f is continuous at p.
(iii) If p is isolated in E, then f is continuous at p.

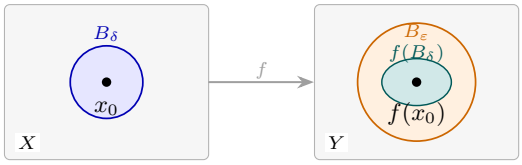


Figure 146: Continuity at x_0: for every $\varepsilon > 0$, choose $\delta > 0$ with $f(B_\delta(x_0)) \subset B_\varepsilon(f(x_0))$.

Theorem 2.3.12. *Continuity as Equality with the Function Limit*
Let X, Y be metric spaces, $E \subseteq X$, $p \in E$, and $f : E \to Y$.
If p is a limit point of E, then f is continuous at p if and only if $\lim_{x \to p} f(x) = f(p)$.

Proof.
($\Rightarrow$) If $0 < d_X(x,p) < \delta$, continuity at p gives $d_Y(f(x), f(p)) < \varepsilon$, hence $\lim_{x \to p} f(x) = f(p)$.
($\Leftarrow$) If $\lim_{x \to p} f(x) = f(p)$, the same δ proves continuity for $x \neq p$; for $x = p$, $d_Y(f(p), f(p)) = 0$.
Structural Commentary: Continuity is the compatibility of the value at p with nearby values.
Key Insight: Limits ignore the point; continuity ties the point back in. □

Remark 2.3.13. *The Sequence-First View*
For functions on subsets of $\mathbb{R}$, continuity at limit points is equivalently sequential: $x_n \to p$ in the domain implies $f(x_n) \to f(p)$; isolated points are automatic by Theorems 2.3.6 and 2.3.12.

Theorem 2.3.14. *Sequential Characterization of Continuity*
Let $(X, d_X), (Y, d_Y)$ be metric spaces, $f : X \to Y$, and $p \in X$.
Then f is continuous at p if and only if $\forall (x_n) \subseteq X$, $x_n \to p \implies f(x_n) \to f(p)$.

Proof.
($\Rightarrow$) Assume f is continuous at p and let $x_n \to p$.
Given $\varepsilon > 0$, choose $\delta > 0$ such that $d_X(x,p) < \delta \implies d_Y(f(x), f(p)) < \varepsilon$.
Since $x_n \to p$, eventually $d_X(x_n, p) < \delta$, hence eventually $d_Y(f(x_n), f(p)) < \varepsilon$. Thus $f(x_n) \to f(p)$.
($\Leftarrow$) Suppose the sequential condition holds. If f is not continuous at p, then $\exists \varepsilon_0 > 0$ such that $\forall \delta > 0$, $\exists x \in X$ with $d_X(x,p) < \delta$ and $d_Y(f(x), f(p)) \geq \varepsilon_0$. Taking $\delta = 1/n$ gives $x_n \in X$ with $d_X(x_n, p) < 1/n$ and $d_Y(f(x_n), f(p)) \geq \varepsilon_0$. Then $x_n \to p$, but $f(x_n) \not\to f(p)$, a contradiction.
Structural Commentary: Failure of epsilon–delta continuity manufactures a bad sequence.
Key Insight: In metric spaces, sequences detect continuity at points. □

Perspective 2.3.15. *Three Views of Continuity*
For metric-space maps, continuity is equivalently ε–δ control, $x_n \to x \Rightarrow f(x_n) \to f(x)$, and $V \subseteq Y$ open $\Rightarrow f^{-1}(V)$ open; see Definition 2.3.11 and Theorems 2.3.14, 2.3.18.
Key Insight: Continuity sends near points forward and pulls open sets back.

Theorem 2.3.16. *Composition of Continuous Functions*
Let $E \subseteq X$ and let $f : E \to Y$, $g : f(E) \to Z$ be maps between metric spaces.
If f is continuous at $p \in E$ and g is continuous at $f(p)$, then $g \circ f$ is continuous at p.

Proof. Given $\varepsilon > 0$, choose $\eta > 0$ for g at $f(p)$, then choose $\delta > 0$ for f at p with tolerance η.
If $x \in E$ and $d_X(x,p) < \delta$, then $d_Y(f(x), f(p)) < \eta$, hence $d_Z(g(f(x)), g(f(p))) < \varepsilon$.
Structural Commentary: The output tolerance for g becomes the input tolerance for f.
Key Insight: Continuity composes by passing epsilon through an intermediate eta. □

Theorem 2.3.17. *The Distance Function Is Continuous*
Let (X, d) be a metric space and $D : X \times X \to \mathbb{R}$, $D(x,y) := d(x,y)$. For $d_\times \in \{d_1, d_2, d_\infty\}$ on $X \times X$, $D : (X \times X, d_\times) \to \mathbb{R}$ is continuous. Equivalently, $x_n \to x$ and $y_n \to y$ in X imply $d(x_n, y_n) \to d(x,y)$.

Proof. For $x, y, u, v \in X$, $|D(x,y) - D(u,v)| = |d(x,y) - d(u,v)| \leq d(x,u) + d(y,v)$. Hence $x_n \to x$ and $y_n \to y$ imply $D(x_n, y_n) \to D(x,y)$. Proposition 2.1.12 and Theorem 2.3.14 give continuity for each product metric.
Structural Commentary: Distance changes continuously when both endpoints move.
Key Insight: The triangle inequality controls distance error by endpoint errors. □

Theorem 2.3.18. *Continuity via Open Preimages*
For metric spaces X, Y, map $f : X \to Y$ is continuous if and only if $f^{-1}(V)$ is open in $X \forall$ open $V \subseteq Y$.

Proof.
($\Rightarrow$) If $p \in f^{-1}(V)$ and V is open, choose $\varepsilon > 0$ with $B_\varepsilon(f(p)) \subseteq V$.
Continuity gives $\delta > 0$ with $B_\delta(p) \subseteq f^{-1}(V)$, so $f^{-1}(V)$ is open.
($\Leftarrow$) For $p \in X$ and $\varepsilon > 0$, the ball $V = B_\varepsilon(f(p))$ is open by Theorem 2.1.16.
Since $p \in f^{-1}(V)$ and $f^{-1}(V)$ is open, some $B_\delta(p) \subseteq f^{-1}(V)$; this is continuity at p.
Structural Commentary: Continuity is exactly preservation of openness under inverse image.
Key Insight: Images are not Boolean-exact, but preimages are; continuity belongs to preimages. □

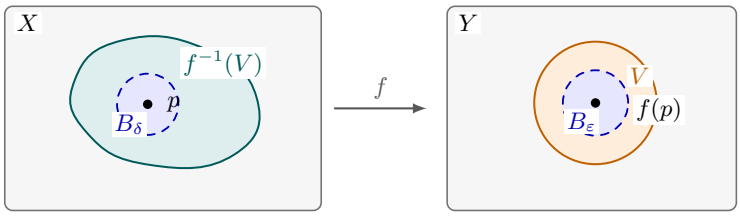


Figure 147: Continuity pulls open sets back to open sets.

Corollary 2.3.19. *Continuity via Closed Preimages*
For metric spaces X, Y, map $f : X \to Y$ is continuous if and only if $f^{-1}(C)$ is closed in $X \forall$ closed $C \subseteq Y$.

Proof.
($\Rightarrow$) If C is closed, then $Y \setminus C$ is open. By Theorem 2.3.18, $f^{-1}(Y \setminus C)$ is open; by Proposition 5.1.2, this is $X \setminus f^{-1}(C)$, so $f^{-1}(C)$ is closed.
($\Leftarrow$) If V is open, then $Y \setminus V$ is closed. Hence $f^{-1}(Y \setminus V)$ is closed, so $f^{-1}(V) = X \setminus f^{-1}(Y \setminus V)$ is open. Theorem 2.3.18 gives continuity.

Structural Commentary: Closed-set continuity is the complement form of open-set continuity.
Key Insight: $f^{-1}(Y \setminus C) = X \setminus f^{-1}(C)$ transfers open and closed formulations. □

Remark 2.3.20. *Preimages, Not Forward Images*
Theorem 2.3.18 is about inverse images. A continuous map need not send open sets to open sets or closed sets to closed sets; compactness is one hypothesis that recovers closedness of images through Theorem 2.3.33.

Definition 2.3.21. *Homeomorphism*
A *homeomorphism* from X to Y is a bijection $f : X \to Y$ such that both f and f^{-1} are continuous.
If such a map exists, X and Y are *homeomorphic*, written $X \simeq Y$.

Proposition 2.3.22. *Homeomorphisms Preserve Topology*
If $f : X \to Y$ is a homeomorphism, then $U \mapsto f(U)$ bijects open subsets of X with open subsets of Y.

Proof. If $U \subseteq X$ is open, then $f(U) = (f^{-1})^{-1}(U)$ is open in Y since f^{-1} is continuous. If $V \subseteq Y$ is open, then $f^{-1}(V)$ is open in X since f is continuous. These two assignments are inverse because f is bijective.
Structural Commentary: A homeomorphism preserves open sets in both directions.
Key Insight: Topological sameness is continuity plus continuous reversibility. □

Definition 2.3.23. *Metric Isometry*
Let $(X, d_X), (Y, d_Y)$ be metric spaces.
(i) A map $f : X \to Y$ is an *isometry* if $\forall x, y \in X$, $d_Y(f(x), f(y)) = d_X(x, y)$.
(ii) X and Y are *isometric* if $\exists$ a bijective isometry $f : X \to Y$.

Theorem 2.3.24. *Basic Properties of Metric Isometries*
Let $f : (X, d_X) \to (Y, d_Y)$ be an isometry.
(i) f is injective.
(ii) f is continuous.
(iii) $f^{-1} : (f(X), d_Y|_{f(X)\times f(X)}) \to (X, d_X)$ is an isometry.

Proof.
(i) If $f(x) = f(y)$, then $d_X(x, y) = d_Y(f(x), f(y)) = 0$, hence $x = y$.
(ii) If $x_n \to x$ in X, then $d_Y(f(x_n), f(x)) = d_X(x_n, x) \to 0$; by Theorem 2.3.14, f is continuous.
(iii) For $u = f(x), v = f(y) \in f(X)$, $d_X(f^{-1}(u), f^{-1}(v)) = d_X(x, y) = d_Y(f(x), f(y)) = d_Y(u, v)$; QED.

Structural Commentary: An isometry preserves all distances, so injectivity and continuity are automatic.
Key Insight: Distance preservation is stronger than topological equivalence. □

Remark 2.3.25. *Metric Versus Linear Isometries*
Metric isometry is purely distance-theoretic. If $T : V \to W$ is linear between inner-product spaces, then $\|Tv\| = \|v\| \ \forall v \in V \iff d_W(Tu, Tv) = d_V(u, v) \ \forall u, v \in V$.

Example 2.3.26. *Continuous Bijection Need Not Be a Homeomorphism*
Let $f : [0, 2\pi) \to S^1$, and $f(t) = (\cos t, \sin t)$. This continuous bijection is not a homeomorphism: points of S^1 approaching $(1, 0)$ from below have preimages tending to 2π, not 0.

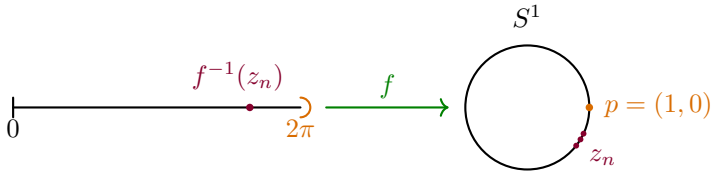


Figure 148: Parametrization $[0, 2\pi) \to S^1$ wraps continuously and bijectively, but inverse jumps at $(1, 0)$.

Remark 2.3.27. *Continuity Only Uses the Domain Metric*
For $f : E \to Y$ with $E \subseteq X$, continuity on E depends only on the restricted metric on E, not on $X \setminus E$.

Definition 2.3.28. *Product Map*
If $f : X \to Y$ and $g : X \to Z$ are maps between metric spaces, their *product map*, or direct sum, is $f \oplus g : X \to Y \times Z$, $(f \oplus g)(x) := (f(x), g(x))$, where $Y \times Z$ carries any of product metrics in Definition 2.1.11.

Proposition 2.3.29. *Continuity of Product Maps*
The product map $f \oplus g : X \to Y \times Z$ is continuous at x_0 iff f and g are continuous at x_0; hence it is continuous on X iff both components are.

Proof. By Proposition 2.1.12, it suffices to use d_{sum}.

($\Rightarrow$) Given $\varepsilon > 0$, choose $\delta > 0$ from continuity of $f \oplus g$ at x_0.
If $d_X(x, x_0) < \delta$, then $d_Y(f(x), f(x_0)), d_Z(g(x), g(x_0)) \le d_{\text{sum}}((f \oplus g)(x), (f \oplus g)(x_0)) < \varepsilon$.
($\Leftarrow$) Given $\varepsilon > 0$, choose δ_f, δ_g for f, g at tolerance $\varepsilon/2$ and set $\delta = \min\{\delta_f, \delta_g\}$.
Then $d_X(x, x_0) < \delta$ gives $d_{\text{sum}}((f \oplus g)(x), (f \oplus g)(x_0)) = d_Y(f(x), f(x_0)) + d_Z(g(x), g(x_0)) < \varepsilon$.

Structural Commentary: Product continuity is exactly componentwise continuity.
Key Insight: A product metric packages two output estimates into one. □

Theorem 2.3.30. *Continuous Algebra and Components*
Let X be a metric space.

(i) If $f, g : X \to \mathbb{C}$ are continuous, so are $f + g$, fg, and f/g where $g \neq 0$.
(ii) For $f = (f_1, \ldots, f_k) : X \to \mathbb{R}^k$, f is continuous if and only if each f_j is continuous.
(iii) If $f, g : X \to \mathbb{R}^k$ are continuous, so are $f + g$ and $f \cdot g$.

Proof.

(i) At isolated points continuity is automatic; at limit points, use Theorems 2.3.9 and 2.3.12.
(ii) By Theorem 1.3.4, convergence in $\mathbb{R}^k$ is exactly coordinatewise convergence; apply this to $f(x_n) \to f(x)$.
(iii) Apply (ii) to the coordinate functions of $f + g$ and $f \cdot g$, using the scalar algebra from (i).

Structural Commentary: Continuity is stable under the algebraic operations already stable for limits.
Key Insight: Componentwise control is equivalent to Euclidean control in finite dimensions. □

Example 2.3.31. *Coordinate, Polynomial, Rational, and Norm Maps*
Coordinate functions on $\mathbb{R}^k$ are continuous; hence polynomials are continuous, and rational functions are continuous where defined. Also $\mathbf{x} \mapsto \|\mathbf{x}\|_2$ is continuous since $\big|\|\mathbf{x}\|_2 - \|\mathbf{y}\|_2\big| \le \|\mathbf{x} - \mathbf{y}\|_2$.
Thus $p \mapsto \|f(p)\|_2$ is continuous whenever $f : X \to \mathbb{R}^k$ is, by Theorem 2.3.16.

Definition 2.3.32. *Bounded Mapping*
A mapping $f : E \to \mathbb{R}^k$ is *bounded* if there exists $M \in \mathbb{R}$ such that $\|f(x)\|_2 \le M$ for every $x \in E$.

Theorem 2.3.33. *Continuous Images of Compact Sets Are Compact*
If $f : X \to Y$ is continuous and X is compact, then $f(X)$ is compact.

Proof. Let $\{V_\alpha\}$ cover $f(X)$. Then $\{f^{-1}(V_\alpha)\}$ is an open cover of X by Theorem 2.3.18.
Choose $\alpha_1, \ldots, \alpha_n$ with $X \subseteq \bigcup_{i=1}^n f^{-1}(V_{\alpha_i})$; then $f(X) \subseteq \bigcup_{i=1}^n V_{\alpha_i}$.
Structural Commentary: Pull back, take a finite subcover, push forward.
Key Insight: Compactness is preserved by continuous images because inverse images preserve openness. □

Corollary 2.3.34. *Continuous Images into $\mathbb{R}^k$ Are Closed and Bounded*
If $f : X \to \mathbb{R}^k$ is continuous and X is compact, then $f(X)$ is closed and bounded. In particular, f is bounded.

Proof. By Theorem 2.3.33, $f(X)$ is compact.
By Theorem 2.2.39, compact subsets of $\mathbb{R}^k$ are closed and bounded; hence f is bounded.
Structural Commentary: Compactness in the domain becomes closed-bounded geometry in Euclidean range.
Key Insight: Continuous functions on compact spaces cannot escape to infinity. □

Theorem 2.3.35. *Extreme Value Theorem*
Let $f : X \to \mathbb{R}$ be continuous, let X be compact.
If $M := \sup_{p\in X} f(p)$, $m := \inf_{p\in X} f(p)$, then $\exists\, p, q \in X$ such that $f(p) = M$ and $f(q) = m$.

Proof. By Corollary 2.3.34, $f(X)$ is closed and bounded in $\mathbb{R}$. Thus $\sup f(X)$ and $\inf f(X)$ exist and belong to $f(X)$ by Theorems 2.1.39 and 4.3.3. Hence $M = f(p)$ and $m = f(q)$ for some $p, q \in X$.
Structural Commentary: Compactness turns suprema and infima into attained values.
Key Insight: Closed bounded images contain their endpoints. □

Theorem 2.3.36. *Compact Bijective Maps Have Continuous Inverses*
If $f : X \to Y$ is a continuous bijection from compact X to a metric space Y, then $f^{-1} : Y \to X$ is continuous.

Proof. It suffices, by Theorem 2.3.18, to show that $f(V)$ is open in Y for every open $V \subseteq X$. If V is open, then V^c is compact by Theorem 2.2.30, so $f(V^c)$ is compact, hence closed by Theorems 2.3.33 and 2.2.29.
Since f is bijective, $f(V) = Y \setminus f(V^c)$ is open.
Structural Commentary: Compactness turns closed complements into closed images.
Key Insight: A continuous bijection from compact domain has no room for inverse discontinuity. □

Definition 2.3.37. *Uniform Continuity*
A mapping $f : X \to Y$ between metric spaces is *uniformly continuous* on X if $\forall\ \varepsilon > 0\ \exists\ \delta > 0$ such that $d_X(p, q) < \delta \implies d_Y(f(p), f(q)) < \varepsilon\ \forall p, q \in X$. The same δ must work for every point of X.
Structural Commentary: Uniform continuity is setwise: for each ε, one δ must work for all pairs in the set.

Proposition 2.3.38. *Sequential Test for Uniform Continuity*
Let X, Y be metric spaces and let $f : X \to Y$. The following are equivalent:

(i) f is uniformly continuous on X;
(ii) whenever (p_n) and (q_n) are sequences in X with $d_X(p_n, q_n) \to 0$, one has $d_Y(f(p_n), f(q_n)) \to 0$.

Proof.

(*i*) $\Rightarrow$ (*ii*) Given $\varepsilon > 0$, choose $\delta > 0$ such that $d_X(x, y) < \delta$ implies $d_Y(f(x), f(y)) < \varepsilon$ for all $x, y \in X$. Since $d_X(p_n, q_n) \to 0$, eventually $d_X(p_n, q_n) < \delta$, hence eventually $d_Y(f(p_n), f(q_n)) < \varepsilon$.

(*ii*) $\Rightarrow$ (*i*) Prove the contrapositive. If f is not uniformly continuous, then $\exists \varepsilon_0 > 0$ such that $\forall n \geq 1$, $\exists p_n, q_n \in X$ with $d_X(p_n, q_n) < 1/n$ and $d_Y(f(p_n), f(q_n)) \geq \varepsilon_0$.
Thus $d_X(p_n, q_n) \to 0$ but $d_Y(f(p_n), f(q_n)) \not\to 0$, contradicting (ii).

Structural Commentary: Uniform continuity preserves vanishing pairwise distances.
Key Insight: Failure of one global δ manufactures two close input sequences with separated outputs. □

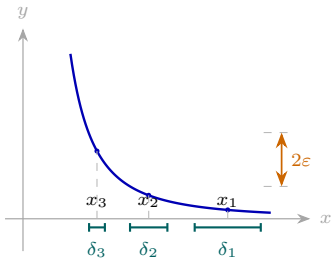


Figure 149: For $f(x) = 1/x^2$ near 0, a fixed vertical tolerance forces smaller pointwise horizontal tolerances.

Theorem 2.3.39. *Continuity on Compact Sets Is Uniform*
If $f : X \to Y$ is continuous and X is compact, then f is uniformly continuous.

Proof. Let $\varepsilon > 0$. For each $a \in X$, choose $r_a > 0$ such that $d_X(x, a) < r_a$ implies $d_Y(f(x), f(a)) < \varepsilon/2$.
Compactness gives a finite cover $X \subseteq \bigcup_{i=1}^{n} B_{r_{a_i}/2}(a_i)$. Set $\delta := \frac{1}{2} \min_i r_{a_i} > 0$. If $d_X(p, q) < \delta$ and $p \in B_{r_{a_i}/2}(a_i)$, then $d_X(q, a_i) < r_{a_i}$ also; hence $f(p)$ and $f(q)$ are each within $\varepsilon/2$ of $f(a_i)$. Thus $d_Y(f(p), f(q)) < \varepsilon$.
Structural Commentary: Compactness converts pointwise radii into finitely many radii with positive minimum.
Key Insight: Uniform continuity is continuity with the point-dependence removed by a finite subcover. □

Theorem 2.3.40. *Compactness Hypotheses Are Essential*
Let E be a noncompact subset of $\mathbb{R}$. Then:

(i) there exists a continuous function on E which is not bounded;
(ii) there exists a continuous bounded function on E which has no maximum;

(iii) if E is bounded, then there exists a continuous function on E which is not uniformly continuous.

Proof.

(i) If E is unbounded, $f(x) = x$ works. If E is bounded, Theorem 2.2.39 gives $x_0 \in \overline{E} \backslash E$, and $f(x) = 1/(x - x_0)$ is continuous on E and unbounded.

(ii) If E is unbounded, $h(x) = x^2/(1 + x^2)$ has $\sup_E h = 1$ but $h < 1$. If E is bounded, with x_0 as above, $g(x) = 1/(1 + (x - x_0)^2)$ is bounded by 1 and has supremum 1 but no maximum.

(iii) If E is bounded, let $f(x) = 1/(x - x_0)$. Given $\delta > 0$, let $x \in E$ with $|x - x_0| < \delta/2$, then $y \in E$ with $|y - x_0| < \min\{\delta/2, 1/(|f(x)| + 1)\}$. Then $|x - y| < \delta$, $|f(x) - f(y)| > 1$, so f is not uniformly continuous.

Structural Commentary: Compactness is essential in the preceding theorems.
Key Insight: Noncompact sets either escape to infinity or miss an accumulation point. □

Theorem 2.3.41. *Continuous Images of Connected Sets Are Connected*
If $f : X \to Y$ is continuous and $E \subseteq X$ is connected, then $f(E)$ is connected.

Proof. If $f(E) = A \cup B$ with A, B nonempty separated, set $G := E \cap f^{-1}(A)$ and $H := E \cap f^{-1}(B)$. Then $E = G \cup H$ and G, H are nonempty. Since $G \subseteq f^{-1}(\overline{A})$ and $f^{-1}(\overline{A})$ is closed, $\overline{G} \subseteq f^{-1}(\overline{A})$; hence $\overline{G} \cap H = \emptyset$. Similarly $G \cap \overline{H} = \emptyset$, so G, H separate E, a contradiction.
Structural Commentary: A separation of the image pulls back to a separation of the domain.
Key Insight: Continuous maps cannot create a gap in the image of a connected set. □

Theorem 2.3.42. *Intermediate Value Theorem*
If $f : [a, b] \to \mathbb{R}$ is continuous and c lies strictly between $f(a)$ and $f(b)$, then $f(x) = c$ for some $x \in (a, b)$.

Proof. The interval $[a, b]$ is connected by Theorem 2.3.2, so $f([a, b])$ is connected by Theorem 2.3.41.
Connected subsets of $\mathbb{R}$ contain every value between any two of their points; hence $c = f(x)$ for some $x \in [a, b]$.
Since $c \neq f(a), f(b)$, $x \in (a, b)$.
Structural Commentary: The intermediate value theorem is connectedness of intervals transported by continuity.
Key Insight: Continuous real functions on intervals cannot jump over a value. □

Definition 2.3.43. *Path-Connected Set*
A set E in a metric space is *path-connected* if for every $p, q \in E$ there is a continuous map $\gamma : [a, b] \to E$ with $\gamma(a) = p$ and $\gamma(b) = q$. Such a map γ is a *path* from p to q.

Theorem 2.3.44. *Path-Connected Sets Are Connected*
Every path-connected set is connected.

Proof. Suppose $E = A \cup B$ with A, B nonempty separated.
Choose $p \in A$, $q \in B$, and a path $\gamma : [a, b] \to E$ from p to q. Since $[a, b]$ is connected, so is $\gamma([a, b])$.
But $\gamma([a, b])$ meets both A and B, which separate it, a contradiction.
Structural Commentary: A path transports connectedness from an interval to the set.
Key Insight: To connect two points by a path, the path image cannot be split into two separated pieces. □

Corollary 2.3.45. *Convex Sets Are Path-Connected*
Every convex subset of $\mathbb{R}^k$ is path-connected.

Proof. For $\mathbf{x}, \mathbf{y} \in E$, set $\gamma(t) = (1-t)\mathbf{x} + t\mathbf{y}$ on $[0, 1]$. Convexity gives $\gamma([0, 1]) \subseteq E$, with $\gamma(0) = \mathbf{x}$ and $\gamma(1) = \mathbf{y}$.
Since the coordinates of γ are affine in t, γ is continuous. Hence E is path-connected.
Structural Commentary: Convexity supplies the whole straight path between any two points.
Key Insight: In Euclidean space, line segments are continuous paths. □

Definition 2.3.46. *One-Sided Limits and Types of Discontinuity*
Let $X \subseteq \mathbb{R}$, $f : X \to \mathbb{R}$, and $x_0 \in \mathbb{R}$.

(i) If x_0 is a limit point of $X \cap (x_0, \infty)$, then $f(x_0+) = q$ means $\lim_{x \to x_0;\, x \in X \cap (x_0, \infty)} f(x) = q$.

(ii) If x_0 is a limit point of $X \cap (-\infty, x_0)$, then $f(x_0-) = q$ means $\lim_{x \to x_0;\, x \in X \cap (-\infty, x_0)} f(x) = q$.

(iii) When X approaches x_0 from both sides, the two-sided punctured limit through X exists if and only if the two one-sided limits exist and are equal.

For functions on an interval, a discontinuity with both one-sided limits is of the *first kind*, or *simple*; otherwise it is of the *second kind*. Simple discontinuities are jumps, $f(x_0+) \neq f(x_0-)$, or removable value mismatches.

Example 2.3.47. *Basic Discontinuity Types*
Dirichlet's function $\mathbf{1}_{\mathbb{Q}}$ has a discontinuity of the second kind at every point.
The function $f(x) = x$ on $\mathbb{Q}$ and $f(x) = 0$ off $\mathbb{Q}$ is continuous only at 0; elsewhere its discontinuities are of the second kind. A jump in a piecewise linear function is a simple discontinuity.
Use Corollary 1.1.17: every interval contains rational and irrational points.

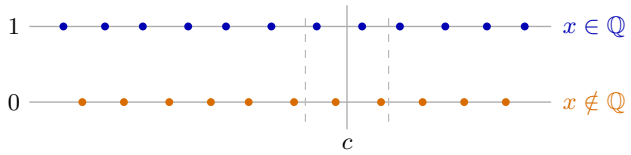


Figure 150: For $g = \mathbf{1}_{\mathbb{Q}}$, interval around c has rational inputs value 1 and irrational inputs value 0.

Definition 2.3.48. *Monotone Functions*
Let f be real-valued on (a, b). The function is *monotonically increasing* if $a < x < y < b$ implies $f(x) \leq f(y)$. Reversing the inequality gives *monotonically decreasing*. A *monotone* function is one of these two types.

Theorem 2.3.49. *One-Sided Limits of Monotone Functions*
If f is increasing on (a, b), then for every $x \in (a, b)$, $\sup_{a<t<x} f(t) = f(x-) \leq f(x) \leq f(x+) = \inf_{x<t<b} f(t)$.
Moreover, $a < x < y < b$ implies $f(x+) \leq f(y-)$. For decreasing f, reverse the inequalities.

Proof. Let $A := \sup_{a<t<x} f(t)$.
For every $\varepsilon > 0$, choose $s < x$ with $A - \varepsilon < f(s) \leq A$; then $s < t < x$ implies $A - \varepsilon < f(t) \leq A$, so $f(x-) = A$.
Let $B := \inf_{x<t<b} f(t)$. For every $\varepsilon > 0$, choose $s > x$ with $B \leq f(s) < B + \varepsilon$; then $x < t < s$ implies $B \leq f(t) < B + \varepsilon$, so $f(x+) = B$. If $x < y$, then $\inf_{x<t<y} f(t) \leq \sup_{x<t<y} f(t)$ gives $f(x+) \leq f(y-)$.
For decreasing f, apply the increasing case to $-f$ and multiply the resulting inequalities by -1.
Structural Commentary: Monotonicity converts one-sided limits into suprema and infima.
Key Insight: A monotone function has jumps, but not oscillatory one-sided failure. □

Corollary 2.3.50. *Monotone Functions Have No Second-Kind Discontinuities*
A monotone real function on an interval has no discontinuities of the second kind.

Proof. At every interior point, Theorem 2.3.49 gives both one-sided limits.
Hence any discontinuity is of the first kind by Definition 2.3.46.
Structural Commentary: Monotonicity rules out oscillatory discontinuities.
Key Insight: All monotone discontinuities are jumps or removable value mismatches. □

Theorem 2.3.51. *Monotone Discontinuities Are Countable*
If f is monotone on (a, b), then the set of points of (a, b) at which f is discontinuous is at most countable.

Proof. Assume f is increasing, and let E be its discontinuity set.
For each $x \in E$, Theorem 2.3.49 gives $f(x-) < f(x+)$; choose $r(x) \in \mathbb{Q}$ with $f(x-) < r(x) < f(x+)$.
If $x_1 < x_2$, then $f(x_1+) \leq f(x_2-)$, so $r(x_1) \neq r(x_2)$. Thus E injects into $\mathbb{Q}$, hence is countable by Corollary 6.1.12. If f is decreasing, apply the increasing case to $-f$, which has the same discontinuity set.
Structural Commentary: Each jump contains a rational marker, and monotonicity keeps the markers distinct.
Key Insight: Countability enters by injecting discontinuities into $\mathbb{Q}$. □

Remark 2.3.52. *Monotone Discontinuities Need Not Be Isolated*
For countable $E = \{x_n : n \in \mathbb{N}\} \subset (a, b)$ and positive summable (c_n), $f(x) = \sum_{x_n<x} c_n$ is increasing, has jumps of size c_n at x_n, and is continuous off E. Thus monotone discontinuities may be dense.

Definition 2.3.53. *Extended Limits and Limits at Infinity*
Neighborhoods of $+\infty$ are intervals $(c, +\infty)$; neighborhoods of $-\infty$ are $(-\infty, c)$.
For $f : E \to \mathbb{R}$ and extended real A, x, $f(t) \to A$ as $t \to x$ means: every neighborhood U of A contains $f(t)$ for all punctured $t \in E$ sufficiently near x, i.e. $\exists V \ni x$ such that $t \in V \cap E,\ t \neq x \implies f(t) \in U$.
For real A, x, this agrees with Definition 2.3.4.
Structural Commentary: Extended limits use same neighborhood test, $\pm\infty$ allowed as input or output limits.

Theorem 2.3.54. *Algebra of Extended Limits*
Let $f, g : E \to \mathbb{R}$ and $f(t) \to A$, $g(t) \to B$ as $t \to x$, with A, B, x extended real. Then, whenever RHS is defined:
(i) if $f(t) \to A'$ as $t \to x$, then $A' = A$;

(ii) $(f+g)(t) \to A+B$;
(iii) $(fg)(t) \to AB$;
(iv) $(f/g)(t) \to A/B$;
The forms $\infty - \infty$, $0 \cdot \infty$, ∞/∞, *and* $A/0$ *are undefined.*

Proof.

(i) If $A' \neq A$, choose disjoint neighborhoods of A and A'.
Definition 2.3.53 would eventually place $f(t)$ in both, a contradiction.
(ii) For any neighborhood W of $A+B$, definedness gives neighborhoods U of A and V of B such that $u+v \in W$ for $u \in U$, $v \in V$. Apply the two limits in a common punctured neighborhood of x.
(iii) The same neighborhood argument applies to multiplication, excluding $0 \cdot \infty$.
(iv) Apply it to division; since $B \neq 0$, choose the neighborhood of B away from 0.

Structural Commentary: Extended limit algebra is ordinary limit algebra in neighborhood form.
Key Insight: The undefined forms are exactly where neighborhood control breaks. □

Method 2.3.55. *Continuity Proof Patterns*

(i) *ε–δ at p*: prove $\forall \varepsilon > 0\, \exists \delta > 0$ such that $x \in E$ and $d_X(x,p) < \delta$ imply $d_Y(f(x), f(p)) < \varepsilon$.
(ii) *Limit at p*: prove same estimate with $0 < d_X(x,p) < \delta$, target q; equivalently test all $x_n \to p$, $x_n \neq p$.
(iii) *Topological continuity*: prove $V \subseteq Y$ open $\Rightarrow f^{-1}(V)$ open, or $C \subseteq Y$ closed $\Rightarrow f^{-1}(C)$ closed.
(iv) *Compactness*: from open cover $\{V_\alpha\}$ of $f(X)$, pulled-back cover $\{f^{-1}(V_\alpha)\}$ of X extract finite subcover.
(v) *Uniform continuity*: choose local radii r_p, cover compact X by $B_{r_p/2}(p)$, and set $\delta = \frac{1}{2}\min_i r_{p_i}$.

Key Insight: Continuity is local control; compactness turns local controls into finite control.

2.4 Selected General Topology

General topology abstracts the open-set behavior of metric spaces. This selective pass records spaces, bases, products, quotients, and countability axioms for later use.

Definition 2.4.1. *Topological Space*
A *topology* on a set X is a collection $\mathcal{T} \subseteq 2^X$ such that

(i) $\emptyset, X \in \mathcal{T}$;
(ii) arbitrary unions of members of $\mathcal{T}$ belong to $\mathcal{T}$;
(iii) finite intersections of members of $\mathcal{T}$ belong to $\mathcal{T}$.

The pair $(X, \mathcal{T})$ is a *topological space*, and the members of $\mathcal{T}$ are the *open sets*. Closed sets are complements of open sets.

Proposition 2.4.2. *Metric Open Sets Form a Topology*
If (X,d) is a metric space, then the open sets from Definition 2.1.14 form a topology on X.

Proof. The empty set and X are open. Arbitrary unions and finite intersections are open by Theorem 2.1.22.
Structural Commentary: Metric balls generate an open-set system with the exact axioms of a topology.
Key Insight: General topology keeps the open-set calculus and forgets the distance function. □

Definition 2.4.3. *Basis*
Let X be a set. A collection $\mathcal{B} \subseteq 2^X$ is a *basis* for a topology if:

(i) for every $x \in X$, some $B \in \mathcal{B}$ contains x;
(ii) if $x \in B_1 \cap B_2$ with $B_1, B_2 \in \mathcal{B}$, then some $B_3 \in \mathcal{B}$ satisfies $x \in B_3 \subseteq B_1 \cap B_2$.

The topology generated by $\mathcal{B}$ consists of all unions of basis elements. A *subbasis* is a collection whose finite intersections form a basis.

Example 2.4.4. *Metric Balls as a Basis*
In a metric space, the collection of all open balls is a basis for the metric topology. On $\mathbb{R}^k$, balls with rational centers and rational radii form a countable basis.

Definition 2.4.5. *Product Topology*
For topological spaces X and Y, the *product topology* on $X \times Y$ is generated by basis sets $U \times V$, where $U \subseteq X$ and $V \subseteq Y$ are open. For a family $\{X_i\}_{i \in I}$, the product topology on $\prod_i X_i$ is generated by finite-coordinate restrictions: sets of the form $\prod_i U_i$ where each U_i is open and $U_i = X_i$ for all but finitely many i.

Proposition 2.4.6. *Product Topology and Coordinate Maps*
A map $f : Z \to X \times Y$ is continuous for the product topology if and only if both coordinate maps $\pi_X \circ f$ and $\pi_Y \circ f$ are continuous.

Proof. If f is continuous, then coordinate maps are continuous because inverse images of open sets are inverse images of product-open cylinders. Conversely, if both coordinates are continuous, then the inverse image of each basis set $U \times V$ is $(\pi_X \circ f)^{-1}(U) \cap (\pi_Y \circ f)^{-1}(V)$, which is open.
Structural Commentary: Product topology is the weakest topology making the projections continuous.
Key Insight: Continuity into a product is checked coordinatewise. □

Definition 2.4.7. *Quotient Topology*
Let $q : X \to Y$ be a surjection from a topological space X onto a set Y.

(i) *Quotient topology* on Y is topology w. open sets are precisely subsets $V \subseteq Y$ s,t, $q^{-1}(V)$ open in X.
(ii) Equivalently, it is the strongest topology on Y for which q is continuous.

Proposition 2.4.8. *Quotient Universal Property*
Let $q : (X, \tau_X) \twoheadrightarrow Y$ and give Y the quotient topology $\tau_Y = \{V \subseteq Y : q^{-1}(V) \in \tau_X\}$.
For any map $g : Y \to (Z, \tau_Z)$, $g : (Y, \tau_Y) \to (Z, \tau_Z)$ is continuous $\iff$ $g \circ q : (X, \tau_X) \to (Z, \tau_Z)$ is continuous.

Proof.

(⇒) If g is continuous, then $g \circ q$ is continuous by composition.
(⇐) Suppose $g \circ q$ is continuous. For $W \in \tau_Z$, $q^{-1}(g^{-1}(W)) = (g \circ q)^{-1}(W) \in \tau_X$.
Thus $g^{-1}(W) \in \tau_Y$, so g is continuous.

Structural Commentary: Continuity from a quotient is tested by precomposition with the quotient map.
Key Insight: Pull back through the quotient map and use the defining open-set test. □

Definition 2.4.9. *Countability Axioms*
A topological space (X, τ) is:

(i) *first countable* if for each $x \in X$ there is a countable neighborhood basis $\mathcal{B}_x$ at x: for every neighborhood U of x, some $B \in \mathcal{B}_x$ satisfies $B \subseteq U$;
(ii) *second countable* if τ has a countable basis;
(iii) *separable* if there is a countable $D \subseteq X$ with $\overline{D} = X$, agreeing with Definition 2.1.37 for metric spaces.

Proposition 2.4.10. *Second Countable Implies Separable*
If (X, τ) is second countable, then there is a countable $D \subseteq X$ with $\overline{D} = X$.

Proof. Let $\mathcal{B} = \{B_n : n \in I\}$ be a countable basis, with $I \subseteq \mathbb{N}$.
For each $n \in I$ with $B_n \neq \emptyset$, choose $x_n \in B_n$ and set $D = \{x_n : n \in I,\ B_n \neq \emptyset\}$. Then D is countable.
If $U \in \tau \setminus \{\emptyset\}$, choose $x \in U$; some $B_n \in \mathcal{B}$ satisfies $x \in B_n \subseteq U$, so $x_n \in D \cap U$. Hence $\overline{D} = X$.
Structural Commentary: $\mathcal{B}$ countable gives countable D meeting every $U \in \tau \setminus \{\emptyset\}$.
Key Insight: D is dense iff $D \cap U \neq \emptyset$ for every $U \in \tau \setminus \{\emptyset\}$. □

3 Convergence and Function Spaces

Convergence and Function Spaces studies series, uniform limits, compactness of functions, power-series approximation, and Fourier orthogonality.

3.1 Infinite Series and Summability

Series are sequences of partial sums. Comparison, condensation, root and ratio tests, rearrangements, products, and double sums all refine that viewpoint.

Definition 3.1.1. *Series and Partial Sums*
Let (a_n) be a complex sequence, and for $p \leq q$ put $\sum_{n=p}^{q} a_n := a_p + a_{p+1} + \cdots + a_q$.

The index n in the displayed finite sum is a bound variable, in the sense of Principle 2.1.1; it may be renamed when its scope is unchanged.

The *series* $\sum_{n=1}^{\infty} a_n$ (or $a_1 + a_2 + \cdots$) is the partial-sum sequence $s_n := \sum_{k=1}^{n} a_k$. If $s_n \to s$, then $\sum_{n=1}^{\infty} a_n = s$ and series *converges*; otherwise it *diverges*. The same convention applies to series starting at $n = 0$.

Motivation 3.1.2. *Partial Sums Remember Order*
An infinite sum is a limit of partial sums. For alternating harmonic series, odd and even partial sums approach the same value from opposite sides; see Corollary 3.1.24 and Theorem 3.1.34.

Perspective 3.1.3. *Formal Algebra Needs Convergence*
Finite sums are algebraic; infinite sums are partial-sum limits. Algebraic manipulations are justified only by convergence control, as in Theorem 3.1.4, Corollary 3.1.5, and Theorems 3.1.31, 3.1.34.
Key Insight: The symbol $\sum_{n=1}^{\infty}$ is a limit claim before it is an algebraic object.

Theorem 3.1.4. *Cauchy Criterion for Series*
The series $\sum a_n$ converges if and only if $\forall \varepsilon > 0$ $\exists$ an integer N s.t. $m \geq n \geq N$ implies $\left|\sum_{k=n}^{m} a_k\right| \leq \varepsilon$.

Proof.
($\Rightarrow$) If $s_n := \sum_{k=1}^{n} a_k$ converges, then (s_n) is Cauchy. Thus, for $m \geq n$ large enough, $\left|\sum_{k=n}^{m} a_k\right| = |s_m - s_{n-1}| \leq \varepsilon$.
($\Leftarrow$) The tail condition gives, for $m > n$ large enough, $|s_m - s_n| = \left|\sum_{k=n+1}^{m} a_k\right| \leq \varepsilon$.
Hence (s_n) is Cauchy and converges by Theorem 2.2.10.
Structural Commentary: A series converges exactly when all sufficiently late finite tails are small.
Key Insight: The Cauchy test measures tails, not the eventual value of the sum. □

Corollary 3.1.5. *Terms of a Convergent Series Tend to Zero*
If $\sum a_n$ converges, then $a_n \to 0$.

Proof. Let $\varepsilon > 0$. By Theorem 3.1.4, there is N such that $m \geq n \geq N$ implies $|\sum_{k=n}^{m} a_k| \leq \varepsilon$.
Taking $m = n$ gives $|a_n| \leq \varepsilon$ for all $n \geq N$, hence $a_n \to 0$.
Structural Commentary: A single term is the shortest possible tail.
Key Insight: The condition $a_n \to 0$ is necessary, but not sufficient for convergence. □

Theorem 3.1.6. *Nonnegative Series and Bounded Partial Sums*
A series of nonnegative real terms converges if and only if its partial sums form a bounded sequence.

Proof.
($\Rightarrow$) Every convergent sequence is bounded, so the partial sums are bounded.
($\Leftarrow$) Since $a_n \geq 0$, $s_n = \sum_{k=1}^{n} a_k$ increases; if bounded, it converges by Theorem 5.5.1.
Structural Commentary: Nonnegative series replace cancellation by monotonicity.
Key Insight: For nonnegative terms, convergence is exactly bounded accumulation. □

Theorem 3.1.7. *Comparison Test*
(i) If $|a_n| \leq c_n$ for $n \geq N_0$, and if $\sum c_n$ converges, then $\sum a_n$ converges.
(ii) If $a_n \geq d_n \geq 0$ for $n \geq N_0$, and if $\sum d_n$ diverges, then $\sum a_n$ diverges.

Proof.
(i) Given $\varepsilon > 0$, choose $N \geq N_0$ so that $m \geq n \geq N$ implies $\sum_{k=n}^{m} c_k \leq \varepsilon$.
Then $\left|\sum_{k=n}^{m} a_k\right| \leq \sum_{k=n}^{m} |a_k| \leq \sum_{k=n}^{m} c_k \leq \varepsilon$, so Theorem 3.1.4 applies.
(ii) If $\sum a_n$ converged, then (i) applied to $d_n \leq a_n$ would force $\sum d_n$ to converge, a contradiction.
Structural Commentary: Absolute tail domination transfers Cauchy convergence.
Key Insight: Comparison is a tail test; finitely many early terms do not matter. □

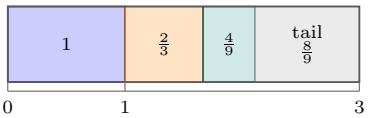


Figure 151: At $x = 2/3$, the strip widths sum to $1 + \frac{2}{3} + \frac{4}{9} + \cdots = 3$.

Theorem 3.1.8. *Geometric Series*
If $0 \leq x < 1$, then $\sum_{n=0}^{\infty} x^n = \frac{1}{1-x}$. If $x \geq 1$, the series diverges.

Proof. For $x \neq 1$, $\sum_{k=0}^{n} x^k = \frac{1-x^{n+1}}{1-x}$. Let $n \to \infty$ and use Theorem 1.3.25.
If $x = 1$, the partial sums are $n+1$, so they diverge.
Structural Commentary: The finite geometric identity converts convergence into the limit of x^{n+1}.
Key Insight: The geometric series is the model comparison series for ratio-like decay. □

Definition 3.1.9. *Real Decimals as Series*
For digits $b_N, \dots, b_0, a_1, a_2, \dots \in \{0, \dots, 9\}$, with $b_N \neq 0$ unless the integer part is 0,
$b_N \cdots b_0.a_1a_2\cdots := \sum_{i=0}^{N} b_i 10^i + \sum_{j=1}^{\infty} a_j 10^{-j}$. The fractional series converges absolutely since $0 \leq a_j 10^{-j} \leq 9 \cdot 10^{-j}$ and $\sum 9 \cdot 10^{-j}$ is geometric. A minus sign denotes the negative of this value.

Proposition 3.1.10. *Terminal Nines Are Nonunique*
The same real number can have two decimal representations: $1.000\ldots = 0.999\ldots$.

Proof. By Definition 3.1.9 and Theorem 3.1.8, $1.000\ldots = 1$; $0.999\ldots = \sum_{j=1}^{\infty} 9 \cdot 10^{-j} = 9 \cdot \frac{10^{-1}}{1-10^{-1}} = 1$.
Structural Commentary: The two digit strings have the same partial-sum limit.
Key Insight: Uniqueness requires excluding terminal 9 tails. □

Perspective 3.1.11. *Decimal Expansions Are Series*
Proposition 1.1.19 gives decimal truncations that converge to a real number, while Definition 3.1.9 interprets the infinite digit string as a series. Proposition 3.1.10 explains the standard terminal-9 ambiguity.

Theorem 3.1.12. *Cauchy Condensation Test*
Suppose $a_1 \geq a_2 \geq a_3 \geq \cdots \geq 0$. *Then* $\sum_{n=1}^{\infty} a_n$ *converges if and only if* $\sum_{k=0}^{\infty} 2^k a_{2^k}$ *converges.*

Proof. Let $s_n := a_1 + \cdots + a_n$ and $t_k := a_1 + 2a_2 + \cdots + 2^k a_{2^k}$.
If $n < 2^k$, then grouping by dyadic blocks gives $s_n \leq a_1 + (a_2 + a_3) + \cdots + (a_{2^k} + \cdots + a_{2^{k+1}-1}) \leq t_k$.
If $n > 2^k$, grouping from the other side gives $2s_n \geq t_k$. Hence the nonnegative partial sums (s_n) and (t_k) are either both bounded or both unbounded; apply Theorem 3.1.6.
Structural Commentary: Monotonicity lets dyadic blocks be compared to their endpoint terms.
Key Insight: A sparse dyadic subsequence detects the whole monotone nonnegative series. □

Corollary 3.1.13. *p-Series Test*
The series $\sum_{n=1}^{\infty} \frac{1}{n^p}$ *converges if* $p > 1$ *and diverges if* $p \leq 1$.

Proof. If $p \leq 0$, the terms do not tend to 0, so Corollary 3.1.5 gives divergence.
If $p > 0$, Theorem 3.1.12 reduces the question to $\sum_{k=0}^{\infty} 2^k \frac{1}{(2^k)^p} = \sum_{k=0}^{\infty} 2^{(1-p)k}$.
This geometric series converges exactly when $2^{1-p} < 1$, i.e. $p > 1$, by Theorem 3.1.8.
Structural Commentary: Condensation turns powers of n into powers of 2.
Key Insight: The boundary is $p = 1$ because the condensed ratio is 2^{1-p}. □

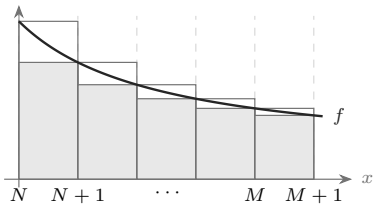


Figure 152: Left and right endpoint estimates for decreasing $f \geq 0$.

Theorem 3.1.14. *Integral Test*
Let $N \in \mathbb{N}$ *and let* $f : [N, \infty) \to [0, \infty)$ *be decreasing and Riemann integrable on compact subintervals.*
Then $\sum_{k=N}^{\infty} f(k)$ *converges iff* $\int_N^{\infty} f(x)\,dx$ *converges.*
For $M \geq N$, $\int_N^{M+1} f(x)\,dx \leq \sum_{k=N}^{M} f(k) \leq f(N) + \int_N^{M} f(x)\,dx$.

Proof. For a monotone function, compact-subinterval Riemann integrability follows from Theorem 4.1.45 with $\alpha(x) = x$; it is kept in the statement to make the improper-integral dependency explicit.

($\Rightarrow$) For $x \in [k, k+1]$, monotonicity gives $f(x) \le f(k)$. Hence $\int_N^{M+1} f(x)\,dx \le \sum_{k=N}^{M} f(k)$. If the series converges, these truncations are bounded, so $\int_N^{\infty} f$ converges by Theorem 4.1.61.

($\Leftarrow$) For $k \ge N+1$ and $x \in [k-1, k]$, monotonicity gives $f(k) \le f(x)$. Hence $\sum_{k=N}^{M} f(k) \le f(N) + \int_N^M f(x)\,dx$. If $\int_N^{\infty} f$ converges, the partial sums are bounded, so the series converges by Theorem 3.1.6.

Structural Commentary: A monotone sequence can be compared to neighboring unit-width areas.
Key Insight: The integral test is comparison of two infinite accumulations with the same tail size. □

Definition 3.1.15. *The Number e*
Define $e := \sum_{n=0}^{\infty} \frac{1}{n!}$, where $0! := 1$; convergence follows by comparison with the geometric series.

Theorem 3.1.16. *A Limit Formula for e*
The limit formula is $\lim_{n\to\infty} \left(1 + \frac{1}{n}\right)^n = e$.

Proof. Let $s_m := \sum_{k=0}^{m} \frac{1}{k!}$ and $t_n := (1 + \frac{1}{n})^n$. The binomial theorem gives $t_n \le s_n$, hence $\limsup t_n \le e$.
For fixed $m \le n$, first m binomial terms give $t_n \ge 1 + 1 + \frac{1}{2!}\left(1 - \frac{1}{n}\right) + \cdots + \frac{1}{m!}\left(1 - \frac{1}{n}\right)\cdots\left(1 - \frac{m-1}{n}\right)$.
Letting $n \to \infty$ with m fixed gives $\liminf t_n \ge s_m$, and then $m \to \infty$ gives $e \le \liminf t_n$.
Structural Commentary: Bound above by the factorial series; bound below by fixed binomial terms.
Key Insight: Limsup and liminf squeeze the binomial limit to the series definition. □

Proposition 3.1.17. *Remainder Bound for the Series of e*
If $n \ge 1$ *and* $s_n = \sum_{k=0}^{n} 1/k!$, *then* $0 < e - s_n < \frac{1}{n!\,n}$.

Proof. Tail is $e - s_n = \frac{1}{(n+1)!} + \frac{1}{(n+2)!} + \cdots < \frac{1}{(n+1)!}\left(1 + \frac{1}{n+1} + \frac{1}{(n+1)^2} + \cdots\right) = \frac{1}{n!\,n}$
Structural Commentary: The factorial tail is bounded by a geometric tail.
Key Insight: Rapid convergence of $\sum \frac{1}{n!}$ gives arithmetic control of e. □

Theorem 3.1.18. *e Is Irrational*
The number e is irrational.

Proof. Suppose $e = p/q$ with p, q positive integers. By Proposition 3.1.17, $0 < q!(e - s_q) < \frac{1}{q}$.
But $q!e$ and $q!s_q$ are integers, so $q!(e - s_q)$ is an integer strictly between 0 and 1, impossible.
Structural Commentary: The sharp tail bound turns rationality into an impossible integer.
Key Insight: A positive integer cannot be smaller than 1. □

Theorem 3.1.19. *Root Test*
For a series $\sum a_n$, *put* $\alpha := \limsup_{n\to\infty} \sqrt[n]{|a_n|}$. *Then:*

(i) if $\alpha < 1$, *then* $\sum a_n$ *converges absolutely;*
(ii) if $\alpha > 1$, *then* $\sum a_n$ *diverges;*
(iii) if $\alpha = 1$, *the test gives no information.*

Proof.

(i) Choose $\alpha < \beta < 1$. Eventually $|a_n| < \beta^n$, so $\sum |a_n|$ converges by comparison with $\sum \beta^n$.
(ii) Since $\alpha > 1$, $|a_n| > 1$ for infinitely many n, so $a_n \not\to 0$; use Corollary 3.1.5.
(iii) $\sum 1/n$ and $\sum 1/n^2$ both have $\alpha = 1$, but behave differently by Corollary 3.1.13.

Structural Commentary: The root test compares term size directly to geometric decay.
Key Insight: Values below 1 force absolute convergence; values above 1 prevent terms from tending to 0. □

Theorem 3.1.20. *Ratio Test*
Assume $a_n \neq 0$ *for all sufficiently large* n.

(i) If $\limsup_{n\to\infty} \left|\frac{a_{n+1}}{a_n}\right| < 1$, *then* $\sum a_n$ *converges absolutely.*
(ii) If $\liminf_{n\to\infty} \left|\frac{a_{n+1}}{a_n}\right| > 1$, *then* $\sum a_n$ *diverges.*
(iii) In the remaining cases, the test gives no information.

Proof.

(i) Choose $\beta < 1$ with $|a_{n+1}/a_n| < \beta$ eventually.
Then $|a_n| = O(\beta^n)$, so $\sum |a_n|$ converges by comparison with a geometric series.

(ii) Choose $\beta > 1$ with $|a_{n+1}/a_n| > \beta$ eventually.
Then the tail $|a_n|$ is eventually increasing and bounded away from 0, so $a_n \not\to 0$.

(iii) The series $\sum 1/n$ and $\sum 1/n^2$ both have ratio limit 1, but behave differently by Corollary 3.1.13.

Structural Commentary: Ratios turn successive terms into a geometric majorant.
Key Insight: Ratios below 1 force geometric decay; ratios above 1 prevent terms from tending to zero. □

Theorem 3.1.21. *Root and Ratio Test Comparison*
For any sequence (c_n) of positive numbers, $\liminf_{n\to\infty} \frac{c_{n+1}}{c_n} \le \liminf_{n\to\infty} \sqrt[n]{c_n} \le \limsup_{n\to\infty} \sqrt[n]{c_n} \le \limsup_{n\to\infty} \frac{c_{n+1}}{c_n}$.

Proof. Let $\alpha = \limsup c_{n+1}/c_n$. If $\alpha < +\infty$ and $\beta > \alpha$, then eventually $c_{n+1} \le \beta c_n$, hence $c_n \le c_N \beta^{n-N}$ for all large n. Taking nth roots gives $\limsup \sqrt[n]{c_n} \le \beta$; let $\beta \downarrow \alpha$. If $\alpha = +\infty$, the rightmost inequality is automatic. Let $\ell = \liminf c_{n+1}/c_n$. If $\ell = 0$, the leftmost inequality is automatic. If $\gamma < \ell$ with $\gamma > 0$, then eventually $c_{n+1} \ge \gamma c_n$, hence $c_n \ge c_N \gamma^{n-N}$ for all large n. Taking nth roots gives $\liminf \sqrt[n]{c_n} \ge \gamma$; let $\gamma \uparrow \ell$.
The middle inequality is the general fact that $\liminf a_n \le \limsup a_n$.
Structural Commentary: A tail ratio bound multiplies into a root bound.
Key Insight: Whenever the ratio test proves convergence, the root test does too. □

Theorem 3.1.22. *Summation by Parts*
Given two sequences (a_n) and (b_n), put $A_n := \sum_{k=0}^{n} a_k$ *and* $A_{-1} := 0$.
If $0 \le p \le q$, then $\sum_{n=p}^{q} a_n b_n = \sum_{n=p}^{q-1} A_n(b_n - b_{n+1}) + A_q b_q - A_{p-1} b_p$.

Proof. Since $a_n = A_n - A_{n-1}$, $\sum_{n=p}^{q} a_n b_n = \sum_{n=p}^{q} A_n b_n - \sum_{n=p}^{q} A_{n-1} b_n = \sum_{n=p}^{q} A_n b_n - \sum_{n=p-1}^{q-1} A_n b_{n+1}$.
Structural Commentary: This is the discrete analogue of integration by parts.
Key Insight: Move the difference from a_n onto b_n by rewriting $a_n = A_n - A_{n-1}$. □

Theorem 3.1.23. *Dirichlet Test*
If the partial sums A_n of $\sum a_n$ are bounded, $b_n \downarrow 0$, and $b_n \ge 0$, then $\sum a_n b_n$ converges.

Proof. Choose M with $|A_n| \le M$. If $M = 0$, the result is immediate. Otherwise choose N with $b_N \le \varepsilon/(2M)$. For $N \le p \le q$, $\left|\sum_{n=p}^{q} a_n b_n\right| \le M\left(\sum_{n=p}^{q-1}(b_n - b_{n+1}) + b_q + b_p\right) = 2Mb_p \le 2Mb_N \le \varepsilon$. Theorem 3.1.4 applies.
Structural Commentary: Bounded oscillation in A_n is damped by a monotone scalar tending to 0.
Key Insight: Summation by parts converts cancellation into a total variation estimate for b_n. □

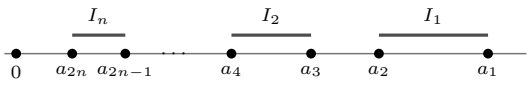


Figure 153: The intervals $I_j = [a_{2j}, a_{2j-1}]$ are disjoint inside $(0, a_1)$.

Corollary 3.1.24. *Alternating Series Test*
Suppose $|c_1| \ge |c_2| \ge |c_3| \ge \cdots$, the signs alternate, and $c_n \to 0$. Then $\sum c_n$ converges.

Proof. After shifting indices to start at 0 if desired, apply Theorem 3.1.23 with $a_n = (-1)^{n+1}$ and $b_n = |c_n|$.
Structural Commentary: Alternation gives bounded sums; decreasing magnitudes damp them.
Key Insight: Alternating convergence is Dirichlet convergence in its simplest form. □

Perspective 3.1.25. *Alternating Partial Sums as Nested Control*
The interval picture explains the boundedness in Corollary 3.1.24: the even partial sums collect disjoint lengths inside a fixed interval, while the odd partial sums differ from the even ones by the next term, which tends to 0.

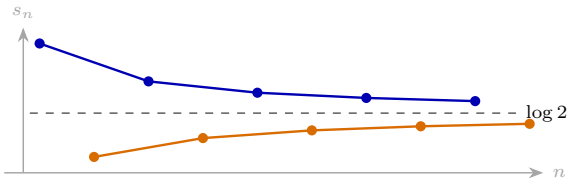


Figure 154: Odd partial sums (blue) decrease and even partial sums (orange) increase toward $\log 2$.

Definition 3.1.26. *Absolute and Nonabsolute Convergence*
The series $\sum a_n$ *converges absolutely* if $\sum |a_n|$ converges.
If $\sum a_n$ converges but $\sum |a_n|$ diverges, then $\sum a_n$ converges *nonabsolutely*.

Theorem 3.1.27. *Absolute Convergence Implies Convergence*
If $\sum a_n$ converges absolutely, then $\sum a_n$ converges.

Proof. For $m \geq n$, $\left|\sum_{k=n}^{m} a_k\right| \leq \sum_{k=n}^{m} |a_k|$. Apply Theorem 3.1.4 to $\sum |a_n|$.
Structural Commentary: Absolute convergence removes cancellation from the convergence question.
Key Insight: Small absolute tails force small signed or complex tails. □

Theorem 3.1.28. *Addition and Scalar Multiplication of Series*
If $\sum a_n = A$ and $\sum b_n = B$, then $\sum (a_n + b_n) = A + B$ and $\sum ca_n = cA$ for every fixed scalar c.

Proof. Apply Theorem 1.3.3 to the corresponding partial sums.
Structural Commentary: Linearity of convergent series is just limit algebra for partial sums.
Key Insight: Add series by adding their partial-sum sequences. □

Definition 3.1.29. *Cauchy Product*
Given $\sum a_n$, $\sum b_n$, define $c_n := \sum_{k=0}^{n} a_k b_{n-k}$, $(n = 0, 1, 2, \ldots)$. Series $\sum c_n$ is *Cauchy product* of given series.

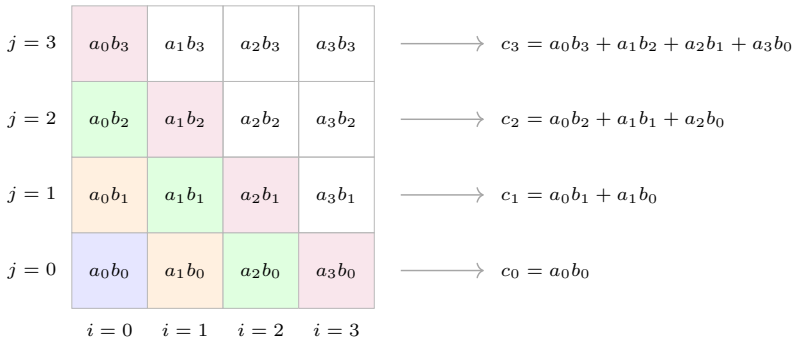


Figure 155: Diagonal sums of $(a_i b_j)$ form the Cauchy product; finite rectangular sums are order-free.

Perspective 3.1.30. *Diagonal Ordering in a Product*
c_n is the diagonal sum of the products $a_i b_j$ with $i + j = n$. Absolute convergence, as in Theorems 3.1.31 and 3.1.34, keeps the value independent of the accumulation order.

Theorem 3.1.31. *Mertens' Theorem for Cauchy Products*
Suppose $\sum a_n$ converges absolutely, $\sum a_n = A$, $\sum b_n = B$, and $c_n = \sum_{k=0}^{n} a_k b_{n-k}$. Then $\sum_{n=0}^{\infty} c_n = AB$.

Proof. Let A_n, B_n, C_n be the partial sums and set $\beta_n := B_n - B$.
Then $C_n = A_n B + \gamma_n$, and $\gamma_n := a_0\beta_n + a_1\beta_{n-1} + \cdots + a_n\beta_0$. Since $A_n B \to AB$, it suffices to show $\gamma_n \to 0$.
Put $\alpha := \sum |a_n|$. If $|\beta_j| \leq \varepsilon$ for $j \geq N$, then $|\gamma_n| \leq |\beta_0 a_n + \cdots + \beta_N a_{n-N}| + \varepsilon\alpha$.
Let $n \to \infty$ with N fixed. The first term tends to 0, so $\limsup |\gamma_n| \leq \varepsilon\alpha$; let $\varepsilon \downarrow 0$.
Structural Commentary: Absolute convergence of one factor controls the convolution error.
Key Insight: One absolute factor makes the Cauchy product converge to AB. □

Definition 3.1.32. *Rearrangement of a Series*
If (k_n) lists each positive integer exactly once and $a'_n := a_{k_n}$, then $\sum a'_n$ is a *rearrangement* of $\sum a_n$.

Theorem 3.1.33. *Riemann Rearrangement Theorem*
If real $\sum a_n$ converges nonabsolutely and $-\infty \le \alpha \le \beta \le \infty$,
then some rearrangement has partial sums s'_n with $\liminf_{n\to\infty} s'_n = \alpha$ and $\limsup_{n\to\infty} s'_n = \beta$.

Proof. Write $a_n = p_n - q_n$, where $p_n = \max(a_n, 0)$ and $q_n = \max(-a_n, 0)$. Since $\sum a_n$ converges but not absolutely, $\sum p_n = \sum q_n = +\infty$, and $p_n, q_n \to 0$. List the positive nonzero p_n as $P_1, P_2, \ldots$ and the positive nonzero q_n as $Q_1, Q_2, \ldots$, preserving their original orders.
Choose finite numbers $\ell_N < u_N$ with $\ell_N \to \alpha$ and $u_N \to \beta$ in the extended sense. Starting from partial sum 0, add unused P terms until the sum is $> u_N$, then add unused terms $-Q$ until the sum is $< \ell_N$; repeat for $N = 1, 2, \ldots$. Divergence of both positive series makes each stopping time finite, and the alternating stages use every nonzero positive and negative term eventually. Insert zero terms anywhere.
After the Nth upward crossing, the partial sum lies between u_N and $u_N + P_j$ for the next positive term used; after the downward crossing, it lies between $\ell_N - Q_j$ and ℓ_N for the next negative term used. Since $P_j, Q_j \to 0$, these overshoots vanish. Hence subsequences of partial sums approach the upper targets and lower targets, while all later partial sums stay within vanishing error of $[\ell_N, u_N]$. Therefore $\liminf s'_n = \alpha$ and $\limsup s'_n = \beta$.
Structural Commentary: Conditional convergence leaves enough mass to steer partial sums.
Key Insight: Vanishing terms make the steering increasingly precise. □

Theorem 3.1.34. *Absolutely Convergent Series Are Rearrangement-Invariant*
Every rearrangement of an absolutely convergent complex series converges to the same sum.

Proof. Let s_n, s'_n be the original and rearranged partial sums.
Given $\varepsilon > 0$, choose N so that $\sum_{i=N}^{m} |a_i| \le \varepsilon$ for $m \ge N$, then choose p so $a_1, \ldots, a_N$ occur among $a'_1, \ldots, a'_p$.
For $n > p$, $s_n - s'_n$ uses only tail terms, hence $|s_n - s'_n| \le \varepsilon$. Thus s'_n has the same limit as s_n.
Structural Commentary: Absolute tails control any finite reordering of the remaining terms.
Key Insight: Absolute convergence is exactly the protection needed against rearrangement. □

Definition 3.1.35. *Absolute Summation over Countable Sets*
Let X be countable, $f : X \to \mathbb{C}$, and $g : \mathbb{N} \to X$ be a bijection, $g_n := g(n)$.
If $\sum_{n=1}^{\infty} |f(g_n)| < \infty$, then $\sum_{x\in X} f(x)$ is *absolutely convergent*, and $\sum_{x\in X} f(x) := \sum_{n=1}^{\infty} f(g_n)$.

Theorem 3.1.36. *Countable Sums Are Order-Free Under Absolute Convergence*
Let X, Y be countable, $f : X \to \mathbb{C}$ be absolutely summable, and $h : Y \to X$ be a bijection.
Then $f \circ h$ is absolutely summable and $\sum_{y\in Y} f(h(y)) = \sum_{x\in X} f(x)$.
In particular, a countable absolutely convergent sum is independent of the enumeration.

Proof. Let $g : \mathbb{N} \to X$ and $k : \mathbb{N} \to Y$ be enumerations, and set $\sigma := g^{-1} \circ h \circ k : \mathbb{N} \to \mathbb{N}$.
Then σ is a bijection and $h(k_n) = g_{\sigma(n)}$. Hence $\sum_{n=1}^{\infty} |f(h(k_n))| = \sum_{n=1}^{\infty} \left|f(g_{\sigma(n)})\right| < \infty$.
Thus $\sum_{y\in Y} f(h(y)) = \sum_{n=1}^{\infty} f(g_{\sigma(n)}) = \sum_{n=1}^{\infty} f(g_n) = \sum_{x\in X} f(x)$, where the middle equality is Theorem 3.1.34.
Taking $Y = \mathbb{N}$ and h to be any enumeration of X gives enumeration independence.
Structural Commentary: Absolute convergence turns an enumeration into harmless bookkeeping.
Key Insight: Once absolute tails are controlled, the order of a countable listing no longer changes the sum. □

Theorem 3.1.37. *Fubini for Absolutely Summable Arrays*
Let $a_{mn} \in \mathbb{C}$ and suppose $L := \sum_{(m,n)\in\mathbb{N}\times\mathbb{N}} |a_{mn}| < \infty$. Then each row and column is absolutely summable.
With $r_m := \sum_{n=1}^{\infty} a_{mn}$ and $c_n := \sum_{m=1}^{\infty} a_{mn}$, the row-total and column-total series are absolutely convergent,
where $\sum_{m=1}^{\infty} |r_m| < \infty$; $\sum_{n=1}^{\infty} |c_n| < \infty$, and $\sum_{m=1}^{\infty} r_m = \sum_{(m,n)\in\mathbb{N}\times\mathbb{N}} a_{mn} = \sum_{n=1}^{\infty} c_n$.

Proof. For finite $F \subseteq \mathbb{N} \times \mathbb{N}$, $\sum_{(m,n)\in F} |a_{mn}| \le L$. Hence each row and column has bounded absolute partial sums, so rows and columns converge absolutely by Theorem 3.1.6.
Moreover, for $M \ge 1$, $\sum_{m=1}^{M} |r_m| \le \sum_{m=1}^{M} \sum_{n=1}^{\infty} |a_{mn}| = \lim_{N\to\infty} \sum_{m=1}^{M} \sum_{n=1}^{N} |a_{mn}| \le L$, so $\sum_m |r_m| < \infty$.
Similarly $\sum_n |c_n| < \infty$. Let $\varepsilon > 0$. Choose finite $F \subseteq \mathbb{N} \times \mathbb{N}$ with $\sum_{(m,n)\in(\mathbb{N}\times\mathbb{N})\setminus F} |a_{mn}| < \varepsilon$, and choose a rectangle $R = \{1, \ldots, M\} \times \{1, \ldots, N\} \supseteq F$. The full sum, row-iterated sum, and column-iterated sum each differ from $\sum_{(m,n)\in R} a_{mn}$ by at most $\sum_{(m,n)\in(\mathbb{N}\times\mathbb{N})\setminus R} |a_{mn}| \le \varepsilon$. Let $\varepsilon \downarrow 0$.

Structural Commentary: Finite rectangles plus absolute tails justify Fubini.
Key Insight: Double sums can be rearranged only after the entire array is absolutely summable. □

Method 3.1.38. *Series Proof Patterns*

(i) *Partial sums*: translate every assertion about $\sum a_n$ into an assertion about $s_n = \sum_{k=1}^{n} a_k$.
(ii) *Cauchy tails*: prove convergence by estimating $|\sum_{k=n}^{m} a_k|$ for $m \geq n$.
(iii) *Nonnegative terms*: replace convergence by boundedness of partial sums.
(iv) *Comparison*: dominate only the tail; finitely many early terms are irrelevant.
(v) *Absolute convergence*: use $\sum |a_n|$ when cancellation should not matter.
(vi) *Countable index sets*: choose a bijection $g : \mathbb{N} \to X$; use Theorem 3.1.36 so $\sum_{x \in X} a_x$ is independent of g.
(vii) *Double sums*: prove identities on $R_{M,N} := \{1, \ldots, M\} \times \{1, \ldots, N\}$; pass to ∞ arrays (Theorem 3.1.37).
(viii) *Conditional convergence*: order matters; use summation by parts or rearrangements.

Key Insight: Series arguments are sequence arguments with special notation for tails.

3.2 Uniform Convergence

Uniform convergence asks when limits of functions preserve structure. The key distinction is whether N may depend on x, which controls continuity, integration, differentiation, and the supremum metric.

Perspective 3.2.1. *Function Limits Before Function-Sequence Limits*
$\lim_{t \to x} f(t)$ denotes the metric-space function limit of Definition 2.3.4.
Uniform convergence is the hypothesis that lets the limits in n and t commute.

Definition 3.2.2. *Pointwise Limits of Functions*
Let $f_n : E \to (Y, d_Y)$.

(i) $f_n \to f$ *pointwise* on E if $d_Y(f_n(x), f(x)) \to 0$ for every $x \in E$.
(ii) If Y is scalar, $\sum f_n \to f$ *pointwise* on E if $s_N(x) := \sum_{n=1}^{N} f_n(x) \to f(x)$ for every $x \in E$.

No metric on E is required; topology matters only for preserving continuity, limits, or integrals.

Example 3.2.3. *Pointwise Limits Can Lose Structure*

(i) $s_{m,n} := m/(m+n)$ gives $\lim_n \lim_m s_{m,n} = 1$ but $\lim_m \lim_n s_{m,n} = 0$.
(ii) For $f_n(x) = x^2/(1+x^2)^n$, $\sum_{n=0}^{\infty} f_n$ equals 0 at 0 and $1 + x^2$ otherwise, by Theorem 3.1.8.
(iii) $f_n(x) = \sin(nx)/\sqrt{n}$ satisfies $f_n \to 0$ pointwise but $f_n'(0) = \sqrt{n} \to +\infty$.
(iv) $f_n(x) = n^2 x(1-x^2)^n$ on $[0,1]$ satisfies $f_n \to 0$ pointwise, while $\int_0^1 f_n \, dx = n^2/(2n+2) \to +\infty$.

Key Insight: Pointwise control can miss global spikes and oscillations.

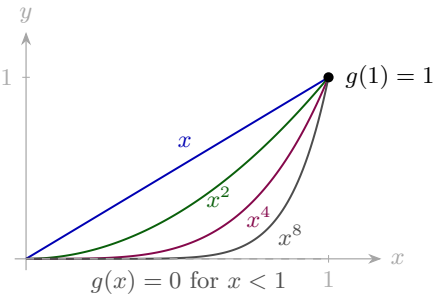


Figure 156: $x^n \to 0$ pointwise on $[0,1)$, but the endpoint prevents uniform convergence on $[0,1]$.

Example 3.2.4. *x^n Converges Pointwise But Not Uniformly*
On $[0,1]$, let $g_n(x) = x^n$. Then $g_n \to g$ pointwise, where $g(x) = \begin{cases} 0, & 0 \leq x < 1, \\ 1, & x = 1. \end{cases}$
The convergence is not uniform: for each n, some $x < 1$ has $|x^n - g(x)| = x^n > 1/2$.

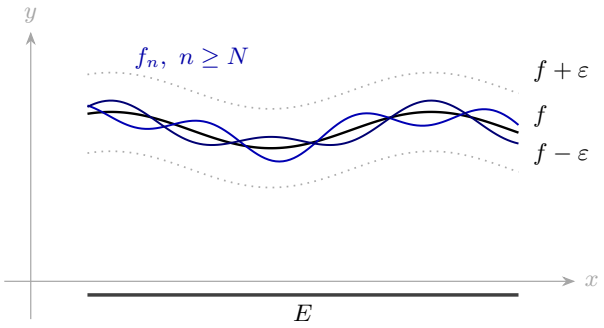


Figure 157: One index N puts all later graphs inside the same ε-strip around f on E.

Definition 3.2.5. *Uniform Convergence*
For $f_n, f : E \to (Y, d_Y)$, $f_n \rightrightarrows f$ on E if $\forall \varepsilon > 0\ \exists\ N \in \mathbb{N}$ such that $n \geq N,\ x \in E \implies d_Y(f_n(x), f(x)) \leq \varepsilon$.
In the scalar-valued case, take $d_Y(u, v) = |u - v|$; $\sum f_n$ converges uniformly on E if its partial sums do.

Perspective 3.2.6. *One Tail, All Points*
Pointwise convergence lets N depend on x; uniform convergence uses one N for all x.
Equivalently, late graphs lie in one ε-strip around f.

Theorem 3.2.7. *Cauchy Criterion for Uniform Convergence*
For complex-valued $f_n : E \to \mathbb{C}$, f_n converges uniformly on E if and only if

$$\forall \varepsilon > 0\ \exists N \in \mathbb{N} : \quad m, n \geq N,\ x \in E \implies |f_n(x) - f_m(x)| \leq \varepsilon$$

Proof.

($\Rightarrow$) If $f_n \rightrightarrows f$, choose N so that $|f_k(x) - f(x)| \leq \varepsilon/2$ for $k \geq N$ and $x \in E$. Then $|f_n(x) - f_m(x)| \leq \varepsilon$.
($\Leftarrow$) For each x, $(f_n(x))$ is Cauchy in $\mathbb{C}$, so define $f(x) := \lim f_n(x)$.
Fix $n \geq N$ and let $m \to \infty$ in the criterion; then $|f_n(x) - f(x)| \leq \varepsilon$ for all $x \in E$, hence $f_n \rightrightarrows f$.

Structural Commentary: The same N gives Cauchy control simultaneously for every $x \in E$.
Key Insight: Uniform Cauchy tails are exactly uniform convergence. □

Theorem 3.2.8. *Supremum Test for Uniform Convergence*
Suppose $f_n(x) \to f(x)\ \forall x \in E$, and $M_n := \sup_{x \in E} |f_n(x) - f(x)|$. Then $f_n \rightrightarrows f$ on E if and only if $M_n \to 0$.

Proof.

($\Rightarrow$) If $f_n \rightrightarrows f$, then $\forall \varepsilon > 0\ \exists\ N$ such that $n \geq N$ implies $|f_n(x) - f(x)| \leq \varepsilon$ for all $x \in E$. Hence $M_n \leq \varepsilon$.
($\Leftarrow$) If $M_n \to 0$, then $\forall \varepsilon > 0\ \exists\ N$ such that $n \geq N$ implies $M_n \leq \varepsilon$. Hence $|f_n(x) - f(x)| \leq \varepsilon$ for all $x \in E$.

Structural Commentary: Uniform convergence is convergence in the supremum error.
Key Insight: The pointwise errors must vanish as one global quantity. □

Theorem 3.2.9. *Weierstrass M-Test*
Let f_n be functions on E and suppose $|f_n(x)| \leq M_n$ for every $x \in E$.
If $\sum M_n$ converges, then $\sum f_n$ converges uniformly on E.

Proof. For $m \geq n$ and $x \in E$, $\left| \sum_{k=n}^{m} f_k(x) \right| \leq \sum_{k=n}^{m} M_k$.
Since $\sum M_k$ converges, Theorem 3.1.4 makes the right-hand side $< \varepsilon$ for all large m, n.
Thus the partial sums of $\sum f_k$ are uniformly Cauchy, so Theorem 3.2.7 applies.
Structural Commentary: Absolute numerical tails dominate functional tails.
Key Insight: A summable uniform majorant gives uniform convergence of a function series. □

Theorem 3.2.10. *Uniform Convergence Interchanges Limits*
Let E lie in a metric space, let x be a limit point of E, and suppose $f_n \rightrightarrows f$ on E.
If $A_n := \lim_{t \to x} f_n(t)$ exists for every n, then (A_n) converges and $\lim_{t \to x} f(t) = \lim_{n \to \infty} A_n$.

Proof. By Theorem 3.2.7, $\forall \varepsilon > 0\ \exists\ N$ such that $m, n \geq N$ and $t \in E$ imply $|f_n(t) - f_m(t)| \leq \varepsilon$.
Letting $t \to x$ gives $|A_n - A_m| \leq \varepsilon$, hence $A_n \to A$ for some $A \in \mathbb{C}$.
For any n, $|f(t) - A| \leq |f(t) - f_n(t)| + |f_n(t) - A_n| + |A_n - A|$. Given $\varepsilon > 0$, choose n so $|f(t) - f_n(t)| < \varepsilon/3$ for all $t \in E$ and $|A_n - A| < \varepsilon/3$; then take t sufficiently close to x so $|f_n(t) - A_n| < \varepsilon/3$. Thus $\lim_{t \to x} f(t) = A$.
Structural Commentary: Uniform convergence controls the outer two terms in the three-term estimate.
Key Insight: Uniformity is precisely what allows the $n \to \infty$ limit to move past the $t \to x$ limit. □

Theorem 3.2.11. *Uniform Limits Preserve Continuity*
If (f_n) is a sequence of continuous functions on E and $f_n \rightrightarrows f$ on E, then f is continuous on E.

Proof. Let $x \in E$. If x is isolated, f is continuous at x.
If x is a limit point, then $\lim_{t\to x} f_n(t) = f_n(x)$ for all n, so Theorem 3.2.10 gives $\lim_{t\to x} f(t) = \lim_{n\to\infty} f_n(x) = f(x)$.
Structural Commentary: Uniform convergence passes continuity to the limit.
Key Insight: Uniform convergence protects continuity from the failure modes in Example 3.2.3. □

Theorem 3.2.12. *Dini's Monotone Uniform Convergence Theorem*
Let K be compact and let $f_n, f : K \to \mathbb{R}$ be continuous. If $f_n(x) \downarrow f(x)$ for every $x \in K$, then $f_n \rightrightarrows f$ on K.

Proof. Set $g_n := f_n - f$. Then g_n is continuous, $g_n \downarrow 0$, and $g_n \geq 0$. Given $\varepsilon > 0$, let $K_n := \{x \in K : g_n(x) \geq \varepsilon\}$.
Then (K_n) is nested compact and $\bigcap_n K_n = \emptyset$, so Corollary 2.2.33 gives $K_N = \emptyset$ for some N.
Thus $0 \leq f_n(x) - f(x) < \varepsilon$ for all $x \in K$ and $n \geq N$.
Structural Commentary: Monotone convergence turns the bad sets into a nested compact sequence.
Key Insight: Compactness turns pointwise eventual escape from every bad set into a uniform escape. □

Definition 3.2.13. *The Supremum Metric on Bounded Continuous Functions*
For a metric space X, let $C_b(X)$ be the bounded complex-valued continuous functions and define $\|f\|_\infty := \sup_{x\in X} |f(x)|$; $d(f,g) := \|f - g\|_\infty$.
Then d is a metric on $C_b(X)$, induced by the norm of Definition 2.1.16; d-convergence is uniform convergence by Theorem 3.2.8. If X is compact, then $C(X) = C_b(X)$ by Corollary 2.3.34.

Perspective 3.2.14. *A Metric Space Whose Points Are Functions*
Uniform convergence is sequence convergence in the metric space $C_b(X)$.

Perspective 3.2.15. *Four Limit Notions*
The four limit notions are sequence limits, function limits, pointwise function limits, and uniform function limits. By Theorem 3.2.8, the last is ordinary sequence convergence in the supremum metric.

Theorem 3.2.16. *Uniform Convergence and Riemann–Stieltjes Integration*
Let α be increasing on $[a,b]$. If $f_n \in \mathcal{R}(\alpha)$ on $[a,b]$ and $f_n \rightrightarrows f$, then $f \in \mathcal{R}(\alpha)$ and $\int_a^b f\,d\alpha = \lim_{n\to\infty} \int_a^b f_n\,d\alpha$.

Proof. Put $\varepsilon_n := \sup_{a\leq x\leq b} |f_n(x) - f(x)|$. Then $\varepsilon_n \to 0$ by Theorem 3.2.8, and $f_n - \varepsilon_n \leq f \leq f_n + \varepsilon_n$.
Upper and lower integrals of f are therefore squeezed between corresponding integrals of $f_n \pm \varepsilon_n$, and gap is at most $2\varepsilon_n[\alpha(b) - \alpha(a)]$. Thus $f \in \mathcal{R}(\alpha)$. Same estimate gives $\left|\int_a^b f\,d\alpha - \int_a^b f_n\,d\alpha\right| \leq \varepsilon_n[\alpha(b) - \alpha(a)] \to 0$.
Structural Commentary: Uniform error becomes an integral error multiplied by the total α-variation.
Key Insight: Uniform convergence is strong enough to pass limits through the Riemann–Stieltjes integral. □

Corollary 3.2.17. *Termwise Integration of Uniformly Convergent Series*
If $f_n \in \mathcal{R}(\alpha)$ on $[a,b]$ and $\sum f_n$ converges uniformly to f, then $\int_a^b f\,d\alpha = \sum_{n=1}^{\infty} \int_a^b f_n\,d\alpha$.

Proof. Let $s_N := \sum_{n=1}^N f_n$. By linearity, $s_N \in \mathcal{R}(\alpha)$ and $\int_a^b s_N\,d\alpha = \sum_{n=1}^N \int_a^b f_n\,d\alpha$.
Since $s_N \rightrightarrows f$, Theorem 3.2.16 gives $\int_a^b f\,d\alpha = \lim_{N\to\infty} \int_a^b s_N\,d\alpha = \sum_{n=1}^{\infty} \int_a^b f_n\,d\alpha$.
Structural Commentary: Reduce the series to its uniformly convergent partial sums.
Key Insight: Uniform convergence lets integration be performed term by term. □

Theorem 3.2.18. *Uniform Convergence and Differentiation*
Let f_n be differentiable on $[a,b]$, suppose $f_n(x_0)$ converges for some $x_0 \in [a,b]$, and suppose $f_n' \rightrightarrows g$ on $[a,b]$. Then $f_n \rightrightarrows f$ for some f, and $f' = g$ on $[a,b]$ (one-sided at endpoints).

Proof. For m, n, MVT 4.1.13 gives $|f_n(x) - f_m(x)| \leq |f_n(x_0) - f_m(x_0)| + (b-a) \sup_{t\in[a,b]} |f_n'(t) - f_m'(t)|$.
Thus (f_n) is uniformly Cauchy, so $f_n \rightrightarrows f$. Fix x and set $\phi_n(t) := \frac{f_n(t)-f_n(x)}{t-x}$, $\phi(t) := \frac{f(t)-f(x)}{t-x}$.
For $t \neq x$, the same estimate gives $|\phi_n(t) - \phi_m(t)| \leq \|f_n' - f_m'\|_\infty$, hence $\phi_n \rightrightarrows \phi$.
By Theorem 3.2.10, $f'(x) = \lim_{t\to x} \phi(t) = \lim_{n\to\infty} f_n'(x) = g(x)$.
Structural Commentary: Derivative control anchors functions, then quotients.
Key Insight: Uniform derivative convergence plus one anchor passes derivatives to the limit. □

Method 3.2.19. *Uniform Convergence Proof Patterns*
(i) *Pointwise*: fix x, then choose $N = N(x, \varepsilon)$.
(ii) *Uniform*: choose $N = N(\varepsilon)$; equivalently $\sup_E |f_n - f| \to 0$.

(iii) *Uniform Cauchy*: prove $\sup_E |f_n - f_m| \to 0$ as $m, n \to \infty$.
(iv) *Preservation*: use a three-term estimate and choose the uniform index first.
(v) *Series of functions*: apply Weierstrass M-test when numerical absolute tails dominate all functional tails.
(vi) *Termwise operations*: require $f_n \to f$ locally uniformly, i.e. $\sup_{x\in K} |f_n - f| \to 0$ for each compact K, before $\lim_n \int_K f_n = \int_K f$; for derivatives, require $f_n' \to g$ locally uniformly and get $(\lim_n f_n)' = g$.

Key Insight: Uniformity means N is chosen before x.

3.3 Function-Space Compactness

Function-space compactness treats functions as points. Supremum metrics preview Banach spaces, Baire category describes generic behavior, and equicontinuity leads to Arzela–Ascoli.

Theorem 3.3.1. *Completeness of the Supremum Metric*
The metric space $C_b(X)$ with metric $d(f,g) = \|f-g\|_\infty$ is complete.

Proof. Let (f_n) be Cauchy in $C_b(X)$. Then Theorem 3.2.7 gives a function f with $f_n \rightrightarrows f$. By Theorem 3.2.11, f is continuous. For some n, $\|f - f_n\|_\infty < 1$, hence f_n is bounded, thus $f \in C_b(X)$ and $\|f_n - f\|_\infty \to 0$.
Structural Commentary: Uniform Cauchy convergence builds the limit inside the same function space.
Key Insight: Completeness of $\mathbb{C}$ plus uniform convergence makes $C_b(X)$ complete. □

Definition 3.3.2. *Bounded Target Metric for All Functions*
For metric spaces (X, d_X) and (Y, d_Y), let $\mathcal{F} = Y^X$ and set $\rho(y,z) := \frac{d_Y(y,z)}{1+d_Y(y,z)}$ and $d_\rho(f,g) := \sup_{x\in X} \rho(f(x), g(x))$. Then $\rho < 1$, so d_ρ is finite on all of $\mathcal{F}$; d_ρ-convergence is uniform convergence measured by d_Y.

Theorem 3.3.3. *Completeness and Closed Subspaces for d_ρ*
If the target metric space (Y, d_Y) is complete, then:

(i) $(\mathcal{F}, d_\rho)$ *is complete;*
(ii) $\mathcal{F}_b := \{f \in \mathcal{F} : f(X)$ *is d_Y-bounded$\}$ is closed in $\mathcal{F}$;*
(iii) $C(X,Y)$ *is closed in $\mathcal{F}$.*

Proof.

(i) Let (f_n) be d_ρ-Cauchy. For each $x \in X$, $(f_n(x))$ is Cauchy in Y; define $f(x) := \lim_n f_n(x)$.
Given $\varepsilon > 0$, choose N with $d_\rho(f_n, f_m) < \varepsilon$ for $m, n \geq N$. Letting $m \to \infty$ gives $d_\rho(f_n, f) \leq \varepsilon$ for $n \geq N$.
(ii) If $f_n \in \mathcal{F}_b$ and $d_\rho(f_n, f) \to 0$, choose N with $d_\rho(f_N, f) < 1/2$.
Then $d_Y(f_N(x), f(x)) < 1$ for all x. Since $f_N(X)$ is bounded, so is $f(X)$.
(iii) If $f_n \in C(X,Y)$ and $d_\rho(f_n, f) \to 0$, then $f_n \to f$ uniformly in d_Y.
Given $x_0 \in X$ and $\varepsilon > 0$, choose N with $d_Y(f_N(x), f(x)) < \varepsilon/3$ for all x, then use continuity of f_N at x_0 and the triangle inequality. Thus f is continuous.

Structural Commentary: Compress the target metric; then close the useful subspaces.
Key Insight: Completeness and uniform control make boundedness and continuity closed. □

Theorem 3.3.4. *A Continuous Nowhere Differentiable Function Exists*
There exists a real continuous function on $\mathbb{R}$ which is nowhere differentiable.

Proof. Let φ be the 2-periodic extension of $|x|$ on $[-1,1]$. Then $0 \leq \varphi \leq 1$ and $|\varphi(s) - \varphi(t)| \leq |s-t|$.
Define $f(x) := \sum_{n=0}^{\infty} \left(\frac{3}{4}\right)^n \varphi(4^n x)$, converges uniformly (Theorem 3.2.9); hence f is continuous (Theorem 3.2.11).
Fix x and $m \in \mathbb{N}$. Choose $\delta_m = \pm\frac{1}{2}4^{-m}$ so that no integer lies between $4^m x$ and $4^m(x+\delta_m)$.
Set $\gamma_{n,m} := \frac{\varphi(4^n(x+\delta_m)) - \varphi(4^n x)}{\delta_m}$. If $n > m$, then $4^n \delta_m$ is an even integer, so $\gamma_{n,m} = 0$.
If $0 \leq n \leq m$, then $|\gamma_{n,m}| \leq 4^n$, and the choice of δ_m gives $|\gamma_{m,m}| = 4^m$.
Hence $\left|\frac{f(x+\delta_m) - f(x)}{\delta_m}\right| = \left|\sum_{n=0}^{m} \left(\frac{3}{4}\right)^n \gamma_{n,m}\right| \geq 3^m - \sum_{n=0}^{m-1} 3^n = \frac{1}{2}(3^m + 1)$
Since $\delta_m \to 0$, the difference quotients are unbounded near x; since x was arbitrary, f is nowhere differentiable.
Structural Commentary: Uniform sums give continuity; scaled cusps force unbounded quotients.
Key Insight: Uniform convergence preserves continuity, not differentiability. □

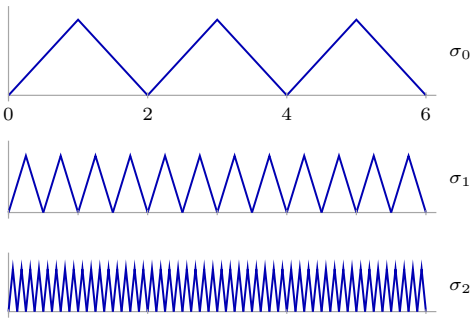


Figure 158: The first compressed sawtooth terms $\sigma_k(x) = (3/4)^k \varphi(4^k x)$: heights shrink, slopes grow.

Perspective 3.3.5. *Explicit Versus Generic Pathology*
The sawtooth series proves existence; Baire category shows nowhere differentiability is residual in $C([a,b],\mathbb{R})$.

Theorem 3.3.6. *Generic Nowhere Differentiability on Compact Intervals*
Let $a < b$, and let $C = C([a,b],\mathbb{R})$ with the supremum metric. There exists a residual set $\mathcal{G} \subseteq C$ such that:

(i) no $f \in \mathcal{G}$ has an available one-sided derivative at any point of $[a,b]$;
(ii) no $f \in \mathcal{G}$ is monotone on any nondegenerate subinterval of $[a,b]$;
(iii) the set $\{f \in C : f$ is differentiable at some point of $[a,b]\}$ is first category.

Proof.

(i) Write $\Delta_h f(x) := f(x+h) - f(x)$, and choose N_0 with $1/N_0 < b - a$. For $n \geq N_0$, define

$$R_n := \left\{ f \in C : \forall x \in [a, b - 1/n] \ \exists h > 0 \text{ with } \left| \frac{\Delta_h f(x)}{h} \right| > n \right\},$$

$$L_n := \left\{ f \in C : \forall x \in [a + 1/n, b] \ \exists h < 0 \text{ with } \left| \frac{\Delta_h f(x)}{h} \right| > n \right\},$$

$$G_n := \{ f \in C : f \text{ is nonmonotone on every subinterval of length } 1/n \}.$$

Each R_n, L_n, G_n is open and dense in C (Weierstrass Approximation Theorem 3.4.15), plus small high-slope sawtooth perturbation; openness follows from compactness and uniform perturbation estimates.

(ii) Set $\mathcal{G} := \bigcap_{n \geq N_0} (R_n \cap L_n \cap G_n)$. By Definition 2.2.47, $\mathcal{G}$ is residual. If $f \in \mathcal{G}$, then the R_n and L_n conditions give one-sided difference quotients of arbitrarily large absolute value; since f is bounded, the corresponding increments tend to 0. Thus no available one-sided derivative exists.

(iii) Let $f \in \mathcal{G}$. Every nondegenerate subinterval contains a subinterval of length $1/n$ for all sufficiently large n, so the G_n conditions make f nonmonotone on every nondegenerate subinterval. If f is differentiable somewhere, then $f \notin \mathcal{G}$; hence the differentiable-somewhere set lies in $C \setminus \mathcal{G}$, which is first category.

Structural Commentary: Baire category turns dense sawtooth perturbability into a generic regularity failure.
Key Insight: The explicit construction shows existence; the category proof shows typicality. □

Definition 3.3.7. *Boundedness and Equicontinuity of Function Families*
Let $\mathcal{F}$ be a family of functions from E into a metric space (Y, d_Y); in scalar cases, take $d_Y(u,v) = |u - v|$.

(i) (f_n) is *pointwise bounded* on E if $(f_n(x))$ is bounded in Y for each $x \in E$.
(ii) A scalar sequence (f_n) is *uniformly bounded* on E if $\exists M$ such that $|f_n(x)| < M$ for all $x \in E, n \in \mathbb{N}$.
(iii) If E lies in a metric space (X, d_X), then $\mathcal{F}$ is *equicontinuous* on E if $\forall \varepsilon > 0 \ \exists \delta > 0$, such that $d_X(x,y) < \delta, \ x, y \in E, \ f \in \mathcal{F} \implies d_Y(f(x), f(y)) < \varepsilon$.
(iv) The family is *pointwise equicontinuous* if for every $p \in E$ and $\varepsilon > 0$ there is $\delta = \delta(p, \varepsilon) > 0$ such that $d_X(p,y) < \delta, \ y \in E, \ f \in \mathcal{F}$ implies $d_Y(f(p), f(y)) < \varepsilon$.

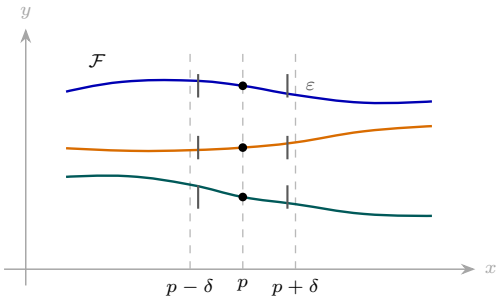


Figure 159: One input scale δ controls every curve in $\mathcal{F}$ within tolerance ε.

Theorem 3.3.8. *Pointwise Equicontinuity Becomes Uniform on Compact Sets*
Every pointwise equicontinuous family $\mathcal{F} \subset C(K, \mathbb{C})$ on compact K is equicontinuous.

Proof. Otherwise, choose $\varepsilon > 0$, $f_n \in \mathcal{F}$, and $x_n, t_n \in K$ with $d(x_n, t_n) < 1/n$ but $|f_n(x_n) - f_n(t_n)| \geq \varepsilon$.
By Theorem 2.2.3, pass to $x_n \to x_0$; then $t_n \to x_0$ too.
Pointwise equicontinuity at x_0 gives $|f_n(x_n) - f_n(t_n)| < \varepsilon$ for large n, a contradiction.
Structural Commentary: Compactness turns moving counterexamples into one base point.
Key Insight: The local $\delta(x, \varepsilon)$ becomes one uniform δ only after compactness prevents escape. □

Example 3.3.9. *Uniform Boundedness Alone Is Not Enough*
On $[0, 1]$, let $f_n(x) := \frac{x^2}{x^2+(1-nx)^2}$. Then $|f_n| \leq 1$ and $f_n \to 0$ pointwise, but $f_n(1/n) = 1$; hence no uniformly convergent subsequence. The missing hypothesis is equicontinuity.

Theorem 3.3.10. *Diagonal Selection on a Countable Set*
Every pointwise bounded sequence (f_n) on countable E has a subsequence converging pointwise on E.

Proof. List $E = \{x_1, x_2, \ldots\}$. Inductively choose nested subsequences S_i so that $S_i(x_i)$ converges.
The diagonal subsequence is eventually contained in each S_i, hence converges at every x_i.
Structural Commentary: The diagonal keeps the ith convergence property after the ith step.
Key Insight: Countability permits diagonal convergence. □

Theorem 3.3.11. *Uniform Convergence Implies Equicontinuity on Compact Sets*
If K is compact, $f_n \in C_b(K)$, and $f_n \rightrightarrows f$ on K, then $\{f_n : n \in \mathbb{N}\}$ is equicontinuous.

Proof. Choose N so $\|f_n - f_N\|_\infty < \varepsilon$ for $n > N$.
By Theorem 2.3.39, choose one $\delta > 0$ giving ε-control for $f_1, \ldots, f_N$.
If $n > N$, $d(x, y) < \delta$, then $|f_n(x) - f_n(y)| \leq |f_n(x) - f_N(x)| + |f_N(x) - f_N(y)| + |f_N(y) - f_n(y)| < 3\varepsilon$.
Together with the finite initial family, $\{f_n\}$ is equicontinuous.
Structural Commentary: A uniformly convergent tail stays close to one uniformly continuous function.
Key Insight: Compactness handles finitely many functions; uniform convergence handles the tail. □

Theorem 3.3.12. *Arzela–Ascoli Selection Theorem*
Let K be compact. If $(f_n) \subset C_b(K)$ is pointwise bounded and equicontinuous on K, then:

(i) (f_n) is uniformly bounded on K;
(ii) (f_n) contains a uniformly convergent subsequence.

Proof.

(i) Choose $\delta > 0$ from equicontinuity with tolerance 1. Compactness gives $K \subseteq \bigcup_{j=1}^{r} B_\delta(p_j)$.
Pointwise boundedness gives M_j such that $|f_n(p_j)| \leq M_j \forall n$.
If $x \in B_\delta(p_j)$, then $|f_n(x)| \leq |f_n(x) - f_n(p_j)| + |f_n(p_j)| < 1 + M_j$, seq uniformly bounded by $1 + \max_j M_j$.

(ii) For each m, compactness gives a finite $1/m$-net in K; the union of these finite nets is a countable dense set $E \subset K$. By Theorem 3.3.10, some subsequence (g_i) converges pointwise on E.
Given $\varepsilon > 0$, choose $\delta > 0$ from equicontinuity with tolerance $\varepsilon/3$, and points $x_1, \ldots, x_r \in E$ with $K \subseteq \bigcup_s B_\delta(x_s)$. As $(g_i(x_s))$ converges for each of finitely many x_s, $\exists N$ s.t. $i, j \geq N$ implies $|g_i(x_s) - g_j(x_s)| < \varepsilon/3 \ \forall s$. If $d(x, x_s) < \delta$, then $|g_i(x) - g_j(x)| \leq |g_i(x) - g_i(x_s)| + |g_i(x_s) - g_j(x_s)| + |g_j(x_s) - g_j(x)| < \varepsilon$.
Thus (g_i) is uniformly Cauchy and converges uniformly by Theorem 3.2.7.

Structural Commentary: Dense convergence plus equicontinuity becomes uniform convergence.
Key Insight: Equicontinuity is exactly the compactness-compatible hypothesis missing from Example 3.3.9. □

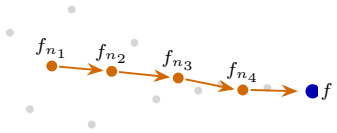


Figure 160: Arzela–Ascoli selects a uniformly convergent subsequence.

Definition 3.3.13. *σ-Compact Metric Space*
A metric space X is *σ-compact* if $X = \bigcup_{j=1}^{\infty} K_j$ for compact subsets $K_j \subseteq X$.

Theorem 3.3.14. *Arzela–Ascoli on σ-Compact Spaces*
Let X be σ-compact, Y a metric space, and (f_n) pointwise equicontinuous maps $X \to Y$. If each orbit $\{f_n(x) : n \in \mathbb{N}\}$ lies in a compact subset of Y, then some subsequence converges uniformly on compact subsets of X.

Proof. Choose an increasing compact exhaustion $X = \bigcup_j K_j$.
By Theorem 3.3.8, each restriction $f_n|_{K_j}$ is uniformly equicontinuous. Applying Theorem 3.3.12 on $K_1, K_2, \dots$ and diagonalizing gives a subsequence (g_m) converging uniformly on every K_j, hence pointwise on X.
If $K \subseteq X$ is compact, then $(g_m|_K)$ is uniformly equicontinuous, and a finite-net argument upgrades pointwise convergence on K to uniform convergence on K.
Structural Commentary: Exhaust X by compact pieces and diagonalize the compact theorem.
Key Insight: Local uniform convergence replaces global uniform convergence. □

Corollary 3.3.15. *Real-Line Local Uniform Subsequence*
If (f_n) is pointwise equicontinuous on $\mathbb{R}$ and bounded at one point, then it has a subsequence converging uniformly on compact subsets of $\mathbb{R}$.

Proof. Let $(f_n(x_0))$ be bounded. On any compact interval containing x_0, Theorem 3.3.8 gives uniform equicontinuity; finitely many δ-steps from x_0 to x bound each orbit $(f_n(x))$. Hence Theorem 3.3.14 applies on $\mathbb{R}$.
Structural Commentary: One bounded anchor propagates across compact intervals.
Key Insight: Pointwise equicontinuity turns local boundedness into compactwise compactness. □

Method 3.3.16. *Function-Space Compactness Proof Patterns*

(i) *Completeness*: prove a function sequence is uniformly Cauchy, build the pointwise limit, and use uniform convergence to keep continuity or boundedness.
(ii) *Baire category*: show a desired property contains a countable intersection of dense open subsets of a complete function space.
(iii) *Compactness/equicontinuity*: pass from finite nets to all of K.
(iv) *Diagonal selection*: get pointwise convergence on a countable dense set, then upgrade using equicontinuity.
(v) *Locally uniform*: exhaust a σ-compact domain by compact sets and diagonalize.

Key Insight: Compactness in function spaces is finite-net control plus equicontinuity.

3.4 Power Series and Approximation

Power series are controlled infinite polynomials, and approximation is uniform closure. This subsection covers radii, Taylor structure, Abel behavior, Weierstrass, Stone–Weierstrass, and elementary analytic functions.

Motivation 3.4.1. *Power Series as Controlled Infinite Polynomials*
Power-series manipulations rest on uniform convergence: Definition 3.1.1, Theorems 3.4.4, 3.2.9, and 3.2.18.

Definition 3.4.2. *Core Objects: Power Series, Approximation, and Gamma*

(i) A real *power series about a* is $\sum_{n=0}^{\infty} c_n(x-a)^n$; a local representation of f by one makes f *analytic* at a.
(ii) A *polynomial approximant* to f on a compact set K is a polynomial P chosen so that $\|f - P\|_{\infty,K}$ is small.
(iii) A function algebra is a collection closed under addition, multiplication, and scalar multiplication; Stone–Weierstrass asks when such algebras are uniformly dense in $C(K)$.
(iv) For $x > 0$, the *gamma function* is $\Gamma(x) := \int_0^{\infty} t^{x-1}e^{-t}\,dt$ when the improper integral converges.

Key Insight: Power series are local analytic objects; approximation theorems turn algebraic families into dense function families; Γ extends factorials by integration.

Definition 3.4.3. *Power Series and Radius of Convergence*
Given a complex sequence (c_n), the series $\sum_{n=0}^{\infty} c_n z^n$ is a *power series*. If $\alpha := \limsup_{n\to\infty} \sqrt[n]{|c_n|}$, define $R := 1/\alpha$, with $R = +\infty$ when $\alpha = 0$ and $R = 0$ when $\alpha = +\infty$. The number R is the *radius of convergence*.

Theorem 3.4.4. *Convergence Inside the Circle*
For $\sum c_n z^n$, radius R: converge if $|z| < R$, diverge if $|z| > R$; boundary $|z| = R$ is undecided.

Proof. For $a_n = c_n z^n$, $\limsup_{n\to\infty} \sqrt[n]{|a_n|} = |z| \limsup_{n\to\infty} \sqrt[n]{|c_n|} = \frac{|z|}{R}$, with extended meaning if $R = 0$ or $R = +\infty$.
If $|z| < R$, this value is < 1; if $|z| > R$, it is > 1. Apply Theorem 3.1.19.
Structural Commentary: Radius controls the interior and exterior, not the boundary.
Key Insight: The root test is the radius-of-convergence theorem in disguise. □

Theorem 3.4.5. *A Boundary Test for Power Series*
If $\sum c_n z^n$ has radius 1, $c_n \downarrow 0$, and $c_n \geq 0$, then it converges for $|z| = 1$, except possibly at $z = 1$.

Proof. For $|z| = 1$ and $z \neq 1$, put $a_n = z^n$ and $b_n = c_n$.
The partial sums of $\sum z^n$ are bounded because $\left|\sum_{m=0}^{n} z^m\right| = \left|\frac{1-z^{n+1}}{1-z}\right| \leq \frac{2}{|1-z|}$. Theorem 3.1.23 applies.
Structural Commentary: Boundary convergence away from $z = 1$ comes from bounded geometric partial sums.
Key Insight: Monotone coefficients can make the unit-circle boundary converge without absolute convergence. □

Perspective 3.4.6. *Centered Series Are Translated Series*
$\sum c_n(x-a)^n$ is $\sum c_n y^n$ with $y = x - a$. Its radius is R from Definition 3.4.3: absolute convergence for $|x-a| < R$, divergence for $|x-a| > R$, and separate endpoint tests for $|x-a| = R$.

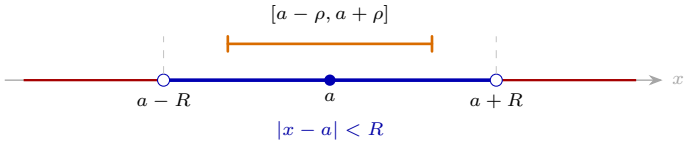


Figure 161: Inside $|x-a| < R$: absolute and locally uniform; outside divergence; endpoints undecided.

Definition 3.4.7. *Real Analytic at a Point*
Let $E \subseteq \mathbb{R}$, let $f : E \to \mathbb{R}$, and let a be an interior point of E. The function f is *real analytic at a* if there are $r > 0$ and coefficients (c_n) such that $(a-r, a+r) \subseteq E$, the power series $\sum c_n(x-a)^n$ has radius of convergence at least r, and $f(x) = \sum_{n=0}^{\infty} c_n(x-a)^n$ for all $x \in (a-r, a+r)$.
If E is open and f is real analytic at every point of E, then f is *real analytic on E*.

Theorem 3.4.8. *Power Series May Be Differentiated Term by Term*
Let $f(x) = \sum_{n\geq 0} c_n x^n$ have radius $R > 0$, and $I_R = (-R, R)$. For $0 < \rho < R$, let $K_\rho = [-\rho, \rho]$, a closed interval contained in I_R. Then $\sum c_n x^n$ converges uniformly on each K_ρ, and $f \in C^\infty(I_R)$.
For $k \geq 1$, $f^{(k)}(x) = \sum_{n\geq k} \frac{n!}{(n-k)!} c_n x^{n-k}$, $(x \in I_R)$, and $f^{(k)}(0) = k! c_k$.

Proof. For $0 < \rho < R$, Theorem 3.4.4 and the root-test argument give $\sum |c_n| \rho^n < \infty$; hence the Weierstrass M-test, Theorem 3.2.9, gives uniform convergence on K_ρ.
Since $\sqrt[n]{n} \to 1$, polynomial factors in n do not change the radius. Thus each differentiated series converges uniformly on K_ρ, and Theorem 3.2.18, iterated, gives the formula and $f^{(k)}(0) = k! c_k$.
Structural Commentary: On each K_ρ, termwise operations are justified.
Key Insight: The derivative coefficients grow only by polynomial factors, which do not change the radius. □

Corollary 3.4.9. *Termwise Calculus for Centered Power Series*
Let $f(x) = \sum_{n=0}^{\infty} c_n(x-a)^n$ have radius $R > 0$. On $(a-R, a+R)$ the series converges locally uniformly and defines a continuous function. For $|x-a| < R$, $f'(x) = \sum_{n=1}^{\infty} n c_n (x-a)^{n-1}$ with local uniform convergence of the derivative series. If $[y, z] \subset (a-R, a+R)$, then $\int_y^z f(x)\,dx = \sum_{n=0}^{\infty} c_n \frac{(z-a)^{n+1} - (y-a)^{n+1}}{n+1}$.

Proof. Set $g(u) = \sum c_n u^n$ and $u = x - a$.
Theorem 3.4.8 applied to g on $(-R, R)$ gives convergence, continuity, derivative formula after translating.
For $[y, z] \subset (a-R, a+R)$, choose $\rho < R$ with $[y, z] \subset [a-\rho, a+\rho]$; then Corollary 3.2.17 lets the integrals of partial sums pass to the uniform limit.
Structural Commentary: All local termwise operations happen on closed intervals strictly inside the radius.
Key Insight: Translation reduces centered power series to the radius-R theorem about $\sum c_n u^n$. □

Corollary 3.4.10. *Analytic Functions Are Their Taylor Series*
If $f(x) = \sum c_n(x-a)^n$ near a, then f is C^∞, expansion f is unique, and $f(x) = f(a) + \sum_{n=1}^{\infty} \frac{f^{(n)}(a)}{n!}(x-a)^n$.

Proof. Corollary 3.4.9, applied repeatedly, gives $c_0 = f(a)$ and $c_k = f^{(k)}(a)/k!$.
Hence two expansions about a have the same coefficients.
Structural Commentary: Derivatives determine the local series.
Key Insight: Once a power-series representation exists, the coefficients cannot vary. □

Perspective 3.4.11. *Taylor Coefficients Are Only the First Half*
Power-series coefficients are forced by derivatives (Corollary 3.4.10); for a function not already known to be analytic, proving the Taylor series represents the function needs a remainder estimate (Theorem 4.1.26).

Theorem 3.4.12. *Abel's Endpoint Theorem*
Let $s_n := \sum_{k=0}^{n} c_k$ with $s_n \to s$, and let $f(x) = \sum_{n\geq 0} c_n x^n$ on $(-1,1)$. Then $f(x) \to s$, for $x \to 1^-$.

Proof. With $s_{-1} = 0$, summation by parts, Theorem 3.1.22, gives for $|x| < 1$, $f(x) = (1-x)\sum_{n=0}^{\infty} s_n x^n$.
Since $(1-x)\sum_{n\geq 0} x^n = 1$, then $f(x) - s = (1-x)\sum_{n\geq 0}(s_n - s)x^n$. Choose N so that $|s_n - s| < \varepsilon/2$ for $n > N$.
The tail is at most $\varepsilon/2$, and the finite part tends to 0 after multiplication by $1-x$.
Structural Commentary: Endpoint convergence becomes one-sided continuity.
Key Insight: The factor $(1-x)$ averages the partial sums s_n with geometric weights. □

Theorem 3.4.13. *Abel Product Theorem for Cauchy Products*
Suppose $\sum a_n = A$, $\sum b_n = B$, $\sum c_n = C$, and $c_n = \sum_{k=0}^{n} a_k b_{n-k}$. Then $C = AB$.

Proof. For $0 \leq x < 1$, set $f(x) := \sum_{n\geq 0} a_n x^n$; $g(x) := \sum_{n\geq 0} b_n x^n$; and $h(x) := \sum_{n\geq 0} c_n x^n$.
Since convergent series have bounded terms, f, g, h converge absolutely for $|x| < 1$. Thus the Cauchy product gives $fg = h$ on $[0,1)$. By Theorem 3.4.12, $f \to A$, $g \to B$, and $h \to C$ as $x \to 1^-$. Hence $AB = C$.
Structural Commentary: Prove $fg = h$ inside, then pass to $x = 1$.
Key Insight: Inside $|x| < 1$, the Cauchy product is safe; Abel then moves the identity to $x = 1$. □

Theorem 3.4.14. *Re-Expansion and Identity for Power Series*
Let $f(x) = \sum_{n\geq 0} c_n x^n$ have radius $R > 0$, and set $S = (-R, R)$.

(i) If $a \in S$, then $f(x) = \sum_{m\geq 0} \frac{f^{(m)}(a)}{m!}(x-a)^m$, $(|x-a| < R - |a|)$

(ii) Let $F = \sum_{n\geq 0} a_n x^n$ and $G = \sum_{n\geq 0} b_n x^n$ on S.
If $\mathcal{Z}_{F,G} := \{x \in S : F(x) = G(x)\}$ has a limit point in S, then $a_n = b_n$ for all $n \geq 0$, hence $F = G$ on S.

Proof.

(i) Write $x^n = ((x-a)+a)^n$. If $|x-a| + |a| < R$, then $\sum_{n\geq 0} |c_n| \left(|x-a| + |a|\right)^n < \infty$, so the resulting double sum may be rearranged. Theorem 3.4.8 gives the coefficients $f^{(m)}(a)/m!$.

(ii) Subtract the series and call the difference f. Let $Z = \{x \in S : f(x) = 0\}$, and let A be the set of limit points of Z in S. Then $A \neq \emptyset$. For $x_0 \in A$, part (i) gives $f(x) = \sum d_n(x-x_0)^n$.
If d_k were the first nonzero coefficient, then $f(x) = (x-x_0)^k g(x)$ with $g(x_0) \neq 0$, contradicting $x_0 \in A$. Hence $f = 0$ near x_0, so A is open. Also $S \setminus A$ is open; by connectedness, Theorem 2.3.2, $A = S$. Since Z is closed, $f = 0$ on S, and evaluating derivatives at 0 gives equality of coefficients.

Structural Commentary: Power-series equality is local-to-global on an interval.
Key Insight: A limit point of zeros forces all Taylor coefficients to vanish. □

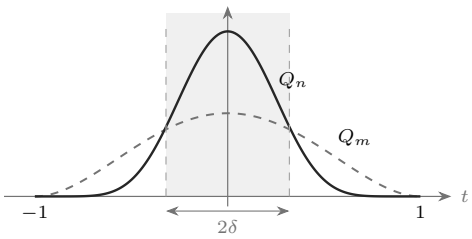


Figure 162: The normalized kernels $Q_n(t) = c_n(1-t^2)^n$ concentrate near 0 while retaining integral 1.

Theorem 3.4.15. *Weierstrass Polynomial Approximation*
If $f : [a,b] \to \mathbb{C}$ is continuous, then polynomials P_n satisfy $P_n \rightrightarrows f$ on $[a,b]$. If f is real-valued, take P_n real.

Proof. By an affine change of variables and subtraction of the endpoint interpolant, assume $[a,b] = [0,1]$ and $f(0) = f(1) = 0$; extend f by 0 outside $[0,1]$. For $n \geq 1$, let $Q_n(x) := c_n(1-x^2)^n$ and $\int_{-1}^{1} Q_n(x)\,dx = 1$.
Source estimates give $c_n < \sqrt{n}$ and $Q_n(t) \leq \sqrt{n}(1-\delta^2)^n$ for $\delta \leq |t| \leq 1$.
Define $P_n(x) := \int_{-1}^{1} f(x+t)Q_n(t)\,dt$, $(0 \leq x \leq 1)$. By change of variables, $P_n(x) = \int_0^1 f(u)Q_n(u-x)\,du$, so P_n is a polynomial in x (real if f is real). Given $\varepsilon > 0$, choose $\delta > 0$ so $|u-v| < \delta$ implies $|f(u) - f(v)| < \varepsilon/2$, and set $M = \sup|f|$. Since $Q_n \geq 0$ and $\int Q_n = 1$, $|P_n(x) - f(x)| \leq \int_{-1}^{1} |f(x+t) - f(x)|\,Q_n(t)\,dt \leq \frac{\varepsilon}{2} + 4M\sqrt{n}(1-\delta^2)^n$.
The last term tends to 0, so $P_n \rightrightarrows f$ on $[0,1]$. Undoing the reductions proves the theorem.
Structural Commentary: The kernels Q_n form polynomial approximate identities.
Key Insight: Uniform continuity controls the near part, and concentration of Q_n kills the far part. □

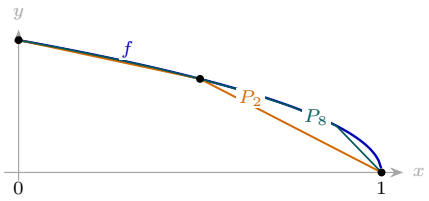


Figure 163: Finer piecewise-linear interpolants approach $f(x) = \sqrt{1-x}$ uniformly on $[0,1]$.

Corollary 3.4.16. *Polynomial Approximation of $|x|$*
For every interval $[-a,a]$, there are real polynomials P_n such that $P_n(0) = 0$ and $P_n \rightrightarrows |x|$ on $[-a,a]$.

Proof. By Theorem 3.4.15, choose real polynomials P_n^* converging uniformly to $|x|$ on $[-a,a]$.
Since $P_n^*(0) \to 0$, the polynomials $P_n(x) := P_n^*(x) - P_n^*(0)$ have $P_n(0) = 0$ and still converge uniformly to $|x|$.
Structural Commentary: Normalize the approximants at the base point.
Key Insight: Subtracting the value at 0 fixes the normalization without changing the uniform limit. □

Definition 3.4.17. *Algebras and Uniform Closure*
Let $\mathcal{A}$ be a family of complex-valued functions on a set E.

(i) $\mathcal{A}$ is an *algebra* if it is closed under $+$, products, and scalar multiplication (real scalars in the real case).
(ii) $\mathcal{A}$ is *uniformly closed* if $f_n \in \mathcal{A}$ and $f_n \rightrightarrows f$ on E imply $f \in \mathcal{A}$.
(iii) The *uniform closure* of $\mathcal{A}$ is $\{f : f_n \in \mathcal{A},\ f_n \rightrightarrows f \text{ on } E\}$.

Theorem 3.4.18. *Uniform Closures Are Closed Algebras*
The uniform closure of an algebra of bounded functions is a uniformly closed algebra.

Proof. If $f, g \in \mathcal{B}$, choose $f_n, g_n \in \mathcal{A}$ with $f_n \rightrightarrows f$ and $g_n \rightrightarrows g$.
Boundedness and uniform convergence give $f_n + g_n \rightrightarrows f + g$; $f_n g_n \rightrightarrows fg$; and $cf_n \rightrightarrows cf$. Thus $\mathcal{B}$ is an algebra.
If $h_n \in \mathcal{B}$ and $h_n \rightrightarrows h$, approximate each h_n uniformly by some $a_n \in \mathcal{A}$ within $1/n$; then $a_n \rightrightarrows h$, so $h \in \mathcal{B}$.
Structural Commentary: Uniform limits preserve algebraic operations for bounded functions.
Key Insight: The uniform closure of an algebra has already absorbed all further uniform limits. □

Definition 3.4.19. *Separation, Nonvanishing, and Self-Adjointness*
Let $\mathcal{A}$ be a family of functions on E.

(i) $\mathcal{A}$ *separates points* on E if for every distinct $x_1, x_2 \in E$ there is $f \in \mathcal{A}$ with $f(x_1) \neq f(x_2)$.
(ii) $\mathcal{A}$ *vanishes at no point* of E if for every $x \in E$ there is $g \in \mathcal{A}$ with $g(x) \neq 0$.
(iii) A complex algebra $\mathcal{A}$ is *self-adjoint* if $f \in \mathcal{A}$ implies $\overline{f} \in \mathcal{A}$, where $\overline{f}(x) := \overline{f(x)}$.

Theorem 3.4.20. *Two-Point Interpolation in a Separating Algebra*
If an algebra $\mathcal{A}$ on E separates points and vanishes nowhere, then for distinct $x_1, x_2 \in E$ and scalars c_1, c_2, there is $f \in \mathcal{A}$ with $f(x_i) = c_i$, $(i = 1, 2)$.

Proof. Choose $g, h, k \in \mathcal{A}$ with $g(x_1) \neq g(x_2)$, $h(x_1) \neq 0$, and $k(x_2) \neq 0$.
Let $u := gk - g(x_1)k$ and $v := gh - g(x_2)h$. Then $u(x_1) = v(x_2) = 0$ and $u(x_2)v(x_1) \neq 0$.
Hence $f := \frac{c_1 v}{v(x_1)} + \frac{c_2 u}{u(x_2)}$ lies in $\mathcal{A}$ and satisfies $f(x_i) = c_i$.
Structural Commentary: Separation and nonvanishing build two functions that isolate the two points.
Key Insight: Once two point-isolating functions are in the algebra, interpolation is linear algebra. □

Perspective 3.4.21. *Interpolation Is Linear Algebra*
The two-point interpolation step can also be written as a 2×2 linear system. Its invertibility is exactly the finite-dimensional criterion from Theorem 1.3.49, with the determinant test supplied by Theorem 4.1.8. The proof above chooses explicit point-isolating functions instead of reopening the calculation.

Theorem 3.4.22. *Stone–Weierstrass Theorem, Real Form*
$\mathcal{A} \subset C(K, \mathbb{R})$ is algebra on compact K that separates points, vanishes nowhere, then uniform closure is $C(K, \mathbb{R})$.

Proof. Let $\mathcal{B}$ be the uniform closure of $\mathcal{A}$; by Theorem 3.4.18, $\mathcal{B}$ is a uniformly closed algebra.
If $f \in \mathcal{B}$, then $|f| \in \mathcal{B}$: for $a = \|f\|_\infty$, approximate $|y|$ on $[-a, a]$ by polynomials p_n with $p_n(0) = 0$ using Corollary 3.4.16; then $p_n(f) \in \mathcal{B}$ and $p_n(f) \rightrightarrows |f|$.
Hence $\max(f, g), \min(f, g) \in \mathcal{B}$ whenever $f, g \in \mathcal{B}$, as $\max(f, g) = \frac{f+g}{2} + \frac{|f-g|}{2}$; $\min(f, g) = \frac{f+g}{2} - \frac{|f-g|}{2}$.
Fix $f \in C(K, \mathbb{R})$, $x \in K$, $\varepsilon > 0$. $\forall y \in K$, Theorem 3.4.20 gives $h_y \in \mathcal{B}$, $h_y(x) = f(x)$, $h_y(y) = f(y)$.
Compactness gives $y_1, \ldots, y_r$ s.t. $g_x := \max(h_{y_1}, \ldots, h_{y_r})$ satisfies $g_x(x) = f(x)$ and $g_x > f - \varepsilon$ on K.
Compactness gives $x_1, \ldots, x_m$ such that $h := \min(g_{x_1}, \ldots, g_{x_m})$ satisfies $|h - f| < \varepsilon$ on K.
Thus every $f \in C(K, \mathbb{R})$ is a uniform limit of members of $\mathcal{B}$, so $f \in \mathcal{B}$.
Structural Commentary: Local supports plus finite max/min give global approximation.
Key Insight: Separation plus compactness gives uniform approximation. □

Theorem 3.4.23. *Stone–Weierstrass Theorem, Complex Form*
If $\mathcal{A} \subset C_b(K)$ is self-adjoint, separates points, and vanishes nowhere on compact K, then $\overline{\mathcal{A}} = C_b(K)$.

Proof. Let $\mathcal{A}_\mathbb{R} := \{f \in \mathcal{A} : f \text{ real-valued}\}$.
If $f = u + iv \in \mathcal{A}$, self-adjointness gives $u = \frac{f+\overline{f}}{2} \in \mathcal{A}_\mathbb{R}$ and $v = \frac{f-\overline{f}}{2i} \in \mathcal{A}_\mathbb{R}$.
Thus $\mathcal{A}_\mathbb{R}$ separates points; rescaling a nonzero value and taking real parts shows it vanishes nowhere.
By Theorem 3.4.22, $\overline{\mathcal{A}_\mathbb{R}} = C(K, \mathbb{R})$. Approximate the real and imaginary parts separately, then combine.
Structural Commentary: Self-adjointness lets the complex algebra expose enough real functions.
Key Insight: The complex theorem is the real theorem applied to real and imaginary parts. □

Theorem 3.4.24. *Exponential and Logarithm from the Power Series*
Define $\mathsf{E}(z) := \sum_{n=0}^{\infty} \frac{z^n}{n!}$ for $z \in \mathbb{C}$.
Then, for $z, w \in \mathbb{C}$, $x \in \mathbb{R}$, $\mathsf{E}(z+w) = \mathsf{E}(z)\mathsf{E}(w)$, $\mathsf{E}(x) > 0$, $\mathsf{E}'(x) = \mathsf{E}(x)$, $\mathsf{E}(x) = e^x$, and E is strictly increasing on $\mathbb{R}$. Its inverse $\mathsf{L} : (0, \infty) \to \mathbb{R}$ satisfies $\mathsf{L}'(y) = \frac{1}{y}$, $\mathsf{L}(uv) = \mathsf{L}(u) + \mathsf{L}(v)$ for $u, v > 0$.
For $x > 0$ and $\alpha \in \mathbb{R}$, set $x^\alpha := \mathsf{E}(\alpha \mathsf{L}(x))$; then $(x^\alpha)' = \alpha x^{\alpha-1}$.

Proof. The ratio test (Theorem 3.1.20), gives convergence for all $z \in \mathbb{C}$.
Absolute convergence and Cauchy product give $\mathsf{E}(z)\mathsf{E}(w) = \sum_{n \geq 0} \frac{(z+w)^n}{n!} = \mathsf{E}(z + w)$, hence $\mathsf{E}(z)\mathsf{E}(-z) = 1$.
Theorem 3.4.8 gives $\mathsf{E}' = \mathsf{E}$ on $\mathbb{R}$. Thus $\mathsf{E}(x) > 0$ $\forall$ real x and $\mathsf{E}' > 0$, so E is strictly increasing.
Since $\mathsf{E}(1) = e$, the addition formula gives $\mathsf{E}(p) = e^p$ for $p \in \mathbb{Q}$; continuity extends this to $x \in \mathbb{R}$.
Let $\mathsf{L} = \mathsf{E}^{-1}$. Differentiating $\mathsf{L}(\mathsf{E}(x)) = x$ gives $\mathsf{L}'(y) = 1/y$.
If $u = \mathsf{E}(x)$ and $v = \mathsf{E}(y)$, then $\mathsf{L}(uv) = \mathsf{L}(\mathsf{E}(x + y)) = \mathsf{L}(u) + \mathsf{L}(v)$.
Finally, the chain rule, Theorem 4.1.6, gives $(x^\alpha)' = \alpha x^{\alpha-1}$.
Structural Commentary: Power series define E; inversion gives L and real powers.
Key Insight: The addition formula is the bridge from coefficients to the logarithm law. □

Perspective 3.4.25. *Real Powers by Rational Approximation*
One route defines x^α for $x > 0$ by choosing $q_n \in \mathbb{Q}$ with $q_n \to \alpha$ and setting $x^\alpha = \lim_{n\to\infty} x^{q_n}$ after proving existence and independence of (q_n). The power-series construction sets $x^\alpha = \mathsf{E}(\alpha \mathsf{L}(x))$; both routes must agree with rational powers and be choice-independent.

Perspective 3.4.26. *Logarithmic Integral Route*
With the Riemann integral, set $L(y) := \int_1^y \frac{1}{t}\,dt$ for $y > 0$. Theorem 4.1.57 gives $L'(y) = 1/y$; endpoint limits make L invertible, and Theorem 4.1.6 then recovers $E' = E$ and the product laws.

Perspective 3.4.27. *From Exponentials to Factorials*
Theorems 3.4.24, 4.1.59 supply real-power laws and integration-by-parts pattern used for Γ in Theorem 3.4.30.

Theorem 3.4.28. *Trigonometric Functions from the Exponential*
For real x, define $\mathsf{C}(x) := \frac{\mathsf{E}(ix)+\mathsf{E}(-ix)}{2}$ and $\mathsf{S}(x) := \frac{\mathsf{E}(ix)-\mathsf{E}(-ix)}{2i}$.
Then $\mathsf{E}(ix) = \mathsf{C}(x) + i\mathsf{S}(x)$, $\mathsf{C}^2 + \mathsf{S}^2 = 1$; $\mathsf{C}' = -\mathsf{S}$, and $\mathsf{S}' = \mathsf{C}$.
If $x_0 = \min\{x > 0 : \mathsf{C}(x) = 0\}$ and $\pi := 2x_0$, then $\mathsf{E}(\pi i) = -1$, $\mathsf{E}(2\pi i) = 1$, $\mathsf{E}(z + 2\pi i) = \mathsf{E}(z)$ for $z \in \mathbb{C}$.
Thus C, S are 2π-periodic, and $\gamma(t) := \mathsf{E}(it)$ traces the unit circle once on $[0, 2\pi]$.

Proof. The definitions give $\mathsf{E}(ix) = \mathsf{C} + i\mathsf{S}$. Since $\mathsf{E}(ix)\mathsf{E}(-ix) = 1$, $|\mathsf{E}(ix)| = 1$, hence $\mathsf{C}^2 + \mathsf{S}^2 = 1$.
Differentiating the definitions and using Theorem 3.4.8 gives $\mathsf{C}' = -\mathsf{S}$ and $\mathsf{S}' = \mathsf{C}$.
If C had no positive zero, then $\mathsf{C} > 0$ and $\mathsf{S}' > 0$ on $(0, \infty)$.
Thus, for $0 < x < y$, $\mathsf{S}(x)(y - x) < \int_x^y \mathsf{S}(t)\,dt = \mathsf{C}(x) - \mathsf{C}(y) \le 2$, impossible for large y. Hence x_0 exists.
On $(0, \pi/2)$, $\mathsf{C} > 0$, so S increases from 0 to $\mathsf{S}(\pi/2) = 1$; thus $\mathsf{E}(\pi i/2) = i$, and the addition formula gives $\mathsf{E}(\pi i) = -1$, $\mathsf{E}(2\pi i) = 1$, and the period.
Finally, $|\mathsf{E}(it)| = 1$ and the source quadrant argument give the one-time traversal of the unit circle.
Structural Commentary: Trigonometry comes from the path $\mathsf{E}(it)$.
Key Insight: The identity $\mathsf{E}(ix)\mathsf{E}(-ix) = 1$ puts the exponential path on the unit circle. □

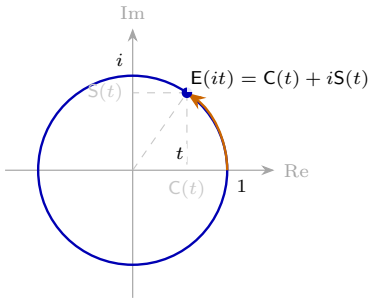


Figure 164: $\mathsf{E}(it) = \mathsf{C}(t) + i\mathsf{S}(t)$ traces the unit circle, defining sine and cosine analytically.

Theorem 3.4.29. *Algebraic Completeness of $\mathbb{C}$*
Every nonconstant complex polynomial has a complex root.

Proof. Let $P(z) = \sum_{j=0}^n a_j z^j$, $n \ge 1$, and normalize $a_n = 1$.
Since $|P(z)| \to \infty$, Theorem 2.3.35 gives z_0 with $\mu := |P(z_0)| = \min_{\mathbb{C}} |P|$. If $\mu = 0$, done.
Assume $\mu > 0$ and set $Q(w) := P(z_0 + w)/P(z_0)$. Then $Q(0) = 1$ and $|Q(w)| \ge 1$.
Write $Q(w) = 1 + \sum_{j=k}^{n} b_j w^j$, where $b_k \ne 0$ and k is minimal. By Theorem 3.4.28, choose θ with $b_k e^{ik\theta} = -|b_k|$.
Then $Q(re^{i\theta}) = 1 - |b_k| r^k + r^{k+1} H(r)$, with H bounded near 0, so $\left|Q(re^{i\theta})\right| < 1$ for small $r > 0$, contradiction.
Structural Commentary: Minimize, translate, then perturb.
Key Insight: Choose the phase that lowers the first nonzero term. □

Theorem 3.4.30. *Gamma Function, Beta Function, and Stirling Skeleton*
For $x > 0$, $\Gamma(x + 1) = x\Gamma(x)$; $\Gamma(n + 1) = n!$ for $n \in \mathbb{N}$; $\log \Gamma$ is convex on $(0, \infty)$.
Moreover, Γ is the unique $f : (0, \infty) \to (0, \infty)$ with $f(1) = 1$, $f(x + 1) = xf(x)$, and $\log f$ convex.
For $x, y > 0$, $\mathrm{B}(x, y) := \int_0^1 t^{x-1}(1 - t)^{y-1}\,dt = \frac{\Gamma(x)\Gamma(y)}{\Gamma(x+y)}$.
Consequently, $\Gamma\left(\frac{1}{2}\right) = \sqrt{\pi}$; $\int_{-\infty}^{\infty} e^{-s^2}\,ds = \sqrt{\pi}$; $\Gamma(x) = \frac{2^{x-1}}{\sqrt{\pi}}\Gamma\left(\frac{x}{2}\right)\Gamma\left(\frac{x+1}{2}\right)$; and $\lim_{x\to\infty} \frac{\Gamma(x+1)}{(x/e)^x\sqrt{2\pi x}} = 1$.

Proof. Integration by parts gives $\Gamma(x + 1) = x\Gamma(x)$; $\Gamma(1) = 1$ gives $\Gamma(n + 1) = n!$.
Holder's inequality gives log-convexity. For uniqueness, let $\varphi = \log f$.
If $0 < x < 1$, convexity and $\varphi(x + 1) = \varphi(x) + \log x$ give $0 \le \varphi(x) - \log \frac{n!\,n^x}{x(x+1)\cdots(x+n)} \le x \log\left(1 + \frac{1}{n}\right) \to 0$.
Thus f is uniquely determined on $(0, 1)$, hence on $(0, \infty)$.
For fixed $y > 0$, $\mathrm{B}(1, y) = 1/y$, $\mathrm{B}(x + 1, y) = \frac{x}{x+y}\mathrm{B}(x, y)$, and Holder gives log-convexity in x.
Hence $F_y(x) := \frac{\Gamma(x+y)}{\Gamma(y)}\mathrm{B}(x, y)$ satisfies the three uniqueness conditions, so $F_y = \Gamma$.
The consequences follow from $t = \sin^2\theta$, $t = s^2$, the duplication identity, and the substitution $t = x(1 + u)$ in

the asymptotic proof of Stirling's formula.
Structural Commentary: Functional equation plus convexity gives uniqueness.
Key Insight: Bohr–Mollerup uniqueness turns beta-function identities into gamma identities. □

Method 3.4.31. *Power-Series and Approximation Proof Patterns*

(i) *Power series*: work on $K_\rho \Subset I_R$; use uniform convergence before termwise limits or derivatives.
(ii) *Endpoint theorems*: prove identities inside the radius and pass to endpoints with Abel-type control.
(iii) *Approximate identities*: use kernels $Q_n \geq 0$, $\int Q_n = 1$, with tails $\to 0$ off 0.
(iv) *Uniform closure*: show $a_n \in \mathcal{A}$ and $a_n \rightrightarrows f$.
(v) *Lattice closure*: derive $|f|$, then $\max, \min$ from algebraic formulas.
(vi) *Stone–Weierstrass*: interpolate locally; use compactness and finite max/min to globalize.
(vii) *Gamma function*: use $\Gamma(x+1) = x\Gamma(x)$, $\Gamma(1) = 1$, and log-convexity for uniqueness.

Key Insight: Uniform convergence and compactness are what make infinite algebra behave like finite algebra.

3.5 Fourier Structure

Fourier structure studies trigonometric series through orthogonality, projection, kernels, and convergence, with inner-product geometry reappearing throughout.

Definition 3.5.1. *Core Fourier Objects*

(i) A *trigonometric polynomial* is a finite sum $a_0 + \sum_{n=1}^{N} (a_n \cos nx + b_n \sin nx)$, equivalently $\sum_{n=-N}^{N} \alpha_n e^{inx}$.
(ii) If f is Riemann integrable on $[-\pi, \pi]$, its *Fourier coefficients* are $\widehat{f}(m) := \frac{1}{2\pi} \int_{-\pi}^{\pi} f(x) e^{-imx}\, dx$, $m \in \mathbb{Z}$, and $f(x) \sim_{\text{coef}} \sum_{m=-\infty}^{\infty} \widehat{f}(m) e^{imx}$ records coefficients, not convergence.
(iii) A sequence (ϕ_n) on $[a, b]$ is *orthonormal* if $\int_a^b \phi_n(x)\overline{\phi_m(x)}\, dx = 0 \quad (n \neq m)$ and $\int_a^b |\phi_n(x)|^2\, dx = 1$. Its Fourier coefficients are $c_n := \int_a^b f(t)\overline{\phi_n(t)}\, dt$.

Key Insight: Fourier coefficients are orthonormal coordinates; convergence is a separate analytic question.

Definition 3.5.2. *Period-One Fourier Normalization*
Let $C(\mathbb{R}/\mathbb{Z}; \mathbb{C})$ denote the continuous functions $f : \mathbb{R} \to \mathbb{C}$ satisfying $f(x+1) = f(x)$ for all x.
For $f, g \in C(\mathbb{R}/\mathbb{Z}; \mathbb{C})$, set $\langle f, g \rangle_{\mathbb{T}} := \int_0^1 f(x)\overline{g(x)}\, dx$; $\|f\|_{L^2(\mathbb{T})} := \sqrt{\langle f, f \rangle_{\mathbb{T}}}$; and $e_n(x) := e^{2\pi inx}$ for $n \in \mathbb{Z}$.
The *period-one Fourier coefficient* is $\widehat{f}_{\mathbb{T}}(n) := \langle f, e_n \rangle_{\mathbb{T}} = \int_0^1 f(x) e^{-2\pi inx}\, dx$.
A period-one trigonometric polynomial is a finite sum $\sum_{n=-N}^{N} c_n e_n$.

Perspective 3.5.3. *Two Equivalent Fourier Clocks*
The period-one and 2π-periodic conventions differ by $t = 2\pi x$; only normalizing constants change.
Key Insight: Changing Fourier clocks changes constants, not the projection geometry.

Proposition 3.5.4. *Characters Are an Orthonormal System*
For $m, n \in \mathbb{Z}$, $\langle e_n, e_m \rangle_{\mathbb{T}} = 1$ if $n = m$ and 0 otherwise.
Hence, if $P = \sum_{n=-N}^{N} c_n e_n$, then $c_n = \langle P, e_n \rangle_{\mathbb{T}}$ and $\|P\|^2_{L^2(\mathbb{T})} = \sum_{n=-N}^{N} |c_n|^2$.

Proof. If $n = m$, the integrand is 1. If $n \neq m$, then $\int_0^1 e^{2\pi i(n-m)x}\, dx = 0$.
The coefficient and norm identities are the finite orthonormal-coordinate identities from Theorems 2.1.48 and 2.1.54, applied inside the finite-dimensional span of $e_{-N}, \ldots, e_N$.
Structural Commentary: Fourier coefficients are orthonormal coordinates in a finite trigonometric subspace.
Key Insight: The exponential characters make cross terms vanish. □

Definition 3.5.5. *Periodic Convolution and Fejer Kernels*
For $f, g \in C(\mathbb{R}/\mathbb{Z}; \mathbb{C})$, define their *periodic convolution* as $(f * g)(x) := \int_0^1 f(y) g(x-y)\, dy$.
For $N \geq 1$, the *Fejer kernel* is $F_N(x) := \sum_{n=-N}^{N} \left(1 - \frac{|n|}{N}\right) e_n(x) = \frac{1}{N} \left| \sum_{n=0}^{N-1} e_n(x) \right|^2$.

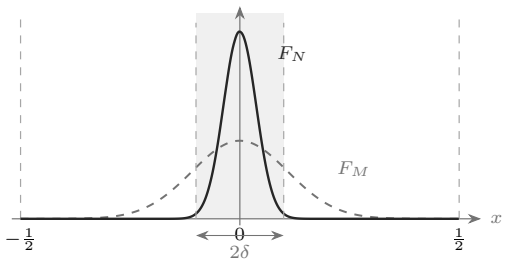


Figure 165: The Fejer kernels are nonnegative trigonometric polynomial approximate identities on $\mathbb{R}/\mathbb{Z}$.

Theorem 3.5.6. *Fejer Approximation on $\mathbb{R}/\mathbb{Z}$*
For $N \geq 1$, F_N is a nonnegative trigonometric polynomial with $\int_0^1 F_N = 1$. For each $0 < \delta < 1/2$, $F_N \to 0$ uniformly on $[\delta, 1-\delta]$. If $f \in C(\mathbb{R}/\mathbb{Z}; \mathbb{C})$, then $f * F_N$ is a trigonometric polynomial and $f * F_N \rightrightarrows f$.

Proof. From Definition 3.5.5, F_N is a finite linear combination of e_n and $F_N = \frac{1}{N}\left|\sum_{n=0}^{N-1} e_n\right|^2 \geq 0$.
Since $\int_0^1 e_n = 0$ for $n \neq 0$, the e_0 coefficient gives $\int_0^1 F_N = 1$.
For $x \notin \mathbb{Z}$, $F_N(x) = \frac{\sin^2(\pi N x)}{N \sin^2(\pi x)}$. If $m_\delta := \min_{\delta \leq x \leq 1-\delta} \sin^2(\pi x) > 0$, then $0 \leq F_N(x) \leq (N m_\delta)^{-1}$ on $[\delta, 1-\delta]$.
Also $f * e_n = \widehat{f}_{\mathbb{T}}(n) e_n$, so $f * F_N$ is a trigonometric polynomial. For uniform convergence, write $(f * F_N)(x) = \int_0^1 f(x-t) F_N(t)\, dt$. Let $M = \|f\|_\infty$ and choose δ so $|f(x-t) - f(x)| < \varepsilon$ whenever $t \in [0, \delta] \cup [1-\delta, 1]$.
By $F_N \geq 0$ and $\int_0^1 F_N = 1$, uniformly in x, $|(f * F_N)(x) - f(x)| \leq \varepsilon + 2M \sup_{\delta \leq t \leq 1-\delta} F_N(t) \to \varepsilon$.
Structural Commentary: Fejer kernels concentrate mass near the identity element of $\mathbb{R}/\mathbb{Z}$.
Key Insight: Positivity plus concentration turns convolution into uniform approximation. □

Perspective 3.5.7. *Fourier and Plancherel Roadmap*
Path: orthonormal characters, Fejer density, best L^2 projection, then Parseval/Plancherel.
Key Insight: Fourier theory is finite-dimensional orthogonal projection plus an approximation theorem.

Perspective 3.5.8. *Fourier Coefficients as Projections*
Fourier coefficients are projections onto orthonormal directions, modeled finitely by Theorems 2.2.4, 2.1.48, and 2.1.52; infinite directions make convergence separate.

Theorem 3.5.9. *Fourier Coefficients Give Best Mean-Square Approximation*
Let $\langle u, v \rangle := \int_a^b u\overline{v}$ and $\|u\|_2^2 := \langle u, u \rangle$.
If (ϕ_n) is orthonormal, $c_m = \langle f, \phi_m \rangle$, $s_n = \sum_{m=1}^{n} c_m \phi_m$, and $t_n = \sum_{m=1}^{n} \gamma_m \phi_m$, then $\|f - s_n\|_2 \leq \|f - t_n\|_2$, with equality iff $\gamma_m = c_m$ for $m = 1, \ldots, n$. In particular, $\sum_{m=1}^{\infty} |c_m|^2 \leq \|f\|_2^2$ for $c_m \to 0$.

Proof. For sums over $m = 1, \ldots, n$, orthonormality gives
$\|f - t_n\|_2^2 = \|f\|_2^2 - 2\,\mathrm{Re}\sum c_m \overline{\gamma_m} + \sum |\gamma_m|^2 = \|f\|_2^2 - \sum |c_m|^2 + \sum |\gamma_m - c_m|^2$.
Thus the minimum occurs exactly when $\gamma_m = c_m$.
Taking $t_n = s_n$ gives $\sum_{m=1}^n |c_m|^2 = \|s_n\|_2^2 \leq \|f\|_2^2$; let $n \to \infty$ to get Bessel and $c_m \to 0$.
Structural Commentary: Orthogonality makes the error split into independent coefficient errors.
Key Insight: Fourier coefficients are the coordinates of the orthogonal projection. □

Theorem 3.5.10. *Dirichlet Kernel and Local Fourier Convergence*
Let f be Riemann integrable and 2π-periodic, with $\widehat{f}(m) := \frac{1}{2\pi}\int_{-\pi}^{\pi} f(t) e^{-imt}\, dt$.
Let $s_N(f; x) := \sum_{m=-N}^{N} \widehat{f}(m) e^{imx}$ and $D_N(t) := \sum_{m=-N}^{N} e^{imt} = \frac{\sin((N+\frac{1}{2})t)}{\sin(t/2)}$.
Then $s_N(f; x) = \frac{1}{2\pi}\int_{-\pi}^{\pi} f(x-t) D_N(t)\, dt$.
If, for some x, $|f(x+t) - f(x)| \leq M|t|$ whenever $|t| < \delta$, then $s_N(f; x) \to f(x)$.
Hence $f = g$ near x implies $s_N(f; x) - s_N(g; x) \to 0$.

Proof. The displayed form of D_N is the finite geometric sum; substituting $\widehat{f}(m)$ gives the convolution formula.
Also $\frac{1}{2\pi}\int_{-\pi}^{\pi} D_N = 1$. Set $q_x(t) := \frac{f(x-t) - f(x)}{\sin(t/2)}$, for $0 < |t| \leq \pi$, $q_x(0) := 0$.
Then $s_N(f; x) - f(x) = \frac{1}{2\pi}\int_{-\pi}^{\pi} q_x(t) \sin\left(N + \frac{1}{2}\right) t\, dt$.
The hypothesis makes q_x bounded near 0; away from 0 it is Riemann integrable.
Since $\sin((N+\frac{1}{2})t) = \sin Nt \cos(t/2) + \cos Nt \sin(t/2)$, Theorem 3.5.9 sends the two Fourier coefficients to 0.
Apply the result to $f - g$ for localization.

Structural Commentary: The Dirichlet kernel rewrites Fourier partial sums as a local oscillatory integral.
Key Insight: Pointwise convergence under a local Lipschitz condition depends only on nearby values of f. □

Corollary 3.5.11. *Uniform Approximation by Trigonometric Polynomials*
If $f \in C(\mathbb{R})$ is 2π-periodic, then for every $\varepsilon > 0$ there is a trigonometric polynomial P such that $\|f - P\|_\infty < \varepsilon$.

Proof. Let $\mathbb{T} := \{z \in \mathbb{C} : |z| = 1\}$ and set $F(e^{ix}) := f(x)$.
The algebra $\mathcal{T} := \operatorname{span}\{z^n : n \in \mathbb{Z}\} \subset C(\mathbb{T})$ is self-adjoint, separates points, and vanishes nowhere.
By Theorem 3.4.23, $\overline{\mathcal{T}} = C(\mathbb{T})$. Pulling approximants back by $z = e^{ix}$ gives the desired P.
Structural Commentary: Periodic approximation is Stone–Weierstrass on the unit circle.
Key Insight: The functions e^{inx} are the algebraic coordinates on the circle. □

Theorem 3.5.12. *Parseval's Theorem for Trigonometric Fourier Series*
Let f, g be Riemann integrable and 2π-periodic. Let $\langle u, v\rangle_2 := \frac{1}{2\pi}\int_{-\pi}^{\pi} u(x)\overline{v(x)}\,dx$ and $\|u\|_2^2 := \langle u, u\rangle_2$.
If $f \sim_{\text{coef}} \sum_{n\in\mathbb{Z}} \widehat{f}(n)e^{inx}$, and $g \sim_{\text{coef}} \sum_{n\in\mathbb{Z}} \widehat{g}(n)e^{inx}$,
then $\|f - s_N(f)\|_2 \to 0$; $\langle f, g\rangle_2 = \sum_{n\in\mathbb{Z}} \widehat{f}(n)\overline{\widehat{g}(n)}$; and $\|f\|_2^2 = \sum_{n\in\mathbb{Z}} \left|\widehat{f}(n)\right|^2$.

Proof. If f is continuous, Corollary 3.5.11 gives trigonometric polynomials P with $\|f - P\|_2$ arbitrarily small.
If P has degree N_0, then Theorem 3.5.9 gives $\|f - s_N(f)\|_2 \le \|f - P\|_2$ for $N \ge N_0$.
For Riemann integrable f, approximate f in $\|\cdot\|_2$ by a continuous 2π-periodic h.
Then $\|f - s_N(f)\|_2 \le \|f - h\|_2 + \|h - s_N(h)\|_2 + \|s_N(h - f)\|_2$ and $\|s_N(h - f)\|_2 \le \|h - f\|_2$ (Theorem 3.5.9).
Hence $s_N(f) \to f$ in $\|\cdot\|_2$. Finally, $\langle s_N(f), g\rangle_2 = \sum_{|n|\le N} \widehat{f}(n)\overline{\widehat{g}(n)}$.
Let $N \to \infty$ and use Schwarz' inequality; the last identity is $g = f$.
Structural Commentary: Density upgrades Bessel to Parseval.
Key Insight: Bessel gives one inequality; density of trigonometric polynomials gives equality. □

Method 3.5.13. *Fourier Proof Patterns*

(i) *Characters*: use orthogonality of e_n to compute coefficients and finite-dimensional projections.
(ii) *Kernels*: rewrite partial sums or averages as convolutions with Dirichlet or Fejer kernels.
(iii) *Approximation*: use Fejer positivity or Stone–Weierstrass to approximate continuous periodic functions by trigonometric polynomials.
(iv) *L^2 control*: project onto finite orthonormal spans, then pass to limits by density.
(v) *Localization*: express the error as an oscillatory integral and use local regularity near the point.

Key Insight: Fourier theory is finite-dimensional orthogonal projection plus an approximation theorem.

4 Differential and Geometric Analysis

4.1 One-Variable Differentiation and Integration

One-variable calculus studies functions on intervals through derivatives, mean-value theorems, Riemann integration, and the FTC, then points toward vector-valued estimates and linear approximation.

Definition 4.1.1. *Derivative of a Real Function*
Let f be real-valued on $[a, b]$ and let $x \in [a, b]$. For $a < t < b$ and $t \ne x$, set $\phi_x(t) := \frac{f(t)-f(x)}{t-x}$.
If $\lim_{t\to x} \phi_x(t)$ exists, denote *derivative* by $f'(x)$; then f is *differentiable at x*.
If this holds for every $x \in E \subseteq [a, b]$, then f is *differentiable on E*.



Figure 166: Domain matters: isolated points give no punctured approach; $[x_0, \infty)$ gives a right-hand one.

Perspective 4.1.2. *Derivative Depends on the Domain*
For $f : X \to \mathbb{R}$, $f'(x_0)$ is the limit of $\frac{f(x)-f(x_0)}{x-x_0}$ as $x \to x_0$ in $X \setminus \{x_0\}$; see Definition 2.3.4. Hence x_0 must be a limit point of X: isolated points give no derivative, while half-interval domains give one-sided derivatives.

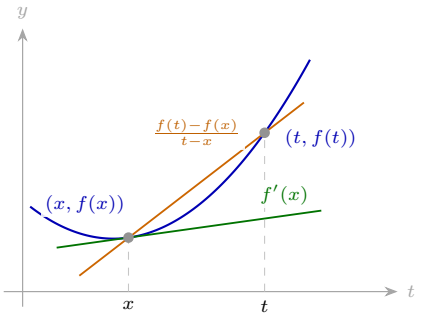


Figure 167: The derivative is the limiting slope of secant lines through $(x, f(x))$.

Proposition 4.1.3. *First-Order Approximation Form*
Let $X \subseteq \mathbb{R}$, let $x_0 \in X$ be a limit point, let $f : X \to \mathbb{R}$, and let $L \in \mathbb{R}$. Then $f'(x_0) = L$ relative to X iff $\forall \varepsilon > 0$ $\exists \delta > 0$ such that $|f(x) - f(x_0) - L(x - x_0)| \leq \varepsilon |x - x_0|$ whenever $x \in X$ and $|x - x_0| < \delta$.

Proof. For $x \neq x_0$, $\frac{|f(x)-f(x_0)-L(x-x_0)|}{|x-x_0|} = \left|\frac{f(x)-f(x_0)}{x-x_0} - L\right|$, while $x = x_0$ makes approx. estimate automatic.

($\Rightarrow$) If $f'(x_0) = L$, then the displayed right-hand side is at most ε for all $x \in X \setminus \{x_0\}$ sufficiently near x_0. Multiplying by $|x - x_0|$ gives the approximation estimate.

($\Leftarrow$) If the approximation estimate holds, then division by $|x - x_0|$ gives $\left|\frac{f(x)-f(x_0)}{x-x_0} - L\right| \leq \varepsilon$ for all $x \in X \setminus \{x_0\}$ sufficiently near x_0. Hence $f'(x_0) = L$.

Structural Commentary: Differentiability is first-order linear approximation.
Key Insight: Multiply or divide by $|x - x_0|$. □

Theorem 4.1.4. *Differentiability Implies Continuity*
If f is differentiable at x, then f is continuous at x.

Proof. For $t \neq x$, $f(t) - f(x) = \frac{f(t)-f(x)}{t-x}(t - x)$. As $t \to x$, the difference quotient tends to $f'(x)$ and $t - x \to 0$. By Theorem 2.3.9, $f(t) - f(x) \to 0$, so $f(t) \to f(x)$. Theorem 2.3.12 gives continuity at x.
Structural Commentary: A differentiable function has first-order control near the point.
Key Insight: The quotient stays near a finite number while $t - x$ shrinks to zero. □

Theorem 4.1.5. *Algebra Rules for Derivatives*
If f and g are differentiable at x, then so are $f+g$ and fg; if $g(x) \neq 0$, then f/g is differentiable at x. Moreover,

$$(f + g)'(x) = f'(x) + g'(x),$$
$$(fg)'(x) = f'(x)g(x) + f(x)g'(x),$$
$$\left(\frac{f}{g}\right)'(x) = \frac{g(x)f'(x) - g'(x)f(x)}{g(x)^2}.$$

Proof. The sum rule is limit algebra. For $h = fg$, $h(t) - h(x) = f(t)\big(g(t) - g(x)\big) + g(x)\big(f(t) - f(x)\big)$. Dividing by $t - x$ and letting $t \to x$ gives the product rule, using Theorem 4.1.4 for $f(t) \to f(x)$.
For $h = f/g$, continuity of g makes $g(t) \neq 0$ near x, and $\frac{h(t)-h(x)}{t-x} = \frac{1}{g(t)g(x)}\left[g(x)\frac{f(t)-f(x)}{t-x} - f(x)\frac{g(t)-g(x)}{t-x}\right]$. Letting $t \to x$ gives the quotient rule.
Structural Commentary: Product and quotient rules are difference-quotient identities plus continuity.
Key Insight: Difference-quotient identities plus continuity give the algebra rules. □

Theorem 4.1.6. *Chain Rule*
Let f be continuous on $[a, b]$ with $f'(x)$ existing at $x \in [a, b]$, and let g be defined on an interval containing $f([a, b])$ and differentiable at $f(x)$. If $h(t) = g(f(t))$, then $h'(x) = g'(f(x))f'(x)$.

Proof. Let $y = f(x)$. For some $u(t) \to 0$ and $v(s) \to 0$,

$$f(t) - f(x) = (t - x)\big(f'(x) + u(t)\big),$$
$$g(s) - g(y) = (s - y)\big(g'(y) + v(s)\big).$$

With $s = f(t)$, continuity gives $s \to y$, and for $t \neq x$, $\frac{h(t)-h(x)}{t-x} = \big(g'(y) + v(f(t))\big)\big(f'(x) + u(t)\big) \to g'(y)f'(x)$.
Structural Commentary: Each differentiability statement is a linear approximation plus a vanishing error.
Key Insight: Composition multiplies the two first-order parts; the error terms disappear. □

Example 4.1.7. *Oscillation at a Point*
Assuming the usual derivative of $\sin x$, define

$$f_1(x) = \begin{cases} x\sin(1/x), & x \neq 0, \\ 0, & x = 0, \end{cases} \qquad f_2(x) = \begin{cases} x^2\sin(1/x), & x \neq 0, \\ 0, & x = 0. \end{cases}$$

For f_1, the quotient at 0 is $\sin(1/t)$, so $f_1'(0)$ does not exist. For f_2, $|(f_2(t) - f_2(0))/t| \leq |t|$, so $f_2'(0) = 0$. However, for $x \neq 0$, $f_2'(x) = 2x\sin(1/x) - \cos(1/x)$, so f_2' is not continuous at 0.

Definition 4.1.8. *Local Extrema*
Let f be a real-valued function on a metric space X. The function has a *local maximum* at $p \in X$ if there is $\delta > 0$ such that $f(q) \leq f(p)$ whenever $d(p, q) < \delta$. Local minima are defined by reversing the inequality.

Remark 4.1.9. *Local and Global Extrema*
Global extrema are local, but local extrema depend on domain: isolated points are automatic extrema. Derivative tests therefore need an interior interval point; compactness supplies global extrema (Theorem 2.3.35).

Theorem 4.1.10. *Interior Extremum Has Zero Derivative*
If f has a local extremum at $x \in (a, b)$ and $f'(x)$ exists, then $f'(x) = 0$.

Proof. For a local maximum, $f(t) \leq f(x)$ near x.
Thus the quotient $(f(t) - f(x))/(t - x)$ is nonnegative for $t < x$ and nonpositive for $t > x$. Since the two-sided limit exists, $f'(x) \geq 0$ and $f'(x) \leq 0$, so $f'(x) = 0$. For a local minimum, apply the local-maximum case to $-f$. Since $(-f)'(x) = -f'(x)$, we get $-f'(x) = 0$, hence $f'(x) = 0$.
Structural Commentary: Interior extrema force opposite one-sided quotient signs.
Key Insight: A two-sided derivative compatible with both signs must be zero. □

Theorem 4.1.11. *Rolle's Theorem*
If $f : [a, b] \to \mathbb{R}$ continuous on $[a, b]$, differentiable on (a, b), and $f(a) = f(b)$, then $\exists x \in (a, b)$ s.t. $f'(x) = 0$.

Proof. If f is constant, any $x \in (a, b)$ works. Otherwise, by Theorem 2.3.35, f attains a nonendpoint extremum at some $x \in (a, b)$. Theorem 4.1.10 gives $f'(x) = 0$.
Structural Commentary: Equal endpoints force any nonconstant height change to turn around inside.
Key Insight: Compactness gives the extremum; differentiability makes the interior extremum stationary. □

Theorem 4.1.12. *Cauchy Mean Value Theorem*
Let f, g be continuous on $[a, b]$, differentiable on (a, b). $\exists\, x \in (a, b)$, $\big(f(b) - f(a)\big)g'(x) = \big(g(b) - g(a)\big)f'(x)$.

Proof. Set $h(t) := \big(f(b) - f(a)\big)g(t) - \big(g(b) - g(a)\big)f(t)$.
Then $h(a) = h(b)$, h is continuous on $[a, b]$, and h is differentiable on (a, b). By Theorem 4.1.11, $h'(x) = 0$ for some $x \in (a, b)$. Since $h'(x) = \big(f(b) - f(a)\big)g'(x) - \big(g(b) - g(a)\big)f'(x)$, the identity follows.
Structural Commentary: Build one auxiliary function whose endpoint values agree.
Key Insight: Equal endpoints force an interior horizontal tangent. □

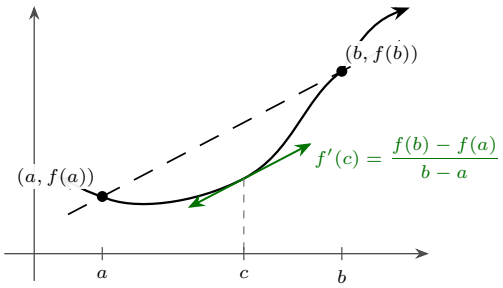


Figure 168: The Mean Value Theorem finds an interior tangent slope equal to the endpoint secant slope.

Corollary 4.1.13. *Mean Value Theorem*
If f is continuous on $[a, b]$, differentiable in (a, b), then $\exists\, x \in (a, b)$ such that $f(b) - f(a) = (b - a)f'(x)$.

Proof. Let $g(t) = t$. Then $g'(x) = 1$, $g(b) - g(a) = b - a$.
Hence Theorem 4.1.12 gives $f(b) - f(a) = (b - a)f'(x)$ for some $x \in (a, b)$.
Structural Commentary: This is Cauchy's theorem with $g(t) = t$.
Key Insight: A secant slope is achieved as a tangent slope somewhere inside. □

Corollary 4.1.14. *Bounded Derivative Gives Lipschitz Control*
If $f : [a, b] \to \mathbb{R}$ is continuous on $[a, b]$, differentiable on (a, b), and $|f'| \leq M$ on (a, b), then f is M-Lipschitz on $[a, b]$, hence uniformly continuous.

Proof. For $x < y$, Corollary 4.1.13 gives $f(y) - f(x) = (y - x)f'(c)$ for some $c \in (x, y)$, so $|f(y) - f(x)| \leq M(y - x)$. Symmetry handles $x > y$, and $x = y$ is trivial.
Structural Commentary: A derivative bound controls every secant slope.
Key Insight: Local slope bounds become global Lipschitz bounds through the Mean Value Theorem. □

Theorem 4.1.15. *Derivative Sign and Monotonicity*
Suppose f is differentiable in (a, b).

(i) If $f'(x) \geq 0$ for all $x \in (a, b)$, then f is monotonically increasing.
(ii) If $f'(x) = 0$ for all $x \in (a, b)$, then f is constant.
(iii) If $f'(x) \leq 0$ for all $x \in (a, b)$, then f is monotonically decreasing.

Proof. For $x_1 < x_2$ in (a, b), Corollary 4.1.13 on $[x_1, x_2]$ gives $f(x_2) - f(x_1) = (x_2 - x_1)f'(c)$ for some $c \in (x_1, x_2)$.

(i) If $f' \geq 0$, then $f(x_2) - f(x_1) \geq 0$.
(ii) If $f' = 0$, then $f(x_2) - f(x_1) = 0$.
(iii) If $f' \leq 0$, then $f(x_2) - f(x_1) \leq 0$.

Structural Commentary: Monotonicity is a global reading of the mean-value identity.
Key Insight: The sign of every tangent slope controls every secant slope. □

Theorem 4.1.16. *Derivatives Have the Intermediate Value Property*
If f is differentiable on $[a, b]$ in the one-sided endpoint sense and $f'(a) < \lambda < f'(b)$, Then $f'(x) = \lambda$ for some $x \in (a, b)$. The same conclusion holds when $f'(b) < \lambda < f'(a)$.

Proof. Set $g(t) = f(t) - \lambda t$. Since $g'(a) < 0 < g'(b)$, there are $t_1, t_2 \in (a, b)$ with $g(t_1) < g(a)$ and $g(t_2) < g(b)$. Hence the minimum of g on $[a, b]$ occurs at some $x \in (a, b)$. By Theorem 4.1.10, $0 = g'(x) = f'(x) - \lambda$. Apply the same argument to $-f$ for the reversed case.
Structural Commentary: A derivative need not be continuous, but it cannot jump over values.
Key Insight: Tilt the graph by subtracting λt and force an interior extremum. □

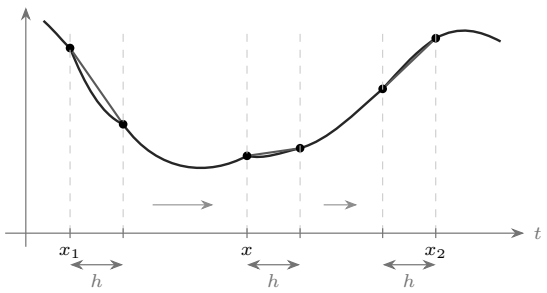


Figure 169: Sliding fixed-width secants in Darboux's theorem.

Perspective 4.1.17. *Sliding Secants*
For fixed small $h > 0$, the secant-slope function $\sigma_h(x) := \frac{f(x+h)-f(x)}{h}$ is continuous. Thus Theorem 2.3.42 finds intermediate secant slopes, and Corollary 4.1.13 converts them to derivative values.
Key Insight: Slide secants, then apply the Mean Value Theorem.

Corollary 4.1.18. *Derivatives Have No Simple Discontinuities*
If f is differentiable on $[a, b]$, then f' cannot have a simple discontinuity on $[a, b]$.

Proof. Suppose a simple discontinuity occurs at an interior point c; endpoint cases are one-sided. If $f'(c-) \neq f'(c+)$, choose λ strictly between these limits, $\lambda \neq f'(c)$. For $x < c < y$ close to c, Theorem 4.1.16 gives $z \in (x, y)$ with $f'(z) = \lambda$, contradicting the one-sided limits. If $f'(c-) = f'(c+) = L \neq f'(c)$, choose λ between L and $f'(c)$ and apply the same theorem on a sufficiently small one-sided interval with endpoint c.
Structural Commentary: Darboux behavior rules out jump discontinuities for derivatives.
Key Insight: Derivatives may oscillate, but they do not jump. □

Theorem 4.1.19. *Derivative of an Inverse Function*
Let $f : I \to \mathbb{R}$ be one-to-one on an open interval I, let $x_0 \in I$, and put $y_0 = f(x_0)$. Suppose f is differentiable at x_0, $f'(x_0) \neq 0$, and $g = f^{-1} : f(I) \to I$ is continuous at y_0. Then g is differentiable at y_0 and $g'(y_0) = \frac{1}{f'(x_0)}$.

Proof. For $y \in f(I)$ with $y \neq y_0$, write $x = g(y)$. Then $x \neq x_0$ and $\frac{g(y)-g(y_0)}{y-y_0} = \frac{x-x_0}{f(x)-f(x_0)} = \left(\frac{f(x)-f(x_0)}{x-x_0}\right)^{-1}$.
As $y \to y_0$, the assumed continuity of g gives $x = g(y) \to x_0$.
The difference quotient for f tends to $f'(x_0) \neq 0$, so the reciprocal tends to $1/f'(x_0)$ by Theorem 2.3.9.
Structural Commentary: Matched inverse quotients are reciprocal.
Key Insight: Continuity of the inverse converts the limit $y \to y_0$ back into $x \to x_0$. □

Corollary 4.1.20. *One-Dimensional Inverse Function Theorem*
If $f : (a,b) \to (c,d)$ is a differentiable bijection and $f' \neq 0$,
then f is a homeomorphism and f^{-1} is differentiable, with $(f^{-1})'(y) = \frac{1}{f'(f^{-1}(y))}$.

Proof. By Theorem 4.1.16, f' has constant sign. Corollary 4.1.13 makes f strictly monotone, hence its continuous interval inverse is continuous. Apply Theorem 4.1.19 pointwise.
Structural Commentary: Darboux gives constant sign; the inverse derivative formula is local.
Key Insight: Nonvanishing derivative is the one-dimensional obstruction to differentiating the inverse. □

Theorem 4.1.21. *L'Hospital's Rule*
Let f, g be differentiable on (a,b), $g' \neq 0$, and $-\infty \leq a < b \leq +\infty$.
If $\frac{f'}{g'} \to A \in \overline{\mathbb{R}}$ as $x \to a$, and either $f, g \to 0$ or $g \to +\infty$, then $\frac{f}{g} \to A$ as $x \to a$.
The same holds as $x \to b$, and for $g \to -\infty$.

Proof. For $a < x < y < b$, Theorem 4.1.12 gives $t \in (x,y)$ with $\frac{f(x)-f(y)}{g(x)-g(y)} = \frac{f'(t)}{g'(t)}$.
If $f, g \to 0$, let $x \to a$ with y fixed. If $g \to +\infty$, use $\frac{f(x)}{g(x)} = \frac{f(x)-f(y)}{g(x)-g(y)}\frac{g(x)-g(y)}{g(x)} + \frac{f(y)}{g(x)}$ and let $x \to a$.
In both cases, bounds on f'/g' near a transfer to f/g, forcing the limit.
Structural Commentary: CMVT transfers derivative bounds to quotient bounds.
Key Insight: Derivative quotients govern the original indeterminate quotient. □

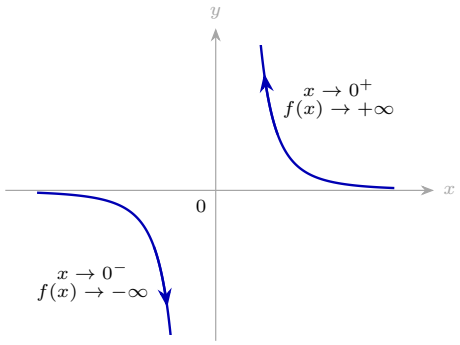


Figure 170: For $f(x) = x^{-3}$, the one-sided limits at 0 are opposite infinities.

Definition 4.1.22. *Higher Derivatives*
If f has derivative f' on an interval and f' is differentiable, then $(f')'$ is denoted f'' and called the *second derivative*. Continuing recursively gives $f^{(n)}$, the derivative of order n.

Definition 4.1.23. *Smoothness Classes*
On an interval, C^0 means continuous, C^r means $f^{(r)}$ exists and continuous, C^∞ means C^r for every finite r. Analytic functions (Definition 3.4.7) form the class C^ω.

Perspective 4.1.24. *Regularity Hierarchy*
The inclusions are $C^\omega \subseteq C^\infty \subseteq \cdots \subseteq C^2 \subseteq C^1 \subseteq C^0$.
The inclusions follow from Theorem 3.4.8 and Definition 4.1.23; they are generally proper.
Key Insight: Higher regularity is extra structure, not just a different notation for continuity.

Theorem 4.1.25. *Higher-Order Leibniz Product Rule*
If $f, g : (a,b) \to \mathbb{R}$ are r times differentiable, $r \geq 1$, then $(fg)^{(r)}(x) = \sum_{k=0}^{r} \binom{r}{k} f^{(k)}(x) g^{(r-k)}(x)$.

Proof. Induct on r. The case $r = 1$ is Theorem 4.1.5.
Differentiating the order-r formula gives $(fg)^{(r+1)} = \sum_{k=0}^{r} \binom{r}{k} f^{(k+1)} g^{(r-k)} + \sum_{k=0}^{r} \binom{r}{k} f^{(k)} g^{(r-k+1)}$.
Reindex and use Pascal's identity.
Structural Commentary: Each derivative chooses whether it lands on f or on g.
Key Insight: The binomial coefficients count those choices. □

Theorem 4.1.26. *Taylor's Theorem with Lagrange Remainder*
Let $f : [a,b] \to \mathbb{R}$, let $n \geq 1$, suppose $f^{(n-1)}$ is continuous on $[a,b]$, and suppose $f^{(n)}$ exists on (a,b).
For distinct $\alpha, \beta \in [a,b]$, set $P(t) := \sum_{k=0}^{n-1} \frac{f^{(k)}(\alpha)}{k!}(t-\alpha)^k$.
Then for some x between α and β, $f(\beta) = P(\beta) + \frac{f^{(n)}(x)}{n!}(\beta-\alpha)^n$.

Proof. Choose M so that $f(\beta) = P(\beta) + M(\beta-\alpha)^n$, and set $G(t) := f(t) - P(t) - M(t-\alpha)^n$.
Then $G(\alpha) = G'(\alpha) = \cdots = G^{(n-1)}(\alpha) = G(\beta) = 0$. Repeated applications of Corollary 4.1.13 give $G^{(n)}(x) = 0$ for some x between α and β. Since $G^{(n)}(x) = f^{(n)}(x) - n!M$, the formula follows.
Structural Commentary: The auxiliary error vanishes at β and to order $n-1$ at α.
Key Insight: Repeated Rolle-style arguments turn endpoint cancellation into a precise remainder. □

Example 4.1.27. *Smooth but Not Analytic*
Define $e(x) = \begin{cases} e^{-1/x}, & x > 0, \\ 0, & x \leq 0. \end{cases}$. By repeated L'Hospital, Theorem 4.1.21, $e \in C^\infty$ and $e^{(r)}(0) = 0$ for all r.
Its Taylor series at 0 is zero, but $e(x) > 0$ for $x > 0$, so e is not analytic at 0 (Definition 3.4.7).
Key Insight: A function can be infinitely tangent to the x-axis without being locally equal to its Taylor series.

Method 4.1.28. *Newton and Secant Iterations*
To solve $f(x) = 0$:

(i) Newton uses tangent-line zero: $x_n = x_{n-1} - \frac{f(x_{n-1})}{f'(x_{n-1})}$, $n \geq 1$ when $f'(x_{n-1}) \neq 0$.
(ii) Secant uses two-point slope: $x_n = x_{n-1} - \frac{f(x_{n-1})(x_{n-2}-x_{n-1})}{f(x_{n-2})-f(x_{n-1})}$, $n \geq 2$ when the denominator is nonzero.

These define iterates only; local convergence is proved with Taylor's theorem (Theorem 4.1.26).



Figure 171: Newton: tangent-line zero. Secant: two-point line zero.

Definition 4.1.29. *Vector-Valued Derivative*
For $\mathbf{f} : [a,b] \to \mathbb{R}^k$, set $\mathbf{f}'(x) = \mathbf{L}$ when $\lim_{t \to x} \left\| \frac{\mathbf{f}(t)-\mathbf{f}(x)}{t-x} - \mathbf{L} \right\|_2 = 0$.
For $\mathbf{f} = (f_1, \ldots, f_k)$, this is equivalent to componentwise differentiability, with $\mathbf{f}'(x) = (f_1'(x), \ldots, f_k'(x))$.

Theorem 4.1.30. *Vector-Valued Mean Value Inequality*
If $\mathbf{f} : [a,b] \to \mathbb{R}^k$ is continuous on $[a,b]$ and differentiable on (a,b).
Then for some $x \in (a,b)$, $\|\mathbf{f}(b) - \mathbf{f}(a)\|_2 \leq (b-a)\,\|\mathbf{f}'(x)\|_2$.

Proof. Set $\mathbf{z} := \mathbf{f}(b) - \mathbf{f}(a)$ and $\varphi(t) := \mathbf{z} \cdot \mathbf{f}(t)$.
Corollary 4.1.13 gives $\varphi(b) - \varphi(a) = (b-a)\varphi'(x) = (b-a)\mathbf{z} \cdot \mathbf{f}'(x)$ for some $x \in (a,b)$. Since $\varphi(b) - \varphi(a) = \|\mathbf{z}\|_2^2$, Theorem 1.1.30 gives $\|\mathbf{z}\|_2^2 \leq (b-a)\,\|\mathbf{z}\|_2\,\|\mathbf{f}'(x)\|_2$, and the result follows by cancellation, with $\mathbf{z} = 0$ trivial.
Structural Commentary: A scalar projection yields the vector norm bound.
Key Insight: Project onto the displacement vector, then use Schwarz to return to norms. □

Method 4.1.31. *Differentiation Proof Patterns*

(i) *Domain*: require x_0 to be a limit point of X; isolated points have no derivative by Perspective 4.1.2.
(ii) *Definition*: factor $f(t) - f(x) = (t-x)Q_x(t)$ and take $Q_x(t) \to f'(x)$.
(iii) *Approximation*: use Proposition 4.1.3 to replace quotient limits by first-order error bounds.
(iv) *Algebra*: prove quotient identities first; use continuity to pass to limits.
(v) *Mean values*: choose h with equal endpoint values and apply Theorem 4.1.10.
(vi) *Taylor*: set $G = f - P - M(t-\alpha)^n$ and iterate the mean value theorem.

Key Insight: Difference quotients and auxiliary functions drive one-variable differentiation.

The integration side of one-variable calculus is organized by Riemann and Riemann–Stieltjes sums; the fundamental theorem connects the two halves.

This subsection develops Riemann and Riemann–Stieltjes integration by partitions, then links integrals to antiderivatives through the Fundamental Theorem of Calculus. Point partitions and interval-envelope partitions are equivalent via common refinement.

Definition 4.1.32. *Partitions and Riemann Upper/Lower Sums*
A *partition* of $[a,b]$ is $\Pi = \{x_0, \dots, x_n\}$ with $a = x_0 \le x_1 \le \cdots \le x_n = b$, $\Delta x_i := x_i - x_{i-1}$.
For bounded real f on $[a,b]$, set $M_i := \sup_{x_{i-1} \le x \le x_i} f(x)$, $m_i := \inf_{x_{i-1} \le x \le x_i} f(x)$,
and $U(\Pi, f) := \sum_{i=1}^{n} M_i \Delta x_i$, $L(\Pi, f) := \sum_{i=1}^{n} m_i \Delta x_i$.
Define $\overline{\int_a^b} f(x)\,dx := \inf_\Pi U(\Pi, f)$, $\underline{\int_a^b} f(x)\,dx := \sup_\Pi L(\Pi, f)$, with Π ranging over all partitions.
If these agree, f is *Riemann integrable*, written $f \in \mathcal{R}$, and the common value is $\int_a^b f(x)\,dx$ or $\int_a^b f$.
For $m \le f \le M$, $m(b-a) \le L(\Pi, f) \le U(\Pi, f) \le M(b-a)$, upper and lower integrals exist for bounded f.

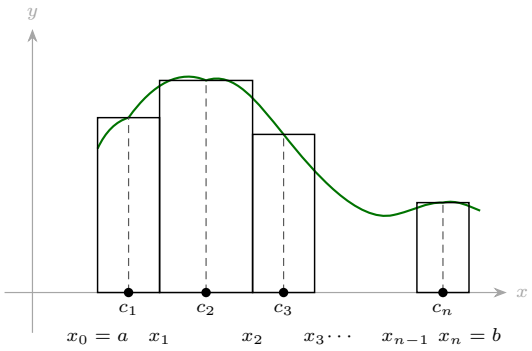


Figure 172: A tagged Riemann sum uses one sample value $f(c_i)$ on each subinterval $[x_{i-1}, x_i]$.

Definition 4.1.33. *Tagged Partitions and Riemann Sums*
A *tagged partition* of $[a,b]$ is a partition $\Pi = \{x_0, \dots, x_n\}$ together with tags $c_i \in [x_{i-1}, x_i]$.
For the ordinary Riemann integral, the corresponding *Riemann sum* is $R(f, \Pi, \{c_i\}) := \sum_{i=1}^{n} f(c_i)\Delta x_i$.
For every tagged partition, $L(\Pi, f) \le R(f, \Pi, \{c_i\}) \le U(\Pi, f)$.

Perspective 4.1.34. *Tagged Sums and Gauges*
Gauge integration keeps tagged sums but replaces one mesh bound δ with a local bound $\delta(c_i)$.
Thus fineness may vary by tag, giving a stronger framework for FTC-type results.

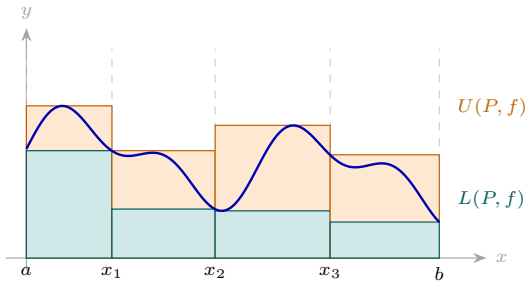


Figure 173: Upper and lower sums record the oscillation of f on each partition interval.

Definition 4.1.35. *Riemann–Stieltjes Integral*
Let α be increasing on $[a,b]$. For $\Pi = \{x_0, \dots, x_n\}$, put $\Delta\alpha_i := \alpha(x_i) - \alpha(x_{i-1}) \ge 0$.
For bounded real f, with M_i, m_i as in Definition 4.1.32, set $U(\Pi, f, \alpha) := \sum_{i=1}^{n} M_i \Delta\alpha_i$; $L(\Pi, f, \alpha) := \sum_{i=1}^{n} m_i \Delta\alpha_i$.
Define $\overline{\int_a^b} f\,d\alpha := \inf_\Pi U(\Pi, f, \alpha)$; $\underline{\int_a^b} f\,d\alpha := \sup_\Pi L(\Pi, f, \alpha)$. If value match, f is *Riemann integrable with respect to* α, written $f \in \mathcal{R}(\alpha)$, and the common value is $\int_a^b f\,d\alpha$, also written $\int_a^b f(x)\,d\alpha(x)$.
Taking $\alpha(x) = x$ gives the ordinary Riemann integral; the variable of integration is dummy.

Remark 4.1.36. *Stieltjes Increments as α-Lengths*
Stieltjes sums replace ordinary length Δx_i by $\Delta\alpha_i$, the length measured by α.
Thus $\alpha(x) = x$ gives Riemann sums, while jumps of α place point mass as in Theorem 4.1.53.

Definition 4.1.37. *Refinements*
A partition Π^* is a *refinement* of Π if $\Pi^* \supseteq \Pi$. Given two partitions Π_1, Π_2, their *common refinement* is $\Pi_1 \cup \Pi_2$.

Theorem 4.1.38. *Refinement Monotonicity*
If Π^* is a refinement of Π, then $L(\Pi, f, \alpha) \le L(\Pi^*, f, \alpha)$ and $U(\Pi^*, f, \alpha) \le U(\Pi, f, \alpha)$.

Proof. It suffices to add one point $x^* \in (x_{i-1}, x_i)$. If w_1, w_2 are infima on $[x_{i-1}, x^*]$ and $[x^*, x_i]$, then $w_1, w_2 \ge m_i$ and $L(\Pi^*, f, \alpha) - L(\Pi, f, \alpha) = (w_1 - m_i)\big[\alpha(x^*) - \alpha(x_{i-1})\big] + (w_2 - m_i)\big[\alpha(x_i) - \alpha(x^*)\big] \ge 0$.
Iterate for additional points. The upper-sum inequality is the same argument with suprema.
Structural Commentary: Refinement raises lower sums and lowers upper sums.
Key Insight: Splitting an interval tightens the local oscillation bounds. □

Perspective 4.1.39. *Piecewise-Constant Envelopes*
A partition Π gives lower and upper step envelopes with values m_i and M_i; their integrals are exactly $L(\Pi, f)$ and $U(\Pi, f)$. Common refinements make these envelope sums compatible across partitions.

Theorem 4.1.40. *Lower Integral Below Upper Integral*
For every bounded real f and increasing α on $[a, b]$, $\underline{\int_a^b f\, d\alpha} \le \overline{\int_a^b f\, d\alpha}$.

Proof. Given partitions Π_1, Π_2, let $\Pi^* = \Pi_1 \cup \Pi_2$.
By Theorem 4.1.38, $L(\Pi_1, f, \alpha) \le L(\Pi^*, f, \alpha) \le U(\Pi^*, f, \alpha) \le U(\Pi_2, f, \alpha)$.
Take the supremum over Π_1 and then the infimum over Π_2.
Structural Commentary: Common refinements compare sums from different partitions.
Key Insight: The lower envelope never rises above the upper envelope. □

Theorem 4.1.41. *Upper–Lower Criterion*
A bounded real f is in $\mathcal{R}(\alpha)$ on $[a, b]$ if and only if $\forall \varepsilon > 0\ \exists$ a partition Π with $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) < \varepsilon$.

Proof.

($\Rightarrow$) If $f \in \mathcal{R}(\alpha)$, choose Π_1, Π_2 whose lower and upper sums are within $\varepsilon/2$ of $\int_a^b f\, d\alpha$.
Their common refinement Π satisfies the displayed inequality by Theorem 4.1.38.

($\Leftarrow$) If such Π exist, then $L(\Pi, f, \alpha) \le \underline{\int_a^b f\, d\alpha} \le \overline{\int_a^b f\, d\alpha} \le U(\Pi, f, \alpha)$ forces upper and lower integrals to agree.

Structural Commentary: Integrability is exactly the ability to make one partition squeeze the oscillation.
Key Insight: The gap $U - L$ is the quantitative defect of integrability. □

Perspective 4.1.42. *Darboux Gaps and Riemann Sums*
For $\alpha(x) = x$, a small Darboux gap $U(\Pi, f) - L(\Pi, f)$ traps every tagged sum for Π near the integral. This is the local reason the Darboux definition above and the tagged Riemann-sum definition give the same number. The same estimate works with Stieltjes weights $\Delta\alpha_i$ whenever the upper–lower gap is small.

Theorem 4.1.43. *Tagged Sums*
Let $\Pi = \{x_0, \dots, x_n\}$ satisfy $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) < \varepsilon$, and choose tags $t_i \in [x_{i-1}, x_i]$.

(i) If $s_i, t_i \in [x_{i-1}, x_i]$, then $\sum_{i=1}^n |f(s_i) - f(t_i)|\, \Delta\alpha_i < \varepsilon$.

(ii) If $f \in \mathcal{R}(\alpha)$, then $\left| \sum_{i=1}^n f(t_i)\Delta\alpha_i - \int_a^b f\, d\alpha \right| < \varepsilon$.

Proof.

(i) Since $|f(s_i) - f(t_i)| \le M_i - m_i$, summing gives $\sum_{i=1}^n |f(s_i) - f(t_i)|\, \Delta\alpha_i \le U(\Pi, f, \alpha) - L(\Pi, f, \alpha) < \varepsilon$.

(ii) The tagged sum and $\int_a^b f\, d\alpha$ both lie between $L(\Pi, f, \alpha)$ and $U(\Pi, f, \alpha)$, so their distance is $< \varepsilon$.

Structural Commentary: A small Darboux gap makes tags and the integral close.
Key Insight: Tags are harmless after oscillation has been controlled. □

Theorem 4.1.44. *Continuous Functions Are Riemann–Stieltjes Integrable*
If f is continuous on $[a, b]$ and α is increasing on $[a, b]$, then $f \in \mathcal{R}(\alpha)$.

Proof. If $\alpha(b) = \alpha(a)$, the Darboux gap is zero.
Otherwise choose $\eta > 0$ with $\eta[\alpha(b) - \alpha(a)] < \varepsilon$. Uniform continuity gives partition Π with $M_i - m_i \le \eta\ \forall i$, so $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) = \sum_{i=1}^n (M_i - m_i)\Delta\alpha_i \le \eta[\alpha(b) - \alpha(a)] < \varepsilon$. Apply Theorem 4.1.41.
Structural Commentary: Uniform continuity controls oscillation uniformly across the partition.
Key Insight: Continuity turns local small intervals into a small total Darboux gap. □

Theorem 4.1.45. *Monotone Functions and Continuous Integrators*
If f is monotone on $[a, b]$ and α is continuous and increasing on $[a, b]$, then $f \in \mathcal{R}(\alpha)$.

Proof. Assume first that f is increasing. If $\alpha(b) = \alpha(a)$, then every $\Delta\alpha_i = 0$, so $U(\Pi, f, \alpha) = L(\Pi, f, \alpha) = 0$ for every partition Π. Choose n and a partition Π such that $\Delta\alpha_i = [\alpha(b) - \alpha(a)]/n$ for all i.
Existence follows from continuity of α and Theorem 2.3.42.
Then $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) = \frac{\alpha(b)-\alpha(a)}{n} \sum_{i=1}^{n} \big(f(x_i) - f(x_{i-1})\big) = \frac{\alpha(b)-\alpha(a)}{n}\big(f(b) - f(a)\big) < \varepsilon$ for n large.
If f is decreasing, use the same partition. Then the suprema and infima on $[x_{i-1}, x_i]$ are $f(x_{i-1})$ and $f(x_i)$, so the same computation gives $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) = \frac{\alpha(b)-\alpha(a)}{n}\big(f(a) - f(b)\big)$, which is $< \varepsilon$ for large n.
Structural Commentary: Equal α-increments convert the monotone oscillations into a telescoping sum.
Key Insight: A monotone function's total oscillation is finite and can be spread thinly. □

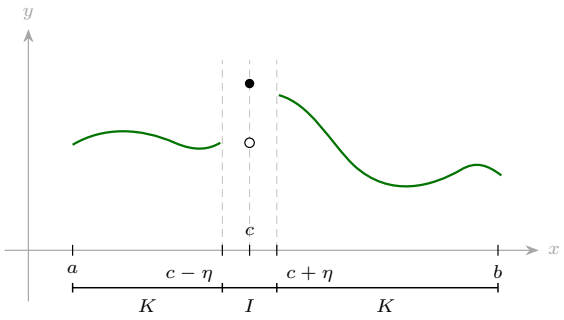


Figure 174: Isolate the discontinuity in a short interval; control the compact pieces outside it.

Theorem 4.1.46. *Finitely Many Discontinuities*
Suppose f is bounded on $[a, b]$, the discontinuity set $E := \{x \in [a, b] : f$ is discontinuous at $x\}$ is finite, and α is continuous at each point of E. Then $f \in \mathcal{R}(\alpha)$.

Proof. Let $M = \sup_{[a,b]} |f|$ and let $\eta > 0$. By continuity of α on the finite set E, choose disjoint intervals $I_j = [u_j, v_j]$ covering E, with each point of $E \cap (a, b)$ in some (u_j, v_j), such that $\sum_j \big(\alpha(v_j) - \alpha(u_j)\big) < \eta$.
Set $K = [a, b] \setminus \bigcup_j (u_j, v_j)$. Then $f|_K$ is uniformly continuous by Theorem 2.3.39.
Choose a partition Π containing all u_j, v_j, with no points inside the intervals (u_j, v_j), and fine enough on K that $M_i - m_i \le \eta$ whenever $[x_{i-1}, x_i] \subset K$.
Hence $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) \le \eta\big[\alpha(b) - \alpha(a)\big] + 2M \sum_j \big(\alpha(v_j) - \alpha(u_j)\big) < \eta\big[\alpha(b) - \alpha(a) + 2M\big]$.
Since η is arbitrary, Theorem 4.1.41 applies.
Structural Commentary: Isolate bad points in small α-intervals; use uniform continuity elsewhere.
Key Insight: Finitely many discontinuities do not matter unless α also jumps there. □

Example 4.1.47. *Dirichlet's Function Is Not Riemann Integrable*
Let $g = \mathbf{1}_{\mathbb{Q}}$ on $[0, 1]$. Every nondegenerate subinterval contains rationals and irrationals by Corollary 1.1.17, so for every partition Π, $M_i = 1$ and $m_i = 0$ on each nonzero-width subinterval.
Thus $U(\Pi, g) = 1$ and $L(\Pi, g) = 0$, hence $g \notin \mathcal{R}$.

Remark 4.1.48. *Small Discontinuity Sets*
The same idea underlies Theorem 5.4.20: small discontinuity sets, later called measure-zero sets in Definition 5.1.66, do not obstruct integrability. Dirichlet's function fails because every point is discontinuous.

Theorem 4.1.49. *Continuous Composition Preserves Integrability*
Let α be increasing on $[a, b]$. Suppose $f \in \mathcal{R}(\alpha)$, $m \le f(x) \le M$ on $[a, b]$, and $\phi : [m, M] \to \mathbb{R}$ is continuous. Define $h : [a, b] \to \mathbb{R}$ by $h(x) := \phi(f(x))$. Then $h \in \mathcal{R}(\alpha)$.

Proof. Let $C := \sup_{y \in [m,M]} |\phi(y)|$ and let $\varepsilon > 0$. Choose $\eta > 0$ such that $\eta[\alpha(b) - \alpha(a) + 2C] < \varepsilon$. Since ϕ is uniformly continuous on $[m, M]$, choose $\gamma > 0$ such that $|s - t| < \gamma \implies |\phi(s) - \phi(t)| < \eta$ $(s, t \in [m, M])$.
Choose a partition $\Pi = \{x_0, \dots, x_n\}$ with $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) < \gamma\eta$.
Let M_i, m_i denote upper and lower values for f on $[x_{i-1}, x_i]$, M_i^h, m_i^h denote corresponding values for h.
Split $A := \{i : M_i - m_i < \gamma\}$, $B := \{i : M_i - m_i \ge \gamma\}$. For $i \in A$, $M_i^h - m_i^h \le \eta$; for $i \in B$, $M_i^h - m_i^h \le 2C$.
Also $\gamma \sum_{i \in B} \Delta\alpha_i \le \sum_{i \in B} (M_i - m_i)\Delta\alpha_i \le U(\Pi, f, \alpha) - L(\Pi, f, \alpha) < \gamma\eta$, so $\sum_{i \in B} \Delta\alpha_i < \eta$.
$U(\Pi, h, \alpha) - L(\Pi, h, \alpha) \le \eta \sum_{i \in A} \Delta\alpha_i + 2C \sum_{i \in B} \Delta\alpha_i < \eta[\alpha(b) - \alpha(a) + 2C] < \varepsilon$. Theorem 4.1.41 gives $h \in \mathcal{R}(\alpha)$.
Structural Commentary: Integrability survives continuous remeasurement of the values of f.
Key Insight: Large oscillation intervals have small total α-weight. □

Theorem 4.1.50. *Basic Properties of the Integral*
Let all functions be bounded and real on $[a,b]$, and let all integrators be increasing.

(i) If $f_1, f_2 \in \mathcal{R}(\alpha)$, $c \in \mathbb{R}$, then $f_1 + f_2, cf_1 \in \mathcal{R}(\alpha)$, and $\int_a^b (f_1+f_2)\,d\alpha = \int_a^b f_1\,d\alpha + \int_a^b f_2\,d\alpha$; $\int_a^b cf_1\,d\alpha = c\int_a^b f_1\,d\alpha$.

(ii) If $f_1 \le f_2$ and both functions are in $\mathcal{R}(\alpha)$, then $\int_a^b f_1\,d\alpha \le \int_a^b f_2\,d\alpha$.

(iii) If $f \in \mathcal{R}(\alpha)$ and $a < c < b$, then f is integrable on $[a,c]$ and $[c,b]$, and $\int_a^c f\,d\alpha + \int_c^b f\,d\alpha = \int_a^b f\,d\alpha$.

(iv) If $f \in \mathcal{R}(\alpha)$ and $|f| \le M$, then $\left|\int_a^b f\,d\alpha\right| \le M\big[\alpha(b) - \alpha(a)\big]$

(v) If $f \in \mathcal{R}(\alpha_1)$ and $f \in \mathcal{R}(\alpha_2)$, then $f \in \mathcal{R}(\alpha_1+\alpha_2)$ and $\int_a^b f\,d(\alpha_1+\alpha_2) = \int_a^b f\,d\alpha_1 + \int_a^b f\,d\alpha_2$. If $c > 0$ and $f \in \mathcal{R}(\alpha)$, then $f \in \mathcal{R}(c\alpha)$ and $\int_a^b f\,d(c\alpha) = c\int_a^b f\,d\alpha$.

Proof.

(i) $L(\Pi, f_1, \alpha) + L(\Pi, f_2, \alpha) \le L(\Pi, f_1+f_2, \alpha) \le U(\Pi, f_1+f_2, \alpha) \le U(\Pi, f_1, \alpha) + U(\Pi, f_2, \alpha)$, where Π is common refinement, and Theorem 4.1.41; scalar multiplication is the same estimate with signs handled by swapping upper and lower sums when $c < 0$.
(ii) If $f_1 \le f_2$, then $L(\Pi, f_2 - f_1, \alpha) \ge 0$ for every Π; apply (i) to $f_2 - f_1$.
(iii) Refine so that $c \in \Pi$. Then the lower and upper sums split into the sums over $[a,c]$ and $[c,b]$.
(iv) Apply (ii) to $-M \le f \le M$ and use the integrals of the constant functions $\pm M$.
(v) Since $\Delta(\alpha_1+\alpha_2)_i = \Delta\alpha_{1,i} + \Delta\alpha_{2,i}$ and $\Delta(c\alpha)_i = c\,\Delta\alpha_i$, corresponding lower and upper sums split exactly.

Structural Commentary: The algebra of integrals is inherited from the algebra of upper and lower sums.
Key Insight: Common refinements let several integrable functions be approximated by the same partition. □

Theorem 4.1.51. *Products and Absolute Values*
If $f, g \in \mathcal{R}(\alpha)$ on $[a,b]$, then $fg \in \mathcal{R}(\alpha)$, $|f| \in \mathcal{R}(\alpha)$, and $\left|\int_a^b f\,d\alpha\right| \le \int_a^b |f|\,d\alpha$.

Proof. By Theorem 4.1.50, $f \pm g \in \mathcal{R}(\alpha)$. Applying Theorem 4.1.49 to $\phi(t) = t^2$ gives $(f \pm g)^2 \in \mathcal{R}(\alpha)$, hence $fg \in \mathcal{R}(\alpha)$ from $4fg = (f+g)^2 - (f-g)^2$.
Applying the same theorem to $\phi(t) = |t|$ gives $|f| \in \mathcal{R}(\alpha)$, and $-|f| \le f \le |f|$ gives the inequality.
Structural Commentary: Products reduce to squares; absolute values to composition.
Key Insight: Once continuous composition is available, the usual algebra closes. □

Definition 4.1.52. *Unit Step Function*
The *unit step function* I is $I(x) = \begin{cases} 0, & x \le 0, \\ 1, & x > 0. \end{cases}$.

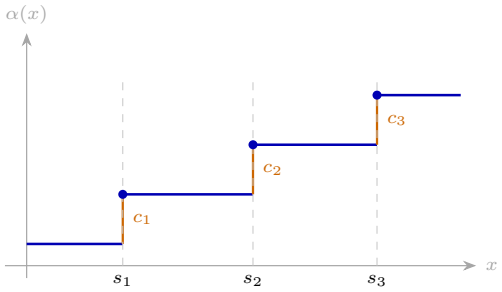


Figure 175: A step integrator concentrates mass at jumps: $\int f\,d\alpha = \sum_i c_i f(s_i)$.

Theorem 4.1.53. *Single Jump Integrator*
If $a < s < b$, f is bounded on $[a,b]$ and continuous at s, and $\alpha(x) = I(x-s)$, then $\int_a^b f\,d\alpha = f(s)$.

Proof. Use partitions $\Pi = \{a, s, x_2, b\}$ with $s < x_2 < b$. The only nonzero $\Delta\alpha_i$ occurs on $[s, x_2]$, so $L(\Pi, f, \alpha)$ and $U(\Pi, f, \alpha)$ are the infimum and supremum of f on $[s, x_2]$. As $x_2 \downarrow s$, continuity at s squeezes both to $f(s)$.
Structural Commentary: A jump of α makes the integral sample f at the jump.
Key Insight: Stieltjes integration includes point masses without changing the Darboux framework. □

Theorem 4.1.54. *Pure Step Integrators Give Series*
Suppose $c_n \ge 0$, $\sum c_n$ converges, and (s_n) is a sequence of distinct points in (a,b).
Define $\alpha(x) := \sum_{n=1}^{\infty} c_n I(x - s_n)$. If f is continuous on $[a,b]$, then $\int_a^b f\,d\alpha = \sum_{n=1}^{\infty} c_n f(s_n)$.

Proof. By comparison, α is increasing. For $N \in \mathbb{N}$, put $\alpha_N(x) := \sum_{n=1}^{N} c_n I(x - s_n)$ and $\beta_N := \alpha - \alpha_N$..

Theorems 4.1.53 and 4.1.50 give $\int_a^b f\, d\alpha_N = \sum_{n=1}^{N} c_n f(s_n)$. Also $f \in \mathcal{R}(\beta_N)$ by Theorem 4.1.44.

If $|f| \le M$, then Theorem 4.1.50 gives $\left|\int_a^b f\, d\beta_N\right| \le M[\beta_N(b) - \beta_N(a)] = M \sum_{n>N} c_n$.

Letting $N \to \infty$ and using additivity in the integrator gives the formula.

Structural Commentary: Finite step integrators sum point samples; summable tails vanish by the integral bound.

Key Insight: A Stieltjes integral can be an ordinary series in disguise. □

Theorem 4.1.55. *Differentiable Integrators Reduce to Riemann Integrals*
Assume α increases monotonically, $\alpha' \in \mathcal{R}$ on $[a, b]$, and f is bounded and real on $[a, b]$.
Then $f \in \mathcal{R}(\alpha) \iff f\alpha' \in \mathcal{R}$, and in that case $\int_a^b f\, d\alpha = \int_a^b f(x)\alpha'(x)\, dx$.

Proof. Let $M = \sup_{[a,b]} |f|$. Choose Π with $U(\Pi, \alpha') - L(\Pi, \alpha') < \varepsilon$.
For any refinement and any tags $s_i \in [x_{i-1}, x_i]$, MVT gives points $t_i \in [x_{i-1}, x_i]$ with $\Delta\alpha_i = \alpha'(t_i)\Delta x_i$.
Hence $|\sum_i f(s_i)\Delta\alpha_i - \sum_i f(s_i)\alpha'(s_i)\Delta x_i| \le M \sum_i |\alpha'(t_i) - \alpha'(s_i)|\, \Delta x_i \le M\varepsilon$. Thus every tagged Stieltjes sum for $f\, d\alpha$ is within $M\varepsilon$ of the corresponding tagged Riemann sum for $f\alpha'\, dx$, and conversely. By the tagged-sum criterion, one family has a limit exactly when the other does, and the two limits are equal.

Structural Commentary: The increment $d\alpha$ is replaced by $\alpha'(x)\, dx$ when α' is Riemann integrable.

Key Insight: The mean value theorem turns Stieltjes increments into ordinary Riemann weights. □

Theorem 4.1.56. *Change of Variable*
Let $\varphi : [A, B] \to [a, b]$ be strictly increasing, continuous, and onto.
Let α be increasing on $[a, b]$, $f \in \mathcal{R}(\alpha)$, and set $\beta(y) := \alpha(\varphi(y))$ and $g(y) := f(\varphi(y))$.
Then $g \in \mathcal{R}(\beta)$ and $\int_A^B g\, d\beta = \int_a^b f\, d\alpha$. If also $\alpha(x) = x$ and $\varphi' \in \mathcal{R}$, then $\int_a^b f(x)\, dx = \int_A^B f(\varphi(y))\varphi'(y)\, dy$.

Proof. Partitions $Q = \{y_i\}$ of $[A, B]$ correspond to partitions $\Pi = \{\varphi(y_i)\}$ of $[a, b]$. On corresponding intervals, g has same range as f and $\Delta\beta_i = \Delta\alpha_i$, so $U(Q, g, \beta) = U(\Pi, f, \alpha)$ and $L(Q, g, \beta) = L(\Pi, f, \alpha)$. The upper–lower criterion gives the Stieltjes formula. If $\alpha(x) = x$, then $\beta = \varphi$; apply Theorem 4.1.55 to the left side.

Structural Commentary: A monotone change of variable relabels the same partition intervals.

Key Insight: Stieltjes change of variable is equality of Darboux sums under reparametrization. □

Theorem 4.1.57. *Integral Functions*
Let $f \in \mathcal{R}$ on $[a, b]$ and define $F(x) := \int_a^x f(t)\, dt$. Then F is continuous.
If f is continuous at x_0, then $F'(x_0) = f(x_0)$, with one-sided derivatives at endpoints.

Proof. Let $|f| \le M$. For $x < y$, Theorem 4.1.50 gives $|F(y) - F(x)| = \left|\int_x^y f(t)\, dt\right| \le M(y - x)$ so F continuous.
If f continuous at x_0, for $x \ne x_0$ near x_0, $\left|\frac{F(x) - F(x_0)}{x - x_0} - f(x_0)\right| \le \sup_{u \in [x, x_0]} |f(u) - f(x_0)| \to 0$. $F'(x_0) = f(x_0)$.

Structural Commentary: Boundedness controls increments; local continuity gives the derivative.

Key Insight: Averaging a continuous function over shrinking intervals returns its value. □

Theorem 4.1.58. *Fundamental Theorem of Calculus*
If $f \in \mathcal{R}$ on $[a, b]$ and $\exists$ a differentiable function F on $[a, b]$ such that $F' = f$, then $\int_a^b f(x)\, dx = F(b) - F(a)$.

Proof. Choose $\Pi = \{x_0, \ldots, x_n\}$ with $U(\Pi, f) - L(\Pi, f) < \varepsilon$.
By MVT, choose $t_i \in [x_{i-1}, x_i]$ so that $F(b) - F(a) = \sum_i \big(F(x_i) - F(x_{i-1})\big) = \sum_i f(t_i)\Delta x_i$.

Theorem 4.1.43 puts this tagged sum within ε of $\int_a^b f$. Let $\varepsilon \to 0$.

Structural Commentary: Mean-value increments convert antiderivatives into tagged Riemann sums.

Key Insight: The integral is the common limit of all sufficiently fine tagged sums. □

Theorem 4.1.59. *Integration by Parts*
Suppose F, G are differentiable on $[a, b]$, $F' = f \in \mathcal{R}$, and $G' = g \in \mathcal{R}$.
Then $\int_a^b F(x)g(x)\, dx = F(b)G(b) - F(a)G(a) - \int_a^b f(x)G(x)\, dx$.

Proof. Set $H := FG$. Since F, G are continuous, $F, G \in \mathcal{R}$; hence $Fg, fG \in \mathcal{R}$ by Theorem 4.1.51.
Thus $H' = Fg + fG \in \mathcal{R}$, and Theorem 4.1.58 gives $F(b)G(b) - F(a)G(a) = H(b) - H(a) = \int_a^b Fg + \int_a^b fG$.

Structural Commentary: Integration by parts is the product rule integrated over the interval.

Key Insight: Move one derivative from G onto F by applying the fundamental theorem to FG. □

Definition 4.1.60. *Improper Riemann Integrals*
Let $f \in \mathcal{R}$ on every compact subinterval of its domain.
Set $\int_a^b f(x)\,dx := \lim_{d\to b^-} \int_a^d f(x)\,dx$ at a right endpoint, and $\int_a^b f(x)\,dx := \lim_{c\to a^+} \int_c^b f(x)\,dx$ at a left endpoint.
If both endpoints are improper, choose $\xi \in (a,b)$ and set $\int_a^b f(x)\,dx := \int_a^\xi f(x)\,dx + \int_\xi^b f(x)\,dx$, provided both one-sided integrals exist and the sum is not indeterminate.
A finite value means the improper integral *converges*; otherwise it diverges to $+\infty$ or $-\infty$.

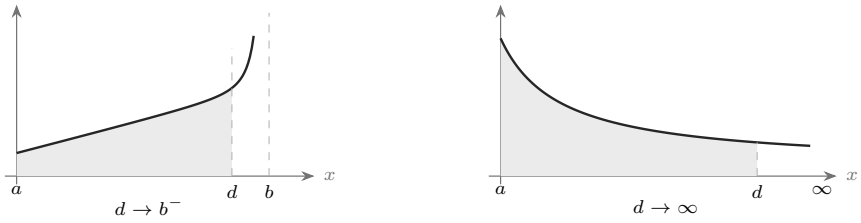


Figure 176: Improper integrals are limits of ordinary integrals over compact truncations.

Theorem 4.1.61. *Nonnegative Improper Integrals*
Let $f \geq 0$ and $f \in \mathcal{R}$ on every $[a,d]$, $a < d < b$. Then $\int_a^b f(x)\,dx = \sup\left\{\int_a^d f(x)\,dx : a < d < b\right\}$.
Thus $\int_a^b f$ exists in $[0,+\infty]$ and converges exactly when this supremum is finite.

Proof. For $a < d_1 < d_2 < b$, $\int_a^{d_2} f - \int_a^{d_1} f = \int_{d_1}^{d_2} f \geq 0$.
Hence $d \mapsto \int_a^d f$ is increasing, so its extended limit exists and equals its supremum.
Structural Commentary: Nonnegative area can only accumulate.
Key Insight: Improper convergence for $f \geq 0$ is boundedness of all finite truncations. □

Perspective 4.1.62. *Principal Value Is Different*
Improper integrals over $(-\infty,\infty)$ require separate one-sided limits.
A symmetric limit such as $\lim_{A\to\infty} \int_{-A}^{A} \sin x\,dx$ may still exist (Cauchy principal value, not Definition 4.1.60).

Definition 4.1.63. *Improper Riemann–Stieltjes Integrals on $\mathbb{R}$*
Let $\alpha : \mathbb{R} \to \mathbb{R}$ be increasing and $f \in \mathcal{R}(\alpha)$ on every finite interval.
Define, when the limits exist, $\int_0^\infty f\,d\alpha := \lim_{b\to\infty} \int_0^b f\,d\alpha$; and $\int_{-\infty}^0 f\,d\alpha := \lim_{a\to-\infty} \int_a^0 f\,d\alpha$.
If their sum is not indeterminate, set $\int_{-\infty}^\infty f\,d\alpha := \int_{-\infty}^0 f\,d\alpha + \int_0^\infty f\,d\alpha$.
Finite value means f is *α-integrable on $\mathbb{R}$*. Write $\alpha(\pm\infty) := \lim_{t\to\pm\infty} \alpha(t)$ when these extended limits are used.

Theorem 4.1.64. *Nonnegative and Bounded-Mass Stieltjes Integrals*
Let $\alpha : \mathbb{R} \to \mathbb{R}$ be increasing, and suppose $f \in \mathcal{R}(\alpha)$ on every finite interval.
(i) If $f \geq 0$, then either f is α-integrable on $\mathbb{R}$ or $\int_{-\infty}^\infty f\,d\alpha = +\infty$.
(ii) If $-\infty < \alpha(-\infty) < \alpha(\infty) < +\infty$ and f is bounded, then f is α-integrable on $\mathbb{R}$.

Proof.
(i) By additivity and order preservation, $\int_0^b f\,d\alpha$ and $\int_a^0 f\,d\alpha$ are monotone under outward truncation. Hence both extended one-sided limits exist, so their sum is finite or $+\infty$.
(ii) Let $|f| \leq M$. For finite $a < b$, Theorem 4.1.50 gives $\left|\int_a^b f\,d\alpha\right| \leq M\big[\alpha(b) - \alpha(a)\big]$.
Since $\alpha(\infty)$ and $\alpha(-\infty)$ are finite, the right and left tails have vanishing α-increment.
Thus the one-sided truncations are Cauchy and have finite limits.

Structural Commentary: Positivity gives monotone limits; finite mass bounds tails.
Key Insight: A bounded increasing integrator behaves like finite total mass. □

Definition 4.1.65. *Vector-Valued Riemann–Stieltjes Integral*
For $\mathbf{f} = (f_1,\dots,f_k) : [a,b] \to \mathbb{R}^k$, define $\mathbf{f} \in \mathcal{R}(\alpha)$ iff $f_j \in \mathcal{R}(\alpha)\ \forall j$, and set $\int_a^b \mathbf{f}\,d\alpha := \left(\int_a^b f_1\,d\alpha,\dots,\int_a^b f_k\,d\alpha\right)$.

Theorem 4.1.66. *Vector-Valued Integral Results*
(i) If $\mathbf{f},\mathbf{F} : [a,b] \to \mathbb{R}^k$, $\mathbf{f} \in \mathcal{R}$, and $\mathbf{F}' = \mathbf{f}$, then $\int_a^b \mathbf{f}(t)\,dt = \mathbf{F}(b) - \mathbf{F}(a)$.
(ii) If $\mathbf{f} : [a,b] \to \mathbb{R}^k$ and $\mathbf{f} \in \mathcal{R}(\alpha)$, then $\|\mathbf{f}\| \in \mathcal{R}(\alpha)$ and $\left\|\int_a^b \mathbf{f}\,d\alpha\right\| \leq \int_a^b \|\mathbf{f}\|\,d\alpha$

Proof.

(i) Write $\mathbf{f} = (f_1, \dots, f_k)$ and $\mathbf{F} = (F_1, \dots, F_k)$. For each j, Theorem 4.1.58 gives $\int_a^b f_j(t)\,dt = F_j(b) - F_j(a)$. The vector identity follows componentwise.

(ii) Write $\mathbf{f} = (f_1, \dots, f_k)$. Since $\|\mathbf{f}\| = (f_1^2 + \cdots + f_k^2)^{1/2}$, Theorems 4.1.51 and 4.1.49 give $\|\mathbf{f}\| \in \mathcal{R}(\alpha)$. Let $\mathbf{y} := \int_a^b \mathbf{f}\,d\alpha$. If $\mathbf{y} = 0$, the inequality is immediate.

Otherwise, $\|\mathbf{y}\|^2 = \sum_{j=1}^{k} y_j \int_a^b f_j\,d\alpha = \int_a^b \sum_{j=1}^{k} y_j f_j\,d\alpha \le \int_a^b \|\mathbf{y}\|\,\|\mathbf{f}\|\,d\alpha = \|\mathbf{y}\| \int_a^b \|\mathbf{f}\|\,d\alpha$, by Theorem 1.1.30 and order preservation. Divide by $\|\mathbf{y}\|$.

Structural Commentary: Coordinatewise integration; Schwarz controls norms.
Key Insight: Integrate the projection of $\mathbf{f}$ onto the resulting vector. □

Definition 4.1.67. *Curves and Rectifiability*
A *curve* in $\mathbb{R}^k$ is a continuous map $\gamma : [a, b] \to \mathbb{R}^k$; it is an *arc* if one-to-one and a *closed curve* if $\gamma(a) = \gamma(b)$.
Curves are maps, not just ranges. For a partition $\Pi = \{x_0, \dots, x_n\}$, set $\ell(\Pi, \gamma) := \sum_{i=1}^{n} \|\gamma(x_i) - \gamma(x_{i-1})\|$.
The *length* is $\ell(\gamma) := \sup_\Pi \ell(\Pi, \gamma)$; if $\ell(\gamma) < \infty$, γ is *rectifiable*.

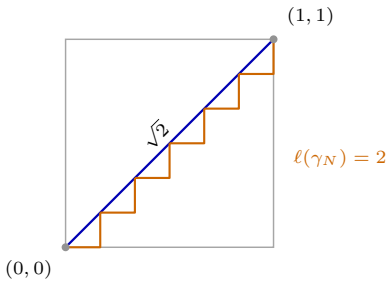


Figure 177: A staircase path can approach the diagonal in position while its polygonal length stays 2, not $\sqrt{2}$.

Perspective 4.1.68. *Length Is Not Just Shape*
Length is a supremum over partitions of one curve. Passing lengths through curve limits needs more than positional convergence; Theorem 4.1.69 uses C^1 speed control.

Theorem 4.1.69. *Length of a Continuously Differentiable Curve*
If $\gamma' : [a, b] \to \mathbb{R}^k$ is continuous, then γ is rectifiable and $\ell(\gamma) = \int_a^b \|\gamma'(t)\|\,dt$.

Proof. For $\Pi = \{x_0, \dots, x_n\}$, Theorem 4.1.66(i) gives $\|\gamma(x_i) - \gamma(x_{i-1})\| = \left\| \int_{x_{i-1}}^{x_i} \gamma'(t)\,dt \right\| \le \int_{x_{i-1}}^{x_i} \|\gamma'(t)\|\,dt$.
Hence $\ell(\Pi, \gamma) \le \int_a^b \|\gamma'\|$. Conversely, choose Π so fine that $\|\gamma'(s) - \gamma'(t)\| < \varepsilon$ on each subinterval.
Choose $t_i \in [x_{i-1}, x_i]$. Then $\int_{x_{i-1}}^{x_i} \|\gamma'(t)\|\,dt \le \|\gamma'(t_i)\|\,\Delta x_i + \varepsilon \Delta x_i$.
Also $\gamma(x_i) - \gamma(x_{i-1}) = \int_{x_{i-1}}^{x_i} \gamma'(t)\,dt$, so $\|\gamma'(t_i)\|\,\Delta x_i \le \|\gamma(x_i) - \gamma(x_{i-1})\| + \varepsilon \Delta x_i$.
Thus $\int_{x_{i-1}}^{x_i} \|\gamma'(t)\|\,dt \le \|\gamma(x_i) - \gamma(x_{i-1})\| + 2\varepsilon \Delta x_i$.
Summing gives $\int_a^b \|\gamma'\| \le \ell(\gamma) + 2\varepsilon(b - a)$; let $\varepsilon \downarrow 0$.
Structural Commentary: Upper bounds are immediate; uniform continuity gives matching lower bounds.
Key Insight: For C^1 curves, length is accumulated speed. □

Method 4.1.70. *Integration Proof Patterns*

(i) *Definitions*: fix $\Pi = \{x_i\}_{i=0}^n$ before defining $M_i, m_i, \Delta x_i, \Delta \alpha_i$.
(ii) *Refinement*: prove the one-point insertion case; iterate.
(iii) *Integrability*: find Π with $U(\Pi, f, \alpha) - L(\Pi, f, \alpha) < \varepsilon$.
(iv) *Bad sets*: cover E by $I_j = [u_j, v_j]$ with $\sum_j [\alpha(v_j) - \alpha(u_j)] < \varepsilon$; control $K = [a, b] \setminus \bigcup_j (u_j, v_j)$ uniformly.
(v) *Properties*: pass to a common refinement before comparing sums.
(vi) *Calculus*: use the mean value theorem to form tagged sums, then apply Theorem 4.1.43.
(vii) *Improper integrals*: truncate endpoints; use monotonicity or tail bounds.

Key Insight: Riemann–Stieltjes integration is Darboux approximation plus a choice of interval weights $\Delta \alpha_i$.

4.2 Several-Variable Differentiation

Several-variable differentiation studies best linear approximations and local geometry, using linear maps, operator norms, determinants, invertibility, and metric estimates.

Definition 4.2.1. *Total Derivative*
Let $E \subseteq \mathbb{R}^n$ be open, $\mathbf{f} : E \to \mathbb{R}^m$, $\mathbf{x} \in E$, and $\mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$ denote space of linear maps (Definition 1.3.1).
If $\exists A \in \mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$ s.t. $\lim_{\mathbf{h}\to\mathbf{0}} \frac{\|\mathbf{f}(\mathbf{x}+\mathbf{h})-\mathbf{f}(\mathbf{x})-A\mathbf{h}\|_2}{\|\mathbf{h}\|_2} = 0$, then $\mathbf{f}$ is *differentiable at* $\mathbf{x}$, and A denoted $D\mathbf{f}(\mathbf{x})$ or $\mathbf{f}'(\mathbf{x})$.
Equivalently, $\mathbf{f}(\mathbf{x}+\mathbf{h}) - \mathbf{f}(\mathbf{x}) = D\mathbf{f}(\mathbf{x})\mathbf{h} + \mathbf{r}(\mathbf{h})$, and $\frac{\|\mathbf{r}(\mathbf{h})\|_2}{\|\mathbf{h}\|_2} \to 0$.
If this holds for every $\mathbf{x} \in E$, then $\mathbf{f}$ is differentiable on E.

Perspective 4.2.2. *Derivative First, Matrix Second*
The total derivative is a linear map first; the Jacobian is its matrix in chosen bases. This keeps coordinate formulas subordinate to the intrinsic estimate $\|A\mathbf{h}\|_2 \le \|A\|_{\mathrm{op}} \|\mathbf{h}\|_2$ from Definition 4.3.4.

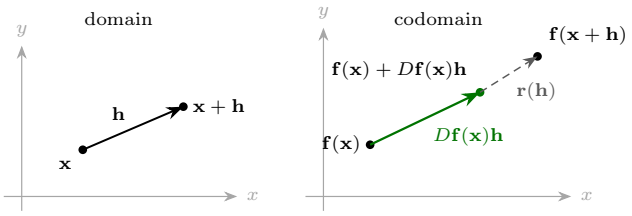


Figure 178: $D\mathbf{f}(\mathbf{x})\mathbf{h}$ is the linear part; $\mathbf{r}(\mathbf{h}) = o(\|\mathbf{h}\|_2)$ is the error.

Theorem 4.2.3. *Basic Consequences of Total Differentiability*
If $\mathbf{f} : E \to \mathbb{R}^m$ is differentiable at $\mathbf{x} \in E$, then $D\mathbf{f}(\mathbf{x})$ is unique and $\mathbf{f}$ is continuous at $\mathbf{x}$.

Proof. If A_1, A_2 both satisfy Definition 4.2.1, set $B = A_1 - A_2$. Then $\|B\mathbf{h}\|_2 / \|\mathbf{h}\|_2 \to 0$ as $\mathbf{h} \to \mathbf{0}$.
For fixed $\mathbf{h} \neq \mathbf{0}$, applying this to $t\mathbf{h}$ and using linearity makes the quotient independent of t, so $B\mathbf{h} = \mathbf{0}$.
Hence $B = 0$. For continuity, write $\mathbf{f}(\mathbf{x}+\mathbf{h}) - \mathbf{f}(\mathbf{x}) = D\mathbf{f}(\mathbf{x})\mathbf{h} + \mathbf{r}(\mathbf{h})$.
The operator-norm estimate gives $\|D\mathbf{f}(\mathbf{x})\mathbf{h}\|_2 \le \|D\mathbf{f}(\mathbf{x})\|_{\mathrm{op}} \|\mathbf{h}\|_2$, and remainder is $o(\|\mathbf{h}\|_2)$, so $\mathbf{f}(\mathbf{x}+\mathbf{h}) \to \mathbf{f}(\mathbf{x})$.
Structural Commentary: The derivative is a rigid first-order linear part.
Key Insight: A remainder smaller than $\|\mathbf{h}\|$ cannot hide a second linear map. □

Example 4.2.4. *Derivative of a Linear Map*
If $A \in \mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$ and $\mathbf{f}(\mathbf{x}) = A\mathbf{x}$, then $D\mathbf{f}(\mathbf{x}) = A$ for every $\mathbf{x} \in \mathbb{R}^n$.
The remainder in Definition 4.2.1 is identically zero.

Theorem 4.2.5. *Chain Rule*
If $\mathbf{f} : E \subseteq \mathbb{R}^n \to \mathbb{R}^m$ is differentiable at $\mathbf{x}_0$, $\mathbf{g} : G \subseteq \mathbb{R}^m \to \mathbb{R}^k$ is differentiable at $\mathbf{f}(\mathbf{x}_0)$, and $\mathbf{f}(E) \subseteq G$, then $\mathbf{g} \circ \mathbf{f}$ is differentiable at $\mathbf{x}_0$ and $D(\mathbf{g} \circ \mathbf{f})(\mathbf{x}_0) = D\mathbf{g}(\mathbf{f}(\mathbf{x}_0))\, D\mathbf{f}(\mathbf{x}_0)$.

Proof. Let $A = D\mathbf{f}(\mathbf{x}_0)$, $B = D\mathbf{g}(\mathbf{f}(\mathbf{x}_0))$, $\mathbf{u}(\mathbf{h}) = \mathbf{f}(\mathbf{x}_0+\mathbf{h}) - \mathbf{f}(\mathbf{x}_0) - A\mathbf{h}$, and $\mathbf{v}(\mathbf{k}) = \mathbf{g}(\mathbf{f}(\mathbf{x}_0)+\mathbf{k}) - \mathbf{g}(\mathbf{f}(\mathbf{x}_0)) - B\mathbf{k}$.
For $\mathbf{k} = \mathbf{f}(\mathbf{x}_0+\mathbf{h}) - \mathbf{f}(\mathbf{x}_0)$, differentiability of $\mathbf{f}$ gives $\mathbf{k} \to \mathbf{0}$ and $(\mathbf{g}\circ\mathbf{f})(\mathbf{x}_0+\mathbf{h}) - (\mathbf{g}\circ\mathbf{f})(\mathbf{x}_0) - BA\mathbf{h} = B\mathbf{u}(\mathbf{h}) + \mathbf{v}(\mathbf{k})$, where $\|\mathbf{u}(\mathbf{h})\|_2 = o(\|\mathbf{h}\|_2)$ and $\|\mathbf{v}(\mathbf{k})\|_2 = o(\|\mathbf{k}\|_2) = o(\|\mathbf{h}\|_2)$. Divide by $\|\mathbf{h}\|_2$.
Structural Commentary: The first-order parts compose as linear maps.
Key Insight: The output increment of $\mathbf{f}$ is small enough to feed into the remainder estimate for $\mathbf{g}$. □

Theorem 4.2.6. *Componentwise Differentiability*
Let $\mathbf{f} = (f_1, \ldots, f_m) : E \subseteq \mathbb{R}^n \to \mathbb{R}^m$ and let $\pi_i : \mathbb{R}^m \to \mathbb{R}$ be the ith coordinate projection. Then $\mathbf{f}$ is differentiable at $\mathbf{x} \in E$ iff each f_i is differentiable at $\mathbf{x}$, and $Df_i(\mathbf{x}) = \pi_i \circ D\mathbf{f}(\mathbf{x})$ when these derivatives exist.

Proof.

($\Rightarrow$) If $\mathbf{f}$ is differentiable, write $\mathbf{f}(\mathbf{x}+\mathbf{h}) - \mathbf{f}(\mathbf{x}) = D\mathbf{f}(\mathbf{x})\mathbf{h} + \mathbf{r}(\mathbf{h})$ with $\|\mathbf{r}(\mathbf{h})\|_2 = o(\|\mathbf{h}\|_2)$.
Applying π_i gives $f_i(\mathbf{x}+\mathbf{h}) - f_i(\mathbf{x}) = \pi_i(D\mathbf{f}(\mathbf{x})\mathbf{h}) + \pi_i(\mathbf{r}(\mathbf{h}))$ and $|\pi_i(\mathbf{r}(\mathbf{h}))| \le \|\mathbf{r}(\mathbf{h})\|_2 = o(\|\mathbf{h}\|_2)$.
Thus $Df_i(\mathbf{x}) = \pi_i \circ D\mathbf{f}(\mathbf{x})$.

($\Leftarrow$) Suppose $Df_i(\mathbf{x}) = A_i \in \mathcal{L}(\mathbb{R}^n, \mathbb{R})$ for each i. Define $A \in \mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$ by $A\mathbf{h} = (A_1\mathbf{h}, \ldots, A_m\mathbf{h})$.
With $r_i(\mathbf{h}) = f_i(\mathbf{x}+\mathbf{h}) - f_i(\mathbf{x}) - A_i\mathbf{h}$, each $r_i(\mathbf{h}) = o(\|\mathbf{h}\|_2)$.
So $\mathbf{r}(\mathbf{h}) = (r_1(\mathbf{h}), \ldots, r_m(\mathbf{h}))$ satisfies $\|\mathbf{r}(\mathbf{h})\|_2 \le \sum_i |r_i(\mathbf{h})| = o(\|\mathbf{h}\|_2)$. Definition 4.2.1 gives $D\mathbf{f}(\mathbf{x}) = A$.

Structural Commentary: Component derivatives assemble into one linear map.
Key Insight: Components test the derivative; the total derivative packages the tests together. □

Definition 4.2.7. *Partial, Gradient, and Directional Derivatives*
Let $E \subseteq \mathbb{R}^n$ be open, $\mathbf{f} = (f_1, \dots, f_m) : E \to \mathbb{R}^m$, and let $\mathbf{e}_1, \dots, \mathbf{e}_n$ be the standard basis of $\mathbb{R}^n$.
If indicated limits exist, *partial derivative* is $(D_j f_i)(\mathbf{x}) := \lim_{t\to 0} \frac{f_i(\mathbf{x}+t\mathbf{e}_j)-f_i(\mathbf{x})}{t}$ for $1 \le i \le m,\ 1 \le j \le n$.
Directional derivative is $(D_{\mathbf{u}} f)(\mathbf{x}) := \lim_{t\to 0} \frac{f(\mathbf{x}+t\mathbf{u})-f(\mathbf{x})}{t}$ for $\mathbf{u} \in \mathbb{R}^n$; when only the direction is meant, one often assumes $\|\mathbf{u}\|_2 = 1$. If $m = 1$ and all partials exist, the *gradient* is $(\nabla f)(\mathbf{x}) := \sum_{j=1}^{n} (D_j f)(\mathbf{x})\mathbf{e}_j$.

Perspective 4.2.8. *Partials Are Coordinate Tests*
Partials and directional derivatives test selected lines; the total derivative controls all small increments by one linear map. Theorem 4.2.9 extracts partials, and Theorem 4.2.14 gives a continuity-based converse.

Theorem 4.2.9. *Matrix of the Total Derivative*
If $\mathbf{f} : E \to \mathbb{R}^m$ is differentiable at $\mathbf{x}$, then every $(D_j f_i)(\mathbf{x})$ exists and, in standard bases,

$$[D\mathbf{f}(\mathbf{x})] = \begin{bmatrix} (D_1 f_1)(\mathbf{x}) & \cdots & (D_n f_1)(\mathbf{x}) \\ \vdots & & \vdots \\ (D_1 f_m)(\mathbf{x}) & \cdots & (D_n f_m)(\mathbf{x}) \end{bmatrix}$$

If $m = 1$ and f is differentiable at $\mathbf{x}$, then $(D_{\mathbf{u}} f)(\mathbf{x}) = (\nabla f)(\mathbf{x}) \cdot \mathbf{u}$ for every $\mathbf{u} \in \mathbb{R}^n$.

Proof. Set $\mathbf{h} = t\mathbf{e}_j$ in Definition 4.2.1; limit is $D\mathbf{f}(\mathbf{x})\mathbf{e}_j$, so components form column j (Proposition 1.3.85).
For $m = 1$, set $\mathbf{h} = t\mathbf{u}$, identify linear functional with dot product against $\nabla f(\mathbf{x})$ (Theorem 2.1.57).
Structural Commentary: Partials are coordinate columns of $D\mathbf{f}(\mathbf{x})$.
Key Insight: Differentiability controls all directions at once; partial derivatives record its coordinate columns. □

Remark 4.2.10. *Row and Column Conventions*
With the coordinate-column convention from Proposition 1.3.85, the Jacobian matrix above is the standard-basis matrix of the linear map $D\mathbf{f}(\mathbf{x})$, and its jth column is $D\mathbf{f}(\mathbf{x})\mathbf{e}_j$.

Theorem 4.2.11. *Mean-Value Inequality on Convex Sets*
Let $E \subseteq \mathbb{R}^n$ be open and convex.
If $\mathbf{f} : E \to \mathbb{R}^m$ is differentiable and $\|D\mathbf{f}\|_{op} \le M$ on E, then for all $\mathbf{a}, \mathbf{b} \in E$, $\|\mathbf{f}(\mathbf{b}) - \mathbf{f}(\mathbf{a})\|_2 \le M \|\mathbf{b} - \mathbf{a}\|_2$.
If $D\mathbf{f} = 0$ on E, then $\mathbf{f}$ is constant.

Proof. Fix $\mathbf{a}, \mathbf{b} \in E$, set $\gamma(t) = (1-t)\mathbf{a} + t\mathbf{b}$, and let $\mathbf{g} = \mathbf{f} \circ \gamma$. Then $\gamma([0,1]) \subseteq E$ and $\|\mathbf{g}'(t)\|_2 = \|D\mathbf{f}(\gamma(t))(\mathbf{b} - \mathbf{a})\|_2 \le M \|\mathbf{b} - \mathbf{a}\|_2$. Apply Theorem 4.1.30; the constant case is $M = 0$.
Structural Commentary: Restricting to a line segment reduces the estimate to one variable.
Key Insight: A uniform bound on the derivative bounds every secant displacement. □

Definition 4.2.12. *C^1 Mapping*
Let $E \subseteq \mathbb{R}^n$ be open. A differentiable map $\mathbf{f} : E \to \mathbb{R}^m$ is *continuously differentiable*, or C^1, if $\mathbf{x} \mapsto D\mathbf{f}(\mathbf{x})$ is continuous from E into $\mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$ in operator norm.

Perspective 4.2.13. *C^1 Uses the Operator-Norm Metric*
C^1 continuity measures $\mathbf{x} \mapsto D\mathbf{f}(\mathbf{x})$ in operator norm. In coordinates this is spectral-norm continuity, equivalent here to continuity of the derivative matrix entries.

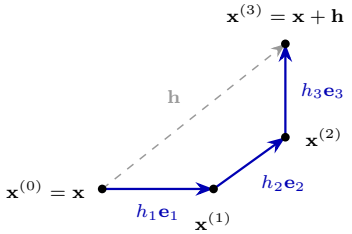


Figure 179: Coordinate increments telescope $\mathbf{x}$ to $\mathbf{x} + \mathbf{h}$.

Theorem 4.2.14. *C^1 Means Continuous Partials*
A map $\mathbf{f} : E \subseteq \mathbb{R}^n \to \mathbb{R}^m$ is C^1 iff all $D_j f_i$ exist and are continuous on E.

Proof. If $\mathbf{f}$ is C^1, Theorem 4.2.9 makes each $D_j f_i$ a continuous matrix entry of $D\mathbf{f}$.
Conversely reduce componentwise to $m = 1$. Fix $\mathbf{x} \in E$ and take $\mathbf{h}$ small enough that the coordinate polygon from $\mathbf{x}$ to $\mathbf{x}+\mathbf{h}$ lies in E. Put $\mathbf{x}^{(0)} = \mathbf{x}$ and $\mathbf{x}^{(j)} = \mathbf{x} + h_1\mathbf{e}_1 + \cdots + h_j\mathbf{e}_j$. By one-variable MVT on each coordinate segment, $f(\mathbf{x}+\mathbf{h}) - f(\mathbf{x}) = \sum_{j=1}^{n} h_j (D_j f)(\xi_j)$ for points ξ_j on corresponding segments. Continuity of partials at $\mathbf{x}$ gives $\left| f(\mathbf{x}+\mathbf{h}) - f(\mathbf{x}) - \sum_{j=1}^{n} h_j (D_j f)(\mathbf{x}) \right| \le \sum_{j=1}^{n} |h_j| \, |(D_j f)(\xi_j) - (D_j f)(\mathbf{x})| = o(\|\mathbf{h}\|_2)$.
Hence $Df(\mathbf{x})$ has row matrix $[(D_1 f)(\mathbf{x}) \ \cdots \ (D_n f)(\mathbf{x})]$. Since operator norm is controlled by the Euclidean norm of the matrix entries, continuity of the partials gives continuity of $\mathbf{x} \mapsto Df(\mathbf{x})$.
Structural Commentary: Continuous partial derivatives assemble into a continuous total derivative.
Key Insight: Coordinate-by-coordinate increments work only because the partials vary continuously. □

Higher derivatives repeat the same idea: once $D\mathbf{f}$ is itself a map into the finite-dimensional vector space $\mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$, we differentiate it again. The linear-algebra viewpoint matters here because the second derivative is naturally bilinear, not merely an array of second partial derivatives.

Definition 4.2.15. *Higher Derivatives and C^r*
Let $E \subseteq \mathbb{R}^n$ be open, $\mathbf{f} : E \to \mathbb{R}^m$, $\mathcal{L}^r(\mathbb{R}^n, \mathbb{R}^m)$ be space of r-linear maps $(\mathbb{R}^n)^r \to \mathbb{R}^m$. Set $D^1\mathbf{f} := D\mathbf{f}$. Recursively, if $D^{r-1}\mathbf{f} : E \to \mathcal{L}^{r-1}(\mathbb{R}^n, \mathbb{R}^m)$ differentiable at $\mathbf{x}$, let $D^r\mathbf{f}(\mathbf{x})(\mathbf{v}_1, \ldots, \mathbf{v}_r) := \big(D(D^{r-1}\mathbf{f})(\mathbf{x})\mathbf{v}_1\big)(\mathbf{v}_2, \ldots, \mathbf{v}_r)$.
For real-valued f, write $D_{ij} f := D_i(D_j f)$ when defined. The map $\mathbf{f}$ is C^r if derivatives through order r exist and $D^r\mathbf{f} : E \to \mathcal{L}^r(\mathbb{R}^n, \mathbb{R}^m)$ is continuous; it is C^∞ if it is C^r for all r.

Theorem 4.2.16. *Second Derivatives and Second Partials*
Let $\mathbf{f} = (f_1, \ldots, f_m) : E \subseteq \mathbb{R}^n \to \mathbb{R}^m$ and $p \in E$.
(⇒) *If $D^2\mathbf{f}(p)$ exists, then each $D_{ij} f_\ell(p)$ exists and $D^2 f_\ell(p)(\mathbf{e}_i, \mathbf{e}_j) = D_{ij} f_\ell(p)$.*
(⇐) *If all $D_{ij} f_\ell$ exist near p and are continuous at p, then $D^2\mathbf{f}(p)$ exists.*

Proof.
(⇒) Apply Theorem 4.2.6 to $D\mathbf{f} : E \to \mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$. The entries of $D\mathbf{f}$ are $D_j f_\ell$ by Theorem 4.2.9; differentiating the (ℓ, j) entry in the $\mathbf{e}_i$ direction gives $D_{ij} f_\ell(p)$ and the displayed formula.
(⇐) Matrix entries of $D\mathbf{f}$ are functions $D_j f_\ell$. Since first partials $D_i(D_j f_\ell) = D_{ij} f_\ell$ exist near p and are continuous at p, coordinate criterion at p gives differentiability of $D\mathbf{f}$ there. Hence $D^2\mathbf{f}(p)$ exists.
Structural Commentary: Second partials are matrix entries of the derivative of the first-derivative matrix.
Key Insight: Continuous second partials make the first-derivative matrix differentiable. □

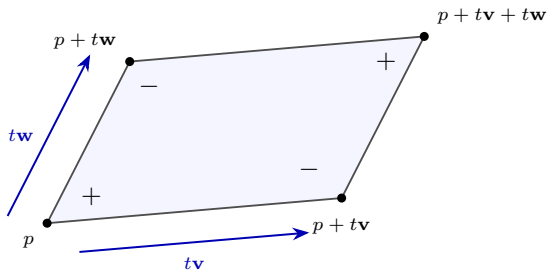


Figure 180: The signed vertex sum on a small parallelogram is symmetric in $\mathbf{v}$ and $\mathbf{w}$.

Theorem 4.2.17. *Symmetry of the Second Derivative*
If $D^2\mathbf{f}(p)$ exists, then $D^2\mathbf{f}(p)(\mathbf{v}, \mathbf{w}) = D^2\mathbf{f}(p)(\mathbf{w}, \mathbf{v})$ for all $\mathbf{v}, \mathbf{w} \in \mathbb{R}^n$.

Proof. Reduce to scalar f by Theorem 4.2.6.
For small t, set $\Delta(t, \mathbf{v}, \mathbf{w}) := f(p + t\mathbf{v} + t\mathbf{w}) - f(p + t\mathbf{v}) - f(p + t\mathbf{w}) + f(p)$.
Then $\Delta(t, \mathbf{v}, \mathbf{w}) = \Delta(t, \mathbf{w}, \mathbf{v})$; see Figure 180. The one-variable mean value theorem and differentiability of Df at p give $D^2 f(p)(\mathbf{v}, \mathbf{w}) = \lim_{t \to 0} \frac{\Delta(t, \mathbf{v}, \mathbf{w})}{t^2}$. The same formula with $\mathbf{v}, \mathbf{w}$ reversed proves symmetry.
Structural Commentary: A pointwise second derivative is already forced to be a symmetric bilinear form.
Key Insight: The parallelogram difference is symmetric in the two directions. □

Corollary 4.2.18. *Mixed Partials from D^2*
If $D^2\mathbf{f}(p)$ exists, then $D_{ij} f_\ell(p) = D_{ji} f_\ell(p)$ for all components and all i, j.

Proof. By Theorem 4.2.16, $D_{ij}f_\ell(p) = D^2f_\ell(p)(\mathbf{e}_i, \mathbf{e}_j)$; and $D_{ji}f_\ell(p) = D^2f_\ell(p)(\mathbf{e}_j, \mathbf{e}_i)$.
Theorem 4.2.17 makes the two right-hand sides equal.
Structural Commentary: Mixed partials are coordinate values of one bilinear map.
Key Insight: Symmetry of $D^2f_\ell(p)$ gives equality of mixed entries. □

Corollary 4.2.19. *Symmetry of Higher Derivatives*
If $D^r\mathbf{f}(p)$ exists, then the r-linear map $D^r\mathbf{f}(p)$ is unchanged by permuting its arguments.

Proof. It suffices to prove invariance under adjacent transpositions. Fix all arguments except positions $k, k+1$. Resulting two-variable second derivative is $(\mathbf{u}, \mathbf{v}) \longmapsto D^r\mathbf{f}(p)(\mathbf{v}_1, \dots, \mathbf{v}_{k-1}, \mathbf{u}, \mathbf{v}, \mathbf{v}_{k+2}, \dots, \mathbf{v}_r)$, so Theorem 4.2.17 allows $\mathbf{u}$ and $\mathbf{v}$ to be swapped. Since adjacent transpositions generate all permutations, $D^r\mathbf{f}(p)$ is symmetric.
Structural Commentary: Higher derivatives are iterated second-derivative symmetries.
Key Insight: Adjacent swaps generate every permutation. □

The next rectangle argument is the classical coordinate route to equality of mixed partials. It is weaker in hypothesis than global C^2, but it proves the same compatibility of coordinate second derivatives.

Theorem 4.2.20. *Rectangle Mean Value Formula*
Let $E \subseteq \mathbb{R}^2$ be open, and suppose D_1f and $D_{21}f$ exist on E. For a closed coordinate rectangle $Q \subseteq E$ with opposite vertices (a, b) and $(a+h, b+k)$, $h, k \neq 0$, set $\Delta(f, Q) := f(a+h, b+k) - f(a+h, b) - f(a, b+k) + f(a, b)$. Then some $(x, y) \in \operatorname{int} Q$ satisfies $\Delta(f, Q) = hk(D_{21}f)(x, y)$.

Proof. Set $u(t) = f(t, b+k) - f(t, b)$.
By Corollary 4.1.13, for some $x \in (a, a+h)$, $\Delta(f, Q) = u(a+h) - u(a) = h\big[(D_1f)(x, b+k) - (D_1f)(x, b)\big]$.
Apply same corollary to $s \mapsto (D_1f)(x, s)$; for some $y \in (b, b+k)$, the bracket equals $k(D_{21}f)(x, y)$.
Hence $(x, y) \in \operatorname{int} Q$ and the formula follows.
Structural Commentary: The rectangle difference is a two-step one-variable mean value calculation.
Key Insight: Mixed second derivatives measure the first-order change of a first partial across the other coordinate. □

Theorem 4.2.21. *Equality of Mixed Partials*
Let (a, b) lie in an open set $E \subseteq \mathbb{R}^2$. If D_1f, D_2f, and $D_{21}f$ exist on E, and $D_{21}f$ is continuous at (a, b), then $D_{12}f(a, b)$ exists and $(D_{12}f)(a, b) = (D_{21}f)(a, b)$. Hence $D_{12}f = D_{21}f$ when $f \in C^2(E)$.

Proof. Let $A = (D_{21}f)(a, b)$.
Given $\varepsilon > 0$, continuity of $D_{21}f$ and Theorem 4.2.20 give, for all small $h, k \neq 0$, $\left|\frac{\Delta(f,Q)}{hk} - A\right| < \varepsilon$.
Fix such an h and let $k \to 0$. Since D_2f exists at $(a+h, b)$ and (a, b), $\frac{\Delta(f,Q)}{hk} \to \frac{(D_2f)(a+h,b) - (D_2f)(a,b)}{h}$.
Thus this last quotient differs from A by at most ε for all sufficiently small $h \neq 0$.
Since ε was arbitrary, the quotient tends to A as $h \to 0$, so $D_{12}f(a, b)$ exists and equals A.
Structural Commentary: Continuity of $D_{21}f$ controls local rectangle quotients.
Key Insight: The missing mixed partial is recovered by taking the two one-variable limits in the reverse order. □

Perspective 4.2.22. *Mixed Partials Need Hypotheses*
$D_{12}f$ and $D_{21}f$ are different iterated limits. Theorems 4.2.17 and 4.2.21 give hypotheses for equality.

Theorem 4.2.23. *Differentiation Under a Riemann–Stieltjes Integral*
Let α be increasing on $[a, b]$, let $\varphi : [a, b] \times [c, d] \to \mathbb{R}$, and put $\varphi^t(x) := \varphi(x, t)$.
Assume $\varphi^t \in \mathcal{R}(\alpha)$ for all t and $D_2\varphi(\cdot, t) \rightrightarrows D_2\varphi(\cdot, s)$ on $[a, b]$ as $t \to s \in (c, d)$.
If $F(t) := \int_a^b \varphi(x, t)\, d\alpha(x)$, then $(D_2\varphi)^s \in \mathcal{R}(\alpha)$ and $F'(s) = \int_a^b (D_2\varphi)(x, s)\, d\alpha(x)$.

Proof. For $t \neq s$, set $\psi_t = (\varphi^t - \varphi^s)/(t - s)$. Corollary 4.1.13 and the uniform hypothesis give $\psi_t \rightrightarrows (D_2\varphi)^s$.
Since $\psi_t \in \mathcal{R}(\alpha)$, Theorem 3.2.16 gives $\frac{F(t)-F(s)}{t-s} = \int_a^b \psi_t\, d\alpha \longrightarrow \int_a^b (D_2\varphi)(x, s)\, d\alpha(x)$.
Structural Commentary: Uniform parameter control lets limits pass under the integral.
Key Insight: The theorem is exactly the commutation of a derivative limit with a Riemann–Stieltjes integral. □

Example 4.2.24. *A Gaussian Fourier Integral*
Define $f(t) = \int_{-\infty}^{\infty} e^{-x^2} \cos(xt)\, dx$ and $g(t) = -\int_{-\infty}^{\infty} xe^{-x^2} \sin(xt)\, dx$.
Then $f'(t) = g(t)$, and integration by parts gives $2f'(t) + tf(t) = 0$. Since $f(0) = \sqrt{\pi}$, $f(t) = \sqrt{\pi} \exp\left(-\frac{t^2}{4}\right)$.

Definition 4.2.25. *Jacobian*
If $\mathbf{f} : E \subseteq \mathbb{R}^n \to \mathbb{R}^n$ is differentiable at $\mathbf{x}$, define $J_\mathbf{f}(\mathbf{x}) := \det D\mathbf{f}(\mathbf{x}) = \frac{\partial(y_1, \dots, y_n)}{\partial(x_1, \dots, x_n)}$ when $f = (y_1, \dots, y_n)$.
By Theorem 4.1.8, $D\mathbf{f}(\mathbf{x})$ is invertible if and only if $J_\mathbf{f}(\mathbf{x}) \neq 0$. The absolute value $|J_\mathbf{f}(\mathbf{x})|$ is the local version of determinant volume scaling from Theorem 4.4.1, used again in the change-of-variables theorem.

4.3 Inverse, Implicit, and Rank Theorems

Theorem 4.3.1. *Finite-Dimensional Linear Maps Are Lipschitz*
Let $A \in \mathcal{L}(\mathbb{R}^n, \mathbb{R}^m)$. Then A is $\|A\|_{op}$-Lipschitz: $\|A\mathbf{x}\|_2 \leq \|A\|_{op} \|\mathbf{x}\|_2$ for $\mathbf{x} \in \mathbb{R}^n$.
If $A \in \mathcal{L}(\mathbb{R}^n)$ is invertible, then $\frac{1}{\|A^{-1}\|_{op}} \|\mathbf{x}\|_2 \leq \|A\mathbf{x}\|_2 \leq \|A\|_{op} \|\mathbf{x}\|_2$, so A is a homeomorphism of $\mathbb{R}^n$.

Proof. The upper estimate follows from Definition 4.3.4: apply it to $\mathbf{x}/\|\mathbf{x}\|_2$ when $\mathbf{x} \neq \mathbf{0}$, while $\mathbf{x} = \mathbf{0}$ is trivial.
If A is invertible, apply the same estimate to A^{-1} with $\mathbf{y} = A\mathbf{x}$: $\|\mathbf{x}\|_2 = \|A^{-1}\mathbf{y}\|_2 \leq \|A^{-1}\|_{\text{op}} \|A\mathbf{x}\|_2$.
Thus A and A^{-1} are continuous.
Structural Commentary: Linear maps have finite stretch; invertible ones have positive lower stretch.
Key Insight: The operator norm turns linear algebra into uniform metric estimates. □

Theorem 4.3.2. *Invertible Operators Form an Open Set*
Let $\Omega \subseteq \mathcal{L}(\mathbb{R}^n)$ be the set of invertible linear operators. If $A \in \Omega$, $B \in \mathcal{L}(\mathbb{R}^n)$, $\|B - A\|_{op} \|A^{-1}\|_{op} < 1$, then $B \in \Omega$. Consequently, Ω is open in $\mathcal{L}(\mathbb{R}^n)$ and the map $A \mapsto A^{-1}$ is continuous on Ω.

Proof. Put $\alpha = 1/\|A^{-1}\|_{\text{op}}$ and $\beta = \|B - A\|_{\text{op}}$.
For every $\mathbf{x} \in \mathbb{R}^n$, $\alpha \|\mathbf{x}\|_2 \leq \|A\mathbf{x}\|_2 \leq \|(A-B)\mathbf{x}\|_2 + \|B\mathbf{x}\|_2 \leq \beta \|\mathbf{x}\|_2 + \|B\mathbf{x}\|_2$.
Since $\beta < \alpha$, this shows $B\mathbf{x} \neq 0$ for $\mathbf{x} \neq 0$; hence B is injective and therefore invertible by Theorem 1.3.49.
Replacing $\mathbf{x}$ by $B^{-1}\mathbf{y}$ gives $\|B^{-1}\|_{\text{op}} \leq (\alpha - \beta)^{-1}$.
Finally, $B^{-1} - A^{-1} = B^{-1}(A - B)A^{-1}$, Proposition 4.3.3 gives the continuity of inversion.
Structural Commentary: Small operator-norm perturbations of an invertible map remain invertible.
Key Insight: The inverse changes continuously because the difference of inverses factors through $A - B$. □

Perspective 4.3.3. *Linear Inverse Estimate*
For invertible $A \in \mathcal{L}(\mathbb{R}^n)$, Theorem 4.3.1 gives $c\|\mathbf{h}\|_2 \leq \|A\mathbf{h}\|_2 \leq \|A\|_{\text{op}} \|\mathbf{h}\|_2$; and $c := \frac{1}{\|A^{-1}\|_{\text{op}}}$.
The inverse function theorem preserves this lower stretch under small derivative changes.

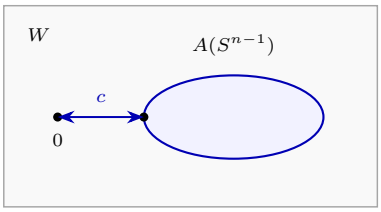


Figure 181: Invertibility keeps $A(S^{n-1})$ a positive distance from 0.

Theorem 4.3.4. *Inverse Function Theorem*
Let $\mathbf{f} : E \subseteq \mathbb{R}^n \to \mathbb{R}^n$ be C^1, let $\mathbf{a} \in E$, and suppose $D\mathbf{f}(\mathbf{a})$ is invertible. With $\mathbf{b} = \mathbf{f}(\mathbf{a})$, $\mathbf{f}$ restricts to a bijection $U \to V$ between open neighborhoods of $\mathbf{a}$ and $\mathbf{b}$; inverse $\mathbf{g} : V \to U$ is C^1, and $D\mathbf{g}(\mathbf{y}) = \big(D\mathbf{f}(\mathbf{g}(\mathbf{y}))\big)^{-1}$, $\mathbf{y} \in V$.

Proof. Put $A = D\mathbf{f}(\mathbf{a})$ and choose $0 < \theta < 1$.
By continuity of $D\mathbf{f}$, choose $r > 0$ with $\overline{B}_r(\mathbf{a}) \subset E$ and $\|I - A^{-1}D\mathbf{f}(\mathbf{x})\|_{\text{op}} \leq \theta$ on this closed ball.
For fixed $\mathbf{y}$, solutions of $\mathbf{f}(\mathbf{x}) = \mathbf{y}$ are fixed points of $\varphi_{\mathbf{y}}(\mathbf{x}) = \mathbf{x} + A^{-1}(\mathbf{y} - \mathbf{f}(\mathbf{x}))$.
Theorem 4.2.11 gives $\|\varphi_{\mathbf{y}}(\mathbf{x}) - \varphi_{\mathbf{y}}(\mathbf{z})\|_2 \leq \theta \|\mathbf{x} - \mathbf{z}\|_2$. If $\|A^{-1}(\mathbf{y} - \mathbf{b})\|_2 \leq (1 - \theta)r$, then $\varphi_{\mathbf{y}}$ maps $\overline{B}_r(\mathbf{a})$ into itself. Since $\overline{B}_r(\mathbf{a})$ is a closed subset of complete $\mathbb{R}^n$, Proposition 2.2.14 makes it complete.
The contraction principle, Theorem 2.2.19, gives a unique $\mathbf{x} \in \overline{B}_r(\mathbf{a})$ with $\mathbf{f}(\mathbf{x}) = \mathbf{y}$. Thus, after shrinking to open neighborhoods U of $\mathbf{a}$ and V of $\mathbf{b}$, $\mathbf{f} : U \to V$ is bijective with inverse $\mathbf{g}$.
For $\mathbf{y} = \mathbf{f}(\mathbf{x})$ and $\mathbf{y} + \mathbf{k} = \mathbf{f}(\mathbf{x} + \mathbf{h})$ in V, the same derivative bound gives $\|\mathbf{h}\|_2 \leq C\|\mathbf{k}\|_2$ for a local constant C. Moreover $\mathbf{g}(\mathbf{y} + \mathbf{k}) - \mathbf{g}(\mathbf{y}) - \big(D\mathbf{f}(\mathbf{x})\big)^{-1}\mathbf{k} = -\big(D\mathbf{f}(\mathbf{x})\big)^{-1}\big[\mathbf{f}(\mathbf{x} + \mathbf{h}) - \mathbf{f}(\mathbf{x}) - D\mathbf{f}(\mathbf{x})\mathbf{h}\big]$.
Since $\mathbf{f}$ is differentiable at $\mathbf{x}$ and $\|\mathbf{h}\|_2 = O(\|\mathbf{k}\|_2)$, division by $\|\mathbf{k}\|_2$ gives $D\mathbf{g}(\mathbf{y}) = (D\mathbf{f}(\mathbf{x}))^{-1}$.
Continuity of $D\mathbf{g}$ follows from continuity of $D\mathbf{f}$, continuity of $\mathbf{g}$, and Theorem 4.3.2.
Structural Commentary: Local inversion is a contraction built from A^{-1}.
Key Insight: Small derivative variation preserves local invertibility. □

Corollary 4.3.5. *Local Diffeomorphisms Are Open*
If $\mathbf{f} : E \subseteq \mathbb{R}^n \to \mathbb{R}^n$ is C^1 and $D\mathbf{f}$ is everywhere invertible, then $\mathbf{f}$ is an open map (locally one-to-one).

Proof. Let $W \subseteq E$ be open and let $\mathbf{y} \in \mathbf{f}(W)$. Choose $\mathbf{x} \in W$ with $\mathbf{f}(\mathbf{x}) = \mathbf{y}$. Applying Theorem 4.3.4 to $\mathbf{f}|_W$ at $\mathbf{x}$ gives open neighborhoods $U \subseteq W$ of $\mathbf{x}$ and V of $\mathbf{y}$ with $\mathbf{f}(U) = V$. Hence $V \subseteq \mathbf{f}(W)$, so $\mathbf{y}$ is interior to $\mathbf{f}(W)$. Since $\mathbf{y}$ was arbitrary, $\mathbf{f}(W)$ is open. Local one-to-one-ness is the same theorem applied pointwise.
Structural Commentary: Openness is checked one image point at a time.
Key Insight: A local inverse turns nearby domain points into a full image neighborhood. □

Definition 4.3.6. *Block Splitting of a Linear Map*
For $A \in \mathcal{L}(\mathbb{R}^{n+m}, \mathbb{R}^n)$ and $(\mathbf{h}, \mathbf{k}) \in \mathbb{R}^n \times \mathbb{R}^m$, define $A_x\mathbf{h} := A(\mathbf{h}, \mathbf{0})$ and $A_y\mathbf{k} := A(\mathbf{0}, \mathbf{k})$.
Then $A_x \in \mathcal{L}(\mathbb{R}^n)$, $A_y \in \mathcal{L}(\mathbb{R}^m, \mathbb{R}^n)$, and $A(\mathbf{h}, \mathbf{k}) = A_x\mathbf{h} + A_y\mathbf{k}$.

Theorem 4.3.7. *Linear Implicit Function Theorem*
If $A \in \mathcal{L}(\mathbb{R}^{n+m}, \mathbb{R}^n)$, A_x invertible, then $\forall \mathbf{k} \in \mathbb{R}^m$ $\exists$ unique $\mathbf{h} \in \mathbb{R}^n$ s.t. $A(\mathbf{h}, \mathbf{k}) = 0$, that $\mathbf{h} = -(A_x)^{-1}A_y\mathbf{k}$.

Proof. Fix $\mathbf{k} \in \mathbb{R}^m$ and set $\mathbf{h}_0 := -(A_x)^{-1}A_y\mathbf{k}$.
By Definition 4.3.6, $A(\mathbf{h}_0, \mathbf{k}) = A_x\mathbf{h}_0 + A_y\mathbf{k} = -A_y\mathbf{k} + A_y\mathbf{k} = 0$, so a solution exists.
If $\mathbf{h} \in \mathbb{R}^n$ also satisfies $A(\mathbf{h}, \mathbf{k}) = 0$, then $A_x\mathbf{h} + A_y\mathbf{k} = 0$, hence $A_x\mathbf{h} = -A_y\mathbf{k}$.
Applying $(A_x)^{-1}$ gives $\mathbf{h} = -(A_x)^{-1}A_y\mathbf{k} = \mathbf{h}_0$. Thus the solution exists and is unique.
Structural Commentary: Splitting A separates the unknown variable from the parameter.
Key Insight: Invertibility of A_x turns constraint $A(\mathbf{h}, \mathbf{k}) = 0$ into a uniquely solvable linear equation for $\mathbf{h}$. □

The implicit theorem below is the nonlinear counterpart of the block equation theorem. The invertible block is the set of variables one can solve for, just as in Theorem 4.3.7.

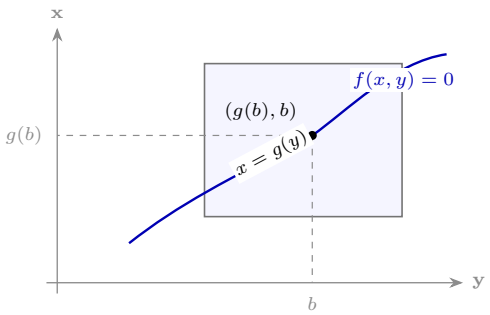


Figure 182: Near a regular point, the zero locus is the graph of an implicit function.

Theorem 4.3.8. *Implicit Function Theorem*
Let $\mathbf{f} : E \subseteq \mathbb{R}^{n+m} \to \mathbb{R}^n$ be C^1, $\mathbf{f}(\mathbf{a}, \mathbf{b}) = 0$, $A = D\mathbf{f}(\mathbf{a}, \mathbf{b})$, and let A_x, A_y be the block maps of Definition 4.3.6. If A_x is invertible, then near $\mathbf{b}$, the equation $\mathbf{f}(\mathbf{x}, \mathbf{y}) = 0$ has unique local solution $\mathbf{x} = \mathbf{g}(\mathbf{y})$, where $\mathbf{g}$ is C^1, $\mathbf{g}(\mathbf{b}) = \mathbf{a}$, $D\mathbf{g}(\mathbf{b}) = -(A_x)^{-1}A_y$.

Proof. Set $\Phi(\mathbf{x}, \mathbf{y}) = (\mathbf{f}(\mathbf{x}, \mathbf{y}), \mathbf{y})$. Then $D\Phi(\mathbf{a}, \mathbf{b})(\mathbf{h}, \mathbf{k}) = (A_x\mathbf{h} + A_y\mathbf{k}, \mathbf{k})$, whose kernel is zero since A_x is invertible; hence $D\Phi(\mathbf{a}, \mathbf{b})$ is invertible by Theorem 1.3.49. Theorem 4.3.4 makes Φ locally invertible.
Points with $\mathbf{f}(\mathbf{x}, \mathbf{y}) = 0$ are exactly those sent to $(0, \mathbf{y})$, so $\Phi^{-1}(0, \mathbf{y}) = (\mathbf{g}(\mathbf{y}), \mathbf{y})$ gives the unique local solution.
Differentiating $\mathbf{f}(\mathbf{g}(\mathbf{y}), \mathbf{y}) = 0$ at $\mathbf{b}$ gives $A_x D\mathbf{g}(\mathbf{b}) + A_y = 0$.
Structural Commentary: Add $\mathbf{y}$ as a coordinate, then invert.
Key Insight: Invertibility of the x-block is precisely the linear ability to solve for x. □

Perspective 4.3.9. *Rank Theorem Linear Model*
Theorem 1.3.41 reduces a linear map to its kernel–range split.
The constant-rank theorem is the local C^1 analogue after coordinate changes.

Theorem 4.3.10. *Rank Theorem*
Let $\mathbf{F} : E \subseteq \mathbb{R}^n \to \mathbb{R}^m$ be C^1 with constant rank r. At $\mathbf{a} \in E$, put $A = D\mathbf{F}(\mathbf{a})$, $Y_1 = \operatorname{range}(A)$, and let P be the orthogonal projection onto Y_1 (Definition 2.2.11). Locally, after a C^1 coordinate change H with C^1 inverse on a small coordinate neighborhood V_0, there is a C^1 map $\varphi : A(V_0) \subseteq Y_1 \to \ker(P)$ such that $\mathbf{F}(H(\mathbf{x})) = A\mathbf{x} + \varphi(A\mathbf{x})$. In coordinates adapted to $\ker A$ and $\operatorname{range}(A)$, this is the same local geometry as a projection onto r coordinates.

Proof. If $r = 0$, Theorem 4.2.11 makes $\mathbf{F}$ locally constant.
Assume $r > 0$. Choose $S : Y_1 \to \mathbb{R}^n$ with $AS = I_{Y_1}$, possible by choosing preimages of a basis of Y_1, and set $G(\mathbf{x}) = \mathbf{x} + SP(\mathbf{F}(\mathbf{x}) - A\mathbf{x})$. Since $DG(\mathbf{a}) = I$, Theorem 4.3.4 gives a local inverse $H = G^{-1}$.
Let $\Psi = \mathbf{F} \circ H$. From $AG = P\mathbf{F}$ we get $P\Psi(\mathbf{x}) = A\mathbf{x}$. Hence $PD\Psi(\mathbf{x}) = A$. The rank of $D\Psi(\mathbf{x})$ is r because $DH(\mathbf{x})$ is invertible and $\mathbf{F}$ has constant rank r; the rank of $PD\Psi(\mathbf{x})$ is also r.

Therefore P is injective on range $D\Psi(\mathbf{x})$, and $D\Psi(\mathbf{x})\mathbf{v} = 0$ whenever $A\mathbf{v} = 0$.
On a small convex coordinate neighborhood, any two points with the same $A\mathbf{x}$ differ by a vector in $\ker(A)$ and can be joined inside that affine fiber. The preceding derivative vanishing and Theorem 4.2.11 show that Ψ is constant on such fibers. Thus $(I - P)\Psi(\mathbf{x})$ depends only on $A\mathbf{x}$; call it $\varphi(A\mathbf{x})$. Then $\mathbf{F}(H(\mathbf{x})) = \Psi(\mathbf{x}) = A\mathbf{x} + \varphi(A\mathbf{x})$, and φ is C^1 because Ψ is C^1 and the dependence is through the coordinate $A\mathbf{x}$.
Structural Commentary: Constant rank gives a local graph normal form.
Key Insight: Same rank, same local level-set geometry. □

Rank-nullity (Theorem 1.3.27) says that the domain directions split into rank directions and fiber directions. The rank theorem says a constant-rank map has, locally, the same split after C^1 changes of coordinates.

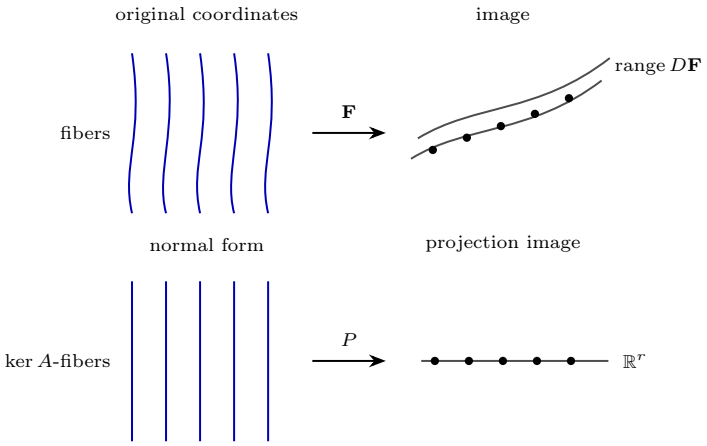


Figure 183: After local coordinate changes, a constant-rank map looks like projection onto the rank directions.

Theorem 4.3.11. *One-Constraint Lagrange Multiplier*
Let $E \subseteq \mathbb{R}^n$ be open, $f, g : E \to \mathbb{R}$ be C^1, and $S = \{\mathbf{x} : g(\mathbf{x}) = c\}$.
If $p \in S$, $\nabla g(p) \neq \mathbf{0}$, and $f|_S$ has a local extremum at p, then $\exists \lambda \in \mathbb{R}$ with $\nabla f(p) = \lambda \nabla g(p)$.

Proof. If $Df(p) = 0$, take $\lambda = 0$. Otherwise suppose $Df(p)$ is not a multiple of $Dg(p)$.
Complete $Dg(p), Df(p)$ to a dual basis $Dg(p), \quad Df(p), \quad \ell_3, \ldots, \ell_n$ by Theorem 1.5.4,
and set $\Phi(\mathbf{x}) = \big(g(\mathbf{x}) - c,\ f(\mathbf{x}) - f(p),\ \ell_3(\mathbf{x} - p), \ldots, \ell_n(\mathbf{x} - p)\big)$.
Then $D\Phi(p)$ is invertible, so Theorem 4.3.4 makes Φ locally invertible. In these coordinates, S is locally the slice with first coordinate 0, while the second coordinate is $f - f(p)$. Local invertibility gives nearby points of S with this second coordinate of both signs, contradicting extremality.
Hence $Df(p) = \lambda Dg(p)$; Riesz representation, Theorem 2.1.57, gives the gradient form.
Structural Commentary: The constraint kills first-order variation along S.
Key Insight: If Df were independent of Dg, f would be a free constraint coordinate. □

Method 4.3.12. *Several-Variable Differentiation Proof Patterns*

(i) *Linear approximation*: write $\mathbf{f}(\mathbf{x} + \mathbf{h}) - \mathbf{f}(\mathbf{x}) = D\mathbf{f}(\mathbf{x})\mathbf{h} + \mathbf{r}(\mathbf{h})$ with $\|\mathbf{r}(\mathbf{h})\|_2 = o(\|\mathbf{h}\|_2)$.
(ii) *Norm estimates*: use $\|A\mathbf{h}\|_2 \leq \|A\|_{\text{op}} \|\mathbf{h}\|_2$ and the product estimate from Proposition 4.3.3.
(iii) *Line-segment reduction*: on convex sets, restrict to $\gamma(t) = (1 - t)\mathbf{a} + t\mathbf{b}$ and apply one-variable results.
(iv) *Local inversion*: if $A = D\mathbf{f}(\mathbf{a})$ is invertible, write $D\mathbf{f}(\mathbf{x}) = A + E_{\mathbf{x}}$ with $\|E_{\mathbf{x}}\|_{\text{op}}$ small and solve by contraction (Theorem 2.2.19); linear block equations use Theorem 4.3.7.
(v) *Coordinates*: read partial derivatives as the entries of the matrix of $D\mathbf{f}(\mathbf{x})$, using Theorem 4.2.9.
(vi) *Higher derivatives*: regard $D^r\mathbf{f}(p)$ as an r-linear map; recover coordinate partials via Theorem 4.2.16.
(vii) *Rank normal forms*: for $A = D\mathbf{F}(\mathbf{a})$, split domain by $\ker A$ and image by range A; use the quotient-by-kernel picture from Theorem 1.3.41 and project with $P_{\text{range } A}$.

Key Insight: Several-variable calculus is one-variable limiting plus finite-dimensional linear algebra.

4.4 Differential Forms and Manifolds

This subsection moves from Euclidean-region integration and change of variables to differential forms, pullbacks, and oriented parametrized surfaces.

Theorem 4.4.1. *Linear Maps Scale Volume by $|\det T|$*
Let $T \in \mathcal{L}(\mathbb{R}^n)$ and let $\Omega \subseteq \mathbb{R}^n$ be a sufficiently regular measurable set. Then $\text{vol}(T(\Omega)) = |\det T| \,\text{vol}(\Omega)$.

Proof. It is enough to prove the formula first for invertible T. Let $T = SP$ by polar decomposition, where S is an isometry and $P = \sqrt{T^*T}$ is positive definite. Isometries preserve Euclidean volume, so $\operatorname{vol}(T(\Omega)) = \operatorname{vol}(P(\Omega))$. By the real spectral theorem, choose orthonormal coordinates with $P = \operatorname{diag}(\lambda_1, \dots, \lambda_n)$ and $\lambda_j > 0$. Rectangles are sent to rectangles with side lengths multiplied by λ_j, hence their volumes are multiplied by $\prod_j \lambda_j = \det P$. Finite unions of disjoint rectangles have the same scaling by additivity, and the usual inner-outer rectangle approximation of a sufficiently regular measurable set gives $\operatorname{vol}(P(\Omega)) = (\det P)\operatorname{vol}(\Omega)$.
If T is singular, then $T(\Omega)$ lies in a proper subspace after an orthogonal change of coordinates, so its n-dimensional volume is 0; also $|\det T| = 0$. Finally, Proposition 4.2.35 gives $\det P = |\det T|$ in the invertible case.
Structural Commentary: Polar decomposition reduces volume scaling to diagonal stretching.
Key Insight: The determinant is the product of independent stretch factors. □

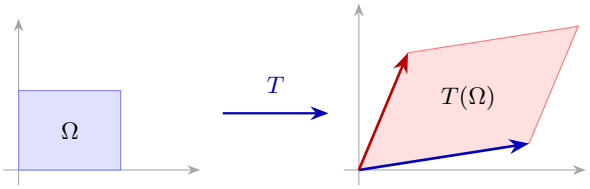


Figure 184: A linear map scales volume by $|\det T|$.

Method 4.4.2. *Change-of-Variables Jacobian Method*
For a differentiable map $\sigma : \Omega \subseteq \mathbb{R}^n \to \mathbb{R}^n$:

(i) compute the derivative matrix $D\sigma(\mathbf{x})$;
(ii) compute the local volume factor $|\det D\sigma(\mathbf{x})|$;
(iii) replace $d\mathbf{y}$ by $|\det D\sigma(\mathbf{x})|\, d\mathbf{x}$ when $\mathbf{y} = \sigma(\mathbf{x})$;
(iv) transform the domain bounds in the new coordinates.

Example 4.4.3. *Polar and Spherical Jacobian Factors*
For polar coordinates $\sigma(r, \theta) = (r\cos\theta, r\sin\theta)$, $D\sigma(r,\theta) = \begin{bmatrix} \cos\theta & -r\sin\theta \\ \sin\theta & r\cos\theta \end{bmatrix}$, and $|\det D\sigma(r,\theta)| = r$.
Thus area integrals acquire the factor $r\, dr\, d\theta$.
For spherical coordinates $\sigma(\rho,\varphi,\theta) = (\rho\sin\varphi\cos\theta, \rho\sin\varphi\sin\theta, \rho\cos\varphi)$, and $|\det D\sigma(\rho,\varphi,\theta)| = \rho^2\sin\varphi$.
Thus volume integrals acquire the factor $\rho^2\sin\varphi\, d\rho\, d\varphi\, d\theta$.

Definition 4.4.4. *k-Cells and Multiple Integrals*
A *k-cell* is $I^k := \prod_{i=1}^{k} [a_i, b_i] \subseteq \mathbb{R}^k$ or $I^j := \prod_{i=1}^{j} [a_i, b_i]$.
For $f \in C(I^k)$, set $f_k = f$ and, for $j = k, \dots, 1$, define $f_{j-1}(x_1, \dots, x_{j-1}) := \int_{a_j}^{b_j} f_j(x_1, \dots, x_j)\, dx_j$.
The *integral of f over I^k* is $\int_{I^k} f(\mathbf{x})\, d\mathbf{x} := f_0$, also written $\int_{I^k} f$.

For bounded functions on rectangles, the Darboux picture from Definition 4.1.32 generalizes by replacing intervals with grid rectangles.

Definition 4.4.5. *Rectangular Grids and Riemann Sums*
For $R = [a,b] \times [c,d] \subseteq \mathbb{R}^2$ and partitions $P : a = x_0 < \cdots < x_m = b$; and $Q : c = y_0 < \cdots < y_n = d$, put $G = P \times Q$, where $R_{ij} := [x_{i-1}, x_i] \times [y_{j-1}, y_j]$ and $|R_{ij}| := (x_i - x_{i-1})(y_j - y_{j-1})$.
For sample points $(s_{ij}, t_{ij}) \in R_{ij}$, $R(f, G, S) := \sum_{i=1}^{m} \sum_{j=1}^{n} f(s_{ij}, t_{ij})\, |R_{ij}|$. For bounded real f, with $m_{ij} := \inf_{R_{ij}} f$ and $M_{ij} := \sup_{R_{ij}} f$, set $L(f, G) := \sum_{i,j} m_{ij}\, |R_{ij}|$ and $U(f, G) := \sum_{i,j} M_{ij}\, |R_{ij}|$.
The lower and upper integrals over R are $\sup_G L(f, G)$ and $\inf_G U(f, G)$. If equal, their common value is $\int_R f$.

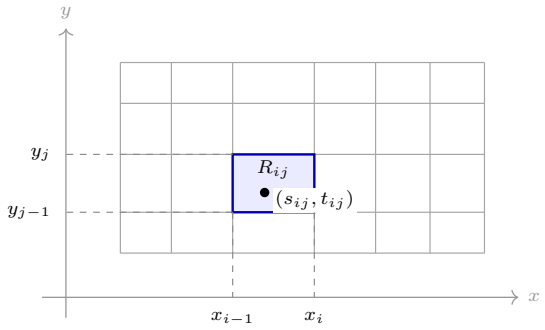


Figure 185: A grid rectangle R_{ij} and sample point (s_{ij}, t_{ij}).

Theorem 4.4.6. *Fubini on a Rectangle, Darboux Form*
Let $R = [a,b] \times [c,d]$ and let $f : R \to \mathbb{R}$ be Riemann integrable.
For $y \in [c,d]$, set $f_y(x) := f(x,y)$ and $\underline{F}(y) := \sup_P L(P, f_y)$; $\overline{F}(y) := \inf_P U(P, f_y)$, where P ranges over partitions of $[a,b]$. Then $\underline{F}, \overline{F}$ are Riemann integrable on $[c,d]$ and $\int_R f = \int_c^d \underline{F}(y)\,dy = \int_c^d \overline{F}(y)\,dy$.
If $\int_a^b f(x,y)\,dx$ exists for every y, this is the usual iterated integral formula.

Proof. For $G = P \times Q$, horizontal slice sums give $L(f,G) \le L(Q,\underline{F}) \le U(Q,\overline{F}) \le U(f,G)$.
As f is integrable on R, the outer terms can be made arbitrarily close. The upper–lower criterion, Theorem 4.1.41, squeezes $\underline{F}$ and $\overline{F}$ to the common value $\int_R f$.
Structural Commentary: Product grids turn area sums into slice sums.
Key Insight: Fubini is a squeeze between lower and upper sums before it is an equality of iterated integrals. □

Theorem 4.4.7. *Order Independence on Cells*
For $f \in C(I^k)$, all iterated orders defining $\int_{I^k} f$ give the same value.

Proof. Let L, L' be two iterated orders. For $h(\mathbf{x}) = h_1(x_1)\cdots h_k(x_k)$, $L(h) = L'(h) = \prod_{i=1}^{k} \int_{a_i}^{b_i} h_i(x_i)\,dx_i$, so $L(g) = L'(g)$ for every finite sum g of such products. These sums are dense in $C(I^k)$ by Theorem 3.4.22.
Choose g with $\|f - g\|_\infty < \varepsilon/V$, where $V = \prod_{i=1}^{k} (b_i - a_i)$.
Since $|L(r)|, |L'(r)| \le V \|r\|_\infty$, then $|L(f) - L'(f)| \le |L(f-g)| + |L'(f-g)| < 2\varepsilon$.
Structural Commentary: Product functions plus uniform approximation prove order independence.
Key Insight: Fubini-type order independence is first an algebraic fact, then a uniform approximation fact. □

Perspective 4.4.8. *Order Interchange Requires Hypotheses*
Finite sums are order-free; limiting processes are not.
Theorems 4.4.6 and 4.4.7 give hypotheses for interchanging integral order.

Definition 4.4.9. *Support and Simplex Integrals*
For $f : \mathbb{R}^k \to \mathbb{R}$ or $\mathbb{C}$, define $\operatorname{supp}(f) := \overline{\{\mathbf{x} : f(\mathbf{x}) \neq 0\}}$.
If f is continuous with compact support and $\operatorname{supp}(f) \subseteq I^k$, set $\int_{\mathbb{R}^k} f := \int_{I^k} f$. This is independent of the containing cell I^k. The *standard k-simplex* is $Q^k := \{\mathbf{x} \in \mathbb{R}^k : x_i \ge 0,\ x_1 + \cdots + x_k \le 1\}$.
For $f \in C(Q^k)$, extend f by 0 to unit cube; resulting order-independent iterated integral is denoted $\int_{Q^k} f$.

Definition 4.4.10. *Primitive Mappings and Flips*
Let $E \subseteq \mathbb{R}^n$ be open.
A map $G : E \to \mathbb{R}^n$ is *primitive* if, for some coordinate m and real function g, $G(\mathbf{x}) = \mathbf{x} + \big(g(\mathbf{x}) - x_m\big)\mathbf{e}_m$.
If g is differentiable at $\mathbf{a}$, then $DG(\mathbf{a})$ is the identity matrix with row m replaced by $\big((D_1 g)(\mathbf{a}), \ldots, (D_n g)(\mathbf{a})\big)$, hence $J_G(\mathbf{a}) = \det DG(\mathbf{a}) = (D_m g)(\mathbf{a})$ by Theorem 4.1.8.
A *flip* is the linear map that swaps two standard-basis vectors and fixes the rest.

Theorem 4.4.11. *Local Decomposition into Primitive Mappings*
Let $\mathbf{F} : E \subseteq \mathbb{R}^n \to \mathbb{R}^n$ be C^1, let $\mathbf{0} \in E$, and suppose $\mathbf{F}(\mathbf{0}) = \mathbf{0}$ and $D\mathbf{F}(\mathbf{0})$ is invertible.
Then, near $\mathbf{0}$, $\mathbf{F}(\mathbf{x}) = B_1 \cdots B_{n-1} G_n \circ \cdots \circ G_1(\mathbf{x})$, where each G_i is a primitive C^1 map with $G_i(\mathbf{0}) = \mathbf{0}$ and $DG_i(\mathbf{0})$ invertible, and each B_i is a flip or the identity.

Proof. Let $P_m \mathbf{x} := x_1\mathbf{e}_1 + \cdots + x_m\mathbf{e}_m$. Put $\mathbf{F}_1 = \mathbf{F}$ and inductively assume that $\mathbf{F}_m$ is defined near $\mathbf{0}$, $\mathbf{F}_m(\mathbf{0}) = \mathbf{0}$, $D\mathbf{F}_m(\mathbf{0})$ is invertible, and $P_{m-1}\mathbf{F}_m = P_{m-1}$. Then $\mathbf{F}_m(\mathbf{x}) = P_{m-1}\mathbf{x} + \sum_{i=m}^{n} \alpha_i(\mathbf{x})\mathbf{e}_i$.

Since $D\mathbf{F}_m(\mathbf{0})\mathbf{e}_m \neq \mathbf{0}$, some $k \in \{m, \dots, n\}$ has $(D_m\alpha_k)(\mathbf{0}) \neq 0$. Let B_m swap coordinates m and k (or be the identity if $k = m$), and set $G_m(\mathbf{x}) := \mathbf{x} + \big(\alpha_k(\mathbf{x}) - x_m\big)\mathbf{e}_m$.
By Definition 4.4.10, G_m is primitive and $DG_m(\mathbf{0})$ is invertible. The inverse function theorem, Theorem 4.3.4, gives a local inverse for G_m, so define $\mathbf{F}_{m+1} := B_m\mathbf{F}_m \circ G_m^{-1}$. Then $P_m\mathbf{F}_{m+1} = P_m$, and the induction continues.
After $m = n-1$, $\mathbf{F}_n$ is primitive with invertible derivative.
Taking $G_n = \mathbf{F}_n$ and using $B_m^2 = I$ gives $\mathbf{F} = \mathbf{F}_1 = B_1 \cdots B_{n-1}G_n \circ G_{n-1} \circ \cdots \circ G_1$ near $\mathbf{0}$.
Structural Commentary: Invertible C^1 maps factor locally into coordinate changes and flips.
Key Insight: Gaussian elimination on $D\mathbf{F}$ becomes nonlinear via the inverse function theorem. □

Theorem 4.4.12. *Partition of Unity on a Compact Set*
Let $K \subseteq \mathbb{R}^n$ be compact and let $\{V_\alpha\}$ cover K by open sets. Then there are $\psi_1, \dots, \psi_s \in C^1(\mathbb{R}^n)$ such that

(i) $0 \leq \psi_i \leq 1$ *for all i;*
(ii) $\operatorname{supp}(\psi_i) \subseteq V_\alpha$ *for some α;*
(iii) $\sum_{i=1}^s \psi_i = 1$ *on K.*

If $f \in C(\mathbb{R}^n)$ and $\operatorname{supp}(f) \subseteq K$, then $f = \sum_{i=1}^s \psi_i f$, with each summand supported in one member of the cover.

Proof.

(i) For each $\mathbf{x} \in K$, choose balls with $\overline{B(\mathbf{x})} \subseteq W(\mathbf{x}) \subseteq \overline{W(\mathbf{x})} \subseteq V_{\alpha(\mathbf{x})}$.
By compactness, Definition 2.2.1, finitely many $B(\mathbf{x}_i)$ cover K.
(ii) Choose $\varphi_i \in C^1(\mathbb{R}^n)$ with $\varphi_i = 1$ on $B(\mathbf{x}_i)$ and $\operatorname{supp}(\varphi_i) \subseteq W(\mathbf{x}_i)$. Set $\psi_i = (1-\varphi_1)\cdots(1-\varphi_{i-1})\varphi_i$.
(iii) Then $0 \leq \psi_i \leq 1$, each ψ_i is supported in some V_α, and $\sum_{i=1}^{s} \psi_i = 1 - \prod_{i=1}^{s}(1-\varphi_i)$, which equals 1 on K.
Multiplying by f gives the final assertion.

Structural Commentary: Compactness finite; product cutoffs sum to one.
Key Insight: A partition of unity turns local change-of-variable information into a global formula. □

The nonlinear change-of-variables theorem is the local volume-scaling formula from Theorem 4.4.1 applied cell by cell after a fine grid has been warped by a C^1 diffeomorphism.

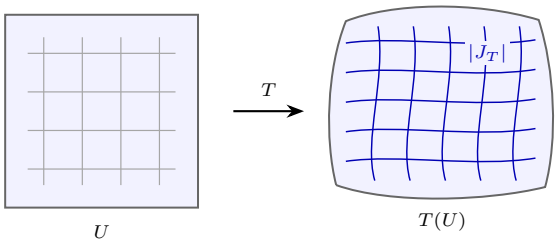


Figure 186: A C^1 change of variables turns a fine grid into curvilinear cells; $|J_T|$ records local volume scaling.

Theorem 4.4.13. *Change of Variables in $\mathbb{R}^k$*
Let $T : E \to \mathbb{R}^k$ be one-to-one and C^1 on open $E \subseteq \mathbb{R}^k$, with $J_T \neq 0$ on E.
If $f \in C(\mathbb{R}^k)$ and $\operatorname{supp}(f)$ is compact in $T(E)$, then $\int_{\mathbb{R}^k} f(\mathbf{y})\,d\mathbf{y} = \int_{\mathbb{R}^k} f(T(\mathbf{x}))\,|J_T(\mathbf{x})|\,d\mathbf{x}$.

Proof. Let $\mathcal{C}$ be the class of maps for which formula holds for all continuous functions with compact support in the image. Primitive maps lie in $\mathcal{C}$ by one-variable substitution, and flips lie in $\mathcal{C}$ by Theorem 4.4.7.
If $P, Q \in \mathcal{C}$, then $P \circ Q \in \mathcal{C}$, since Theorems 4.2.5 and 4.1.8 give $J_{P\circ Q}(\mathbf{x}) = J_P(Q(\mathbf{x}))J_Q(\mathbf{x})$.
After translating coordinates, Theorem 4.4.11 therefore gives the formula on a neighborhood of each point of E. By Theorem 4.3.4, these neighborhoods have open images covering $\operatorname{supp}(f)$.
Choose a partition of unity from Theorem 4.4.12, write $f = \sum_i \psi_i f$ with each $\psi_i f$ supported in one such image, apply the local formula to each summand, and sum.
Structural Commentary: Model cases plus partitions of unity prove the theorem.
Key Insight: The absolute Jacobian is local volume scaling, globalized by compact support. □

Method 4.4.14. *Multiple-Integration Proof Patterns*
For a multiple-integration identity:

(i) *Products*: prove product cases, then pass by uniform approximation;
(ii) *Grids*: compare product-grid upper and lower sums with one-variable slice sums;
(iii) *Support*: enclose compact support in a cell and check cell-independence;
(iv) *Local form*: reduce C^1 coordinate changes to primitive maps and flips;
(v) *Determinants*: use multiplicativity and Theorem 4.1.8;

(vi) *Globalize*: apply the partition of unity, Theorem 4.4.12.

Key Insight: Change of variables is local determinant geometry plus compactness.

Differential forms package oriented integration on parametrized cells. They use determinant sign and laws (Definition 4.1.1, Theorems 4.1.8, 4.1.15) and change of variables (Theorem 4.4.13); conceptually, forms test oriented cells.

Perspective 4.4.15. *Forms as Functionals*
The notation $\int f\,dx$ treats dx as directed coordinate variation. Differential forms make this precise: a 1-form such as $f\,dx + g\,dy$ assigns a number to an oriented path; dx records net x-variation, and $f\,dx$ weights it by f.

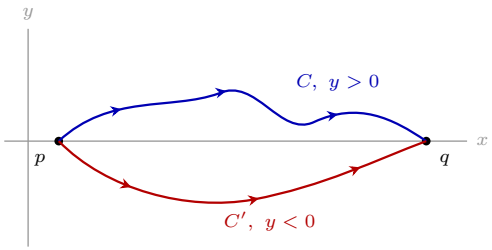


Figure 187: $y\,dx$ records signed, y-weighted x-variation.

Definition 4.4.16. *k-Surfaces and Differential Forms*
Let $E \subseteq \mathbb{R}^n$ be open. A *k-surface* in E is a C^1 map $\Phi : D \to E$ from a compact *parameter domain* $D \subseteq \mathbb{R}^k$, here a k-cell or the simplex Q^k of Definition 4.4.9.
For compact D, C^r means extension to an open neighborhood of D.
A *k-form* in E, $k \geq 1$, is $\omega = \sum a_{i_1 \cdots i_k}(\mathbf{x})\,dx_{i_1} \wedge \cdots \wedge dx_{i_k}$, with continuous real coefficients on E, and assigns to each k-surface Φ the number $\int_D \sum a_{i_1 \cdots i_k}(\Phi(\mathbf{u}))\frac{\partial(x_{i_1},\ldots,x_{i_k})}{\partial(u_1,\ldots,u_k)}\,d\mathbf{u}$.
A 0-form is a continuous function; a form is C^r if all coefficients are C^r.

Remark 4.4.17. *Parametrized Cells, Not Just Images*
Parametrizations may fail to be injective or regular. Forms therefore act on parametrized cells, not just images; this matters for orientation, multiplicity, and cancellation.

Proposition 4.4.18. *Basic Forms and Standard Presentation*
Swapping neighboring factors in $dx_{i_1} \wedge \cdots \wedge dx_{i_k}$ changes sign, and repeated indices give 0. For an increasing k-index $I = (i_1, \ldots, i_k)$, write $dx_I := dx_{i_1} \wedge \cdots \wedge dx_{i_k}$.
Every k-form has a unique standard presentation $\omega = \sum_I b_I(\mathbf{x})\,dx_I$, where I ranges over increasing k-indices.

Proof. The sign rules follow from the determinant row-swap and repeated-row rules in Theorem 4.1.8, applied to the Jacobian in Definition 4.4.16. The sign needed to sort an index tuple is exactly the permutation sign of Definition 3.2.1; hence every term either vanishes or reduces to an increasing index, giving existence.
For uniqueness, suppose $\omega = \sum_I b_I(\mathbf{x})\,dx_I$ is a zero k-form in standard presentation. If $b_J(\mathbf{v}) \neq 0$ for some increasing $J = (j_1, \ldots, j_k)$, continuity gives a small cell $D = [-h, h]^k$ such that, for $\Phi(\mathbf{u}) = \mathbf{v} + \sum_{r=1}^{k} u_r \mathbf{e}_{j_r}$, the coefficient $b_J(\Phi(\mathbf{u}))$ has constant nonzero sign on D. In the defining integral, the J-Jacobian is 1 and every other standard Jacobian is 0, so $0 = \int_\Phi \omega = \int_D b_J(\Phi(\mathbf{u}))\,d\mathbf{u}$, a contradiction. Thus every b_I is zero.
Structural Commentary: Wedge signs are determinant signs.
Key Insight: Standard presentation removes ordered-index redundancy. □

Theorem 4.4.19. *Naturality Under Reparametrization*
Let $\Phi : D \to E \subseteq \mathbb{R}^n$ be a C^1 k-surface and let $R : D' \to D$ be a C^1 diffeomorphism.
For every k-form ω, $\int_{\Phi \circ R} \omega = \operatorname{sgn}(\det DR) \int_\Phi \omega$, with constant sign on D'.

Proof. It suffices to take $\omega = f\,dx_I$. By Theorem 4.1.15, the relevant minor of $D(\Phi \circ R)$ contributes the corresponding minor of $D\Phi$ times $\det DR$. Change of variables (Theorem 4.4.13) cancels $|\det DR|$ and leaves $\operatorname{sgn}(\det DR)$; linearity handles general forms.
Structural Commentary: Reparametrization changes coordinates, not the oriented cell.
Key Insight: The same determinant sign controls orientation and the sign of the form integral. □

Definition 4.4.20. *Product and Exterior Derivative*
For basic forms, define $dx_I \wedge dx_J$ by concatenating indices and reducing to standard presentation by Proposition 4.4.18; extend bilinearly. For a C^1 0-form f, set $df := \sum_{i=1}^{n} (D_i f)(\mathbf{x})\, dx_i$.
For a C^1 k-form $\omega = \sum_I b_I(\mathbf{x})\, dx_I$, set $d\omega := \sum_I (db_I) \wedge dx_I$.

Theorem 4.4.21. *Exterior Derivative Laws*
If ω, λ are C^1 forms in E and ω is a k-form, then $d(\omega \wedge \lambda) = (d\omega) \wedge \lambda + (-1)^k \omega \wedge d\lambda$. If ω is C^2, then $d^2\omega = 0$.

Proof. By bilinearity, it suffices to take $\omega = f\, dx_I$ and $\lambda = g\, dx_J$.
The ordinary product rule gives $d(fg) = f\, dg + g\, df$; moving dg past k factors in dx_I gives the displayed sign.
For a C^2 0-form f, $d^2 f = \sum_{i,j} (D_{ij} f)(\mathbf{x})\, dx_i \wedge dx_j = 0$, by Theorem 4.2.21 and antisymmetry.
The general case follows from $\omega = \sum_I b_I\, dx_I$ and the product rule just proved.
Structural Commentary: The exterior derivative is built to make mixed partial cancellation visible.
Key Insight: Symmetric second derivatives meet antisymmetric wedge products, so $d^2 = 0$. □

Definition 4.4.22. *Change of Variables for Forms*
Let $T : E \subseteq \mathbb{R}^n \to V \subseteq \mathbb{R}^m$ be C^1, with $T = (t_1, \ldots, t_m)$. For a k-form $\omega = \sum_I b_I(\mathbf{y})\, dy_I$ in V, define its *pullback* by $\omega_T := T^*\omega := \sum_I b_I(T(\mathbf{x}))\, dt_{i_1} \wedge \cdots \wedge dt_{i_k}$, and $dt_i := \sum_{j=1}^{n} (D_j t_i)(\mathbf{x})\, dx_j$.

Pullback is the differential-form analogue of the dual map in Definition 1.5.11: cells move forward by composition with T, while forms move backward so they can be evaluated before or after that forward motion.

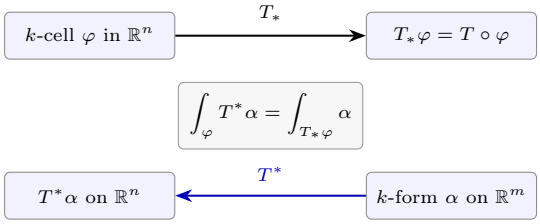


Figure 188: Cells push forward with T; forms pull back with T^*, so $T^*\alpha(\varphi) = \alpha(T_*\varphi)$.

Theorem 4.4.23. *Pullback Commutes with the Basic Operations*
With $T : E \to V$ as in Definition 4.4.22,

(i) $(\omega + \lambda)_T = \omega_T + \lambda_T$ *for forms of the same order;*
(ii) $(\omega \wedge \lambda)_T = \omega_T \wedge \lambda_T$*;*
(iii) $d(\omega_T) = (d\omega)_T$ *when ω is C^1 and T is C^2;*
(iv) if $S : V \to W$ is C^1, then $(\omega_S)_T = \omega_{S \circ T}$.

Moreover, if Φ is a k-surface in E and ω is a k-form in V, then $\int_{T \circ \Phi} \omega = \int_\Phi \omega_T$.

Proof.

(i) Addition is the linearity of Definition 4.4.22.
(ii) Pullback replaces each dy_i by dt_i; bilinearity gives $(\omega \wedge \lambda)_T = \omega_T \wedge \lambda_T$.
(iii) Let $\omega = \sum_I b_I\, dy_I$. Since $d(b_I \circ T) = (db_I)_T$ (Theorem 4.2.5) and $d(dt_i) = d^2 t_i = 0$, $d(\omega_T) = (d\omega)_T$.
(iv) If $S = (s_1, \ldots, s_p)$, both $(\omega_S)_T$ and $\omega_{S \circ T}$ replace dz_q by $d(s_q \circ T)$, by the chain rule.

For the integral identity, expand both sides over D; the Jacobian minors agree by Theorem 4.1.15, since $D(T \circ \Phi) = DT(\Phi) D\Phi$.
Structural Commentary: Pullback translates forms before integration.
Key Insight: Forms are defined so that composition of parameterizations becomes an algebraic substitution. □

4.5 Stokes Theory

Stokes theory relates exterior differentiation to oriented boundary, with Green's theorem, classical Stokes, and divergence as special cases.

Definition 4.5.1. *Chains and Boundaries*
For $\mathbf{p}_0, \ldots, \mathbf{p}_k \in \mathbb{R}^n$, oriented affine k-simplex $\sigma = [\mathbf{p}_0, \ldots, \mathbf{p}_k]$ is k-surface on Q^k: $\sigma(\mathbf{u}) = \mathbf{p}_0 + \sum_{i=1}^{k} u_i(\mathbf{p}_i - \mathbf{p}_0)$.
Permuting vertices multiplies σ by the permutation sign. An *affine k-chain* is a finite formal sum of oriented

affine k-simplices; integration is extended linearly.

For $k \geq 1$, define $\partial_{\mathrm{ch}}[\mathbf{p}_0, \ldots, \mathbf{p}_k] := \sum_{j=0}^{k} (-1)^j [\mathbf{p}_0, \ldots, \mathbf{p}_{j-1}, \mathbf{p}_{j+1}, \ldots, \mathbf{p}_k]$, and extend ∂_{ch} linearly. If $\Phi = T \circ \sigma$ is a C^2 oriented simplex, set $\partial_{\mathrm{ch}} \Phi := T(\partial_{\mathrm{ch}} \sigma)$, again extended linearly to chains.

Theorem 4.5.2. *Stokes' Theorem*
If Ψ is a C^2 k-chain in open $V \subseteq \mathbb{R}^m$ and ω is a C^1 $(k-1)$-form in V, then $\int_\Psi d\omega = \int_{\partial_{\mathrm{ch}} \Psi} \omega$.

Proof. By linearity, take $\Psi = \Phi = T \circ \sigma$, where σ is the standard oriented k-simplex.
By Theorem 4.4.23, $\int_\Phi d\omega = \int_\sigma (d\omega)_T = \int_\sigma d(\omega_T)$; and $\int_{\partial_{\mathrm{ch}} \Phi} \omega = \int_{\partial_{\mathrm{ch}} \sigma} \omega_T$, remains to prove formula on σ.
By linearity, reduce to $\lambda = f \, dx_1 \wedge \cdots \wedge \widehat{dx_r} \wedge \cdots \wedge dx_k$.
The boundary formula in Definition 4.5.1 leaves only the two faces in the x_r direction, and the resulting identity is the one-variable fundamental theorem of calculus inside the iterated integral.
Structural Commentary: Pullback reduces Stokes' theorem to a computation on the standard simplex.
Key Insight: The boundary signs are exactly the signs produced by the missing wedge factor. □

Definition 4.5.3. *Closed and Exact Forms*
A k-form ω in E is *exact* if $\omega = d\lambda$ for some $(k-1)$-form λ in E, and *closed* if ω is C^1 with $d\omega = 0$.
By Theorem 4.4.21, exact C^1 forms are closed.

Perspective 4.5.4. *Closed Is Local, Exact Is Global*
Exact $\Rightarrow$ closed is just $d^2 = 0$; the converse depends on the domain.
On $\mathbb{R}^2 \setminus \{\mathbf{0}\}$ the standard warning example is $\omega = \frac{-y}{x^2+y^2} \, dx + \frac{x}{x^2+y^2} \, dy$, which is closed but not exact.
Holes can obstruct global antiderivatives even when local derivative tests vanish.

Theorem 4.5.5. *Poincare Lemma on Convex Sets*
If $E \subseteq \mathbb{R}^n$ is convex and open, then every closed C^1 k-form in E, $k \geq 1$, is exact.
The same holds on every open set C^2-equivalent to a convex open set.

Proof. Choose $\mathbf{a} \in E$ and let $H_t(\mathbf{x}) = \mathbf{a} + t(\mathbf{x} - \mathbf{a})$, $0 \leq t \leq 1$. Convexity gives $H_t(E) \subseteq E$. For a k-form ω, define the $(k-1)$-form as $(K\omega)_{\mathbf{x}}(\mathbf{v}_1, \ldots, \mathbf{v}_{k-1}) := \int_0^1 t^{k-1} \omega_{H_t(\mathbf{x})}(\mathbf{x} - \mathbf{a}, \mathbf{v}_1, \ldots, \mathbf{v}_{k-1}) \, dt$.
A coordinate differentiation under the integral gives the homotopy identity $\omega - H_0^* \omega = dK\omega + Kd\omega$. Since $k \geq 1$, the constant map H_0 pulls every k-form back to 0. If $d\omega = 0$, then $\omega = d(K\omega)$, so ω is exact.
If E is C^2-equivalent to a convex open set C by $\Phi : C \to E$, then $\Phi^* \omega$ is closed by Theorem 4.4.23. The convex case gives $\Phi^* \omega = d\eta$ on C. Applying $(\Phi^{-1})^*$ and using the same pullback law gives $\omega = d((\Phi^{-1})^* \eta)$ on E.
Structural Commentary: Convexity lets coefficient-level antiderivatives be chosen consistently.
Key Insight: Closed is local symmetry; convexity supplies the global primitive. □

Theorem 4.5.6. *Vector Analysis Dictionary in $\mathbb{R}^3$*
For $\mathbf{F} = F_1 \mathbf{e}_1 + F_2 \mathbf{e}_2 + F_3 \mathbf{e}_3$ on open $E \subseteq \mathbb{R}^3$,
set $\lambda_{\mathbf{F}} := F_1 \, dx + F_2 \, dy + F_3 \, dz$ and $\omega_{\mathbf{F}} := F_1 \, dy \wedge dz + F_2 \, dz \wedge dx + F_3 \, dx \wedge dy$. Then
$\mathbf{F} = \nabla u \iff \lambda_{\mathbf{F}} = du$; and $\nabla \times \mathbf{F} = 0 \iff d\lambda_{\mathbf{F}} = 0$; and
$\mathbf{F} = \nabla \times \mathbf{G} \iff \omega_{\mathbf{F}} = d\lambda_{\mathbf{G}}$; and $\nabla \cdot \mathbf{F} = 0 \iff d\omega_{\mathbf{F}} = 0$.
Hence $\nabla \times (\nabla u) = 0$ and $\nabla \cdot (\nabla \times \mathbf{G}) = 0$; on sets C^2-equivalent to convex open sets, the converses hold.

Proof. Expand $d\lambda_{\mathbf{F}}$ and $d\omega_{\mathbf{F}}$ by Definition 4.4.20; their coefficients are exactly $\nabla \times \mathbf{F}$ and $\nabla \cdot \mathbf{F}$. Potentials and vector potentials are exactness of $\lambda_{\mathbf{F}}$ and $\omega_{\mathbf{F}}$; the converses follow from Theorem 4.5.5.
Structural Commentary: Gradient, curl, and divergence are exterior derivative in degrees $0, 1, 2$.
Key Insight: The familiar vector identities are $d^2 = 0$ written in $\mathbb{R}^3$ notation. □

Theorem 4.5.7. *Green, Classical Stokes, and Divergence*
The following are special cases of Theorem 4.5.2.

(i) If $\Omega \subseteq \mathbb{R}^2$ has positively oriented boundary, $\alpha, \beta \in C^1$ near Ω, then $\int_{\mathrm{bd}\,\Omega} (\alpha \, dx + \beta \, dy) = \int_\Omega \left(\frac{\partial \beta}{\partial x} - \frac{\partial \alpha}{\partial y} \right) dA$
(ii) If Φ is a C^2 2-surface in $\mathbb{R}^3$ and $\mathbf{F} \in C^1$ near Φ, then $\int_\Phi (\nabla \times \mathbf{F}) \cdot \mathbf{n} \, dA = \int_{\partial_{\mathrm{ch}} \Phi} \mathbf{F} \cdot \mathbf{t} \, ds$.
(iii) If $\Omega \subseteq \mathbb{R}^3$ has positively oriented boundary and $\mathbf{F} \in C^1$ near Ω, then $\int_\Omega (\nabla \cdot \mathbf{F}) \, dV = \int_{\mathrm{bd}\,\Omega} \mathbf{F} \cdot \mathbf{n} \, dA$

Proof.
(i) Put $\lambda = \alpha \, dx + \beta \, dy$. Then $d\lambda = (\partial \beta / \partial x - \partial \alpha / \partial y) \, dA$.
(ii) Put $\lambda = \lambda_{\mathbf{F}}$. By Theorem 4.5.6, $d\lambda = \omega_{\nabla \times \mathbf{F}}$.
(iii) Put $\omega = \omega_{\mathbf{F}}$. Then $d\omega = (\nabla \cdot \mathbf{F}) \, dx \wedge dy \wedge dz$.

Theorem 4.5.2, together with the line, area, and volume element translations, gives the three displayed formulas.
Structural Commentary: Forms translate Stokes into the classical vector formulas.
Key Insight: Boundary plus d is the content; vector notation is translation. □

Method 4.5.8. *Differential-Forms Proof Patterns*

(i) *Signs*: compute wedge signs by determinant swaps and permutation parity.
(ii) *Standard presentation*: remove repeated indices and sort the rest.
(iii) *Pullback*: substitute dt_i for dy_i and apply the chain rule.
(iv) *Stokes*: pull back to the standard simplex, prove the coordinate case, extend linearly.
(v) *Closed versus exact*: use $d^2 = 0$ for necessity; use convexity or C^2-equivalence for primitives.
(vi) *Vector analysis*: translate fields to forms, apply Stokes, translate back.

Key Insight: Forms package signs, orientation, and parameter changes.

5 Measure and Integration Structure

5.1 Sigma-Algebras and Measures

Measure theory first builds measurable sets, then measurable functions and integrals. Outer measure covers from outside; measurability splits test sets without losing measure.

Perspective 5.1.1. *Why Abstract Measure Comes First*
Lebesgue integration is built so that, under hypotheses such as domination, $f_n \to f$ implies $\int f_n \, d\mu \to \int f \, d\mu$; see Theorem 5.5.4. The abstract data separate into a measurable space $(X, \mathcal{M})$, a countably additive measure $\mu : \mathcal{M} \to [0, +\infty]$, and an integral defined by approximation from simple functions.
Key Insight: Countable additivity is the set-level structure behind the convergence theorems.

Definition 5.1.2. *σ-Algebras and Measurable Spaces*
Let X be a set. A family $\mathcal{M} \subseteq \mathcal{P}(X)$ is a *σ-algebra* on X if:

(i) $\emptyset, X \in \mathcal{M}$;
(ii) $E \in \mathcal{M}$ implies $X \setminus E \in \mathcal{M}$;
(iii) $E_n \in \mathcal{M}$ for $n \in \mathbb{N}$ implies $\bigcup_{n=1}^{\infty} E_n \in \mathcal{M}$.

The pair $(X, \mathcal{M})$ is a *measurable space*; the members of $\mathcal{M}$ are *measurable sets*.

Proposition 5.1.3. *Closure and Generation for σ-Algebras*
Let $\mathcal{M} \subseteq \mathcal{P}(X)$ be a σ-algebra.

(i) $(E_n)_{n\in\mathbb{N}} \subseteq \mathcal{M} \Rightarrow \bigcap_{n=1}^{\infty} E_n \in \mathcal{M}$.
(ii) If $(\mathcal{M}_\alpha)_{\alpha\in A}$ is a nonempty family of σ-algebras on X, then $\bigcap_{\alpha\in A} \mathcal{M}_\alpha$ is a σ-algebra on X.
(iii) For $\mathcal{A} \subseteq \mathcal{P}(X)$, let $\mathscr{S}_{\mathcal{A}}$ be the family of σ-algebras $\mathcal{N}$ on X with $\mathcal{A} \subseteq \mathcal{N}$. Then $\sigma(\mathcal{A}) := \bigcap_{\mathcal{N}\in\mathscr{S}_{\mathcal{A}}} \mathcal{N}$ is the smallest σ-algebra containing $\mathcal{A}$.

Proof.

(i) Since $X \setminus E_n \in \mathcal{M}$, Theorem 1.4.3 gives $\bigcap_n E_n = X \setminus \bigcup_n (X \setminus E_n) \in \mathcal{M}$.
(ii) Set $\mathcal{N} := \bigcap_{\alpha\in A} \mathcal{M}_\alpha$. Then $\emptyset, X \in \mathcal{N}$, and $E \in \mathcal{N}$ implies $X \setminus E \in \mathcal{N}$. Likewise $(E_n) \subseteq \mathcal{N}$ implies $\bigcup_n E_n \in \mathcal{N}$.
(iii) The indexing class in the displayed definition is nonempty since $\mathcal{P}(X)$ belongs to it. By (ii), its intersection is a σ-algebra; it contains $\mathcal{A}$ and is contained in every σ-algebra containing $\mathcal{A}$.

Structural Commentary: Generation is intersection over all admissible σ-algebras.
Key Insight: Intersections preserve universal closure properties. □

Example 5.1.4. *Basic σ-Algebras*
Standard models:

(i) Extremes: $\{\emptyset, X\} \subseteq \mathcal{M} \subseteq \mathcal{P}(X)$ for every σ-algebra $\mathcal{M}$ on X.
(ii) For $E \subseteq X$, $\sigma(\{E\}) = \{\emptyset, E, X \setminus E, X\}$, with repetitions omitted.
(iii) On $\mathbb{N}$, if $O := \{2n - 1 : n \in \mathbb{N}\}$, then $\sigma(\{O\}) = \{\emptyset, O, \mathbb{N} \setminus O, \mathbb{N}\}$.
(iv) If X is uncountable, $\mathcal{C}_X := \{E \subseteq X : E \text{ is countable or } X \setminus E \text{ is countable}\}$ is the countable–cocountable σ-algebra.

Key Insight: A σ-algebra records exactly which distinctions the theory is allowed to observe.

Definition 5.1.5. *Borel and Extended Borel σ-Algebras*

(i) For $n \in \mathbb{N}$, the *Borel σ-algebra* on $\mathbb{R}^n$ is $\mathcal{B}(\mathbb{R}^n) := \sigma(\{U \subseteq \mathbb{R}^n : U \text{ is open}\})$.
On $\mathbb{R}$, it is also generated by open intervals, closed intervals, half-open intervals, or open half-rays $(a, +\infty)$.
(ii) The *extended Borel σ-algebra* on $\overline{\mathbb{R}}$ is $\overline{\mathcal{B}} := \{E \cup F : E \in \mathcal{B}(\mathbb{R}),\ F \subseteq \{-\infty, +\infty\}\}$.

Definition 5.1.6. *Monotone Classes*
A nonempty $\mathcal{C} \subseteq \mathcal{P}(X)$ is a *monotone class* iff $(E_n) \subseteq \mathcal{C},\ E_n \uparrow E \Rightarrow E \in \mathcal{C}$ and $(F_n) \subseteq \mathcal{C},\ F_n \downarrow F \Rightarrow F \in \mathcal{C}$.
For $\mathcal{A} \subseteq \mathcal{P}(X)$, define $\mathfrak{M}_{\mathcal{A}} := \{\mathcal{C} \subseteq \mathcal{P}(X) : \mathcal{C} \text{ is a monotone class},\ \mathcal{A} \subseteq \mathcal{C}\}$; $\operatorname{mon}(\mathcal{A}) := \bigcap \mathfrak{M}_{\mathcal{A}}$.
Then $\operatorname{mon}(\mathcal{A})$ is the least monotone class containing $\mathcal{A}$, and $\operatorname{mon}(\mathcal{A}) \subseteq \sigma(\mathcal{A})$.
Key Insight: Monotone closure captures limits, not finite unions.

Definition 5.1.7. *Algebras of Sets and Premeasures*
Let X be a set.

(i) An *algebra of sets* (or *field of sets*) on X is $\mathcal{A} \subseteq \mathcal{P}(X)$ such that $\emptyset, X \in \mathcal{A}$; $E \in \mathcal{A} \Rightarrow X \setminus E \in \mathcal{A}$; and $E_1, \ldots, E_n \in \mathcal{A} \Rightarrow \bigcup_{j=1}^{n} E_j \in \mathcal{A}$. Equivalently, $\mathcal{A}$ is a ring of sets containing X.
(ii) A *premeasure* (or *measure on an algebra*) on $\mathcal{A}$ is $\mu : \mathcal{A} \to [0, +\infty]$ such that $\mu(\emptyset) = 0$ and, whenever $(E_n)_{n\in\mathbb{N}} \subseteq \mathcal{A}$ is pairwise disjoint with $\bigcup_{n=1}^{\infty} E_n \in \mathcal{A}$, then $\mu\left(\bigcup_{n=1}^{\infty} E_n\right) = \sum_{n=1}^{\infty} \mu(E_n)$.

Key Insight: Premeasures are countably additive where the algebra permits unions.

Theorem 5.1.8. *Monotone Class Theorem for Algebras*
Let $\mathcal{A} \subseteq \mathcal{P}(X)$ be an algebra and set $\mathcal{C} := \operatorname{mon}(\mathcal{A})$. Then $\mathcal{C} = \sigma(\mathcal{A})$.

Proof. Since $\sigma(\mathcal{A})$ is monotone and contains $\mathcal{A}$, $\mathcal{C} \subseteq \sigma(\mathcal{A})$. It suffices to show that $\mathcal{C}$ is a σ-algebra.
For complements, set $\mathcal{D} := \{E \in \mathcal{C} : E^c \in \mathcal{C}\}$. Then $E_n \uparrow E$ in $\mathcal{D}$ implies $E_n^c \downarrow E^c$, and $E_n \downarrow E$ implies $E_n^c \uparrow E^c$; hence $\mathcal{D}$ is monotone. Since $\mathcal{A} \subseteq \mathcal{D}$, minimality gives $\mathcal{C} \subseteq \mathcal{D}$.
For intersections, fix $B \in \mathcal{C}$ and set $\mathcal{D}_B := \{E \in \mathcal{C} : E \cap B \in \mathcal{C}\}$. This is monotone because intersections preserve $\uparrow, \downarrow$ limits. If $B \in \mathcal{A}$, then $\mathcal{A} \subseteq \mathcal{D}_B$, so $\mathcal{C} \subseteq \mathcal{D}_B$; hence $E \cap B \in \mathcal{C}$ for $E \in \mathcal{C}$, $B \in \mathcal{A}$. Fixing $E \in \mathcal{C}$ now gives $\mathcal{A} \subseteq \mathcal{D}_E$, hence $\mathcal{C} \subseteq \mathcal{D}_E$. Thus $\mathcal{C}$ is closed under finite intersections.
Complements and finite intersections give finite unions. If $(E_n) \subseteq \mathcal{C}$, then $F_N := \bigcup_{n=1}^{N} E_n \in \mathcal{C}$ and $F_N \uparrow \bigcup_n E_n$, so $\bigcup_n E_n \in \mathcal{C}$. Hence $\mathcal{C}$ is a σ-algebra containing $\mathcal{A}$, and $\sigma(\mathcal{A}) \subseteq \mathcal{C}$.
Structural Commentary: Finite Boolean closure plus monotone limits gives σ-closure.
Key Insight: Prove closure properties by showing the sets that have the property form a monotone class. □

Corollary 5.1.9. *Monotone-Class Test*
Let $\mathcal{A} \subseteq \mathcal{P}(X)$ be an algebra and let $\mathcal{C} \subseteq \mathcal{P}(X)$ be a monotone class. If $\mathcal{A} \subseteq \mathcal{C}$, then $\sigma(\mathcal{A}) \subseteq \mathcal{C}$.

Proof. Since $\mathcal{A} \subseteq \mathcal{C}$ and $\mathcal{C}$ is monotone, $\operatorname{mon}(\mathcal{A}) \subseteq \mathcal{C}$. By Theorem 5.1.8, $\operatorname{mon}(\mathcal{A}) = \sigma(\mathcal{A})$.
Structural Commentary: Algebra checks plus monotone stability control $\sigma(\mathcal{A})$.
Key Insight: Prove on the algebra; extend by monotone closure. □

Definition 5.1.10. *Rings, σ-Rings, and Set Functions*
Let $\mathcal{A} \subseteq \mathcal{P}(X)$.

(i) $\mathcal{A}$ is a *ring of sets* iff $\emptyset \in \mathcal{A}$ and $A, B \in \mathcal{A} \Rightarrow A \cup B,\ A \setminus B \in \mathcal{A}$
(ii) $\mathcal{A}$ is a *σ-ring* iff it is a ring and $(A_n) \subseteq \mathcal{A} \Rightarrow \bigcup_{n=1}^{\infty} A_n \in \mathcal{A}$
(iii) A *set function* on $\mathcal{A}$ is $\phi : \mathcal{A} \to [-\infty, +\infty]$ with $\{+\infty, -\infty\} \nsubseteq \phi(\mathcal{A})$.
(iv) ϕ is *additive* iff $A \cap B = \emptyset \Rightarrow \phi(A \cup B) = \phi(A) + \phi(B)$
(v) ϕ is *countably additive* iff $(A_n) \subseteq \mathcal{A}$ pairwise disjoint, $\bigcup_n A_n \in \mathcal{A} \Rightarrow \phi\left(\bigcup_{n=1}^{\infty} A_n\right) = \sum_{n=1}^{\infty} \phi(A_n)$.

Structural Commentary: Rings support finite operations; σ-rings support countable unions.

Definition 5.1.11. *Measures on Measurable Spaces*
Let $(X, \mathcal{M})$ be a measurable space. A *measure* on $(X, \mathcal{M})$ is a map $\mu : \mathcal{M} \to [0, +\infty]$ such that:

(i) $\mu(\emptyset) = 0$;
(ii) if $(E_n) \subseteq \mathcal{M}$ is pairwise disjoint, then $\mu\left(\bigcup_{n=1}^{\infty} E_n\right) = \sum_{n=1}^{\infty} \mu(E_n)$

The triple $(X, \mathcal{M}, \mu)$ is a *measure space*.
Key Insight: Measure is countably additive size on the chosen observable sets.

Definition 5.1.12. *Finite and σ-Finite Measures*
Let $(X, \mathcal{M}, \mu)$ be a measure space.

(i) For nonnegative terms, $\sum_{n=1}^{\infty} a_n := \lim_{N\to\infty} \sum_{n=1}^{N} a_n \in [0, +\infty]$.
(ii) μ is *finite* iff $\mu(E) < +\infty$ for every $E \in \mathcal{M}$.
(iii) μ is *σ-finite* iff there exist $E_n \in \mathcal{M}$ such that $X = \bigcup_{n=1}^{\infty} E_n$, $\mu(E_n) < +\infty$ for $n \in \mathbb{N}$

Definition 5.1.13. *Charges*
Let $(X, \mathcal{M})$ be a measurable space. A *charge* is a map $\lambda : \mathcal{M} \to \mathbb{R}$ with $\lambda(\emptyset) = 0$ such that for every pairwise disjoint sequence $(E_n) \subseteq \mathcal{M}$, the series $\sum_{n=1}^{\infty} \lambda(E_n)$ converges in $\mathbb{R}$ and $\lambda(\bigcup_{n=1}^{\infty} E_n) = \sum_{n=1}^{\infty} \lambda(E_n)$.
Key Insight: Charges are real-valued signed measures with unconditional sums.

Example 5.1.14. *First Abstract Measures*

(i) The zero set function $\mu(E) = 0$ on $\mathcal{M}$ is a finite measure.
(ii) If $X \neq \emptyset$ and $\mathcal{M} = 2^X$, $\mu(\emptyset) = 0$ and $\mu(E) = +\infty$ for $E \neq \emptyset$ give a measure, neither finite nor σ-finite.
(iii) For $p \in X$, the *unit mass* $\delta_p(E) := \mathbf{1}_E(p)$ on 2^X is a finite measure.
(iv) On $2^{\mathbb{N}}$, *counting measure* is $\#(E)$ for finite E and $+\infty$ otherwise; it is σ-finite.
More generally, $(a_n) \subseteq [0, +\infty]$ gives $\mu(E) := \sum_{n \in E} a_n$, and every measure on $2^{\mathbb{N}}$ is uniquely of this form.
(v) If X is uncountable and $\mathcal{M} = 2^X$, $\mu(E) = 0$ for countable E and $\mu(E) = +\infty$ otherwise defines a measure.
(vi) On $2^{\mathbb{N}}$, $\mu(E) = 0$ for finite E and $\mu(E) = +\infty$ otherwise is not a measure, since $\mathbb{N} = \bigcup_{n=1}^{\infty} \{n\}$ gives $+\infty \neq 0$; the constant $+\infty$ rule fails at $\emptyset$.
(vii) Lebesgue measure on $\mathcal{B}(\mathbb{R})$ is σ-finite and gives length $b-a$ to standard intervals; Borel–Stieltjes measures come from increasing right-continuous functions as in Theorem 5.1.27.

Structural Commentary: Measures may be zero, atomic, counting-type, cardinality-based, or geometric.
Key Insight: Countable additivity turns disjoint decompositions into nonnegative series.

Definition 5.1.15. *Outer Measure Generated by a Premeasure*
Let μ be a premeasure on an algebra $\mathcal{A} \subseteq \mathcal{P}(X)$. The *outer measure generated by* μ is the set function $\mu^* : \mathcal{P}(X) \to [0, +\infty]$ defined by $\mu^*(B) := \inf\{\sum_{j=1}^{\infty} \mu(E_j) : E_j \in \mathcal{A},\ B \subseteq \bigcup_{j=1}^{\infty} E_j\}$ for $B \subseteq X$.

Lemma 5.1.16. *Generated Outer-Measure Properties*
Let μ be a premeasure on an algebra $\mathcal{A} \subseteq \mathcal{P}(X)$, and let μ^* be as in Definition 5.1.15.

(i) $\mu^*(\emptyset) = 0$.
(ii) $C \subseteq D \subseteq X \Rightarrow \mu^*(C) \leq \mu^*(D)$.
(iii) $(B_n) \subseteq \mathcal{P}(X) \Rightarrow \mu^*(\bigcup_{n=1}^{\infty} B_n) \leq \sum_{n=1}^{\infty} \mu^*(B_n)$.
(iv) $E \in \mathcal{A} \Rightarrow \mu^*(E) = \mu(E)$.

Proof.

(i) The cover $(\emptyset)_j$ gives $\mu^*(\emptyset) \leq 0$, while $\mu^* \geq 0$ by definition.
(ii) Every algebra cover of D covers C; taking infima gives monotonicity.
(iii) If $\sum_n \mu^*(B_n) = +\infty$, there is nothing to prove.
Otherwise, for $\varepsilon > 0$ choose covers $(E_{n,k})_{k \geq 1} \subseteq \mathcal{A}$ of B_n with $\sum_{k=1}^{\infty} \mu(E_{n,k}) \leq \mu^*(B_n) + \varepsilon 2^{-n}$.
Then $(E_{n,k})_{n,k}$ covers $\bigcup_n B_n$, hence $\mu^*(\bigcup_n B_n) \leq \sum_{n=1}^{\infty} \sum_{k=1}^{\infty} \mu(E_{n,k}) \leq \sum_n \mu^*(B_n) + \varepsilon$; let $\varepsilon \downarrow 0$.
(iv) For $E \in \mathcal{A}$, the cover $E, \emptyset, \emptyset, \ldots$ gives $\mu^*(E) \leq \mu(E)$.
Conversely, if $E \subseteq \bigcup_n F_n$ with $F_n \in \mathcal{A}$, set $G_n := E \cap F_n \setminus \bigcup_{k<n} F_k$. Then (G_n) is disjoint in $\mathcal{A}$, $\bigcup_n G_n = E$, and $\mu(E) = \sum_n \mu(G_n) \leq \sum_n \mu(F_n)$. Taking infima over covers gives $\mu(E) \leq \mu^*(E)$.

Structural Commentary: Generated outer measures extend premeasures and are monotone and subadditive.
Key Insight: Intersect an arbitrary cover with E and disjointize inside the algebra. □

Proposition 5.1.17. *Single-Set Envelope Comparison*
Let μ be a premeasure on an algebra $\mathcal{A} \subseteq \mathcal{P}(X)$, and let μ^* be generated by μ. For $B \subseteq X$, set $\mu^\sharp(B) := \inf\{\mu(C) : C \in \mathcal{A},\ B \subseteq C\}$. Then $\mu^\sharp|_{\mathcal{A}} = \mu$ and $\mu^* \leq \mu^\sharp$.

Proof. If $E \in \mathcal{A}$, then E is admissible, so $\mu^\sharp(E) \leq \mu(E)$.
Conversely, $E \subseteq C \in \mathcal{A}$ implies $\mu(E) \leq \mu(C)$; taking infima gives $\mu(E) \leq \mu^\sharp(E)$.
If $B \subseteq C \in \mathcal{A}$, then $C, \emptyset, \emptyset, \ldots$ covers B, so $\mu^*(B) \leq \mu(C)$. Taking infima over C gives $\mu^*(B) \leq \mu^\sharp(B)$.
Structural Commentary: One-set envelopes dominate generated outer measures.
Key Insight: A single algebra superset is a countable algebra cover. □

Example 5.1.18. *Single-Set Envelopes Need Not Be Outer Measures*
Let X be uncountable, let $\mathcal{A}$ be the finite–cofinite algebra on X, and set $\mu(E) := 0$ for finite E, $\mu(E) := +\infty$ for cofinite E. Then μ is a premeasure, $\mu^*(B) = 0$ iff B is countable, and $\mu^\sharp(B) = 0$ iff B is finite.
For any countably infinite $S \subseteq X$, $\mu^\sharp(S) = +\infty$ while $S = \bigcup_{x \in S} \{x\}$ and $\sum_{x \in S} \mu^\sharp(\{x\}) = 0$, QED.
Structural Commentary: Countable covers can be strictly smaller than one-set covers.
Key Insight: Outer measure is built from countable covers precisely to recover countable subadditivity.

Theorem 5.1.19. *Generated Caratheodory Extension*
Let μ be a premeasure on an algebra $\mathcal{A} \subseteq \mathcal{P}(X)$, let μ^* be generated by μ,
and set $\mathcal{C}_\mu := \{E \subseteq X : \forall B \subseteq X, \mu^*(B) = \mu^*(B \cap E) + \mu^*(B \setminus E)\}$.
Then $\mathcal{A} \subseteq \mathcal{C}_\mu$, $\mathcal{C}_\mu$ is a complete σ-algebra, and $\mu^*|_{\mathcal{C}_\mu}$ is a measure extending μ.

Proof. Lemma 5.1.16 makes μ^* an outer measure; Theorem 5.1.56 makes $\mathcal{C}_\mu$ a σ-algebra, makes $\mu^*|_{\mathcal{C}_\mu}$ a measure, and gives completeness through zero sets. Let $E \in \mathcal{A}$ and $B \subseteq X$.
Subadditivity gives $\mu^*(B) \le \mu^*(B \cap E) + \mu^*(B \setminus E)$.
Conversely, if $B \subseteq \bigcup_n F_n$ with $F_n \in \mathcal{A}$, then $F_n \cap E, F_n \setminus E \in \mathcal{A}$ and $\mu(F_n) = \mu(F_n \cap E) + \mu(F_n \setminus E)$, hence $\mu^*(B \cap E) + \mu^*(B \setminus E) \le \sum_n \mu(F_n)$. Infimizing over covers gives the reverse inequality, so $E \in \mathcal{C}_\mu$.
Thus $\mathcal{A} \subseteq \mathcal{C}_\mu$, and Lemma 5.1.16 gives $\mu^*|_{\mathcal{A}} = \mu$.
Structural Commentary: Caratheodory equality is the additivity domain.
Key Insight: Every algebra set is a clean splitter for the outer measure it generated. □

Proposition 5.1.20. *Caratheodory Splitters for Finitely Subadditive Size*
Let $\alpha : \mathcal{P}(X) \to [0, +\infty)$ satisfy $0 \le \alpha(E) \le \alpha(E \cup F) \le \alpha(E) + \alpha(F)$ for all $E, F \subseteq X$, and set $\mathcal{S} := \{E \subseteq X : (\forall A \subseteq X)\alpha(A) = \alpha(A \cap E) + \alpha(A \setminus E)\}$.

(i) If $\mathcal{S} \neq \emptyset$, then $\mathcal{S}$ is an algebra on X.
(ii) If $E, F \in \mathcal{S}$ and $E \cap F = \emptyset$, then $\alpha(E \cup F) = \alpha(E) + \alpha(F)$.
(iii) Nonemptiness is necessary: $\alpha \equiv 1$ gives $\mathcal{S} = \emptyset$.

Proof.

(i) Choose $E_0 \in \mathcal{S}$. Taking $A = \emptyset$ gives $\alpha(\emptyset) = 0$, hence $\emptyset, X \in \mathcal{S}$. If $E \in \mathcal{S}$, the defining equality is symmetric in E and $X \setminus E$, so $X \setminus E \in \mathcal{S}$. If $E, F \in \mathcal{S}$ and $A \subseteq X$, applying F to A and $A \setminus (E \cap F)$ and applying E to $A \cap F$ gives $\alpha(A) = \alpha(A \cap E \cap F) + \alpha(A \setminus (E \cap F))$; hence $E \cap F \in \mathcal{S}$. Complements give finite unions.
(ii) Apply the defining equality for E to $A = E \cup F$: $\alpha(E \cup F) = \alpha(E) + \alpha(F)$.
(iii) If $\alpha \equiv 1$ and $E \in \mathcal{S}$, taking $A = \emptyset$ gives $1 = 1 + 1$, impossible.

Structural Commentary: The Caratheodory splitting equation is already Boolean at the finite level.
Key Insight: Nonempty splitters force $\alpha(\emptyset) = 0$; without that, the clean-splitting class may vanish. □

Theorem 5.1.21. *σ-Finite Generated Extensions Are Unique*
Let μ be a premeasure on an algebra $\mathcal{A} \subseteq \mathcal{P}(X)$, let μ^* be generated by μ, and let $\mathcal{C}_\mu$ be as in Theorem 5.1.19. If $X = \bigcup_n F_n$ with $F_n \in \mathcal{A}$ and $\mu(F_n) < +\infty$, then $\mu^*|_{\mathcal{C}_\mu}$ is the unique measure ν on $\mathcal{C}_\mu$ satisfying $\nu|_{\mathcal{A}} = \mu$.

Proof. Existence is Theorem 5.1.19. For uniqueness, let ν be a measure on $\mathcal{C}_\mu$ with $\nu|_{\mathcal{A}} = \mu$.
If $\mu(X) < +\infty$, then for $E \in \mathcal{C}_\mu$ and every algebra cover $E \subseteq \bigcup_n A_n$, $\nu(E) \le \sum_n \mu(A_n)$, hence $\nu(E) \le \mu^*(E)$.
Applying this to $X \setminus E$ and using $\nu(E) + \nu(X \setminus E) = \nu(X) = \mu(X) = \mu^*(E) + \mu^*(X \setminus E)$ forces $\nu(E) = \mu^*(E)$.
In general, set $H_n := \bigcup_{j=1}^n F_j$. The finite case on $\mathcal{A}|_{H_n} := \{A \cap H_n : A \in \mathcal{A}\}$ gives $\nu(E \cap H_n) = \mu^*(E \cap H_n)$; since $E \cap H_n \uparrow E$, continuity from below gives $\nu(E) = \mu^*(E)$.
Structural Commentary: Uniqueness reduces to finite-mass exhaustions.
Key Insight: Existence is Caratheodory; uniqueness needs a finite-mass exhaustion. □

Corollary 5.1.22. *Uniqueness on the Generated σ-Algebra*
Let μ be a premeasure on an algebra $\mathcal{A} \subseteq \mathcal{P}(X)$. If $X = \bigcup_n F_n$ with $F_n \in \mathcal{A}$ and $\mu(F_n) < +\infty$, then any measures ν_1, ν_2 on $\sigma(\mathcal{A})$ with $\nu_1|_{\mathcal{A}} = \nu_2|_{\mathcal{A}} = \mu$ satisfy $\nu_1 = \nu_2$.

Proof. For $E \in \sigma(\mathcal{A})$, repeat the proof of Theorem 5.1.21 with ν_1, ν_2 in place of ν, μ^*; finite-mass comparison on $F_1 \cup \cdots \cup F_n$ and continuity from below give $\nu_1(E) = \nu_2(E)$.
Structural Commentary: σ-finite algebra data determine the generated measure.
Key Insight: Finite algebra exhaustions force uniqueness. □

Example 5.1.23. *Why σ-Finiteness Matters*
Let $X = \mathbb{Q} \cap (0, 1]$, and let $\mathcal{A}$ be the algebra of finite unions of sets $(a, b] \cap X$, where $0 \le a \le b \le 1$ and $a, b \in \mathbb{Q}$. Then every nonempty $A \in \mathcal{A}$ is infinite and $\sigma(\mathcal{A}) = \mathcal{P}(X)$, since each singleton $\{r\}$ is a countable intersection of sets $(a_n, r] \cap X$ with $a_n \uparrow r$. If μ_1 is counting measure and $\mu_2 = 2\mu_1$ on $\mathcal{P}(X)$, then $\mu_1|_{\mathcal{A}} = \mu_2|_{\mathcal{A}}$, but $\mu_1(\{r\}) \neq \mu_2(\{r\})$. Thus uniqueness may fail without σ-finiteness.
Key Insight: If the starting algebra has no finite nonempty pieces, it cannot detect the later singleton scale.

Proposition 5.1.24. *Half-Open Interval Algebra on $\mathbb{R}$*
Let $\mathcal{F}$ be the finite-union algebra generated by $(a, b]$, $(-\infty, b]$, $(a, +\infty)$, and $\mathbb{R}$, $a, b \in \mathbb{R}$. Define $\ell((a, b]) := b - a$ and $\ell((-\infty, b]) = \ell((a, +\infty)) = \ell(\mathbb{R}) := +\infty$, extended additively over finite disjoint unions. Then

(i) $\mathcal{F}$ is an algebra on $\mathbb{R}$ and $\sigma(\mathcal{F}) = \mathcal{B}(\mathbb{R})$.
(ii) ℓ is a premeasure on $\mathcal{F}$.

Proof.

(i) Generator complements lie in $\mathcal{F}$, hence $\mathcal{F}$ is an algebra.
Since $\mathcal{F} \subseteq \mathcal{B}(\mathbb{R})$ and $(a, b) = \bigcup_{n > 1/(b-a)} (a, b - 1/n]$, $\sigma(\mathcal{F}) = \mathcal{B}(\mathbb{R})$.

(ii) Normal forms give well-defined finite additivity. For countable additivity, reduce to $(a,b] = \bigsqcup_j (a_j, b_j]$: finite subunions give $\sum_j (b_j - a_j) \le b - a$, while compactness after enlarging to $(a_j - \varepsilon_j, b_j + \varepsilon_j)$ gives $b - a \le \sum_j (b_j - a_j) + \varepsilon$. Let $\varepsilon \downarrow 0$. Half-lines follow by truncation, finite disjoint unions componentwise.

Structural Commentary: Lebesgue measure starts as length on a finite interval algebra, not on all subsets.
Key Insight: Compactness supplies the reverse inequality in countable additivity of interval length. □

Remark 5.1.25. *Why the Half-Open Convention Matters*
Finite unions generated by (a,b), $(-\infty,b)$, $(a,+\infty)$, and $\mathbb{R}$ still generate $\mathcal{B}(\mathbb{R})$, but are not an algebra, where $(0,1)^c = (-\infty,0] \cup [1,+\infty)$. Half-open endpoints make complements remain in the starting family.

Definition 5.1.26. *Borel–Stieltjes and Lebesgue–Stieltjes Measures*
Let $g : \mathbb{R} \to \mathbb{R}$ be increasing and right-continuous, set $g(\pm\infty) := \lim_{x \to \pm\infty} g(x)$ in $\overline{\mathbb{R}}$, and let $\mathcal{F}$ be the algebra of finite unions generated by $(a,b]$, $(-\infty,b]$, $(a,+\infty)$, and $\mathbb{R}$. Define μ_g on generators by $\mu_g((a,b]) := g(b) - g(a)$, $\mu_g((-\infty,b]) := g(b) - g(-\infty)$, $\mu_g((a,+\infty)) := g(+\infty) - g(a)$, and $\mu_g(\mathbb{R}) := g(+\infty) - g(-\infty)$; extend to finite disjoint unions by summation. Its unique extension to $\mathcal{B}(\mathbb{R})$ is the *Borel–Stieltjes measure generated by g*; the completion of this Borel measure is the *Lebesgue–Stieltjes measure generated by g*.

Theorem 5.1.27. *Borel–Stieltjes Measure Existence*
Let $g : \mathbb{R} \to \mathbb{R}$ be increasing and right-continuous, $g(\pm\infty) := \lim_{x \to \pm\infty} g(x)$,
$\mathcal{I} := \{(a,b] : a < b\} \cup \{(-\infty,b] : b \in \mathbb{R}\} \cup \{(a,+\infty) : a \in \mathbb{R}\} \cup \{\mathbb{R}\}$, and let $\mathcal{F}$ be the finite-union algebra generated by $\mathcal{I}$. Define $G_g((a,b]) := g(b) - g(a)$, $G_g((-\infty,b]) := g(b) - g(-\infty)$, $G_g((a,+\infty)) := g(+\infty) - g(a)$, $G_g(\mathbb{R}) := g(+\infty) - g(-\infty)$, and $\mu_g(\bigsqcup_{j=1}^m I_j) := \sum_{j=1}^m G_g(I_j)$ for disjoint $I_j \in \mathcal{I}$. Then:

(i) μ_g is a σ-finite premeasure on $\mathcal{F}$ and has a unique extension to $\mathcal{B}(\mathbb{R})$, again denoted μ_g.
(ii) If $g = \mathrm{id}_{\mathbb{R}}$, then μ_g is Lebesgue measure.
(iii) If g is continuous, then $\mu_g(J) = g(b) - g(a)$ for $J \in \{(a,b),[a,b),(a,b],[a,b]\}$, $a < b$.

Proof.

(i) Finite normal-form refinement makes μ_g well defined and finitely additive: adjacent g-increments telescope and monotonicity gives positivity.
If $(a,b] = \bigsqcup_{j \ge 1} (a_j, b_j]$, finite additivity gives $\sum_{j=1}^N G_g((a_j,b_j]) \le G_g((a,b])$; conversely, covering $[a+\varepsilon, b]$ by finitely many $(a_j, b_j + \varepsilon_j)$ gives $g(b) - g(a+\varepsilon) \le \sum_{j \in K} (g(b_j + \varepsilon_j) - g(a_j))$ for finite $K \subseteq \mathbb{N}$, and right-continuity lets $\varepsilon, \varepsilon_j \downarrow 0$. Hence bounded generators are countably additive.
Half-lines follow by truncation and finite unions componentwise, so μ_g is premeasure. As $\mathbb{R} = \bigcup_{n \ge 1} (-n, n]$, $\mu_g((-n,n]) = g(n) - g(-n) < +\infty$, μ_g is σ-finite; Theorem 5.1.21 gives uniqueness.
(ii) If $g = \mathrm{id}_{\mathbb{R}}$, then $G_g((a,b]) = b - a$, so μ_g agrees with half-open length on $\mathcal{F}$.
Uniqueness identifies its Borel extension with Lebesgue measure.
(iii) If g is continuous, then $(x - 1/n, x] \downarrow \{x\}$ gives $\mu_g(\{x\}) = \lim_n (g(x) - g(x - 1/n)) = 0$.

Structural Commentary: Stieltjes measure replaces ordinary length $b - a$ by the increment $g(b) - g(a)$.
Key Insight: Right-continuity aligns the half-open interval convention with countable covers. □

Proposition 5.1.28. *Finite Borel Measures Give Stieltjes Generators*
Let ν be finite on $\mathcal{B}(\mathbb{R})$ and $g(x) := \nu((-\infty,x])$. Then g is increasing and right-continuous, $g(-\infty) = 0$, $g(+\infty) = \nu(\mathbb{R})$, $\nu((a,b]) = g(b) - g(a)$ for $a \le b$, and $\nu = \mu_g$ on $\mathcal{B}(\mathbb{R})$.

Proof. For $a \le b$, inclusion gives $g(a) \le g(b)$ and $(-\infty,b] = (-\infty,a] \sqcup (a,b]$, hence $\nu((a,b]) = g(b) - g(a)$.
If $x_n \downarrow x$, then $(-\infty, x_n] \downarrow (-\infty, x]$, so finite continuity from above gives $g(x_n) \downarrow g(x)$; thus g is right-continuous.
Also $(-\infty,-n] \downarrow \emptyset$ and $(-\infty,n] \uparrow \mathbb{R}$, whence $g(-\infty) = 0$ and $g(+\infty) = \nu(\mathbb{R})$.
Thus ν and μ_g agree on the generators of $\mathcal{F}$; Theorem 5.1.27 and uniqueness give $\nu = \mu_g$ on $\mathcal{B}(\mathbb{R})$.
Structural Commentary: A finite real-line measure is encoded by its cumulative mass function.
Key Insight: Right-continuity of distribution functions is continuity from above for measures. □

Example 5.1.29. *Basic Borel–Stieltjes Generators*
The following generators illustrate the interval rule in Theorem 5.1.27.

(i) If $g(x) = 2x$, then $\mu_g = 2m$ on $\mathcal{B}(\mathbb{R})$.
(ii) If $g(x) = \arctan x$, then $\mu_g((a,b]) = \arctan b - \arctan a = \int_a^b (1 + x^2)^{-1}\, dx$ and $\mu_g(\mathbb{R}) = \pi$.
(iii) If $g(x) = 0$ for $x < 0$ and $g(x) = 1$ for $x \ge 0$, then $\mu_g = \delta_0$.
(iv) If $g(x) = 0$ for $x < 0$ and $g(x) = x$ for $x \ge 0$, then $\mu_g(E) = m(E \cap [0,+\infty))$ for Borel E.

Proposition 5.1.30. *Elementary Measure Constructions*
Let $(X, \mathcal{M})$ be a measurable space.

(i) If μ is a measure on $\mathcal{M}$ and $A \in \mathcal{M}$, then $\mu_A(E) := \mu(A \cap E)$ defines a measure μ_A on $\mathcal{M}$.

(ii) If $\mu_1, \dots, \mu_m$ are measures on $\mathcal{M}$ and $a_1, \dots, a_m \geq 0$, then $\lambda := \sum_{j=1}^{m} a_j \mu_j$ is a measure on $\mathcal{M}$.
(iii) If $(\mu_n)_{n\geq 1}$ are measures on $\mathcal{M}$ with $\mu_n(X) = 1$, then $\lambda := \sum_{n=1}^{\infty} 2^{-n}\mu_n$ is a measure on $\mathcal{M}$ and $\lambda(X) = 1$.

Proof.

(i) $\mu_A(\emptyset) = 0$. If (E_k) is disjoint in $\mathcal{M}$, then $A \cap \bigsqcup_k E_k = \bigsqcup_k (A \cap E_k)$, whence $\mu_A(\bigsqcup_k E_k) = \sum_k \mu_A(E_k)$.
(ii) $\lambda(\emptyset) = 0$. For disjoint (E_k), finite summation gives $\lambda(\bigsqcup_k E_k) = \sum_{j=1}^{m} a_j \sum_k \mu_j(E_k) = \sum_k \lambda(E_k)$.
(iii) $\lambda(\emptyset) = 0$. For disjoint (E_k), nonnegative reordering gives $\lambda(\bigsqcup_k E_k) = \sum_n 2^{-n} \sum_k \mu_n(E_k) = \sum_k \lambda(E_k)$. Also $\lambda(X) = \sum_n 2^{-n}\mu_n(X) = \sum_n 2^{-n} = 1$.

Structural Commentary: New measures come from restriction, superposition, and countable convex mixtures.
Key Insight: Nonnegative series make the bookkeeping order-independent. □

Proposition 5.1.31. *Basic Laws for Additive Set Functions*
Let ϕ be additive on a ring $\mathcal{A}$. Then:

(i) $\phi(\emptyset) = 0$;
(ii) $\phi(A_1 \cup \cdots \cup A_n) = \phi(A_1) + \cdots + \phi(A_n)$ for pairwise disjoint $A_i \in \mathcal{A}$;
(iii) $\phi(A \cup B) + \phi(A \cap B) = \phi(A) + \phi(B)$;
(iv) if ϕ is nonnegative and $A \subseteq B$, then $\phi(A) \leq \phi(B)$;
(v) if $B \subseteq A$ and $|\phi(B)| < +\infty$, then $\phi(A \setminus B) = \phi(A) - \phi(B)$.

Proof.

(i) Apply additivity to the empty disjoint union.
(ii) Induct on n, applying additivity to $(A_1 \cup \cdots \cup A_{n-1}) \cup A_n$.
(iii) By disjoint decompositions, $A \cup B = (A \setminus B) \cup (B \setminus A) \cup (A \cap B)$; $A = (A \setminus B) \cup (A \cap B)$; $B = (B \setminus A) \cup (A \cap B)$
(iv) If $A \subseteq B$, then $B = A \cup (B \setminus A)$ disjointly, so nonnegativity gives $\phi(A) \leq \phi(B)$.
(v) If $B \subseteq A$, then $A = B \cup (A \setminus B)$ disjointly; since $|\phi(B)| < +\infty$, subtract $\phi(B)$.

Structural Commentary: Finite additivity plus disjoint decompositions controls all elementary identities.
Key Insight: Rewrite overlaps as disjoint pieces before applying ϕ. □

Theorem 5.1.32. *Continuity from Below*
Suppose ϕ is countably additive on a ring $\mathcal{A}$.
If $A_1 \subseteq A_2 \subseteq A_3 \subseteq \cdots$, and $A = \bigcup_{n=1}^{\infty} A_n$ with all $A_n, A \in \mathcal{A}$, then $\phi(A_n) \to \phi(A)$.

Proof. Put $B_1 = A_1$ and $B_n = A_n \setminus A_{n-1}$ for $n \geq 2$. Then B_n are pairwise disjoint, $A_n = B_1 \cup \cdots \cup B_n$, and $A = \bigcup_n B_n$. Countable additivity turns $\phi(A_n)$ and $\phi(A)$ into partial sums and full sum of same series.
Structural Commentary: Increasing unions are disjoint tail decompositions in disguise.
Key Insight: Countable additivity converts set convergence into series convergence. □

Definition 5.1.33. *Cells, Intervals, and Volume in $\mathbb{R}^p$*

(i) *Cells in $\mathbb{R}$*. For $a \leq b$, a cell with endpoints a, b is $I \in \{(a,b), [a,b], (a,b], [a,b)\}$, respectively open, closed, right-half-open, left-half-open; if $a = b$, only $[a,a]$ is nonempty.
(ii) *Length*. For a cell I with endpoints $a \leq b$, $\ell(I) := b - a$, independent of endpoint inclusion; $(-\infty, a)$, $(-\infty, a]$, $[b, +\infty)$, $(b, +\infty)$, and $\mathbb{R}$ are intervals, not cells, with length $+\infty$.
(iii) *Cells in $\mathbb{R}^p$*. For $p \geq 1$, a cell is $I = \prod_{i=1}^{p} I_i$, each I_i a cell in $\mathbb{R}$; its type is open, closed, or half-open, $\prod_i [a_i, b_i)$ is *southwest*, $\prod_i (a_i, b_i]$ is *northeast*, and a *cube* satisfies $\ell(I_1) = \cdots = \ell(I_p)$.
(iv) *Volume*. For $I = \prod_{i=1}^{p} I_i$, $\mathrm{vol}_p(I) := \prod_{i=1}^{p} \ell(I_i)$ and $\mathrm{vol}_p(\emptyset) := 0$.
If endpoints are $a_i \leq b_i$, then $\mathrm{vol}_p(I) = \prod_{i=1}^{p} (b_i - a_i)$.
(v) *Unbounded products*. For cells I, $\mathrm{vol}_p(I^\circ) = \mathrm{vol}_p(\overline{I}) = \mathrm{vol}_p(I)$.
For products $I = \prod_i I_i$ of possibly unbounded intervals, $\exists j : \ell(I_j) = 0$ gives volume 0, while $\ell(I_j) > 0$ for all j and $\exists j : \ell(I_j) = +\infty$ gives volume $+\infty$.

Remark 5.1.34. *Cells versus Intervals*
Cells are precisely bounded intervals with finite endpoints, so every cell is an interval but not conversely. Open cells are open in $\mathbb{R}^p$, closed cells are closed in $\mathbb{R}^p$, and nonempty half-open cells are neither.
In $\mathbb{R}$, $(-\infty, a)$ and $(b, +\infty)$ are open, $(-\infty, a]$ and $[b, +\infty)$ are closed, and $\mathbb{R} = (-\infty, +\infty)$ is clopen.
Key Insight: Topology likes open and closed endpoints; measure bookkeeping often likes half-open endpoints.

Proposition 5.1.35. *Half-Open Cell Calculus*
Fix $\sigma \in \{$southwest, northeast$\}$, and let $\mathcal{H}_\sigma$ be the half-open cells in $\mathbb{R}^p$ with orientation σ.

(i) Any finite Boolean combination of $H_1, \dots, H_n \in \mathcal{H}_\sigma$ is $\bigsqcup_{r=1}^{N} K_r$ with $K_r \in \mathcal{H}_\sigma$.
(ii) If $H \in \mathcal{H}_\sigma$, then $\mathbb{R}^p \setminus H = \bigsqcup_{r=1}^{\infty} K_r$ for some $K_r \in \mathcal{H}_\sigma$.

(iii) For $H \in \mathcal{H}_\sigma$ and $\delta > 0$, $H = \bigsqcup_{r=1}^{N} K_r$ for some $K_r \in \mathcal{H}_\sigma$ with $\operatorname{diam} K_r < \delta$.

Proof.

(i) For coordinate endpoints $E_j = \{c_{j,0} < \cdots < c_{j,m_j}\}$ of $H_1, \ldots, H_n$, set $[a,b)_\sigma = [a,b)$ if $\sigma =$ *southwest* and $[a,b)_\sigma = (a,b]$ if $\sigma =$ *northeast*, $\mathcal{P}_j = \{[c_{j,q-1}, c_{j,q})_\sigma : 1 \le q \le m_j\}$, and $\mathcal{G} = \{\prod_{j=1}^{p} P_j : P_j \in \mathcal{P}_j\}$; every finite Boolean combination is $\bigsqcup_{G \in \mathcal{G}'} G$, $\mathcal{G}' \subseteq \mathcal{G}$.

(ii) Write $H = \prod_{j=1}^{p} H_j$ and $D_s = (\prod_{j<s} H_j) \times (\mathbb{R} \setminus H_s) \times \mathbb{R}^{p-s}$; then $\mathbb{R}^p \setminus H = \bigsqcup_{s=1}^{p} D_s$. Decomposing each one-dimensional ray and $\mathbb{R}$ into countably many bounded σ-half-open intervals and taking products gives the required union; countability follows from Corollary 6.1.11.

(iii) Choose m with $2^{-m}(\sum_{j=1}^{p} \ell(H_j)^2)^{1/2} < \delta$. After m bisections per coordinate, each subcell has diameter $< \delta$; disjointness and volume preservation follow coordinatewise.

Structural Commentary: Half-open cells behave like grid atoms: finite set operations become finite grid selection.
Key Insight: Half-open endpoints prevent overlap while still partitioning boxes cleanly. □

Definition 5.1.36. *Translation of a Set*
If $A \subseteq \mathbb{R}^p$ and $t \in \mathbb{R}^p$, the translate of A by t is $t + A := \{t + a : a \in A\}$.

Proposition 5.1.37. *Cell Volume Is Translation Invariant*
Let $I = \prod_{i=1}^{p} I_i$ be a cell in $\mathbb{R}^p$ and $t = (t_i)_{i=1}^{p} \in \mathbb{R}^p$. Then $t + I = \prod_{i=1}^{p} (t_i + I_i)$ is a cell with the same endpoint pattern as I, and $\operatorname{vol}_p(t + I) = \operatorname{vol}_p(I)$.

Proof. Let $a_i \le b_i$ be the endpoints of I_i. Then $t_i + I_i$ has endpoints $t_i + a_i \le t_i + b_i$, unchanged endpoint inclusion, and $\ell(t_i + I_i) = (t_i + b_i) - (t_i + a_i) = \ell(I_i)$. Hence $\operatorname{vol}_p(t + I) = \prod_{i=1}^{p} \ell(t_i + I_i) = \prod_{i=1}^{p} \ell(I_i) = \operatorname{vol}_p(I)$.
Structural Commentary: Translation moves location but not side lengths.
Key Insight: Coordinate volume depends only on side lengths. □

Definition 5.1.38. *Elementary Sets and Regularity in $\mathbb{R}^p$*

(i) *Elementary sets.* $\mathcal{E} := \{\bigcup_{r=1}^{n} I_r : n \in \mathbb{N}\}$, where each $I_r \subseteq \mathbb{R}^p$ is a bounded cell.

(ii) *Elementary volume.* If $A = \bigsqcup_{r=1}^{n} I_r$ with cells I_r, set $m(A) := \sum_{r=1}^{n} \operatorname{vol}_p(I_r)$.

(iii) *Regularity.* A nonnegative additive $\mu : \mathcal{E} \to [0, +\infty]$ is regular iff $\forall A \in \mathcal{E}$, $\forall \varepsilon > 0$, $\exists F, G \in \mathcal{E}$ with $F \subseteq A \subseteq G$, F closed, G open, and $\mu(G) - \varepsilon \le \mu(A) \le \mu(F) + \varepsilon$.

$\mathcal{E}$ is a ring; every $A \in \mathcal{E}$ admits a finite disjoint cell decomposition; m is well defined, additive, and regular.

Perspective 5.1.39. *Box Volume and Linear Algebra*
For a box $B = \prod_i (a_i, b_i)$ or $\prod_i [a_i, b_i]$, $m(B) = \prod_i (b_i - a_i)$, the diagonal case of determinant volume scaling. Measure extends this product rule beyond boxes.
Key Insight: Lebesgue measure starts from determinant-style coordinate volume.

Perspective 5.1.40. *Outer Measure as Overestimation*
Elementary $A = \bigsqcup_{r=1}^{n} I_r$ have $m(A) = \sum_{r=1}^{n} \operatorname{vol}_p(I_r)$.
General $E \subseteq \mathbb{R}^p$, e.g. $D = \{(x, y) \in \mathbb{R}^2 : x^2 + y^2 \le 1\}$, need not be a finite disjoint cell union.
Replace decompositions by countable cell covers $E \subseteq \bigcup_{k=1}^{\infty} I_k$ and minimize $\sum_{k=1}^{\infty} \operatorname{vol}_p(I_k)$.
Key Insight: Outer covers size every $E \subseteq \mathbb{R}^p$ before measurability.

Definition 5.1.41. *Cell-Cover Outer Measure*
For $E \subseteq \mathbb{R}^p$, let $\mathcal{C}(E)$ be the countable cell covers $(I_k)_{k \ge 1}$ with $E \subseteq \bigcup_{k=1}^{\infty} I_k$. Define $m_p^*(E) := \inf_{(I_k) \in \mathcal{C}(E)} \sum_{k=1}^{\infty} \operatorname{vol}_p(I_k)$.

Here $\mathcal{C}(E) \neq \emptyset$ since $E \subseteq \bigcup_{j=1}^{\infty} [-j, j]^p$, and sum is nonnegative extended-real series, hence order-independent.

Proposition 5.1.42. *Equivalent Covering Families*
For $E \subseteq \mathbb{R}^p$, $\delta > 0$, and fixed half-open orientation σ, let $\mathcal{C}_{\mathsf{cl}}(E)$, $\mathcal{C}_{\mathsf{op}}(E)$, $\mathcal{C}_\sigma(E)$, and $\mathcal{C}_\delta(E)$ be the subfamilies of $\mathcal{C}(E)$ with cells respectively closed, open, σ-half-open, and satisfying $\operatorname{diam} I_k < \delta$.
If $\alpha(\mathcal{A}) := \inf_{(I_k) \in \mathcal{A}} \sum_{k=1}^{\infty} \operatorname{vol}_p(I_k)$, then $m_p^*(E) = \alpha(\mathcal{C}_{\mathsf{cl}}(E)) = \alpha(\mathcal{C}_{\mathsf{op}}(E)) = \alpha(\mathcal{C}_\sigma(E)) = \alpha(\mathcal{C}_\delta(E))$.

Proof. Let $\mathcal{C}_*(E)$ be any restricted class above. Since $\mathcal{C}_*(E) \subseteq \mathcal{C}(E)$, $\alpha(\mathcal{C}_*(E)) \ge m_p^*(E)$.
Conversely, for $(I_k) \in \mathcal{C}(E)$ and $\varepsilon > 0$, closures give $\overline{I_k} \supseteq I_k$ and $\operatorname{vol}_p(\overline{I_k}) = \operatorname{vol}_p(I_k)$; open cells $J_k \supseteq I_k$ may be chosen with $\operatorname{vol}_p(J_k) \le \operatorname{vol}_p(I_k) + \varepsilon 2^{-k}$; each J_k admits a countable disjoint σ-half-open decomposition with total volume $\operatorname{vol}_p(J_k)$; each half-open cell has a finite disjoint subdivision into cells of diameter $< \delta$ with unchanged total volume. Thus $\alpha(\mathcal{C}_*(E)) \le \sum_k \operatorname{vol}_p(I_k) + \varepsilon$; take infima over (I_k) and let $\varepsilon \downarrow 0$.
Structural Commentary: Near-optimal covers ignore endpoint convention.
Key Insight: Closures, small enlargements, and fine subdivisions preserve the infimum. □

Theorem 5.1.43. *Basic Properties of Cell-Cover Outer Measure*
For $E, F, E_n \subseteq \mathbb{R}^p$,

(i) Null empty set. $0 \le m_p^*(E) \le +\infty$ *and* $m_p^*(\emptyset) = 0$.
(ii) Monotonicity. $E \subseteq F \Longrightarrow m_p^*(E) \le m_p^*(F)$.
(iii) Subadditivity. $m_p^*(\bigcup_{n=1}^{\infty} E_n) \le \sum_{n=1}^{\infty} m_p^*(E_n)$.

Proof.

(i) Since all cover sums lie in $[0, +\infty]$, $0 \le m_p^*(E) \le +\infty$; the cover $(\emptyset, \emptyset, \ldots)$ gives $m_p^*(\emptyset) = 0$.
(ii) If $E \subseteq F$, then $\mathcal{C}(F) \subseteq \mathcal{C}(E)$, hence $\inf_{\mathcal{C}(E)} \le \inf_{\mathcal{C}(F)}$.
(iii) If some $m_p^*(E_n) = +\infty$, the claim is trivial.
Otherwise, for $\varepsilon > 0$, choose $(I_k^n)_{k\ge1} \in \mathcal{C}(E_n)$ with $\sum_k \mathrm{vol}_p(I_k^n) \le m_p^*(E_n) + \varepsilon 2^{-n}$.
Then $(I_k^n)_{(n,k)\in\mathbb{N}^2}$ countable, covers $\bigcup_n E_n$, and $m_p^*(\bigcup_n E_n) \le \sum_{n,k} \mathrm{vol}_p(I_k^n) \le \sum_n m_p^*(E_n) + \varepsilon$.

Structural Commentary: Countable subadditivity is proved by merging near-optimal covers.
Key Insight: Nonnegative double sums may be scanned in one sequence. □

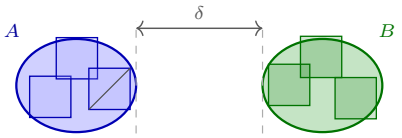


Figure 189: Positive distance separates cells with $\operatorname{diam} I_k < \delta$.

Theorem 5.1.44. *Additivity at Positive Distance*
Let $A, B \subseteq \mathbb{R}^p$ be nonempty and disjoint, and suppose $\operatorname{dist}(A, B) := \inf\{\|a - b\|_2 : a \in A,\ b \in B\} > 0$.
Then $m_p^*(A \cup B) = m_p^*(A) + m_p^*(B)$.

Proof. The inequality $\le$ follows from Theorem 5.1.43. For $\ge$, set $\delta := \operatorname{dist}(A, B)$. If $m_p^*(A \cup B) = +\infty$, done. Otherwise, given $\varepsilon > 0$, choose $(I_n) \in \mathcal{C}(A \cup B)$ with $\operatorname{diam} I_n < \delta$ and $\sum_n \mathrm{vol}_p(I_n) \le m_p^*(A \cup B) + \varepsilon$.
Put $N_A := \{n : I_n \cap A \ne \emptyset\}$ and $N_B := \{n : I_n \cap B \ne \emptyset\}$. Since $N_A \cap N_B = \emptyset$, $(I_n)_{n\in N_A}$ covers A, $(I_n)_{n\in N_B}$ covers B, and $m_p^*(A) + m_p^*(B) \le \sum_{n\in N_A} \mathrm{vol}_p(I_n) + \sum_{n\in N_B} \mathrm{vol}_p(I_n) \le \sum_n \mathrm{vol}_p(I_n) \le m_p^*(A \cup B) + \varepsilon$. Let $\varepsilon \downarrow 0$.
Structural Commentary: Positive separation lets a sufficiently fine cover be sorted unambiguously.
Key Insight: Countable subadditivity becomes equality when no small cell can touch both pieces. □

Theorem 5.1.45. *Outer Measure of a Cell*
If I is a cell in $\mathbb{R}^p$, then $m_p^*(I) = \mathrm{vol}_p(I)$.

Proof. The cover $(I, \emptyset, \ldots)$ gives $m_p^*(I) \le \mathrm{vol}_p(I)$. For $\ge$, the case $\mathrm{vol}_p(I) = 0$ follows by nonnegativity.
Assume $\mathrm{vol}_p(I) > 0$ and fix $\varepsilon > 0$. Choose a closed cell $J \subseteq I$ with $\mathrm{vol}_p(I) - \varepsilon < \mathrm{vol}_p(J)$, taking $J = I$ when I is closed. For any open-cell cover (G_k) of I, compactness gives $J \subseteq \bigcup_{k=1}^m G_k$.
Let $\mathcal{Q}$ be the finite rectangular grid determined by the faces of $J, \overline{G_1}, \ldots, \overline{G_m}$. Up to zero-volume boundary atoms, each $Q \in \mathcal{Q}$ with $Q \subseteq J$ lies in some G_k, hence $\mathrm{vol}_p(J) \le \sum_{k=1}^m \mathrm{vol}_p(G_k) \le \sum_{k=1}^{\infty} \mathrm{vol}_p(G_k)$.
Infimize over open-cell covers, use Proposition 5.1.42, and let $\varepsilon \downarrow 0$.
Structural Commentary: Compact inner cells reduce covers to finite grids.
Key Insight: Outer measure is calibrated so that cells keep their original volume. □

Theorem 5.1.46. *Translation Invariance of Outer Measure*
For $E \subseteq \mathbb{R}^p$ and $t \in \mathbb{R}^p$, $m_p^*(t + E) = m_p^*(E)$.

Proof. The map $\mathcal{C}(E) \to \mathcal{C}(t + E)$, $(I_k) \mapsto (t + I_k)$, is bijective with inverse $(J_k) \mapsto (-t + J_k)$.
By Proposition 5.1.37, $\sum_k \mathrm{vol}_p(t + I_k) = \sum_k \mathrm{vol}_p(I_k)$, so the two infima are equal.
Structural Commentary: Translate covers forward and backward.
Key Insight: Translation invariance of cell volume lifts directly to translation invariance of outer measure. □

Definition 5.1.47. *Outer Measure and μ-Measurable Sets*
Let μ be additive, regular, nonnegative, and finite on $\mathcal{E}$.
For $E \subseteq \mathbb{R}^p$, define $\mu^*(E) := \inf \sum_{n=1}^{\infty} \mu(A_n)$, where infimum is over open elementary covers $E \subseteq \bigcup_{n=1}^{\infty} A_n$.
For $A, B \subseteq \mathbb{R}^p$, set $A \triangle B := (A \setminus B) \cup (B \setminus A)$ and $d_\mu(A, B) := \mu^*(A \triangle B)$.
A set is *finitely μ-measurable* if it is a d_μ-limit of elementary sets.
It is *μ-measurable* if it is a countable union of such sets; denote the classes by $\mathcal{M}_F(\mu)$ and $\mathcal{M}(\mu)$.

Definition 5.1.48. *Box-Cover Outer Measure Model*
An *open box* in $\mathbb{R}^n$ is $B = \prod_{i=1}^{n} (a_i, b_i) = \{(x_1, \ldots, x_n) \in \mathbb{R}^n : x_i \in (a_i, b_i)\ \forall i\}$, $b_i \geq a_i$, $\mathrm{vol}_n(B) := \prod_{i=1}^{n} (b_i - a_i)$.
For $\Omega \subseteq \mathbb{R}^n$, the *box-cover outer measure* is $m_n^*(\Omega) := \inf\left\{\sum_{j\in J} \mathrm{vol}_n(B_j) : (B_j)_{j\in J} \text{ covers } \Omega,\ J \text{ at most countable}\right\}$.

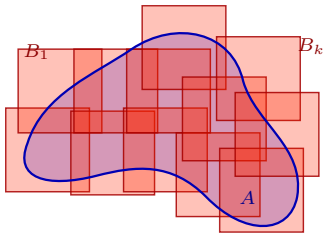


Figure 190: Cover $A \subseteq \bigcup_k B_k$ by squares; minimize $\sum_k |B_k|$.

Perspective 5.1.49. *Outer Measure as Outside Approximation*
Outer measure covers A from outside; $d_\mu(A, B) = \mu^*(A \triangle B)$ measures closeness.
Key Insight: Countable subadditivity collects near-optimal covers.

Method 5.1.50. *Lebesgue Measure Construction*

(i) *Cover from outside*
For $A \subseteq \mathbb{R}^n$, define $m_n^*(A) := \inf\left\{\sum_{j\in J} \mathrm{vol}_n(B_j) : (B_j)_{j\in J} \text{ covers } A,\ J \text{ at most countable}\right\}$.
This is defined for every set, but it is only countably subadditive in general.

(ii) *Keep the clean splitters*
Call A measurable when every test set $X \subseteq \mathbb{R}^n$ splits additively: $m_n^*(X) = m_n^*(X \cap A) + m_n^*(X \cap A^c)$.
These measurable sets form a σ-algebra.

(iii) *Restrict the outer measure*
Define Lebesgue measure on measurable sets by $m(A) := m_n^*(A)$.
On this restricted class, m is countably additive.

Key Insight: Outer measure sizes all sets; clean splitting restores additivity.

Theorem 5.1.51. *Extension to a Measure*
Let μ be additive, regular, nonnegative, and finite on $\mathcal{E}$. Form μ^* from open elementary covers, set $d_\mu(A, B) = \mu^*(A \triangle B)$, and let $\mathcal{M}(\mu)$ be the countable unions of d_μ-limits of elementary sets. Then:

(i) $\mu^*(A) = \mu(A)$ for $A \in \mathcal{E}$;
(ii) μ^* is countably subadditive;
(iii) $\mathcal{M}(\mu)$ is a σ-ring;
(iv) μ^* is countably additive on $\mathcal{M}(\mu)$.

Proof.

(i) Regularity gives $\mu^*(A) \leq \mu(A)$; compact closed approximants inside A give the reverse inequality.
(ii) Cover each E_n within $2^{-n}\varepsilon$ of $\mu^*(E_n)$; the combined cover proves subadditivity.
(iii) If $A_k \to A$, $B_k \to B$ in d_μ, $A_k, B_k \in \mathcal{E}$, then $d_\mu(A_k \diamond B_k, A \diamond B) \leq d_\mu(A_k, A) + d_\mu(B_k, B) \to 0$, $\diamond \in \{\cup, \cap, \setminus\}$.
Thus $\mathcal{M}_F(\mu)$ is a ring. If $A = \bigcup_k A_k$ with $A_k \in \mathcal{M}_F(\mu)$, then $A \setminus B = \bigcup_k (A_k \setminus B)$, so $\mathcal{M}(\mu)$ is a σ-ring.
(iv) For $A = \bigsqcup_n A_n$, $A_n \in \mathcal{M}_F(\mu)$, subadditivity gives $\leq$,

$$\text{while finite additivity gives } \mu^*(A) \geq \mu^*\left(\bigsqcup_{n=1}^{N} A_n\right) = \sum_{n=1}^{N} \mu^*(A_n).$$

Let $N \to \infty$, then decompose general disjoint families in $\mathcal{M}(\mu)$ into finite-measurable pieces.

Structural Commentary: Elementary sets are completed under the distance measured by symmetric difference.
Key Insight: Approximate by elementary sets, prove identities there, then pass through $d_\mu(A, B)$. □

Definition 5.1.52. *Extended and Lebesgue Measure*
The *extended measure* is $\mu^*|_{\mathcal{M}(\mu)}$, denoted again by μ.
For the elementary volume m, it is *Lebesgue measure* on $\mathbb{R}^p$.

Definition 5.1.53. *Abstract Outer Measure and Clean Division*
An *abstract outer measure* on M is a monotone, countably subadditive function $\omega : 2^M \to [0, +\infty]$ with $\omega(\emptyset) = 0$.
A set $E \subseteq M$ is *ω-measurable* if it cleanly splits every $X \subseteq M$: $\omega(X) = \omega(X \cap E) + \omega(X \cap E^c)$.
For Lebesgue outer measure, this is Caratheodory measurability.

Lemma 5.1.54. *Finite-Measure Test for Caratheodory Measurability*
Let $\omega : 2^M \to [0, +\infty]$ be an outer measure. Then $E \subseteq M$ is ω-measurable iff, for all $X \subseteq M$, $\omega(X) < +\infty \Longrightarrow \omega(X) \geq \omega(X \cap E) + \omega(X \cap E^c)$.

Proof.

($\Rightarrow$) If E is ω-measurable, then $\omega(X) = \omega(X \cap E) + \omega(X \cap E^c)$ $\forall X \subseteq M$, hence finite-test inequality holds.
($\Leftarrow$) For all $X \subseteq M$, $X = (X \cap E) \sqcup (X \cap E^c)$ gives $\omega(X) \leq \omega(X \cap E) + \omega(X \cap E^c)$. If $\omega(X) < +\infty$, the hypothesis gives $\geq$; if $\omega(X) = +\infty$, $\geq$ is automatic in $[0, +\infty]$. Thus equality holds for every X.

Structural Commentary: To verify a clean split, only finite outer-measure test sets need real work.
Key Insight: Subadditivity supplies one inequality for free. □

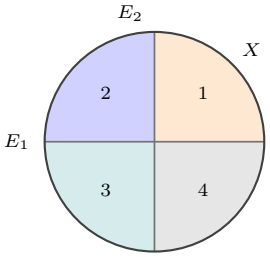


Figure 191: Two cuts split X into four regions; $E_1 \setminus E_2$ is region 2.

Definition 5.1.55. *Zero Sets*
For outer measure $\omega : 2^M \to [0, +\infty]$, *$\omega$-zero set* is $Z \subseteq M$ with $\omega(Z) = 0$; $\mathcal{Z}_\omega := \{Z \subseteq M : \omega(Z) = 0\}$.

Theorem 5.1.56. *Caratheodory Measurable Sets*
Let $\omega : 2^M \to [0, +\infty]$ be an outer measure, and let $\mathcal{M}_\omega$ be the class of ω-measurable subsets of M. Then $\mathcal{M}_\omega$ is a σ-algebra, $\omega|_{\mathcal{M}_\omega}$ is countably additive, $\mathcal{Z}_\omega \subseteq \mathcal{M}_\omega$, and $E \in \mathcal{M}_\omega$, $Z \in \mathcal{Z}_\omega$ imply $E \cup Z, E \setminus Z, E \triangle Z \in \mathcal{M}_\omega$ and $\omega(E \cup Z) = \omega(E \setminus Z) = \omega(E \triangle Z) = \omega(E)$.

Proof. If $Z \in \mathcal{Z}_\omega$, then $\omega(X) \leq \omega(X \cap Z) + \omega(X \cap Z^c) \leq 0 + \omega(X)$, so $Z \in \mathcal{M}_\omega$. Complements are immediate.
If $E, F \in \mathcal{M}_\omega$, iterating the two clean cuts gives $\omega(X) = \sum_{\epsilon,\eta \in \{0,1\}} \omega(X \cap E^\epsilon \cap F^\eta)$, where $E^1 = E$, $E^0 = E^c$; hence all finite Boolean combinations of E, F are measurable.
For disjoint $E_i \in \mathcal{M}_\omega$, set $E = \bigsqcup_i E_i$ and $F_N = (\bigcup_{i=1}^N E_i)^c$.
Finite splitting gives $\omega(X) = \sum_{i=1}^N \omega(X \cap E_i) + \omega(X \cap F_N) \geq \sum_{i=1}^N \omega(X \cap E_i) + \omega(X \cap E^c)$.
Letting $N \to \infty$ and using subadditivity for the reverse inequality yields $\omega(X) = \sum_{i=1}^\infty \omega(X \cap E_i) + \omega(X \cap E^c)$; hence $E \in \mathcal{M}_\omega$, and $X = E$ gives countable additivity.
General countable unions follow from $F_1 = E_1$, $F_n = E_n \setminus \bigcup_{i<n} E_i$.
Finally, if $E \in \mathcal{M}_\omega$ and $Z \in \mathcal{Z}_\omega$, then $E = (E \setminus Z) \sqcup (E \cap Z)$ and $E \cup Z = E \sqcup (Z \setminus E)$ give $\omega(E \setminus Z) = \omega(E)$ and $\omega(E \cup Z) = \omega(E)$; also $E \triangle Z = (E \setminus Z) \sqcup (Z \setminus E)$ has measure $\omega(E)$.
Structural Commentary: Measurable sets preserve countable additivity.
Key Insight: Measurable cuts turn subadditivity into equality. □

Definition 5.1.57. *Lebesgue σ-Algebra via Caratheodory*

(i) $\mathcal{L}^n := \{E \subseteq \mathbb{R}^n : E \text{ is } m_n^*\text{-measurable}\}$ is the *Lebesgue σ-algebra*.
(ii) $m := m_n^*|_{\mathcal{L}^n}$ is *Lebesgue measure*; for $E \in \mathcal{L}^n$, write $m(E) = m_n^*(E)$.

Proposition 5.1.58. *Cells Are Lebesgue Measurable*
For every coordinate cell $I \subseteq \mathbb{R}^n$, $I \in \mathcal{L}^n$ and $m(I) = \mathrm{vol}_n(I)$.

Proof. Since $I \in \mathcal{E}$, Theorem 5.1.51 gives $I \in \mathcal{L}^n$ and $m(I) = m_n^*(I) = \mathrm{vol}_n(I)$.
For the open-cell split, fix $A \subseteq \mathbb{R}^n$ with $m_n^*(A) < +\infty$.
Set $I_k := \{x \in I : \mathrm{dist}(x, I^c) > k^{-1}\}$, $\eta_k := m_n^*(I \setminus I_k) \to 0$. Since $\mathrm{dist}(A \cap I_k, A \setminus I) > 0$, Theorem 5.1.44 gives $m_n^*(A) \geq m_n^*(A \cap I_k) + m_n^*(A \setminus I) \geq m_n^*(A \cap I) + m_n^*(A \setminus I) - \eta_k$; let $k \to \infty$ and apply Lemma 5.1.54.
Structural Commentary: Cells are the bridge between coordinate volume and Lebesgue measure.
Key Insight: Shrink, split, let the collar vanish. □

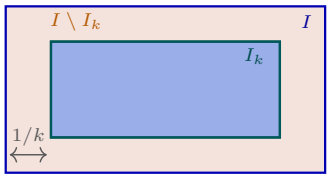


Figure 192: Inner cells I_k exhaust I; the collar $I \setminus I_k$ has vanishing volume.

Lemma 5.1.59. *Open Sets Are Countable Unions of Rational Boxes*
For every open $G \subseteq \mathbb{R}^n$, there is a countable $\Sigma \subseteq \{\prod_{i=1}^n (a_i, b_i) : a_i, b_i \in \mathbb{Q},\ a_i < b_i\}$ such that $G = \bigcup_{B\in\Sigma} B$.

Proof. Let $\mathcal{B}_{\mathbb{Q}} := \{\prod_{i=1}^n (a_i, b_i) : a_i, b_i \in \mathbb{Q},\ a_i < b_i\}$; then $\mathcal{B}_{\mathbb{Q}}$ is countable since $\mathbb{Q}^{2n}$ is countable.
Set $\Sigma := \{B \in \mathcal{B}_{\mathbb{Q}} : B \subseteq G\}$. The inclusion $\bigcup_{B\in\Sigma} B \subseteq G$ is immediate.
If $x \in G$, choose $r > 0$ with $B_r(x) \subseteq G$, then choose $a_i, b_i \in \mathbb{Q}$ with $x_i - r/n < a_i < x_i < b_i < x_i + r/n$.
Thus $x \in B := \prod_i (a_i, b_i) \subseteq B_r(x) \subseteq G$, so $x \in \bigcup_{B\in\Sigma} B$.
Structural Commentary: Rational boxes give a countable Euclidean base.
Key Insight: Openness supplies a ball; density of $\mathbb{Q}$ supplies rational box endpoints inside it. □

Theorem 5.1.60. *Lebesgue Measure Basic Properties*
Let $m = m_n^*|_{\mathcal{L}^n}$.

(i) $\mathcal{B}(\mathbb{R}^n) \subseteq \mathcal{L}^n$, and $\mathcal{L}^n$ is closed under $E \mapsto E^c$, countable $\cup$, and countable $\cap$.
(ii) For $E, F, E_k \in \mathcal{L}^n$, $0 \le m(E)$, $E \subseteq F \Rightarrow m(E) \le m(F)$, and $m(\bigcup_k E_k) \le \sum_k m(E_k)$.
(iii) If $(E_k) \subseteq \mathcal{L}^n$ are pairwise disjoint, then $m(\bigsqcup_k E_k) = \sum_k m(E_k)$.
(iv) $m([0,1]^n) = 1$, and $E \in \mathcal{L}^n$, $t \in \mathbb{R}^n$ imply $E + t \in \mathcal{L}^n$ and $m(E + t) = m(E)$.

Proof.

(i) By Theorems 5.1.51, 5.1.56, $\mathcal{L}^n$ is a σ-algebra, hence closed under complements and countable unions/intersections. Open boxes are elementary, hence in $\mathcal{L}^n$; by Lemma 5.1.59, each open $G = \bigcup_j B_j$ with rational open boxes B_j, so $G \in \mathcal{L}^n$. Thus $\mathcal{B}(\mathbb{R}^n) = \sigma(\mathcal{O}_{\mathbb{R}^n}) \subseteq \mathcal{L}^n$.
(ii) Since $m = m_n^*|_{\mathcal{L}^n}$, nonnegativity follows from $m_n^* \ge 0$. If $E \subseteq F$, then $m(E) \le m(F)$ by outer-measure monotonicity. Also $m(\bigcup_k E_k) = m_n^*(\bigcup_k E_k) \le \sum_k m_n^*(E_k) = \sum_k m(E_k)$ by outer-measure subadditivity.
(iii) If $(E_k) \subseteq \mathcal{L}^n$ pairwise disjoint, by Theorem 5.1.56, m countably additive on $\mathcal{L}^n$, so $m(\bigsqcup_k E_k) = \sum_k m(E_k)$.
(iv) By Theorem 5.1.45, since $[0,1]^n$ is a cell, $m([0,1]^n) = \mathrm{vol}_n([0,1]^n) = 1$.
For translation, let $A \subseteq \mathbb{R}^n$. By Theorem 5.1.46 and
$E \in \mathcal{L}^n$, $m_n^*(A) = m_n^*(A-t) = m_n^*((A-t)\cap E) + m_n^*((A-t)\cap E^c) = m_n^*(A\cap(E+t)) + m_n^*(A\cap(E+t)^c)$;
hence $E + t \in \mathcal{L}^n$, and $m(E + t) = m_n^*(E + t) = m_n^*(E) = m(E)$.

Structural Commentary: The wish list is recovered only after restricting outer measure to measurable sets.
Key Insight: The box-cover construction gives size; the measurable-set condition restores additivity. □

Theorem 5.1.61. *Uniqueness from Open Cells*
Let ν be a measure on $\mathcal{L}^n$ with $\nu(I) = \mathrm{vol}_n(I)$ for every open cell $I \subseteq \mathbb{R}^n$. Then $\nu = m$ on $\mathcal{L}^n$.

Proof. Fix $N \in \mathbb{N}$, $Q_N = (-N, N)^n$. If $E \in \mathcal{L}^n$ and $E \subseteq Q_N$, then any open-cell cover $E \subseteq \bigcup_j J_j$ gives $\nu(E) \le \sum_j \nu(J_j) = \sum_j \mathrm{vol}_n(J_j)$, hence $\nu(E) \le m(E)$ by infimum.
Applying this to E and $Q_N \setminus E$, and using $\nu(Q_N) = m(Q_N) < +\infty$, yields $\nu(E)+\nu(Q_N\setminus E) = m(E)+m(Q_N\setminus E)$ with each left summand $\le$ its right counterpart; hence $\nu(E) = m(E)$.
For arbitrary $E \in \mathcal{L}^n$, set $A_1 = E \cap Q_1$, $A_N = E \cap (Q_N \setminus Q_{N-1})$ for $N > 1$.
Then $E = \bigsqcup_N A_N$, each $A_N \subseteq Q_N$, and countable additivity gives $\nu(E) = \sum_N \nu(A_N) = \sum_N m(A_N) = m(E)$.
Structural Commentary: Agreement on coordinate cells determines the whole Lebesgue measure.
Key Insight: Covers give $\nu \le m$ locally; finite complements force equality. □

Method 5.1.62. *Translation Invariance by Moving the Test Set*

(i) For $t \in \mathbb{R}^n$ and $A \subseteq \mathbb{R}^n$, translated box covers give $m_n^*(A + t) = m_n^*(A)$.
(ii) If $E \in \mathcal{L}^n$, then for all $A \subseteq \mathbb{R}^n$, $(A \cap (E + t)) - t = (A - t) \cap E$ and $(A \cap (E + t)^c) - t = (A - t) \cap E^c$.
(iii) Apply the Caratheodory split for E to $A - t$ to get $E + t \in \mathcal{L}^n$ and $m(E + t) = m(E)$.

Key Insight: Translation invariance of outer measure transports the clean-splitting identity.

Proposition 5.1.63. *First Borel Hierarchy Examples*
If $G = \bigcap_n U_n$ with U_n open or $F = \bigcup_n C_n$ with C_n closed, then $G, F \in \mathcal{B}(\mathbb{R}^p) \subseteq \mathcal{L}^p$. More generally, countable $\cup, \cap$ operations starting from open or closed sets stay Borel. In $\mathbb{R}$, $\{0\} = \bigcap_{n\ge1}(-1/n, 1/n)$ is G_δ but not open, and $(0, 1) = \bigcup_{n\ge2}[1/n, 1 - 1/n]$ is F_σ but not closed.

Proof. Since $\mathcal{B}(\mathbb{R}^p) = \sigma(\mathcal{O}_{\mathbb{R}^p})$, it contains open sets, closed sets by complements, and countable $\cup, \cap$ by σ-algebra closure; Proposition 5.1.3 gives the stated permanence, and Theorem 5.1.60 gives $\mathcal{B}(\mathbb{R}^p) \subseteq \mathcal{L}^p$. The identities follow from $x = 0 \iff |x| < 1/n$ for all n and $0 < x < 1 \iff \exists n \geq 2 : 1/n \leq x \leq 1 - 1/n$; $\{0\}$ is not open and $(0, 1)$ is not closed in the usual topology.
Structural Commentary: Borel sets are generated by countable topological operations.
Key Insight: Countable intersections and unions quickly leave the class of open and closed sets. □

Perspective 5.1.64. *Why Not Measure Every Subset?*
Vitali-type sets show m_1^* is not additive on $2^{\mathbb{R}}$: disjoint unions and refinements can break additivity.
Thus measure is m^* restricted to clean splitters E, where $m^*(X) = m^*(X \cap E) + m^*(X \cap E^c)$ for every test set X.
Key Insight: Outer measure is universal; measure lives on clean splitters.

Corollary 5.1.65. *Measure Continuity*
Let μ be a measure. If $E_k \uparrow E$, then $\mu(E_k) \uparrow \mu(E)$. If $F_k \downarrow F$ and $\mu(F_1) < +\infty$, then $\mu(F_k) \downarrow \mu(F)$.

Proof. Disjointize $E_k \uparrow E$ by $E_1' = E_1$ and $E_k' = E_k \setminus E_{k-1}$; countable additivity makes $\mu(E_k)$ the kth partial sum for $\mu(E)$. For $F_k \downarrow F$, apply this to $F_1 \setminus F_k \uparrow F_1 \setminus F$ and subtract from finite $\mu(F_1)$.
Structural Commentary: Monotone set limits are disjoint series or complements of such series.
Key Insight: Downward continuity needs one finite anchor so subtraction is legal. □

Definition 5.1.66. *Null Sets and Real-Line Measure Zero*
In a measure space, N is *null* if it is measurable and $\mu(N) = 0$.
A set $A \subseteq \mathbb{R}$ has *measure zero* if, for every $\varepsilon > 0$, open intervals O_n cover A and satisfy $\sum_{n=1}^{\infty} |O_n| \leq \varepsilon$.
For Lebesgue measure, null sets are also *zero sets*.
A subset of a compact interval has *content zero* if the same condition holds with finitely many intervals.

Proposition 5.1.67. *Outer-Null Sets Are Lebesgue Measurable*
If $Z \subseteq \mathbb{R}^n$ and $m_n^*(Z) = 0$, then $Z \in \mathcal{L}^n$, $m(Z) = 0$, and $\mathcal{P}(Z) \subseteq \mathcal{L}^n$ with $m(W) = 0$ for every $W \subseteq Z$.

Proof. For $A \subseteq \mathbb{R}^n$, monotonicity gives $m_n^*(A \cap Z) = 0$.
Hence $m_n^*(A) \leq m_n^*(A \cap Z) + m_n^*(A \cap Z^c) = m_n^*(A \cap Z^c) \leq m_n^*(A)$, Z satisfies Caratheodory split and $m(Z) = 0$.
If $W \subseteq Z$, then $m_n^*(W) \leq m_n^*(Z) = 0$, and the same argument gives $W \in \mathcal{L}^n$, $m(W) = 0$.
Structural Commentary: Lebesgue measure is complete because outer-null subsets split every test set trivially.
Key Insight: A zero piece cannot disturb additivity. □

Proposition 5.1.68. *Basic Measure-Zero Examples on $\mathbb{R}$*
Every finite $A \subseteq \mathbb{R}$ has content 0; every countable $A \subseteq \mathbb{R}$ has measure 0; Cantor set P has content 0.

Proof. Let $\varepsilon > 0$. If $A = \{a_1, \ldots, a_N\}$ with $N \geq 1$, choose open intervals $I_j \ni a_j$ with $|I_j| < \varepsilon/N$; then $A \subseteq \bigcup_{j=1}^N I_j$ and $\sum_{j=1}^N |I_j| < \varepsilon$. The case $A = \emptyset$ is trivial.
If $A = \{a_j : j \in \mathbb{N}\}$, choose open intervals $I_j \ni a_j$ with $|I_j| < \varepsilon 2^{-j}$; then $A \subseteq \bigcup_j I_j$ and $\sum_j |I_j| < \varepsilon$.
For P, choose N with $(2/3)^N < \varepsilon/2$. Write $E_N = \bigcup_{k=1}^{2^N} J_k$, where each J_k is closed and $|J_k| = 3^{-N}$. Choose open $I_k \supseteq J_k$ with $|I_k| < 3^{-N} + \varepsilon/2^{N+1}$. Since $P \subseteq E_N$, the I_k cover P, and $\sum_{k=1}^{2^N} |I_k| < \left(\frac{2}{3}\right)^N + \frac{\varepsilon}{2} < \varepsilon$.
Structural Commentary: Zero size means arbitrarily small interval covers.
Key Insight: Lists split ε by terms; Cantor stages have length $(2/3)^N \to 0$. □

Corollary 5.1.69. *Uncountable Null Sets Exist*
There exists $P \subseteq \mathbb{R}$ with $|P| = 2^{\aleph_0}$ and $m(P) = 0$.

Proof. Let $P := \{\sum_{k=1}^{\infty} \varepsilon_k 3^{-k} : \varepsilon_k \in \{0, 2\}\}$. The map $(\varepsilon_k) \mapsto \sum_k \varepsilon_k 3^{-k}$ identifies P with $\{0, 2\}^{\mathbb{N}}$, so $|P| = 2^{\aleph_0}$.
At stage N, P is covered by 2^N closed intervals of length 3^{-N}, hence $m_1^*(P) \leq (2/3)^N$ for all N; letting $N \to \infty$ gives $m_1^*(P) = 0$. By outer-null completeness, $P \in \mathcal{L}^1$ and $m(P) = 0$.
Structural Commentary: Cantor coding separates cardinality from length.
Key Insight: Null need not mean countable. □

Corollary 5.1.70. *Real-Line Lebesgue Measure Checks*
Let m be Lebesgue measure on $\mathbb{R}$.

(i) $a < b \Rightarrow m((a, b)) = m((a, b]) = m([a, b)) = m([a, b]) = b - a$.
(ii) $A \subseteq \mathbb{R}$ *countable* $\Rightarrow A \in \mathcal{B}(\mathbb{R})$ *and* $m(A) = 0$.
(iii) $\emptyset \neq G \subseteq \mathbb{R}$ *open* $\Rightarrow m(G) > 0$.
(iv) $K \subseteq \mathbb{R}$ *compact* $\Rightarrow m(K) < +\infty$.

Proof.

(i) By cell volume, $m([a,b]) = b-a$; endpoint sets finite null; add or delete $\{a\}, \{b\}$ does not change measure.
(ii) If $A = \{a_k : k \in J\}$ with J finite or countable, then $A = \bigcup_{k\in J}\{a_k\} \in \mathcal{B}(\mathbb{R})$; $m(A) \le \sum_{k\in J} m(\{a_k\}) = 0$.
(iii) Choose $(a,b) \subseteq G$ with $a < b$; by monotonicity and (i), $m(G) \ge m((a,b)) = b-a > 0$.
(iv) Compact K is bounded, so $K \subseteq [-R,R]$ for some $R > 0$; hence $m(K) \le m([-R,R]) = 2R < +\infty$.

Structural Commentary: On $\mathbb{R}$, length, null endpoints, and compact boundedness give quick measure estimates.
Key Insight: Nonempty open sets contain intervals; compact sets fit inside one. □

Example 5.1.71. *A Positive-Measure Cantor-Type Set with No Interval*
Let $S_0 = [0,1]$ and, for $n \ge 1$, let R_n be the union of the middle open intervals of length 2^{-2n} removed from the 2^{n-1} components of S_{n-1}; set $S_n := S_{n-1} \setminus R_n$ and $S := \bigcap_{n=0}^{\infty} S_n$.
Then $m(R_n) = 2^{-n-1}$, $\sum_{n\ge1} m(R_n) = 1/2$, so $S = [0,1] \setminus \bigcup_n R_n$ is compact with $m(S) = 1/2$.
Since every component of S_n has length $\le 2^{-n}$, no interval (a,b) can lie in S; hence $\operatorname{int} S = \emptyset$.
Structural Commentary: Summable removals leave measure; repeated splitting kills interior.
Key Insight: Total removed length controls measure; repeated splitting controls interior.

Remark 5.1.72. *Borel Sets, Null Sets, and Regular Approximation*
Open, closed, and Borel subsets of $\mathbb{R}^p$ are measurable. Each $A \in \mathcal{M}(\mu)$ admits closed–open approximations $F \subseteq A \subseteq G$ with $\mu(G \setminus A)$ and $\mu(A \setminus F)$ arbitrarily small. In Lebesgue measure, one may take $F_\sigma \subseteq A \subseteq G_\delta$ with $\mu(G_\delta \setminus F_\sigma) = 0$. Null sets form a σ-ring, and countable subsets of $\mathbb{R}^p$ are null.
Structural Commentary: Measurable sets are Borel sets modified by null sets.
Key Insight: Regularity explains why open and closed approximations are enough for integration.

Corollary 5.1.73. *Lebesgue Regularity and Borel Completion on $\mathbb{R}$*
For $A \in \mathcal{L}^1$,

(i) $m(A) = \inf\{m(G) : A \subseteq G,\ G$ *open in* $\mathbb{R}\}$.
(ii) $m(A) = \sup\{m(K) : K \subseteq A,\ K$ *compact in* $\mathbb{R}\}$, *finite-measure case applied to* $A \cap (-n,n]$; $m(A) = +\infty$.
(iii) There exists $B \in \mathcal{B}(\mathbb{R})$ *with* $A \subseteq B$ *and* $m(B \setminus A) = 0$; *equivalently,* $\mathcal{L}^1$ *is the completion of* $\mathcal{B}(\mathbb{R})$.

Proof.

(i) For $\varepsilon > 0$, regular approximation gives open $G \supseteq A$ with $m(G \setminus A) < \varepsilon$, hence $m(G) \le m(A) + \varepsilon$. Monotonicity gives the reverse inequality.
(ii) If $m(A) < +\infty$, choose closed $F \subseteq A$ with $m(A \setminus F) < \varepsilon$. Since $F \cap [-n,n] \uparrow F$, continuity from below gives compact $K = F \cap [-n,n]$ with $m(F) - m(K) < \varepsilon$, so $m(A) - m(K) < 2\varepsilon$. If $m(A) = +\infty$, apply the finite case to $A_n = A \cap (-n,n]$ and let $n \to \infty$.
(iii) Choose open $G_k \supseteq A$ with $m(G_k \setminus A) < 1/k$, and set $B = \bigcap_k G_k$. Then $B \in \mathcal{B}(\mathbb{R})$, $A \subseteq B$, and $m(B \setminus A) = 0$. Thus every Lebesgue set is Borel plus a null subset.
Conversely, Borel sets and null subsets are Lebesgue measurable.

Structural Commentary: Lebesgue sets are Borel sets plus null-set ambiguity.
Key Insight: Open outer approximation and compact inner approximation are the usable form of regularity. □

Proposition 5.1.74. *Local Measurability on Unit Intervals*
Let $I_n := (n, n+1]$, $n \in \mathbb{Z}$. For $E \subseteq \mathbb{R}$, $E \in \mathcal{L}^1$ iff $E \cap I_n \in \mathcal{L}^1$ for all $n \in \mathbb{Z}$.

Proof.

($\Rightarrow$) Since $I_n \in \mathcal{B}(\mathbb{R}) \subseteq \mathcal{L}^1$, $E \in \mathcal{L}^1$ implies $E \cap I_n \in \mathcal{L}^1$ for all n.
($\Leftarrow$) If $E \cap I_n \in \mathcal{L}^1$ for all n, then $E = \bigcup_{n\in\mathbb{Z}}(E \cap I_n) \in \mathcal{L}^1$.

Structural Commentary: A countable Borel exhaustion localizes measurability.
Key Insight: Local measurable pieces patch by countable unions. □

Proposition 5.1.75. *Finite Measure Meeting Every Unit Interval*
With $I_n := (n, n+1]$, there exists $E \in \mathcal{B}(\mathbb{R})$ such that $0 < m(E \cap I_n)$ for all $n \in \mathbb{Z}$ and $m(E) < +\infty$.

Proof. Set $\varepsilon_n := 2^{-|n|-3}$ and $E := \bigcup_{n\in\mathbb{Z}}(n, n+\varepsilon_n]$. Then $(n, n+\varepsilon_n] \subseteq I_n$ are pairwise disjoint Borel intervals, so $E \in \mathcal{B}(\mathbb{R})$ and $m(E \cap I_n) = \varepsilon_n > 0$. Countable additivity gives $m(E) = \sum_{n\in\mathbb{Z}} \varepsilon_n = \frac{3}{8} < +\infty$.
Structural Commentary: Finite measure can be spread over countably many cells.
Key Insight: Positive local masses can be summable globally. □

Definition 5.1.76. *Symmetric Difference*
For $A, B \subseteq X$, define $A \triangle B := (A \setminus B) \cup (B \setminus A) = (A \cup B) \setminus (A \cap B)$.
Equivalently, $x \in A \triangle B$ iff x belongs to exactly one of A, B.

Proposition 5.1.77. *Null Sets Form a Sigma-Ideal*
Let $(X, \mathcal{M}, \mu)$ be a measure space and $\mathcal{Z} := \{N \in \mathcal{M} : \mu(N) = 0\}$.
Then $\mathcal{Z}$ is a σ-ring, and $N \in \mathcal{Z}$, $E \in \mathcal{M}$ imply $N \cap E \in \mathcal{Z}$.

Proof. If $(N_k) \subseteq \mathcal{Z}$, then $\mu(\bigcup_k N_k) \leq \sum_k \mu(N_k) = 0$, so $\bigcup_k N_k \in \mathcal{Z}$.
If $A, B \in \mathcal{Z}$, then $A \setminus B \subseteq A$, hence $A \setminus B \in \mathcal{Z}$. Finally, $N \cap E \subseteq N$ gives $N \cap E \in \mathcal{Z}$.
Structural Commentary: Null sets are stable under countable unions and measurable cuts.
Key Insight: Monotonicity makes every measurable piece of a null set null. □

Proposition 5.1.78. *Completed Sigma-Algebra*
Let $(X, \mathcal{M}, \mu)$ be measure space, $\mathcal{Z} := \{N \in \mathcal{M} : \mu(N) = 0\}$; $\overline{\mathcal{M}}_\mu := \{A \subseteq X : \exists E \in \mathcal{M},\ \exists N \in \mathcal{Z},\ A \triangle E \subseteq N\}$.
(i) $\mathcal{M} \subseteq \overline{\mathcal{M}}_\mu$, and $\overline{\mathcal{M}}_\mu$ is a σ-algebra;
(ii) $A \subseteq N \in \mathcal{Z} \Longrightarrow A \in \overline{\mathcal{M}}_\mu$;
(iii) $\overline{\mathcal{M}}_\mu = \{E \cup Z : E \in \mathcal{M},\ \exists N \in \mathcal{Z},\ Z \subseteq N\}$.

Proof.
(i) For $E \in \mathcal{M}$, $E \triangle E = \emptyset$, so $\mathcal{M} \subseteq \overline{\mathcal{M}}_\mu$. If $A \triangle E \subseteq N$, then $A^c \triangle E^c = A \triangle E \subseteq N$.
If $A_k \triangle E_k \subseteq N_k$, then $(\bigcup_k A_k) \triangle (\bigcup_k E_k) \subseteq \bigcup_k N_k \in \mathcal{Z}$. Hence $\overline{\mathcal{M}}_\mu$ is a σ-algebra.
(ii) If $A \subseteq N \in \mathcal{Z}$, then $A \triangle \emptyset \subseteq N$, so $A \in \overline{\mathcal{M}}_\mu$.
(iii) If $A \triangle E \subseteq N$, then $A = (E \setminus N) \cup (A \cap N)$, with $E \setminus N \in \mathcal{M}$ and $A \cap N \subseteq N$.
Conversely, if $A = E \cup Z$ and $Z \subseteq N \in \mathcal{Z}$, then $A \triangle E \subseteq Z \subseteq N$.
Structural Commentary: Completion adds all null modifications of measurable sets.
Key Insight: Equality modulo null sets is the membership criterion. □

Proposition 5.1.79. *Completed Measure*
Let $(X, \mathcal{M}, \mu)$ be a measure space, $\mathcal{Z} := \{N \in \mathcal{M} : \mu(N) = 0\}$,
and $\overline{\mathcal{M}}_\mu := \{A \subseteq X : \exists E \in \mathcal{M},\ \exists N \in \mathcal{Z},\ A \triangle E \subseteq N\}$. For $A \in \overline{\mathcal{M}}_\mu$, define $\overline{\mu}(A) := \mu(E)$ whenever $E \in \mathcal{M}$ and $A \triangle E \subseteq N \in \mathcal{Z}$. Then $\overline{\mu}$ is a well-defined complete measure on $\overline{\mathcal{M}}_\mu$, and $\overline{\mu}|_{\mathcal{M}} = \mu$.

Proof. If $A \triangle E \subseteq N_1$ and $A \triangle F \subseteq N_2$, with $E, F \in \mathcal{M}$ and $N_1, N_2 \in \mathcal{Z}$, then $E \triangle F \subseteq N_1 \cup N_2$, so $\mu(E) = \mu(F)$.
Hence $\overline{\mu}$ is well defined; since $E \triangle E = \emptyset$, $\overline{\mu}|_{\mathcal{M}} = \mu$.
For disjoint $(A_k) \subseteq \overline{\mathcal{M}}_\mu$, choose $E_k \in \mathcal{M}$ and $N_k \in \mathcal{Z}$ with $A_k \triangle E_k \subseteq N_k$, and set $N := \bigcup_k N_k$.
Then $E_i \cap E_j \subseteq N$ for $i \neq j$. Put $F_1 := E_1$ and $F_k := E_k \setminus \bigcup_{j<k} E_j$.
Then (F_k) disjoint, $A_k \triangle F_k \subseteq N$, $(\bigcup_k A_k) \triangle (\bigcup_k F_k) \subseteq N$, so $\overline{\mu}(\bigcup_k A_k) = \mu(\bigcup_k F_k) = \sum_k \mu(F_k) = \sum_k \overline{\mu}(A_k)$.
If $A \subseteq B \in \overline{\mathcal{M}}_\mu$ and $\overline{\mu}(B) = 0$, choose $E \in \mathcal{M}$ and $N \in \mathcal{Z}$ with $B \triangle E \subseteq N$.
Then $\mu(E) = 0$, so $A \subseteq B \subseteq E \cup N \in \mathcal{Z}$; hence $A \in \overline{\mathcal{M}}_\mu$ and $\overline{\mu}(A) = 0$.
Structural Commentary: The completed measure ignores null errors.
Key Insight: Null perturbations do not change countable additivity. □

Theorem 5.1.80. *Lebesgue Measurable Non-Borel Sets Exist*
There exists $E \subseteq \mathbb{R}$ such that $E \in \mathcal{L}^1$ and $E \notin \mathcal{B}(\mathbb{R})$; equivalently, $\mathcal{B}(\mathbb{R}) \subsetneq \mathcal{L}^1$.

Proof. By Lemma 5.1.59, $\mathcal{B}(\mathbb{R})$ is countably generated, hence $|\mathcal{B}(\mathbb{R})| \leq \mathfrak{c}$. Let C be the Cantor set.
Since $|C| = \mathfrak{c}$ and $m(C) = 0$, Proposition 5.1.67 gives $\mathcal{P}(C) \subseteq \mathcal{L}^1$, so $|\mathcal{L}^1| \geq 2^{\mathfrak{c}}$.
Also $\mathcal{L}^1 \subseteq \mathcal{P}(\mathbb{R})$, so $|\mathcal{L}^1| \leq 2^{\mathfrak{c}}$; hence $|\mathcal{L}^1| = 2^{\mathfrak{c}}$. By Cantor's theorem, $\mathfrak{c} < 2^{\mathfrak{c}}$, so $\mathcal{B}(\mathbb{R}) \subsetneq \mathcal{L}^1$.
Structural Commentary: Countable generation is smaller than completion by null subsets.
Key Insight: Completeness creates more measurable sets than topology can generate. □

Definition 5.1.81. *Variation Candidates of a Charge*
Let λ be a charge on $\mathcal{M}$. For $E \in \mathcal{M}$, define
(i) $\mathfrak{P}_{\mathcal{M}}(E) := \{(A_j)_{j=1}^n : n \in \mathbb{N},\ E = \bigsqcup_{j=1}^n A_j,\ A_j \in \mathcal{M}\}$, the finite measurable partitions of E;
(ii) $\pi_\lambda(E) := \sup\{\lambda(A) : A \subseteq E,\ A \in \mathcal{M}\}$, the *positive variation* candidate;
(iii) $\nu_\lambda(E) := \sup\{\sum_{j=1}^n |\lambda(A_j)| : (A_j)_{j=1}^n \in \mathfrak{P}_{\mathcal{M}}(E)\}$, the *total variation* candidate.
For $E = \emptyset$, the empty partition is allowed and contributes 0.

Proposition 5.1.82. *Charge Operations and Variation Measures*
Let $\lambda, \lambda_1, \ldots, \lambda_m$ be charges on $\mathcal{M}$.
(i) If $a_1, \ldots, a_m \in \mathbb{R}$, then $\Lambda := \sum_{j=1}^m a_j \lambda_j$ is a charge.
(ii) If $f \in \mathscr{L}^1(\mu)$ on $(X, \mathcal{M}, \mu)$, then $\lambda_f(A) := \int_A f \, d\mu$ is a charge.
(iii) The set functions π_λ and ν_λ are measures on $\mathcal{M}$.

Proof.

(i) If $E = \bigsqcup_n E_n$, then $\Lambda(E) = \sum_{j=1}^m a_j \sum_n \lambda_j(E_n) = \sum_n \Lambda(E_n)$.
(ii) By Theorem 5.4.11, $A \mapsto \int_A f\,d\mu$ is countably additive, and $|\lambda_f(A)| \le \int_X |f|\,d\mu < +\infty$.
(iii) Let $E = \bigsqcup_n E_n$. For $A \subseteq E$, $A = \bigsqcup_n (A \cap E_n)$, so $\lambda(A) \le \sum_n \pi_\lambda(E_n)$; by suprema, $\pi_\lambda(E) \le \sum_n \pi_\lambda(E_n)$. Conversely, for $N \in \mathbb{N}$ choose $A_n \subseteq E_n$ with $\lambda(A_n) > \pi_\lambda(E_n) - \varepsilon 2^{-n}$ for $n \le N$.
Then $\pi_\lambda(E) \ge \sum_{n=1}^N \pi_\lambda(E_n) - \varepsilon$. Let $N \to \infty$ and $\varepsilon \downarrow 0$. For ν_λ, every $(A_j)_{j=1}^r \in \mathfrak{P}_{\mathcal{M}}(E)$ satisfies $\sum_{j=1}^r |\lambda(A_j)| \le \sum_n \nu_\lambda(E_n)$ after refining by $A_j \cap E_n$, hence $\nu_\lambda(E) \le \sum_n \nu_\lambda(E_n)$. Conversely, combine near-optimal finite partitions of $E_1, \dots, E_N$, add $\bigcup_{n>N} E_n$ as one part; let $N \to \infty$.

Structural Commentary: Variation turns signed additivity into measure.
Key Insight: Refine partitions against a disjoint decomposition. □

Definition 5.1.83. *Limit Inferior and Superior of Sets*
For $(E_n) \subseteq \mathcal{P}(X)$, define $\limsup_n E_n := \bigcap_{m=1}^\infty \bigcup_{n \ge m} E_n$ and $\liminf_n E_n := \bigcup_{m=1}^\infty \bigcap_{n \ge m} E_n$.

(i) *Infinitely often*: $x \in \limsup_n E_n$ iff $\{n : x \in E_n\}$ is infinite;
(ii) *Eventually always*: $x \in \liminf_n E_n$ iff $\{n : x \notin E_n\}$ is finite;
(iii) *Set limit*: if $\limsup_n E_n = \liminf_n E_n$, write $\lim_n E_n$ for the common set.

Proposition 5.1.84. *Basic Laws for Set Limits*
Let $(E_n) \subseteq \mathcal{P}(X)$.

(i) $\emptyset \subseteq \liminf_n E_n \subseteq \limsup_n E_n \subseteq X$.
(ii) $E_n \uparrow E \implies \liminf_n E_n = \limsup_n E_n = E = \bigcup_n E_n$.
(iii) $E_n \downarrow E \implies \liminf_n E_n = \limsup_n E_n = E = \bigcap_n E_n$.
(iv) $(E_n) \subseteq \mathcal{M} \implies \liminf_n E_n, \limsup_n E_n \in \mathcal{M}$.

Proof.

(i) Since $\bigcap_m \bigcup_{n \ge m} E_n \subseteq \bigcup_{n \ge 1} E_n \subseteq X$ and $\bigcap_{n \ge m} E_n \subseteq \bigcup_{n \ge m} E_n$ for all m, the inclusions follow.
(ii) If $E_n \uparrow E$, then $\{n : x \in E_n\}$ is nonempty iff it is cofinite, hence $\liminf_n E_n = \limsup_n E_n = \bigcup_n E_n = E$.
(iii) If $E_n \downarrow E$, then $\{n : x \in E_n\}$ is infinite iff it is all of $\mathbb{N}$, hence $\liminf_n E_n = \limsup_n E_n = \bigcap_n E_n = E$.
(iv) If $(E_n) \subseteq \mathcal{M}$, then $\bigcup_{n \ge m} E_n, \bigcap_{n \ge m} E_n \in \mathcal{M}\ \forall m$; countable closure gives $\limsup_n E_n, \liminf_n E_n \in \mathcal{M}$.

Structural Commentary: Set limsup means infinitely often; set liminf means eventually always.
Key Insight: Translate the definitions into pointwise membership statements. □

Proposition 5.1.85. *Measure Bounds for Set Limits*
Let (E_n) be measurable in a measure space $(X, \mathcal{M}, \mu)$.

(i) $\mu(\liminf_n E_n) \le \liminf_n \mu(E_n)$.
(ii) $\mu(\bigcup_{n=1}^\infty E_n) < +\infty \implies \limsup_n \mu(E_n) \le \mu(\limsup_n E_n)$.

Proof.

(i) Set $A_m := \bigcap_{n \ge m} E_n$. Then $A_m \uparrow \liminf_n E_n$ and $A_m \subseteq E_n$ for all $n \ge m$, so $\mu(A_m) \le \inf_{n \ge m} \mu(E_n)$. By continuity from below, $\mu(\liminf_n E_n) = \lim_m \mu(A_m) \le \sup_m \inf_{n \ge m} \mu(E_n) = \liminf_n \mu(E_n)$.
(ii) Set $B_m := \bigcup_{n \ge m} E_n$. $B_m \downarrow \limsup_n E_n$, $\mu(B_1) < +\infty$, $E_n \subseteq B_m\ \forall n \ge m$, so $\sup_{n \ge m} \mu(E_n) \le \mu(B_m)$. By continuity from above, $\limsup_n \mu(E_n) = \inf_m \sup_{n \ge m} \mu(E_n) \le \lim_m \mu(B_m) = \mu(\limsup_n E_n)$.

Structural Commentary: Tail intersections increase to liminf; tail unions decrease to limsup.
Key Insight: Compare each tail envelope to the corresponding tail bounds of the measures. □

5.2 Regular Approximation, Additivity, and Nonmeasurable Sets

The Caratheodory definition is intrinsic to outer measure, but in $\mathbb{R}^p$ it has a topological shadow: measurable sets are exactly those that can be approximated by open, closed, compact, and cell-built sets up to arbitrarily small error. That bridge is also the route to the first nonmeasurable and non-Borel examples.

Lemma 5.2.1. *Outer Approximation by Open and G_δ Sets*
For $A \subseteq \mathbb{R}^p$:

(i) for every $\varepsilon > 0$, there exists open $G \subseteq \mathbb{R}^p$ with $A \subseteq G$ and $m(G) \le m_p^*(A) + \varepsilon$;
(ii) $m_p^*(A) = \inf\{m(G) : A \subseteq G,\ G \subseteq \mathbb{R}^p \text{ open}\}$;
(iii) there exists G_δ set $H \supseteq A$ with $m(H) = m_p^*(A)$.

Proof.

(i) If $m_p^*(A) = +\infty$, take $G = \mathbb{R}^p$.
If $m_p^*(A) < +\infty$, choose an open-cell cover $A \subseteq \bigcup_k I_k$ with $\sum_k \mathrm{vol}_p(I_k) \le m_p^*(A) + \varepsilon$, and set $G := \bigcup_k I_k$.
Then G is open, $A \subseteq G$, and $m(G) \le \sum_k m(I_k) = \sum_k \mathrm{vol}_p(I_k) \le m_p^*(A) + \varepsilon$.

(ii) By (i), $\inf\{m(G) : A \subseteq G,\ G \subseteq \mathbb{R}^p \text{ open}\} \le m_p^*(A) + \varepsilon$ for all $\varepsilon > 0$; let $\varepsilon \downarrow 0$.
Conversely, $A \subseteq G$ gives $m_p^*(A) \le m_p^*(G) = m(G)$ for open G; take the infimum over G.

(iii) If $m_p^*(A) = +\infty$, take $H = \mathbb{R}^p$.
If $m_p^*(A) < +\infty$, choose open $G_n \supseteq A$ with $m(G_n) \le m_p^*(A) + 1/n$ and set $H := \bigcap_n G_n$.
Then H is G_δ, $A \subseteq H$, and $m_p^*(A) \le m(H) \le m(G_n) \le m_p^*(A) + 1/n$ for all n, so $m(H) = m_p^*(A)$.

Structural Commentary: Outer measure can be realised by a Borel envelope from outside.
Key Insight: Infima over cell covers can be repackaged as open-set approximation. □

Corollary 5.2.2. *Null Sets Sit Inside Borel Null Sets*
If $Z \subseteq \mathbb{R}^p$ and $m_p^*(Z) = 0$, then there exists $B \in \mathcal{B}(\mathbb{R}^p)$ such that $Z \subseteq B$ and $m(B) = 0$.

Proof. By Lemma 5.2.1, choose G_δ set $B \supseteq Z$ with $m(B) = m_p^*(Z) = 0$. Since $G_\delta \subseteq \mathcal{B}(\mathbb{R}^p)$, B is Borel.
Structural Commentary: Null sets admit Borel envelopes of the same measure.
Key Insight: Outer approximation turns arbitrary null sets into Borel null supersets. □

Theorem 5.2.3. *Open and G_δ Characterizations of Measurability*
For $E \subseteq \mathbb{R}^p$, the following are equivalent:

(i) $E \in \mathcal{L}^p$;
(ii) $\forall \varepsilon > 0$, $\exists$ *open* $G \subseteq \mathbb{R}^p$ *such that* $E \subseteq G$ *and* $m_p^*(G \setminus E) < \varepsilon$;
(iii) $\exists$ G_δ *set* $H \subseteq \mathbb{R}^p$ *such that* $E \subseteq H$ *and* $m_p^*(H \setminus E) = 0$;
(iv) $\exists$ G_δ *set* $H \subseteq \mathbb{R}^p$ *and* $Z \in \mathcal{L}^p$ *such that* $Z \subseteq H$, $m(Z) = 0$, *and* $E = H \setminus Z$.

Proof.

(i) $\Rightarrow$ (ii). If $m(E) < +\infty$, choose open $G \supseteq E$ with $m(G) \le m(E) + \varepsilon/2$.
Since $E \in \mathcal{L}^p$, $m_p^*(G \setminus E) = m(G) - m(E) < \varepsilon$. If $m(E) = +\infty$, set $E_1 := E \cap \{x : \|x\|_2 \le 1\}$ and $E_n := E \cap \{x : n-1 < \|x\|_2 \le n\}$ for $n \ge 2$. For each n, choose open $G_n \supseteq E_n$ with $m_p^*(G_n \setminus E_n) < \varepsilon 2^{-n}$; then $G := \bigcup_n G_n$ is open, $E \subseteq G$, and $G \setminus E \subseteq \bigcup_n (G_n \setminus E_n)$, so $m_p^*(G \setminus E) < \varepsilon$.

(ii) $\Rightarrow$ (iii). Choose open $G_n \supseteq E$ with $m_p^*(G_n \setminus E) < 1/n$, and set $H := \bigcap_n G_n$.
Then H is G_δ, $E \subseteq H$, and $H \setminus E \subseteq G_n \setminus E$ for all n; hence $m_p^*(H \setminus E) = 0$.

(iii) $\Rightarrow$ (i). Since H is G_δ, $H \in \mathcal{B}(\mathbb{R}^p) \subseteq \mathcal{L}^p$. Also $m_p^*(H \setminus E) = 0$ implies $H \setminus E \in \mathcal{L}^p$ by Proposition 5.1.67.
Thus $E = H \setminus (H \setminus E) \in \mathcal{L}^p$.

(iii) $\Leftrightarrow$ (iv). If (iii) holds, set $Z := H \setminus E$; then $Z \subseteq H$, $m(Z) = 0$, and $E = H \setminus Z$.
Conversely, if (iv) holds, then $E \subseteq H$ and $H \setminus E = Z$, so $m_p^*(H \setminus E) = 0$.

Structural Commentary: Lebesgue measurable sets are G_δ sets with null material removed.
Key Insight: Small open excess becomes a null G_δ excess after countably many approximations. □

Corollary 5.2.4. *Open Outer Regularity for Measurable Sets*
If $E \in \mathcal{L}^p$, then

(i) $\forall \varepsilon > 0$, $\exists$ *open* $G \supseteq E$ *such that* $m(G) \le m(E) + \varepsilon$;
(ii) $m(E) = \inf\{m(G) : E \subseteq G,\ G \subseteq \mathbb{R}^p \text{ open}\}$.

Proof.

(i) By Theorem 5.2.3, choose open $G \supseteq E$ with $m_p^*(G \setminus E) < \varepsilon$.
If $m(E) < +\infty$, then $m(G) = m(E) + m(G \setminus E) \le m(E) + \varepsilon$; if $m(E) = +\infty$, the inequality is trivial.

(ii) From (i), $\inf_G m(G) \le m(E)$. Conversely, $E \subseteq G$ gives $m(E) \le m(G)$ by monotonicity; take $\inf_G$.

Structural Commentary: Measurable sets are measured from outside by open supersets.
Key Insight: Open excess controls the outer regular infimum. □

Theorem 5.2.5. *Closed and F_σ Characterizations of Measurability*
For $E \subseteq \mathbb{R}^p$, the following are equivalent:

(i) $E \in \mathcal{L}^p$;
(ii) $\forall \varepsilon > 0$, $\exists$ *closed* $F \subseteq E$ *such that* $m_p^*(E \setminus F) < \varepsilon$;
(iii) $\exists$ F_σ *set* $K \subseteq E$ *such that* $m_p^*(E \setminus K) = 0$;
(iv) $\exists$ F_σ *set* $K \subseteq E$ *and* $Z \in \mathcal{L}^p$ *such that* $Z \subseteq E$, $m(Z) = 0$, *and* $E = K \cup Z$.

Proof.

(*i*) $\Rightarrow$ (*ii*). Since $E^c \in \mathcal{L}^p$, Theorem 5.2.3 gives open $G \supseteq E^c$ with $m_p^*(G \setminus E^c) < \varepsilon$. Set $F := G^c$.
Then F is closed, $F \subseteq E$, and $E \setminus F = G \setminus E^c$, so $m_p^*(E \setminus F) < \varepsilon$.

(*ii*) $\Rightarrow$ (*iii*). Choose closed $F_n \subseteq E$ with $m_p^*(E \setminus F_n) < 1/n$, and set $K := \bigcup_n F_n$.
Then K is F_σ, $K \subseteq E$, and $E \setminus K \subseteq E \setminus F_n$ for all n; hence $m_p^*(E \setminus K) = 0$.

(*iii*) $\Rightarrow$ (*i*). Since K is F_σ, $K \in \mathcal{B}(\mathbb{R}^p) \subseteq \mathcal{L}^p$. Also $m_p^*(E \setminus K) = 0$ implies $E \setminus K \in \mathcal{L}^p$ by Proposition 5.1.67.
Thus $E = K \cup (E \setminus K) \in \mathcal{L}^p$.

(*iii*) $\Leftrightarrow$ (*iv*). If (iii) holds, set $Z := E \setminus K$; then $Z \subseteq E$, $m(Z) = 0$, and $E = K \cup Z$.
Conversely, if (iv) holds, then $E \setminus K \subseteq Z$, so $m_p^*(E \setminus K) = 0$.

Structural Commentary: Inside approximation is the complement of outside approximation.
Key Insight: Measurable sets are also F_σ sets with null material added. □

Corollary 5.2.6. *Closed Inner Regularity*
If $E \in \mathcal{L}^p$, then

(*i*) *$\forall \varepsilon > 0$, $\exists$ closed $F \subseteq E$ such that $m(E) \leq m(F) + \varepsilon$;*
(*ii*) *$m(E) = \sup\{m(F) : F \subseteq E,\ F \subseteq \mathbb{R}^p \text{ closed}\}$.*

Proof.

(i) By Theorem 5.2.5, choose closed $F \subseteq E$ with $m_p^*(E \setminus F) < \varepsilon$.
Since $F \in \mathcal{L}^p$ and $E = F \sqcup (E \setminus F)$, $m(E) = m(F) + m(E \setminus F) \leq m(F) + \varepsilon$.

(ii) From (i), $m(E) \leq \sup_F m(F)$. Conversely, $F \subseteq E$ gives $m(F) \leq m(E)$; take $\sup_F$.

Structural Commentary: Measurable sets are measured from inside by closed subsets.
Key Insight: Small discarded measure controls the inner regular supremum. □

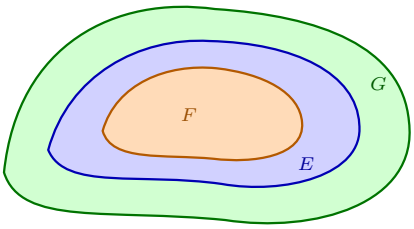


Figure 193: Regular approximation sandwiches E between F and G with small excess.

Theorem 5.2.7. *Compact Approximation at Finite Outer Measure*
Let $E \subseteq \mathbb{R}^p$ and $m_p^*(E) < +\infty$. Then $E \in \mathcal{L}^p \iff \forall \varepsilon > 0,\ \exists C \subseteq$ compact with $m_p^*(E \setminus C) < \varepsilon$.

Proof.

($\Rightarrow$) Set $E_N := E \cap \{x : \|x\|_2 \leq N\}$. Then $E_N \uparrow E$ and $m(E) < +\infty$, so choose N with $m(E \setminus E_N) < \varepsilon/2$. By Theorem 5.2.5, choose closed $C \subseteq E_N$ with $m(E_N \setminus C) < \varepsilon/2$.
Then C is closed and bounded, hence compact, and $m_p^*(E \setminus C) \leq m(E \setminus E_N) + m(E_N \setminus C) < \varepsilon$.

($\Leftarrow$) Choose compact $C_n \subseteq E$ with $m_p^*(E \setminus C_n) < 1/n$, and set $K := \bigcup_n C_n$.
Then K is F_σ, $K \subseteq E$, and $E \setminus K \subseteq E \setminus C_n$ $\forall n$, so $m_p^*(E \setminus K) = 0$. By Theorem 5.2.5, $E \in \mathcal{L}^p$.

Structural Commentary: Finite measure permits bounded truncation before inner closure.
Key Insight: Compact approximation is closed approximation plus a finite-measure tail estimate. □

Theorem 5.2.8. *Finite Cell Approximation in Symmetric Difference*
If $E \in \mathcal{L}^p$, $m(E) < +\infty$, then $\forall \varepsilon > 0$ $\exists$ bounded open cells $I_1, \ldots, I_N$ s.t. $K_N := \bigcup_{j=1}^N I_j$; $m(E \triangle K_N) < \varepsilon$.

Proof. Choose bounded open cells (J_k) with $E \subseteq G := \bigcup_k J_k$ and $\sum_k \mathrm{vol}_p(J_k) < m(E) + \varepsilon/2$.
Since $G \in \mathcal{L}^p$ and $E \subseteq G$, $m(G \setminus E) = m(G) - m(E) \leq \sum_k \mathrm{vol}_p(J_k) - m(E) < \varepsilon/2$..
By Theorem 5.2.7, choose compact $C \subseteq E$ with $m(E \setminus C) < \varepsilon/2$.
Since $C \subseteq \bigcup_k J_k$, compactness gives $I_1, \ldots, I_N \in \{J_k\}$ with $C \subseteq K_N := \bigcup_{j=1}^N I_j \subseteq G$.
Hence $E \triangle K_N = (E \setminus K_N) \cup (K_N \setminus E) \subseteq (E \setminus C) \cup (G \setminus E)$. Thus $m(E \triangle K_N) \leq m(E \setminus C) + m(G \setminus E) < \varepsilon$.
Structural Commentary: Compactness turns a countable open cover into a finite geometric approximant.
Key Insight: Approximate from outside by cells and from inside by a compact core. □

Corollary 5.2.9. *Cell Approximation Variant Choices*
If $E \in \mathcal{L}^p$, $m(E) < +\infty$, and $\varepsilon > 0$, then there exist bounded cells $I_1, \dots, I_N$ such that $K := \bigcup_{j=1}^{N} I_j$ satisfies $m(E \triangle K) < \varepsilon$, with the cells chosen as any one of:

(i) $I_i \cap I_j = \emptyset$ for $i \neq j$, and each I_j is open;
(ii) each I_j is closed;
(iii) each I_j is half-open of a fixed orientation.

Proof.

(i) Choose compact $C \subseteq E$ and open $G \supseteq E$ with $m(E \setminus C) + m(G \setminus E) < \varepsilon$.
Cover C by bounded open cells $J_1, \dots, J_M$ with $\overline{J_j} \subseteq G$, and set $U := \bigcup_{j=1}^{M} J_j$.
Then $C \subseteq U \subseteq G$, so $m(E \triangle U) \leq m(E \setminus C) + m(G \setminus E) < \varepsilon$.
(ii) Let $\mathcal{G}$ be finite grid from coordinate endpoints of $J_1, \dots, J_M$. For $\mathcal{G}' \subseteq \mathcal{G}$ selecting U, $\bigsqcup_{Q \in \mathcal{G}'} Q$ gives open or fixed-orientation half-open cells, differs from U only on finitely many faces, hence by a null set.
(iii) Let $K_c := \bigcup_{j=1}^{M} \overline{J_j}$. Then K_c is a finite union of closed cells and $C \subseteq K_c \subseteq G$.
Hence $m(E \triangle K_c) \leq m(E \setminus C) + m(G \setminus E) < \varepsilon$.

Structural Commentary: Cell format changes only boundary-null pieces.
Key Insight: A compact core inside an open envelope permits finite cell variants. □

Theorem 5.2.10. *Outer Measure Additivity with One Measurable Set*
Let $E \in \mathcal{L}^p$ and $F \subseteq \mathbb{R}^p$. Then $m_p^*(E \cup F) + m_p^*(E \cap F) = m(E) + m_p^*(F)$.

Proof. Since $E \in \mathcal{L}^p$, Caratheodory splitting: $m_p^*(E \cup F) = m(E) + m_p^*(F \setminus E)$; $m_p^*(F) = m_p^*(E \cap F) + m_p^*(F \setminus E)$.
Hence $m_p^*(E \cup F) + m_p^*(E \cap F) = m(E) + m_p^*(F)$
Structural Commentary: One measurable cut splits arbitrary sets exactly.
Key Insight: Split both $E \cup F$ and F along the same measurable cut E. □

Corollary 5.2.11. *Disjoint Additivity and Difference*
Let $E \in \mathcal{L}^p$ and $F \subseteq \mathbb{R}^p$.

(i) $E \cap F = \emptyset \implies m_p^*(E \cup F) = m(E) + m_p^*(F)$.
(ii) $E \subseteq F$, $m(E) < +\infty \implies m_p^*(F \setminus E) = m_p^*(F) - m(E)$.

Proof.

(i) By Theorem 5.2.10, $m_p^*(E \cup F) + m_p^*(E \cap F) = m(E) + m_p^*(F)$.
Since $E \cap F = \emptyset$ and $m_p^*(\emptyset) = 0$, $m_p^*(E \cup F) = m(E) + m_p^*(F)$.
(ii) Since $E \subseteq F$, Caratheodory splitting gives $m_p^*(F) = m(E) + m_p^*(F \setminus E)$; subtract $m(E) < +\infty$.

Structural Commentary: Exact splitting yields disjoint sums and finite differences.
Key Insight: Disjointness removes the overlap; containment isolates the remainder. □

Theorem 5.2.12. *Finite Outer Measure Inner Test*
If $m_p^*(E) < +\infty$, then $E \in \mathcal{L}^p \iff \exists B \in \mathcal{L}^p$, $B \subseteq E$, $m(B) = m_p^*(E)$.

Proof.

($\Rightarrow$) Set $B := E$. Then $B \in \mathcal{L}^p$, $B \subseteq E$, and $m(B) = m(E) = m_p^*(E)$.
($\Leftarrow$) Assume $B \in \mathcal{L}^p$, $B \subseteq E$, and $m(B) = m_p^*(E)$.
Since $m(B) < +\infty$, the difference formula gives $m_p^*(E \setminus B) = m_p^*(E) - m(B) = 0$.
Thus $E \setminus B \in \mathcal{L}^p$, and $E = B \cup (E \setminus B) \in \mathcal{L}^p$.

Structural Commentary: Exact inner mass leaves only a null remainder.
Key Insight: Equality of outer mass converts containment into measurability. □

Theorem 5.2.13. *Single Finite Container Test*
Let $A \in \mathcal{L}^p$, $m(A) < +\infty$, and $E \subseteq A$. Then $E \in \mathcal{L}^p \iff m(A) = m_p^*(E) + m_p^*(A \setminus E)$.

Proof.

($\Rightarrow$) Since $E \in \mathcal{L}^p$ and $E \subseteq A$,
Caratheodory splitting of A by E gives $m(A) = m(E) + m(A \setminus E) = m_p^*(E) + m_p^*(A \setminus E)$.
($\Leftarrow$) Choose G_δ set $H \supseteq A \setminus E$ with $m(H) = m_p^*(A \setminus E)$. Since $A \setminus E \subseteq A \cap H \subseteq H$, $m(A \cap H) = m_p^*(A \setminus E)$.
Set $B := A \setminus H$. $B \in \mathcal{L}^p$, $B \subseteq E$; $m(B) = m(A) - m(A \cap H) = m_p^*(E)$. By Theorem 5.2.12, $E \in \mathcal{L}^p$.

Structural Commentary: Inside a finite measurable container, Caratheodory only needs to be tested once.
Key Insight: Outer-approximate the complement; the remainder has full outer mass. □

Corollary 5.2.14. *Bounded Inner Measure Test*
For bounded cell $Q \subseteq \mathbb{R}^p$ and $E \subseteq Q$, define $m_{*,Q}(E) := m(Q) - m_p^*(Q \setminus E)$. Then:

(i) $E \in \mathcal{L}^p \iff m_{*,Q}(E) = m_p^*(E)$;
(ii) for arbitrary $E \subseteq \mathbb{R}^p$, $E \in \mathcal{L}^p \iff E \cap [-n,n]^p \in \mathcal{L}^p$ *for all* $n \in \mathbb{N}$.

Proof.

(*i*) ($\Rightarrow$) If $E \in \mathcal{L}^p$, then $m(Q) = m(E) + m_p^*(Q \setminus E)$, hence $m_{*,Q}(E) = m_p^*(E)$.
(*i*) ($\Leftarrow$) If $m_{*,Q}(E) = m_p^*(E)$, then $m(Q) = m_p^*(E) + m_p^*(Q \setminus E)$, so $E \in \mathcal{L}^p$ by the finite-container test.
(*ii*) ($\Rightarrow$) If $E \in \mathcal{L}^p$, then $E \cap [-n,n]^p \in \mathcal{L}^p$ for all n, since $[-n,n]^p$ is Borel.
(*ii*) ($\Leftarrow$) If $E \cap [-n,n]^p \in \mathcal{L}^p$ for all n, then $E = \bigcup_{n=1}^{\infty}(E \cap [-n,n]^p) \in \mathcal{L}^p$.

Structural Commentary: Bounded inner measure tests one finite container.
Key Insight: Equality of inner and outer measure is a bounded version of Caratheodory. □

Definition 5.2.15. *Difference Sets and Vitali Selectors*

(i) *Difference set*: for $A \subseteq \mathbb{R}^p$, $A - A := \{x - y : x, y \in A\}$; $A \subseteq B \Longrightarrow A - A \subseteq B - B$.
(ii) *Rational equivalence*: $x \sim_{\mathbb{Q}^p} y \iff x - y \in \mathbb{Q}^p$; equivalently, $[x]_{\mathbb{Q}^p} = x + \mathbb{Q}^p$.
(iii) *Vitali selector*: choose $\mathcal{V} \subseteq \mathbb{R}^p$ such that $|\mathcal{V} \cap (x + \mathbb{Q}^p)| = 1$ for every $x \in \mathbb{R}^p$.
(iv) *Rational translate*: for $q \in \mathbb{Q}^p$, $\mathcal{V}_q := q + \mathcal{V}$.

Lemma 5.2.16. *Steinhaus Lemma for Compact Sets*
If $K \subseteq \mathbb{R}^p$ is compact and $m(K) > 0$, then $\exists \delta > 0$ such that $B(0,\delta) \subseteq K - K$.

Proof. Choose open $G \supseteq K$ with $m(G) < 2m(K)$. Since K is compact and G^c is closed, set $\delta := \operatorname{dist}(K, G^c) > 0$.
If $\|x\|_2 < \delta$, then $x + K \subseteq G$. If $(x + K) \cap K = \emptyset$, translation invariance and additivity give
$2m(K) = m(x + K) + m(K) = m((x + K) \cup K) \leq m(G) < 2m(K)$, a contradiction.
Hence $(x + K) \cap K \neq \emptyset$, so $x = k_1 - k_2 \in K - K$. Thus $B(0,\delta) \subseteq K - K$.
Structural Commentary: A positive compact set overlaps all sufficiently small translates of itself.
Key Insight: Disjoint small translates cannot fit inside a tight open envelope. □

Theorem 5.2.17. *Steinhaus Theorem*
If $E \in \mathcal{L}^p$ and $m(E) > 0$, then $\exists \delta > 0$ such that $B(0,\delta) \subseteq E - E$.

Proof. Set $E_n := E \cap B(0,n)$. Since $E_n \uparrow E$ and $m(E) > 0$, choose n with $0 < m(E_n) < +\infty$.
By compact inner approximation, choose compact $K \subseteq E_n$ with $m(K) > 0$.
Lemma 5.2.16 gives $\delta > 0$ with $B(0,\delta) \subseteq K - K \subseteq E - E$.
Structural Commentary: Positive measurable sets contain positive compact cores.
Key Insight: Difference-set thickness follows from compact inner regularity. □

Lemma 5.2.18. *Rational Translates of a Vitali Selector Partition $\mathbb{R}^p$*
Let $(q_i)_{i \geq 1}$ enumerate $\mathbb{Q}^p$.
Then $\mathbb{R}^p = \bigsqcup_{i=1}^{\infty} \mathcal{V}_{q_i}$; equivalently, each $x \in \mathbb{R}^p$ has a unique representation $x = q_i + v$ with $i \in \mathbb{N}$ and $v \in \mathcal{V}$.

Proof. For $x \in \mathbb{R}^p$, choose the unique $v \in \mathcal{V} \cap (x + \mathbb{Q}^p)$; then $x - v = q_i$ for some i, so $x \in \mathcal{V}_{q_i}$. If $q + v = q' + v'$ with $q, q' \in \mathbb{Q}^p$ and $v, v' \in \mathcal{V}$, then $v - v' = q' - q \in \mathbb{Q}^p$, hence $v = v'$ by selector uniqueness and then $q = q'$.
Structural Commentary: Rational translates enumerate the equivalence classes selected by $\mathcal{V}$.
Key Insight: One representative per class makes the rational translates disjoint and exhaustive. □

Theorem 5.2.19. *Vitali Selectors Are Nonmeasurable*
If $\mathcal{V} \subseteq \mathbb{R}^p$ is a Vitali selector, then $\mathcal{V} \notin \mathcal{L}^p$.

Proof. Assume $\mathcal{V} \in \mathcal{L}^p$. If $m(\mathcal{V}) > 0$, Steinhaus gives $\delta > 0$ with $B(0,\delta) \subseteq \mathcal{V} - \mathcal{V}$. Choose $0 \neq q \in \mathbb{Q}^p \cap B(0,\delta)$; then $q = v_1 - v_2$ for $v_1, v_2 \in \mathcal{V}$, so $v_1 \sim_{\mathbb{Q}^p} v_2$, forcing $v_1 = v_2$ and $q = 0$, a contradiction.
If $m(\mathcal{V}) = 0$, then each $\mathcal{V}_{q_i}$ is null by translation invariance, and $\mathbb{R}^p = \bigsqcup_i \mathcal{V}_{q_i}$ would imply every bounded cell has measure 0, contradicting positive cell volume.
Structural Commentary: Positive and zero measure are both impossible.
Key Insight: Steinhaus forbids positive measure; countable translated partition forbids zero measure. □

Theorem 5.2.20. *Nonmeasurable Subsets Are Everywhere*
If $E \subseteq \mathbb{R}^p$ and $m_p^*(E) > 0$, then $\exists N \subseteq E$ with $N \notin \mathcal{L}^p$.

Proof. Let (q_i) enumerate $\mathbb{Q}^p$ and set $E_i := E \cap \mathcal{V}_{q_i}$. If $E_i \in \mathcal{L}^p$ and $m(E_i) > 0$, then $E_i - E_i$ contains a ball about 0, so $0 \neq q \in \mathbb{Q}^p \cap (E_i - E_i)$. But $E_i \subseteq q_i + \mathcal{V}$, so $q = (q_i + v_1) - (q_i + v_2) = v_1 - v_2$ with $v_1, v_2 \in \mathcal{V}$, contradicting selector uniqueness. Hence each measurable E_i is null. If every E_i were measurable, then $E = \bigcup_i E_i$ would satisfy $m_p^*(E) = 0$, contradiction. Thus some $E_i \subseteq E$ is nonmeasurable.
Structural Commentary: Any positive outer-measure set meets the Vitali partition in a nonmeasurable piece.
Key Insight: Vitali pieces cannot be measurable and positive. □

Corollary 5.2.21. *Nonadditive Decompositions*

(i) If $E \in \mathcal{L}^p$ and $0 < m(E) < +\infty$, then $\exists B, C \notin \mathcal{L}^p$ such that $E = B \sqcup C$ and $m(E) < m_p^*(B) + m_p^*(C)$.
(ii) If $B \notin \mathcal{L}^p$, $m_p^*(B) < +\infty$, and H is G_δ with $B \subseteq H$, $m(H) = m_p^*(B)$, then $C := H \setminus B \notin \mathcal{L}^p$ and $m(H) = m_p^*(B \cup C) < m_p^*(B) + m_p^*(C)$.

Proof.

(i) Pick $B \subseteq E$, $B \notin \mathcal{L}^p$, and put $C := E \setminus B$. Then $C \notin \mathcal{L}^p$, else $B = E \setminus C \in \mathcal{L}^p$.
Subadditivity gives $\leq$; equality would force $B \in \mathcal{L}^p$ by the finite-container test in E.

(ii) Since $H \in \mathcal{L}^p$, $C \in \mathcal{L}^p$ would give $B = H \setminus C \in \mathcal{L}^p$, contradiction. Also $m_p^*(C) > 0$, else $C \in \mathcal{L}^p$.
Thus $B \cup C = H$ and $m(H) = m_p^*(B)$ give $m_p^*(B \cup C) = m(H) < m_p^*(B) + m_p^*(C)$.

Structural Commentary: Nonmeasurability appears exactly as strict failure of additive splitting.
Key Insight: Equality in a finite measurable container is a measurability test. □

Theorem 5.2.22. *A Nonmeasurable Set with Only Null Measurable Pieces*
There exists $\mathcal{U} \subseteq \mathbb{R}^p$ with $\mathcal{U} \notin \mathcal{L}^p$ such that, for every $F \in \mathcal{L}^p$,

$$F \subseteq \mathcal{U} \Longrightarrow m(F) = 0,$$
$$F \subseteq \mathcal{U}^c \Longrightarrow m(F) = 0.$$

Proof. Let $e := (1, \dots, 1) \in \mathbb{R}^p$, $G := \mathbb{Q}^p + \sqrt{2}\,\mathbb{Z}e$, $G_0 := \mathbb{Q}^p + (2\sqrt{2})\,\mathbb{Z}e$, and $G_1 := \sqrt{2}\,e + G_0$. Then $G = G_0 \sqcup G_1$, and G_0, G_1 are dense. Choose $\mathcal{E} \subseteq \mathbb{R}^p$ with one representative from each class of $x - y \in G$, and set $\mathcal{U} := G_0 + \mathcal{E}$.
Unique representatives give $\mathbb{R}^p = (G_0 + \mathcal{E}) \sqcup (G_1 + \mathcal{E})$, hence $\mathcal{U}^c = G_1 + \mathcal{E} = \sqrt{2}\,e + \mathcal{U}$.
Let $F \in \mathcal{L}^p$, $F \subseteq \mathcal{U}$, and $m(F) > 0$. Then $B(0, r) \subseteq F - F$ for some $r > 0$; choose $g_1 \in G_1 \cap B(0, r)$.
Since $F \subseteq G_0 + \mathcal{E}$, $g_1 = (g_0 + s_1) - (g_0' + s_2)$ for some $g_0, g_0' \in G_0$ and $s_1, s_2 \in \mathcal{E}$.
Thus $s_1 - s_2 = g_1 - g_0 + g_0' \in G$, so $s_1 = s_2$ and $g_1 = g_0 - g_0' \in G_0$, impossible.
Hence every measurable $F \subseteq \mathcal{U}$ is null; translation gives the same for $\mathcal{U}^c$.
If $\mathcal{U} \in \mathcal{L}^p$, then $\mathcal{U}$ and $\mathcal{U}^c$ are null, so $m([0,1]^p) = 0$, contradiction.
Structural Commentary: This construction partitions representatives using two dense cosets.
Key Insight: Positive measurable pieces force forbidden coset differences. □

Definition 5.2.23. *Cantor–Lebesgue Singular Function*
Let $P \subseteq [0,1]$ be the Cantor set. For $x = (0.c_1c_2\cdots)_3 \in P$ with $c_k \in \{0, 2\}$, set $\varphi(x) := (0.(c_1/2)(c_2/2)\cdots)_2$.
If (a, b) is a component of $[0,1] \setminus P$, extend by $\varphi(x) := \varphi(a) = \varphi(b)$ for $x \in (a, b)$.
The resulting $\varphi : [0,1] \to [0,1]$ is the *Cantor–Lebesgue singular function*.

Proposition 5.2.24. *Properties of the Singular Function*
Let $P \subseteq [0,1]$ be the Cantor set and let $\varphi : [0,1] \to [0,1]$ be the Cantor–Lebesgue singular function. Then φ is continuous, nondecreasing, onto, nonconstant, and $\varphi' = 0$ on $[0,1] \setminus P$.

Proof. On P, the digit map $\tau((0.c_1c_2\cdots)_3) := (0.(c_1/2)(c_2/2)\cdots)_2$ is nondecreasing and onto $[0,1]$.
The only duplications are endpoint pairs of removed middle-third intervals, where the binary values agree; constant extension across those intervals is therefore nondecreasing with no jumps, hence continuous.
Each component of $[0,1] \setminus P$ is an interval of constancy, so $\varphi' = 0$ there. Since $\varphi([0,1]) = [0,1]$, φ is nonconstant.
Structural Commentary: Variation is confined to P; the range is $[0,1]$.
Key Insight: Ternary 0/2 digits are re-read as binary digits. □

Theorem 5.2.25. *Lebesgue Measurable Non-Borel Sets Exist*
There exist $A \in \mathcal{L}^1 \setminus \mathcal{B}(\mathbb{R})$ and a homeomorphism $\psi : [0,1] \to [0,2]$ such that $\psi(A) \notin \mathcal{L}^1$.

Proof. Let $\psi(x) := x + \varphi(x)$, where φ is the Cantor–Lebesgue singular function. Since φ is continuous and nondecreasing, ψ is continuous, strictly increasing, maps $[0,1]$ onto $[0,2]$, and hence is a homeomorphism.
Let P be the Cantor set and $U := [0,1] \setminus P$. On each component I of U, φ is constant, so $m(\psi(I)) = m(I)$; the intervals $\psi(I)$ are disjoint. Hence $m(\psi(U)) = m(U) = 1$.
Since $[0,2] = \psi(P) \sqcup \psi(U)$ and $\psi(P)$ is compact, $m(\psi(P)) = 1$. Choose $W \subseteq \psi(P)$ with $W \notin \mathcal{L}^1$, and set $A := \psi^{-1}(W) \subseteq P$. Since $m(P) = 0$, $A \in \mathcal{L}^1$. If $A \in \mathcal{B}(\mathbb{R})$, then $W = \psi(A)$ is Borel in $[0,2]$, hence Lebesgue

measurable, contradiction. Thus $A \in \mathcal{L}^1 \setminus \mathcal{B}(\mathbb{R})$ and $\psi(A) = W \notin \mathcal{L}^1$.
Structural Commentary: Completion by null sets creates Lebesgue sets invisible to the Borel hierarchy.
Key Insight: A null Cantor set can be homeomorphically stretched to positive measure. □

5.3 Measurable Functions

Definition 5.3.1. *Borel, Lebesgue, and Function Measurability*

(i) *Borel*: $\mathcal{B}(\mathbb{R}^p) := \sigma(\mathcal{O}_{\mathbb{R}^p})$; its elements are *Borel sets*.
(ii) *Lebesgue*: $\mathcal{B}(\mathbb{R}^p) \subseteq \mathcal{L}^p$, and $\mathcal{L}^p$ is obtained by adjoining subsets of Lebesgue null sets.
(iii) *Maps*: measurable spaces $(X, \mathcal{M})$, $(Y, \mathcal{N})$, $f : X \to Y$ is *$(\mathcal{M},\mathcal{N})$-measurable* iff $f^{-1}(E) \in \mathcal{M} \forall E \in \mathcal{N}$.
(iv) *Borel maps*: $f : \mathbb{R}^p \to \mathbb{R}^q$ is *Borel measurable* iff $f^{-1}(B) \in \mathcal{B}(\mathbb{R}^p)$ for all $B \in \mathcal{B}(\mathbb{R}^q)$.
(v) *Thresholds*: write $\{f \bowtie a\} := \{x \in X : f(x) \bowtie a\}$ for $\bowtie \in \{>, \geq, <, \leq\}$; extended-real f is *$\mathcal{M}$-measurable* iff $\{f > a\} \in \mathcal{M}$ for all $a \in \mathbb{R}$.

Structural Commentary: Borel measurability tracks topology; Lebesgue measurability adds null-set completion.
Key Insight: Measurability is inverse-image stability.

Proposition 5.3.2. *Generator Test for Measurable Maps*
Let $f : X \to Y$, $\mathcal{A} \subseteq \mathcal{P}(Y)$, and $\mathcal{N} := \sigma(\mathcal{A})$.

(i) If $f^{-1}(A) \in \mathcal{M}$ for all $A \in \mathcal{A}$, then $f : X \to (Y, \mathcal{N})$ is measurable.
(ii) For real-valued f, $\forall a \in \mathbb{R}$, $\{f > a\} \in \mathcal{M} \iff \forall B \in \mathcal{B}(\mathbb{R})$, $f^{-1}(B) \in \mathcal{M}$.

Proof.

(i) Let $\mathcal{G} := \{E \subseteq Y : f^{-1}(E) \in \mathcal{M}\}$. Since $f^{-1}(E^c) = f^{-1}(E)^c$ and $f^{-1}(\bigcup_n E_n) = \bigcup_n f^{-1}(E_n)$, $\mathcal{G}$ is a σ-algebra. Thus $\mathcal{A} \subseteq \mathcal{G} \Longrightarrow \mathcal{N} = \sigma(\mathcal{A}) \subseteq \mathcal{G}$.
(ii) ($\Rightarrow$) Put $\mathcal{A} := \{(a, +\infty) : a \in \mathbb{R}\}$. Since $f^{-1}((a, +\infty)) = \{f > a\} \in \mathcal{M}$ for all a and $\sigma(\mathcal{A}) = \mathcal{B}(\mathbb{R})$, (i) gives $f^{-1}(B) \in \mathcal{M}$ for all $B \in \mathcal{B}(\mathbb{R})$.
(ii) ($\Leftarrow$) Assume $f^{-1}(B) \in \mathcal{M}$ for every $B \in \mathcal{B}(\mathbb{R})$.
For each $a \in \mathbb{R}$, $(a, +\infty) \in \mathcal{B}(\mathbb{R})$, hence $\{f > a\} = f^{-1}((a, +\infty)) \in \mathcal{M}$.

Structural Commentary: Measurability can be checked on any generating family of target sets.
Key Insight: Pullbacks turn generated σ-algebras into generated collections of preimages. □

Proposition 5.3.3. *Open-Set Test for Euclidean Targets*
Let $\Omega \subseteq \mathbb{R}^n$ be Lebesgue measurable and let $f : \Omega \to \mathbb{R}^m$. The following tests are equivalent:

(i) $f^{-1}(V)$ is measurable for every open $V \subseteq \mathbb{R}^m$;
(ii) $f^{-1}(B)$ is measurable for every open box $B \subseteq \mathbb{R}^m$;
(iii) every coordinate function $f_j : \Omega \to \mathbb{R}$ is threshold-measurable.

In particular, every continuous $f : \Omega \to \mathbb{R}^m$ is measurable.

Proof.

(i) $\Rightarrow$ *(ii)*. Open boxes are open.
(ii) $\Rightarrow$ *(i)*. By Lemma 5.1.59, write $V = \bigcup_k B_k$ with rational boxes B_k.
Then $f^{-1}(V) = \bigcup_k f^{-1}(B_k)$ is measurable.
(ii) $\Rightarrow$ *(iii)*. By the previous implication, $f^{-1}(U)$ is measurable for every open $U \subseteq \mathbb{R}^m$.
Hence $\{f_j > a\} = f^{-1}(\{y : y_j > a\})$ is measurable.
(iii) $\Rightarrow$ *(ii)*. If $B = \prod_{j=1}^m (a_j, b_j)$, then $f^{-1}(B) = \bigcap_{j=1}^m \big(\{f_j > a_j\} \cap \{f_j < b_j\}\big)$,
where $\{f_j < b_j\} = \Omega \setminus \bigcap_{\ell=1}^{\infty} \{f_j > b_j - 1/\ell\}$. Thus each factor is measurable, hence so is $f^{-1}(B)$.

If f is continuous, Theorem 2.3.18 makes $f^{-1}(V)$ open relative to Ω.
By Theorem 2.1.41, it is $W \cap \Omega$ for some open $W \subseteq \mathbb{R}^n$, hence measurable by Theorem 5.1.60.
Structural Commentary: Euclidean-valued measurability is just countable box algebra applied to inverse images.
Key Insight: Open sets reduce to rational boxes; vector functions reduce to coordinate threshold tests. □

Example 5.3.4. *Basic Measurable Functions*
Let $(X, \mathcal{M})$ be a measurable space.

(i) Constant real-valued functions are measurable, since each strict superlevel set is either X or $\emptyset$.
(ii) For $E \subseteq X$, the indicator $\mathbf{1}_E$ is measurable iff $E \in \mathcal{M}$.
(iii) Continuous maps between Euclidean measurable subsets are measurable by Proposition 5.3.3.
(iv) Monotone functions $\mathbb{R} \to \mathbb{R}$ are Borel measurable, since their strict superlevel sets are half-lines, $\mathbb{R}$, or $\emptyset$.

Key Insight: Constants, indicators of measurable sets, continuous functions, and monotone functions form the first reusable stock of measurable functions.

Theorem 5.3.5. *Closure Properties of Measurable Functions*
Let $(X, \mathcal{M})$ be a measurable space.

(i) The four tests $\{f > a\}, \{f \geq a\}, \{f < a\}, \{f \leq a\} \in \mathcal{M}$ for all $a \in \mathbb{R}$ are equivalent.
(ii) f measurable implies $|f|$ measurable.
(iii) If f_n are measurable, then $\sup_n f_n, \inf_n f_n, \limsup_n f_n, \liminf_n f_n$ are measurable.
(iv) If f, g real-valued measurable, $F : \mathbb{R}^2 \to \mathbb{R}$ continuous, then $F(f, g)$ measurable; so are $f + g$ and fg.

Proof.

(i) Use $\{f \geq a\} = \bigcap_{m \geq 1} \{f > a - 1/m\}$; $\{f < a\} = X \setminus \{f \geq a\}$; $\{f \leq a\} = \bigcap_{m \geq 1} \{f < a + 1/m\}$.
(ii) For $a > 0$, $\{|f| < a\} = \{f < a\} \cap \{f > -a\}$; for $a \leq 0$ the set is empty.
(iii) $\{\sup_n f_n > a\} = \bigcup_n \{f_n > a\}$. Also $\limsup_n f_n = \inf_m \sup_{n \geq m} f_n$; $\liminf_n f_n = \sup_m \inf_{n \geq m} f_n$; inf follow from (i).
(iv) For $G_a = \{(u, v) : F(u, v) > a\}$, write $G_a = \bigcup_k (\alpha_k, \beta_k) \times (\gamma_k, \delta_k)$.

Then $\{F(f, g) > a\} = \bigcup_k \big(\{\alpha_k < f < \beta_k\} \cap \{\gamma_k < g < \delta_k\}\big) \in \mathcal{M}$.

Structural Commentary: Measurability is countable set algebra applied to level sets.
Key Insight: Once threshold sets are measurable, ordinary limiting and algebraic operations stay measurable. □

Definition 5.3.6. *Positive and Negative Parts*
For an extended-real-valued function f, define $f^+(x) := \sup\{f(x), 0\}$; $f^-(x) := \sup\{-f(x), 0\}$.
When f is finite-valued, $f = f^+ - f^-$ and $|f| = f^+ + f^-$; equivalently $f^+ = \frac{|f| + f}{2}$; $f^- = \frac{|f| - f}{2}$.

Corollary 5.3.7. *Measurability of Positive and Negative Parts*
If $f : X \to [-\infty, +\infty]$ is measurable, then $f^+, f^-, |f|$ are measurable.
If $f : X \to \mathbb{R}$, then f measurable $\iff$ f^+, f^- measurable.

Proof.

($\Rightarrow$) Since $\{-f > a\} = \{f < -a\}$, $-f$ is measurable.
Thus $f^+ = f \vee 0$, $f^- = (-f) \vee 0$, and $|f| = f \vee (-f)$ are measurable by finite sup-closure.
($\Leftarrow$) If $f : X \to \mathbb{R}$ and f^+, f^- are measurable, then $f = f^+ - f^-$ is measurable.

Structural Commentary: Sign decomposition is compatible with threshold-set measurability.
Key Insight: Reduce signs to suprema and finite algebra. □

Lemma 5.3.8. *Extended-Valued Measurability via the Finite Part*
Let $f : X \to \overline{\mathbb{R}}$ and set $A := f^{-1}(\{+\infty\})$, and $B := f^{-1}(\{-\infty\})$. Define $f_0(x) := \begin{cases} f(x), & x \notin A \cup B, \\ 0, & x \in A \cup B. \end{cases}$
Then f is measurable iff $A, B \in \mathcal{M}$ and $f_0 : X \to \mathbb{R}$ is measurable.

Proof.

($\Rightarrow$) If f is measurable, then $A = \bigcap_{n=1}^{\infty} \{f > n\}$; and $B = X \setminus \bigcup_{n=1}^{\infty} \{f > -n\}$, so $A, B \in \mathcal{M}$.

Also, $\{f_0 > \alpha\} = \begin{cases} \{f > \alpha\} \setminus A, & \alpha \geq 0, \\ \{f > \alpha\} \cup B, & \alpha < 0, \end{cases}$, so f_0 is measurable.

($\Leftarrow$) If $A, B \in \mathcal{M}$, f_0 is measurable, then $\{f > \alpha\} = \begin{cases} \{f_0 > \alpha\} \cup A, & \alpha \geq 0, \\ \{f_0 > \alpha\} \setminus B, & \alpha < 0, \end{cases}$, so $\{f > \alpha\} \in \mathcal{M}$ $\forall \alpha \in \mathbb{R}$.

Structural Commentary: The infinite-level sets are measurable exactly when the finite core is measurable.
Key Insight: Assign harmless finite values at $\pm\infty$, then put the infinite fibers back by set algebra. □

Corollary 5.3.9. *Extended-Valued Algebra Conventions*
Let $f, g : X \to \overline{\mathbb{R}}$ be measurable.

(i) For $c \in \mathbb{R}$, cf, $|f|$, f^+, and f^- are measurable, with $0(\pm\infty) = 0$.
(ii) Let $E_1 := \{f = -\infty,\ g = +\infty\}$, $E_2 := \{f = +\infty,\ g = -\infty\}$, and $F := E_1 \cup E_2$.
If $h = f + g$ on $X \setminus F$ and $h = 0$ on F, then h is measurable.
(iii) With $0(\pm\infty) = 0$, fg is measurable.

Proof.

(i) By Lemma 5.3.8, $A_{\pm}(f) := \{f = \pm\infty\} \in \mathcal{M}$ and f_0 is measurable. For $c \in \mathbb{R}$, the infinite fibers of cf are among $A_{\pm}(f)$ and $(cf)_0 = cf_0$, so cf is measurable. Corollary 5.3.7 gives $|f|, f^+, f^-$.

(ii) $E_1, E_2 \in \mathcal{M}$. Set $f_n := \max\{-n, \min(f, n)\}$ and $g_n := \max\{-n, \min(g, n)\}$. Then f_n, g_n are real-valued measurable and $(f_n + g_n)\mathbf{1}_{X\setminus F} \to h$ pointwise; hence h is measurable.

(iii) With f_n, g_n as in (ii), $f_n g_n$ measurable, $f_n g_n \to fg$ pointwise under $0(\pm\infty) = 0$; hence fg measurable.

Structural Commentary: Extended algebra is finite algebra plus measurable exceptional fibers.
Key Insight: Truncate first, use real-valued closure, then pass to pointwise limits. □

Definition 5.3.10. *Simple Functions*
Let $(X, \mathcal{M})$ be a measurable space.

(i) *Indicator*: for $E \subseteq X$, $\mathbf{1}_E(x) := 1$ on E and $\mathbf{1}_E(x) := 0$ on $X \setminus E$.

(ii) *Simple*: a real-valued s is simple iff $s(X)$ is finite.

(iii) *Standard representation*: if $s(X) = \{c_1, \ldots, c_n\}$ and $E_i := \{s = c_i\} \neq \emptyset$, then $X = \bigsqcup_{i=1}^{n} E_i$; $s = \sum_{i=1}^{n} c_i \mathbf{1}_{E_i}$. The pairs (c_i, E_i) are unique up to reordering.

(iv) *Measurability*: s measurable iff $E_i \in \mathcal{M}$ $\forall i$; step functions are special case $X \subseteq \mathbb{R}$ with interval level sets.

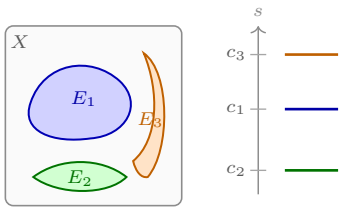


Figure 194: A simple function takes value c_i on each measurable piece E_i.

Proposition 5.3.11. *Simple Functions Form a Lattice Algebra*
Let $s, t : X \to \mathbb{R}$ be simple and $a \in \mathbb{R}$. Then as, $s + t$, st, $s \vee t$, and $s \wedge t$ are simple.
If s, t are measurable, so are these five functions.

Proof. Let $s(X) = \{c_i\}_{i=1}^m$ and $t(X) = \{d_j\}_{j=1}^n$. Each displayed range is contained in the corresponding finite set $\{ac_i\}$, $\{c_i + d_j\}$, $\{c_i d_j\}$, $\{c_i \vee d_j\}$, or $\{c_i \wedge d_j\}$, hence is finite.
If s, t are measurable, the atoms $P_{ij} := \{s = c_i\} \cap \{t = d_j\}$ are measurable and form a finite partition of X. Each displayed function is constant on every P_{ij}, so its level sets are finite unions of P_{ij}.
Structural Commentary: Algebra and lattice operations reduce to finite tables.
Key Insight: Refine two simple functions by the common partition of their level sets. □

Theorem 5.3.12. *Approximation by Simple Functions*
Every real f on X is the pointwise limit of simple functions s_n.
If f is measurable, choose s_n measurable; if $f \geq 0$, choose $0 \leq s_n \uparrow f$.

Proof. For $f \geq 0$, set $E_{ni} = \left\{x : \frac{i-1}{2^n} \leq f(x) < \frac{i}{2^n}\right\}$; $F_n = \{f \geq n\}$; and $s_n = \sum_{i=1}^{n2^n} \frac{i-1}{2^n} \mathbf{1}_{E_{ni}} + n\mathbf{1}_{F_n}$.
Then s_n is simple, measurable when f is, and $0 \leq s_n \uparrow f$. For general f, approximate f^+ and f^- and subtract.
Structural Commentary: Simple functions quantize the range of f.
Key Insight: Dyadic range partitions give pointwise approximation while preserving measurability. □

Definition 5.3.13. *Complex-Valued Measurable Functions*
A map $f : X \to \mathbb{C}$ is *measurable* iff $\operatorname{Re} f, \operatorname{Im} f : X \to \mathbb{R}$ are measurable.
Equivalently, $f = \operatorname{Re} f + i \operatorname{Im} f$ with measurable real components.

Proposition 5.3.14. *Complex Measurable Closure*
If $f, g : X \to \mathbb{C}$ are measurable, $a \in \mathbb{C}$, and $h_n : X \to \mathbb{C}$ are measurable with $h_n \to h$ pointwise, then $f + g$, fg, af, and h are measurable.

Proof. Write $f = u + iv$, $g = r + is$, and $a = \alpha + i\beta$. Then $f + g = (u + r) + i(v + s)$, $af = (\alpha u - \beta v) + i(\alpha v + \beta u)$, and $fg = (ur - vs) + i(us + vr)$, so real closure gives measurability.
If $h_n \to h$, then $\operatorname{Re} h_n \to \operatorname{Re} h$ and $\operatorname{Im} h_n \to \operatorname{Im} h$, so pointwise-limit closure gives measurability.
Structural Commentary: Complex measurability is two real measurability statements.
Key Insight: Split into real and imaginary parts. □

Definition 5.3.15. *Almost Everywhere and In Measure*
Let E be measurable.

(i) *A.E. equality*: $f \sim_{\text{a.e.}} g$ on E iff $\mu(\{x \in E : f(x) \neq g(x)\}) = 0$.
(ii) *A.E. property*: a property holds a.e. on E iff its failure set in E is null.
(iii) *Null invariance*: $\sim_{\text{a.e.}}$ is an equivalence relation, and $f \sim_{\text{a.e.}} g$ on E preserves integrals over measurable $A \subseteq E$ whenever they exist.
(iv) *A.E. convergence*: $f_n \to f$ a.e. on E iff $\exists N \subseteq E$ null such that $f_n(x) \to f(x)$ for all $x \in E \setminus N$.
(v) *In measure*: $f_n \to f$ in measure on E iff $\forall \varepsilon > 0$, $\mu(\{x \in E : |f_n(x) - f(x)| > \varepsilon\}) \to 0$.

If $f \in \mathscr{L}^1(\mu)$ on E, then f is finite a.e. on E.

Remark 5.3.16. *Completion and Almost-Everywhere Limits*
If $E \subseteq N \in \mathcal{M}$, $\mu(N) = 0$, and $E \notin \mathcal{M}$, then $f_n \equiv 0$ is measurable and $f_n \to \mathbf{1}_E$ a.e.,
since $\{x : f_n(x) \not\to \mathbf{1}_E(x)\} \subseteq N$. Yet $\mathbf{1}_E$ is not $\mathcal{M}$-measurable.
In the completion, every $E \subseteq N$ is measurable, so null modifications preserve measurability.
Key Insight: Almost-everywhere statements are most stable on complete measure spaces.

Proposition 5.3.17. *Completed Measurable Functions Have Original Representatives*
Let $(X, \mathcal{M}, \mu)$ be a measure space with completion $(X, \overline{\mathcal{M}}_\mu, \overline{\mu})$.
If $f : X \to \overline{\mathbb{R}}$ is $\overline{\mathcal{M}}_\mu$-measurable, then $\exists$ $\mathcal{M}$-measurable $g : X \to \overline{\mathbb{R}}$ with $f = g$ μ-a.e.

Proof. For $q \in \mathbb{Q}$, let $A_q := \{f > q\} \in \overline{\mathcal{M}}_\mu$, let $A_q = E_q \cup Z_q$, with $E_q \in \mathcal{M}$ and $Z_q \subseteq N_q \in \mathcal{M}$, $\mu(N_q) = 0$.
Put $N := \bigcup_{q \in \mathbb{Q}} N_q$, so $N \in \mathcal{M}$ and $\mu(N) = 0$, and define $g(x) := \begin{cases} f(x), & x \notin N, \\ 0, & x \in N. \end{cases}$.
For $q \in \mathbb{Q}$, $\{g > q\} = (E_q \setminus N) \cup \begin{cases} N, & q < 0, \\ \emptyset, & q \geq 0, \end{cases}$ belongs to $\mathcal{M}$. Hence, for $a \in \mathbb{R}$, $\{g > a\} = \bigcup_{\substack{q \in \mathbb{Q} \\ q > a}} \{g > q\} \in \mathcal{M}$.
Thus g is $\mathcal{M}$-measurable and $f = g$ on $X \setminus N$.
Structural Commentary: Completion only changes measurable functions on null sets.
Key Insight: Replace completed rational levels off one null set. □

Method 5.3.18. *Measure Proof Patterns*

(i) *Disjointize*. If $A_n \uparrow$, set $B_1 := A_1$, $B_n := A_n \setminus A_{n-1}$; then $A_n = \bigsqcup_{k \leq n} B_k$ and $\bigcup_n A_n = \bigsqcup_n B_n$.
(ii) *Approximate*. Prove on $\mathcal{E}$; if $d_\mu(A_n, A), d_\mu(B_n, B) \to 0$ with $d_\mu(C, D) := \mu^*(C \triangle D)$, pass limits through $\cup, \cap, \setminus$ via $\triangle$ estimates.
(iii) *Threshold*. For real f, show $\{f > a\} \in \mathcal{M}$ for all $a \in \mathbb{R}$; for $f : X \to \mathbb{R}^m$, use coordinate/open-set tests.
(iv) *Generators*. For $f : X \to (Y, \sigma(\mathcal{A}))$, check $f^{-1}(A) \in \mathcal{M}$ for all $A \in \mathcal{A}$.
(v) *Truncate*. For $f : X \to \overline{\mathbb{R}}$, prove results for $f_n := (-n) \vee (f \wedge n)$, then let $n \to \infty$.
(vi) *Topological bridge*. Continuous maps satisfy $f^{-1}(V)$ open; reduce Borel tests to countable interval or rectangle preimages.

Key Insight: Measure theory alternates between countable set algebra and approximation by elementary sets.

5.4 Lebesgue Integration

Lebesgue integration starts with simple functions, then extends by monotone approximation and positive–negative decomposition.

Definition 5.4.1. *Lebesgue Integral*
Let $(X, \mathcal{M}, \mu)$ be a measure space and $E \in \mathcal{M}$.

(i) *Simple*: if $0 \leq s = \sum_{i=1}^n c_i \mathbf{1}_{E_i}$ with $c_i \geq 0$, $E_i \in \mathcal{M}$ disjoint, set $\mathcal{I}_E(s) := \sum_{i=1}^n c_i \mu(E \cap E_i)$; $0 \cdot (+\infty) := 0$.
(ii) *Nonnegative*: measurable $f \geq 0$, $\mathcal{S}^+(f) := \{s : 0 \leq s \leq f,\ s \text{ simple measurable}\}$; $\int_E f \, d\mu := \sup_{s \in \mathcal{S}^+(f)} \mathcal{I}_E(s)$
(iii) *Signed*: real measurable f, $f^+ := f \vee 0$, $f^- := (-f) \vee 0$.
If at least one of $\int_E f^+ \, d\mu$, $\int_E f^- \, d\mu$ is finite, define $\int_E f \, d\mu := \int_E f^+ \, d\mu - \int_E f^- \, d\mu$.
(iv) *Integrable*: f is Lebesgue integrable on E, written $f \in \mathscr{L}^1(\mu)$ on E, iff $\int_E f^+ \, d\mu, \int_E f^- \, d\mu < +\infty$.
(v) *Whole space*: if $E = X$, this gives the real-valued integrable, or summable, functions on $(X, \mathcal{M}, \mu)$.

Proposition 5.4.2. *Signed Decompositions Give the Same Integral*
Let $E \in \mathcal{M}$ and $f = f_1 - f_2$ on E, with $f_i \geq 0$ measurable and $\int_E f_i \, d\mu < +\infty$ for $i = 1, 2$.
Then $f \in \mathscr{L}^1(\mu)$ on E and $\int_E f \, d\mu = \int_E f_1 \, d\mu - \int_E f_2 \, d\mu$.

Proof. On E, $f^+ \le f_1$, $f^- \le f_2$, hence $\int_E f^\pm \, d\mu < +\infty$.
Also $f^+ + f_2 = f_1 + f^-$, so nonnegative additivity gives $\int_E f^+ \, d\mu + \int_E f_2 \, d\mu = \int_E f_1 \, d\mu + \int_E f^- \, d\mu$.
Subtract the finite terms $\int_E f^- \, d\mu$ and $\int_E f_2 \, d\mu$.
Structural Commentary: Any finite nonnegative decomposition gives the same signed value.
Key Insight: Move all positive pieces to one side, use nonnegative additivity, then subtract finite numbers. □

Simple functions are finite linear combinations of indicators of measurable sets: coefficients carry size, while the level sets carry measurability.

Lemma 5.4.3. *Refining a Simple Integral*
Let $s \ge 0$ be simple measurable and suppose $s = \sum_{k=1}^m b_k \mathbf{1}_{F_k}$ with $b_k \ge 0$ and $F_k \in \mathcal{M}$.
Then, for every $E \in \mathcal{M}$, $\mathcal{I}_E(s) = \sum_{k=1}^m b_k \mu(E \cap F_k)$.

Proof. For $\eta \in \{0,1\}^m$, set $F_k^1 := F_k$, $F_k^0 := F_k^c$, discard empty atoms $G_\eta := \bigcap_{k=1}^m F_k^{\eta_k}$. On G_η, $s = c_\eta := \sum_{\eta_k = 1} b_k$.
Group equal c_η values to recover the standard representation; finite additivity gives
$\mathcal{I}_E(s) = \sum_\eta c_\eta \mu(E \cap G_\eta) = \sum_{k=1}^m b_k \sum_{\eta : \eta_k = 1} \mu(E \cap G_\eta) = \sum_{k=1}^m b_k \mu(E \cap F_k)$.
Structural Commentary: Refine nonstandard sums to one finite partition.
Key Insight: Refine first, then regroup by equal function values. □

Proposition 5.4.4. *Simple Integral Algebra*
If $s, t \ge 0$ are simple measurable, $c \ge 0$, and $E \in \mathcal{M}$, then $\mathcal{I}_E(cs) = c\mathcal{I}_E(s)$; $\mathcal{I}_E(s+t) = \mathcal{I}_E(s) + \mathcal{I}_E(t)$.
For fixed $s \ge 0$, $\nu_s(A) := \mathcal{I}_A(s)$ is a measure on $\mathcal{M}$.

Proof. Let $s = \sum_i a_i \mathbf{1}_{A_i}$ and $t = \sum_j b_j \mathbf{1}_{B_j}$ be standard.
Homogeneity follows from $\mathcal{I}_E(cs) = \sum_i (ca_i)\mu(E \cap A_i)$ and $0 \cdot (+\infty) = 0$.
On the common refinement, $s + t = \sum_{i,j} (a_i + b_j) \mathbf{1}_{A_i \cap B_j}$.
By Lemma 5.4.3, $\mathcal{I}_E(s+t) = \sum_{i,j} (a_i + b_j) \mu(E \cap A_i \cap B_j) = \sum_i a_i \mu(E \cap A_i) + \sum_j b_j \mu(E \cap B_j)$.
Finally, $\nu_s = \sum_i a_i \, \mu(\cdot \cap A_i)$ is a finite nonnegative linear combination of measures.
Structural Commentary: The integral is additive first on finite measurable partitions.
Key Insight: Add simple functions on the common refinement of their level sets. □

Proposition 5.4.5. *Monotonicity for Nonnegative Integrals*
Let $f, g \ge 0$ be measurable and let $E, F \in \mathcal{M}$.
(i) If $f \le g$ on E, then $\int_E f \, d\mu \le \int_E g \, d\mu$.
(ii) If $E \subseteq F$, then $\int_E f \, d\mu \le \int_F f \, d\mu$.

Proof.
(i) If $s \in \mathcal{S}^+(f)$, then $s\mathbf{1}_E \in \mathcal{S}^+(g)$, $\mathcal{I}_E(s) = \mathcal{I}_E(s\mathbf{1}_E) \le \int_E g \, d\mu$. Taking $\sup_{s \in \mathcal{S}^+(f)}$ gives $\int_E f \, d\mu \le \int_E g \, d\mu$.
(ii) If $E \subseteq F$ and $s \in \mathcal{S}^+(f)$, then $\mathcal{I}_E(s) \le \mathcal{I}_F(s)$. Taking $\sup_{s \in \mathcal{S}^+(f)}$ gives $\int_E f \, d\mu \le \int_F f \, d\mu$.
Structural Commentary: The nonnegative integral is a supremum over simple functions below the integrand.
Key Insight: More room below the integrand, or over a larger set, can only increase the supremum. □

Method 5.4.6. *Splitting a Simple Minorant Under a Sum*
Let $E \in \mathcal{M}$, $f, g \ge 0$ measurable, and $\omega \in \mathcal{S}^+(f+g)$ simple.
(i) *Grid*: for $n \ge 1$, use $D_n := \{0, \frac{1}{n}, \dots, 1\}$.
(ii) *Split*: define $\varphi_n(x) := \max\{\theta\omega(x) : \theta \in D_n,\ \theta\omega(x) \le f(x)\}$; $\psi_n(x) := \big((1 - \frac{1}{n})\,\omega(x) - \varphi_n(x)\big)^+$.
(iii) *Bounds*: φ_n, ψ_n are simple, $\varphi_n \le f$, $\psi_n \le g$, and $(1 - \frac{1}{n})\omega \le \varphi_n + \psi_n$.
(iv) *Use*: simple additivity gives $(1 - \frac{1}{n})\mathcal{I}_E(\omega) \le \int_E f \, d\mu + \int_E g \, d\mu$; then let $n \to \infty$ and take $\sup_{\omega \in \mathcal{S}^+(f+g)}$.
Key Insight: Grid-split ω between f and g, losing only $1/n$.

Proposition 5.4.7. *Homogeneity and Additivity for Nonnegative Integrals*
Let $f, g \ge 0$ be measurable, let $c \ge 0$, and let $E \in \mathcal{M}$. Then $\int_E cf \, d\mu = c \int_E f \, d\mu$; and $\int_E (f+g) \, d\mu = \int_E f \, d\mu + \int_E g \, d\mu$, where the identities are interpreted in $[0, +\infty]$ with $0 \cdot (+\infty) = 0$.

Proof.
(i) If $c = 0$, both sides are 0. If $c > 0$, $s \mapsto cs$ bijects $\mathcal{S}^+(f)$ with $\mathcal{S}^+(cf)$ and $\mathcal{I}_E(cs) = c\mathcal{I}_E(s)$; take suprema.

(ii) If $s \in \mathcal{S}^+(f)$ and $t \in \mathcal{S}^+(g)$, then $s+t \in \mathcal{S}^+(f+g)$ and $\mathcal{I}_E(s)+\mathcal{I}_E(t) = \mathcal{I}_E(s+t) \le \int_E (f+g)\,d\mu$. Taking $\sup_{s,t}$ gives $\int_E f\,d\mu + \int_E g\,d\mu \le \int_E (f+g)\,d\mu$. Conversely, for $\omega \in \mathcal{S}^+(f+g)$, choose simple $\varphi_n \le f$, $\psi_n \le g$ with $(1-\frac{1}{n})\omega \le \varphi_n + \psi_n$. Then $\left(1-\frac{1}{n}\right)\mathcal{I}_E(\omega) \le \mathcal{I}_E(\varphi_n) + \mathcal{I}_E(\psi_n) \le \int_E f\,d\mu + \int_E g\,d\mu$. Let $n \to \infty$, then take $\sup_{\omega \in \mathcal{S}^+(f+g)}$.

Structural Commentary: The nonnegative integral is linear before signed integrability is introduced.
Key Insight: Scale minorants; split sum-minorants on a finite grid. □

Perspective 5.4.8. *Step Functions versus Simple Functions*
Riemann: $\sup \sum_i m_i \ell(I_i)$ over interval partitions. Lebesgue: $\sup\{\mathcal{I}(s) : 0 \le s \le f,\ s = \sum_i c_i \mathbf{1}_{E_i},\ E_i \in \mathcal{M}\}$.
Structural Commentary: Riemann partitions intervals; Lebesgue partitions level sets.
Key Insight: Same finite constants, measurable pieces.

Perspective 5.4.9. *Undergraph Picture for Nonnegative Functions*
For measurable $f : \mathbb{R} \to [0,+\infty]$, set $\mathcal{U}_f := \{(x,y) \in \mathbb{R} \times [0,+\infty) : 0 \le y < f(x)\}$.
Geometrically, $\int f\,dm = m_2(\mathcal{U}_f)$; analytically, the integral is built as $\sup \mathcal{I}(s)$ over simple $0 \le s \le f$.

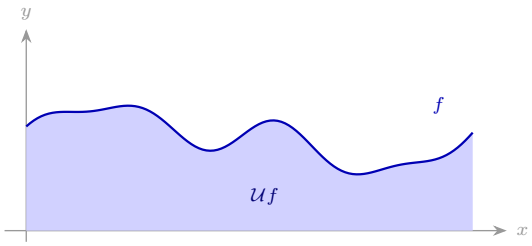


Figure 195: The integral of $f \ge 0$ is modeled by the measure of its undergraph.

Proposition 5.4.10. *Basic Integral Properties*
Let $E \in \mathcal{M}$.

(i) If f is measurable and bounded on E, and $\mu(E) < +\infty$, then $f \in \mathscr{L}^1(\mu)$ on E.
(ii) If $a \le f(x) \le b$ on E and $\mu(E) < +\infty$, then $a\mu(E) \le \int_E f\,d\mu \le b\mu(E)$.
(iii) If $f, g \in \mathscr{L}^1(\mu)$ on E and $f \le g$ on E, then $\int_E f\,d\mu \le \int_E g\,d\mu$.
(iv) If $f \in \mathscr{L}^1(\mu)$ on E, then $cf \in \mathscr{L}^1(\mu)$ on E for finite c, and $\int_E cf\,d\mu = c\int_E f\,d\mu$.
(v) If $\mu(E) = 0$ and f is measurable, then $\int_E f\,d\mu = 0$.
(vi) If $f \in \mathscr{L}^1(\mu)$ on E, $A \in \mathcal{M}$, and $A \subseteq E$, then $f \in \mathscr{L}^1(\mu)$ on A.

Proof.

(i) If $|f| \le M$ on E, then $0 \le f^{\pm} \le M\mathbf{1}_E$ on E, so $\int_E f^{\pm}\,d\mu \le M\mu(E) < +\infty$.
(ii) Since $a\mathbf{1}_E \le f \le b\mathbf{1}_E$ on E, order for the positive integral gives $a\mu(E) \le \int_E f\,d\mu \le b\mu(E)$.
(iii) Since $g - f \ge 0$ on E, $\int_E (g-f)\,d\mu \ge 0$, hence $\int_E f\,d\mu \le \int_E g\,d\mu$.
(iv) For $c \ge 0$, scale the simple approximants. For $c < 0$, use $(cf)^+ = (-c)f^-$ and $(cf)^- = (-c)f^+$.
(v) If $\mu(E) = 0$, then $\mathcal{I}_E(s) = 0$ for every $s \ge 0$ simple, so $\int_E f^{\pm}\,d\mu = 0$.
(vi) Since $A \subseteq E$, $\mathcal{I}_A(s) \le \mathcal{I}_E(s)$ for $s \ge 0$ simple; hence $\int_A f^{\pm}\,d\mu \le \int_E f^{\pm}\,d\mu < +\infty$.

Structural Commentary: Estimate $f^{\pm}$ by simpler functions.
Key Insight: Monotonicity makes bounds, null sets, and restrictions immediate. □

Theorem 5.4.11. *The Integral as a Measure or Charge*
If $f \ge 0$ is measurable, then $\phi_f : \mathcal{M} \to [0,+\infty]$, $\phi_f(A) := \int_A f\,d\mu$, is a measure.
If $f \in \mathscr{L}^1(\mu)$, then $\lambda_f : \mathcal{M} \to \mathbb{R}$, $\lambda_f(A) := \int_A f\,d\mu$, is finite-valued and countably additive.

Proof. For simple $s \ge 0$, $A \mapsto \mathcal{I}_A(s)$ is a measure.
If $A = \dot{\bigcup}_n A_n$ and $s \in \mathcal{S}^+(f)$, then $\mathcal{I}_A(s) = \sum_n \mathcal{I}_{A_n}(s) \le \sum_n \phi_f(A_n)$; taking $\sup_s$ gives $\phi_f(A) \le \sum_n \phi_f(A_n)$.
Conversely, for every N, finite additivity gives $\sum_{n=1}^N \phi_f(A_n) = \phi_f(\bigcup_{n=1}^N A_n) \le \phi_f(A)$; let $N \to \infty$.
Hence ϕ_f is countably additive. If $f \in \mathscr{L}^1(\mu)$, then $\phi_{f^+}(X), \phi_{f^-}(X) < +\infty$ and $\lambda_f = \phi_{f^+} - \phi_{f^-}$.
Thus, for $A = \dot{\bigcup}_n A_n$, $\lambda_f(A) = \sum_n \phi_{f^+}(A_n) - \sum_n \phi_{f^-}(A_n) = \sum_n \lambda_f(A_n)$. *Structural Commentary*: Setwise integration turns densities into measures.
Key Insight: Sup simple measures; subtract finite signs. □

Corollary 5.4.12. *Null and Finiteness Consequences for $f \ge 0$*
Let $f : X \to [0,+\infty]$ be measurable and $\phi_f(A) := \int_A f\,d\mu$.

(i) $\mu(A) = 0 \Longrightarrow \phi_f(A) = 0$ for $A \in \mathcal{M}$.

(ii) $\int_X f\,d\mu = 0 \iff f = 0$ *a.e.*
(iii) $\int_X f\,d\mu < +\infty \implies \mu\{f = +\infty\} = 0$ *and* $\{f > 0\}$ *is* σ*-finite.*

Proof.

(i) If $\mu(A) = 0$ and $0 \le s \le f$ is simple, then $\mathcal{I}_A(s) = 0$; hence $\phi_f(A) = \sup_s \mathcal{I}_A(s) = 0$.
(ii) ($\Rightarrow$) Set $E_n := \{f > 1/n\}$. Since $(1/n)\mathbf{1}_{E_n} \le f$, $0 \le n^{-1}\mu(E_n) \le \int_X f\,d\mu = 0$.
Thus $\mu(E_n) = 0$ for all n, and $\{f > 0\} = \bigcup_n E_n$ is null.
(ii) ($\Leftarrow$) If $N := \{f > 0\}$ null, every simple $0 \le s \le f$ vanishes on $X \setminus N$, so $\mathcal{I}_X(s) = \mathcal{I}_N(s) = 0$; take $\sup_s$.
(iii) Let $F_n := \{f \ge n\}$. Then $n\mu(F_n) \le \int_X f\,d\mu$, so $\mu\{f = +\infty\} \le \mu(F_n) \le n^{-1}\int_X f\,d\mu \to 0$.
Also $\{f > 0\} = \bigcup_n \{f > 1/n\}$ and $\mu\{f > 1/n\} \le n\int_X f\,d\mu < +\infty$.

Structural Commentary: Finite nonnegative integrals control threshold sets.
Key Insight: Threshold sets convert integral bounds into measure bounds. □

Theorem 5.4.13. *Absolute Integrability and Domination*
If $f \in \mathscr{L}^1(\mu)$ *on* E*, then* $|f| \in \mathscr{L}^1(\mu)$ *and* $\left|\int_E f\,d\mu\right| \le \int_E |f|\,d\mu$.
If f *is measurable and* $|f| \le g \in \mathscr{L}^1(\mu)$ *on* E*, then* $f \in \mathscr{L}^1(\mu)$ *on* E.

Proof. Use $|f| = f^+ + f^-$, $f \le |f|$, and $-f \le |f|$. If $|f| \le g$, then $0 \le f^\pm \le g$, so $\int_E f^\pm\,d\mu < +\infty$.
Structural Commentary: Lebesgue integrability is absolute integrability.
Key Insight: A single integrable majorant controls both signs. □

Proposition 5.4.14. *Linearity of the Lebesgue Integral*
Let $E \in \mathcal{M}$*,* $f, g \in \mathscr{L}^1(\mu)$ *on* E*, and* $a, b \in \mathbb{R}$.
Then $af + bg \in \mathscr{L}^1(\mu)$ *on* E *and* $\int_E (af + bg)\,d\mu = a\int_E f\,d\mu + b\int_E g\,d\mu$.

Proof. For $f, g \ge 0$, choose simple $s_n \uparrow f$ and $t_n \uparrow g$.
Then $s_n + t_n \uparrow f + g$ and $\int_E (s_n + t_n)\,d\mu = \int_E s_n\,d\mu + \int_E t_n\,d\mu \to \int_E f\,d\mu + \int_E g\,d\mu$.
Thus nonnegative additivity holds; homogeneity follows by scaling simple minorants.
For signed f, g, put $a = a^+ - a^-$, $b = b^+ - b^-$ and $P := a^+f^+ + a^-f^- + b^+g^+ + b^-g^-$;
$Q := a^+f^- + a^-f^+ + b^+g^- + b^-g^+$. Then $P, Q \ge 0$, $\int_E P\,d\mu, \int_E Q\,d\mu < +\infty$, and $af + bg = P - Q$.
By Proposition 5.4.2 and nonnegative linearity, the displayed identity follows.
Structural Commentary: Simple additivity plus monotone convergence gives linearity.
Key Insight: Reduce additivity to simple functions before splitting signs. □

Proposition 5.4.15. *L^1 Level Sets, Null Changes, and Indefinite Integrals*
Let $f, g \in \mathscr{L}^1(\mu)$ *on* X*, and write* $\lambda_u(A) := \int_A u\,d\mu$.

(i) *For* $a > 0$, $\mu\{|f| \ge a\} \le a^{-1}\int_X |f|\,d\mu < +\infty$*; hence* $\{f \ne 0\}$ *is* σ*-finite.*
(ii) *If measurable* $h = 0$ *a.e., then* $h \in \mathscr{L}^1(\mu)$ *and* $\lambda_h \equiv 0$.
(iii) $\lambda_f \ge 0$ *on* $\mathcal{M} \implies f \ge 0$ *a.e.;* $\lambda_f \equiv 0 \implies f = 0$ *a.e.*
(iv) $\lambda_f = \lambda_g$ *on* $\mathcal{M} \iff f = g$ *a.e.*

Proof.

(i) Since $a\mathbf{1}_{\{|f| \ge a\}} \le |f|$, integration gives $a\mu\{|f| \ge a\} \le \int_X |f|\,d\mu$. Also $\{f \ne 0\} = \bigcup_{k \ge 1}\{|f| \ge 1/k\}$.
(ii) Let $N := \{h \ne 0\}$. For $A \in \mathcal{M}$, $h^\pm\mathbf{1}_A$ vanish off $A \cap N$, a null set.
Hence $\int_A h^+\,d\mu = \int_A h^-\,d\mu = 0$, so $\lambda_h(A) = 0$.
(iii) For $k \ge 1$, set $N_k := \{f \le -1/k\}$. By (i), $\mu(N_k) < +\infty$, and $\lambda_f(N_k) \le -k^{-1}\mu(N_k)$.
If $\lambda_f \ge 0$, then $\mu(N_k) = 0$, so $\{f < 0\} = \bigcup_k N_k$ is null. If $\lambda_f \equiv 0$, apply this to f and $-f$.
(iv) If $\lambda_f = \lambda_g$, then $\lambda_{f-g} \equiv 0$, so $f - g = 0$ a.e. by (iii). Conversely, if $f = g$ a.e., apply (ii) to $h = f - g$.

Structural Commentary: An L^1 function is detected by its integrals over all measurable sets, up to null sets.
Key Insight: Test sign on the level sets where the function is bounded away from 0. □

Proposition 5.4.16. *L^1 Approximation and Bounded Multipliers*
Let $f \in \mathscr{L}^1(\mu)$ *on* X.

(i) $\forall \varepsilon > 0$, $\exists$ *measurable simple* φ *with* $\int_X |f - \varphi|\,d\mu < \varepsilon$.
(ii) *If* g *is bounded measurable, then* $fg \in \mathscr{L}^1(\mu)$.

Proof.

(i) Choose simple $0 \le s_n \uparrow f^+$ and $0 \le t_n \uparrow f^-$. Then $\int_X (f^+ - s_n)\,d\mu \to 0$ and $\int_X (f^- - t_n)\,d\mu \to 0$.
For large n, $\varphi := s_n - t_n$ satisfies $\int_X |f - \varphi|\,d\mu \le \int_X (f^+ - s_n)\,d\mu + \int_X (f^- - t_n)\,d\mu < \varepsilon$.
(ii) If $|g| \le M$, then fg is measurable and $\int_X |fg|\,d\mu \le M\int_X |f|\,d\mu < +\infty$.

Structural Commentary: L^1 is simple-closed and bounded-multiplier stable.
Key Insight: Approximate signs separately; dominate products by $M\,|f|$. □

Proposition 5.4.17. *Finite-Measure Indicators by Intervals and Continuous Functions*
Let $A \subseteq \mathbb{R}$ be Lebesgue measurable with $m(A) < +\infty$. For every $\varepsilon > 0$:

(i) $\exists G = \bigcup_{j=1}^{N} I_j$, *$I_j$ open intervals, such that* $m(A \triangle G) < \varepsilon$.
(ii) $\exists f \in C(\mathbb{R})$ *such that* $0 \le f \le 1$ *and* $\int_{\mathbb{R}} |\mathbf{1}_A - f|\, dm < \varepsilon$.

Proof.

(i) Choose open $O \supseteq A$ with $m(O \setminus A) < \varepsilon/2$ and $m(O) < +\infty$. Write $O = \bigsqcup_{j\ge 1} I_j$ with open intervals I_j, choose N with $m(O \setminus G) < \varepsilon/2$ for $G := \bigcup_{j=1}^{N} I_j$. Then $A \setminus G \subseteq O \setminus G$ and $G \setminus A \subseteq O \setminus A$, so $m(A \triangle G) < \varepsilon$.
(ii) Apply (i) with $\varepsilon/2$ to get finite open G. Choose compact $K \subseteq G$ with $m(G \setminus K) < \varepsilon/2$, and choose piecewise-linear f with $0 \le f \le 1$, $f = 1$ on K, $f = 0$ on $\mathbb{R} \setminus G$. Then $\int_{\mathbb{R}} |\mathbf{1}_G - f|\, dm \le m(G \setminus K) < \varepsilon/2$, hence $\int_{\mathbb{R}} |\mathbf{1}_A - f|\, dm \le m(A \triangle G) + \int_{\mathbb{R}} |\mathbf{1}_G - f|\, dm < \varepsilon$.

Structural Commentary: Finite-measure sets can be made geometrically simple up to small L^1 error.
Key Insight: Approximate by intervals, then smooth the indicator. □

Example 5.4.18. *Counting Measure Gives Absolute Series*
Let μ be counting measure on $(\mathbb{N}, 2^{\mathbb{N}})$. For $f : \mathbb{N} \to \mathbb{R}$,
$f \in \mathscr{L}^1(\mu) \iff \sum_{n=1}^{\infty} |f(n)| < +\infty$; $\int_{\mathbb{N}} f\, d\mu = \sum_{n=1}^{\infty} f(n)$ when these equivalent conditions hold.

Example 5.4.19. *L^1 Is Not Closed Under Squaring*
On $X = \mathbb{N}$ with $\mathcal{M} = 2^{\mathbb{N}}$, define $\mu(E) := \sum_{n\in E} 4^{-n}$ and $f(n) := 2^n/\sqrt{n}$.
Then $\int_X |f|\, d\mu = \sum_{n=1}^{\infty} \frac{1}{2^n\sqrt{n}} < +\infty$; $\int_X f^2\, d\mu = \sum_{n=1}^{\infty} \frac{1}{n} = +\infty$.
The first series is dominated by the geometric series $\sum 2^{-n}$, while the second diverges by Corollary 3.1.13.
Thus $f \in \mathscr{L}^1(\mu)$ need not imply $f^2 \in \mathscr{L}^1(\mu)$.

Theorem 5.4.20. *Comparison with the Riemann Integral*
On $[a, b]$ with Lebesgue measure μ, $f \in \mathcal{R}$ implies $f \in \mathscr{L}^1(\mu)$ and the integrals agree.
If f is bounded with discontinuity set D, then $f \in \mathcal{R} \iff \mu(D) = 0$.

Proof. For refining partitions, let $L_n \uparrow L$ and $U_n \downarrow U$ be lower and upper step envelopes.
Monotone convergence identifies $\int L$ and $\int U$ with the lower and upper Riemann integrals; if $f \in \mathcal{R}$, then $L = U$ a.e., so the common Lebesgue integral equals the Riemann integral.
For bounded f, put $\omega_f(x) := \lim_{r\downarrow 0} \operatorname{diam} f\big([a, b] \cap (x - r, x + r)\big)$, where $D_k := \{x : \omega_f(x) \ge 1/k\}$.
Then f is continuous at x iff $\omega_f(x) = 0$, and its discontinuity set is $D = \bigcup_{k\ge 1} D_k$.

($\Rightarrow$) Fix k and $\varepsilon > 0$. Choose P with $U(P, f) - L(P, f) < \varepsilon/k$.
Intervals of P whose interiors meet D_k have total length $< \varepsilon$; adding the finitely many partition points, null by Proposition 5.1.68, covers D_k by arbitrarily short intervals.
Thus $\mu(D_k) = 0$ by Definition 5.1.66, and $\mu(D) = 0$.
($\Leftarrow$) Suppose $\mu(D) = 0$ and $|f| \le B$. If $B = 0$, there is nothing to prove. Given $\varepsilon > 0$, choose k with $(b - a)/k < \varepsilon/2$ and cover D_k by intervals of total length $< \varepsilon/(4B)$.
On the compact complement, $\omega_f < 1/k$, so a finite refinement makes the good intervals contribute $< \varepsilon/2$ to $U - L$; the covered intervals contribute $< 2B \cdot \varepsilon/(4B) = \varepsilon/2$.
Hence $U(P, f) - L(P, f) < \varepsilon$, and Theorem 4.1.41 with $\alpha(x) = x$ gives $f \in \mathcal{R}$.

Structural Commentary: Darboux gap measures oscillation sets.
Key Insight: Lebesgue measure turns "bad discontinuities" into a null-set condition. □

Corollary 5.4.21. *Riemann Integrable Functions as Monotone Step Limits*
Let $f : [a, b] \to \mathbb{R}$ be Riemann integrable. Then $\exists$ step functions $(\varphi_n), (\psi_n)$ such that

(i) $\varphi_n \uparrow$, $\psi_n \downarrow$, *and* $\varphi_n \le f \le \psi_n$ *on* $[a, b]$;
(ii) $\lim_n \int_a^b \varphi_n\, dx = \lim_n \int_a^b \psi_n\, dx = \int_a^b f\, dx$;
(iii) $\varphi_n, \psi_n \to f$ *a.e.; hence f is Lebesgue measurable and* $\int_{[a,b]} f\, dm = \int_a^b f\, dx$.

Proof.

(i) Choose nested partitions P_n with $\max_{I\in P_n} \ell(I) \to 0$ and $U(P_n, f) - L(P_n, f) \to 0$.
Let φ_n, ψ_n be the lower and upper step envelopes on P_n; refinement gives $\varphi_n \uparrow$, $\psi_n \downarrow$, and $\varphi_n \le f \le \psi_n$.
(ii) Since $\int_a^b \varphi_n\, dx = L(P_n, f)$ and $\int_a^b \psi_n\, dx = U(P_n, f)$, by Darboux gap condition, both limits equal $\int_a^b f\, dx$.

(iii) If $x \notin D_f$, then $\max_{I \in P_n} \ell(I) \to 0$ forces $\varphi_n(x), \psi_n(x) \to f(x)$.
Since $m(D_f) = 0$, convergence holds a.e.; hence f is measurable and $\int_{[a,b]} f\, dm = \int_a^b f\, dx$.

Structural Commentary: Riemann integrability is monotone step squeezing.
Key Insight: Refining partitions make Darboux sums into pointwise a.e. step approximations. □

Corollary 5.4.22. *Nonnegative Continuous Functions and Riemann Integrals*

(i) If $f \geq 0$ is continuous on $[a,b]$, then $\int_{[a,b]} f\, dm = \int_a^b f(x)\, dx$.

(ii) If $f \geq 0$ is continuous on $[0,+\infty)$, then $\int_{[0,+\infty)} f\, dm = \lim_{b \to +\infty} \int_0^b f(x)\, dx$ in $[0,+\infty]$.

Proof.

(i) Continuous f on $[a,b]$ is Riemann integrable; by Riemann–Lebesgue comparison, $\int_{[a,b]} f\, dm = \int_a^b f(x)\, dx$.

(ii) Put $f_b := f\mathbf{1}_{[0,b]}$. By (i), $\int_{[0,b]} f\, dm = \int_0^b f(x)\, dx$. As $b \uparrow +\infty$, $f_b \uparrow f$; monotone convergence gives the left limit, and the right limit is the nonnegative improper Riemann integral.

Structural Commentary: Compact comparison plus monotone truncation.
Key Insight: Truncate to compact intervals and pass upward. □

Perspective 5.4.23. *Lebesgue's Criterion as a Roadmap*
For bounded $f : [a,b] \to \mathbb{R}$ with discontinuity set D_f, $f \in \mathcal{R}[a,b] \iff m(D_f) = 0$.
Thus finite, countable, and null Cantor-type D_f are allowed, while $D_f = [a,b]$ is not.

5.5 Convergence Theorems

This subsection collects the main limit-interchange tools: monotone convergence, Fatou, and dominated convergence.

Theorem 5.5.1. *Monotone Convergence*
Let $E \in \mathcal{M}$, let f_n, f be measurable on E, and assume $0 \leq f_n \uparrow f$. Then $\int_E f_n\, d\mu \uparrow \int_E f\, d\mu$.

Proof. By monotonicity, $\int_E f_n\, d\mu \uparrow L \leq \int_E f\, d\mu$.
For $s \in \mathcal{S}^+(f)$ and $0 < c < 1$, set $E_n := \{x \in E : f_n(x) \geq cs(x)\}$.
Since $f_n \uparrow f$ and $s \leq f$, $E_n \uparrow E$, and $\int_E f_n\, d\mu \geq \int_{E_n} f_n\, d\mu \geq c \int_{E_n} s\, d\mu$.
Continuity from below for $\phi_s(A) := \int_A s\, d\mu$ gives $L \geq c \int_E s\, d\mu$. Let $c \uparrow 1$, then take $\sup_{s \in \mathcal{S}^+(f)}$.
Structural Commentary: Nonnegative integration commutes with monotone limits.
Key Insight: Simple minorants of the limit are eventually captured below the sequence. □

Theorem 5.5.2. *Series Form of Monotone Convergence*
Let $E \in \mathcal{M}$, let $f_n \geq 0$ be measurable, and set $f := \sum_{n=1}^{\infty} f_n$ on E. Then $\int_E f\, d\mu = \sum_{n=1}^{\infty} \int_E f_n\, d\mu$.

Proof. Set $S_N := \sum_{n=1}^{N} f_n$. Then $0 \leq S_N \uparrow f$.

By Theorem 5.5.1, $\int_E f\, d\mu = \lim_{N \to \infty} \int_E S_N\, d\mu = \lim_{N \to \infty} \sum_{n=1}^{N} \int_E f_n\, d\mu = \sum_{n=1}^{\infty} \int_E f_n\, d\mu$.
Structural Commentary: Nonnegative series are monotone partial sums.
Key Insight: Integrate partial sums, then pass upward. □

Theorem 5.5.3. *Fatou's Lemma*
Let $E \in \mathcal{M}$ and let $f_n \geq 0$ be measurable. Then $\int_E \liminf_{n \to \infty} f_n\, d\mu \leq \liminf_{n \to \infty} \int_E f_n\, d\mu$

Proof. Set $g_n := \inf_{i \geq n} f_i$. Then $0 \leq g_n \uparrow \liminf_i f_i$ and, for every $i \geq n$, $g_n \leq f_i$.
Hence $\int_E g_n\, d\mu \leq \inf_{i \geq n} \int_E f_i\, d\mu$. By Theorem 5.5.1, let $n \to \infty$ to obtain the claim.
Structural Commentary: Fatou compares the liminf to tail lower envelopes.
Key Insight: Apply monotone convergence to $g_n = \inf_{i \geq n} f_i$. □

Theorem 5.5.4. *Dominated Convergence*
Let $E \in \mathcal{M}$, f_n measurable, $f_n \to f$ a.e. on E, and $|f_n| \leq g \in \mathscr{L}^1(\mu)$ on E.
Then $f \in \mathscr{L}^1(\mu)$ on E and $\lim_{n \to \infty} \int_E f_n\, d\mu = \int_E f\, d\mu$.

Proof. Choose $N \subseteq E$ null such that $f_n \to f$ on $E \setminus N$; replacing all functions by their restrictions off N does not change integrals. Thus assume $f_n \to f$ everywhere. Then $|f| \le g$, so $f \in \mathscr{L}^1(\mu)$. Put $a_n := \int_E f_n \, d\mu$.
For $g + f_n \ge 0$, by Fatou, $\int_E (g+f)\, d\mu \le \liminf_n \int_E (g+f_n)\, d\mu = \int_E g\, d\mu + \liminf_n a_n$; $\int_E f\, d\mu \le \liminf_n a_n$.
For $g - f_n \ge 0$, by Fautou, $\int_E (g-f)\, d\mu \le \liminf_n \int_E (g-f_n)\, d\mu = \int_E g\, d\mu - \limsup_n a_n$; $\limsup_n a_n \le \int_E f\, d\mu$.
Thus $\liminf_n a_n = \limsup_n a_n = \int_E f\, d\mu$.
Structural Commentary: Domination turns signed convergence into two Fatou inequalities.
Key Insight: Apply Fatou to $g + f_n$ and $g - f_n$ to trap liminf and limsup. □

Corollary 5.5.5. *Monotone Convergence Almost Everywhere*
Let $E \in \mathcal{M}$ and $0 \le f_n \uparrow f$ a.e. on E. Then $\int_E f_n \, d\mu \uparrow \int_E f \, d\mu$.

Proof. Choose $N \subseteq E$ with $\mu(N) = 0$ and $f_n \uparrow f$ on $E \setminus N$. Set $g_n := f_n \mathbf{1}_{E\setminus N}$ and $g := f\mathbf{1}_{E\setminus N}$.
Then $0 \le g_n \uparrow g$, so $\int_E g_n \, d\mu \uparrow \int_E g \, d\mu$. Since $f_n = g_n$ and $f = g$ off N, the integrals over N vanish.
Structural Commentary: Monotone convergence ignores null exceptional sets.
Key Insight: Remove the exceptional set, prove the theorem there, then add back zero. □

Example 5.5.6. *Counting Measure Integrals Are Series*
Let μ be counting measure on $(\mathbb{N}, 2^{\mathbb{N}})$. For $f : \mathbb{N} \to [0, +\infty]$, set $s_N := \sum_{n=1}^{N} f(n)\mathbf{1}_{\{n\}}$; then $0 \le s_N \uparrow f$, with the case $f(k) = +\infty$ giving both sides $+\infty$. Hence $\int_{\mathbb{N}} f \, d\mu = \lim_N \int_{\mathbb{N}} s_N \, d\mu = \lim_N \sum_{n=1}^{N} f(n) = \sum_{n=1}^{\infty} f(n)$.
Structural Commentary: Counting measure turns integration into summation.
Key Insight: Finite-support truncations are the simple-function partial sums.

Example 5.5.7. *Boundary Tests for Monotone Convergence and Fatou*
On $([0, +\infty), \mathcal{B}, m)$:
(i) $f_n := \mathbf{1}_{[0,n]} \uparrow \mathbf{1}_{[0,+\infty)}$ and $\int f_n \, dm = n \uparrow +\infty$; monotone convergence allows $+\infty$.
(ii) $g_n := n^{-1}\mathbf{1}_{[n,+\infty)} \downarrow 0$ uniformly, but $\int g_n \, dm = +\infty \ \forall n$; decreasing convergence needs finite mass.
(iii) $w_n := -n^{-1}\mathbf{1}_{[0,n]} \rightrightarrows 0$, but $\int w_n \, dm = -1$; Fatou can fail without $f_n \ge 0$ or $f_n \ge -h \in L^1$.

Theorem 5.5.8. *Fatou with an Integrable Lower Bound*
Let $h \ge 0$, $h \in \mathscr{L}^1(\mu)$, and let f_n be measurable with $f_n \ge -h$.
Then, whenever the signed integrals are defined in $\overline{\mathbb{R}}$, $\int \liminf_{n\to\infty} f_n \, d\mu \le \liminf_{n\to\infty} \int f_n \, d\mu$.

Proof. Set $F_n := f_n + h \ge 0$. Fatou gives $\int \liminf_n F_n \, d\mu \le \liminf_n \int F_n \, d\mu$.
Since $\liminf_n F_n = h + \liminf_n f_n$ and $\int F_n \, d\mu = \int f_n \, d\mu + \int h \, d\mu$, subtract the finite term $\int h \, d\mu$.
Structural Commentary: Fatou survives controlled negative parts.
Key Insight: Add an integrable floor to make the sequence nonnegative. □

Remark 5.5.9. *Why the Negative Example Has No Floor*
For $w_n := -n^{-1}\mathbf{1}_{[0,n]}$, any $h \in L^1$ with $-h \le w_n \ \forall n$ must satisfy $h \ge n^{-1}$ on $[0, n]$.
Thus $h \ge k^{-1}$ on $[k-1, k]$, and $\int_0^N h \, dm \ge H_N := \sum_{k=1}^{N} \frac{1}{k} \to +\infty$, contradicting $h \in L^1$.

Perspective 5.5.10. *Domination Is the Safeguard*
$f_n \to f$ a.e. need not imply $\int f_n \to \int f$; domination $|f_n| \le g \in L^1$ restores the limit interchange.

Example 5.5.11. *A Nonuniform Limit with Vanishing Integrals*
For $x > 0$, set $f_n(x) := e^{-nx}x^{-1/2}$. Then $f_n \to 0$ pointwise, but $\sup_{0<x<\delta} f_n(x) = +\infty$ for every $\delta > 0$.
For $n \ge 1$, $0 \le f_n \le g := e^{-x}x^{-1/2}$; $g \le x^{-1/2}\mathbf{1}_{(0,1]} + e^{-x}\mathbf{1}_{[1,\infty)} \in \mathscr{L}^1(0, \infty)$.
Thus dominated convergence gives $I_n := \int_0^\infty e^{-nx}x^{-1/2}\, dx \to 0$.
Structural Commentary: Uniform convergence is unnecessary once one integrable envelope controls the sequence.
Key Insight: Domination handles the singular behavior at 0 and the tail at infinity simultaneously.

Corollary 5.5.12. *Continuity Under a Lebesgue Integral*
Let $f_t(x) := f(x, t)$ be measurable for each $t \in [a, b]$. Fix $t_0 \in [a, b]$.
If $f_t \to f_{t_0}$ a.e. as $t \to t_0$ and $|f_t| \le g \in \mathscr{L}^1(\mu)$ for all t, then $\lim_{t\to t_0} \int_X f_t \, d\mu = \int_X f_{t_0} \, d\mu$.
If $t \mapsto f(x, t)$ is continuous on $[a, b]$ for a.e. x, then $F(t) := \int_X f_t \, d\mu$ is continuous on $[a, b]$.

Proof. For $t_n \to t_0$, set $u_n := f_{t_n}$.
Then $u_n \to f_{t_0}$ a.e. and $|u_n| \le g$, so dominated convergence gives $\int_X u_n \, d\mu \to \int_X f_{t_0} \, d\mu$.
The sequential criterion gives the first claim; applying it for every t_0 gives continuity of F.
Structural Commentary: Parameter continuity reduces to dominated convergence along sequences.
Key Insight: The parameter disappears once t_n is frozen into a sequence of functions. □

Theorem 5.5.13. *Differentiation Under a Lebesgue Integral*
Let $f_t(x) := f(x,t)$ be measurable for each $t \in [a,b]$. Suppose $f_{t_0} \in \mathscr{L}^1(\mu)$ for some $t_0 \in [a,b]$, $\partial_t f$ exists on $X \times (a,b)$,
and $|\partial_t f(x,t)| \le g(x)$ for some $g \in \mathscr{L}^1(\mu)$ and all $(x,t) \in X \times (a,b)$.
Then $f_t \in \mathscr{L}^1(\mu)$ for every $t \in [a,b]$, and, with $F(t) := \int_X f_t \, d\mu$, $F'(t) = \int_X \partial_t f(x,t)\, d\mu$, $t \in (a,b)$.

Proof. For fixed t, $x \mapsto \partial_t f(x,t)$ is measurable as the pointwise limit of measurable difference quotients.
By the mean value theorem, $|f_t - f_{t_0}| \le |t - t_0|\, g$, so $|f_t| \le |f_{t_0}| + |t - t_0|\, g \in \mathscr{L}^1(\mu)$.
If $t_n \to t \in (a,b)$ and $t_n \ne t$, then $\frac{F(t_n)-F(t)}{t_n - t} = \int_X \frac{f_{t_n} - f_t}{t_n - t}\, d\mu$.
Integrand tends pointwise to $\partial_t f(\cdot, t)$ and has modulus $\le g$; dominated convergence gives derivative formula.
Structural Commentary: Bounded parameter derivatives dominate the difference quotients.
Key Insight: Differentiating the integral is just applying dominated convergence to secant slopes. □

Corollary 5.5.14. *A Pointwise Difference-Quotient Test*
Let $f_t(x) := f(x,t)$ be measurable for each $t \in [a,b]$, let $f_{t_0} \in \mathscr{L}^1(\mu)$ for some $t_0 \in [a,b]$, and fix $t_1 \in [a,b]$.
Suppose $\partial_t f(x,t_1)$ exists for every x and $\exists g \in \mathscr{L}^1(\mu)$ such that, for $t \ne t_1$, $Q_t(x) := \frac{f_t(x) - f_{t_1}(x)}{t - t_1}$, $|Q_t| \le g$.
Then $F(t) := \int_X f_t \, d\mu$ is differentiable at t_1 and $F'(t_1) = \int_X \partial_t f(x,t_1)\, d\mu$.

Proof. If $t_0 \ne t_1$, then $|f_{t_1}| \le |f_{t_0}| + |t_0 - t_1|\, g$; otherwise this is given.
Hence $f_{t_1} \in \mathscr{L}^1$, and $|f_t| \le |f_{t_1}| + |t - t_1|\, g$ gives $f_t \in \mathscr{L}^1$. For any $t_n \to t_1$, $t_n \ne t_1$, $Q_{t_n} \to \partial_t f(\cdot, t_1)$ pointwise and $|Q_{t_n}| \le g$, so dominated convergence gives $[F(t_n) - F(t_1)]/(t_n - t_1) \to \int_X \partial_t f(x,t_1)\, d\mu$.
Structural Commentary: One parameter point needs only domination of the secants into that point.
Key Insight: A global derivative bound is sufficient, but not necessary, for a one-point formula. □

Corollary 5.5.15. *One Lebesgue and One Riemann Integral*
Let $f_t(x) := f(x,t)$ be measurable for each $t \in [a,b]$, let $t \mapsto f(x,t)$ be continuous on $[a,b]$ for a.e. x, and suppose $|f_t| \le g \in \mathscr{L}^1(\mu)$ $\forall t$. Then, with Riemann integrals in t, $\int_a^b \left[\int_X f_t \, d\mu\right] dt = \int_X \left[\int_a^b f(x,t)\, dt\right] d\mu(x)$.

Proof. Set $h_t(x) := \int_a^t f(x,s)\, ds$. Then h_t is measurable, $|h_t| \le (b-a)g$, and $\partial_t h_t = f_t$.
By differentiation under the integral, $H(t) := \int_X h_t \, d\mu$ satisfies $H'(t) = \int_X f_t \, d\mu$.
Hence $\int_a^b \left[\int_X f_t \, d\mu\right] dt = H(b) - H(a) = \int_X \left[\int_a^b f(x,t)\, dt\right] d\mu(x)$. *Structural Commentary*: Finite parameter integrals move through by differentiating primitives.
Key Insight: Integrate first in the parameter up to t, then differentiate the resulting Lebesgue integral. □

Corollary 5.5.16. *Improper Parameter Interchange*
Let $f_t(x) := f(x,t)$ be measurable for each $t \in \mathbb{R}$, and let $t \mapsto f(x,t)$ be continuous for a.e. x.
Suppose $\exists g, h \in \mathscr{L}^1(\mu)$ such that $|f_t| \le g$ $\forall t$ and $\int_{\mathbb{R}} |f(x,t)|\, dt \le h(x)$ for a.e. x.
Put $G(t) := \int_X f_t \, d\mu$; $H(x) := \int_{\mathbb{R}} f(x,t)\, dt$. Then $H \in \mathscr{L}^1(\mu)$; $\int_{\mathbb{R}} G(t)\, dt = \int_X H \, d\mu$, $\int_{\mathbb{R}} G := \lim_{R \to +\infty} \int_{-R}^R G(t)\, dt$.

Proof. For $R > 0$, finite-interval interchange gives $\int_{-R}^R G(t)\, dt = \int_X H_R \, d\mu$ and $H_R(x) := \int_{-R}^R f(x,t)\, dt$.
Since $H_R \to H$ a.e. and $|H_R| \le h$, dominated convergence gives $\int_X H_R \, d\mu \to \int_X H \, d\mu$.
Structural Commentary: Finite interchanges pass to improper limits under an L^1 bound.
Key Insight: Truncate in t, interchange on finite intervals, then pass $R \to +\infty$ by domination in x. □

Example 5.5.17. *Differentiating Under the Integral Sign*
For $t > 0$, set $F(t) := \int_0^\infty x^2 e^{-tx}\, dx$. For $t_0 > 0$ and $|h| < t_0/2$, put $Q_h(x) := \frac{x^2 e^{-(t_0+h)x} - x^2 e^{-t_0 x}}{h}$
Then $Q_h \to -x^3 e^{-t_0 x}$ pointwise and $|Q_h| \le x^3 e^{-(t_0/2)x} \in \mathscr{L}^1(0,\infty)$. Hence $F'(t_0) = -\int_0^\infty x^3 e^{-t_0 x}\, dx$.
Structural Commentary: Parameter differentiation is a limit-interchange problem.
Key Insight: Prove a local-in-parameter integrable majorant for the difference quotients.

Example 5.5.18. *Laplace Transform Powers and Factorials*
For $t > 0$, set $I(t) := \int_0^\infty e^{-tx}\, dx$. Then $I(t) = \lim_{b \to +\infty} \frac{1 - e^{-tb}}{t} = t^{-1}$.
For $[a,b] \subset (0, +\infty)$ and $k \ge 0$, $0 \le x^k e^{-tx} \le x^k e^{-ax} \in \mathscr{L}^1(0,\infty)$, so repeated differentiation under the integral gives $I^{(n)}(t) = (-1)^n \int_0^\infty x^n e^{-tx}\, dx$. Since $I^{(n)}(t) = (-1)^n n! t^{-n-1}$, for every $n \in \mathbb{N} \cup \{0\}$, $\int_0^\infty x^n e^{-x}\, dx = n!$.
Structural Commentary: A parameter integral turns powers of x into derivatives in t.
Key Insight: The exponential envelope e^{-ax} controls all nearby parameter values $t \ge a$.

Proposition 5.5.19. *Uniform Convergence on Finite Measure Spaces*
Assume $\mu(X) < +\infty$.

(i) If $0 \le f_n$ are measurable and $f_n \rightrightarrows f$, then $f \ge 0$ is measurable and $\int_X f \, d\mu = \lim_n \int_X f_n \, d\mu$ in $[0, +\infty]$.
(ii) If $f_n \in \mathscr{L}^1(\mu)$ and $f_n \rightrightarrows f$, then $f \in \mathscr{L}^1(\mu)$ and $\int_X f_n \, d\mu \to \int_X f \, d\mu$ in $\mathbb{R}$.

Proof.

(i) Measurability and $f \geq 0$ pass to pointwise limits. For $\varepsilon > 0$ and $n \gg 1$, $f \leq f_n + \varepsilon \mathbf{1}_X$ and $f_n \leq f + \varepsilon \mathbf{1}_X$. Thus $\int_X f\,d\mu \leq \int_X f_n\,d\mu + \varepsilon\mu(X)$ and $\int_X f_n\,d\mu \leq \int_X f\,d\mu + \varepsilon\mu(X)$. Let $\varepsilon \downarrow 0$.

(ii) Choose N with $|f - f_N| \leq 1$. Then $f = f_N + (f - f_N) \in \mathscr{L}^1(\mu)$, and $\left|\int_X f_n\,d\mu - \int_X f\,d\mu\right| \leq \int_X |f_n - f|\,d\mu \leq \mu(X) \sup_X |f_n - f| \to 0$.

Structural Commentary: On finite measure spaces, uniform error controls integral error.
Key Insight: Finiteness of $\mu(X)$ turns a supremum bound into an integral bound. □

Example 5.5.20. *Why the Finite-Measure Hypothesis Matters*
On $(\mathbb{R}, m)$, let $f_n := n^{-1}\mathbf{1}_{[0,n]}$. Then $\|f_n\|_\infty = 1/n \to 0$, but $\int_{\mathbb{R}} f_n\,dm = 1\ \forall n$.
Thus uniform convergence need not commute with integrals when $m(\mathbb{R}) = +\infty$; Fatou only yields $0 \leq 1$.

Example 5.5.21. *Why Domination Cannot Be Dropped*
On $(\mathbb{R}, m)$, let $f_n := n\mathbf{1}_{[1/n,2/n]}$. Then $f_n \to 0$ pointwise, but $\int_{\mathbb{R}} f_n\,dm = 1\ \forall n$.
Thus pointwise convergence without domination need not commute with integrals; Fatou only yields $0 \leq 1$.

Theorem 5.5.22. *Absolutely Summable L^1 Series*
Let $f_n \in \mathscr{L}^1(\mu)$ and $\sum_{n\geq 1} \|f_n\|_1 < +\infty$.
Then $\sum_n f_n$ converges absolutely a.e. to some $f \in \mathscr{L}^1(\mu)$, and $\int_X f\,d\mu = \sum_{n\geq 1} \int_X f_n\,d\mu$.

Proof. Set $G := \sum_{n\geq 1} |f_n|$. Then $\int_X G\,d\mu = \sum_{n\geq 1} \|f_n\|_1 < +\infty$, so $G < +\infty$ a.e. and $G \in \mathscr{L}^1$.
Let $f := \sum_{n\geq 1} f_n$ a.e. and $S_N := \sum_{n=1}^N f_n$.
Since $|S_N| \leq G$ and $S_N \to f$ a.e., dominated convergence gives $\int_X f\,d\mu = \lim_N \int_X S_N\,d\mu = \sum_{n\geq 1} \int_X f_n\,d\mu$.
Structural Commentary: Absolute summability in L^1 gives a.e. absolute convergence and termwise integration.
Key Insight: First integrate the nonnegative majorant $\sum |f_n|$. □

Proposition 5.5.23. *L^1 Convergence Controls Absolute Integrals*
If $f_n, f \in \mathscr{L}^1(\mu)$ and $\|f_n - f\|_1 \to 0$, then $\|f_n\|_1 \to \|f\|_1$; $\int_X |f|\,d\mu = \lim_n \int_X |f_n|\,d\mu$.

Proof. The reverse triangle inequality gives $||f_n| - |f|| \leq |f_n - f|$; hence $|\|f_n\|_1 - \|f\|_1| \leq \|f_n - f\|_1 \to 0$.
Structural Commentary: L^1 convergence is strong enough to converge the L^1 sizes.
Key Insight: The reverse triangle inequality converts function convergence into norm convergence. □

Proposition 5.5.24. *Finite Integral Localizes to Finite Measure*
If $f : X \to [0, +\infty]$ is measurable, $\int_X f\,d\mu < +\infty$, then $\forall \varepsilon > 0\ \exists\ E \in \mathcal{M}:\quad \mu(E) < +\infty$; $\int_{X\setminus E} f\,d\mu < \varepsilon$.

Proof. Set $E_n := \{1/n < f \leq n\}$. Then $\mu(E_n) \leq n \int_X f\,d\mu < +\infty$, and $f\mathbf{1}_{E_n} \uparrow f$ a.e.
Hence $\int_{E_n} f\,d\mu \uparrow \int_X f\,d\mu$ by monotone convergence, so $\int_{X\setminus E_n} f\,d\mu \downarrow 0$; choose $E = E_n$.
Structural Commentary: Finite integral mass can be captured on a finite-measure part of the space.
Key Insight: Keep the levels where f is bounded above and away from zero. □

Proposition 5.5.25. *Setwise Convergence from Total Mass Convergence*
Let $f_n, f : X \to [0, +\infty]$ be measurable, $f_n \to f$ pointwise, and $\int_X f\,d\mu = \lim_{n\to\infty} \int_X f_n\,d\mu < +\infty$.
Then $\forall E \in \mathcal{M}$, $\int_E f\,d\mu = \lim_{n\to\infty} \int_E f_n\,d\mu$.

Proof. Fix $E \in \mathcal{M}$ and set $E^c := X \setminus E$. By Fatou, $\int_E f\,d\mu \leq \liminf_n \int_E f_n\,d\mu$ and $\int_{E^c} f\,d\mu \leq \liminf_n \int_{E^c} f_n\,d\mu$.
Since $\int_X f_n\,d\mu = \int_E f_n\,d\mu + \int_{E^c} f_n\,d\mu \to \int_X f\,d\mu < +\infty$, $\limsup_n \int_E f_n\,d\mu \leq \int_X f\,d\mu - \int_{E^c} f\,d\mu = \int_E f\,d\mu$.
Thus $\liminf_n \int_E f_n\,d\mu = \limsup_n \int_E f_n\,d\mu = \int_E f\,d\mu$.
Structural Commentary: Total mass plus Fatou on complements fixes each part.
Key Insight: A lower bound on the complement is an upper bound on the chosen set. □

Definition 5.5.26. *Function Truncates*
For $f : X \to \mathbb{R}$ and $n \in \mathbb{N}$, define the *nth truncate*

$$T_n f(x) := (-n) \vee (f(x) \wedge n) = \begin{cases} -n, & f(x) < -n, \\ f(x), & -n \leq f(x) \leq n, \\ n, & f(x) > n. \end{cases}$$

Proposition 5.5.27. *Truncates Characterize Integrability*
Let $f : X \to \mathbb{R}$ be measurable and $f_n := T_n f$. Then $f \in \mathscr{L}^1(\mu) \iff \sup_n \|f_n\|_1 < +\infty$.
Moreover, if $f \in \mathscr{L}^1(\mu)$, then $\int_X f_n\,d\mu \to \int_X f\,d\mu$.

Proof. The map $T_n : \mathbb{R} \to [-n, n]$ is continuous, so f_n is measurable; also $f_n \to f$, $|f_n| \le |f|$, and $|f_n| \uparrow |f|$.
If $f \in \mathscr{L}^1(\mu)$, dominated convergence gives $\int_X f_n \, d\mu \to \int_X f \, d\mu$ and $\sup_n \|f_n\|_1 \le \|f\|_1$.
Conversely, monotone convergence gives $\|f\|_1 = \lim_{n\to\infty} \|f_n\|_1 = \sup_n \|f_n\|_1 < +\infty$, so $f \in \mathscr{L}^1(\mu)$.

Structural Commentary: Truncation reduces L^1 questions to bounded approximants.
Key Insight: Dominated convergence gives limits; monotone convergence reads integrability from $|f_n| \uparrow |f|$. □

Example 5.5.28. *Why the Finite Total Mass Assumption Matters*
Let $X = \mathbb{N} \times \{0, 1\}$ with counting measure, set $E := \mathbb{N} \times \{1\}$.
Define $f := \mathbf{1}_{\mathbb{N}\times\{0\}}$; $f_n := \mathbf{1}_{\{1,\dots,n\}\times\{0\}} + n\mathbf{1}_{\{n\}\times\{1\}}$.
Then $f_n \to f$ pointwise and $\int_X f \, d\mu = \lim_n \int_X f_n \, d\mu = +\infty$; $\int_E f \, d\mu = 0$, $\int_E f_n \, d\mu = n$.
Thus Proposition 5.5.25 can fail when the common total mass is infinite.

5.6 Product Measures and Fubini Theory

Product measure builds measurable rectangles and reads higher-dimensional integrals through slices. The abstract construction comes first; the Euclidean statements below are its geometric specialization.

Definition 5.6.1. *Product Measurable Spaces and Sections*
Let $(X, \mathcal{M})$, $(Y, \mathcal{N})$ be measurable spaces and $Z := X \times Y$.

(i) A *measurable rectangle* is $A \times B$ with $A \in \mathcal{M}$ and $B \in \mathcal{N}$; set $\mathcal{Z}_0 := \{\bigcup_{j=1}^n (A_j \times B_j) : A_j \in \mathcal{M},\ B_j \in \mathcal{N}\}$ and $\mathcal{M} \otimes \mathcal{N} := \sigma\{A \times B : A \in \mathcal{M},\ B \in \mathcal{N}\}$.
(ii) For $E \subseteq Z$, the *sections* are $E_x := \{y \in Y : (x, y) \in E\}$ and $E^y := \{x \in X : (x, y) \in E\}$.
(iii) For $F : Z \to \overline{\mathbb{R}}$, the *function sections* are $F_x(y) := F(x, y)$ and $F^y(x) := F(x, y)$, where defined.

Lemma 5.6.2. *Finite Rectangle Algebra*

(i) If $R_j = A_j \times B_j$, then $R_1 \cap R_2 = (A_1 \cap A_2) \times (B_1 \cap B_2)$ and $R_1 \setminus R_2 = \big[(A_1 \cap A_2) \times (B_1 \setminus B_2)\big] \sqcup \big[(A_1 \setminus A_2) \times B_1\big]$.
(ii) $\mathcal{Z}_0$ is an algebra on Z.
(iii) Every $E \in \mathcal{Z}_0$ admits $E = \bigsqcup_{j=1}^n (A_j \times B_j)$ with $A_j \in \mathcal{M}$ and $B_j \in \mathcal{N}$.

Proof.

(i) For $(x, y) \in Z$, $(x, y) \in R_1 \cap R_2 \iff x \in A_1 \cap A_2$ and $y \in B_1 \cap B_2$. Also, $(x, y) \in R_1 \setminus R_2$ iff $x \in A_1$, $y \in B_1$, and either $x \notin A_2$ or $y \notin B_2$; splitting by $x \in A_2$ gives the disjoint union in (i).
(ii) If $E = \bigcup_j R_j$ and $F = \bigcup_k S_k$, then $E \cap F = \bigcup_{j,k} (R_j \cap S_k) \in \mathcal{Z}_0$ by (i).
Repeatedly applying (i) gives $R_j \setminus F \in \mathcal{Z}_0$, hence $E \setminus F = \bigcup_j (R_j \setminus F) \in \mathcal{Z}_0$. Since $Z = X \times Y \in \mathcal{Z}_0$, complements are differences from Z.
(iii) If $E = \bigcup_{j=1}^n R_j$, set $S_j := R_j \setminus \bigcup_{i<j} R_i$. Then $S_j \in \mathcal{Z}_0$ by (ii), the S_j are disjoint, and $E = \bigsqcup_j S_j$. Writing each S_j as a finite union of rectangles and relabeling gives the required finite disjoint representation.

Structural Commentary: Product rectangles behave like elementary sets for the product space.
Key Insight: Before extending to a σ-algebra, make all finite overlaps into disjoint rectangle pieces. □

Theorem 5.6.3. *Product Measure Theorem*
For measure spaces $(X, \mathcal{M}, \mu)$ and $(Y, \mathcal{N}, \nu)$, there exists a measure π on $\mathcal{M} \otimes \mathcal{N}$ with $\pi(A \times B) = \mu(A)\nu(B)$ for $A \in \mathcal{M}$ and $B \in \mathcal{N}$, where $0 \cdot (+\infty) = 0$. If μ, ν are σ-finite, π is unique; write $\pi = \mu \otimes \nu$.

Proof. For $E = \bigsqcup_{j=1}^n (A_j \times B_j) \in \mathcal{Z}_0$, set $\pi_0(E) := \sum_{j=1}^n \mu(A_j)\nu(B_j)$.
If $A \times B = \bigsqcup_{j\ge1} (A_j \times B_j)$, then $\mathbf{1}_{A\times B} = \sum_{j\ge1} \mathbf{1}_{A_j \times B_j}$, so monotone convergence in y, then x, gives $\mu(A)\nu(B) = \sum_{j\ge1} \mu(A_j)\nu(B_j)$. Refining two finite disjoint rectangle decompositions by intersections gives equal π_0-values; hence π_0 is well defined. If $E = \bigsqcup_k E_k \in \mathcal{Z}_0$ with $E_k \in \mathcal{Z}_0$, write $E = \bigsqcup_i R_i$ and $E_k = \bigsqcup_\ell R_{k\ell}$. Since $R_i = \bigsqcup_{k,\ell} (R_i \cap R_{k\ell})$, the rectangle comparison yields $\pi_0(E) = \sum_k \pi_0(E_k)$; hence π_0 is a premeasure.
Extend π_0 to $\sigma(\mathcal{Z}_0) = \mathcal{M} \otimes \mathcal{N}$. If μ, ν are σ-finite, then $\mathcal{Z}_0$ is π_0-σ-finite, so the extension is unique.
Structural Commentary: Rectangles fix the product measure.
Key Insight: Monotone convergence makes rectangle mass respect countable disjoint refinements. □

Lemma 5.6.4. *Measurable Sections*

(i) $E \in \mathcal{M} \otimes \mathcal{N}$ implies $E_x \in \mathcal{N}$ for all $x \in X$ and $E^y \in \mathcal{M}$ for all $y \in Y$.
(ii) If $F : Z \to \overline{\mathbb{R}}$ is $\mathcal{M} \otimes \mathcal{N}$-measurable, then F_x is $\mathcal{N}$-measurable for all x and F^y is $\mathcal{M}$-measurable for all y.

Proof.

(i) Fix $x \in X$ and set $\mathcal{C}_x := \{E \subseteq Z : E_x \in \mathcal{N}\}$. Since $(E^c)_x = (E_x)^c$ and $(\bigcup_n E_n)_x = \bigcup_n (E_n)_x$, $\mathcal{C}_x$ is a σ-algebra. Also $(A \times B)_x = B$ if $x \in A$ and $\emptyset$ otherwise; hence $\mathcal{M} \otimes \mathcal{N} \subseteq \mathcal{C}_x$. The E^y proof is identical.

(ii) For $a \in \mathbb{R}$, set $H_a := \{(u,v) \in Z : F(u,v) > a\} \in \mathcal{M} \otimes \mathcal{N}$. By (i), $(H_a)_x \in \mathcal{N}$ and $(H_a)^y \in \mathcal{M}$. Since $(H_a)_x = \{y : F_x(y) > a\}$ and $(H_a)^y = \{x : F^y(x) > a\}$, all threshold sets are measurable.

Structural Commentary: Product measurability is stable under freezing one coordinate.
Key Insight: Section operations preserve the set algebra used to generate the product σ-algebra. □

Proposition 5.6.5. *Section Measures Integrate to Product Measure*
Let $(X, \mathcal{M}, \mu)$ and $(Y, \mathcal{N}, \nu)$ be σ-finite, $\pi := \mu \otimes \nu$, and $E \in \mathcal{M} \otimes \mathcal{N}$. Set $f_E(x) := \nu(E_x)$ and $g_E(y) := \mu(E^y)$. Then f_E, g_E are measurable and $\int_X f_E \, d\mu = \pi(E) = \int_Y g_E \, d\nu$.

Proof. Assume $\mu(X), \nu(Y) < +\infty$ and let $\mathcal{C}$ be the class of E satisfying the claim.
For $E = A \times B$, $f_E = \nu(B)\mathbf{1}_A$ and $g_E = \mu(A)\mathbf{1}_B$, so $A \times B \in \mathcal{C}$; finite additivity gives $\mathcal{Z}_0 \subseteq \mathcal{C}$.
If $E_n \uparrow E$ in $\mathcal{C}$, then $(E_n)_x \uparrow E_x$ and $(E_n)^y \uparrow E^y$; MCT gives $E \in \mathcal{C}$.
If $E_n \downarrow E$, finite total mass gives downward continuity and DCT, hence $E \in \mathcal{C}$.
Thus $\mathcal{C}$ is a monotone class containing $\mathcal{Z}_0$, so $\mathcal{C} = \mathcal{M} \otimes \mathcal{N}$. For σ-finite μ, ν, choose $X_n \uparrow X$, $Y_n \uparrow Y$ with $\mu(X_n), \nu(Y_n) < +\infty$, apply the finite case to $E_n := E \cap (X_n \times Y_n)$, then let $n \to \infty$ by MCT.
Structural Commentary: Cavalieri's principle is already true for abstract σ-finite products.
Key Insight: Prove rectangles, close under monotone limits, then exhaust by finite-measure rectangles. □

Theorem 5.6.6. *Tonelli's Theorem for σ-Finite Products*
Let $(X, \mathcal{M}, \mu)$ and $(Y, \mathcal{N}, \nu)$ be σ-finite, $\pi := \mu \otimes \nu$, and $F : Z \to [0, +\infty]$ be measurable. Set $I_F(x) := \int_Y F_x \, d\nu$ and $J_F(y) := \int_X F^y \, d\mu$. Then I_F, J_F are measurable and $\int_X I_F \, d\mu = \int_Z F \, d\pi = \int_Y J_F \, d\nu$.

Proof. For $F = \mathbf{1}_E$, the claim is Proposition 5.6.5. By finite linearity, it holds for simple $\Phi = \sum_{j=1}^n a_j \mathbf{1}_{E_j} \geq 0$.
Choose simple $0 \leq \Phi_n \uparrow F$. Then $(\Phi_n)_x \uparrow F_x$, $(\Phi_n)^y \uparrow F^y$, $I_{\Phi_n} \uparrow I_F$, and $J_{\Phi_n} \uparrow J_F$.
MCT in X, Y, Z passes the identity from Φ_n to F.
Structural Commentary: Tonelli extends sections from indicators to $F \geq 0$.
Key Insight: Nonnegativity lets all three integrals pass through monotone approximation together. □

Theorem 5.6.7. *Fubini's Theorem for σ-Finite Products*
Let $(X, \mathcal{M}, \mu)$ and $(Y, \mathcal{N}, \nu)$ be σ-finite, $\pi := \mu \otimes \nu$, and $F \in \mathscr{L}^1(\pi)$.

(i) $F_x \in \mathscr{L}^1(\nu)$ for μ-a.e. x and $F^y \in \mathscr{L}^1(\mu)$ for ν-a.e. y.
(ii) With $f(x) := \int_Y F_x \, d\nu$ and $g(y) := \int_X F^y \, d\mu$ on those full-measure sets, and 0 elsewhere, $f \in \mathscr{L}^1(\mu)$, $g \in \mathscr{L}^1(\nu)$, and $\int_X f \, d\mu = \int_Z F \, d\pi = \int_Y g \, d\nu$.

Proof.

(i) Tonelli on $|F|$ gives $\int_X (\int_Y |F_x| \, d\nu) d\mu = \int_Z |F| \, d\pi = \int_Y (\int_X |F^y| \, d\mu) d\nu < +\infty$, so the inner integrals are finite a.e.
(ii) Tonelli on F^+ and F^- gives finite identities; subtract them to obtain $\int_X f \, d\mu = \int_Z F \, d\pi = \int_Y g \, d\nu$. Moreover $|f| \leq \int_Y |F_x| \, d\nu$ and $|g| \leq \int_X |F^y| \, d\mu$, so (i) gives $f, g \in \mathscr{L}^1$.

Structural Commentary: Fubini is Tonelli plus subtraction.
Key Insight: $\int |F| < +\infty$ makes the two nonnegative Tonelli identities finite enough to subtract. □

Proposition 5.6.8. *Product Borel and Measurable Building Blocks*
The following hold.

(i) $\mathcal{B}(\mathbb{R}) \otimes \mathcal{B}(\mathbb{R}) = \mathcal{B}(\mathbb{R}^2)$.
(ii) If $f : X \to \mathbb{R}$ is $\mathcal{M}$-measurable and $g : Y \to \mathbb{R}$ is $\mathcal{N}$-measurable, then $(x, y) \mapsto f(x)g(y)$ is $\mathcal{M} \otimes \mathcal{N}$-measurable.
(iii) If $u : \mathbb{R} \to \mathbb{R}$ is Borel measurable, then $(x, y) \mapsto u(x - y)$ is $\mathcal{B}(\mathbb{R}) \otimes \mathcal{B}(\mathbb{R})$-measurable.

Proof.

(i) Open rectangles give $\mathcal{B}(\mathbb{R}) \otimes \mathcal{B}(\mathbb{R}) \subseteq \mathcal{B}(\mathbb{R}^2)$.
Conversely, every open $O \subseteq \mathbb{R}^2$ is a countable union of rational rectangles, each in $\mathcal{B}(\mathbb{R}) \otimes \mathcal{B}(\mathbb{R})$.
(ii) For $\tilde{f}(x, y) := f(x)$ and $\tilde{g}(x, y) := g(y)$, $\{\tilde{f} > a\} = f^{-1}((a, +\infty)) \times Y$ and $\{\tilde{g} > a\} = X \times g^{-1}((a, +\infty))$.
Thus $\tilde{f}, \tilde{g}$ are measurable, so $\tilde{f}\tilde{g}$ is measurable.
(iii) The map $s(x, y) := x - y$ is continuous on $\mathbb{R}^2$, hence measurable by (i); therefore $u \circ s$ is measurable.

Structural Commentary: Product measurability supports two-variable formulas.
Key Insight: Check coordinate pullbacks on rectangles, then use ordinary measurable-function closure. □

Proposition 5.6.9. *Counting Factor Products*
Let $X = \mathbb{N}$, $\mathcal{M} = 2^{\mathbb{N}}$, μ be counting measure, and $(Y, \mathcal{N}, \nu)$ be a measure space.

(i) For $E \subseteq \mathbb{N} \times Y$, $E \in 2^{\mathbb{N}} \otimes \mathcal{N} \iff E_n \in \mathcal{N}$ for all n, and then $E = \bigsqcup_{n \geq 1} (\{n\} \times E_n)$.

(ii) *The unique product measure is* $\pi = \mu \otimes \nu$ *with* $\pi(E) = \sum_{n\geq 1} \nu(E_n)$ *for* $E \in 2^{\mathbb{N}} \otimes \mathcal{N}$.
(iii) *For* $F : \mathbb{N} \times Y \to \mathbb{R}$, F *is product-measurable* $\iff$ F_n *is* $\mathcal{N}$*-measurable for all* n.
(iv) *For product-measurable* F, $F \in \mathscr{L}^1(\pi) \iff \sum_{n\geq 1} \int_Y |F_n|\, d\nu < +\infty$;
then $\int_{\mathbb{N}\times Y} F\, d\pi = \sum_{n\geq 1} \int_Y F_n\, d\nu = \int_Y \sum_{n\geq 1} F_n\, d\nu$, *with* $\sum_n F_n$ *absolutely convergent* ν*-a.e.*

Proof.

(*i*) ($\Rightarrow$) If $E \in 2^{\mathbb{N}} \otimes \mathcal{N}$, Lemma 5.6.4 gives $E_n \in \mathcal{N}$ for all n.
(*i*) ($\Leftarrow$) If $E_n \in \mathcal{N}$ for all n, then each $\{n\} \times E_n \in 2^{\mathbb{N}} \otimes \mathcal{N}$ and $E = \bigsqcup_{n\geq 1}(\{n\} \times E_n)$, hence $E \in 2^{\mathbb{N}} \otimes \mathcal{N}$.
(*ii*) Define $\pi(E) := \sum_{n\geq 1} \nu(E_n)$.
If $(E^k)_k$ are disjoint, then $(\bigsqcup_k E^k)_n = \bigsqcup_k E^k_n$, so $\pi(\bigsqcup_k E^k) = \sum_n \sum_k \nu(E^k_n) = \sum_k \pi(E^k)$.
Also $\pi(\{n\} \times A) = \nu(A)$, and (i) forces uniqueness.
(*iii*) ($\Rightarrow$) If F is product-measurable, Lemma 5.6.4 gives each F_n measurable.
(*iii*) ($\Leftarrow$) If each F_n is measurable, then for every $a \in \mathbb{R}$, $H_a := \{F > a\}$ satisfies $(H_a)_n = \{F_n > a\} \in \mathcal{N}$,
By (i), $H_a \in 2^{\mathbb{N}} \otimes \mathcal{N}$, so F is product-measurable.
(*iv*) For $G \geq 0$, (ii) gives the identity for $G = \mathbf{1}_E$; linearity and MCT give $\int_{\mathbb{N}\times Y} G\, d\pi = \sum_{n\geq 1} \int_Y G_n\, d\nu$.
(*iv*) Apply the preceding identity to $G = |F|$: $F \in \mathscr{L}^1(\pi) \iff \sum_{n\geq 1} \int_Y |F_n|\, d\nu < +\infty$.
If finite, then $\int_Y \sum_n |F_n|\, d\nu < +\infty$, so $\sum_n F_n$ converges absolutely a.e.; applying the nonnegative identity to F^+ and F^- and subtracting gives the integral formulas.

Structural Commentary: A product with counting measure is a measurable sequence of fibers.
Key Insight: Counting atoms turn products into sums. □

Example 5.6.10. *A Diagonal Failure Without* σ*-Finiteness*
Let $X = Y = [0,1]$, $\mathcal{M} = \mathcal{N} = \mathcal{B}([0,1])$, $\mu = m$, $\nu = \#$, and $D := \{(x,y) \in [0,1]^2 : x = y\}$.
Then $D \in \mathcal{M} \otimes \mathcal{N}$, $D_x = \{x\}$, and $D^y = \{y\}$, so $\nu(D_x) = 1$ and $\mu(D^y) = 0$.
Hence $\int_X \nu(D_x)\, d\mu = 1$ but $\int_Y \mu(D^y)\, d\nu = 0$, so Tonelli/Fubini can fail without σ-finite factors.

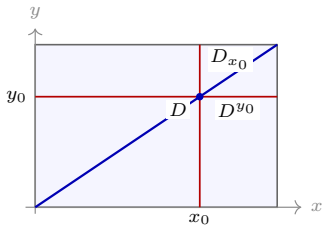


Figure 196: Diagonal sections: $\nu(D_x) = 1$ but $\mu(D^y) = 0$.

Example 5.6.11. *Double Series as Product Integrals*
With counting measure on $\mathbb{N}^2$, Tonelli gives, for $a_{mn} \geq 0$,
$\sum_m \sum_n a_{mn} = \sum_{(m,n)\in\mathbb{N}^2} a_{mn} = \sum_n \sum_m a_{mn} \in [0, +\infty]$.
Fubini needs integrability: set $a_{nn} = 1$, $a_{n,n+1} = -1$, and $a_{mn} = 0$ otherwise.
Then $\sum_m \sum_n a_{mn} = 0$, while $\sum_n \sum_m a_{mn} = 1$, and $\sum_{m,n} |a_{mn}| = +\infty$.

Proposition 5.6.12. *Separated Products*
Let $(X, \mathcal{M}, \mu)$ *and* $(Y, \mathcal{N}, \nu)$ *be* σ*-finite. If* $f \in \mathscr{L}^1(\mu)$, $g \in \mathscr{L}^1(\nu)$, *and* $H(x,y) := f(x)g(y)$, *then* $H \in \mathscr{L}^1(\mu\otimes\nu)$ *and* $\int_{X\times Y} H\, d(\mu \otimes \nu) = (\int_X f\, d\mu)(\int_Y g\, d\nu)$.

Proof. Measurability follows from Proposition 5.6.8.
Tonelli gives $\int_{X\times Y} |H|\, d(\mu \otimes \nu) = \int_X |f(x)|\, (\int_Y |g|\, d\nu) d\mu = (\int_X |f|\, d\mu)(\int_Y |g|\, d\nu) < +\infty$, so $H \in \mathscr{L}^1(\mu \otimes \nu)$.
Fubini applied to H gives the displayed identity.
Structural Commentary: Product-measure integration separates products of one-variable integrable functions.
Key Insight: Absolute integrability factors by Tonelli; signed equality follows by Fubini. □

Corollary 5.6.13. *Equal Section Measures*
Let $(X, \mathcal{M}, \mu)$ *and* $(Y, \mathcal{N}, \nu)$ *be* σ*-finite.*
If $E, F \in \mathcal{M} \otimes \mathcal{N}$ *and* $\nu(E_x) = \nu(F_x)$ *for all* $x \in X$, *then* $(\mu \otimes \nu)(E) = (\mu \otimes \nu)(F)$.

Proof. Set $f_E(x) := \nu(E_x)$ and $f_F(x) := \nu(F_x)$.
Since $f_E = f_F$, Proposition 5.6.5 gives $(\mu \otimes \nu)(E) = \int_X f_E\, d\mu = \int_X f_F\, d\mu = (\mu \otimes \nu)(F)$.
Structural Commentary: Product mass is recovered from vertical section masses.
Key Insight: Equal fibers give equal product mass. □

Definition 5.6.14. *Convolution*
For Borel $f, g : \mathbb{R} \to \mathbb{R}$, if $y \mapsto f(x-y)g(y)$ is integrable for a.e. x,
define the *convolution* $(f * g)(x) := \int_{\mathbb{R}} f(x-y)g(y)\, dm(y)$ on that full-measure set and 0 elsewhere.

Proposition 5.6.15. *Convolution of L^1 Functions*
If Borel $f, g \in \mathscr{L}^1(m)$, then $f * g$ is a.e. defined, $f * g \in \mathscr{L}^1(m)$, and $\int_{\mathbb{R}} |f * g|\, dm \le (\int_{\mathbb{R}} |f|\, dm)(\int_{\mathbb{R}} |g|\, dm)$.

Proof. Set $\Phi(x, y) := f(x-y)g(y)$. By Proposition 5.6.8, Φ is product-measurable. For fixed y, translation invariance gives $\int_{\mathbb{R}} |f(x-y)|\, dm(x) = \int_{\mathbb{R}} |f|\, dm$. Hence Tonelli yields $\int_{\mathbb{R}^2} |\Phi|\, d(m \otimes m) = (\int_{\mathbb{R}} |f|\, dm)(\int_{\mathbb{R}} |g|\, dm) < +\infty$. Fubini gives the a.e. existence of $f * g$, and $\int_{\mathbb{R}} |f * g|\, dm \le \int_{\mathbb{R}^2} |\Phi|\, d(m \otimes m)$ gives the bound.
Structural Commentary: Convolution is Fubini applied to the difference map $(x, y) \mapsto x - y$.
Key Insight: Translation invariance makes every shifted copy of $|f|$ have the same integral. □

Example 5.6.16. *Nonunique Products Without σ-Finiteness*
Let $X = \mathbb{R}$, $\mathcal{M} = 2^{\mathbb{R}}$, and $\mu(A) = 0$ if A is countable, $\mu(A) = +\infty$ otherwise.
Define $\lambda_i : \mathcal{M} \otimes \mathcal{M} \to \{0, +\infty\}$ by $\lambda_i(E) = 0$ exactly on the following null ideals:
$\mathcal{I}_1 = \{E : E = G \cup H,\ \mathrm{pr}_x\, G \text{ countable},\ \mathrm{pr}_y\, H \text{ countable}\}$
and $\mathcal{I}_2 = \{E : E = G \cup H \cup K,\ \mathrm{pr}_x\, G \text{ countable},\ \mathrm{pr}_y\, H \text{ countable},\ \mathrm{pr}_{y=x}\, K \text{ countable}\}$.
Both ideals are closed under countable unions, hence λ_i are measures.
For every rectangle $A \times B$, both give $\mu(A)\mu(B)$: countable sides are null, and uncountable-by-uncountable rectangles are not in either ideal.
For $E = \{(x, y) : x + y = 0\}$, $\mathrm{pr}_{y=x}\, E$ is a point, so $\lambda_2(E) = 0$ but $\lambda_1(E) = +\infty$.
Structural Commentary: Without σ-finiteness, rectangle values need not determine a product measure.
Key Insight: Extra null lines preserve rectangles but change non-rectangles.

Definition 5.6.17. *Products and Slices*
For $A \subseteq \mathbb{R}^n$, $B \subseteq \mathbb{R}^k$, set $A \times B \subseteq \mathbb{R}^{n+k}$.
For $E \subseteq \mathbb{R}^n \times \mathbb{R}^k$ define $E_x := \{y \in \mathbb{R}^k : (x, y) \in E\}$ as *x-slice*, and $E^y := \{x \in \mathbb{R}^n : (x, y) \in E\}$ as *y-slice*.
For $f : E \to \mathbb{R}$, write $f_x(y) = f^y(x) := f(x, y)$ where defined.

Theorem 5.6.18. *Measurable Products*
If $A \subseteq \mathbb{R}^n$, $B \subseteq \mathbb{R}^k$ Lebesgue measurable, then $A \times B$ measurable, $m(A \times B) = m(A)m(B)$; $0 \cdot \infty = \infty \cdot 0 = 0$.

Proof. For boxes I, J, $m(I \times J) = m(I)m(J)$; disjoint box decompositions extend this to open sets.
First assume A, B are bounded with finite measure.
By regularity choose closed $F_A \subseteq A$, $F_B \subseteq B$ and open $G_A \supseteq A$, $G_B \supseteq B$ such that $m(G_A \setminus F_A)$, $m(G_B \setminus F_B) < \varepsilon$.
Then $F_A \times F_B \subseteq A \times B \subseteq G_A \times G_B$, and the open/closed product formula squeezes $m(A \times B)$ to $m(A)m(B)$.
For general A, B, set $A_N = A \cap [-N, N]^n$ and $B_N = B \cap [-N, N]^k$. Since $A_N \times B_N \uparrow A \times B$, continuity from below gives $m(A \times B) = \lim_N m(A_N \times B_N) = \lim_N m(A_N)m(B_N) = m(A)m(B)$ with $0 \cdot \infty = \infty \cdot 0 = 0$.
Structural Commentary: Product measure comes from boxes and regularity.
Key Insight: Regularity turns rough measurable products into box products plus zero-set errors. □

Theorem 5.6.19. *Zero Slices*
If $E \subseteq \mathbb{R}^n \times \mathbb{R}^k$ is Lebesgue measurable, then E is null iff E_x is null for a.e. x.

Proof.

(⇒) If $m(E) = 0$, cover E by an open W with small measure.
For $X_\alpha(W) = \{x : m(W_x) > \alpha\}$, $m(W) \ge \alpha m(X_\alpha(W))$; letting $m(W) \to 0$ gives $m(E_x) = 0$ for a.e. x.

(⇐) Remove the exceptional null base so every slice is null. For compact $K \subseteq E$, cover each K_x by a small-measure slice neighborhood; compactness gives finitely many product neighborhoods covering K, whose total measure is arbitrarily small by Theorem 5.6.18. Thus $m(K) = 0$, hence $m(E) = 0$ by regularity.

Structural Commentary: Nullity can be checked fiberwise after removing an exceptional base set.
Key Insight: Product measure turns small slice covers into small total measure. □

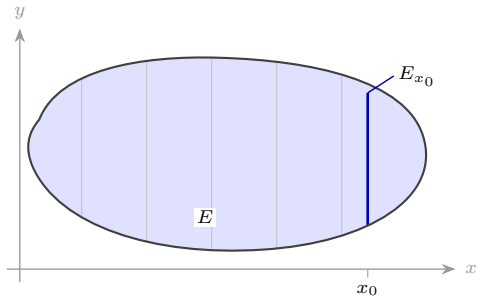


Figure 197: A planar set can be read by vertical slices; Fubini integrates the slice data over the base variable.

Theorem 5.6.20. *Cavalieri's Principle for Lebesgue Measure*
If $E \in \mathcal{L}^{n+k}$, then $E_x \in \mathcal{L}^k$ for a.e. x, and the a.e.-defined function $f_E(x) := m_k(E_x)$ has a measurable representative with $m_{n+k}(E) = \int_{\mathbb{R}^n} f_E \, dm_n$. The analogous statement holds for E^y.

Proof. Let $P(E)$ denote the displayed assertion. For $E = A \times B$, $f_E = m_k(B)\mathbf{1}_A$, so $P(E)$ is Theorem 5.6.18. Fix a bounded box $Q = I \times J$. The class $\{E \subseteq Q : P(E)\}$ contains finite unions of rectangles and is monotone: if $E_j \uparrow E$, then $f_{E_j} \uparrow f_E$ and monotone convergence applies; if $E_j \downarrow E$, then $f_{E_j} \downarrow f_E \le m_k(J)\mathbf{1}_I$ and dominated convergence applies. Hence the monotone class theorem gives $P(E)$ for every Borel $E \subseteq Q$. Exhaustion by $Q_N = [-N, N]^{n+k}$ and monotone convergence give all Borel E.
For Lebesgue E, choose Borel B with $m_{n+k}(E \triangle B) = 0$ (Remark 5.1.72). Then $m_k((E \triangle B)_x) = 0$ for a.e. x (Theorem 5.6.19), so $f_E = f_B$ a.e. and $m_{n+k}(E) = m_{n+k}(B)$. The y-slice statement is identical.
Structural Commentary: Slicing turns a measure in $\mathbb{R}^{n+k}$ into a measurable family of measures in $\mathbb{R}^k$.
Key Insight: Prove boxes, close under monotone limits, then ignore null symmetric differences. □

Perspective 5.6.21. *Fubini Proof Ladder*
Fubini ladder: $\mathbf{1}_E \to$ simple $\Phi \ge 0 \to f \ge 0$ by MCT, then $f = f^+ - f^-$ when $\int |f| < +\infty$.
Theorems 5.6.6 and 5.6.7 are the abstract form.
Key Insight: Fubini is a slice theorem, proof runs through indicators, simple functions, monotone convergence.

Proposition 5.6.22. *Undergraphs and Level-Set Measurability*
Let $f : \mathbb{R}^n \to [0, \infty]$, set $F_a := \{x : f(x) > a\}$, and write $\mathcal{U}f = \{(x, t) \in \mathbb{R}^n \times [0, \infty) : 0 < t < f(x)\}$. Then f is Lebesgue measurable iff $F_q \in \mathcal{L}^n$ $\forall q \in \mathbb{Q}_{>0}$. In that case, $\mathcal{U}f \in \mathcal{L}^{n+1}$ and $(\mathcal{U}f)^t = F_t$, $t > 0$.

Proof. For $a \ge 0$, $F_a = \bigcup_{\substack{q \in \mathbb{Q} \\ q > a}} F_q$; for $a < 0$, $F_a = \mathbb{R}^n$.
Hence rational positive thresholds are equivalent to measurability.
If they are measurable, then $\mathcal{U}f = \bigcup_{\substack{q,r \in \mathbb{Q} \\ 0 < q < r}} F_r \times (q, r) \in \mathcal{L}^{n+1}$. Finally, $(\mathcal{U}f)^t = \{x : 0 < t < f(x)\} = F_t$ for $t > 0$.
Structural Commentary: A nonnegative function may be studied as the region lying beneath its graph.
Key Insight: Horizontal slices of an undergraph are exactly strict superlevel sets. □

Theorem 5.6.23. *Fubini–Tonelli for Nonnegative Functions*
Let $f : \mathbb{R}^n \times \mathbb{R}^k \to [0, \infty]$ be Lebesgue measurable. Then f_x is measurable for a.e. x; with $F(x) := \int_{\mathbb{R}^k} f_x \, dm_k$ there and arbitrary values elsewhere, F is measurable and $\int_{\mathbb{R}^{n+k}} f \, dm_{n+k} = \int_{\mathbb{R}^n} F \, dm_n$.
The same holds for $G(y) := \int_{\mathbb{R}^n} f^y \, dm_n$, with $\int_{\mathbb{R}^{n+k}} f \, dm_{n+k} = \int_{\mathbb{R}^k} G \, dm_k$.

Proof. Put $U := \mathcal{U}f \subseteq \mathbb{R}^{n+k} \times [0, \infty)$. Proposition 5.6.22 gives $U \in \mathcal{L}^{n+k+1}$. By Cavalieri, Theorem 5.6.20, U_x is measurable for a.e. x, $H(x) := m_{k+1}(U_x)$ is measurable, and $m_{n+k+1}(U) = \int_{\mathbb{R}^n} H \, dm_n$.
For those x, $U_x = \mathcal{U}(f_x)$, so f_x is measurable and $H(x) = \int_{\mathbb{R}^k} f_x \, dm_k$.
Also $m_{n+k+1}(U) = \int_{\mathbb{R}^{n+k}} f \, dm_{n+k}$ by the undergraph formula. The y-formula is identical.
Structural Commentary: Tonelli's theorem is Cavalieri's principle applied one dimension higher.
Key Insight: $(\mathcal{U}f)_x = \mathcal{U}(f_x)$. □

Theorem 5.6.24. *Fubini for L^1 Functions*
If $f \in L^1(\mathbb{R}^2)$, then $f(x, \cdot) \in L^1(\mathbb{R})$ for a.e. x and $f(\cdot, y) \in L^1(\mathbb{R})$ for a.e. y. Defining $F(x) := \int_{\mathbb{R}} f(x, y) \, dy$; $G(y) := \int_{\mathbb{R}} f(x, y) \, dx$ on those sets and as 0 elsewhere gives $F, G \in L^1(\mathbb{R})$ and $\int_{\mathbb{R}} F(x) \, dx = \int_{\mathbb{R}^2} f = \int_{\mathbb{R}} G(y) \, dy$

Proof. Let $E_x = \{y : (x, y) \in E\}$ and $E^y = \{x : (x, y) \in E\}$. For measurable $E \subseteq \mathbb{R}^2$, the box/regularity argument and Theorems 5.6.18–5.6.19 give $\int_{\mathbb{R}} m(E_x) \, dx = m(E) = \int_{\mathbb{R}} m(E^y) \, dy$.
Hence Fubini holds for $\mathbf{1}_E$, and by linearity for every nonnegative simple function.
Let $0 \le h$ be measurable and choose $0 \le s_n \uparrow h$ by Theorem 5.3.12.

Monotone convergence gives $\int_{\mathbb{R}}\big(\int_{\mathbb{R}} h(x,y)\,dy\big)\,dx = \lim_n \int_{\mathbb{R}}\big(\int_{\mathbb{R}} s_n(x,y)\,dy\big)\,dx = \lim_n \int_{\mathbb{R}^2} s_n = \int_{\mathbb{R}^2} h$ and the same argument with x, y interchanged. Now let $f \in L^1(\mathbb{R}^2)$.
Applying the nonnegative case to $|f|$ yields $\int_{\mathbb{R}}\big(\int_{\mathbb{R}} |f(x,y)|\,dy\big)\,dx = \int_{\mathbb{R}^2} |f| < \infty$, so $f(x,\cdot) \in L^1(\mathbb{R})$ for a.e. x; similarly $f(\cdot,y) \in L^1(\mathbb{R})$ for a.e. y. With $f = f^+ - f^-$, define, up to null-set choices, $F = F^+ - F^-$, $F^{\pm}(x) = \int_{\mathbb{R}} f^{\pm}(x,y)\,dy$, and analogously $G = G^+ - G^-$. Then $\int_{\mathbb{R}} |F| \le \int_{\mathbb{R}}(F^+ + F^-) = \int_{\mathbb{R}^2} |f| < \infty$, and similarly $G \in L^1(\mathbb{R})$. Subtracting the nonnegative identities for f^+ and f^- gives $\int_{\mathbb{R}} F = \int_{\mathbb{R}^2} f = \int_{\mathbb{R}} G$.
Structural Commentary: Prove identity for $\mathbf{1}_E$, extend by linearity and monotone convergence, apply to $f^{\pm}$.
Key Insight: $\int |f| < \infty$ makes the positive and negative slice integrals finite a.e., so subtraction is legitimate. □

5.7 Radon-Nikodym and Differentiation

This subsection first decomposes finite signed set functions into honest measures, then compares measures by absolute continuity and singularity. The differentiation results later in the subsection are the local Euclidean counterpart of the Radon–Nikodym viewpoint.

Definition 5.7.1. *Positive, Negative, and Null Sets for a Charge*
Let λ be a charge on $(X, \mathcal{M})$.

(i) $\mathcal{P}_\lambda := \{P \in \mathcal{M} : \forall E \in \mathcal{M}, \lambda(E \cap P) \ge 0\}$, the *positive* sets.
(ii) $\mathcal{N}_\lambda := \{N \in \mathcal{M} : \forall E \in \mathcal{M}, \lambda(E \cap N) \le 0\}$, the *negative* sets.
(iii) $\mathcal{Z}_\lambda := \{M \in \mathcal{M} : \forall E \in \mathcal{M}, \lambda(E \cap M) = 0\}$, the *null* sets.

Proposition 5.7.2. *Positive and Negative Set Closure*
For a charge λ:

(i) $A \in \mathcal{M}$ and $A \subseteq P \in \mathcal{P}_\lambda$ imply $A \in \mathcal{P}_\lambda$; $A \subseteq N \in \mathcal{N}_\lambda$ implies $A \in \mathcal{N}_\lambda$.
(ii) $\mathcal{P}_\lambda$ and $\mathcal{N}_\lambda$ are closed under finite unions.

Proof.

(i) If $A \subseteq P \in \mathcal{P}_\lambda$, then $\lambda(E \cap A) = \lambda((E \cap A) \cap P) \ge 0$ for all $E \in \mathcal{M}$, hence $A \in \mathcal{P}_\lambda$.
The negative case reverses the inequality.
(ii) If $P_1, P_2 \in \mathcal{P}_\lambda$, then $\lambda(E \cap (P_1 \cup P_2)) = \lambda(E \cap P_1) + \lambda(E \cap P_1^c \cap P_2) \ge 0$ $\forall E \in \mathcal{M}$, since $E \cap P_1^c \cap P_2 \subseteq P_2$.
Thus $P_1 \cup P_2 \in \mathcal{P}_\lambda$; induction gives finite unions, negative case is identical with reversed inequalities.

Structural Commentary: Sign sets are hereditary and finite-union stable.
Key Insight: Test sets can always be disjointized against the first positive set. □

Theorem 5.7.3. *Hahn Decomposition Theorem*
Let λ be a charge on $(X, \mathcal{M})$. Then there exist $P, N \in \mathcal{M}$ such that

(i) $P \cap N = \emptyset$ and $X = P \sqcup N$.
(ii) $P \in \mathcal{P}_\lambda$.
(iii) $N \in \mathcal{N}_\lambda$.

Such a pair (P, N) is a Hahn decomposition of X for λ.

Proof.

(i) Put $\alpha = \sup\{\lambda(A) : A \in \mathcal{P}_\lambda\}$. If $\alpha = +\infty$, choose $B_m \in \mathcal{P}_\lambda$, $\lambda(B_m) > m$, and set $C_m = \bigcup_{j=1}^m B_j$; then $C_m \in \mathcal{P}_\lambda$, $C_m \uparrow C$, and $\lambda(C_m) \ge \lambda(B_m) > m$, contradicting continuity from below and $\lambda(C) \in \mathbb{R}$.
Choose $A_n \in \mathcal{P}_\lambda$, $\lambda(A_n) \to \alpha$, replace A_n by $\bigcup_{j=1}^n A_j$, and set $P = \bigcup_{n\ge1} A_n$, $N = X \setminus P$.
Hence $P, N \in \mathcal{M}$, $P \cap N = \emptyset$, and $X = P \sqcup N$.
(ii) For $E \in \mathcal{M}$, $E \cap A_n \uparrow E \cap P$ and $\lambda(E \cap A_n) \ge 0$, so continuity from below gives $\lambda(E \cap P) \ge 0$.
Thus $P \in \mathcal{P}_\lambda$, and $\lambda(P) = \alpha$.
(iii) Suppose $E \in \mathcal{M}$, $E \subseteq N$, and $\lambda(E) > 0$. Then $E \notin \mathcal{P}_\lambda$, else $P \cup E \in \mathcal{P}_\lambda$ and $\lambda(P \cup E) > \alpha$.
Set $R_0 = E$; while $R_{k-1} \notin \mathcal{P}_\lambda$, choose n_k minimal with some $H \in \mathcal{M}$, $H \subseteq R_{k-1}$, $\lambda(H) \le -1/n_k$, put $E_k = H$, and set $R_k = R_{k-1} \setminus E_k$. If $R_m \in \mathcal{P}_\lambda$, then $\lambda(R_m) = \lambda(E) - \sum_{k=1}^m \lambda(E_k) > 0$, so $P \cup R_m \in \mathcal{P}_\lambda$ and $\lambda(P \cup R_m) > \alpha$; hence the construction is infinite.
For $F_m = \bigcup_{k=1}^m E_k$ and $F = \bigcup_{k\ge1} E_k$, $\lambda(F_m) = \sum_{k=1}^m \lambda(E_k) \le -\sum_{k=1}^m 1/n_k$; since $F_m \uparrow F$ and $\lambda(F) \in \mathbb{R}$, $\sum_k 1/n_k < +\infty$ and $n_k \to +\infty$. If $G \in \mathcal{M}$ and $G \subseteq E \setminus F$, then $G \subseteq R_{k-1}$ for every k, so minimality gives $\lambda(G) > -1/(n_k - 1)$ for all large k, hence $\lambda(G) \ge 0$ and $E \setminus F \in \mathcal{P}_\lambda$. But $\lambda(E \setminus F) = \lambda(E) - \lambda(F) > 0$, so $P \cup (E \setminus F) \in \mathcal{P}_\lambda$ and $\lambda(P \cup (E \setminus F)) > \alpha$, contradiction. Thus $N \in \mathcal{N}_\lambda$.

Structural Commentary: A charge splits X into nonnegative-piece and nonpositive-piece regions.
Key Insight: Maximize positive mass; purify any leftover positive piece. □

Lemma 5.7.4. *Hahn Parts Are Unique up to Null Sets*
Let (P_i, N_i), $i = 1, 2$, be Hahn decompositions for a charge λ.
Then $\forall E \in \mathcal{M}$, $\lambda(E \cap P_1) = \lambda(E \cap P_2)$ and $\lambda(E \cap N_1) = \lambda(E \cap N_2)$.

Proof. Fix $E \in \mathcal{M}$. For $\{i, j\} = \{1, 2\}$, $E \cap (P_i \setminus P_j) \subseteq P_i \cap N_j$, hence $\lambda(E \cap (P_i \setminus P_j)) = 0$.
Thus $\lambda(E \cap P_i) = \lambda(E \cap P_1 \cap P_2)$ for $i = 1, 2$.
Similarly, $E \cap (N_i \setminus N_j) \subseteq N_i \cap P_j$, hence $\lambda(E \cap N_i) = \lambda(E \cap N_1 \cap N_2)$ for $i = 1, 2$.
Structural Commentary: Hahn decompositions are not unique as sets, but the measures they define are unique.
Key Insight: Positive-and-negative overlap is necessarily null. □

Definition 5.7.5. *Jordan Variations of a Charge*
Let λ be a charge, (P, N) a Hahn decomposition, and $E \in \mathcal{M}$.

(i) $\lambda^+(E) := \lambda(E \cap P)$, the *positive variation*.
(ii) $\lambda^-(E) := -\lambda(E \cap N)$, the *negative variation*.
(iii) $|\lambda|(E) := \lambda^+(E) + \lambda^-(E)$, the *total variation*.

Lemma 5.7.4 makes these independent of (P, N).

Theorem 5.7.6. *Jordan Decomposition Theorem*
For every charge λ, $\lambda = \lambda^+ - \lambda^-$. If $\lambda = \mu - \nu$ for finite measures μ, ν, then $\lambda^+ \leq \mu$ and $\lambda^- \leq \nu$.

Proof. Let (P, N) be a Hahn decomposition and $E \in \mathcal{M}$. Then $\lambda(E) = \lambda(E \cap P) + \lambda(E \cap N) = \lambda^+(E) - \lambda^-(E)$.
If $\lambda = \mu - \nu$, then $\lambda^+(E) = \mu(E \cap P) - \nu(E \cap P) \leq \mu(E)$ and $\lambda^-(E) = \nu(E \cap N) - \mu(E \cap N) \leq \nu(E)$.
Structural Commentary: Jordan decomposition is the minimal finite-measure difference.
Key Insight: Positive and negative Hahn pieces isolate the irreducible positive masses. □

Corollary 5.7.7. *Variation by Subsets and Null Sets*
For a charge λ and $E, M \in \mathcal{M}$:

(i) $\lambda^+(E) = \sup\{\lambda(F) : F \in \mathcal{M}, F \subseteq E\}$ *and* $\lambda^-(E) = -\inf\{\lambda(F) : F \in \mathcal{M}, F \subseteq E\}$.
(ii) $M \in \mathcal{Z}_\lambda \Leftrightarrow |\lambda|(M) = 0$.
(iii) $\pi_\lambda = \lambda^+$ *and* $\nu_\lambda = |\lambda|$.

Proof.

(i) Let (P, N) be Hahn. If $F \in \mathcal{M}$ and $F \subseteq E$, then $\lambda(F) = \lambda^+(F) - \lambda^-(F) \leq \lambda^+(E)$, with equality at $F = E \cap P$; similarly $\lambda(F) \geq -\lambda^-(E)$, with equality at $F = E \cap N$.
(ii) ($\Rightarrow$) If $M \in \mathcal{Z}_\lambda$, then $\lambda^+(M) = \lambda(M \cap P) = 0$ and $\lambda^-(M) = -\lambda(M \cap N) = 0$, so $|\lambda|(M) = 0$.
(ii) ($\Leftarrow$) If $|\lambda|(M) = 0$, then $\lambda^\pm(G \cap M) = 0$ for every $G \in \mathcal{M}$, hence $\lambda(G \cap M) = 0$ and $M \in \mathcal{Z}_\lambda$.
(iii) The identity $\pi_\lambda = \lambda^+$ is (i). For finite partitions (A_j) of E, $\sum_j |\lambda(A_j)| \leq \sum_j |\lambda|(A_j) = |\lambda|(E)$, and $E = (E \cap P) \sqcup (E \cap N)$ attains equality; hence $\nu_\lambda = |\lambda|$.

Structural Commentary: Variation is intrinsic to λ.
Key Insight: The positive part is the largest charge obtainable inside a set. □

Theorem 5.7.8. *Variations of Integral Charges*
Let $f \in \mathscr{L}^1(\mu)$ and $\lambda(E) := \int_E f \, d\mu$.
Then $\forall E \in \mathcal{M}$, $\lambda^+(E) = \int_E f^+ \, d\mu$, $\lambda^-(E) = \int_E f^- \, d\mu$, and $|\lambda|(E) = \int_E |f| \, d\mu$.

Proof. Set $P = \{f \geq 0\}$ and $N = \{f < 0\}$.
For $A \in \mathcal{M}$, $\lambda(A \cap P) = \int_{A \cap P} f \, d\mu \geq 0$ and $\lambda(A \cap N) = \int_{A \cap N} f \, d\mu \leq 0$, so (P, N) is Hahn.
Thus $\lambda^+(E) = \lambda(E \cap P) = \int_E f^+ \, d\mu$, $\lambda^-(E) = -\lambda(E \cap N) = \int_E f^- \, d\mu$, and $|\lambda|(E) = \lambda^+(E) + \lambda^-(E) = \int_E |f| \, d\mu$.
Structural Commentary: Measure variations match function parts.
Key Insight: Sign sets of the density are Hahn sets. □

Definition 5.7.9. *Absolute Continuity and Singularity*

(i) *Absolute continuity for measures*: $\lambda \ll \mu \Leftrightarrow \forall E \in \mathcal{M}, \mu(E) = 0 \Rightarrow \lambda(E) = 0$.
(ii) *Absolute continuity for charges*: $\lambda \ll \mu \Leftrightarrow |\lambda| \ll \mu$.
(iii) *Mutual singularity*: $\lambda \perp \mu \Leftrightarrow \exists A, B \in \mathcal{M}$ with $A \cap B = \emptyset$, $X = A \cup B$, $\lambda(A) = 0$, and $\mu(B) = 0$.

Proposition 5.7.10. *Absolute Continuity Calculus*
Let λ, μ, ν be measures and η a charge.

(i) $\lambda \ll \lambda$ *and* $\lambda \ll \mu \ll \nu \Rightarrow \lambda \ll \nu$.
(ii) $\ll$ *is not symmetric.*
(iii) $(\forall E \in \mathcal{M}, \mu(E) = 0 \Rightarrow \eta(E) = 0) \Rightarrow \eta^+, \eta^-, |\eta| \ll \mu$.

Proof.

(i) Reflexivity is immediate. If $\lambda \ll \mu \ll \nu$ and $\nu(E) = 0$, then $\mu(E) = 0$, hence $\lambda(E) = 0$.
(ii) On $X = \{0, 1\}$, $\mathcal{M} = 2^X$, $\delta_0 \ll \#$, but $\# \not\ll \delta_0$ since $\delta_0(\{1\}) = 0$ and $\#\{1\} = 1$.
(iii) If $\mu(E) = 0$, then $\mu(F) = 0$ for every $F \in \mathcal{M}$, $F \subseteq E$, so $\eta(F) = 0$.
Hence $E \in \mathcal{Z}_\eta$, and Corollary 5.7.7 gives $|\eta|(E) = 0$; therefore $\eta^{\pm}(E) \le |\eta|(E) = 0$.

Structural Commentary: Absolute continuity is null-set order; charges reduce to variation.
Key Insight: Null sets pass to measurable subsets, so all signed pieces vanish there. □

Example 5.7.11. *Dominating a Countable Family of Measures*
Let (μ_n) be measures with $\mu_n(X) \le 1$, and set $\lambda(E) := \sum_{n\ge1} 2^{-n}\mu_n(E)$.
Then λ is finite, since $\lambda(X) \le 1$, and countably additive by Tonelli for nonnegative series.
Also $2^{-n}\mu_n(E) \le \lambda(E)$, hence $\lambda(E) = 0 \Rightarrow \mu_n(E) = 0$ and $\mu_n \ll \lambda$ for all n.
Key Insight: A summable positive mixture packages countably many null-set tests into one reference measure.

Lemma 5.7.12. *Epsilon–Delta Form of Absolute Continuity*
Let λ be finite and μ a measure on $(X, \mathcal{M})$. Then $\lambda \ll \mu \Leftrightarrow \forall \varepsilon > 0, \exists \delta > 0, \forall E \in \mathcal{M}, \mu(E) < \delta \Rightarrow \lambda(E) < \varepsilon$.

Proof.

($\Rightarrow$) Suppose the ε-δ condition fails. Then $\exists \varepsilon > 0$ and $E_n \in \mathcal{M}$ with $\mu(E_n) < 2^{-n}$ and $\lambda(E_n) \ge \varepsilon$.
Let $F_n = \bigcup_{k\ge n} E_k$ and $F = \bigcap_n F_n$. Then $\mu(F_n) \le 2^{-n+1}$, so $\mu(F) = 0$.
Since $F_n \downarrow F$, $\lambda(F_n) \ge \varepsilon$, and λ is finite, continuity from above gives $\lambda(F) \ge \varepsilon$, contradicting $\lambda \ll \mu$.
($\Leftarrow$) If $\mu(E) = 0$, then $\mu(E) < \delta$ for every $\delta > 0$; hence $\lambda(E) < \varepsilon$ for every $\varepsilon > 0$, so $\lambda(E) = 0$.

Structural Commentary: For finite target measure, absolute continuity is uniform smallness on small μ-sets.
Key Insight: A failing sequence is converted into a decreasing null tail with positive λ-mass. □

Example 5.7.13. *Finite Target Is Needed*
On $(\mathbb{N}, 2^{\mathbb{N}})$, let $\lambda = \#$ and $\mu(E) := \sum_{n\in E} 2^{-n}$. Since $\mu(E) = 0 \Leftrightarrow E = \emptyset$, $\lambda \ll \mu$.
But for $E_n = \{n\}$, $\mu(E_n) = 2^{-n} \to 0$ while $\lambda(E_n) = 1$, so the ε-δ condition fails for $\varepsilon = 1$.

Example 5.7.14. *Mutual Absolute Continuity Need Not Preserve Finiteness*
On $(\mathbb{N}, 2^{\mathbb{N}})$, set $\mu(E) := \sum_{n\in E} 2^{-n}$ and $\lambda(\emptyset) = 0$, $\lambda(E) = +\infty$ for $E \ne \emptyset$.
Then $\mu(\mathbb{N}) = 1$, while λ is not σ-finite. Since both measures have null ideal $\{\emptyset\}$, $\lambda \ll \mu$ and $\mu \ll \lambda$.
Key Insight: Absolute continuity compares null sets; it does not compare finiteness, σ-finiteness, or size.

Theorem 5.7.15. *Radon–Nikodym Theorem*
Let λ, μ be σ-finite measures on $(X, \mathcal{M})$ with $\lambda \ll \mu$. Then there is a measurable $f : X \to [0, +\infty]$, unique μ-a.e. and finite μ-a.e., such that $\lambda(E) = \int_E f\, d\mu$ for all $E \in \mathcal{M}$. Write $f = d\lambda/d\mu$.

Proof. Assume first $\lambda(X), \mu(X) < +\infty$.
For $c > 0$, let (P_k, N_k) be Hahn for $\lambda - kc\mu$, put $A_k = N_k \setminus \bigcup_{j<k} A_j$, and $B = X \setminus \bigcup_{k\ge1} A_k$.
Then $A_k \subseteq N_k$ and, for $k \ge 2$, $A_k \subseteq P_{k-1}$; hence $E \subseteq A_k$ implies $(k-1)c\,\mu(E) \le \lambda(E) \le kc\,\mu(E)$.
Also $B \subseteq P_k$ for all k, so $\mu(B) = \lambda(B) = 0$.
Define $f_c := \sum_{k\ge1}(k-1)c\,\mathbf{1}_{A_k}$. Then $\int_E f_c\, d\mu \le \lambda(E) \le \int_E f_c\, d\mu + c\mu(X)$ for all $E \in \mathcal{M}$.
With $c_n = 2^{-n}$ and $f_n = f_{c_n}$, $\|f_n - f_m\|_1 \le (2^{-n} + 2^{-m})\mu(X) \le 2^{-n+1}\mu(X)$ for $m \ge n$, so $f_n \to f$ in L^1 for some $f \ge 0$. Letting $n \to \infty$ gives $\lambda(E) = \int_E f\, d\mu$.
For σ-finite λ, μ, choose disjoint $Y_r \in \mathcal{M}$ with $X = \bigsqcup_r Y_r$ and $\lambda(Y_r), \mu(Y_r) < +\infty$.
Apply the finite case to $\lambda|_{Y_r}, \mu|_{Y_r}$, get f_r, and set $f = \sum_r f_r \mathbf{1}_{Y_r}$. Then $\lambda(E) = \sum_r \lambda(E \cap Y_r) = \int_E f\, d\mu$.
Since $\int_{Y_r} f\, d\mu = \lambda(Y_r) < +\infty$, $f < +\infty$ μ-a.e. If f, g both represent λ, then $f = g$ μ-a.e. on each Y_r by Proposition 5.4.15; hence $f = g$ μ-a.e. on X.
Structural Commentary: Hahn decompositions of $\lambda - c\mu$ build lower step approximations to the desired density.
Key Insight: The density is recovered by measuring where λ lies between consecutive multiples of μ. □

Example 5.7.16. *A Non-σ-Finite Failure*
Let $X = [0, 1]$, $\mathcal{M} = \mathcal{B}([0, 1])$, $\mu = \#$, and $\lambda = m$. Since $\mu(E) = 0 \Leftrightarrow E = \emptyset$, $\lambda \ll \mu$.
If $\lambda(E) = \int_E f\, d\mu$ for some $f \ge 0$, then $0 = \lambda(\{x\}) = \int_{\{x\}} f\, d\mu = f(x)$ for every x, so $\lambda([0, 1]) = 0$, contradiction.
Thus RN can fail when μ is not σ-finite.

Example 5.7.17. *Another Non-σ-Finite Failure*
Let X be uncountable and $\mathcal{M} := \{E \subseteq X : E \text{ is countable or } X \setminus E \text{ is countable}\}$.
Set $\mu(E) = \#E$ for finite E and $\mu(E) = +\infty$ otherwise; set $\lambda(E) = 0$ for countable E and $\lambda(E) = +\infty$ otherwise. Since $\mu(E) = 0 \Leftrightarrow E = \emptyset$, $\lambda \ll \mu$. If $\lambda(E) = \int_E f\, d\mu$ with $f \ge 0$, then $0 = \lambda(\{x\}) = \int_{\{x\}} f\, d\mu = f(x)$ for all x, so $\lambda(X) = 0$, contradiction.
Structural Commentary: This is the countable–cocountable version of Example 5.7.16.
Key Insight: Non-σ-finite references can erase all singleton tests.

Remark 5.7.18. *When the Radon–Nikodym Derivative Is Integrable*
In Theorem 5.7.15, $\frac{d\lambda}{d\mu} \in L^1(\mu) \Leftrightarrow \lambda(X) < +\infty$, since $\int_X (d\lambda/d\mu)\, d\mu = \lambda(X)$.

Theorem 5.7.19. *Extended-Valued Radon–Nikodym Theorem*
Let μ be σ-finite, let λ be a measure on $(X, \mathcal{M})$, and assume $\lambda \ll \mu$.
Then $\exists$ measurable $f : X \to [0, +\infty]$ such that $\lambda(E) = \int_E f\, d\mu$ for all $E \in \mathcal{M}$; any two such f agree μ-a.e.

Proof. Assume first $\mu(X) < +\infty$. Let (P_n, N_n) be Hahn for $\lambda - n\mu$, and put $P = \bigcap_{n\geq 1} P_n$, $N = \bigcup_{n\geq 1} N_n$.
If $E \subseteq N_n$, then $\lambda(E) \leq n\mu(E)$; hence $\lambda|_N$ is σ-finite and Theorem 5.7.15 gives $g : N \to [0, +\infty)$ with $\lambda(E \cap N) = \int_{E\cap N} g\, d\mu$. If $E \subseteq P$, then $\lambda(E) \geq n\mu(E)$ for all n, so $\mu(E) > 0 \Rightarrow \lambda(E) = +\infty$, while $\mu(E) = 0 \Rightarrow \lambda(E) = 0$. Thus $f = g\mathbf{1}_N + \infty \cdot \mathbf{1}_P$ represents λ.
For general σ-finite μ, choose $X = \bigsqcup_r Y_r$ with $\mu(Y_r) < +\infty$, apply the finite-μ case to each restriction, and set $f = \sum_r f_r \mathbf{1}_{Y_r}$. Then $\lambda(E) = \sum_r \lambda(E \cap Y_r) = \int_E f\, d\mu$. If f, h represent λ, then on each Y_r, $\int_A f\, d\mu = \int_A h\, d\mu$ for all $A \subseteq Y_r$; applying this to $\{f < h\}$ and $\{h < f\}$ gives $f = h$ μ-a.e.
Structural Commentary: If λ is not σ-finite, the density may have a genuinely infinite part.
Key Insight: Hahn decompositions isolate finite and infinite density regions. □

Theorem 5.7.20. *Lebesgue Decomposition Theorem*
Let λ, μ be σ-finite measures on $(X, \mathcal{M})$. Then $\exists!$ measures λ_s, λ_a with $\lambda = \lambda_s + \lambda_a$, $\lambda_s \perp \mu$, and $\lambda_a \ll \mu$.

Proof. Set $\rho = \lambda + \mu$. Choose $f = d\lambda/d\rho$ and $g = d\mu/d\rho$.
Let $A = \{g = 0\}$, $B = \{g > 0\}$, $\lambda_s(E) := \lambda(E \cap A)$, and $\lambda_a(E) := \lambda(E \cap B)$. Since $\mu(A) = \int_A g\, d\rho = 0$, $\lambda_s \perp \mu$.
If $\mu(E) = 0$, then $g = 0$ ρ-a.e. on E, so $\rho(E \cap B) = 0$ and $\lambda_a(E) = 0$; hence $\lambda_a \ll \mu$.
For uniqueness, let $(\tilde\lambda_s, \tilde\lambda_a)$ be another decomposition.
Since $\lambda_s \perp \mu$ and $\tilde\lambda_a \ll \mu$, also $\lambda_s \perp \tilde\lambda_a$; choose $X = C \cup D$ with $\lambda_s(C) = 0$ and $\tilde\lambda_a(D) = 0$.
Then $\lambda_s(E) = \lambda_s(E \cap D) \leq \lambda(E \cap D) = \tilde\lambda_s(E \cap D) \leq \tilde\lambda_s(E)$. By symmetry, $\lambda_s = \tilde\lambda_s$, hence $\lambda_a = \tilde\lambda_a$.
Structural Commentary: Every measure decomposes into a μ-visible part and a part living on a μ-null set.
Key Insight: Differentiate both measures with respect to their sum, then split where the μ-density vanishes. □

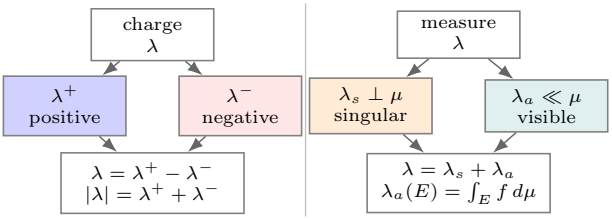


Figure 198: Charges split by sign; measures split into $\lambda_s \perp \mu$ and $\lambda_a \ll \mu$ parts.

Proposition 5.7.21. *Radon–Nikodym Calculus*
Let all measures below be σ-finite.

(i) $\lambda \ll \mu$, $g \geq 0$ *measurable* $\Rightarrow \int_X g\, d\lambda = \int_X g(d\lambda/d\mu)\, d\mu$.
(ii) $\nu \ll \lambda \ll \mu \Rightarrow d\nu/d\mu = (d\nu/d\lambda)(d\lambda/d\mu)$ μ-*a.e.*
(iii) $\lambda_1, \lambda_2 \ll \mu \Rightarrow d(\lambda_1 + \lambda_2)/d\mu = d\lambda_1/d\mu + d\lambda_2/d\mu$ μ-*a.e.*
(iv) $\lambda \ll \mu \ll \lambda \Rightarrow d\lambda/d\mu = 1/(d\mu/d\lambda)$ μ-*a.e.*

Proof.

(i) For $g = \mathbf{1}_E$, $\int_X g\, d\lambda = \lambda(E) = \int_E (d\lambda/d\mu)\, d\mu = \int_X g(d\lambda/d\mu)\, d\mu$.
By linearity the identity holds for nonnegative simple g; for measurable $g \geq 0$, take simple $g_n \uparrow g$ and apply monotone convergence to both sides.

(ii) Let $h := (d\nu/d\lambda)(d\lambda/d\mu)$.
For every $E \in \mathcal{M}$, (i) with $g = \mathbf{1}_E(d\nu/d\lambda)$ gives $\nu(E) = \int_E (d\nu/d\lambda)\, d\lambda = \int_E h\, d\mu$.
Hence h represents ν relative to μ, so $h = d\nu/d\mu$ μ-a.e.

(iii) Let $h := d\lambda_1/d\mu + d\lambda_2/d\mu$. For every $E \in \mathcal{M}$, $\int_E h\, d\mu = \lambda_1(E) + \lambda_2(E) = (\lambda_1 + \lambda_2)(E)$.
Hence h represents $\lambda_1 + \lambda_2$ relative to μ, so uniqueness gives the identity.

(iv) Apply (ii) to $\mu \ll \lambda \ll \mu$. Since $d\mu/d\mu = 1$, $1 = (d\mu/d\lambda)(d\lambda/d\mu)$ μ-a.e.
Also $\mu \ll \lambda$ gives $d\mu/d\lambda > 0$ μ-a.e.; divide to obtain $d\lambda/d\mu = 1/(d\mu/d\lambda)$ μ-a.e.

Structural Commentary: Radon–Nikodym derivatives obey chain and linearity rules.
Key Insight: Prove identities by integrating over arbitrary measurable sets, then invoke uniqueness. □

Corollary 5.7.22. *Absolutely Continuous and Singular Forces Zero*
For measures λ, μ, $\lambda \ll \mu$ and $\lambda \perp \mu$ imply $\lambda = 0$. For a charge η, $|\eta| \perp \mu$ implies $\eta^+ \perp \mu$ and $\eta^- \perp \mu$.

Proof. If $\lambda \perp \mu$, choose $X = A \cup B$ with $\lambda(A) = 0$ and $\mu(B) = 0$. Since $\lambda \ll \mu$, $\lambda(B) = 0$, hence $\lambda(X) = 0$.
If $|\eta| \perp \mu$, choose $X = A \cup B$ with $|\eta|(A) = 0$ and $\mu(B) = 0$. Since $\eta^{\pm} \le |\eta|$, $\eta^{\pm}(A) = 0$, so $\eta^{\pm} \perp \mu$.
Structural Commentary: Absolute continuity and singular support are incompatible except at zero.
Key Insight: Visible and singular forces zero. □

Corollary 5.7.23. *Radon–Nikodym for Charges*
Let μ be σ-finite and let λ be a charge with $\lambda \ll \mu$. Then $\exists f \in \mathscr{L}^1(\mu)$ such that, for all $E \in \mathcal{M}$, $\lambda(E) = \int_E f\,d\mu$, $\lambda^+(E) = \int_E f^+\,d\mu$, $\lambda^-(E) = \int_E f^-\,d\mu$, and $|\lambda|(E) = \int_E |f|\,d\mu$.

Proof. Since $|\lambda| \ll \mu$, $\lambda^{\pm} \ll \mu$. By Radon–Nikodym, choose $f_{\pm} \ge 0$ with $\lambda^{\pm}(E) = \int_E f_{\pm}\,d\mu$. Since $\lambda^+ \perp \lambda^-$, $f_+ f_- = 0$ μ-a.e. Set $f = f_+ - f_-$. Then $\int_X |f|\,d\mu = \int_X (f_+ + f_-)\,d\mu = |\lambda|(X) < +\infty$, so $f \in \mathscr{L}^1(\mu)$, $f^+ = f_+$, $f^- = f_-$ μ-a.e., and $\lambda(E) = \lambda^+(E) - \lambda^-(E) = \int_E f\,d\mu$.
Structural Commentary: Signed absolute continuity is reduced to the two positive Jordan pieces.
Key Insight: Differentiate λ^+ and λ^- separately, then subtract. □

Example 5.7.24. *Radon–Nikodym Classification of Basic Stieltjes Generators*
Let m be Lebesgue measure on $\mathcal{B}(\mathbb{R})$, and let λ_i be the Borel–Stieltjes measure generated by g_i.

(i) $g_1(x) = 2x$: $\lambda_1 = 2m$, so $\lambda_1 \ll m$, $d\lambda_1/dm = 2$, $m \ll \lambda_1$, $dm/d\lambda_1 = 1/2$, and $\lambda_1(\mathbb{R}) = +\infty$.
(ii) $g_2(x) = \arctan x$: $\lambda_2(E) = \int_E (1+x^2)^{-1}\,dm$, so $d\lambda_2/dm = (1+x^2)^{-1}$, $\lambda_2(\mathbb{R}) = \pi$, $m \ll \lambda_2$, and $dm/d\lambda_2 = 1 + x^2$.
(iii) $g_3 = \mathbf{1}_{[0,+\infty)}$: $\lambda_3 = \delta_0$,
so $\lambda_3 \perp m$, $\lambda_3(\mathbb{R}) = 1$, $\lambda_3 \not\ll m$, and $m \not\ll \lambda_3$.
(iv) $g_4(x) = x\mathbf{1}_{[0,+\infty)}(x)$: $\lambda_4(E) = m(E \cap [0,+\infty))$, so $\lambda_4 \ll m$, $d\lambda_4/dm = \mathbf{1}_{[0,+\infty)}$, $m \not\ll \lambda_4$, and $\lambda_4(\mathbb{R}) = +\infty$.

Thus $\lambda_1, \lambda_2, \lambda_4 \ll m$, $\lambda_3 \perp m$, only λ_2, λ_3 are finite, and $m \ll \lambda_i$ exactly for $i = 1, 2$.
Structural Commentary: Stieltjes generators split naturally into density, atom, and restriction examples.
Key Insight: A positive density makes null sets equivalent; an atom lives on a Lebesgue-null set.

Theorem 5.7.25. *Charges Form a Banach Space*
Let $\mathrm{ca}(X, \mathcal{M}) := \{\lambda : \mathcal{M} \to \mathbb{R} : \lambda$ *is a charge*$\}$ *and* $\|\lambda\|_{\mathrm{var}} := |\lambda|(X)$.
With pointwise operations, $(\mathrm{ca}(X, \mathcal{M}), \|\cdot\|_{\mathrm{var}})$ *is Banach.*

Proof. Linearity follows from charge linearity.
Definiteness follows from $|\lambda|(X) = 0 \Rightarrow |\lambda|(E) = 0\ \forall E$, homogeneity from $|a\lambda| = |a|\,|\lambda|$, and the triangle inequality from $\sum_j |(\lambda+\eta)(A_j)| \le \sum_j |\lambda(A_j)| + \sum_j |\eta(A_j)|$ over finite measurable partitions (A_j).
Let (λ_n) be variation-Cauchy. Since $|\lambda_n(E) - \lambda_m(E)| \le \|\lambda_n - \lambda_m\|_{\mathrm{var}}$, define $\lambda(E) := \lim_n \lambda_n(E)$; then λ is finitely additive. For fixed n and finite disjoint (A_j), $\sum_j |(\lambda - \lambda_n)(A_j)| \le \limsup_m \|\lambda_m - \lambda_n\|_{\mathrm{var}}$.
For disjoint (E_k), set $E = \bigcup_k E_k$ and $R_K = \bigcup_{k>K} E_k$. Choose n with the preceding limsup $< \varepsilon$.
Since $\lambda_n(R_K) \to 0$, for large K, $|\lambda(R_K)| \le |\lambda_n(R_K)| + |(\lambda - \lambda_n)(R_K)| < 2\varepsilon$. Hence $\lambda(R_K) \to 0$, so λ is countably additive. Taking suprema over finite partitions gives $\|\lambda - \lambda_n\|_{\mathrm{var}} \le \limsup_m \|\lambda_m - \lambda_n\|_{\mathrm{var}} \to 0$.
Structural Commentary: Total variation is the signed-measure analogue of an L^1 norm.
Key Insight: Variation-Cauchy controls every finite partition. □

Definition 5.7.26. *Averages and Absolute Continuity*

(i) *Average*: for $f \in \mathscr{L}^1(E)$ and $A \subseteq E$, $m(A) > 0$, set $[f : A] := m(A)^{-1} \int_A f\,dm$.
(ii) *Density ratio*: for measurable B and cube or ball Q, $m(Q) > 0$, set $[B : Q] := m(B \cap Q)/m(Q)$.
(iii) *Indefinite integral*: for $f \in \mathscr{L}^1([a,b])$, $F(x) = F(a) + \int_a^x f(t)\,dt$.
(iv) *Absolute continuity*: $F : [a,b] \to \mathbb{R}$ is absolutely continuous if $\forall \varepsilon > 0, \exists \delta > 0$,
finite disjoint $(a_j, b_j) \subseteq [a,b]$ satisfy $\sum_j (b_j - a_j) < \delta \Rightarrow \sum_j |F(b_j) - F(a_j)| < \varepsilon$.

Definition 5.7.27. *Bounded Variation*

(i) *Total variation*: for $F : [a,b] \to \mathbb{R}$, $V_a^b(F) := \sup\{\sum_{j=1}^N |F(x_j) - F(x_{j-1})| : a = x_0 < \cdots < x_N = b\}$.
(ii) *Bounded variation*: F has bounded variation on $[a,b]$ if $V_a^b(F) < +\infty$.

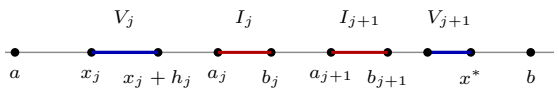


Figure 199: Interval bookkeeping for absolute continuity.

Perspective 5.7.28. *Absolute Continuity as Interval Bookkeeping*
If $F(x) = \int_a^x f(t)\,dt$, by absolute continuity, disjoint intervals with small total length force small F total change.

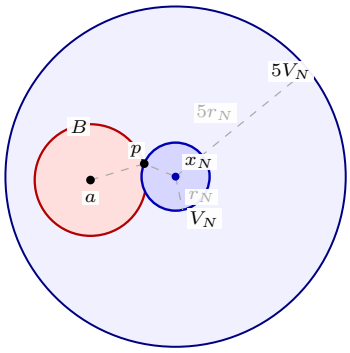


Figure 200: A comparable missed ball lies inside the fivefold dilation of the selected ball.

Perspective 5.7.29. *Vitali Covering Control*
Select balls disjointly by decreasing radius.
If B is rejected, then $B \cap B_j \neq \emptyset$ for some earlier B_j with $r(B) \leq r(B_j)$, hence $B \subseteq 5B_j$.
Key Insight: Disjoint balls add; fixed dilations cover misses.

Definition 5.7.30. *Vitali Coverings and Density Points*

(i) For $A \subseteq M$, $\mathcal{V}$ is a *Vitali covering* of A if $\forall p \in A, \forall r > 0, \exists V \in \mathcal{V}$ with $p \in V$ and $\operatorname{diam}(V) < r$.
(ii) For measurable $E \subseteq \mathbb{R}^n$, its *density-point set* is $\operatorname{dp}(E) := \{p \in \mathbb{R}^n : \lim_{Q \downarrow p}[E : Q] = 1\}$, with Q cubes, equivalently balls, shrinking to p.

Theorem 5.7.31. *Vitali Covering Lemma for Balls and Cubes*
Let $E \subseteq \mathbb{R}^n$, $m^*(E) < +\infty$, and let $\mathcal{V}$ be a Vitali cover of E by closed balls or axis-parallel cubes. Then $\exists$ pairwise disjoint $(V_j) \subseteq \mathcal{V}$ with $m^*(E \setminus \bigcup_{j \geq 1} V_j) = 0$.

Proof. Fix open $O \supseteq E$ with $m(O) < m^*(E) + \varepsilon$, and replace $\mathcal{V}$ by $\{V \in \mathcal{V} : V \subseteq O\}$.
Greedily choose a maximal disjoint (V_j) by dyadic diameter scales.
For a fixed dilation $V \mapsto V^*$, if $W \cap V \neq \emptyset$ and $\operatorname{diam} W \leq 2 \operatorname{diam} V$, then $W \subseteq V^*$ and $m(V^*) \leq C_n m(V)$.
Let $R = E \setminus \bigcup_j V_j$. Fix N and $p \in R$.
Since $V_1, \ldots, V_N$ are closed, choose $W \in \mathcal{V}$ with $p \in W$ and $W \cap \bigcup_{j \leq N} V_j = \emptyset$. Maximality gives some $j > N$ with $W \cap V_j \neq \emptyset$ and comparable diameter, hence $p \in W \subseteq V_j^*$. Thus $R \subseteq \bigcup_{j > N} V_j^*$.
Therefore $m^*(R) \leq C_n \sum_{j > N} m(V_j) \to 0$, because $\sum_j m(V_j) = m(\bigsqcup_j V_j) \leq m(O) < +\infty$. Hence $m^*(R) = 0$.
Structural Commentary: Vitali selection gives a disjoint cover modulo null sets.
Key Insight: $R \subseteq \bigcup_{j > N} V_j^*$ and disjoint tails vanish. □

Theorem 5.7.32. *Lebesgue Density Theorem*
For measurable $E \subseteq \mathbb{R}^n$, $\lim_{Q \downarrow p}[E : Q] = \mathbf{1}_E(p)$ for a.e. p, where Q ranges over cubes, equivalently balls. Equivalently, $m(E \triangle \operatorname{dp}(E)) = 0$.

Proof. Assume $m(E) < +\infty$. For $a \in (0,1)$ set $E_a := \{p \in E : \liminf_{Q \downarrow p}[E : Q] < a\}$.
Cubes $Q \ni p$ with $[E : Q] < a$ form a Vitali cover of E_a. For open $O \supseteq E_a$ with $m(O) < m(E_a) + \varepsilon$, Theorem 5.7.31 gives disjoint $(Q_j) \subseteq O$ covering E_a modulo null sets.
Hence $m(E_a) \leq \sum_j m(E_a \cap Q_j) \leq \sum_j m(E \cap Q_j) < a \sum_j m(Q_j) \leq a\, m(O) < a(m(E_a) + \varepsilon)$.
Let $\varepsilon \downarrow 0$; since $a < 1$, $m(E_a) = 0$. Thus $p \in E \setminus \bigcup_{a \in \mathbb{Q} \cap (0,1)} E_a$ has density 1. Apply same argument to bounded pieces of E^c to get density 0 a.e. on E^c; then exhaust arbitrary E by bounded pieces.
Structural Commentary: Measurable sets are locally 0 or 1 a.e.
Key Insight: Vitali turns bad density into $m(E_a) \leq a\, m(E_a)$. □

Corollary 5.7.33. *Density Points Commute with Slicing Almost Everywhere*
Let $E \subseteq \mathbb{R}^n \times \mathbb{R}^k$ be measurable. For a.e. $x \in \mathbb{R}^n$, $\operatorname{dp}(E_x) = \{y : (x,y) \in \operatorname{dp}(E)\}$ up to a null subset of $\mathbb{R}^k$.

Proof. Set $D = \operatorname{dp}(E)$ and $N = E \triangle D$.
By Theorem 5.7.32, $m_{n+k}(N) = 0$; hence Theorem 5.6.19 gives $m_k(N_x) = 0$ for a.e. x.
Intersecting with the Fubini good set, E_x, D_x are measurable and $m_k(E_x \triangle D_x) = m_k((E \triangle D)_x) = 0$.
Applying Theorem 5.7.32 in $\mathbb{R}^k$ gives $m_k(E_x \triangle \operatorname{dp}(E_x)) = 0$.
Therefore $m_k(\operatorname{dp}(E_x) \triangle D_x) = 0$; and $m_k(\operatorname{dp}(E_x) \triangle D_x) = 0$. This is the claimed equality modulo null sets.
Structural Commentary: Almost-everywhere local fullness is compatible with almost-everywhere slicing.
Key Insight: Null ambient errors have null slices for almost every base point. □

Lemma 5.7.34. *Large Averages Have Small Outer Measure*
For a cube $Q_0 \subseteq \mathbb{R}^n$, $g \geq 0$ in $\mathscr{L}^1(Q_0)$, and $\lambda > 0$,
set $A_\lambda := \{p \in Q_0 : \forall r > 0\ \exists$ cube $Q \subseteq Q_0,\ p \in Q,\ \operatorname{diam} Q < r,\ [g : Q] > \lambda\}$ for $m^*(A_\lambda) \leq \frac{1}{\lambda}\int_{Q_0} g$.

Proof. Let $\mathcal{V}_\lambda = \{Q \subseteq Q_0 : Q$ is a cube and $[g : Q] > \lambda\}$. By definition, $\mathcal{V}_\lambda$ is a Vitali cover of A_λ; choose disjoint $(Q_j) \subseteq \mathcal{V}_\lambda$ with $m^*(A_\lambda \setminus \bigcup_j Q_j) = 0$. Then $\lambda m^*(A_\lambda) \leq \lambda \sum_j m(Q_j) < \sum_j \int_{Q_j} g \leq \int_{Q_0} g$.
Structural Commentary: Weak type for local averages.
Key Insight: Disjoint Vitali cubes turn pointwise large averages into one integral bound. □

Theorem 5.7.35. *Lebesgue Average Value Theorem*
For $f \in \mathscr{L}^1_{\mathrm{loc}}(\mathbb{R}^n)$, $\lim_{Q \downarrow p}[f : Q] = f(p)$ for a.e. $p \in \mathbb{R}^n$ with Q cubes, equivalently balls, shrinking to p.

Proof. Fix a cube Q_0; it suffices to prove the claim for a.e. $p \in Q_0^\circ$.
For $\alpha > 0$, let $B_\alpha = \left\{p \in Q_0^\circ : \limsup_{Q \downarrow p} |[f : Q] - f(p)| > \alpha\right\}$.
Choose $h \in C(Q_0)$ with $\|f - h\|_{L^1(Q_0)} < \eta$, and set $g = |f - h|$. Since $[h : Q] \to h(p)$, $B_\alpha \subseteq \{g > \alpha/3\} \cup A_{\alpha/3}(g)$.
By Markov and Lemma 5.7.34, $m^*(B_\alpha) \leq \frac{3}{\alpha}\int_{Q_0} g + \frac{3}{\alpha}\int_{Q_0} g < \frac{6\eta}{\alpha}$.
Let $\eta \downarrow 0$; then $m(B_\alpha) = 0$. Count over $\alpha \in \mathbb{Q}_{>0}$ and exhaust $\mathbb{R}^n$ by cube interiors.
Structural Commentary: Local averages recover the original integrable function almost everywhere.
Key Insight: Continuous approximation handles the good part; Vitali's estimate makes the bad part small. □

Theorem 5.7.36. *Lebesgue's Antiderivative Theorem*
For $f \in \mathscr{L}^1([a,b])$, define $F(x) = c + \int_a^x f(t)\,dt$. Then $F \in AC([a,b])$ and $F'(x) = f(x)$ for a.e. $x \in [a,b]$.

Proof. Given $\varepsilon > 0$, choose $\delta > 0$ so that $m(A) < \delta$ implies $\int_A |f| < \varepsilon$.
If (α_j, β_j) are disjoint and $\sum_j (\beta_j - \alpha_j) < \delta$, then $\sum_j |F(\beta_j) - F(\alpha_j)| \leq \sum_j \int_{\alpha_j}^{\beta_j} |f| = \int_{\bigcup_j (\alpha_j, \beta_j)} |f| < \varepsilon$, so $F \in AC([a,b])$. For $x, x+h \in [a,b]$, let $I_h(x) = (\min\{x, x+h\}, \max\{x, x+h\})$.
Then $\frac{F(x+h)-F(x)}{h} = [f : I_h(x)]$. Extending f by 0 outside $[a,b]$, Theorem 5.7.35 in one dimension gives $[f : I_h(x)] \to f(x)$ for a.e. x. Hence $F' = f$ a.e.
Structural Commentary: The indefinite Lebesgue integral is the correct antiderivative for integrable functions.
Key Insight: Differentiating the integral is the same as shrinking the averaging interval. □

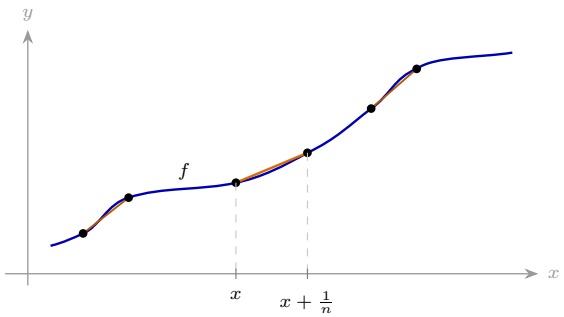


Figure 201: $g_n(x)$ is the slope of the short right secant from x to $x + \frac{1}{n}$.

Perspective 5.7.37. *Right Secants and Differentiation*
Right quotients $g_n(x) = n\big(f(x + 1/n) - f(x)\big)$ encode local slope data.
Key Insight: Secants turn differentiation into slope estimates.

Theorem 5.7.38. *Lebesgue's Last Theorem and Consequences*
On $[a,b]$:

(i) F monotone $\Rightarrow F'$ exists a.e.;
(ii) F Lipschitz $\Rightarrow F'$ exists a.e.;
(iii) $F \in BV([a,b]) \Rightarrow F'$ exists a.e.;
(iv) $F \in AC([a,b]) \Rightarrow F \in BV([a,b])$, hence F' exists a.e.

Proof.

(i) Assume F nondecreasing. Let $D^+F(x) = \limsup_{h \downarrow 0} \frac{F(x+h)-F(x)}{h}$ and $D_-F(x) = \liminf_{h \downarrow 0} \frac{F(x+h)-F(x)}{h}$.
For $\alpha < \beta$ in $\mathbb{Q}$, set $A_{\alpha,\beta} = \{D_-F < \alpha < \beta < D^+F\}$. The right-secant Vitali selection of Figure 201

gives $m(A_{\alpha,\beta}) = 0$: otherwise selected intervals with slopes $< \alpha$ and $> \beta$ contradict monotonicity of total rise. Also $m^*(\{D^+F > M\}) \le \frac{F(b)-F(a)}{M} \to 0$, so $D^+F = D_-F < \infty$ a.e.; hence F' exists a.e.

(ii) If F is L-Lipschitz, set $G_+(x) = F(x) + Lx$ and $G_-(x) = Lx - F(x)$. Then $G_\pm$ are nondecreasing, by (i) $\exists$ null sets $N_\pm$ s.t. $G'_\pm$ exist off $N_\pm$. For $x \notin N_+ \cup N_-$, $F'(x) = \frac{G'_+(x)-G'_-(x)}{2}$, so F' exists a.e.

(iii) If $F \in BV([a,b])$, write its Jordan decomposition $F = F(a) + V^+ - V^-$ with $V^\pm$ nondecreasing. By (i), $\exists$ null sets $N_\pm$ s.t. $(V^\pm)'$ exist off $N_\pm$. For $x \notin N_+ \cup N_-$, $F'(x) = (V^+)'(x) - (V^-)'(x)$, so F' exists a.e.

(iv) If $F \in AC([a,b])$, choose δ for $\varepsilon = 1$, and partition $a = x_0 < \cdots < x_N = b$ with $x_i - x_{i-1} < \delta$. Grouping any refinement by these subintervals gives $V_a^b(F) \le N$, so $F \in BV([a,b])$. By (iii), $\exists Z \subset [a,b]$ with $m(Z) = 0$ s.t. $F'(x)$ exists for all $x \in [a,b] \setminus Z$.

Structural Commentary: Monotonicity drives all cases.
Key Insight: Vitali controls secants; decompose to monotone. □

5.8 Lp Spaces

Lebesgue spaces turn integrability into normed linear geometry. The essential point is that an integral norm cannot distinguish functions that differ only on null sets, so the actual vectors are almost-everywhere equivalence classes. Their bounded linear functionals are themselves represented by integration against functions in the conjugate space.

Definition 5.8.1. *Norms and Seminorms*
Let $\mathcal{V}$ be a vector space over $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$.

(i) A *seminorm* is $N : \mathcal{V} \to [0,\infty)$ with $N(\alpha v) = |\alpha|\, N(v)$ and $N(u+v) \le N(u) + N(v)$ $\forall u, v \in \mathcal{V}$, $\alpha \in \mathbb{F}$.
(ii) A *norm* is a seminorm with $N(v) = 0 \iff v = 0$.
(iii) A norm N induces $d_N(u,v) := N(u-v)$.

Example 5.8.2. *Norms and Seminorms Before Quotienting*

(i) On $\mathbb{F}^n$, $\|u\|_p := (\sum_{j=1}^n |u_j|^p)^{1/p}$ for $1 \le p < \infty$ and $\|u\|_\infty := \max_{1\le j\le n} |u_j|$ are norms. Minkowski gives $\|u+v\|_p \le \|u\|_p + \|v\|_p$.
(ii) On sequences, $\ell^p := \{a : \sum_n |a_n|^p < +\infty\}$ with $\|a\|_p := (\sum_n |a_n|^p)^{1/p}$, and $\ell^\infty := \{a : \sup_n |a_n| < +\infty\}$ with $\|a\|_\infty := \sup_n |a_n|$.
(iii) On $B(X)$, $\|f\|_\infty := \sup_{x\in X} |f(x)|$ is a norm. On $C[0,1]$, $N_0(f) := |f(0)|$ and $N_1(f) := \sup_{0\le x\le 1/2} |f(x)|$ are seminorms. On $C^1[a,b]$, $N_D(f) := \sup_{a\le x\le b} |f'(x)|$ is a seminorm with kernel the constant functions.

Key Insight: Seminorms become norms after quotienting out their zero directions.

Remark 5.8.3. *Integral Norm on Continuous Functions*
On $C[0,1]$, $N(f) := \int_0^1 |f(x)|\, dx$ is a norm: if $f(x_0) \neq 0$, continuity gives $N(f) > 0$. It is not complete, since continuous ramps can be Cauchy in N and converge in L^1 to a discontinuous step function.

Definition 5.8.4. *L^p Spaces and Mean Convergence*
Let $(X, \mathcal{M}, \mu)$ be a measure space.

(i) For $1 \le p < \infty$, the *p-integrable class* is $\mathscr{L}^p(\mu) := \{f : X \to \mathbb{F} : f \text{ measurable}, \int_X |f|^p\, d\mu < +\infty\}$.
(ii) The *L^p space* and *L^p norm* are $L^p(\mu) := \mathscr{L}^p(\mu)/=_{\text{a.e.}}$ and $\|f\|_p := (\int_X |f|^p\, d\mu)^{1/p}$.
(iii) *L^p convergence*, or *mean convergence of order p*, means $f_n \to f \iff \|f_n - f\|_p \to 0$. *$L^p$-Cauchy* means $\|f_m - f_n\|_p \to 0$ as $m, n \to \infty$.
(iv) *Essentially bounded* means $|f| \le M$ a.e. for some $M < +\infty$; $L^\infty(\mu)$ is the a.e.-quotient of such functions, with *essential sup norm* $\|f\|_\infty := \inf\{M \ge 0 : |f| \le M \text{ a.e.}\} = \inf_{\mu(N)=0} \sup_{x\in X\setminus N} |f(x)|$.

Proposition 5.8.5. *L^1 as a Quotient of a Seminorm*
On $\mathscr{L}^1(\mu)$ set $N_\mu(f) := \int_X |f|\, d\mu$. Then N_μ is a seminorm and $N_\mu(f) = 0 \iff f = 0$ a.e. Hence $L^1(\mu) := \mathscr{L}^1(\mu)/=_{\text{a.e.}}$ with $\|[f]\|_1 := \int_X |f|\, d\mu$ is normed; write f for $[f]$.

Proof. $\mathscr{L}^1(\mu)$ is linear since $|af + bg| \le |a|\,|f| + |b|\,|g|$. Hence $N_\mu(\alpha f) = |\alpha|\, N_\mu(f)$ and $N_\mu(f+g) \le N_\mu(f) + N_\mu(g)$.

($\Rightarrow$) If $N_\mu(f) = 0$, set $E_n := \{|f| \ge 1/n\}$.
Then $(1/n)\mu(E_n) \le \int_{E_n} |f|\, d\mu \le 0$, so $\mu(E_n) = 0$ and $\{f \neq 0\} = \bigcup_n E_n$ is null.
($\Leftarrow$) If $f = 0$ a.e., then $N_\mu(f) = \int_X |f|\, d\mu = 0$.

The quotient kills exactly $\ker N_\mu$, so $\|\cdot\|_1$ is well defined and definite. *Structural Commentary*: The seminorm kernel is exactly the a.e.-zero functions.
Key Insight: Quotient by null differences before calling the integral a norm. □

Definition 5.8.6. *Complex Integrals and L^2*

(i) For $f = u + iv$, *complex measurability* means measurability of u, v.
(ii) On $E \in \mathcal{M}$, *complex integrability* means $u, v \in \mathscr{L}^1(\mu)$ on E; set $\int_E f\,d\mu := \int_E u\,d\mu + i\int_E v\,d\mu$.
(iii) $L^2(\mu)$ consists of complex measurable f with $\int_X |f|^2\,d\mu < +\infty$, modulo a.e. equality; $\|f\|_2 := (\int_X |f|^2\,d\mu)^{1/2}$.
(iv) The L^2 *inner product* is $\langle f, g\rangle_{L^2} := \int_X f\overline{g}\,d\mu$; it induces $\|f\|_2$.

Proposition 5.8.7. *Complex L^1 and Dominated Convergence*

(i) For complex measurable f on E, f is integrable $\iff |f| \in \mathscr{L}^1(\mu)$ on E; then $\left|\int_E f\,d\mu\right| \le \int_E |f|\,d\mu$.
(ii) If $f_n \to f$ a.e. on E and $|f_n| \le g \in \mathscr{L}^1(\mu)$, then f is integrable and $\int_E f\,d\mu = \lim_n \int_E f_n\,d\mu$.

Proof.

(*i*) ($\Rightarrow$) Write $f = u + iv$. If $u, v \in \mathscr{L}^1$, then $|f| \le |u| + |v|$, so $|f| \in \mathscr{L}^1$.
(*i*) ($\Leftarrow$) If $|f| \in \mathscr{L}^1$, then $|u|, |v| \le |f|$, so $u, v \in \mathscr{L}^1$ and f is integrable.
If $\int_E f\,d\mu = re^{i\theta}$ and $h = e^{-i\theta}f$, then $\left|\int_E f\,d\mu\right| = r = \operatorname{Re}\int_E h\,d\mu = \int_E \operatorname{Re} h\,d\mu \le \int_E |h|\,d\mu = \int_E |f|\,d\mu$.
(*ii*) Write $f_n = u_n + iv_n$ and $f = u + iv$. Since $|u_n|, |v_n| \le g$ and $u_n \to u$, $v_n \to v$ a.e., real DCT gives $\int_E u_n\,d\mu \to \int_E u\,d\mu$ and $\int_E v_n\,d\mu \to \int_E v\,d\mu$; hence the complex identity.

Structural Commentary: Complex integration is two real integrals.
Key Insight: Rotate the integral to make its value real and nonnegative before applying the real inequality. □

Theorem 5.8.8. *L^2 Geometry*
For $f, g \in L^2(\mu)$, $\int_X |fg|\,d\mu \le \|f\|_2\,\|g\|_2$, $f + g \in L^2(\mu)$, and $\|f + g\|_2 \le \|f\|_2 + \|g\|_2$.
Modulo equality a.e., $(L^2(\mu), \|\cdot\|_2)$ is complete.

Proof. For $\lambda \in \mathbb{R}$, $0 \le Q(\lambda) := \int_X (|f| + \lambda\,|g|)^2\,d\mu = \|f\|_2^2 + 2\lambda\int_X |fg|\,d\mu + \lambda^2\,\|g\|_2^2$; hence $\Delta(Q) \le 0$ and Schwarz follows. Thus $\|f + g\|_2^2 \le \int_X (|f| + |g|)^2\,d\mu \le (\|f\|_2 + \|g\|_2)^2$.
Let (f_n) be Cauchy. Choose $n_k \uparrow \infty$ so $m, n \ge n_k$ imply $\|f_m - f_n\|_2 \le 2^{-k}$, and set $h_N := \sum_{k<N} \left|f_{n_{k+1}} - f_{n_k}\right|$.
Then $\|h_N\|_2 \le \sum_{k<N} 2^{-k}$, so Fatou gives $h := \lim_N h_N \in L^2$; hence $f_{n_k} \to f$ a.e. for some f.
For fixed n, Fatou gives $\|f - f_n\|_2 \le \liminf_k \|f_{n_k} - f_n\|_2$.
The Cauchy property sends the right side to 0 as $n \to \infty$, so $f_n \to f$ in L^2.
Structural Commentary: L^2 is Hilbert-space geometry in measure form.
Key Insight: The same Schwarz mechanism survives because the integral is positive and additive. □

Lemma 5.8.9. *Young's Inequality*
If $A, B \ge 0$, $1 < p < \infty$, and $p^{-1} + q^{-1} = 1$, then $AB \le A^p/p + B^q/q$, with equality iff $A^p = B^q$.

Proof. If $AB = 0$, the claim is immediate. For $A, B > 0$, let $\alpha = 1/p$.
Since $\varphi(t) = \alpha t + 1 - \alpha - t^\alpha \ge 0$ for $t \ge 0$, with equality iff $t = 1$, apply this to $t = A^p/B^q$ and multiply by B^q.
$AB \le A^p/p + B^q/q$, with equality iff $A^p = B^q$.
Structural Commentary: Young's inequality is weighted AM–GM in conjugate exponents.
Key Insight: Rescale to A^p and B^q, then use convexity. □

Theorem 5.8.10. *Hölder's Inequality*
Let $1 < p < \infty$ and $p^{-1} + q^{-1} = 1$. Then $f \in L^p(\mu)$ and $g \in L^q(\mu)$ imply $fg \in L^1(\mu)$ and $\|fg\|_1 \le \|f\|_p\,\|g\|_q$.

Proof. If $\|f\|_p\,\|g\|_q = 0$, then $|f|^p = 0$ or $|g|^q = 0$ a.e., so $fg = 0$ a.e.
Otherwise set $A = |f| / \|f\|_p$ and $B = |g| / \|g\|_q$. Young gives $|fg| / (\|f\|_p\,\|g\|_q) \le A^p/p + B^q/q$.
The right side lies in L^1, so $fg \in L^1$, and integration yields $\|fg\|_1 / (\|f\|_p\,\|g\|_q) \le p^{-1} + q^{-1} = 1$.
Structural Commentary: Normalize both factors to unit L^p and L^q size.
Key Insight: Products are controlled by conjugate powers. □

Definition 5.8.11. *Conjugate Indices*
For $p, q \in (1, \infty)$, p and q are *conjugate indices* if $p^{-1} + q^{-1} = 1$, equivalently $p + q = pq$.
The unique self-conjugate index is $p = q = 2$.

Corollary 5.8.12. *Cauchy–Bunyakovsky–Schwarz in L^2*
If $f, g \in L^2(\mu)$, then $fg \in L^1(\mu)$ and $\left|\int_X fg\,d\mu\right| \le \int_X |fg|\,d\mu \le \|f\|_2\,\|g\|_2$.

Proof. Hölder with $p = q = 2$ gives $fg \in L^1$ and $\int_X |fg|\,d\mu \le \|f\|_2\,\|g\|_2$.
The first inequality is the complex integral estimate applied to fg.
Structural Commentary: The self-conjugate exponent turns Hölder into Cauchy–Schwarz.
Key Insight: Pair two L^2 factors through the 2–2 conjugacy. □

Theorem 5.8.13. *Minkowski's Inequality*
If $1 \le p < \infty$, then $f, h \in L^p(\mu)$ imply $f + h \in L^p(\mu)$ and $\|f+h\|_p \le \|f\|_p + \|h\|_p$.

Proof. For $p = 1$, integrate $|f + h| \le |f| + |h|$. Let $p > 1$ and $q = p/(p-1)$.
Since $|f+h|^p \le 2^{p-1}(|f|^p + |h|^p)$, $f + h \in L^p$ and $|f+h|^{p-1} \in L^q$.
From $|f+h|^p \le (|f| + |h|)\,|f+h|^{p-1}$, Hölder gives $\|f+h\|_p^p \le (\|f\|_p + \|h\|_p)\,\|f+h\|_p^{p-1}$.
If $\|f+h\|_p > 0$, divide by $\|f+h\|_p^{p-1}$; otherwise the inequality is trivial.
Structural Commentary: Hölder turns the pointwise triangle inequality into a norm triangle inequality.
Key Insight: Minkowski is exactly the closure and triangle estimate for L^p. □

Theorem 5.8.14. *L^p Is Normed*
For $1 \le p < \infty$, $(L^p(\mu), \|\cdot\|_p)$ is a normed linear space.

Proof. Pointwise operations descend to a.e.-classes, so $L^p(\mu)$ is linear. For $\alpha \in \mathbb{F}$, $\|\alpha f\|_p = |\alpha|\,\|f\|_p$.
If $\|f\|_p = 0$, then $\int_X |f|^p\,d\mu = 0$, hence $|f|^p = 0$ a.e. and $f = 0$ in L^p. The triangle axiom is Minkowski.
Structural Commentary: Quotienting by a.e. equality turns pointwise operations into vector-space operations.
Key Insight: Minkowski supplies the nontrivial norm axiom. □

Definition 5.8.15. *Completeness Language*
A normed space $(V, \|\cdot\|)$ is *complete*, or a *Banach space*, if every Cauchy sequence $(x_n) \subset V$ has $x \in V$ with $\|x_n - x\| \to 0$. Thus $L^p(\mu)$ is complete iff every L^p-Cauchy sequence converges in $\|\cdot\|_p$.

Lemma 5.8.16. *Convergent Implies Cauchy*
If $f_n \to f$ in $L^p(\mu)$, then (f_n) is Cauchy in $L^p(\mu)$.

Proof. Given $\varepsilon > 0$, choose N so $n \ge N$ implies $\|f_n - f\|_p < \varepsilon/2$.
Then $m, n \ge N$ gives $\|f_m - f_n\|_p \le \|f_m - f\|_p + \|f - f_n\|_p < \varepsilon$ by Minkowski.
Structural Commentary: A convergent sequence has uniformly small tails around its limit.
Key Insight: Use the common limit as the bridge between two tails. □

Theorem 5.8.17. *Riesz–Fischer Completeness Theorem*
For $1 \le p < \infty$, $(L^p(\mu), \|\cdot\|_p)$ is complete.

Proof. Let (f_n) be Cauchy. Choose $n_k \uparrow \infty$ so $m, n \ge n_k$ imply $\|f_m - f_n\|_p \le 2^{-k}$, and set $g_k = f_{n_k}$.
For $G_N := |g_1| + \sum_{k<N} |g_{k+1} - g_k|$, Minkowski gives $\|G_N\|_p \le \|g_1\|_p + \sum_{k<N} 2^{-k} \le \|g_1\|_p + 1$.
Fatou gives $G := \lim_N G_N \in L^p$. Hence $\sum_k |g_{k+1} - g_k| < +\infty$ a.e., so $f := g_1 + \sum_k (g_{k+1} - g_k)$ exists a.e. and $g_k \to f$ a.e. Since $|g_k - f| \le G$ a.e., DCT gives $\|g_k - f\|_p \to 0$.
Finally, for $\varepsilon > 0$, choose k with $\|g_k - f\|_p < \varepsilon/2$ and then $n \ge n_k$ with $\|f_n - g_k\|_p < \varepsilon/2$; hence $\|f_n - f\|_p < \varepsilon$.
Structural Commentary: A rapid subsequence converges a.e.; DCT gives norm convergence.
Key Insight: Fatou finds the pointwise limit, dominated convergence brings it back to L^p. □

Proposition 5.8.18. *Essential Supremum Test*
If $f \in L^\infty(\mu)$, then $|f| \le \|f\|_\infty$ a.e.; if $A < \|f\|_\infty$, then $\mu(\{|f| > A\}) > 0$.

Proof. For each k, $|f| \le \|f\|_\infty + 1/k$ a.e.; intersect the full-measure sets and let $k \to \infty$.
If $A < \|f\|_\infty$ and $\mu(\{|f| > A\}) = 0$, then A is an essential bound, contradiction.
Structural Commentary: The essential supremum is the least a.e. upper bound.
Key Insight: The sharp bound cannot be lowered on a full-measure set. □

Theorem 5.8.19. *L^∞ Is Banach*
$(L^\infty(\mu), \|\cdot\|_\infty)$ is a Banach space.

Proof. Linearity is pointwise modulo a.e. equality; homogeneity is immediate. The essential bound test gives definiteness, and off the union of two null sets $|f + g| \le \|f\|_\infty + \|g\|_\infty$, giving the triangle axiom.
Let (f_n) be L^∞-Cauchy. For each m, n, choose null N_{mn} with $|f_m - f_n| \le \|f_m - f_n\|_\infty$ on $X \setminus N_{mn}$, and put $N = \bigcup_{m,n} N_{mn}$. On $X \setminus N$, (f_n) is uniformly Cauchy, hence uniformly converges to a measurable f.
Thus $\|f_n - f\|_\infty \to 0$; for fixed N_0, $|f| \le |f_{N_0}| + \sup_{n \ge N_0} \|f_n - f_{N_0}\|_\infty$ a.e., so $f \in L^\infty$.
Structural Commentary: L^∞ completeness is uniform convergence after discarding one null set.
Key Insight: Essential bounds become ordinary uniform bounds off a common null set. □

Proposition 5.8.20. *Simple Approximation, Support, and Tails*
Let $1 \le p < \infty$ and $f \in L^p(\mu)$.

(i) Simple measurable functions are dense in $L^p(\mu)$; also in $L^\infty(\mu)$ for $\|\cdot\|_\infty$.
(ii) $\{|f| > 0\}$ is σ-finite.

(iii) *For* $E_n := \{|f| \geq n\}$, $\mu(E_n) \to 0$.
(iv) *For every* $\varepsilon > 0$, *there is* $E_\varepsilon \in \mathcal{M}$ *with* $\mu(E_\varepsilon) < +\infty$ *such that* $F \cap E_\varepsilon = \emptyset$ *implies* $\|f\mathbf{1}_F\|_p < \varepsilon$.

Proof.

(i) Choose simple $s_n \to f$ a.e. with $|s_n| \leq |f|$ in the real case; treat real and imaginary parts separately in the complex case. Then $|f - s_n|^p \leq 2^p |f|^p$, so DCT gives $\|f - s_n\|_p \to 0$. For $f \in L^\infty$, discard null N with $|f| \leq \|f\|_\infty + \eta$ on $X \setminus N$, then quantize that bounded range.
(ii) Put $A_k := \{|f| \geq 1/k\}$. Then $\mu(A_k) \leq k^p \|f\|_p^p < +\infty$ and $\{|f| > 0\} = \bigcup_k A_k$.
(iii) $n^p \mu(E_n) \leq \int_{E_n} |f|^p \, d\mu \leq \|f\|_p^p$, hence $\mu(E_n) \leq n^{-p} \|f\|_p^p \to 0$.
(iv) With A_k as in (ii), $A_k \uparrow \{|f| > 0\}$ and $|f|^p \mathbf{1}_{X \setminus A_k} \to 0$ a.e. DCT gives $\int_{X \setminus A_k} |f|^p \, d\mu \to 0$.
Choose $E_\varepsilon = A_k$ with this integral $< \varepsilon^p$; if $F \cap E_\varepsilon = \emptyset$, then $\|f\mathbf{1}_F\|_p^p \leq \int_{X \setminus E_\varepsilon} |f|^p \, d\mu < \varepsilon^p$.

Structural Commentary: L^p mass is simple-approximable and finite-piece capturable.
Key Insight: Integrability turns level sets into finite-measure exhaustions. □

Proposition 5.8.21. *Shell Criterion on Finite Measure Spaces*
Let $\mu(X) < +\infty$, f *be measurable, and* $E_n := \{n - 1 \leq |f| < n\}$ *for* $n \geq 1$.
For $1 \leq p < \infty$, $f \in L^p(\mu) \iff \sum_{n \geq 1} n^p \mu(E_n) < +\infty$.

Proof. If $\sum_n n^p \mu(E_n) < +\infty$, then $|f|^p \leq \sum_n n^p \mathbf{1}_{E_n}$, so $f \in L^p$.
Conversely, if $f \in L^p$, then $\mu(E_1) \leq \mu(X) < +\infty$ and, for $n \geq 2$, $n^p \mathbf{1}_{E_n} \leq 2^p |f|^p \mathbf{1}_{E_n}$.
Hence $\sum_n n^p \mu(E_n) \leq \mu(X) + 2^p \int_X |f|^p \, d\mu < +\infty$.
Structural Commentary: The finite base shell handles small values; weighted shells control tails.
Key Insight: Shell weights characterize L^p mass. □

Proposition 5.8.22. *Inclusions, Interpolation, and Multipliers*

(i) *If* $\mu(X) < +\infty$ *and* $1 \leq r \leq p < \infty$, *then* $L^p(\mu) \subseteq L^r(\mu)$ *and* $\|f\|_r \leq \|f\|_p \mu(X)^{1/r - 1/p}$.
(ii) *On* $(\mathbb{N}, 2^{\mathbb{N}}, \#)$, $1 \leq p \leq s < \infty$ *implies* $\ell^p \subseteq \ell^s$ *and* $\|a\|_s \leq \|a\|_p$.
(iii) *If* $1 \leq p_1 < p_2 < \infty$ *and* $f \in L^{p_1}(\mu) \cap L^{p_2}(\mu)$, *then* $f \in L^p(\mu)$ *for* $p_1 \leq p \leq p_2$.
(iv) *If* $1 \leq p \leq \infty$, *then* $L^p(\mu) \cdot L^\infty(\mu) \subseteq L^p(\mu)$ *and* $\|fg\|_p \leq \|f\|_p \|g\|_\infty$.
(v) $L^\infty(\mu) \subseteq L^1(\mu) \iff \mu(X) < +\infty$. *If* $\mu(X) = 1$, *then* $\|f\|_\infty = \lim_{p \to \infty} \|f\|_p$ *for* $f \in L^\infty(\mu)$.

Proof.

(i) For $r = p$, $L^p \subseteq L^p$ and the estimate reads $\|f\|_p \leq \|f\|_p$, i.e. equality.
For $r < p$, Hölder applied to $|f|^r \cdot 1$ gives $\int_X |f|^r \, d\mu \leq \|f\|_p^r \mu(X)^{1 - r/p}$; take rth roots.
(ii) For $a \in \ell^p$, $|a_n| \leq \|a\|_p$ for all n, hence $\sum_n |a_n|^s \leq \|a\|_p^{s-p} \sum_n |a_n|^p \leq \|a\|_p^s$.
(iii) For $p \in [p_1, p_2]$, $|f|^p \leq |f|^{p_1} \mathbf{1}_{\{|f| \leq 1\}} + |f|^{p_2} \mathbf{1}_{\{|f| > 1\}}$.
(iv) For $p < \infty$, $|fg|^p \leq \|g\|_\infty^p |f|^p$ a.e.; for $p = \infty$, $|fg| \leq \|f\|_\infty \|g\|_\infty$ a.e.
(v) $(\Rightarrow)$ If $L^\infty(\mu) \subseteq L^1(\mu)$, then $\mathbf{1}_X \in L^\infty(\mu)$ lies in $L^1(\mu)$, so $\mu(X) = \int_X \mathbf{1}_X \, d\mu < +\infty$.
(v) $(\Leftarrow)$ If $\mu(X) < +\infty$ and $f \in L^\infty(\mu)$, then $\int_X |f| \, d\mu \leq \|f\|_\infty \mu(X) < +\infty$. If $\mu(X) = 1$, then $\|f\|_p \leq \|f\|_\infty$.
For $A < \|f\|_\infty$, $\mu(\{|f| > A\}) > 0$ and $\|f\|_p \geq A \mu(\{|f| > A\})^{1/p} \to A$. Let $A \uparrow \|f\|_\infty$.

Structural Commentary: Inclusions depend on domain size and pointwise bounds.
Key Insight: Check whether size comes from the function values or from the measure of the domain. □

Example 5.8.23. *Membership Models*

(i) On $(\mathbb{N}, 2^{\mathbb{N}}, \#)$, $a_n = 1/n$ satisfies $\|a\|_p^p = \sum_{n \geq 1} n^{-p}$; hence $a \in \ell^p \iff p > 1$, while $a \in \ell^\infty \setminus \ell^1$.
(ii) On $\mathbb{N}$ with $\lambda(\{n\}) = n^{-2}$, $\lambda(\mathbb{N}) < +\infty$ and $f(n) = n^{1/p_0}$ satisfies $\int_{\mathbb{N}} |f|^p \, d\lambda = \sum_{n \geq 1} n^{p/p_0 - 2}$.
Hence $f \in L^p(\lambda) \iff 1 \leq p < p_0$. In particular, $f(n) = \sqrt{n}$ gives $p_0 = 2$.
(iii) On $(0, \infty)$, $f(x) = x^{-1/2}(1 + |\log x|)^{-1}$ gives $\int_0^\infty |f|^p \, dx = \int_{-\infty}^\infty e^{(1-p/2)t}(1 + |t|)^{-p} \, dt$ under $x = e^t$.
If $p < 2$ it diverges at $+\infty$; if $p > 2$ at $-\infty$; if $p = 2$ it is $\int (1 + |t|)^{-2} \, dt < +\infty$.

Structural Commentary: L^p membership reduces to series and endpoint tests.
Key Insight: Different ends impose different p-conditions.

Proposition 5.8.24. *Hölder Dual Norm Attainment*
Let $1 < p < \infty$, $p^{-1} + q^{-1} = 1$, *and* $f \in L^p(\mu)$ *be real-valued. Then* $\sup_{\|g\|_q \leq 1} \left| \int_X fg \, d\mu \right| = \|f\|_p$.
If $f \neq 0$, *the supremum is attained by* $g_0 = \|f\|_p^{-p/q} \operatorname{sign}(f) |f|^{p-1}$.

Proof. Hölder gives $\left|\int_X fg\,d\mu\right| \le \|f\|_p \|g\|_q$, hence the supremum is $\le \|f\|_p$.
Since $(p-1)q = p$, $\|g_0\|_q^q = \|f\|_p^{-p} \int_X |f|^p\,d\mu = 1$, and $\int_X fg_0\,d\mu = \|f\|_p^{-p/q} \int_X |f|^p\,d\mu = \|f\|_p$.
Structural Commentary: Equality in Hölder is obtained by matching the conjugate powers.
Key Insight: The L^q test function points in the power direction of f. □

Definition 5.8.25. *Bounded Linear Functionals on L^p*

(i) A *linear functional* on real $L^p(\mu)$ is a linear map $G : L^p(\mu) \to \mathbb{R}$.
(ii) G is *bounded* if $\exists M < +\infty$ such that $|G(f)| \le M \|f\|_p$ for all $f \in L^p(\mu)$.
(iii) The *operator norm* is $\|G\| := \sup\{|G(f)| : f \in L^p(\mu),\ \|f\|_p \le 1\}$.
(iv) G is *positive* if $f \ge 0$ a.e. implies $G(f) \ge 0$.

Lemma 5.8.26. *Positive Split of a Bounded Functional*
Let $1 \le p < +\infty$ and $G \in (L^p(\mu))^*_{\mathbb{R}}$. Then $G = G^+ - G^-$ for positive bounded linear functionals $G^\pm$.

Proof. For $f \ge 0$, set $G^+(f) := \sup\{G(g) : 0 \le g \le f\}$.
This is finite since $|G(g)| \le \|G\| \|g\|_p \le \|G\| \|f\|_p$, and positive homogeneity is immediate.
For $f_1, f_2 \ge 0$, $G^+(f_1) + G^+(f_2) \le G^+(f_1 + f_2)$; conversely, if $0 \le h \le f_1 + f_2$, set $h_1 = (h - f_2)^+$ and $h_2 = h - h_1$, so $h = h_1 + h_2$ with $0 \le h_j \le f_j$, hence $G(h) \le G^+(f_1) + G^+(f_2)$.
Thus G^+ is additive on the positive cone. Extend by $G^+(f) = G^+(f^+) - G^+(f^-)$; then G^+ is positive, linear, and bounded. Let $G^- = G^+ - G$. Since $G^+(f) \ge G(f)$ for $f \ge 0$, G^- is positive and bounded.
Structural Commentary: Function-lattice functionals split into positive parts.
Key Insight: Split any competitor $h \le f_1 + f_2$ into the portions assigned to f_1 and f_2. □

Theorem 5.8.27. *Riesz Representation for L^1*
Let $(X, \mathcal{M}, \mu)$ be σ-finite and $G \in (L^1(\mu))^*_{\mathbb{R}}$.

(i) There is $g \in L^\infty(\mu)$ such that $G(f) = \int_X fg\,d\mu$ for all $f \in L^1(\mu)$.
(ii) $\|G\| = \|g\|_\infty$.
(iii) If G is positive, then $g \ge 0$ a.e.

Proof.

(i) First let $\mu(X) < +\infty$ and $G \ge 0$. Set $\lambda(E) := G(\mathbf{1}_E)$. Finite additivity is linearity; if $E_n \uparrow E$, then $\mathbf{1}_{E_n} \to \mathbf{1}_E$ in L^1, hence $\lambda(E_n) \to \lambda(E)$. Thus λ is finite, and $\lambda \ll \mu$ since $\mu(E) = 0$ gives $\mathbf{1}_E = 0$ in L^1. Radon–Nikodym gives $g \ge 0$ with $\lambda(E) = \int_E g\,d\mu$, so simple functions are represented.
By (ii), $g \in L^\infty$, hence L^1-density of simple functions extends the identity to all $f \in L^1$. For σ-finite μ, apply the finite case to $f \mapsto G(f\mathbf{1}_{X_n})$ on $X_n \uparrow X$, $\mu(X_n) < +\infty$, and patch densities on overlaps. For general G, use the positive split $G = G^+ - G^-$ and subtract the two densities.
(ii) The pairing estimate gives $\|G\| \le \|g\|_\infty$. Conversely, if $A > \|G\|$ and $E \subseteq \{|g| > A\}$ has $\mu(E) < +\infty$, then with $f = \operatorname{sign}(g)\mathbf{1}_E$, $A\mu(E) < \int_E |g|\,d\mu = G(f) \le \|G\|\,\mu(E)$, so $\mu(E) = 0$.
A σ-finite exhaustion gives $\mu(\{|g| > A\}) = 0$ for every $A > \|G\|$, hence $\|g\|_\infty \le \|G\|$.
(iii) If G is positive and $B = \{g < 0\}$, then for every finite-measure $E \subseteq B$, $0 \le G(\mathbf{1}_E) = \int_E g\,d\mu \le 0$ only if $\int_E (-g)\,d\mu = 0$. Hence $\mu(E) = 0$ on $\{g < -1/k\}$ for all k, and a σ-finite exhaustion gives $\mu(B) = 0$.

Structural Commentary: A bounded functional on L^1 is integration against an essentially bounded density.
Key Insight: Indicators give a measure; Radon–Nikodym gives the density. □

Theorem 5.8.28. *Riesz Representation for L^p, $1 < p < +\infty$*
Let $(X, \mathcal{M}, \mu)$ be any measure space, $1 < p < +\infty$, and $q = p/(p-1)$. For every $G \in (L^p(\mu))^*_{\mathbb{R}}$, there is a unique $g \in L^q(\mu)$ modulo a.e. equality such that $G(f) = \int_X fg\,d\mu$ for all $f \in L^p(\mu)$, and $\|G\| = \|g\|_q$.

Proof. On finite measure spaces, repeat the L^1 proof with $\mathbf{1}_{E_n} \to \mathbf{1}_E$ in L^p and the positive split.
Testing truncations of $\operatorname{sign}(g)\,|g|^{q-1}$ gives $g \in L^q$ and the reverse Hölder inequality, hence $\|G\| = \|g\|_q$; uniqueness follows from testing finite-measure sets.
On σ-finite spaces, patch the finite-measure representatives along $X_n \uparrow X$, $\mu(X_n) < +\infty$.
For arbitrary μ, choose $\|f_n\|_p = 1$ with $G(f_n) \to \|G\|$; the supports of the f_n lie in a σ-finite X_0.
If $E \cap X_0 = \emptyset$ and $\mu(E) < +\infty$, boundedness on $f_n \pm t\mathbf{1}_E$, followed by $n \to \infty$ and $t \downarrow 0$, gives $G(\mathbf{1}_E) = 0$.
Thus G vanishes off X_0; represent it on X_0 and extend the density by 0.
Structural Commentary: For $1 < p < \infty$, every bounded functional comes from the conjugate space L^q.
Key Insight: Localize to σ-finite support, then patch. □

Remark 5.8.29. *L^2 Riesz Gives Another Radon–Nikodym Route*
Let λ, μ be finite and $\lambda \ll \mu$. Put $\rho = \lambda + \mu$ and $G(h) := \int h\,d\lambda$ on $L^2(\rho)$.
If $G(h) = \int hg\,d\rho$, indicator tests give $0 \le g \le 1$ ρ-a.e. and $\mu(\{g = 1\}) = 0$.

Since $\int h(1-g)\,d\lambda = \int hg\,d\mu$, take $E_N = E \cap \{1-g \geq 1/N\}$ and $h = \mathbf{1}_{E_N}/(1-g) \in L^2(\rho)$ to get $\lambda(E_N) = \int_{E_N} g/(1-g)\,d\mu$. Let $N \to \infty$ by monotone convergence and $\lambda \ll \mu$: $\lambda(E) = \int_E g/(1-g)\,d\mu$.
Key Insight: Riesz representation on $L^2(\lambda+\mu)$ encodes the density in the ratio $g/(1-g)$.

Lemma 5.8.30. *Absolute Continuity of L^p Mass*
For $1 \leq p < \infty$ and $f \in L^p(\mu)$, $\forall \varepsilon > 0\, \exists \delta > 0$ such that $\mu(E) < \delta$ implies $\|f\mathbf{1}_E\|_p < \varepsilon$.

Proof. Let $h = |f|^p \in L^1$. Choose $M > 0$ with $\int_{\{h>M\}} h\,d\mu < \varepsilon^p/2$, and set $\delta = \varepsilon^p/(2M)$.
If $\mu(E) < \delta$, then $\|f\mathbf{1}_E\|_p^p = \int_E h\,d\mu \leq M\mu(E) + \int_{\{h>M\}} h\,d\mu < \varepsilon^p$.
Structural Commentary: L^p mass is absolutely continuous with respect to measure.
Key Insight: Truncate the tail; small measure controls the bounded part. □

Proposition 5.8.31. *Local Mass Limits for L^p Cauchy Sequences*
For $1 \leq p < \infty$ and $f_n \in L^p(\mu)$, set $\beta_n(E) := \|f_n\mathbf{1}_E\|_p$ for $E \in \mathcal{M}$. Then $|\beta_n(E) - \beta_m(E)| \leq \|f_n - f_m\|_p$. Hence, if (f_n) is L^p-Cauchy, then $\beta(E) := \lim_n \beta_n(E)$ exists for every $E \in \mathcal{M}$.

Proof. Minkowski on E gives $\beta_n(E) \leq \beta_m(E) + \|(f_n - f_m)\mathbf{1}_E\|_p \leq \beta_m(E) + \|f_n - f_m\|_p$; swap m, n.
Thus L^p-Cauchy implies $(\beta_n(E))$ is Cauchy in $\mathbb{R}$.
Structural Commentary: Restricting to E is a 1-Lipschitz operation on L^p mass.
Key Insight: Local mass inherits global Cauchy control. □

Proposition 5.8.32. *Uniform Local Mass Control for L^p Cauchy Sequences*
Let $1 \leq p < \infty$, f_n be L^p-Cauchy, and $\beta_n(E) := \|f_n\mathbf{1}_E\|_p$. For every $\varepsilon > 0$:
(i) $\exists E_\varepsilon \in \mathcal{M}$ with $\mu(E_\varepsilon) < +\infty$ such that $F \cap E_\varepsilon = \emptyset$ implies $\sup_n \beta_n(F) < \varepsilon$;
(ii) $\exists \delta > 0$ such that $\mu(E) < \delta$ implies $\sup_n \beta_n(E) < \varepsilon$.

Proof. Choose N so $n \geq N$ implies $\|f_n - f_N\|_p < \varepsilon/3$.
(i) Choose finite-measure E_j for $1 \leq j \leq N$ with $F \cap E_j = \emptyset \Rightarrow \beta_j(F) < \varepsilon/3$, and set $E_\varepsilon = \bigcup_{j \leq N} E_j$.
For $n \leq N$ this works; for $n > N$, $\beta_n(F) \leq \|f_n - f_N\|_p + \beta_N(F) < \varepsilon$.
(ii) Choose $\delta_j > 0$ for $1 \leq j \leq N$ with $\mu(E) < \delta_j \Rightarrow \beta_j(E) < \varepsilon/3$, and set $\delta = \min_{j \leq N} \delta_j$.
For $n \leq N$ this works; for $n > N$, $\beta_n(E) \leq \|f_n - f_N\|_p + \beta_N(E) < \varepsilon$.
Structural Commentary: Cauchy families are uniformly tight and absolutely continuous.
Key Insight: Control finitely many early terms directly and the tail by Cauchy closeness. □

Method 5.8.33. *Lebesgue Proof Patterns*
(i) *Approximate*: for $f \geq 0$, use $s \in \mathcal{S}^+(f)$ and $\int_E f\,d\mu = \sup_{s \in \mathcal{S}^+(f)} \mathcal{I}_E(s)$.
(ii) *Promote*: prove for $\mathbf{1}_A$, extend to $s = \sum_i c_i \mathbf{1}_{A_i}$, then pass through $s_n \uparrow f$.
(iii) *Ignore null sets*: if $f \sim_{\text{a.e.}} g$ on E, replace f by g inside integrals over measurable $A \subseteq E$.
(iv) *Split signs*: use $f = f^+ - f^-$ and $h = u + iv$.
(v) *Exchange limits*: reduce Fatou and dominated convergence to $0 \leq f_n \uparrow f$.
(vi) *Slice*: prove boxes first; use regularity and null slices to pass to products and Fubini.
Key Insight: Reduce to simple functions, monotone convergence, null sets, and product boxes.

5.9 Modes of Convergence

Throughout this subsection $(X, \mathcal{M}, \mu)$ is a fixed measure space and, unless otherwise stated, functions are real-valued and measurable. The point is not to rank convergence modes linearly, but to know which hypotheses convert one mode into another.

Definition 5.9.1. *Pointwise, Uniform, and Almost-Everywhere Convergence*
Let (f_n) and f be functions on X.
(i) $f_n \to f$ *pointwise* if $\forall x \in X$, $\varepsilon > 0$, $\exists N(\varepsilon, x)$ s.t. $n \geq N(\varepsilon, x)$ implies $|f_n(x) - f(x)| < \varepsilon$.
(ii) $f_n \to f$ *uniformly* if $\forall \varepsilon > 0\ \exists N(\varepsilon)$ s.t. $n \geq N(\varepsilon)$ implies $|f_n(x) - f(x)| < \varepsilon\ \forall x \in X$.
(iii) $f_n \to f$ *almost everywhere* if $\exists A \in \mathcal{M}$ with $\mu(A) = 0$ s.t. $\forall x \in X \setminus A$, $\varepsilon > 0$, $\exists N(\varepsilon, x)$ s.t. $n \geq N(\varepsilon, x)$ implies $|f_n(x) - f(x)| < \varepsilon$.

Proposition 5.9.2. *First Implications and First Failures*
Let $f_n, f : X \to \mathbb{R}$.
(i) $f_n \rightrightarrows f \Rightarrow f_n \to f$ pointwise $\Rightarrow f_n \to f$ a.e.; converses fail.
(ii) $|X| < +\infty$ and $f_n \to f$ pointwise $\Rightarrow f_n \rightrightarrows f$.
(iii) $\mu(A) = 0 \Rightarrow A = \emptyset$ and $f_n \to f$ a.e. $\Rightarrow f_n \to f$ pointwise.

Proof.

(i) If $f_n \rightrightarrows f$, then $\forall \varepsilon > 0\ \exists N_\varepsilon$ such that $n \geq N_\varepsilon$ implies $|f_n(x) - f(x)| < \varepsilon$ for all $x \in X$; hence pointwise convergence with $N(\varepsilon, x) = N_\varepsilon$. Pointwise convergence is a.e. convergence with $A = \emptyset$. Conversely, on $([0,1], \lambda)$, $x^n \to \mathbf{1}_{\{1\}}$ pointwise but not uniformly, and $\mathbf{1}_{\{0\}} \to 0$ a.e. but not pointwise.

(ii) If X is finite and $f_n \to f$ pointwise, then for each $\varepsilon > 0$ and $x \in X$ choose $N(\varepsilon, x)$. Set $N_\varepsilon = \max_{x \in X} N(\varepsilon, x)$; then $n \geq N_\varepsilon$ implies $|f_n(x) - f(x)| < \varepsilon$ for all $x \in X$, so $f_n \rightrightarrows f$.

(iii) If $f_n \to f$ a.e., then $\exists A \in \mathcal{M}$ with $\mu(A) = 0$ and $f_n(x) \to f(x)$ for all $x \in X \setminus A$. Since every null set is empty, $A = \emptyset$, so $f_n \to f$ pointwise.

Structural Commentary: Uniformity is bounded x-dependence of $N(\varepsilon, x)$.
Key Insight: Uniformity controls indices; a.e. convergence controls null exceptions. □

Proposition 5.9.3. *Uniform Convergence Gives L^p Convergence on Finite Measure Spaces*
Let $1 \leq p < \infty$, $\mu(X) < +\infty$, and $f_n \in L^p(\mu)$. If $f_n \rightrightarrows f$ on X, then $f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$.

Proof. For $\varepsilon > 0$, choose N with $n \geq N \Rightarrow |f_n - f| \leq \varepsilon$ on X.
Then $n \geq N \Rightarrow \|f_n - f\|_p \leq (\int_X \varepsilon^p \, d\mu)^{1/p} = \varepsilon \mu(X)^{1/p}$, so $\|f_n - f\|_p \to 0$.
Taking $\varepsilon = 1$, choose n_0 with $|f_{n_0} - f| \leq 1$ on X; hence $f - f_{n_0} \in L^p(\mu)$ and $f = f_{n_0} + (f - f_{n_0}) \in L^p(\mu)$.
Structural Commentary: Uniform error becomes L^p error after multiplication by $\mu(X)^{1/p}$.
Key Insight: Finite measure converts a pointwise sup-bound into an integral bound. □

Theorem 5.9.4. *Dominated Almost-Everywhere Convergence Gives L^p Convergence*
Let $1 \leq p < \infty$, $f_n \in L^p(\mu)$, and $f_n \to f$ a.e.
If $\exists g \in L^p(\mu)$ such that $|f_n| \leq g$ a.e. for all n, then $f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$.

Proof. Off a null set, $f_n \to f$ and $|f_n| \leq g$ for all n; hence $|f| \leq g$, so $f \in L^p(\mu)$. Also $|f_n - f|^p \to 0$ a.e. and $|f_n - f|^p \leq (2g)^p \in \mathscr{L}^1(\mu)$ a.e.; Theorem 5.5.4 gives $\|f_n - f\|_p^p = \int_X |f_n - f|^p \, d\mu \to 0$.
Structural Commentary: Apply ordinary dominated convergence to the pth-power error.
Key Insight: A single L^p envelope controls both membership of the limit and convergence in mean. □

Corollary 5.9.5. *Uniform Boundedness on Finite Measure Spaces*
Let $\mu(X) < +\infty$, $0 \leq K < +\infty$, and $f_n \to f$ a.e.
If $|f_n| \leq K$ a.e. for all n, then for every $1 \leq p < \infty$, $f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$.

Proof. Fix $1 \leq p < \infty$ and set $g := K\mathbf{1}_X$. Then $\int_X g^p \, d\mu = K^p \mu(X) < +\infty$, so $g \in L^p(\mu)$, and $|f_n| \leq g$ a.e. for all n. Since $f_n \to f$ a.e., Theorem 5.9.4 gives $f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$.
Structural Commentary: Finite measure turns a uniform a.e. bound into an L^p envelope.
Key Insight: Bounded a.e. convergence on finite measure is dominated convergence. □

Example 5.9.6. *Uniform Need Not Imply L^p on Infinite Measure*
On $(\mathbb{R}, \mathcal{B}(\mathbb{R}), \lambda)$, fix $1 \leq p < \infty$ and set $f_n := n^{-1/p} \mathbf{1}_{[0,n]}$. Then $\|f_n\|_\infty = n^{-1/p} \to 0$, so $f_n \rightrightarrows 0$.
However $\|f_n\|_p^p = \int_0^n n^{-1} \, d\lambda = 1$, hence $\|f_n\|_p \nrightarrow 0$.
Structural Commentary: The support grows exactly enough to preserve L^p mass.
Key Insight: Uniform smallness alone does not control integrals on infinite measure spaces.

Example 5.9.7. *Almost Everywhere and In-Measure Need Not Imply L^p*
On $[0,1]$ with Lebesgue measure, set $f_n := n\mathbf{1}_{[1/n,\, 2/n]}$.
For each $x \in [0,1]$, $x \in [1/n, 2/n]$ for only finitely many n, hence $f_n(x) \to 0$.
If $\alpha > 0$ and $n \geq \alpha$, then $\lambda\{|f_n| \geq \alpha\} = 1/n \to 0$, so $f_n \to 0$ in measure.
Yet for $1 \leq p < \infty$, $\|f_n\|_p^p = n^p \lambda([1/n, 2/n]) = n^{p-1} \nrightarrow 0$.
Given $\delta > 0$, choose $0 < \eta < \delta$; on $[0,1] \setminus [0, \eta]$, $f_n = 0$ for all $n > 2/\eta$, so $f_n \to 0$ almost uniformly.
Thus a.e., in-measure, and almost-uniform convergence do not imply L^p convergence.

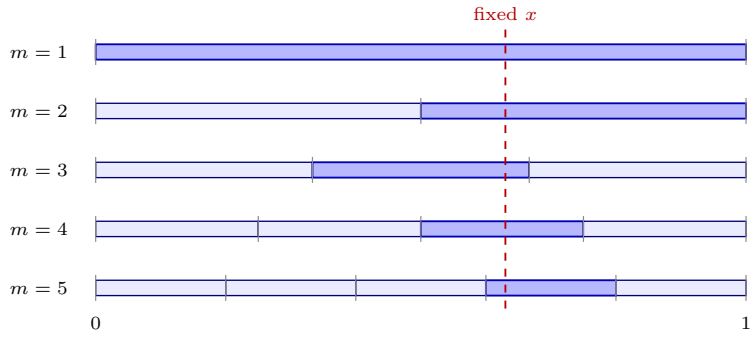


Figure 202: Typewriter intervals shrink, but each fixed x is hit once per partition.

Example 5.9.8. *L^p Convergence Without Pointwise Convergence*
Let $X = [0,1]$ with Lebesgue measure and list $I_{m,k} := [(k-1)/m, k/m]$, $1 \le k \le m$, row by row in m.
Let $f_n := \mathbf{1}_{I_n}$, I_n is nth interval. If I_n in row $m(n)$, then $m(n) \to \infty$, $\|f_n\|_p^p = \lambda(I_n) = 1/m(n) \to 0$; hence $f_n \to 0$ in L^p, and therefore in measure by Proposition 5.9.10.
$\forall x \in [0,1]$, every row contains some $I_{m,k}$ with $x \in I_{m,k}$, giving a subsequence with value 1 at x; for infinitely many rows, some $I_{m,k}$ omits x, giving a subsequence with value 0 at x. Thus $(f_n(x))$ converges nowhere.
The subfamily $I_{m,2} = [1/m, 2/m]$, $m \ge 2$, gives a subsequence converging pointwise to 0 on all of X.

Definition 5.9.9. *Convergence and Cauchy Convergence in Measure*
Let (f_n) be measurable functions on X.

(i) $f_n \to f$ *in measure* if $\forall \alpha > 0$, $\mu\{x \in X : |f_n(x) - f(x)| \ge \alpha\} \to 0$.
(ii) (f_n) is *Cauchy in measure* if $\forall \alpha > 0$, $\mu\{x \in X : |f_m(x) - f_n(x)| \ge \alpha\} \to 0$ as $m, n \to \infty$.

The same definitions result with $>$ in place of $\ge$.

Proposition 5.9.10. *L^p Convergence Implies Convergence in Measure*
Let $1 \le p < \infty$. If $\|f_n - f\|_p \to 0$, then $f_n \to f$ in measure.

Proof. For $\alpha > 0$, set $E_n(\alpha) := \{|f_n - f| \ge \alpha\}$.
Since $\|f_n - f\|_p^p \ge \int_{E_n(\alpha)} |f_n - f|^p \, d\mu \ge \alpha^p \mu(E_n(\alpha))$, $\mu(E_n(\alpha)) \le \alpha^{-p} \|f_n - f\|_p^p \to 0$.
Structural Commentary: Large-error sets are controlled by the integral of the pth-power error.
Key Insight: This is Chebyshev's inequality in deterministic measure form. □

Theorem 5.9.11. *Riesz Subsequence Principle for Convergence in Measure*
If (f_n) is Cauchy in measure, then there exist a subsequence (g_k) and a measurable f such that:

(i) $g_k \to f$ a.e.
(ii) $g_k \to f$ in measure.
(iii) For every $\varepsilon > 0$, there is $E \in \mathcal{M}$ with $\mu(E) < \varepsilon$ such that $g_k \to f$ uniformly on $X \setminus E$.

Proof.

(i) Choose $n_1 < n_2 < \cdots$ so that $m, n \ge n_k \Rightarrow \mu\{|f_m - f_n| \ge 2^{-k}\} < 2^{-k}$, and set $g_k := f_{n_k}$.
For $E_k := \{|g_{k+1} - g_k| \ge 2^{-k}\}$, $\mu(E_k) < 2^{-k}$. Set $F_k := \bigcup_{j=k}^{\infty} E_j$.
Then $F_k \downarrow$, $\mu(F_k) \le \sum_{j=k}^{\infty} 2^{-j} = 2^{-(k-1)}$, and for $i > j \ge k$, $x \notin F_k$,
$|g_i(x) - g_j(x)| \le \sum_{r=j}^{i-1} |g_{r+1}(x) - g_r(x)| < 2^{-(j-1)}$. Hence (g_j) is uniformly Cauchy on $X \setminus F_k$ $\forall k$.
Let $F := \bigcap_k F_k$, so $\mu(F) = 0$. Define $f = \lim_j g_j$ on $X \setminus F$, $f = 0$ on F. f is measurable, $g_j \to f$ a.e.
(ii) Letting $i \to \infty$ gives $j \ge k$, $x \notin F_k \Rightarrow |f(x) - g_j(x)| \le 2^{-(j-1)} \le 2^{-(k-1)}$. Given $\alpha, \varepsilon > 0$, choose k with $2^{-(k-1)} < \alpha$ and $\mu(F_k) < \varepsilon$; then $j \ge k \Rightarrow \{|f - g_j| \ge \alpha\} \subseteq F_k$, so $g_j \to f$ in measure.
(iii) Given $\varepsilon > 0$, choose k with $\mu(F_k) < \varepsilon$.
The estimate in (ii) gives $\sup_{x \in X \setminus F_k} |f(x) - g_j(x)| \le 2^{-(j-1)} \to 0$, so $g_j \to f$ uniformly on $X \setminus F_k$.

Structural Commentary: Cauchy in measure lets us select summably small exceptional sets.
Key Insight: Outside the tail union of exceptional sets, the subsequence is uniformly Cauchy. □

Corollary 5.9.12. *Completeness for Convergence in Measure*
(f_n) is Cauchy in measure iff $\exists$ measurable f with $f_n \to f$ in measure.

Proof.

($\Rightarrow$) If (f_n) is Cauchy in measure, Theorem 5.9.11 gives $f_{n_k} \to f$ in measure for some measurable f.
For $\alpha > 0$, $\{|f_n - f| \ge \alpha\} \subseteq \{|f_n - f_{n_k}| \ge \alpha/2\} \cup \{|f_{n_k} - f| \ge \alpha/2\}$.
Choose k large, then n large, to get $f_n \to f$ in measure.

($\Leftarrow$) For $\alpha > 0$, $\{|f_m - f_n| \geq \alpha\} \subseteq \{|f_m - f| \geq \alpha/2\} \cup \{|f_n - f| \geq \alpha/2\}$.
Hence $\mu\{|f_m - f_n| \geq \alpha\} \to 0$ as $m, n \to \infty$.

Structural Commentary: The subsequential in-measure limit attracts the whole Cauchy sequence.
Key Insight: Convergence in measure is complete after measurable limits are admitted. □

Corollary 5.9.13. *Uniqueness for Convergence in Measure*
If $f_n \to f$ and $f_n \to g$ in measure, then $f = g$ a.e.

Proof. For $\alpha > 0$, $\{|f - g| \geq \alpha\} \subseteq \{|f - f_n| \geq \alpha/2\} \cup \{|f_n - g| \geq \alpha/2\}$, so $\mu\{|f - g| \geq \alpha\} = 0$.
Hence $\mu\{|f - g| > 0\} = 0$ and $f = g$ a.e.
Structural Commentary: Two in-measure limits cannot differ on a positive-measure set.
Key Insight: In-measure limits are unique modulo null sets. □

Proposition 5.9.14. *Subsequence and Indicator Consequences*

(i) $f_n \to f$ in measure $\Rightarrow$ every subsequence $f_{n_k} \to f$ in measure; if (f_n) is Cauchy in measure, then every (f_{n_k}) is Cauchy in measure.
(ii) If $\|f_n - f\|_p \to 0$ and $\|f_{n_k} - g\|_p \to 0$, then $f = g$ a.e.
(iii) If $f_n = \mathbf{1}_{E_n}$ and $\|f_n - f\|_p \to 0$, then $f = \mathbf{1}_E$ a.e. for some $E \in \mathcal{M}$.

Proof.

(i) For $\alpha > 0$, $\mu\{|f_{n_k} - f| \geq \alpha\} \to 0$ since $n_k \to \infty$; the Cauchy claim is identical with $m_k, n_\ell \to \infty$.
(ii) Proposition 5.9.10 gives $f_n \to f$ and $f_{n_k} \to g$ in measure; also $f_{n_k} \to f$ in measure by (i). Uniqueness in Corollary 5.9.13 gives $f = g$ a.e.
(iii) By Proposition 5.9.10 and Theorem 5.9.11, some $f_{n_k} \to f$ a.e. Since $f_{n_k} \in \{0, 1\}$, $f \in \{0, 1\}$ a.e.; with $E := \{f = 1\}$, $E \in \mathcal{M}$ and $f = \mathbf{1}_E$ a.e.

Structural Commentary: Subsequence permanence and a.e. uniqueness transfer mode limits to representatives.
Key Insight: In-measure subsequences preserve limits; indicator limits remain indicators a.e. □

Theorem 5.9.15. *Dominated Convergence in Measure Gives L^p Convergence*
Let $1 \leq p < \infty$, $f_n \in L^p(\mu)$, $f_n \to f$ in measure, and $|f_n| \leq g$ a.e. for some $g \in L^p(\mu)$.
Then $f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$.

Proof. Since $f_n \to f$ in measure, (f_n) is Cauchy in measure.
By Theorem 5.9.11, some subsequence (u_j) satisfies $u_j \to h$ a.e. and in measure; uniqueness in Corollary 5.9.13 gives $h = f$ a.e. Since $|u_j| \leq g$ a.e., Theorem 5.9.4 gives $f \in L^p(\mu)$. If $\|f_n - f\|_p \not\to 0$, then $\exists \varepsilon > 0$ and a subsequence (g_k) with $\|g_k - f\|_p \geq \varepsilon$. Since $g_k \to f$ in measure, the same argument gives a subsubsequence (h_j) with $h_j \to f$ a.e.; Theorem 5.9.4 yields $\|h_j - f\|_p \to 0$, contradicting $\|h_j - f\|_p \geq \varepsilon$.
Structural Commentary: Reduce convergence in measure to almost-everywhere convergence along a subsequence.
Key Insight: Domination turns every in-measure subsequence into an L^p-convergent subsubsequence. □

Definition 5.9.16. *Almost Uniform Convergence*
Let (f_n) be measurable functions on X.

(i) $f_n \to f$ *almost uniformly* if $\forall \delta > 0\ \exists E_\delta \in \mathcal{M}$ with $\mu(E_\delta) < \delta$ and $f_n \rightrightarrows f$ on $X \setminus E_\delta$.
(ii) (f_n) is *almost uniformly Cauchy* if $\forall \delta > 0\ \exists E_\delta \in \mathcal{M}$ with $\mu(E_\delta) < \delta$ s.t. (f_n) uniformly Cauchy on $X \setminus E_\delta$.

Lemma 5.9.17. *Almost Uniform Cauchy Sequences Have Almost Uniform Limits*
If (f_n) is almost uniformly Cauchy, then $\exists$ measurable f such that $f_n \to f$ almost uniformly and a.e.

Proof. Choose $E_k \in \mathcal{M}$ with $\mu(E_k) < 2^{-k}$ and (f_n) uniformly Cauchy on $X \setminus E_k$. Set $F_k := \bigcup_{j=k}^{\infty} E_j$ and $F := \bigcap_k F_k$; then $F_k \downarrow$, $\mu(F_k) < 2^{-(k-1)}$, $\mu(F) = 0$, and (f_n) is uniformly Cauchy on $X \setminus F_k$. By Theorem 3.2.7, f_n has a uniform limit on each $X \setminus F_k$; these limits agree on overlaps, so define measurable f on $X \setminus F$, and set $f = 0$ on F. Hence $f_n \to f$ a.e. and uniformly on each $X \setminus F_k$; since $\mu(F_k) \to 0$, $f_n \to f$ almost uniformly.
Structural Commentary: Tail unions convert many small exceptional sets into nested uniform domains.
Key Insight: Almost uniform Cauchy means uniformly Cauchy after discarding a tail-small exceptional set. □

Theorem 5.9.18. *Almost Uniform Convergence and Convergence in Measure*
If $f_n \to f$ almost uniformly, then $f_n \to f$ in measure.
If $f_n \to f$ in measure, then $\exists (f_{n_k})$ with $f_{n_k} \to f$ almost uniformly.

Proof. If $f_n \rightrightarrows f$ on $X \setminus E_\delta$ with $\mu(E_\delta) < \delta$, then for each $\alpha > 0$ and all large n, $\{|f_n - f| \geq \alpha\} \subseteq E_\delta$; hence $\limsup_n \mu\{|f_n - f| \geq \alpha\} \leq \delta$, and $\delta \downarrow 0$ gives convergence in measure. Conversely, $f_n \to f$ in measure implies (f_n) is Cauchy in measure; Theorem 5.9.11 gives $f_{n_k} \to h$ almost uniformly and in measure. By uniqueness, $h = f$ a.e.; absorbing $\{h \neq f\}$ into the exceptional sets gives $f_{n_k} \to f$ almost uniformly.
Structural Commentary: Almost uniform convergence is uniform convergence off arbitrarily small sets.
Key Insight: Convergence in measure always hides an almost-uniform subsequence. □

Theorem 5.9.19. *Egoroff's Theorem*
If $\mu(X) < +\infty$ and $f_n \to f$ a.e., with all functions finite-valued off a null set, then $f_n \to f$ almost uniformly and in measure.

Proof. Absorb the null exceptional set into every later E_δ, and assume $f_n(x) \to f(x)$ for all $x \in X$.
For $m, n \in \mathbb{N}$, set $E_n(m) := \bigcup_{k=n}^{\infty} \{|f_k - f| \geq 1/m\}$. Then $E_n(m) \downarrow \emptyset$, so $\mu(E_n(m)) \to 0$ by Corollary 5.1.65.
Given $\delta > 0$, choose k_m with $\mu(E_{k_m}(m)) < \delta 2^{-m}$ and set $E_\delta := \bigcup_{m \geq 1} E_{k_m}(m)$.
Then $\mu(E_\delta) < \delta$, and $x \notin E_\delta$, $k \geq k_m$ imply $|f_k(x) - f(x)| < 1/m$; hence $f_n \rightrightarrows f$ on $X \setminus E_\delta$.
The in-measure conclusion follows from Theorem 5.9.18.
Structural Commentary: Finite measure turns decreasing bad sets with empty intersection into small bad sets.
Key Insight: A.e. convergence is almost uniform on finite measure spaces. □

Corollary 5.9.20. *Finite-Measure Almost-Everywhere Convergence Implies In-Measure Convergence*
If $\mu(X) < +\infty$ and $f_n \to f$ a.e., with f_n, f finite-valued off a null set, then $f_n \to f$ in measure.

Proof. By Theorem 5.9.19, $f_n \to f$ almost uniformly; Theorem 5.9.18 gives $f_n \to f$ in measure.
Structural Commentary: Finite measure upgrades a.e. convergence through almost uniform convergence.
Key Insight: On finite measure spaces, a.e. convergence implies in-measure convergence. □

Example 5.9.21. *Egoroff Needs a Finite-Valued Limit*
On $([0,1], \lambda)$, set $f_n := n\mathbf{1}_{[0,1]}$. Then $f_n(x) \to +\infty$ for all x, so the pointwise limit is not finite-valued. If $f_n \to f$ almost uniformly for finite-valued measurable f, then for some E with $\lambda(E) < 1/2$, $f_n \rightrightarrows f$ on $[0,1] \setminus E$; this nonempty set has $f_n(x) = n \to +\infty$, contradicting finite-valued pointwise convergence.

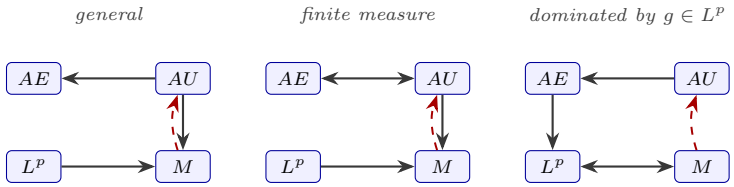


Figure 203: Modes of convergence. Solid arrows are implications; dashed arrows give a convergent subsequence.

Method 5.9.22. *Counterexample Dictionary*

(i) On $(\mathbb{R}, \lambda)$, $f_n = n^{-1/p}\mathbf{1}_{[0,n]}$ satisfies $f_n \rightrightarrows 0$ but $\|f_n\|_p = 1$.
(ii) On $([0,1], \lambda)$, $f_n = n\mathbf{1}_{[1/n,2/n]}$ satisfies $f_n \to 0$ a.e., in measure, almost uniformly, but $\|f_n\|_p^p = n^{p-1} \nrightarrow 0$.
(iii) Typewriter intervals on $[0,1]$: if $f_n = \mathbf{1}_{I_n}$ and $\lambda(I_n) \to 0$ rowwise, then $\|f_n\|_p \to 0$ and $f_n \to 0$ in measure, but $(f_n(x))$ converges nowhere.
(iv) On $(\mathbb{R}, \lambda)$, $f_n = \mathbf{1}_{[n,n+1]}$ satisfies $f_n \to 0$ pointwise but $\lambda\{|f_n| \geq 1/2\} = 1$ for all n, so $f_n \nrightarrow 0$ in measure.
(v) On $(\mathbb{R}, \lambda)$, $f_n = \mathbf{1}_{[n,n+1]}$ is not almost uniformly convergent to 0: if $\lambda(E) < 1$ and $f_n \rightrightarrows 0$ on $\mathbb{R} \setminus E$, then $[n, n+1] \subseteq E$ for all large n, impossible.

Proposition 5.9.23. *An Integral Metric for Convergence in Measure on Finite Measure Spaces* *Let $\mu(X) < +\infty$ and $r(h) := \int_X |h|\,(1+|h|)^{-1}\,d\mu$. Then $f_n \to f$ in measure iff $r(f_n - f) \to 0$.*

Proof.
($\Rightarrow$) Set $h_n := f_n - f$. For $\eta > 0$ and $A_n := \{|h_n| \geq \eta\}$, $r(h_n) \leq \mu(A_n) + \eta\mu(X)$.
Let $n \to \infty$, then $\eta \downarrow 0$.
($\Leftarrow$) For $\alpha > 0$, $\alpha(1+\alpha)^{-1}\mathbf{1}_{\{|h_n| \geq \alpha\}} \leq |h_n|\,(1+|h_n|)^{-1}$, hence $\mu\{|h_n| \geq \alpha\} \leq (1+\alpha)\alpha^{-1} r(h_n) \to 0$.
Structural Commentary: The bounded error transform metrizes convergence in measure on finite measure spaces.
Key Insight: Small integral of the bounded error transform is equivalent to small large-error sets. □

Corollary 5.9.24. *Fatou and Dominated Convergence with In-Measure Hypotheses*

(i) $f_n \geq 0$ and $f_n \to f$ in measure $\Rightarrow \int_X f\,d\mu \leq \liminf_n \int_X f_n\,d\mu$.

(ii) $f_n \to f$ *in measure and* $|f_n| \le g \in \mathscr{L}^1(\mu) \Rightarrow \int_X f_n\,d\mu \to \int_X f\,d\mu$.

Proof.

(i) Choose f_{n_k} with $\int f_{n_k}\,d\mu \to \liminf_n \int f_n\,d\mu$. By Theorem 5.9.11 and uniqueness, some $f_{n_{k_j}} \to f$ a.e.; Fatou gives $\int f\,d\mu \le \liminf_j \int f_{n_{k_j}}\,d\mu = \liminf_n \int f_n\,d\mu$.

(ii) Theorem 5.9.15 with $p = 1$ gives $\|f_n - f\|_1 \to 0$, hence $\left|\int_X (f_n - f)\,d\mu\right| \le \|f_n - f\|_1 \to 0$.

Structural Commentary: In-measure hypotheses reduce to a.e. subsequences or L^1 convergence.
Key Insight: Fatou and DCT survive convergence in measure through subsequences and domination. □

Theorem 5.9.25. *Vitali Convergence Criterion for L^p*
Let $1 \le p < \infty$, $f_n \in L^p(\mu)$, and let f be measurable. The following are equivalent:

(i) $f \in L^p(\mu)$ *and* $\|f_n - f\|_p \to 0$.
(ii) The following hold:
  (a) $f_n \to f$ *in measure.*
  (b) $\forall \varepsilon > 0\ \exists E_\varepsilon \in \mathcal{M}$ *with* $\mu(E_\varepsilon) < +\infty$ *such that* $F \cap E_\varepsilon = \emptyset \Rightarrow \sup_n \int_F |f_n|^p\,d\mu < \varepsilon^p$.
  (c) $\forall \varepsilon > 0\ \exists \delta(\varepsilon) > 0$ *such that* $\mu(E) < \delta(\varepsilon) \Rightarrow \sup_n \int_E |f_n|^p\,d\mu < \varepsilon^p$.

Proof.

(i) Assume (i). Proposition 5.9.10 gives (ii)(a). Also (f_n) is L^p-Cauchy.
By Proposition 5.8.32, for each $\varepsilon > 0$ there are $E_\varepsilon \in \mathcal{M}$ with $\mu(E_\varepsilon) < +\infty$ and $\delta(\varepsilon) > 0$ such that $F \cap E_\varepsilon = \emptyset \Rightarrow \sup_n \|f_n \mathbf{1}_F\|_p < \varepsilon$ and $\mu(E) < \delta(\varepsilon) \Rightarrow \sup_n \|f_n \mathbf{1}_E\|_p < \varepsilon$; raising to p gives (ii)(b),(ii)(c).

(ii) Assume (ii). Fix $\eta > 0$, choose E_η from (ii)(b) and $\delta(\eta)$ from (ii)(c), and put $F = X \setminus E_\eta$.
Then $\|(f_n - f_m)\mathbf{1}_F\|_p \le \|f_n \mathbf{1}_F\|_p + \|f_m \mathbf{1}_F\|_p < 2\eta$.
If $\mu(E_\eta) > 0$, set $\alpha := \eta \mu(E_\eta)^{-1/p}$ and $H_{nm} := E_\eta \cap \{|f_n - f_m| \ge \alpha\}$.
By (ii)(a), (f_n) is Cauchy in measure, so $\mu(H_{nm}) < \delta(\eta)$ for large m, n. On $E_\eta \setminus H_{nm}$, $|f_n - f_m| < \alpha$; on H_{nm}, use $|f_n - f_m| \le |f_n| + |f_m|$. Hence $\left\|(f_n - f_m)\mathbf{1}_{E_\eta}\right\|_p \le \alpha \mu(E_\eta)^{1/p} + \|f_n \mathbf{1}_{H_{nm}}\|_p + \|f_m \mathbf{1}_{H_{nm}}\|_p < 3\eta$ by (ii)(c). If $\mu(E_\eta) = 0$, the E_η term is 0. Thus $\|f_n - f_m\|_p < 5\eta$ eventually, so (f_n) is L^p-Cauchy.
Riesz–Fischer gives $f_n \to h$ in L^p; Proposition 5.9.10 and Corollary 5.9.13 give $h = f$ a.e., so (i) holds.

Structural Commentary: Vitali = in-measure convergence + tightness + uniform small-set control.
Key Insight: Tightness and small-set control convert in-measure convergence into an L^p Cauchy estimate. □

Corollary 5.9.26. *Domination Gives Tail and Small-Set Control*
Let $1 \le p < \infty$, $g \in L^p(\mu)$, and $|f_n| \le g$ a.e. for all n. Then, for every $\varepsilon > 0$:

(i) $\exists E_\varepsilon \in \mathcal{M}$ *with* $\mu(E_\varepsilon) < +\infty$ *such that* $F \cap E_\varepsilon = \emptyset \Rightarrow \sup_n \int_F |f_n|^p\,d\mu < \varepsilon^p$.
(ii) $\exists \delta(\varepsilon) > 0$ *such that* $\mu(E) < \delta(\varepsilon) \Rightarrow \sup_n \int_E |f_n|^p\,d\mu < \varepsilon^p$.

Proof.

(i) Proposition 5.8.20 gives E_ε with $\mu(E_\varepsilon) < +\infty$ and $F \cap E_\varepsilon = \emptyset \Rightarrow \int_F |g|^p\,d\mu < \varepsilon^p$.
Since $|f_n|^p \le |g|^p$ a.e., $\sup_n \int_F |f_n|^p\,d\mu < \varepsilon^p$.

(ii) Lemma 5.8.30 gives $\delta(\varepsilon) > 0$ with $\mu(E) < \delta(\varepsilon) \Rightarrow \int_E |g|^p\,d\mu < \varepsilon^p$.
Since $|f_n|^p \le |g|^p$ a.e., $\sup_n \int_E |f_n|^p\,d\mu < \varepsilon^p$.

Structural Commentary: One L^p envelope supplies both tail tightness and small-set control.
Key Insight: Domination transfers both estimates from g to every f_n. □

Proposition 5.9.27. *Composition and Pointwise Modes*
Let $\varphi : \mathbb{R} \to \mathbb{R}$.

(i) φ *continuous,* $f_n \to f$ *a.e.* $\Rightarrow \varphi \circ f_n \to \varphi \circ f$ *a.e.*
(ii) φ *uniformly continuous preserves uniform, a.u., and in-measure convergence.*
(iii) If φ *is discontinuous, composition need not preserve a.e. convergence.*
(iv) Non-uniform continuity can destroy uniform, even in-measure, convergence.

Proof.

(i) Outside a null set, $f_n(x) \to f(x)$; continuity gives $\varphi(f_n(x)) \to \varphi(f(x))$.

(ii) Given $\varepsilon > 0$, choose $\delta > 0$ with $|s - t| < \delta \Rightarrow |\varphi(s) - \varphi(t)| < \varepsilon$. Uniform and almost-uniform convergence use $\sup |f_n - f| < \delta$ on X or $X \setminus E_\eta$; in measure, $\{|\varphi(f_n) - \varphi(f)| \ge \varepsilon\} \subseteq \{|f_n - f| \ge \delta\}$.

(iii) If φ is discontinuous at a, choose $t_n \to a$ and $\alpha > 0$ with $|\varphi(t_n) - \varphi(a)| \ge \alpha$.
On $X = \{*\}$, $f_n(*) = t_n \to a = f(*)$, but $\varphi(f_n(*)) \not\to \varphi(f(*))$.

(iv) If φ is not uniformly continuous, choose a_n, b_n with $|a_n - b_n| \to 0$ and $|\varphi(a_n) - \varphi(b_n)| \geq \alpha > 0$.
On $(\mathbb{N}, \#)$, set $f(k) = b_k$ and $f_n(k) = a_n$ for $k = n$, $f_n(k) = b_k$ for $k \neq n$.
Then $\|f_n - f\|_\infty \to 0$, but $\#\{|\varphi(f_n) - \varphi(f)| \geq \alpha\} = 1$ for all n.

Structural Commentary: Continuity is pointwise; uniform continuity gives one global δ.
Key Insight: Composition needs the same uniformity as the convergence mode. □

Proposition 5.9.28. *Composition and Convergence in Measure*
Let $f_n, f : X \to \mathbb{R}$ be measurable with $f_n \to f$ in measure, and let $\varphi : \mathbb{R} \to \mathbb{R}$.

(i) If φ is uniformly continuous, then $\mu\{|\varphi(f_n) - \varphi(f)| \geq \varepsilon\} \to 0$ for every $\varepsilon > 0$.
(ii) If $\mu(X) < +\infty$ and φ is continuous, then $\mu\{|\varphi(f_n) - \varphi(f)| \geq \varepsilon\} \to 0$ for every $\varepsilon > 0$.

Proof.

(i) Fix $\varepsilon > 0$ and choose $\delta > 0$ with $|s - t| < \delta \Rightarrow |\varphi(s) - \varphi(t)| < \varepsilon$.
Then $\{|\varphi(f_n) - \varphi(f)| \geq \varepsilon\} \subseteq \{|f_n - f| \geq \delta\}$, so the measures tend to 0.

(ii) Fix $\varepsilon, \eta > 0$. Since $\mu(X) < +\infty$ and f is finite-valued, choose M with $\mu\{|f| > M\} < \eta$. Uniform continuity of φ on $[-M-1, M+1]$ gives $0 < \delta < 1$ such that $s, t \in [-M-1, M+1]$ and $|s - t| < \delta \Rightarrow |\varphi(s) - \varphi(t)| < \varepsilon$.
Hence $\{|\varphi(f_n) - \varphi(f)| \geq \varepsilon\} \subseteq \{|f| > M\} \cup \{|f_n - f| \geq \delta\}$, so $\limsup_n \mu\{|\varphi(f_n) - \varphi(f)| \geq \varepsilon\} \leq \eta \downarrow 0$.

Structural Commentary: Use global uniformity, or truncate to a bounded range.
Key Insight: On finite measure spaces, measurable finite-valued functions are almost bounded. □

Example 5.9.29. *Continuity Alone Is Not Enough on Infinite Measure*
On $(\mathbb{R}, \mathcal{B}(\mathbb{R}), \lambda)$, set $f(x) = x$, $f_n(x) = x + 1/n$, and $\varphi(t) = t^2$.
Then $\|f_n - f\|_\infty = 1/n \to 0$, hence $f_n \rightrightarrows f$ and $f_n \to f$ in measure.
However $\varphi(f_n) - \varphi(f) = 2x/n + 1/n^2$; for $E_n := \{|\varphi(f_n) - \varphi(f)| \geq 1\}$, $[n/2, +\infty) \subseteq E_n$, so $\lambda(E_n) = +\infty$.
Thus $\varphi \circ f_n \not\to \varphi \circ f$ in measure.

Proposition 5.9.30. *Growth Conditions for L^p Composition*
Let $1 \leq p < \infty$, let $\varphi : \mathbb{R} \to \mathbb{R}$ be measurable, and let $f_n, f : X \to \mathbb{R}$ be measurable.

(i) If $|\varphi(t)| \leq A|t| + B$ for all t and some $A, B \geq 0$, $\mu(X) < +\infty$, and $f \in L^p(\mu)$, then $\varphi \circ f \in L^p(\mu)$.
(ii) If $|\varphi(t)| \leq A|t|$ for all t and some $A \geq 0$, and $f \in L^p(\mu)$, then $\varphi \circ f \in L^p(\mu)$.
(iii) If $\mu(X) < +\infty$, φ is uniformly continuous, $f_n, f \in L^p(\mu)$, and $\|f_n - f\|_p \to 0$, then $\|\varphi \circ f_n - \varphi \circ f\|_p \to 0$.
(iv) If φ is Lipschitz, $\varphi(0) = 0$, $f_n, f \in L^p(\mu)$, and $\|f_n - f\|_p \to 0$, then $\|\varphi \circ f_n - \varphi \circ f\|_p \to 0$.

Proof.

(i) $|\varphi(f)|^p \leq 2^{p-1}(A^p |f|^p + B^p)$, so $\|\varphi \circ f\|_p^p \leq 2^{p-1}(A^p \|f\|_p^p + B^p \mu(X)) < +\infty$.

(ii) $|\varphi(f)|^p \leq A^p |f|^p$, so $\|\varphi \circ f\|_p \leq A \|f\|_p < +\infty$.

(iii) Uniform continuity gives constants $C, D \geq 0$ with $|\varphi(t)| \leq C|t| + D$ and $|\varphi(s) - \varphi(t)| \leq C|s - t| + D$.
Hence all compositions lie in L^p by (i). Let $h_n := |f_n - f|$.
Given $\eta > 0$, choose $\delta > 0$ with $|s - t| < \delta \Rightarrow |\varphi(s) - \varphi(t)| < \eta$ and set $H_n := \{h_n \geq \delta\}$.
Then $\int_X |\varphi(f_n) - \varphi(f)|^p \, d\mu \leq \eta^p \mu(X) + 2^{p-1} \int_{H_n} (C^p h_n^p + D^p) \, d\mu$.
Since $\int h_n^p \, d\mu \to 0$ and $\mu(H_n) \leq \delta^{-p} \int h_n^p \, d\mu \to 0$, take $\limsup_n$ and then $\eta \downarrow 0$.

(iv) If L is Lipschitz, $|\varphi(f)| \leq L|f|$ and $\|\varphi \circ f_n - \varphi \circ f\|_p \leq L \|f_n - f\|_p \to 0$.

Structural Commentary: Growth gives membership; continuity estimates transfer L^p convergence.
Key Insight: Finite measure makes small uniform errors integrable; Lipschitz bounds give norm control. □

Example 5.9.31. *Uniform Continuity Does Not Suffice on Infinite Measure*
On $(\mathbb{R}, \lambda)$, fix $1 \leq p < \infty$, set $f = 0$, $f_n := n^{-1} \mathbf{1}_{[0, n^{p/2}]}$, and $\varphi(t) := \sqrt{|t|}$.
Then φ is uniformly continuous, $\varphi(0) = 0$, $\|f_n - f\|_p^p = n^{-p/2} \to 0$, but $\|\varphi \circ f_n - \varphi \circ f\|_p^p = 1$ for all n.
Thus uniform continuity need not preserve L^p convergence on infinite measure spaces.

Lemma 5.9.32. *L^p Convergence Allows Continuous Composition in Measure*
Let $1 \leq p < \infty$, $f_n, f \in L^p(\mu)$, $\|f_n - f\|_p \to 0$, and let $\varphi : \mathbb{R} \to \mathbb{R}$ be continuous.
Then $\varphi \circ f_n \to \varphi \circ f$ in measure.

Proof. Fix $\alpha, \eta > 0$. Choose M with $\mu\{|f| > M\} < \eta$. By uniform continuity of φ on $[-M - 1, M + 1]$, choose $0 < \delta < 1$ such that $s, t \in [-M - 1, M + 1]$ and $|s - t| < \delta \Rightarrow |\varphi(s) - \varphi(t)| < \alpha$.
Then $\{|\varphi(f_n) - \varphi(f)| \geq \alpha\} \subseteq \{|f| > M\} \cup \{|f_n - f| \geq \delta\}$, so $\limsup_n \mu\{|\varphi(f_n) - \varphi(f)| \geq \alpha\} \leq \eta$. Let $\eta \downarrow 0$.
Structural Commentary: L^p convergence localizes to a large bounded range where continuity is uniform.
Key Insight: Norm convergence gives in-measure control; bounded-range continuity transfers it. □

Theorem 5.9.33. *Composition Growth Criteria in L^p*
Let $1 \le p < \infty$ and let $\varphi : \mathbb{R} \to \mathbb{R}$ be continuous.

(i) If $\mu(X) < +\infty$ and $\exists K > 0$ such that $|\varphi(t)| \le K\,|t|$ for $|t| \ge K$, then $f \in L^p(\mu) \Rightarrow \varphi \circ f \in L^p(\mu)$.
(ii) Under (i), if $f_n, f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$, then $\|\varphi \circ f_n - \varphi \circ f\|_p \to 0$.
(iii) If $\exists K \ge 0$ such that $|\varphi(t)| \le K\,|t|$ for all t, then $f \in L^p(\mu) \Rightarrow \varphi \circ f \in L^p(\mu)$ on any measure space.
(iv) Under (iii), if $f_n, f \in L^p(\mu)$ and $\|f_n - f\|_p \to 0$, then $\|\varphi \circ f_n - \varphi \circ f\|_p \to 0$ on any measure space.

Proof.

(i) Let $M := \sup_{|t| \le K} |\varphi(t)|$ and $C := K + M$.
Then $|\varphi(t)| \le C(1 + |t|)$ for all t. hence $|\varphi(f)|^p \le 2^{p-1}C^p(1 + |f|^p) \in L^1(\mu)$.
(ii) Let C be as in (i). Lemma 5.9.32 gives $\varphi \circ f_n \to \varphi \circ f$ in measure.
Since (f_n) is L^p-Cauchy, Proposition 5.8.32 and $|\varphi(f_n)|^p \le 2^{p-1}C^p(1 + |f_n|^p)$ give the finite-tail estimate with $E_\varepsilon = X$ and the small-set estimate. Theorem 5.9.25 gives $\|\varphi \circ f_n - \varphi \circ f\|_p \to 0$.
(iii) $|\varphi(f)|^p \le K^p\,|f|^p \in L^1(\mu)$.
(iv) Lemma 5.9.32 gives $\varphi \circ f_n \to \varphi \circ f$ in measure. Since (f_n) is L^p-Cauchy and $|\varphi(f_n)|^p \le K^p\,|f_n|^p$, Proposition 5.8.32 gives the finite-tail and small-set estimates for $(\varphi \circ f_n)$. Theorem 5.9.25 gives $\|\varphi \circ f_n - \varphi \circ f\|_p \to 0$.

Structural Commentary: Local growth needs finite measure; global growth does not.
Key Insight: Vitali converts in-measure composition plus inherited mass control into L^p convergence. □

Method 5.9.34. *Atomic Counterexamples for Composition Growth*

(i) If finite-measure tail growth fails, choose $|t_n| \ge n$ with $|\varphi(t_n)| > n^2\,|t_n|$. On $\mathbb{N}$, set $\mu(\{n\}) = (n^2\,|t_n|^p)^{-1}$ and $f(n) = t_n$. Then $\mu(\mathbb{N}) < +\infty$, $\int |f|^p\,d\mu = \sum_n n^{-2} < +\infty$, but $\int |\varphi(f)|^p\,d\mu \ge \sum_n n^{2p-2} = +\infty$.
(ii) If the same tail growth fails, choose $|t_n| \ge n$ with $|\varphi(t_n) - \varphi(0)| \ge n\,|t_n|$. With $\mu(\{n\}) = (n^p\,|t_n|^p)^{-1}$, $f = 0$, and $f_n = t_n \mathbf{1}_{\{n\}}$, one has $\mu(\mathbb{N}) < +\infty$, $\|f_n - f\|_p = 1/n \to 0$, but $\|\varphi \circ f_n - \varphi \circ f\|_p \ge 1$.
(iii) If global linear growth fails and $\varphi(0) = 0$, the constructions in (i)–(ii), without requiring $\mu(\mathbb{N}) < +\infty$, give membership and convergence counterexamples on arbitrary measure spaces.
(iv) If $\varphi(0) \ne 0$, then on $(\mathbb{N}, \#)$, $0 \in L^p$ but $\varphi \circ 0$ is a nonzero constant and hence not in L^p; the constant sequence $f_n = 0$ gives the convergence failure.

6 Functional and Operator Analysis

Functional and Operator Analysis treats infinite-dimensional linear algebra through Banach spaces, Hilbert spaces, duality, weak topologies, compact operators, spectral theory, and distributions.

6.1 Normed and Banach Spaces

Normed spaces make vector spaces metric; Banach spaces add completeness. This is the setting for infinite series, fixed points, bounded linear maps, and operator norms.

Definition 6.1.1. *Normed and Banach Spaces*
Let $\mathcal{X}$ be a vector space over $\mathbb{F} \in \{\mathbb{R}, \mathbb{C}\}$. A *norm* is a function $\|\cdot\| : \mathcal{X} \to [0, \infty)$ such that, for all $x, y \in \mathcal{X}$ and $\lambda \in \mathbb{F}$, $\|x\| = 0 \iff x = 0$; $\|\lambda x\| = |\lambda|\,\|x\|$; and $\|x + y\| \le \|x\| + \|y\|$.
The metric $d(x, y) := \|x - y\|$ makes $\mathcal{X}$ a *normed space*. If this metric is complete, $\mathcal{X}$ is a *Banach space*.

Definition 6.1.2. *Bounded Linear Maps and Operator Norm*
For normed spaces $\mathcal{X}, \mathcal{Y}$, a linear map $T : \mathcal{X} \to \mathcal{Y}$ is *bounded* if there exists $C \ge 0$ such that $\|Tx\|_{\mathcal{Y}} \le C\,\|x\|_{\mathcal{X}}$ for all $x \in \mathcal{X}$. The *operator norm* is $\|T\|_{\mathrm{op}} := \sup_{\|x\|_{\mathcal{X}} \le 1} \|Tx\|_{\mathcal{Y}} = \sup_{x \ne 0} \frac{\|Tx\|_{\mathcal{Y}}}{\|x\|_{\mathcal{X}}}$.

Theorem 6.1.3. *Bounded Equals Continuous for Linear Maps*
For a linear map $T : \mathcal{X} \to \mathcal{Y}$ between normed spaces, the following are equivalent:

(i) T is bounded.
(ii) T is continuous at 0.
(iii) T is continuous everywhere.

Proof.

(i) $\Rightarrow$ *(iii)*. If T is bounded, then $\|T(x - y)\| \le \|T\|_{\mathrm{op}}\,\|x - y\|$, so T is continuous everywhere.
(iii) $\Rightarrow$ *(ii)*. Since $0 \in \mathcal{X}$, continuity everywhere gives continuity at 0.
(ii) $\Rightarrow$ *(i)*. Choose $\delta > 0$ such that $\|x\| < \delta$ implies $\|Tx\| < 1$.
For $x \ne 0$, apply this to $\delta x/(2\,\|x\|)$ to get $\|Tx\| \le 2\delta^{-1}\,\|x\|$; the case $x = 0$ is trivial.

Structural Commentary: Linearity turns local continuity at one point into uniform norm control.
Key Insight: In normed linear spaces, continuity and boundedness are the same operator property. □

Example 6.1.4. *Standard Banach Space Examples*

(i) Every complete supremum space from Theorem 3.3.1 is Banach under $\|\cdot\|_\infty$.
(ii) The sequence spaces ℓ^∞ and ℓ^p, $1 \le p < \infty$, are Banach by Example 2.2.17.
(iii) The Lebesgue spaces $L^p(\mu)$, $1 \le p < \infty$, and $L^\infty(\mu)$ are Banach by Theorems 5.8.17 and 5.8.19.
(iv) In particular, L^2 is a Hilbert space by Theorem 5.8.8.

6.2 Hilbert Spaces

Hilbert spaces are complete inner-product spaces. They extend the orthogonal geometry of Definition 2.1.3 to function and sequence spaces, supporting projection, Fourier expansion, and Riesz representation.

Definition 6.2.1. *Hilbert Space*
An inner-product space $\mathcal{H}$ is a *Hilbert space* if the norm $\|x\| := \sqrt{\langle x, x\rangle}$ makes $\mathcal{H}$ complete.
A subset $M \subseteq \mathcal{H}$ is *orthonormal* if $\langle e_i, e_j\rangle = \delta_{ij}$ for all indexed vectors in M.

Proposition 6.2.2. *Continuity of the Inner Product*
Let $\mathcal{V}$ be an inner-product space. For sequences $(x_n), (y_n) \subseteq \mathcal{V}$ and $x, y \in \mathcal{V}$,
$\|x_n - x\| \to 0$, $\|y_n - y\| \to 0 \implies \langle x_n, y_n\rangle \to \langle x, y\rangle$, $x_n \to x \implies \|x_n\| \to \|x\|$.

Proof.
(i) $|\|x_n\| - \|x\|| \le \|x_n - x\| \to 0$, hence $\|x_n\| \to \|x\|$.
(ii) Since $y_n \to y$, choose $M < \infty$ with $\|y_n\| \le M$; Cauchy–Schwarz gives $|\langle x_n, y_n\rangle - \langle x, y\rangle| \le |\langle x_n - x, y_n\rangle| + |\langle x, y_n - y\rangle| \le M\|x_n - x\| + \|x\|\|y_n - y\| \to 0$.

Structural Commentary: Cauchy–Schwarz turns norm convergence in either coordinate into scalar convergence.
Key Insight: Inner products are continuous because their errors are controlled by products of norms. □

Theorem 6.2.3. *Hilbert Completion*
For every inner-product space $\mathcal{V}$, $\exists$ Hilbert space $\widehat{\mathcal{V}}$, linear isometry $J : \mathcal{V} \to \widehat{\mathcal{V}}$ such that $\overline{J(\mathcal{V})} = \widehat{\mathcal{V}}$.
If $(\mathcal{K}, I)$ is another such pair, then $\exists!$ surjective linear isometry $U : \widehat{\mathcal{V}} \to \mathcal{K}$ with $U \circ J = I$.

Proof. Let $C(\mathcal{V})$ be Cauchy sequences, $(x_n) \sim (y_n) \iff \|x_n - y_n\| \to 0$, $\widehat{\mathcal{V}} := C(\mathcal{V})/\sim$.
Set $[x_n] + [y_n] := [x_n + y_n]$, $a[x_n] := [ax_n]$, $Jx := [(x)]$, and $\langle [x_n], [y_n]\rangle_{\widehat{\mathcal{V}}} := \lim_{n\to\infty} \langle x_n, y_n\rangle_{\mathcal{V}}$.
As Cauchy sequences bounded, $|\langle x_n, y_n\rangle - \langle x_m, y_m\rangle| \le M\|x_n - x_m\| + N\|y_n - y_m\|$;
and $|\langle x_n, y_n\rangle - \langle x'_n, y'_n\rangle| \le M\|x_n - x'_n\| + N\|y_n - y'_n\|$ for suitable $M, N < \infty$, hence limit exists, and is representative-independent. Sesquilinearity and positivity pass to limits, and $\|[x_n] - [y_n]\| = \lim_{n\to\infty}\|x_n - y_n\|$.
Thus induced metric is the completion metric, so $\widehat{\mathcal{V}}$ is Hilbert, J is a linear isometry, and $\overline{J(\mathcal{V})} = \widehat{\mathcal{V}}$.
For another pair $(\mathcal{K}, I)$, metric-completion uniqueness gives a unique surjective isometry $U : \widehat{\mathcal{V}} \to \mathcal{K}$ with $U \circ J = I$; density gives $U(ax + by) = aU(x) + bU(y)$, and Theorem 2.1.23 gives $\langle Ux, Uy\rangle_{\mathcal{K}} = \langle x, y\rangle_{\widehat{\mathcal{V}}}$.
Structural Commentary: Complete the metric first, then recover the inner product by taking limits.
Key Insight: Cauchy–Schwarz makes the lifted inner product well defined. □

Theorem 6.2.4. *Complete, Closed, and Finite-Dimensional Subspaces*
Let $\mathcal{S}$ be a subspace of an inner-product space $\mathcal{V}$.

(i) If $\mathcal{S}$ is complete in the inherited metric, then $\mathcal{S}$ is closed in $\mathcal{V}$.
(ii) If $\mathcal{V}$ is Hilbert, then $\mathcal{S}$ is Hilbert if and only if $\mathcal{S}$ is closed.
(iii) If $\mathcal{S}$ is finite-dimensional, then $\mathcal{S}$ is complete and closed in $\mathcal{V}$.

Proof.
(i) Let $(s_n) \subseteq \mathcal{S}$ and $s_n \to v \in \mathcal{V}$.
Then (s_n) is Cauchy in $\mathcal{S}$, so $\exists s \in \mathcal{S} : s_n \to s$, $v = s \in \mathcal{S}$ by uniqueness, and $\mathcal{S}$ is closed.
(ii) Assume $\mathcal{V}$ is Hilbert; if $\mathcal{S}$ is closed and $(s_n) \subseteq \mathcal{S}$ is Cauchy.
Then $\exists v \in \mathcal{V} : s_n \to v$, closedness gives $v \in \mathcal{S}$, and the converse follows from (i).
(iii) Let $\dim \mathcal{S} = m$ and choose an orthonormal basis $e_1, \dots, e_m$.
If $(s_n) \subseteq \mathcal{S}$ is Cauchy and $s_n = \sum_{j=1}^m a_{n,j} e_j$, then $|a_{n,j} - a_{k,j}| = |\langle s_n - s_k, e_j\rangle| \le \|s_n - s_k\|$.
So $a_{n,j} \to a_j \in \mathbb{F}$ and $s_n \to \sum_{j=1}^m a_j e_j \in \mathcal{S}$, hence $\mathcal{S}$ is complete and closed by (i).

Structural Commentary: Closedness is the ambient form of internal completeness.
Key Insight: Finite orthonormal coordinates reduce completeness to completeness of $\mathbb{F}$. □

Definition 6.2.5. *Series in an Inner-Product Space*
Let $(x_k) \subseteq \mathcal{V}$.

(i) The *partial sums* are $s_n := \sum_{k=1}^{n} x_k$.
(ii) The *series* $\sum_{k=1}^{\infty} x_k$ *converges* to $s \in \mathcal{V}$ iff $s_n \to s$.
(iii) The series is *absolutely convergent* iff $\sum_{k=1}^{\infty} \|x_k\| < \infty$.

Theorem 6.2.6. *Completeness via Absolutely Convergent Series*
$\mathcal{V}$ is complete $\iff$ every series $\sum_{k=1}^{\infty} x_k$ in $\mathcal{V}$ with $\sum_{k=1}^{\infty} \|x_k\| < \infty$ converges in $\mathcal{V}$.

Proof.
($\Rightarrow$) Let $s_n = \sum_{k=1}^{n} x_k$ and $\sum_{k=1}^{\infty} \|x_k\| < \infty$.
For $n > m$, $\|s_n - s_m\| \le \sum_{k=m+1}^{n} \|x_k\| \to 0$, so (s_n) is Cauchy and hence convergent.
($\Leftarrow$) Let (x_n) be Cauchy and choose $N_1 < N_2 < \cdots$ with $i, j \ge N_k \Longrightarrow \|x_i - x_j\| < 2^{-k}$.
Then $\sum_{k=1}^{\infty} \|x_{N_{k+1}} - x_{N_k}\| \le \sum_{k=1}^{\infty} 2^{-k} < \infty$, so $\sum_{k=1}^{\infty} (x_{N_{k+1}} - x_{N_k})$ converges, whence (x_{N_k}) converges, and Lemma 2.2.11 gives $x_n \to x$.

Structural Commentary: Absolute convergence controls tails; subsequences recover Cauchy limits.
Key Insight: Completeness is exactly the ability to sum every norm-controlled vector tail. □

Example 6.2.7. *ℓ^2 and the Parallelogram Test*
Let $\ell^2 := \{x = (x_n) : \sum_{n=1}^{\infty} |x_n|^2 < \infty\}$. With $\langle x, y\rangle := \sum_{n=1}^{\infty} x_n \overline{y_n}$ and $\|x\|_2 = \langle x, x\rangle^{1/2}$, Cauchy coordinate limits and finite-set estimates give completeness, so ℓ^2 is Hilbert. For $p \neq 2$, $\|\cdot\|_p$ is not inner-product-induced: with $x = (1, 1, 0, \ldots)$ and $y = (1, -1, 0, \ldots)$, $\|x + y\|_p^2 + \|x - y\|_p^2 = 8$ while $2\|x\|_p^2 + 2\|y\|_p^2 = 4 \cdot 2^{2/p}$, so the parallelogram identity holds only for $p = 2$.
Structural Commentary: Completeness alone does not make a Banach space Hilbert.
Key Insight: The parallelogram identity detects whether a norm really came from an inner product.

Example 6.2.8. *The Supremum Norm Is Not Hilbertian*
For $a < b$, $C([a, b], \mathbb{F})$ with $\|h\|_\infty := \sup_{t \in [a,b]} |h(t)|$ is Banach but not Hilbertian: take $f(t) = 1$ and $g(t) = (t - a)/(b - a)$, so $\|f\|_\infty = \|g\|_\infty = 1$, $\|f + g\|_\infty = 2$, $\|f - g\|_\infty = 1$, and $\|f + g\|_\infty^2 + \|f - g\|_\infty^2 = 5 \neq 4 = 2\|f\|_\infty^2 + 2\|g\|_\infty^2$, violating the parallelogram identity.
Structural Commentary: Banach need not mean Hilbert.
Key Insight: One failed parallelogram identity rules out every possible inner product inducing the norm.

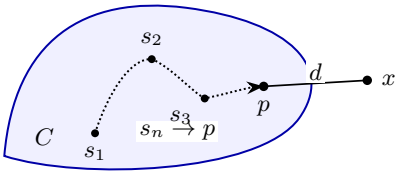


Figure 204: A minimizing sequence in C converges to the nearest point p to x.

Theorem 6.2.9. *Hilbert Projection Theorem*
Let $C \subseteq \mathcal{H}$ be nonempty, closed, and convex. For every $x \in \mathcal{H}$ there is a unique $p \in C$ minimizing $\|x - p\|$. If $M \subseteq \mathcal{H}$ is a closed subspace, then $x = p + (x - p)$ with $p \in M$ and $x - p \in M^\perp$.

Proof. Let $d = \inf_{q \in C} \|x - q\|$ and choose $q_n \in C$ with $\|x - q_n\| \to d$.
Since C is convex, $(q_n + q_m)/2 \in C$. The parallelogram identity gives
$\|q_n - q_m\|^2 = 2\|x - q_n\|^2 + 2\|x - q_m\|^2 - 4\left\|x - \frac{q_n + q_m}{2}\right\|^2 \le 2\|x - q_n\|^2 + 2\|x - q_m\|^2 - 4d^2$, so (q_n) is Cauchy.
Completeness and closedness give $q_n \to p \in C$, and continuity of the norm gives $\|x - p\| = d$. If p_1, p_2 both minimize, the same identity with $q_n = p_1$, $q_m = p_2$ gives $\|p_1 - p_2\| = 0$. Now let $C = M$ and $r = x - p$.
For $y \in M$ and $t \in \mathbb{R}$, minimality of $p + ty \in M$ gives $0 \le \|r - ty\|^2 - \|r\|^2 = -2t \operatorname{Re}\langle r, y\rangle + t^2\|y\|^2$, hence $\operatorname{Re}\langle r, y\rangle = 0$. Replacing y by iy in the complex case gives $\langle r, y\rangle = 0$. Thus $r \in M^\perp$ and $x = p + r$.
Structural Commentary: Convexity makes minimizing sequences Cauchy; closedness keeps the limit.
Key Insight: The nearest point is characterized by an orthogonal error. □

Definition 6.2.10. *Closed Linear Span and Total Sets*
For $A \subseteq \mathcal{H}$, define the *closed linear span* $\operatorname{cspan}(A) := \overline{\operatorname{span}(A)}$. The set A is *total* in $\mathcal{H}$ if $\operatorname{cspan}(A) = \mathcal{H}$.

Theorem 6.2.11. *Double Orthogonal Complements in Hilbert Spaces*
Let $\mathcal{H}$ be a Hilbert space.

(i) $\forall A \subseteq \mathcal{H}$, $\operatorname{cspan}(A) = A^{\perp\perp}$.
(ii) $\forall \mathcal{S} \leq \mathcal{H}$, $\overline{\mathcal{S}} = \mathcal{S}^{\perp\perp}$.
(iii) $\forall \mathcal{K} \leq \mathcal{H}$, $\mathcal{K}$ *closed* $\iff \mathcal{K} = \mathcal{K}^{\perp\perp}$.
(iv) $\forall A \subseteq \mathcal{H}$, $\operatorname{span}(A)$ *dense in* $\mathcal{H} \iff A^{\perp} = \{\mathbf{0}\}$.

Proof.

(i) Let $B = \operatorname{cspan}(A)$. Then B is closed, so Theorem 6.2.9 gives $\mathcal{H} = B \odot B^{\perp}$ and $B^{\perp\perp} = B$.
Also $B^{\perp} = A^{\perp}$, since $A \subseteq B$ and $x \in A^{\perp} \Rightarrow x \perp \operatorname{span}(A) \Rightarrow x \perp B$ by Proposition 6.2.2.
Hence $A^{\perp\perp} = B^{\perp\perp} = B = \operatorname{cspan}(A)$.
(ii) Apply (i) with $A = \mathcal{S}$: $\mathcal{S}^{\perp\perp} = \operatorname{cspan}(\mathcal{S}) = \overline{\operatorname{span}(\mathcal{S})} = \overline{\mathcal{S}}$.
(iii) By (ii), $\mathcal{K}^{\perp\perp} = \overline{\mathcal{K}}$; hence $\mathcal{K}$ closed $\iff \overline{\mathcal{K}} = \mathcal{K} \iff \mathcal{K}^{\perp\perp} = \mathcal{K}$.
(iv) $\operatorname{span}(A)$ dense $\iff \operatorname{cspan}(A) = \mathcal{H} \iff A^{\perp\perp} = \mathcal{H}$. If $A^{\perp} = \{\mathbf{0}\}$, then $A^{\perp\perp} = \mathcal{H}$.
Conversely $A^{\perp\perp} = \mathcal{H} \Rightarrow A^{\perp} \subseteq A^{\perp} \cap A^{\perp\perp} = \{\mathbf{0}\}$.

Structural Commentary: In a Hilbert space, closure is encoded by taking the orthogonal complement twice.
Key Insight: A missing direction from a closed span is detected by a nonzero orthogonal vector. □

Example 6.2.12. *Completeness Is Needed for Double Complements*
Let $\mathcal{V} := c_{00}(\mathbb{N};\mathbb{R}) \subseteq \ell^2$ with inherited inner product, $Lx := \sum_{n=1}^{\infty} x_n/n$, and $\mathcal{K} := \ker L$.
Then $|Lx| \leq \|x\|_2 \left(\sum_{n=1}^{\infty} n^{-2}\right)^{1/2}$, so L is continuous and $\mathcal{K}$ is a closed subspace of $\mathcal{V}$.
If $y \in \mathcal{K}^{\perp}$, then $e_1 - ne_n \in \mathcal{K}$ for $n \geq 2$, hence $0 = \langle y, e_1 - ne_n\rangle = y_1 - ny_n$.
Since $y \in c_{00}$, choose $n \geq 2$ with $y_n = 0$, so $y_1 = 0$ and then $y_n = 0$ for every n.
Thus $\mathcal{K}^{\perp} = \{\mathbf{0}\}$ and $\mathcal{K}^{\perp\perp} = \mathcal{V}$. But $e_1 \notin \mathcal{K}$, hence $\mathcal{K} \neq \mathcal{V} = \mathcal{K}^{\perp\perp}$.
Structural Commentary: Closedness need not recover double complements without completeness.
Key Insight: The identity $\mathcal{K} = \mathcal{K}^{\perp\perp}$ is genuinely Hilbert-space geometry.

Theorem 6.2.13. *Hilbert Basis Characterizations*
Let $\mathcal{O} \subseteq \mathcal{H}$ be orthonormal. The following are equivalent:

(i) $\mathcal{O}$ *is maximal orthonormal.*
(ii) $\mathcal{O}^{\perp} = \{\mathbf{0}\}$.
(iii) $\operatorname{cspan}(\mathcal{O}) = \mathcal{H}$.

Proof.

(i) $\Rightarrow$ *(ii)* If $0 \neq v \in \mathcal{O}^{\perp}$, then $\mathcal{O} \cup \{v/\|v\|\}$ is orthonormal, contradicting maximality; hence $\mathcal{O}^{\perp} = \{\mathbf{0}\}$.
(ii) $\Rightarrow$ *(iii)* By Theorem 6.2.11, $\operatorname{cspan}(\mathcal{O}) = \mathcal{O}^{\perp\perp} = (\mathcal{O}^{\perp})^{\perp} = \{\mathbf{0}\}^{\perp} = \mathcal{H}$.
(iii) $\Rightarrow$ *(i)* If $\mathcal{O}$ is not maximal, $\exists$ unit $u \in \mathcal{O}^{\perp}$; but (iii) and Theorem 6.2.11 give $\mathcal{O}^{\perp} = \{\mathbf{0}\}$, a contradiction.

Structural Commentary: Maximality means no orthogonal direction remains.
Key Insight: Hilbert bases span by closure, not usually by finite linear combinations. □

Definition 6.2.14. *Arbitrary Sums in a Hilbert Space*
Let $(x_k)_{k\in K} \subseteq \mathcal{H}$.

(i) *Finite partial sums*: $\mathcal{F}(K) := \{J \subseteq K : |J| < \infty\}$, ordered by $\subseteq$, and $s_J := \sum_{k\in J} x_k$.
(ii) *Arbitrary convergence*: $\sum_{k\in K} x_k = x \iff \forall \varepsilon > 0\, \exists S \in \mathcal{F}(K)\, \forall T \in \mathcal{F}(K) : S \subseteq T \Rightarrow \|s_T - x\| \leq \varepsilon$.
(iii) *Nonnegative scalar sum*: for $(a_k)_{k\in K} \subseteq [0,\infty)$, $\sum_{k\in K} a_k < \infty \iff \sup_{J\in\mathcal{F}(K)} \sum_{k\in J} a_k < \infty$.

Proposition 6.2.15. *Basic Properties of Arbitrary Sums*
Let $(x_k)_{k\in K}, (y_k)_{k\in K} \subseteq \mathcal{H}$ with $\sum_{k\in K} x_k = x$ and $\sum_{k\in K} y_k = y$. Then $\forall a, b \in \mathbb{F}$, $\sum_{k\in K}(ax_k + by_k) = ax + by$.
Moreover, $\forall z \in \mathcal{H}$, $\sum_{k\in K} \langle x_k, z\rangle = \langle x, z\rangle$ and $\sum_{k\in K} \langle z, x_k\rangle = \langle z, x\rangle$.

Proof. For $J \in \mathcal{F}(K)$, set $s_J := \sum_{k\in J} x_k$ and $t_J := \sum_{k\in J} y_k$.
Then $s_J \to x$ and $t_J \to y$, hence $as_J + bt_J \to ax + by$, proving $\sum_{k\in K}(ax_k + by_k) = ax + by$.
For $z \in \mathcal{H}$, Proposition 6.2.2 gives $\langle s_J, z\rangle \to \langle x, z\rangle$ and $\langle z, s_J\rangle \to \langle z, x\rangle$,
i.e. $\sum_{k\in K} \langle x_k, z\rangle = \langle x, z\rangle$ and $\sum_{k\in K} \langle z, x_k\rangle = \langle z, x\rangle$.
Structural Commentary: Arbitrary sums are ordinary finite sums plus a convergence net.
Key Insight: Once the vector sum converges, continuous linear tests may pass through the sum. □

Theorem 6.2.16. *Arbitrary Orthonormal Sums*
Let $(u_k)_{k\in K} \subseteq \mathcal{H}$ be orthonormal.

(i) $\forall (r_k)_{k\in K} \subseteq \mathbb{F}$, $\sum_{k\in K} r_k u_k$ *converges* $\iff \sup_{J\in\mathcal{F}(K)} \sum_{k\in J} |r_k|^2 < \infty$; *then* $\left\|\sum_{k\in K} r_k u_k\right\|^2 = \sum_{k\in K} |r_k|^2$.

(ii) $\forall (x_k)_{k\in K} \subseteq \mathcal{H}$, $\sum_{k\in K} x_k$ *converges* $\Longrightarrow |\{k \in K : x_k \neq \mathbf{0}\}| \leq \aleph_0$.

Proof.

(i) ($\Rightarrow$) Let $s_J := \sum_{k\in J} r_k u_k$. If $s_J \to x$, choose $S \in \mathcal{F}(K)$ such that $S \subseteq T \Rightarrow \|s_T - x\| \leq 1$. For any $J \in \mathcal{F}(K)$, $\sum_{k\in J} |r_k|^2 \leq \sum_{k\in J\cup S} |r_k|^2 = \|s_{J\cup S}\|^2 \leq (\|x\|+1)^2$.

(i) ($\Leftarrow$) Let $M := \sup_{J\in\mathcal{F}(K)} \sum_{k\in J} |r_k|^2 < \infty$. Given $\varepsilon > 0$, choose $S \in \mathcal{F}(K)$ with $M - \sum_{k\in S} |r_k|^2 < \varepsilon^2$. If $S \subseteq J, L$, then $\|s_J - s_L\|^2 \leq \sum_{k\in(J\cup L)\setminus S} |r_k|^2 < \varepsilon^2$, so (s_J) is Cauchy and converges in $\mathcal{H}$. Passing $\|s_J\|^2 = \sum_{k\in J} |r_k|^2$ to limits gives $\left\|\sum_{k\in K} r_k u_k\right\|^2 = \sum_{k\in K} |r_k|^2$.

(ii) If $\sum_{k\in K} x_k$ converges, then $\forall n \geq 1$ choose $I_n \in \mathcal{F}(K)$ such that $F \in \mathcal{F}(K \setminus I_n) \Rightarrow \left\|\sum_{k\in F} x_k\right\| \leq 1/n$. With $I := \bigcup_{n\geq 1} I_n$, I is countable; if $k \notin I$, then $\|x_k\| \leq 1/n$ for all n, hence $x_k = \mathbf{0}$.

Structural Commentary: Orthonormal sums reduce vector convergence to scalar square-summability.
Key Insight: A genuinely convergent arbitrary sum can hide no uncountable collection of nonzero terms. □

Theorem 6.2.17. *Fourier Expansion in a Hilbert Space*
Let $\mathcal{O} = (u_k)_{k\in K} \subseteq \mathcal{H}$ be orthonormal. For $x \in \mathcal{H}$, set $\widehat{x}_{\mathcal{O}} := \sum_{k\in K} \langle x, u_k\rangle u_k$.

(i) *$\widehat{x}_{\mathcal{O}}$ converges in $\mathcal{H}$ and $\|\widehat{x}_{\mathcal{O}}\|^2 = \sum_{k\in K} |\langle x, u_k\rangle|^2 \leq \|x\|^2$.*
(ii) *$\widehat{x}_{\mathcal{O}}$ is the unique minimizer of $\|x - y\|$ over $y \in \operatorname{cspan}(\mathcal{O})$.*
(iii) *$\|\widehat{x}_{\mathcal{O}}\| = \|x\| \Longleftrightarrow x \in \operatorname{cspan}(\mathcal{O})$.*

Proof.

(i) Finite Bessel gives $\sup_{J\in\mathcal{F}(K)} \sum_{k\in J} |\langle x, u_k\rangle|^2 \leq \|x\|^2$. With $r_k := \langle x, u_k\rangle$, Theorem 6.2.16 applies, so $\widehat{x}_{\mathcal{O}} = \sum_{k\in K} r_k u_k$ converges and $\|\widehat{x}_{\mathcal{O}}\|^2 = \sum_{k\in K} |r_k|^2 = \sum_{k\in K} |\langle x, u_k\rangle|^2 \leq \|x\|^2$.

(ii) Proposition 6.2.15 gives $\langle x - \widehat{x}_{\mathcal{O}}, u_k\rangle = 0$ for all k, hence $r := x - \widehat{x}_{\mathcal{O}} \in \operatorname{cspan}(\mathcal{O})^{\perp}$. Let $M := \operatorname{cspan}(\mathcal{O})$. For $y \in M$, $x - y = r + (\widehat{x}_{\mathcal{O}} - y)$ orthogonally, so $\|x - y\|^2 = \|r\|^2 + \|\widehat{x}_{\mathcal{O}} - y\|^2 \geq \|r\|^2 = \|x - \widehat{x}_{\mathcal{O}}\|^2$, with equality iff $y = \widehat{x}_{\mathcal{O}}$.

(iii) ($\Rightarrow$) If $\|\widehat{x}_{\mathcal{O}}\| = \|x\|$, then $x = \widehat{x}_{\mathcal{O}} + (x - \widehat{x}_{\mathcal{O}})$ orthogonally, so $\|x - \widehat{x}_{\mathcal{O}}\|^2 = \|x\|^2 - \|\widehat{x}_{\mathcal{O}}\|^2 = 0$ and $x = \widehat{x}_{\mathcal{O}} \in \operatorname{cspan}(\mathcal{O})$.

(iii) ($\Leftarrow$) If $x \in \operatorname{cspan}(\mathcal{O})$, then the unique minimizer in (ii) is x, so $\widehat{x}_{\mathcal{O}} = x$ and $\|\widehat{x}_{\mathcal{O}}\| = \|x\|$.

Structural Commentary: Fourier expansion is projection onto the closed span.
Key Insight: Bessel bounds coefficients; completeness sums them. □

Corollary 6.2.18. *Parseval Characterization of Hilbert Bases*
For an orthonormal family $\mathcal{O} = \{u_k : k \in K\}$ in $\mathcal{H}$, the following are equivalent:

(i) *$\mathcal{O}$ is a Hilbert basis.*
(ii) *$x = \sum_{k\in K} \langle x, u_k\rangle u_k$ for every $x \in \mathcal{H}$.*
(iii) *Parseval's identity $\langle x, y\rangle = \sum_{k\in K} \langle x, u_k\rangle \overline{\langle y, u_k\rangle}$ holds for all $x, y \in \mathcal{H}$.*
(iv) *$\|x\|^2 = \sum_{k\in K} |\langle x, u_k\rangle|^2$ for every $x \in \mathcal{H}$.*

Proof.

(i) $\Rightarrow$ *(ii)* By Theorem 6.2.13, $\operatorname{cspan}(\mathcal{O}) = \mathcal{H}$. For $x \in \mathcal{H}$, both x, $\widehat{x}_{\mathcal{O}}$ minimize $\|x - y\|$ over $y \in \operatorname{cspan}(\mathcal{O})$, so Theorem 6.2.17 gives $x = \widehat{x}_{\mathcal{O}} = \sum_{k\in K} \langle x, u_k\rangle u_k$.

(ii) $\Rightarrow$ *(iii)* For $x, y \in \mathcal{H}$, (ii) and Proposition 6.2.15 give $\langle x, y\rangle = \sum_{k\in K} \langle x, u_k\rangle \langle u_k, y\rangle = \sum_{k\in K} \langle x, u_k\rangle \overline{\langle y, u_k\rangle}$.

(iii) $\Rightarrow$ *(iv)* Setting $y = x$ in (iii) gives $\|x\|^2 = \langle x, x\rangle = \sum_{k\in K} \langle x, u_k\rangle \overline{\langle x, u_k\rangle} = \sum_{k\in K} |\langle x, u_k\rangle|^2$ for all $x \in \mathcal{H}$.

(iv) $\Rightarrow$ *(i)* By Theorem 6.2.17, (iv) implies $\|\widehat{x}_{\mathcal{O}}\| = \|x\|$, hence $x \in \operatorname{cspan}(\mathcal{O})$ for all $x \in \mathcal{H}$. Therefore $\operatorname{cspan}(\mathcal{O}) = \mathcal{H}$ and Theorem 6.2.13 gives (i).

Structural Commentary: Completeness is detected by expansion, norms, or inner products.
Key Insight: Parseval says the Hilbert basis loses no length or angle information. □

Definition 6.2.19. *Hilbert Dimension*
For a Hilbert space $\mathcal{H}$ with Hilbert basis $\mathcal{B}$, define $\operatorname{hdim}(\mathcal{H}) := |\mathcal{B}|$.

Theorem 6.2.20. *Hilbert Dimension and $\ell^2(K)$ Models*
Let K be a set.

(i) *Every Hilbert space has a Hilbert basis $(\delta_k)_{k\in K}$; if $\mathcal{B}, \mathcal{C}$ are Hilbert bases of $\mathcal{H}$, then $|\mathcal{B}| = |\mathcal{C}|$.*
(ii) *$\ell^2(K) := \{f : K \to \mathbb{F} : \sum_{k\in K} |f(k)|^2 < \infty\}$ is Hilbert under $\langle f, g\rangle = \sum_{k\in K} f(k)\overline{g(k)}$, with $(\delta_k)_{k\in K}$.*
(iii) *For Hilbert spaces $\mathcal{H}_1, \mathcal{H}_2$, $\mathcal{H}_1 \cong_{\text{iso}} \mathcal{H}_2 \Longleftrightarrow \operatorname{hdim}(\mathcal{H}_1) = \operatorname{hdim}(\mathcal{H}_2)$.*

Proof.

(*i*) Zorn's lemma, Principle 6.5.2, gives a maximal orthonormal set, hence a Hilbert basis by Theorem 6.2.13. If $\mathcal{B} = \{b_i : i \in I\}$ and $\mathcal{C} = \{c_j : j \in J\}$ are infinite Hilbert bases, set $J_i := \{j \in J : \langle b_i, c_j\rangle \neq 0\}$. Theorem 6.2.16 gives $|J_i| \leq \aleph_0$, and totality of $\mathcal{B}$ gives $J = \bigcup_{i\in I} J_i$; hence $|J| \leq \aleph_0 |I| = |I|$, and symmetry with Theorem 6.3.1 gives $|I| = |J|$. The finite case is ordinary dimension.

(*ii*) Finite-subset Cauchy–Schwarz makes $\langle f, g\rangle$ well defined on $\ell^2(K)$. If (f_n) is Cauchy, $f(k) := \lim_n f_n(k)$ exists for each k, and finite-subset estimates give $f \in \ell^2(K)$ and $f_n \to f$; hence $\ell^2(K)$ is Hilbert. The family $(\delta_k)_{k\in K}$ is orthonormal, and $(\delta_k)_{k\in K}^{\perp} = \{\mathbf{0}\}$, so it is a Hilbert basis by Theorem 6.2.13.

(*iii*) ($\Rightarrow$) If $U : \mathcal{H}_1 \to \mathcal{H}_2$ is a surjective linear isometry and $\mathcal{B}$ is a Hilbert basis of $\mathcal{H}_1$, then $U\mathcal{B}$ is a Hilbert basis of $\mathcal{H}_2$; hence $\operatorname{hdim}(\mathcal{H}_1) = \operatorname{hdim}(\mathcal{H}_2)$.

(*iii*) ($\Leftarrow$) If $\operatorname{hdim}(\mathcal{H}_1) = \operatorname{hdim}(\mathcal{H}_2)$, choose Hilbert bases $(b_i)_{i\in I}$ and $(c_i)_{i\in I}$. Parseval gives isometric isomorphisms $\Phi_1 x = (i \mapsto \langle x, b_i\rangle)$ and $\Phi_2 y = (i \mapsto \langle y, c_i\rangle)$ onto $\ell^2(I)$. Thus $\Phi_2^{-1}\Phi_1 : \mathcal{H}_1 \to \mathcal{H}_2$ is a surjective linear isometry.

Structural Commentary: Hilbert bases classify Hilbert spaces by cardinality.
Key Insight: Every Hilbert space is a coordinate $\ell^2(K)$ after choosing a Hilbert basis. □

Corollary 6.2.21. *Separable Hilbert Spaces*
$\mathcal{H}$ separable $\iff \operatorname{hdim}(\mathcal{H}) \leq \aleph_0$.

Proof.

($\Rightarrow$) Let $D \subseteq \mathcal{H}$ be countable dense and let $\mathcal{O}$ be a Hilbert basis. For each $u \in \mathcal{O}$ choose $d_u \in D$ with $\|u - d_u\| < 1/3$. If $u \neq v$ and $d_u = d_v$, then $\sqrt{2} = \|u - v\| \leq \|u - d_u\| + \|d_v - v\| < 2/3$, impossible. Hence $u \mapsto d_u$ injects $\mathcal{O}$ into D, so $\operatorname{hdim}(\mathcal{H}) = |\mathcal{O}| \leq \aleph_0$.

($\Leftarrow$) Let $\mathcal{O}$ be a Hilbert basis with $|\mathcal{O}| \leq \aleph_0$. The finite $\mathbb{Q}$-span of $\mathcal{O}$ over $\mathbb{R}$, respectively finite $(\mathbb{Q} + i\mathbb{Q})$-span over $\mathbb{C}$, is countable and dense in $\operatorname{cspan}(\mathcal{O}) = \mathcal{H}$.

Structural Commentary: Countable bases and countable dense sets force each other.
Key Insight: Separability of a Hilbert space is exactly countability of its orthogonal coordinate directions. □

Corollary 6.2.22. *Infinite Hilbert Bases Are Not Hamel Bases*
If $\operatorname{hdim}(\mathcal{H})$ is infinite, no Hilbert basis of $\mathcal{H}$ is a Hamel basis.

Proof. Let $\mathcal{O}$ be a Hilbert basis and choose pairwise distinct $(u_n) \subseteq \mathcal{O}$.
Since $\sum_{n=1}^{\infty} 4^{-n} < \infty$, Theorem 6.2.16 gives $x := \sum_{n=1}^{\infty} 2^{-n} u_n \in \mathcal{H}$. Proposition 6.2.15 gives $\langle x, u_n\rangle = 2^{-n} \neq 0$ for all n, so x has infinitely many nonzero $\mathcal{O}$-coordinates and $x \notin \operatorname{span}(\mathcal{O})$.
Structural Commentary: Infinite Hilbert coordinates allow square-summable infinite expansions.
Key Insight: A Hilbert basis spans by limits, while a Hamel basis spans by finite sums. □

Perspective 6.2.23. *Bounded Functionals*
Here $\mathcal{H}^*$ denotes the continuous dual of $\mathcal{H}$. By Definition 6.1.2 and Theorem 6.1.3, $\mathcal{H}^* = \{\varphi : \mathcal{H} \to \mathbb{F} : \varphi \text{ linear}, \ \|\varphi\|_{\mathrm{op}} < \infty\}$.

Proposition 6.2.24. *Operator Norm Algebra*
Let $\mathcal{B}(\mathcal{H})$ be the bounded linear operators $\mathcal{H} \to \mathcal{H}$.
For $S, T \in \mathcal{B}(\mathcal{H})$ and $a \in \mathbb{F}$, $\|aT\|_{op} = |a| \, \|T\|_{op}$, $\|S + T\|_{op} \leq \|S\|_{op} + \|T\|_{op}$, and $\|ST\|_{op} \leq \|S\|_{op} \, \|T\|_{op}$.
Hence $\mathcal{B}(\mathcal{H})$ is a normed algebra.

Proof. Taking $\sup_{\|x\|\leq 1}$ in $\|aTx\| = |a| \, \|Tx\|$ and $\|(S+T)x\| \leq \|Sx\| + \|Tx\| \leq (\|S\|_{\mathrm{op}} + \|T\|_{\mathrm{op}}) \, \|x\|$ gives homogeneity and subadditivity. Also $\|STx\| \leq \|S\|_{\mathrm{op}} \, \|Tx\| \leq \|S\|_{\mathrm{op}} \, \|T\|_{\mathrm{op}} \, \|x\|$, so $\|ST\|_{\mathrm{op}} \leq \|S\|_{\mathrm{op}} \, \|T\|_{\mathrm{op}}$.
Structural Commentary: Composition multiplies stretch bounds.
Key Insight: Bounded operators behave like matrices with a compatible norm. □

Theorem 6.2.25. *Riesz Representation for Hilbert Spaces*
For every bounded linear functional $\varphi : \mathcal{H} \to \mathbb{F}$, $\exists! y \in \mathcal{H}$ such that $\forall x \in \mathcal{H}$, $\varphi(x) = \langle x, y\rangle$; moreover $\|\varphi\|_{op} = \|y\|$.

Proof. If $\varphi = 0$, take $y = 0$. If $\varphi \neq 0$, let $M := \ker \varphi$ and choose $a \notin M$. Since M is closed, Theorem 6.2.9 gives $a = p + u$ with $p \in M$ and $0 \neq u \in M^{\perp}$; hence $\varphi(u) = \varphi(a) \neq 0$. For $x \in \mathcal{H}$, $x - \varphi(x)\varphi(u)^{-1} u \in M$, so $\langle x, u\rangle = \varphi(x)\varphi(u)^{-1} \, \|u\|^2$. Put $y := \overline{\varphi(u)} \, \|u\|^{-2} \, u$; conjugate linearity in the second slot gives $\langle x, y\rangle = \varphi(x)$.
If z also represents φ, then $\langle x, y - z\rangle = 0$ for all x, so $x = y - z$ gives $y = z$.
Finally, $|\varphi(x)| \leq \|x\| \, \|y\|$ gives $\|\varphi\|_{\mathrm{op}} \leq \|y\|$, while $y \neq 0$ gives $\|\varphi\|_{\mathrm{op}} \geq |\varphi(y)| \, / \, \|y\| = \|y\|$; $y = 0$ is immediate.
Structural Commentary: The kernel is a hyperplane; its orthogonal direction represents the functional.
Key Insight: Bounded functionals are exactly inner products with fixed vectors. □

Corollary 6.2.26. *Riesz Identifies the Continuous Dual*
The Riesz map $R : \mathcal{H} \to \mathcal{H}^*$, $Ry := (x \mapsto \langle x, y\rangle)$, is an isometric bijection.
It is conjugate-linear over $\mathbb{C}$ and linear over $\mathbb{R}$.

Proof. Theorem 6.2.25 gives $\forall \varphi \in \mathcal{H}^* \, \exists! y \in \mathcal{H} : \varphi = Ry$ and $\|Ry\|_{\text{op}} = \|y\|$, so R is bijective and isometric.
Also $R(\alpha y + \beta z) = \overline{\alpha} Ry + \overline{\beta} Rz$ over $\mathbb{C}$, while $R(\alpha y + \beta z) = \alpha Ry + \beta Rz$ over $\mathbb{R}$.
Structural Commentary: Continuous functionals are represented by vectors.
Key Insight: Riesz identifies $\mathcal{H}^*$ with $\mathcal{H}$, up to complex conjugation. □

Theorem 6.2.27. *Orthonormal Systems in L^2*
Let $(\phi_n)_{n\geq 1} \subseteq L^2(\mu)$ be orthonormal, $\langle \phi_n, \phi_m\rangle = \delta_{nm}$, and $c_n(f) := \langle f, \phi_n\rangle = \int_X f\overline{\phi_n}\, d\mu$.

(i) $\forall f \in L^2(\mu)$, $\sum_{n=1}^{\infty} |c_n(f)|^2 \leq \|f\|_2^2$ and $f \sim_{\text{coef}} \sum c_n(f)\phi_n$.
(ii) If $(c_n) \in \ell^2$, then $s_N := \sum_{k=1}^{N} c_k \phi_k$ converges in $L^2(\mu)$.
(iii) If (ϕ_n) is complete, then $\forall f \in L^2(\mu)$, $\|f\|_2^2 = \sum_{n=1}^{\infty} |c_n(f)|^2$.

For the trigonometric system on $[-\pi, \pi]$, this is L^2 Parseval and a.e. uniqueness.

Proof.

(i) Finite orthonormality gives $\left\| f - \sum_{k=1}^{N} c_k(f)\phi_k \right\|_2^2 = \|f\|_2^2 - \sum_{k=1}^{N} |c_k(f)|^2 \geq 0$.
Hence $\sum_{n=1}^{\infty} |c_n(f)|^2 \leq \|f\|_2^2$ and $f \sim_{\text{coef}} \sum c_n(f)\phi_n$.
(ii) For $m < N$, $\|s_N - s_m\|_2^2 = \sum_{k=m+1}^{N} |c_k|^2 \to 0$ since $(c_n) \in \ell^2$.
Completeness of $L^2(\mu)$ gives $s_N \to g \in L^2(\mu)$.
(iii) If (ϕ_n) is complete, then (i) gives $(c_n(f)) \in \ell^2$, so (ii) gives $s_N \to g$. For every j, $\langle f - g, \phi_j\rangle = c_j(f) - \lim_N \langle s_N, \phi_j\rangle = 0$. Completeness gives $f = g$, hence $\|f\|_2^2 = \lim_N \|s_N\|_2^2 = \sum_{n=1}^{\infty} |c_n(f)|^2$.

Structural Commentary: Complete orthonormal systems coordinatize L^2 by square-summable sequences.
Key Insight: Parseval is Pythagoras after passing from finite coordinates to an L^2 limit. □

Method 6.2.28. *Hilbert Space Proof Patterns*

(i) *Complete*: Cauchy tails, completions, and convergent subsequences force limits.
(ii) *Parallelogram*: test Hilbertian norms and make distance-minimizing sequences Cauchy.
(iii) *Project*: for closed M, write $x = p + r$, $p \in M$, $r \in M^{\perp}$.
(iv) *Separate*: density/equality via $M^{\perp} = \{\mathbf{0}\}$ and $M^{\perp\perp} = \overline{M}$.
(v) *Sum*: arbitrary convergence is net convergence over $\mathcal{F}(K)$; control finite tails.
(vi) *Expand*: orthonormal finite identities + Bessel bounds + completeness $\Rightarrow$ Fourier expansions.
(vii) *Classify*: Hilbert bases give hdim, separability, and $\mathcal{H} \cong \ell^2(K)$.
(viii) *Dualize*: $\mathcal{B}(\mathcal{H})$ norm algebra and Riesz $\mathcal{H}^* \cong \mathcal{H}$.

Key Insight: Hilbert-space arguments are finite orthogonal geometry plus completeness.

6.3 Core Functional Analytic Theorems

The core theorems explain why complete normed spaces behave rigidly under linear maps. Baire category turns pointwise control, surjectivity, and closed graphs into operator statements.

Theorem 6.3.1. *Hahn–Banach Extension Theorem*
Let $\mathcal{X}$ be a real normed space, let $M \subseteq \mathcal{X}$ be a linear subspace, and let $p : \mathcal{X} \to \mathbb{R}$ be sublinear. If $\varphi : M \to \mathbb{R}$ is linear and $\varphi(x) \leq p(x)$ for all $x \in M$, then φ extends to a linear functional $\Phi : \mathcal{X} \to \mathbb{R}$ with $\Phi(x) \leq p(x)$ for all $x \in \mathcal{X}$. In the normed complex form, every bounded functional on a subspace extends with the same norm.

Theorem 6.3.2. *Positive Riesz Representation on $C([a, b])$*
Let $J = [a, b]$ and let $G \in C(J)^*$ be positive, i.e. $f \geq 0$ on J implies $G(f) \geq 0$.
Then $\exists$ a finite positive Borel measure γ on $\mathbb{R}$ with $\operatorname{supp}\gamma \subseteq J$ s.t. $G(f) = \int_J f\, d\gamma$ $\forall f \in C(J)$.
Moreover, $\|G\|_{op} = G(\mathbf{1}_J) = \gamma(J)$.

Proof. For $a \leq t < b$ and n large, set

$$\varphi_{t,n}(s) := \begin{cases} 1, & a \leq s \leq t, \\ n(t + 1/n - s), & t < s < t + 1/n, \\ 0, & t + 1/n \leq s \leq b. \end{cases}$$

Then $\varphi_{t,n} \downarrow \mathbf{1}_{[a,t]}$ pointwise, hence $G(\varphi_{t,n})$ decreases.
Define $g(t) := \lim_{n\to\infty} G(\varphi_{t,n})$; $g(t) = 0$ for $t < a$; and $g(t) = G(\mathbf{1}_J)$ for $t \geq b$.
Ramp comparisons give g increasing and right-continuous.

By Theorem 5.1.27, g induces a finite Borel–Stieltjes measure γ with $\operatorname{supp}\gamma \subseteq J$ and $\gamma(J) = g(b) = G(\mathbf{1}_J)$.
Fix $f \in C(J)$ and $\varepsilon > 0$. Choose $P : a = t_0 < \cdots < t_m = b$ with oscillation $\omega_f(P) \le \varepsilon$.
Put $\Delta g_k := g(t_k) - g(t_{k-1})$. For n large, ramp increments $\varphi_{t_k,n} - \varphi_{t_{k-1},n}$ give continuous approximant $f_{P,n}$ whose G-increments approximate $\sum_k f(t_k)\Delta g_k$ within ε, while step approximant s_P satisfies $\|f - s_P\|_{L^1(\gamma)} \le \varepsilon\gamma(J)$.
Hence $\left|G(f) - \int_J f\,d\gamma\right| \le \varepsilon(2\,\|G\|_{\mathrm{op}} + 1)$ after absorbing constants into ε. Letting $\varepsilon \downarrow 0$ gives $G(f) = \int_J f\,d\gamma$.
Finally, positivity gives $|G(f)| \le G(\mathbf{1}_J)\,\|f\|_\infty = \gamma(J)\,\|f\|_\infty$, and $f = \mathbf{1}_J$ gives $\|G\|_{\mathrm{op}} \ge \gamma(J)$.
Structural Commentary: Positivity turns ramp limits into interval masses, hence a Stieltjes measure.
Key Insight: Ramps turn order on $C(J)$ into interval masses. □

Corollary 6.3.3. *Signed Riesz Representation by a Charge*
Let $J = [a,b]$ and let $G \in C(J)^*$ admit $G = G^+ - G^-$, $G^\pm \in C(J)^*$; for $G^\pm \ge 0$.
Then $\exists$ a charge $\lambda : \mathcal{B}(\mathbb{R}) \to \mathbb{R}$ with $\operatorname{supp}\lambda \subseteq J$ such that $G(f) = \int_J f\,d\lambda\ \forall f \in C(J)$.

Proof. By Theorem 6.3.2, $\exists$ finite positive Borel measures $\gamma^\pm$ with $\operatorname{supp}\gamma^\pm \subseteq J$; $G^\pm(f) = \int_J f\,d\gamma^\pm\ \forall f \in C(J)$.
Set $\lambda := \gamma^+ - \gamma^-$, i.e. $\lambda(E) := \gamma^+(E) - \gamma^-(E)$ for $E \in \mathcal{B}(\mathbb{R})$. Then λ is a charge by Proposition 5.1.82, $\operatorname{supp}\lambda \subseteq J$, and $G(f) = G^+(f) - G^-(f) = \int_J f\,d\gamma^+ - \int_J f\,d\gamma^- = \int_J f\,d\lambda$.
Structural Commentary: Positive representations subtract to a charge representation.
Key Insight: Signed representation is positive representation plus subtraction. □

Theorem 6.3.4. *Uniform Boundedness Principle*
Let $\mathcal{X}$ be Banach and let $\mathcal{Y}$ be normed. If $\mathcal{F} \subseteq \mathcal{L}(\mathcal{X},\mathcal{Y})$ is pointwise bounded, meaning $\sup_{T\in\mathcal{F}} \|Tx\| < \infty$ for every $x \in \mathcal{X}$, then $\sup_{T\in\mathcal{F}} \|T\|_{op} < \infty$.

Theorem 6.3.5. *Open Mapping Theorem*
A surjective bounded linear map $T : \mathcal{X} \to \mathcal{Y}$ between Banach spaces is open. Equivalently, T maps neighborhoods of 0 onto neighborhoods of 0.

Theorem 6.3.6. *Closed Graph Theorem*
Let $\mathcal{X}$ and $\mathcal{Y}$ be Banach spaces. A linear map $T : \mathcal{X} \to \mathcal{Y}$ is bounded if and only if its graph

$$\Gamma(T) := \{(x, Tx) : x \in \mathcal{X}\}$$

is closed in $\mathcal{X} \times \mathcal{Y}$.

6.4 Weak Topologies and Duality

Weak topologies replace norm control with scalar tests by functionals, exposing compactness and duality that strong topology can hide.

Definition 6.4.1. *Dual Space and Weak Convergence*
For a normed space $\mathcal{X}$, its continuous dual is $\mathcal{X}^* := \mathcal{L}(\mathcal{X},\mathbb{F})$. A sequence (x_n) *converges weakly* to x, written $x_n \rightharpoonup x$, if

$$\varphi(x_n) \to \varphi(x) \qquad \text{for every } \varphi \in \mathcal{X}^*.$$

Definition 6.4.2. *Weak-* Convergence*
If $(\varphi_n) \subseteq \mathcal{X}^*$ and $\varphi \in \mathcal{X}^*$, then φ_n *converges weak-** to φ if

$$\varphi_n(x) \to \varphi(x) \qquad \text{for every } x \in \mathcal{X}.$$

This is weak convergence in the dual tested only against evaluation at points of the original space.

Definition 6.4.3. *Reflexivity*
The canonical map $J : \mathcal{X} \to \mathcal{X}^{**}$ is defined by $Jx(\varphi) = \varphi(x)$. A Banach space is *reflexive* if J is onto.

Theorem 6.4.4. *Banach–Alaoglu Theorem*
The closed unit ball of $\mathcal{X}^*$ is compact in the weak- topology. Consequently, bounded nets of functionals have weak-* cluster points.*

6.5 Compact Operators

Compact operators carry bounded sets to relatively compact sets. Their spectral and solvability theory therefore resembles matrix theory inside infinite-dimensional spaces.

Definition 6.5.1. *Compact and Finite-Rank Operators*
Let $\mathcal{X},\mathcal{Y}$ be normed spaces. A linear map $T : \mathcal{X} \to \mathcal{Y}$ is *compact* if $T(B)$ has compact closure in $\mathcal{Y}$ whenever $B \subseteq \mathcal{X}$ is bounded.
The map has *finite rank* if $\dim \operatorname{range}(T) < \infty$.

Proposition 6.5.2. *Finite Rank Operators Are Compact*
If $\mathcal{X}$ is normed, $\mathcal{Y}$ is Banach, and $T : \mathcal{X} \to \mathcal{Y}$ is bounded with finite-dimensional range, then T is compact.

Theorem 6.5.3. *Norm Limits of Compact Operators*
If $T_n : \mathcal{X} \to \mathcal{Y}$ are compact operators and $T_n \to T$ in operator norm, with $\mathcal{Y}$ Banach, then T is compact.

Theorem 6.5.4. *Fredholm Alternative*
Let $K : \mathcal{X} \to \mathcal{X}$ be compact on a Banach space. For $\lambda \neq 0$, the equation

$$(\lambda I - K)x = y$$

has a solution for every y exactly when the homogeneous equation $(\lambda I - K)x = 0$ has only the zero solution. In this case $(\lambda I - K)^{-1}$ is bounded.

6.6 Spectral Theory of Operators

For bounded operators, spectrum replaces the finite-dimensional eigenvalue list. The resolvent tracks invertibility, while the spectrum records every obstruction to it.

Definition 6.6.1. *Spectrum and Resolvent*
Let $T \in \mathcal{L}(\mathcal{X})$ be a bounded operator on a Banach space. The *resolvent set* is

$$\rho(T) := \{\lambda \in \mathbb{C} : \lambda I - T \text{ is bijective and } (\lambda I - T)^{-1} \in \mathcal{L}(\mathcal{X})\}.$$

The *spectrum* is $\sigma(T) := \mathbb{C} \setminus \rho(T)$.

Definition 6.6.2. *Spectral Radius*
The *spectral radius* of $T \in \mathcal{L}(\mathcal{X})$ is

$$r(T) := \sup\{|\lambda| : \lambda \in \sigma(T)\}.$$

Theorem 6.6.3. *Spectrum of a Complex Bounded Operator*
If $\mathcal{X}$ is a nonzero complex Banach space and $T \in \mathcal{L}(\mathcal{X})$, then $\sigma(T)$ is a nonempty compact subset of $\mathbb{C}$ and $\sigma(T) \subseteq \{\lambda : |\lambda| \leq \|T\|_{op}\}$.

Theorem 6.6.4. *Spectral Radius Formula*
For $T \in \mathcal{L}(\mathcal{X})$ on a complex Banach space,

$$r(T) = \lim_{n\to\infty} \|T^n\|_{op}^{1/n}.$$

Perspective 6.6.5. *Hilbert Operator Spectral Theory*
For bounded self-adjoint and normal operators on Hilbert spaces, spectral theory recovers a functional-calculus form of diagonalization. Compact self-adjoint operators have an orthonormal eigenbasis away from the kernel, giving the closest infinite-dimensional analogue of the finite spectral theorem.

6.7 Distributions and Sobolev Preview

Distributions test against smooth compactly supported functions; Sobolev spaces package weak derivatives in normed or Hilbert spaces as a bridge to PDE.

Definition 6.7.1. *Test Functions and Distributions*
Let $\Omega \subseteq \mathbb{R}^n$ be open. A *test function* is a smooth function $\varphi \in C_c^\infty(\Omega)$. A *distribution* on Ω is a continuous linear functional on the test-function space, written $u \in \mathcal{D}'(\Omega)$ and paired as $\langle u, \varphi \rangle$.

Part VI
Probability and Stochastic Structure

1 Probability Spaces and Convergence

1.1 Probability Spaces and Random Variables

Probability turns a data-generating experiment into numerical statements about its outcomes. The elementary language is set-theoretic; the measure-theoretic form is obtained by choosing which subsets are measurable events.

Definition 1.1.1. *Sample Spaces, Events, and Probability Spaces*

(i) A *sample space* is a set Ω of possible outcomes.
(ii) An *outcome* or *realization* is a point $\omega \in \Omega$.
(iii) An *event* is a subset $A \subseteq \Omega$; the admissible events form a *σ-algebra* $\mathcal{F} \subseteq \mathcal{P}(\Omega)$.
(iv) A *probability measure* on $(\Omega, \mathcal{F})$ is a measure $\mathbb{P} : \mathcal{F} \to [0,1]$ with $\mathbb{P}(\Omega) = 1$.
(v) The triple $(\Omega, \mathcal{F}, \mathbb{P})$ is a *probability space*.

Key Insight: Probability is countably additive size normalized so that the whole sample space has mass 1.

Method 1.1.2. *Sample Space Construction Method*
Choose Ω as the set of complete outcomes and encode the target condition as an event $A \subseteq \Omega$:

(i) *Finite product method*: for m discrete trials with state set S, use $\Omega = S^m$ and coordinate constraints. For two coin tosses, $S = \{H, T\}$, $\Omega = S^2$, and first toss heads gives $A = \{\omega \in \Omega : \omega_1 = H\} = \{HH, HT\}$.
(ii) *Interval method*: for real-valued measurements, use $\Omega = \mathbb{R}$ and translate inequalities into intervals. $10 < \omega \leq 23$ gives $A = (10, 23]$.
(iii) *Cylinder method*: for infinite sequences with state set S, use $\Omega = S^{\mathbb{N}}$ and constrain finitely many coordinates; for coin tosses, $S = \{H, T\}$ and first head on toss 3 gives $A = \{\omega \in \Omega : \omega_1 = T, \omega_2 = T, \omega_3 = H\}$.

Definition 1.1.3. *Event Algebra, Indicators, and Monotone Events*

(i) For $A \in \mathcal{F}$, the *complement* is $A^c := \Omega \setminus A$.
(ii) For $A, B \in \mathcal{F}$, the *union* and *intersection* are $A \cup B$ and $A \cap B$.
(iii) For $A, B \in \mathcal{F}$, the *difference* is $A \setminus B := A \cap B^c$.
(iv) *Commutativity*: $A \cup B = B \cup A$ and $A \cap B = B \cap A$.
(v) *Associativity*: $(A \cup B) \cup C = A \cup (B \cup C)$ and $(A \cap B) \cap C = A \cap (B \cap C)$.
(vi) *Distributivity*: $A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$ and $A \cup (B \cap C) = (A \cup B) \cap (A \cup C)$.
(vii) *De Morgan laws*: $(A \cup B)^c = A^c \cap B^c$ and $(A \cap B)^c = A^c \cup B^c$.
(viii) Events $(A_i)_{i \in I}$ are *disjoint* or *mutually exclusive* if $A_i \cap A_j = \emptyset$ for $i \neq j$.
(ix) Events $(A_i)_{i \in I}$ form a *partition* of Ω if $A_i \cap A_j = \emptyset$ for $i \neq j$ and $\bigcup_{i \in I} A_i = \Omega$.
(x) The *indicator* of $A \in \mathcal{F}$ is $\mathbf{1}_A : \Omega \to \{0, 1\}$, where $\mathbf{1}_A(\omega) = 1$ for $\omega \in A$ and $\mathbf{1}_A(\omega) = 0$ for $\omega \notin A$.
(xi) A sequence (A_n) is *increasing* to A, written $A_n \uparrow A$, if $A_1 \subseteq A_2 \subseteq \cdots$ and $A = \bigcup_{n \geq 1} A_n$.
(xii) A sequence (A_n) is *decreasing* to A, written $A_n \downarrow A$, if $A_1 \supseteq A_2 \supseteq \cdots$ and $A = \bigcap_{n \geq 1} A_n$.

Example 1.1.4. *Endpoint Sensitivity in Monotone Limits*
Let $\Omega = \mathbb{R}$.

(i) If $A_i = [0, 1/i]$, then $\bigcup_{i=1}^{\infty} A_i = [0, 1]$ and $\bigcap_{i=1}^{\infty} A_i = \{0\}$.
(ii) If $A_i = (0, 1/i)$, then $\bigcup_{i=1}^{\infty} A_i = (0, 1)$ and $\bigcap_{i=1}^{\infty} A_i = \emptyset$.

Key Insight: Monotone set limits remember whether boundary points are included.

Proposition 1.1.5. *First Consequences of the Probability Axioms*
For events $A, B \in \mathcal{F}$:

(i) Null event: $\mathbb{P}(\emptyset) = 0$.
(ii) Monotonicity: $A \subseteq B \Rightarrow \mathbb{P}(A) \leq \mathbb{P}(B)$.
(iii) Bounds: $0 \leq \mathbb{P}(A) \leq 1$.
(iv) Complement: $\mathbb{P}(A^c) = 1 - \mathbb{P}(A)$.
(v) Disjoint additivity: $A \cap B = \emptyset \Rightarrow \mathbb{P}(A \cup B) = \mathbb{P}(A) + \mathbb{P}(B)$.
(vi) Two-set inclusion-exclusion: $\mathbb{P}(A \cup B) = \mathbb{P}(A) + \mathbb{P}(B) - \mathbb{P}(A \cap B)$.

Proof.

(i) Countable additivity on $\Omega = \Omega \sqcup \emptyset \sqcup \emptyset \sqcup \cdots$ gives $1 = 1 + \sum_{n \geq 1} \mathbb{P}(\emptyset)$, hence $\mathbb{P}(\emptyset) = 0$.

(ii) If $A \subseteq B$, then $B = A \sqcup (B \setminus A)$, so $\mathbb{P}(B) = \mathbb{P}(A) + \mathbb{P}(B \setminus A) \geq \mathbb{P}(A)$.
(iii) Since $\emptyset \subseteq A \subseteq \Omega$, (i) and (ii) give $0 = \mathbb{P}(\emptyset) \leq \mathbb{P}(A) \leq \mathbb{P}(\Omega) = 1$.
(iv) Since $\Omega = A \sqcup A^c$, $1 = \mathbb{P}(\Omega) = \mathbb{P}(A) + \mathbb{P}(A^c)$, so $\mathbb{P}(A^c) = 1 - \mathbb{P}(A)$.
(v) If $A \cap B = \emptyset$, then $A \cup B = A \sqcup B$, so $\mathbb{P}(A \cup B) = \mathbb{P}(A) + \mathbb{P}(B)$.
(vi) Decompose $A \cup B = (A \setminus B) \sqcup (A \cap B) \sqcup (B \setminus A)$, $A = (A \setminus B) \sqcup (A \cap B)$, and $B = (A \cap B) \sqcup (B \setminus A)$. Hence $\mathbb{P}(A \cup B) = \mathbb{P}(A) + \mathbb{P}(B) - \mathbb{P}(A \cap B)$.

Structural Commentary: The algebraic identities are set partitions in disguise.
Key Insight: Split overlapping sets into disjoint pieces before adding probabilities. □

Example 1.1.6. *Two Coin Tosses and Inclusion-Exclusion*
Let $\Omega = \{HH, HT, TH, TT\}$ with $\mathbb{P}(\{\omega\}) = 1/4$ for all $\omega \in \Omega$. Set $H_1 = \{HH, HT\}$ and $H_2 = \{HH, TH\}$. Since $H_1 \cap H_2 = \{HH\}$, $\mathbb{P}(H_1 \cup H_2) = \mathbb{P}(H_1) + \mathbb{P}(H_2) - \mathbb{P}(H_1 \cap H_2) = \frac{1}{2} + \frac{1}{2} - \frac{1}{4} = \frac{3}{4}$.

Theorem 1.1.7. *Continuity of Probabilities*

(i) Continuity from below: if $A_n \uparrow A$, then $\lim_{n\to\infty} \mathbb{P}(A_n) = \mathbb{P}(A)$.
(ii) Continuity from above: if $A_n \downarrow A$, then $\lim_{n\to\infty} \mathbb{P}(A_n) = \mathbb{P}(A)$.

Proof.

(i) Set $B_1 = A_1$ and $B_n = A_n \setminus A_{n-1}$ for $n \geq 2$. Then (B_n) is disjoint, $A_n = \bigsqcup_{j=1}^{n} B_j$, and $A = \bigsqcup_{j=1}^{\infty} B_j$, so countable additivity gives $\mathbb{P}(A_n) = \sum_{j=1}^{n} \mathbb{P}(B_j) \to \sum_{j=1}^{\infty} \mathbb{P}(B_j) = \mathbb{P}(A)$.
(ii) Set $C_n = A_1 \setminus A_n$. Then $C_n \uparrow A_1 \setminus A$, so (i) gives $\mathbb{P}(C_n) \to \mathbb{P}(A_1 \setminus A)$; since $A_1 = A_n \sqcup C_n$ and $A_1 = A \sqcup (A_1 \setminus A)$, finite additivity gives $\mathbb{P}(A_n) = \mathbb{P}(A_1) - \mathbb{P}(C_n) \to \mathbb{P}(A_1) - \mathbb{P}(A_1 \setminus A) = \mathbb{P}(A)$.

Structural Commentary: Monotone limits convert countable additivity into limit rules.
Key Insight: Increasing uses increments; decreasing uses complements. □

Definition 1.1.8. *Events Occurring Infinitely Often and Eventually Always*
For events $(A_n)_{n\geq 1}$:

(i) *Tail union*: $B_n := \bigcup_{i=n}^{\infty} A_i$.
(ii) *Tail intersection*: $C_n := \bigcap_{i=n}^{\infty} A_i$.
(iii) *Limsup*: $\limsup_{n\to\infty} A_n := \bigcap_{n=1}^{\infty} B_n = \bigcap_{n=1}^{\infty} \bigcup_{i=n}^{\infty} A_i$.
(iv) *Liminf*: $\liminf_{n\to\infty} A_n := \bigcup_{n=1}^{\infty} C_n = \bigcup_{n=1}^{\infty} \bigcap_{i=n}^{\infty} A_i$.
(v) *Infinitely often*: $\omega \in \limsup_n A_n$ iff $\omega \in A_n$ for infinitely many n.
(vi) *Eventually always*: $\omega \in \liminf_n A_n$ iff $\omega \in A_n$ for all but finitely many n.

Proposition 1.1.9. *Tail Monotonicity*
For events $(A_n)_{n\geq 1}$ with tail union B_n and tail intersection C_n:

(i) Tail unions: $B_{n+1} \subseteq B_n$ for all n, hence $B_n \downarrow \bigcap_{n\geq 1} B_n$.
(ii) Tail intersections: $C_n \subseteq C_{n+1}$ for all n, hence $C_n \uparrow \bigcup_{n\geq 1} C_n$.

Proof.

(i) Since $\{i : i \geq n+1\} \subseteq \{i : i \geq n\}$, $B_{n+1} = \bigcup_{i=n+1}^{\infty} A_i \subseteq \bigcup_{i=n}^{\infty} A_i = B_n$.
(ii) Since $\{i : i \geq n+1\} \subseteq \{i : i \geq n\}$, $C_n = \bigcap_{i=n}^{\infty} A_i \subseteq \bigcap_{i=n+1}^{\infty} A_i = C_{n+1}$.

Structural Commentary: Tail shifts turn recurrence and permanence into monotone limits.
Key Insight: Tail unions track recurrence; tail intersections track permanence. □

Proposition 1.1.10. *Union Bound and Countable Intersections of Almost-Sure Events*
For events $(A_n)_{n\geq 1}$:

(i) Union bound: $\mathbb{P}(\bigcup_{n=1}^{\infty} A_n) \leq \sum_{n=1}^{\infty} \mathbb{P}(A_n)$.
(ii) Countable almost-sure intersection: if $\mathbb{P}(A_n) = 1$ for all n, then $\mathbb{P}(\bigcap_{n=1}^{\infty} A_n) = 1$.

Proof.

(i) Set $B_n := A_n \setminus \bigcup_{j<n} A_j$. Then (B_n) is disjoint, $\bigcup_n B_n = \bigcup_n A_n$, and $B_n \subseteq A_n$. Hence $\mathbb{P}(\bigcup_n A_n) = \sum_n \mathbb{P}(B_n) \leq \sum_n \mathbb{P}(A_n)$.
(ii) If $\mathbb{P}(A_n) = 1$ for all n, then $\mathbb{P}(A_n^c) = 0$ for all n. By (i), $\mathbb{P}((\bigcap_n A_n)^c) = \mathbb{P}(\bigcup_n A_n^c) \leq \sum_n \mathbb{P}(A_n^c) = 0$, so $\mathbb{P}(\bigcap_n A_n) = 1$.

Structural Commentary: Countable subadditivity is countable additivity after disjointization.
Key Insight: Almost sure statements survive countable intersections. □

Definition 1.1.11. *Uniform Probability on a Countable Space*
Let Ω be countable.

(i) A probability measure $\mathbb{P}$ on $(\Omega, \mathcal{P}(\Omega))$ is *uniform* if $\exists c \in [0,1]$ such that $\mathbb{P}(\{\omega\}) = c$ for all $\omega \in \Omega$.
(ii) The constant c is the *common atom mass*.

Proposition 1.1.12. *No Uniform Probability on a Countably Infinite Set*
Let $\Omega = \{0, 1, 2, \ldots\}$. No uniform probability measure exists on $(\Omega, \mathcal{P}(\Omega))$.

Proof. Assume uniform $\mathbb{P}$ exists, and let c be its common atom mass.
Since $\Omega = \bigsqcup_{n=0}^{\infty} \{n\}$, countable additivity gives $1 = \mathbb{P}(\Omega) = \sum_{n=0}^{\infty} \mathbb{P}(\{n\}) = \sum_{n=0}^{\infty} c$.
If $c = 0$, the sum is 0; if $c > 0$, the sum diverges. Both contradict $1 = \mathbb{P}(\Omega)$.
Structural Commentary: Countable additivity forces total mass to equal the sum of the atom masses.
Key Insight: Equal mass on countably many atoms is incompatible with total mass 1. □

Method 1.1.13. *Finite Uniform Probability by Counting*
Let $0 < |\Omega| < \infty$ and assume the *uniform model* $\mathbb{P}(\{\omega\}) = |\Omega|^{-1}$ for all $\omega \in \Omega$.
For $A \subseteq \Omega$, $\mathbb{P}(A) = \sum_{\omega \in A} \mathbb{P}(\{\omega\}) = |A|/|\Omega|$, so probability reduces to counting.

Definition 1.1.14. *Independent Events*

(i) Events $A, B \in \mathcal{F}$ are *independent*, written $A \perp\!\!\!\perp B$, if $\mathbb{P}(A \cap B) = \mathbb{P}(A)\mathbb{P}(B)$.
(ii) A family $(A_i)_{i \in I} \subseteq \mathcal{F}$ is *independent* if $\mathbb{P}(\bigcap_{i \in J} A_i) = \prod_{i \in J} \mathbb{P}(A_i)$ for every finite $J \subseteq I$.

Key Insight: Independence is multiplicative probability for simultaneous occurrence.

Proposition 1.1.15. *Basic Independence Tests*
Let $A, B \in \mathcal{F}$.

(i) If $A \cap B = \emptyset$ and $\mathbb{P}(A)\mathbb{P}(B) > 0$, then $A \not\perp\!\!\!\perp B$.
(ii) If $A \perp\!\!\!\perp B$, then $A^c \perp\!\!\!\perp B$, $A \perp\!\!\!\perp B^c$, and $A^c \perp\!\!\!\perp B^c$.
(iii) If $\mathbb{P}(A) \in \{0, 1\}$, then $A \perp\!\!\!\perp C$ for every $C \in \mathcal{F}$.
(iv) If $A \perp\!\!\!\perp A$, then $\mathbb{P}(A) \in \{0, 1\}$.

Proof.

(i) Since $A \cap B = \emptyset$, $\mathbb{P}(A \cap B) = 0 \neq \mathbb{P}(A)\mathbb{P}(B)$.
(ii) If $A \perp\!\!\!\perp B$, then $\mathbb{P}(A^c \cap B) = \mathbb{P}(B) - \mathbb{P}(A \cap B) = (1 - \mathbb{P}(A))\mathbb{P}(B) = \mathbb{P}(A^c)\mathbb{P}(B)$
and $\mathbb{P}(A \cap B^c) = \mathbb{P}(A) - \mathbb{P}(A \cap B) = \mathbb{P}(A)(1 - \mathbb{P}(B)) = \mathbb{P}(A)\mathbb{P}(B^c)$.
Applying the first identity to $A \perp\!\!\!\perp B^c$ gives $A^c \perp\!\!\!\perp B^c$.
(iii) Let $C \in \mathcal{F}$. If $\mathbb{P}(A) = 0$, then $0 \leq \mathbb{P}(A \cap C) \leq \mathbb{P}(A) = 0$, so $\mathbb{P}(A \cap C) = 0 = \mathbb{P}(A)\mathbb{P}(C)$.
If $\mathbb{P}(A) = 1$, then $\mathbb{P}(A^c) = 0$, hence $A^c \perp\!\!\!\perp C$ by the previous case and $A \perp\!\!\!\perp C$ by (ii).
(iv) Let $p = \mathbb{P}(A)$. From $A \perp\!\!\!\perp A$, $p = \mathbb{P}(A \cap A) = p^2$, hence $p \in \{0, 1\}$.

Structural Commentary: Independence is probabilistic, not visual; disjointness usually means strong dependence.
Key Insight: Complements preserve independence because B splits into $A \cap B$ and $A^c \cap B$. □

Example 1.1.16. *Independence, Complements, and First Success*

(i) *Independence by counting*: for a fair die, $\Omega = \{1, \ldots, 6\}$, $A = \{2, 4, 6\}$, and $B = \{1, 2, 3, 4\}$.
Then $\mathbb{P}(A \cap B) = 2/6 = (3/6)(4/6) = \mathbb{P}(A)\mathbb{P}(B)$, hence $A \perp\!\!\!\perp B$.
(ii) *Complement method*: for independent fair tosses $X_1, \ldots, X_{10}$, let $A = \{\exists j \leq 10 : X_j = H\}$
and $T_j = \{X_j = T\}$. Then $A^c = \bigcap_{j=1}^{10} T_j$ and $\mathbb{P}(A) = 1 - \mathbb{P}(A^c) = 1 - \prod_{j=1}^{10} \mathbb{P}(T_j) = 1 - (1/2)^{10}$.
(iii) *First-success decomposition*: let $p_1 = 1/3$, $p_2 = 1/4$.
Let A_j be player 1's first success on turn j before player 2 succeeds.
Then $E = \bigsqcup_{j \geq 1} A_j$, $\mathbb{P}(A_j) = ((1 - p_1)(1 - p_2))^{j-1} p_1 = (1/2)^{j-1}/3$, and $\mathbb{P}(E) = \sum_{j \geq 1} \mathbb{P}(A_j) = 2/3$.

Key Insight: Verify independence by products, use complements for at-least-one events, and decompose first successes into disjoint geometric pieces.

Method 1.1.17. *Independence, Complement, and Decomposition Methods*

(i) *Verify independence*: compute $\mathbb{P}(A)$, $\mathbb{P}(B)$, and $\mathbb{P}(A \cap B)$; conclude $A \perp\!\!\!\perp B$ iff $\mathbb{P}(A \cap B) = \mathbb{P}(A)\mathbb{P}(B)$.
(ii) *Use complements*: if A^c is simpler, compute $\mathbb{P}(A) = 1 - \mathbb{P}(A^c)$; if $A^c = \bigcap_{j=1}^{n} C_j$ with independent C_j.
Then $\mathbb{P}(A) = 1 - \prod_{j=1}^{n} \mathbb{P}(C_j)$.
(iii) *Disjoint decomposition*: write $E = \bigsqcup_j A_j$ with disjoint A_j, then $\mathbb{P}(E) = \sum_j \mathbb{P}(A_j)$.
(iv) *Geometric summation*: if $\mathbb{P}(A_j) = ar^{j-1}$ for $j \geq 1$ and $|r| < 1$, then $\mathbb{P}(E) = \sum_{j \geq 1} ar^{j-1} = a/(1 - r)$.

Definition 1.1.18. *Random Variables, Laws, and Distribution Functions*

(i) A *real-valued random variable* on $(\Omega, \mathcal{F}, \mathbb{P})$ is map $X : \Omega \to \mathbb{R}$ s.t. $\{\omega \in \Omega : X(\omega) \in B\} \in \mathcal{F}$ $\forall B \in \mathcal{B}(\mathbb{R})$.

(ii) Its *law* or *distribution* is $\mathcal{L}(X) : \mathcal{B}(\mathbb{R}) \to [0,1]$ given by $\mathcal{L}(X)(B) := \mathbb{P}(X^{-1}(B)) = \mathbb{P}(X \in B)$.
(iii) Its *cumulative distribution function* is $F_X : \mathbb{R} \to [0,1]$ given by $F_X(x) := \mathbb{P}(X \le x) = \mathcal{L}(X)((-\infty, x])$.

Definition 1.1.19. *Distribution Functions and Densities*

(i) *Distribution function* is nondecreasing right-continuous $F : \mathbb{R} \to [0,1]$, $\lim_{x\to-\infty} F(x) = 0$, $\lim_{x\to\infty} F(x) = 1$.
(ii) A *density* for F is an integrable $g : \mathbb{R} \to [0,\infty)$ with $\int_{-\infty}^{\infty} g(x)\,dx = 1$ and $F(t) = \int_{-\infty}^{t} g(x)\,dx$ $\forall t \in \mathbb{R}$. If g is continuous, then $F'(t) = g(t)$.

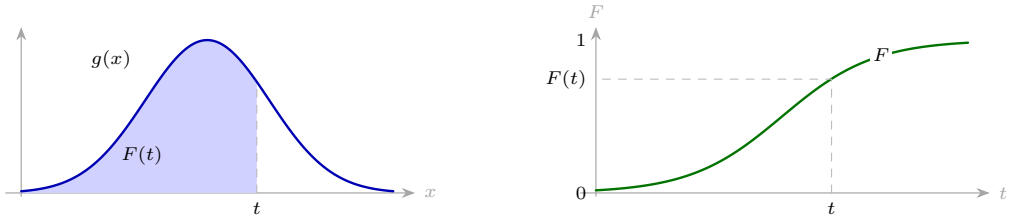


Figure 205: A normal density accumulates area into the distribution function $\Phi(t) = \int_{-\infty}^{t} g(x)\,dx$.

Definition 1.1.20. *Discrete and Continuous Random Variables*

(i) X is *discrete* if $X(\Omega) \subseteq S$ for some finite or countable $S = \{x_1, x_2, \ldots\} \subseteq \mathbb{R}$.
(ii) The *probability mass function* of discrete X is $f_X : S \to [0,1]$, $f_X(x) := \mathbb{P}(X = x)$, $\sum_{x\in S} f_X(x) = 1$.
(iii) The *discrete CDF* is $F_X(x) = \mathbb{P}(X \le x) = \sum_{x_i \le x} f_X(x_i)$.
(iv) X is *continuous* if its law has a density $f_X : \mathbb{R} \to [0,\infty)$.
(v) The *continuous probabilities* satisfy $\mathbb{P}(a < X < b) = \int_a^b f_X(t)\,dt$ and $F_X(x) = \int_{-\infty}^{x} f_X(t)\,dt$.
(vi) The *point-mass distinction*: for continuous X, $\mathbb{P}(X = x) = 0$ for every fixed x, so $f_X(x) \neq \mathbb{P}(X = x)$.

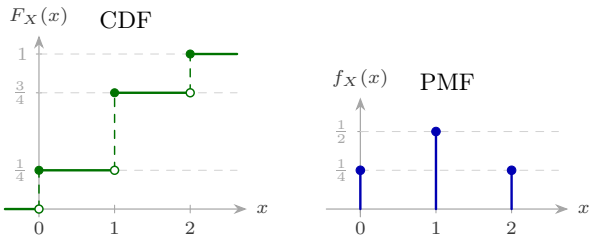


Figure 206: CDF and PMF for the number of heads in two independent fair coin tosses.

Example 1.1.21. *Counting Heads in Two Tosses*
Let $\Omega = \{H,T\}^2$ with $\mathbb{P}(\{\omega\}) = 1/4$, and set $X(\omega) = |\{j \in \{1,2\} : \omega_j = H\}|$.
Then $f_X(k) = \mathbb{P}(X = k) = \binom{2}{k}2^{-2}$ for $k \in \{0,1,2\}$, so $(f_X(0), f_X(1), f_X(2)) = (1/4, 1/2, 1/4)$.
Hence $F_X(x) = 0$ for $x < 0$, $F_X(x) = 1/4$ for $0 \le x < 1$, $F_X(x) = 3/4$ for $1 \le x < 2$, $F_X(x) = 1$ for $x \ge 2$.
Key Insight: A CDF is defined for every real x, even when a discrete random variable takes only a few values.

Proposition 1.1.22. *CDF Probability Calculus*
Let X have CDF F and define $F(x^-) := \lim_{y\uparrow x} F(y)$.

(i) Point mass: $\mathbb{P}(X = x) = F(x) - F(x^-)$.
(ii) Half-open interval: if $x < y$, then $\mathbb{P}(x < X \le y) = F(y) - F(x)$.
(iii) Right tail: $\mathbb{P}(X > x) = 1 - F(x)$.
(iv) Continuous endpoints: if X is continuous and $a < b$, then $\mathbb{P}(a < X < b) = \mathbb{P}(a \le X < b) = \mathbb{P}(a < X \le b) = \mathbb{P}(a \le X \le b)$.

Proof.

(i) Let $E_n = \{x - 1/n < X \le x\}$. Then $E_n \downarrow \{X = x\}$ and $\mathbb{P}(E_n) = F(x) - F(x - 1/n)$. Hence $\mathbb{P}(X = x) = \lim_n \mathbb{P}(E_n) = F(x) - F(x^-)$.
(ii) Since $\{X \le y\} = \{X \le x\} \sqcup \{x < X \le y\}$, $F(y) = F(x) + \mathbb{P}(x < X \le y)$.
(iii) Since $\{X > x\} = \{X \le x\}^c$, $\mathbb{P}(X > x) = 1 - \mathbb{P}(X \le x) = 1 - F(x)$.

(iv) If X is continuous, then $\mathbb{P}(X = t) = 0$ for all $t \in \mathbb{R}$. For $I_1 = (a, b)$, $I_2 = [a, b)$, $I_3 = (a, b]$, and $I_4 = [a, b]$, each $X^{-1}(I_j) \setminus X^{-1}(I_1) \subseteq \{X = a\} \cup \{X = b\}$, hence all four probabilities are equal.

Structural Commentary: CDF calculus subtracts nested lower-level events.
Key Insight: CDF jumps are point masses; continuous laws have none. □

Definition 1.1.23. *Quantiles and Equality in Distribution*

(i) The *quantile function* of a CDF F is $F^{-1}(q) := \inf\{x \in \mathbb{R} : F(x) \geq q\}$ for $q \in [0, 1]$.
(ii) If F is strictly increasing and continuous, then *inverse CDF* satisfies $F(F^{-1}(q)) = q$ and $F^{-1}(F(x)) = x$.
(iii) Random variables X, Y are *equal in distribution*, written $X \overset{\text{dist}}{=} Y$, if $F_X(x) = F_Y(x)$ for all $x \in \mathbb{R}$.

Definition 1.1.24. *Core Discrete Distributions*
Common discrete laws include:

(i) *Point mass*: $X \sim \delta_a$ for $a \in \mathbb{R}$ if $\mathbb{P}(X = a) = 1$.
(ii) *Discrete uniform*: $X \sim \text{Unif}\{1, \ldots, k\}$ for $k \in \mathbb{N}$ if $\mathbb{P}(X = x) = 1/k$ for $x = 1, \ldots, k$.
(iii) *Bernoulli*: $X \sim \text{Bernoulli}(p)$ for $p \in [0, 1]$ if $\mathbb{P}(X = 1) = p$ and $\mathbb{P}(X = 0) = 1 - p$.
(iv) *Binomial*: $X \sim \text{Binomial}(n, p)$ for $n \in \mathbb{N}$ and $p \in [0, 1]$ if $\mathbb{P}(X = x) = \binom{n}{x} p^x (1 - p)^{n-x}$ for $x = 0, \ldots, n$.
(v) *Geometric*: $X \sim \text{Geom}(p)$ for $p \in (0, 1]$ if $\mathbb{P}(X = k) = p(1 - p)^{k-1}$ for $k \geq 1$.
(vi) *Poisson*: $X \sim \text{Poisson}(\lambda)$ for $\lambda \geq 0$ if $\mathbb{P}(X = x) = e^{-\lambda}\lambda^x / x!$ for $x = 0, 1, 2, \ldots$.

Definition 1.1.25. *Core Continuous Distributions*
Common continuous laws include:

(i) *Uniform*: $X \sim \text{Uniform}(a, b)$ for $a < b$ if $f_X(x) = (b - a)^{-1}\mathbf{1}_{[a,b]}(x)$.
(ii) *Normal*: $X \sim N(\mu, \sigma^2)$ for $\mu \in \mathbb{R}$, $\sigma > 0$, if $f_X(x) = (\sigma\sqrt{2\pi})^{-1} \exp(-(x - \mu)^2/(2\sigma^2))$ for $x \in \mathbb{R}$. $N(0, 1)$ has density ϕ and CDF Φ.
(iii) *Exponential*: $X \sim \text{Exp}(\beta)$ for $\beta > 0$ if $f_X(x) = \beta^{-1} e^{-x/\beta}\mathbf{1}_{(0,\infty)}(x)$.
(iv) *Gamma*: $X \sim \text{Gamma}(\alpha, \beta)$ for $\alpha, \beta > 0$ if $f_X(x) = \beta^{-\alpha}\Gamma(\alpha)^{-1} x^{\alpha-1} e^{-x/\beta}\mathbf{1}_{(0,\infty)}(x)$; $\text{Exp}(\beta) = \text{Gamma}(1, \beta)$.
(v) *Beta*: $X \sim \text{Beta}(\alpha, \beta)$ for $\alpha, \beta > 0$ if $f_X(x) = \Gamma(\alpha + \beta)(\Gamma(\alpha)\Gamma(\beta))^{-1} x^{\alpha-1}(1 - x)^{\beta-1}\mathbf{1}_{(0,1)}(x)$.
(vi) *Cauchy*: $X \sim \text{Cauchy}(0, 1)$ if $f_X(x) = (\pi(1 + x^2))^{-1}$ for $x \in \mathbb{R}$.
(vii) χ^2: $X \sim \chi^2_p$ for $p > 0$ if $f_X(x) = (\Gamma(p/2)2^{p/2})^{-1} x^{p/2-1} e^{-x/2}\mathbf{1}_{(0,\infty)}(x)$.

Proposition 1.1.26. *Standardizing Normal Variables*
For $\sigma, \sigma_i > 0$:

(i) Standardization: if $X \sim N(\mu, \sigma^2)$, then $(X - \mu)/\sigma \sim N(0, 1)$.
(ii) Affine reconstruction: if $Z \sim N(0, 1)$, then $\mu + \sigma Z \sim N(\mu, \sigma^2)$.
(iii) Independent sums: if $X_i \sim N(\mu_i, \sigma_i^2)$ are independent, then $\sum_{i=1}^n X_i \sim N(\sum_{i=1}^n \mu_i, \sum_{i=1}^n \sigma_i^2)$.
(iv) Interval probabilities: if $X \sim N(\mu, \sigma^2)$ and $a < b$, then $\mathbb{P}(a < X < b) = \Phi((b - \mu)/\sigma) - \Phi((a - \mu)/\sigma)$.

Proof.

(i) Let $Y = (X - \mu)/\sigma$. By change of variables, $f_Y(z) = \sigma f_X(\mu + \sigma z) = (2\pi)^{-1/2} e^{-z^2/2}$, so $Y \sim N(0, 1)$.
(ii) Let $Y = \mu + \sigma Z$. Then $f_Y(y) = \sigma^{-1}\phi((y - \mu)/\sigma) = (\sigma\sqrt{2\pi})^{-1} \exp(-(y - \mu)^2/(2\sigma^2))$, hence $Y \sim N(\mu, \sigma^2)$.
(iii) Let $S = X + Y$, $m_s = (v_2\mu_1 + v_1(s - \mu_2))/(v_1 + v_2)$, and $w = v_1 v_2/(v_1 + v_2)$.
Since $(u - \mu_1)^2/v_1 + (s - u - \mu_2)^2/v_2 = (s - \mu_1 - \mu_2)^2/(v_1 + v_2) + (u - m_s)^2/w$,
convolution gives $f_S(s) = (2\pi(v_1 + v_2))^{-1/2} \exp(-(s - \mu_1 - \mu_2)^2/(2(v_1 + v_2)))$; induction gives (iii).
(iv) Let $Z = (X - \mu)/\sigma$. Since $\sigma > 0$, $\{a < X < b\} = \{(a - \mu)/\sigma < Z < (b - \mu)/\sigma\}$.
Hence $\mathbb{P}(a < X < b) = \Phi((b - \mu)/\sigma) - \Phi((a - \mu)/\sigma)$.

Structural Commentary: Normal laws are stable under affine maps and independent addition.
Key Insight: Standardize to $N(0, 1)$, compute with Φ, then transform back. □

Example 1.1.27. *Normal Tail and Quantile Standardization*
Let $X \sim N(3, 5)$ and $Z = (X - 3)/\sqrt{5} \sim N(0, 1)$. Then $\mathbb{P}(X > 1) = \mathbb{P}(Z > -2/\sqrt{5}) = 1 - \Phi(-2/\sqrt{5}) \approx 0.81$.
If $\mathbb{P}(X < q) = 0.2$, then $\Phi((q - 3)/\sqrt{5}) = 0.2$, so $q = 3 + \sqrt{5}\,\Phi^{-1}(0.2) \approx 3 - 0.8416\sqrt{5} \approx 1.1181$.

Definition 1.1.28. *Joint and Marginal Distributions*

(i) The *joint mass function* of discrete X, Y is $f_{X,Y}(x, y) := \mathbb{P}(X = x, Y = y)$.
(ii) The *discrete marginals* are $f_X(x) = \sum_y f_{X,Y}(x, y)$ and $f_Y(y) = \sum_x f_{X,Y}(x, y)$.
(iii) If (X, Y) has joint density $f_{X,Y}$, the *continuous marginals* are $f_X(x) = \int_{-\infty}^{\infty} f_{X,Y}(x, y)\,dy$ and $f_Y(y) = \int_{-\infty}^{\infty} f_{X,Y}(x, y)\,dx$.
(iv) The *joint CDF* is $F_{X,Y}(x, y) := \mathbb{P}(X \leq x, Y \leq y)$.

Definition 1.1.29. *Independent Random Variables and IID Samples*

(i) Random variables $X_1, \dots, X_n$ are *independent* if $\mathbb{P}(X_1 \in B_1, \dots, X_n \in B_n) = \prod_{i=1}^n \mathbb{P}(X_i \in B_i)$ for all $B_i \in \mathcal{B}(\mathbb{R})$.
(ii) They are *identically distributed* if $\mathcal{L}(X_1) = \cdots = \mathcal{L}(X_n)$.
(iii) They are *iid* if they are independent and identically distributed.
(iv) The *random vector* $X = (X_1, \dots, X_n)$ is a *random sample* when $X_1, \dots, X_n$ are iid from a common law.

Proposition 1.1.30. *Factorization Tests for Independence*

(i) Discrete factorization: discrete X, Y are independent iff $f_{X,Y}(x,y) = f_X(x) f_Y(y)$ for all x, y.
(ii) Density factorization: if (X, Y) has joint density $f_{X,Y}$ and marginals f_X, f_Y, then X, Y are independent iff $f_{X,Y}(x,y) = f_X(x) f_Y(y)$ a.e.
(iii) Rectangular support test: if $f_{X,Y}(x,y) = g(x)h(y)\mathbf{1}_{A\times B}(x,y)$ for measurable A, B, then X, Y are independent after normalization.

Proof.

(i) ($\Rightarrow$) If $X \perp\!\!\!\perp Y$, then $f_{X,Y}(x,y) = \mathbb{P}(X = x, Y = y) = \mathbb{P}(X = x)\mathbb{P}(Y = y) = f_X(x) f_Y(y)$.
(i) ($\Leftarrow$) If $f_{X,Y} = f_X f_Y$, then for A, B, $\mathbb{P}(X \in A, Y \in B) = \sum_{x\in A}\sum_{y\in B} f_X(x) f_Y(y) = \mathbb{P}(X \in A)\mathbb{P}(Y \in B)$.
(ii) ($\Rightarrow$) If $X \perp\!\!\!\perp Y$, the product law has density $f_X f_Y$, so uniqueness of densities gives $f_{X,Y} = f_X f_Y$ a.e.
(ii) ($\Leftarrow$) If $f_{X,Y} = f_X f_Y$ a.e., then for Borel A, B, $\mathbb{P}(X \in A, Y \in B) = \int_A \int_B f_X(x) f_Y(y)\, dy\, dx = \mathbb{P}(X \in A)\mathbb{P}(Y \in B)$.
(iii) Let $G = \int_A g$ and $H = \int_B h$. Normalization gives $GH = 1$, while $f_X(x) = g(x)H\mathbf{1}_A(x)$ and $f_Y(y) = h(y)G\mathbf{1}_B(y)$, hence $f_X(x) f_Y(y) = f_{X,Y}(x,y)$ a.e.; apply (ii).

Structural Commentary: Independence is product structure of the joint law.
Key Insight: Independence means the joint law factors. □

Example 1.1.31. *A Non-Rectangular Bivariate Density*

Let $S = \{(x,y) : x^2 \le y \le 1\}$ and $f_{X,Y}(x,y) = cx^2 y \mathbf{1}_S(x,y)$.
Since $S = \{(x,y) : -1 \le x \le 1,\ x^2 \le y \le 1\}$, $1 = c\int_{-1}^1 \int_{x^2}^1 x^2 y\, dy\, dx = 4c/21$, so $c = 21/4$.
For $E = \{X \ge Y\}$, $S \cap E = \{(x,y) : 0 \le x \le 1,\ x^2 \le y \le x\}$, hence $\mathbb{P}(E) = (21/4)\int_0^1 \int_{x^2}^x x^2 y\, dy\, dx = 3/20$.
Key Insight: In bivariate integration, pick one variable first and let the other range over the corresponding slice.

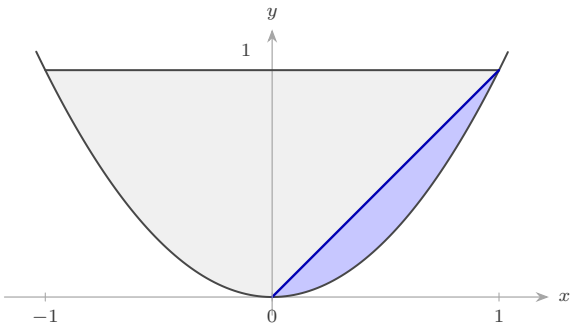


Figure 207: Gray: support $x^2 \le y \le 1$; blue: event $X \ge Y$ ($x^2 \le y \le x$).

Definition 1.1.32. *Two Multivariate Families*

(i) *Multinomial*: for $n \in \mathbb{N}$ and $p \in [0,1]^k$ with $\sum_{j=1}^k p_j = 1$, $X \sim \text{Multinomial}(n,p)$ if $\mathbb{P}(X = x) = n!(x_1! \cdots x_k!)^{-1} \prod_{j=1}^k p_j^{x_j}$ for $x \in \mathbb{N}_0^k$ with $\sum_{j=1}^k x_j = n$; then $X_j \sim \text{Binomial}(n, p_j)$.
(ii) *Multivariate normal*: for $\mu \in \mathbb{R}^k$ and symmetric positive definite $\Sigma \in \mathbb{R}^{k\times k}$, $X \sim N(\mu, \Sigma)$ if $f_X(x) = ((2\pi)^k \det \Sigma)^{-1/2} \exp(-(x-\mu)^t \Sigma^{-1}(x-\mu)/2)$ for $x \in \mathbb{R}^k$. $N(0, I_k)$ has independent $N(0,1)$ coordinates.

Proposition 1.1.33. *Multivariate Normal Affine Standardization*

Let $\mu \in \mathbb{R}^k$ and $\Sigma \in \mathbb{R}^{k\times k}$ with $\Sigma \succ 0$. For a random vector $X \in \mathbb{R}^k$, $X \sim N(\mu, \Sigma)$ iff $Z := \Sigma^{-1/2}(X - \mu) \sim N(0, I_k)$; equivalently, if $Z \sim N(0, I_k)$, then $\mu + \Sigma^{1/2} Z \sim N(\mu, \Sigma)$.

Proof.

($\Rightarrow$) Set $Z = \Sigma^{-1/2}(X - \mu)$, so $x = \mu + \Sigma^{1/2} z$ and $|\det \Sigma^{1/2}| = (\det \Sigma)^{1/2}$.
By change of variables, $f_Z(z) = f_X(\mu + \Sigma^{1/2} z)(\det \Sigma)^{1/2} = (2\pi)^{-k/2} \exp(-z^t z/2)$, hence $Z \sim N(0, I_k)$.

($\Leftarrow$) Let $X = \mu + \Sigma^{1/2}Z$ with $Z \sim N(0, I_k)$. Since $z = \Sigma^{-1/2}(x-\mu)$ and $|\det \Sigma^{-1/2}| = (\det \Sigma)^{-1/2}$, $f_X(x) = f_Z(\Sigma^{-1/2}(x-\mu))(\det \Sigma)^{-1/2} = ((2\pi)^k \det \Sigma)^{-1/2} \exp(-(x-\mu)^t\Sigma^{-1}(x-\mu)/2)$, so $X \sim N(\mu, \Sigma)$.

Structural Commentary: Positive definite covariance is an affine deformation of the standard normal.
Key Insight: Standardize by $\Sigma^{-1/2}$; reconstruct by $\Sigma^{1/2}$. □

Proposition 1.1.34. *Basic Multivariate Normal Calculus*

Let $X = (X_a, X_b) \sim N(\mu, \Sigma)$ *with* $X_a \in \mathbb{R}^r$, $X_b \in \mathbb{R}^s$, $\mu = (\mu_a, \mu_b)$, *and* $\Sigma = \begin{pmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{pmatrix} \succ 0$.

(i) Marginal: $X_a \sim N(\mu_a, \Sigma_{aa})$.
(ii) Conditional: $X_b \mid X_a = x_a \sim N(\mu_b + \Sigma_{ba}\Sigma_{aa}^{-1}(x_a - \mu_a), \Sigma_{bb} - \Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{ab})$.
(iii) Linear form: *for any vector* a, $a^tX \sim N(a^t\mu, a^t\Sigma a)$.

Proof.

(i) Let $P_aX = X_a$. By Proposition 1.1.33, $X = \mu + \Sigma^{1/2}Z$ with $Z \sim N(0, I)$.
Hence $X_a = P_a\mu + P_a\Sigma^{1/2}Z \sim N(\mu_a, P_a\Sigma P_a^t) = N(\mu_a, \Sigma_{aa})$.

(ii) Set $S := \Sigma_{bb} - \Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{ab}$ and $m(x_a) := \mu_b + \Sigma_{ba}\Sigma_{aa}^{-1}(x_a - \mu_a)$.
The block determinant and inverse identities give $\det \Sigma = \det \Sigma_{aa} \det S$ and
$(x-\mu)^t\Sigma^{-1}(x-\mu) = (x_a - \mu_a)^t\Sigma_{aa}^{-1}(x_a - \mu_a) + (x_b - m(x_a))^tS^{-1}(x_b - m(x_a))$.
Thus $f_{b|a}(x_b \mid x_a) = f_{a,b}(x_a, x_b)/f_a(x_a)$ is the $N(m(x_a), S)$ density.

(iii) With $X = \mu + \Sigma^{1/2}Z$, $a^tX = a^t\mu + (\Sigma^{1/2}a)^tZ$.
Since Z has independent $N(0,1)$ coordinates, $(\Sigma^{1/2}a)^tZ \sim N(0, a^t\Sigma a)$, hence $a^tX \sim N(a^t\mu, a^t\Sigma a)$.

Structural Commentary: Marginals, conditionals, and linear projections of Gaussian vectors remain Gaussian.
Key Insight: Gaussian calculus is affine calculus plus Schur complements. □

Method 1.1.35. *One-Variable Transformations*
Let $Y = r(X)$.

(i) *Discrete pushforward*: if X is discrete, then $f_Y(y) = \mathbb{P}(Y = y) = \mathbb{P}(X \in r^{-1}(\{y\})) = \sum_{x:r(x)=y} f_X(x)$.
(ii) *CDF method*: if X has density f_X, set $A_y := \{x : r(x) \le y\}$, compute $F_Y(y) = \mathbb{P}(Y \le y) = \int_{A_y} f_X(x)\,dx$, and use $f_Y(y) = F_Y'(y)$ where differentiable.
(iii) *Monotone change of variables*: if r is strictly monotone with inverse $s = r^{-1}$, then $f_Y(y) = f_X(s(y))|s'(y)|$ on $r(\operatorname{supp} X)$.

Key Insight: The CDF method works because it translates the transformed event back into an event about X.

Example 1.1.36. *Log Transform of an Exponential Variable*
Let $X \sim \text{Exp}(1)$ and $Y = \log X$. Since $r(x) = \log x$ has inverse $s(y) = e^y$ with $s'(y) = e^y$, the change-of-variables formula gives $f_Y(y) = f_X(e^y)e^y = e^ye^{-e^y}$ for $y \in \mathbb{R}$.

Example 1.1.37. *A Nonmonotone Square Transform*
Let $X \sim \text{Uniform}(-1, 3)$ and $Y = X^2$. Since $f_X = 1/4$ on $[-1, 3]$, $r^{-1}(\{y\}) \cap [-1, 3] = \{-\sqrt{y}, \sqrt{y}\}$ for $0 < y < 1$ and $\{\sqrt{y}\}$ for $1 < y < 9$. Hence $f_Y(y) = 1/(4\sqrt{y})\mathbf{1}_{(0,1)}(y) + 1/(8\sqrt{y})\mathbf{1}_{(1,9)}(y)$.
Key Insight: Non-one-to-one transformations require splitting the range into cases.

Method 1.1.38. *Several-Variable Transformations*
For $Z = r(X, Y)$ with joint density $f_{X,Y}$:

(i) *Sublevel set*: define $A_z := \{(x, y) : r(x, y) \le z\}$.
(ii) *CDF integral*: compute $F_Z(z) = \mathbb{P}(Z \le z) = \iint_{A_z} f_{X,Y}(x, y)\,dx\,dy$.
(iii) *Density*: differentiate where possible, $f_Z(z) = F_Z'(z)$.

Example 1.1.39. *Sum of Two Independent Uniforms*
Let $X_1, X_2 \sim \text{Uniform}(0, 1)$ be independent and $Y = X_1 + X_2$. Since $f_{X_1,X_2} = \mathbf{1}_{[0,1]^2}$, $F_Y(y)$ is the area of $A_y = \{(x_1, x_2) \in [0, 1]^2 : x_1 + x_2 \le y\}$, so $F_Y(y) = (y^2/2)\mathbf{1}_{[0,1]}(y) + (1 - (2-y)^2/2)\mathbf{1}_{(1,2]}(y) + \mathbf{1}_{(2,\infty)}(y)$. Hence $f_Y(y) = y\mathbf{1}_{(0,1)}(y) + (2-y)\mathbf{1}_{(1,2)}(y)$.

1.2 Expectation and Integration in Probability

Expectation is the averaging operation attached to a probability law. In the notation of Section 1.1.1, if X has law F_X, then integration against F_X is the common notation that covers both sums for discrete variables and ordinary integrals for variables with densities.

Definition 1.2.1. *Expectation Through the Law*
Let X be real-valued with distribution function F_X.

(i) *Expectation*: if $\int_{\mathbb{R}} |x|\,dF_X(x) < \infty$, then $\mathbb{E}(X) := \int_{\mathbb{R}} x\,dF_X(x)$.

(ii) *Discrete form*: if X has support S and mass function f_X, then $\mathbb{E}(X) = \sum_{x\in S} x f_X(x)$ when $\sum_{x\in S} |x| f_X(x) < \infty$.
(iii) *Density form*: if X has density f_X, then $\mathbb{E}(X) = \int_{\mathbb{R}} x f_X(x)\,dx$ when $\int_{\mathbb{R}} |x| f_X(x)\,dx < \infty$.
(iv) *Function form*: for measurable $r : \mathbb{R} \to \mathbb{R}$, if $\int_{\mathbb{R}} |r(x)|\,dF_X(x) < \infty$, then $\mathbb{E}(r(X)) := \int_{\mathbb{R}} r(x)\,dF_X(x)$.

Example 1.2.2. *Expectation of the Two-Toss Head Count*
Let X count the number of heads in two fair coin tosses. From Example 1.1.21,

$$f_X(0) = \frac{1}{4}, \qquad f_X(1) = \frac{1}{2}, \qquad f_X(2) = \frac{1}{4}.$$

Thus

$$\mathbb{E}(X) = 0 \cdot \frac{1}{4} + 1 \cdot \frac{1}{2} + 2 \cdot \frac{1}{4} = 1.$$

Definition 1.2.3. *Variance and Covariance*
Assume the displayed expectations exist. The *variance* of X is

$$\operatorname{Var}(X) := \mathbb{E}\big[(X - \mathbb{E}X)^2\big],$$

and the *standard deviation* is $\sqrt{\operatorname{Var}(X)}$. For two random variables X and Y, their *covariance* is

$$\operatorname{Cov}(X, Y) := \mathbb{E}\big[(X - \mathbb{E}X)(Y - \mathbb{E}Y)\big].$$

Proposition 1.2.4. *Indicators Turn Probabilities into Expectations*
For an event A, let $\mathbf{1}_A$ denote its indicator:

$$\mathbf{1}_A(\omega) = \begin{cases} 1, & \omega \in A, \\ 0, & \omega \notin A. \end{cases}$$

Then

$$\mathbb{E}(\mathbf{1}_A) = \mathbb{P}(A).$$

Proof. The indicator takes value 1 on A and value 0 on A^c, so its expectation is

$$1 \cdot \mathbb{P}(A) + 0 \cdot \mathbb{P}(A^c) = \mathbb{P}(A).$$

□

Remark 1.2.5. *Source Scope*
The source pages supplied here develop expectation first as integration against a distribution, then use it to organize variance, covariance, concentration inequalities, and moment generating functions. The measure-theoretic shorthand $\int r(x)\,dF_X(x)$ is interpreted through the law of X; the abstract Lebesgue construction is the one developed in the measure and integration part.

Example 1.2.6. *First Expectation Computations*
(i) If $X \sim \text{Bernoulli}(p)$, then $\mathbb{E}X = p$.
(ii) If $X \sim \text{Uniform}(-1, 3)$, then $\mathbb{E}X = \frac{1}{4}\int_{-1}^{3} x\,dx = 1$.
(iii) If X has Cauchy density $f_X(x) = 1/(\pi(1 + x^2))$, then

$$\int_{\mathbb{R}} |x|\, f_X(x)\,dx = \frac{2}{\pi}\int_0^{\infty} \frac{x}{1 + x^2}\,dx = +\infty.$$

Thus the Cauchy mean does not exist; sample averages from a Cauchy law need not settle down around a finite number.

Theorem 1.2.7. *Expectation of a Function*
Let X be a random variable and let r be measurable. Whenever $r(X)$ is integrable,

$$\mathbb{E}(r(X)) = \int r(x)\,dF_X(x).$$

For a function of several random variables, the same rule uses the joint law:

$$\mathbb{E}(r(X, Y)) = \iint r(x, y)\,dF_{X,Y}(x, y).$$

Proof. This is the defining property of the pushforward law $\mathcal{L}(X)$, applied to r. The joint version is the same statement for the random vector (X, Y). □

Example 1.2.8. *Using the Function Rule*
(i) If $X \sim \text{Uniform}(0, 1)$ and $Y = e^X$, then

$$\mathbb{E}Y = \int_0^1 e^x\,dx = e - 1.$$

(ii) Break a unit stick at a uniformly chosen point X. If $Y = \max\{X, 1-X\}$ is the longer piece, then

$$\mathbb{E}Y = \int_0^{1/2} (1-x)\,dx + \int_{1/2}^1 x\,dx = \frac{3}{4}.$$

(iii) If (X,Y) is uniform on $[0,1]^2$ and $Z = X^2 + Y^2$, then

$$\mathbb{E}Z = \int_0^1 \int_0^1 (x^2+y^2)\,dx\,dy = \frac{2}{3}.$$

Definition 1.2.9. *Moments, Correlation, and Sample Statistics*
The *kth moment* of X is $\mathbb{E}(X^k)$, assuming $\mathbb{E}\,|X|^k < +\infty$. If the standard deviations are positive, the *correlation* of X and Y is

$$\operatorname{Corr}(X,Y) := \frac{\operatorname{Cov}(X,Y)}{\sqrt{\operatorname{Var}(X)}\sqrt{\operatorname{Var}(Y)}}.$$

For random variables $X_1, \dots, X_n$, define the *sample mean* and *sample variance* by

$$\overline{X}_n := \frac{1}{n}\sum_{i=1}^n X_i, \qquad S_n^2 := \frac{1}{n-1}\sum_{i=1}^n (X_i - \overline{X}_n)^2.$$

Theorem 1.2.10. *Expectation, Variance, and Covariance Algebra*
Assume the displayed quantities exist.

(i) $\mathbb{E}(\sum_i a_i X_i) = \sum_i a_i \mathbb{E}X_i$.
(ii) If $X_1, \dots, X_n$ *are independent, then* $\mathbb{E}(\prod_i X_i) = \prod_i \mathbb{E}X_i$.
(iii) $\operatorname{Var}(X) = \mathbb{E}(X^2) - (\mathbb{E}X)^2$ *and* $\operatorname{Var}(aX+b) = a^2 \operatorname{Var}(X)$.
(iv) $\operatorname{Cov}(X,Y) = \mathbb{E}(XY) - \mathbb{E}X\,\mathbb{E}Y$.
(v) For constants a_i,

$$\operatorname{Var}\left(\sum_i a_i X_i\right) = \sum_i a_i^2 \operatorname{Var}(X_i) + 2\sum_{i<j} a_i a_j \operatorname{Cov}(X_i, X_j).$$

Proof. Linearity is the linearity of the integral. The product rule follows by factoring the joint law under independence. The remaining identities come from expanding squares and bilinear products. □

Example 1.2.11. *Indicator Sums as a Method*

(i) If $X \sim \text{Binomial}(n,p)$, write $X = \sum_{i=1}^n X_i$ with independent Bernoulli(p) indicators. Then

$$\mathbb{E}X = np, \qquad \operatorname{Var}(X) = np(1-p).$$

(ii) If a particle moves one unit left with probability p and one unit right with probability $1-p$, then its position after n jumps has

$$\mathbb{E}X_n = n(1-2p), \qquad \operatorname{Var}(X_n) = 4np(1-p).$$

(iii) If a fortune starts at c and each play independently doubles or halves it with equal probability, then

$$\mathbb{E}(\text{fortune after } n \text{ plays}) = c\left(\frac{5}{4}\right)^n.$$

Proposition 1.2.12. *Zero Variance Means Almost Surely Constant*
For a square-integrable random variable X, $\operatorname{Var}(X) = 0$ *if and only if* $X = \mathbb{E}X$ *almost surely.*

Proof. The forward direction follows because $(X - \mathbb{E}X)^2 \geq 0$ has integral 0, so it vanishes almost everywhere. □

Proposition 1.2.13. *Tail Integral Formula*
If $X \geq 0$ *has density* f_X, *distribution function* F_X, *and finite expectation, then*

$$\mathbb{E}X = \int_0^\infty \mathbb{P}(X > x)\,dx.$$

Proof. For $b > 0$, integration by parts gives

$$\int_0^b x f_X(x)\,dx = -b(1 - F_X(b)) + \int_0^b (1 - F_X(x))\,dx.$$

The boundary term tends to 0 under the finite-expectation hypothesis. □

Example 1.2.14. *Maximum of Uniform Samples*
If $X_1, \dots, X_n$ are independent Uniform$(0,1)$ variables and $Y_n = \max_i X_i$, then $\mathbb{P}(Y_n \le y) = y^n$ for $0 \le y \le 1$. Hence

$$\mathbb{E}Y_n = \int_0^1 (1 - y^n)\, dy = \frac{n}{n+1}.$$

Theorem 1.2.15. *Mean and Variance of the Basic Sample Statistics*
Let $X_1, \dots, X_n$ be IID with mean μ and variance $\sigma^2 < +\infty$. Then

$$\mathbb{E}\overline{X}_n = \mu, \qquad \mathrm{Var}(\overline{X}_n) = \frac{\sigma^2}{n}, \qquad \mathbb{E}S_n^2 = \sigma^2.$$

Proof. The first two identities use Theorem 1.2.10. For S_n^2, use

$$\sum_{i=1}^n (X_i - \overline{X}_n)^2 = \sum_{i=1}^n (X_i - \mu)^2 - n(\overline{X}_n - \mu)^2$$

and take expectations. □

Definition 1.2.16. *Mean Vector and Covariance Matrix*
For a random vector $X = (X_1, \dots, X_k)^t$, define

$$\mathbb{E}X = (\mathbb{E}X_1, \dots, \mathbb{E}X_k)^t, \qquad \mathrm{Var}(X) = \big[\mathrm{Cov}(X_i, X_j)\big]_{i,j=1}^k.$$

Proposition 1.2.17. *Linear Transformations of Random Vectors*
If X has mean μ and covariance matrix Σ, then

$$\mathbb{E}(a^t X) = a^t \mu, \qquad \mathrm{Var}(a^t X) = a^t \Sigma a,$$

and, for a matrix A,

$$\mathbb{E}(AX) = A\mu, \qquad \mathrm{Var}(AX) = A\Sigma A^t.$$

Example 1.2.18. *Multinomial Covariance Structure*
If $X \sim \mathrm{Multinomial}(n, \mathbf{p})$ with $\mathbf{p} = (p_1, \dots, p_k)$, then

$$\mathbb{E}X = n\mathbf{p}, \qquad \mathrm{Var}(X_i) = np_i(1 - p_i), \qquad \mathrm{Cov}(X_i, X_j) = -np_ip_j \quad (i \ne j).$$

Compare $\mathrm{Var}(X_i + X_j)$ with the binomial variance formula to get the covariance.

Method 1.2.19. *Common Means and Variances*

Distribution	$\mathbb{E}X$	$\mathrm{Var}(X)$
point mass at a	a	0
Bernoulli(p)	p	$p(1-p)$
Binomial(n,p)	np	$np(1-p)$
Geometric(p)	$1/p$	$(1-p)/p^2$
Poisson(λ)	λ	λ
Uniform(a,b)	$(a+b)/2$	$(b-a)^2/12$
$N(\mu, \sigma^2)$	μ	σ^2
Exp(β)	β	β^2
Gamma(α, β)	$\alpha\beta$	$\alpha\beta^2$
Beta(α, β)	$\dfrac{\alpha}{\alpha+\beta}$	$\dfrac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$
χ^2_p	p	$2p$

Theorem 1.2.20. *Markov and Chebyshev Inequalities*
If $X \ge 0$ and $\mathbb{E}X < +\infty$, then for $t > 0$,

$$\mathbb{P}(X > t) \le \frac{\mathbb{E}X}{t}.$$

If $\mu = \mathbb{E}X$ and $\sigma^2 = \mathrm{Var}(X) < +\infty$, then for $t > 0$,

$$\mathbb{P}(|X - \mu| \ge t) \le \frac{\sigma^2}{t^2}.$$

Proof. For Markov, $X \ge t\mathbf{1}_{\{X > t\}}$. Chebyshev applies Markov to $(X - \mu)^2$. □

Example 1.2.21. *A Distribution-Free Error-Rate Bound*
Let X_i indicate an error made by a prediction method on the ith test case. If the true error rate is p, then

$$\mathbb{P}(|\overline{X}_n - p| > \varepsilon) \le \frac{p(1-p)}{n\varepsilon^2} \le \frac{1}{4n\varepsilon^2}.$$

Theorem 1.2.22. *Hoeffding's Bounded-Sum Bound*
Let $Y_1, \dots, Y_n$ be independent with $\mathbb{E}Y_i = 0$ and $a_i \le Y_i \le b_i$. For every $t > 0$ and $\varepsilon > 0$,

$$\mathbb{P}\left(\sum_{i=1}^{n} Y_i \ge \varepsilon\right) \le e^{-t\varepsilon} \prod_{i=1}^{n} \exp\left(\frac{t^2(b_i - a_i)^2}{8}\right).$$

Proof idea. Apply Markov's inequality to $\exp(t\sum_i Y_i)$, use independence to factor the MGF, and bound each bounded zero-mean exponential moment by the convexity/Taylor estimate

$$\mathbb{E}e^{tY_i} \le \exp\left(\frac{t^2(b_i - a_i)^2}{8}\right).$$

□

Corollary 1.2.23. *Hoeffding for Bernoulli Averages*
If $X_1, \dots, X_n$ are independent Bernoulli(p) *variables and $\overline{X}_n = n^{-1}\sum_i X_i$, then*

$$\mathbb{P}(\left|\overline{X}_n - p\right| > \varepsilon) \le 2e^{-2n\varepsilon^2}.$$

Proof. Apply Theorem 1.2.22 to $Y_i = (X_i - p)/n$, optimize the one-sided bound at $t = 4n\varepsilon$, and repeat for $p - X_i$. □

Corollary 1.2.24. *Hoeffding Confidence Interval for a Bernoulli Parameter*
Let $\widehat{p}_n = \overline{X}_n$ and fix $\alpha > 0$. With

$$\varepsilon_n = \sqrt{\frac{1}{2n}\log\left(\frac{2}{\alpha}\right)}, \qquad C_n = (\widehat{p}_n - \varepsilon_n, \widehat{p}_n + \varepsilon_n),$$

we have $\mathbb{P}(p \in C_n) \ge 1 - \alpha$. The interval length is at most 0.05 whenever

$$n \ge 800 \log\left(\frac{2}{\alpha}\right).$$

Theorem 1.2.25. *Cauchy–Schwarz and Jensen*
If $X, Y \in L^2$, then

$$\mathbb{E}\,|XY| \le \sqrt{\mathbb{E}(X^2)\mathbb{E}(Y^2)}.$$

If g is convex and the expectations exist, then

$$\mathbb{E}g(X) \ge g(\mathbb{E}X),$$

with the reverse inequality for concave g.

Proof. The first inequality is Cauchy–Schwarz in L^2. For Jensen, use a supporting line L to g at $\mathbb{E}X$: then $g \ge L$ and

$$\mathbb{E}g(X) \ge \mathbb{E}L(X) = L(\mathbb{E}X) = g(\mathbb{E}X).$$

□

Definition 1.2.26. *Moment Generating Function*
The *moment generating function* of X is

$$\psi_X(t) := \mathbb{E}(e^{tX}),$$

for values of t where the expectation is finite. If it is finite near 0, then

$$\psi_X^{(k)}(0) = \mathbb{E}(X^k).$$

Proposition 1.2.27. *Basic MGF Properties*
Assume the MGFs are finite near 0.

(i) If $Y = aX + b$, then $\psi_Y(t) = e^{bt}\psi_X(at)$.
(ii) If $X_1, \dots, X_n$ are independent and $Y = \sum_i X_i$, then $\psi_Y(t) = \prod_i \psi_{X_i}(t)$.
(iii) If two MGFs agree on an open interval around 0, then the corresponding laws agree.

Example 1.2.28. *MGF Computations* If $X \sim \mathrm{Exp}(1)$, then $\psi_X(t) = 1/(1-t)$ for $t < 1$, so $\mathbb{E}X = 1$ and $\mathrm{Var}(X) = 1$. If $X \sim \mathrm{Binomial}(n, p)$, then

$$\psi_X(t) = \left(pe^t + 1 - p\right)^n.$$

Method 1.2.29. *Common MGFs*

Distribution	$\psi_X(t)$
Bernoulli(p)	$pe^t + (1-p)$
Binomial(n,p)	$\left(pe^t + (1-p)\right)^n$
Poisson(λ)	$e^{\lambda(e^t-1)}$
$N(\mu,\sigma^2)$	$\exp\left(\mu t + \frac{\sigma^2 t^2}{2}\right)$
Gamma(α,β)	$(1-\beta t)^{-\alpha} \quad (t < 1/\beta)$

1.3 Conditional Probability and Conditional Expectation

Definition 1.3.1. *Conditional Probability*
If $B \in \mathcal{F}$ and $\mathbb{P}(B) > 0$, the *conditional probability* of A given B is

$$\mathbb{P}(A \mid B) := \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)}.$$

Think of this as restricting attention to the repetitions in which B occurs.

Proposition 1.3.2. *Conditioning on a Fixed Event Gives a Probability Measure*
Fix $B \in \mathcal{F}$ with $\mathbb{P}(B) > 0$. Then $A \mapsto \mathbb{P}(A \mid B)$ is a probability measure on $\mathcal{F}$.

Proof. Nonnegativity is immediate and $\mathbb{P}(\Omega \mid B) = \mathbb{P}(B)/\mathbb{P}(B) = 1$. If (A_i) are disjoint, then $(A_i \cap B)$ are disjoint, so

$$\mathbb{P}\left(\bigcup_i A_i \mid B\right) = \frac{\mathbb{P}((\bigcup_i A_i) \cap B)}{\mathbb{P}(B)} = \frac{\sum_i \mathbb{P}(A_i \cap B)}{\mathbb{P}(B)} = \sum_i \mathbb{P}(A_i \mid B).$$

Key Insight: Once B is fixed, conditional probability is ordinary probability on the restricted world. □

Remark 1.3.3. *Conditional Probability Is Not Symmetric*
In general, $\mathbb{P}(A \mid B) \neq \mathbb{P}(B \mid A)$. The distinction is a frequent source of base-rate errors: a test can have high probability of being positive among sick people while the probability of sickness among people with positive tests remains small if the disease is rare.

Proposition 1.3.4. *Multiplication and Chain Rules*
If the displayed conditional probabilities are defined, then

$$\mathbb{P}(A \cap B) = \mathbb{P}(A \mid B)\mathbb{P}(B) = \mathbb{P}(B \mid A)\mathbb{P}(A),$$

and

$$\mathbb{P}(A \cap B \cap C) = \mathbb{P}(A \mid B \cap C)\mathbb{P}(B \mid C)\mathbb{P}(C).$$

Proof. The first formula is Definition 1.3.1 rearranged. For the second,

$$\mathbb{P}(A \cap B \cap C) = \mathbb{P}(A \mid B \cap C)\mathbb{P}(B \cap C) = \mathbb{P}(A \mid B \cap C)\mathbb{P}(B \mid C)\mathbb{P}(C).$$

Key Insight: A joint probability can be built one conditional factor at a time. □

Example 1.3.5. *Cards Without Replacement*
Draw two cards from a standard deck without replacement. Let A be the event that the first card is the ace of clubs and let B be the event that the second card is the queen of diamonds. Then

$$\mathbb{P}(A \cap B) = \mathbb{P}(A)\mathbb{P}(B \mid A) = \frac{1}{52} \cdot \frac{1}{51}.$$

The second factor is conditional because, once the first card is fixed, only 51 cards remain.

Proposition 1.3.6. *Independence Through Conditioning*
If $\mathbb{P}(A) > 0$ and $\mathbb{P}(B) > 0$, then A and B are independent iff

$$\mathbb{P}(A \mid B) = \mathbb{P}(A)$$

iff $\mathbb{P}(B \mid A) = \mathbb{P}(B)$.

Proof. Each statement is equivalent, after multiplying by the relevant positive denominator, to $\mathbb{P}(A \cap B) = \mathbb{P}(A)\mathbb{P}(B)$. □

Theorem 1.3.7. *Law of Total Probability*
Let $A_1, \ldots, A_k$ be a partition of Ω with $\mathbb{P}(A_i) > 0$. Then, for every event B,

$$\mathbb{P}(B) = \sum_{i=1}^{k} \mathbb{P}(B \mid A_i)\mathbb{P}(A_i).$$

Proof. The sets $B \cap A_i$ are disjoint and their union is B, so

$$\mathbb{P}(B) = \sum_i \mathbb{P}(B \cap A_i) = \sum_i \mathbb{P}(B \mid A_i)\mathbb{P}(A_i).$$

Key Insight: Partition by hidden cases, compute within each case, and add. □

Theorem 1.3.8. *Bayes' Theorem*
Let $A_1, \dots, A_k$ be a partition of Ω with $\mathbb{P}(A_i) > 0$. If $\mathbb{P}(B) > 0$, then

$$\mathbb{P}(A_i \mid B) = \frac{\mathbb{P}(B \mid A_i)\mathbb{P}(A_i)}{\sum_{j=1}^{k} \mathbb{P}(B \mid A_j)\mathbb{P}(A_j)}.$$

The probabilities $\mathbb{P}(A_i)$ are prior probabilities; the probabilities $\mathbb{P}(A_i \mid B)$ are posterior probabilities.

Proof. Apply Definition 1.3.1, the multiplication rule, and Theorem 1.3.7:

$$\mathbb{P}(A_i \mid B) = \frac{\mathbb{P}(A_i \cap B)}{\mathbb{P}(B)} = \frac{\mathbb{P}(B \mid A_i)\mathbb{P}(A_i)}{\sum_j \mathbb{P}(B \mid A_j)\mathbb{P}(A_j)}.$$

Key Insight: Bayes' theorem is the multiplication rule plus total probability. □

Corollary 1.3.9. *Posterior Mass Must Balance*
Let $A_1, \dots, A_k$ be a partition and assume $\mathbb{P}(B) > 0$. If $\mathbb{P}(A_1 \mid B) < \mathbb{P}(A_1)$, then $\mathbb{P}(A_i \mid B) > \mathbb{P}(A_i)$ for some $i \in \{2, \dots, k\}$.

Proof. Both the priors and posteriors sum to 1. If every posterior for $i \geq 2$ were at most its prior, then summing would give

$$1 = \sum_i \mathbb{P}(A_i \mid B) < \mathbb{P}(A_1) + \sum_{i=2}^{k} \mathbb{P}(A_i) = 1,$$

a contradiction. □

Example 1.3.10. *Base Rates in a Medical Test*
Let D be disease and let $+$ be a positive test. Suppose the joint probabilities are

	D	D^c
$+$	0.0081	0.0900
$-$	0.0009	0.9010

Then

$$\mathbb{P}(+ \mid D) = \frac{0.0081}{0.0081 + 0.0009} = 0.9, \qquad \mathbb{P}(- \mid D^c) = \frac{0.9010}{0.9010 + 0.0900} \approx 0.91.$$

The test looks accurate in those directions. But after a positive test,

$$\mathbb{P}(D \mid +) = \frac{0.0081}{0.0081 + 0.0900} \approx 0.083.$$

Key Insight: The posterior can be small even when sensitivity and specificity are high, because the prior disease rate is small.

Example 1.3.11. *Monty Hall*
Suppose the player always initially chooses door 1. Let an outcome be (ω_1, ω_2), where ω_1 is the prize door and ω_2 is the door Monty opens. Monty never opens the prize door or door 1, and if the prize is behind door 1 he chooses randomly between doors 2 and 3. Thus

$$\mathbb{P}(1,2) = \mathbb{P}(1,3) = \frac{1}{6}, \qquad \mathbb{P}(2,3) = \mathbb{P}(3,2) = \frac{1}{3}.$$

If Monty opens door 2, then

$$\mathbb{P}(\text{win by staying} \mid \omega_2 = 2) = \frac{1/6}{1/6 + 1/3} = \frac{1}{3},$$

while switching wins with probability 2/3. The same conclusion holds if Monty opens door 3.
Key Insight: The sample space must encode Monty's rule; after that, switching has conditional probability 2/3.

Example 1.3.12. *Three Cards and Conditional Sampling*
There are three cards: green-green, red-red, and green-red. Choose a card uniformly, then choose one side uniformly. Given that the visible side is green, the probability that the hidden side is also green is 2/3.
Indeed, among the equally likely visible sides, there are three green sides: two come from the green-green card and one comes from the green-red card. Two of those three have green on the other side.

Example 1.3.13. *Conditioning Changes the Reference Class*
Suppose each child independently has blue eyes with probability $1/4$, and let a family have five children.

(i) If it is known that at least one child has blue eyes, then

$$\mathbb{P}(\text{at least three blue-eyed children} \mid \text{at least one}) = \frac{\binom{5}{3}(1/4)^3(3/4)^2 + \binom{5}{4}(1/4)^4(3/4) + (1/4)^5}{1-(3/4)^5} = \frac{106}{781}.$$

(ii) If it is known that the youngest child has blue eyes, then we need at least two blue-eyed children among the remaining four:

$$\mathbb{P}(\text{at least three blue-eyed children} \mid \text{youngest is blue-eyed}) = 1 - \left(\frac{3}{4}\right)^4 - 4\left(\frac{1}{4}\right)\left(\frac{3}{4}\right)^3 = \frac{67}{256}.$$

Key Insight: Similar-sounding information can condition on different events.

Example 1.3.14. *Bayes with Several Categories*
Suppose computer owners use Mac, Windows, and Linux with probabilities $0.30, 0.50, 0.20$. Suppose infection rates are $0.65, 0.82, 0.50$, respectively. If a randomly selected owner's system is infected, then

$$\mathbb{P}(\text{Windows} \mid \text{infected}) = \frac{0.82 \cdot 0.50}{0.65 \cdot 0.30 + 0.82 \cdot 0.50 + 0.50 \cdot 0.20} = \frac{82}{141}.$$

Example 1.3.15. *Posterior and Predictive Probabilities*
A box contains five coins with head probabilities

$$p_1 = 0, \qquad p_2 = \frac{1}{4}, \qquad p_3 = \frac{1}{2}, \qquad p_4 = \frac{3}{4}, \qquad p_5 = 1.$$

Choose a coin uniformly and toss it. If the first toss is heads, then

$$\mathbb{P}(C_i \mid H_1) = \frac{p_i}{p_1+p_2+p_3+p_4+p_5} = \left(0, \frac{1}{10}, \frac{1}{5}, \frac{3}{10}, \frac{2}{5}\right)_i.$$

The probability that the second toss is also heads is the posterior average of p_i:

$$\mathbb{P}(H_2 \mid H_1) = \sum_i p_i \, \mathbb{P}(C_i \mid H_1) = \frac{3}{4}.$$

If instead we toss until the first head occurs and B_4 is the event that this first head occurs on toss 4, then

$$\mathbb{P}(C_i \mid B_4) \propto (1-p_i)^3 p_i,$$

so the posterior vector is

$$\left(0, \frac{27}{46}, \frac{16}{46}, \frac{3}{46}, 0\right).$$

Key Insight: Bayes' theorem updates a prior over hidden parameters; prediction averages over the posterior.

Definition 1.3.16. *Conditional Mass and Density Functions*
For discrete random variables X, Y, the *conditional probability mass function* is

$$f_{X\mid Y}(x \mid y) = \mathbb{P}(X = x \mid Y = y) = \frac{f_{X,Y}(x,y)}{f_Y(y)},$$

when $f_Y(y) > 0$.
For continuous random variables, the *conditional density* is defined by the same ratio

$$f_{X\mid Y}(x \mid y) = \frac{f_{X,Y}(x,y)}{f_Y(y)}$$

when $f_Y(y) > 0$, and probabilities are obtained by integration:

$$\mathbb{P}(X \in A \mid Y = y) = \int_A f_{X\mid Y}(x \mid y)\, dx.$$

Source TODO: The screenshots support the elementary conditional-density treatment below; the abstract σ-algebra version of conditional expectation is deferred to the later filtrations and martingales material.

Example 1.3.17. *Uniform Square*
If (X, Y) is uniform on the unit square, then $f_{X,Y}(x,y) = 1$ on $0 \le x, y \le 1$, and $f_Y(y) = 1$ for $0 \le y \le 1$. Hence

$$f_{X\mid Y}(x \mid y) = 1, \qquad 0 \le x \le 1,$$

so $X \mid Y = y \sim \text{Uniform}(0,1)$.

Example 1.3.18. *Conditioning in the Parabolic Region*

For the density in Example 1.1.31,

$$f(x,y) = \frac{21}{4}x^2 y, \qquad x^2 \le y \le 1.$$

The marginal density of X is

$$f_X(x) = \frac{21}{8}x^2(1-x^4), \qquad -1 \le x \le 1.$$

Thus, for $x^2 \le y \le 1$,

$$f_{Y|X}(y \mid x) = \frac{(21/4)x^2 y}{(21/8)x^2(1-x^4)} = \frac{2y}{1-x^4}.$$

At $x = 1/2$, this gives $f_{Y|X}(y \mid 1/2) = 32y/15$, and

$$\mathbb{P}(Y \ge 3/4 \mid X = 1/2) = \int_{3/4}^{1} \frac{32y}{15}\,dy = \frac{7}{15}.$$

Conditional expectation is ordinary expectation after the available information has been fixed. In this first pass we use the elementary conditional-mass and conditional-density viewpoint; the measure-theoretic version appears later through filtrations and martingales.

Definition 1.3.19. *Conditional Expectation Given a Random Variable*
Let X and Y be random variables. In the discrete case,

$$\mathbb{E}(X \mid Y = y) := \sum_x x\, p_{X|Y}(x \mid y),$$

and in the density case,

$$\mathbb{E}(X \mid Y = y) := \int_{\mathbb{R}} x f_{X|Y}(x \mid y)\,dx,$$

assuming the relevant sum or integral exists. More generally, for a measurable function r,

$$\mathbb{E}(r(X,Y) \mid Y = y) = \begin{cases} \sum_x r(x,y) p_{X|Y}(x \mid y), & \text{discrete case,} \\ \int_{\mathbb{R}} r(x,y) f_{X|Y}(x \mid y)\,dx, & \text{density case.} \end{cases}$$

The random variable $\mathbb{E}(X \mid Y)$ is the function of Y whose value at $Y = y$ is $\mathbb{E}(X \mid Y = y)$.

Example 1.3.20. *A Conditional Uniform Model*
Draw $X \sim \text{Uniform}(0,1)$. Given $X = x$, draw $Y \mid X = x \sim \text{Uniform}(x,1)$. Then

$$f_{Y|X}(y \mid x) = \frac{1}{1-x}, \qquad x < y < 1,$$

and hence

$$\mathbb{E}(Y \mid X = x) = \frac{1}{1-x}\int_x^1 y\,dy = \frac{1+x}{2}.$$

Thus, as a random variable, $\mathbb{E}(Y \mid X) = \frac{1+X}{2}$.

Theorem 1.3.21. *Tower Property*
Assume the expectations exist. Then

$$\mathbb{E}\big[\mathbb{E}(Y \mid X)\big] = \mathbb{E}Y.$$

More generally,

$$\mathbb{E}\big[\mathbb{E}(r(X,Y) \mid X)\big] = \mathbb{E}(r(X,Y)).$$

Proof. In the density case, write $f(x,y) = f_X(x) f_{Y|X}(y \mid x)$. Then Fubini's theorem gives

$$\mathbb{E}\big[\mathbb{E}(Y \mid X)\big] = \int \left(\int y f_{Y|X}(y \mid x)\,dy\right) f_X(x)\,dx = \int\int y f(x,y)\,dy\,dx = \mathbb{E}Y.$$

The same computation with sums proves the discrete case, and replacing y by $r(x,y)$ gives the general identity. □

Example 1.3.22. *Using the Tower Property*
In Example 1.3.20,

$$\mathbb{E}Y = \mathbb{E}\left(\frac{1+X}{2}\right) = \frac{1+\mathbb{E}X}{2} = \frac{3}{4}.$$

This avoids integrating over the triangular joint region directly.

Proposition 1.3.23. *Pulling Out What Is Known*
Let r be a function of X and s a function of Y. When the expectations exist,

$$\mathbb{E}\big(r(X)s(Y) \mid X\big) = r(X)\mathbb{E}\big(s(Y) \mid X\big), \qquad \mathbb{E}\big(r(X) \mid X\big) = r(X).$$

Proof. Condition on $X = x$. The factor $r(x)$ is constant inside the conditional expectation, so it pulls outside the sum or integral. The second identity is the special case $s \equiv 1$. □

Definition 1.3.24. *Conditional Variance*
For random variables X, Y, define

$$\operatorname{Var}(Y \mid X = x) := \mathbb{E}\left((Y - \mu(x))^2 \mid X = x\right), \qquad \mu(x) := \mathbb{E}(Y \mid X = x).$$

The random variable $\operatorname{Var}(Y \mid X)$ is obtained by evaluating this function at X.

Theorem 1.3.25. *Law of Total Variance*
For square-integrable random variables X, Y,

$$\operatorname{Var}(Y) = \mathbb{E}\big[\operatorname{Var}(Y \mid X)\big] + \operatorname{Var}\big(\mathbb{E}(Y \mid X)\big).$$

Proof. Let $m = \mathbb{E}Y$ and $b(X) = \mathbb{E}(Y \mid X)$. Then $\mathbb{E}b(X) = m$ by Theorem 1.3.21. Write

$$Y - m = (Y - b(X)) + (b(X) - m).$$

After squaring and taking expectations, the cross term vanishes because

$$\mathbb{E}\big[(Y - b(X))(b(X) - m)\big] = \mathbb{E}\left[(b(X) - m)\mathbb{E}(Y - b(X) \mid X)\right] = 0$$

by Proposition 1.3.23. The remaining two terms are exactly the displayed identity. □

Example 1.3.26. *Random Binomial Proportion*
Choose a county at random, and let Q be the proportion of people in that county with a certain disease. Given $Q = q$, suppose $X \mid Q = q \sim \text{Binomial}(n, q)$. If $Q \sim \text{Uniform}(0, 1)$, then

$$\mathbb{E}X = \mathbb{E}[\mathbb{E}(X \mid Q)] = \mathbb{E}(nQ) = \frac{n}{2}.$$

For the variance,

$$\mathbb{E}[\operatorname{Var}(X \mid Q)] = \mathbb{E}[nQ(1 - Q)] = n\int_0^1 q(1 - q)\,dq = \frac{n}{6},$$

and

$$\operatorname{Var}(\mathbb{E}(X \mid Q)) = \operatorname{Var}(nQ) = \frac{n^2}{12}.$$

Thus $\operatorname{Var}(X) = \frac{n}{6} + \frac{n^2}{12}$.

Corollary 1.3.27. *Constant Conditional Mean Implies Uncorrelated*
If $\mathbb{E}(X \mid Y) = c$ is constant and X, Y are square-integrable, then $\operatorname{Cov}(X, Y) = 0$.

Proof. By the tower property, $c = \mathbb{E}X$. Then

$$\mathbb{E}(XY) = \mathbb{E}\big[\mathbb{E}(XY \mid Y)\big] = \mathbb{E}\big[Y\mathbb{E}(X \mid Y)\big] = c\,\mathbb{E}Y = \mathbb{E}X\,\mathbb{E}Y.$$

Hence $\operatorname{Cov}(X, Y) = \mathbb{E}(XY) - \mathbb{E}X\,\mathbb{E}Y = 0$. □

1.4 Modes of Probabilistic Convergence

Pointwise convergence is too strict for most random sequences: two continuous random variables can have the same law and still be unequal with probability one. The useful replacement is to ask whether probabilities, distributions, or moments of the errors become small.

Definition 1.4.1. *Convergence in Probability and in Distribution*
Let $X_1, X_2, \ldots$ and X be real-valued random variables, and let F_n and F be the distribution functions of X_n and X.

(i) X_n *converges to X in probability*, written $X_n \xrightarrow{\mathbb{P}} X$, if for every $\varepsilon > 0$,

$$\mathbb{P}(|X_n - X| > \varepsilon) \to 0.$$

(ii) X_n *converges to X in distribution*, written $X_n \rightsquigarrow X$, if

$$F_n(t) \to F(t)$$

at every continuity point t of F.

If $X = c$ almost surely, we also write $X_n \xrightarrow{\mathbb{P}} c$ and $X_n \rightsquigarrow c$.

Definition 1.4.2. *L^p, Quadratic Mean, and Almost Sure Convergence*
For $p \geq 1$, X_n *converges to X in L^p* if

$$\mathbb{E}\,|X_n - X|^p \to 0.$$

The case $p = 2$ is also called *convergence in quadratic mean*, and is written $X_n \xrightarrow{qm} X$.
The sequence X_n *converges almost surely* to X, written $X_n \xrightarrow{\text{a.s.}} X$, if

$$\mathbb{P}\big(\{\omega : X_n(\omega) \to X(\omega)\}\big) = 1.$$

Example 1.4.3. *A Shrinking Normal Law*
Let $X_n \sim N(0, 1/n)$ and let F be the distribution function of a point mass at 0. Since $\sqrt{n}X_n \sim N(0,1)$,

$$F_n(t) = \mathbb{P}(X_n \leq t) = \Phi(\sqrt{n}\,t).$$

Thus $F_n(t) \to 0$ for $t < 0$ and $F_n(t) \to 1$ for $t > 0$. The only discontinuity point of F is 0, so $X_n \rightsquigarrow 0$.

Theorem 1.4.4. *Basic Implications Between Convergence Modes*
For random variables X_n and X:

(i) $X_n \xrightarrow{qm} X$ *implies* $X_n \xrightarrow{\mathbb{P}} X$.
(ii) $X_n \xrightarrow{\mathbb{P}} X$ *implies* $X_n \rightsquigarrow X$.
(iii) If $X_n \rightsquigarrow c$ *for a constant c, then* $X_n \xrightarrow{\mathbb{P}} c$.
(iv) $X_n \xrightarrow{\text{a.s.}} X$ *implies* $X_n \xrightarrow{\mathbb{P}} X$.
(v) $X_n \xrightarrow{L^1} X$ *implies* $X_n \xrightarrow{\mathbb{P}} X$, *and* $X_n \xrightarrow{qm} X$ *implies* $X_n \xrightarrow{L^1} X$.

In general, the reverse implications fail except for the point-mass case in (iii).

Proof. For (i), Markov's inequality from Theorem 1.2.20 gives

$$\mathbb{P}(|X_n - X| > \varepsilon) \leq \frac{\mathbb{E}(X_n - X)^2}{\varepsilon^2} \to 0.$$

For (ii), if x is a continuity point of F, then for every $\varepsilon > 0$,

$$F(x - \varepsilon) - \mathbb{P}(|X_n - X| > \varepsilon) \leq F_n(x) \leq F(x + \varepsilon) + \mathbb{P}(|X_n - X| > \varepsilon),$$

and then let $n \to \infty$ and $\varepsilon \downarrow 0$. For (iii),

$$\mathbb{P}(|X_n - c| > \varepsilon) \leq F_n(c - \varepsilon) + 1 - F_n(c + \varepsilon) \to 0.$$

Almost sure convergence implies probability convergence by dominated convergence applied to $\mathbf{1}_{\{|X_n - X| > \varepsilon\}}$. The L^1 claim is Markov's inequality applied to $|X_n - X|$, and the last implication follows from Cauchy–Schwarz. □

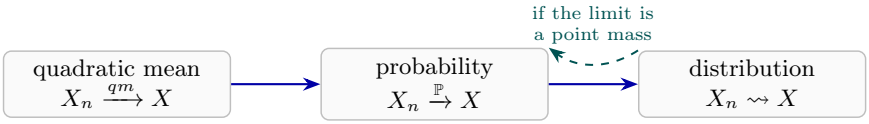


Figure 208: The main one-way implications between quadratic-mean, probability, and distribution convergence.

Example 1.4.5. *Why the Reverse Arrows Fail*

(i) Let $U \sim \text{Uniform}(0,1)$ and $X_n = \sqrt{n}\,\mathbf{1}_{(0,1/n)}(U)$. Then $\mathbb{P}(|X_n| > \varepsilon) = 1/n \to 0$, so $X_n \xrightarrow{\mathbb{P}} 0$, but $\mathbb{E}X_n^2 = 1$ for every n.
(ii) Let $X \sim N(0,1)$ and set $X_n = -X$. Then X_n has the same distribution as X for every n, hence $X_n \rightsquigarrow X$, but $\mathbb{P}(|X_n - X| > \varepsilon) = \mathbb{P}(2\,|X| > \varepsilon)$ does not tend to 0.

Remark 1.4.6. *A Warning About Expectations*
Convergence in probability does not by itself imply convergence of expectations. If

$$\mathbb{P}(X_n = n^2) = \frac{1}{n}, \qquad \mathbb{P}(X_n = 0) = 1 - \frac{1}{n},$$

then $X_n \xrightarrow{\mathbb{P}} 0$, but $\mathbb{E}X_n = n \to +\infty$.

Definition 1.4.7. *Asymptotic Uniform Integrability*
A sequence X_n is *asymptotically uniformly integrable* if

$$\lim_{M \to \infty} \limsup_{n \to \infty} \mathbb{E}\big(|X_n|\,\mathbf{1}_{\{|X_n| > M\}}\big) = 0.$$

Proposition 1.4.8. *Recovering Expectations from Probability Convergence*
If $X_n \xrightarrow{\mathbb{P}} b$ for a constant b and the sequence X_n is asymptotically uniformly integrable, then $\mathbb{E}X_n \to b$.

Proof idea. Truncate X_n at level M. Probability convergence gives convergence of the bounded truncated expectations, while asymptotic uniform integrability makes the discarded tails uniformly small for large M. □

Theorem 1.4.9. *Continuous Mapping and Slutsky Rules*
Let g be continuous.

(i) If $X_n \xrightarrow{\mathbb{P}} X$, then $g(X_n) \xrightarrow{\mathbb{P}} g(X)$.
(ii) If $X_n \rightsquigarrow X$, then $g(X_n) \rightsquigarrow g(X)$.
(iii) If $X_n \rightsquigarrow X$ and $Y_n \xrightarrow{\mathbb{P}} c$, then $X_n + Y_n \rightsquigarrow X + c$ and $X_n Y_n \rightsquigarrow cX$.
(iv) If $X_n \xrightarrow{\mathbb{P}} X$ and $Y_n \xrightarrow{\mathbb{P}} Y$, then $X_n + Y_n \xrightarrow{\mathbb{P}} X + Y$ and $X_n Y_n \xrightarrow{\mathbb{P}} XY$.

Proof idea. The probability statements follow by bounding the bad event for sums and products and using continuity of g on high-probability bounded sets. The distribution statements are the corresponding continuous-mapping theorem and its constant-limit specialization, usually called Slutsky's theorem. □

1.5 Laws of Large Numbers

The law of large numbers formalizes the idea that a large sample average is close to the mean of the distribution.

Perspective 1.5.1. *Sampling Distribution of a Statistic*
If $X_1, \dots, X_n$ are IID with mean μ and variance σ^2, then

$$\overline{X}_n = \frac{1}{n}\sum_{i=1}^{n} X_i$$

is a random variable. Its law is the *sampling distribution* of the statistic. Theorem 1.2.15 gives

$$\mathbb{E}\overline{X}_n = \mu, \qquad \operatorname{Var}(\overline{X}_n) = \frac{\sigma^2}{n}.$$

Thus the sampling distribution becomes more concentrated as n grows.

Theorem 1.5.2. *Weak Law of Large Numbers*
Let $X_1, X_2, \dots$ be IID with $\mathbb{E}X_1 = \mu$ and $\operatorname{Var}(X_1) = \sigma^2 < +\infty$. Then

$$\overline{X}_n \xrightarrow{\mathbb{P}} \mu.$$

Proof. By Theorem 1.2.15, $\operatorname{Var}(\overline{X}_n) = \sigma^2/n$. Chebyshev's inequality gives

$$\mathbb{P}(|\overline{X}_n - \mu| > \varepsilon) \le \frac{\sigma^2}{n\varepsilon^2} \to 0.$$

□

Example 1.5.3. *Coin Proportions and a Chebyshev Sample Size*
Flip a coin with head probability p and let X_i be the head indicator on toss i. Then $\overline{X}_n$ is the fraction of heads and $\overline{X}_n \xrightarrow{\mathbb{P}} p$. For a fair coin,

$$\mathbb{P}(0.4 \le \overline{X}_n \le 0.6) \ge 1 - \frac{1}{4n(0.1)^2} = 1 - \frac{25}{n}.$$

The lower bound is at least 0.7 once $n \ge 84$.

Theorem 1.5.4. *Strong Law of Large Numbers*
Let $X_1, X_2, \dots$ be IID and assume $\mathbb{E}|X_1| < +\infty$. If $\mu = \mathbb{E}X_1$, then

$$\overline{X}_n \xrightarrow{\text{a.s.}} \mu.$$

Corollary 1.5.5. *Strong Implies Weak*
Under the hypotheses of Theorem 1.5.4,

$$\overline{X}_n \xrightarrow{\mathbb{P}} \mu.$$

Proof. Almost-sure convergence implies convergence in probability by Theorem 1.4.4. □

Proposition 1.5.6. *Consistency of the Sample Variance*
Let $X_1, X_2, \dots$ be IID with mean μ and variance $\sigma^2 < +\infty$. Then

$$S_n^2 \xrightarrow{\mathbb{P}} \sigma^2.$$

Proof. The identity

$$S_n^2 = \frac{n}{n-1}\left(\frac{1}{n}\sum_{i=1}^{n} X_i^2 - \overline{X}_n^2\right)$$

reduces the claim to the strong law for X_i^2 and X_i, followed by the continuous mapping rules in Theorem 1.4.9. □

Remark 1.5.7. *Why Integrability Matters*
The Cauchy example in Example 1.2.6 has no first absolute moment. Its sample averages do not settle down around a finite mean, so it lies outside the hypotheses of the laws of large numbers above.

1.6 Central Limit Theory

The law of large numbers says that $\overline{X}_n$ is close to μ for large n. The central limit theorem describes the scale and shape of the remaining fluctuation.

Theorem 1.6.1. *Central Limit Theorem*
Let $X_1, X_2, \dots$ be IID with mean μ and variance $0 < \sigma^2 < +\infty$. Define

$$Z_n := \frac{\sqrt{n}(\overline{X}_n - \mu)}{\sigma}.$$

Then $Z_n \rightsquigarrow Z$, where $Z \sim N(0,1)$. Equivalently,

$$\mathbb{P}(Z_n \le z) \to \Phi(z) = \int_{-\infty}^{z} \frac{1}{\sqrt{2\pi}} e^{-x^2/2}\, dx.$$

Method 1.6.2. *Equivalent CLT Approximation Forms*
The statement $Z_n \rightsquigarrow N(0,1)$ is often used as

$$\overline{X}_n \approx N\left(\mu, \frac{\sigma^2}{n}\right), \qquad \sqrt{n}(\overline{X}_n - \mu) \approx N(0, \sigma^2).$$

The approximation is for probability statements about $\overline{X}_n$.

Lemma 1.6.3. *MGF Convergence Criterion*
Let Z_n have MGF ψ_n and let Z have MGF ψ. If $\psi_n(t) \to \psi(t)$ for every t in some open interval around 0, then $Z_n \rightsquigarrow Z$.

Proof idea for Theorem 1.6.1 under an MGF hypothesis. Assume $Y_i = (X_i - \mu)/\sigma$ has an MGF ψ finite near 0. Then $\psi(t) = 1 + t^2/2 + o(t^2)$. Since $Z_n = n^{-1/2}\sum_i Y_i$, Proposition 1.2.27 gives

$$\psi_{Z_n}(t) = \left[\psi\left(\frac{t}{\sqrt{n}}\right)\right]^n = \left(1 + \frac{t^2}{2n} + o\left(\frac{1}{n}\right)\right)^n \to e^{t^2/2},$$

the MGF of $N(0,1)$. □

Example 1.6.4. *A Poisson Error-Count Approximation*
Suppose the number of errors in a computer program has a Poisson(5) distribution, and $X_1, \dots, X_{125}$ are independent program error counts. Since $\mu = \sigma^2 = 5$,

$$\mathbb{P}(\overline{X}_{125} < 5.5) \approx \mathbb{P}\left(Z < \frac{\sqrt{125}(5.5-5)}{\sqrt{5}}\right) = \mathbb{P}(Z < 2.5) \approx 0.9938.$$

Theorem 1.6.5. *Studentized Central Limit Theorem*
Under the hypotheses of Theorem 1.6.1, let S_n^2 be the sample variance from Definition 1.2.9. Then

$$\frac{\sqrt{n}(\overline{X}_n - \mu)}{S_n} \rightsquigarrow N(0,1).$$

Proof. By Proposition 1.5.6, $S_n \xrightarrow{\mathbb{P}} \sigma$. Combine this with Theorem 1.6.1 and Slutsky's rule from Theorem 1.4.9. □

Theorem 1.6.6. *Berry–Esseen Rate*
Under the hypotheses of Theorem 1.6.1, assume also that $\mathbb{E}\,|X_1 - \mu|^3 < +\infty$. Then

$$\sup_{z \in \mathbb{R}} |\mathbb{P}(Z_n \le z) - \Phi(z)| \le \frac{33}{4}\frac{\mathbb{E}\,|X_1 - \mu|^3}{\sqrt{n}\sigma^3}.$$

Theorem 1.6.7. *Multivariate Central Limit Theorem*
Let $X_1, \ldots, X_n$ be IID random vectors in $\mathbb{R}^k$ with mean vector μ and covariance matrix Σ. If $\overline{X} = n^{-1}\sum_i X_i$, then

$$\sqrt{n}(\overline{X} - \mu) \rightsquigarrow N(0, \Sigma).$$

Theorem 1.6.8. *Delta Method*
Suppose

$$\frac{\sqrt{n}(Y_n - \mu)}{\sigma} \rightsquigarrow N(0,1)$$

and $g : \mathbb{R} \to \mathbb{R}$ is differentiable at μ with $g'(\mu) \neq 0$. Then

$$\frac{\sqrt{n}(g(Y_n) - g(\mu))}{|g'(\mu)|\,\sigma} \rightsquigarrow N(0,1).$$

Equivalently,

$$g(Y_n) \approx N\left(g(\mu), \frac{(g'(\mu))^2\sigma^2}{n}\right).$$

Example 1.6.9. *Exponentiating a Sample Mean*
Let $X_1, \ldots, X_n$ be IID with finite mean μ and finite variance σ^2, and set $W_n = e^{\overline{X}_n}$. Since $g(s) = e^s$, the delta method gives

$$W_n \approx N\left(e^{\mu}, \frac{e^{2\mu}\sigma^2}{n}\right).$$

Theorem 1.6.10. *Multivariate Delta Method*
Let Y_n be random vectors in $\mathbb{R}^k$ such that $\sqrt{n}(Y_n - \mu) \rightsquigarrow N(0, \Sigma)$. If $g : \mathbb{R}^k \to \mathbb{R}$ is differentiable at μ, then

$$\sqrt{n}(g(Y_n) - g(\mu)) \rightsquigarrow N\big(0, \nabla g(\mu)^t \Sigma \nabla g(\mu)\big).$$

Example 1.6.11. *Product and Ratio of Sample Means*
Let $(X_{1i}, X_{2i})^t$ be IID with mean $(\mu_1, \mu_2)^t$ and covariance matrix $\Sigma = (\sigma_{ij})$. For $\overline{X}_j = n^{-1}\sum_i X_{ji}$,

$$\sqrt{n}(\overline{X}_1\overline{X}_2 - \mu_1\mu_2) \rightsquigarrow N\left(0, \mu_2^2\sigma_{11} + 2\mu_1\mu_2\sigma_{12} + \mu_1^2\sigma_{22}\right).$$

If $\mu_2 \neq 0$ and $a = (1/\mu_2, -\mu_1/\mu_2^2)^t$, then

$$\sqrt{n}\left(\frac{\overline{X}_1}{\overline{X}_2} - \frac{\mu_1}{\mu_2}\right) \rightsquigarrow N(0, a^t\Sigma a).$$

1.7 Probability on Metric Spaces

2 Stochastic Processes

2.1 Process-Level Probability

Definition 2.1.1. *Stochastic Process, State Space, and Index Set*
A *stochastic process* is a family of random variables

$$\{X_t : t \in T\}.$$

The set T is the *index set*; in applications it is often time. The variables X_t take values in a *state space* $\mathcal{X}$. The index set may be discrete, such as $T = \{0, 1, 2, \ldots\}$, or continuous, such as $T = [0, \infty)$.

Example 2.1.2. *IID Sequences as Processes*
An IID sequence $X_1, X_2, \ldots$ is a stochastic process with discrete index set $T = \{1, 2, \ldots\}$. This is the process-level way to regard the IID data used throughout the earlier probability and statistics sections.

Example 2.1.3. *Discrete States and Discrete Time*
If $\mathcal{X} = \{\text{sunny}, \text{cloudy}\}$ and X_n denotes the weather on day n, then

$$X_1, X_2, \ldots$$

is a stochastic process with a discrete state space and a discrete index set.

Example 2.1.4. *The Empirical CDF as a Process*
Let $X_1, \ldots, X_n$ be IID with CDF F on $[0,1]$, and let

$$\widehat{F}_n(t) = \frac{1}{n}\sum_{i=1}^{n} \mathbf{1}\{X_i \leq t\}.$$

For each fixed t, $\widehat{F}_n(t)$ is a random variable. The whole family

$$\{\widehat{F}_n(t) : t \in [0,1]\}$$

is a stochastic process indexed by t. This example has a continuous index set and values in $[0,1]$.

Remark 2.1.5. *Sample Paths*
One realization of a process over time is a sample path. A stock-price record, a sequence of daily weather states, and the count process in Definition 2.7.1 are different kinds of sample paths; the state space and index set can each be discrete or continuous independently.

Proposition 2.1.6. *Finite-Dimensional Chain Rule*
For random variables $X_1, \dots, X_n$ with joint mass function or density,

$$f(x_1, \dots, x_n) = f(x_1) \prod_{i=2}^{n} f(x_i \mid x_1, \dots, x_{i-1}).$$

Proof. Repeatedly apply the product rule $f(x_1, \dots, x_i) = f(x_i \mid x_1, \dots, x_{i-1}) f(x_1, \dots, x_{i-1})$. □

Markov chains impose a one-step simplification on Proposition 2.1.6: the next conditional factor depends only on the current state. This is formalized in Definition 2.4.5.

2.2 Filtrations and Adapted Processes

2.3 Martingales in Discrete and Continuous Time

2.4 Markov Chains and Markov Processes

Definition 2.4.1. *Probability Vectors and Markov Matrices*
A vector $\mathbf{x} \in \mathbb{R}^n$ is a *probability vector* if its entries are nonnegative and sum to 1. Let $\mathbf{e} := (1, \dots, 1)^t$. A matrix $A = [a_{ij}] \in \mathrm{Mat}(n, n, \mathbb{R})$ is a *Markov matrix* if $a_{ij} \geq 0$ for all i, j and $A\mathbf{e} = \mathbf{e}$; equivalently, every row sum is 1.

Lemma 2.4.2. *Basic Properties of Markov Matrices*
If $A \in \mathrm{Mat}(n, n, \mathbb{R})$ is Markov, then $1 \in \mathrm{spec}(A)$, and each A^p is Markov with entries in $[0,1]$.

Proof. Since $A\mathbf{e} = \mathbf{e}$, 1 is an eigenvalue. Inductively, A^p has nonnegative entries and $A^p\mathbf{e} = \mathbf{e}$, so A^p is Markov. Each nonnegative row sums to 1, so all entries lie in $[0,1]$.
Structural Commentary: Markov structure is closed under powers.
Key Insight: Nonnegative row sums of 1 bound every entry. □

Theorem 2.4.3. *Unit-Circle Spectrum of Markov Matrices*
If $A \in \mathrm{Mat}(n, n, \mathbb{R})$ is Markov, then $\rho(A) = 1$, and $\dim E_\lambda(A) = m_A(\lambda)$ for every $\lambda \in \mathrm{spec}(A)$ with $|\lambda| = 1$.

Proof. Lemma 2.4.2 makes A power bounded. Theorem 4.1.95 gives $\rho(A) \leq 1$ and $\dim E_\lambda(A) = m_A(\lambda)$ on $|\lambda| = 1$. Since $A\mathbf{e} = \mathbf{e}$, $1 \in \mathrm{spec}(A)$, so $\rho(A) \geq 1$. Thus $\rho(A) = 1$.
Structural Commentary: Bounded Markov powers rule out unit-circle Jordan growth.
Key Insight: $A\mathbf{e} = \mathbf{e}$ forces the eigenvalue 1. □

Theorem 2.4.4. *Positive Markov Matrices Converge to a Stationary Distribution*
Let $n \geq 2$ and let $A \in \mathrm{Mat}(n, n, \mathbb{R})$ be a Markov matrix with positive entries. Then there is a unique probability vector $\mathbf{x}$ with $A^t\mathbf{x} = \mathbf{x}$; its entries are positive and $\lim_{p \to \infty} A^p = \mathbf{e}\mathbf{x}^t$.

Proof. Let $A\mathbf{y} = \lambda\mathbf{y}$ with $|\lambda| = 1$, and choose k with $|y_k| = \|\mathbf{y}\|_\infty$.
Positivity and the kth row sum give $\|\mathbf{y}\|_\infty = |\lambda y_k| = \left| \sum_{j=1}^{n} a_{kj} y_j \right| \leq \sum_{j=1}^{n} a_{kj} |y_j| \leq \|\mathbf{y}\|_\infty$.
Because every a_{kj} is positive, equality in the last inequality forces $|y_j| = \|\mathbf{y}\|_\infty$ for all j. Equality in the triangle inequality then forces all nonzero y_j to have the same argument. Hence $\mathbf{y} = c\mathbf{e}$, and $A\mathbf{e} = \mathbf{e}$ gives $\lambda = 1$. Thus $E_1(A) = \mathrm{span}(\mathbf{e})$ and no other eigenvalue lies on the unit circle.
By Theorem 2.4.3, the Jordan form is $[1] \oplus J$ with $\rho(J) < 1$. Choose $S = [\mathbf{e}\ Y]$ such that $A = S([1] \oplus J)S^{-1}$, and let $\mathbf{x}^t$ be the first row of S^{-1}. Then $\mathbf{x}^t A = \mathbf{x}^t$ and $\mathbf{x}^t\mathbf{e} = 1$. Since $J^p \to 0$, $A^p = S([1] \oplus J^p)S^{-1} \to \mathbf{e}\mathbf{x}^t$.
The nonnegative powers A^p imply $\mathbf{e}\mathbf{x}^t \geq 0$, so $\mathbf{x} \geq 0$. Since $A^t\mathbf{x} = \mathbf{x}$, $\mathbf{x} \neq 0$, and A^t has positive entries, in fact $\mathbf{x} > 0$. Finally, $\dim E_1(A^t) = \dim E_1(A) = 1$ because $A - I_n$ and $(A - I_n)^t$ have the same rank, so $\mathbf{x}^t\mathbf{e} = 1$ gives uniqueness among probability vectors.
Structural Commentary: Positivity leaves one simple unit-circle block; all other blocks decay.
Key Insight: The limit is the rank-one projection $\mathbf{e}\mathbf{x}^t$. □

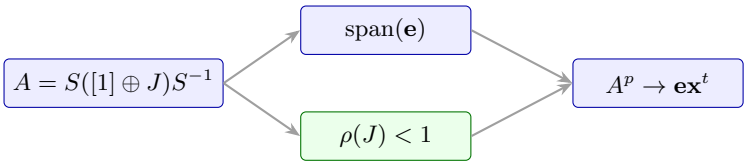


Figure 209: The line span(**e**) remains; all other Jordan directions decay, so $A^p \to \mathbf{e}\mathbf{x}^t$.

The preceding results are matrix facts. In probability language, a row-stochastic matrix becomes a transition matrix for a discrete-time chain.

Definition 2.4.5. *Markov Chain and Transition Matrix*
Let $\mathcal{X}$ be a countable state space and let $T = \{0, 1, 2, \ldots\}$. A process $\{X_n : n \in T\}$ is a *Markov chain* if

$$\mathbb{P}(X_n = x \mid X_0, \ldots, X_{n-1}) = \mathbb{P}(X_n = x \mid X_{n-1})$$

for all admissible n and $x \in \mathcal{X}$. It is *homogeneous* if

$$p_{ij} := \mathbb{P}(X_{n+1} = j \mid X_n = i)$$

does not depend on n. The matrix $P = [p_{ij}]$ is the *transition matrix*; it has nonnegative entries and each row sums to 1.

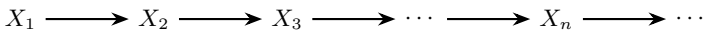


Figure 210: A Markov chain remembers the current state, not the full past.

Example 2.4.6. *Random Walk with Absorbing Barriers*
Let $\mathcal{X} = \{1, \ldots, N\}$. From an interior state, move one step right with probability p and one step left with probability $q = 1 - p$. If the chain hits 1 or N, it stays there. The transition matrix has $p_{11} = p_{NN} = 1$, interior entries $p_{i,i+1} = p$ and $p_{i,i-1} = q$, and all other entries 0.

Definition 2.4.7. *n-Step Transition Probabilities*
For a homogeneous Markov chain, define

$$p_{ij}(n) := \mathbb{P}(X_{m+n} = j \mid X_m = i),$$

and let $P_n = [p_{ij}(n)]$.

Theorem 2.4.8. *Chapman-Kolmogorov Equations*
For all $m, n \geq 0$,

$$p_{ij}(m + n) = \sum_k p_{ik}(m) p_{kj}(n).$$

Consequently,

$$P_{m+n} = P_m P_n, \qquad P_n = P^n.$$

Proof. Condition on the state at the intermediate time m and use the Markov property:

$$\mathbb{P}(X_{m+n} = j \mid X_0 = i) = \sum_k \mathbb{P}(X_{m+n} = j \mid X_m = k)\mathbb{P}(X_m = k \mid X_0 = i).$$

This is exactly the displayed sum. Matrix multiplication gives the matrix form, and $P_1 = P$ gives $P_n = P^n$ by induction. □

Lemma 2.4.9. *Marginal Evolution*
Let μ_n be the row vector with entries $\mu_n(i) = \mathbb{P}(X_n = i)$. If μ_0 is the initial distribution, then

$$\mu_n = \mu_0 P^n.$$

Proof. For each j,

$$\mu_n(j) = \sum_i \mathbb{P}(X_n = j \mid X_0 = i)\mathbb{P}(X_0 = i) = \sum_i \mu_0(i) p_{ij}(n).$$

□

Definition 2.4.10. *Communication Classes*
State i *reaches* state j, written $i \to j$, if $p_{ij}(n) > 0$ for some $n \geq 0$. If $i \to j$ and $j \to i$, write $i \leftrightarrow j$ and say that i and j *communicate*. A chain is *irreducible* if all states communicate. A set of states is *closed* if, once the chain enters the set, it cannot leave. A one-state closed class is an *absorbing state*.

Proposition 2.4.11. *Communication Classes Partition the State Space*
The relation $\leftrightarrow$ is reflexive, symmetric, and transitive, so the state space is a disjoint union of communication classes.

Proof. Reflexivity follows from $p_{ii}(0) = 1$; symmetry is built into the definition. If $i \to j$ in m steps and $j \to k$ in n steps, then Theorem 2.4.8 gives $p_{ik}(m+n) \geq p_{ij}(m)p_{jk}(n) > 0$. Apply the same argument in the reverse direction. □

Definition 2.4.12. *Recurrent and Transient States*
A state i is *recurrent* or *persistent* if

$$\mathbb{P}(X_n = i \text{ for some } n \geq 1 \mid X_0 = i) = 1.$$

Otherwise, i is *transient*.

Theorem 2.4.13. *Recurrence Criterion*
A state i is recurrent if and only if

$$\sum_{n=0}^{\infty} p_{ii}(n) = \infty.$$

It is transient if and only if this sum is finite.

Proof. Let $I_n = \mathbf{1}\{X_n = i\}$ and $Y = \sum_{n=0}^{\infty} I_n$, the total number of visits to i. Then

$$\mathbb{E}(Y \mid X_0 = i) = \sum_{n=0}^{\infty} p_{ii}(n).$$

If i is recurrent, the chain returns to i with probability 1 and then restarts from i, so this expected number of visits is infinite. If i is transient, the probability of ever returning is some $a_i < 1$, and the number of visits has a geometric tail, hence finite mean. □

Definition 2.4.14. *Period and Ergodicity*
For a recurrent state i, let

$$T_i = \min\{n > 0 : X_n = i\}$$

be the first return time. The state is *positive recurrent* if $\mathbb{E}(T_i \mid X_0 = i) < \infty$ and *null recurrent* otherwise. The *period* of state i is

$$d(i) = \gcd\{n \geq 1 : p_{ii}(n) > 0\}.$$

The state is *aperiodic* if $d(i) = 1$. A state is *ergodic* if it is recurrent, positive recurrent, and aperiodic. A chain is ergodic if all of its states are ergodic.

Definition 2.4.15. *Stationary and Limiting Distributions*
A probability row vector π is a *stationary* or *invariant* distribution if

$$\pi = \pi P.$$

A finite chain has *limiting distribution* π if every row of P^n converges to π; equivalently, $p_{ij}(n) \to \pi_j$ independently of the starting state i.

Theorem 2.4.16. *Finite Ergodic Chain Convergence*
An irreducible, ergodic finite Markov chain has a unique stationary distribution π, and P^n converges to the matrix whose rows are all π. If g is bounded, then with probability 1,

$$\frac{1}{N}\sum_{n=1}^{N} g(X_n) \to \sum_{j} g(j)\pi_j.$$

Definition 2.4.17. *Detailed Balance*
A probability vector π satisfies *detailed balance* for P if

$$\pi_i p_{ij} = \pi_j p_{ji}$$

for all states i, j.

Theorem 2.4.18. *Detailed Balance Implies Stationarity*
If π satisfies detailed balance for P, then π is stationary.

Proof. The jth entry of πP is

$$\sum_{i} \pi_i p_{ij} = \sum_{i} \pi_j p_{ji} = \pi_j \sum_{i} p_{ji}.$$

Since $\sum_i p_{ji} = 1$ is the sum of row j of the row-stochastic matrix P, this equals π_j. □

Remark 2.4.19. *Stationary Need Not Mean Convergent*
The deterministic three-cycle

$$P = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{bmatrix}$$

has stationary distribution $(1/3, 1/3, 1/3)$, but P^n keeps cycling and does not converge. Periodicity is the obstruction.

Method 2.4.20. *Estimating a Finite Transition Matrix*
For a finite homogeneous chain observed as $X_0, \dots, X_n$, let n_{ij} be the number of observed transitions from state i to state j, and let $n_i = \sum_j n_{ij}$. Maximizing the multinomial transition likelihood row by row gives

$$\widehat{p}_{ij} = \frac{n_{ij}}{n_i}$$

when $n_i > 0$. If state i is never observed as a departure state, the corresponding row is not identifiable from the sample without an added convention or prior.

Proposition 2.4.21. *Two-State Chain Limit*
Let

$$P = \begin{bmatrix} 1-a & a \\ b & 1-b \end{bmatrix}, \qquad 0 < a < 1, \quad 0 < b < 1.$$

Then

$$\lim_{n\to\infty} P^n = \begin{bmatrix} \frac{b}{a+b} & \frac{a}{a+b} \\ \frac{b}{a+b} & \frac{a}{a+b} \end{bmatrix}.$$

Proof. The row vector

$$\pi = \left(\frac{b}{a+b}, \frac{a}{a+b} \right)$$

satisfies $\pi P = \pi$. The second eigenvalue is $1 - a - b$, whose absolute value is < 1 because $0 < a, b < 1$. Thus the component transverse to the stationary direction decays, leaving both rows equal to π. □

2.5 Stationary Processes and Ergodicity

2.6 Gaussian Processes

2.7 Point Processes and Renewal Structure

Counting processes model random event arrivals over time. The Poisson process is the basic model in which disjoint time intervals contribute independent counts and short intervals usually contain either zero or one event.

Definition 2.7.1. *Counting Process*
A *counting process* is a stochastic process $\{X(t) : t \geq 0\}$ with state space $\{0, 1, 2, \dots\}$ such that $X(t)$ records the number of events observed up to and including time t.

Definition 2.7.2. *Poisson Process with Intensity Function*
A counting process $\{X(t) : t \geq 0\}$ is a *Poisson process* with *intensity function* $\lambda(t)$ if:

(i) $X(0) = 0$.

(ii) For $0 = t_0 < t_1 < \cdots < t_n$, the increments

$$X(t_1) - X(t_0), \; X(t_2) - X(t_1), \dots, \; X(t_n) - X(t_{n-1})$$

are independent.

(iii) As $h \downarrow 0$,

$$\mathbb{P}(X(t+h) - X(t) = 1) = \lambda(t)h + o(h), \qquad \mathbb{P}(X(t+h) - X(t) \geq 2) = o(h).$$

Theorem 2.7.3. *Increment Distribution*
Let $X(t)$ be a Poisson process with intensity function $\lambda(t)$ and define

$$m(t) = \int_0^t \lambda(s)\, ds.$$

Then

$$X(s+t) - X(s) \sim \text{Poisson}(m(s+t) - m(s)).$$

In particular, $X(t) \sim \text{Poisson}(m(t))$, so

$$\mathbb{E}(X(t)) = m(t), \qquad \mathbb{V}(X(t)) = m(t).$$

Definition 2.7.4. *Homogeneous Poisson Process*
If $\lambda(t) \equiv \lambda > 0$, then $X(t)$ is a *homogeneous Poisson process* with *rate* λ. In this case,

$$X(t) \sim \text{Poisson}(\lambda t).$$

Definition 2.7.5. *Waiting and Interarrival Times*
For a homogeneous Poisson process, let W_n be the time at which the nth event occurs and set $W_0 = 0$. The variables $W_1, W_2, \ldots$ are *waiting times*. The differences

$$S_n = W_{n+1} - W_n$$

are *sojourn* or *interarrival times*.

Theorem 2.7.6. *Poisson Interarrival Times*
For a homogeneous Poisson process with rate λ, the interarrival times $S_0, S_1, \ldots$ are IID exponential random variables with density

$$f(s) = \lambda e^{-\lambda s}, \qquad s \geq 0.$$

Moreover,

$$W_n \sim \text{Gamma}(n, 1/\lambda),$$

so $\mathbb{E}(W_n) = n/\lambda$ and $\mathbb{V}(W_n) = n/\lambda^2$.

Proof. For the first interarrival time,

$$\mathbb{P}(S_0 > t) = \mathbb{P}(X(t) = 0) = e^{-\lambda t}.$$

Thus S_0 is exponential with rate λ. By independent increments, after an event occurs the future waiting time has the same distribution and is independent of the past. Repeating this argument gives IID exponential interarrivals. The waiting time W_n is a sum of n IID exponential variables, hence has the displayed Gamma distribution. □

Remark 2.7.7. *Uniform Event-Time Diagnostic*
If a homogeneous Poisson process is observed on $[0, T]$, then conditional on $X(T) = N$, the event times behave like a sample from Uniform$(0, T)$. This gives a simple diagnostic: if the observed event times are visibly nonuniform over $[0, T]$, the homogeneous Poisson assumption is suspect.

Proposition 2.7.8. *Conditional Binomial Count*
Let $X(t)$ be a homogeneous Poisson process with rate λ. For $s, t > 0$,

$$X(t) \mid X(t+s) = n \sim \text{Binomial}\left(n, \frac{t}{t+s}\right).$$

Proof. The increments $X(t)$ and $X(t+s) - X(t)$ are independent Poisson variables with means λt and λs. For $0 \leq k \leq n$,

$$\mathbb{P}(X(t) = k \mid X(t+s) = n) = \frac{\frac{(\lambda t)^k}{k!}\frac{(\lambda s)^{n-k}}{(n-k)!}}{\frac{(\lambda(t+s))^n}{n!}},$$

which is the Binomial mass function with success probability $t/(t+s)$. □

Proposition 2.7.9. *Odd Poisson Count*
For a homogeneous Poisson process,

$$\mathbb{P}(X(t) \text{ is odd}) = \frac{1 - e^{-2\lambda t}}{2}.$$

Proof. Let $\mu = \lambda t$. Since $X(t) \sim \text{Poisson}(\mu)$,

$$\mathbb{P}(X(t) \text{ is odd}) = e^{-\mu} \sum_{k=0}^{\infty} \frac{\mu^{2k+1}}{(2k+1)!} = e^{-\mu} \sinh(\mu) = \frac{1 - e^{-2\mu}}{2}.$$

□

Proposition 2.7.10. *Nearest Neighbor in a Planar Poisson Process*
Consider a planar Poisson point process with intensity λ, so the number of points in any region A is Poisson with mean λ area(A) and disjoint regions have independent counts. If R is the distance from a fixed point x_0 to the nearest random point, then

$$\mathbb{P}(R > t) = e^{-\lambda \pi t^2}, \qquad \mathbb{E}(R) = \frac{1}{2\sqrt{\lambda}}.$$

Proof. The event $\{R > t\}$ is the event that no process point falls in the disk of radius t around x_0. That disk has area πt^2, so the no-point probability is $e^{-\lambda \pi t^2}$. Since $R \geq 0$,

$$\mathbb{E}(R) = \int_0^\infty \mathbb{P}(R > t)\, dt = \int_0^\infty e^{-\lambda \pi t^2}\, dt = \frac{1}{2\sqrt{\lambda}}.$$

□

3 Stochastic Calculus and Dynamics

3.1 Brownian Motion

3.2 Stochastic Integration

3.3 Itô Calculus

3.4 Stochastic Differential Equations

3.5 Change of Measure and Girsanov Theory

3.6 Diffusions and PDE Connections

3.7 Stochastic Flows and Advanced Dynamics

Part VII
Dynamical and PDE Systems

1 Dynamical Systems and ODEs

1.1 ODE Foundations and Flows

1.2 Linear Systems and Stability

Method 1.2.1. *Solving $\mathbf{x}'(t) = A\mathbf{x}(t)$ by Jordan Form*
To solve $\mathbf{x}'(t) = A\mathbf{x}(t)$ with $\mathbf{x}(0) = \mathbf{x}_0$:

(i) Put $A = SJS^{-1}$, where $J = \bigoplus_{i=1}^{r} J_{k_i}(\lambda_i)$.
(ii) Set $\mathbf{y}(t) = S^{-1}\mathbf{x}(t)$ and $\mathbf{y}_0 = S^{-1}\mathbf{x}_0$. Then $\mathbf{y}'(t) = J\mathbf{y}(t)$ and $\mathbf{x}(t) = S\mathbf{y}(t)$.
(iii) Split $\mathbf{y} = (\mathbf{y}_1, \dots, \mathbf{y}_r)$ according to the Jordan blocks. Then $\mathbf{y}_i'(t) = J_{k_i}(\lambda_i)\mathbf{y}_i(t)$.
(iv) For each block, use Proposition 3.5.18: $\mathbf{y}_i(t) = e^{tJ_{k_i}(\lambda_i)}\mathbf{y}_i(0) = e^{\lambda_i t}\sum_{q=0}^{k_i-1} \frac{t^q}{q!} J_{k_i}(0)^q \mathbf{y}_i(0)$.
(v) Reassemble $\mathbf{y}(t)$ and transform back: $\mathbf{x}(t) = S\mathbf{y}(t)$.

Thus a size-k block with eigenvalue λ contributes entries of the form $e^{\lambda t}p(t)$, where $\deg p < k$.

Perspective 1.2.2. *Jordan Functions and Polynomial Factors*
For a Jordan block $J_k(\lambda) = \lambda I_k + N$ with $N^k = 0$, any polynomial f satisfies $f(J_k(\lambda)) = \sum_{r=0}^{k-1} \frac{f^{(r)}(\lambda)}{r!} N^r$.
Nilpotent part adds derivative terms; exponentials gives polynomial factors (Proposition 3.5.18, Method 1.2.1).

Example 1.2.3. *Nilpotent Jordan Block Evolution*
$N = J_3(0) = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$. Since $N^3 = 0$, $e^{tN} = I_3 + tN + \frac{t^2}{2}N^2 = \begin{bmatrix} 1 & t & \frac{t^2}{2} \\ 0 & 1 & t \\ 0 & 0 & 1 \end{bmatrix}$.
Nilpotent Jordan blocks create polynomial, not exponential, growth along chains.

1.3 Nonlinear Stability Theory

1.4 Bifurcation Theory

1.5 Dynamical Systems Structure

1.6 Semigroups and Evolution Viewpoint

2 PDE and Evolution Equations

2.1 First-Order PDE and Characteristics

2.2 Elliptic Equations

2.3 Parabolic Equations

2.4 Hyperbolic Equations

2.5 Sobolev Spaces for PDE

Definition 2.5.1. *Weak Derivative*
Let $u \in L^1_{\mathrm{loc}}(\Omega)$ and let α be a multi-index. A function $v \in L^1_{\mathrm{loc}}(\Omega)$ is the *weak derivative* $D^\alpha u$ if

$$\int_\Omega u\, D^\alpha \varphi \, dx = (-1)^{|\alpha|} \int_\Omega v\, \varphi \, dx \qquad \text{for every } \varphi \in C_c^\infty(\Omega).$$

Definition 2.5.2. *Sobolev Spaces*
For $1 \le p \le \infty$ and an integer $k \ge 0$, the Sobolev space $W^{k,p}(\Omega)$ consists of functions $u \in L^p(\Omega)$ whose weak derivatives $D^\alpha u$ lie in $L^p(\Omega)$ for every $|\alpha| \le k$. The norm is

$$\|u\|_{W^{k,p}} := \sum_{|\alpha| \le k} \|D^\alpha u\|_{L^p}$$

with the usual supremum modification when $p = \infty$.

2.6 Weak Solutions and Variational Methods

2.7 Evolution Equations

Part VIII
Numerical and Computational Mathematics

1 Numerical Analysis

1.1 Floating Point and Error Analysis

1.2 Nonlinear Equations and Optimization Basics

1.3 Interpolation and Approximation

1.4 Numerical Differentiation and Quadrature

1.5 Numerical Stability and Convergence

2 Numerical Differential Systems

2.1 Numerical ODEs

2.2 Geometric Numerical Integration

2.3 Finite Difference PDE Methods

2.4 Finite Element Methods

2.5 Numerical SDEs

2.6 Computational Stability Across ODE/PDE/SDE

3 Monte Carlo and Stochastic Computation

3.1 Monte Carlo Foundations

Monte Carlo methods turn integrals into averages of simulated random variables. This is especially useful when the integral is high-dimensional or when the target density is known only up to a normalizing constant.

Remark 3.1.1. *Bayesian Integrals as a Motivation*
Given a prior density $f(\theta)$ and likelihood $\mathcal{L}(\theta)$, the posterior density is

$$f(\theta \mid X^n) = \frac{\mathcal{L}(\theta) f(\theta)}{\int \mathcal{L}(u) f(u)\, du}.$$

Posterior summaries are often integrals, for example

$$\mathbb{E}(\theta \mid X^n) = \frac{\int \theta\, \mathcal{L}(\theta) f(\theta)\, d\theta}{\int \mathcal{L}(\theta) f(\theta)\, d\theta}.$$

For multidimensional $\theta = (\theta_1, \dots, \theta_k)$, marginal posterior densities require integrating out the other coordinates, so direct analytic calculation can become infeasible.

Method 3.1.2. *Basic Monte Carlo Integration*
Let

$$I = \int h(x) f(x)\, dx$$

where f is a probability density. If $X_1, \dots, X_N$ are IID draws from f, then

$$\widehat{I}_N = \frac{1}{N} \sum_{i=1}^{N} h(X_i)$$

satisfies $\widehat{I}_N \to I$ in probability by the law of large numbers. If $Y_i = h(X_i)$ and

$$s^2 = \frac{1}{N-1} \sum_{i=1}^{N} (Y_i - \widehat{I}_N)^2,$$

then a simulation standard error estimate is

$$\widehat{\mathrm{se}} = \frac{s}{\sqrt{N}}.$$

An approximate $1-\alpha$ Monte Carlo interval is

$$\widehat{I}_N \pm z_{\alpha/2}\widehat{\mathrm{se}}.$$

Example 3.1.3. *Integrating on an Interval*
For $I = \int_a^b h(x)\,dx$, take $X \sim \mathrm{Uniform}(a,b)$ and write

$$I = \mathbb{E}\big((b-a)h(X)\big).$$

Thus the estimator is

$$\widehat{I}_N = \frac{b-a}{N}\sum_{i=1}^{N} h(X_i), \qquad X_i \sim \mathrm{Uniform}(a,b).$$

Example 3.1.4. *Estimating a Normal CDF*
Let $Z \sim N(0,1)$. To estimate $\Phi(x) = \mathbb{P}(Z \le x)$ without a table, take

$$h(s) = \mathbf{1}\{s \le x\}$$

and simulate $Z_1, \dots, Z_N \sim N(0,1)$. Then

$$\widehat{\Phi}(x) = \frac{1}{N}\sum_{i=1}^{N} \mathbf{1}\{Z_i \le x\}$$

is the fraction of simulated observations below x.

Example 3.1.5. *Two Binomial Proportions*
Suppose $X \sim \mathrm{Binomial}(n, p_1)$ and $Y \sim \mathrm{Binomial}(m, p_2)$ independently, and use independent flat priors on p_1, p_2. Then

$$p_1 \mid X \sim \mathrm{Beta}(X+1, n-X+1), \qquad p_2 \mid Y \sim \mathrm{Beta}(Y+1, m-Y+1).$$

For $\delta = p_2 - p_1$, draw posterior samples

$$P_1^{(i)} \sim \mathrm{Beta}(X+1, n-X+1), \qquad P_2^{(i)} \sim \mathrm{Beta}(Y+1, m-Y+1),$$

set $\delta^{(i)} = P_2^{(i)} - P_1^{(i)}$, and estimate posterior means or intervals from the simulated $\delta^{(i)}$ values.

Method 3.1.6. *Simulating a Posterior Function*
If a posterior target is a function $\delta = g(p_1, \dots, p_k)$, simulate posterior draws

$$(P_1^{(i)}, \dots, P_k^{(i)}), \qquad i = 1, \dots, N,$$

and compute $\delta^{(i)} = g(P_1^{(i)}, \dots, P_k^{(i)})$. The empirical distribution of the $\delta^{(i)}$ values estimates the posterior distribution of δ. This is the simulation principle behind the dose-response example in the source pages.

3.2 Variance Reduction

Method 3.2.1. *Importance Sampling*
To estimate

$$I = \int h(x) f(x)\,dx,$$

choose a proposal density g that is easy to sample from and is positive wherever $h(x)f(x)$ is nonzero. Then

$$I = \int \frac{h(x)f(x)}{g(x)} g(x)\,dx = \mathbb{E}_g\left(\frac{h(X)f(X)}{g(X)}\right).$$

With $X_1, \dots, X_N \sim g$, the importance-sampling estimator is

$$\widehat{I}_N = \frac{1}{N}\sum_{i=1}^{N} \frac{h(X_i)f(X_i)}{g(X_i)}.$$

Remark 3.2.2. *Choosing the Proposal*
Importance sampling can have very large or even infinite variance if g is too small where $|h|\,f$ is large. A practical rule supported by the source discussion is to choose g with a shape similar to $|h|\,f$ but with tails at least as thick as the target integrand.

Theorem 3.2.3. *Optimal Importance Density*
Among proposal densities g, the density that minimizes the variance of the estimator in Method 3.2.1 is

$$g^*(x) = \frac{|h(x)|\,f(x)}{\int |h(s)|\,f(s)\,ds}.$$

Proof. The variance differs from

$$\mathbb{E}_g\left[\left(\frac{h(X)f(X)}{g(X)}\right)^2\right] = \int \frac{h(x)^2 f(x)^2}{g(x)}\,dx$$

by the constant I^2, so it is enough to minimize this second moment. By Cauchy-Schwarz,

$$\left(\int |h(x)|\, f(x)\,dx\right)^2 = \left(\int \frac{|h(x)|\, f(x)}{\sqrt{g(x)}}\sqrt{g(x)}\,dx\right)^2 \le \int \frac{h(x)^2 f(x)^2}{g(x)}\,dx.$$

Equality holds when g is proportional to $|h|\, f$, which gives the displayed density. □

Example 3.2.4. *Rare Tail Probability*
To estimate $I = \mathbb{P}(Z > 3)$ for $Z \sim N(0,1)$, basic Monte Carlo uses $h(x) = \mathbf{1}\{x > 3\}$ and wastes most simulated draws far from the right tail. A proposal such as a Normal distribution centered near the tail makes more simulated points relevant; the estimator then weights each draw by $f(X_i)/g(X_i)$ to remain unbiased for the original probability.

3.3 Markov Chain Monte Carlo

Monte Carlo integration in Method 3.1.2 uses independent draws from the target density. Markov chain Monte Carlo instead constructs a Markov chain $X_0, X_1, \ldots$ whose stationary density is the target f. When the chain mixes well, sample averages from the chain are used as approximations to expectations under f.

Method 3.3.1. *Metropolis-Hastings*
Let f be the target density, known up to a normalizing constant, and let $q(y \mid x)$ be a proposal density that is easy to sample from. Choose X_0 arbitrarily. Given $X_i = x$, generate X_{i+1} as follows.

1. Draw a proposal $Y \sim q(y \mid x)$.
2. Compute
$$r(x, Y) = \min\left\{\frac{f(Y)q(x \mid Y)}{f(x)q(Y \mid x)}, 1\right\}.$$
3. Set
$$X_{i+1} = \begin{cases} Y, & \text{with probability } r(x,Y), \\ x, & \text{with probability } 1 - r(x,Y). \end{cases}$$

Remark 3.3.2. The acceptance ratio only uses ratios of f, so an unknown normalizing constant in f cancels. This is one reason the method is useful for posterior densities whose normalizing integrals are difficult to compute.

Proposition 3.3.3. *Detailed Balance for Metropolis-Hastings*
The transition rule in Method 3.3.1 satisfies detailed balance with respect to f:

$$f(x)p(x,y) = f(y)p(y,x),$$

where $p(x,y)$ denotes the off-diagonal transition density of the chain. Consequently, f is a stationary density for the chain.

Proof. For $x \neq y$, the probability density of proposing and accepting a move from x to y is

$$p(x,y) = q(y \mid x)\min\left\{\frac{f(y)q(x \mid y)}{f(x)q(y \mid x)}, 1\right\}.$$

If $f(x)q(y \mid x) \ge f(y)q(x \mid y)$, then

$$f(x)p(x,y) = f(y)q(x \mid y).$$

In the same case the reverse move has acceptance probability 1, so $f(y)p(y,x) = f(y)q(x \mid y)$. The other inequality is symmetric. Thus detailed balance holds, and stationarity follows by integrating both sides over x. □

Method 3.3.4. *Random-Walk Metropolis-Hastings*
A common proposal has the form

$$Y = X_i + \varepsilon_i,$$

where the increment has a symmetric density. Then $q(y \mid x) = q(x \mid y)$ and the acceptance probability simplifies to

$$r(x,y) = \min\left\{\frac{f(y)}{f(x)}, 1\right\}.$$

For example, one may take $\varepsilon_i \sim N(0, b^2)$. Very small b gives small moves and slow exploration; very large b proposes many points in low density regions and causes frequent rejection. The tuning parameter should be chosen so that the chain mixes reasonably well.

Example 3.3.5. *Cauchy Target*
For the Cauchy density

$$f(x) = \frac{1}{\pi(1+x^2)}$$

and a symmetric random-walk proposal, the Metropolis-Hastings acceptance probability is

$$r(x, y) = \min\left\{\frac{1+x^2}{1+y^2}, 1\right\}.$$

This example is a useful reminder that proposal scale affects efficiency rather than the target stationary distribution.

Method 3.3.6. *Independence Metropolis-Hastings*
If the proposal does not depend on the current state, so $Y \sim g$ and $q(y \mid x) = g(y)$, then

$$r(x, y) = \min\left\{\frac{f(y)g(x)}{f(x)g(y)}, 1\right\}.$$

This is close in spirit to importance sampling: g should be a tractable approximation to the target with enough mass in the important parts of the space.

Method 3.3.7. *Gibbs Sampling*
Suppose (X, Y) has joint density $f_{X,Y}$ and that the conditional densities $f_{X|Y}$ and $f_{Y|X}$ are easy to sample from. Starting at (X_0, Y_0), the Gibbs sampler repeats

$$X_{n+1} \sim f_{X|Y}(x \mid Y_n), \qquad Y_{n+1} \sim f_{Y|X}(y \mid X_{n+1}).$$

The same coordinatewise idea extends to higher dimensions. Gibbs sampling turns one high-dimensional sampling problem into a sequence of conditional one-dimensional or lower-dimensional sampling problems.

Method 3.3.8. *Metropolis within Gibbs*
If a full conditional in a Gibbs sampler is not directly simulable, replace that coordinate update by a Metropolis-Hastings step while holding the other coordinates fixed. In the bivariate case, one may update X by proposing $Z \sim q(z \mid X_n)$ and accepting it with probability

$$\min\left\{\frac{f(Z, Y_n)q(X_n \mid Z)}{f(X_n, Y_n)q(Z \mid X_n)}, 1\right\},$$

then update Y similarly, conditional on the newly updated X.

3.4 Sequential Monte Carlo

3.5 Rare-Event Simulation

3.6 Stochastic Optimization

Part IX
Statistical and Learning Systems

1 Statistical Learning

1.1 Statistical Inference Foundations

Statistical inference reverses the direction of a probability calculation. Probability starts with a data-generating process and asks what outcomes it produces; inference starts with observed data and asks what can be said about the process that produced them.

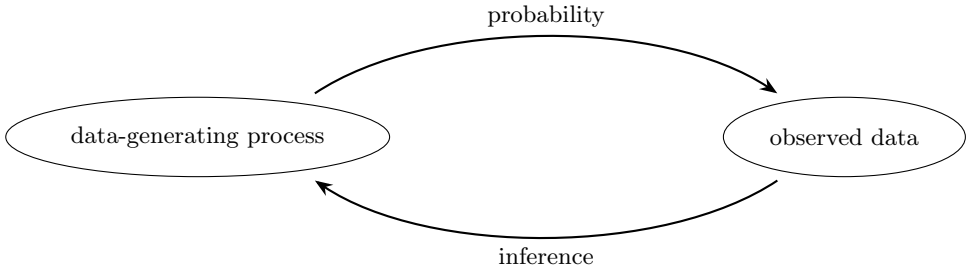


Figure 211: Probability pushes a data-generating process toward observed data; inference uses observed data to learn about the process.

Perspective 1.1.1. *Probability Versus Inference*
The basic probability question is: given a data-generating process, what are the properties of the outcomes? The basic inference question is the inverse one: given the outcomes, what can be said about the process? Prediction, classification, clustering, and estimation are all inference problems once the observed data are being used to learn something not directly observed.

Definition 1.1.2. *Model, Parameter, Functional, and Likelihood*
A *statistical model* is a set of candidate probability laws, distribution functions, density functions, or regression functions for the data. A *parameter* θ indexes a member or subfamily of a model, and $\widehat{\theta}$ denotes an estimate of θ. A *statistical functional* is a numerical target of the form $T(F)$, such as a mean or quantile, viewed as a function of a distribution F. For a parametric model, the *likelihood* $\mathcal{L}_n(\theta)$ records the joint mass or density of the observed data as a function of θ.

Remark 1.1.3. *Statistics and Data-Mining Vocabulary*
The same objects are often named differently across statistics, data mining, and machine learning. The following dictionary records the recurring translations from the source introduction.

Statistics	Computer science	Meaning
estimation	learning	use data to estimate an unknown quantity
classification	supervised learning	predict a discrete outcome Y from covariates $X \in \mathcal{X}$
clustering	unsupervised learning	put observations into groups
data	training sample	observations such as $(X_1, Y_1), \ldots, (X_n, Y_n)$
covariates	features	the observed predictors X_i
classifier	hypothesis	a map from covariates to outcomes
hypothesis	–	subset of a parameter space Θ
confidence interval	–	interval containing the unknown target with a prescribed long-run frequency
directed acyclic graph	Bayes net	multivariate model with specified conditional-independence relations
Bayesian inference	Bayesian inference	update subjective beliefs using data
frequentist inference	–	produce point estimates and confidence intervals with frequency guarantees
large-deviation bounds	PAC learning	uniform bounds on probabilities of error

Remark 1.1.4. *Source Scope*
The introductory screenshots motivate the probability/inference loop and the vocabulary dictionary. The remaining source material integrated here develops estimators, confidence intervals, empirical distribution functions, bootstrap procedures, parametric likelihood methods, hypothesis testing, Bayesian inference, and decision-theoretic comparison before the notes return to risk minimization and learning models.

1.1.1 Models, Estimators, and Confidence

Definition 1.1.5. *Parametric and Nonparametric Models*
A *parametric model* is a family $\mathcal{F} = \{f(x;\theta) : \theta \in \Theta\}$ with finite-dimensional parameter space $\Theta \subseteq \mathbb{R}^k$. A *nonparametric model* is not indexed by finitely many real parameters, such as the class of all distribution functions on $\mathbb{R}$. If $\theta = (\psi, \lambda)$ and only ψ is the target, then λ is a *nuisance parameter*.

Definition 1.1.6. *Estimator, Bias, Standard Error, and MSE*
For a quantity of interest θ, a *point estimator* is a statistic

$$\widehat{\theta}_n = g(X_1, \dots, X_n).$$

Its bias, standard error, and mean squared error are

$$\operatorname{bias}_\theta(\widehat{\theta}_n) = \mathbb{E}_\theta \widehat{\theta}_n - \theta, \qquad \operatorname{se}(\widehat{\theta}_n) = \sqrt{\operatorname{Var}_\theta(\widehat{\theta}_n)},$$

$$\operatorname{MSE}_\theta(\widehat{\theta}_n) = \mathbb{E}_\theta(\widehat{\theta}_n - \theta)^2.$$

The distribution of $\widehat{\theta}_n$ is its *sampling distribution*; the estimator is *consistent* if $\widehat{\theta}_n \xrightarrow{\mathbb{P}} \theta$.

Theorem 1.1.7. *MSE Equals Bias Squared Plus Variance*
If $\mathbb{E}_\theta \widehat{\theta}_n^2 < +\infty$, *then*

$$\operatorname{MSE}_\theta(\widehat{\theta}_n) = \operatorname{bias}_\theta(\widehat{\theta}_n)^2 + \operatorname{Var}_\theta(\widehat{\theta}_n).$$

Proof. Let $\overline{\theta}_n = \mathbb{E}_\theta \widehat{\theta}_n$. Expand

$$\widehat{\theta}_n - \theta = (\widehat{\theta}_n - \overline{\theta}_n) + (\overline{\theta}_n - \theta)$$

and take expectation. The cross term is zero. □

Example 1.1.8. *Bernoulli Sample Proportion*
If $X_i \sim \text{Bernoulli}(p)$ and $\widehat{p}_n = n^{-1}\sum_i X_i$, then

$$\mathbb{E}\widehat{p}_n = p, \qquad \operatorname{se}(\widehat{p}_n) = \sqrt{\frac{p(1-p)}{n}}, \qquad \widehat{\operatorname{se}}(\widehat{p}_n) = \sqrt{\frac{\widehat{p}_n(1-\widehat{p}_n)}{n}}.$$

Definition 1.1.9. *Confidence Interval*
A random interval $C_n = (a(X^n), b(X^n))$ is a $1-\alpha$ *confidence interval* for θ if

$$\mathbb{P}_\theta(\theta \in C_n) \geq 1 - \alpha \qquad \text{for all } \theta \in \Theta.$$

When

$$\frac{\widehat{\theta}_n - \theta}{\widehat{\operatorname{se}}} \rightsquigarrow N(0,1),$$

the normal-based interval

$$\widehat{\theta}_n \pm z_{\alpha/2}\widehat{\operatorname{se}}$$

has asymptotic coverage $1-\alpha$. For 95% intervals, $z_{\alpha/2} \approx 1.96$.

Example 1.1.10. *Two Bernoulli Intervals*
For Bernoulli data, the normal interval is

$$\widehat{p}_n \pm z_{\alpha/2}\sqrt{\frac{\widehat{p}_n(1-\widehat{p}_n)}{n}}.$$

Hoeffding's inequality gives the conservative finite-sample interval

$$\widehat{p}_n \pm \sqrt{\frac{1}{2n}\log\left(\frac{2}{\alpha}\right)}.$$

Definition 1.1.11. *Hypothesis Testing*
A *null hypothesis* H_0 is the default model or parameter restriction; an *alternative hypothesis* H_1 is the competing claim. For a coin, testing

$$H_0 : p = \frac{1}{2} \qquad \text{against} \qquad H_1 : p \neq \frac{1}{2}$$

can be based on $T = |\widehat{p}_n - 1/2|$, with large values giving evidence against H_0.

1.1.2 Empirical CDFs and Functionals

Definition 1.1.12. *Empirical Distribution Function*
For IID real-valued observations $X_1, \dots, X_n \sim F$, the *empirical distribution function* is

$$\widehat{F}_n(x) = \frac{1}{n} \sum_{i=1}^{n} \mathbf{1}_{\{X_i \le x\}}.$$

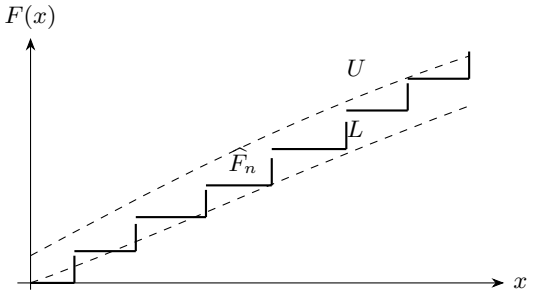


Figure 212: A schematic empirical CDF with a uniform confidence band. The source nerve-data plot has this structure, but the raw observations were not supplied.

Proposition 1.1.13. *Fixed-Point Moments of the Empirical CDF*
For fixed x,

$$\mathbb{E}\widehat{F}_n(x) = F(x), \qquad \operatorname{Var}(\widehat{F}_n(x)) = \frac{F(x)(1-F(x))}{n}.$$

Proof. The indicators $\mathbf{1}_{\{X_i \le x\}}$ are IID Bernoulli variables with success probability $F(x)$. □

Theorem 1.1.14. *Glivenko–Cantelli and DKW Bounds*
For IID data from F,

$$\sup_x |\widehat{F}_n(x) - F(x)| \xrightarrow{\mathbb{P}} 0.$$

Moreover,

$$\mathbb{P}\left(\sup_x |\widehat{F}_n(x) - F(x)| > \varepsilon \right) \le 2e^{-2n\varepsilon^2}.$$

Hence, for $\varepsilon_n = \sqrt{(2n)^{-1} \log(2/\alpha)}$, the functions

$$L(x) = \max\{\widehat{F}_n(x) - \varepsilon_n, 0\}, \qquad U(x) = \min\{\widehat{F}_n(x) + \varepsilon_n, 1\}$$

form a nonparametric $1-\alpha$ confidence band.

Definition 1.1.15. *Plug-in Estimator*
For a statistical functional $\theta = T(F)$, the *plug-in estimator* is $T(\widehat{F}_n)$. If $T(F) = \int r(x)\, dF(x)$ is linear, then

$$T(\widehat{F}_n) = \frac{1}{n} \sum_i r(X_i).$$

Thus the plug-in estimators of the mean, variance, skewness, correlation, and quantiles are the usual sample mean, empirical variance, empirical skewness, sample correlation, and empirical quantile.

Proposition 1.1.16. *Covariance of Empirical CDF Values*
If $x \le y$, then

$$\operatorname{Cov}(\widehat{F}_n(x), \widehat{F}_n(y)) = \frac{F(x)(1-F(y))}{n}.$$

Proof. Let $I_i = \mathbf{1}_{\{X_i \le x\}}$ and $J_i = \mathbf{1}_{\{X_i \le y\}}$. Since $I_i J_i = I_i$ and cross-observation covariances vanish, the covariance is $n^{-1}(F(x) - F(x)F(y))$. □

Example 1.1.17. *Interval Probability and Two Proportions*
For $a < b$, the interval probability $\theta = F(b) - F(a)$ has plug-in estimator

$$\widehat{\theta} = \widehat{F}_n(b) - \widehat{F}_n(a) = \frac{1}{n} \sum_i \mathbf{1}_{\{a < X_i \le b\}},$$

with estimated standard error $\sqrt{\widehat{\theta}(1-\widehat{\theta})/n}$. For independent Bernoulli samples, an estimated standard error for $\widehat{p}-\widehat{q}$ is

$$\sqrt{\frac{\widehat{p}(1-\widehat{p})}{n}+\frac{\widehat{q}(1-\widehat{q})}{m}}.$$

In the antibiotic exercise, $\widehat{p}_1=.90$, $\widehat{p}_2=.85$, $\widehat{\theta}=.05$, and $\widehat{\mathrm{se}}\approx .0466$, giving approximate 80% and 95% intervals $(-.010,.110)$ and $(-.041,.141)$.

1.1.3 Bootstrap

Definition 1.1.18. *Bootstrap Sample*
For a statistic $T_n=g(X_1,\dots,X_n)$, a *bootstrap sample* $X_1^*,\dots,X_n^*$ is obtained by drawing n observations with replacement from the observed data. The corresponding replication is $T_n^*=g(X_1^*,\dots,X_n^*)$.

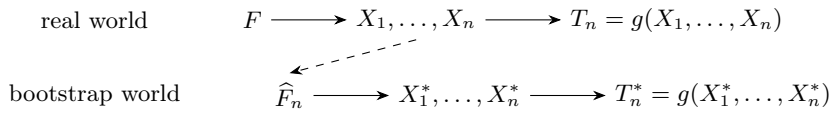


Figure 213: The bootstrap replaces the unknown law F by $\widehat{F}_n$ and recomputes the statistic on resampled data.

Method 1.1.19. *Bootstrap Standard Error*
From bootstrap replications $T_{n,1}^*,\dots,T_{n,B}^*$, estimate the standard error by

$$\widehat{\mathrm{se}}_{\mathrm{boot}}=\left[\frac{1}{B}\sum_{b=1}^{B}\left(T_{n,b}^*-\frac{1}{B}\sum_{r=1}^{B}T_{n,r}^*\right)^2\right]^{1/2}.$$

Proposition 1.1.20. *Bootstrap Mean Conditional Moments*
Given data $X_1,\dots,X_n$, if $\overline{X}_n^*$ is the bootstrap mean, then

$$\mathbb{E}(\overline{X}_n^*\mid X_1,\dots,X_n)=\overline{X}_n,\qquad \mathrm{Var}(\overline{X}_n^*\mid X_1,\dots,X_n)=\frac{1}{n}\left(\frac{1}{n}\sum_i(X_i-\overline{X}_n)^2\right).$$

Proposition 1.1.21. *Counting Distinct Bootstrap Samples*
If $X_1,\dots,X_n$ are distinct, then the number of distinct unordered bootstrap samples of size n is

$$\binom{2n-1}{n}.$$

Proof. An unordered bootstrap sample is a count vector $(N_1,\dots,N_n)$ with nonnegative entries and $\sum_i N_i=n$. Stars and bars gives the displayed count. □

Definition 1.1.22. *Bootstrap Confidence Intervals*
Let $\widehat{\theta}=T_n$ and let $\widehat{\theta}_\beta^*$ denote the empirical β quantile of the bootstrap replications. The normal, percentile, and pivotal bootstrap intervals are

$$\widehat{\theta}\pm z_{\alpha/2}\widehat{\mathrm{se}}_{\mathrm{boot}},\qquad (\widehat{\theta}_{\alpha/2}^*,\widehat{\theta}_{1-\alpha/2}^*),$$

$$(2\widehat{\theta}-\widehat{\theta}_{1-\alpha/2}^*,2\widehat{\theta}-\widehat{\theta}_{\alpha/2}^*).$$

Example 1.1.23. *Source Numerical Bootstrap Summaries*
For the nerve waiting-time data, the source reports plug-in skewness 1.76 and bootstrap standard error about .16 from $B=1000$ replications. For the plasma cholesterol difference of medians, it reports point estimate 18.5, bootstrap standard error 7.42, and approximate 95% intervals

method	95% interval
normal	$(3.7,33.3)$
percentile	$(5.0,33.3)$
pivotal	$(5.0,34.0)$.

Raw data were not supplied, so this pass records the reusable method and reported numerical summaries only.

Example 1.1.24. *Endpoint Failure of the Nonparametric Bootstrap*
For $X_i\sim\mathrm{Uniform}(0,\theta)$, the MLE is $\widehat{\theta}=X_{(n)}$. A parametric bootstrap from $\mathrm{Uniform}(0,\widehat{\theta})$ has $\mathbb{P}(\widehat{\theta}^*=\widehat{\theta}\mid X^n)=0$, but the nonparametric bootstrap has

$$\mathbb{P}(\widehat{\theta}^*=\widehat{\theta}\mid X^n)=1-\left(1-\frac{1}{n}\right)^n\to 1-e^{-1}\approx .632.$$

1.1.4 Parametric Inference

Method 1.1.25. *Method of Moments*
For $\theta = (\theta_1, \dots, \theta_k)$, define $\alpha_j(\theta) = \mathbb{E}_\theta(X^j)$. The method of moments estimator solves

$$\alpha_j(\widehat{\theta}) = \frac{1}{n}\sum_i X_i^j, \qquad j = 1, \dots, k.$$

For Gamma(α, β) with mean $\alpha\beta$ and variance $\alpha\beta^2$,

$$\widehat{\alpha} = \frac{\overline{X}_n^2}{n^{-1}\sum_i (X_i - \overline{X}_n)^2}, \qquad \widehat{\beta} = \frac{n^{-1}\sum_i (X_i - \overline{X}_n)^2}{\overline{X}_n}.$$

Definition 1.1.26. *Likelihood and MLE*
For IID data with mass or density $f(x;\theta)$,

$$\mathcal{L}_n(\theta) = \prod_i f(X_i;\theta), \qquad \ell_n(\theta) = \log \mathcal{L}_n(\theta).$$

A *maximum likelihood estimator* is a maximizer of $\mathcal{L}_n$.

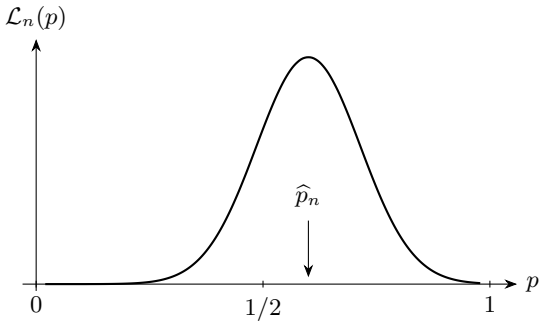


Figure 214: Schematic Bernoulli likelihood as a function of p, peaking at the observed sample proportion.

Example 1.1.27. *Basic MLEs*
For Bernoulli, Poisson, normal, and uniform endpoint models,

$$\widehat{p} = \frac{1}{n}\sum_i X_i, \qquad \widehat{\lambda} = \overline{X}_n, \qquad \widehat{\mu} = \overline{X}_n, \qquad \widehat{\sigma}^2 = \frac{1}{n}\sum_i (X_i - \overline{X}_n)^2, \qquad \widehat{\theta} = X_{(n)}.$$

For $X_i \sim \text{Uniform}(0, \theta)$, $\widehat{\theta} = X_{(n)}$ is consistent because

$$\mathbb{P}(|X_{(n)} - \theta| > \varepsilon) = \left(1 - \frac{\varepsilon}{\theta}\right)^n \to 0.$$

Definition 1.1.28. *Kullback–Leibler Distance*
For densities f and g,

$$D(f, g) = \int f(x) \log\left(\frac{f(x)}{g(x)}\right) dx.$$

It is nonnegative and $D(f, f) = 0$, but it is not a metric. In regular identifiable models, the MLE is consistent because the normalized log-likelihood converges to a limit uniquely maximized at the true parameter.

Definition 1.1.29. *Score and Fisher Information*
For a one-parameter model,

$$s(X;\theta) = \frac{\partial}{\partial\theta}\log f(X;\theta), \qquad I(\theta) = \text{Var}_\theta(s(X;\theta)).$$

Under regularity conditions, $I(\theta) = -\mathbb{E}_\theta \partial_\theta^2 \log f(X;\theta)$ and $I_n(\theta) = nI(\theta)$.

Theorem 1.1.30. *MLE Normal Approximation and Delta Method*
Under regularity conditions,

$$\frac{\widehat{\theta}_n - \theta}{\sqrt{1/I_n(\widehat{\theta}_n)}} \rightsquigarrow N(0, 1).$$

Hence $\widehat{\theta}_n \pm z_{\alpha/2}/\sqrt{I_n(\widehat{\theta}_n)}$ *is an approximate confidence interval. If* $\tau = g(\theta)$*, then*

$$\widehat{\text{se}}(g(\widehat{\theta})) = |g'(\widehat{\theta})|\widehat{\text{se}}(\widehat{\theta}).$$

Example 1.1.31. *Information and Delta Examples*
For Bernoulli, normal mean with known variance, and Poisson models,

$$I(p) = \frac{1}{p(1-p)}, \qquad I(\theta) = \frac{1}{\sigma^2}, \qquad I(\lambda) = \frac{1}{\lambda}.$$

For Bernoulli log-odds $\psi = \log(p/(1-p))$,

$$\widehat{\mathrm{se}}(\widehat{\psi}) = \frac{1}{\sqrt{n\widehat{p}(1-\widehat{p})}}.$$

For $X_i \sim N(\mu, \sigma^2)$ with known μ and $\psi = \log \sigma$, the delta-method standard error is $1/\sqrt{2n}$.

Method 1.1.32. *Parametric Bootstrap*
Fit the model to get $\widehat{\theta}$, simulate $X_1^*, \dots, X_n^* \sim f(x; \widehat{\theta})$, and recompute the estimator on each simulated sample. The spread of the replications estimates standard errors and confidence intervals.

Definition 1.1.33. *Sufficiency and Exponential Families*
A statistic $T(X^n)$ is sufficient if the conditional distribution of X^n given $T(X^n)$ does not depend on θ. By the factorization theorem, this is equivalent to

$$f(x^n; \theta) = g(T(x^n), \theta) h(x^n).$$

A one-parameter exponential family has the form

$$f(x; \theta) = h(x) \exp\{\eta(\theta) T(x) - B(\theta)\},$$

and $\sum_i T(X_i)$ is sufficient for IID data.

Theorem 1.1.34. *Rao–Blackwell Improvement*
If $\widehat{\theta}$ is an estimator and T is sufficient, then $\widetilde{\theta} = \mathbb{E}(\widehat{\theta} \mid T)$ has mean squared error no larger than that of $\widehat{\theta}$. For two Bernoulli trials, Rao–Blackwellizing X_1 by $T = X_1 + X_2$ gives $(X_1 + X_2)/2$.

1.1.5 Hypothesis Testing and p-values

Definition 1.1.35. *Hypotheses and Rejection Region*
A hypothesis test partitions a parameter space as $\Theta = \Theta_0 \sqcup \Theta_1$ and compares $H_0 : \theta \in \Theta_0$ against $H_1 : \theta \in \Theta_1$. A *rejection region* is a set R of sample outcomes for which H_0 is rejected. Usually $R = \{x : T(x) > c\}$, where T is a *test statistic* and c is a *critical value*.

Definition 1.1.36. *Power, Size, and Level*
For rejection region R, the *power function* is $\beta(\theta) = \mathbb{P}_\theta(X \in R)$. The *size* is

$$\alpha_* = \sup_{\theta \in \Theta_0} \beta(\theta),$$

and the test has *level* α if $\alpha_* \le \alpha$. Rejecting H_0 when H_0 is true is a type I error; failing to reject H_0 when H_1 is true is a type II error.

Example 1.1.37. *One-Sided Normal Mean Test*
Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ with σ known. To test $H_0 : \mu \le 0$ against $H_1 : \mu > 0$, reject when $\overline{X}_n > c$. Then

$$\beta(\mu) = 1 - \Phi\left(\frac{\sqrt{n}(c - \mu)}{\sigma}\right),$$

so the size is attained at $\mu = 0$. A level-α test uses $c = \sigma z_\alpha / \sqrt{n}$.

Definition 1.1.38. *Wald Test and p-value*
For $H_0 : \theta = \theta_0$ against $H_1 : \theta \neq \theta_0$, define

$$W = \frac{\widehat{\theta} - \theta_0}{\widehat{\mathrm{se}}}.$$

When $W \rightsquigarrow N(0,1)$ under H_0, the two-sided level-α test rejects when $|W| > z_{\alpha/2}$, and its p-value is approximately $2\Phi(-|w|)$ for the observed value w.

Theorem 1.1.39. *Wald Size and Confidence-Interval Equivalence*
If $W \rightsquigarrow N(0,1)$ under $\theta = \theta_0$, then $\mathbb{P}_{\theta_0}(|W| > z_{\alpha/2}) \to \alpha$. The Wald test rejects $H_0 : \theta = \theta_0$ exactly when θ_0 is outside

$$\left(\widehat{\theta} - z_{\alpha/2}\widehat{\mathrm{se}}, \widehat{\theta} + z_{\alpha/2}\widehat{\mathrm{se}}\right).$$

Example 1.1.40. *Common Wald Comparisons*
For independent binomial samples, test $p_1 - p_2 = 0$ with

$$W = \frac{\widehat{p}_1 - \widehat{p}_2}{\sqrt{\widehat{p}_1(1-\widehat{p}_1)/m + \widehat{p}_2(1-\widehat{p}_2)/n}}.$$

For paired data, use $W = \overline{D}/(S_D/\sqrt{n})$. For two independent means, use

$$W = \frac{\overline{X} - \overline{Y}}{\sqrt{s_1^2/m + s_2^2/n}}.$$

Definition 1.1.41. *p-value*
For level-α rejection regions R_α, the *p-value* is $\inf\{\alpha : T(x^n) \in R_\alpha\}$. If a test rejects exactly when $T(X^n) \geq c_\alpha$, then for observed $t = T(x^n)$,

$$\text{p-value} = \sup_{\theta \in \Theta_0} \mathbb{P}_\theta(T(X^n) \geq t).$$

Under a simple continuous null, the p-value is Uniform$(0, 1)$.

Remark 1.1.42. *p-value Warnings*
Small p-values are evidence against H_0, but they do not measure effect size or practical importance. A p-value is not $\mathbb{P}(H_0 \mid \text{data})$; posterior probabilities require a Bayesian prior.

Remark 1.1.43. *Interpreting Small p-values*
The source uses the following rough reporting scale for evidence against H_0:

p-value	informal evidence against H_0
$< .01$	very strong
.01–.05	strong
.05–.10	weak
$> .10$	little or none.

This is only a convention. A large p-value may occur because H_0 is true or because the test has low power, and statistical significance can coexist with a negligible practical effect.

Example 1.1.44. *Testing a Uniform Endpoint*
Let $X_i \sim \text{Uniform}(0, \theta)$ and $Y = \max_i X_i$. Test $H_0 : \theta = 1/2$ against $H_1 : \theta > 1/2$ by rejecting when $Y > c$. Since $\mathbb{P}_\theta(Y \leq y) = (y/\theta)^n$, the power is $\beta(\theta) = 1 - (c/\theta)^n$. Size .05 gives $c = \frac{1}{2}(.95)^{1/n}$. For $n = 20$, the p-value is $1 - (2y)^{20}$ for $y \leq 1/2$ and 0 for $y > 1/2$, so $y = .48$ gives about .56 and $y = .52$ gives 0.

Definition 1.1.45. *Pearson's χ^2 Statistic*
For $X = (X_1, \ldots, X_k) \sim \text{Multinomial}(n, p)$ and $H_0 : p = p_0$, put $O_j = X_j$ and $E_j = np_{0j}$. Pearson's statistic is

$$T = \sum_{j=1}^{k} \frac{(O_j - E_j)^2}{E_j}.$$

Under H_0, $T \rightsquigarrow \chi^2_{k-1}$.

Example 1.1.46. *Mendel's Peas*
For Mendel's model $p_0 = (9/16, 3/16, 3/16, 1/16)$ and observed counts $(315, 101, 108, 32)$ with $n = 556$, the expected counts are $(312.75, 104.25, 104.25, 34.75)$ and Pearson's statistic is about .47. Since $.47 < 7.815$, the 5% χ^2_3 critical value, the data do not reject the model.

Method 1.1.47. *Permutation Test*
For independent samples from F_X and F_Y, the null $F_X = F_Y$ makes labels exchangeable after conditioning on the pooled values. For a statistic T, the exact permutation p-value is

$$\frac{1}{N!} \sum_{j=1}^{N!} \mathbf{1}_{\{T_j > t_{\text{obs}}\}}, \qquad N = m + n,$$

or the corresponding average over sampled random permutations.

Example 1.1.48. *A Toy Permutation Calculation*
For $(X_1, X_2, Y_1) = (1, 9, 3)$ and $T = |\overline{X} - \overline{Y}|$, the observed statistic is 2. The six ordered reallocations give statistic values $2, 2, 7, 7, 5, 5$, so the permutation p-value for $T > 2$ is $4/6$.

Method 1.1.49. *Bonferroni and Benjamini–Hochberg*
Bonferroni rejects H_{0i} when $P_i < \alpha/m$, controlling the probability of at least one false rejection by the union bound. For false discoveries, let

$$\text{FDP} = \begin{cases} V/R, & R > 0, \\ 0, & R = 0, \end{cases} \qquad \text{FDR} = \mathbb{E}(\text{FDP}).$$

The Benjamini–Hochberg rule orders p-values $P_{(1)} \leq \cdots \leq P_{(m)}$, uses thresholds $\ell_i = i\alpha/(C_m m)$, and rejects up to the last i with $P_{(i)} \leq \ell_i$. It satisfies $\text{FDR} \leq (m_0/m)\alpha \leq \alpha$.

Proposition 1.1.50. *Bonferroni Controls Familywise Error*
If each true null is tested by rejecting only when $P_i < \alpha/m$, then the probability of falsely rejecting at least one true null is at most α.

Proof. Let R_i be the event that the ith true null is falsely rejected. Under the null p-value calibration recorded above,

$$\mathbb{P}(R_i) \le \frac{\alpha}{m}.$$

If there are $m_0 \le m$ true nulls, the union bound gives

$$\mathbb{P}\left(\bigcup_{i=1}^{m_0} R_i\right) \le \sum_{i=1}^{m_0} \mathbb{P}(R_i) \le m_0 \frac{\alpha}{m} \le \alpha.$$

□

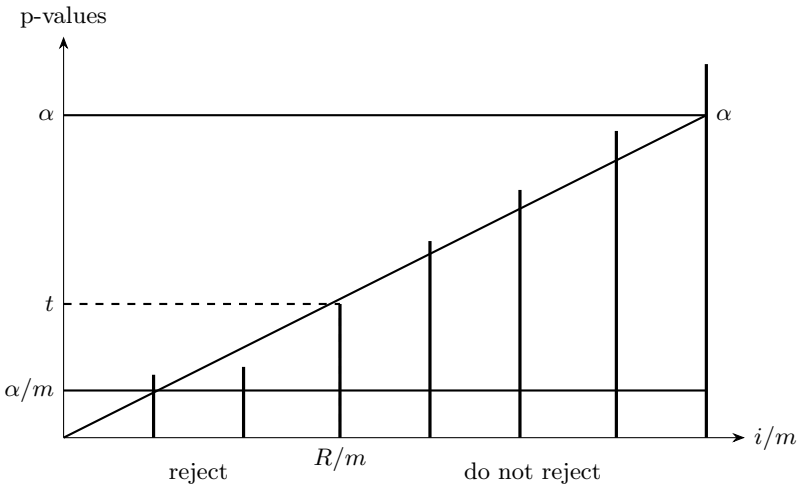


Figure 215: Schematic Benjamini–Hochberg threshold. The last ordered p-value under the sloped line determines the rejection threshold.

Method 1.1.51. *Technical Testing Tools*
The source also records the Neyman–Pearson lemma for most powerful tests of simple hypotheses, the one-sample t statistic $\sqrt{n}(\overline{X}_n - \mu_0)/S_n$ for normal means with unknown variance, and the likelihood-ratio statistic

$$2 \log \left(\frac{\sup_{\theta \in \Theta} \mathcal{L}(\theta)}{\sup_{\theta \in \Theta_0} \mathcal{L}(\theta)} \right),$$

which is asymptotically χ^2 with degrees of freedom equal to the number of constraints in regular nested models.

1.1.6 Bayesian Inference

Method 1.1.52. *Bayesian Update*
Choose a prior $\pi(\theta)$ and model $f(x \mid \theta)$. For data x^n,

$$\pi(\theta \mid x^n) = \frac{\mathcal{L}_n(\theta)\pi(\theta)}{\int \mathcal{L}_n(u)\pi(u)\,du} \propto \mathcal{L}_n(\theta)\pi(\theta).$$

Definition 1.1.53. *Posterior Summaries*
The posterior mean is $\overline{\theta}_n = \int \theta\, \pi(\theta \mid x^n)\, d\theta$. A set C is a *$1-\alpha$ posterior interval* if $\mathbb{P}(\theta \in C \mid x^n) = 1 - \alpha$.

Example 1.1.54. *Bernoulli–Beta Updating*
If $X_i \sim \text{Bernoulli}(p)$, $s = \sum_i X_i$, and $p \sim \text{Beta}(\alpha, \beta)$, then

$$p \mid x^n \sim \text{Beta}(\alpha + s, \beta + n - s).$$

The posterior mean is a weighted average of $\widehat{p} = s/n$ and the prior mean $\alpha/(\alpha+\beta)$.

Example 1.1.55. *Normal Mean with a Normal Prior*
Let $X_i \sim N(\theta, \sigma^2)$ with σ known and prior $\theta \sim N(a, b^2)$. Then $\theta \mid X^n \sim N(\overline{\theta}, \tau^2)$, where

$$\overline{\theta} = w\overline{X}_n + (1-w)a, \qquad w = \frac{1/\,\text{se}^2}{1/\,\text{se}^2 + 1/b^2}, \qquad \frac{1}{\tau^2} = \frac{1}{\text{se}^2} + \frac{1}{b^2}.$$

Method 1.1.56. *Conjugacy, Priors, and Simulation*
A prior is *conjugate* if the posterior stays in the same family. Flat priors can be improper and are not invariant under reparameterization; Jeffreys' one-parameter rule uses $\pi(\theta) \propto I(\theta)^{1/2}$. For $\tau = g(\theta)$, draw $\theta^{(b)}$ from the posterior and use $g(\theta^{(b)})$ to approximate the posterior of τ.

Remark 1.1.57. *Flat Priors Depend on Coordinates*
The source highlights a common trap: a prior that is flat in one coordinate need not be flat after reparameterization. For Bernoulli p, the flat prior $\pi(p) = 1$ induces on $\psi = \log(p/(1-p))$ the density

$$\pi_\Psi(\psi) = \frac{e^\psi}{(1+e^\psi)^2},$$

so flatness is not transformation invariant. Jeffreys' rule is introduced as a coordinate-invariant alternative; in several parameters the analogous expression is proportional to $\sqrt{\det I(\theta)}$ when the Fisher information matrix is available.

Theorem 1.1.58. *Large-Sample Bayesian Normal Approximation*
Under regularity conditions, if $\widehat{\theta}_n$ is the MLE and $\widehat{\mathrm{se}} = \{nI(\widehat{\theta}_n)\}^{-1/2}$, then the posterior is approximately $N(\widehat{\theta}_n, \widehat{\mathrm{se}}^{\,2})$.

Example 1.1.59. *Two Exercise-Derived Bayesian Posteriors*
If $X_i \sim \mathrm{Uniform}(0,\theta)$, $\pi(\theta) \propto 1/\theta$, and $y = X_{(n)}$, then

$$\pi(\theta \mid x^n) = n y^n \theta^{-(n+1)} \mathbf{1}_{\{\theta \ge y\}}.$$

If $X_i \sim \mathrm{Poisson}(\lambda)$, $s = \sum_i X_i$, and $\lambda \sim \mathrm{Gamma}(\alpha,\beta)$ with density proportional to $\lambda^{\alpha-1}e^{-\lambda/\beta}$, then

$$\lambda \mid x^n \sim \mathrm{Gamma}\left(\alpha + s, \frac{1}{n + 1/\beta}\right).$$

Jeffreys' prior for the Poisson mean is $\pi(\lambda) \propto \lambda^{-1/2}$.

Method 1.1.60. *Multiparameter Marginalization*
For $\theta = (\theta_1, \dots, \theta_p)$, inference about θ_1 uses

$$\pi(\theta_1 \mid x^n) = \int \cdots \int \pi(\theta_1, \dots, \theta_p \mid x^n)\, d\theta_2 \cdots d\theta_p.$$

Posterior simulation avoids computing this integral explicitly. For two binomial proportions with flat prior, independent beta posterior draws for p_1 and p_2 give draws from the posterior of $p_2 - p_1$ by subtraction.

Perspective 1.1.61. *Different Questions*
A posterior interval answers a conditional question about θ after choosing a prior and observing data. A confidence interval answers a long-run question about the random procedure that produced the interval.

1.1.7 Statistical Decision Theory

Definition 1.1.62. *Loss, Risk, Bayes Risk, and Minimax Risk*
A loss $L(\theta, a)$ measures the cost of action a when the truth is θ. The risk of $\widehat{\theta}(X)$ is

$$R(\theta, \widehat{\theta}) = \mathbb{E}_\theta L(\theta, \widehat{\theta}(X)).$$

The maximum risk and Bayes risk are

$$\overline{R}(\widehat{\theta}) = \sup_\theta R(\theta, \widehat{\theta}), \qquad r(\pi, \widehat{\theta}) = \int R(\theta, \widehat{\theta}) \pi(\theta)\, d\theta.$$

A Bayes rule minimizes Bayes risk; a minimax rule minimizes maximum risk.

Definition 1.1.63. *Posterior Risk*
Given observed data x, the posterior risk of an action or rule value $\widehat{\theta}(x)$ is

$$r(\widehat{\theta} \mid x) = \int L(\theta, \widehat{\theta}(x))\, \pi(\theta \mid x)\, d\theta.$$

If $m(x) = \int f(x \mid \theta)\pi(\theta)\, d\theta$ is the marginal density of the data, then

$$r(\pi, \widehat{\theta}) = \int r(\widehat{\theta} \mid x) m(x)\, dx.$$

Thus minimizing posterior risk pointwise in x minimizes Bayes risk.

Proposition 1.1.64. *Squared-Error Risk*
Under squared-error loss,

$$R(\theta, \widehat{\theta}) = \mathrm{Var}_\theta(\widehat{\theta}) + \mathrm{bias}_\theta(\widehat{\theta})^2.$$

Example 1.1.65. *Risk Functions Can Cross*
If $X \sim N(\theta, 1)$, $\widehat{\theta}_1 = X$, and $\widehat{\theta}_2 = 3$, then

$$R(\theta, \widehat{\theta}_1) = 1, \qquad R(\theta, \widehat{\theta}_2) = (3 - \theta)^2.$$

Neither estimator uniformly dominates the other.

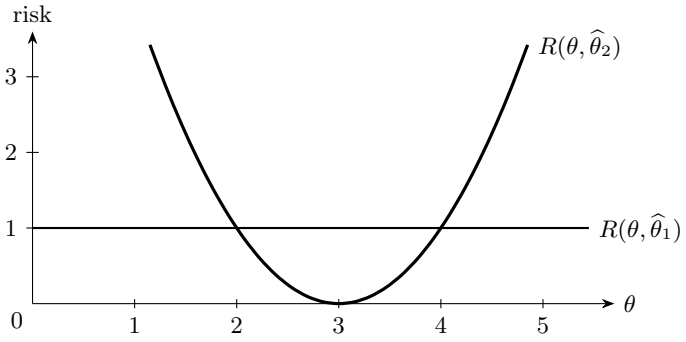


Figure 216: Two risk functions can cross, so pointwise risk comparison need not choose a single uniformly better estimator.

Theorem 1.1.66. *Bayes Rules and Common Losses*
A rule minimizing posterior risk pointwise in the observed data is a Bayes rule. Under squared-error loss the Bayes estimator is the posterior mean; under absolute-error loss it is a posterior median; under zero-one loss on a finite parameter space it is a posterior mode.

Theorem 1.1.67. *Minimax and Admissibility Tools*
If a Bayes rule $\widehat{\theta}^{\pi}$ satisfies $R(\theta, \widehat{\theta}^{\pi}) \le r(\pi, \widehat{\theta}^{\pi})$ for all θ, then it is minimax. In particular, a Bayes rule with constant risk is minimax. A rule is inadmissible if another rule has no larger risk everywhere and smaller risk somewhere; Bayes rules with full-support priors and finite Bayes risk are admissible under the source regularity assumptions.

Remark 1.1.68. *Admissibility and Full Support*
For one-dimensional $\Theta \subseteq \mathbb{R}$, the source calls a prior full-support if every neighborhood of every parameter value has positive prior mass. Under continuity and finite Bayes risk assumptions, a full-support Bayes rule cannot be improved everywhere: an improvement at one point would improve it on a small neighborhood, lowering the Bayes risk and contradicting Bayes optimality.

Example 1.1.69. *A Bernoulli Minimax Shrinkage Rule*
For Bernoulli data with squared-error loss,

$$\widehat{p}(X^n) = \frac{\sum_i X_i + \sqrt{n}/4}{n + \sqrt{n}}$$

is the posterior mean for a Beta$(\sqrt{n}/4, \sqrt{n}/4)$ prior and has constant risk $n/\{4(n+\sqrt{n})^2\}$. Hence it is minimax.

Corollary 1.1.70. *Normal Sample Mean Is Minimax*
For $X_1, \dots, X_n \sim N(\theta, 1)$ under squared-error loss, $\widehat{\theta} = \overline{X}_n$ is minimax: its risk is the constant $1/n$, and the source records that it is admissible.

Perspective 1.1.71. *Stein's Paradox and Shrinkage*
For $k \ge 3$ independent normal means, the usual estimator $X = (X_1, \dots, X_k)$ is inadmissible under squared-error loss. A James–Stein type estimator shrinks coordinates toward 0:

$$\widehat{\theta}_i^{S}(X) = \left(1 - \frac{k-2}{\sum_j X_j^2}\right)^{+} X_i.$$

The lesson is that shrinkage can reduce total risk in many-parameter problems.

1.2 Risk Minimization

Remark 1.2.1. *Optimisation*
A cost function $J(\boldsymbol{\theta})$ is reduced to improve some performance measure P.

Optimisation algorithms include some specialisation on specific structure of machine learning objective functions. Cost function is average over training set,

$$J(\boldsymbol{\theta}) = \mathbb{E}_{(\boldsymbol{x},y)\sim\hat{p}_{\text{data}}} L(f(\boldsymbol{x};\boldsymbol{\theta}), y)$$

where L is per-example loss function, $f(\boldsymbol{x};\boldsymbol{\theta})$ is predicted output given input $\boldsymbol{x}$, $\hat{p}_{\text{data}}$ is empirical distribution. For supervised learning, y is target output.
The goal is to usually to minimise objective function where expectation is taken on data generating distribution p_{data} rather than over the finite training set:

$$J^*(\boldsymbol{\theta}) = \mathbb{E}_{(\boldsymbol{x},y)\sim p_{\text{data}}} L(f(\boldsymbol{x};\boldsymbol{\theta}), y)$$

Note that machine learning algorithm usually minimises a surrogate loss function, and halts when convergence criterion based on early stopping is satisfied.

Definition 1.2.2. *Empirical Risk Minimisation*
To minimise expected loss on training set, where $\hat{p}(\boldsymbol{x}, y)$ is the empirical distribution. The empirical risk is

$$\mathbb{E}_{\boldsymbol{x},y\sim\hat{p}_{\text{data}}(\boldsymbol{x},y)}[L(f(\boldsymbol{x},\boldsymbol{\theta}), y)] = \frac{1}{m}\sum_{i=1}^{m} L(f(\boldsymbol{x}^{(i)};\boldsymbol{\theta}), y^{(i)})$$

where m is number of training samples. Note that empirical risk minimisation is prone to overfitting, as models with high capacity may memorise the training set.

Definition 1.2.3. *Classification Rules and Error Rates*
For supervised classification, the data are IID pairs

$$(X_1, Y_1), \ldots, (X_n, Y_n), \qquad X_i \in \mathcal{X} \subseteq \mathbb{R}^d, \quad Y_i \in \mathcal{Y},$$

where $\mathcal{Y}$ is finite. A *classification rule* is a function $h : \mathcal{X} \to \mathcal{Y}$. Its *true error rate* is

$$L(h) := \mathbb{P}\{h(X) \neq Y\},$$

and its *empirical* or *training error rate* is

$$\widehat{L}_n(h) := \frac{1}{n}\sum_{i=1}^{n} \mathbf{1}\{h(X_i) \neq Y_i\}.$$

Definition 1.2.4. *Bayes Classifier*
Assume first that $\mathcal{Y} = \{0, 1\}$. The *regression function* for classification is

$$r(x) := \mathbb{E}(Y \mid X = x) = \mathbb{P}(Y = 1 \mid X = x).$$

The *Bayes classification rule* is

$$h^*(x) = \begin{cases} 1, & r(x) > 1/2, \\ 0, & r(x) \leq 1/2. \end{cases}$$

The set

$$\mathcal{D} = \{x : \mathbb{P}(Y = 1 \mid X = x) = \mathbb{P}(Y = 0 \mid X = x)\}$$

is the *decision boundary*.

Remark 1.2.5. *Bayes Rule Versus Bayesian Inference*
The Bayes classification rule is named because it uses conditional probabilities and Bayes' theorem, not because the estimation procedure must be Bayesian. For class-conditional densities $f_k(x) = f(x \mid Y = k)$ and prior probabilities $\pi_k = \mathbb{P}(Y = k)$, Theorem 1.3.8 gives

$$\mathbb{P}(Y = 1 \mid X = x) = \frac{\pi_1 f_1(x)}{\pi_1 f_1(x) + \pi_0 f_0(x)}.$$

Theorem 1.2.6. *Optimality of the Bayes Classifier*
For every classifier $h : \mathcal{X} \to \{0, 1\}$,

$$L(h^*) \leq L(h).$$

Proof. Condition on $X = x$. If a rule predicts 1, its conditional error is $\mathbb{P}(Y = 0 \mid X = x) = 1 - r(x)$. If it predicts 0, its conditional error is $r(x)$. The smaller of these two numbers is obtained by predicting 1 exactly when $r(x) > 1/2$. Integrating the pointwise conditional error over X gives the claim. □

Proposition 1.2.7. *Multiclass Bayes Rule*
If $\mathcal{Y} = \{1, \ldots, K\}$, *the optimal classifier is*

$$h^*(x) \in \operatorname*{argmax}_{1\leq k\leq K} \mathbb{P}(Y = k \mid X = x) = \operatorname*{argmax}_{1\leq k\leq K} \pi_k f_k(x).$$

Proof. Given $X = x$, predicting class a has conditional error $1 - \mathbb{P}(Y = a \mid X = x)$. This is minimized by maximizing the posterior class probability. The density form follows from Bayes' theorem and the denominator is common to all classes. □

Method 1.2.8. *Three Ways to Approximate the Bayes Rule*
The Bayes rule depends on unknown population quantities. The source chapter isolates three standard approximation routes.

(i) *Empirical risk minimization*: choose a class $\mathcal{H}$ and minimize an estimate of $L(h)$ over $h \in \mathcal{H}$.
(ii) *Regression plug-in*: estimate $r(x) = \mathbb{P}(Y = 1 \mid X = x)$ and classify by thresholding $\widehat{r}(x)$ at $1/2$.
(iii) *Density plug-in*: estimate π_k and $f_k(x) = f(x \mid Y = k)$, then use the multiclass rule in Proposition 1.2.7.

Method 1.2.9. *Naive Bayes Classifier*
Naive Bayes is a density plug-in classifier that makes the simplifying assumption that, within each class,

$$f_k(x_1, \dots, x_d) = \prod_{j=1}^{d} f_{kj}(x_j),$$

where f_{kj} is the conditional density or mass function of the jth covariate given $Y = k$. Estimate each one-dimensional f_{kj} using only the observations with $Y_i = k$, estimate

$$\widehat{\pi}_k = \frac{1}{n} \sum_{i=1}^{n} \mathbf{1}\{Y_i = k\},$$

and classify by

$$\widehat{h}(x) \in \operatorname*{argmax}_k \widehat{\pi}_k \prod_{j=1}^{d} \widehat{f}_{kj}(x_j).$$

The conditional independence assumption is often false, but it reduces a high-dimensional density estimation problem to many one-dimensional ones and can still yield useful classifiers.

Example 1.2.10. *Linear Decision Boundary*
For two covariates $x = (x_1, x_2)$, a linear classifier has the form

$$h(x) = \begin{cases} 1, & a + b_1 x_1 + b_2 x_2 > 0, \\ 0, & a + b_1 x_1 + b_2 x_2 \le 0. \end{cases}$$

The line $a + b_1 x_1 + b_2 x_2 = 0$ is the decision boundary. Real data need not separate cleanly, so the boundary is usually a compromise rather than a perfect separator.

1.3 Linear and Generalized Linear Models

Remark 1.3.1. *Linear Models*
Linear models may fit data efficiently and reliably, either in closed form or with convex optimisation.
Model capacity limited to linear functions, does not model interaction between any two input variables

Definition 1.3.2. *Gaussian Discriminant Functions*
Assume $Y \in \{1, \dots, K\}$, $\pi_k = \mathbb{P}(Y = k)$, and

$$X \mid Y = k \sim N(\mu_k, \Sigma_k).$$

The Gaussian density plug-in rule in Method 1.2.8 uses the discriminant score

$$\delta_k(x) = -\frac{1}{2} \log |\Sigma_k| - \frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k) + \log \pi_k,$$

and classifies by $h(x) \in \operatorname{argmax}_k \delta_k(x)$.

Proposition 1.3.3. *QDA and LDA*
Under Definition 1.3.2, the decision boundary between classes j and k is quadratic in x when the covariance matrices may differ. This is quadratic discriminant analysis (QDA). If instead $\Sigma_1 = \cdots = \Sigma_K = \Sigma$, terms common to all classes drop out and one may use the linear discriminant

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k.$$

This is linear discriminant analysis (LDA).

Proof. The Bayes rule chooses the class maximizing $\pi_k f_k(x)$. Taking logarithms of the Gaussian density gives the displayed δ_k. If the matrices differ, the quadratic forms $(x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k)$ remain class-dependent. If all covariances are equal, the term $-\frac{1}{2} x^T \Sigma^{-1} x$ and the determinant term are common to all classes and can be removed. □

Method 1.3.4. *Sample Discriminant Estimates*
For binary labels $Y_i \in \{0,1\}$, write

$$n_0 = \sum_i (1 - Y_i), \qquad n_1 = \sum_i Y_i.$$

A conservative plug-in implementation uses

$$\widehat{\pi}_0 = \frac{n_0}{n}, \quad \widehat{\pi}_1 = \frac{n_1}{n}, \quad \widehat{\mu}_0 = \frac{1}{n_0} \sum_{Y_i=0} X_i, \quad \widehat{\mu}_1 = \frac{1}{n_1} \sum_{Y_i=1} X_i,$$

and within-class covariance estimates

$$S_j = \frac{1}{n_j} \sum_{Y_i=j} (X_i - \widehat{\mu}_j)(X_i - \widehat{\mu}_j)^T.$$

QDA inserts S_0, S_1 separately. LDA uses the pooled covariance

$$S = \frac{n_0 S_0 + n_1 S_1}{n_0 + n_1}.$$

Method 1.3.5. *Fisher Linear Discriminant*
For two classes, Fisher's method projects X to a scalar $U = w^T X$ and chooses w to make the projected class means far apart relative to within-class spread. With

$$S_B = (\overline{X}_0 - \overline{X}_1)(\overline{X}_0 - \overline{X}_1)^T, \qquad S_W = \frac{(n_0 - 1)S_0 + (n_1 - 1)S_1}{(n_0 - 1) + (n_1 - 1)},$$

the Rayleigh quotient is

$$\widehat{J}(w) = \frac{w^T S_B w}{w^T S_W w}.$$

The maximizing direction is proportional to

$$w = S_W^{-1}(\overline{X}_0 - \overline{X}_1).$$

Classification is then done by thresholding $w^T x$, often at the midpoint between the projected sample means. With equal priors and common covariance, this agrees with the Bayes linear discriminant in Proposition 1.3.3.

Definition 1.3.6. *Logistic Regression Classifier*
For binary classification, logistic regression models

$$r(x) = \mathbb{P}(Y = 1 \mid X = x) = \frac{\exp(\beta_0 + \sum_{j=1}^d \beta_j x_j)}{1 + \exp(\beta_0 + \sum_{j=1}^d \beta_j x_j)}.$$

Equivalently,

$$\operatorname{logit} \mathbb{P}(Y = 1 \mid X = x) = \log \frac{\mathbb{P}(Y = 1 \mid X = x)}{\mathbb{P}(Y = 0 \mid X = x)} = \beta_0 + \sum_{j=1}^d \beta_j x_j.$$

Given data, the likelihood is

$$\mathcal{L}(\beta) = \prod_{i=1}^n p_i(\beta)^{Y_i} (1 - p_i(\beta))^{1 - Y_i}.$$

The fitted classifier uses the plug-in rule $\widehat{h}(x) = 1$ when $\widehat{r}(x) > 1/2$.

Remark 1.3.7. *Richer Logistic Models*
Logistic regression can be made more flexible by adding interactions and polynomial terms, for example

$$\operatorname{logit} \mathbb{P}(Y = 1 \mid X = x) = \beta_0 + \sum_j \beta_j x_j + \sum_{j,k} \beta_{jk} x_j x_k.$$

Adding higher-order terms can reduce training error, but it also increases model complexity and can worsen out-of-sample error. This is the classification version of the bias-variance pressure formalized for estimators in Theorem 1.1.7.

Remark 1.3.8. *Linear Regression as a Plug-In Classifier*
One can also regress the binary outcome on the covariates by least squares and threshold the fitted value at $1/2$. This may produce useful predictions, but the fitted regression function is not constrained to lie in $[0,1]$. Logistic regression enforces this range by modeling the log-odds instead.

Proposition 1.3.9. *LDA and Logistic Regression Have the Same Linear Boundary*
In the two-class Gaussian model with common covariance Σ, the posterior log-odds are affine in x:

$$\log \frac{\mathbb{P}(Y = 1 \mid X = x)}{\mathbb{P}(Y = 0 \mid X = x)} = \alpha_0 + \alpha^T x.$$

Thus LDA and logistic regression both produce linear decision boundaries; they differ in how the parameters are estimated.

Proof. Using the Gaussian discriminants from Proposition 1.3.3,

$$\log \frac{\mathbb{P}(Y=1 \mid X=x)}{\mathbb{P}(Y=0 \mid X=x)} = \delta_1(x) - \delta_0(x).$$

With common Σ, the quadratic terms cancel and only an affine function of x remains. □

Remark 1.3.10. *Prediction Is Not Causation*
The source example with heart-disease covariates emphasizes that a small coefficient or a surprising sign in a multivariable logistic regression is not a causal conclusion. The coefficient measures predictive association relative to the other variables in the fitted model; unobserved confounding can change causal interpretation.

1.4 Linear Regression

Regression studies how a response Y changes with covariates X. The central target is the conditional mean

$$r(x) := \mathbb{E}(Y \mid X = x),$$

which may also be written as $\int y f_{Y|X}(y \mid x)\, dy$ when a conditional density is available. The linear model below is a parametric approximation to this regression function. The source first allows $\mathrm{Var}(\epsilon_i \mid X_i = x) = \sigma^2(x)$ and then makes the homoscedastic simplification $\sigma^2(x) = \sigma^2$.

Definition 1.4.1. *Simple Linear Regression Model*
The simple linear regression model assumes

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i, \qquad \mathbb{E}(\epsilon_i \mid X_i) = 0, \qquad \mathrm{Var}(\epsilon_i \mid X_i) = \sigma^2.$$

For estimates $\widehat{\beta}_0, \widehat{\beta}_1$, the fitted values and residuals are

$$\widehat{Y}_i = \widehat{\beta}_0 + \widehat{\beta}_1 X_i, \qquad \widehat{\epsilon}_i = Y_i - \widehat{Y}_i,$$

and the residual sum of squares is $\mathrm{RSS} = \sum_i \widehat{\epsilon}_i^2$.

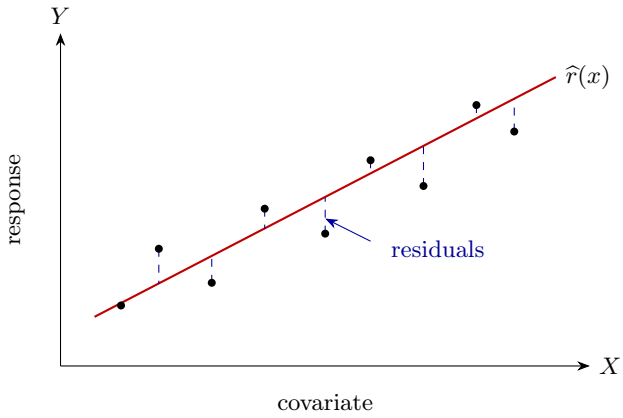


Figure 217: A fitted regression line with vertical residuals. Least squares minimizes the sum of squared residual lengths.

Theorem 1.4.2. *Simple Least-Squares Estimates*
Let

$$\overline{X} = \frac{1}{n}\sum_i X_i, \quad \overline{Y} = \frac{1}{n}\sum_i Y_i, \quad S_{XX} = \sum_i (X_i - \overline{X})^2, \quad S_{XY} = \sum_i (X_i - \overline{X})(Y_i - \overline{Y}).$$

If $S_{XX} > 0$, then

$$\widehat{\beta}_1 = \frac{S_{XY}}{S_{XX}}, \qquad \widehat{\beta}_0 = \overline{Y} - \widehat{\beta}_1 \overline{X}.$$

An unbiased estimate of σ^2 is

$$\widehat{\sigma}^2 = \frac{1}{n-2}\sum_i \widehat{\epsilon}_i^2.$$

Proof idea. Differentiate RSS with respect to β_0 and β_1. The resulting normal equations are the scalar version of Theorem 2.3.4. □

Proposition 1.4.3. *Normal Errors Make Least Squares Maximum Likelihood*
If, conditionally on the covariates, the errors are independent $N(0,\sigma^2)$ variables, then maximizing the conditional likelihood over (β_0,β_1) is equivalent to minimizing RSS. The maximum likelihood estimate of the error variance is

$$\widehat{\sigma}^2_{\mathrm{MLE}} = \frac{1}{n}\sum_i \widehat{\epsilon}_i^2.$$

Proof idea. The joint likelihood factors into the marginal likelihood of the covariates and the conditional likelihood of Y given X. The latter has log-likelihood, up to constants independent of (β_0,β_1,σ),

$$\ell(\beta_0,\beta_1,\sigma) = -n\log\sigma - \frac{1}{2\sigma^2}\sum_i (Y_i-\beta_0-\beta_1 X_i)^2 .$$

For fixed σ, maximizing ℓ over (β_0,β_1) is exactly minimizing RSS. Maximizing in σ after substituting the least-squares residuals gives the displayed MLE. □

Theorem 1.4.4. *Conditional Variances and Wald Inference*
Conditioning on $X^n=(X_1,\dots,X_n)$, the least-squares estimator is unbiased and

$$\mathrm{Var}\Big((\widehat{\beta}_0,\widehat{\beta}_1)^t \mid X^n\Big) = \frac{\sigma^2}{n s_X^2}\begin{pmatrix} n^{-1}\sum_i X_i^2 & -\overline{X} \\ -\overline{X} & 1\end{pmatrix}, \qquad s_X^2 = \frac{1}{n}\sum_i (X_i-\overline{X})^2.$$

Thus

$$\widehat{\mathrm{se}}(\widehat{\beta}_0) = \frac{\widehat{\sigma}}{s_X\sqrt{n}}\sqrt{\frac{1}{n}\sum_i X_i^2}, \qquad \widehat{\mathrm{se}}(\widehat{\beta}_1) = \frac{\widehat{\sigma}}{s_X\sqrt{n}}.$$

Approximate Wald intervals are $\widehat{\beta}_j \pm z_{\alpha/2}\widehat{\mathrm{se}}(\widehat{\beta}_j)$.

Method 1.4.5. *Wald Tests for Linear Contrasts*
For $H_0: c^t\beta = a$, use

$$W = \frac{c^t\widehat{\beta}-a}{\sqrt{c^t\widehat{\mathrm{Var}}(\widehat{\beta}\mid X^n)c}}.$$

The source exercise $H_0:\beta_1 = 17\beta_0$ is the contrast $c=(-17,1)^t$ with $a=0$.

Theorem 1.4.6. *Prediction Interval*
For a new covariate value x_*, set $\widehat{Y}_* = \widehat{\beta}_0+\widehat{\beta}_1 x_*$ and

$$\widehat{\xi}_n^2 = \widehat{\sigma}^2\left(\frac{\sum_i (X_i-x_*)^2}{n\sum_i (X_i-\overline{X})^2}+1\right).$$

An approximate $1-\alpha$ prediction interval is

$$\widehat{Y}_* \pm z_{\alpha/2}\widehat{\xi}_n.$$

Remark 1.4.7. *Prediction Versus Mean-Response Intervals*
A confidence interval for the mean response $\beta_0+\beta_1 x_*$ uses only the variability of $\widehat{Y}_* = \widehat{\beta}_0+\widehat{\beta}_1 x_*$. A prediction interval for a future observation must also cover

$$Y_* = \beta_0+\beta_1 x_* + \epsilon_*,$$

so its variance has the extra σ^2 term. This is the reason the interval in Theorem 1.4.6 is wider than the corresponding mean-response interval.

Example 1.4.8. *Source Regression Examples*
For the star data, the source reports $\widehat{\beta}_0 = 3.58$ and $\widehat{\beta}_1 = .166$, so $\widehat{r}(x) = 3.58 + .166x$. For the election data on the log scale,

$$\log(\text{Buchanan}) = -2.3298 + .7303\log(\text{Bush}),$$

with $\widehat{\mathrm{se}}(\widehat{\beta}_1) = .0358$, giving the approximate 95% slope interval $(.66,.80)$. For Palm Beach, the source prediction interval on the log scale is $(6.200, 6.578)$, corresponding roughly to $(493, 717)$ votes, far below the observed 3467.

Proposition 1.4.9. *Regression Through the Origin*
For $Y_i = \beta X_i + \epsilon_i$ with $\mathbb{E}(\epsilon_i \mid X_i)=0$ and $\mathrm{Var}(\epsilon_i\mid X_i)=\sigma^2$,

$$\widehat{\beta} = \frac{\sum_i X_iY_i}{\sum_i X_i^2}, \qquad \mathrm{Var}(\widehat{\beta}\mid X_1,\dots,X_n) = \frac{\sigma^2}{\sum_i X_i^2}.$$

Hence $\widehat{\beta}$ is consistent whenever the conditional variances are bounded and $\sum_i X_i^2\to\infty$.

Definition 1.4.10. *Multiple Linear Regression*
In matrix form, multiple regression is

$$Y = X\beta + \epsilon, \qquad \mathbb{E}(\epsilon \mid X) = 0.$$

An intercept is included by making one column of X equal to 1. If X^tX is invertible, then

$$\widehat{\beta} = (X^tX)^{-1}X^tY, \qquad \mathrm{Var}(\widehat{\beta} \mid X) = \sigma^2(X^tX)^{-1}.$$

The estimated regression function is

$$\widehat{r}(x) = \sum_{j=1}^{k} \widehat{\beta}_j x_j.$$

An unbiased variance estimate is

$$\widehat{\sigma}^2 = \frac{1}{n-k}\sum_i \widehat{\epsilon}_i^2, \qquad \widehat{\epsilon} = Y - X\widehat{\beta}.$$

An approximate confidence interval for β_j is

$$\widehat{\beta}_j \pm z_{\alpha/2}\widehat{\mathrm{se}}(\widehat{\beta}_j),$$

where $\widehat{\mathrm{se}}(\widehat{\beta}_j)^2$ is the jth diagonal entry of $\widehat{\sigma}^2(X^tX)^{-1}$.

Remark 1.4.11. *Regression Output Needs Context*
The source crime-data regression table reports estimates, standard errors, Wald t-statistics, and p-values for many covariates. Such a table answers a conditional predictive question inside the chosen model. It does not by itself decide which variables belong in the model, addressed in Method 1.6.3, or whether a coefficient has a causal interpretation, addressed later through counterfactual causal inference.

1.5 Regularization and Model Selection

Definition 1.5.1. *L^2 (Ridge) Regularisation*
The regularisation term added to objective function is

$$\Omega(\boldsymbol{\theta}) = \frac{1}{2}\|\boldsymbol{w}\|_2^2$$

Remark 1.5.2. *Behaviour of Weight Decay (L^2) Regularisation*
Assuming no bias parameter, a model have the following objective function and parameter gradient:

$$\tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = \frac{\alpha}{2}\boldsymbol{w}^T\boldsymbol{w} + J(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$$
$$\nabla_{\boldsymbol{w}}\tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = \alpha\boldsymbol{w} + \nabla_{\boldsymbol{w}}J(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$$

On a single gradient step, the update is as follows:

$$\boldsymbol{w} \leftarrow (1 - \epsilon\alpha)\boldsymbol{w} - \epsilon\nabla_{\boldsymbol{w}}J(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$$

The addition of weight decay has modified learning rule to multiplicatively shrink weight vector by constant factor on each step before gradient update.
Using quadratic approximation to objective function at minimal unregularised training cost, approximation is

$$\hat{J}(\overline{\theta}) = J(\boldsymbol{w}^*) + \frac{1}{2}(\boldsymbol{w} - \boldsymbol{w}^*)\boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*), \qquad \boldsymbol{w}^* = \arg\min_{\boldsymbol{w}} J(\boldsymbol{w})$$

where $\boldsymbol{H}$ is Hessian matrix of J with respect to $\boldsymbol{w}$ evaluated at $\boldsymbol{w}^*$. Minimum of $\hat{J}$ occurs where gradient is

$$\nabla_{\boldsymbol{w}}\hat{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*) = \boldsymbol{0}$$

Adding weight decay gradient, solve for minimum of regularised version of $\hat{J}$. Let $\tilde{\boldsymbol{w}}$ be minimum, then

$$\alpha\tilde{\boldsymbol{w}} + \boldsymbol{H}(\tilde{\boldsymbol{w}} - \boldsymbol{w}^*) = 0$$
$$(\boldsymbol{H} + \alpha\boldsymbol{I})\tilde{\boldsymbol{w}} = \boldsymbol{H}\boldsymbol{w}^*$$
$$\tilde{\boldsymbol{w}} = (\boldsymbol{H} + \alpha\boldsymbol{I})^{-1}\boldsymbol{H}\boldsymbol{w}^*$$

As $\alpha \to 0$, regularised solution $\tilde{\boldsymbol{w}} \to \boldsymbol{w}^*$. Note $\boldsymbol{H}$ is real and symmetric, hence decompose into diagonal matrix $\boldsymbol{A}$ and orthonormal basis of eigenvectors $\boldsymbol{Q}$ such that $\boldsymbol{H} = \boldsymbol{Q}\boldsymbol{\Lambda}\boldsymbol{Q}^T$, to get

$$\tilde{\boldsymbol{w}} = (\boldsymbol{Q}\boldsymbol{\Lambda}\boldsymbol{Q}^T + \alpha\boldsymbol{I})^{-1}\boldsymbol{Q}\boldsymbol{\Lambda}\boldsymbol{Q}^T\boldsymbol{w}^* = \boldsymbol{Q}(\boldsymbol{\Lambda} + \alpha\boldsymbol{I})^{-1}\boldsymbol{\Lambda}\boldsymbol{Q}^T\boldsymbol{w}^*$$

The effect of weight decay rescale $\boldsymbol{w}^*$ along axes defined by eigenvectors of $\boldsymbol{H}$. The component of $\boldsymbol{w}^*$ aligned with i-th eigenvector of $\boldsymbol{H}$ is rescaled by factor $\frac{\lambda_i}{\lambda_i + \alpha}$. For components where $\lambda_i \gg \alpha$, the effects of regularisation is relatively small. For components where $\lambda_i \ll \alpha$, components will be shrunk to nearly zero magnitude.

Definition 1.5.3. *L^1 Regularisation*
The L^1 regularisation is the sum of absolute values of individual parameters.

$$\Omega(\boldsymbol{\theta}) = \|\boldsymbol{w}\|_1 = \sum_i |w_i|$$

Definition 1.5.4. *Behaviour of L^1 Regularisation*
Assuming no bias parameter, a model has the following objective function and parameter gradient:

$$\tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = \alpha\|\boldsymbol{w}\|_1 + J(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$$
$$\nabla_{\boldsymbol{w}}\tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = \alpha\ \text{sign}(\boldsymbol{w}) + \nabla_{\boldsymbol{w}}J(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$$

Note that the regularisation contribution to gradient is a constant factor with sign equal to $\text{sign}(w_i)$. Minimum of $\tilde{J}$ occurs at where

$$\nabla_{\boldsymbol{w}}\tilde{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

Assuming the Hessian is diagonal, $\boldsymbol{H} = \text{diag}([H_{1,1}, \ldots, H_{n,n}])$, where each $H_{i,i} > 0$. The quadratic approximation of regularised objective function is then

$$\tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{w}^*; \boldsymbol{X}, \boldsymbol{y}) + \sum_i \left[\frac{1}{2}H_{i,i}(\boldsymbol{w}_i - \boldsymbol{w}_i^*)^2 + \alpha\,|w_i|\right]$$
$$w_i = \text{sign}(w_i^*)\max\left\{|w_i^*| - \frac{\alpha}{H_{i,i}}, 0\right\}$$

For the case where $w_i^* > 0$ for all i, then

i. if $w_i^* \leq \frac{\alpha}{H_{i,i}}$, the optimal value is $w_i = 0$ as contribution of $J(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$ to regularised objective $\tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y})$ is overwhelmed in direction i by L^1 regularisation which pushes w_i to zero.
ii. if $w_i^* > \frac{\alpha}{H_{i,i}}$, regularisation shifts optimal value of w_i in the direction by $\frac{\alpha}{H_{i,i}}$.

For the case where $w_i^* < 0$, this happens similarly, but with L^1 penalty decreasing w_i by $\frac{\alpha}{H_{i,i}}$, with min value 0. Note that L^1 produces a more sparse solution, which is a feature selection mechanism.

1.6 Regression Model Selection

Definition 1.6.1. *Prediction Risk and Training Error*
For a subset S of covariates, let $\widehat{Y}_i(S)$ be the fitted value from the model using only S. If Y_i^* is a future response at the same covariate value, define

$$R(S) = \sum_i \mathbb{E}\left(\widehat{Y}_i(S) - Y_i^*\right)^2, \qquad \widehat{R}_{\text{tr}}(S) = \sum_i (\widehat{Y}_i(S) - Y_i)^2.$$

Training error is optimistic because the same observations are used both to fit and to evaluate the model:

$$\mathbb{E}\widehat{R}_{\text{tr}}(S) - R(S) = -2\sum_i \text{Cov}(\widehat{Y}_i(S), Y_i).$$

Remark 1.6.2. *Bias–Variance Pressure in Subset Choice*
Adding covariates usually makes the fitted model more flexible: prediction bias can decrease, but estimation variance and overfitting risk can increase. The source frames model selection as two linked tasks: choose a score that estimates prediction risk, and search the set of candidate models for a good score.

Method 1.6.3. *Model Scores*
Common scores balance fit and complexity.

(i) Mallows' C_p:

$$\text{RSS}(S) + 2|S|\widehat{\sigma}^2,$$

where $\widehat{\sigma}^2$ is estimated from the full model.
(ii) AIC: maximize $\ell_S - |S|$, or equivalently minimize $-2\ell_S + 2|S|$.
(iii) BIC: maximize $\ell_S - \frac{|S|}{2}\log n$, or equivalently minimize $-2\ell_S + |S|\log n$.
(iv) Leave-one-out cross-validation:

$$\widehat{R}_{\text{CV}}(S) = \sum_i \left(\frac{Y_i - \widehat{Y}_i(S)}{1 - U_{ii}(S)}\right)^2, \quad U(S) = X_S(X_S^tX_S)^{-1}X_S^t.$$

With normal errors and a common variance estimate, maximizing AIC is equivalent to minimizing Mallows' C_p.

Remark 1.6.4. *Gaussian AIC and Mallows' C_p*
Under normal errors, and using a common variance estimate across the candidate models,

$$\ell_S = \text{constant} - \frac{\text{RSS}(S)}{2\widehat{\sigma}^2}.$$

Thus maximizing $\ell_S - |S|$ is equivalent to minimizing $\text{RSS}(S) + 2|S|\widehat{\sigma}^2$, which is Mallows' C_p in Method 1.6.3.

Method 1.6.5. *Stepwise Search*
If there are k optional covariates, exhaustive subset search compares 2^k models and is only practical for modest k. Forward stepwise search starts with no covariates and adds the one giving the best score improvement. Backward stepwise search starts with the largest model and drops one covariate at a time. Both are greedy searches and need not find the globally best subset.

Example 1.6.6. *Backward AIC on the Source Crime Data*
The source crime example starts from a full model with AIC about 310.37. Backward AIC first deletes population, then the southern-state indicator, and eventually leaves

$$\text{Crime} = 1.2\,\text{Age} + .75\,\text{Education} + .87\,\text{Expenditure} + .34\,\text{Males} - .86\,U1 + 2.31\,U2.$$

This is a predictive model-selection statement, not a causal conclusion.

Definition 1.6.7. *Lasso*
After putting covariates on a common scale, the lasso minimizes

$$\sum_i (Y_i - \widehat{Y}_i)^2 + \lambda \sum_{j=1}^{k} |\beta_j|, \qquad \lambda > 0.$$

The penalty can set some coefficients exactly to zero, so estimation and variable selection happen simultaneously. The tuning parameter λ is commonly chosen by cross-validation.

Example 1.6.8. *Lasso on the Source Crime Data*
For the source crime data, the leave-one-out cross-validation curve is minimized near $\lambda = .7$. At that value the source reports that all fitted coefficients remain nonzero, so the lasso selects the full model in that example. The useful lesson is methodological: the lasso searches over shrinkage levels rather than over discrete subsets directly.

Example 1.6.9. *AIC Selection in a Two-Normal-Model Exercise*
Compare $\mathcal{M}_0 : X_i \sim N(0,1)$ with $\mathcal{M}_1 : X_i \sim N(\theta, 1)$ and use $\widehat{\theta} = \overline{X}_n$. With $\text{AIC}_0 = \ell_n(0)$ and $\text{AIC}_1 = \ell_n(\widehat{\theta}) - 1$,

$$\ell_n(\widehat{\theta}) - \ell_n(0) = \frac{n\overline{X}_n^2}{2},$$

so $\mathcal{M}_0$ is selected exactly when $n\overline{X}_n^2 < 2$. If $\theta = 0$, the limiting probability of selecting $\mathcal{M}_0$ is $\mathbb{P}(\chi_1^2 < 2) \neq 1$; if $\theta \neq 0$, this probability tends to 0. Nevertheless the selected density is KL-consistent, because the possible overfit under $\theta = 0$ uses $N(\overline{X}_n, 1)$ with $\overline{X}_n \to 0$.

1.7 Kernels and Nonlinear Methods

Remark 1.7.1. *Nonlinear Models*
To represent nonlinear functions of $\boldsymbol{x}$, apply linear model to transformed input $\phi(\boldsymbol{x})$ where ϕ is a nonlinear transformation. Kernel trick may be applied to obtain nonlinear learning algorithm. To choose mapping ϕ:

i. Choose generic ϕ, such as that used by kernel machines based on RBF kernel.
If $\phi(\boldsymbol{x})$ is of high enough dimension, can find enough capacity to fit training set, but generalisations to test set remains poor. Mappings are based on principle of local smoothness and do not encode enough prior information to solve advanced problems.
ii. Manually engineer ϕ. Requires decades of human effort for each task, and practitioners specialising in different domains (i.e, speech recognition, computer vision) with little transfer between domains.
iii. Learn ϕ through deep learning. Model is $y = f(\boldsymbol{x}; \boldsymbol{\theta}, \boldsymbol{w}) = \phi(\boldsymbol{x}; \boldsymbol{\theta})^T \boldsymbol{w}$, where parameters $\boldsymbol{\theta}$ can be used to learn ϕ from broad class of functions, and parameters $\boldsymbol{w}$ that map from $\phi(\boldsymbol{x})$ to desired output.
Do not require training problem to be convex, and only require finding the right general function.

Definition 1.7.2. *Linear Separability and Margins*
For support vector machines it is convenient to code the labels as $Y_i \in \{-1, 1\}$. A linear classifier is

$$h(x) = \text{sign}(H(x)), \qquad H(x) = a_0 + a^T x.$$

If the data are linearly separable, the separating hyperplane can be rescaled so that

$$Y_i H(X_i) \geq 1, \qquad i = 1, \ldots, n.$$

The *margin* is the distance from the separating hyperplane to the nearest data point. Points lying on the margin are *support vectors*.

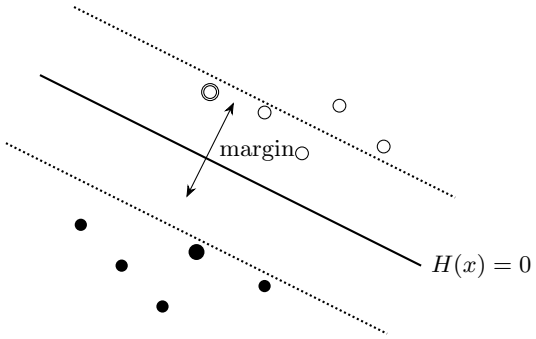


Figure 218: A maximum-margin classifier chooses the separating hyperplane whose closest supporting points are as far away as possible.

Method 1.7.3. *Hard-Margin Support Vector Machine*
Among all scaled separating hyperplanes satisfying $Y_i H(X_i) \geq 1$, the hard-margin SVM chooses the one minimizing

$$\frac{1}{2}\|a\|^2.$$

The dual form maximizes

$$\sum_{i=1}^{n} \alpha_i - \frac{1}{2}\sum_{i=1}^{n}\sum_{k=1}^{n} \alpha_i \alpha_k Y_i Y_k \langle X_i \rangle X_k$$

subject to $\alpha_i \geq 0$ and $\sum_i \alpha_i Y_i = 0$. The fitted normal vector is

$$\widehat{a} = \sum_{i=1}^{n} \widehat{\alpha}_i Y_i X_i,$$

so only observations with $\widehat{\alpha}_i > 0$ affect the boundary.

Remark 1.7.4. *Soft Margins*
When no perfect separating hyperplane exists, slack variables $\xi_i \geq 0$ allow overlap:

$$Y_i H(X_i) \geq 1 - \xi_i.$$

The tuning parameter in the constrained or penalized optimization controls how much margin violation is allowed. This is the SVM analogue of trading fit against model complexity.

Method 1.7.5. *Kernelization*
Kernelization maps $x \in \mathcal{X}$ to $\phi(x) \in \mathcal{Z}$ and applies a linear method in $\mathcal{Z}$. If the algorithm only needs inner products, and if

$$K(x, \widetilde{x}) = \langle \phi(x) \rangle \, \phi(\widetilde{x}),$$

then one can replace every inner product by K without explicitly constructing ϕ. Common examples include

$$K(x, \widetilde{x}) = (\langle x \rangle \, \widetilde{x} + a)^r, \qquad K(x, \widetilde{x}) = \tanh(a \langle x \rangle \, \widetilde{x} + b),$$

and the Gaussian kernel

$$K(x, \widetilde{x}) = \exp\left(-\frac{\|x - \widetilde{x}\|^2}{2\sigma^2}\right).$$

Example 1.7.6. *Quadratic Feature Map*
For $x = (x_1, x_2)$, the map

$$\phi(x) = (x_1^2, \sqrt{2}x_1 x_2, x_2^2)$$

satisfies

$$\langle \phi(x) \rangle \, \phi(\widetilde{x}) = (\langle x \rangle \, \widetilde{x})^2.$$

Thus a linear boundary in the feature coordinates can correspond to a quadratic boundary in the original coordinates.

Example 1.7.7. *A One-Dimensional Kernel Separation*
Suppose $X_i \in \mathbb{R}$ and $Y_i = 1$ whenever $|X_i| \leq 1$, while $Y_i = 0$ whenever $|X_i| > 1$. No threshold classifier on the

original line can perfectly separate both tails from the middle interval. With the feature map

$$Z_i = (X_i, X_i^2),$$

the classes are linearly separable: the line $z_2 = 1$ separates points with $z_2 \le 1$ from points with $z_2 > 1$.

1.8 Trees, Ensembles, and Boosting

Remark 1.8.1. *Bootstrap Aggregating (Bagging)*
Construct k different datasets, each with same number of examples as original dataset, but constructed by sampling with replacement from original dataset. Each dataset may miss some examples but also contain duplicate examples. Each model i is then trained on the dataset i. Differences in random initialisation, random selection of mini-batches, differences in hyperparameters, or different outcomes of non-deterministic implementation of neural networks are enough to allow each model to make partially independent errors.

Definition 1.8.2. *Classification Trees and Gini Impurity*
A classification tree partitions the covariate space $\mathcal{X}$ into disjoint leaves and assigns a class label to each leaf. For binary labels and a candidate split of a node into A_1 and A_2, let

$$\widehat{p}_s(j) = \frac{\sum_{i=1}^n \mathbf{1}\{Y_i = j,\ X_i \in A_s\}}{\sum_{i=1}^n \mathbf{1}\{X_i \in A_s\}}, \qquad s = 1, 2, \quad j = 0, 1.$$

The *Gini impurity* of leaf A_s is

$$\gamma_s = 1 - \sum_{j=0}^{1} \widehat{p}_s(j)^2,$$

and the impurity of the split is $\gamma_1 + \gamma_2$. For K classes, replace the sum over $j = 0, 1$ by a sum over $j = 1, \dots, K$.

Method 1.8.3. *Greedy Tree Construction*
A basic tree algorithm repeatedly chooses the covariate and split point that gives the lowest impurity among the candidate splits. Splitting stops when a stopping rule is met, such as a minimum leaf size. Each terminal leaf is assigned the majority class among the observations falling in that leaf. Large trees can overfit, so one often grows a large tree and then prunes it using validation or cross-validation.

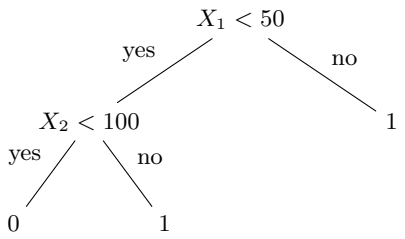


Figure 219: A small classification tree can be read as nested threshold rules.

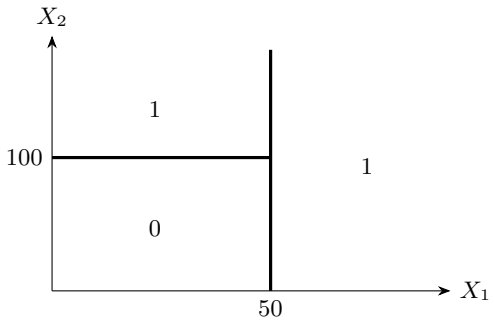


Figure 220: The same tree induces a rectangular partition of the covariate space.

Remark 1.8.4. *Boosting*
Construct k different datasets from original dataset. Using ith dataset, train ith weak learner. Test the ith weak learner with validation dataset, and each datapoint with wrong prediction is sent to $(i+1)$th dataset. The $(i+1)$th dataset is then used to train the $(i+1)$th weak learner. Continue until reaching the total number of subsets for the total prediction.

Method 1.8.5. *Bagging Classifiers*
Draw B bootstrap samples from the training data. Let h_b be the classifier trained on bootstrap sample b. For binary labels in $\{0,1\}$, the bagged classifier is the majority vote

$$\widehat{h}(x) = \begin{cases} 1, & B^{-1}\sum_{b=1}^{B} h_b(x) \geq 1/2, \\ 0, & \text{otherwise.} \end{cases}$$

Bagging is most useful for high-variance classifiers such as trees.

Method 1.8.6. *Boosting by Reweighting*
For labels $Y_i \in \{-1,1\}$ and weak learners $h_j : \mathcal{X} \to \{-1,1\}$, a basic boosting scheme is:

(i) Start with weights $w_i = 1/n$.

(ii) At step j, train h_j using the current weights and compute

$$\widehat{L}_j = \frac{\sum_{i=1}^{n} w_i \mathbf{1}\{Y_i \neq h_j(X_i)\}}{\sum_{i=1}^{n} w_i}.$$

(iii) Set $\alpha_j = \log\{(1-\widehat{L}_j)/\widehat{L}_j\}$ and increase the weights of misclassified observations by

$$w_i \leftarrow w_i \exp\{\alpha_j \mathbf{1}(Y_i \neq h_j(X_i))\}.$$

(iv) Return the weighted vote

$$\widehat{h}(x) = \operatorname{sign}\left(\sum_{j=1}^{J} \alpha_j h_j(x)\right).$$

In the source discussion, bagging mainly reduces variance, while boosting starts with a simple high-bias learner and gradually reduces bias.

1.9 Generalization Theory

The training error in Definition 1.2.3 is usually optimistic: after a classifier has been chosen to make $\widehat{L}_n(h)$ small, the same data no longer give an unbiased assessment of its true error. The source chapter develops two conservative remedies: hold-out validation and probability inequalities.

Method 1.9.1. *Validation and K-Fold Cross-Validation*
A validation split divides the data into a training set $\mathcal{T}$ and a validation set $\mathcal{V}$. The classifier is trained on $\mathcal{T}$ and assessed by

$$\widehat{L}_{\mathcal{V}}(h) = \frac{1}{|\mathcal{V}|}\sum_{i \in \mathcal{V}} \mathbf{1}\{h(X_i) \neq Y_i\}.$$

In K-fold cross-validation, the data are split into K approximately equal folds. For each k, fit the classifier on all folds except k, compute the error $\widehat{L}_{(k)}$ on fold k, and report

$$\widehat{L}_{\mathrm{CV}} = \frac{1}{K}\sum_{k=1}^{K} \widehat{L}_{(k)}.$$

Remark 1.9.2. *Validation, Test Sets, and Tree Size*
When refitting a classifier many times is too expensive, one can reserve a held-out test set, fit all candidate models on the training set, and compare their prediction errors on the held-out data. For classification trees, a common workflow is to grow a large tree, prune it to several smaller sizes, and use validation or cross-validation to choose the number of terminal nodes.

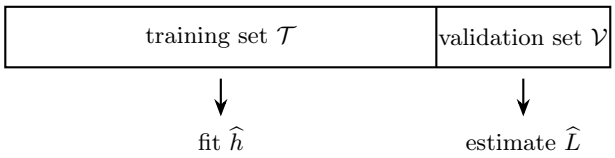


Figure 221: Validation estimates error on data not used to fit the classifier.

Theorem 1.9.3. *Finite-Class Uniform Convergence*
Let $\mathcal{H} = \{h_1, \dots, h_m\}$ be a finite class of classifiers. Then for every $\epsilon > 0$,

$$\mathbb{P}\left(\max_{h \in \mathcal{H}} \left|\widehat{L}_n(h) - L(h)\right| > \epsilon\right) \leq 2m\exp(-2n\epsilon^2).$$

Consequently, with probability at least $1-\alpha$, every $h \in \mathcal{H}$ satisfies

$$\left|\widehat{L}_n(h) - L(h)\right| \le \sqrt{\frac{1}{2n}\log\frac{2m}{\alpha}}.$$

Proof. For a fixed h, the variables $\mathbf{1}\{h(X_i) \neq Y_i\}$ are IID Bernoulli with mean $L(h)$, so Hoeffding's inequality gives $\mathbb{P}(\left|\widehat{L}_n(h) - L(h)\right| > \epsilon) \le 2e^{-2n\epsilon^2}$. Apply the union bound over the m classifiers. □

Definition 1.9.4. *Shatter Coefficients and VC Dimension*
Let $\mathcal{A}$ be a class of subsets of a space $\mathcal{X}$. For a finite set $F = \{x_1, \dots, x_n\}$, define

$$N_{\mathcal{A}}(F) = \#\{F \cap A : A \in \mathcal{A}\}.$$

The *shatter coefficient* is

$$s(\mathcal{A}, n) = \max_{|F|=n} N_{\mathcal{A}}(F).$$

The *VC dimension* of $\mathcal{A}$ is the largest k such that $s(\mathcal{A}, k) = 2^k$, if such a largest k exists; otherwise it is ∞. For a class of binary classifiers $\mathcal{H}$, take

$$\mathcal{A}_{\mathcal{H}} = \{\{x : h(x) = 1\} : h \in \mathcal{H}\}$$

and set $\mathrm{VC}(\mathcal{H}) = \mathrm{VC}(\mathcal{A}_{\mathcal{H}})$.

Theorem 1.9.5. *VC Uniform Bound*
For a class $\mathcal{H}$ of binary classifiers,

$$\mathbb{P}\left\{\sup_{h\in\mathcal{H}}\left|\widehat{L}_n(h) - L(h)\right| > \epsilon\right\} \le 8s(\mathcal{H}, n)\exp(-n\epsilon^2/32).$$

Thus a $1-\alpha$ error band is obtained by choosing

$$\epsilon_n^2 = \frac{32}{n}\log\left(\frac{8s(\mathcal{H}, n)}{\alpha}\right).$$

Example 1.9.6. *Basic VC Dimensions*
The class $\{(-\infty, a] : a \in \mathbb{R}\}$ has VC dimension 1. Closed intervals on the real line have VC dimension 2. Axis-parallel rectangles in the plane have VC dimension 4. Linear half-spaces in $\mathbb{R}^d$ have VC dimension $d+1$. Therefore, for linear classifiers in d dimensions, the source bound gives the conservative radius

$$\epsilon_n^2 = \frac{32}{n}\log\left(\frac{8(n^{d+1}+1)}{\alpha}\right).$$

1.10 Association and Independence Inference

This subsection treats dependence as an inferential target. The probability prerequisites are the definitions of independence and conditioning in Proposition 1.3.6, the multinomial and multivariate-normal families in Definition 1.1.32, and the delta method in Theorem 1.6.10.

Method 1.10.1. *Fisher z Interval for a Correlation*
For paired observations (X_{1i}, X_{2i}), let $\widehat{\rho}$ be the sample correlation. Define

$$f(r) = \frac{1}{2}\left(\log(1+r) - \log(1-r)\right), \qquad g(z) = f^{-1}(z) = \frac{e^{2z}-1}{e^{2z}+1}.$$

Set $\widehat{\theta} = f(\widehat{\rho})$. The source-supported normal approximation uses

$$\widehat{\mathrm{se}}(\widehat{\theta}) = \frac{1}{\sqrt{n-3}}.$$

Thus an approximate $1-\alpha$ interval for $\theta = f(\rho)$ is

$$\left(\widehat{\theta} - \frac{z_{\alpha/2}}{\sqrt{n-3}}, \widehat{\theta} + \frac{z_{\alpha/2}}{\sqrt{n-3}}\right),$$

and transforming the endpoints back gives a confidence interval for ρ.

Remark 1.10.2. *Covariance Matrices Must Be Positive Semidefinite*
A proposed covariance matrix must be positive semidefinite. One source exercise displays

$$\begin{pmatrix} 2 & 6 \\ 6 & 2 \end{pmatrix}$$

as a covariance matrix for a bivariate normal simulation. As written, this cannot be a covariance matrix because

$$(1, -1)\begin{pmatrix} 2 & 6 \\ 6 & 2 \end{pmatrix}\begin{pmatrix} 1 \\ -1 \end{pmatrix} = -8 < 0.$$

Equivalently, the implied correlation would have absolute value 3. A corrected covariance matrix is needed before that simulation exercise can be carried out.

Definition 1.10.3. *Two-by-Two Tables and Odds Ratio*
For binary variables $Z, Y \in \{0,1\}$, write

	$Y=0$	$Y=1$	
$Z=0$	X_{00}	X_{01}	$X_{0\cdot}$
$Z=1$	X_{10}	X_{11}	$X_{1\cdot}$
	$X_{\cdot 0}$	$X_{\cdot 1}$	$n = X_{\cdot\cdot}$

for counts, and use the same dotted-subscript convention for probabilities $p_{ij} = \mathbb{P}(Z=i, Y=j)$. Assuming positive cell probabilities, the *odds ratio* and *log odds ratio* are

$$\psi = \frac{p_{00}p_{11}}{p_{01}p_{10}}, \qquad \gamma = \log \psi.$$

Proposition 1.10.4. *Odds Ratio and Independence*
For a two-by-two table with positive cell probabilities, the following are equivalent:

(i) $Y \perp\!\!\!\perp Z$.
(ii) $p_{ij} = p_{i\cdot}p_{\cdot j}$ *for all* $i, j \in \{0,1\}$.
(iii) $\psi = 1$.
(iv) $\gamma = 0$.

Proof. The equivalence of (i) and (ii) is the discrete factorization test for independence. If (ii) holds, then

$$\psi = \frac{p_{0\cdot}p_{\cdot 0}p_{1\cdot}p_{\cdot 1}}{p_{0\cdot}p_{\cdot 1}p_{1\cdot}p_{\cdot 0}} = 1.$$

Conversely, $\psi = 1$ says $p_{00}p_{11} = p_{01}p_{10}$, so the 2×2 probability matrix has determinant 0 and rank 1. Since its entries sum to 1, it factors as its row marginals times its column marginals. Finally, $\gamma = \log \psi$, so $\psi = 1$ iff $\gamma = 0$. □

Theorem 1.10.5. *Likelihood-Ratio and Pearson Tests for a Two-by-Two Table*
For testing

$$H_0 : Y \perp\!\!\!\perp Z \qquad \text{against} \qquad H_1 : Y \not\perp\!\!\!\perp Z,$$

the likelihood-ratio statistic is

$$T = 2\sum_{i=0}^{1}\sum_{j=0}^{1} X_{ij} \log\left(\frac{X_{ij}X_{\cdot\cdot}}{X_{i\cdot}X_{\cdot j}}\right),$$

with the usual convention that zero-count terms contribute 0. *Pearson's statistic is*

$$U = \sum_{i=0}^{1}\sum_{j=0}^{1} \frac{(X_{ij} - E_{ij})^2}{E_{ij}}, \qquad E_{ij} = \frac{X_{i\cdot}X_{\cdot j}}{n}.$$

Under H_0, *both statistics are approximately* χ^2_1 *for large samples, so a level* α *test rejects for large values.*

Method 1.10.6. *Estimating an Odds Ratio*
The plug-in estimates are

$$\widehat{\psi} = \frac{X_{00}X_{11}}{X_{01}X_{10}}, \qquad \widehat{\gamma} = \log \widehat{\psi}.$$

The delta method gives

$$\widehat{\mathrm{se}}(\widehat{\gamma}) = \sqrt{\frac{1}{X_{00}} + \frac{1}{X_{01}} + \frac{1}{X_{10}} + \frac{1}{X_{11}}}, \qquad \widehat{\mathrm{se}}(\widehat{\psi}) = \widehat{\psi}\,\widehat{\mathrm{se}}(\widehat{\gamma}).$$

An approximate interval for γ is

$$\widehat{\gamma} \pm z_{\alpha/2}\widehat{\mathrm{se}}(\widehat{\gamma}),$$

and exponentiating its endpoints gives the usually preferable interval for ψ.

Remark 1.10.7. *Small-Sample Stabilization*
If some cell counts are small, $\widehat{\psi}$ and $\widehat{\gamma}$ can have large variance or even be undefined. A common stabilizing modification is

$$\widetilde{\psi} = \frac{(X_{00} + 1/2)(X_{11} + 1/2)}{(X_{01} + 1/2)(X_{10} + 1/2)}.$$

Example 1.10.8. *Tonsillectomy and Hodgkin Disease*
The source table records

	Hodgkin disease	No disease	
Tonsillectomy	90	165	255
No tonsillectomy	84	307	391
	174	472	646

The source reports the likelihood-ratio statistic $T = 14.75$ and Pearson statistic $U = 14.96$, both with p-value about .0001. The odds-ratio estimate is

$$\widehat{\psi} = \frac{90 \cdot 307}{165 \cdot 84} \approx 1.99, \qquad \widehat{\gamma} \approx .69,$$

and

$$\widehat{\mathrm{se}}(\widehat{\gamma}) = \sqrt{\frac{1}{90} + \frac{1}{84} + \frac{1}{165} + \frac{1}{307}} \approx .18.$$

Thus a 95% interval for ψ from the log scale is approximately $(1.39, 2.86)$. This establishes association in the table, not causation.

Proposition 1.10.9. *The Same Odds-Ratio Estimate Across Common Sampling Designs*
For a two-by-two exposure–disease table, the estimate

$$\widehat{\psi} = \frac{X_{00}X_{11}}{X_{01}X_{10}}$$

arises under multinomial sampling, prospective cohort sampling, and retrospective case-control sampling.

Proof idea. Under multinomial sampling, estimate all cell probabilities by X_{ij}/n and plug into Definition 1.10.3. Under prospective sampling, estimate the two conditional disease probabilities within exposed and unexposed groups; the same cross-product ratio results. Under case-control sampling, estimate the exposure odds among diseased and non-diseased groups; the same cross-product ratio results again. □

Proposition 1.10.10. *Rare Disease Approximation*
Let

$$\xi = \frac{\mathbb{P}(D \mid E)}{\mathbb{P}(D \mid E^c)}$$

be the relative risk and let

$$\psi = \frac{\mathrm{odds}(D \mid E)}{\mathrm{odds}(D \mid E^c)}$$

be the odds ratio. If both disease probabilities tend to 0*, then*

$$\frac{\psi}{\xi} \to 1.$$

Proof. Write $p_1 = \mathbb{P}(D \mid E)$ and $p_0 = \mathbb{P}(D \mid E^c)$. Then

$$\psi = \frac{p_1/(1-p_1)}{p_0/(1-p_0)} = \xi \, \frac{1-p_0}{1-p_1}.$$

As $p_0, p_1 \to 0$, the final factor tends to 1. □

Example 1.10.11. *Death Sentencing and Race Exercise*
The source exercise gives the following Maryland death-sentencing table:

	Death sentence	No death sentence
Black victim	14	641
White victim	62	594

With this table orientation,

$$\widehat{\psi} = \frac{14 \cdot 594}{641 \cdot 62} \approx .209, \qquad \widehat{\gamma} \approx -1.56,$$

and

$$\widehat{\mathrm{se}}(\widehat{\gamma}) \approx .301.$$

The corresponding log-scale 95% interval exponentiates to approximately $(.116, .378)$. Pearson's statistic is about 32.10 and the likelihood-ratio statistic is about 34.53, so the aggregate association is very strong. As the source warns, these counts alone do not justify a causal conclusion because they do not adjust for other case characteristics.

Theorem 1.10.12. *Independence Tests for an I-by-J Table*
For counts X_{ij} in an I-by-J table, the likelihood-ratio statistic

$$T = 2\sum_{i=1}^{I}\sum_{j=1}^{J} X_{ij}\log\left(\frac{X_{ij}X_{\cdot\cdot}}{X_{i\cdot}X_{\cdot j}}\right)$$

and Pearson statistic

$$U = \sum_{i=1}^{I}\sum_{j=1}^{J}\frac{(X_{ij}-E_{ij})^2}{E_{ij}}, \qquad E_{ij} = \frac{X_{i\cdot}X_{\cdot j}}{n},$$

are approximately χ^2_ν under independence, where

$$\nu = (I-1)(J-1).$$

Example 1.10.13. *Treatment Response and Histological Type*
For the source Hodgkin-disease table with four histological types and three response categories, the source reports Pearson statistic 75.89 and likelihood-ratio statistic 68.30, each with

$$(4-1)(3-1) = 6$$

degrees of freedom. Both p-values are essentially 0, giving strong evidence of association between response and histological type.

Definition 1.10.14. *A Discrete Dependence Strength*
For discrete variables Y and Z, define

$$\delta(Y,Z) := \max_{A,B}\left|\mathbb{P}(Y\in A, Z\in B) - \mathbb{P}(Y\in A)\mathbb{P}(Z\in B)\right|,$$

where the maximum is over all pairs of events A, B. For a finite table,

$$\delta(Y,Z) = \frac{1}{2}\sum_{i=1}^{I}\sum_{j=1}^{J}\left|p_{ij} - p_{i\cdot}p_{\cdot j}\right|,$$

and the plug-in estimate replaces $p_{ij}, p_{i\cdot}, p_{\cdot j}$ by their empirical frequencies. Then $\delta = 0$ iff Y and Z are independent.

Method 1.10.15. *Bootstrap Interval for δ*
To estimate uncertainty in $\widehat{\delta}$:

(i) Draw a bootstrap table $X^* \sim \text{Multinomial}(n, \widehat{p})$.
(ii) Compute the bootstrap empirical probabilities and $\widehat{\delta}^*$.
(iii) Repeat many times.
(iv) Use a bootstrap interval method for the resulting distribution.

The Wald method is unreliable near independence because $\delta = 0$ lies on the boundary of the parameter space.

Method 1.10.16. *Continuous Association: Correlation and Spearman Rank Correlation*
For two continuous variables, the plug-in estimator of correlation is the sample correlation

$$\widehat{\rho} = \frac{\sum_{i=1}^{n}(X_{1i}-\overline{X}_1)(X_{2i}-\overline{X}_2)}{\sqrt{\sum_i (X_{1i}-\overline{X}_1)^2\sum_i (X_{2i}-\overline{X}_2)^2}}.$$

Large-sample intervals and tests can use bootstrap standard errors, while normal-model intervals can use Method 1.10.1. For small samples or monotone nonlinear dependence, replace each variable by its ranks and compute the correlation of the ranks; this is Spearman's rank correlation $\widehat{\rho}_S$. A null distribution for $\widehat{\rho}_S$ can be simulated by fixing one rank ordering and randomly permuting the other.

Example 1.10.17. *Blood Pressure Correlation*
In the source blood-pressure example, the sample correlation between systolic and diastolic pressure is reported as $\widehat{\rho} = .88$. A bootstrap standard error is .046, giving a very large Wald statistic and a pivotal bootstrap 95% interval about $(.78, .94)$. The rank correlation is reported as $\widehat{\rho}_S = .94$, with a simulation-based null p-value near 0.

Theorem 1.10.18. *One Discrete and One Continuous Variable*
Let $Y \in \{1,\dots,I\}$ be discrete and let Z be continuous. Write

$$F_i(z) = \mathbb{P}(Z \le z \mid Y = i).$$

Then $Y \perp\!\!\!\perp Z$ if and only if

$$F_1 = \cdots = F_I.$$

For $I = 2$, the two-sample Kolmogorov–Smirnov statistic is

$$D = \sup_z \left| \widehat{F}_1(z) - \widehat{F}_2(z) \right|.$$

If

$$H(t) = 1 - 2\sum_{j=1}^{\infty} (-1)^{j-1} e^{-2j^2 t^2},$$

then under $F_1 = F_2$,

$$\lim_{n\to\infty} \mathbb{P}\left(\sqrt{\frac{n_1 n_2}{n_1 + n_2}}\, D \le t \right) = H(t).$$

Thus an approximate level α test rejects when

$$\sqrt{\frac{n_1 n_2}{n_1 + n_2}}\, D > H^{-1}(1 - \alpha).$$

Remark 1.10.19. *Source Data Exercises Left Open*
Several source exercises request analyses using external datasets for car mileage, Montana age and financial status, United States temperatures, and calcium intake. Those datasets are not present in the attached screenshots, so this pass records the methods but does not invent numerical conclusions for those exercises.

1.11 Undirected Graphs and Log-Linear Models

Graphical models give a compact language for multivariate independence constraints. For k binary variables the unconstrained joint law has $2^k - 1$ free cell probabilities, so sparsity appears quickly. An undirected graph trades some of that large multinomial surface for conditional-independence structure, and log-linear models give the corresponding algebra for discrete tables.

Definition 1.11.1. *Conditional Independence*
Let X, Y, Z be random vectors. We write

$$X \perp\!\!\!\perp Y \mid Z$$

when, after conditioning on Z, knowing Y gives no further information about X. In the discrete case this means

$$\mathbb{P}(X = x, Y = y \mid Z = z) = \mathbb{P}(X = x \mid Z = z)\mathbb{P}(Y = y \mid Z = z)$$

for all relevant x, y, z. In the continuous case the corresponding density statement is

$$f_{X,Y|Z}(x, y \mid z) = f_{X|Z}(x \mid z) f_{Y|Z}(y \mid z).$$

This is the conditional version of the factorization tests in Proposition 1.1.30; the conditional mass and density notation is as in Definition 1.3.16.

Proposition 1.11.2. *Conditional Factorization*
Let (X_a, X_b, X_c) have a positive joint mass function or density. Then

$$X_b \perp\!\!\!\perp X_c \mid X_a$$

if and only if the joint law can be written in the form

$$f(x_a, x_b, x_c) = g(x_a, x_b) h(x_a, x_c)$$

for nonnegative functions g, h.

Proof. If conditional independence holds, then

$$f(x_a, x_b, x_c) = f_{b|a}(x_b \mid x_a) f_{c|a}(x_c \mid x_a) f_a(x_a),$$

which has the displayed form after absorbing the factor f_a into either side. Conversely, a factorization by the pairs (a, b) and (a, c) normalizes inside each fixed value of x_a to give a product conditional law. □

Theorem 1.11.3. *Basic Conditional-Independence Rules*
The following implications hold for random vectors of compatible dimensions:

(i) $X \perp\!\!\!\perp Y \mid Z \Rightarrow Y \perp\!\!\!\perp X \mid Z$.
(ii) If $X \perp\!\!\!\perp Y \mid Z$ *and* $U = h(X)$, *then* $U \perp\!\!\!\perp Y \mid Z$.
(iii) If $X \perp\!\!\!\perp Y \mid Z$ *and* $U = h(X)$, *then* $X \perp\!\!\!\perp Y \mid (Z, U)$.
(iv) If $X \perp\!\!\!\perp Y \mid Z$ *and* $X \perp\!\!\!\perp W \mid (Y, Z)$, *then* $X \perp\!\!\!\perp (Y, W) \mid Z$.
(v) If all relevant probabilities are positive, then

$$X \perp\!\!\!\perp Y \mid Z \quad \text{and} \quad X \perp\!\!\!\perp Z \mid Y \quad \Rightarrow \quad X \perp\!\!\!\perp (Y, Z).$$

The last implication is the intersection rule; positivity prevents zero-probability conditioning events from hiding dependence.

Definition 1.11.4. *Undirected Separation and Markov Properties*
An *undirected graph* $\mathcal{G} = (V, E)$ has a finite vertex set V and unordered edges E. Vertices X and Y are *adjacent*, written $X \sim Y$, if $\{X, Y\} \in E$. A *path* is a sequence $X_0, \dots, X_n$ with $X_{i-1} \sim X_i$ for every i, and a graph is *complete* if every pair of distinct vertices is adjacent.

For disjoint subsets $A, B, C \subseteq V$, we say that C *separates* A and B if every path from a vertex in A to a vertex in B intersects C. A distribution on the variables V has the *pairwise Markov property* for $\mathcal{G}$ if every missing edge encodes

$$X \perp\!\!\!\perp Y \mid V \setminus \{X, Y\}.$$

It has the *global Markov property* if

$$C \text{ separates } A \text{ and } B \quad \Rightarrow \quad X_A \perp\!\!\!\perp X_B \mid X_C.$$

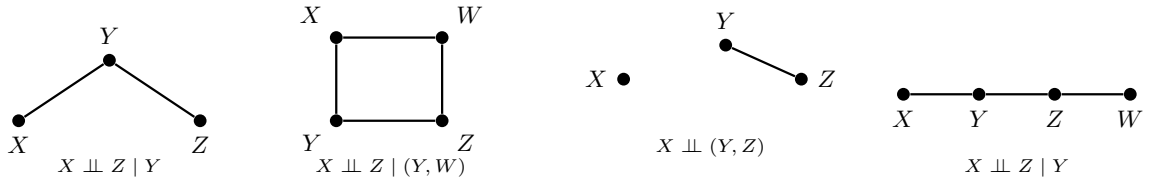


Figure 222: Small undirected graphs used to read conditional-independence statements from separation.

Theorem 1.11.5. *Pairwise and Global Markov Equivalence*
For a positive discrete distribution on a finite vertex set V, the pairwise and global Markov properties for an undirected graph $\mathcal{G}$ are equivalent.

Proof. The global property implies the pairwise property because nonadjacent vertices are separated by all remaining vertices. The converse is obtained by repeatedly applying the conditional-independence rules in Theorem 1.11.3: missing-edge statements are combined along separated paths, while the positivity assumption supplies the intersection step needed to remove variables from conditioning sets. □

Example 1.11.6. *Reading a Graph*
Figure 222 illustrates the main patterns.

(i) In the path $X - Y - Z$, the vertex Y separates X and Z, so $X \perp\!\!\!\perp Z \mid Y$.
(ii) In a complete triangle on X, Y, Z, there are no missing edges and no nontrivial conditional-independence statements forced by the graph.
(iii) In the four-cycle $X - W - Z - Y - X$, opposite vertices are separated by the two remaining vertices:
$$X \perp\!\!\!\perp Z \mid (Y, W), \qquad Y \perp\!\!\!\perp W \mid (X, Z).$$
(iv) In the chain $X - Y - Z - W$, pairwise Markov deletion gives, for instance, $X \perp\!\!\!\perp Z \mid (Y, W)$. The global Markov property sharpens this to $X \perp\!\!\!\perp Z \mid Y$; similarly, $Y \perp\!\!\!\perp W \mid Z$.
(v) If two components are disconnected, the empty set separates them, so the corresponding variables are marginally independent.

Method 1.11.7. *Drawing a Pairwise Markov Graph*
For a finite collection of variables, start with the complete graph and delete the edge $\{X_i, X_j\}$ exactly when the model asserts

$$X_i \perp\!\!\!\perp X_j \mid V \setminus \{X_i, X_j\}.$$

For three variables this gives the following source-exercise answers:

(a) $X_1 \perp\!\!\!\perp X_3 \mid X_2$ gives the path $X_1 - X_2 - X_3$.
(b) $X_1 \perp\!\!\!\perp X_2 \mid X_3$ and $X_1 \perp\!\!\!\perp X_3 \mid X_2$ leave only the edge $X_2 - X_3$, with X_1 isolated.
(c) All three pairwise conditional independences give the empty graph.

Example 1.11.8. *Minimal Separators*
A conditional independence $X_i \perp\!\!\!\perp X_j \mid C$ read from a graph is *minimal* if no variable can be removed from C while preserving separation. The exercise figures in the source give these useful patterns.

(i) For the star centered at X_2:
$$X_1 \perp\!\!\!\perp X_3 \mid X_2, \qquad X_1 \perp\!\!\!\perp X_4 \mid X_2, \qquad X_3 \perp\!\!\!\perp X_4 \mid X_2.$$
(ii) For the path $X_1 - X_2 - X_3 - X_4$:
$$X_1 \perp\!\!\!\perp X_3 \mid X_2, \qquad X_2 \perp\!\!\!\perp X_4 \mid X_3,$$
$$X_1 \perp\!\!\!\perp X_4 \mid X_2, \qquad X_1 \perp\!\!\!\perp X_4 \mid X_3.$$

(iii) For the four-cycle:

$$X_1 \perp\!\!\!\perp X_3 \mid (X_2, X_4), \qquad X_2 \perp\!\!\!\perp X_4 \mid (X_1, X_3).$$

For the six-vertex triangular graph in the source, the nonadjacent pairs and minimal separators are

$$\begin{array}{l|l} \{X_1, X_4\} & \{X_2, X_3\},\ \{X_2, X_5\} \\ \{X_1, X_5\} & \{X_2, X_3\} \\ \{X_1, X_6\} & \{X_2, X_3\},\ \{X_3, X_5\} \\ \{X_2, X_6\} & \{X_3, X_5\} \\ \{X_3, X_4\} & \{X_2, X_5\} \\ \{X_4, X_6\} & \{X_2, X_5\},\ \{X_3, X_5\}. \end{array}$$

Definition 1.11.9. *Baseline Log-Linear Expansion*
Let $X = (X_1, \dots, X_m)$ be a positive discrete random vector, let $S = \{1, \dots, m\}$, and write $x_A = (x_j : j \in A)$. After choosing a baseline level 0 for each coordinate, the joint probability function can be written as

$$\log f(x) = \sum_{A \subseteq S} \psi_A(x),$$

where $\psi_\emptyset$ is constant, ψ_A depends only on x_A, and $\psi_A(x) = 0$ whenever some coordinate in A is at its baseline level. The functions ψ_A are the *log-linear terms*; terms with $|A| \geq 2$ are interactions.

Example 1.11.10. *Small Log-Linear Expansions*
For $X \sim \text{Bernoulli}(p)$ with $p_1 = p$ and $p_0 = 1 - p$,

$$\log f(x) = \psi_\emptyset(x) + \psi_1(x), \qquad \psi_\emptyset(x) = \log p_0, \qquad \psi_1(x) = x \log\left(\frac{p_1}{p_0}\right).$$

For a 2-by-3 table with cells p_{ij}, $i \in \{0, 1\}$ and $j \in \{0, 1, 2\}$, one baseline expansion is

$$\log f(x_1, x_2) = \psi_\emptyset + \psi_1 + \psi_2 + \psi_{12},$$

with

$$\begin{aligned} \psi_\emptyset &= \log p_{00}, \\ \psi_1 &= x_1 \log\left(\frac{p_{10}}{p_{00}}\right), \\ \psi_2 &= \mathbf{1}_{\{x_2=1\}} \log\left(\frac{p_{01}}{p_{00}}\right) + \mathbf{1}_{\{x_2=2\}} \log\left(\frac{p_{02}}{p_{00}}\right), \\ \psi_{12} &= \mathbf{1}_{\{x_1=1,x_2=1\}} \log\left(\frac{p_{11}p_{00}}{p_{01}p_{10}}\right) + \mathbf{1}_{\{x_1=1,x_2=2\}} \log\left(\frac{p_{12}p_{00}}{p_{02}p_{10}}\right). \end{aligned}$$

Thus the interaction terms are log odds ratios against the baseline cell.

Proposition 1.11.11. *Conditional Independence by Missing Interactions*
Let (a, b, c) be a partition of $\{1, \dots, m\}$. In a positive discrete log-linear model,

$$X_b \perp\!\!\!\perp X_c \mid X_a$$

if and only if every log-linear term ψ_A with at least one index in b and at least one index in c is identically zero.

Proof. If all such cross terms vanish, then the log-density splits as

$$\log f(x) = \sum_{A \subseteq a \cup b} \psi_A(x) + \sum_{A \subseteq a \cup c} \psi_A(x) - \sum_{A \subseteq a} \psi_A(x),$$

so $f(x_a, x_b, x_c) = g(x_a, x_b)h(x_a, x_c)$. Proposition 1.11.2 gives the conditional independence. Reversing the argument gives the converse. □

Definition 1.11.12. *Graphical Log-Linear Model*
Let $\mathcal{G}$ be an undirected graph on $\{1, \dots, m\}$. A subset A is a *clique* if every pair of distinct vertices in A is adjacent. The graphical log-linear model for $\mathcal{G}$ is the model in which the structurally allowed log-linear terms are exactly the clique terms:

$$\psi_A \text{ may be present} \iff A \text{ is a clique of } \mathcal{G}.$$

Equivalently, if A contains a nonedge of $\mathcal{G}$, then ψ_A is structurally omitted.

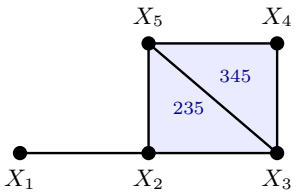


Figure 223: A graphical log-linear model whose largest clique terms are ψ_{235} and ψ_{345}.

Example 1.11.13. *Graphical Versus Extra Constraints*
For the graph in Figure 223, the graphical model allows main effects, all edge interactions, and the two three-way clique interactions ψ_{235} and ψ_{345}. Any term containing a nonedge, such as ψ_{15} or ψ_{1245}, is omitted. If we also omit ψ_{235} and ψ_{345}, the graph of pairwise conditional independences is unchanged, but the model is no longer the graphical model for that graph; it has additional non-independence constraints.

Definition 1.11.14. *Hierarchical Log-Linear Models*
A log-linear model is *hierarchical* if inclusion of a higher-order term forces inclusion of all lower-order terms contained in it. Equivalently, if ψ_A is structurally omitted and $A \subseteq B$, then ψ_B is structurally omitted.

Proposition 1.11.15. *Graphical Models Are Hierarchical*
Every graphical log-linear model is hierarchical, but not every hierarchical log-linear model is graphical.

Proof. If B is a clique, then every $A \subseteq B$ is also a clique, so graphical models include lower-order terms automatically. For the converse, the three-variable model with all main effects and all two-way interactions but no ψ_{123} is hierarchical. Its pairwise graph is complete, and the graphical model for a complete graph is saturated, so this hierarchical model is not graphical. □

Example 1.11.16. *Three Small Models*
The model

$$\log f = \psi_\emptyset + \psi_1 + \psi_2 + \psi_3 + \psi_{12} + \psi_{13}$$

is graphical and hierarchical; its graph is the path $X_2 - X_1 - X_3$. The model

$$\log f = \psi_\emptyset + \psi_1 + \psi_2 + \psi_3 + \psi_{12} + \psi_{13} + \psi_{23}$$

is hierarchical but not graphical, since the complete graph would also allow ψ_{123}. The model

$$\log f = \psi_\emptyset + \psi_3 + \psi_{12}$$

is not hierarchical because it includes ψ_{12} without the corresponding main effects.

Definition 1.11.17. *Model Generators*
A hierarchical log-linear model can be specified by its maximal included terms, called *generators*. For example,

$$M = 12 + 13$$

means that ψ_{12} and ψ_{13} are included, together with all lower-order terms:

$$\log f = \psi_\emptyset + \psi_1 + \psi_2 + \psi_3 + \psi_{12} + \psi_{13}.$$

The generator $M = 123$ is the saturated three-variable model, while $M = 1 + 2 + 3$ is mutual independence.

saturated (graphical)		123	
two-way		12 + 13 + 23	
graphical	12 + 13	12 + 23	13 + 23
graphical	12 + 3	13 + 2	23 + 1
mutual independence (graphical)		1 + 2 + 3	
conditional uniform	12	13	23
independence	1 + 2	1 + 3	2 + 3
	1	2	3
uniform		∅	

Figure 224: A compact lattice of hierarchical models for three variables, ordered by inclusion of log-linear terms.

Method 1.11.18. *Fitting Log-Linear Models*
For table counts x_j and model cell probabilities $p_j(\beta)$, the log-likelihood is

$$\ell(\beta) = \sum_j x_j \log p_j(\beta).$$

The saturated model has every log-linear term and reproduces the empirical cell proportions. A submodel M is commonly compared with the saturated model using the deviance

$$\operatorname{dev}(M) = 2\left(\widehat{\ell}_{\text{sat}} - \widehat{\ell}_M\right).$$

Under the null model M, this is approximately χ^2_ν, where ν is the difference between the number of parameters in the saturated model and the number in M. Penalized likelihood instead scores

$$\operatorname{AIC}(M) = -2\left(\widehat{\ell}(M) - |M|\right) = -2\widehat{\ell}(M) + 2\,|M|\,,$$

where $|M|$ is the number of fitted parameters; smaller AIC is preferred.

Remark 1.11.19. *AIC and KL Risk*
For a candidate density g, the discrete Kullback–Leibler loss from the true law f is

$$D(f, g) = \sum_x f(x) \log\left(\frac{f(x)}{g(x)}\right).$$

Since $\sum_x f(x) \log f(x)$ does not depend on g, minimizing KL risk is equivalent to maximizing the expected log score $\sum_x f(x) \log g(x)$. AIC corrects the fitted log-likelihood by a parameter-count penalty to reduce the optimism from evaluating fit on the same data used to estimate the parameters.

Example 1.11.20. *Breast-Cancer Table and Selected Graph*
The source table records diagnostic center X_1, nuclear grade X_2, and survival X_3:

	malignant		benign	
	died	survived	died	survived
Boston	35	59	47	112
Glamorgan	42	77	26	76.

The unconstrained multinomial MLE is the vector of cell counts divided by 474. For a benign tumor at the Glamorgan clinic, the estimated death probability is

$$\frac{26}{26+76} \approx .255, \qquad \widehat{\text{se}} \approx \sqrt{\frac{.255(1-.255)}{102}} \approx .043.$$

The source log-linear analysis selects a model with center–grade and grade–survival interactions, but no center–survival edge and no three-way interaction. Thus the selected graph is

$$\text{Center} - \text{Grade} - \text{Survival},$$

so the model encodes

$$X_1 \perp\!\!\!\perp X_3 \mid X_2.$$

The reported deviance is .6 with $8 - 6 = 2$ degrees of freedom, giving a reported p-value about .74 and no evidence of poor fit.

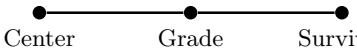


Figure 225: Selected graph for the breast-cancer log-linear example.

Example 1.11.21. *Useful Exercise Solutions*
Several source exercises are best kept as reusable checks.

(i) In the 2-by-3 expansion of Example 1.11.10, if

$$\beta_1 = \log p_{00}, \qquad \beta_2 = \log\left(\frac{p_{10}}{p_{00}}\right), \qquad \beta_3 = \log\left(\frac{p_{01}}{p_{00}}\right),$$
$$\beta_4 = \log\left(\frac{p_{02}}{p_{00}}\right), \quad \beta_5 = \log\left(\frac{p_{11}p_{00}}{p_{01}p_{10}}\right), \quad \beta_6 = \log\left(\frac{p_{12}p_{00}}{p_{02}p_{10}}\right),$$

then the inverse map is proportional to

$$p_{00} = e^{\beta_1}, \quad p_{10} = e^{\beta_1+\beta_2}, \quad p_{01} = e^{\beta_1+\beta_3},$$
$$p_{02} = e^{\beta_1+\beta_4}, \quad p_{11} = e^{\beta_1+\beta_2+\beta_3+\beta_5}, \quad p_{12} = e^{\beta_1+\beta_2+\beta_4+\beta_6},$$

with a final normalization so that the six probabilities sum to 1.

(ii) The model

$$\log f = \psi_\emptyset + \psi_{12} + \psi_{13} + \psi_{24} + \psi_{34}$$

has pairwise graph with edges $12, 13, 24, 34$. Its separation statements include

$$X_1 \perp\!\!\!\perp X_4 \mid (X_2, X_3), \qquad X_2 \perp\!\!\!\perp X_3 \mid (X_1, X_4).$$

As written, however, it is neither hierarchical nor graphical because it contains two-way interactions without their main effects.

(iii) If the binary cell probabilities for (X_1, X_2, X_3) are proportional to

	00	01	10	11
$x_1 = 0$	2	8	4	16
$x_1 = 1$	16	128	32	256,

where the columns are (x_2, x_3), then up to the normalizing constant the nonzero baseline terms are

$$\psi_1 = \log 8, \qquad \psi_2 = \log 2, \qquad \psi_3 = \log 4, \qquad \psi_{13} = \log 2.$$

The terms ψ_{12}, ψ_{23}, and ψ_{123} vanish, so the model is graphical and hierarchical with the single edge $1 - 3$ and isolated vertex 2.

(iv) For four binary variables, main effects only give the empty graph and a graphical hierarchical model. Main effects plus $\psi_{23}, \psi_{24}, \psi_{34}, \psi_{234}$ give the graphical hierarchical model with a triangle on $\{2, 3, 4\}$ and 1 isolated. Main effects plus $\psi_{23}, \psi_{34}, \psi_{14}, \psi_{12}$ give the graphical hierarchical four-cycle $1 - 2 - 3 - 4 - 1$. A model with only a four-way interaction has complete pairwise graph but is neither hierarchical nor graphical.

Remark 1.11.22. *Source Limits*
The attached screenshots support the definitions, Markov-separation rules, log-linear expansions, breast-cancer table, and selected-model summaries above. The synthetic five-variable simulation is recorded only at the level of its setup and reported average risks, so this pass keeps its qualitative message: AIC and BIC can improve on the saturated multinomial MLE in sparse high-dimensional discrete tables.

1.12 Counterfactual Causal Inference

The preceding association tests decide whether variables move together. Causal inference asks a stronger question: what would happen to the distribution of an outcome if a treatment were externally changed? The key additional object is the unobserved response a subject would have had under a treatment value not actually received.

Definition 1.12.1. *Binary Potential Outcomes and Consistency*
Let $X \in \{0, 1\}$ be a binary treatment, with $X = 1$ denoting treatment and $X = 0$ denoting no treatment. For each subject introduce potential outcomes

$$C_0, \qquad C_1,$$

where C_0 is the outcome under $X = 0$ and C_1 is the outcome under $X = 1$. The observed response satisfies the *consistency relation*

$$Y = \begin{cases} C_0, & X = 0, \\ C_1, & X = 1, \end{cases} \qquad \text{equivalently} \qquad Y = C_X.$$

Only one potential outcome is observed for each subject; the other is counterfactual.

Example 1.12.2. *Binary Response Types*
When $C_0, C_1 \in \{0, 1\}$ there are four hidden subject types:

type	C_0	C_1
survivor	1	1
responder	0	1
anti-responder	1	0
doomed	0	0

Observed data reveal C_0 only for untreated subjects and C_1 only for treated subjects.

Definition 1.12.3. *Average Causal Effect and Association*
The *average causal effect*, or *average treatment effect*, is

$$\theta := \mathbb{E}(C_1) - \mathbb{E}(C_0).$$

The observed *association* between treatment and outcome is

$$\alpha := \mathbb{E}(Y \mid X = 1) - \mathbb{E}(Y \mid X = 0).$$

For binary potential outcomes, other causal summaries include the causal odds ratio

$$\frac{\mathbb{P}(C_1 = 1)}{\mathbb{P}(C_1 = 0)} \Big/ \frac{\mathbb{P}(C_0 = 1)}{\mathbb{P}(C_0 = 0)}$$

and the causal relative risk $\mathbb{P}(C_1 = 1)/\mathbb{P}(C_0 = 1)$, when the denominators are positive.

Proposition 1.12.4. *Association Need Not Equal Causation*
In general,

$$\theta \neq \alpha.$$

Proof. By consistency,

$$\mathbb{E}(Y \mid X = 1) = \mathbb{E}(C_1 \mid X = 1), \qquad \mathbb{E}(Y \mid X = 0) = \mathbb{E}(C_0 \mid X = 0).$$

The causal effect compares population means $\mathbb{E}C_1$ and $\mathbb{E}C_0$; association compares means in the subpopulations that did and did not receive treatment. □

Example 1.12.5. *No Effect but Positive Association*
Suppose eight subjects have

i	1	2	3	4	5	6	7	8
X_i	0	0	0	0	1	1	1	1
C_{0i}	0	0	0	0	1	1	1	1
C_{1i}	0	0	0	0	1	1	1	1.

Then $C_{0i} = C_{1i}$ for every subject, so $\theta = 0$. But all treated subjects have $Y = 1$ and all untreated subjects have $Y = 0$, so $\alpha = 1$. The association is created by selection into treatment.

Example 1.12.6. *Opposite Signs*
The source asks for an example with $\alpha > 0$ but $\theta < 0$. One construction is

i	1	2	3	4	5	6
X_i	1	1	0	0	0	0
C_{0i}	1	1	1	1	1	0
C_{1i}	1	1	0	0	0	0.

The treated subjects both have observed outcome 1, while untreated subjects have observed average $3/4$, so $\alpha = 1/4 > 0$. However

$$\theta = \frac{1+1+0+0+0+0}{6} - \frac{1+1+1+1+1+0}{6} = -\frac{1}{2}.$$

The observed association has the wrong causal sign because the anti-responders were left untreated.

Theorem 1.12.7. *Random Assignment Identifies the Average Treatment Effect*
Suppose $\mathbb{P}(X = 0) > 0$, $\mathbb{P}(X = 1) > 0$, *and*

$$X \perp\!\!\!\perp (C_0, C_1).$$

Then $\alpha = \theta$. A consistent estimator is

$$\widehat{\theta} = \overline{Y}_1 - \overline{Y}_0 = \frac{1}{n_1} \sum_{i:X_i=1} Y_i - \frac{1}{n_0} \sum_{i:X_i=0} Y_i,$$

where $n_1 = \sum_i X_i$ and $n_0 = \sum_i (1 - X_i)$.

Proof. Random assignment gives $\mathbb{E}(C_1) = \mathbb{E}(C_1 \mid X = 1)$ and $\mathbb{E}(C_0) = \mathbb{E}(C_0 \mid X = 0)$. Consistency then gives

$$\theta = \mathbb{E}(Y \mid X = 1) - \mathbb{E}(Y \mid X = 0) = \alpha.$$

The estimator is the sample analogue, justified by the law of large numbers within the two treatment groups. □

Definition 1.12.8. *Counterfactual Process and Causal Regression*
For a non-binary treatment $X \in \mathcal{X}$, replace (C_0, C_1) by a *counterfactual process*

$$\{C(x) : x \in \mathcal{X}\}.$$

Here $C(x)$ is the outcome the subject would have under dose x, and consistency becomes $Y = C(X)$. The *causal regression function* is

$$\theta(x) := \mathbb{E}(C(x)),$$

whereas the ordinary regression function is

$$r(x) := \mathbb{E}(Y \mid X = x).$$

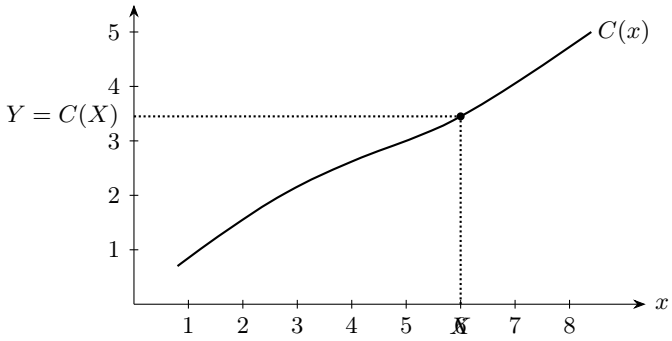


Figure 226: A counterfactual process $C(x)$ records every dose response for one subject, while only $Y = C(X)$ is observed.

Theorem 1.12.9. *Random Assignment Identifies the Causal Regression*
In general $\theta(x) \neq r(x)$. If X is independent of the process $\{C(t) : t \in \mathcal{X}\}$, then for any x with $\mathbb{P}(X = x) > 0$,

$$\theta(x) = r(x).$$

Proof. Under independence and consistency,

$$\theta(x) = \mathbb{E}C(x) = \mathbb{E}(C(x) \mid X = x) = \mathbb{E}(Y \mid X = x) = r(x).$$

□

Example 1.12.10. *No Causal Effect but Nonconstant Regression*
Suppose four subjects have constant response curves

$$C_1(x) = 4, \qquad C_2(x) = 3, \qquad C_3(x) = 2, \qquad C_4(x) = 1.$$

Changing a dose changes no subject's outcome, so

$$\theta(x) = \frac{4 + 3 + 2 + 1}{4} = \frac{5}{2}$$

is constant. If the high-response subjects tend to receive high doses and the low-response subjects tend to receive low doses, the observed regression $r(x)$ can still increase with x.

Definition 1.12.11. *No Unmeasured Confounding*
In an observational study, treatment is not randomly assigned. Let Z denote observed confounders. The *no unmeasured confounding* condition is

$$\{C(t) : t \in \mathcal{X}\} \perp\!\!\!\perp X \mid Z.$$

It says that, within levels of Z, treatment assignment is independent of the counterfactual response process.

Theorem 1.12.12. *Adjustment Formula*
Assume no unmeasured confounding and consistency. Then

$$\theta(x) = \int \mathbb{E}(Y \mid X = x, Z = z)\, dF_Z(z).$$

Consequently, if $\widehat{r}(x,z)$ estimates $r(x,z) = \mathbb{E}(Y \mid X = x, Z = z)$, a plug-in adjustment estimator is

$$\widehat{\theta}(x) = \frac{1}{n}\sum_{i=1}^{n} \widehat{r}(x, Z_i).$$

In the linear case $r(x,z) = \beta_0 + \beta_1 x + \beta_2 z$,

$$\widehat{\theta}(x) = \widehat{\beta}_0 + \widehat{\beta}_1 x + \widehat{\beta}_2 \overline{Z}_n.$$

Proof. Conditioning on $Z = z$,

$$\mathbb{E}(C(x) \mid Z = z) = \mathbb{E}(C(x) \mid X = x, Z = z) = \mathbb{E}(Y \mid X = x, Z = z),$$

where the first equality uses no unmeasured confounding and the second uses consistency. Integrate over the marginal law of Z and use the tower property from Theorem 1.3.21. □

Remark 1.12.13. *Adjusted Versus Observed Regression*
Adjustment averages $\mathbb{E}(Y \mid X = x, Z = z)$ over the marginal law F_Z. The observational regression averages over the conditional law $F_{Z|X=x}$:

$$\mathbb{E}(Y \mid X = x) = \int \mathbb{E}(Y \mid X = x, Z = z)\, dF_{Z|X=x}(z).$$

The difference between these averaging measures is the algebraic trace of confounding.

Example 1.12.14. *Simpson's Paradox as a Translation Error*
It can happen that

$$\mathbb{P}(Y = 1 \mid X = 1) < \mathbb{P}(Y = 1 \mid X = 0),$$

while within each stratum $Z = z$,

$$\mathbb{P}(Y = 1 \mid X = 1, Z = z) > \mathbb{P}(Y = 1 \mid X = 0, Z = z).$$

The apparent paradox comes from translating association statements into causal statements. For binary outcomes, "treatment is harmful" means

$$\mathbb{P}(C_1 = 1) < \mathbb{P}(C_0 = 1),$$

not the marginal association inequality above.

Proposition 1.12.15. *No Causal Simpson Reversal Under the Same Stratification*
If Z is discrete and

$$\mathbb{P}(C_1 = 1 \mid Z = z) > \mathbb{P}(C_0 = 1 \mid Z = z)$$

for every z with $\mathbb{P}(Z = z) > 0$, then

$$\mathbb{P}(C_1 = 1) > \mathbb{P}(C_0 = 1).$$

Proof. Use the law of total probability:

$$\mathbb{P}(C_1 = 1) = \sum_z \mathbb{P}(C_1 = 1 \mid Z = z)\mathbb{P}(Z = z) > \sum_z \mathbb{P}(C_0 = 1 \mid Z = z)\mathbb{P}(Z = z) = \mathbb{P}(C_0 = 1).$$

□

Proposition 1.12.16. *Observable Bounds for a Binary Outcome*
Suppose $X, Y, C_0, C_1 \in \{0,1\}$ and only the law of (X,Y) is observed. Let

$$p_x = \mathbb{P}(X = x), \qquad m_x = \mathbb{E}(Y \mid X = x),$$

with $p_0, p_1 > 0$. Then

$$p_1 m_1 - p_0 m_0 - p_1 \le \theta \le p_1 m_1 - p_0 m_0 + p_0,$$

and the interval has width 1.

Proof. By consistency,

$$\mathbb{E}C_1 = p_1 m_1 + p_0\, \mathbb{E}(C_1 \mid X = 0), \qquad \mathbb{E}C_0 = p_0 m_0 + p_1\, \mathbb{E}(C_0 \mid X = 1).$$

The two unobserved conditional means lie in $[0,1]$. Taking their extreme values gives the bounds, whose difference is $p_0 + p_1 = 1$. □

Example 1.12.17. *Positive Individual Slopes, Decreasing Regression*
Let $X \sim \text{Uniform}(1,2)$ and set $\beta = 1/X^2$. Each subject has potential response curve $C(t) = \beta t$, so every individual slope is positive. But

$$Y = C(X) = \frac{1}{X}, \qquad r(x) = \mathbb{E}(Y \mid X = x) = \frac{1}{x},$$

which is decreasing. High-dose subjects have systematically smaller slopes, so the cross-sectional regression reverses the individual causal direction.

Remark 1.12.18. *Evidence and Exercises*
Observational adjustment is not proof that all confounding has been removed. The source emphasizes replication, plausible adjustment variables, and external scientific mechanism as supporting evidence. Exercise content with short source-supported answers was formalized above as Examples 1.12.6 and 1.12.17 and Proposition 1.12.16.

1.13 Directed Graphs and Interventions

Directed graphs encode conditional-independence structure with arrows. For causal use, arrows are not merely predictive: they specify how the data-generating process changes under intervention.

Definition 1.13.1. *Directed Graphs and DAGs*
A *directed graph* $\mathcal{G} = (V, E)$ consists of vertices V and directed edges $E \subseteq V \times V$. If $(X, Y) \in E$, write $X \to Y$ and call X a *parent* of Y and Y a *child* of X. The set of parents of X is denoted π_X or $\pi(X)$.
A *directed path* follows arrow directions. An *undirected path* ignores arrow directions. If there is a directed path from X to Y, then X is an *ancestor* of Y and Y is a *descendant* of X. A directed path that returns to its starting vertex is a *cycle*; a directed graph with no cycles is a *directed acyclic graph*, or *DAG*.

Definition 1.13.2. *Colliders*
On a three-node path, a node is a *collider* if the adjacent arrows both point into it:

$$X \to Y \leftarrow Z.$$

The other three basic configurations

$$X \to Y \to Z, \qquad X \leftarrow Y \leftarrow Z, \qquad X \leftarrow Y \to Z$$

are *non-colliders* at the middle node.

Definition 1.13.3. *Markov Factorization for a DAG*
Let $\mathcal{G}$ be a DAG with vertices $V = (X_1, \ldots, X_k)$. A distribution with mass or density p is *Markov to* $\mathcal{G}$ if

$$p(v) = \prod_{i=1}^{k} p(x_i \mid \pi_i),$$

where π_i denotes the values of the parents of X_i. The set of distributions Markov to $\mathcal{G}$ is denoted $\mathcal{M}(\mathcal{G})$.

Example 1.13.4. *A Four-Node Factorization*
For

$$X \to Z \leftarrow Y, \qquad Z \to W,$$

the Markov factorization is

$$p(x, y, z, w) = p(x)p(y)p(z \mid x, y)p(w \mid z).$$

It records $X \perp\!\!\!\perp Y$ and $W \perp\!\!\!\perp \{X, Y\} \mid Z$.

Theorem 1.13.5. *Markov Condition*
A distribution belongs to $\mathcal{M}(\mathcal{G})$ if and only if each variable W is conditionally independent of all its non-parent non-descendants given its parents:

$$W \perp\!\!\!\perp \widetilde{W} \mid \pi_W,$$

where $\widetilde{W}$ denotes the other variables except the parents and descendants of W.

Example 1.13.6. *Independences from a DAG*
For

$$A \to B \to D \to E, \qquad A \to C \to D,$$

the factorization is

$$p(a, b, c, d, e) = p(a)p(b \mid a)p(c \mid a)p(d \mid b, c)p(e \mid d).$$

The Markov condition implies

$$B \perp\!\!\!\perp C \mid A, \qquad D \perp\!\!\!\perp A \mid \{B, C\}, \qquad E \perp\!\!\!\perp \{A, B, C\} \mid D.$$

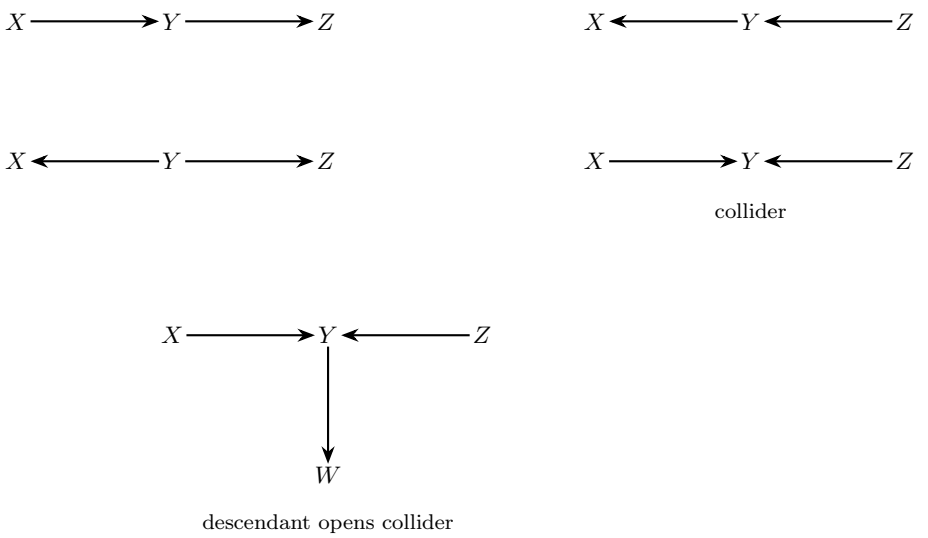


Figure 227: Non-collider paths are blocked by conditioning on the middle node; collider paths are opened by conditioning on the collider or one of its descendants.

Definition 1.13.7. *D-Separation*
Let X and Y be distinct vertices and let W be a set of vertices not containing either of them. An undirected path from X to Y is *active given W* if

(i) every collider on the path has a descendant in W, and
(ii) no non-collider on the path lies in W.

The vertices X and Y are *d-separated given W* if there is no active path between them given W; otherwise they are *d-connected*.

Method 1.13.8. *Rules of D-Separation*
For a path segment with endpoints X and Z and middle node Y:

(i) If Y is a non-collider, the path is active unless one conditions on Y.
(ii) If Y is a collider, the path is blocked unless one conditions on Y or on a descendant of Y.
(iii) Conditioning on a descendant of a collider opens the path just as conditioning on the collider does.

Example 1.13.9. *Three-Node Patterns*
In each non-collider pattern

$$X \to Y \to Z, \qquad X \leftarrow Y \leftarrow Z, \qquad X \leftarrow Y \to Z,$$

the endpoints X and Z are d-connected marginally and d-separated given Y. In the collider pattern

$$X \to Y \leftarrow Z,$$

the endpoints are d-separated marginally but d-connected given Y.

Theorem 1.13.10. *D-Separation and Conditional Independence*
If a distribution is Markov to a DAG $\mathcal{G}$, then d-separation in $\mathcal{G}$ implies the corresponding conditional independence statement. If the distribution is faithful to $\mathcal{G}$, then the conditional independences are exactly those implied by d-separation.

Example 1.13.11. *Opening a Collider Through Descendants*
In the DAG

$$X \to U \leftarrow V \to W \leftarrow Y, \qquad U \to S_1, \quad W \to S_2,$$

the variables X and Y are d-separated given the empty set. They become d-connected given $\{S_1, S_2\}$, because conditioning on descendants opens both colliders. They are d-separated again given $\{S_1, S_2, V\}$, because the non-collider V is conditioned on.

Definition 1.13.12. *Markov Equivalence*
Let $\mathcal{I}(\mathcal{G})$ be the set of conditional independence statements implied by a DAG $\mathcal{G}$. Two DAGs on the same variables are *Markov equivalent* if

$$\mathcal{I}(\mathcal{G}_1) = \mathcal{I}(\mathcal{G}_2).$$

The *skeleton* of $\mathcal{G}$ is the undirected graph obtained by replacing each arrow with an undirected edge.

Theorem 1.13.13. *Markov Equivalence Criterion*
Two DAGs are Markov equivalent if and only if they have the same skeleton and the same colliders.

Example 1.13.14. *Three Markov-Equivalent Paths*
The three non-collider DAGs on X, Y, Z with skeleton $X - Y - Z$,

$$X \to Y \to Z, \qquad X \leftarrow Y \leftarrow Z, \qquad X \leftarrow Y \to Z,$$

are Markov equivalent. The collider $X \to Y \leftarrow Z$ has the same skeleton but is not Markov equivalent to them.

Method 1.13.15. *Maximum Likelihood for a Discrete DAG*
Assume $V = (X_1, \dots, X_k)$ is discrete and the DAG structure is fixed. Regard the parent vector of X_i as a single discrete variable $\widetilde{X}_i$. For IID observations,

$$\widehat{p}(v) = \prod_{i=1}^{k} \widehat{p}(x_i \mid \widetilde{x}_i),$$

where

$$\widehat{p}(x_i \mid \widetilde{x}_i) = \frac{\#\{j : X_{ji} = x_i,\ \widetilde{X}_{ji} = \widetilde{x}_i\}}{\#\{j : \widetilde{X}_{ji} = \widetilde{x}_i\}}$$

whenever the denominator is positive.

Method 1.13.16. *Parametric Likelihood for a Continuous DAG*
For continuous or mixed variables, specify a parametric conditional density $p(x_j \mid \pi_j; \theta_j)$ for each node. For observations $V_i = (X_{i1}, \dots, X_{im})$,

$$\mathcal{L}(\theta) = \prod_{i=1}^{n} p(V_i; \theta) = \prod_{i=1}^{n} \prod_{j=1}^{m} p(X_{ij} \mid \pi_{ij}; \theta_j).$$

A fixed DAG reduces the multivariate likelihood into nodewise conditional models.

Remark 1.13.17. *Structure Learning Caution*
One may try to learn DAG structure from data by comparing likelihoods or information criteria, but reliable confidence about graph structure can require enormous sample sizes. These notes use DAGs primarily as a language for conditional independence and causal assumptions supplied from context.

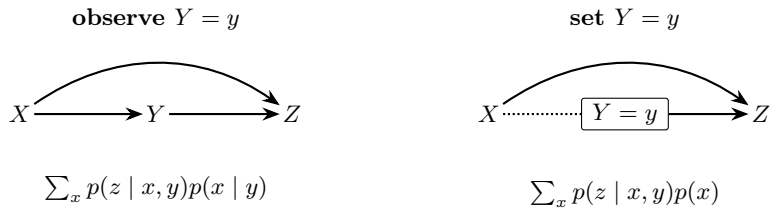


Figure 228: Conditioning on $Y = y$ leaves the original graph intact; intervening to set $Y = y$ removes arrows into Y.

Definition 1.13.18. *Passive and Active Conditioning*
For the DAG

$$X \to Y \to Z, \qquad X \to Z,$$

with factorization $p(x, y, z) = p(x)p(y \mid x)p(z \mid x, y)$, ordinary conditioning gives

$$\mathbb{P}(Z = z \mid Y = y) = \sum_x p(z \mid x, y)p(x \mid y).$$

Intervening to set $Y = y$ deletes the mechanism that generated Y and gives

$$\mathbb{P}(Z = z \mid Y := y) = \sum_x p(z \mid x, y)p(x).$$

The first is *passive conditioning*; the second is *active conditioning*.

Method 1.13.19. *Intervention by Truncated Factorization*
For a DAG factorization $p(v) = \prod_i p(x_i \mid \pi_i)$, the intervention "set $X_j = x$" is represented by:

(i) deleting all arrows into X_j in the graph;
(ii) deleting the factor $p(x_j \mid \pi_j)$ and fixing $X_j = x$ in the remaining factors.

The modified graph and distribution represent the intervened system.

Example 1.13.20. *Rain and Wet Lawn*
The graphs Rain → Wet Lawn and Rain ← Wet Lawn can both factor the same non-independent joint distribution. Causally, only the first is sensible: setting the lawn to be wet should not change the marginal probability of rain. Under the first graph, intervening on wetness leaves the rain distribution $p(r)$ unchanged; under the reversed causal graph, the same intervention would incorrectly use $p(r \mid w)$.

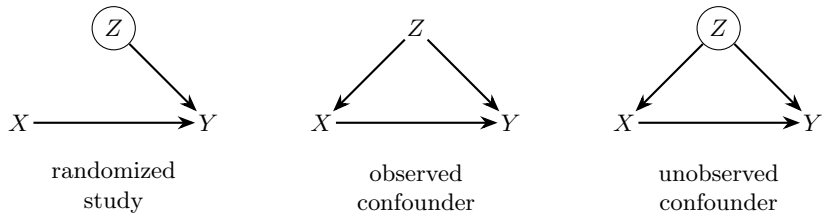


Figure 229: Randomization breaks the arrow from confounders into treatment. Observational adjustment requires measuring the confounders that affect both treatment and outcome.

Proposition 1.13.21. *DAG View of Confounding*
Suppose X is a treatment, Y is an outcome, and Z is a confounder with arrows into both X and Y.

(i) In a randomized study, the arrow $Z \to X$ is broken, and $\mathbb{E}(Y \mid X := x) = \mathbb{E}(Y \mid X = x)$ under the randomized graph.

(ii) In an observational study with all confounders observed,

$$\mathbb{E}(Y \mid X := x) = \int \mathbb{E}(Y \mid X = x, Z = z)\, dF_Z(z),$$

matching Theorem 1.12.12.

(iii) If the relevant Z is unobserved, the adjustment formula involves an unobserved variable and is not estimable from (X, Y) alone.

Remark 1.13.22. *Counterfactual Types as Confounders*
For binary X and Y, the potential-outcome type can be encoded by

$$Z = \begin{cases} 1, & (C_0, C_1) = (0, 0), \\ 2, & (C_0, C_1) = (0, 1), \\ 3, & (C_0, C_1) = (1, 0), \\ 4, & (C_0, C_1) = (1, 1). \end{cases}$$

This records the connection between the DAG confounder and the counterfactual pair (C_0, C_1).

Example 1.13.23. *Discrete Conditioning Versus Intervention*
Let

$$X \sim \text{Bernoulli}\left(\frac{1}{2}\right), \qquad Y \mid X = x \sim \text{Bernoulli}\left(\frac{e^{4x-2}}{1 + e^{4x-2}}\right),$$

and

$$Z \mid X = x, Y = y \sim \text{Bernoulli}\left(\frac{e^{2(x+y)-2}}{1 + e^{2(x+y)-2}}\right).$$

Write $\ell(t) = e^t/(1 + e^t)$. Then

$$\mathbb{P}(Z = 1 \mid Y = y) = \sum_{x=0}^{1} \ell(2(x + y) - 2)\mathbb{P}(X = x \mid Y = y),$$

whereas

$$\mathbb{P}(Z = 1 \mid Y := y) = \frac{1}{2} \sum_{x=0}^{1} \ell(2(x + y) - 2).$$

For $y = 1$, with $a = \ell(-2) = 1/(1 + e^2)$ and $b = \ell(2) = e^2/(1 + e^2)$,

$$\mathbb{P}(Z = 1 \mid Y = 1) = \frac{a}{2} + b^2, \qquad \mathbb{P}(Z = 1 \mid Y := 1) = \frac{1}{4} + \frac{b}{2}.$$

Example 1.13.24. *Gaussian Structural Equations*
Let

$$X \sim N(0, 1), \qquad Y \mid X = x \sim N(\alpha x, 1), \qquad Z \mid X = x, Y = y \sim N(\beta y + \gamma x, 1).$$

Equivalently, $Y = \alpha X + \varepsilon_Y$ and $Z = \beta Y + \gamma X + \varepsilon_Z$ with independent standard normal errors. Since

$$\mathbb{E}(X \mid Y = y) = \frac{\alpha}{1+\alpha^2} y,$$

ordinary conditioning gives

$$Z \mid Y = y \sim N\left(\left(\beta + \frac{\gamma\alpha}{1+\alpha^2}\right) y,\, 1 + \frac{\gamma^2}{1+\alpha^2}\right).$$

Intervening to set $Y = y$ instead leaves $X \sim N(0,1)$ and gives

$$Z \mid Y := y \sim N(\beta y, 1+\gamma^2).$$

Proposition 1.13.25. *Marginal Correlation Can Mislead*
In Example 1.13.24,

$$\mathrm{Cov}(Y, Z) = \beta(1+\alpha^2) + \alpha\gamma.$$

Thus marginal correlation can be nonzero when $\beta = 0$, due entirely to the unobserved common cause X. It can also vanish while $\beta \neq 0$, for instance when $\alpha \neq 0$ and $\gamma = -\beta(1+\alpha^2)/\alpha$.

Example 1.13.26. *Randomization Restores the Correlation Reading*
If Y is randomized independently of X, for example

$$X \sim N(0,1), \qquad Y \sim N(\alpha, 1), \qquad Z \mid X = x, Y = y \sim N(\beta y + \gamma x, 1),$$

then $\mathrm{Cov}(Y, Z) = \beta$, and

$$Z \mid Y := y \sim N(\beta y, 1+\gamma^2).$$

The interventional law is independent of y iff $\beta = 0$, matching the zero-correlation criterion in the randomized model.

Remark 1.13.27. *Exercise Integration Note*
The source exercises on three-node d-separation are integrated into Example 1.13.9. The simulation-and-bootstrap finite-DAG exercise is left as a source note because the screenshots give the task but no canonical numerical distribution or simulation output.

1.14 Nonparametric Curve Estimation

Nonparametric curve estimation studies unknown functions such as a density f or a regression function $r(x) = \mathbb{E}(Y \mid X = x)$. Unlike the empirical distribution function from Definition 1.1.12, these targets cannot be estimated well at fine resolution without smoothness. The practical question is therefore how much to smooth the data.

Definition 1.14.1. *Curve Estimator, ISE, and MISE*
Let g be an unknown function and let $\widehat{g}_n$ be an estimator of g. At a fixed point x, $\widehat{g}_n(x)$ is random because it is a function of the sample, while x is the evaluation point. The *integrated squared error* is

$$L(g, \widehat{g}_n) := \int (g(u) - \widehat{g}_n(u))^2 \, du.$$

The *risk*, or *mean integrated squared error*, is

$$R(g, \widehat{g}_n) := \mathbb{E}\left[L(g, \widehat{g}_n)\right].$$

Theorem 1.14.2. *Integrated Bias–Variance Decomposition*
Assume the following integrals are finite and that expectation may be interchanged with integration. Define the pointwise bias and variance by

$$b(x) := \mathbb{E}\widehat{g}_n(x) - g(x), \qquad v(x) := \mathrm{Var}(\widehat{g}_n(x)).$$

Then

$$R(g, \widehat{g}_n) = \int b(x)^2 \, dx + \int v(x) \, dx.$$

Proof. Apply the scalar identity in Theorem 1.1.7 at each fixed x:

$$\mathbb{E}\left(\widehat{g}_n(x) - g(x)\right)^2 = b(x)^2 + v(x).$$

Integrating the identity over x gives the result. □

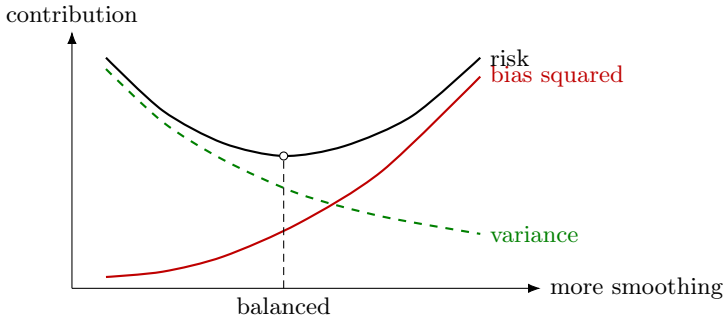


Figure 230: A smoothing parameter trades variance against squared bias. Too little smoothing has high variance; too much smoothing has high bias.

Principle 1.14.3. *Smoothing Trades Variance for Bias*
The source chapter repeatedly uses the same heuristic. Increasing the smoothing parameter averages over more data and reduces variance, but it also erases local features and increases bias. Choosing a smoothing parameter means balancing the two terms in Theorem 1.14.2.

1.14.1 Histograms

Definition 1.14.4. *Regular Histogram Estimator*
Let $X_1, \dots, X_n$ be IID observations from a density f on $[0,1]$. For an integer $m \geq 1$, set $h = 1/m$ and define bins

$$B_j = \left[\frac{j-1}{m}, \frac{j}{m}\right), \qquad j = 1, \dots, m-1, \qquad B_m = \left[\frac{m-1}{m}, 1\right].$$

Let ν_j be the number of observations in B_j,

$$\widehat{p}_j := \frac{\nu_j}{n}, \qquad p_j := \int_{B_j} f(u)\, du.$$

The *histogram density estimator* is

$$\widehat{f}_n(x) := \sum_{j=1}^{m} \frac{\widehat{p}_j}{h} \mathbf{1}_{\{x \in B_j\}}.$$

Proposition 1.14.5. *Histogram Mean and Variance*
Fix m and x, and let B_j be the bin containing x. Then

$$\mathbb{E}\widehat{f}_n(x) = \frac{p_j}{h}, \qquad \operatorname{Var}(\widehat{f}_n(x)) = \frac{p_j(1-p_j)}{nh^2}.$$

Proof. The count ν_j is Binomial(n, p_j), so $\mathbb{E}\widehat{p}_j = p_j$ and $\operatorname{Var}(\widehat{p}_j) = p_j(1-p_j)/n$. Divide by h for the mean and by h^2 for the variance. □

Theorem 1.14.6. *Histogram Risk and Optimal Binwidth*
Suppose $\int_0^1 (f'(u))^2\, du < +\infty$. For small h and large n,

$$R(f, \widehat{f}_n) \approx \frac{h^2}{12} \int_0^1 (f'(u))^2\, du + \frac{1}{nh}.$$

The binwidth that minimizes the displayed approximation is

$$h^* = n^{-1/3} \left(\frac{6}{\int_0^1 (f'(u))^2\, du} \right)^{1/3}.$$

With this choice, the approximate risk has order $n^{-2/3}$.

Proof. This is the source asymptotic calculation. For $x \in B_j$, a first-order Taylor expansion gives

$$\frac{p_j}{h} - f(x) \approx f'(x) \left(h\left(j - \frac{1}{2}\right) - x \right).$$

After squaring and integrating over each bin, the integrated squared bias is approximately

$$\frac{h^2}{12} \int_0^1 (f'(u))^2\, du.$$

For small h, $p_j \approx f(x)h$ and $1 - p_j \approx 1$, so the integrated variance is approximately $1/(nh)$. Minimize the resulting expression in h. □

Method 1.14.7. *Choosing Histogram Binwidth by Cross-Validation*
The population loss, as a function of h, is

$$L(h) = \int \left(\widehat{f}_n(x) - f(x)\right)^2 dx.$$

Since $\int f^2$ does not depend on h, minimizing risk is equivalent to minimizing the expectation of

$$J(h) := \int \widehat{f}_n(x)^2\, dx - 2\int \widehat{f}_n(x) f(x)\, dx.$$

The leave-one-out cross-validation score estimates this by

$$\widehat{J}(h) := \int \widehat{f}_n(x)^2\, dx - \frac{2}{n}\sum_{i=1}^{n} \widehat{f}_{-i}(X_i),$$

where $\widehat{f}_{-i}$ is the histogram estimator formed after deleting X_i. For a regular histogram this can be computed without refitting n histograms:

$$\widehat{J}(h) = \frac{2}{(n-1)h} - \frac{n+1}{(n-1)h}\sum_{j=1}^{m} \widehat{p}_j^{\,2}.$$

Proof. The first term satisfies $\int \widehat{f}_n^2 = h^{-1}\sum_j \widehat{p}_j^{\,2}$. If $X_i \in B_j$, then $\widehat{f}_{-i}(X_i) = (\nu_j - 1)/((n-1)h)$. Hence

$$\sum_i \widehat{f}_{-i}(X_i) = \frac{1}{(n-1)h}\sum_{j=1}^{m} \nu_j(\nu_j - 1).$$

Substitution gives the displayed shortcut. □

Method 1.14.8. *Histogram Confidence Envelope*
For a histogram with m bins, define the histogramized target

$$\widetilde{f}(x) := \frac{p_j}{h} \qquad x \in B_j.$$

Set

$$c := \frac{z_{\alpha/(2m)}}{2}\sqrt{\frac{m}{n}}, \qquad \ell(x) := \left(\max\left\{\sqrt{\widehat{f}_n(x)} - c, 0\right\}\right)^2, \qquad u(x) := \left(\sqrt{\widehat{f}_n(x)} + c\right)^2.$$

Under the asymptotic normal approximation in the source, and if $m = m(n)$ grows slowly enough that $m \log n/n \to 0$,

$$\mathbb{P}\left(\ell(x) \le \widetilde{f}(x) \le u(x) \text{ for all } x\right) \approx 1 - \alpha.$$

Proof Idea: Use the central limit theorem for $\widehat{p}_j$, apply the delta method to $\sqrt{\widehat{p}_j}$, and use a union bound over the m bins. The band is for the histogramized density $\widetilde{f}$, not the exact density f.

1.14.2 Kernel Density Estimation

Definition 1.14.9. *Kernel Density Estimator*
A *kernel* is a smooth function K such that

$$K(x) \ge 0, \qquad \int K(x)\, dx = 1, \qquad \int xK(x)\, dx = 0, \qquad \sigma_K^2 := \int x^2 K(x)\, dx > 0.$$

Given $h > 0$, the *bandwidth*, the kernel density estimator is

$$\widehat{f}_n(x) := \frac{1}{n}\sum_{i=1}^{n} \frac{1}{h} K\left(\frac{x - X_i}{h}\right).$$

Two source examples are the Gaussian kernel $K(x) = (2\pi)^{-1/2} e^{-x^2/2}$ and the Epanechnikov kernel

$$K(x) = \begin{cases} \frac{3}{4}(1 - x^2/5)/\sqrt{5}, & |x| < \sqrt{5}, \\ 0, & \text{otherwise.} \end{cases}$$

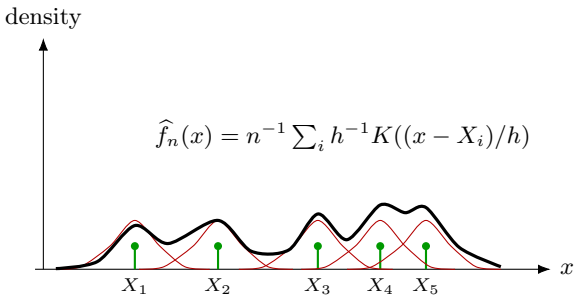


Figure 231: A kernel density estimator averages one scaled kernel bump around each data point.

Theorem 1.14.10. *Kernel Density Risk Rate*
Under the smoothness assumptions used in the source asymptotic calculation,

$$R(f,\widehat{f}_n) \approx \frac{1}{4}\sigma_K^4 h^4 \int (f''(x))^2\,dx + \frac{\int K(x)^2\,dx}{nh}.$$

Writing

$$c_1 = \int x^2 K(x)\,dx, \qquad c_2 = \int K(x)^2\,dx, \qquad c_3 = \int (f''(x))^2\,dx,$$

the bandwidth that minimizes this approximation is

$$h^* = c_1^{-2/5} c_2^{1/5} c_3^{-1/5} n^{-1/5}.$$

At this bandwidth, the approximate risk has order $n^{-4/5}$.

Proof. The source proof writes $K_h(x,X) = h^{-1}K((x-X)/h)$, so

$$\mathbb{E}\widehat{f}_n(x) = \mathbb{E}K_h(x,X), \qquad \operatorname{Var}(\widehat{f}_n(x)) = n^{-1}\operatorname{Var}(K_h(x,X)).$$

A Taylor expansion and the moment conditions on K give bias approximately $\frac{1}{2}\sigma_K^2 h^2 f''(x)$, while the leading variance term is approximately $f(x)\int K^2/(nh)$. Integrating squared bias plus variance yields the display. □

Proposition 1.14.11. *Boxcar Kernel Exercise*
Let $K(t) = \mathbf{1}_{\{-1/2<t<1/2\}}$ *and let*

$$\widehat{f}_n(x) = \frac{1}{n}\sum_{i=1}^{n} \frac{1}{h} K\left(\frac{x-X_i}{h}\right).$$

Set

$$p_h(x) := \int_{x-h/2}^{x+h/2} f(y)\,dy.$$

Then

$$\mathbb{E}\widehat{f}_n(x) = \frac{p_h(x)}{h}, \qquad \operatorname{Var}(\widehat{f}_n(x)) = \frac{1}{nh^2}\left(p_h(x) - p_h(x)^2\right).$$

If f *is continuous at* x, $h \to 0$, *and* $nh \to +\infty$, *then* $\widehat{f}_n(x) \xrightarrow{\mathbb{P}} f(x)$.

Proof. The summand $K((x-X_i)/h)$ is Bernoulli with success probability $p_h(x)$. This gives the moments. Continuity gives $p_h(x)/h \to f(x)$, and the variance is $O((nh)^{-1})$. □

Method 1.14.12. *Kernel Density Cross-Validation*
For bandwidth selection, use the leave-one-out score

$$\widehat{J}(h) = \int \widehat{f}_n(x)^2\,dx - \frac{2}{n}\sum_{i=1}^{n} \widehat{f}_{-i}(X_i),$$

where $\widehat{f}_{-i}$ omits X_i. The source gives the computational shortcut

$$\widehat{J}(h) \approx \frac{1}{hn^2}\sum_i\sum_j K^*\left(\frac{X_i - X_j}{h}\right) + \frac{2}{nh}K(0),$$

where

$$K^*(z) = K^{(2)}(z) - 2K(z), \qquad K^{(2)}(z) = \int K(z-y)K(y)\,dy.$$

For a standard Gaussian kernel, $K^{(2)}$ is the $N(0,2)$ density.

Theorem 1.14.13. *Stone's Cross-Validation Oracle Property*
Let $\widehat{f}_h$ be the kernel estimator with bandwidth h, and let h_n be chosen by cross-validation. Under the boundedness condition stated in the source,

$$\frac{\int \left(f(x) - \widehat{f}_{h_n}(x)\right)^2 dx}{\inf_h \int \left(f(x) - \widehat{f}_h(x)\right)^2 dx} \xrightarrow{\mathbb{P}} 1.$$

Method 1.14.14. *Grid Confidence Band for KDE*
To build an approximate confidence band, choose a grid $\mathcal{V} = \{v_1, \dots, v_N\}$. For each $v \in \mathcal{V}$, define

$$Y_i(v) := \frac{1}{h} K\left(\frac{v - X_i}{h}\right), \qquad \widehat{f}_n(v) = \overline{Y}_n(v).$$

Estimate the standard error by

$$\operatorname{se}(v) = \frac{s(v)}{\sqrt{n}}, \qquad s^2(v) = \frac{1}{n-1} \sum_{i=1}^{n} \left(Y_i(v) - \overline{Y}_n(v)\right)^2.$$

The source then uses an effective sample size

$$\operatorname{ESS}(v) = \frac{\sum_{i=1}^{n} K((v - X_i)/h)}{K(0)}$$

to choose an approximate number m of independent blocks. With

$$q = \Phi^{-1}\left(\frac{1 + (1-\alpha)^{1/m}}{2}\right),$$

set

$$\ell(v) = \max\{\widehat{f}_n(v) - q \operatorname{se}(v), 0\}, \qquad u(v) = \widehat{f}_n(v) + q \operatorname{se}(v).$$

The intended coverage statement is approximate and grid-based:

$$\mathbb{P}\{\ell(v) \le f(v) \le u(v) \text{ for all } v \in \mathcal{V}\} \approx 1 - \alpha.$$

1.14.3 Nonparametric Regression

Definition 1.14.15. *Nadaraya–Watson Kernel Regression*
For regression data $(x_1, Y_1), \dots, (x_n, Y_n)$ satisfying

$$Y_i = r(x_i) + \varepsilon_i, \qquad \mathbb{E}\varepsilon_i = 0,$$

the *Nadaraya–Watson kernel estimator* is

$$\widehat{r}(x) = \sum_{i=1}^{n} w_i(x) Y_i, \qquad w_i(x) = \frac{K((x - x_i)/h)}{\sum_{j=1}^{n} K((x - x_j)/h)}.$$

The estimator is a local weighted average, with higher weight assigned to observations near x.

Proposition 1.14.16. *Binned Regression Risk*
Suppose $0 \le x_i \le 1$, define bins B_j as in Definition 1.14.4, and estimate

$$\widehat{r}_n(x) = \overline{Y}_j \qquad x \in B_j,$$

where $\overline{Y}_j$ is the average of the Y_i with $x_i \in B_j$. If the design is approximately regular, $\operatorname{Var}(\varepsilon_i) = \sigma^2$, and r is smooth, then

$$R(r, \widehat{r}_n) \approx \frac{h^2}{12} \int_0^1 (r'(u))^2 \, du + \frac{\sigma^2}{nh}.$$

Thus the optimal binwidth is of order $n^{-1/3}$, and the corresponding risk is of order $n^{-2/3}$.

Proof. This formalizes the source exercise by repeating the histogram calculation for the regression function. The bias comes from approximating r by its bin average, and the variance is the noise variance divided by the number of observations in a bin. □

Theorem 1.14.17. *Nadaraya–Watson Risk Rate*
Under the smoothness conditions in the source calculation, the Nadaraya–Watson estimator has the same bandwidth order as one-dimensional kernel density estimation:

$$h^* \asymp n^{-1/5}, \qquad R(\widehat{r}_n, r) \asymp n^{-4/5}.$$

More precisely, the source gives an approximation of the form

$$R(\widehat{r}_n, r) \approx \textit{constant} \cdot h^4 + \textit{constant} \cdot \frac{1}{nh}.$$

Remark 1.14.18. *Risk Constant Scope*
The screenshots give the full Nadaraya–Watson risk expression, involving r'', r', the design density f, the noise variance σ^2, and kernel moments. I have recorded the rate and structure here, rather than reproduce every constant, because the surrounding notes do not yet develop the design-density regularity assumptions needed to make that formula self-contained.

Method 1.14.19. *Regression Bandwidth Cross-Validation*
Choose the regression bandwidth by minimizing

$$\widehat{J}(h) = \sum_{i=1}^{n} \left(Y_i - \widehat{r}_{-i}(x_i)\right)^2,$$

where $\widehat{r}_{-i}$ omits observation i. For kernel smoothers, the source gives the computable approximation

$$\widehat{J}(h) \approx \sum_{i=1}^{n} \left(Y_i - \widehat{r}(x_i)\right)^2 \left(1 - \frac{K(0)}{\sum_{j=1}^{n} K((x_i - x_j)/h)}\right)^{-2}.$$

Proposition 1.14.20. *Adjacent-Difference Estimate of Noise Variance*
Assume the x_i are ordered, r is smooth, and $\mathrm{Var}(\varepsilon_i) = \sigma^2$. Define

$$\widehat{\sigma}^2 := \frac{1}{2(n-1)} \sum_{i=1}^{n-1} (Y_{i+1} - Y_i)^2.$$

Then $\widehat{\sigma}^2$ is a natural consistent estimator of σ^2 under the source smoothness heuristic.

Proof. If adjacent design points are close, then $r(x_{i+1}) - r(x_i) \approx 0$, so

$$Y_{i+1} - Y_i \approx \varepsilon_{i+1} - \varepsilon_i.$$

The variance of the right-hand side is $2\sigma^2$. Averaging the squared adjacent differences and dividing by 2 gives the estimator. □

Method 1.14.21. *Kernel Regression Confidence Band*
For a grid point v, use the same independent-block correction as in Method 1.14.14, but replace the standard error by

$$\mathrm{se}(v) = \widehat{\sigma} \left(\sum_{i=1}^{n} w_i(v)^2\right)^{1/2},$$

where $w_i(v)$ are the Nadaraya–Watson weights and $\widehat{\sigma}$ is defined in Proposition 1.14.20. This gives a confidence envelope for the smoothed regression curve rather than for every fine feature of the true r.

Remark 1.14.22. *Confidence Bands and Bias*
The confidence bands above are best understood as bands for smoothed targets. In a scalar parametric problem, if $\widehat{\theta}$ has bias b_n and standard error s_n, the usual normal interval has nominal asymptotic coverage when $b_n/s_n \to 0$. For kernel density estimation at a fixed point,

$$\widehat{f}(x) \approx N\left(f(x) + b_n(x), \frac{c_2 f(x)}{nh}\right), \qquad b_n(x) \approx \frac{1}{2} h^2 f''(x) c_1,$$

where $c_1 = \int x^2 K(x)\,dx$ and $c_2 = \int K(x)^2\,dx$. At the risk-optimal bandwidth $h \asymp n^{-1/5}$, the bias and standard error are of the same order, so a naive band for the unsmoothed f may undercover.

Definition 1.14.23. *Additive Nonparametric Regression*
For multiple regressors $X = (X_1, \ldots, X_p)$, a fully nonparametric regression function suffers from the same curse of dimensionality as multivariate density estimation. An *additive regression model* restricts the regression function to

$$Y = \sum_{j=1}^{p} r_j(X_j) + \varepsilon.$$

One may enrich the model by adding selected interactions, for example

$$Y = \sum_{j=1}^{p} r_j(X_j) + \sum_{j<k} r_{jk}(X_j X_k) + \varepsilon.$$

The gain is that the main-effect additive model only asks us to estimate p one-dimensional functions.

Method 1.14.24. *Backfitting*
The source describes the standard iterative fitting method for additive regression.

(i) Initialize $r_1, \ldots, r_p$.

(ii) For $j = 1, \dots, p$, form partial residuals

$$e_i^{(j)} = Y_i - \sum_{s \neq j} r_s(X_{is}),$$

then update r_j by nonparametrically regressing the partial residuals $e_i^{(j)}$ on the jth covariate.
(iii) Stop when the component functions have converged; otherwise repeat the previous step.

Additive models reduce dimensionality, but they are restrictive: the true regression function need not decompose as a sum of one-dimensional component functions.

Remark 1.14.25. *Source Examples and Gaps*
The source uses astronomy distance data, cosmic microwave background data, and glass-fragment data to illustrate smoothing, cross-validation, and confidence bands. The screenshots include plots and narrative conclusions but not the raw observations, so I integrated the mathematical procedures and qualitative lessons rather than reproducing data-dependent numerical figures.

1.15 Orthogonal Series and Wavelet Smoothing

Orthogonal series methods replace local averaging by coordinates in an orthonormal basis. The analytic backbone is already in the Fourier and Hilbert-space notes: Fourier coefficients give best mean-square approximation in Theorem 3.5.9, and complete orthonormal systems in L^2 satisfy Parseval identities in Theorem 6.2.27.

1.15.1 Orthogonal Coordinates in L^2

Definition 1.15.1. *Orthogonal Series Coordinates*
Let $(\phi_j)_{j \geq 1}$ be a complete orthonormal system for $L^2[a,b]$ over the real scalars. For $f \in L^2[a,b]$, define

$$\beta_j := \int_a^b f(x)\phi_j(x)\,dx.$$

Then, in L^2,

$$f(x) = \sum_{j=1}^{\infty} \beta_j \phi_j(x), \qquad \int_a^b f(x)^2\,dx = \sum_{j=1}^{\infty} \beta_j^2.$$

The finite sum $f_J(x) = \sum_{j=1}^{J} \beta_j \phi_j(x)$ is the best J-term approximation among sums using the first J basis functions.

Example 1.15.2. *Cosine and Legendre Bases*
On $[0,1]$, the source uses the cosine basis

$$\phi_1(x) = 1, \qquad \phi_j(x) = \sqrt{2}\cos((j-1)\pi x), \quad j \geq 2.$$

On $[-1,1]$, the Legendre polynomials are

$$P_\ell(x) = \frac{1}{2^\ell \ell!}\frac{d^\ell}{dx^\ell}(x^2-1)^\ell, \qquad \ell = 0,1,2,\dots,$$

with

$$\int_{-1}^{1} P_\ell(x)^2\,dx = \frac{2}{2\ell+1}.$$

Thus $\phi_\ell(x) = \sqrt{(2\ell+1)/2}\,P_\ell(x)$ gives an orthonormal basis after indexing by $\ell = 0,1,\dots$. The first few are

$$P_0(x) = 1, \qquad P_1(x) = x, \qquad P_2(x) = \frac{1}{2}(3x^2-1), \qquad P_3(x) = \frac{1}{2}(5x^3-3x),$$

and they satisfy the recursion

$$P_{\ell+1}(x) = \frac{(2\ell+1)xP_\ell(x) - \ell P_{\ell-1}(x)}{\ell+1}.$$

Principle 1.15.3. *Smoothness Means Small High-Frequency Coefficients*
In an oscillatory basis, high-index coefficients measure high-frequency structure. For the cosine basis, differentiating the jth term multiplies its size by a power of $\pi(j-1)$. Thus finiteness of an energy such as $\int_0^1 (f^{(k)}(x))^2\,dx$ forces the weighted coefficient sequence to be summable, and therefore forces β_j to be small for large j.

$$\int_0^1 (f^{(k)}(x))^2\,dx \asymp \sum_{j=1}^{\infty} \beta_j^2 \big(\pi(j-1)\big)^{2k},$$

up to the harmless normalization constants attached to the chosen cosine basis.

Example 1.15.4. *A Finite Orthonormal Basis Check*
The source begins with finite-dimensional coordinates. The vectors

$$\psi_1 = \left(\frac{1}{\sqrt{3}}, \frac{1}{\sqrt{3}}, \frac{1}{\sqrt{3}}\right), \quad \psi_2 = \left(\frac{1}{\sqrt{2}}, -\frac{1}{\sqrt{2}}, 0\right), \quad \psi_3 = \left(\frac{1}{\sqrt{6}}, \frac{1}{\sqrt{6}}, -\frac{2}{\sqrt{6}}\right)$$

have norm 1 and pairwise dot products 0. Hence every $v \in \mathbb{R}^3$ has the expansion

$$v = \sum_{j=1}^{3} \langle \psi_j, v \rangle \psi_j.$$

Example 1.15.5. *Two Cosine-Basis Exercise Answers*
For the cosine basis on $[0,1]$,

$$\sqrt{2}\cos(3\pi x) = \phi_4(x),$$

so its only nonzero coefficient is $\beta_4 = 1$. For $f(x) = \sin(\pi x)$,

$$\beta_1 = \frac{2}{\pi},$$

and for $m = j - 1 \geq 1$,

$$\beta_{m+1} = \begin{cases} 0, & m \text{ odd}, \\ -\dfrac{2\sqrt{2}}{\pi(m^2-1)}, & m \text{ even}. \end{cases}$$

This formalizes the analytic part of the source exercise asking for cosine-basis expansions.

1.15.2 Orthogonal Series Density Estimation

Definition 1.15.6. *Orthogonal Series Density Estimator*
Let $X_1, \ldots, X_n$ be IID from a density f on $[0,1]$, and suppose

$$f(x) = \sum_{j=1}^{\infty} \beta_j \phi_j(x)$$

in an orthonormal basis. Since $\beta_j = \mathbb{E}\phi_j(X_1)$, estimate

$$\widehat{\beta}_j := \frac{1}{n} \sum_{i=1}^{n} \phi_j(X_i).$$

The J-term orthogonal series estimator is

$$\widehat{f}_J(x) := \sum_{j=1}^{J} \widehat{\beta}_j \phi_j(x).$$

Theorem 1.15.7. *Density Coefficient Moments and Risk*
Let

$$\sigma_j^2 := \operatorname{Var}(\phi_j(X_1)) = \int (\phi_j(x) - \beta_j)^2 f(x)\, dx.$$

Then

$$\mathbb{E}\widehat{\beta}_j = \beta_j, \qquad \operatorname{Var}(\widehat{\beta}_j) = \frac{\sigma_j^2}{n}.$$

Moreover, the integrated risk of $\widehat{f}_J$ is

$$R(J) = \sum_{j=1}^{J} \frac{\sigma_j^2}{n} + \sum_{j=J+1}^{\infty} \beta_j^2.$$

Proof. The moment identities follow because $\widehat{\beta}_j$ is a sample mean of $\phi_j(X_i)$. Orthogonality makes the squared L^2 error split into retained-coordinate variance plus omitted-coordinate squared bias. □

Method 1.15.8. *Orthogonal Density Risk Estimation*
Restrict $1 \leq J \leq p$, with the source choice $p = \sqrt{n}$. Estimate

$$\widehat{\sigma}_j^2 := \frac{1}{n-1} \sum_{i=1}^{n} \left(\phi_j(X_i) - \widehat{\beta}_j\right)^2.$$

The source risk estimate is

$$\widehat{R}(J) = \sum_{j=1}^{J} \frac{\widehat{\sigma}_j^2}{n} + \sum_{j=J+1}^{p} \left(\widehat{\beta}_j^2 - \frac{\widehat{\sigma}_j^2}{n} \right)_+, \qquad a_+ := \max\{a, 0\}.$$

Choose $\widehat{J}$ to minimize $\widehat{R}(J)$ and set

$$\widehat{f}(x) = \sum_{j=1}^{\widehat{J}} \widehat{\beta}_j \phi_j(x).$$

If a genuine density is required, truncate the negative part and normalize:

$$\widehat{f}^*(x) = \frac{\max\{\widehat{f}(x), 0\}}{\int_0^1 \max\{\widehat{f}(u), 0\}\, du}.$$

Theorem 1.15.9. *Orthogonal Density Confidence Band*
Let

$$\widetilde{f}(x) = \sum_{j=1}^{J} \beta_j \phi_j(x)$$

be the J-term smoothed density. An approximate $1-\alpha$ band for $\widetilde{f}$ is

$$\ell(x) = \widehat{f}_J(x) - c, \qquad u(x) = \widehat{f}_J(x) + c,$$

where

$$c = \frac{JK^2}{\sqrt{n}} \sqrt{1 + \frac{\sqrt{2z_\alpha}}{\sqrt{J}}}, \qquad K = \max_{1 \le j \le J} \max_x |\phi_j(x)|.$$

For the cosine basis, $K = \sqrt{2}$. The band is for the smoothed target $\widetilde{f}$, not for all fine-scale features of f.

1.15.3 Orthogonal Series Regression

Definition 1.15.10. *Orthogonal Series Regression Estimator*
Consider fixed-design regression

$$Y_i = r(x_i) + \varepsilon_i, \qquad i = 1, \dots, n,$$

with independent errors of mean 0 and variance σ^2. In the source regular-grid case, $x_i = i/n$ and

$$r(x) = \sum_{j=1}^{\infty} \beta_j \phi_j(x).$$

Estimate

$$\widehat{\beta}_j := \frac{1}{n} \sum_{i=1}^{n} Y_i \phi_j(x_i), \qquad \widehat{r}_J(x) := \sum_{j=1}^{J} \widehat{\beta}_j \phi_j(x).$$

Theorem 1.15.11. *Orthogonal Regression Coefficients and Risk*
Under the regular-grid approximation used in the source,

$$\widehat{\beta}_j \approx N\left(\beta_j, \frac{\sigma^2}{n} \right),$$

and the J-term risk is approximately

$$R(J) = \mathbb{E} \int (\widehat{r}_J(x) - r(x))^2\, dx = \frac{J\sigma^2}{n} + \sum_{j=J+1}^{\infty} \beta_j^2.$$

Proof. The expectation of $\widehat{\beta}_j$ is a Riemann-sum approximation to $\int r(x)\phi_j(x)\, dx = \beta_j$. The variance is approximately

$$\frac{\sigma^2}{n^2} \sum_{i=1}^{n} \phi_j(x_i)^2 \approx \frac{\sigma^2}{n} \int_0^1 \phi_j(x)^2\, dx = \frac{\sigma^2}{n}.$$

Orthogonality gives the risk decomposition. □

Method 1.15.12. *REACT Risk Estimation*
The source estimates σ^2 from high-frequency coefficients. With $k \approx n/4$,

$$\widehat{\sigma}^2 := \frac{n}{k} \sum_{j=n-k+1}^{n} \widehat{\beta}_j^2.$$

Then estimate risk by

$$\widehat{R}(J) = J\frac{\widehat{\sigma}^2}{n} + \sum_{j=J+1}^{n}\left(\widehat{\beta}_j^2 - \frac{\widehat{\sigma}^2}{n}\right)_+.$$

Choose $\widehat{J}$ minimizing $\widehat{R}(J)$ and set

$$\widehat{r}(x) = \sum_{j=1}^{\widehat{J}} \widehat{\beta}_j \phi_j(x).$$

This is the orthogonal series regression procedure that the source identifies as REACT.

Theorem 1.15.13. *Orthogonal Regression Confidence Band*
Suppose $\widehat{r}$ is based on J terms, $\widehat{\sigma}^2$ is defined as in Method 1.15.12, and $J < n - k + 1$. An approximate $1-\alpha$ confidence band for the smoothed regression function is

$$\ell(x) = \widehat{r}(x) - K\sqrt{J}\sqrt{\frac{z_\alpha \widehat{\tau}}{\sqrt{n}} + \frac{J\widehat{\sigma}^2}{n}},$$

$$u(x) = \widehat{r}(x) + K\sqrt{J}\sqrt{\frac{z_\alpha \widehat{\tau}}{\sqrt{n}} + \frac{J\widehat{\sigma}^2}{n}},$$

where

$$K = \max_{1\le j\le J}\max_x |\phi_j(x)|, \qquad \widehat{\tau}^2 = \frac{2J\widehat{\sigma}^4}{n}\left(1 + \frac{J}{k}\right).$$

For the cosine basis, $K = \sqrt{2}$.

Remark 1.15.14. *Source Data Examples*
The source illustrates these methods with the Doppler test function, simulated noisy observations, and confidence envelopes for different truncation levels. The screenshots do not provide the simulated observations, so the notes record the formulas and interpretation rather than reproduce the numerical plots.

1.15.4 Haar Wavelets

Perspective 1.15.15. *Why Wavelets*
Cosine and polynomial bases are global: keeping more terms improves resolution, but can introduce oscillation across the whole interval. This is inefficient for a function that is mostly smooth but has a sharp local feature. Wavelets are localized basis functions, so a large coefficient can add a local feature without forcing high-frequency oscillation everywhere else. This is the source motivation for using wavelets on spatially inhomogeneous signals.

Definition 1.15.16. *Haar Father and Mother Wavelets*
The *Haar father wavelet*, or scaling function, is

$$\phi(x) = \begin{cases} 1, & 0 \le x < 1, \\ 0, & \text{otherwise.} \end{cases}$$

The *Haar mother wavelet* is

$$\psi(x) = \begin{cases} -1, & 0 \le x \le \frac{1}{2}, \\ 1, & \frac{1}{2} < x \le 1, \\ 0, & \text{otherwise.} \end{cases}$$

For integers $j \ge 0$ and k, define

$$\phi_{j,k}(x) := 2^{j/2}\phi(2^j x - k), \qquad \psi_{j,k}(x) := 2^{j/2}\psi(2^j x - k).$$

Increasing j localizes the wavelet on a shorter interval.

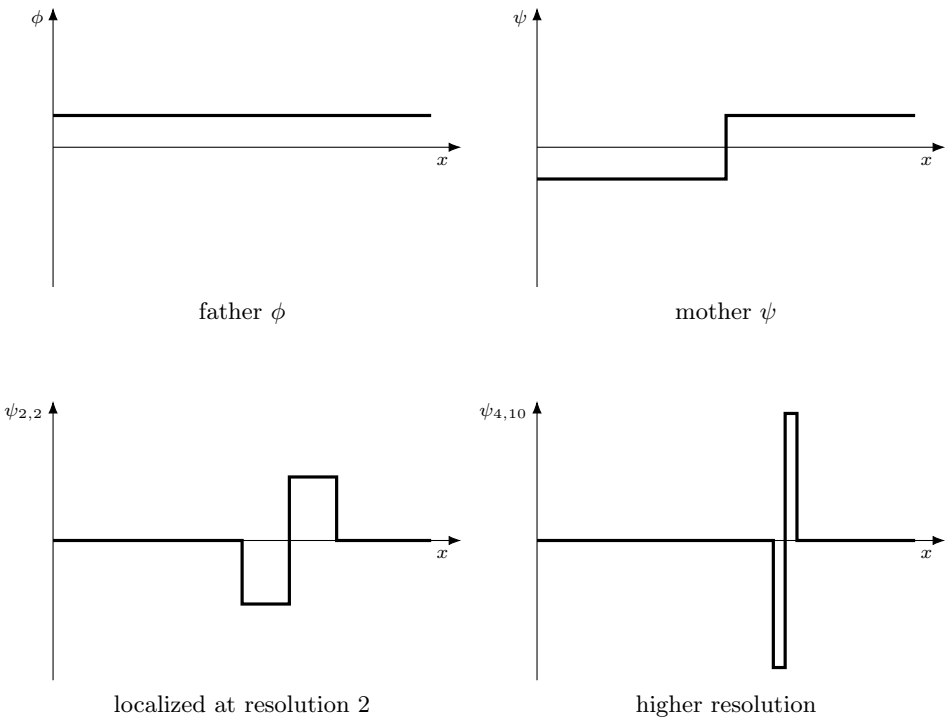


Figure 232: Haar wavelets are localized basis functions: higher resolution means narrower support and larger amplitude.

Theorem 1.15.17. *Haar Basis Expansion*
Let

$$W_j = \{\psi_{j,k} : k = 0, 1, \ldots, 2^j - 1\}.$$

The set

$$\{\phi, W_0, W_1, W_2, \ldots\}$$

is an orthonormal basis for $L^2(0,1)$*. Hence every* $f \in L^2(0,1)$ *has the expansion*

$$f(x) = \alpha\phi(x) + \sum_{j=0}^{\infty} \sum_{k=0}^{2^j-1} \beta_{j,k}\psi_{j,k}(x),$$

where

$$\alpha = \int_0^1 f(x)\phi(x)\,dx, \qquad \beta_{j,k} = \int_0^1 f(x)\psi_{j,k}(x)\,dx.$$

The resolution-J approximation is

$$\widetilde{f}_J(x) = \alpha\phi(x) + \sum_{j=0}^{J-1} \sum_{k=0}^{2^j-1} \beta_{j,k}\psi_{j,k}(x),$$

with 2^J *total terms.*

Proposition 1.15.18. *Why Haar Wavelets Are Orthonormal*
The Haar functions in Theorem 1.15.17 have norm 1 *and are pairwise orthogonal.*

Proof. Each $\psi_{j,k}$ is supported on one dyadic interval of length 2^{-j} and has absolute value $2^{j/2}$ there, so its squared integral is $2^j 2^{-j} = 1$. Disjoint supports give orthogonality. If supports are nested, the finer wavelet has equal positive and negative mass on each half of the coarser interval, so the product integrates to 0. Orthogonality with ϕ follows because each mother wavelet has integral 0. □

Method 1.15.19. *Haar Wavelet Regression by Universal Thresholding*
Assume $n = 2^J$ and $x_i = i/n$. For regression data $Y_i = r(x_i) + \sigma\varepsilon_i$, estimate the scaling and detail coefficients by

$$\widehat{\alpha} = \frac{1}{n}\sum_i \phi(x_i)Y_i, \qquad D_{j,k} = \frac{1}{n}\sum_i \psi_{j,k}(x_i)Y_i, \qquad 0 \le j \le J-1.$$

Estimate the noise level by the median absolute deviation at the finest scale:

$$\widehat{\sigma} = \sqrt{n}\,\frac{\operatorname{median}\{|D_{J-1,k}| : k = 0, \dots, 2^{J-1}-1\}}{0.6745}.$$

Apply hard universal thresholding:

$$\widehat{\beta}_{j,k} = \begin{cases} D_{j,k}, & |D_{j,k}| > \widehat{\sigma}\sqrt{2\log n/n}, \\ 0, & \text{otherwise.} \end{cases}$$

Then set

$$\widehat{r}(x) = \widehat{\alpha}\phi(x) + \sum_{j=0}^{J-1}\sum_{k=0}^{2^j-1} \widehat{\beta}_{j,k}\psi_{j,k}(x).$$

Method 1.15.20. *Discrete Haar Wavelet Transform*
The source notes that the coefficients are usually computed by the discrete wavelet transform. Let y be the vector of observations, $J = \log_2 n$, and set $\text{temp} = y/\sqrt{n}$. For $j = J-1, J-2, \dots, 0$, form

$$D_j(I) = \frac{\text{temp}(2I) - \text{temp}(2I-1)}{\sqrt{2}}, \qquad I = 1, \dots, 2^j,$$

and replace

$$\text{temp}(I) = \frac{\text{temp}(2I) + \text{temp}(2I-1)}{\sqrt{2}}.$$

The recursion stores detail coefficients while replacing adjacent pairs by their local averages.

Theorem 1.15.21. *Universal Threshold Removes Pure Noise*
Suppose $\beta_{j,k} = 0$ for all j, k, and let $\widehat{\beta}_{j,k}$ be the universal-threshold estimator in Method 1.15.19. Then

$$\mathbb{P}\left(\widehat{\beta}_{j,k} = 0 \text{ for all } j,k\right) \to 1.$$

Proof. The source proof assumes σ known for clarity. Then $D_{j,k} \approx N(0, \sigma^2/n)$. With $\lambda = \sigma\sqrt{2\log n/n}$, Mills' inequality and a union bound give

$$\mathbb{P}\left(\max_{j,k}|D_{j,k}| > \lambda\right) \le \sum_{j,k}\mathbb{P}\left(\frac{\sqrt{n}\,|D_{j,k}|}{\sigma} > \sqrt{2\log n}\right) \to 0.$$

Therefore all pure-noise coefficients are thresholded away with probability tending to 1. □

Proposition 1.15.22. *Haar Density Exercise Formalization*
Let $X_1, \dots, X_n$ have density f on $[0,1]$, and write the Haar expansion

$$f(x) = \phi(x) + \sum_{j=0}^{\infty}\sum_{k=0}^{2^j-1}\beta_{j,k}\psi_{j,k}(x),$$

where the scaling coefficient is 1 for a density on $[0,1]$. The coefficient estimator

$$\widehat{\beta}_{j,k} := \frac{1}{n}\sum_{i=1}^{n}\psi_{j,k}(X_i)$$

is unbiased for $\beta_{j,k}$. The resolution-B Haar histogram

$$\widehat{f}_B(x) = \phi(x) + \sum_{j=0}^{B}\sum_{k=0}^{2^j-1}\widehat{\beta}_{j,k}\psi_{j,k}(x)$$

has mean integrated squared error

$$\sum_{j=0}^{B}\sum_{k=0}^{2^j-1}\frac{\operatorname{Var}(\psi_{j,k}(X_1))}{n} + \sum_{j=B+1}^{\infty}\sum_{k=0}^{2^j-1}\beta_{j,k}^2.$$

Proof. Unbiasedness is $\mathbb{E}\psi_{j,k}(X_1) = \int_0^1 \psi_{j,k}(x)f(x)\,dx = \beta_{j,k}$. The risk formula is the orthogonal-coordinate variance plus omitted-tail bias decomposition. □

Remark 1.15.23. *Why the Constant 0.6745 Appears*
If $Z \sim N(0,1)$, then $\operatorname{median}(|Z|) \approx 0.6745$. Dividing a median absolute deviation by this number converts it to a standard-deviation estimate under a normal-noise approximation. The source uses the finest Haar details because they are designed to be relatively insensitive to isolated sharp peaks in the signal.

2 Deep Learning and Representation Systems

2.1 Neural Network Foundations

2.1.1 Deep Feedforward Networks

Remark 2.1.1. *Deep Feedforward Networks (DFNs)*
Defines a mapping $\boldsymbol{y} = f(\boldsymbol{x}; \boldsymbol{\theta})$ and learns value of parameters $\boldsymbol{\theta}$ that results in best function approximation.

Definition 2.1.2. *Cost Functions*

i. Learning Conditional Distributions with Maximum Likelihood: cost function is negative log-likelihood, which is the cross-entropy between training data and model distribution:

$$J(\boldsymbol{\theta}) = -\mathbb{E}_{\boldsymbol{x},\boldsymbol{y}\sim\hat{p}_{\text{data}}} \log p_{\text{model}}(\boldsymbol{y} \mid \boldsymbol{x})$$

Specific form of cost function changes from model to model, depending on form of $\log p_{\text{model}}$.
Method removes the burden of designing cost functions for each model, as specifying a model $p(\boldsymbol{y} \mid \boldsymbol{x})$ automatically determines a cost function $\log p(\boldsymbol{y} \mid \boldsymbol{x})$.
Gradient of cost function for neural networks must be large and predictable enough to serve as good guide for the learning algorithm. The negative log-likelihood helps to avoid saturation of the function.

ii. Learning Conditional Statistics: to learn just one conditional statistic of $\boldsymbol{y}$ given $\boldsymbol{x}$.
With sufficiently powerful neural network, this can represent any function f from a wide class of function, limited by features of continuity and boundedness. Hence the cost function is a functional. Learning is choosing a functional rather than a set of parameters.
By calculus of variation , solving the optimisation problem yields the below function,

$$f^* = \arg\min_f \mathbb{E}_{\boldsymbol{x},\boldsymbol{y}\sim p_{\text{data}}} \|\boldsymbol{y} - f(\boldsymbol{x})\|^2$$

$$f^* = \mathbb{E}_{\boldsymbol{y}\sim p_{\text{data}}(\boldsymbol{y} \mid \boldsymbol{x})}[\boldsymbol{y}]$$

If infinitely many samples from the true data-generating distribution could be trained, then minimising the mean squared error cost function gives a function that predicts mean of $\boldsymbol{y}$ for each value of $\boldsymbol{x}$.
A second result from calculus of variations is:

$$f^* = \arg\min_f \mathbb{E}_{\boldsymbol{x},\boldsymbol{y}\sim p_{\text{data}}} \|\boldsymbol{y} - f(\boldsymbol{x})\|_1$$

which is a function that predicts the median value of $\boldsymbol{y}$ for each $\boldsymbol{x}$. This is the mean absolute error.
Note that mean squared error and mean absolute error often lead to poor results when used with gradient-based optimisation. Output units that saturate may produce very small gradients when combined with these cost functions. Hence the reason that cross-entropy cost function is more popular.

Definition 2.1.3. *Output Units*

i. Linear Units: base on affine transformation with no nonlinearity.
Given features $\boldsymbol{h}$, a layer of output units produces vector $\hat{\boldsymbol{y}} = \boldsymbol{W}^T\boldsymbol{h} + \boldsymbol{b}$. Linear output layers are used to produce mean of conditional Gaussian distribution:

$$p(\boldsymbol{y} \mid \boldsymbol{x}) = \mathcal{N}(\boldsymbol{y}; \hat{\boldsymbol{y}}, \boldsymbol{I})$$

Maximising log-likelihood is equivalent to minimising mean squared error.
Linear units do not saturate, hence may be used for wide variety of optimisation algorithms.

ii. Sigmoid Units: define Bernoulli distribution y conditioned on $\boldsymbol{x}$. Neural net to predict $P(y = 1 \mid \boldsymbol{x})$, which lies in interval $[0, 1]$. The sigmoid unit is defined by

$$\hat{y} = \sigma(\boldsymbol{w}^T\boldsymbol{h} + b)$$

where σ is the logistic sigmoid function.
Note that the cost function used with maximum likelihood is $-\log P(y \mid \boldsymbol{x})$, preventing saturation.
The loss function for MLE of Bernoulli parametrised by sigmoid is

$$J(\boldsymbol{\theta}) = -\log P(y \mid \boldsymbol{x}) = -\log \sigma((2y-1)z) = \zeta((1-2y)z)$$

Function saturates only when $(1-2y)z$ is very negative, i.e., when model has the right answer.

iii. Softmax Units: used to represent probability distribution over n different classes.
A linear layer predicts unnormalised log probabilities:

$$\boldsymbol{z} = \boldsymbol{W}^T\boldsymbol{h} + \boldsymbol{b}$$

where $z_i = \log \tilde{P}(y = i \mid \boldsymbol{x})$. Softmax function is then

$$\text{softmax}(\boldsymbol{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

The function is to maximise $\log P(y = i; \boldsymbol{z}) = \log \text{softmax}(\boldsymbol{z})_i = z_i - \log \sum_j \exp(z_j)$.
Note that the input z_i always has direct contribution to cost function, and the term cannot saturate.
Un-regularised maximum likelihood will drive the model to learn parameters that drive the softmax to predict fraction of counts for each outcome observed in the training set:

$$\text{softmax}(z(\boldsymbol{x}; \boldsymbol{\theta}))_i \approx \frac{\sum_{j=1}^{m} \mathbf{1}_{y^{(j)}=i, \boldsymbol{x}^{(j)}=x}}{\sum_{j=1}^{m} \mathbf{1}_{\boldsymbol{x}^{(j)}=x}}$$

Note that objective functions other than log-likelihood does not work well with softmax function. Squared error is poor loss function for softmax units, and can fail to train the model to change its output.
Softmax function can saturate, and many functions based on softmax also saturate, unless they are able to invert the saturating activating function.

iv. Other Output Types: generally, given a conditional distribution $p(\boldsymbol{y} \mid \boldsymbol{x};\ \boldsymbol{\theta})$, principle of maximum likelihood suggests using $-\log p(\boldsymbol{y} \mid \boldsymbol{x};\ \boldsymbol{\theta})$ as the cost function.
Neural networks represent a function $f(\boldsymbol{x}; \boldsymbol{\theta}) = \boldsymbol{\omega}$. The outputs of the function are not direct predictions of value $\boldsymbol{y}$, but the parameters for distribution over $\boldsymbol{y}$. The loss function is then $-\log p(\boldsymbol{y}; \boldsymbol{\omega}(\boldsymbol{x}))$.

Remark 2.1.4. *Learning Distribution Parameters*

i. Heteroscedastic Model: to predict different variance in $\boldsymbol{y}$ for different values of $\boldsymbol{x}$, formulate the Gaussian distribution using precision rather than variance. In multi-variate case, the diagonal precision matrix is used, $\text{diag}(\boldsymbol{\beta})$. The log-likelihood of Gaussian distribution parametrised by $\boldsymbol{\beta}$ involves only multiplication by β_i and addition of $\log \boldsymbol{\beta}_i$. The gradient of these operations are well-behaved.
Let $\boldsymbol{\alpha}$ be raw activation of the model used to determine diagonal precision. The softplus function may be used to obtain a positive precision vector $\boldsymbol{\beta} = \zeta(\boldsymbol{\alpha})$. Same strategy applies equally if using variance variance or standard deviation rather than precision.
If covariance is full and conditional, then parametrisation must be chosen that guarantees positive-definiteness of predicted covariance matrix.

$$\boldsymbol{\Sigma}(\boldsymbol{x}) = \boldsymbol{B}(\boldsymbol{x})\boldsymbol{B}^T(\boldsymbol{x})$$

where $\boldsymbol{B}$ is unconstrained square matrix. Note that if matrix is full rank, then computing likelihood requires $O(d^3)$ a $d \times d$ matrix for the determinant and inverse of $\boldsymbol{\Sigma}(\boldsymbol{x})$.

ii. Mixture Density Networks: to perform multimodal regression (predict real values that come from conditional distribution $p(\boldsymbol{y} \mid \boldsymbol{x})$) that can have several different peaks in $\boldsymbol{y}$ for the same $\boldsymbol{x}$.

$$p(\boldsymbol{y} \mid \boldsymbol{x}) = \sum_{i=1}^{n} p(c = i \mid \boldsymbol{x}) \mathcal{N}(\boldsymbol{y}; \boldsymbol{\mu}^{(i)}(\boldsymbol{x}), \boldsymbol{\Sigma}^{(i)}(\boldsymbol{x}))$$

The neural network will have three outputs:

1. Mixture components $p(c = i \mid \boldsymbol{x})$, forming a multinoulli distribution over n different components with latent variable c, obtained by softmax over n-dimensional vector.
2. Means $\boldsymbol{\mu}^{(i)}(\boldsymbol{x})$
3. Covariances $\boldsymbol{\Sigma}^{(i)}(\boldsymbol{x})$

Gradient-based optimisation of conditional Gaussian mixtures can be unreliable as the divisions can be numerically unstable. May be solved by clipping gradients, or scaling gradients heuristically.

Definition 2.1.5. *Hidden Units*

Even if hidden units are not differentiable at all input points, gradient descent still performs well enough as the training algorithms do not usually arrive at local minimum of cost function, but reduce its value significantly.
Most hidden units has input vector $\boldsymbol{x}$, computes an affine transformation $\boldsymbol{z} = \boldsymbol{W}^T\boldsymbol{x} + \boldsymbol{b}$, then applying element-wise non-linear function $g(\boldsymbol{z})$.

i. Rectified Linear Units (ReLUs): uses activation function $g(z) = \max\{0, z\}$.
Typically used on top of an affine transformation $\boldsymbol{h} = g(\boldsymbol{W}^T\boldsymbol{x} + \boldsymbol{b})$. In initialisation, set all elements of $\boldsymbol{b}$ to small positive values, so that ReLUs will be initially active for most inputs in training set.
Generalisations of ReLUs have non-zero slope α_i for $z_i < 0$: $h_i = \max(0, z_i) + \alpha_i \min(0, z_i)$.

1. Absolute Value Rectification: sets $\alpha_i = -1$ to obtain $g(z) = |z|$. Used for object recognition from images, to seek features that are invariant under polarity reversal of input illumination.
2. Leaky ReLU: fixes α_i to small positive value
3. Parametric ReU: treats α_i as a learnable parameter

ii. Maxout Units: divide $\boldsymbol{z}$ into groups of k values. Each maxout unit then outputs maximum element of one of these groups: $g(\boldsymbol{z})_i = \max_{j \in \mathbb{G}^{(i)}} z_j$, where $\mathbb{G}^{(i)}$ is indices of inputs for group i, which is $\{(i-1)k+1, \ldots, ik\}$.

This allows learning of piecewise linear function that responds to multiple directions in input $\boldsymbol{x}$ space .
Maxout units learn the activation function itself. With large k, maxout unit can learn to approximate any convex function. Each maxout unit is parametrised by k weight vectors, hence need more regularisation than ReLUs. Benefits include:

1. Can work well without regularisation if training set is large and number of pieces per unit is low.
2. Can gain statistical and computation advantages by requiring fewer parameters.
3. Have redundancy that resists catastrophic forgetting, where neural networks forgot how to perform tasks that were trained on in the past.

iii. Logistic Sigmoid and Hyperbolic Tangent: the logistic sigmoid activation function is $g(z) = \sigma(z)$, and the hyperbolic tangent activation function is $g(z) = \tanh(z)$. Note that $\tanh(z) = 2\sigma(2z) - 1$.
Sigmoidal units saturate across most of the domain, which makes gradient-based learning very difficult. Hence the use in hidden units in feedforward networks is now discouraged.
If sigmoidal activation function must be used, hyperbolic tangent activation function performs better, as it resembles identity function more closely. Training $\hat{y} = \boldsymbol{w}^T \tanh(\boldsymbol{U}^T \tanh(\boldsymbol{V}^T \boldsymbol{x}))$ resembles training a linear model $\hat{y} = \boldsymbol{w}^T \boldsymbol{U}^T \boldsymbol{V}^T \boldsymbol{x}$ as long as the activations of the network can be kept small.
Sigmoidal functions are more common in recurrent networks, probabilities models, auto-encoders

iv. Linear Unit: consider neural network layer with n inputs, p outputs, $\boldsymbol{h} = g(\boldsymbol{W}^T \boldsymbol{x} + \boldsymbol{b})$. Replace with two layers, one using weight matrix $\boldsymbol{U}$ and the other using weight matrix $\boldsymbol{V}$. If first layer has no activation function, then the factored approach is to compute $\boldsymbol{h} = g(\boldsymbol{V}^T \boldsymbol{U}^T \boldsymbol{x} + \boldsymbol{b})$. If $\boldsymbol{U}$ produces q outputs, then $\boldsymbol{U}$ and $\boldsymbol{V}$ together only contains $(n + p)q$ parameters, while $\boldsymbol{W}$ contains np parameters. For small q, this is considerable saving in parameters, while the cost is constraining linear transformation to be low rank. This is an efficient way of reducing number of parameters in the model.

v. Softmax: naturally represent probability distribution over discrete variable with k possible values. Used only in more advanced architectures that explicitly learn to manipulate memory.

vi. Radial Basis Function (RBF): function becomes more active as $\boldsymbol{x}$ approaches a template $\boldsymbol{W}_{:,i}$. As it saturates to 0 for most $\boldsymbol{x}$, it can be difficult to optimise.

$$h_t = \exp\left(-\frac{1}{\sigma_i^2}\|\boldsymbol{W}_{:,i} - \boldsymbol{x}\|^2\right)$$

vii. Softplus: smooth version of rectifier for functional approximation and for conditional distributions of undirected probabilistic models. Usage is generally discouraged.

$$g(a) = \zeta(a) = \log(1 + e^a)$$

viii. Hard tanh: shaped similarly to tanh but bounded.

$$g(a) = \max(-1, \min(1, a))$$

Theorem 2.1.6. *Universal Approximation Theorem*
A feedforward network with linear output layer and at least one hidden layer with any 'squashing' activation function can approximate any Borel measurable function from one finite-dimensional space to another with any desired non-zero amount of error, provided the network is given enough hidden units.

Method 2.1.7. *Architecture Design*
Neural network layers are arranged in a chain structure, with each layer being a function of preceding layer.

$$\boldsymbol{h}^{(1)} = g^{(1)}(\boldsymbol{W}^{(1)T} \boldsymbol{x} + \boldsymbol{b}^{(1)})$$
$$\boldsymbol{h}^{(k)} = g^{(k)}(\boldsymbol{W}^{(k)T} \boldsymbol{h}^{(k-1)} + \boldsymbol{b}^{(k)}), \quad k \geq 2$$

The main considerations are the depth of network and width of each layer. Deeper networks use far fewer units per layers and far fewer parameters, and often generalise to the test set, but are harder to optimise.

Remark 2.1.8. *Depth of Network and Universal Approximation Theorem*
A feedforward network with single layer is sufficient to represent any function, but the layer may be infeasibly large and fail to learn and generalise correctly. Using deeper models can reduce number of units required to represent the desired function and can reduce generalisation error.
Shallow networks with broad family of non-polynomial activation functions have universal approximation properties. Piecewise linear networks can represent functions with number of regions that is exponential to depth. The number of linear regions carved out by deep rectifier network with d inputs, depth l, and n units per hidden layer is $O(\binom{n}{d}^{d(d-1)} n^d)$. For maxout networks with k filters per unit, this is $O(k^{(k-1)+d})$.

Remark 2.1.9. *Backpropagation*
To calculate the gradient of loss function with respect to each of individual parameters of the neural network. Model training begins with random initialisation of weights and biases.

i. Forward pass: input is sampled from training data. Nodes receive input vector and passes their value

(multiplied by random initial weight) to nodes of first hidden layer. The hidden units take weighted sum of these output values as an input to the activation function, whose output is used for next hidden layer.

ii. Error computation: the final output of the network is compared to the ground truth, difference is calculated for the error value.

iii. Backwards pass: the error value computed earlier is used to compute the gradient of loss function. The gradient is then propagated back through the network, and the weights are updated according to their contribution to the error. The learning rate determines the size of weight updates.

iv. Weights update: the weights are updated in opposite direction of the gradient

Algorithm 1: Backpropagation Learning Algorithm

Require:
A set of training examples $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^{n}$
A multilayer network with L layers, weights w_{ij}^{l}, and activation function f
Loss function $J(y, o)$
Learning rate $0 < \alpha < 1$
Number of epochs *epochs*.

Initialize all weights:
for w_{ij}^{l} in the network **do**
  $w_{ij}^{l} \leftarrow$ small random number
end for

for $i = 1$ to *epochs* **do**
  for all training examples $(x, y) \in \mathcal{D}$ **do**

    Propagate the inputs forward to compute the outputs
    for all neurons i in the input layer **do**
      $a_i^0 \leftarrow x_i$
    end for

    for $l = 2$ to L **do**
      for all neurons i in layer l **do**
        $z_i^l \leftarrow \sum_j w_{ij}^l \, a_j^{l-1}$
        $a_i^l \leftarrow f(z_i^l)$
      end for
    end for

    Propagate deltas backward from output layer to input layer
    for all neurone i in the output layer **do**
      Let $o_i \equiv a_i^L$ *be output of neurone* i
      $\delta_i^L \leftarrow \dfrac{\partial J(y_i, o_i)}{\partial o_i} f'(z_i^L)$
    end for

    for $l = L - 1$ to 1 **do**
      for all neurons i in layer l **do**
        $\delta_i^l \leftarrow f'(z_i^l) \sum_j (w_{ji}^{l+1} \, \delta_j^{l+1})$
      end for
    end for

    Update weights using the deltas
    for all w_{ij}^{l} in the network **do**
      $w_{ij}^l \leftarrow w_{ij}^l - \alpha \, \delta_i^l \, a_j^{l-1}$
    end for
  end for
end for

Example 2.1.10. *Backpropagation Computation Example*

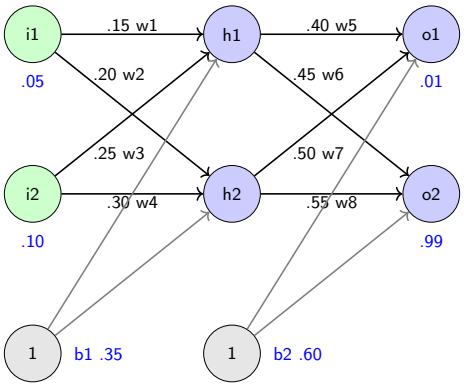


Figure 233: Backpropagation Example

For the forward pass, compute the total net input to each hidden layer neurone, squash total net input with activation function (logistic function), then repeat process with output layer neurones.
The total net input for h_1 is:

$$\begin{aligned}\text{net}_{h1} &= w_1 \cdot i_1 + w_2 \cdot i_2 + b_1 \cdot 1 \\ &= 0.15 \cdot 0.05 + 0.2 \cdot 0.1 + 0.35 \cdot 1 = 0.3775\end{aligned}$$

Squash with logistic function to get output of h_2:

$$\text{out}_{h1} = 1/(1 + e^{-\text{net}_{h1}}) = 1/(1 + e^{-0.3775}) = 0.593269992$$

For h_2, repeating the same process, the output is then $\text{out}_{h2} = 0.596884378$.
Repeating the process for outer layer neurones,

$$\begin{aligned}\text{net}_{o1} &= w_5 \cdot \text{out}_{h1} + w_6 \cdot \text{out}_{h2} + b_2 \cdot 1 \\ &= 0.4 \cdot 0.593269992 + 0.45 \cdot 0.596884378 + 0.6 \cdot 1 = 1.105905967 \\ \text{out}_{o1} &= 1/(1 + e^{-\text{net}_{o1}}) = 1/(1 + e^{-1.105905967}) = 0.75136507\end{aligned}$$

For o_2, repeating the same process, the output is then $\text{out}_{o2} = 0.772928465$.
The total error for each output neurone may be computed using squared error function $E_{\text{total}} = \sum \frac{1}{2}(Y - \hat{Y})^2$.
For o_1, the target output is 0.01, neurone output is 0.75136507, hence error is

$$E_{o1} = \frac{1}{2}(0.01 - 0.75136507)^2 = 0.274811083$$

Repeating the process for o_2, the total error for neural network is then

$$E_{\text{total}} = E_{o1} + E_{o2} = 0.274811083 + 0.023560026 = 0.298371109$$

For the backpropagation step, update each of the weights in network to reduce the error for each output neurone.

$$\begin{aligned}E_{\text{total}} &= \frac{1}{2}(Y_{o1} - \hat{Y}_{o1})^2 + \frac{1}{2}(Y_{o2} - \hat{Y}_{o2})^2 \\ \frac{\partial E_{\text{total}}}{\partial \text{out}_{o1}} &= 2 \cdot \frac{1}{2}(Y_{o1} - \hat{Y}_{o1}) \cdot -1 + 0 = -(0.01 - 0.75136507) = 0.74136507\end{aligned}$$

Also, note that

$$\begin{aligned}\text{out}_{o1} &= 1/(1 + e^{-\text{net}_{o1}}) \\ \frac{\partial \text{out}_{o1}}{\partial \text{net}_{o1}} &= \text{out}_{o1}(1 - \text{out}_{o1}) = 0.75136507(1 - 0.75136507) = 0.186815602 \\ \text{net}_{o1} &= w_5 \cdot \text{out}_{h1} + w_6 \cdot \text{out}_{h2} + b_2 \cdot 1 \\ \frac{\text{net}_{o1}}{w_5} &= 1 \cdot \text{out}_{h1} \cdot w_5^0 + 0 + 0 = \text{out}_{h1} = 0.593269992\end{aligned}$$

Putting it all together,

$$\begin{aligned}\frac{\partial E_{\text{total}}}{\partial w_5} &= \frac{\partial E_{\text{total}}}{\partial \text{out}_{o1}} \cdot \frac{\partial \text{out}_{o1}}{\partial \text{net}_{o1}} \cdot \frac{\partial \text{out}_{o1}}{\partial w_5} \\ &= -(Y_{o1} - \hat{Y}_{o1}) \cdot \text{out}_{o1}(1 - \text{out}_{o1}) \cdot textout_{o1} \\ &= 0.74136507 \cdot 0.186815602 \cdot 0.593269992 = 0.0821267041\end{aligned}$$

To decrease the error, subtract the value from current weight multiplied by learning rate (i.e., $\alpha = 0.5$).

$$w_5^+ = w_5 - \alpha \frac{\partial E_{\text{total}}}{\partial w_5} = 0.4 - 0.5 \cdot 0.0821267041$$

Repeating for the rest of the weights, we get

$$w_6^+ = 0.408666186, \quad w_7^+ = 0.511301270, \quad w_8^+ = 0.561370121$$

For the hidden layers, note that

$$\begin{aligned}\frac{\partial E_{\text{total}}}{\partial w_1} &= \frac{\partial E_{\text{total}}}{\partial \text{out}_{h1}} \cdot \frac{\partial \text{out}_{h1}}{\partial \text{net}_{h1}} \cdot \frac{\partial \text{net}_{h1}}{\partial w_1} \\ &= \left(\frac{\partial E_{o1}}{\partial \text{out}_{h1}} + \frac{\partial E_{o2}}{\partial \text{out}_{h1}}\right) \cdot \frac{\partial \text{out}_{h1}}{\partial \text{net}_{h1}} \cdot \frac{\partial \text{net}_{h1}}{\partial w_1}\end{aligned}$$

Note that out_{h1} affects both out_{o1} and out_{o2}.

$$\begin{aligned}\frac{\partial E_{o1}}{\partial \text{out}_{h1}} &= \frac{\partial E_{o1}}{\partial \text{net}_{h1}} \cdot \frac{\partial \text{net}_{h1}}{\partial \text{out}_{h1}} \\ \frac{\partial E_{o1}}{\partial \text{net}_{h1}} &= \frac{\partial E_{o1}}{\partial \text{out}_{h1}} \cdot \frac{\partial \text{out}_{o1}}{\partial \text{net}_{h1}} = 0.74136507 \cdot 0.186815602 = 0.138498562 \\ \frac{\partial \text{net}_{h1}}{\partial \text{out}_{h1}} &= w_5 = 0.4 \\ \frac{\partial E_{o1}}{\partial \text{out}_{h1}} &= 0.138498562 \cdot 0.4 = 0.055399425\end{aligned}$$

Following same process, $\frac{\partial E_{o2}}{\partial \text{out}_{h1}} = -0.019049119$. Hence

$$\frac{\partial E_{\text{total}}}{\partial \text{out}_{h1}} = \frac{\partial E_{o1}}{\partial \text{out}_{h1}} + \frac{\partial E_{o2}}{\partial \text{out}_{h1}} = 0.05539425 - 0.019049119 = 0.036350306$$

Next, figure out $\frac{\partial \text{out}_{h1}}{\partial \text{net}_{h1}}$ and $\frac{\partial \text{net}_{h1}}{\partial w_1}$.

$$\begin{aligned}\frac{\partial \text{out}_{h1}}{\partial \text{net}_{h1}} &= \text{out}_{h1} \cdot (1 - \text{out}_{h1}) = 0.59326999(1 - 0.59326999) = 0.2410.241300709 \\ \frac{\partial \text{net}_{h1}}{\partial w_1} &= i_1 = 0.05\end{aligned}$$

Putting it together,

$$\frac{\partial E_{\text{total}}}{\partial w_1} = \frac{\partial E_{\text{total}}}{\partial \text{out}_{h1}} \cdot \frac{\partial \text{out}_{h1}}{\partial \text{net}_{h1}} \cdot \frac{\partial \text{net}_{h1}}{\partial w_1} = 0.036350306 \cdot 0.241300709 \cdot 0.05 = 0.000438658$$

This is also written as

$$\frac{\partial E_{\text{total}}}{\partial w_1} = \left(\sum_o \frac{\partial E_{\text{total}}}{\partial \text{out}_o} \cdot \frac{\partial \text{out}_o}{\partial \text{net}_o} \cdot \frac{\partial \text{net}_o}{\partial \text{out}_{h_1}}\right) \cdot \frac{\partial \text{out}_{h1}}{\partial \text{net}_{h1}} \cdot \frac{\partial \text{net}_{h1}}{\partial w_1}$$

Next, update w_1 and rest of the weights

$$\begin{aligned}w_1^+ &= w_1 - \alpha \frac{\partial E_{\text{total}}}{\partial w_1} = 0.15 - 0.5 \cdot 0.000438658 = 0.149780716 \\ w_2^+ &= 0.19956143, \quad w_3^+ = 0.24975114, \quad w_4^+ = 0.29950229\end{aligned}$$

Hence all the weights has been updated. After the first round of backpropagation, total error has decreased to 0.291027924. Repeating the process $10,000$ times, the error reduces to 0.0000351085.

2.2 Optimization for Deep Learning

Remark 2.2.1. *Batch and Minibatch Algorithms*
Objective algorithms for machine learning computes each update to the parameters based on expected value of cost function using only a subset of terms of full cost function.
For maximum likelihood estimation problems, note that

$$\begin{aligned}\boldsymbol{\theta}_{\text{ML}} &= \arg\max_{\boldsymbol{\theta}} \sum_{i=1}^{m} \log p_{\text{model}}(\boldsymbol{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}) \\ J(\boldsymbol{\theta}) &= \mathbb{E}_{\boldsymbol{x}, y \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\boldsymbol{x}, y; \boldsymbol{\theta})\end{aligned}$$

The gradient is then

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \mathbb{E}_{\boldsymbol{x}, y \sim \hat{p}_{\text{data}}} \nabla_{\boldsymbol{\theta}} \log \log p_{\text{model}}(\boldsymbol{x}, y; \boldsymbol{\theta})$$

If the optimisation algorithm uses the entire training set, then it is a batch gradient method.
If a single sample is used at a time, then this is a stochastic/online gradient method.
For algorithms that uses more than one but less than all samples, then this is a minibatch gradient method.

Remark 2.2.2. *Challenges: Ill-Conditioning of Hessian Matrix H*
May cause stochastic gradient descent to be stuck where very small step increases the cost function.

Note that a gradient descent step of $-\epsilon \boldsymbol{g}$ will add

$$\frac{1}{2}\epsilon^2 \boldsymbol{g}^T \boldsymbol{H}\boldsymbol{g} - \epsilon \boldsymbol{g}^T \boldsymbol{g}$$

to the cost. This will be a problem when $\frac{1}{2}\epsilon^2 \boldsymbol{g}^T \boldsymbol{H}\boldsymbol{g} > \epsilon \boldsymbol{g}^T \boldsymbol{g}$.
Monitor squared gradient norm $\boldsymbol{g}^T \boldsymbol{g}$ and $\boldsymbol{g}^T \boldsymbol{H}\boldsymbol{g}$ term. When ill-conditioning happens, gradient norm does not shrink significantly but $\boldsymbol{g}^T \boldsymbol{H}\boldsymbol{g}$ term grows more than order of magnitude, and learning rate shrinks.

Remark 2.2.3. *Challenges: Local Minimums*
Neural networks and models with multiple parametrised latent variables have multiple local minima due to *model identifiability* problem. Model is identifiable if a sufficiently large training set can rule out all but one setting of model parameters. Models with latent variables are not identifiable as equivalent models are obtained by exchanging latent variables with each other (*weight space symmetry*). If there are m layers with n units each, then there are $n!^m$ ways to arrange the hidden units, making the model unidentifiable.
There may be an extremely large or uncountably infinite amount of local minima in a cost function. If the local minima have a high cost compared to global minimum, then this is problematic.

Remark 2.2.4. *Challenges: Plateaus, Saddle Points, Flat Regions*
For a function $f : \mathbb{R}^n \to \mathbb{R}$, expected ratio of saddle points to local minima grows exponentially with n. Hessian matrix at a saddle point has a mixture of positive and negative eigenvalues.
Degenerate regions with flat gradient may correspond to high value of objective function.

Remark 2.2.5. *Challenges: Cliffs and Exploding Gradients*
On extremely steep cliff structure, gradient update step may move parameters extremely far, losing most of the optimisation work done previously. The consequences may be mitigated using gradient clipping heuristics, which reduces the step size to be small enough.

Remark 2.2.6. *Challenges: Long-Term Dependencies*
Occurs when computational graph becomes very deep. Let a computational graph contain a path that repeatedly multiply a matrix $\boldsymbol{W}$. After t steps, this is $\boldsymbol{W}^t$. Let the eigen-decomposition of $\boldsymbol{W}$ be $\boldsymbol{W} = \boldsymbol{V}\text{diag}(\boldsymbol{\lambda})\boldsymbol{V}^{-1}$. Then we can see that $\boldsymbol{W}^t = (\boldsymbol{V}\text{diag}(\boldsymbol{\lambda})\boldsymbol{V}^{-1})^t = \boldsymbol{V}\text{diag}(\lambda)^t\boldsymbol{V}^{-1}$.
There will be a vanishing and exploding gradient problem if gradients on the graph are scaled according to $\text{diag}(\boldsymbol{\lambda})^t$. Note that recurrent networks uses same $\boldsymbol{W}$ at each time step, hence may face the problem.

Remark 2.2.7. *Challenges: Inexact Gradients*
Objective function for minimisation is intractable due to noisy/biased estimate of these quantities; in this case its gradient is intractable as well. Choose surrogate loss function that is easier to approximate than true loss.

Remark 2.2.8. *Challenges: Poor Correspondence between Local and Global Structure*
As gradient descent algorithms make small local moves, optimisation on local downhill moves may fail if local surface does not point toward the global solution. In higher dimensional space, learning algorithms may incur high costs or long training time to circumnavigate the manifold. Hence to find good initial points for problems.

Remark 2.2.9. *Stochastic Gradient Descent (SGD)*
Crucial parameter is the learning rate, which may be changed over time; at iteration k this is ϵ_k.
The SGD gradient estimation introduces a source of noise that does not vanish even at minimum. A sufficient condition to guarantee converge of SGD is

$$\sum_{k=1}^{\infty} \epsilon_k = \infty, \quad \sum_{k=1}^{\infty} \epsilon_k^2 < \infty$$

In practice, it is common to decay learning rate linearly until iteration τ:

$$\epsilon_k = (1-\alpha)\epsilon_0 + \alpha\epsilon_\tau, \quad \alpha = \frac{k}{\tau}$$

After iteration τ, leave ϵ constant.
Compute time per update does not growth with number of training examples.

Algorithm 2: Stochastic Gradient Descent (SGD) Update at Training Iteration k

Require:
Learning rate ϵ_k, Initial parameter $\boldsymbol{\theta}$

while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Compute gradient estimate $\hat{\boldsymbol{g}} \leftarrow +\frac{1}{m}\nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \epsilon\hat{\boldsymbol{g}}$

end while

Remark 2.2.10. *Momentum Algorithm*
Algorithm introduces variable $\boldsymbol{v}$ that is set to an exponentially decaying average of negative gradient.
Hyperparameter $\alpha \in [0, 1)$ determines how quickly the contributions of previous gradients exponentially decay.
The update rule is given by

$$\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \epsilon \nabla_{\boldsymbol{\theta}} \left(\frac{1}{m} \sum_{i=1}^{m} L(\boldsymbol{f}(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)}) \right)$$
$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$$

The larger α is relative to ϵ, the more previous gradients affect current direction.
If gradient is $\boldsymbol{g}$, it will accelerate in direction of $-\boldsymbol{g}$ until reaching terminal velocity where size of each step is

$$\frac{\epsilon \|\boldsymbol{g}\|}{1 - \alpha}$$

Algorithm 3: Stochastic Gradient Descent (SGD) with Momentum

Require:
Learning rate ϵ_k, Momentum parameter α

while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Compute gradient estimate $\boldsymbol{g} \leftarrow +\frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
Compute velocity update: $\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \epsilon \boldsymbol{g}$
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$
end while

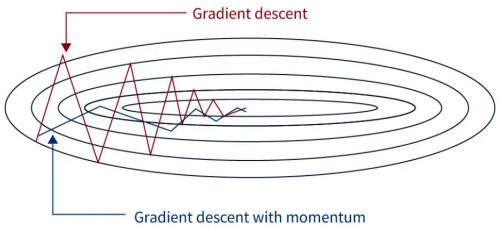


Figure 234: Stochastic Gradient Descent (SGD) Algorithm with and without Momentum

Remark 2.2.11. *Nesterov Momentum*
Inspired by Nesterov's accelerated gradient method, the update rules are:

$$\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \epsilon \nabla_{\boldsymbol{\theta}} \left(\frac{1}{m} \sum_{i=1}^{m} L(\boldsymbol{f}(\boldsymbol{x}^{(i)}; \boldsymbol{\theta} + \alpha \boldsymbol{v}), \boldsymbol{y}^{(i)}) \right)$$
$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$$

Note, gradient is evaluated after current velocity is applied.
Hence, Nesterov momentum is adding a correction factor to the standard method of momentum.

Algorithm 4: Stochastic Gradient Descent (SGD) with Nesterov Momentum

Require:
Learning rate ϵ_k, Momentum parameter α
Initial parameter $\boldsymbol{\theta}$, Initial velocity $\boldsymbol{v}$

while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Apply interim update: $\tilde{\boldsymbol{\theta}} \leftarrow \boldsymbol{\theta} + \alpha \boldsymbol{v}$
Compute gradient (at interim point): $\boldsymbol{g} \leftarrow \frac{1}{m} \nabla_{\tilde{\boldsymbol{\theta}}} \sum_i L(f(\boldsymbol{x}^{(i)}; \tilde{\boldsymbol{\theta}}), \boldsymbol{y}^{(i)})$
Compute velocity update: $\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \epsilon \boldsymbol{g}$
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$
end while

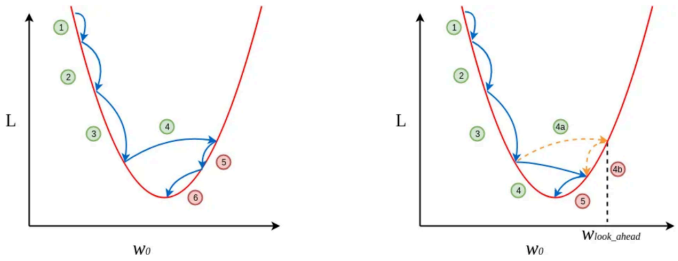


Figure 235: Momentum vs Nesterov Gradient Descent

Remark 2.2.12. *Parameter Initialisation Challenges*
Training deep neural networks is difficult as most algorithms are strongly affected by choice of initialisation.
Designing initialisation strategies is difficult as neural network optimisation is not yet well understood, as the nice properties from initialisation may not be preserved as training proceeds.
Some initial points may be beneficial for optimisation, but detrimental for generalisation.

Remark 2.2.13. *Parameter Initialisation Strategies*
Initial parameters need to 'break symmetry' between different units. May explicitly search for large set of basis functions mutually different from each other, but this may incur noticeable computational cost.
Random initialisation from high-entropy distribution over high-dimensional space is computationally cheaper.
Larger initial weights will yield stronger symmetry breaking effect and avoid redundant units. However, if weights are set too large, may result in chaos (extreme sensitivity to small perturbations of input). Exploding gradient may be mitigated by gradient clipping. The bias (i.e., parameters encoding conditional variance of a prediction) for each unit is set to heuristically chosen constraints.

Remark 2.2.14. *Weight Initialisation Heuristics*
Initialising parameters of neural network $\boldsymbol{\theta}$ to $\boldsymbol{\theta}_0$ is similar to imposing a Gaussian prior $p(\boldsymbol{\theta})$ with mean $\boldsymbol{\theta}_0$, hence it make sense to choose $\boldsymbol{\theta}_0$ near 0. Initialisation of $\boldsymbol{\theta}_0$ to large values means prior specifies which units should interact with each other, and how they should interact.
For a fully connected neural network with m inputs and n outputs, weights may be initialised by:

i. Sampling from $U\left(-\frac{1}{\sqrt{m}}, \frac{1}{\sqrt{m}}\right)$
ii. Sampling from normalised initialisation $W_{i,j} \sim U\left(-\sqrt{\frac{6}{m+n}}, \sqrt{\frac{6}{m+n}}\right)$ (Glorot and Bengio). Compromise between goal of initialising all layers to have same activation variance and that of same gradient variance. Derived from assumption that network is chain of matrix multiplications with no nonlinearities.
iii. Sampling from random orthogonal matrices, with scaling or gain factor g that accounts for nonlinearity at each layer, with specific values for different types of nonlinear activation functions. Assumes neural network is sequence of matrix multiplications without nonlinearities. Guarantees the total number of training iterations required to reach convergence is independent of depth.
Increasing scaling factor g pushes network toward regime where activations increase in norm in forward propagation, and gradients increase in norm in backward propagation.
For scaling rules that set all of initial weights to have same standard deviation, individual weight becomes extremely small when layers become large.
iv. Sparse initialisation (Martens), where each unit is initialised to have exactly k non-zero weight. Keep total number of input to unit independent from number of inputs m without making magnitude of individual weight elements shrink with m.
Imposes very strong prior on weights that are chosen to have large Gaussian values.

If computational resources allow, treat initial scale of weights for each layer as a hyper-parameter, and use hyper-parameter search algorithm. Choice of dense or sparse initialisation may also be a hyper-parameter.

Remark 2.2.15. *Bias Initialisation Heuristics*
Setting bias to zero is compatible with most weight initialisation schemes. The few other situations are:

i. If bias is for output unit, initialise bias to obtain right marginal statistics of the output. Assume initial weights are small enough that output of unit is determined only by bias. The bias will be inverse of activation function applied to marginal statistics of the output in training set.
If output is a highly skewed distribution over classes with marginal probability of class i given by element c_i of some vector $\boldsymbol{c}$, then bias vector $\boldsymbol{b}$ is such that $\text{softmax}(\boldsymbol{b}) = \boldsymbol{c}$. The models have output layers that should resemble input data $\boldsymbol{x}$, hence bias should match marginal distribution over $\boldsymbol{x}$.
ii. If bias is to be set to avoid causing too much saturation at initialisation (i.e., with ReLU). Not compatible with weight initialisation schemes that do not expect strong input form bias (i.e., with random walk init).

iii. If a unit controls whether other units are able to participate in a function. There is unit with output u and another unit $h \in [0, 1]$, multiplied together to produce output uh. Then h is a gate that determines if $uh \approx u$ or $uh \approx 0$. Hence, to set bias for h such that $h \approx 1$ at initialisation, else u does not learn.

Remark 2.2.16. *Variance or Precision Parameter Initialisation Heuristics*
For linear regression with conditional variance estimate using model

$$p(y \mid \boldsymbol{x}) = \mathcal{N}(y \mid \boldsymbol{w}^T \boldsymbol{x} + b, 1/\beta)$$

where β is a precision parameter. The variance of precision parameters may be initialised to 1.
Another approach is to assume initial weights are close enough to zero, and bias may be set while ignoring effect of the weights, then set bias to produce the correct marginal mean of the output, and set variance parameters to the marginal variance of the output in the training set.

Remark 2.2.17. *Model Parameter Initialisation with Machine Learning*
Initialise supervised model with parameters learned by unsupervised model trained on same inputs.
Alternatively, perform supervised training on a related task, which may yield initialisation that offers faster convergence than random initialisation. Information on distribution is encoded into initial parameters of model.

Remark 2.2.18. *Adaptive Gradient (AdaGrad) Algorithm*
Algorithm individually adapts learning rates of all model parameters by scaling them inversely proportional to square root of sum of all of their historical squared values. Has greater progress in more gently sloped directions of parameter space. On deep neural networks, accumulation of squared gradients from beginning of training can result in premature and excessive decrease in effective learning rate.

Algorithm 5: Adaptive Gradient (AdaGrad) Algorithm

Require:
Global Learning rate ϵ
Initial parameter $\boldsymbol{\theta}$
Small constant δ, perhaps 10^{-7} for numerical stability

Initialise gradient accumulation variable $\boldsymbol{r} = \boldsymbol{0}$
while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Compute gradient: $\boldsymbol{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
Accumulate squared gradient: $\boldsymbol{r} \leftarrow \boldsymbol{r} + \boldsymbol{g} \odot \boldsymbol{g}$
Compute update: $\Delta\boldsymbol{\theta} \leftarrow -\frac{\epsilon}{\delta + \sqrt{\boldsymbol{r}}} \odot \boldsymbol{g}$ (Division and square root applied element-wise)
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \Delta\boldsymbol{\theta}$
end while

Remark 2.2.19. *Root Mean Squared Propagation (RMSProp) Algorithm*
Modifies AdaGrad to perform better in non-convex setting by changing gradient accumulation into exponentially weighted moving average, which discards history from extreme past so it can converge rapidly after finding a convex bowl. Use of moving average introduces a new hyper-parameter ρ that controls length scale of the moving average. RMSProp is effective and practical for deep neural networks.

Algorithm 6: Root Mean Squared Propagation (RMSProp) Algorithm

Require:
Global Learning rate ϵ
Decay rate ρ
Initial parameter $\boldsymbol{\theta}$
Small constant δ, perhaps 10^{-6} to stabilise division by small numbers

Initialise accumulation variable $\boldsymbol{r} = \boldsymbol{0}$
while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Compute gradient: $\boldsymbol{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
Accumulate squared gradient: $\boldsymbol{r} \leftarrow \rho\boldsymbol{r} + (1 - \rho)\boldsymbol{g} \odot \boldsymbol{g}$
Compute update: $\Delta\boldsymbol{\theta} \leftarrow -\frac{\epsilon}{\sqrt{\delta + \boldsymbol{r}}} \odot \boldsymbol{g}$ ($\frac{\epsilon}{\sqrt{\delta + \boldsymbol{r}}}$ applied element-wise)
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \Delta\boldsymbol{\theta}$
end while

Algorithm 7: Root Mean Squared Propagation (RMSProp) Algorithm

Require:
Global Learning rate ϵ
Decay rate ρ
Momentum coefficient α
Initial parameter $\boldsymbol{\theta}$
Initial velocity $\boldsymbol{v}$

Initialise accumulation variable $\boldsymbol{r} = \boldsymbol{0}$
while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Compute interim update $\tilde{\boldsymbol{\theta}} \leftarrow \boldsymbol{\theta} + \alpha \boldsymbol{v}$
Compute gradient: $\boldsymbol{g} \leftarrow \frac{1}{m} \nabla_{\tilde{\boldsymbol{\theta}}} \sum_i L(f(\boldsymbol{x}^{(i)}; \tilde{\boldsymbol{\theta}}), \boldsymbol{y}^{(i)})$
Accumulate gradient: $\boldsymbol{r} \leftarrow \rho \boldsymbol{r} + (1 - \rho) \boldsymbol{g} \odot \boldsymbol{g}$
Compute velocity update: $\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \frac{\epsilon}{\sqrt{\boldsymbol{r}}} \odot \boldsymbol{g}$ ($\frac{1}{\sqrt{\boldsymbol{r}}}$ applied element-wise)
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$
end while

Remark 2.2.20. *Adaptive Moment Estimation (Adam) Algorithm*
Variant on combination of RMSProp and momentum. Note, momentum is incorporated directly as an estimate of first order moment of gradient; bias corrections are made to estimates of both first-order moments and second-order moments to account for initialisation at origin.
Algorithm is fairly robust to choice of hyper-parameters.

Algorithm 8: Adaptive Moment Estimation (Adam) Algorithm

Require:
Step size ϵ (Suggested default: 0.001)
Exponential decay rates for moment estimates ρ_1, ρ_2 in $[0, 1)$ (Suggested defaults: 0.9 and 0.999)
Small constant δ used for numerical stabilisation (Suggested default: 10^{-8})
Initial parameters $\boldsymbol{\theta}$

Initialise 1st and 2nd momentum variables $\boldsymbol{s} = \boldsymbol{0}$, $\boldsymbol{r} = \boldsymbol{0}$
Initialise time step $t = 0$
while stopping criterion not met **do**
Sample a minibatch of m samples from the training set $\{\boldsymbol{x}^{(1)}, \ldots, \boldsymbol{x}^{(m)}\}$ with targets $\boldsymbol{y}^{(i)}$
Compute gradient: $\boldsymbol{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
Update time step $t \leftarrow t + 1$
Update biased first and second moment estimate: $\boldsymbol{s} \leftarrow \rho_1 \boldsymbol{s} + (1 - \rho_1) \boldsymbol{g}$ and $\boldsymbol{r} \leftarrow \rho_2 \boldsymbol{r} + (1 - \rho_2) \boldsymbol{g} \odot \boldsymbol{g}$
Correct bias in first and second moment: $\hat{\boldsymbol{s}} \leftarrow \frac{\boldsymbol{s}}{1 - \rho_1^t}$ and $\hat{\boldsymbol{r}} \leftarrow \frac{\boldsymbol{r}}{1 - \rho_2^t}$
Compute update: $\Delta \boldsymbol{\theta} = -\epsilon \frac{\hat{\boldsymbol{s}}}{\sqrt{\hat{\boldsymbol{r}} + \delta}}$ (Operations applied element-wise)
Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \Delta \boldsymbol{\theta}$
end while

Remark 2.2.21. *Newton's Method*
Based on second-order Taylor's series approximating $J(\boldsymbol{\theta})$ near the point $\boldsymbol{\theta}_0$, ignoring higher derivatives

$$J(\boldsymbol{\theta}) \approx J(\boldsymbol{\theta}_0) + (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^T \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)^T \boldsymbol{H} (\boldsymbol{\theta} - \boldsymbol{\theta}_0)$$

where $\boldsymbol{H}$ is the Hessian of J with respect to $\boldsymbol{\theta}$ evaluated at $\boldsymbol{\theta}_0$.
Solving for critical point will get the Newton parameter update rule

$$\boldsymbol{\theta}^* = \boldsymbol{\theta}_0 - \boldsymbol{H}^{-1} \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0)$$

Note that for deep learning, the surface of objective function is non-convex. Regularising the Hessian will solve many of the issues such as saddle points etc.

$$\boldsymbol{\theta}^* = \boldsymbol{\theta}_0 - [H(f(\boldsymbol{\theta}_0)) + \alpha \boldsymbol{I}]^{-1} \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_0)$$

Due to requirement for inversion of the $k \times k$ matrix with time complexity $O(k^3)$ at every training iteration, only networks with very small number of parameters may be trained with this method.

Algorithm 9: Newton's Method with Objective $J(\boldsymbol{\theta}) = \frac{1}{m} \sum_{i=1}^{m} L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), y^{(i)})$

Require:
Initial parameter $\boldsymbol{\theta}_0$
Training set of m examples

while stopping criterion not met **do**
  Compute gradient $\boldsymbol{g} \leftarrow \frac{1}{m}\nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
  Compute Hessian: $\boldsymbol{H} \leftarrow \frac{1}{m}\nabla^2_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
  Compute Hessian Inverse: $\boldsymbol{H}^{-1}$
  Compute update: $\Delta\boldsymbol{\theta} = -\boldsymbol{H}^{-1}\boldsymbol{g}$
  Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \Delta\boldsymbol{\theta}$
end while

Remark 2.2.22. *Conjugate Gradient Method*
Method effectively avoids computation of inversion Hessian by descending conjugate directions, which is a search direction conjugate to the previous line search direction (will not undo progress made in that direction).
Let current and previous search directions be $\boldsymbol{d}_t, \boldsymbol{d}_{t-1}$. At training iteration t, the next search direction $\boldsymbol{d}_t$ is

$$\boldsymbol{d}_t = \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) + \beta_t \boldsymbol{d}_{t-1}$$

note that β_t may be computed via Fletcher Reeves:

$$\beta_t = \frac{\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_t)^T \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_t)}{\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_{t-1})^T \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_{t-1})}$$

or may be computed with Polak-Ribiere:

$$\beta_t = \frac{(\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_t) - \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_{t-1}))^T \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_t)}{\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_{t-1})^T \nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}_{t-1})}$$

The non-linear conjugate gradient method is a variation which uses occasional resets where the method is restarted with line search along with the unaltered gradient.

Algorithm 10: Conjugate Gradient Method

Require:
  Initial parameter $\boldsymbol{\theta}_0$
  Training set of m examples
  Initialise $\boldsymbol{\rho}_0 = 0$
  Initialise $g_0 = 0$
  Initialise $t = 1$

while stopping criterion not met **do**
  Initialise the gradient $\boldsymbol{g}_t = \boldsymbol{0}$
  Compute gradient $\boldsymbol{g} \leftarrow \frac{1}{m}\nabla_{\boldsymbol{\theta}} \sum_i L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}), \boldsymbol{y}^{(i)})$
  Compute Polak-Ribiere metric: $\beta_t = \frac{(\boldsymbol{g}_t - \boldsymbol{g}_{t-1})^T \boldsymbol{g}_t}{\boldsymbol{g}_{t-1}^T \boldsymbol{g}_{t-1}}$
  (Nonlinear conjugate gradient: optionally reset $\beta_t = 0$, i.e., if t is a multiple of some constant k)
  Compute search direction: $\boldsymbol{\rho}_t = -\boldsymbol{g}_t + \beta_t \boldsymbol{\rho}_{t-1}$
  Perform line search to find: $\epsilon^* = \arg\min_\epsilon \frac{1}{m} \sum_{i=1}^m L(f(\boldsymbol{x}^{(i)}; \boldsymbol{\theta}_t + \epsilon\boldsymbol{\rho}_t), \boldsymbol{y}^{(i)})$
  (On truly quadratic cost function, analytically solve for ϵ^*)
  Apply update: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \epsilon^\theta \boldsymbol{\rho}_t$
  Apply time step $t \leftarrow t + 1$
end while

Remark 2.2.23. *Nonlinear Conjugate Gradients*
Method of conjugate gradients is still applicable for non-linear objective function.
Conjugate directions are no longer assumed to remain at minimum of the objective for previous directions. Has occasional resets where method of conjugate gradients is restarted with line search along unaltered gradient.

Remark 2.2.24. *Broyden-Fletcher-Goldfarb-Shanno (BGFS) Algorithm*
Similar to conjugate gradient, but uses more direct approach to approximation of newton's update.
Method approximate the inverse of hessian $\boldsymbol{H}$ with matrix $\boldsymbol{M}_t$, iteratively refined by low rank updates to better approximate $\boldsymbol{H}^{-1}$. Direction of gradient descent is then determined by $\boldsymbol{\rho}_t = \boldsymbol{M}_t \boldsymbol{g}_t$. Line search is performed in this direction to determine step size ϵ^*. Final update of parameter is then $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \epsilon^* \boldsymbol{\rho}_t$.
Note, BGFS store matrix $\boldsymbol{M}$ which requires $O(n^2)$ memory, making it impractical for most modern models.

Remark 2.2.25. *Limited Memory BFGS (L-BFGS) Algorithm*
Memory costs of BFGS may be decreased by not storing complete inverse Hessian approximation $\boldsymbol{M}$. L-BFGS computes approximation $\boldsymbol{M}$ with same method as BFGS, but assumes at beginning that $\boldsymbol{M}^{(t-1)}$ is the identity matrix, rather than storing the approximation between steps.

If used with exact line searches, directions defined by L-BFGS are mutually conjugate, and this procedure remains well-behaved when minimum of line search is reached only approximately.
L-BFGS may be generalised to store some vectors used to update $\boldsymbol{M}$ at each time step, costing $O(n)$ per step.

Remark 2.2.26. *Meta-Algorithm: Batch Normalisation*
An update from optimisation algorithm may result in multiple-order effects, and in very deep networks, even higher-order interactions can be significant.
Batch normalisation provides an elegant way of re-parametrising any deep network, which reduces the problem of coordinating updates across multiple layers. Batch normalisation may be applied to any input/hidden layer. Let $\boldsymbol{H}$ be minibatch of activations of layer to normalise, with activations for each example appearing in a row. To normalise $\boldsymbol{H}$, replace it with $\boldsymbol{H}' = \frac{\boldsymbol{H}-\boldsymbol{\mu}}{\boldsymbol{\sigma}}$ where $\boldsymbol{\mu}$ is a vector of mean of each unit and σ is a vector of standard deviation of each unit. Then H_{ij} is normalised by subtracting μ_j and dividing σ_j. Rest of network then operates on $\boldsymbol{H}'$ just as in original network on $\boldsymbol{H}$.
At training, the mean and standard deviations are computed as

$$\boldsymbol{\mu} = \frac{1}{m}\sum_i \boldsymbol{H}_{i,:}, \quad \boldsymbol{\sigma} = \sqrt{\delta + \frac{1}{m}\sum_i (\boldsymbol{H}-\boldsymbol{\mu})_i^2}$$

where δ is a small positive value (i.e., 10^{-8}) to avoid undefined gradient of $\sqrt{z}$ at $z = 0$.
These are then back-propagated to normalise $\boldsymbol{H}$, hence gradient will never increase the standard deviation or mean of h_i. Batch normalisation re-parametrises the model to make some units always standardised.
At test time, μ and σ may be replaced by running averages from training time. This allows model to be evaluated on a single example, without needing to use definitions of μ and $\boldsymbol{\sigma}$ that depends on entire minibatch.

Remark 2.2.27. *Batch Normalisation but Maintain Network Express Power*
Replace batch of hidden activations $\boldsymbol{H}$ with $\boldsymbol{\gamma}\boldsymbol{H}' + \boldsymbol{\beta}$ rather than $\boldsymbol{H}'$, where $\boldsymbol{\gamma}$ and $\boldsymbol{\beta}$ are learned parameters to allow new variable to have any mean and standard deviation. With original $\boldsymbol{H}'$, mean of $\boldsymbol{H}$ was determined by complicated interaction between parameters in layers below $\boldsymbol{H}$. With $\boldsymbol{\gamma}\boldsymbol{H}' + \boldsymbol{\beta}$, the mean is determined solely by $\boldsymbol{\beta}$, which is much easier to learn with gradient descent.
Most neural network layers are of form $\phi(\boldsymbol{X}\boldsymbol{W} + \boldsymbol{b})$, where ϕ is some fixed nonlinear activation function. Then $\boldsymbol{X}\boldsymbol{W} + \boldsymbol{b})$ should be replaced by normalised version of $\boldsymbol{X}\boldsymbol{W}$, and bias term omitted as it becomes redundant with β parameter applied from batch normalisation re-parametrisation. Statistics of input are non-Gaussian and less amenable to standardisation by linear operations.

Remark 2.2.28. *Meta-Algorithm: Coordinate Descent*
Solving optimisation problems by minimising the objective function $f(\boldsymbol{x})$ with respect to a single variable x_i, and rapidly cycling through all variables. *Block coordinate descent* refers to minimising with respect to a subset of variables simultaneously.
Coordinate descent works when different variables can be clearly separated into groups that play relatively isolated roles, or when optimisation with respect to one group of variables is significantly more efficient than optimisation with respect to all of the variables.

Remark 2.2.29. *Polyak Averaging*
Average several points in trajectory through parameter space visited by optimisation algorithm.
If t iterations of gradient descent visit points $\boldsymbol{\theta}^{(1)}, \dots, \boldsymbol{\theta}^{(t)}$, then the output is $\hat{\boldsymbol{\theta}}^{(t)} = \frac{1}{t}\sum_i \boldsymbol{\theta}^{(i)}$.
The optimisation algorithm may leap back and forth across a valley several times without visiting the minimum point, and the average of all points on either side of the valley should be close to minimum value.
This is developed solely for stochastic gradient descent (SGD) to stabilise the estimates.
For non-convex problems, an exponentially decaying running average is used:

$$\hat{\boldsymbol{\theta}}^{(t)} = \alpha\hat{\boldsymbol{\theta}}^{(t-1)} + (1-\alpha)\boldsymbol{\theta}^{(t)}$$

Remark 2.2.30. *Continuation Methods*
Family of strategies that can make optimisation easier by choosing initial points to ensure that local optimisation spends most of its time in well-behaved regions. Method constructs series of objective functions $\{J^{(0)}, \dots, J^{(n)}\}$ which are increasingly difficult to minimise (well-behaved over less regions of $\boldsymbol{\theta}$ space) for the true cost function $J(\boldsymbol{\theta})$. The cost functions are designed such that a solution to one is a good initial point for the next.
Continuation methods aim to reach a global minimum despite many local minima. Easier cost functions are constructed through the blurring operation, which approximates

$$J^{(i)}(\boldsymbol{\theta}) = \mathbb{E}_{\theta' \sim \mathcal{N}(\boldsymbol{\theta}', \boldsymbol{\theta}, \sigma^{(i)2})} J(\boldsymbol{\theta})'$$

via sampling. Some non-convex functions become approximately convex when blurred, but preserved enough information about local global minimum. However, approach might not work when

i. so many incremental cost functions are required that cost of entire procedure remains high;

ii. NP-hard optimisation problems remain NP-hard, even when continuation methods are applicable;
iii. function might not become convex, no matter how much it is blurred;
iv. function may become convex due to blurring, but minimum of this blurred function may track a local rather than global minimum of original cost function.

Remark 2.2.31. *Curriculum Learning*
Planning a learning process to begin by learning simple concepts, and progress to learning more complex topics that depend on these simpler concepts. Earlier objective functions $J^{(i)}$ are made easier by increasing the influence of simpler examples.
Stochastic curriculum is a variation of the approach where a random mix of easy and difficult examples is always presented to the learner, but average proportion of more difficult examples is gradually increased.

2.3 Regularization and Generalization

Definition 2.3.1. *Regularisation* refers to adding a parameter norm penalty $\Omega(\boldsymbol{\theta})$ to the objective function J. The regularised objective function is then

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha\Omega(\boldsymbol{\theta})$$

where $\alpha in [0, \infty)$ is a hyper-parameter that weights the contribution of norm penalty term.
For neural networks, the parameter norm penalty Ω is chosen such that it penalises only the weights of the affine transformation at each layer and leaves the biases un-regularised.

Remark 2.3.2. *Norm Penalties as Constrained Optimisation*
Let cost function regularised by parameter norm penalty be

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha\Omega(\boldsymbol{\theta})$$

Construct a generalised generalised Lagrange function,

$$\mathcal{L}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha(\Omega(\boldsymbol{\theta}) - k)$$

The solution to the constrained problem is then

$$\boldsymbol{\theta}^* = \arg\min_{\boldsymbol{\theta}} \max_{\alpha, \alpha \geq 0} \mathcal{L}(\boldsymbol{\theta}, \alpha)$$

Note that optimal value α^* will shrink $\Omega(\boldsymbol{\theta})$, but not such that it is less than k.
Fixing α^*, the problem is then a function of $\boldsymbol{\theta}$,

$$\boldsymbol{\theta}^* = \arg\min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}, \alpha^*) = \arg\min_{\boldsymbol{\theta}} J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha^*\Omega(\boldsymbol{\theta})$$

This is the regularised training problem of minimising $\tilde{J}$. The parameter norm penalty is imposing a constraint on the weights. If explicit constraints are to be used rather than penalties, use stochastic gradient descent on $J(\boldsymbol{\theta})$, then take the projected $\boldsymbol{\theta}$ to the nearest point that satisfies $\Omega(\boldsymbol{\theta}) < k$. This method is used if k is predefined by user, and time is not to be spent on searching for α value that corresponds to this k.
Explicit constraints and re-projection work better in circumstances where non-convex optimisation is involved, and this may avoid getting stuck in local minima. Explicit constraints also impose stability, and together with high learning rate allows rapid exploration of parameter space.

Remark 2.3.3. *Dataset Augmentation*
Effective technique for object classification. Operations such as translation of images by few pixels in each direction can greatly improved generalisation, even if the model has been designed to be partially translation invariant using convolution and pooling techniques. Rotation and scaling of images are also quite effective.
For speech recognition tasks, noice injection will improve performance. To improve robustness, train models with random noise applied to inputs. Noise injection also works when applied to the hidden units (data augmentation at multiple levels of abstraction). Enabling dropout constructs new inputs by multiplying by noise.

Remark 2.3.4. *Noise Robustness*
For some models, addition of noise with infinitesimal variance at input of model is equivalent to imposing penalty on norm of the weights. Generally, noise injection can be much more powerful than simply shrinking the parameters, especially when noise is added to hidden units.
Noise may also be added to weights, which is primarily used in recurrent neural networks. This is a stochastic implementation of Bayesian inference over the weights (model weights are uncertain and representable via a probability distribution that reflects this uncertainty).
Adding noise to weights may also result in more stability in function to be learned. Let the cost function be

$$J = \mathbb{E}_{p(x,y)}[(\hat{y}(\boldsymbol{x}) - y)^2]$$

where $\hat{y}(\boldsymbol{x})$ is function to be trained that maps set of features $\boldsymbol{x}$ to a scalar using least-squares cost function J. Assume each input presentation include a random perturbation $\epsilon_{\boldsymbol{W}} \sim \mathcal{N}(\boldsymbol{\epsilon}; \boldsymbol{0}, \eta \boldsymbol{I})$ in the network weight of a standard l-layer MLP. Let the perturbed model be $\hat{y}_{\boldsymbol{\epsilon_W}}(\boldsymbol{x})$. With injection o noise, objective function becomes

$$\tilde{J}_{\boldsymbol{W}} = \mathbb{E}_{p(\boldsymbol{x},y,\boldsymbol{\epsilon_W})}\left[(\hat{y}_{\boldsymbol{\epsilon_W}}(\boldsymbol{x}) - y)^2\right] = \mathbb{E}_{p(\boldsymbol{x},y,\boldsymbol{\epsilon_W})}\left[\hat{y}^2_{\boldsymbol{\epsilon_W}}(\boldsymbol{x}) - 2y\hat{y}_{\boldsymbol{\epsilon_W}} + y^2\right]$$

For small η, the minimisation of J with added noise (with covariance $\eta \boldsymbol{I}$) is equivalent to minimisation of J with additional regularisation term $\eta \mathbb{E}_{p(\boldsymbol{x}),y}[\|\nabla_{\boldsymbol{W}} \hat{y}(\boldsymbol{x})\|^2]$. This encourages parameters to go to regions of parameter space where small perturbations of weights have relatively small influence on output.

Remark 2.3.5. *Label Smoothing*
Datasets may have some mistake in y labels. Regularise a model based on softmax with k output values by replacing hard 0 and 1 classification targets with targets of $\frac{\epsilon}{k-1}$ and $1-\epsilon$ respectively, assuming that the training set label y is correct with probability $1-\epsilon$. Cross-entropy can then be used on these soft targets.

Remark 2.3.6. *Semi-Supervised Learning*
Unlabelled examples $P(\boldsymbol{x})$ and $P(\boldsymbol{x}, \boldsymbol{y})$ are used to estimate $P(\boldsymbol{y} \mid \boldsymbol{x})$ or predict $\boldsymbol{y}$ from $\boldsymbol{x}$.
Construct a generative model of either $P(\boldsymbol{x})$ or $P(\boldsymbol{x}, \boldsymbol{y})$ which shares parameters with a discriminative model of $P(\boldsymbol{y} \mid \boldsymbol{x})$. The supervised criterion $\log P(\boldsymbol{y} \mid \boldsymbol{x})$ may then be traded with the unsupervised or generative one ($\log P(\boldsymbol{x})$ or $\log P(\boldsymbol{x}, \boldsymbol{y})$). The generative criterion then expresses a particular form of prior belief about the solution to the supervised learning problem (structure of $P(\boldsymbol{x})$ is connected to structure of $P(\boldsymbol{y} \mid \boldsymbol{x})$ such that is is captured by the shared parametrisation).

Remark 2.3.7. *Multi-Task Learning*
Improve generalisations by pooling examples (soft constraints imposed on parameters) from several tasks.
Different supervised tasks (predicting $\boldsymbol{y}^{(i)}$ given $\boldsymbol{x}$) share the same input $\boldsymbol{x}$, as well as some intermediate-level representation $\boldsymbol{h}^{(\text{shared})}$ capturing a common pool of factors. Model can be divided in two parts:

i. Task-specific parameters
ii. Generic parameters, shared across all the tasks

Improved generalisation and associated error bounds can be achieved due to shared parameters, but this only happens if some assumptions about statistical relationship between different tasks are valid.

Remark 2.3.8. *Early Stopping*
To obtain a model with better validation set error by returning to the parameter setting at the point in time with lowest validation set error. Model validation is done in parallel to model training.
Early stopping has the effect of restricting the optimisation procedure to a relatively small volume of parameter space in neighbourhood of initial parameter value $\boldsymbol{\theta}_o$. Assume τ optimisation steps are taken with learning rate ϵ. The product $\epsilon\tau$ is a measure of effective capacity. Assuming gradient is bounded, restricting $\epsilon\tau$ limits the volume of parameter space reachable from $\boldsymbol{\theta}_o$, hence $\epsilon\tau$ behaves as the reciprocal of weight decay coefficients.

Algorithm 11: Early Stopping Algorithm

Require:
Number of steps n between evaluations
Patience p, the number of times to observe worsening validation error before stopping
Initial parameters $\boldsymbol{\theta}_o$

$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta}_o$, $i \leftarrow 0$, $j \leftarrow 0$, $v \leftarrow \infty$, $\boldsymbol{\theta}^* \leftarrow \boldsymbol{\theta}$, $i^* \leftarrow i$
while $j < p$ **do**
  Update $\boldsymbol{\theta}$ by running training algorithm for n steps
  $i \leftarrow i + n$, $v' \leftarrow$ ValidationSetError($\boldsymbol{\theta}$)
  if v' < v **then**
    $j \leftarrow 0$, $\boldsymbol{\theta}^* \leftarrow \boldsymbol{\theta}$, $i^* \leftarrow i$, $v \leftarrow v'$
  end if
  $j \leftarrow j + 1$
end while
Best parameters are $\boldsymbol{\theta}^*$, best number of training steps is i^*

Remark 2.3.9. *Parameter Tying and Parameter Sharing*
Used to express prior knowledge on suitable values of model parameters (precise values are not known, but there should be some dependencies between parameters). To regularise parameters to be close, may use:

i. Parameter norm penalty: assume model A with parameters $\boldsymbol{w}^{(A)}$ and model B with parameters $\boldsymbol{w}^{(B)}$. Two models map input to two different but related outputs $\hat{y}^{(A)} = f(\boldsymbol{w}^{(A)}, \boldsymbol{x})$ and $\hat{y}^{(B)} = g(\boldsymbol{w}^{(B)}, \boldsymbol{x})$. The model parameters $w_i^{(A)}$ should be close to $w_i^{(B)}$ $\forall i$. Using regularisation, a parameter norm penalty may be imposed, such as using L^2 penalty to get $\Omega(\boldsymbol{w}^{(A)}, \boldsymbol{w}^{(B)}) = \left\|\boldsymbol{w}^{(A)} - \boldsymbol{w}^{(B)}\right\|_2^2$.

ii. Parameter sharing: using constraints. Only a subset of the parameters (input set) need to be stored in memory. May lead to significant reduction in memory footprint of the model (i.e., in CCNs)

Remark 2.3.10. *Dropout*
Dropout trains the ensemble consisting of all sub-networks that can be formed by removing non-output units from the base network by multiplying its output value by zero. To train a neural network with dropout, use minibatch-based learning algorithms (i.e., Stochastic Gradient Descent). Each time an example is loaded into minibatch, randomly (and independently) sample a different binary mask $\boldsymbol{\mu}$ to apply to all of input and hidden units in the network, where the probability of sampling a mask value of one is the dropout rate.
Let $J(\boldsymbol{\theta}, \boldsymbol{\mu})$ be cost of model defined by parameters $\boldsymbol{\theta}$ and mask $\boldsymbol{\mu}$. Dropout training is minimising $\mathbb{E}_{\boldsymbol{\mu}} J(\boldsymbol{\theta}, \boldsymbol{\mu})$. Each model inherits different subset of parameters from parent neural network, and a tiny fraction of possible sub-networks are trained for a single step. Parameter sharing allows sub-networks to arrive at good parameter settings. Each sub-network given by $\boldsymbol{\mu}$ defines a probability distribution $p(y \mid \boldsymbol{x}, \boldsymbol{\mu})$. The arithmetic mean over all masks is then given by

$$\sum_{\boldsymbol{\mu}} p(\boldsymbol{\mu}) p(y \mid \boldsymbol{x}, \boldsymbol{\mu})$$

where $p(\boldsymbol{\mu})$ is the distribution used to sample $\boldsymbol{\mu}$ during training.
Using dropout as a regularisation technique reduces effective capacity of a model. To offset this, increase the size of the model. Typically optimal validation error is much lower when using dropout, but this comes at cost of larger model and more iterations of training.
For very large datasets, costs of using dropout and larger models may outweigh benefits of regularisation. For sparse datasets (extremely few labelled training examples available), Bayesian neural networks outperform dropout. When additional unlabelled data is available, unsupervised learning may gain advantage over dropout.

Remark 2.3.11. *Adversarial Training*
Adversarial example may be constructed by optimisation to search for an input $\boldsymbol{x}'$ near $\boldsymbol{x}$ such that the model output is very different at $\boldsymbol{x}'$.
The primary cause is excessive linearity in some neural networks. Changing input by ϵ causes linear function with weights $\boldsymbol{w}$ to change by as much as $\epsilon \|\boldsymbol{w}\|_1$, which can be very large if $\boldsymbol{w}$ is high-dimensional.
Adversarial examples assists in semi-supervised learning. At point $\boldsymbol{x}$ that is not associated with a label, model assigns label $\hat{y}$ (may not be true label). Seek adversarial example $\boldsymbol{x}'$ that causes classifier to output label $y' \neq \hat{y}$. Classifier may then be trained to assign same label to $\boldsymbol{x}$ and $\boldsymbol{x}'$, learning function robust to small changes.

Remark 2.3.12. *Tangent Distance Algorithm*
Non-parametric nearest-neighbour algorithm, where examples on the same manifold share the same category. The nearest-neighbour distance between points $\boldsymbol{x}_1$ and $\boldsymbol{x}_2$ is the distance between manifolds M_1 and M_2. Computationally cheap alternative is to approximate M_i by its tangent plane at $\boldsymbol{x}_i$ and measure the distance between two tangents, or between a tangent plane and a point.

Remark 2.3.13. *Tangent Propagation Algorithm*
Trains neural net classifier with extra penalty to make each output $f(\boldsymbol{x})$ of neural net locally invariant to factors of variation, which correspond to movement along manifold which examples from same class concentrate. Local invariance is achieved by requiring $\nabla_{\mathbf{x}} f(\boldsymbol{x})$ to be orthogonal to known manifold tangent vectors $\boldsymbol{v}^{(i)}$ at $\boldsymbol{x}$, or equivalently the directional derivative of f at $\boldsymbol{x}$ in directions $\boldsymbol{v}^{(i)}$ be small by adding regularisation penalty:

$$\Omega(f) = \sum_i \left(\left(\nabla_{\boldsymbol{x}} f(\boldsymbol{x}) \right)^T \boldsymbol{v}^{(i)} \right)^2$$

Tangent vectors are derived a priori from effect of transformations (i.e., translation, rotation, scaling).
Note this regularises model to resist infinitesimal perturbation, and the infinitesimal approach poses difficulties for models based on rectified linear units.
Data augmentation is the non-infinitesimal version of Tangent Propagation Algorithm, where network is explicitly trained to correctly classify distinct inputs created from more than infinitesimal transformation.

Remark 2.3.14. *Double Back-propagation Algorithm*
Algorithm regularises Jacobian to be small. Require model to be invariant to all directions of change in input so long as the change is small. Adversarial training is the non-infinitesimal version of Double Back-propagation Algorithm, where inputs were found near original inputs and model is trained to produce same input.

Remark 2.3.15. *Manifold Tangent Classifier*
Eliminates need to know tangent vectors a priori. Auto-encoders may estimate manifold tangent vectors, and the classifier makes use of this technique to avoid requiring user-specified tangent vectors. The steps are:

i. Use auto-encoder to learn manifold structure by unsupervised learning
ii. Use these tangents to regularise a neural net classifier as in Tangent Propagation Algorithm.

2.4 Representation Learning

Remark 2.4.1. *Sparse Representations*
Representation sparsity is where elements of the representation are zero, accomplished by same sort of mechanisms used in parameter regularisation.
Given a linear regression $\boldsymbol{y} = \boldsymbol{B}\boldsymbol{h}$, where $\boldsymbol{y} \in \mathbb{R}^m, \boldsymbol{B} \in \mathbb{R}^{m \times n}, \boldsymbol{h} \in \mathbb{R}^n$, and $\boldsymbol{h}$ is the representation of data $\boldsymbol{x}$, add to loss function J a norm penalty on the representation ($\Omega(\boldsymbol{h})$):

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha\Omega(\boldsymbol{h})$$

where $\alpha \in [0, \infty)$ weights relative contribution of norm penalty, with larger values of α for more regularisation. L^1 penalty on elements of representation induces sparsity, $\Omega(\boldsymbol{h}) = \|\boldsymbol{h}\|_1 = \sum_i |h_i|$. For representations with elements constrained to lie on unit interval, may use Student-t prior or KL divergence penalties.
Other approaches such as hard constraint on activation values may also work. *Orthogonal matching pursuit* encodes an input $\boldsymbol{x}$ with representation $\boldsymbol{h}$ that solves the constrained optimisation problem

$$\arg \min_{\boldsymbol{h}, \|\boldsymbol{h}\|_0 < k} \|\boldsymbol{x} - \boldsymbol{W}\boldsymbol{h}\|^2$$

where $\|\boldsymbol{h}\|_0$ is the number of non-zero entries of $\boldsymbol{h}$. If $\boldsymbol{W}$ is constrained to be orthogonal, this can be solved efficiently with OMP-k method, where k is number of non-zero features allowed.

Example 2.4.2. *Sparse Encoding with Coordinate Descent*
Consider the cost function

$$J(\boldsymbol{H}, \boldsymbol{W}) = \sum_{i,j} |H_{ij}| + \sum_{i,j} (\boldsymbol{X} - \boldsymbol{W}^T\boldsymbol{H})^2_{i,j}$$

To find weight matrix $\boldsymbol{W}$ that can linearly decode a matrix of activation values $\boldsymbol{H}$ to reconstruct training set $\boldsymbol{X}$. Most applications also include weight decay or constraint on norms of columns of $\boldsymbol{W}$ to prevent the pathological solution with extremely small $\boldsymbol{H}$ and large $\boldsymbol{W}$.
Note J is not convex, but inputs may be divided into the dictionary parameters $\boldsymbol{W}$ and the code representations $\boldsymbol{H}$. Minimising object function with respect to either set is a convex problem.

2.5 Sequence Models

2.6 Generative Models

2.7 Modern Architecture and Systems View

Remark 2.7.1. *Supervised Pre-training*
Pre-training refers to training simple models on simple tasks before training the desired model on desired task. Greedy algorithms break problem into many components, then solve for optimal version of each component in isolation, which is must cheaper than algorithms that solve for best joint solution. Fine-tuning on joint optimisation algorithm searches for optimal solution to the full problem.
Greedy supervised training breaks supervised learning problems into other simpler supervised learning problems. Transfer learning refers to pre-training a deep network on a set of tasks, then initialise a same-size network with first k layers of first net. The layers of second networks (upper layers initialised randomly) are then jointly trained to perform a different set of tasks with fewer training examples than for first set of tasks.
FitNets approach begins by training a network with low depth and high width, which becomes a teacher for a second network (the student), which is much deeper and thinner. Student network is trained to not only predict the output on original task, but also on value of middle layer of teacher network.

Remark 2.7.2. *Neural Network Design Choice: Linear Transformation*
Model units have generally moved towards using more linear functions, which allow model to have nice properties that make optimisation easier. The model's local gradient information corresponds reasonably well to moving towards a distant solution.

Remark 2.7.3. *Neural Network Design Choice: Linear Paths and Skip Connections*
Linear paths and skip connections between layers reduce length of shortest path from lower layer's parameters to the output, mitigating vanishing gradient problem.
Adding extra copies of output that are attached to intermediate hidden layers is an alternative to pre-training, where these 'auxiliary heads' are trained to perform the same task as primary output at the top of the network to ensure that the lower layers receive a large gradient. After training, these may be discarded.

2.7.1 Convolutional Neural Networks

Definition 2.7.4. *Convolutional Neural Networks (CNNs)*
Neural network for processing data with a known, grid-like topology. Uses convolution instead of general matrix

multiplication in at least one of their layers.

Definition 2.7.5. *Convolution*
The convolution of two functions f and g, denoted by $*$, is the integral of product of two functions after one is reflected about the y-axis and shifted.

$$(f * g)(t) = \int_{-\infty}^{\infty} f(\tau)g(t-\tau)d\tau$$

The first argument $f(\tau)$ is the input, and second argument $g(t-\tau)$ is the kernel. The output is feature map.

Definition 2.7.6. *Discrete Convolution*
For complex-valued functions f and g defined on set of $\mathbb{Z}$ integers, the discrete convolution of f and g is

$$(f * g)[n] = \sum_{m=-\infty}^{\infty} f[m]g[n-m]$$

If function g_N is periodic with period N, then the convolution is also periodic and identical to

$$(f * g_N)[n] = \sum_{m=0}^{N-1} \left(\sum_{k=-\infty}^{\infty} f[m+kN] \right) g_N[n-m]$$

Note that convolution is commutative, as the kernel is flipped relative to the input (as m increases, the index into input increases, but index into kernel decreases), but this is not an important property of neural network implementation. Neural network libraries implement cross-correlation instead.

Definition 2.7.7. *Cross-Correlation*
A measure of similarity of two series as a function of the displacement of one relative to the other. Also known as *sliding inner or dot product*. This is defined as

$$(f \star g)(\tau) = \int_{-\infty}^{\infty} \overline{f(\tau)}g(t+\tau)dt$$

This is related to convolution by

$$[f(t) \star gt](t) = [\overline{f(-t)} * g(t)](t)$$

Remark 2.7.8. *Motivation: Sparse Interactions*
The kernel is smaller than input, but is able to detect meaningful features. Hence fewer parameters are to be stored, reducing memory requirements of the model. Units in deeper layers may indirectly interact with a larger portion of the input, hence efficiently describe complicated interactions between many variables.

Remark 2.7.9. *Motivation: Parameter Sharing*
Each member of kernel is used at every position of the input except at boundary, which depends on design decisions regarding the boundary. Only one set of parameter is required to be learned. The runtime of forward propagation is at $O(k \times n)$ for $k \lll m$ connections for each input, given m input and n outputs.

Remark 2.7.10. *Motivation: Equivariant Representations*
An *equivariant* function is such that if the input changes, the output changes in the same way.
Applying a transformation, then convolution, is equivalent to applying convolution, then the transformation.

Remark 2.7.11. *Typical CNN Layers*
The layers of a typical CNN are:

i. Convolution: layer performs several convolutions in parallel to produce set of linear activations
ii. Detector: each linear activation is run through a non-linear activation function
iii. Pooling: modify the output of the layer further with summary statistics of nearby outputs

Example 2.7.12. *Example Types of Pooling*

i. Max Pooling: reports maximum output for each region of the feature map
ii. Average Pooling: reports average for each region of the feature map
iii. L^2 Pooling: calculates Euclidean norm (L^2 norm) for each region of the feature map

Remark 2.7.13. *Pooling on Invariance*
Pooling makes the representation approximately invariant to small translations of input.
Use of pooling is analogous to adding an infinitely strong prior that the function the layer learns must be invariant to small translations. As pooling summarises responses over whole neighbourhood, fewer pooling units than detector units may be used. To handle inputs of varying size, the size of offset between pooling regions may be changed so that the classification layer always receives the same number of summary statistics.

Remark 2.7.14. *Multi-Channel Operations*
Not commutative unless each operation has the same number of output and input channels.
Assume there exists 4-D kernel tensor $\mathcal{K}$ with element $K_{i,j,k,l}$ giving connection between strength between a unit in channel i of the output and a unit in channel j of the input, with offset of k rows and l columns between the output unit and input unit. Assume input consists of observed data $\mathcal{V}$ with element $V_{i,j,k}$ giving the value of input unit within channel i at row j and column k. Assume output consists of $\mathcal{Z}$ in same format as $\mathcal{V}$. If $\mathbb{Z}$ is produced by convolving $\mathcal{K}$ across $\mathcal{V}$ without flipping $\mathcal{K}$, then

$$Z_{i,j,k} = \sum_{l,m,n} V_{l,j+m-1,k+n-1} K_{i,l,m,n}$$

where summation over l, m, n is over all values for which tensor indexing operations inside summation is valid.

Remark 2.7.15. *Downsampled Convolution*
To reduce computational cost by sampling only every s pixels (*stride*) in each direction of output, then

$$Z_{i,j,k} = c(\mathcal{K}, \mathcal{V}, s)_{i,j,k} = \sum_{l,m,n} [\mathcal{V}_{l,(j-1)\times s+m,(k-1)\times s+n} \mathcal{K}_{i,l,m,n}]$$

Remark 2.7.16. *Zero Padding*
Zero-padding is adding zeroes to end of input sequence $\mathcal{V}$ so that the total number of samples is equal to the next higher multiple of two.

i. Valid Convolution: no zero-padding used. If input image has width m and kernel has width k, the output will be width $m - k + 1$. Rate of shrinkage can be dramatic if kernel used are large.
ii. Same Convolution: just enough zero-padding is added to keep size of output equal to size of input. Border pixels may be underrepresented in the model.
iii. Full Convolution: zeroes are added for every pixel to be visited k times in each direction, resulting in an input image of width $m + k - 1$.

Remark 2.7.17. *Locally Connected Layers (Unshared Convolution)*
Similar operation to discrete convolution with small kernel, but without parameter sharing. Adjacency matrix of MLP is the same, but every connection has its own weight as specified by 6-D tensor $\mathcal{W}$, where the indices of $\mathcal{W}$ are the output channel i, the output row j, the output column k, the input channel l, the row offset within the the input m, and the column offset within the input n. The liner part of the layer is then

$$Z_{i,j,k} = \sum_{l,m,n} [\mathcal{V}_{l,j+m-1,k+n-1} w_{i,j,k,l,m,n}]$$

Locally connected layers are useful when it is known that each feature should be function of a small part of space but would not occur across all of space.

Remark 2.7.18. *Tiled Convolution*
Compromise between convolutional layer and locally connected layer. A set of kernels will be rotated through space. Immediate neighbouring locations will have different filters, but memory requirements for storing the parameters will increase only by a factor of the size of this set of kernels.
Let k be a 6-D tensor, where two of the dimensions correspond to different locations in output matrix. Output locations cycle through a set of t different choices of kernel stack in each direction.

$$Z_{i,j,k} = \sum_{l,m,n} \mathcal{V}_{l,j+m-1,k+n-1} \mathcal{K}_{i,l,m,n,j \bmod t+1,k \bmod t+1}$$

3 Time Series and Sequential Systems

These notes develop the standard time-series framework.

3.1 Time Series Foundations

3.1.1 Univariate Time Series Models

Trading activities through an exchange can be described by a sequence of time stamps ('ticks') $t_0 < t_1 < \cdots < t_n$, and 'marks' y_i at time t_i, where t_i denote market open, t_n denote market close. The marks y_i is a characteristic of the order book at time of ith activity. Events with marks associated with the ticks can be described mathematically as a marked point process.

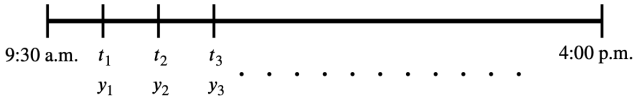


Figure 236: Ticks and Marks

A data aggregation method is where aggregation is conducted when there is a change in marker. An alternative aggregation method would be to divide the time span T for exchange hours into K intervals, so regularly spaced intervals are of size $\Delta t = T/K$. The time series method to be discussed will apply to all aggregated data.

Let $p_{it} = \ln P_{it}$ denote price of ith asset at time t; $p_t = (p_{1t}, p_{2t}, \ldots, p_{nt})$ denote price vector for n assets; y_{it} denote vector of characteristics of ith asset at time t. These quantities are aggregated from high frequency data. Consider r factors $f_t = (f_{1t}, f_{2t}, \ldots, f_{rt})$ that may include market and industry factors, and asset characteristics. Trading rules can be broadly grouped as follows:

i. *Statistical Arbitrage*: $E(p_{i,t+1} \mid p_{i,t}, p_{i,t-1}, \ldots, y_{i,t}, y_{i,t-1}, \ldots)$, that predicts the price of ith stock at $t+1$ based on past trading information (also known as time series momentum)
ii. *Momentum*: $E(p_{t+1} \mid p_t, p_{t-1}, \ldots, y_t, y_{t-1}, \ldots)$, that predicts the cross-sectional momentum of a subset of stocks based on past trading characteristics. For portfolio formation and rebalancing, pairs trading
iii. *Fair Value*: $E(p_{t+1} \mid p_t, p_{t-1}, \ldots, y_t, y_{t-1}, \ldots, f_t, f_{t-1}, \ldots)$, predicts price using all relevant quantities. Factors normally include market, Fama-French; at a more macro level than timescale considered for price prediction, but may still be useful.

Hence price and volatility prediction can be formulated as a time series prediction problem. Autocorrelations and partial autocorrelations can be used to build autoregressive and ARCH models with some predictive power.

Let $Y_1, Y_2, \ldots, Y_T$ be a sequence of random variables with a joint probability distribution. A sequence of observations of stochastic process $\{Y_t, t = 1, \ldots, T\}$ is a realisation of the process. A time series $\{Y_t\}$ is *stationary* if for every integer m, the set of variables $Y_{t_1}, Y_{t_2}, \ldots, Y_{t_m}$ depends only on the distance between times $t_1, t_2, \ldots, t_m$. Thus $E(Y_t) = \mu$, $Var(Y_t) = \sigma^2$ are constant for all t. The *auto-covariance* function is defined as

$$\gamma(s) = Cov(Y_t, Y_{t-s}) = E[(Y_t - \mu)(Y_{t-s} - \mu)] \quad \forall s = 0, \pm 1, \ldots$$

The *auto-correlation* of process at lag s is defined as

$$\rho(s) = Corr(Y_t, Y_{t-s}) = \frac{Cov(Y_t, Y_{t-s})}{(Var(Y_t)\ Var(Y_{t-s}))^{1/2}} = \frac{\gamma(s)}{\gamma(0)}, \quad \forall s = 0, \pm 1, \ldots$$

Some examples of stationary stochastic processes are as follows:

Example 3.1.1. *(White Noise)* Sequence of discrete independent random variables (r.v.) $\{\epsilon_t\}$ with $E[\epsilon_t] = 0$ and $E[\epsilon_t^2] = \sigma^2$. Set $Y_t = \mu + \epsilon_t$, then $E[Y_t] = \mu$, $Cov(Y_t, Y_{t-s}) = Var(Y_t) = \sigma^2$ if $s = 0$, and $Cov(Y_t, Y_{t-s}) = 0$ if $s \neq 0$. Hence the process is stationary.

Some examples of non-stationary stochastic processes are as follows:

Example 3.1.2. *(Random Walk with Drift)* Let $\{\epsilon_t\}$ be sequence of independent r.v. with $E[\epsilon_t] = 0$, $E[\epsilon_t^2] = \sigma^2$, and define process $\{Y_t\}$ by $Y_t = Y_{t-1} + \delta + \epsilon_t$ with $Y_0 = 0$. Then the process can be summarised as

$$Y_t = \delta t + \sum_{j=1}^{t} \epsilon_j$$

Note that $E[Y_t] = \delta t$ and $Var[Y_t] = t\sigma^2$, hence process $\{Y_t\}$ is not stationary.

Given T observations from stationary process $\{Y_t\}$, the *sample mean* is $\overline{Y} = \frac{1}{T}\sum_{t=1}^{T} Y_t$. The *sample auto-covariance* function is defined by $\hat{\gamma}(s) = \frac{1}{T}\sum_{t=1}^{T-s}(Y_t - \overline{Y})(Y_{T-s} - \overline{Y})$ for $s = 0, 1, \ldots$. The *sample auto-correlation* function (ACF) is defined as $\hat{\rho}(s) = \frac{\hat{\gamma}(s)}{\hat{\gamma}(0)} = r(s)$. The *sample variance* is $Var(\overline{Y}) = \frac{\gamma(0)}{T}\left[1 + 2\sum_{u=1}^{Y-1}\left(\frac{T-u}{T}\right)\rho(u)\right]$; note that it has to account for auto-correlations.

Definition 3.1.3. A stochastic process $\{Y_t\}$ is a *linear process* if it can be represented as

$$Y_t = \mu + \sum_{j=0}^{\infty} \Psi_j \epsilon_{t-j}$$

where ϵ_t are independent with mean 0, variance σ_ϵ^2, and $\sum_{j=0}^{\infty} < \infty$.

3.2 ARMA and ARIMA Structure

3.2.1 Stationary Time Series

Definition 3.2.1. A *linear first-order difference equation* is defined as

$$y_t = \phi y_{t-1} + w_t$$

where y_t is the target variable (with y_{t-1} the lag 1 variable), w_t is an input variable at time t

The difference equation may be solved by recursive substitution to arrive at

$$y_t = \phi^{t+1} y_{-1} + \phi^t w_0 + \phi^{t-1} w_1 + \phi^{t-2} w_2 + \cdots + \phi w_{t-1} + w_t$$

The *dynamic multiplier* calculates effect of w_t on y_{t+j}, and is given by

$$\frac{\partial y_{t+j}}{\partial w_t} = \phi^j$$

If $|\phi| < 1$, the system is stable. If $|\phi| > 1$, then the system is explosive.

We may generalise the process to *p-th order difference equation*, i.e.,

$$y_t = \phi_t y_{t-1} + \phi_2 t_{t-2} + \cdots + \phi_p y_{t-p} + w_t$$

We may rewrite this as a first-order difference equation in a $(p \times 1)$ vector $\boldsymbol{\xi}_t$:

$$\boldsymbol{\xi}_t = \begin{bmatrix} y_t \\ y_{t-1} \\ \vdots \\ y_{t-p+1} \end{bmatrix}$$

Define the $(p \times p)$ matrix $\boldsymbol{F}$ by

$$\boldsymbol{F} = \begin{bmatrix} \phi_1 & \phi_2 & \phi_3 & \cdots & \phi_{p-1} & \phi_p \\ 1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 1 & 0 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & 1 & 0 \end{bmatrix}$$

Finally, define the $(p \times 1)$ vector $\boldsymbol{v}_t$ by

$$\boldsymbol{v}_t = \begin{bmatrix} w_t \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

Then the system of p equations $\boldsymbol{\xi}_t = \boldsymbol{F}\boldsymbol{\xi}_{t-1} + \boldsymbol{v}_t$ is identical to the first order difference equation, and by recursive substitution, we have the following for the case of p-th order difference equation:

$$\boldsymbol{\xi}_{t+j} = \boldsymbol{F}^{j+1}\boldsymbol{\xi}_{t-1} + \boldsymbol{F}^j \boldsymbol{v}_t + \boldsymbol{F}^{j-1}\boldsymbol{v}_{t+1} + \cdots + \boldsymbol{F}\boldsymbol{v}_{t+j-1} + \boldsymbol{v}_{t+j}$$

Hence the dynamic multiplier is then $\frac{\partial y_{t+j}}{\partial w_t} f_{11}^{(j)}$, where $f_{11}^{(j)}$ is the $(1,1)$ element of $\boldsymbol{F}^j$.

Example 3.2.2. *(Moving Average)* Let $\{\epsilon_t\}$ be independent r.v., with process $\{Y_t\}$ where $Y_t = \mu + \epsilon_t + \epsilon_{t_1}$ for $t = 0, 1, 2, \ldots$, and μ is constant. Then $E[Y_t] = \mu \ \forall t$, and

$$Cov(Y_t, Y_{t-s}) = \gamma(s) = \begin{cases} 2\sigma^2 & \text{if } s = 0 \\ \sigma^2 & \text{if } s = 1 \\ 0 & \text{if } s > 1 \end{cases}$$

Hence $\{Y_t\}$ is stationary, with $\rho(s) = \gamma(s)/\gamma(0)$ such that $\rho(0) = 1$, $\rho(1) = 1/2$, $\rho(s) = 0$ for $|s| > 1$.

3.3 Spectral Methods

3.4 State-Space Models

3.5 Nonlinear and Volatility Models

3.6 Long Memory and Multiscale Structure

3.7 Sequential Inference and Filtering

References

Abbott, S. (2015). *Understanding Analysis* (2 ed.). Undergraduate Texts in Mathematics. New York, NY: Springer.

Axler, S. (2024). *Linear Algebra Done Right* (4 ed.). Undergraduate Texts in Mathematics. Cham: Springer.

Bartle, R. G. (1995). *The Elements of Integration and Lebesgue Measure*. Wiley Classics Library. New York, NY: Wiley-Interscience.

Garcia, S. R. and R. A. Horn (2017). *A Second Course in Linear Algebra*. Cambridge: Cambridge University Press.

Granziol, D., X. Wan, and S. Roberts (2020). Gadam: Combining adaptivity with iterate averaging gives greater generalisation. *arXiv preprint arXiv:2003.01247*.

Grinberg, R. (2017). *The Real Analysis Lifesaver: All the Tools You Need to Understand Proofs*. Princeton, NJ: Princeton University Press.

Hamilton, J. D. (1994). *Time Series Analysis*. Princeton University Press.

Mazur, M. (2015). A step by step backpropagation example.

Pugh, C. C. (2015). *Real Mathematical Analysis* (2 ed.). Undergraduate Texts in Mathematics. Cham: Springer.

Roman, S. (2008). *Advanced Linear Algebra* (3 ed.), Volume 135 of *Graduate Texts in Mathematics*. New York, NY: Springer.

Ross, K. A. (2013). *Elementary Analysis: The Theory of Calculus* (2 ed.). Undergraduate Texts in Mathematics. New York, NY: Springer.

Rudin, W. (1976). *Principles of Mathematical Analysis* (3 ed.). New York: McGraw-Hill.

Strang, G. (2006). *Linear Algebra and Its Applications* (4 ed.). Belmont, CA: Thomson Brooks/Cole.

Tao, T. (2016a). *Analysis I* (3 ed.), Volume 37 of *Texts and Readings in Mathematics*. Singapore: Springer.

Tao, T. (2016b). *Analysis II* (3 ed.), Volume 38 of *Texts and Readings in Mathematics*. Singapore: Springer.